

IBM DB2 10.5
for Linux, UNIX, and Windows

Message Reference Volume 1



IBM DB2 10.5
for Linux, UNIX, and Windows

Message Reference Volume 1



Note

Before using this information and the product it supports, read the general information under Appendix B, "Notices," on page 957.

Edition Notice

This document contains proprietary information of IBM. It is provided under a license agreement and is protected by copyright law. The information contained in this publication does not include any product warranties, and any statements provided in this manual should not be interpreted as such.

You can order IBM publications online or through your local IBM representative.

- To order publications online, go to the IBM Publications Center at <http://www.ibm.com/shop/publications/order>
- To find your local IBM representative, go to the IBM Directory of Worldwide Contacts at <http://www.ibm.com/planetwide/>

To order DB2 publications from DB2 Marketing and Sales in the United States or Canada, call 1-800-IBM-4YOU (426-4968).

When you send information to IBM, you grant IBM a nonexclusive right to use or distribute the information in any way it believes appropriate without incurring any obligation to you.

© Copyright IBM Corporation 1993, 2013.

US Government Users Restricted Rights – Use, duplication or disclosure restricted by GSA ADP Schedule Contract with IBM Corp.

Contents

| | |
|---|--|
| Part 1. Introduction to messages . . . 1 | Chapter 24. ADM12000 - ADM12499 . . . 85 |
| Part 2. ADM Messages 5 | Chapter 25. ADM12500 - ADM12999 . . . 89 |
| Chapter 1. ADM0000 - ADM0499 7 | Chapter 26. ADM13000 - ADM13499 . . . 95 |
| Chapter 2. ADM0500 - ADM0999 9 | Chapter 27. ADM13500 - ADM13999 . . . 97 |
| Chapter 3. ADM1000 - ADM1499 11 | Chapter 28. ADM14000 - ADM14499 . . . 99 |
| Chapter 4. ADM1500 - ADM1999 13 | Chapter 29. ADM14500 - ADM14999 . . . 103 |
| Chapter 5. ADM2000 - ADM2499 21 | Chapter 30. ADM15000 - ADM15499 . . . 105 |
| Chapter 6. ADM2500 - ADM2999 23 | Chapter 31. ADM15500 - ADM15999 . . . 107 |
| Chapter 7. ADM3000 - ADM3499 25 | Chapter 32. ADM16000 - ADM16499 . . . 109 |
| Chapter 8. ADM3500 - ADM3999 27 | Chapter 33. ADM16500 - ADM16999 . . . 111 |
| Chapter 9. ADM4000 - ADM4499 29 | Part 3. AMI Messages 113 |
| Chapter 10. ADM4500 - ADM4999 33 | Chapter 34. AMI0000 - AMI0499 115 |
| Chapter 11. ADM5500 - ADM5999 35 | Part 4. Replication Messages 117 |
| Chapter 12. ADM6000 - ADM6499 41 | Chapter 35. ASN0000 - ASN0499 119 |
| Chapter 13. ADM6500 - ADM6999 55 | Chapter 36. ASN0500 - ASN0999 133 |
| Chapter 14. ADM7000 - ADM7499 57 | Chapter 37. ASN1000 - ASN1499 157 |
| Chapter 15. ADM7500 - ADM7999 59 | Chapter 38. ASN1500 - ASN1999 171 |
| Chapter 16. ADM8000 - ADM8499 65 | Chapter 39. ASN2000 - ASN2499 211 |
| Chapter 17. ADM8500 - ADM8999 67 | Chapter 40. ASN2500 - ASN2999 263 |
| Chapter 18. ADM9000 - ADM9499 69 | Chapter 41. ASN3500 - ASN3999 269 |
| Chapter 19. ADM9500 - ADM9999 71 | Chapter 42. ASN4000 - ASN4499 273 |
| Chapter 20. ADM10000 - ADM10499 . . . 75 | Chapter 43. ASN5000 - ASN5499 283 |
| Chapter 21. ADM10500 - ADM10999 . . . 77 | Chapter 44. ASN6000 - ASN6499 297 |
| Chapter 22. ADM11000 - ADM11499 . . . 81 | Chapter 45. ASN7000 - ASN7499 299 |
| Chapter 23. ADM11500 - ADM11999 . . . 83 | Chapter 46. ASN7500 - ASN7999 331 |

| | |
|---|------------|
| Part 5. AUD Messages | 375 |
| Chapter 47. AUD0000 - AUD0499 . . . | 377 |
| Part 6. CCA Messages | 379 |
| Chapter 48. CCA1000 - CCA1499 . . . | 381 |
| Chapter 49. CCA2000 - CCA2499 . . . | 383 |
| Chapter 50. CCA3000 - CCA3499 . . . | 387 |
| Chapter 51. CCA5000 - CCA5499 . . . | 393 |
| Part 7. CIE Messages | 395 |
| Chapter 52. CIE0000 - CIE0499 . . . | 397 |
| Chapter 53. CIE0500 - CIE0999 . . . | 417 |
| Part 8. CLI Messages | 427 |
| Chapter 54. CLI0000 - CLI0499 . . . | 429 |
| Chapter 55. CLI0500 - CLI0999 . . . | 447 |
| Part 9. Net Search Extender Messages | 453 |
| Chapter 56. CTE0000 - CTE0499 . . . | 455 |
| Chapter 57. CTE0500 - CTE0999 . . . | 475 |
| Chapter 58. CTE5000 - CTE5499 . . . | 477 |
| Part 10. DB2 Messages | 483 |
| Chapter 59. DB20000 - DB20499 . . . | 485 |
| Chapter 60. DB21000 - DB21499 . . . | 487 |
| Chapter 61. DB21500 - DB21999 . . . | 497 |
| Chapter 62. DB22000 - DB22499 . . . | 503 |
| Chapter 63. DB29000 - DB29499 . . . | 509 |
| Chapter 64. DB29500 - DB29999 . . . | 511 |
| Chapter 65. DB210000 - DB210499 . . . | 513 |
| Chapter 66. DB216000 - DB216499 . . . | 517 |

| | |
|--|------------|
| Chapter 67. DB250000 - DB250499 . . . | 519 |
| Chapter 68. DB250500 - DB250999 . . . | 527 |
| Chapter 69. DB255000 - DB255499 . . . | 529 |
| Part 11. DBA Messages | 537 |
| Chapter 70. DBA0000 - DBA0499 . . . | 539 |
| Chapter 71. DBA0500 - DBA0999 . . . | 547 |
| Chapter 72. DBA1000 - DBA1499 . . . | 549 |
| Chapter 73. DBA1500 - DBA1999 . . . | 561 |
| Chapter 74. DBA2000 - DBA2499 . . . | 565 |
| Chapter 75. DBA3000 - DBA3499 . . . | 575 |
| Chapter 76. DBA4000 - DBA4499 . . . | 579 |
| Chapter 77. DBA4500 - DBA4999 . . . | 583 |
| Chapter 78. DBA5000 - DBA5499 . . . | 585 |
| Chapter 79. DBA5500 - DBA5999 . . . | 587 |
| Chapter 80. DBA6000 - DBA6499 . . . | 591 |
| Chapter 81. DBA7000 - DBA7499 . . . | 595 |
| Chapter 82. DBA7500 - DBA7999 . . . | 601 |
| Chapter 83. DBA8000 - DBA8499 . . . | 607 |
| Part 12. DBI Messages | 609 |
| Chapter 84. DBI1000 - DBI1499 . . . | 611 |
| Chapter 85. DBI1500 - DBI1999 . . . | 669 |
| Chapter 86. DBI20000 - DBI20499 . . . | 699 |
| Part 13. DBT Messages | 719 |
| Chapter 87. DBT1000 - DBT1499 . . . | 721 |
| Chapter 88. DBT2000 - DBT2499 . . . | 729 |
| Chapter 89. DBT3000 - DBT3499 . . . | 735 |
| Chapter 90. DBT3500 - DBT3999 . . . | 737 |

| | |
|--|------------|
| Chapter 91. DBT4000 - DBT4499 . . . | 753 |
| Chapter 92. DBT5000 - DBT5499 . . . | 759 |
| Chapter 93. DBT5500 - DBT5999 . . . | 765 |
| Chapter 94. DBT6000 - DBT6499 . . . | 773 |
| Chapter 95. DBT7000 - DBT7499 . . . | 777 |
| Part 14. DQP Messages | 793 |
| Chapter 96. DQP0000 - DQP0499 . . . | 795 |
| Chapter 97. DQP1000 - DQP1499 . . . | 799 |
| Chapter 98. DQP2000 - DQP2499 . . . | 801 |
| Chapter 99. DQP2500 - DQP2999 . . . | 811 |
| Chapter 100. DQP3000 - DQP3499 . . . | 815 |
| Part 15. EXP Messages | 817 |
| Chapter 101. EXP0000 - EXP0499 . . . | 819 |
| Part 16. GSE Messages | 833 |
| Part 17. ICM Messages | 879 |
| Chapter 102. ICM0000 - ICM0499 . . . | 881 |
| Chapter 103. ICM0500 - ICM0999 . . . | 891 |
| Chapter 104. ICM10000 - ICM10499 . . . | 899 |
| Chapter 105. ICM10500 - ICM10999 . . . | 903 |
| Part 18. LIC Messages | 905 |

| | |
|--|------------|
| Chapter 106. LIC1000 - LIC1499 . . . | 907 |
| Part 19. MQL Messages | 913 |
| Chapter 107. MQL0000 - MQL0499 . . . | 915 |
| Part 20. SAT Messages | 919 |
| Chapter 108. SAT1000 - SAT1499 . . . | 921 |
| Chapter 109. SAT2000 - SAT2499 . . . | 923 |
| Chapter 110. SAT3000 - SAT3499 . . . | 927 |
| Chapter 111. SAT4000 - SAT4499 . . . | 933 |
| Part 21. SPM Messages | 935 |
| Chapter 112. SPM0000 - SPM0499 . . . | 937 |
| Part 22. SPM protocol violation records | 947 |
| Part 23. Appendixes | 949 |
| Appendix A. Overview of the DB2 technical information | 951 |
| DB2 technical library in hardcopy or PDF format | 951 |
| Displaying SQL state help from the command line processor | 954 |
| Accessing different versions of the DB2 Information Center | 954 |
| Terms and conditions. | 954 |
| Appendix B. Notices | 957 |
| Index | 961 |

Part 1. Introduction to messages

It is assumed that you are familiar with the functions of the operating system where DB2® is installed. You can use the information contained in the following chapters to identify an error or problem and resolve the problem by using the appropriate recovery action. This information can also be used to understand where messages are generated and logged.

Message Structure

Message help describes the cause of a message and describes any action you should take in response to the message.

Message identifiers consist of a three character message prefix, followed by a four or five digit message number, followed by a single letter suffix. For example, *SQL1042C*. For a list of message prefixes, see “Invoking message help” on page 2 and “Other DB2 Messages” on page 2. The single letter suffix describes the severity of the error message.

In general, message identifiers ending with a *C* are for severe messages; those ending with an *E* indicate urgent messages; those ending with an *N* indicate error messages; those ending with a *W* indicate warning messages; and those ending with an *I* indicate informational message.

For ADM messages, message identifiers ending with a *C* indicate severe messages; those ending with an *E* indicate urgent messages; those ending with a *W* indicate important messages; and those ending with an *I* indicate informational messages.

For SQL messages, message identifiers ending with a *C* indicate critical system errors; those ending with an *N* indicate error messages; those ending with a *W* indicate warning or informational messages.

Some messages include tokens, sometimes also called message variables. When a message containing tokens is generated by DB2, each token is replaced by a value specific to the error condition that was encountered, to help the user diagnose the cause of the error message. For example, the DB2 message *SQL0107N* is as follows:

- from the command line processor:
SQL0107N The name "<name>" is too long. The maximum length is "<length>".
- from the DB2 information center:
SQL0107N The name *name* is too long. The maximum length is *length*.

This message includes the two tokens "<name>" and "<length>". When this message is generated at runtime, the message tokens would be replaced by the actual name of the object that caused the error, and the maximum length allowed for that type of object, respectively.

In some cases a token is not applicable for a specific instance of an error, and the value *N is returned instead, for example:

SQL20416N The value provided ("*N") could not be converted to a security label. Labels for the security policy with a policy ID of "1" should be "8" characters long. The value is "0" characters long. SQLSTATE=23523

Invoking message help

To invoke message help, open the command line processor and enter:

```
? XXXnnnnnn
```

where *XXX* represents a valid message prefix and *nnnnn* represents a valid message number.

The message text associated with a given SQLSTATE value can be retrieved by issuing:

```
? nnnnn
```

or

```
? nn
```

where *nnnnn* is a five digit SQLSTATE (alphanumeric) and *nn* is the two digit SQLSTATE class code (first two digits of the SQLSTATE value).

Note: The message identifier accepted as a parameter of the **db2** command is not case sensitive. Also, the single letter suffix is optional and is ignored.

Therefore, the following commands will produce the same result:

- ? SQL0000N
- ? sql0000
- ? SQL0000w

To invoke message help on the command line of a UNIX based system, enter:

```
db2 "? XXXnnnnnn"
```

where *XXX* represents a valid message prefix
and *nnnnn* represents a valid message number.

If the message text is too long for your screen, use the following command (on Unix-based systems and others which support 'more'):

```
db2 "? XXXnnnnnn" | more
```

Other DB2 Messages

Some DB2 components return messages that are not available online or are not described in this manual. Some of the message prefixes might include:

AUD Messages generated by the DB2 Audit facility.

DIA Diagnostics messages generated by many DB2 components. These messages are written in the DB2 diagnostic (db2diag) log files, and are intended to provide additional information for users and DB2 service personnel when investigating errors.

In most cases, these messages provide sufficient information to determine the cause of the warning or error. For more information on the command or utility that generated the messages, please refer to the appropriate manual where the command or utility is documented.

Other Message Sources

When running other programs on the system, you may receive messages with prefixes other than those mentioned in this reference.

For information on these messages, refer to the information available for that program product.

Part 2. ADM Messages

This section contains the Administration Notification (ADM) messages. The messages are listed in numeric sequence.

Chapter 1. ADM0000 - ADM0499

ADM0001C A severe error has occurred. Examine the administration notification log and contact IBM Support if necessary.

Chapter 2. ADM0500 - ADM0999

ADM0500E The DB2 Service does not have the necessary authority to complete the command. If a user account has been associated with the DB2 Service, ensure that the user account has the correct access rights on the local system and access to the Domain Security Database for user authentication and group enumeration.

ADM0501C A stack overflow exception has occurred. The DB2 instance has terminated abnormally. To remedy this problem, you may increase the default stack size for db2syscs.exe using the db2hdr utility as follows: db2hdr \sqllib\bin\ db2syscs.exe /s:<stackreserve>[,<stackcommit>] Note that increasing the default stack size will consume virtual memory space and may reduce the maximum number of concurrent connections. Contact IBM Support for further assistance.

ADM0502C The DB2 instance has terminated abnormally. To remedy this problem, increase the AGENT_STACK_SZ DBM configuration parameter. Contact IBM Support for further assistance.

ADM0503C An unexpected internal processing error has occurred. All DB2 processes associated with this instance have been shutdown. Diagnostic information has been recorded. Contact IBM Support for further assistance.

ADM0504C An unexpected internal processing error has occurred. All DB2 processes associated with this instance have been suspended. Diagnostic information has been recorded. Contact IBM Support for further assistance.

ADM0505E DB2 received a SIGDANGER signal from the operating system. This signal indicates that the system is running low on paging space. If the paging space gets too low, the operating system will forcibly terminate user processes. Contact your system administrator to increase your paging space.

ADM0506I DB2 has automatically updated the *parameter* kernel parameter from *originalValue* to the recommended value *recommendedValue*.

ADM0507W DB2 could not automatically update the *parameter* kernel parameter to the recommended value *recommendedValue*. Update this kernel parameter manually.

ADM0508E DB2 was unable to load the Java interpreter library *libraryPath*. This error commonly occurs because of restrictions imposed by the operating system. Consult the IBM DB2 documentation for a solution. If this problem persists, contact IBM support.

ADM0509E A non root capable instance of DB2 has been detected. Limited functionality will be available.

ADM0510E AGENTPRI configuration variable or Agent Priority Resource Configuration changes are not possible without root capabilities.

ADM0511E Raw device access using the character device controller has been deprecated by the operating system and might be removed in the future. The DB2 database manager processes can access the same device by using the block device interface; however db2relocatedb will need to be run on each individual database to change the name of the affected device in any table space containers and raw log containers.

Explanation: When working with containers to store data, the database manager supports direct disk (raw) access using the block device interface (that is, raw I/O).

User response: Change the name of the affected device in any table space containers and raw log containers by using the db2relocate command.

ADM0512W The database manager instance does not have permission to pin a database memory segment. Contact your system administrator to grant permission for the database manager instance owner to pin memory. The database manager will continue to function with unpinned database memory.

ADM0513W db2start succeeded. However, no I/O completion port (IOCP) is available.

Explanation: This error is returned when there are no I/O completion ports with status "Available" on the computer where DB2 database is installed. An I/O completion port is an operating system channel for I/O requests. Using I/O completion ports can improve the performance of very large databases.

User response: You do not have to respond to this error, unless you want to use I/O completion ports to improve performance. To, configure I/O completion ports, follow the instructions in the topic called "Configuring IOCP on AIX" in the DB2 Information Center.

ADM0514W The system Network Time Protocol (NTP) process was not detected to be up and running or properly synchronized.

Explanation: To ensure that clock-sensitive operations like error-logging, monitoring and point-in-time recovery run optimally across a distributed environment, it is highly recommended that the Network Time Protocol (NTP) service be running and synchronized to the same system peer across all members.

User response: Ensure that the NTP service is up and running and that all members are synchronized to the same host.

ADM0515W The system clock on this member is out of synchronization by more than the allowed threshold with member coordinator-member. The differential is timestamp-diff minutes.

Explanation: Timestamps from each member are periodically compared to ensure that the time discrepancy between members is less than the predefined allowable threshold set by the MAX_TIME_DIFF database manager configuration parameter. The differential between the local system clock on this host and that of the indicated remote member has been detected to exceed this limit.

User response: Verify that the differentials between the system clocks on all members are within the limit specified by MAX_TIME_DIFF.

ADM0516W CPU binding information: cluster caching facility (CF) processes are bound to *number* cores.

Explanation: When a cluster caching facility and a DB2 member coexist on a single host, the CF processes will be assigned to a subset of the total cores available. For more specific information on the logical processors, see the db2diag.log file.

User response: No user response is required.

ADM0517W CPU binding information: DB2 member processes are bound to *number* cores.

Explanation: When a DB2 member and a cluster caching facility (CF) coexist on a single host, the member processes will be assigned to a subset of the total cores available. For more specific information on the logical processors, see the db2diag.log file.

User response: No user response is required.

Chapter 3. ADM1000 - ADM1499

ADM1010I Reconcile started on table *tableName*.

ADM1011I Reconcile has completed successfully on table *tableName*.

ADM1012W Reconcile has failed on table *tableName*.
The table will now be placed into
DATALINK Reconcile Not Possible
(DRNP) state.

ADM1013W Reconcile for table *tableName* is
successful on Data Links Managers
(DLMs) that were up and is pending on
DLMs that were down. The table will
now be placed in DATALINK Reconcile
Pending (DRP) state.

ADM1014E Reconcile has failed on table *tableName*.

ADM1022W The metadata information for the
DATALINK column(s) is missing on
DB2 Data Links Manager *serverName*.

ADM1023W Reconcile processing is PENDING on
Data Links Manager *serverName*.

ADM1024W Reconcile processing is COMPLETE on
Data Links Manager *serverName*.

Chapter 4. ADM1500 - ADM1999

ADM1500W DB2 is unable to locate log file *filename*.

ADM1501W Log file compression failed on log file *log-file* during archiving for *log-archive-method* for database *database* on member *member-number*.

Explanation: A log file could not be compressed before it was archived. The warning message is logged in the db2diag.log file. The log file is still archived without being compressed.

Should the archive of the log file fail, an ADM1848W message is also logged in the db2diag.log file.

User response: None.

ADM1502C A command has set BACKUP_PENDING state off. Data integrity or the referential integrity of the data is at risk.

Explanation: There are multiple events that can turn on BACKUP_PENDING state for a database. In general, when a database is in BACKUP_PENDING state, you should perform a backup operation to turn the BACKUP_PENDING state off.

You can manually change the BACKUP_PENDING database state to OFF by using the db2dart tool. Using the db2dart tool to change the database state should be done only in rare circumstances, because changing the database state in this way puts data integrity at risk.

This message is returned when the BACKUP_PENDING state has been manually turned off and as a result data integrity or the referential integrity of the data cannot be guaranteed if a full, offline database backup is not taken.

User response: To ensure data integrity or the referential integrity of the data, perform a full offline database backup.

ADM1503W Application *appl-name* with application handle *appl-handle* and application id *appl-id* running under authentication ID *auth_id* has been forced off of the database because caching facility (CF) has run out of memory used for locking. The unit of work will be rolled back.

Explanation: The CF memory used for locking for this particular database has been exhausted. This application has been forced off the database in order to make CF memory available for locking purposes. The unit of work will be rolled back.

This condition indicates that the LOCKLIST database configuration parameter is too large relative to the CF_LOCK_SZ database configuration parameter. Reducing the LOCKLIST database configuration parameter or increasing the CF_LOCK_SZ parameter can help avoid this condition.

User response: Take the following action:

- Decrease the LOCKLIST database configuration parameter
- Increase the CF_LOCK_SZ database configuration parameter.

ADM1504W The caching facility (CF) memory used for locking has run out of memory. The amount of CF memory used for locking will be increased and the amount of CF memory used for group buffer pool usage will be decreased. The performance of database *db-name* may be impacted.

Explanation: The caching facility (CF) has run out of memory used for locking. To make more CF memory available for locking purpose, the amount of the CF memory used for group buffer pool purpose will be reduced.

This condition indicates that the LOCKLIST database configuration parameter is too large relative to the CF_LOCK_SZ database configuration parameter. If this condition persists, reduce the LOCKLIST database configuration parameter or increase the CF_LOCK_SZ parameter.

User response: No action required.

ADM1510W A file sharing violation has occurred while accessing log file *fileName*. Another process may be using this file. DB2 will try to delete it later.

ADM1511W The error *error* has occurred while deleting log file *filename*. DB2 will try to delete this file later.

ADM1512E The error *error* has occurred while deleting log file *filename*. User must manually delete this file.

ADM1513W The log file *fileName* has been deleted.

ADM1514W The log file *fileName* no longer exists.

ADM1524I Member crash recovery has been initiated.

Explanation: The database modified by a DB2 member is being returned to a consistent and usable state. Any incomplete transactions that were in memory when the DB2 member crashed will be rolled back, and any committed transactions that were in memory when the DB2 member crashed will be completed. The database is not available on this DB2 member until member crash recovery successfully completes; however, other DB2 members can continue to access the database while the recovery is in progress.

User response: No user action required.

ADM1525I Member crash recovery has completed successfully.

Explanation: The database modified by this DB2 member has been returned to a consistent state. Any incomplete transactions that were in memory when the crash occurred were rolled back and any committed transactions were completed.

User response: No user action required.

ADM1526E Member crash recovery has failed with SQLCODE *SQLCODE*

Explanation: Member crash recovery has failed. The database is not available on the failed DB2 member until member crash recovery successfully completes. Transactions on other DB2 members may be affected if there were any retained locks held on the member at the time of the DB2 member's abnormal termination.

User response: See the DB2 Information Center for information about the given SQLCODE.

ADM1527I Group crash recovery has been initiated.

Explanation: The database is being returned to a consistent and usable state. Any incomplete transactions that were in memory when the crash occurred will be rolled back and any committed transactions will be completed. The database is not available until group crash recovery successfully completes.

User response: No user action required.

ADM1528I Group crash recovery has completed successfully.

Explanation: The database has been returned to a consistent and usable state. Any incomplete transactions that were in memory when the crash occurred were rolled back and any committed transactions were completed.

User response: No user action required.

ADM1529E Group crash recovery has failed with SQLCODE *SQLCODE*

Explanation: Group crash recovery has failed. The database is not available until group crash recovery successfully completes.

User response: See the DB2 Information Center for information about the given SQLCODE.

ADM1530I Crash recovery has been initiated.

Explanation: The database is being returned to a consistent and usable state. Any incomplete transactions that were in memory when the crash occurred will be rolled back, and any committed transactions that were in memory when the crash occurred will be completed. The database is not available until crash recovery successfully completes.

User response: No user action required.

ADM1531I Crash recovery has completed successfully.

Explanation: The database has been returned to a consistent and usable state. Any incomplete transactions that were in memory when the crash occurred were rolled back and any committed transactions were completed.

User response: No user action required.

ADM1532E Crash recovery has failed with SQLCODE *SQLCODE*.

Explanation: Crash recovery has failed. The database is not available until crash recovery successfully completes.

User response: See the DB2 Information Center for information about the given SQLCODE.

ADM1533W Database has recovered. However, one or more table spaces are offline.

ADM1534W Table space *tablespaceName* could not be reduced in size because there are used pages at the end of the table space.

ADM1540W Application *appl_name* with application handle *appl_handle* and application id *appl_id* executing under authentication id *auth_id* has used more log space than allowed by database configuration parameter MAX_LOG (current value *MAX_LOG_value*). Terminate this application by COMMIT, ROLLBACK or FORCE APPLICATION.

ADM1541W Application *appl_name* with application handle *appl_handle* and application id *appl_id* executing under authentication id *auth_id* has been forced off of the database for violating database configuration parameter **NUM_LOG_SPAN** (current value *NUM_LOG_SPAN_value*). The unit of work will be rolled back.

ADM1542W Application *appl_name* with application handle *appl_handle* and application id *appl_id* executing under authentication id *auth_id* will be forced off of the database for violating database configuration parameter **MAX_LOG** (current value *MAX_LOG_value*). The unit of work will be rolled back.

ADM1550W The active log space exceeds the **LOGPRIMARY** DB configuration parameter. **ROLLBACK** may be slow if log files have to be retrieved from archive.

ADM1551W DB2 is retrieving the active log file *fileName* from archive. This may result in a delay.

ADM1552E DB2 is unable to open active log file *filename*. This could be caused by problem in retrieving the log file from archive. DB2 will try again every 5 minutes.

ADM1600W There was not enough memory to enable parallel recovery; however, recovery is proceeding. Increase either the **DBHEAP** or **UTIL_HEAP_SZ** configuration parameters by *blockSize* for any future recoveries.

ADM1601E A recovery or rollforward operation on database *database-name* cannot continue because of a missing log file *log-file-name* on database partition *dbpartitionnum* and log stream *log-stream-ID*.

Explanation: The recovery or rollforward operation cannot find the specified log file in the archive, database log directory, or overflow log directory for the given log stream.

If the operation was a crash recovery, the database is left in an inconsistent state. If the operation was a rollforward operation, the operation has stopped and the database is left in rollforward pending state.

User response: Recover the missing log file by taking

one of the following actions:

- Move the specified log file into the database log directory and restart the operation.
- If an overflow log path can be specified, restart the operation with the overflow log path, specifying the path that contains the log file.

If the missing log file cannot be found, determine whether one of the following special cases applies:

- If the operation is a **ROLLFORWARD DATABASE** command to maintain a standby system through log shipping, this error might be normal, because some files that are available on the primary site might not yet be available on the standby system. To ensure that your standby system is up-to-date, issue a **ROLLFORWARD DATABASE** command with the **QUERY STATUS** option after each rollforward operation to verify that the log replay is progressing properly. If you find that a rollforward operation on the standby system is not making progress over an extended period of time, determine why the log file that is reported as missing is not available on the standby system, and correct the problem. Note that the **ARCHIVE LOG** command can be used to truncate currently active log files on the primary system, making them eligible for archiving and subsequent replay on the standby system.
- If a **ROLLFORWARD DATABASE** command with the **TO END OF LOGS** option was issued following a restore operation from an online backup image in which the only available logs are those that were included in the backup image, there are two possible scenarios to consider:
 - Scenario 1: All of the log files contained in the backup image are found by the rollforward operation. However, the rollforward operation still looks for log files that were updated following the original backup operation. Bring the database to a consistent state by issuing the **ROLLFORWARD DATABASE** command with the **STOP** option (without the **TO END OF LOGS** option). To avoid this scenario in the future, use the **END OF BACKUP** option instead of the **END OF LOGS** option, because the rollforward operation will not look for log files that were updated after the backup was taken.
 - Scenario 2: One or more log files that were contained in the backup image were not found by the rollforward operation. These log files are required to bring the database to a consistent state. Attempting to bring the database to a consistent state by issuing the **ROLLFORWARD DATABASE** command with the **STOP** option (without the **TO END OF LOGS** option) will fail with **SQL1273N**; recover the missing log file, as described earlier in this section.

If the missing log file cannot be recovered:

- If the operation is a ROLLFORWARD DATABASE command, you can issue the ROLLFORWARD DATABASE command again with the STOP option (without the END OF LOGS option or the END OF BACKUP option) to bring the database to a consistent state. If this consistency point (immediately prior to the missing log file) is not acceptable, you can restore the database and roll forward to any point in time that is prior to the missing log file by providing an earlier time stamp to the ROLLFORWARD DATABASE command.
- If the operation is a ROLLFORWARD DATABASE command with the STOP or COMPLETE option (without the END OF LOGS option or the END OF BACKUP option), the missing log file is needed to bring the database to a consistent state. Because you cannot recover the missing log file, you must restore and roll forward to an earlier point in time (as long as that point in time is not prior to the minimum recovery time).

ADM1602W Rollforward recovery has been initiated.

ADM1603I DB2 is invoking the forward phase of the database rollforward recovery.

ADM1604I DB2 is invoking the forward phase of the table space rollforward recovery.

ADM1605I DB2 is invoking the backward phase of database rollforward recovery.

ADM1606I DB2 is invoking the backward phase of table space rollforward recovery.

ADM1607I DB2 is invoking the completion phase of the database rollforward recovery.

ADM1608I DB2 is invoking the completion phase of the table space rollforward recovery.

ADM1609W DB2 is canceling the database rollforward recovery.

ADM1610W DB2 is canceling the table space rollforward recovery.

ADM1611W The rollforward recovery phase has been completed.

ADM1612W Online table space rollforward recovery could not complete due to the high volume of database activity. Either reduce the database activity, increase the MAXAPPLS configuration parameter, or perform an offline table space rollforward.

ADM1613W The table space *tablespaceName* (ID *tablespaceID*) was previously rolled forward to *timestampPIT* at *timestampRFWD*.

ADM1614W The table space *tablespaceName* (ID *tablespaceID*) has been placed into restore pending state. The rest of the table spaces have been rolled forward to the end of logs.

ADM1615W DB2 cannot roll forward table space *tablespaceName* (ID *tablespaceID*) because it is still used by an active transaction. The application handle is *appHandle*. Terminate this application by COMMIT, ROLLBACK, or FORCE APPLICATION.

ADM1616E DB2 could not find the dropped table log record that matches the given dropped table with ID *droppedTableID* during recovery.

ADM1617W The table space *tablespaceName* (ID *tablespaceID*) that was included in the last rollforward has not yet completed. It has not been included in the table space list provided for this rollforward. It has been placed into restore pending state.

ADM1618W DB2 cannot recover table space *tablespaceName* (ID *tablespaceID*) during database rollforward. The table space has been placed in rollforward pending state. To recover, perform a table space rollforward after the database rollforward is complete.

ADM1619W Rollforward cannot stop here. Rollforward to LSN *rollforwardLSN* in log file *logFilename*.

ADM1620W Table partition with objectID *objectID* in tablespace *tablespaceID* was skipped during drop table recovery. To recover the data from this partition include the tablespace in the rollforward list.

ADM1700W DB2 is unable to use the log path specified by the backup image; Switching to the default log path.

ADM1701W The USEREXIT log retention logging is not enabled because the LOGARCHMETH1 DB configuration parameter is set to LOGRETAIN.

User response:

ADM1710C The database is not accessible because the DB2 database manager cannot write to the log control file *file-name* on database partition *database-partition-num*. Possible reasons for this include: the file does not exist; the file is not accessible to the database manager because of network, file system, or OS problems; or incorrect file permissions are preventing the database manager from writing to the file. Ensure that the file is accessible to the database manager and that the file system is functioning properly, and restart or reconnect to the database.

ADM1711W The log control file *log-control-file-name* on database partition *database-partition-num* was missing or corrupt. This might be a result of the database previously being stopped abnormally. The database manager recreated this file.

ADM1712C The log control file *file-name* on database partition *database-partition-num* was missing or corrupt. The DB2 database manager attempted to recreate it, but was not successful. Possible reasons for this include: the file is not accessible to the database manager because of network, file system, or OS problems; or incorrect file permissions are preventing the database manager from reading from or writing to the file. Ensure that the file is accessible to the database manager and that the file system is functioning properly.

ADM1713C The DB2 database manager cannot start the database on database partition *database-partition-num* because the database manager can open neither the primary log control file *primary-log-file* nor the secondary log control file *secondary-log-file*.

ADM1800E DB2 was unable to confirm log *logNumber* was archived on the old log path. Take a database backup to ensure that the database may be recovered.

ADM1801W DB2 could not access more than 2GB of file storage on the raw device. Therefore, only 2GB of file storage will be used for logging.

ADM1802W Database was put into offline backup pending state because the LOGFILSIZ DB configuration parameter has changed while using raw device for logging.

ADM1803W There is not enough space left on raw device *path* for log files. The actual device size is *actualDevice* 4K pages. The minimum required device size is *minimumDevice* 4K pages.

ADM1804W The raw device is too small to support active log space. *actualDeviceSize* 4K pages are available, however, *minimumDeviceSize* 4K pages are required. Use a larger device or reduce the LOGPRIMARY and/or the LOGFILSIZ DB configuration parameters.

ADM1805E DB2 is unable to use the NEWLOGPATH DB configuration parameter because the raw device is already being used as a log or table space container elsewhere.

ADM1806E Unable to use the NEWLOGPATH DB configuration parameter *newLogPath* DB2 will continue to use the current log path.

ADM1807E The new DB configuration parameter values for NEWLOGPATH and MIRRORLOGPATH are identical, this is not allowed. Ensure that the values are different for the two parameters.

ADM1810E The new value of the MIRRORLOGPATH DB configuration parameter is invalid or cannot be used. DB2 will ignore the new value. Confirm the new path value and ensure that it is valid.

ADM1811E DB2 will now switch to the default log path *logpath*.

ADM1812E The new value of the MIRRORPATH DB configuration parameter *newMirrorPath* is not usable. DB2 will continue to use the existing mirror path *defaultMirrorPath*.

ADM1813E The current mirror path *currentMirrorPath* is invalid.

ADM1814E The current log path *currentLogPath* is invalid.

ADM1815E An error occurred while DB2 was trying to remove log files from *logFilePath*.

ADM1817E DB2 was unable to execute the user exit program when archiving log file *logFilename* from *dirPath* for database *DBName*. The error code was *returnCode*. Verify user exit program by manually running it.

ADM1818E An error was received from the user exit program. DB2 will not call the user exit program for this database for 5 minutes.

ADM1819C DB2 was unable to switch to the new log file size because an error occurred while archiving an old log on the raw device. As a result, the database cannot be accessed.

ADM1820W For USEREXIT to be enabled, you cannot have both DB configuration parameters LOGRETAIN set to CAPTURE and USEREXIT turned on. Therefore, USEREXIT is disabled.

ADM1821E The log path has been reset to the default value.

ADM1822W The active transaction log is being held by dirty pages. Database performance may be impacted.

Explanation: The rate at which the database work load is generating dirty pages has exceeded the rate at which the page-cleaner agents are writing dirty pages to disk. Dirty pages that have not yet been written to disk are holding the transaction log.

User response:

1. Reduce the database work load, if possible.

2. If this problem persists, take the following action:

- Decrease the value of the PAGE_AGE_TRGT_MCR and PAGE_AGE_TRGT_GCR database configuration parameters
 - Increase the value of the NUM_IOCLEANERS database configuration parameter
-

ADM1823E The active log is full and is held by application handle *handle*. Terminate this application by COMMIT, ROLLBACK or FORCE APPLICATION.

ADM1824W DB2 cannot delete the log file *filename*. User must manually delete the log file.

ADM1825W DB2 cannot create the next log file because of a user exit problem.

ADM1826E DB2 cannot continue because the disk used for logging is full.

ADM1827E There is no space left on raw device used for logging. User must make an offline backup of the database.

ADM1828C DB2 will attempt to create the log file again in 5 minutes.

ADM1829E The active log is full and is held by an indoubt transaction.

Explanation: This message is returned when the database manager must release the current active log file and use a different log file, but cannot release the current active log file because an indoubt transaction has a lock on the current active log file.

User response: Resolve the indoubt transaction using the LIST INDOUBT TRANSACTIONS WITH PROMPTING command.

ADM1830E The active log path is set to the default value.

ADM1831E DB2 was unable to execute user exit program when retrieving log file *logFilename* to *dirPath* for database *DBName*. The error code was *returnCode*. Verify the user exit program by manually running it.

ADM1832E DB2 was unable to find the user exit program when archiving log file *logFilename* from *dirPath* for database *DBName*. The error code was *returnCode*.

ADM1833E The user exit program returned an error when archiving log file *logFilename* from *dirPath* for database *DBName*. The error code was *returnCode*.

ADM1834E DB2 was unable to find the user exit program when retrieving log file *logFilename* to *dirPath* for database *DBName*. The error code was *returnCode*.

ADM1835E The user exit program returned an error when retrieving log file *logFilename* to *dirPath* for database *DBName*. The error code that was returned by the user exit program to the database manager was: *returnCode*.

Explanation: You can automate log file archiving and retrieval by creating a user exit program that the DB2 database manager calls to carry out the archiving or retrieval operation.

When the DB2 database manager invokes your user exit program, the following happens:

- The database manager passes control to the user exit program
- The database manager passes parameters to the user exit program
- On completion, the user exit program passes a return code back to the database manager

The DB2 database manager can only handle specific error codes. However, a user exit program might encounter many different kinds of error conditions, such as operating system errors. A user exit program must map the error conditions it encounters to error codes that the database manager can handle.

This message is returned when a user exit program fails and returns the specified return code to the DB2 database manager.

User response:

1. Refer to the documented list of standard user exit error codes.
 2. Refer to any error handling information for your user exit program.
-

ADM1836W The table *tableID* (ID *tableID*) on *tablespaceName* (ID *tablespaceName*) is in DATALINK Reconcile Pending (DRP) state.

ADM1837W The table *tableName* (ID *tableID*) on *tablespaceName* (ID *tablespaceID*) is in DATALINK Reconcile Not Possible (DRNP) state.

ADM1838W An application is waiting for a lock held by an indoubt transaction. This will cause the application to wait indefinitely. Use the LIST INDOUBT TRANSACTIONS command to investigate and resolve the indoubt transactions.

ADM1839W An error occurred while DB2 was writing log data to log file *logFile* on path *logPath1*. The log file is still available on log path *logPath2*. DB2 will attempt to use both paths for subsequent log files. In the meantime, check that the specified log path in which the error occurred exists and is accessible. Also check that there is space available in the file system.

ADM1840W The error on log path *logPath* has been resolved. DB2 will now write log files to this path.

ADM1841W DB2 was unable to locate log *logNumber* while attempting to archive it. Your existing recovery strategy may not work if the missing log file is required during recovery. A database backup is required to ensure the recoverability of the database. However, the backup should be taken after the First Active Log file (LOGHEAD) DB configuration parameter is beyond the specified log. You might want to consider deactivating the database now to have the First Active Log file (LOGHEAD) move up, and then take a backup.

User response:

ADM1842W The database configuration update was successful, but the database could not be made recoverable.

Explanation: Although the update was successful, the database is not recoverable. Possible reasons for this are:

- At least one table space is in a table space state other than "Normal"
- Member crash recovery is pending on at least one member in this DB2 pureScale instance.

The database will continue to use circular logging until all of the table spaces are in a "Normal" state.

In environments using DB2 pureScale, all members are consistent.

User response: If this error is being returned because one or more table spaces are not in "Normal" state, respond to this error by performing the following steps:

1. Determine which table spaces are not in the "Normal" state using the MON_GET_TABLESPACE table function.
2. For each table spaces that is not in "Normal" state, resolve the issue that corresponds to the state that it is in.

If this error is being caused by a pending member crash recovery, perform member crash recovery on each member that is inconsistent. In most cases, this requires no user response because member crash recovery is initiated automatically. If member crash recovery is not initiated automatically, issue a RESTART DATABASE command.

The database will be recoverable when it is next started or activated. A full database backup will be required at that time.

ADM1843I Started retrieve for log file *logFilename*.

ADM1844I Started archive for log file *logFilename*

ADM1845I Completed retrieve for log file *logFilename* on chain *chain* from *destPath*.

ADM1846I Completed archive for log file *logFilename* to *destPath* from *dirPath*.

ADM1847W Failed to retrieve log file *logFilename* on chain *chain* to *destPath*.

ADM1848W Failed to archive log file *logFilename* to *destPath* from *dirPath*.

ADM1849C The database has reached *current-LSN/LSO/LFS-value* for log record identifier type *LSN/LSO/LFS-type*, which is approaching its maximum value. After that maximum value is reached, you will no longer be able to use the database.

Explanation: The database manager identifies database log records using various non-decreasing identifiers: LSN, LFS, and LSO.

Your database has used nearly all of the possible values for at least one of these identifiers.

User response: Before your database runs out of unique LSN/LSO/LFS values, perform the following steps:

1. Unload all of the data from the database.
2. Drop and recreate the database.
3. Reload the data.

The LSN/LSO/LFS values will restart at 0 after these steps have been completed.

For additional assistance in responding to this error, contact IBM software support.

ADM1850C The database has run out of log record identifiers. Log record identifier type *LSN/LSO/LFS-type* has reached *current-LSN/LSO/LFS-value*. The database cannot process any transactions that require additional log records to be written.

Explanation: The database manager identifies database log records using various non-decreasing identifiers: LSN, LFS, and LSO.

Your database has used nearly all of the possible values for the *LSN/LSO/LFS-type* identifier. No more log records can be written.

User response: Perform the following steps:

1. Unload all of the data from the database.
2. Drop and recreate the database.
3. Reload the data.

The LSN/LSO/LFS values will restart at 0 after these steps have been completed.

For additional assistance in responding to this error, contact IBM software support.

Chapter 5. ADM2000 - ADM2499

ADM2000E The Event Monitor *eventMonitor* was deactivated because it encountered an I/O error.

ADM2001W The Event Monitor *eventMonitor* was deactivated because the MAXFILES and MAXFILESIZE CREATE EVENT MONITOR parameters' limits have been reached.

ADM2002E The Event Monitor *eventMonitor* was deactivated when the process reading from the target pipe had disconnected.

ADM2003W The Event Monitor *eventMonitor* was not activated because the limit on the number of active event monitors has already been reached.

ADM2004E The Event Monitor *eventMonitor* was not activated because there was not enough storage space in the database monitor heap. To remedy this problem, increase the MON_HEAP_SZ DBM configuration parameter.

ADM2005W The database monitor heap is exhausted. Increase the MON_HEAP_SZ DBM configuration parameter.

ADM2006W The database monitor heap is exhausted. Statement elements may be unreliable. Increase the MON_HEAP_SZ DBM configuration parameter.

ADM2007W The event monitor *eventMonitor* was deactivated due to PCTDEACTIVE limit reached. To reactivate the named event monitor, either increase the table space size or increase the PCTDEACTIVE threshold.

ADM2009C The Event Monitor *eventMonitor* detected on table *tableName* (ID *tableID*) that the column name *colName* is invalid.

ADM2010C The Event Monitor *eventMonitor* detected on table *tableName* (ID *tableID*) that the column *colName* is not allowed.

ADM2011C The Event Monitor *eventMonitor* detected on table *tableName* (ID *tableID*) that the column *colName* has an invalid *dataType* data type.

ADM2012C For table *tableName*, (ID *tableID*) the row size is bigger than the page size.

ADM2013C The Event Monitor *eventMonitor* detected on table *tableName* (ID *tableID*) that the column *colName* is incorrect. The first column must be PARTITION_KEY.

ADM2014W The Event Monitor *eventMonitor* detected on table *tableName* (ID *tableID*) that the size of the column *colName* is smaller than the default size of *defaultColumnSize*. Therefore, contents will be truncated to the user specified size.

ADM2015W The SQLCODE received from the remote target node is *SQLCODE*.

ADM2016W The table *tableName* (ID *tableID*) was not found.

ADM2017C The Event Monitor *monitor-name* has reached its file capacity. Delete the files in the target directory *directory* or move them to another directory.

ADM2018W The Event Monitor *monitor-name* has been deactivated. Event monitor data collection cannot be performed.

Explanation: The event monitor deactivated due to some error in processing that is described in probe information found in the administration notification and db2diag log files.

User response: Look in the administration notification or db2diag log file for probe information that indicates why the event monitor deactivated, correct the error, and then reactivate the event monitor by issuing the following statement:

SET EVENT MONITOR *monitor-name* STATE 1

Chapter 6. ADM2500 - ADM2999

ADM2500E A serious error has occurred at the database server *DRDAApplicationServer* which has made continued processing impossible. A dump has been generated. If the remote server is DB2 UDB for OS/390,zOS, check the console log for information about this error. If the remote server is DB2 UDB for iSeries, the job log of the server job, and/or a first failure data capture (FFDC) spooled file is usually necessary to determine the cause of the error. If the remote server is DB2 Database for Linux, UNIX, and Windows, check the remote database server's administration log for information regarding this error.

ADM2501C The amount of data received for a data type at the database server did not match the amount which was expected. A dump has been generated. If the remote server is DB2 UDB for OS/390,zOS, check the console log for information about this error. If the remote server is DB2 UDB for iSeries, the job log of the server job, and/or a first failure data capture (FFDC) spooled file is usually necessary to determine the cause of the error. If the remote server is DB2 Database for Linux, UNIX, and Windows, check the remote database server's administration log for information regarding this error.

ADM2502C The remote database server has encountered an error.

Explanation: The client has detected a problem on the remote server.

User response: If the remote server is DB2 UDB for OS/390,zOS, check the console log for information about this error.

If the remote server is DB2 UDB for iSeries, the job log of the server job, and/or a first failure data capture (FFDC) spooled file is usually necessary to determine the cause of the error.

If the remote server is DB2 Database for Linux, UNIX, and Windows, check the remote database server's administration log for information regarding this error.

If the remote server is IBM Informix Dynamic Server, check the online.log or use onstat (-m) for information about this error.

ADM2503C A datastream parsing error has been detected. A dump has been generated.

ADM2504C The LOB data type is not supported on this database server. APAR fixes are required. If this database server is DB2 for OS/390 v6 and v7, apply the fixes for APARs PQ50016 and PQ50017. If the database server is DB2 for iSeries, any release prior to V5.1 is not supported. For DB2 for iSeries v5.1, the PTF for APAR 9A00531 must be applied. If the database server is DB2 for VM and VSE, contact IBM Support.

ADM2505E While DB2 was attempting to utilize sysplex support, a connection could not be established to the database server at SNA address *SNAAddress* due to an unknown cpic symbolic destination name *destName*. Check your SNA configuration and validate it with network support personnel or disable sysplex support at the DB2 Connect Server.

User response:

ADM2506W While an application was attempting to issue a deferred SET statement, the server returned a non-zero sqlcode. The sqlcode is *SQLCODE*. Check your application, and validate the deferred SET statement.

ADM2507I Client Reroute failed because of product signature violation. Original product signature: *product signature*. Retry product signature: *product signature*.

ADM2508I Client Reroute successful. Hostname/IP Address: *hostname/IP address* and Service name/Port number: *service name/port number*.

ADM2509I Database connection is successful. Connecting to --> Hostname/IP Address: *hostname/IP address* and Service name/Port number: *service name/port number*.

Chapter 7. ADM3000 - ADM3499

ADM3000C The network host entry *networkHostEntry* at line *lineNumber* of *db2nodes.cfg* cannot be resolved.

ADM3001C DB2 cannot create the intermediate socket directory *socketDirectory*. Check the permissions of this path.

ADM3003C The node *node* is no longer present in *db2nodes.cfg*, even though this node had existed when FCM had started up. This may indicate a communication failure.

ADM3006C The retry limit was reached while attempting to establish a connection to node *node*. DB2 will now initiate node recovery.

ADM3008C The connection with node *node* has been unexpectedly severed. DB2 will now initiate node recovery.

ADM3019C The link between node *node1* and node *node2* is broken. Check your switch and cables.

ADM3020C The number of FCM buffers is too small. This will impact communication between DB2 agents and lead to runtime errors. If this message persists, you should adjust the FCM_NUM_BUFFERS DBM configuration parameter.

ADM3021W VI is enabled on this node.

Chapter 8. ADM3500 - ADM3999

ADM3500W The group IDs on the database partitions do not match. In a partitioned database, each partition must have the same set of users and groups defined. If the definitions are not the same, a user may be authorized to perform different actions on different partitions. Consistent user and group definitions across all partitions is recommended.

Chapter 9. ADM4000 - ADM4499

ADM4000W A catalog cache overflow condition has occurred. There is no error but this indicates that the catalog cache has exceeded the configured maximum size. If this condition persists, you may want to adjust the CATALOGCACHE_SZ DB configuration parameter.

ADM4001I A failure has occurred while regenerating the view *viewName*.

ADM4002W The Event Monitor target table *targetTableName* (table schema *tableSchema*) already exists.

ADM4003E The UPGRADE DATABASE command failed to upgrade the DB2 Text Search catalogs or indexes due to an error in stored procedure *stored-procedure-name*.

Explanation: The database is enabled for DB2 Text Search and has been upgraded successfully. However, an error occurred while upgrading the text search catalogs or indexes.

User response: Call the stored procedures to upgrade the DB2 Text Search catalogs or indexes.

- If SYSPROC.SYSTS_UPGRADE_INDEX failed to complete successfully, reissue the stored procedure.
 - If SYSPROC.SYSTS_UPGRADE_CATALOG failed to complete successfully, then call these upgrade procedures again in this specific order:
 1. SYSPROC.SYSTS_UPGRADE_CATALOG
 2. SYSPROC.SYSTS_UPGRADE_INDEX
-

ADM4004W The db2ckupgrade utility found that the database contains the following types of objects, which are not supported by the version of DB2 database to which upgrading is being considered: XML global variables; or compiled SQL functions that use XML parameters or that return XML types.

Explanation: You can verify that a given database can successfully be upgraded to a later version of DB2 database using the db2ckupgrade utility.

This message is returned when the db2ckupgrade utility finds database objects which are not supported in the version of DB2 database to which you are upgrading. Specifically, this message is returned when the following database objects are in a database being checked:

- Global variables of type XML
- Compiled SQL functions that use parameters of type XML or that return XML types

If you proceed to upgrade the database to the target version of DB2 database, these objects will be invalidated during database upgrade. You will be unable to use these database objects until you upgrade to a version of DB2 database that supports those database objects.

User response: To be able to use the database objects that are invalidated during upgrade, upgrade to a release and fix pack of DB2 database that supports XML global variables and compiled SQL functions that use XML parameters or that return XML types. When you upgrade to a fixpack that supports those database objects, the objects will be automatically revalidated the first time they are referenced after the database has been upgraded.

ADM4005W The upgrade procedure marked the following types of database objects as "invalidated" because those objects are not supported in the upgraded version of DB2 database: XML global variables; and compiled SQL functions that use XML parameters or that return XML types.

Explanation: You can upgrade a database to a later version of DB2 database using the UPGRADE DATABASE command.

This message is returned when a database that contains the following database objects is upgraded to a version of DB2 database that does not support those types of database objects:

- Global variables of type XML
- Compiled SQL functions that use parameters of type XML or that return XML types

The upgrade operation marked all of these types of database objects as "invalidated". You will be unable to use these database objects until you upgrade to a version of DB2 database that supports those database objects.

User response: To be able to use the database objects that were invalidated during upgrade, upgrade to a release and fix pack of DB2 database that supports XML global variables and compiled SQL functions that use XML parameters or that return XML types. When you upgrade to a fixpack that supports those database objects, the objects will be automatically revalidated the

first time they are referenced after the database has been upgraded.

ADM4014N The ALTER TABLE ATTACH operation failed because either there are no indexes on the source table, or the indexes on the source table *source-table-name* do not match the partitioned index *index-name* on the target table *target-table-name*. Reason: *reason-code*.

Explanation: See the applicable reason code for more information:

1

The indexes on the source table do not match the unique partitioned index on the target table.

2

The indexes on the source table do not match the partitioned XML pattern values index that was created with the REJECT INVALID VALUES on the target table.

3

The indexes on the source table do not match the partitioned indexes on the target table, and the ALTER TABLE ATTACH statement is defined with the REQUIRE MATCHING INDEXES clause.

User response: Respond according to the reason code, and then rerun the ALTER TABLE ATTACH operation.

1

Create a unique index on the source table that matches the unique partitioned index on the target table.

2

Create an XML pattern values index on the source table (using the REJECT INVALID VALUES clause) to match the index on the target table.

3

Either create an index on the source table that matches the partitioned index on the target table, or remove the REQUIRE MATCHING INDEXES clause from the ALTER TABLE ATTACH statement.

ADM4015I No index on the *source-table* source table matches the partitioned index *index-name* on the target table *target-table*. ALTER TABLE ATTACH processing continues.

Explanation: After the ATTACH operation completes successfully, the index partition will be built on the first

access to the newly attached table partition. Usually, the first access is from a SET INTEGRITY statement.

To maximize roll-in efficiency, create indexes on the source table that match the partitioned indexes on the target table before attaching the source table to the target table. For more information, see the topic on attaching data partitions in the DB2 Information Center.

User response: No response is required.

ADM4016I The index *indexName* on the source table *source-table* does not match any partitioned indexes on the target table *target-table*. ALTER TABLE ATTACH processing continues.

Explanation: A successful ATTACH operation will drop this index on the source table.

To maximize roll-in efficiency, drop the indexes on the source table that do not match the indexes on the target table before attaching the source table to the target table. For more information, see the topic on attaching data partitions in the DB2 Information Center.

User response: No response is required.

ADM4100W The db2ckupgrade or UPGRADE DATABASE command identified external routines or user-defined wrappers that could have potential incompatibility with the multi-threaded database manager. See the file *generated-file* that contains the list of routines and wrappers identified, or a set of statements to restore the routine or wrapper definitions altered by the UPGRADE DATABASE command.

Explanation: Beginning in DB2 Version 9.5, the database manager is now multi-threaded as opposed to multi-processed on Linux and UNIX operating systems. The execution of NOT FENCED and NOT THREADSAFE external routines or user-defined wrappers, which are NOT FENCED in the multi-threaded database manager, could lead to incorrect results, database corruption, or abnormal termination of the database manager. As a result, all NOT FENCED routines and all NOT FENCED user-defined wrappers must be THREADSAFE.

During database upgrade, all external NOT FENCED routines that have no dependency on the DB2 engine library are altered to FENCED and NOT THREADSAFE. Also, the DB2_FENCED option is set to 'Y' for all user-defined wrappers. The UPGRADE DATABASE command generates a script *generated-file* that contains the statements to restore the routine or wrapper definitions.

When running db2ckupgrade, external NOT FENCED routines that have no dependency on the DB2 engine library are identified. These routines will be altered to

FENCED and NOT THREADSAFE during database upgrade. The file *generated-file* is generated with a list of all the identified external NOT FENCED routines and user-defined wrappers with the DB2_FENCED option is set to 'N'.

User response: When upgrading the database, verify that all the affected routines and user-defined wrappers can safely be run as NOT FENCED and THREADSAFE. Once verified, they can be altered back to NOT FENCED and THREADSAFE, by running the file *generated-file* to alter all of the routines and user-defined wrappers to NOT FENCED. Modify this file to contain only the statements that should be executed and run the CLP script after the database has been upgraded.

After running db2ckupgrade and before upgrading your database, alter the identified routines, listed in the generated file *generated-file*, to FENCED and NOT THREADSAFE and set the DB2_FENCED option to 'Y' for user-defined wrappers.

ADM4101W The UPGRADE DATABASE command failed to automatically collect statistics on the *table-name* system catalog table. The following RUNSTATS command returned SQLCODE *sqlcode* with tokens *tokens* : command.

Explanation: After successfully completing the database upgrade, the UPGRADE DATABASE command was unable to collect the statistics on the *table-name* system catalog table. The RUNSTATS command returned SQLCODE *sqlcode* with tokens *tokens*.

User response:

1. Determine the appropriate user response based on the SQLCODE *sqlcode* returned.
2. Correct the problem and re-issue the RUNSTATS command on the *table-name* system catalog table to collect statistics.

ADM4102W The database contains one or more identifiers called NULL. To avoid conflict with the NULL keyword, you should qualify or delimit with double quotes any identifiers called NULL in SQL statements.

Explanation: An untyped NULL specification can occur anywhere in an expression. If an identifier called NULL is used in an SQL statement without being fully qualified or delimited, the identifier specification might resolve to the NULL keyword instead of the identifier reference. This would result in a change in behavior from previous releases.

User response: Determine if you have identifiers called NULL in the database by issuing the following statements:

- SELECT TABSCHEMA, TABNAME, COLNAME FROM SYSCAT.COLUMNS WHERE COLNAME = 'NULL';
- SELECT ROUTINESCHEMA, ROUTINENAME, PARMNAME FROM SYSCAT.ROUTINEPARMS WHERE PARMNAME = 'NULL';
- For Version 9.5 databases or later, SELECT VARSHEMA, VARNAME FROM SYSCAT.VARIABLES WHERE VARNAME = 'NULL'.

To avoid conflict with the NULL keyword, qualify or delimit with double quotes the identifiers called NULL in SQL statements.

ADM4103W The connection attribute *attributename* contains asterisks (*) in the workload *workloadname*. A single asterisk (*) will be replaced with two asterisks () during database upgrade. Reason code = *reason-code*.**

Explanation: Starting with DB2 Version 9.7, you can use a single asterisk (*) as a wildcard character and two asterisks (**) to represent one literal asterisk (*) in some workload attributes.

The db2ckupgrade command identifies asterisks (*) in the connection attribute and the UPGRADE DATABASE command replaces the single asterisk (*) with two asterisks (**) if the type of connection attribute is one of the following values: 1 (APPLNAME) 6 (CURRENT CLIENT_USERID) 7 (CURRENT CLIENT_APPLNAME) 8 (CURRENT CLIENT_WRKSTNNAME) 9 (CURRENT CLIENT_ACCTNG)

Reason codes are as follows:

1

The connection attribute was truncated because it reached the maximum length when a single asterisk (*) is replaced with two asterisks (**) during database upgrade.

2

The connection attribute was not truncated when a single asterisk (*) is replaced with two asterisks (**) during database upgrade.

User response: Replace asterisks (*) in the connection attribute with another character if possible.

ADM4104E One or more databases are enabled for XML Extender. You must remove the XML Extender functionality from the instance and databases before upgrading.

Explanation: Starting with DB2 Version 9.7, XML Extender is discontinued. Possible reasons for this error are:

- The instance that you specified for upgrade has XML Extender functionality enabled, and the implicit call to db2ckupgrade fails because one or more databases are enabled for XML Extender.
- The database that you are checking for upgrade is enabled for XML Extender.

User response: Remove the XML Extender functionality from the instance and disable databases for XML Extender. Then re-issue the db2iupgrade or db2ckupgrade command.

Refer to the DB2 Information Center for details on the steps to upgrade XML Extender, including how to disable XML Extender in databases.

ADM4105W The Database is enabled for DB2 WebSphere MQ functions. The set of functions defined for XML Extender will be dropped during database upgrade.

Explanation: Starting with DB2 Version 9.7, XML Extender is discontinued. The set of DB2 WebSphere MQ functions that are defined for XML Extender will be dropped during database upgrade. These functions declare parameters using XML Extender user-defined data types.

User response: After the database upgrade, if you want to use DB2 WebSphere MQ functions with xml data type parameters, run the enable_MQFunctions command with the -force and -xml parameters to create new MQ functions for XML data type and re-create the existing MQ functions. Refer to the DB2 Information Center for details on how to run the enable_MQFunctions command.

ADM4106W The upgraded database was enabled for XML Extender. XML Extender functionality was not disabled during database upgrade. However, this functionality will fail to run.

Explanation: Starting with DB2 Version 9.7, XML Extender has been discontinued. The database that you restored from a backup image created in a previous release is enabled for XML Extender.

The database was upgraded successfully. However, the XML Extender functionality such as routines will fail to run.

User response: Restore the database to a DB2 copy running a DB2 database product version that supports XML Extender, then follow the steps to upgrade from XML Extender.

Refer to the DB2 Information Center for details on the steps to upgrade from XML Extender.

ADM4200N The data type of the column could not be changed because of the data in the row with row identifier rowid.

Explanation: You can change the data type of a table column by using the ALTER TABLE statement with the ALTER COLUMN clause and the SET DATA TYPE clause.

This message is returned when an attempt to change the data type of a column fails because of an incompatibility between data in one of the rows of the table and the data type to which the column is being changed.

User response:

1. Review the db2diag logs for more detailed information about why the data in the named row prevented the data type of the column from being changed.
2. Resolve the incompatibility by doing one of the following:
 - Modify the data in the row so that the data type of the column can be changed.
 - Change the data type of the column to one that is compatible with the existing data.

ADM4201N The data type of the column could not be changed because of the default values for the column named column-name in the table named table-name.

Explanation: You can change the data type of a table column by using the ALTER TABLE statement with the ALTER COLUMN clause and the SET DATA TYPE clause.

This message is returned when an attempt to change the data type of a column fails because of an incompatibility with the default values for the column being altered.

User response:

1. Review the db2diag logs for more detailed information about why the default values for the column prevented the data type of the column from being changed.
2. Resolve the incompatibility by doing one of the following:
 - Modify or drop the default values for the column so that the data type of the column can be changed.
 - Change the data type of the column to one that is compatible with the existing data.

Chapter 10. ADM4500 - ADM4999

ADM4500W A package cache overflow condition has occurred. There is no error but this indicates that the package cache has exceeded the configured maximum size. If this condition persists, you should perform additional monitoring to determine if you need to change the PCKCACHESZ DB configuration parameter. You could also set it to AUTOMATIC.

Chapter 11. ADM5500 - ADM5999

ADM5500W The database manager is performing lock escalation. The affected application is named *appl_name*, and is associated with the workload name *workload_name* and application ID *appl_id* at member *member_num*. The total number of locks currently held is *locksHeld*, and the target number of locks to hold is *targetNumber*. Reason code: *reason_code*

Explanation: The reason code indicates why lock escalation is happening:

1

The portion of the lock list that is being held by an application has exceeded the value specified for the maxlocks database configuration parameter.

2

The lock list is full.

3

The global lock list is full.

User response: To reduce the likelihood of lock escalation, take one or more of the following actions:

- Perform commit operations more frequently.
- Configure the lock list size to automatically adjust by setting the locklist configuration parameter to AUTOMATIC.
- Configure the global lock list size to automatically adjust by setting the cf_lock_size configuration parameter to AUTOMATIC.
- If you are making several updates, consider locking the entire table using the LOCK TABLE statement. Locking the table uses only one lock, as opposed to the several locks that might be required when updating multiple rows.
 - You can also specify that only table locks be used for a table using the LOCKSIZE option of the ALTER TABLE statement.
 - If you use the repeatable read isolation level, the optimizer might elect to acquire a single table-level lock, if it appears that lock escalation is likely.
- Use the cursor stability isolation level whenever possible to decrease the number of shared locks held. If application integrity requirements are not compromised, use the uncommitted read isolation level instead of cursor stability to further decrease the amount of locking.

ADM5501I DB2 is performing lock escalation. The affected application is named *appl_name*, and is associated with the workload name *workload_name* and application ID *appl_id* at member *member_num*. The total number of locks currently held is *locksHeld*, and the target number of locks to hold is *targetNumber*. The current statement being executed is *currentStatement*. Reason code: *reason_code*

Explanation: The possible reason codes that explain why lock escalation is happening are as follows:

1

The portion of lock list held by an application has exceeded the value specified for the maxlocks database configuration parameter.

2

The lock list is full.

3

The global lock list is full.

User response: To reduce the likelihood of lock escalation, you can take one or more of the following actions:

- Perform commit operations more frequently.
- Set the locklist configuration parameter to AUTOMATIC to have the lock list size adjusted automatically.
- Set the cf_lock_sz configuration parameter to AUTOMATIC to have the global lock list size adjusted automatically.
- If you are making several updates, consider locking the entire table using the LOCK TABLE statement. Locking the table uses only one lock, as opposed to the several locks that might be required when updating multiple rows.
 - You can also specify that only table locks be used for a table using the LOCKSIZE option of the ALTER TABLE statement.
 - If you use the Repeatable Read (RR) isolation level, the optimizer might elect to acquire a single table-level lock, if it appears that lock escalation is likely.
- Use the Cursor Stability isolation level whenever possible to decrease the number of share locks held. If application integrity requirements are not compromised, use the Uncommitted Read isolation level instead of Cursor Stability to further decrease the amount of locking.

ADM5502W The escalation of *numLocks* locks on table *tableName* to lock intent *lockIntent* was successful.

Explanation: The token *numLocks* lists the total number of locks currently held by all applications that are accessing the table.

User response: No user response is required.

ADM5503E The escalation of *numLocks* locks on table *tableName* to lock intent *lockIntent* has failed. The **SQLCODE** is *SQLCODE*.

ADM5504W The escalation of *numLocks* locks on **DATAPARTITIONID** *datapartitionid* of table *tableName* to lock intent *lockIntent* was successful.

ADM5505E The escalation of *numLocks* locks on **DATAPARTITIONID** *datapartitionid* of table *tableName* to lock intent *lockIntent* has failed. The **SQLCODE** is *SQLCODE*.

ADM5506W The current unit of work was involved in an unresolved contention for use of an object. The type of the event is: *event-type*. The identifier of the lock on which this event happened is: *lock-ID*. The timestamp of the event is: *timestamp*. The identifier of the member at which the event happened, and the identifier of the event is: *member-ID-event-ID*. The affected application is named *application-name*. The application is associated with the workload named *workload-name*. The application identifier is: *application-ID*. The identifier of the member on which the application is running is: *app-member-id*. The role that this application plays with respect to this lock is: *role*.

Explanation: The current unit of work was involved in an unresolved contention for use of an object.

Lock escalation is the process of replacing row locks with table locks, reducing the number of locks in the list.

User response: To investigate this event further, create an event monitor using the CREATE EVENT MONITOR FOR LOCKING statement which should help to narrow down the source of the problem.

To help avoid a deadlock or lock timeout, issue frequent COMMIT operations, if possible, for a long-running application, or for an application likely to encounter a deadlock.

Deadlocks are often normal or expected while processing certain combinations of SQL statements. It is

recommended that you design applications to avoid deadlocks to the extent possible.

If a deadlock state was reached because of a queuing threshold such as the CONCURRENTDBCOORDACTIVITIES threshold, increase the value of the queuing threshold.

For more detailed information about preventing deadlocks or lock timeouts, search the DB2 Information Center using phrases such as "deadlock prevention", and terms such as "deadlocks" and "lock timeouts".

To prevent lock escalations, adjust the **locklist** and **maxlocks** configuration parameters. For assistance, see the DB2 Information Center.

ADM5507W The current unit of work was involved in an unresolved contention for use of an object. The type of the event is: *event-type*. The identifier of the lock on which this event happened is: *lock-ID*. The timestamp of the event is: *event-timestamp*. The identifier of the member at which the event happened is: *event-source-member-id*. The identifier of the event is: *event-ID*. The identifier of the member at which the application is holding the lock is: *app-member-id*.

Explanation: The current unit of work was involved in an unresolved contention for use of an object.

Lock escalation is the process of replacing row locks with table locks, reducing the number of locks in the list.

User response: To investigate this event further, create an event monitor using the CREATE EVENT MONITOR FOR LOCKING statement which should help to narrow down the source of the problem.

To help avoid a deadlock or lock timeout, issue frequent COMMIT operations, if possible, for a long-running application, or for an application likely to encounter a deadlock.

Deadlocks are often normal or expected while processing certain combinations of SQL statements. It is recommended that you design applications to avoid deadlocks to the extent possible.

If a deadlock state was reached because of a queuing threshold such as the CONCURRENTDBCOORDACTIVITIES threshold, increase the value of the queuing threshold.

For more detailed information about preventing deadlocks or lock timeouts, search the DB2 Information Center using phrases such as "deadlock prevention", and terms such as "deadlocks" and "lock timeouts".

To prevent lock escalations, adjust the **locklist** and **maxlocks** configuration parameters. For assistance, see the DB2 Information Center.

ADM5508I One or more partitioned indexes on data partition *DATAPARTITIONID* of the table named *tableName* are marked invalid and require rebuilding.

Explanation: The DB2 data server will automatically rebuild the invalid index partitions on this data partition. A super exclusive Z lock will be held on the data partition from the time that the index rebuild begins until the end of unit of work in which the rebuild occurs.

User response: No response is required.

ADM5509I Rebuilding *count* index partitions on data partition *DATAPARTITIONID* of the table *tableName*.

Explanation: The data server is rebuilding index partitions for the table on the specified data partition.

User response: No response is required.

ADM5510E The table space *tableSpace* (ID *tableSpace*) is full. There is no more room in the internal object table of this table space.

ADM5511E The object *object* with id *object-id* in tablespace *tblspace-id* for table *tableName* has reached the maximum possible size.

ADM5512N Rebuild of index partitions on partition *DATAPARTITIONID* of table *tableName* failed with *SQLCODE* of *SQLCODE*.

Explanation: The rebuild of the index partitions failed for the specified data partition.

User response: See the DB2 Information Center for information about the given *SQLCODE*.

ADM5513I Index rebuild on the data partition *DATAPARTITIONID* of table *tableName* completed successfully.

Explanation: Although the data server successfully rebuilt the index for the data partition, other partitioned indexes on the data partition might still require rebuilding. This index rebuild will occur during the current unit of work.

User response: No response is required.

ADM5514I Rebuilding index partition with IID *indexIID* in object *indexObjectID* and table space *indexTablespaceID* on data partition *DATAPARTITIONID* of table *tableName*.

Explanation: The data server is rebuilding the specified index partition.

User response: No response is required.

ADM5515I Index partitions on data partition *DATAPARTITIONID* of table *tableName* are successfully rebuilt.

Explanation: The data server has rebuilt the index partitions for the table on the specified data partition.

User response: No response is required.

ADM5520E This database version is not supported for database upgrade.

Explanation: The database version that you are trying to upgrade is unsupported in the DB2 copy from which you try to upgrade the database.

User response: Determine what releases are supported for database upgrade in the DB2 copy from which you want to upgrade the database. Upgrade the database to one of those releases. Then re-try upgrading the database to the DB2 copy.

Refer to the DB2 Information Center for details on what versions are supported for database upgrade.

ADM5521C The table *tableName* is missing or it was truncated and was not migrated. Database migration will continue, but this table will be inaccessible.

ADM5530W The COMMIT processing of table *tableName* that used NOT LOGGED INITIALLY has been initiated. It is recommended that you take a backup of this table's table space(s).

ADM5540W Rebuilding *count* indexes on table *tableName*.

Explanation: The rebuild is either for nonpartitioned indexes on a partitioned table, or for indexes on a nonpartitioned table.

User response: No response is required.

ADM5541W Rebuilding index with IID *indexIID* in object with ID *indexObjectID* and table space ID *indexTablespaceID* on table *tablename*.

Explanation: The rebuild is either for nonpartitioned indexes on a partitioned table, or for indexes on a nonpartitioned table.

User response: No response is required.

ADM5542W Indexes on table *tableName* are successfully rebuilt.

Explanation: The rebuild was either for nonpartitioned indexes on a partitioned table, or for indexes on a nonpartitioned table.

User response: No response is required.

ADM5543E Rebuild of indexes on table *tablename* failed with SQLCODE of *SQLCODE*.

Explanation: The rebuild was either for nonpartitioned indexes on a partitioned table, or for indexes on a nonpartitioned table.

User response: See the referenced SQLCODE for more information.

ADM5550C The table space *tablespaceName* (ID *tablespaceID*) is being removed from the rollforward set. The SQLCODE is *SQLCODE*.

ADM5560C DB2 is unable to redo the reorganization of a table unless both the data table space *dataTablespaceName* (ID *dataTablespaceID*) and the long table space *LongTablespaceName* (ID *dataTablespaceID*) are being rolled forward together. Ensure that both of the table spaces are being rolled forward together, or restore a backup image taken after the reorganization to eliminate the need for reorganization to be redone.

ADM5561C DB2 is unable to redo the reorganization of a table unless both the data table space *dataTablespaceName* (ID *dataTablespaceID*) and the long table space *longTablespaceName* (ID *dataTablespaceID*) are at the same point in time. Ensure that a backup of both table spaces are either from before the table reorganization (so that it can be redone) or from after the table reorganization (so that it does not have to be redone).

ADM5562C DB2 is unable to undo the in-place reorganization of a table unless both the data table space *dataTablespaceName* (ID *dataTablespaceID*) and the index table space *indexTablespaceName* (ID *dataTablespaceID*) are being rolled forward together. Ensure that both table spaces are being rolled forward together.

ADM5570W Access was attempted on an unavailable object with id *object-id* in tablespace *tbspace-id* for table *tableName*. If the object is a table it will have to be dropped. If the object is a partition it will have to be detached. If the object is a non-partitioned index the index will have to be dropped.

ADM5571W The *object* object with ID *object-id* in table space *tbspace-id* for table *tableName* is being marked as unavailable.

Explanation: The object cannot be accessed.

User response: If the object is a table or a nonpartitioned index, drop it. If the object is a data partition of a partitioned table, detach it.

ADM5572I One or more indexes on table *tableName* are marked invalid and require rebuilding.

Explanation: The DB2 data server will automatically rebuild the invalid indexes on this table. The rebuild is either for nonpartitioned indexes on a partitioned table or for indexes on a nonpartitioned table.

A super exclusive Z lock will be held on the table for the duration of the index rebuild and the unit of work in which the rebuild occurs.

User response: No response is required.

ADM5580W The table space with ID *tablespaceID* either does not exist or is not in the set of table spaces being recovered. Recovery will continue, but filtering of this table space or tables/objects in this table space will be ignored.

ADM5581W DB2 has successfully filtered table space with ID *tablespaceID*, object ID *object-ID*.

ADM5582C Internal error *error* occurred while attempting to filter table space ID *tablespaceID*, object ID *objectID*.

ADM5583W The total allotted time of 16.7 hours per Data Links Manager has been exceeded.

ADM5590E The specified INPLACE table reorganization action on table *tableName* is not allowed on this node because of SQLCODE -2219 reason code *reasonCode*.

ADM5591W A new compression dictionary could not be built for the object of type *object-type* and object ID *objectID* because insufficient data was found. The object is in the table named *table-name*, which is in the table space: *table-space-id*.

Explanation: When Automatic Dictionary Creation (ADC) occurs, a compression dictionary is created for a table when the table is enabled for data row compression. A dictionary is created to compress data in the rows of the database table. For a tables with columns of type XML, an additional, separate dictionary is created to compress the data in the XML storage object of the table.

A compression dictionary for the table *table-name* was not created. If the *object-type* is DATA, a dictionary could not be created for rows in the table *table-name*. If the *object-type* is XML, a dictionary could not be created for the XML storage object of the table.

If a dictionary exists and a new compression dictionary cannot be built, the following occurs:

- When not using LOAD, the existing dictionary is retained and used.
- When using LOAD, the existing dictionary is not retained.

User response: No user action required.

ADM5592I A compression dictionary for the *object-type* object with ID *objectID* in tablespace *table-spaceID* for table *table-name* was built by dictionary-creator processing.

Explanation: When Automatic Dictionary Creation (ADC) occurs, a compression dictionary is created for a table when the table is enabled for data row compression. A dictionary is created to compress data in the rows of the database table. For a tables with columns of type XML, an additional, separate dictionary is created to compress the data in the XML storage object of the table.

If the *object-type* is DATA, a compression dictionary was created for the table object for the table *table-name*. If the *object-type* is XML, a compression dictionary was created for the XML storage object of the table.

User response: No user action required.

ADM5593I Automatic Dictionary Creation (ADC) processing for the *object-type* object with ID *objectID* in tablespace *table-spaceID* for table *table-name* has been temporarily disabled. ADC processing will be enabled again once the instance is restarted.

Explanation: When Automatic Dictionary Creation (ADC) occurs, a compression dictionary is created for a

table when the table is enabled for data row compression. A dictionary is created to compress data in the rows of the database table. For a tables with columns of type XML, an additional, separate dictionary is created to compress the data in the XML storage object of the table.

ADC processing has been disabled for the object *objectID*. If the *object-type* is DATA, ADC processing is disabled for the table rows for the table *table-name*. If the *object-type* is XML, ADC processing is disabled for the XML storage object of the table.

Even if the table contains sufficient table row data to create a dictionary, a dictionary might not be created if the data is fragmented .

User response: To defragment table row data, use the REORG TABLE command to reorganize the table *table-name*. If *object-type* is XML use the LONGLOBDATA option to reorganize the XML storage object of the table

To enable ADC processing, restart the instance.

ADM5594I Automatic Dictionary Creation (ADC) processing for *object-type* objects has been temporarily disabled for the database.

Explanation: When Automatic Dictionary Creation (ADC) occurs, a compression dictionary is created for a table when the table is enabled for data row compression. A dictionary is created to compress data in the rows of the database table. For a tables with columns of type XML, an additional, separate dictionary is created to compress the data in the XML storage object of the table.

This message is returned when ADC processing is disabled for the database. If the *object-type* is DATA, ADC processing is disabled for the data in tables. If the *object-type* is XML, ADC processing is disabled for data in the XML storage objects of tables.

User response: To enable ADC processing, restart the database.

ADM5595E An index data inconsistency is detected on table *schema-name. table-name*. Please run "INSPECT CHECK TABLE NAME *table-name* SCHEMA *schema-name* INDEXDATA RESULTS KEEP *table-name_resfile.out*" on the node that failed, and then contact DB2 support team to report the problem.

ADM5600I

ADM5600I Scan sharing is temporarily limited due
 to memory constraints.

User response: No response is required.

Chapter 12. ADM6000 - ADM6499

ADM6000E DB2 encountered a read error while reading page *page-number* from tablespace *tbspace-id* for object *object-id* (located at offset *offset* of container *container-path*). DB2 was able to bypass the error by re-reading the page, but this may indicate the presence of a serious problem that could result in a future outage. You may wish to begin an investigation of the hardware (such as disk, controllers, network) and/or the filesystem involved in accessing the container.

ADM6001I A container was renamed based on the rules defined in the path rename configuration file *configFile*. The container *oldName* has been renamed to *newName*.

ADM6002E A container path of an invalid length was specified. The path specified was *containerName*. Refer to the documentation for SQLCODE -297.

ADM6003I A container was not renamed because it was not affected by the rules defined in the path rename configuration file *configFile*. Its name remains as *contname*.

ADM6004N The SET WRITE SUSPEND command failed for the database. Database name: *database-name*.

Explanation: An error occurred during an attempt to suspend the write operations for the database.

User response:

1. Investigate the cause of the failure by reviewing the DB2 diagnostic log files.
 2. Correct the problem.
 3. Reissue the SET WRITE SUSPEND command again.
-

ADM6005N The SET WRITE RESUME command failed for the database. Database name: *database-name*.

Explanation: An error occurred during an attempt to resume the write operations for the database.

User response:

1. Investigate the cause of the failure by reviewing the DB2 diagnostic log files.

2. Correct the problem.
 3. Reissue the SET WRITE RESUME command again.
-

ADM6006E DB2 encountered an error while reading page *page-number* from table space *tbspace-id* for object *object-id* (located at offset *offset* of container *container-path*).

Explanation: DB2 was not able to complete the operation, but the database remains accessible. This may indicate the presence of a serious problem that could result in a future outage.

User response: You can begin an investigation of the hardware (such as disk, controllers, network) and/or the file system involved in accessing the container.

If it is suspected that the DB2 data itself is in error, contact IBM Software Support and they will guide you through the proper corrective actions.

ADM6007C DB2 detected an error while processing page *page-number* from table space *tbspace-id* for object *object-id* of object type *object-type*.

Explanation: DB2 was not able to complete the operation, but the database remains accessible. This may indicate the presence of a serious problem that could result in a future outage.

User response: Contact IBM Software Support and they will guide you through the proper corrective actions.

ADM6008I Extents within table space *tablespace_name* (ID *tablespace_id*) have been moved. Reason code = *reason-code*.

Explanation: Here are the reasons why the extent movement terminated:

1. Another utility attempted to work with the table space and interrupted the extent movement.
2. A pending delete state is preventing the movement of the high water mark.
3. There is no remaining free space in the table space to move the remaining extents.
4. The extent movement has completed.

User response: You should consider the following actions based on the reason code:

1. Re-run the extent movement when another utility cannot interrupt the move operation.
2. Correct the pending delete state before re-running the extent movement operation.

3. No action is required.
4. No action is required.

ADM6009W The database manager detected that the amount of restart light memory (rstrt_light_mem) is too small for optimal recovery performance.

Explanation: The database manager automatically reserves some memory on each member host to be used for restart light operation. This reserved memory is used to accommodate failed members that need to be restarted in restart light mode, on a host other than their home host. The rstrt_light_mem database manager configuration parameter specifies the maximum amount of memory that is allocated and reserved on a host for restart light recovery purposes.

This message is returned when the amount of memory reserved for restart light purposes on a host is not large enough for optimal recovery performance.

User response: Optional: To improve recovery performance, increase the value of the following database manager configuration parameters:

- RSTRT_LIGHT_MEM
- INSTANCE_MEMORY

ADM6010I Rebalance for table space *tsname* (ID *tsid*) has been suspended.

Explanation: The rebalance operation for the table space has been manually suspended using the ALTER TABLESPACE REBALANCE SUSPEND statement.

User response: Manually resume the table space rebalance operation by running the following statement:

```
ALTER TABLESPACE tablespace-name REBALANCE RESUME
```

ADM6011E Write operations could not be suspended or resumed on this DB2 member because the database is in the process of shutting down or stopping, or the database is quiesced.

Explanation: You can suspend and resume write operations for a database using the SET WRITE command or the db2SetWriteForDB API. In a DB2 pureScale environment, when you execute the SET WRITE command or call the db2SetWriteForDB API on any member, write operations are suspended or resumed on all members in the DB2 cluster.

Write operations can be successfully suspend or resumed only when the database is either active or stopped. This message is returned when an attempt is made to suspend or resume write operations for a database that is neither active nor stopped. For example, this message can be returned when the SET

WRITE command is executed in the following kinds of scenarios:

- The STOP DATABASE command has been executed for the database but the database has not completely stopped.
- The QUIESCE DATABASE command has been executed.
- The database encountered an error and is entering crash recovery.

This message is returned in DB2 pureScale environments only. Specifically, this message is printed to the administration notification log for the member that detected the problem with the status of the database.

User response: To resolve the problem that is preventing the SET WRITE command from succeeding on the member that is reporting this error, wait until the database makes the transition into being either active or stopped.

After resolving the problem on this member, review the administration notification log on other members in the DB2 cluster.

ADM6012W The total combined maximum number of authorization IDs that can be specified in both of the DB2_HI_PRI_PREFETCH_AUTHID and DB2_LO_PRI_PREFETCH_AUTHID registry variables *maxAuthIds* has been surpassed. Ensure that no more than this specified number of authorization IDs is specified by removing authorization IDs from one or both of the registry variables. The changes made will not take effect until the instance is restarted.

ADM6013W Authorization ID *authId* was specified multiple times in either one or both of the DB2_HI_PRI_PREFETCH_AUTHID and DB2_LO_PR_PREFETCH_AUTHID registry variables. As a result, the highest priority was chosen for this authorization ID. If this behaviour is undesirable, it is necessary to set the registry variables accordingly ensuring that there are no duplicate authorization IDs among them. The changes made will not take effect until the instance is restarted.

ADM6014E Write operations could not be suspended or resumed on this DB2 member because one or more table spaces are not in NORMAL state.

Explanation: You can suspend and resume write

operations for a database using the SET WRITE command or the db2SetWriteForDB API. In a DB2 pureScale environment, when you execute the SET WRITE command or call the db2SetWriteForDB API on any member, write operations are suspended or resumed on all members in the DB2 cluster.

Write operations can be successfully suspend or resumed only when all table spaces in the database are in the NORMAL state.

This message is returned in DB2 pureScale environments only. Specifically, this message is printed to the administration notification log for the member that detected the problem with the state of the table spaces.

User response: To resolve the problem that is preventing the SET WRITE command from succeeding on the DB2 member that is reporting this error, perform the following steps:

1. Identify which table spaces are not in NORMAL state using the MON_GET_TABLESPACE table function. The TBSP_STATE column lists the table space state.
2. Change any table spaces in the database that are not in a NORMAL state into a NORMAL state by performing the required troubleshooting steps.

After resolving the problem on this member, review the administration notification log on other members in the DB2 cluster.

ADM6015E Write operations could not be suspended on this DB2 member because an internal error was encountered while database logging was being suspended.

Explanation: You can suspend and resume write operations for a database using the SET WRITE command or the db2SetWriteForDB API. In a DB2 pureScale environment, when you execute the SET WRITE command or call the db2SetWriteForDB API on any member, write operations are suspended or resumed on all members in the DB2 cluster.

This message is returned in DB2 pureScale environments only. Specifically, this message is printed to the administration notification log for the member that detected the internal problem processing the suspend operation.

User response: To resolve the problem that is preventing the SET WRITE command from succeeding on the DB2 member that is reporting this error, perform the following troubleshooting steps:

1. Collect additional diagnostic information from the DB2 diagnostic (db2diag) log files on this member.
2. Identify and resolve the root cause of the failure of the suspend operation using the additional diagnostic information from the db2diag log files.

After resolving the problem on this member, review the administration notification log on other members in the DB2 cluster.

ADM6016E Write operations could not be resumed on this DB2 member because an internal error was encountered while database logging was being resumed.

Explanation: You can suspend and resume write operations for a database using the SET WRITE command or the db2SetWriteForDB API. In a DB2 pureScale environment, when you execute the SET WRITE command or call the db2SetWriteForDB API on any member, write operations are suspended or resumed on all members in the DB2 cluster.

This message is returned in DB2 pureScale environments only. Specifically, this message is printed to the administration notification log for the member that detected the internal problem processing the resume operation.

User response: To resolve the problem that is preventing the SET WRITE command from succeeding on the DB2 member that is reporting this error, perform the following troubleshooting steps:

1. Collect additional diagnostic information from the DB2 diagnostic (db2diag) log files on this member.
2. Identify and resolve the root cause of the failure of the resume operation using the additional diagnostic information from the db2diag log files.

After resolving the problem on this member, review the administration notification log on other members in the DB2 cluster.

ADM6017E The following table space is full. Table space name: *table-space-name*. Table space identifier: *table-space-ID*. Container path: *container-path*. Container identifier: *container-ID*.

Explanation: This message can be returned for many different reasons, including the following reasons:

- The underlying file system is full.
- The maximum allowed space usage for the file system has been reached.
- User limits on maximum file size have been reached.
- User limits on maximum number of open files have been reached.

User response:

1. Review information in diagnostic logs to determine which limits have been reached.
2. Depending on the limits that were reached, free the necessary resources. For example, free memory, allocate more memory, or close open files.

ADM6018I Unmounting filesystem at *directory-path*.

ADM6019E All pages in buffer pool *bpname* (ID *bpid*) are in use. Refer to the documentation for SQLCODE -1218.

ADM6020I Leaving mounted filesystem used for *contPath*.

ADM6021I Multi-page file allocation is not currently enabled. Disabling it has no effect.

ADM6022W Asynchronous drop of a temporary table failed. The resources associated with this table will not be freed until next database startup. You are advised to restart the database for this reason. You are also advised to contact IBM support to determine the reason for the failure.

ADM6023I The table space *tablespaceName* (ID *tablespaceID*) is in state *tablespaceState*. The table space cannot be accessed. Refer to the documentation for SQLCODE -290.

User response:

ADM6024C The database cannot be restarted. Database name: *database-name*.

Explanation: The database cannot be restarted because write operations for the database are suspended or are being suspended.

User response: To restart the database if the database write operations are suspended, specify the RESTART DATABASE command with the WRITE RESUME parameter.

To restart the database if the database write operations are being suspended, wait until the SET WRITE SUSPEND operation that is being processed completes, then reissue the RESTART DATABASE command with the WRITE RESUME parameter.

ADM6025I The table space *tsname* (ID *tsid*) is in state *state*. operation is not possible. Refer to the documentation for SQLCODE -290.

User response:

ADM6026E Write operations could not be suspended or resumed on this DB2 member because an internal error occurred during the suspend or resume process.

Explanation: You can suspend and resume write operations for a database using the SET WRITE command or the db2SetWriteForDB API. In a DB2 pureScale environment, when you execute the SET WRITE command or call the db2SetWriteForDB API on any member, write operations are suspended or resumed on all members in the DB2 cluster.

This message is returned in DB2 pureScale environments only. Specifically, this message is printed to the administration notification log for the member that detected the internal problem processing the suspend or resume operation.

User response: To resolve the problem that is preventing the SET WRITE command from succeeding on the DB2 member that is reporting this error, perform the following troubleshooting steps:

1. Collect additional diagnostic information from the DB2 diagnostic (db2diag) log files on this member.
2. Identify and resolve the root cause of the failure of the suspend or resume operation using the additional diagnostic information from the db2diag log files.

After resolving the problem on this member, review the administration notification log on other members in the DB2 cluster.

ADM6027E Error writing page *page-number* from table space *tblspace-id* for object *object-id* (writing to offset *offset* of container *container-path*). The operation did not complete.

Explanation: DB2 was not able to complete the operation, but the database remains accessible. This might indicate a serious problem that could result in a future outage.

User response: Investigate the hardware such as disk, controllers, and network, and the file system involved in accessing the container. If the problem is being caused by errors in the DB2 data, contact DB2 Support for guidance through the proper corrective actions.

ADM6028W The registry variable setting *regvarstr* was ignored because of an out of memory condition that did not permit the processing of the registry variable setting.

ADM6029W The registry variable setting (*regvarstr*) was ignored because the specified setting is invalid.

ADM6031W EXTENDED STORAGE was configured for use with buffer pool *bpname* (ID *bpid*) but this configuration will be ignored because AWE is enabled for the database. EXTENDED STORAGE should be disabled and buffer pools should be configured not to use it.

ADM6034W EXTENDED STORAGE was configured for use with buffer pool *bpname* (ID *bpid*) but this configuration will be ignored because the DB2_OVERRIDE_BPF registry variable is set. This registry variable is only meant to be used under the direction of IBM support.

ADM6035W Scattered read could not be used because the registry variable DB2NTNOCACHE was not set. Set the DB2NTNOCACHE registry variable to enable scattered read.

ADM6036W EXTENDED STORAGE is being used for multiple page sizes. Performance may not be optimal. Refer to the DB2 documentation for more information regarding the use of EXTENDED STORAGE.

ADM6037W Container *path* was created to be *userBytes* KB in size on a device that is *userBytes* KB in size. Extra storage will be wasted. The container can be extended to use the wasted space by using ALTER TABLESPACE.

ADM6038E Unable to perform an operation on container *path* because it does not exist in table space *tsname* (ID *tsid*). Refer to the documentation for SQLCODE -298.

ADM6039E Invalid stripe set *sset* specified in ALTER TABLESPACE for table space *tsname* (ID *tsid*). Maximum stripe set in table space is *maxsset*.

ADM6040E Invalid container size specified for container *contpath*. The number of pages specified for use is too large. Refer to the documentation for SQLCODE -1422.

ADM6041E A table space operation failed because the resulting size of the table space would have exceeded the defined maximum size for the table space. Table space name: *table-space-name*. Table space identifier: *table-space-id*. Defined maximum size: *max-size*.

Explanation: Rows of table data are organized into blocks called pages. For database-managed spaces (DMS), temporary DMS and nontemporary automatic storage table spaces, the page size you choose for your database using the PAGESIZE database configuration parameter determines the upper limit for the table space size.

This message is returned when a table space operation on a REGULAR or USER TEMPORARY DMS table space causes the size of the table space to increase beyond the defined maximum allowed size, as determined by the PAGESIZE database configuration parameter.

User response: Respond to this error in one of the following ways:

- Increase the maximum allowable size of the table space using the ALTER TABLESPACE statement with the MAXSIZE clause or the EXTEND clause, and then execute the table space operation again.
 - Modify the table space operation so that it does not increase the size of the table space beyond the maximum allowed size, as determined by the PAGESIZE database configuration parameter, and then execute the table space operation again.
-

ADM6042E An operation for table space *tsname* (ID *tsid*) was not successful because the size of the current table space is too big. The size of a REGULAR table space is limited to 0x0100 0000 (16777216) pages while the size of LARGE and TEMPORARY table spaces are limited to 0x7FFF FEFF (2147483391) pages. Refer to the documentation for SQLCODE -1139.

ADM6043W The registry variable DB2_NO_MPFA_FOR_NEW_DB is set to an invalid value (*regVarVal*). As a result, the newly created database will have multi-page file allocation enabled. If this is not the desired result, set the DB2_NO_MPFA_FOR_NEW_DB registry variable to YES and recreate the database. To disable multi-page file allocation for all newly created databases, set the DB2_NO_MPFA_FOR_NEW_DB registry variable to YES.

ADM6044E The DMS table space *tsname* (ID *tsid*) is full. If this is an autoresize or automatic storage DMS tablespace, the maximum table space size may have been reached or the existing containers or storage paths cannot grow any more. Additional space can be added to the table space by either adding new containers or extending existing ones using the ALTER TABLESPACE SQL statement. If this is an autoresize or automatic storage DMS table space, additional space can be added by adding containers to an autoresize table space or by adding new storage paths to the storage group it is using.

User response:

ADM6045I The database is no longer in the WRITE SUSPEND state. Database name: *database-name*.

Explanation: The WRITE SUSPEND state has been removed because write operations were resumed for the database by either the db2DatabaseRestart API or the RESTART DATABASE command with the WRITE RESUME parameter.

User response: No user action is required.

ADM6046I Multi-page file allocation already enabled. Enabling it again has no effect.

ADM6047W The table space *tsname* (ID *tsid*) is in the DROP_PENDING state. The table space will be kept OFFLINE. The table space state is *state*. This table space is unusable and should be dropped.

User response:

ADM6048I The table space *tsname* (ID *tsid*), which was previously OFFLINE has been brought back ONLINE. The table space state is *state*.

User response:

ADM6049E The database cannot be restarted because one or more table spaces cannot be brought online. To restart the database specify the "DROP PENDING TABLESPACES" option on the RESTART DATABASE command. Putting a table space into the drop pending state means that no further access to the table space will be allowed. The contents of the table space will be inaccessible throughout the remainder of the life of the table space; and the only operation that will be allowed on that table space is "DROP TABLE SPACE". There is no way in which it can be brought back. It is important that you consider the consequences of this action as data can be lost as a result. Before proceeding consult the DB2 documentation and contact IBM support if necessary. The table spaces to specify in the DROP PENDING TABLESPACES list are: *tsnames*.

ADM6050W The ALTER BUFFERPOOL statement for buffer pool *bpname* (ID *bpid*) was successful but could not be performed immediately because of insufficient memory. The change will take effect on the next database startup. Refer to the documentation for SQLCODE 20189.

ADM6051E The REORG command failed during RESTART DATABASE (crash recovery).

Explanation: An index or index partition cannot be recreated because the associated table or data partition is in a table space that is in the DROP_PENDING state.

User response: Set the INDEXREC database configuration parameter to "ACCESS" to delay index recreation until RESTART DATABASE completes.

ADM6052E Invalid container size specified for container *contpath* in the CREATE TABLESPACE or ALTER TABLESPACE statement. The number of pages specified for use is too small.

Explanation: You can specify the size of database managed space (DMS) table space containers when you create the table space using the CREATE TABLESPACE command. You can also change the size of DMS table space containers using the ALTER TABLESPACE command.

This message is returned when the specified number of pages is too small, when compared to the value of EXTENTSIZE that was specified when the table space was created.

For more information about the size of DMS table spaces, refer to the following topics in the DB2 Information Center: "CREATE TABLESPACE statement" and "Resizing DMS containers."

User response: If you are executing the CREATE TABLESPACE statement, execute the statement again, specifying a value for the container size that is at least twice as large as the value specified for EXTENTSIZE.

If you are executing the ALTER TABLESPACE statement, execute the statement again, specifying a value for the container size that is at least twice as large as the value for EXTENTSIZE that was specified when the table was created.

ADM6053W The CREATE BUFFERPOOL statement for buffer pool *bpname* (ID *bpid*) could not be performed immediately because not enough free memory existed in the database shared memory. The bufferpool will be created on the next database restart. Refer to the documentation for SQLCODE 20189.

ADM6054I DB2DART encountered table space *tsname* (ID *tsid*) which is currently not accessible. Since DB2DART is a diagnostic utility it will continue working on this table space.

ADM6055I Page *pagenum* was requested from table space *tsname* (ID *tsid*) but is not yet available because a rebalance which is in progress on this table space has not yet made that space useable. The requestor will wait for the page to become available.

ADM6056I Prefetch queue full has been encountered. A prefetch request has not been queued as a result and performance will not be optimal. Reconsider the prefetcher configuration (NUM_IOSERVERS, PREFETCHSIZE) and number of containers in each table space.

ADM6057I Prefetch queue full has been encountered. Agent will wait for space to free up on the queue and thus performance will not be optimal. Reconsider the prefetcher configuration (NUM_IOSERVERS, PREFETCHSIZE) and number of containers in each table space.

ADM6058I Rebalancer for table space *tsname* (ID *tsid*) was started.

ADM6059I Rebalancer for table space *tsname* (ID *tsid*) was restarted.

ADM6060I The rebalancer utility started successfully for table space *tsname* (ID *tsid*) but due to an out of memory condition the ability to throttle the execution of this utility has been disabled. Restarting the instance may resolve the memory issue and will allow throttling to be enabled.

ADM6061I Rebalance for table space *tsname* (ID *tsid*) stopping. The last extent moved by the rebalance was *lastext*.

ADM6062I Rebalance for table space *tsname* (ID *tsid*) has been completed.

ADM6063I Rebalance for table space *tsname* (ID *tsid*) has been paused.

ADM6064I Rollforward is waiting on rebalance of table space *tsname* (ID *tsid*) to complete.

ADM6065I An attempt was made to remove the OFFLINE state from table space *tsname* (ID *tsid*) but it was not offline. This command will be ignored. The table space state is *tsstate*.

User response:

ADM6066I An attempt was made to remove the OFFLINE state from table space *tsname* (ID *tsid*) but the action was not successful and so the table space will remain in this state. The table space state is *tsstate*. Refer to the documentation for SQLCODE -293.

User response:

ADM6069W Buffer pool *bpname* (ID *bpid*) has a NUMBLOCKPAGES value of *oldNumBlockPages* which exceeds the maximum supported value for a buffer pool of this size. The NUMBLOCKPAGES value for this buffer pool has been reduced to *newNumBlockPages* for its current activation. The value of NUMBLOCKPAGES can be permanently changed by using the ALTER BUFFERPOOL SQL statement.

ADM6070W Buffer pool *bpname* (ID *bpid*) has a NUMBLOCKPAGES value of *numBlockPages*. This value is smaller than the smallest allowable value for NUMBLOCKPAGES which is the BLOCKSIZE of *blockSize*. As a result the buffer pool will not be made block based for its current activation. The value for NUMBLOCKPAGES can be permanently changed by using ALTER BUFFERPOOL SQL statement.

ADM6071I The maximum number of pinned pages allowed concurrently was reached in buffer pool *bpname* (ID *bpid*). As a result performance may not be optimal. Increasing the SORTHEAP database configuration parameter may reduce the chances of this condition occurring in the future.

ADM6072W Disk full was encountered when writing out a temporary page from buffer pool *bpname* (ID *bpid*). The buffer pool cannot be decreased in size until all necessary pages have been written to disk. Refer to previous messages in the administration notification log for details regarding the disk full condition. DB2 will continue to retry writing this page to disk until the disk full condition is resolved.

ADM6073W The table space *tsname* (ID *tsid*) is configured to use buffer pool ID *ondiskBP*, but this buffer pool is not active at this time. In the interim the table space will use buffer pool ID *RTBP*. The inactive buffer pool should become available at next database startup provided that the required memory is available.

ADM6074I The SET WRITE SUSPEND command had no effect because the specified database is already in the WRITE SUSPENDED state. Database name: *database-name*.

Explanation: An attempt was made to suspend write operations on the specified database, but this database is already in the WRITE SUSPENDED state.

User response: No user action is required.

ADM6075W The database has been placed in the WRITE SUSPENDED state. Database name: *database-name*.

Explanation: You can suspend I/O write operations for a database by calling the SET WRITE SUSPEND command or by using the db2SetWriteForDB API. This message is returned when a database has been placed in WRITE SUSPENDED state.

User response: No user action is required.

Optional: To move the database out of WRITE SUSPENDED state, issue a SET WRITE RESUME command or a RESTART DATABASE command with the WRITE RESUME parameter.

ADM6076W The database is no longer in the WRITE SUSPEND state. Database name: *database-name*.

Explanation: The database, which was previously in the WRITE SUSPEND state, is no longer in that state. Write operations have resumed for the database.

User response: No user action is required.

ADM6077I The database is not in the WRITE SUSPEND state. Database name: *database-name*.

Explanation: An attempt was made to remove the database from the WRITE SUSPEND state, but the database is not in the WRITE SUSPEND state. Specifying WRITE RESUME has no effect.

User response: If the command that failed is the SET WRITE RESUME command, no user action is required.

If the command that failed is the RESTART DATABASE command with the WRITE RESUME parameter, issue the RESTART DATABASE command without the WRITE RESUME parameter. If this fails and the database is not in the WRITE SUSPEND state:

1. Wait until all SET WRITE SUSPEND operations are completed.
 2. Reissue the RESTART DATABASE command with the WRITE RESUME parameter.
-

ADM6078W The following table spaces were specified on the RESTART DATABASE command to be placed in the DROP_PENDING state: *tsnames*.

ADM6079E The table space *tsname* (ID *tsid*), was in the ROLLFORWARD_IN_PROGRESS state but can no longer be accessed. The table space has been taken offline and put into the RESTORE_PENDING state.

ADM6080E The table space *tsname* (ID *tsid*), was put OFFLINE and in ROLLFORWARD_PENDING. Tablespace state is *tsstate*.

User response:

ADM6081W The table space *tsname* (ID *tsid*) is in the OFFLINE state and is not accessible. The table space state is *tsstate*. Refer to the documentation for SQLCODE -293.

User response:

ADM6082W The current transaction is attempting to do work that is not allowed to fail. However, this work could not be completed as there are no free pages available in the buffer pool. Further attempts will be made to find free pages but in the future this situation can be avoided by increasing the size of buffer pool *bpname* (ID *bpid*).

ADM6083E An error occurred while redoing an alter tablespace operation against table space *tsname* (ID *tsid*) This error will be temporarily ignored while the remainder of the transaction is replayed. If the alter operation is eventually rolled back then the error will be discarded. However, if the operation is committed then this error will be returned, stopping recovery against the table space.

ADM6084E An attempt is being made to commit an alter operation against table space *tsname* (ID *tsid*) but a previous error is preventing this from being done. Resolve the original error before attempting the recovery again.

ADM6085I An out of memory condition was encountered while resizing buffer pool hash buckets. As a result of this condition, performance may not be optimal. It is recommended that database be shut down and started again so that the buffer pool can start up with an optimally sized hash table.

ADM6086W The attempt to acquire and reset the phantom quiesce state for table space *tsname* (ID *tsid*) did not change the state even though the command returned successfully. The authorization ID of the current user does not match the quiescer authorization ID and these must be the same to successfully acquire and reset a phantom quiesce. Determine the correct quiescer authorization ID by taking a table space snapshot and retry the quiesce reset command using that ID.

ADM6087I An attempt has been made to reset the quiesce state of table space *tsname* (ID *tsid*) but the table space is not currently in the quiesce state.

ADM6088W Due to an out of memory condition, the current agent is unable to monitor table space *tsname* (ID *tsid*). As a result, monitor output may be unreliable. Increase the DBHEAP configuration parameter to avoid this problem in the future.

ADM6089I The rebalancer utility started successfully for table space *tsname* (ID *tsid*), but due to an out of memory condition the progress monitor service is disabled. Restarting the instance may resolve the memory issue and will enable the progress monitor service again.

ADM6090W A file handle limit, system wide or process/thread specific, was reached by a page cleaner. The page cleaner will respond by closing all file handles opened by the cleaner that received the error condition. This error will only be logged once per page cleaner that receives this error. Recommended actions are: check operating system file handle limits, then reduce the value of the maxfilop configuration parameter. You must restart the database (database deactivation or last connection termination), for any changes to these

parameters to take effect.

User response:

ADM6091W An attempt was made to automatically increase the size of auto-resize table space *tsname* (ID *tsid*) but the table space's maximum size (*maxsize* bytes) has been reached. Note that this value may be lower than expected due to uncommitted ALTER TABLESPACE statements. Also, because DB2 attempts to extend containers at the same rate and extending must occur in a multiple of extents, it may not actually be possible to reach the maximum size exactly. In this case, the current size (*currentsize* bytes) will have a smaller value than the maximum size. The MAXSIZE clause of the ALTER TABLESPACE statement can be used to increase the maximum size for this table space.

ADM6092W An attempt was made to automatically increase the size of automatic storage table space *tsname* (ID *tsid*). This failed because all of the storage paths associated with the storage group are full or they do not have enough space to create a new container on them. New storage paths can be added to the storage group using the ALTER STOGROUP statement.

User response:

ADM6093W An attempt was made to automatically extend container *container* in auto-resize table space *tsname* (ID *tsid*) but there is no space on the filesystem. As a result, the table space cannot increase in size any further. One way to resolve this is to make more space available on the filesystem. Alternately, a new stripe set can be added to the table space using the BEGIN NEW STRIPE SET clause of the ALTER TABLESPACE statement. When subsequent attempts to auto-resize the table space occur, only those newly added containers will be extended and the existing ones will remain as-is.

ADM6094W A RESTORE DATABASE or ROLLFORWARD DATABASE command resulted in storage changes for auto-resize table space *tsname* (ID *tsid*). The effect is that the table space's maximum size had to be increased from *oldmaxsize* bytes to *newmaxsize* bytes on this partition. If there are multiple database partitions then the maximum size for the table space is now inconsistent between the database partitions. This will not cause any problems but it may not be desired. To resolve this use the ALTER TABLESPACE statement to set a new MAXSIZE value that is greater than or equal to the current size for each of the database partitions. Alternately, use the MAXSIZE NONE option to indicate that there is no maximum size.

ADM6095W A rollforward operation encountered a log record associated with the ADD STORAGE clause of the ALTER DATABASE statement or the ADD clause of the ALTER STOGROUP statement. However, this log record was not replayed because the storage paths associated with the database were redefined during the previous database restore and it is assumed that the full set of storage paths were defined at that point. As a result, storage path *storagepath* was not added to the database.

User response:

ADM6096W The rollforward operation encountered a log record that is attempting to alter the maximum size for auto-resize table space *tsname* (ID *tsid*). While this was successful at run-time, a new container configuration has been established during the process of restoring and rolling forward the database or table space and this maximum size is less than the current size of the table space. The effect is that the table space's maximum size is being set to *currentsize* bytes instead of *maxsize* bytes. If there are multiple database partitions then the maximum size for the table space is now inconsistent between them. This will not cause any problems but it may not be desired. This can be resolved after the rollforward has completed by using the ALTER TABLESPACE statement to set a new MAXSIZE value that is greater than or equal to the current size for each of the database partitions. Alternately, use the MAXSIZE NONE option to indicate that there is no maximum size for the table space.

ADM6097I Table space *tsname* (ID *tsid*) was successfully extended by *extend* bytes in a new stripe set.

ADM6098W Table space *tsname* (ID *tsid*) could not be extended because the container map became too complex.

ADM6099W Table space *tsname* (ID *tsid*) has reached the maximum size (*maxsize* bytes) for a tablespace of this type and page size. To add more storage to the database, add a new table space.

ADM6100W While attempting to autoresize table space *tsname* (ID *tsid*), container *cname* could not be grown.

ADM6101W While attempting to extend table space *tsname* (ID *tsid*), there were less than *freebytes* bytes free on *path*. This space is reserved for use by DB2 and/or the operating system.

ADM6102I Table space *tsname* (ID *tsid*) was successfully extended by *extend* bytes.

ADM6103W Table *table-name* cannot allocate a new page because the index with identifier *index-id* does not yet support large RIDs. The table space in which this table resides was converted to a large table space via the CONVERT TO LARGE clause of the ALTER TABLESPACE statement. The table itself cannot support large RIDs until all previously existing indexes on the table have been reorganized or rebuilt to support large RIDs. The index(es) must be reorganized or rebuilt to support future growth of this table.

Explanation: The indexes can be reorganized using the rebuild mode of the REORG INDEXES ALL FOR TABLE *table-name* command. For partitioned tables ALLOW NO ACCESS must be specified. Alternatively, the table can be reorganized (classic REORG, not INPLACE) which will not only rebuild all indexes, but will enable the table to support greater than 255 rows per page.

User response:

ADM6104W Table space *tblspace-name* is being converted from REGULAR to LARGE. Indexes on tables in this table space must be reorganized or rebuilt to support large RIDs. The table space being converted can, following COMMIT, support a storage capacity larger than that of a regular table space. The maximum page number for a data page in a regular table space is 0x00FFFFFF. A table which is allocated a page number above 0x00FFFFFF must have support from the indexes on the table to support such a page number. Until the indexes on a table are reorganized or rebuilt to support such a page number, allocation of such a page number to the table will result in an error.

Explanation: The documentation for the ALTER TABLESPACE statement in the SQL Reference specifies the best practices when using the CONVERT TO LARGE option. Please follow these recommendations to be pro-active in reorganizing or rebuilding all indexes on all tables in this table space to prevent possible future failures to grow the tables.

ADM6105E The storage paths associated with this storage group are inconsistent between database partition *X* and database partition *Y*. The database manager attempts to keep the storage paths consistent across all database partitions but there is now a discrepancy. While the database continues to function, it is recommended that the storage paths be made consistent across all of the partitions in the database. This is accomplished by backing up the database on each partition (unless backup images already exist) and restoring them. The catalog partition should be restored first, specifying the list of paths that each partition should use (using the RESTORE REDIRECT and SET STOGROUP PATHS commands). If the database is recoverable then a rollforward needs to be done after the database is restored on every partition.

User response:

ADM6106E Table space *name* (ID = *id*) could not be created during the rollforward operation. The most likely cause is that there is not enough space to create the containers associated with the table space. Connect to the database after the rollforward operation completes and use the SET TABLESPACE CONTAINERS command to assign containers to the table space. Then, issue another ROLLFORWARD DATABASE command to complete recovery of this table space.

ADM6107E Automatic storage table space *name* (ID = *id*) could not be created during the rollforward operation. The most likely cause is that there is not enough space on the storage group's storage paths to create the *size* byte table space. If this is the case then make more space available on the existing storage paths or add new storage paths using the ALTER STOGROUP SQL statement before issuing another ROLLFORWARD DATABASE command to recover this table space.

User response:

ADM6108I A request was made to drop storage path *storage-path* from the storage group. Multiple instances of this path exist on this database partition and all instances of it are being dropped.

Explanation: When a storage group on a database partition contains multiple instances of the specified storage path, a request to drop this storage path, using the ALTER STOGROUP statement, results in the dropping of all of the instances of this storage path in that storage group.

User response: No response is necessary.

ADM6109I A request was made to drop storage path *storage-path* from the storage group. The storage path will be dropped immediately.

Explanation: No table space uses the storage path on this database partition. The storage path is being dropped immediately.

User response: No response is necessary.

ADM6110I A request was made to drop storage path *storage-path* from the storage group. The storage path will not be removed from this database partition until all containers residing on it have been dropped.

Explanation: One or more automatic storage table spaces have containers on this storage path on this database partition. Therefore, the storage path is being placed into the "drop pending" state.

User response: One of the following actions:

- Drop and recreate temporary automatic storage table spaces.
- Use the REBALANCE clause of the ALTER TABLESPACE statement to move data from the storage paths that are in the "drop pending" state.
- Drop table spaces that you do not need.

ADM6111I No table space containers reside on storage path *storage-path* on this database partition. The storage path is in the "drop pending" state and will be removed from this database partition.

Explanation: Because the storage path does not contain any data, it will be dropped immediately.

User response: No response is necessary.

ADM6112I Storage path *storage-path* was added to the storage group. The new storage path will not be used automatically.

Explanation: Temporary, regular, or large automatic storage table spaces cannot automatically use new storage paths.

User response: If you want temporary table spaces to use the new storage path, the database must be stopped and started again:

- If the database was explicitly activated, first deactivate the database by using the DEACTIVATE DATABASE command, then reactivate it by using the ACTIVATE DATABASE command.
- If the database was not explicitly activated, disconnect or force all users from the database and then connect to the database again.

When the database is started, the temporary table spaces will use the new storage path.

Alternatively, you can drop the temporary table space by using the DROP TABLESPACE statement, then recreate it by using the CREATE TABLESPACE statement. When it is recreated, the temporary table space will use the new storage path.

Regular and large table spaces do not use a new storage path until a "disk full" condition occurs for an existing container. If you want to stripe the table space over all the storage paths (including new paths), use the ALTER TABLESPACE statement with the REBALANCE clause. This statement creates containers for each of the table space stripe sets, one for each path that does not already contain a container.

You can determine the list of automatic storage table spaces in the database by issuing the following SQL statement:

```
SELECT TBSP_NAME, TBSP_CONTENT_TYPE
FROM SYSIBMADM.SNAPTbsp
WHERE TBSP_USING_AUTO_STORAGE = 1
ORDER BY TBSP_ID
```

ADM6113I Table space *tablespace* exceeded its maximum size.

Explanation: A rebalance operation added data from one or more dropped storage paths to this table space. The rebalance operation was successful; however, the table space has exceeded its maximum size.

User response: No user response is required.

Chapter 13. ADM6500 - ADM6999

ADM6500W The connection for the DATALINK file server *fileServer* has failed. The server is blocked.

ADM6501W The connection for the DATALINK file server *fileServer* has failed to restart.

ADM6502W The connection for the DATALINK file server *fileServer* has failed.

ADM6510W The DATALINK file server *fileServer* has not been registered.

ADM6511W The DATALINK file server *fileServer* is not active.

ADM6512W DB2 is unable to read from the DATALINK configuration file.

ADM6513W DB2 is unable to create the DATALINK configuration file *fileName*.

ADM6514W An error has occurred while DB2 was writing to the DATALINK configuration file.

ADM6515W An error has occurred while DB2 was closing the DATALINK configuration file.

ADM6516W DB2 is unable to open the DATALINK configuration file *fileName*.

Chapter 14. ADM7000 - ADM7499

| | | | |
|----------|---|----------|--|
| ADM7000W | An invalid SEARCH discovery protocol, <i>fileName</i> , was specified in the DISCOVER_COMM registry variable. | ADM7011E | The Sync Point Manager does not support the version of Microsoft SNA Server installed on this machine. Minimum requirement is Microsoft SNA Server V4 Service Pack 3. |
| ADM7001E | The DB2 function sqlinstancepath has failed. The instance path was not set. If you are on UNIX, check the DB2INSTANCE registry variable. If you are on Windows, check the DB2INSTANCE, DB2PATH, and DB2INSTPROF registry variables. | ADM7012E | The SSL_SVR_KEYDB DBM configuration parameter was not configured. Update the SSL_SVR_KEYDB configuration parameter. |
| ADM7004E | An invalid value (<i>value</i>) was specified for the DB2COMM registry variable. | ADM7013E | The SSL_SVR_STASH DBM configuration parameter was not configured. Update the SSL_SVR_STASH configuration parameter. |
| ADM7005W | The DISCOVER mode was set to SEARCH. However, the DISCOVER_COMM registry variable was not configured with any protocols. | ADM7014E | The SSL_SVCENAME DBM configuration parameter was not configured. Update the SSL_SVCENAME configuration parameter using the service name defined in the TCP/IP services file. |
| ADM7006E | The SVCENAME DBM configuration parameter was not configured. Update the SVCENAME configuration parameter using the service name defined in the TCP/IP services file. | ADM7015I | The connection is using SSL version <i>version</i> and cipher spec <i>cipherspec</i> . |
| ADM7007E | The SVCENAME DBM configuration parameter, <i>socketAddress</i> , is configured with a port or a service name. When it is configured with a service name, the TCP/IP services files is used to map the service name to a port number. The port specified in this field is being used by another process. Resolve this problem by either deleting the process using the port or use another port. | ADM7016I | The allowed SSLv3 cipher specs are <i>cipherspecs</i> . The allowed TLSv1 cipher specs are <i>cipherspecs</i> . |
| ADM7008W | The DB2TCPCONNMGERS registry variable <i>value</i> , <i>userValue</i> , is invalid. A valid value is from 1 to 8. A default DB2TCPCONNMGERS value of <i>defaultValue</i> has been used. | ADM7017E | TCP/IP and SSL cannot be configured on the same port number. |
| ADM7009E | An error was encountered in the <i>protocolTCPIP</i> protocol support. A possible cause is that the maximum number of agents has been exceeded. | | |

Chapter 15. ADM7500 - ADM7999

ADM7500W A request was made to quiesce an instance with the following quiesce mode: *quiesce-mode*

Explanation: You can force off all users of a database manager instance and put the instance and all databases in the instance into quiesce mode using the QUIESCE INSTANCE or the START DATABASE MANAGER commands (or the db2InstanceQuiesce or db2InstanceStart APIs). This message is returned to acknowledge and log the quiesce mode of the request to quiesce an instance.

User response: No user response is required.

ADM7501W Instance quiesce request has completed successfully.

ADM7502E Instance quiesce request has failed. The SQLCODE is *SQLCODE*.

ADM7503W Instance unquiesce request has completed successfully.

ADM7504W Instance unquiesce has been requested.

ADM7505E Instance unquiesce request has failed. The SQLCODE is *SQLCODE*.

ADM7506W A request was made to quiesce a database with the following quiesce mode: *quiesce-mode*

Explanation: The database was requested to be put into quiesce mode using the QUIESCE DATABASE command or the db2DatabaseQuiesce API. This message is returned to acknowledge and log the quiesce mode of the request to quiesce a database.

User response: No user response is required.

ADM7507W Database quiesce request has completed successfully.

ADM7508E Database quiesce request has failed. The SQLCODE is *SQLCODE*.

ADM7509W Database unquiesce request has completed successfully.

ADM7510W Database unquiesce has been requested.

ADM7511E Database unquiesce request has failed. The SQLCODE is *SQLCODE*.

ADM7512E A transaction resynchronization error has occurred, contact IBM Support for assistance.

ADM7513W Database manager has started.

ADM7514W Database manager has stopped.

ADM7515W The concentrator has been enabled without directory caching support which may result in a performance penalty. To correct this problem, enable directory caching in the database manager configuration.

Explanation: Directory caching is a feature with relatively low overhead that is enabled by default in the engine. There are only a few valid reasons why a user may want to disable it. The concentrator configurations are particularly sensitive to having directory caching turned off because the dispatcher processes, which are the bottleneck for concentrator, could be potentially required to go to disk with every new connection.

User response: Enable directory caching, or run in a non-concentrator configuration.

ADM7516N During automatic restart, DB2 cluster services failed to recover the database named *database-name* on the following DB2 member: *member-ID*

Explanation: In a DB2 pureScale environment, if a database member fails, member crash recovery is automatically initiated for that member. In addition, if both cluster caching facilities fail at the same time, group crash recovery is automatically initiated.

During member crash recovery, no connections are allowed against the database through that member. Attempted connections to the named database through the named DB2 member will fail until the database has been successfully recovered.

This message is returned when DB2 cluster services fails to recover a database for a failed member during an automatic restart. An ALERT is set for the member on the host where the recovery failed.

User response:

1. Collect additional, related troubleshooting and diagnostic information by searching in the db2diag log files for this message identifier.
2. Determine whether the database is in a backup pending, rollforward pending, or restore pending state by running the GET DB CONFIG command
3. If the database is in a pending state, take the required actions to move the database out of the pending state.
4. Manually recover the failed database by running the RESTART DATABASE command. The alert will automatically be cleared when the database is successfully recovered.

ADM7517W When the concentrator is enabled, the maximum number of applications per application group in the database *databaseName* will be limited to *maxApp*. To correct this problem, decrease the value of the APPGROUP_MEM_SZ configuration parameter.

ADM7518C The database manager has shut down the following database because a severe error has occurred: *database-name*.

Explanation: This error is returned when there is a sudden, severe problem that might compromise the integrity of data during database access. Here is one example of the kind of scenario that could cause this error to be returned:

- If a database object unexpectedly becomes inaccessible during a transaction, all write operations to the database logs would stop and all operations on this database would be shut down. The database would need to undergo crash recovery when it is reactivated.

User response: Collect the following diagnostic information and then contact IBM software support using tools such as the db2support tool:

- Contents of the directory identified by the diagpath database manager configuration parameter including the contents of the FODC_ForceDBShutdown directory.

ADM7519W DB2 could not allocate an agent. The SQLCODE is *SQLCODE*.

ADM7520E Mount unmounted filesystem(s) has failed on node *db-partition-number*. Manual remount is required before retrying migration. Check *errorFile* and remount the filesystems before retrying migration.

ADM7521E Unmount mounted filesystem(s) has failed on node *db-partition-number*. Remount the unmounted filesystem(s) and try the database migration again.

ADM7522E Remount unmounted filesystem(s) to new mount points has failed on node *db-partition-number*. Restart database migration is needed.

ADM7523E DB2 was unable to allocate the application group memory set. This will cause database activation to fail. Reduce the APPGROUP_MEM_SZ DB configuration parameter and retry.

ADM7524W The database manager has increased the amount of memory that is allocated for restart light recovery purposes to the minimum required size that is needed for starting the DB2 idle process.

Explanation: The current calculated restart light memory size is too small to allocate enough memory to start the DB2 idle process. The restart light memory size is calculated using the instance_memory configuration parameter and the rstrt_light_mem configuration parameter, which specifies a percentage of instance_memory. The restart light memory size must be greater than or equal to the minimum required restart light memory size.

The DB2 idle process needs to be started for DB2START processing to succeed, so the restart light memory size has been increased to the minimum required size.

User response:

1. Check the db2diag.log to find the values for:
 - calculated restart light memory size
 - minimum required restart light memory size
2. Check the current values of the instance_memory and rstrt_light_mem configuration parameters and increase one (or both) of them so that the calculated restart light memory size is greater than or equal to the minimum required restart light memory size.

ADM7525E The following member was unable to access the secondary CF to initiate catch up: *member-id*. Member host name: *host-name*. Secondary CF host name: *CF-host-name*.

Explanation: After a secondary cluster caching facility, also known as CF, is added, and after a secondary CF host computer is restarted, the secondary CF must go through a process to join the cluster. Specifically, when the first database connection is made after a secondary CF joins a cluster, the secondary CF copies information from the primary CF. This copying of information is

known as "catch up", and while the secondary CF is copying data from the primary CF, the secondary CF is in CATCHUP state.

This message is returned when the member that is processing the first database connection after a secondary CF joins a cluster cannot access the secondary CF to initiate the catch up process.

Possible reasons that a member cannot initiate the catch up process for a secondary CF include the following reasons:

- Network errors, such as problems with uDAPL, prevent the member from accessing the secondary CF
- The host on which the secondary CF is located is unavailable
- The secondary CF server processes are not running

User response: Perform the following troubleshooting steps:

- Verify that the member *member-id* can access the host on which the secondary CF is located, *CF-host-name* using the ping utility, for example.
- Look for messages related to the secondary CF server processes in the following diagnostic files:
 - cfdump.out* files
 - cfdiag.log files
 - core files
- Investigate whether there are any uDAPL communication problems.

ADM7526E The database activation failed because the secondary CF failed to allocate required structures.

Explanation: The cluster caching facility, also known as CF, contains the following structures:

- Group Buffer Pool (GBP)
- CF Lock Manager (LOCK)
- Shared Communication Area (SCA)

The size of each of these structures, GBP, SCA, and LOCK can be configured online.

One reason that this failure can happen is that an invalid size was calculated for one or more of the structures of the secondary CF during database activation. When this error occurs, detailed diagnostic information about this error is printed in the DB2 diagnostic (db2diag) log files.

User response:

1. Collect detailed diagnostic information about this error from the db2diag log files.
2. Correct the cause of the problem.
3. Retry the database activation.

ADM7527E The database activation failed because the secondary CF server had insufficient memory to allocate one or more required structures.

Explanation: The cluster caching facility, also known as CF, contains the following structures:

- Group Buffer Pool (GBP)
- CF Lock Manager (LOCK)
- Shared Communication Area (SCA)

The size of each of these structures, GBP, SCA, and LOCK can be configured online.

One reason that this failure can happen is that non-DB2 applications on the secondary CF host used a large amount of memory while the secondary CF was allocating memory for structures.

User response: Respond to this error by performing one or more of the following steps:

- If there are non-DB2 applications running on the secondary CF server host in addition to the secondary CF itself, ensure there is sufficient memory to handle the secondary CF workload before activating the secondary CF.

For example, configure the amount of memory specified by the CF_MEM_SZ database manager parameter for the secondary CF to be as large as or larger than the value of CF_MEM_SZ on the primary CF.

- Configure the secondary CF to use less total memory by specifying a smaller size for the following database configuration parameters:
 - CF_SCA_SZ
 - CF_GBP_SZ
 - CF_LOCK_SZ
 - CF_DB_MEM_SZ
 - CF_MEM_SZ

ADM7528N The database manager started with fewer Host Channel Adapter Adapters (HCAs) than configured for the cluster caching facilities (CFs).

Explanation: The database manager encountered a non-critical error when starting. Communication could not be established with one or more HCAs on a CF, however each cluster caching facility (CF) is connected by at least one HCA. With fewer HCAs the CF has less throughput capacity. There is also a greater risk of downtime due to the reduced redundancy.

User response: Review cluster alerts by running 'db2cluster -cm -list -alert' to see which HCAs are not responding. Follow the corrective actions given by the alerts to rectify any reported problem.

ADM7529I Cluster caching facility (CF) has stopped.

Explanation: The CF has successfully stopped.

User response: No response is required.

ADM7530W An unexpected cluster caching facility (CF) error occurred. CF recovery will be attempted.

Explanation: The CF stopped because of an unexpected error. DB2 will perform a cleanup and restart the CF.

User response: To avoid receiving this message in the future, check the db2diag.log and the cfdiag.cfid.log files to see if there are any messages around the time of the recovery attempt.

ADM7531I A CF changed state. CF: *identifier*. New state: *state*.

Explanation: The indicated CF is now in the indicated state.

User response: No response is required.

ADM7532E Cluster caching facility (CF) with identifier *identifier* is now in ERROR state.

Explanation: The indicated CF could not be restarted.

User response: Resolve the problem and manually clear the alert. You can use the db2instance and db2cluster commands to help you with this. After the problem has been resolved, attempt to start the CF.

ADM7533I Cluster caching facility (CF) has started.

Explanation: The CF has successfully started.

User response: No response is required.

ADM7534W During database upgrade the rebind operation failed for one or more packages. Log file: *log-file*.

Explanation: The REBINDALL option was specified on the UPGRADE DATABASE command to rebind all the packages during database upgrade. The database was upgraded successfully, however, one or more packages failed to rebind. The errors and the package names are recorded in the generated log file.

User response:

1. Resolve the source of the problem by following the actions indicated by the messages that are recorded in the generated log file.
 2. Rebind the packages using the db2rbind command.
-

ADM7535W The UPGRADE DATABASE command failed to refresh the table space attributes in the catalog table. However the database was upgraded successfully.

Explanation: In DB2 V10 some new attributes are added in the catalog tables for each table space in the database. The database upgrade process failed to initialize these new attributes with appropriate values. See the file db2diag.log to determine the cause of failure. The database performance may not be optimal until the table space attributes are properly initialized.

User response: Resolve the errors that caused the initialization of the tablespace attributes to fail during database upgrade. Execute the DATABASE UPGRADE command again to initialize the table space attributes.

ADM7536I The status of adapter *adapter-name* was found to be: *ONLINE-or-OFFLINE-status*.

Explanation: This message is returned to show whether adapters are online or offline.

User response: No response is required.

ADM7537I The status of the following adapter changed. Adapter name: *adapter-name*. New status: *ONLINE-or-OFFLINE*. Number of adapters that are currently online: *number*. Host name: *host-name*.

Explanation: The cluster caching facility (CF) monitors adapters to determine whether those adapters are online or offline.

This message is returned when the CF detects that an adapter has changed from online to offline or from offline to online.

User response: No response is required.

ADM7538E The database upgrade operation completed successfully. However there are tables in the upgraded database that are in load pending state. These tables cannot be used by the upgraded database.

Explanation: This message is returned if the database upgrade operation found that a previous LOAD attempt on one or more tables in the database did not complete successfully. These tables cannot be accessed in the upgraded database.

User response:

1. To find out what tables are in a load pending state run the following command:

```
select <tablename> from SYSIBMADM.ADMINTABINFO
where load_status is not NULL
```
-

2. In the previous release environment, restart or terminate the failed LOAD operation by issuing the LOAD command with the RESTART or TERMINATE option.
3. Reload the table by performing the following steps:
 - a. Unload the table from the previous release database.
 - b. Drop and recreate the table in the upgraded database.
 - c. Reload the table in the upgraded database, using the data that was unloaded from the previous release database.

ADM7539I The database manager has completed processing the topology change that resulted in a backup pending state.

Explanation: The database manager has completed processing a topology change that resulted in the dropping of some members, or a configuration change between ESE and DB2 pureScale instances. In both cases, for those members that were in the previous member topology for the database but are not in the current member topology, this might have resulted in the deletion of files from those member database directories.

User response: No action required.

Chapter 16. ADM8000 - ADM8499

ADM8000C Backup has been terminated. The SQLCODE returned is *SQLCODE*.

ADM8001W Incremental backup was not enabled for this database because the TRACKMOD DB configuration parameter was not enabled.

ADM8002W This backup image cannot be used for ROLLFORWARD because the logs associated with this backup has been overwritten on the raw device. Use a more recent backup image.

ADM8003C Restore has been terminated. The SQLCODE returned is *SQLCODE*.

ADM8004W Incremental backup was not enabled for table space *tablespaceName* (ID *tablespaceID*) because the TRACKMOD configuration parameter was not enabled.

ADM8005W Incremental backup was not enabled for table space *tablespaceName* (ID *tablespaceID*). A non-incremental backup of this table space is required.

ADM8006W DB2 cannot use the specified restore buffer size of *restoreBufferSize* 4K pages. A restore buffer size must be multiple of backup buffer size of *backupBufferSize* 4K pages. The restore operation will continue with the default buffer size.

ADM8007W DB2 cannot perform multiple concurrent incremental restores.

ADM8008W DB2 could not find and/or delete the online reorganization state files for all table spaces during restore. Manual intervention may be needed to remove the file(s).

ADM8009W Could not find and/or delete the online reorganization state files for table space *tablespaceName* (ID *tablespaceID*) during restore. Manual intervention may be needed to remove the file(s).

ADM8010E Backup was unable to copy requested log file *logfileName* for inclusion in the backup image. The backup has been aborted.

ADM8011W The database backup succeeded. However, the DB2 database server was unable to create part of the incremental chain of images during the backup, and will be unable to use the affected images during restore. Specifically, you will not be able to use the backup image with timestamp *timestamp* for incremental restores involving table space *table-space-name*.

ADM8012W The database backup succeeded. However, the entry in the recovery history file corresponding to the backup image with timestamp *timestamp* will not be well-formed, because a write error occurred while updating the recovery history file itself. See the db2diag log file for more information.

Chapter 17. ADM8500 - ADM8999

ADM8500W DB2 has failed to read from the history file because of a possible data corruption. Ensure that the file exists and is intact.

ADM8501W DB2 has failed to write to the history file because the disk is full.

ADM8502W The history file is corrupt. An unrecoverable error has been detected with the file. The existing file has been deleted and a backup has been made. If you would like to determine the root cause of the problem, contact IBM Support. Otherwise, no further action is needed.

ADM8503W DB2 was unable to record a history entry for operation *operation*.

ADM8504I Successfully deleted the backup image with timestamp *timestamp*.

ADM8505I Successfully deleted the load copy image with timestamp *timestamp*.

ADM8506I Successfully deleted the following database logs *log-list* in log chain *log-chain*.

ADM8507N Unable to delete the backup image with timestamp *timestamp*.

Explanation: The DB2 database manager attempted to delete the given backup image, but was unsuccessful.

User response: Verify that the DB2 database manager has access to the storage manager or directory where the backup images are stored. See the db2diag log file for more information.

ADM8508N Unable to delete the load copy image with timestamp *timestamp*.

Explanation: The DB2 database manager attempted to delete the given load copy image, but was unsuccessful.

User response: Verify that the DB2 database manager has access to the storage manager or directory where the load copy images are stored. See the db2diag log file for more information.

ADM8509N Unable to delete the database logs *log-list* in log chain *log-chain*.

Explanation: The DB2 database manager attempted to delete the given database logs, but was unsuccessful.

User response: Verify that DB2 has access to the storage manager or directory where the log files are stored. See the db2diag log file for more information.

Chapter 18. ADM9000 - ADM9499

ADM9000W Prefetching was disabled during sort merge; performance may be suboptimal. If this message persists, consider increasing the buffer pool size for temporary table space *tablespaceName* (ID *tablespaceID*) or increase the value of the SORTHEAP DB configuration parameter to reduce the extent of sort spilling.

Chapter 19. ADM9500 - ADM9999

ADM9500W Too many concurrent updates had occurred on table *tableName* (ID *tableID*) and table space *tablespaceName* (ID *tablespaceID*) during online index create/reorganization. Thus, it will take a longer time to finish online index create/reorganization. You may want to increase your UTIL_HEAP_SZ DB configuration parameter.

ADM9501W Index reorganization has started for table *tableName* (ID *tableID*) and table space *tablespaceName* (ID *tablespaceID*).

Explanation: The data server is reorganizing the indexes for the specified table. The re-organization is either for nonpartitioned indexes on a partitioned table or for indexes on a nonpartitioned table.

User response: No response is required.

ADM9502W Index reorganization is complete for table *tableName* (ID *tableID*) and table space *tablespaceName* (ID *tablespaceID*).

Explanation: The data server reorganized the indexes for the specified table. The index reorganization was either for nonpartitioned indexes on a partitioned table or for indexes on a nonpartitioned table.

User response: No response is required.

ADM9503W Reorganizing index IID *indexIID* (OBJECTID *indexObjectID*) in table space *indexTablespaceName* (ID *indexTablespaceID*) for table *tableName* (ID *tableID*) in table space *tablespaceName* (ID *tablespaceID*).

Explanation: The data server is reorganizing the specified index. The reorganization is either for nonpartitioned indexes on a partitioned table or for indexes on a nonpartitioned table.

User response: No response is required.

ADM9504W Index reorganization on table *tableName* (ID *tableID*) and table space *tablespaceName* (ID *tablespaceID*) failed on this database partition with SQLCODE *SQLCODE* reason code *reasonCode*.

Explanation: Index reorganization failed on this database partition for the reason described by the SQLCODE. The indexes referred to in this context are either nonpartitioned indexes on a partitioned table or indexes on a nonpartitioned table

User response: Correct the problem that is described by the SQLCODE, and then retry the REORG INDEXES command on the database partition.

ADM9505W Online index reorganization on table *tableName* (ID *tableID*) and table space *tablespaceName* (ID *tablespaceID*) has been switched to offline mode because the indexes are marked for rebuild. These indexes may have been marked for rebuild during a roll forward through an index creation and/or recreation. If this is the case, consider setting the INDEXREC database manager configuration parameter to RESTART. This will cause indexes that are marked for rebuild during roll forward to be rebuilt during RESTART DATABASE processing.

ADM9506W HADR is enabled, but full logging is disabled for any index creation, recreation, or reorganization on table *table-name* (table object id: *object-id*) in tablespace *tablespace-name* (tablespace id: *tablespace-id*), since you have explicitly requested to disable it. As a result, any index build operations on this table will not be recovered immediately on the secondary database server using HADR. Indexes on the secondary database server will be recreated implicitly at the end of the HADR takeover process or after the HADR takeover process when the underlying tables are to be accessed. If this is not the desired behavior, enable the full logging on the table before any create, recreate, or reorg index is performed.

ADM9507W When HADR is enabled, it is recommended that the database configuration parameter `LOGINDEXBUILD` is set to ON, on both the HADR primary database server and the HADR secondary database server. Otherwise, you may not log your index creation, recreation, or reorganization on current or future HADR primary database server. Any non-fully logged index creation, recreation, or reorganization on the primary database server will not be recovered on the secondary database server using HADR. Those indexes which cannot be recovered will be marked as invalid, and will be rebuilt implicitly at the end of the HADR takeover process or after the HADR takeover process when the underlying tables are accessed. If this is not the desired behavior, enable the full logging or use the default setting for this configuration parameter before any index build operations take place.

ADM9508W When HADR is enabled, it is recommended that the database or database manager configuration parameter `INDEXREC` is set to either `RESTART` or `ACCESS` in order to enable redo of any index creation, recreation or reorganization. Otherwise, any fully logged index creation, recreation, or reorganization on the primary database server will not be recovered on the secondary database server using HADR. Those indexes which cannot be recovered will be marked as invalid and will be rebuilt implicitly at the end of the HADR takeover process or after the HADR takeover process when the underlying tables are accessed. If this is not the desired behavior, update `INDEXREC` or use the default setting for this configuration parameter before any index build operations take place.

ADM9509W It is recommended that the database configuration parameter `LOGINDEXBUILD` is set to ON before HADR is started. Otherwise, any index creation, recreate, or reorganization on the current or future primary database server may not be recovered on the current or future secondary database server using HADR. Those indexes which cannot be recovered will be marked as invalid and will be rebuilt implicitly either at the end of the HADR takeover process or after the HADR takeover process when the underlying tables are to be accessed. If this is not the desired behavior, update the database configuration parameter `LOGINDEXBUILD` to ON.

ADM9510W An error (sqlcode *sqlcode*) occurred which prevented the completion of the index rebuild process. Any invalid indexes that have not been rebuilt when the process terminated will be recreated on the first table access. The index rebuild process was invoked either during an explicit or implicit restarting of the database, or at the end of the HADR takeover.

ADM9511W Index reorganization proceeds on index *indexname* (IID *indexIID*, OBJECTID *indexObjectID*) in table space *indexTablespaceName* (ID *indexTablespaceID*) for table *tableName* (ID *tableID*) in table space ID *tablespaceID*.

ADM9512W Index reorganization for index *indexname* (IID *indexIID*, OBJECTID *indexObjectID*) in table space *indexTablespaceName* (ID *indexTablespaceID*) for table *tableName* (ID *tableID*) in table space ID *tablespaceID* failed on this node with SQLCODE *SQLCODE* reason code *reasonCode*. To resolve this problem, re-submit the REORG INDEX command on the failing node(s).

ADM9513W Online index reorganization on table *tableName* (ID *tableID*) in table space *tablespaceName* (ID *tablespaceID*) found one or more indexes that are marked invalid and cannot proceed until they are rebuilt.

Explanation: The data server will automatically rebuild the indexes on this table. If any nonpartitioned indexes are being rebuilt, the data server will obtain a super exclusive Z table lock for the duration of the

rebuild. If only partitioned indexes are being rebuilt, the data server will obtain a super exclusive Z partition lock on each data partition that has invalid indexes for the duration of the rebuild. After the rebuild is complete, the online index reorganization will proceed (using the original lock modes) for any indexes that were specified for reorganization by the current command and have not yet been rebuilt.

User response: No response is required.

ADM9514I **BEGIN async index cleanup on table** *tableName* **(ID** *tableID***) and table space** *tableName* **(ID** *tablespaceID***).**

ADM9515I **END async index cleanup on table** *tableName* **(ID** *tableID***) and table space** *tableName* **(ID** *tablespaceID***).**

ADM9516W **Indexes on the table** *table_identifier* **were marked to be rebuilt while upgrading the database.**

Explanation: The *table_identifier* is shown in one of the following forms:

- *TBSPACEID=table_space_id.TABLEID=table_id*
- *schema_name.table_name*

The indexes on the identified table must be rebuilt because one of the following situations was encountered during a database upgrade:

- A root page has insufficient space
- A type-1 index was detected
- One or more non-severe errors occurred converting the index page.

User response: Indexes will be rebuilt automatically after upgrading the database in one of the following ways:

- If the **indexrec** database configuration parameter is set to **RESTART** or **RESTART_NO_REDO**, then issuing the **RESTART DATABASE** command will trigger the rebuild of the indexes.
- If the **indexrec** database configuration parameter is set to **ACCESS** or **ACCESS_NO_REDO**, then the indexes will be rebuilt on first access to the table on which the indexes are defined. The **ADMIN_GET_TAB_INFO** function can be used to identify which tables have indexes that require rebuilding.

There will be a one-time performance impact while the indexes are rebuilt.

ADM9518I **Index reorganization has started for the partitioned indexes on data partition** *dataPartitionID* **of table** *tableName* **(ID** *tableID***) and table space** *tableName* **(ID** *tablespaceID***).**

Explanation: The data server is reorganizing the partitioned indexes for the specified data partition.

User response: No response is required.

ADM9519I **Index reorganization is complete for partitioned indexes on data partition** *dataPartitionID* **of table** *tableName* **(ID** *tableID***) and table space** *tableName* **(ID** *tablespaceID***).**

Explanation: The data server has reorganized the partitioned indexes for the specified data partition.

User response: No response is required.

ADM9520I **Reorganizing partitioned index IID** *indexIID* **(OBJECTID** *indexObjectID***) in table space** *indexTableName* **(ID** *indexTablespaceID***) for data partition** *dataPartitionID* **of table** *tableName* **(ID** *tableID***) in table space** *tableName* **(ID** *tablespaceID***).**

Explanation: The data server is reorganizing the index partition on the specified data partition for the specified partitioned index.

User response: No response is required.

ADM9521W **Index reorganization for partitioned indexes on data partition** *dataPartitionID* **of table** *tableName* **(ID** *tableID***) and table space** *tableName* **(ID** *tablespaceID***) failed on this database partition with** **SQLCODE** *SQLCODE* **reason code** *reasonCode*.

Explanation: Index reorganization of the partitioned indexes failed on this database partition for the reason described by the **SQLCODE**.

User response: Correct the problem that is described by the **SQLCODE**, and then retry the **REORG INDEXES** command on the database partition.

ADM9522I **Index reorganization is complete for partitioned indexes on data partition** *dataPartitionID* **of table** *tableName* **(ID** *tableID***) and table space** *tableName* **(ID** *tablespaceID***).**

Explanation: The data server has reorganized the partitioned indexes for the specified data partition. Some partitioned indexes on the data partition might still need to be reorganized. This index reorganization will occur later during the reorganization.

User response: No response is required.

Chapter 20. ADM10000 - ADM10499

ADM10000W A Java exception has been caught. The
Java stack traceback has been written to
the db2diag log file.

Chapter 21. ADM10500 - ADM10999

ADM10500E Health indicator *Health-Indicator-Short-Description (Health-Indicator-Short-Name)* **breached the Threshold-Bound-Name alarm threshold of Threshold-Bound-Value with value Health-Indicator-Value on Monitored-Object-Type Monitored-Object-Name. Calculation:** *Formula-String = Formula-with-Values = Health-Indicator-Value. History (Timestamp, Value, Formula): Health-Indicator-History-List*

Explanation: The health monitor generated an alert because the alarm threshold for this health indicator was breached. This situation should be addressed immediately as it can lead to a degradation in database performance or an interruption in operation.

User response: You can use CLP commands to obtain recommendations, and in some cases take actions, to resolve this alert.

From the CLP, you can obtain the Health indicator description and recommended actions by executing the following commands:

- GET RECOMMENDATIONS FOR HEALTH INDICATOR *Health-Indicator-Short-Name*
- GET DESCRIPTION FOR HEALTH INDICATOR *Health-Indicator-Short-Name*

ADM10501W Health indicator *Health-Indicator-Short-Description (Health-Indicator-Short-Name)* **breached the Threshold-Bound-Name warning threshold of Threshold-Bound-Value with value Health-Indicator-Value on Monitored-Object-Type Monitored-Object-Name. Calculation:** *Formula-String = Formula-with-Values = Health-Indicator-Value. History (Timestamp, Value, Formula): Health-Indicator-History-List*

Explanation: The health monitor generated an alert because the warning threshold for this health indicator was breached. This condition does not necessarily require immediate attention, but rather it may lead to a degradation in database performance or an interruption in operation if the condition worsens over time.

User response: You can use CLP commands to obtain recommendations, and in some cases take actions, to resolve this alert.

From the CLP, you can obtain the Health indicator description and recommended actions by executing the following commands:

- GET RECOMMENDATIONS FOR HEALTH INDICATOR *Health-Indicator-Short-Name*
- GET DESCRIPTION FOR HEALTH INDICATOR *Health-Indicator-Short-Name*

ADM10502W Health indicator *Health-Indicator-Short-Description (Health-Indicator-Short-Name)* **is in state Health-Indicator-Value on Monitored-Object-Type Monitored-Object-Name.**

Explanation: The health monitor generated an alert because the state value of this health indicator was non-normal. This condition does not necessarily require immediate attention, but rather it will depend on the expected state of the database given operations being performed on it at the time, and the prevailing workload.

User response: You can use CLP commands to obtain recommendations, and in some cases take actions, to resolve this alert.

From the CLP, you can obtain the Health indicator description and recommended actions by executing the following commands:

- GET RECOMMENDATIONS FOR HEALTH INDICATOR *Health-Indicator-Short-Name*
- GET DESCRIPTION FOR HEALTH INDICATOR *Health-Indicator-Short-Name*

ADM10503I The health monitor has initiated an alert action, running Alert-Action-Type Alert-Action-Name on system System-Name, because the Health Indicator Health-Indicator-Short-Description (Health-Indicator-Short-Name) is in the Alert-State alert state on Monitored-Object-Type Monitored-Object-Name.

Explanation: The health monitor was configured to initiate the action when the health indicator is in this alert state. This message is an indication that the action was indeed initiated.

User response: No action is required.

ADM10504E The health monitor failed, with sqlcode *SQLCODE*, to initiate an alert action, running *Alert-Action-Type* *Alert-Action-Name* on system *System-Name*, when the Health Indicator *Health-Indicator-Short-Description* (*Health-Indicator-Short-Name*) went into the *Alert-State* alert state on *Monitored-Object-Type* *Monitored-Object-Name*.

Explanation: The health monitor was configured to initiate the action when the health indicator is in this alert state, but received this SQLCODE when it called the API to execute the action. The alert action was not initiated.

User response: Check the First Failure Service log (db2diag log file) for a record detailing the failure.

ADM10505E The DB2 Service does not have the necessary authority to run the Health Monitor. The Health Monitor has been shut down. If the service is configured to log on using the Local System account (SYSTEM), then it must be changed to log on with a particular user account. If it is configured to log on with a particular user account, then you must ensure that the user account is valid and has the necessary access rights to run the DB2 service. Once the log on configuration has been corrected, it is necessary to restart the DB2 service to start the Health Monitor.

ADM10506E The health monitor is unable to send an alert notification because the SMTP Server (smtp_server) DB2 Administration Server configuration parameter is not set. Update the smtp_server configuration parameter with the name of a valid SMTP server.

Explanation: The health monitor was configured to send notifications upon alert occurrence, but was unable to send the notification because no SMTP server name was specified for the SMTP Server DAS configuration parameter.

User response: Update the smtp_server configuration parameter with the name of a valid SMTP server.

ADM10507E The health monitor was unable to send an alert notification because the server *SMTP-Server-Name*, specified in the SMTP Server DB2 Administration Server configuration parameter (smtp_server), does not appear to be an SMTP server. Ensure that a valid SMTP server name is specified in the DB2

Administration Server configuration.

Explanation: The health monitor was configured to send notifications upon alert occurrence, but was unable to send the notification because the server specified in the DB2 Administration Server configuration does not have SMTP server functionality.

User response: Ensure that a valid SMTP server name is specified in the DB2 Administration Server configuration for the smtp_server parameter.

ADM10508E The health monitor was unable to send an alert notification because invalid recipient(s) were specified in the health notification list which contains *Notification-List*. Update the Contact record with the invalid address.

Explanation: The health monitor was configured to send notifications upon alert occurrence, but was unable to send the notification because one or more addresses for the contacts specified on the health notification list are invalid.

User response: Check the Contact record for the contacts specified for health notification and update the invalid recipient address.

ADM10509E The health monitor was unable to send an alert notification because the notification was sent by an invalid Sender with the address *Sender-Address*. Look at the SMTP server configuration. If all the settings are correct, contact DB2 Support.

Explanation: The health monitor was configured to send notifications upon alert occurrence, but was unable to send the notification because the Sender address was rejected as unacceptable by the SMTP server. The Sender address has the format <instance name>@<host>, where 'instance' is running on 'host'.

User response: Look at the SMTP server configuration. If all the settings are correct, contact DB2 Support.

ADM10510E The health monitor was unable to send an alert notification because the SMTP server issued the following error: *SMTP_ERROR*. Check the SMTP server documentation for information on the error code that was returned. If the problem cannot be remedied, contact DB2 Support.

Explanation: The health monitor was configured to send notifications upon alert occurrence, but was unable to send the notification because the SMTP server encountered an error.

User response: Check the SMTP server documentation

for information on the error code that was returned. If the problem cannot be remedied, contact DB2 Support.

ADM10511E The health monitor was unable to send an alert notification because there was a communication error with the SMTP server. Check the First Failure Service log (db2diag log file) for a record detailing the failure.

Explanation: The health monitor was configured to send notifications upon alert occurrence, but was unable to send the notification because there was a communication error when trying to reach the SMTP server.

User response: Check the First Failure Service log (db2diag log file) for a record detailing the failure.

ADM10512W Health indicator *Health-Indicator-Short-Description (Health-Indicator-Short-Name)* is in state *Health-Indicator-Value* on *Monitored-Object-Type Monitored-Object-Name*. Collection (Object Name, Timestamp, Value, Detail): *Collection*.

Explanation: The health monitor generated an alert because the state value of this health indicator was non-normal. This condition does not necessarily require immediate attention but indicates that a non-optimal situation prevails with respect to the health of one or many object collected under this health indicator. The situation may get resolved automatically if the corresponding automatic maintenance utility was turned on and the state is automated .

User response: You can use CLP commands to obtain recommendations, and in some cases take actions, to resolve this alert.

From the CLP, you can obtain the Health indicator description and recommended actions by executing the following commands:

- GET RECOMMENDATIONS FOR HEALTH INDICATOR *Health-Indicator-Short-Name*
- GET DESCRIPTION FOR HEALTH INDICATOR *Health-Indicator-Short-Name*

ADM10513I Automatic *Utility-Name* has completed on table *Table-Name* in database *Database-Name* with a return code of *SQL-Return-Code*. The utility started at *Start-Timestamp* and completed at *End-Timestamp*.

ADM10514I Automatic BACKUP has completed on database *Database-Name* with a return code of *SQL-Return-Code*. The utility started at *Start-Timestamp* and completed at *End-Timestamp*. The timestamp for the backup image is *Backup-Image-Timestamp*.

ADM10515I The automatic maintenance policy *Policy-Name* has been updated in database *Database-Name*. The policy options have been updated from *Original-Policy-Options-List* to *New-Policy-Options-List*.

ADM10516I The automatic maintenance policy *Policy-Name* has been updated in database *Database-Name*. The options now being used for this policy are: *Policy-Options-List*.

Chapter 22. ADM11000 - ADM11499

ADM11000E DB2 is unable to create or attach to the memory segment used for fenced routine communications. Decrease the amount of database shared memory used by your instance, and retry.

ADM11001E DB2 did not create a memory segment for running fenced routines. This was specified by the use of DB2_FMP_COMM_HEAPSZ registry variable.

ADM11002E Insufficient shared memory available for communication with the db2fmp process. Use the DB2_FMP_COMM_HEAPSZ registry variable to increase the amount of shared memory available for fenced routines.

ADM11003E DB2 failed to create the memory segment used for communication with fenced routines. If restarting DB2, ensure that no db2fmp processes are active on the instance prior to start. Otherwise, you can adjust the value through the DB2_FMP_COMM_HEAPSZ registry variable, or you can decrease the value of ASLHEAPSZ in the database manager configuration.

Chapter 23. ADM11500 - ADM11999

ADM11500W MQListener generated a message.
Message code = *MQL-msgcode*. Refer to
the documentation for information
about the message.

Chapter 24. ADM12000 - ADM12499

ADM12000C DB2START processing failed; a valid product license was not found. If you have licensed this product, ensure the license key is properly installed. You can install the license using the db2licm command. The license file can be obtained from your licensed product CD.

User response:

ADM12001C DB2 connect processing failed; a valid product license was not found. If you have licensed this product, ensure the license key is properly installed. You can install the license using the db2licm command. The license file can be obtained from your licensed product CD.

User response:

ADM12002C Connect processing failed; a valid product license was not found. If you are accessing host or iSeries database server, make sure you have a licensed DB2 Connect product or DB2 Connect server support component installed. DB2 Connect server support component is included in DB2 Enterprise edition.

User response:

ADM12006E The product *product-name* does not have a valid license key registered. If you have licensed this product, ensure the license key is properly registered. You can register the license using the db2licm command. The license key can be obtained from your licensed product CD. If a license key is not registered, this product will be enabled for a *num-days* day evaluation period. Use of the product for the evaluation period constitutes acceptance of the terms of the IBM license agreement located in the installation path of this product in the license directory.

User response:

ADM12007E There are *num-days* day(s) left in the evaluation period for the product *product-name*. For evaluation license terms and conditions, refer to the IBM License Acceptance and License Information document located in the license directory in the installation path of this product. If you have licensed this product, ensure the license key is properly registered. You can register the license using the db2licm command line utility. The license file can be obtained from your licensed product CD.

User response:

ADM12008C The product *product-name* does not have a valid license key installed and the evaluation period has expired. Functions specific to this product are not enabled. If you have licensed this product, ensure the license key is properly installed. You can install the license using the db2licm command. The license file can be obtained from your licensed product CD.

User response:

ADM12009E The number of concurrent users of the DB2 Workgroup product has exceeded the defined entitlement of *entitlement*. Concurrent user count is *user-count*. You should purchase additional user based entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12010E The number of concurrent users of the DB2 Connect product has exceeded the defined entitlement of *entitlement*. Concurrent user count is *user-count*. You should purchase additional user based entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12011C One or more database partitions does not have a valid DB2 license key installed for the *product-name* product. Install a valid license key on each physical partition using the db2licm command.

User response:

ADM12012E The number of concurrent users of the DB2 Enterprise product has exceeded the defined entitlement of *entitlement*. Concurrent user count is *user-count*. You should purchase additional user based entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12013E The number of concurrent database connections to the DB2 Connect product has exceeded the defined entitlement of *entitlement*. Database connection count is *num-connections*.

ADM12014C The version of the DB2 Connect product you are using is not licensed for use with TCP/IP protocol. Upgrade to a full function DB2 Connect product to use TCP/IP.

ADM12015C The version of the DB2 Connect product you are using is not licensed for updating multiple databases in the same transaction. Upgrade to a full function DB2 Connect product to update multiple databases in the same transaction.

ADM12017E The number of processors on this machine exceeds the defined entitlement of *entitlement* for the product *product-name*. The number of processors on this machine is *num-cpu*.

Explanation: You can purchase DB2 Connect and DB2 database products either per processor (priced by Processor Value Unit (PVU)) or per user.

If you purchase your DB2 database product license per processor, you will see License Type: "CPU Option" in the License Center or in the output from the command "db2licm -l".

Prior to DB2 V9.5 Fix Pack 4, this message is returned by DB2 Connect and DB2 database products when it is determined that there are more processors on the current machine than are entitled to be used with the named product. The DB2 product will continue to

function even when this message is returned. However, using the DB2 product with more processors than the defined entitlement allows will cause problems with business audit processes and will complicate support processes, if you need to access IBM support.

This message is not returned by DB2 Connect and DB2 database products starting with DB2 V9.5 Fix Pack 4 and later because the DB2 product will only use the machine resources for which it is licensed to use.

User response:

1. Purchase additional processor-based entitlements from your IBM representative or authorized dealer.
2. Update your license using the db2licm command.

ADM12018E The number of concurrent users for this product has exceeded the defined entitlement of *entitlement*. Concurrent user count is *user-count*. You should purchase additional user based entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12020E The number of connectors has exceeded the defined entitlement of *entitlement*. Current number of connectors is *num-connectors*. You should purchase additional connector entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12022E DB2 has detected that database partitioning feature is being used without database partitioning license. Purchase database partitioning entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12023E The number of concurrent users of *product-name* product has exceeded the defined entitlement of *entitlement*. Concurrent user count is *user-count*. You should purchase additional user based entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12024E A valid license key was not found for the requested function. The current license key for *product_name* product does not allow the requested functionality. Purchase the license key for this function from your IBM representative or authorized dealer and update your license using the db2licm command.

User response:

ADM12025E The amount of memory on this machine exceeds the defined limit of *limit* (MB) for the product *product-name*. The amount of memory on this machine is *memory* (MB).

Explanation: This product has a defined memory limit that has been exceeded. The memory limit cannot be changed via DB2 licensing tools.

User response: Contact your IBM representative or authorized dealer to obtain a product that can be licensed to run on this system.

ADM12026W The DB2 server has detected that a valid license for the product *product-name* has not been registered.

Explanation: Registration of a valid license key is required in order to comply with the terms and conditions of your License Agreement. The license keys for this product are located on this product's activation CD in the 'license' directory.

User response: Use the db2licm command (Run db2licm -a *license-file-name* from sqllib\bin on Windows or sqllib/adm on Unix and Linux) to register the appropriate licenses that you have purchased. The License Agreement text is located in the 'license' directory in the installation directory of this product.

ADM12027E The amount of memory on this server now exceeds the defined limit of *limit* (GB) for the product *product-name*. The amount of memory on this server is *memory* (GB).

Explanation: The DB2 instance is attempting to use more memory allocated for DB2 than specified in your product license. DB2 limits itself to the licensed amount of memory.

User response: To take full advantage of your server's memory capacity, contact your IBM representative or authorized dealer to obtain an edition of DB2 with a higher licensed memory limit.

Chapter 25. ADM12500 - ADM12999

ADM12500E The HADR standby database cannot be made consistent with the primary database. The log stream of the standby database is incompatible with that of the primary database. To use this database as a standby, it must be recreated from a backup image or split mirror of the primary database.

ADM12501E Unable to establish HADR primary-standby connection because the operating systems do not match the primary and standby databases. Move the primary or standby databases to a different host or upgrade the operating system of one host to match the other host.

ADM12502E Unable to establish HADR primary-standby connection because the DB2 versions do not match the primary and standby databases. Upgrade the DB2 software to the same release and FixPak on either the HADR primary or HADR standby database to match the other.

ADM12503E Unable to establish HADR primary-standby connection because the DB2 bit level (32-bit vs. 64-bit) do not match the primary and standby databases. Change the bit level of the primary or standby database to match each other.

ADM12504E Unable to establish HADR primary-standby connection because the value of HADR_REMOTE_INST at one of the instances does not match the actual instance name of the other instance. This is a sanity check to ensure that only the intended database pairing occurs. If any of the HADR_REMOTE_INST configuration parameters or instance names is set incorrectly, you may correct it and try again to start HADR.

ADM12505E Unable to establish HADR primary-standby connection because the database names do not match. Correct the database name so it matches on the HADR primary and HADR standby.

ADM12506E Unable to establish HADR primary-standby connection because the primary and standby databases did not originate from the same database. Recreate the standby from a backup image or split mirror of the primary database.

ADM12507E Unable to establish HADR primary-standby connection because the HADR configuration parameters do not match. Change the configuration parameters HADR_TIMEOUT and HADR_SYNCMODE on either the HADR primary or standby so that they match the other system's values, and ensure that HADR_LOCAL_HOST and HADR_REMOTE_HOST match the host name of the local and remote machines.

ADM12508W Log receiving has been suspended on the HADR standby database because of a disk full condition. If the primary and standby databases are in peer state and in SYNC, NEARSYNC, or ASYNC HADR synchronization mode, transactions on the primary might be blocked.

Explanation: The standby database is unable to receive any more log data from the primary database because of a disk full condition on the standby. The standby frees disk space automatically after it finishes replaying logs, as long as the logs are archived on the primary and there are no uncommitted transactions started in or before the given log file.

User response: One of the following:

- Wait for the standby to catch up with the primary. Log receiving will resume on the standby as old logs are released after they have been replayed on the standby and archived on the primary.
- Provide more disk space on the standby for logs. Do not delete any logs because DB2 handles that automatically.

If the primary was blocked because of a disk full situation on the standby, consider using the SUPERASNYC HADR synchronization mode to avoid encountering this error in the future.

ADM12509E HADR encountered an abnormal condition. Reason code: *reason-code*

Explanation: The explanation corresponding to the reason code is:

1

The HADR primary database cannot find a log file that was requested by the standby. HADR will attempt to recover from the condition by disconnecting and then reconnecting the primary and the standby. Each reconnection will attempt to access the file. These retry attempts will help avoid transient errors such as a conflict with other log file usage.

2

The HADR standby received a bad log page. HADR will attempt to recover from the condition by disconnecting and then reconnecting the primary and standby. These retry attempts will help avoid transient errors.

3

The HADR primary has broken its connection to the standby because the primary reach its peer limit. This condition indicates that the standby or network cannot keep up with the primary's workload. The primary has disconnected from the standby to unblock itself. After the standby has replayed the logs it received, it will reconnect to the primary and the pair will start from remote catchup state, which does not block primary log writing.

4

The HADR standby is shutting down because its log writing disk is full and it cannot reclaim any space. The standby log device is not large enough to hold the logs produced by the primary.

5

The HADR standby was unable to access the needed log file or log files. Possible causes of the access failure include user removal of log files, or I/O errors on the log device. The HADR standby automatically manages log files, so always ensure that they are not manually removed or altered. The HADR standby will attempt to recover from the condition by starting remote catch up from the missing log file.

6

Detected another primary database.

7

The standby database is shutting down after repeated failures to start the replay member in DB2 pureScale environment.

8

In a DB2 pureScale environment, an address listed in `hadr_target_list` does not actually belong to the remote cluster.

9

HADR does not support column-organized tables.

10

A previous takeover by force was performed on this standby database but it did not complete.

99

Fatal error encountered. The database is shutting down.

User response: Either re-create the standby using a newer backup image from the primary, or follow the user response corresponding to the reason code:

1

If the attempts to access the log file keep failing, check the `db2diag.log` file for the file number and determine why the file cannot be found. This can occur if the file is very old and has been removed from the archive, or if there has been a media failure. Depending on the cause, possible responses are:

- Put the file back in the log path or archive.
- Repair the media that failed.

2

If the error condition persists, determine the root cause, make the appropriate corrections, and attempt to start HADR on the database.

3

One of the following actions:

- Increase the setting for the `DB2_HADR_PEER_WAIT_LIMIT` registry variable.
- Reduce the workload on the primary.
- Tune or upgrade the network.
- Tune or upgrade the standby.

4

Increase the log device capacity and restart the standby.

5

The standby should be able to recover by requesting the files from the primary. If the error condition persists, determine the root cause, and make the appropriate corrections.

6

Determine which database should be the primary. Drop the other primary database or

| | | | |
|---|--|---|--|
| | try converting it to a standby. If both should be primary, they should be removed from each other's <code>hadr_target_list</code> , so that they are independent. | | for the <code>hadr_local_host</code> and <code>hadr_local_service</code> configuration parameters on the remote database. |
| 7 | | 3 | |
| | Determine the cause of the failures, correct the problem, then restart the standby database. If the problem is limited to a particular member, you can designate another member as the preferred replay member by running the <code>START HADR</code> command with the <code>AS STANDBY</code> option on another member. | 4 | The <code>hadr_target_list</code> configuration parameter is not configured on all of the HADR databases. |
| 8 | | | The local database cannot be started as an HADR standby because its <code>hadr_target_list</code> configuration parameter setting does not specify the HADR primary database's local host and service names. |
| | This is not a fatal error. HADR remains functional. Remove the bad address from the <code>hadr_target_list</code> at the earliest convenient moment. | 5 | |
| | | | The setting for the <code>hadr_timeout</code> configuration parameter is not the same on the local and the remote databases. |
| 9 | | 6 | |
| | If you still want to use HADR, drop any column-organized tables on the primary database, and then reinitialize the standby database. | | The setting for the <code>hadr_syncmode</code> configuration parameter is not the same on the local and the remote databases. |
| 10 | | 7 | |
| | Do one of the following: | | The setting for the <code>hadr_peer_window</code> configuration parameter is not the same on the local and the remote databases. |
| | <ul style="list-style-type: none"> • Reissue takeover by force to complete the takeover. • Deactivate the database, stop HADR, and then roll forward the database to convert it into a standard database. • Deactivate, then drop the database. Re-create it as a standby if needed. | 8 | |
| 99 | | | Delayed replay cannot be enabled because the standby database is not running in superAsync synchronization mode. When you have the <code>hadr_target</code> list configured, the effective synchronization mode of the principal standby is determined by the setting of the <code>hadr_syncmode</code> configuration parameter on the primary. |
| | Determine the root cause, make the appropriate corrections, and attempt to start the database. | | User response: The user response corresponding to the reason code is: |
| <hr/> | | | |
| ADM12510E Unable to establish HADR primary-standby connection. Reason code: <i>reason-code</i> | | | |
| Explanation: The explanation corresponding to the reason code is: | | | |
| 1 | | 1 | |
| | The settings for the <code>hadr_remote_host</code> and <code>hadr_remote_service</code> configuration parameters on the remote database do not match the settings for the <code>hadr_local_host</code> and <code>hadr_local_service</code> configuration parameters on the local database. | | On the remote database, set the <code>hadr_remote_host</code> configuration parameter to the local database's local host name (specified by the <code>hadr_local_host</code> configuration parameter) and set the <code>hadr_remote_svc</code> configuration parameter to the local database's local service name (specified by the <code>hadr_local_host</code> configuration parameter). |
| 2 | | 2 | |
| | The settings for the <code>hadr_remote_host</code> and <code>hadr_remote_service</code> configuration parameters on the local database do not match the settings | | On the local database, set the <code>hadr_remote_host</code> configuration parameter to the remote database's local host name (specified by the <code>hadr_local_host</code> configuration parameter) and set the <code>hadr_remote_svc</code> configuration parameter to the remote |

database's local service name (specified by the `hadr_local_host` configuration parameter).

3

Ensure that the `hadr_target_list` configuration parameter is configured on all of the HADR databases or not configured on any of the HADR databases.

4

Ensure that the local database specifies the HADR primary database's local host and service names in its `hadr_target_list` setting.

5

Ensure that the setting for the `hadr_timeout` configuration parameter is the same on all HADR databases.

6

Ensure that the setting for the `hadr_syncmode` configuration parameter is the same on the primary and standby HADR databases.

7

Ensure that the setting for the `hadr_peer_window` configuration parameter is the same on the primary and standby HADR databases.

8

If you have multiple standbys configured, ensure that the `hadr_replay_delay` configuration parameter is configured on an auxiliary standby.

ADM12511W The HADR standby database was not able to access the log archive when attempting to retrieve log file *log-file-name*.

Explanation: An error occurred when the standby database attempted to retrieve a log file from the log archive. This error is not fatal. The standby database will attempt to acquire the file from the primary database by entering remote catchup state.

User response: Investigate why the standby database is not able to access the log archive:

- Make sure that LOGARCHMETH1 and LOGARCHMETH2 are configured correctly.
- Make sure that the standby machine and instance owner can access the log archive device.
- Make sure that the archive contains log files on the correct log chain.

It is important to resolve this issue, especially if this database might have to perform takeover to become a primary database in the future. After it is a primary database, it must be able to access the archive device to avoid running out of disk space in the active log path.

ADM12512W Log replay on the HADR standby has stopped on table space *tablespace_name* (ID *tablespace_id*) because it has been put into *tablespace_state* state.

Explanation: The standby database can no longer replay log records on this table space due to error. Data in this table space will not be available if this database takes over the primary role. The standby database will continue replaying logs shipped from primary on other table spaces.

User response: Investigate and resolve the possible causes, such as:

- File system is not mounted.
- The standby is unable to access some containers, or the container is not big enough.
- The standby is unable to locate or access a load copy image.

ADM12513E Unable to establish HADR primary-standby connection because HADR configurations are incompatible on primary and standby. Reason code: *reason-code*

Explanation:

1

The settings for the `hadr_local_host` and `hadr_local_svc` configuration parameters on the local database must match the settings for the `hadr_remote_host` and `hadr_remote_svc` configuration parameters on the remote database.

2

The settings for the `hadr_local_host` and `hadr_local_svc` configuration parameters on the remote database must match the settings for the `hadr_remote_host` and `hadr_remote_svc` configuration parameters on the local database.

3

At least one of the HADR databases does not have its `hadr_target_list` configuration parameter set. With the `hadr_target_list` parameter, all of the participating databases must have it set or unset.

4

The `hadr_target_list` configuration parameter on the local database does not include entries that correspond to the `hadr_local_host` and `hadr_local_svc` configuration parameter settings for the remote database members.

5

| | | | |
|----|---|----|---|
| | The local log stream ID does not match the remote log stream ID. | 1 | |
| 6 | | | Configure the <code>hadr_local_host</code> and <code>hadr_local_svc</code> configuration parameters on the local database to match the settings for the <code>hadr_remote_host</code> and <code>hadr_remote_svc</code> configuration parameters on the remote database. |
| | The local database member ID does not match the remote database member ID ID. | | |
| 7 | | | |
| | The setting for the <code>hadr_timeout</code> configuration parameter on the local database must match the <code>hadr_timeout</code> setting on the remote database. | 2 | |
| 8 | | | Configure the <code>hadr_local_host</code> and <code>hadr_local_svc</code> configuration parameters on the remote database to match the settings for the <code>hadr_remote_host</code> and <code>hadr_remote_svc</code> configuration parameters on the local database. |
| | If the <code>hadr_target_list</code> configuration parameter is not configured on your HADR databases, the setting for the <code>hadr_syncmode</code> configuration parameter on the local database must match the <code>hadr_syncmode</code> setting on the remote database. | 3 | |
| 9 | | | Either configure the <code>hadr_target_list</code> configuration parameter on all of the databases or on none of the databases. |
| | The <code>hadr_replay_delay</code> configuration parameter is set to a non-zero value on the standby, but the effective synchronization mode (which derived from the setting on the primary) is not <code>SUPERASYNC</code> . | 4 | |
| 10 | | | Add at least one value from the <code>hadr_local_host</code> and <code>hadr_local_svc</code> configuration parameter settings for a remote database member to the <code>hadr_target_list</code> configuration parameter on the local database. |
| | The setting for the <code>hadr_peer_window</code> configuration parameter on the local database must match the <code>hadr_peer_window</code> setting on the remote database. | 5 | |
| 11 | | | Because the two databases are incompatible with each other, HADR cannot be set up. Reinitialize the standby with a backup or split mirror image from the primary. |
| | The primary and standby database types are incompatible; for example, if one of the databases uses the DB2 pureScale feature, the other must as well. | 6 | |
| 12 | | | Because the two databases are incompatible with each other, HADR cannot be set up. Reinitialize the standby with a backup or split mirror image from the primary. |
| | The primary and standby database topologies are incompatible. The two databases must have same number of members and each pair of primary-standby matching members must have same member id and log stream id. HADR cannot be set up if the two databases do not meet this requirement. | 7 | |
| 13 | | | Configure the <code>hadr_timeout</code> configuration parameter to the same value on the primary and standby databases. |
| | The two databases are incompatible because they do not have the same replay type (physical replay). | 8 | |
| 99 | | | Configure the <code>hadr_syncmode</code> configuration parameter to the same value on the primary and standby databases. |
| | Other incompatibility errors. | 9 | |
| | User response: Refer to the explanation of the error message that follows this message and take any needed corrective action. | 10 | |
| | | | Set the <code>hadr_replay_delay</code> configuration parameter to 0, or change the effective syncmode to <code>SUPERASYNC</code> by setting the <code>hadr_syncmode</code> configuration parameter on the primary database. |

Configure the `hadr_peer_window` configuration parameter to the same value on the primary and standby databases.

11

Use the same database type for both the HADR primary and the HADR standby.

12

Reinitialize the standby with a backup or split mirror image from the primary so that standby has the same topology as the primary.

13

Use the physical replay type.

99

Generic incompatibility error. See the `db2diag.log` for more information.

ADM12514E Log replay on the HADR standby database has stopped because member *member-id* was added on the HADR primary but has not been added on the standby. The standby database has been deactivated due to this error.

Explanation: When a member is added to the HADR primary, it must also be added to the standby or the standby will be deactivated.

User response: Add the new member on the standby, then reactivate the database on the standby to restart log replay. You should also ensure that the `hadr_local_host` and `hadr_local_svc` configuration parameters have been set on the new member.

ADM12515E The HADR primary and standby databases are no longer compatible because the primary database has fewer members than the standby.

Explanation: This can occur in the following scenarios:

- A member was dropped on the primary.
- A member was added on the primary; then a forced takeover happened before the add member event was replicated to the standby; then the old primary attempted to reintegrate as a new standby, but it has more members than the new primary.

User response:

1. Cause the the instance of the primary database to have the same number of members as the instance of the standby database by performing one of the following actions:
 - Add another member to the instance of the primary database.
 - Drop a member from the instance of the standby database.

2. Reinitialize the standby database by performing the following steps:

- a. Create a backup image of the primary database.
- b. Restore the image of the primary database to the standby database.

Chapter 26. ADM13000 - ADM13499

ADM13000E Plug-in *plugin-name* received error code *error-code* from the GSS (Generic Security Service) API *GSS-API-name* with the error message *error-message*.

ADM13001E Plug-in *plugin-name* received error code *error-code* from the DB2 security plug-in API *GSS-API-name* with the error message *error-message*.

ADM13002E Unable to unload plug-in *plugin-name*.
No further action is required.

ADM13003E The principal name *principal-name* used for *plugin-name* is invalid. Ensure that the principal name is valid and that it is in a format that is recognized by the security plug-in.

ADM13004E The plug-in name *plugin-name* is invalid. Ensure that a valid plug-in name is specified.

ADM13005E Unable to load plug-in *plugin-name*. Verify that the plug-in exists and that the directory location and file permissions are valid.

ADM13006E Plug-in *plugin-name* encountered an unexpected error. Contact IBM Support for assistance.

Chapter 27. ADM13500 - ADM13999

ADM13500E An agent executing an asynchronous background task processor encountered an unrecoverable error. The task processor has been suspended and diagnostic information was written to the db2diag log file. Contact IBM Support for assistance. The task processor context is *address*. The task processor description is *description*.

Chapter 28. ADM14000 - ADM14499

ADM14000E DB2 is unable to open diagnostic log file *filename*. Run the command "db2diag -rc rcList" to find out more.

ADM14001C An unexpected and critical error has occurred: *error-type*. The instance may have been shutdown as a result. *capture-type* FODC (First Occurrence Data Capture) has been invoked and diagnostic information has been recorded in directory *directory-name*. Please look in this directory for detailed evidence about what happened and contact IBM support if necessary to diagnose the problem.

ADM14002C *capture-type* FODC has been invoked for symptom *error-type* and diagnostic information has been recorded in directory *directory-name*. Please look in this directory for detailed evidence about what happened and contact IBM support if necessary to diagnose the problem.

ADM14003W FODC has been invoked by the user from db2fodc tool for symptom *symptom* and diagnostic information has been recorded in directory *directory*. Please look in this directory for detailed evidence about what happened and contact IBM support if necessary to diagnose the problem.

ADM14004C EDU Database *database-name* marked bad. *capture-type* FODC has been invoked and diagnostic information has been recorded in directory *parameter*. Please look in this directory for detailed evidence about what happened and contact IBM support if necessary to diagnose the problem.

ADM14005E The following error occurred: *symptom*. First Occurrence Data Capture (FODC) has been invoked in the following mode: *capture-mode*. Diagnostic information has been recorded in the directory named *directory-name*.

Explanation: First occurrence data capture (FODC) is a general term applied to the set of diagnostic information the DB2 administration server captures

automatically when errors occur.

For more information about FODC, search the DB2 Information Center for "First occurrence data capture (FODC)".

User response: Review the diagnostic information, such as log files, dump files, or trap files, in the named directory.

ADM14010C An unexpected and critical error has occurred: *error-type*. *capture-type* First Occurrence Data Capture has been invoked and diagnostic information has been recorded in directory *directory-name*.

Explanation: One or more DB2 threads associated with this instance have been suspended, but the instance process is still running. The DB2 instance might become unstable and must be stopped and restarted.

User response: To restore the stability of the DB2 instance, stop and restart the instance by running the following commands at a command prompt:

db2_kill

db2start

If at all possible, wait until the DB2 instance is no longer accessed by any applications before issuing the db2_kill command. db2_kill may result in crash recovery processing upon subsequent db2start.

Look in the named directory for detailed evidence about what happened and, if required, contact IBM Software Support to diagnose the problem.

ADM14011C A critical failure has caused the following type of error: *error-type*. The DB2 database manager cannot recover from the failure. First Occurrence Data Capture (FODC) was invoked in the following mode: *capture-type*. FODC diagnostic information is located in the following directory: *directory-name*.

Explanation: First occurrence data capture (FODC) is a general term applied to the set of diagnostic information the DB2 administration server captures automatically when errors occur.

The DB2 engine on all supported platforms has trap resilience capabilities that enable it to respond to certain code traps or segmentation violations in order to keep the DB2 instance running.

This message is returned when a trap is caused by a critical error from which the database manager cannot

recover, despite trap resilience being enabled.

The DB2 database manager will stop the instance.

User response:

1. Collect the FODC diagnostic information from the named directory.
2. Contact IBM Software Support to diagnose the problem.

ADM14012C A critical failure has caused the following type of error: *error-type*. The DB2 database manager will attempt to recover from the failure. First Occurrence Data Capture (FODC) was invoked in the following mode: *capture-type*. FODC diagnostic information is located in the following directory: *directory-name*.

Explanation: First occurrence data capture (FODC) is a general term applied to the set of diagnostic information the DB2 administration server captures automatically when errors occur. For more information about FODC, refer to the topic called "First occurrence data capture (FODC)" in the DB2 Information Center.

The DB2 engine on all supported platforms has trap resilience capabilities that enable it to survive certain code traps or segmentation violations in order to keep the DB2 instance running. For more information about sustained traps, refer to the topic called "Recovering from sustained traps" in the DB2 Information Center.

This message is returned when a trap is caused by a critical error, and the DB2 database manager will attempt to recover from the failure because trap resilience is enabled.

The DB2 database manager will stop the instance.

User response: Although the DB2 database manager will attempt to recover from the failure, it is important to diagnose why the failure happened by performing the following steps:

1. Collect the FODC diagnostic information from the named directory.
2. Contact IBM Software Support to diagnose the problem.

ADM14013C The following type of critical error occurred: *error-type*. This error occurred because one or more threads that are associated with the current DB2 instance have been suspended, but the instance process is still running. First Occurrence Data Capture (FODC) was invoked in the following mode: *capture-type*. FODC diagnostic information is located in the following directory: *directory-name*.

Explanation: First occurrence data capture (FODC) is a general term applied to the set of diagnostic

information the DB2 administration server captures automatically when errors occur. For more information about FODC, refer to the topic called "First occurrence data capture (FODC)" in the DB2 Information Center.

For more information about how the DB2 database manager uses operating system threads and processes, refer to the topic called "The DB2 process model" in the DB2 Information Center. For more information about sustained traps, refer to the topic called "Recovering from sustained traps" in the DB2 Information Center.

User response:

1. To terminate all active applications which issue a COMMIT or ROLLBACK within the timeout period, which minimizes the recovery window for crash recovery when the db2start command is run, issue the following command:

```
db2 quiesce instance <instance_name>
   user <user_name>
   defer with timeout <minutes>
```

2. [Optional] To terminate any applications that did not COMMIT or ROLLBACK during the timeout period in Step 1 and any new applications which accessed the database after the timeout period completed, issue the following command:

```
db2 quiesce instance <instance_name>
   user <user_name>
   immediate
```

3. Forcefully shut down the instance and suspended EDUs by executing the following command:

```
db2_kill
```

Note: Issuing the db2stop command will not complete when an instance has sustained a trap.

4. Restart the DB2 instance using either one of the following commands:

```
db2start
```

or

```
START DATABASE MANAGER
```

ADM14014C An unexpected and critical error has occurred: *error-type*. This error might affect the availability of the *resource-or-service-name* resource or service on member *member-id*.

Explanation: One or more DB2 threads encountered an unexpected error, but the DB2 member is still running. The error might affect the availability of the specified resource or service for the affected database. The database is still available; however, it might not provide optimal service, especially if applications have a dependency on the specified resource or service.

User response: To restore the stability of the database, stop all applications on this member and force a restart of the DB2 instance on this member with the following commands:

1. `db2stop member <member-id>`
`quiesce <timeout_minutes>`

Existing applications have up to <timeout_minutes> to reach the next commit or rollback point before being forced off. After the specified <timeout_minutes>, all remaining applications are terminated.

2. `db2start member <member-id>`

Look in the DIAGPATH for the affected DB2 member for any additional evidence about what has occurred and, if required, contact IBM Software Support to diagnose the problem.

Chapter 29. ADM14500 - ADM14999

ADM14500E Unable to allocate memory required for deferred index cleanup on table *schema.table*. If you wish to use immediate cleanup rollout for the statement, either set the DB2_MDC_ROLLOUT registry variable to ON, or use the CURRENT ROLLOUT MODE special register, and rerun the statement.

Chapter 30. ADM15000 - ADM15499

ADM15000E The file *logfileName* is not accessible for reading. Verify the access permissions for this file and its associated device.

ADM15001E An error was returned trying to access file *logfileName*. Make sure that the file exists and that the device or file is accessible.

Chapter 31. ADM15500 - ADM15999

ADM15500E An index data inconsistency is detected on table *schema-name.table-name* during INSPECT command. Please contact DB2 support team to report the problem.

ADM15501W The administrative task scheduler encountered a temporary resource constraint that prevented the execution of task *task-id*. The scheduler will retry every *retry-interval* seconds.

ADM15502W The administrative task scheduler could not execute task *task-id* because the database is inactive.

ADM15503E The administrative task scheduler detected a security error on database *database-name*. No scheduled tasks will be executed on this database. To resume task execution, drop the SYSTOOLS.ADMINTASKS and SYSTOOLS.ADMINTASKSTATUS table, and recreate all the scheduled tasks on this database using the SYSPROC.ADMIN_TASK_ADD stored procedure.

ADM15510E The INSPECT command found inconsistent row contents in a block in the multidimensional clustering (MDC) table named *schema-name.table-name*.

Explanation: In the context of MDC tables, a block is a set of contiguous pages on disk. In MDC tables, rows of table data that contain the same indexes are clustered together on disk, in these blocks, to facilitate faster searching and improved performance.

The INSPECT command was verifying the block indexes of the named table, and discovered rows in one or more blocks that are not valid according to the block index entries. This can happen for several reasons, including disk error or data corruption.

User response: Refer to the db2diag log file for more information about this error.

Contact IBM software support for assistance.

Chapter 32. ADM16000 - ADM16499

ADM16000W Following an attach operation, a SET INTEGRITY...ALL IMMEDIATE UNCHECKED statement running against a partitioned table that has a nonpartitioned index defined on it executes as though it were a SET INTEGRITY...ALLOW WRITE ACCESS IMMEDIATE CHECKED statement. The SET INTEGRITY operation updates the index to include data for the newly attached partition. It also performs constraints checking, range validation, and maintenance of any other nonpartitioned indexes that are defined on the table.

Chapter 33. ADM16500 - ADM16999

ADM16500E The specified member subset references a member that does not exist in the instance. **Member subset:** *subset_name*.
Member: *member_id*.

Explanation: A member subset references a member that does not exist in the instance. The member subset definition was adjusted in-memory to exclude the specified member. The member subset definition in the system catalogs was not modified.

User response: Either remove the member from the member subset, or drop the member subset. To remove the member from the member subset definition, use system routine SYSPROC.WLM_ALTER_MEMBER_SUBSET. To drop the member subset, use system routine SYSPROC.WLM_DROP_MEMBER_SUBSET.

ADM16501E The specified exclusive member subset does not include any members that exist in member subset *member-subset-token*.

Explanation: The in-memory definition of exclusive member subset does not contain any members that exist in the instance. The in-memory definition of the subset was adjusted; the subset was disabled. Connections assigned to this member subset are rejected with SQLCODE SQL1717N. The system catalog definition of the subset is not modified.

User response: To add members which exist in the instance to the member subset definition, use the SYSPROC.WLM_ALTER_MEMBER_SUBSET system routine. Alternatively, to drop the member subset, use the SYSPROC.WLM_DROP_MEMBER_SUBSET system routine.

ADM16502I The specified database alias used by the member subset is not cataloged in the system database directory. The database manager will catalog the database alias.
Database alias: *database_alias*. **Member subset:** *member_subset*.

Explanation: The database alias is managed by a member subset but is not currently cataloged in the system database directory. The database manager will automatically catalog this database alias. Applications can connect to the database alias and be assigned to the member subset. The definition of a member subset can be viewed by querying the system catalog view SYSCAT.MEMBERSUBSETS.

User response: No action required.

ADM16503E An attempt to automatically catalog the specified database alias by the database manager has failed. **Database alias:** *database_alias*. **Member subset:** *member_subset*. **SQLCODE:** *sqlcode*.

Explanation: An attempt was made by the database manager to catalog a database alias that is managed by a member subset. The attempt failed. Connections to the database alias will fail to resolve the database.

User response: Explicitly catalog the database alias using the CATALOG DATABASE DATABASE_NAME AS DATABASE_ALIAS command. Alternatively, to drop the member subset, use system routine SYSPROC.WLM_DROP_MEMBER_SUBSET.

Part 3. AMI Messages

Chapter 34. AMI0000 - AMI0499

AMI0016E Insufficient memory.

Explanation: There is not enough memory available to complete the requested operation.

User response: Make sure that sufficient memory is allocated and available for the Message Query (MQ), MQ Application Messaging Interface (AMI), and DB2 programs.

AMI0018E Service not found.

Explanation: The specified service is not found. The request is not completed. The service might be one of the following:

- sender
- receiver
- distribution list
- publisher
- subscriber

User response: Make sure that the Application Messaging Interface (AMI) repository file that you used contains the service definition.

AMI0019E Message not found.

Explanation: The specified message is not found. The request is not completed.

User response: Make sure that the message is available for use.

AMI0020E Policy not found.

Explanation: The specified policy is not found. The request is not completed.

User response: Make sure that the Application Messaging Interface (AMI) repository file that you used contains the policy definition.

AMI0035E No message available.

Explanation: No message is available for a receive request after the specified wait time. This happens when the target queue is empty.

User response: Make sure that the message is present in the target queue.

AMI0049E Transport error.

Explanation: The underlying (MQSeries) message transport layer is reporting an error. You can obtain the message transport reason code by the secondary reason

code value returned from a GetLastError request for the specific AMI object. For more information, see "Common causes of problems" in *Application Messaging Interface*.

User response: Make sure that you have created the underlying MQ objects (system queues, user queues, among others) to which the Application Messaging Interface (AMI) service or policy is referring. Notice that the names of these objects are case sensitive. A queue named q1 is not the same as Q1. Also, ensure that the attributes of the queue are sufficient to handle the message. For example, check that MAXMSGL is big enough to handle the message size. If Publish or Subscribe functions fail, make sure that the MQ message broker is started.

AMI0109E User is not authorized.

Explanation: The user is not authorized by the underlying transport layer to perform the specified request.

User response: Make sure that the user of the UDFs has the appropriate authorization. This normally means that the user ID has to be a member of group mqm.

AMI0110E Transport is not available.

Explanation: The underlying transport layer is not available.

User response: Make sure that the queue manager is running.

AMI0402E Host file not found.

Explanation: A local host file with the specified name is not found.

User response: Make sure that file amthost.xml exists and in the right directory

AMI0405E The policy is not in the repository.

Explanation: You created a policy using a definition name that is not found in the repository. The policy is created using default values.

User response: Make sure that the Application Messaging Interface (AMI) policy is defined in the repository file.

AMI0406E The sender is not in the repository.

Explanation: You created a sender using a definition name that is not found in the repository. The sender is created using default values.

User response: Make sure that the Application Messaging Interface (AMI) sender service is defined in the repository file.

AMI0407E The receiver is not in the repository.

Explanation: You created a receiver using a definition name that is not found in the repository. The receiver is created using default values.

User response: Make sure that the Application Messaging Interface (AMI) receiver service is defined in the repository file.

AMI0409E The publisher is not in the repository.

Explanation: You created a publisher using a definition name that is not found in the repository. The publisher is created using default values.

User response: Make sure that the Application Messaging Interface (AMI) publisher service is defined in the repository file.

AMI0410E The subscriber is not in the repository.

Explanation: You created a subscriber using a definition name that is not found in the repository. The subscriber is created using default values.

User response: Make sure that the Application Messaging Interface (AMI) subscriber service is defined in the repository file.

AMI0416E Repository error.

Explanation: An error was returned when initializing or accessing the repository. This can occur for any of the following reasons:

- The repository XML file (for example, amt.xml) contains data that is not valid.
- The DTD file (*.dtd) is not found or contains data that is not valid.
- The files that are needed to initialize the repository can not be found. These file are usually in the directory intlFiles and the directory locales.

User response: This is generally the result of XML parser errors. Make sure that the right amt.dtd file is used and that it matches with the amt.xml file.

AMI0418E Repository not found.

Explanation: The repository file is not found.

User response: Make sure that the repository file is in the default location or in the path specified by environment variable AMT_DATA_PATH.

AMI0419E Transport library error.

Explanation: An error occurred loading the transport library.

User response: Make sure that MQSeries and the AMI libraries are installed correctly. There are sample programs shipped with MQ and AMI that can be used to verify the successful installation of this software.

AMI0424E DTD file not found.

Explanation: An AMI dtd file (amt.dtd) is not found with the XML repository file in the same directory.

User response: Make sure that the file amt.dtd is in the same directory as the repository file amt.xml. You can define the location with the environment variable AMT_DATA_PATH.

Part 4. Replication Messages

Chapter 35. ASN0000 - ASN0499

ASN0004E **CAPTURE** *capture_schema*. The Capture program could not start the trace. The return code is *return_code*. The reason code is *reason_code*.

Explanation: An error occurred when the START TRACE DB2 command was issued, or when the Capture program read the DB2 log.

User response: See the DB2 Codes section in the messages and codes documentation of the DB2 database manager on your operating system to find the appropriate reason code. For more information, see either of the following administration documentation: the Call Attachment Facility (CAF) for START TRACE DB2 errors, or the Instrumentation Facility Interface (IFI) for DB2 log read errors, or contact your DBA. If the CAF or the IFI returned a message, it is also printed on the system display console.

ASN0005E **CAPTURE** *capture_schema*. The Capture program encountered an error when reading the DB2 log. The log sequence number is *lsn*, the SQLCODE is *sql_return_code*, and the reason code is *reason_code*.

Explanation: An error occurred when the Capture program read the DB2 log. There might be an SQL error.

- For DB2 Replication, the *sqlcode* value is for the Asynchronous Read Log API.
- For Capture for VSE, the *sqlcode* is for the VSE/VSAM GET macro.
- For Capture for VM, the *sqlcode* is for Diagnose X'A4'.

User response: See the DB2 Codes section in the messages and codes documentation of the DB2 database manager on your operating system for the appropriate reason code, as suggested in the following list:

- For Capture program for z/OS, see the Instrumentation Facility Interface (IFI) section in the administration documentation of the DB2 database manager on your operating system, or contact your DBA.
- For Capture for VSE, see the VSE/VSAM Commands and Macros, VSE/ESA System Macro Reference, and VSE/ESA V2R3 Messages and Codes manuals for more information.
- For VM/ESA, see the VM/ESA Programming Services for more information.

- For Capture on Linux, Windows, and UNIX, see the active and archived database logs administration documentation for DB2, or contact IBM Software Support.

ASN0006E **CAPTURE** *capture_schema*. The Capture program encountered an unexpected log error of unknown log variation.

Explanation: An unexpected log error occurred when the Capture program was processing the DB2 log records, and was not reported by either of the following interfaces:

- The Instrumentation Facility Interface (IFI) for Capture program for z/OS
- The Asynchronous Read Log API for the Capture program

The Capture program could not determine which type of SQL update was associated with the log record.

User response: Contact IBM Software Support.

ASN0008I **CAPTURE** *capture_schema*. The Capture program was stopped.

Explanation: The Capture program has stopped.

User response: This message is for your information only, and no action is required.

ASN0009E **CAPTURE** *capture_schema*. The registered source table *src_owner.src_table* does not have the DATA CAPTURE CHANGES attribute.

Explanation: When attempting to initialize a registration, the Capture program encountered a source table that is incorrectly defined. The Capture program cannot process the log records associated with a source table if the DATA CAPTURE CHANGES attribute of the source table is not set. This message is issued during a Capture program warm start or when the first CAPSTART signal is received for a subscription set against this registration. If this error occurs during the processing of a CAPSTART signal, the Capture program does not activate the registration. If this error occurs during a reinitialization (from a warm start or a reinit capture command), the Capture program places the registration in the "stopped" state, indicating that you must repair the registration before the Apply program can resynchronize the associated subscription sets.

User response:

1. Alter the source table to turn on data capture changes. For example:
 - *alter table regress.table3 data capture changes*
2. If the registration has been deactivated by the Capture program (state = stopped), update the registration to set the state to inactive.
3. Use the Replication Center to force the Apply program to perform a full refresh for all subscription sets that replicate from this source table.

ASN0011E **CAPTURE** *capture_schema*. **The Capture program log read failed because the DB2 compression dictionary that was used to create the compressed log record no longer exists. The log record that could not be read was for the registered source table *table_owner.table_name*. The reason code is *reason_code*.**

Explanation: The Capture program received an error from the DB2 log read interface. The reason code is from DB2 and indicates that the data on a log record cannot be processed because the compression dictionary for the corresponding DB2 table or table space is permanently unavailable. The compressed table or table space that contains this source table was probably reorganized by the REORG utility that ran without the KEEPDICTIONARY option. The Capture program cannot read the remaining compressed log records from source changes that occurred before the reorganization. The Capture program follows the option that was specified for Stop Capture on Error (Replication Center) or STOP ON ERROR (ASNCLP command-line program) for this registration.

User response: Take one of the following actions, depending on whether the registration specifies that Capture stops on error:

No

No action is required. The Capture program deactivates the registration. The Apply program reactivates the registration and performs a full refresh of the target table.

Yes

Deactivate the registration and restart the Capture program. The Apply program reactivates the registration and performs a full refresh of the target table.

ASN0013E **CAPTURE** *capture_schema*. **The Capture program required a column that was not defined in the change-data (CD) table. The table name is *table_name*.**

Explanation: A required column in the CD table is not defined.

User response: Ensure that the CD table definition is correct.

ASN0019E **CAPTURE** *capture_schema*. **The Capture program did not start because its program libraries are not authorized for the Authorized Program Facility (APF).**

Explanation: The Capture program must be APF authorized.

User response: Authorize the Capture link library for APF, and restart the program.

ASN0020I **CAPTURE** *capture_schema*. **Netview Generic Alerts Interface failure. The Netview return code is *return_code*.**

Explanation: The Network Major Vector Transport (NMVT) could not be sent to Netview by the program because the program interface failed. This is a secondary informational message.

User response: See the Netview programming documentation for a description of the return code to determine the interface error. The Capture program alerts are not received by the System Services Control Point (SSCP) until the error is corrected.

ASN0021I **CAPTURE** *capture_schema*. **Netview Program to Program Interface unavailable. The Netview return code is *return_code*.**

Explanation: Netview is unavailable. This is a secondary informational message.

User response: See the Netview programming documentation for a description of the return code to determine the Netview problem. For example, the subsystem might not have been started.

ASN0023I **CAPTURE** *capture_schema*. **The Capture program has been reinitialized and is capturing changes for *number* registrations. Stopped *number* registrations are in a stopped state. Inactive *number* registrations are in an inactive state.**

Explanation: A REINIT command was issued to the Capture program. The Capture program then tried to refresh all the internal control information for all the registrations.

User response: If the Capture program is capturing changes for all the registrations, no action is required. Otherwise, examine the preceding error messages to determine the cause of the failure, and follow the suggested user response to repair the failing registration definition. After you repair the registration, re-issue the REINIT command to the Capture program.

ASN0028I **CAPTURE** *capture_schema*. **The Capture program is suspended by an operator command.**

Explanation: An operator command has suspended the Capture program, and the program has entered a wait state.

User response: This message is for your information only, and no action is required.

ASN0029I **CAPTURE** *capture_schema*. **The Capture program is resumed by an operator command.**

Explanation: An operator command has resumed the Capture program from a suspended state, and the Capture program has continued running.

User response: This message is for your information only, and no action is required.

ASN0031E **CAPTURE** *capture_schema*. **The program parameter table IBMSNAP_CAPPARMS can have only one row.**

Explanation: The Capture program parameters table was not defined correctly or has been updated with rows that are not valid.

User response: Make sure that there is only one row in the IBMSNAP_CAPPARMS table. Refer to the Table structures documentation in the *SQL Replication Guide and Reference* for additional information.

ASN0035E **CAPTURE** *capture_schema*. **A row with an unsupported architecture level was found in the table IBMSNAP_REGISTER. The row is not valid and specifies CD table *cd_owner.cd_table*, and the architecture level is *arch_level*.**

Explanation: The Capture program tried to initialize a registration and found that the registration definition contains an architecture level that is not valid. The Capture program can use only registrations that are at the compatible architecture level. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Verify that the value in the ARCH_LEVEL column in the IBMSNAP_REGISTER table at the source server is correct. If the architecture level is at a lower level than what is expected by the Capture program, migrate the Capture control tables to a level compatible with the Capture program.

ASN0049I **CAPTURE** *capture_schema*. **A row for the SIGNAL_SUBTYPE CAPSTOP was inserted into the table IBMSNAP_SIGNAL.**

Explanation: The Capture program received a signal to stop capturing data. The Capture program commits current work in progress and terminates.

User response: This message is for your information only, and no action is required.

ASN0055E **CAPTURE** *capture_schema*. **The Capture program encountered a column *column_name* with an unsupported SQLTYPE in table *table_name*.**

Explanation: The Capture program tried to initialize a registration and found that the registration definition contains an unsupported SQL type such as an abstract type. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Use the replication administration tools to change the registration so that it does not contain the unsupported SQL type, or drop the registration.

ASN0057E **CAPTURE** *capture_schema*. **The Capture program encountered error *errno* on operation for file *filename*.**

Explanation: An error occurred when the Capture program was handling files. The Capture program terminates.

User response: Ensure that the Capture program has correct access and security permissions for all required paths and files. Also, ensure that adequate space is available on your system. If you believe that this message was issued because of product failure, contact IBM Software Support for assistance.

ASN0058W **CAPTURE** *capture_schema*. **The MAP_ID *mapid* in a CAPSTART row in the IBMSNAP_SIGNAL table does not correspond to any entry in the IBMSNAP_PRUNCNTL table.**

Explanation: The value for the MAP_ID that is specified by the CAPSTART signal does not match any current value in the MAP_ID column of the IBMSNAP_PRUNCNTL table. The subscription set might have been deleted, or a user might have inserted the CAPSTART signal incorrectly.

User response: If this CAPSTART was issued by a user, check that the MAP_ID for the Signal table insert is correct, and try again. If this CAPSTART signal was issued by the Apply program, verify that the subscription set still exists.

ASN0059W CAPTURE *capture_schema*. The SYNCHPOINT field in the IBMSNAP_PRUNCNTL table is not zeros for the CAPSTART of subscription with MAP_ID *map_id*.

Explanation: When the Apply program signals a full refresh to the Capture program, the Apply program inserts a row for the CAPSTART signal in the IBMSNAP_SIGNAL table. At the same time, the SYNCHPOINT column of the IBMSNAP_PRUNCNTL table is set to hex zeroes. The Capture program then responds to the Apply program to confirm that the Capture program received the CAPSTART signal, as follows: the Capture program sets the value for the SYNCHPOINT column in the IBMSNAP_PRUNCNTL table to the number of the log sequence that corresponds to the CAPSTART log record. Because the Apply program set the value in the SYNCHPOINT column to hex zeroes, the Apply program checks if a nonzero value has been inserted by the Capture program. The Capture program updates the value for SYNCHPOINT, even if the value was not hex zeroes. However, if the value for SYNCHPOINT is not hex zeroes, the Capture program issues this warning that the value it found was not expected.

This warning can occur if you issue the APPLY CAPSTART signal yourself and do not completely simulate the actions of the Apply program.

User response: This message is for your information only, and no action is required.

ASN0060E CAPTURE *capture_schema*. The Capture program encountered an internal error *error_code*.

Explanation: An unexpected error occurred in the Capture program. The Capture program terminates.

User response: Contact IBM Software Support for assistance.

ASN0061E CAPTURE *capture_schema*. A registration that is not valid has been found. Source table *source_owner.source_table* does not exist in the system catalog tables.

Explanation: The Capture program tried to initialize a registration and found that the source table specified in the registration was not in the source system catalog. This message is issued during Capture warm start or when the Apply program issues the first CAPSTART signal for a subscription set that contains a registration that is not valid. This error does not cause the Capture program to terminate. The values for the columns *source_owner* and *source_table* in the Capture control table IBMSNAP_REGISTER might have been incorrectly specified, or the source table was dropped and no longer exists.

User response: If the registration is in error, correct

the values for the columns *source_owner* and *source_table*. If the source table no longer exists, then the registration is no longer valid and can be removed.

ASN0062E CAPTURE *capture_schema*. A registration that is not valid has been found. CD table *phys_change_owner.phys_change_table* does not exist in the system catalog tables.

Explanation: The Capture program tried to initialize a registration and found that the CD table specified in the registration was not in the source system catalog tables. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate. The values for the columns *phys_change_owner* and *phys_change_table* in the Capture control table IBMSNAP_REGISTER might have been incorrectly specified, or the CD table was dropped and no longer exists.

User response: If the registration is in error, correct the values in the columns *phys_change_owner* and *phys_change_table*. If the CD table no longer exists, then the registration is no longer valid and can be removed.

ASN0063E CAPTURE *capture_schema*. The source table *source_owner.source_table* associated with the subscription having MAP_ID *mapid* does not exist in the system catalog tables. The Capture program could not start capturing changes for this subscription.

Explanation: The Capture program tried to respond to a CAPSTART signal for a subscription and found that the source table which corresponds to the subscription was not in the source system catalog tables. This error message is issued when the first CAPSTART signal is received for a subscription that is not valid. This error does not cause the Capture program to terminate. The values for the columns *source_owner* and *source_table* in the Capture control table IBMSNAP_PRUNCNTL might have been incorrectly specified, or the source table was dropped and no longer exists.

User response: If the subscription is in error, correct the values for the columns *source_owner* and *source_table*. If the source table no longer exists, then the subscription is no longer valid and can be removed.

ASN0064E CAPTURE *capture_schema*. The registration is not valid for an associated subscription having MAP_ID *mapid*. The Capture program cannot start capturing change data for this subscription.

Explanation: The Capture program tried to initialize a registration associated with a particular subscription and found that the registration contains one or more

column values that are not valid. This message is issued when the first CAPSTART signal for a subscription is received against this registration. This error does not cause the Capture program to terminate. The values for the columns *phys_change_owner* and *phys_change_table* in the Capture control table IBMSNAP_REGISTER might have been incorrectly specified, or the CD table was dropped and no longer exists.

User response: If the registration is in error, correct the values for the columns *phys_change_owner* and *phys_change_table*. If the registration is no longer needed, you can remove it.

ASN0065E CAPTURE *capture_schema*. A registration that is not valid has been found. The source table *source_owner.source_table* is not a local physical table.

Explanation: The Capture program tried to initialize a registration and found that the source table for the registration is not a local physical table, but is instead on a non-DB2 relational server that is used as a source and called by a nickname. When a non-DB2 relational server is used as a source, data from each server is captured through a trigger program. Each source table must be in its own register table built on the non-DB2 relational server. This message is issued during a Capture warm start or when the first CAPSTART signal against this registration is received for a subscription. This error does not cause the Capture program to terminate. The registration has been created incorrectly in an IBMSNAP_REGISTER table in a DB2 database.

User response: This registration must be rebuilt and made valid at the correct non-DB2 relational server.

ASN0066E CAPTURE *capture_schema*. A registration that is not valid has been found. The CD table *phys_change_owner.phys_change_table* is not a local physical table.

Explanation: The Capture program tried to initialize a registration and found that the CD table for the registration that corresponds to the subscription is a nickname for a non-DB2 relational database used as a source. This message is issued when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Non-DB2 relational sources are captured through trigger programs, and must be in their own register table in the non DB2 relational source system. CCD tables for such sources are also created in the non DB2 relational source system. Somehow the non-DB2 relational source table registration has been incorrectly registered in an IBMSNAP_REGISTER control table. This registration must be rebuilt at the correct non-DB2 relational server.

ASN0067E CAPTURE *capture_schema*. The view registration associated with the subscription having MAP_ID *map_id* was not found in the table IBMSNAP_REGISTER. The Capture program could not start capturing change data for this subscription.

Explanation: The Capture program tried to initialize a registration and found that the view registration that corresponds to the subscription does not exist. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate. The values of the columns *source_owner*, *source_table*, and *source_view_qual* in the IBMSNAP_REGISTER or IBMSNAP_PRUNCNTL Capture control table could have been incorrectly specified. Therefore, either no match was found or the registration was dropped and no longer exists.

User response: If the subscription or the registration is in error, correct the values in the columns *source_owner*, *source_table*, and *source_view_qual*. If the registration no longer exists, then the subscription is no longer valid and can be removed.

ASN0068E CAPTURE *capture_schema*. The insert statement is too long for CD table *phys_chg_owner.phys_chg_tbl*.

Explanation: The number of columns in the CD table is too large; the SQL INSERT statement exceeds the 32K Capture coding limit.

User response: If all of the table columns are defined in the registration but only a subset of these columns are needed at the target, reduce the number of columns for that registration. Alternatively, split the table over two registrations so that each registration has a different subset of the table columns.

ASN0069E CAPTURE *capture_schema*. SQLCODE *sqlcode* was returned during an insert into the CD table *phys_chg_owner.phys_chg_tbl*. The CD table appears to have been dropped.

Explanation: The Capture program tried to insert a row into a CD table, and DB2 returned a SQLCODE indicating that the CD table no longer exists. The CD table could inadvertently have been dropped, or the whole registration could have been dropped. If there are still rows in the IBMSNAP_REGISTER table that refer to this CD table, the Capture program deactivates these registrations by setting the value of the CD_OLD_SYNCHPOINT column to NULL and no longer attempts to capture changes for this CD table. This error does not cause the Capture program to terminate.

User response: If the CD table no longer exists and is

no longer required, the registration is no longer valid and should be removed. It is recommended that you deactivate the registration before you remove it. Any subscription sets associated with registrations that use this CD table should also be deactivated. Additionally, the associated subscription-set members should be removed so that these subscription sets can be activated and can run successfully.

ASN0070E **CAPTURE** *capture_schema*. The combination of column name *column_name* in the CD table *phys_chg_owner.phys_chg_tbl*, and the value of BEFORE_IMG_PREFIX *before_img_prefix* in the table IBMSNAP_REGISTER for this registration matches multiple column names in the source table. The ambiguity in the registration definition cannot be resolved by the Capture program.

Explanation: The Capture program tried to initialize a registration and found that a column within the CD table for the registration is ambiguous. The column could refer to either a before image for one source column or to an after image for another source column. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: In the IBMSNAP_REGISTER table, change the current value in the BEFORE_IMG_PREFIX column to a character value that does not produce this ambiguity.

ASN0071E **CAPTURE** *capture_schema*. The data type attribute of the column *column_name* in the CD table *phys_chg_owner.phys_chg_tbl* is not compatible with the data type attribute of the corresponding source column.

Explanation: The Capture program tried to initialize a registration and found that a column within the CD table for the registration is not compatible with the corresponding source column. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Correct the CD table for this registration.

ASN0072E **CAPTURE** *capture_schema*. The before-image column *column_name* in the CD table *phys_chg_owner.phys_chg_tbl* must allow NULL values.

Explanation: The Capture program tried to initialize a

registration and found that the before-image column within the CD table for the registration was not defined to accept null values. This message is issued during Capture warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Correct the CD table for this registration.

ASN0073E **CAPTURE** *capture_schema*. The specification *input_in* describing the CD table on a CAPSTOP signal is not valid.

Explanation: The Capture program found that the INPUT_IN value specified on the CAPSTOP signal is not in a valid format of *phys_change_owner.phys_change_table*. This error does not cause the Capture program to terminate, and no action is taken for this signal.

User response: Ensure that the value of INPUT_IN matches the name of the CD table associated with the registration that you want to deactivate. Insert a new row into the IBMSNAP_SIGNAL table.

ASN0074E **CAPTURE** *capture_schema*. There is no row in the IBMSNAP_REGISTER table that corresponds to *src_owner.src_table* that is specified on a CAPSTOP signal.

Explanation: The Capture program found that the INPUT_IN value specified on the CAPSTOP signal is in a valid format, and there is no match for the value of *source_owner.source_table* in the registration table. This error does not cause the Capture program to terminate.

User response: Correct the value of INPUT_IN, and insert the signal again.

ASN0075W **CAPTURE** *capture_schema*. The registration corresponding to the INPUT_IN, *src_owner.src_table* on a CAPSTOP signal was not capturing changes. No action is taken.

Explanation: The Capture program found that the INPUT_IN value specified on the CAPSTOP signal is in a valid format and matches the value of a *source_owner.source_table* in the registration table, but this registration is already inactive. This error does not cause the Capture program to terminate, and the Capture program takes no action for the signal.

User response: This message is for your information only, and no action is required.

ASN0076I CAPTURE *capture_schema*. Capture has stopped capturing changes for source table *source_owner.source_table* in response to a CAPSTOP signal.

Explanation: The Capture program successfully deactivated a registration that was specified in a CAPSTOP signal.

User response: This message is for your information only, and no action is required.

ASN0077E CAPTURE *capture_schema*. The values specified for the PHYS_CHANGE_OWNER and PHYS_CHANGE_TABLE columns in the IBMSNAP_PRUNCNTL table where MAP_ID = *mapid* are NULL or do not match a valid row in the IBMSNAP_REGISTER table.

Explanation: The Capture program tried to initialize a registration and found that the column values of PHYS_CHANGE_OWNER and PHYS_CHANGE_TABLE within the IBMSNAP_PRUNCNTL table for the subscription do not match a registration row in the IBMSNAP_REGISTER table. The message was issued during Capture warm start or when the first CAPSTART signal was received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Correct the values in the IBMSNAP_PRUNCNTL table for this subscription. If this subscription was built using the Replication Center, contact IBM Software Support to report a potential administration problem.

ASN0078E CAPTURE *capture_schema*. The before-image column *column_name* in the CD table *phys_owner.phys_table* has no corresponding after-image column in the CD table for this registration. The registration is not valid.

Explanation: The Capture program tried to initialize a registration and found a before-image column within the CD table of the registration that has no corresponding after-image column. This message is issued during a Capture program warm start or when the first CAPSTART signal is received for a subscription against this registration. This error does not cause the Capture program to terminate.

User response: Correct the CD table for this registration by ensuring that before-image columns are included in the CD table only when the corresponding after-image column is also included.

ASN0079E CAPTURE *capture_schema*. SQLCODE *sqlcode* was returned during an update to the IBMSNAP_REGISTER table for the registrations associated with the CD table *phys_chg_owner.phys_chg_tbl*. The rows might have been deleted.

Explanation: The Capture program tried to update the IBMSNAP_REGISTER table to indicate that data has been captured for the named CD table, and DB2 returned a SQLCODE indicating that the rows no longer exist. The registrations could have been dropped. This error does not cause the Capture program to terminate.

User response: If the registration has been dropped, no further action is required for the registrations. When dropping registrations, it is recommended that you deactivate the registrations first. If the rows in the IBMSNAP_REGISTER table were inadvertently deleted, drop the associated CD table and rebuild the registrations. Deactivate any subscription sets that are associated with the registrations. If a registrations must be dropped, remove the associated subscription-set members so that these subscription sets can be activated and can run successfully. If the registrations are rebuilt, a signal is sent to the Apply program indicating that a full refresh should be performed for the associated subscription sets.

ASN0080E CAPTURE *capture_schema*. A table space full condition has been encountered for CD table *phys_chg_owner.phys_chg_tbl*, which is associated with the registration for source table *source_owner.source_table*.

Explanation: The Capture program tried to process an insert into the named CD table but was unable to process the insert due to a table space full condition. Typically this condition results from insufficient space allocation for CD table spaces, infrequent pruning, or ineffective pruning. This error causes the Capture program to terminate.

User response: Take the following steps to determine the cause of the table space full condition:

1. Ensure that sufficient space has been allocated to the table space for this CD table in order to accommodate normal processing conditions.
 2. Ensure that pruning is performed often enough to reduce the storage requirements for the Capture control tables.
 3. Ensure that the Apply programs are running often enough to accommodate normal pruning processing.
 4. Verify that no subscription sets have been deactivated for a long period of time without taking the additional steps necessary for normal pruning.
-

ASN0082W **CAPTURE** *capture_schema*. **The Capture program encountered a registration with a column *column_name* in the CD table *phys_chg_owner.phys_chg_tbl* with a column length *CD_column_length* that is shorter than the length of the corresponding column in the source table *source_owner.source_table*, with a length of *src_column_length*.**

Explanation: During the initialization of a registration, the Capture program found that the registration definition contains a column in the CD table with a column length that is shorter than the corresponding column length in the source table. The registration definition is allowed, but a warning message is issued to inform you that the captured source table data might not fit within the defined CD table column. This message is issued during a Capture program warm start or when the first CAPSTART signal is received for a subscription against this registration. The registration initializes successfully.

User response: Unless there is a specific reason why you need to define the registration in this manner (for example, if you are certain that the length of the changed data will never be larger than the length of the CD table column), you should define the registration so that the source table and the CD table data definitions match exactly.

ASN0083E **CAPTURE** *capture_schema*. **SQLCODE *sqlcode* was returned when trying to process an insert into the CD table *phys_chg_owner.phys_chg_tbl*. The CD table column *column_name* is too short and cannot contain the captured data from the corresponding column in the source table, *source_owner.source_table*. The registration has been stopped by the Capture program.**

Explanation: The Capture program tried to process an insert into a CD table and encountered a SQLCODE from DB2 that indicates that the CD table contains a column that is shorter than the length of the corresponding column in the source table. This error does not cause the Capture program to terminate, but the registration is placed in the stopped state.

User response: Re-evaluate this registration definition. Either alter the registration so that the lengths of the source table column and the CD table column match or add a trigger to the CD table to truncate the data.

ASN0084E **CAPTURE** *capture_schema*. **The registration with the source table *source_owner.source_table* and the CD table *phys_chg_owner.phys_chg_tbl* has been stopped by the Capture program.**

Explanation: This error message is issued any time

that a registration is placed in the stopped state (with the STATE column set to a value of 'S' in the IBMSNAP_REGISTER table) by the Capture program. The reason for this action is described in one or more of the preceding messages.

User response: Examine the preceding error messages to determine the cause of the failure, and follow the suggested user response to repair the failing registration definition. After you repair the registration definition, you must manually set the value of the STATE column to 'I' in the IBMSNAP_REGISTER table to indicate that the registration can be used again by the Apply program.

ASN0100I **CAPTURE** *capture_schema* : **The Capture version *version_number* program initialized successfully.**

Explanation: This message informs you that the Capture program started, and the version of the Capture program.

User response: This message is for your information only. No action is required.

ASN0101W **CAPTURE** *capture_schema*. **The Capture program warm start failed because existing data is too old. A cold start will be attempted.**

Explanation: The data in the change data tables is too old. A cold start is performed.

User response: See "Operating the Capture program for SQL Replication" in the DB2 Information Center for more details.

ASN0102W **CAPTURE** *capture_schema*. **The Capture program switches to cold start because the warm start information is insufficient.**

Explanation: A problem occurred during the retrieval of the restart information. The restart table data is not valid. A cold start is performed.

- For DB2, an Asynchronous Read Log API error occurred during warm start while DB2 was reading the log.
- For z/OS, an Instrumentation Facility Information (IFI) error occurred during warm start while DB2 was reading the log.

User response: See "Operating the Capture program for SQL Replication" in the DB2 Information Center for more details.

ASN0104I **CAPTURE** *capture_schema*. **In response to a CAPSTART signal with MAP_ID *mapid*, change capture has started for the source table *source_owner.source_table* for changes found on the log beginning with log sequence number *log_sequence_number*.**

Explanation: The Capture program successfully processed a CAPSTART signal. If this is the first CAPSTART signal associated with a particular source table, this message indicates that the Capture program is now capturing updates to the source table. If this is a CAPSTART signal for a table for which changes are already being captured, this message indicates that the Capture program received the signal and performed the required processing to allow the Apply program to start receiving changes for the subscription set that is associated with the input MAP_ID value.

User response: This message is for your information only, and no action is required.

ASN0105I **CAPTURE** *capture_schema*. ***n* rows have been pruned from the table *table_owner.table_name* at *timestamp*.**

Explanation: The Capture program pruned records from a CD, UOW, TRACE, MONITOR, or SIGNAL table.

User response: This message is for your information only, and no action is required.

ASN0109I **CAPTURE** *capture_schema*. **The Capture program has successfully initialized and is capturing data changes for *number* registrations. *Stopped_number* registrations are in a stopped state. *Inactive_number* registrations are in an inactive state.**

Explanation: This message is issued when the Capture program completes the reinitialization of registration entries. The reinitialization can occur during a warm start, during the processing of a CAPSTART signal, or in response to a Capture REINIT command.

User response: If the Capture program is capturing the changes for all the registrations, no action is required. Otherwise, examine the preceding error messages to determine the cause of the failure, and follow the suggested user responses to repair the failing registration definition. After you have repaired the registration definition, issue the asncmd command with the reinit parameter.

ASN0111I **CAPTURE** *capture_schema*. **The pruning cycle started at *timestamp*.**

Explanation: This message is issued at the beginning of each pruning cycle.

User response: This message is for your information only, and no action is required.

ASN0112I **CAPTURE** *capture_schema*. **The pruning cycle ended at *timestamp*.**

Explanation: This message is issued at the termination of each pruning cycle.

User response: This message is for your information only, and no action is required.

ASN0113W **CAPTURE** *capture_schema*. **The pruning retention limit will be reached in the next 24 hours.**

Explanation: This warning message is issued when the values in the IBMSNAP_UOW table indicate that retention limit pruning could occur in the next day.

User response: Check to see why regular pruning is not occurring. Usually this is because one or more Apply programs have not been run for a period of many days, and therefore the CD and UOW tables cannot be effectively pruned. Another potential hazard is the removal or deactivation of a subscription set, without the removal or reset of the corresponding synchpoint value in the table IBMSNAP_PRUNE_SET. The Replication Analyzer tool can be used to provide a detailed analysis of the situation.

ASN0114E **CAPTURE** *capture_schema*. **Pruning has failed with SQL code *sqlcode* when pruning the table *table_owner.table_name*.**

Explanation: This error message is issued when pruning fails with an unexpected SQL error code. Pruning terminates and tries again at the next interval or command invocation. This error does not cause the Capture program to terminate.

User response: If this SQL code indicates a temporary error, then no action is required. Otherwise, take action as indicated for the SQL error in the "ASN - Replication" messages section of the DB2 Information Center.

ASN0121E **CAPTURE** *capture_schema*. **The Capture program warm start failed because existing data is too old. The Capture program will terminate.**

Explanation: The lag_limit parameter represents the number of minutes that the Capture program can lag in processing records from the DB2 log. Capture cannot warm start because it will process transactions older than is allowed by the user specified lag_limit.

User response: Determine why the Capture program is behind in reading the log. If you are in a test environment, where you have no practical use for the lag limit parameter, you might want to set the lag limit higher and try starting the Capture program again. Alternatively, if you have very little data in the source table in your test environment, you might want to cold start the Capture program and fully refresh the data in all the target tables.

ASN0122E CAPTURE *capture_schema*. An error occurred while reading the restart information or DB2 log. The Capture program will terminate.

Explanation: A problem occurred while retrieving the restart information. The restart table data was not valid or for z/OS, an Instrumentation Facility Interface (IFI) error occurred while reading the log during a restart. When the error is resolved, you can restart using the warm start option.

User response: Determine why the restart information could not be retrieved by Capture from IBMSNAP_RESTART. Look at previous messages for clues as to what may have gone wrong. Correct the problem, if possible and restart Capture warm. If the problem cannot be corrected, restart Capture cold.

ASN0123I CAPTURE *capture_schema*. At program termination, the highest log sequence number of a successfully captured log record is *max_commitseq* and the lowest log sequence number of a record still to be committed is *min_inflightseq*.

Explanation: The Capture program terminates and records the values of the restart table at that time for auditing purposes.

User response: This message is for your information only, and no action is required.

ASN0133I CAPTURE *capture_schema*. The Capture program has reached the end of the active log and will terminate because the AUTOSTOP feature is specified.

Explanation: The Capture program terminated when it reached the end of the active log as requested by the user option of AUTOSTOP.

User response: This message is for your information only, and no action is required.

ASN0142E CAPTURE *capture_schema*. The Capture program is unable to perform an insert operation on the monitor table IBMSNAP_CAPMON. The SQL code is *sqlcode*. The monitoring information for this interval will be skipped.

Explanation: This error message is issued when the monitoring thread has failed with an unexpected SQL code. Monitor functions for this interval are skipped, and the program tries again at the next interval. This error does not cause the Capture program to terminate.

User response: If this SQL code indicates a temporary error, then no action is required. Otherwise, take action as indicated for the SQL error in the "ASN - Replication" messages section of the DB2 Information Center.

ASN0143W CAPTURE *capture_schema*. The program detected that the source database *src_db_name* has been restored or rolled forward. The Capture program has switched from a warm start to a cold start.

Explanation: The Capture program started with a startmode of warmsa or warmsi. When the Capture program attempted to warm start, it received a return code from the DB2 log read API that indicates that the source database has been restored or rolled forward and that log sequence numbers have been reused; the state of the source database and the state of the captured data are no longer consistent. The Capture program switched to a cold start.

User response: This message is for your information only, and no action is required.

ASN0144E CAPTURE *capture_schema*. The program detected that the source database *src_db_name* has been restored or rolled forward. A cold start is recommended to restore consistency.

Explanation: The Capture program started with a startmode of warmns or warmsi. When the Capture program attempted to warm start, it received a return code from the DB2 log read API that indicates that the source database has been restored or rolled forward and that log sequence numbers have been reused; the state of the source database and the state of the captured data are no longer consistent. The Capture program terminates and does not automatically switch to a cold start.

User response: If you are certain that it is safe to perform a Capture program warm start, restart the Capture program; it will not terminate on a second attempt. If you are not certain whether the captured data will be in a consistent state after a Capture program warm start, it is recommended that you perform a Capture program cold start.

ASN0180W CAPTURE *capture_schema*. The table IBMSNAP_SIGNAL is not an EBCDIC table as required by capture. The signal has been processed.

Explanation: The Capture program detected that the IBMSNAP_SIGNAL table is not defined as an EBCDIC table. Additional processing is required to translate signals to EBCDIC to process them properly. The additional processing requires a small performance degradation.

User response: At your earliest convenience, perform the following steps:

1. Stop the Capture program.
2. Drop and re-create the IBMSNAP_SIGNAL table with EBCDIC encoding.
3. Restart the Capture program.

ASN0181W CAPTURE *capture_schema*. The row for the signal with timestamp *signal_time* no longer exists in the IBMSNAP_SIGNAL table. The signal has been processed.

Explanation: The Capture program processed the request from the signal but could not update the SIGNAL_STATE and SIGNAL_LSN. Therefore, the issuer of the signal cannot determine that the Capture program received the signal.

User response: Determine if another process is expecting the update for the signal from the Capture program, and if necessary, re-send the signal.

ASN0182W CAPTURE *capture_schema*. The row for signal with timestamp *signal_time* no longer exists in IBMSNAP_SIGNAL table and the table is not EBCDIC. The signal will be ignored by capture.

Explanation: An initialization failure occurred, because the Capture program received a signal that was not encoded in EBCDIC. The Capture program could not translate the signal to EBCDIC, because the row in the IBMSNAP_SIGNAL table no longer exists. The Capture program cannot determine what signal was sent and so ignores it.

User response: Determine what signal was sent, and re-send the signal.

At your earliest convenience, perform the following steps:

1. Stop the Capture program.
2. Drop and re-create the IBMSNAP_SIGNAL table with EBCDIC encoding.
3. Restart the Capture program.

ASN0183E CAPTURE *capture_schema*. The Capture program detected an inconsistency between the IBMSNAP_PARTITIONINFO table and DB2 partition information.

Explanation: This error message occurred due to one of the following reasons:

- A new database partition was added to the database.
- The IBMSNAP_PARTITIONINFO control table is corrupted.

User response: If a new partition was added, restart the Capture program with the add_partition=Y option.

If the IBMSNAP_PARTITIONINFO control table is corrupted, cold start the Capture program or call IBM Software Support.

ASN0184I CAPTURE *capture_schema*. The Capture program started on partition *partition_num* for changes found in the log beginning with log sequence number *log_sequence_number*.

Explanation: A new partition was added using the ADD_PARTITION option. The Capture program started processing for that partition.

User response: This message is for your information only, and no action is required.

ASN0185I CAPTURE *capture_schema*. The Capture program did not find a required control table so it created the control table *control_table* and its associated index *index* in the Capture control server.

Explanation: The Capture program requires the specified control table. If the Capture program does not find the control table, it creates the control table and its associated index on the appropriate Capture control server.

User response: This message is for your information only, and no action is required.

ASN0186W CAPTURE *capture_schema*. The Capture program cannot find the source database *database* on partition *partition_ID*. The Capture program cannot process the log for this partition.

Explanation: This partition is not known to the source database. The Capture program captures data only from the partitions that are known to the source database.

User response: Add the partition to the database and restart the Capture program using the ADD_PARTITION=Y option. If the partition is not needed, remove it.

ASN0187W **CAPTURE** *capture_schema*. **The Capture program cannot find partition**
partition_id **previously known in DB2.**

Explanation: The Capture program cannot find a previously known partition in DB2. The partition is no longer part of the multipartition source server.

User response: Verify that the missing partition was properly removed. No further action is required.

ASN0188I **CAPTURE** *capture_schema*. ***n* rows have been pruned from the table**
src_owner.table **at timestamp due to retention limit pruning.**

Explanation: The Capture program pruned records from a CD, UOW, TRACE, MONITOR, or SIGNAL table due to retention limit pruning.

User response: This message is for your information only, and no action is required.

ASN0189I **CAPTURE** *capture_schema*. **The Capture program is terminating because an error occurred for registered table**
src_owner.table **and the registration is configured to stop on error.**

Explanation: The registration has STOP_ON_ERROR = Y. The Capture program is terminating due to an error that is identified in a previous message.

User response: Take one or both of the following actions and then restart the Capture program:

- Correct the error identified in the previous message.
- In the register (IBMSNAP_REGISTER) table, set STOP_ON_ERROR=N.

ASN0190I **CAPTURE** *capture_schema*. **The registration for source table**
src_owner.table **was deactivated due to an error. The Capture program did not terminate.**

Explanation: The Capture program deactivated the registration for the specified source table due to an error that is identified in a previous message. The Capture program did not terminate; it is still running because the registration has STOP_ON_ERROR = N.

User response: Correct the error for this source table that is identified in the STATE column of the register (IBMSNAP_REGISTER) table and then reactivate the registration. When you reactivate the registration, a full-refresh is performed.

To avoid a full-refresh in the future, set STOP_ON_ERROR to Y so that the Capture program terminates if it encounters the error.

ASN0191E **Capture** *capture_schema* : **Log record with LSN** *LSN* **could not be processed because it is an unknown log variation. Its DBID is** *DBID*. **Its OBID is** *OBID*. **The transaction ID is** *ID*.

Explanation: The Capture program could not process a log record because the log record has an unknown log variation type.

User response: If Capture stopped because of this error, you will need to disable the registration or subscription with the matching DBID and OBID so that Capture can ignore this log record. If Capture did not stop because of this error, make sure this log record does not have data that you need replicated.

ASN0192E **Capture** *capture_schema* : **The Capture program is unable to decode log records for the source table**
table_owner.table_name. **The program stopped.**

Explanation: The Capture program encountered a different version of the table because the table was altered. The Capture program was not able to determine the correct version by reading from the IBMQREP_COLVERSION and IBMQREP_TABVERSION tables. The tables might have been dropped.

User response: Ensure that the IBMQREP_COLVERSION and IBMQREP_TABVERSION tables have not been dropped. If they are missing, restore these tables and restart the Capture program.

ASN0193W **CAPTURE** : *capture_schema* : **The log reader thread is still initializing. The program continues to wait.**

Explanation: During initialization, the Q Capture or Capture program tries to read from a restart point (either a specific Log Sequence Number when the Q Capture or Capture program is started in warm mode, or the current active LSN for cold starts). The initialization might take excessive time if the recovery log or the database is unavailable. In a DB2 multiple-partitioned environment or Oracle RAC environment, the log reader thread also takes longer to initialize.

User response: Continue waiting for the log reader to initialize.

ASN0194W **CAPTURE : *capture_schema* : The Capture program detected that a partition was attached to table *table_owner.table_name*. Capture had not read log records up to the point that corresponds to the time that the registration was activated. Capture might have incorrectly captured rows for the table before it was attached as a partition.**

Explanation: When capture activates a registration for a table, it gathers table information from the system catalogs. If a partition was attached to the table before Capture activates the registration, the system catalog information might not correctly reflect where capture is in the log. For a newly attached partition, capture might incorrectly capture rows from the table before it was attached as a partition.

User response: Stop capturing for the registration and perform a full refresh for the subscription set.

ASN0195I **Capture *capture_schema* : The transaction with ID *transaction_identifier* was ignored as requested.**

Explanation: The Capture program did not capture the transaction because the transaction ID was specified in the IGNORE_TRANSID parameter when Capture was started, or information for ignoring the transaction was inserted into the IBMQREP_IGNTRAN table.

User response: This message is for your information only. No action is required.

ASN0196W **Capture : *capture_schema* : The contents of the registered source table *table_owner.table_name* were altered by the DB2 utility program *utility_program*.**

Explanation: The source table was altered by one of the following DB2 online utility programs:

- LOAD SHRLEVEL NONE RESUME YES
- LOAD SHRLEVEL NONE REPLACE
- REORG TABLESPACE DISCARD
- CHECK DATA DELETE YES LOG NO
- RECOVER PIT
- EXCHANGE DATA ON CLONE TABLESPACE

The changes to the source table are not automatically replicated to the target table, and the two tables could be out of synch.

User response: To maintain synchronization between the source and target, take one of the following actions:

- Perform a full refresh of the target table.
- Use the asntdiff program to identify any differences between the source and target tables, and then use the asntrep program to repair the differences.

ASN0197W **Capture *capture_schema* : The registered source table *table_name* was dropped.**

Explanation: The Capture program detected that the table space for the source table was dropped, and therefore the table itself was dropped.

User response: If you no longer want to replicate data from the source table, use the replication administration tools to drop the registration. If you plan to recover the table, reinitialize the Capture program after the table is recovered, or stop and start the Capture program in warm mode.

ASN0198E **Capture *capture_schema*: The number of database partitions exceeds the maximum number of partitions that SQL Replication supports. The Capture program will stop.**

Explanation: The Capture program supports up to 16 partitions in a multiple-partitioned environment.

User response: Reduce the number of partitions to 16 or fewer and start the Capture program.

ASN0199E **Capture *capture_schema* : The Capture program could not activate the registration for table *table_owner.table_name*. The table was altered. The table version is *version*.**

Explanation: When a table is altered, it must be reorganized before the Capture program can begin processing log records for it. For the specified table, one of the following situations is likely:

- The table was altered but never reorganized.
- The table was altered after its most recent reorganization.

User response: Reorganize the table and activate the registration.

ASN0200E **CAPTURE *capture_schema*. The Capture program log read failed because the DB2 compression dictionary that was used to create the compressed log record is temporarily unavailable. The log record that could not be read was for the registered source table *table_owner.table_name*. The reason code is *reason_code*.**

Explanation: The Capture program received an error from the DB2 log read. The error indicates that the data on a log record cannot be processed because the compression dictionary for the corresponding DB2 table or table space is temporarily unavailable. For z/OS, the reason code is a z/OS diagnostic code. For Linux, UNIX, and Windows, the reason code is an SQL code. One cause for the problem could be that the compressed table space is in the STOPPED state when

the DB2 Log Read Interface attempts to read the compression dictionary. DB2 takes a latch on the source compressed table space to access the dictionary and the latch does not work if the table space is stopped. This error prompts the Capture program to stop.

User response: The problem will be resolved when the compression dictionary becomes available. Restart the Capture program.

ASN0201W **CAPTURE : *capture_schema* : The Capture program detected that a partition was detached from table *table_owner.table_name*. Capture was behind in the log when the registration for this table was activated. Capture might have incorrectly missed rows from the partition before it was detached from the table.**

Explanation: When Capture activates a registration for a table, it gathers table information from the system catalogs. If a partition was detached from the table before Capture activates the registration, the system catalog information might not correctly reflect where Capture is in the log. For a newly detached partition, capture might incorrectly miss rows from the partition before it was detached.

User response: Stop capturing for the registration and perform a full refresh on the subscription set.

ASN0202W **CAPTURE : *capture_schema* : An ALTER TABLE ALTER COLUMN statement was detected for column *column_name* of table *table_name* with a new data type of *data_type*. Column *column_name* of CD table *table_name* was altered to a data type of *data_type*. The target table is not altered automatically.**

Explanation: The Capture program detected that a column in a registered source table was altered. Capture automatically altered the CD table column to match the new data type.

User response: If you want to replicate data from the altered column, you must alter the matching column in the target table.

ASN0203E **CAPTURE *capture_schema* : The value *value* for the COMPATIBILITY column in the IBMSNAP_CAPPARMS table is not supported for a DB2 pureScale system at version *version* with two or more members.**

Explanation: A DB2 pureScale system at Version 10 with two or more members requires a value of 1001 or higher in the COMPATIBILITY column of the IBMSNAP_CAPPARMS table.

User response: Before upgrading the compatibility

setting for the Capture program, review the effect of that change on any Apply programs or Apply control servers that Capture works with. You might need to migrate the Apply control tables to Version 10. If required, update the COMPATIBILITY column to 1001 or higher and restart the Capture program.

ASN0204E **CAPTURE : *capture_schema* : The log sequence number value *LSN* in the IBMSNAP_RESTART table is not expected. The Capture program will stop.**

Explanation: The LSN value that Capture uses to find the correct restart point in the DB2 recovery log is not in the correct format. This problem can occur when you upgrade DB2 to Version 10 but the IBMSNAP_RESTART table still contains LSN values from a previous version. The LSN values that replication uses changed for DB2 for Linux, UNIX, and Windows V10 and are not compatible with earlier saved LSN values.

User response: Restart the Capture program with the `migrate=y` parameter, which updates the restart information to the correct format.

ASN0205E **CAPTURE *capture_schema*. The connection that the Capture program uses to read DB2 log records was terminated because the DB2 subsystem has terminated or is in the process of terminating. The log sequence number is *lsn*. The return code is *return_code*. The reason code is *reason_code*.**

Explanation: The Capture program connection with DB2 terminated while the Capture program was reading the DB2 log.

User response: Look up the reason code from the DB2 for z/OS instrumentation facility interface (IFI) to verify why the connection was terminated. Fix any problems and restart the Capture program if necessary.

Chapter 36. ASN0500 - ASN0999

ASN0500E *pgmname : program_qualifier : The parameter input input_value supplied for parameter name parameter_name is not valid.*

Explanation: The program or a command program has been invoked with a specified input parameter that is not valid. The message indicates the name of the program that is reporting the error, along with the parameter name and the parameter value.

User response: Check the documentation on valid invocation parameters, correct the input, and resubmit the task or command.

ASN0501E *pgmname : program_qualifier : The value input_value supplied for the parameter parameter_name is not the correct data type.*

Explanation: The program or a command program was invoked with an input value with an associated data type that is not valid. The message indicates the name of the program that is reporting the error, the incorrect input value, and the name of the parameter for which this input value was specified.

User response: Correct the invocation to include the correct data type for the parameter input and resubmit it.

ASN0502E *pgmname : program_qualifier : The value input_value of length invalid_string_length, supplied for parameter parameter_name, is greater than the maximum allowed string length of allowed_string_length.*

Explanation: The program or a command program has been invoked using an input value with a string length that is not valid. The message indicates the name of the program that is reporting the error, what input value is incorrect, and for which parameter this input value was specified.

User response: Correct the invocation to include the correct string length for the parameter input and resubmit it.

ASN0503E *pgmname : program_qualifier : The integer value input_value, supplied for parameter parameter_name, is outside the supported range for this parameter.*

Explanation: The program or a command program was invoked with an input value specified which is outside the supported range. The message indicates the name of the program that is reporting the error, which

input value is incorrect, and for which parameter this input value was specified.

User response: Correct the invocation to include the correct range value for the parameter input and resubmit it.

ASN0504E *pgmname : program_qualifier : The program did not recognize the invocation parameter incorrect_input.*

Explanation: The program or a command program has been invoked with a specified parameter or command that is not valid. The message indicates which program issued this message, and the invocation input that is unrecognized.

User response: Check the documentation on valid input parameters, correct the input and resubmit the task or command.

ASN0505E *pgmname : program_qualifier : The program was unable to get or set an IPC key.*

Explanation: The program or a command program was unable to initialize the inter-process communications needed to process commands. This error causes the failing program to terminate.

User response: Retry the failing program or command. Contact IBM Software Support if the problem persists.

ASN0506E *program_name : program_identifier : The command was not processed because the target replication program program_name was not running or because you entered an incorrect parameter value (for example, a server name, schema, or qualifier might be misspelled).*

Explanation: This problem could also have occurred because of an issue with the IPC message queue that the program uses to receive commands.

User response: Verify that the parameter values are correct and that they identify a running replication program. Retry the command.

ASN0507E *pgmname : program_qualifier : The program could not create the replication communications message queue. The program did not terminate for this failure, but the command was not executed.*

Explanation: The program or a command program encountered an internal error while trying to process a user command.

User response: Retry the failing command. See "Troubleshooting problems with IPC message queues" for more information.

ASN0508E *pgmname : program_qualifier : The program could not send a message to the replication communications message queue.*

Explanation: The program or a command program encountered an internal error while trying to process a user command. The program did not terminate for this failure, but the command did not get executed.

User response: Retry the failing command. Contact IBM Software Support if the problem persists.

ASN0509E *pgmname : program_qualifier : The program could not process a received message because of an incorrect message version.*

Explanation: The program or a command program encountered an internal error while trying to process a user command. The program did not terminate for this failure, but the command did not get executed.

User response: Retry the failing command. Contact IBM Software Support if the problem persists.

ASN0510E *pgmname : program_qualifier : The program encountered a timeout while waiting for reply messages.*

Explanation: The command program encountered an internal error while trying to process a user command. The program did not terminate for this failure, but the command did not get executed.

User response: Retry the failing command. Contact IBM Software Support if the problem persists.

ASN0511E *pgmname : program_qualifier : The program was unable to process the received message because of an unknown message function.*

Explanation: The program encountered an internal error while trying to process a user command. The program did not terminate for this failure, but the command did not get executed.

User response: Retry the failing command. Contact IBM Software Support if the problem persists.

ASN0512E *pgmname : program_qualifier : The program could not read from its replication communications message queue.*

Explanation: The program encountered an internal error while trying to process a user command. The program did not terminate for this failure, but the command did not get executed.

User response: Retry the failing command. Contact IBM Software Support if the problem persists.

ASN0513E *pgmname : program_qualifier : The program could not open the message file named msg_file.*

Explanation: This message file used by the program has been installed incorrectly, or the language environment variables are not set correctly.

User response: Refer to the documentation for information about installation and configuration.

ASN0514E *pgmname : program_qualifier : The program could not open the log file log_file.*

Explanation: The program encountered an internal error while trying to open a file for its own program message log, and terminates abnormally because of this failure. This problem might have occurred because the file was inadvertently deleted, or because the userid associated with this process does not have the sufficient authority to open the file.

User response: Verify that sufficient authority is provided to the processing userid. If the file was inadvertently deleted, restart the program to create a new log file.

ASN0515E *pgmname : program_qualifier : The program could not close the log file.*

Explanation: The program encountered an internal error while trying to close the file used for its own program message log. The file might have been deleted inadvertently before the program tried to terminate. Final termination messages might not be issued.

User response: If the file was inadvertently deleted, restart the program to create a new log file.

ASN0516E *pgmname : program_qualifier : The program could not close the message catalog.*

Explanation: The program encountered an internal error while trying to close the message catalog file. The file might have been deleted inadvertently before the program tried to terminate. Final termination messages might not be issued.

User response: If the message file has been deleted, it needs to be reinstalled.

ASN0517E *pgmname : program_qualifier : The program has recovered the ability to read from its replication communications message queue.*

Explanation: The program was able to reinitialize its read capability from the message queue needed to process commands after a previous failure.

User response: This message is for your information only, and no action is required.

ASN0518E *pgmname : program_qualifier : The program does not accept multiple commands.*

Explanation: The command program was invoked with multiple commands specified. Each command invocation must be performed with a single input command, along with any other required command input. Note: The CHGPARMs command allows multiple parameters to be changed with one invocation of the CHGPARMs command.

User response: Correct the command input, and resubmit the command.

ASN0519E *pgmname : program_qualifier : The parameter input parameter_value supplied for CHGPARMs parameter parameter_name is not valid.*

Explanation: The CHGPARMs command was invoked with incorrect parameter input.

User response: Correct the command input and resubmit the command.

ASN0520I *pgmname : program_qualifier : The STATUS command response: thread_type thread is in the status_condition state.*

Explanation: In response to the status command, one of these messages will be issued for each of the threads associated with the program that received the command, in each case providing the current state of that thread.

User response: This message is for your information only, and no action is required.

ASN0521I *pgmname : program_qualifier : The QRYPARMS command response: parameter_name was set to parameter_value by the following method: method.*

Explanation: In response to the QRYPARMS command, a message will be issued for each of the program parameters. For each parameter, the message

provides the name of the parameter, the current setting of the parameter, and the method (by default, by changing the IBMSNAP_CAPPARMs table, by the startup option, or by the use of the CHGPARMs command) that was employed by the user to set the value of the parameter.

User response: This message is for your information only, and no action is required.

ASN0522I *pgmname : program_qualifier : The program received the command_type command.*

Explanation: The program received a command to be processed.

User response: This message is for your information only, and no action is required.

ASN0523I *pgmname : program_qualifier : The CHGPARMs command response: parameter_name has been set to parameter_value.*

Explanation: In response to the CHGPARMs command, one of these messages will be issued for each of the program parameters that was changed. For each parameter, the message provides the new value for the parameter.

User response: This message is for your information only, and no action is required.

ASN0524E *pgmname : program_qualifier : The program required parameter parameter_name was not specified.*

Explanation: The program or a command program was invoked without one of the required parameters specified. If the parameter missing is the *capture_server* or *control_server*, the program or command also tried accessing the database name implicitly through the DB2DBDFT environment variable setting, where applicable, and this also was not successful.

User response: Correct the invocation to include the appropriate parameter and its corresponding input value.

ASN0525E *pgmname : program_qualifier : The program could not read from its external communications message queue.*

Explanation: The program or a command program was unable to initialize its read capability from the external communications message queue needed to process commands.

User response: Retry the failing command, and if the problem persists, contact IBM Software Support.

ASN0526E *pgmname : program_qualifier : The program was invoked without any command input.*

Explanation: The command program was invoked without a command to process. No command processing is performed.

User response: Resubmit the command with all the required input.

ASN0527E *pgmname : program_qualifier : The program was invoked without any CHGPARMs command input.*

Explanation: The commands program was invoked with the CHGPARMs command but without any command input to process. No command processing is performed.

User response: Resubmit the command with all the required input.

ASN0528E *pgmname : program_qualifier : The program will terminate because the required control table tableowner.tablename does not exist.*

Explanation: The Capture or the Apply program tried to execute an SQL operation against a required Capture control table. The program received a *not found* return code from DB2. The return code occurs either if the migration has not been completed or if a required Capture control table has been accidentally dropped from the environment.

User response: See the message text for the name of the missing control table. Corrective action for this problem depends on which table is missing. For example, if the table is IBMSNAP_PRUNE_LOCK, then the table can simply be re-created, and the Capture program can be restarted. However, if the table is IBMSNAP_RESTART, and if the correct table contents cannot be restored, then the table needs to be re-created, and the Capture program requires a cold start.

ASN0529I *pgmname : program_qualifier : The value of parameter_name was set to parameter_value at startup by the following method: method.*

Explanation: The program started, and the program parameters were initialized based on the combination of startup options that are specified and the existing contents of the parameters table. The parameters were set by one of the methods: by default, by changing the parameters table, or by the startup option.

User response: This message is for your information only, and no action is required.

ASN0530E *pgmname : program_qualifier : The program could not connect to database database_name with USERID user_id . The SQLCODE is sql_code.*

Explanation: An error occurred when the program issued one of the following functions:

- A CONNECT function to DB2 for VSE and VM
- A CONNECT function to DB2 Call Attachment Facility (CAF)
- An implicit connect to DB2

User response: See the DB2 codes in the messages and codes publication of the DB2 database manager on your operating system for the appropriate reason code.

For DB2 for z/OS errors, see the section in the administration guide that describes the Call Attachment Facility. Contact your DBA for questions and diagnosis.

If you are running the program under DB2 for Linux or UNIX, ensure that the LIBPATH environment variable is set to the same environment in which the program starts.

ASN0531E *pgmname : program_qualifier : The program could not open the plan. The SQL return code is return_code, the reason code is reason_code, the subsystem name is DB2_subsystem, and the plan name is plan_name.*

Explanation: An error occurred when the program tried to open the plan, ASNLPLAN.

User response: See the DB2 Codes section in the messages and codes publication of the DB2 database manager on your operating system to find the appropriate reason code. See the section in the administration guide that describes the Call Attachment Facility.

ASN0532E *pgmname : program_qualifier : DB2 release release_number is not supported.*

Explanation: The program does not support this release of DB2.

User response: Run the program with the appropriate release of DB2.

ASN0533E *pgmname : program_qualifier : DB2 was terminated abnormally.*

Explanation: DB2 was terminated while the program was still active.

For z/OS, VSE/ESA, or VM/ESA, DB2 was terminated while program was active and the user did not specify the NOTERM invocation parameter.

User response: Start DB2 and start the program.

ASN0534E *pgmname : program_qualifier : DB2 database cannot be used, because it is in the state state.*

Explanation: DB2 was terminated while the program was still active. The database is in one of the following states: UNDETERMINED, TERMINATED, QUIESCED, ROLLWARD, or ACTIVE.

User response: Start DB2, and then start the program.

ASN0535E *pgmname : program_qualifier : The program could not disconnect from the database db_server. The return code is return_code, and the reason code is reason_code.*

Explanation: While terminating the connection to DB2, the program received an error code from the Call Attachment Facility (CAF).

User response: Restart the program.

ASN0536E *pgmname : program_qualifier : An error was returned while getting the instance name. The SQLCODE is sqlcode.*

Explanation: The SQLEGINs API of DB2 returned an error.

User response: See the DB2 API Reference for information about the SQLEGINs API to determine the error, or contact IBM Software Support.

ASN0537E *pgmname : program_qualifier : The program could not connect to database database_name, the return code is return_code, and the reason code is reason_code.*

Explanation: An error occurred when the program issued one of the following functions:

- A CONNECT function to DB2 for VSE and VM
- A CONNECT function to DB2 Call Attachment Facility (CAF)
- An implicit connect to DB2

User response: See the DB2 codes in the messages and codes publication of the DB2 database manager on your operating system for the appropriate reason code.

For DB2 for z/OS errors, see the section in the administration guide that describes the Call Attachment Facility. Contact your DBA for questions and diagnosis.

If you are running the program under DB2 for Linux or UNIX, ensure that the LIBPATH environment variable is set to the same environment in which the program starts.

ASN0538I *pgmname : program_qualifier : The program is waiting for DB2 to start.*

Explanation: When the program is initially started, if DB2 is stopped the program waits until DB2 starts. After DB2 starts, the Capture program makes the connection and begins to capture changes.

If the TERM=N option is specified in the Capture invocation parameters and DB2 stops smoothly, the Capture program waits for DB2 to start.

User response: This message is for your information only, and no action is required.

ASN0539E *program_name : program_identifier : The Capture program did not start. You must first configure the logarchmeth1 or logarchmeth2 database configuration parameter for database database_name to support archival logging.*

Explanation: The Capture program tried to start but the source database was not configured properly for Capture to use the log read interface.

User response: Take the following steps:

1. Review the settings for the **logarchmeth1** and **logarchmeth2** database configuration parameters. If both parameters are set to OFF, enable recoverable logging by setting the **logarchmeth1** or **logarchmeth2** database configuration parameter to a value other than OFF.
 2. Use the backup command to back up the database before using it with any application.
-

ASN0540E *pgmname : program_qualifier : The program was not able to execute the autobind operation successfully on package pkg_name from the file path_filename. The SQLSTATE sqlstate was returned.*

Explanation: The program discovered that a bind or rebind is required in order to run. The program attempted to autobind, and the autobind was unsuccessful. The program failed to initialize.

User response: Check for corresponding database messages that might provide additional details about the source of the autobind failure. Possible problems include authorization failures, missing or incorrect control tables, or bind files that do not match the program. Correct the situation, and restart the program.

ASN0541E *pgmname : program_qualifier : An incorrect value column_value was supplied for column column_name of the program parameter table parms_table.*

Explanation: This message is issued during initialization of the Capture program if the validation

of the IBMSNAP_CAPPARMS table found that one of the parameter value settings is not valid. The Capture program terminates with this error.

User response: Check the documentation for permitted parameter values allowed in this table. Correct the values and restart the Capture program.

ASN0542E *pgmname : program_qualifier : The maximum number of lock time-out or deadlock retries has been reached.*

Explanation: The program has internally retried a time-out or deadlock condition multiple times. If the program task that receives the persistent lock condition is critical, such as a worker thread, the whole program terminates. If the program task is not critical, such as pruning or monitoring, then the task will be retried at a later time, and the program remains active.

User response: Check for corresponding database messages that might provide additional detail about the source of the locking contention. Correct the situation if the error is a user error, such as a user held lock. If the condition persists, contact IBM Software Support for assistance.

ASN0543E *program_name : program_identifier : The program cannot obtain number bytes of storage for a object. The program terminates.*

Explanation: The program is unable to obtain memory for a necessary in-memory storage structure.

User response: Consider allocating more memory to the program, then restart the program.

ASN0544E *pgmname : program_qualifier : The program is already active.*

Explanation: An attempt was made to start more than one Capture program per DB2 subsystem or database.

User response: Ensure that the schema or qualifier is specified correctly.

- For DB2 for z/OS subsystems, either run only one instance of the Capture program for all subsystems that are members of a data-sharing group, or run only one instance of the Capture program on any standalone system. Display the ENQ resource to determine the unique resource name violation.
- For other DB2 database platforms, run only one Capture program per database using a given schema.

ASN0545E *pgmname : program_qualifier : The program started with the startup parameter PWDFILE, but the password file password_file was not found.*

Explanation: The program cannot find the password file. The user specifies the password filename through

the PWDFILE parameter. If the user specifies a path startup parameter, the password file should reside in the specified directory. If the user did not specify a path startup parameter, the password file should reside in the current directory in which the program is running.

User response: Ensure that the password file name is correctly specified and is located in the proper directory.

ASN0546W *pgmname : program_qualifier : The program call issued to the Automatic Restart Manager failed. The invoked IXCARM macro is arm_call, the return code is return_code, and the reason code is reason_code.*

Explanation: The Capture or Apply program cannot connect to, disconnect from, or receive a ready status indication from the Automatic Restart Manager (ARM). The message displays the unsuccessful call and the return or reason code that returned to the program from the ARM. The program does not terminate but cannot continue ARM processing.

User response: Check the Automatic Restart Manager documentation for more information about the cause of this failure.

ASN0547I *pgmname : program_qualifier : The number of substitution variables passed, nbr_vars, does not match the number of tokens, nbr_tokens, in the text of message number msg_nbr.*

Explanation: The program code and the program message file do not match; the release level of the program and the message file catalog might not match.

User response: Verify that the program message file is correctly installed with the appropriate file permission settings.

ASN0548I *pgmname : program_qualifier : The program received an operator stop command.*

Explanation: This informational message indicates that a stop command was issued to the program.

User response: This message is for your information only, and no action is required.

ASN0552E *pgmname : program_qualifier : The program encountered an SQL error. The server name is server_name. The SQL request is sql_request. The table name is table_name. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sql_tokens. The SQLERRP is error_module.*

Explanation: A nonzero SQLCODE returned when the Capture, Apply, or Monitor program issued an EXEC SQL statement or CLI call. This SQLCODE might be caused by a DB2 problem that needs to be investigated, such as, an out of space condition, or DB2 is unavailable for use by applications. This message is sometimes followed by a second message that provides information about what replication was doing when this SQLCODE was encountered.

User response: See the messages and codes documentation of the DB2 database manager on your operating system for an explanation of this SQLCODE and for information about corrective actions that might need to be taken in DB2. If replication issued another message immediately following this one, see the explanation and user response for that message.

ASN0553W *pgmname : program_qualifier : Internal error error_number occurred for message number msg_number containing num_tokens substitution fields: sub_tokens.*

Explanation: The *error_number* is a decimal internal error number which is defined as:

- | | |
|---|-----------------------|
| 1 | Instance is not valid |
| 2 | Access denied |
| 3 | No files |
| 4 | No message |
| 5 | Locale is not valid |
| 6 | System error |
| 7 | Not enough memory |

The *msg_number* is the message that the program was trying to issue. The *num_tokens* is the number of substitution tokens given for the message (not including the *pgmname* and *program_qualifier* tokens). The *sub_tokens* is the substitution tokens for the message in error separated by commas.

User response: Take any corrective action possible based on the error code given. For example, if the message file was not found or could not be accessed, you should also see message ASN0513 which gives you the file name. Verify that the message file exists with the correct permissions. If you get error code 4, you might have an old message file.

ASN0554E *pgmname : program_qualifier : The program encountered a DB2 log full condition on server server_name.*

Explanation: The program tried to process an insert or update which was denied by DB2 because the DB2 transaction log is full. The program will stop processing.

User response: Check the amount of space remaining on the file system that contains your database files. Consider increasing the maximum log size in the database configuration file.

ASN0555W *pgmname : program_qualifier : The program cannot register with Automatic Resource Manager (ARM) because it is not APF authorized.*

Explanation: The Capture, Apply or Monitor program cannot register to use Automatic Resource Manager services because the program libraries are not APF authorized.

User response: If you desire Capture, Apply, or Monitor program to register with Automatic Resource manager, authorize the program libraries for APF and restart the program.

ASN0556E *pgmname : program_qualifier : A registration that is not valid was found. The CD table phys_chg_owner.phys_chg_tbl does not have any columns that match the base table source_owner.source_table.*

Explanation: The Capture program tried to initialize a registration and found that the CD table does not have any columns that match the source table. Nothing can be captured for this registration, and therefore it is considered not valid. The registration remains inactive if the STOP_ON_ERROR column in the register (IBMSNAP_REGISTER) table for this registration is set to N, otherwise, the Capture program stops.

User response: Make sure that there is at least one column in the CD table that matches the source table for the registration.

ASN0557E *pgmname : program_qualifier : The value column_value for column column_name of owner. table with source table source_owner.source_table, is not valid.*

Explanation: The Capture program tried to initialize a registration and found a value that is not valid for a column in the IBMSNAP_REGISTER table.

User response: Correct the value for the column in error.

ASN0558E *pgmname : program_qualifier : The parameter table parameter_table can have only one row corresponding to program qualifier program_qualifier.*

Explanation: The program parameters table was not defined correctly or was updated with rows that are not valid.

User response: Ensure that there is only one row in the parameter table corresponding to the

program_qualifier. Ensure that the parameters table has a unique index on the program qualifier.

ASN0559W *pgm_name : program_qualifier : The job was started with a CPU time limit of xx seconds. The program will terminate when the time limit expires.*

Explanation: For z/OS only. The job was started with a CPU time limit of xx seconds. The program will terminate when the time limit expires.

User response: To run the program continuously, cancel the current job. Resubmit the job specifying NOLIMIT or 1440 as the new job limit.

ASN0560E *pgm_name : program_qualifier : The program is unable to create db_object object_name in database db_name*

Explanation: The program is unable to create the database object in the database. The database object is being created in the same object table space as the server control tables.

User response: Review the SQL error code that is related to this error message and take the appropriate action.

ASN0561W *program_name : program_identifier : The program's application code page application_code_page is not the same as the code page database_code_page of database database_name.*

Explanation: For Linux, UNIX, Windows, and iSeries: The Capture program's application code page is different from the code page of the source database. Unless the two code pages are compatible, this difference can result in corrupted data or unexpected errors when the Capture program inserts data into CD tables.

User response: If the code pages are compatible, no action is required. If the code pages are incompatible, stop the Capture program, change the Capture program's application code page to the code page of the database, and restart the Capture program.

ASN0562E *pgmname : program_qualifier : An error occurred when checking to see whether the program is already running. The Return code is return_code. The error message is error_message. The operation is operation.*

Explanation: An error occurred while checking to see whether the program is already running. The following values are valid return codes:

- 1 On Windows, the program encountered an error while trying to create a semaphore.

- 2 On UNIX, the HOME environment variable was not found.
- 3 The mkdir command failed trying to create the HOME/sqllib directory.
- 4 The mkdir command failed trying to create the HOME/sqllib/dpropr directory.
- 5 The fgets operation failed to read a row from the pid file.
- 6 The ps command failed.
- 7 An error occurred while removing the grp file.
- 8 An error occurred while reading the grp file.
- 9 An error occurred while opening the pid file.
- 10 The fput command finished in error for the pid file.

User response: Review the reason codes in the explanation, and respond with the following options:

- 1 This is an internal error on Windows. See the Windows Reference.
- 2 On UNIX, verify that the HOME environment variable is set to the correct value
- 3-10 Ensure that you have the right authorizations to create directories and files in the current HOME path.

ASN0563W *pgmname : program_qualifier : The parameter parameter-1 is not compatible with parameter parameter-2. The parameter parameter-1 will be ignored.*

Explanation: The program was started with both parameters specified in the command or in the parameters table. These parameters are not compatible. One parameter was ignored when the program started and continued to run.

User response: The next time you start the program, specify only the parameter that you want the program to use.

ASN0564I *pgmname : program_qualifier : The program could not perform the sql_request operation on the control table control_table_name to support long schema, owner, and table names. The server name is server_name. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sql_token. The SQLERRP is module_name.*

Explanation: The Capture, Apply, or Monitor control table is defined on a Version 8 new-function mode DB2 subsystem. The control table includes columns that do not support long schema, owner, and table names. The program attempted to convert the control table to support long schema, owner, and table names and the

operation was unsuccessful. The replication program continues to run.

User response: See the DB2 for z/OS messages and codes documentation for an explanation of the SQLCODE and for information about corrective actions that might need to be taken in DB2. If the SQLCODE is -551, do one of the following:

- Grant the user ID that starts the Capture, Apply, or Monitor program the ALTER privilege on the control table
- Run AASNSAMP member (ASNM2V8) to ALTER all replication control tables to support long schema, owner, and table names.

ASN0565E *program_name : program_identifier : The program cannot insert statistics into the table table_name. The SQL return code is sqlcode. The data for this interval will be skipped and included in the next interval.*

Explanation: The program maintains statistics about its activities, but it could not save this data in the table due to an unexpected SQL code. The program will attempt to insert the data again at the next interval.

User response: If the SQL return code indicates a temporary error, no action is required. Otherwise, take action as indicated for the SQL error in the DB2 Information Center.

ASN0566E *program_name : program_identifier : Pruning of table table_name failed with SQL return code sqlcode.*

Explanation: Pruning failed with an unexpected SQL error code. This error does not cause the program to terminate. The program will try to prune the table again at the next pruning interval or when the program is restarted.

User response: If this SQL return code indicates a temporary error, then no action is required. Otherwise, take action as indicated for the SQL error in the DB2 Information Center.

ASN0567I *program_name : program_identifier : number rows were pruned from table table_name.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN0568E *program_name : program_identifier : An ICU error occurred while translating character data from code page code_page. The ICU function is function_name using converter for CCSID CCSID. The ICU return code is return_code. The expected ICU version is version_number. Details: details.*

Explanation: Incompatible versions of ICU (International Components for Unicode) were encountered. The ICU function is the ICU API name. The return code was returned by this API. The details about the ICU operation are provided by the Q Capture program.

User response: Ensure you have the correct version of ICU and try restarting the program.

ASN0569E *program_name : program_identifier : The program encountered an internal error error_code. The program will terminate.*

Explanation: An unexpected error occurred in the program.

User response: See "Collecting data for InfoSphere Replication Server and InfoSphere Data Event Publisher" for information on how you can assist in troubleshooting this problem.

ASN0571E *program_name : program_identifier : The program cannot start because the parameter table table_name is empty or it contains more than one row.*

Explanation: The program parameter table must have exactly one row.

User response: Use the sample migration scripts in the sqllib/samples/repl directory for the DDL statements to redefine the content of the parameter table so that it contains exactly one row.

ASN0572I *program_name : program_identifier : The version program initialized successfully.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN0573I *program_name : program_identifier : The program was stopped.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN0574E *program_name : program_identifier : The WebSphere MQ queue manager queue_manager_name is not available or it was not started. The program will terminate.*

Explanation: The program cannot connect to the WebSphere MQ queue manager. This error causes the program to terminate.

User response: Start the queue manager if it is not running, for example by using the strmqm command. You can also look at "2059 (080B) (RC2059): MQRC_Q_MGR_NOT_AVAILABLE" in the WebSphere MQ Information Center for other details.

ASN0575E *program_name : program_identifier : The program encountered a WebSphere MQ error reason_code while issuing the WebSphere MQ command command on object name.*

Explanation: A nonzero WebSphere MQ reason code was returned when the program issued a WebSphere MQ command. The reason code indicates that there is a WebSphere MQ problem (for example, storage is not available for WebSphere MQ objects) or that WebSphere MQ is not available.

The ERROR_ACTION value determines how the program behaves due to this error. The value for the Q Capture program is stored in the IBMQREP_SENDQUEUES table. The value for the Q Apply program is stored in the IBMQREP_TARGETS table.

User response: See the reason code documentation in the WebSphere MQ Application Programming Reference for an explanation of this WebSphere MQ reason code and for information about actions that you might need to take. The following list provides information about possible actions that are specific to Q Replication.

2003 (X'07D3') MQRC_BACKED_OUT (Linux, UNIX, Windows)

If Q Capture receives this reason code, it might indicate that the size of your queue manager log is not large enough to handle the workload that Q Capture generates, especially if large transactions are being replicated. The queue manager log is used for recovery; the queue manager also has an error log. To determine if queue manager log size is a problem, check the queue manager error log for messages that indicate that the queue manager log became full (for example, message AMQ7469). By default, the queue manager error log is in a directory that has the following path: MQ_install_directory/queue_manager_name/errors. The error log files have names like AMQERR01.LOG. You

can increase the size of the queue manager log by increasing the values of the following three queue manager properties: LogPrimaryFiles, LogSecondaryFiles, LogFilePages. For example, you might want to make sure that these properties are no smaller than the following:

- LogPrimaryFiles: 15
- LogSecondaryFiles: 10
- LogFilePages: 1024

ASN0576E *program_name : program_identifier : The program cannot access the WebSphere MQ queue queue_name.*

Explanation: Either the queue does not exist or it is specified incorrectly in the control tables. The program terminated as a result of this error.

User response: Ensure that the queue name is specified correctly in the control tables and that it is accessible to application programs. Create the queue if it does not exist.

ASN0580E *program_name : program_identifier : The license for product product_name was not found.*

Explanation: The program terminated because it cannot run without a license.

User response: Install the program license or contact your IBM representative.

ASN0581W *program_name : program_identifier : The program was not able to initialize a connection to name because Recoverable Resource Manager Services (RRS) is not started. The program will attempt to use Call Attach Facility (CAF) instead.*

Explanation: With RRS/AF, you must first issue an Identify request. The request failed because the RRS service in z/OS was not started.

User response: No action is required because the program will run correctly with CAF. However, if you want the program to use RRS/AF, you must first make sure RRS is started, and then restart the program.

ASN0582I *program_name : program_identifier : The program was suspended by an operator command.*

Explanation: An operator command has suspended the program, and the program has entered a wait state.

User response: This message is for your information. No action is required.

ASN0583I *program_name : program_identifier : The program was resumed by an operator command.*

Explanation: An operator command has resumed the program from a suspended state, and the program has continued running.

User response: This message is for your information. No action is required.

ASN0584E *program_name : program_identifier : An error occurred while the program was dynamically loading the WebSphere MQ library library_name. Error code: error_code, error_description. Environment variable ASNUSEMQCLIENT is set to value.*

Explanation: An error occurred when trying to dynamically load a WebSphere MQ library.

User response: If the environment variable ASNUSEMQCLIENT is set to TRUE, the WebSphere MQ client libraries are used; otherwise the WebSphere MQ server libraries are used. The server libraries are also used if the environment variable ASNUSEMQCLIENT is not set. If the server configuration is used, then ensure that the WebSphere MQ server is running on the local host. If the client configuration is used, then make sure that this host can communicate with the remote host that the server is running on.

If you did not install WebSphere MQ in the default location, ensure that the library path variable (for example LIBPATH on AIX, LD_LIBRARY_PATH on SUN OS, SHLIB_PATH on HP/UX, and PATH on Windows) points to the correct installed path. If you have multiple versions of WebSphere MQ, make sure the variable points to the correct version.

ASN0585I *program_name : program_identifier : The program successfully loaded the WebSphere MQ library library_name. Environment variable ASNUSEMQCLIENT is set to value.*

Explanation: The dynamic loading of the WebSphere MQ library was successful

User response: This message is for your information. No action is required.

ASN0586E *program_name : program_identifier. An error occurred while the program was getting the function address for function_name from a dynamically loaded library. Error code error_code, error_description.*

Explanation: An error occurred while retrieving the

function address for the specified WebSphere MQ API command.

User response: This error is likely to occur if there are differences in the API signatures for WebSphere MQ. Please check that you are using the correct versions of the WebSphere MQ server or client and its libraries. If the environment variable ASNUSEMQCLIENT is set to TRUE, the client libraries are used, otherwise the server libraries are used. The WebSphere MQ server libraries are also used if the environment variable ASNUSEMQCLIENT is not set. If the server configuration is used, then ensure that the WebSphere MQ server is running on the local host. If the client configuration is used, then make sure that this host can communicate with the remote host that the server is running on.

If you did not install WebSphere MQ in the default location, ensure that the library path variable (for example LIBPATH on AIX, LD_LIBRARY_PATH on SUN OS, SHLIB_PATH on HP/UX, and PATH on Windows) points to the correct installed path. If you have multiple versions of WebSphere MQ, make sure the variable points to the correct version.

ASN0587I *program_name : program_identifier The thread is retrying to connect with the WebSphere MQ queue manager.*

Explanation: The WebSphere MQ queue manager is not currently available, The thread is retrying to connect.

User response: Restart the WebSphere MQ queue manager if it was not active.

ASN0588I *program_name : program_identifier The program timed out while it was attempting to connect to database database_name.*

Explanation: This message is issued when the program is initializing if the program cannot connect to the server with its control tables. The program shuts down.

User response: Check for database messages that might explain why the program was unable to establish a database connection. Contact the database administrator for questions and diagnosis.

ASN0589I *program_name : program_identifier The program received return code return_code from routine routine.*

Explanation: The program received the stated return code from one of its routines. The return code and routine information might help IBM determine the cause of an error.

User response: This message is for your information. This message is typically preceded by error messages,

which you can view in the program's log file.

ASN0590I *program_name : program_identifier* **The thread *thread_name* received return code *return_code* from the exiting thread *thread_name*.**

Explanation: The program received this return code when it stopped one of its threads. The following thread exit return codes are the most common:

0

Thread stopped successfully. There is no reason code.

n

The Apply program encountered *n* failed cycles.

2001

Thread was stopped by either a stop command or signal.

2009

Thread was stopped by the recovery manager.

2010

Thread was stopped by recursive calls to the recovery manager.

2011

Thread stopped itself without an error.

2012

Thread stopped because of an error.

User response: This message is for your information. If this message is displayed with error messages, look in the program's log file that contains the messages.

ASN0591I *program_name : program_identifier* **The thread *thread_name* received action *signal* *signal_name*.**

Explanation: The thread received this signal before the thread shut down. The ACTION parameter value is either handled or not handled. The thread recovery function handles signals that it expects and does not handle signals that it does not expect. The HoldLThread normally receives a SIGUSR2 signal before the initial thread shuts the HoldLThread down. Other replication threads receive a SIGUSR1 signal if they are terminated by the initial thread.

User response: This message is for your information. If this message is displayed with error messages, look in the program's log file that contains the messages.

ASN0592I *program_name : program_identifier* **The program attached to the IPC queue with keys *key_values*.**

Explanation: The program attached to an IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate.

User response: This message is for your information. No action is required.

ASN0593I *program_name : program_identifier* **The program detached from the IPC queue with keys *key_values*.**

Explanation: The program detached from an IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate.

User response: This message is for your information. No action is required.

ASN0594I *program_name : program_identifier* **The program created an IPC queue with keys *key_values*.**

Explanation: The program created an IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate.

User response: This message is for your information. No action is required.

ASN0595I *program_name : program_identifier* **The program removed an IPC queue with keys *key_values*.**

Explanation: The program removed an IPC command message queue while the program was initializing or terminating. Replication command programs use message queues to communicate with the programs that they operate. Every time the program starts, the replication program creates a new message queue. If the message queue exists, it removes it and creates a new one. Every time the replication program shuts down, it tries to remove the message queue.

User response: This message is for your information. No action is required.

ASN0596I *program_name : program_identifier* **The program could not create an IPC queue with keys *key_values* for path *path_name*. OSSE reason is *reason*.**

Explanation: The program is not able to create its IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate. Every time the program starts, the replication program tries to create a new

message queue. If the message queue exists, it tries to remove it and create a new one.

User response: If the reason is "Shared memory ID already exists with the given resource ID", perform the following steps to remove the file and its associated message queue manually:

1. Note the key values. The key values will look similar to these:
(0x300667f7, 0x310667f7, 0x320667f7)
2. Log into the server using the user ID that was used to start the replication program.
3. Enter the following command for each of the keys in the ASN0594I message:
ipcs grep | 0x300667f7

If a key is located, note its IPC resource type, ID, and owner. The IPC resource type (m, q, or s) will be in the first column of the IPCS output on UNIX and z/OS servers. The IPC resource type will be identified by a header line on LINUX servers.

4. If the keys are not located, log into the server using the superuser ID and enter the IPCS command in step 3 for each of the keys.
5. Log into the server using the user ID that owns the IPC resources that was located in the steps above. If the key identifies a shared memory segment, enter the following command to remove the shared memory segment:

```
ipcrm -m ID
```

If the key identifies a semaphore, enter the following command to remove the semaphore:

```
ipcrm -s ID
```

If the key identifies a message queue, enter the following command to remove the message queue:

```
ipcrm -q ID
```

6. Log into the server using the user ID that owns the file specified by the PATHNAME parameter value. Enter a command to remove the file, as in this example:
rm /tmp/dpropr5.SRCDB.TIMING03.APP.IPC

ASN0597I *program_name : program_identifier* **The program could not attach to an IPC queue with keys *key_values* for path *path_name*. OSSE reason is *reason*.**

Explanation: The program is not able to attach to a replication IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate. The program that owned the message queue might not be running.

User response: This message is for your information. No action is required.

ASN0598W *program_name : program_identifier* : **The program connected to database *database_name* but only single-byte characters will be supported. The SQLCODE is +863.**

Explanation: DB2 has indicated that the server database and client application are using code pages for different language types and any characters outside the seven-bit ASCII range cannot be guaranteed (SQLCODE +863).

User response: Verify that you have configured the operating system and database managers correctly for the code pages that are being used. See the *DB2 Message Reference, Volume 2* for more details on SQLCODE +863.

ASN0599E *program_name : program_identifier* : **The program detected an unsupported architecture level *bad_arch_level*. The only supported level is *good_arch_level*.**

Explanation: The replication program is at a different product level than the control tables for the given program identifier. The program determines this by checking the architecture level that is set in the ARCH_LEVEL column of the IBMQREP_APPLYPARMS table, IBMQREP_CAPPARMS table, or IBMSNAP_APPLEVEL table.

User response: Verify that you specified the correct database alias when you started the program. If the replication control tables are older than the program, migrate the tables to the same product level as the program.

ASN0600I *program_name : program_identifier* : **Program *program_name-version* is starting.**

User response: This message is for your information. No action is required.

ASN0601I *program_name : program_identifier* : **The program could not send a message to an IPC queue with keys *keys* for path *path*. Error code is *error_code*.**

Explanation: The program is not able to send a message to a replication IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate. The program that owned the message queue may no longer be running.

This message is issued as a follow-on message to error message ASN0508E and contains the error code that provides additional diagnostic information.

User response: This message is for your information. No action is required.

ASN0602I *program_name : program_identifier : The program could not read from its IPC queue with keys keys for path path. Error code is error_code.*

Explanation: The program is not able to read a message from its IPC command message queue. Replication command programs use message queues to communicate with the programs that they operate. The program that owned the message queue may no longer be running.

This message is issued as a follow-on message to error message ASN0508E and contains the error code that provides additional diagnostic information.

User response: This message is for your information. No action is required.

ASN0603E *program_name : program_identifier : Active Q subscriptions exist, but the license for WebSphere Replication Server was not found.*

Explanation: The Q Capture program terminated because active Q subscriptions exist, but a license for replication was not found. This situation can happen for one of the following reasons:

- The program was migrated from Version 8 to Version 9 and a Version 9 license was not installed.
- A try-and-buy license expired.

User response: Install a license for WebSphere Replication Server.

ASN0604E *program_name : program_identifier : Active publications exist, but the license for WebSphere Event Publisher was not found.*

Explanation: The Q Capture program terminated because active publications exist, but a license for event publishing was not found. This situation can happen for one of the following reasons:

- The program was migrated from Version 8 to Version 9 and a Version 9 license was not installed.
- A try-and-buy license expired.

User response: Install a license for WebSphere Event Publisher and restart Q Capture before activating any publications.

ASN0605E *program_name : program_identifier : Neither the license for WebSphere Replication Server nor the license for WebSphere Event Publisher was found.*

Explanation: The Q Capture program terminated because a valid license was not found. This situation can happen for one of the following reasons:

- A license was never installed.

- The program was migrated from Version 8 to Version 9 and a Version 9 license was not installed.
- A try-and-buy license expired.

User response: Install a license for WebSphere Event Publisher to have Q Capture process publications. Install a license for WebSphere Replication Server to have Q Capture process Q subscriptions.

ASN0606I *program_name : program_identifier : The program successfully loaded the library library_name.*

Explanation: This message appears after a successful start of the replication program.

User response: This message is for your information only, and no action is required.

ASN0607E *program_name : program_identifier : Unable to load the library library_name. Error code is error_code. Error message is error_message.*

Explanation: The program encountered an error while loading the replication library.

User response: Verify that the library exists in the specified path.

ASN0608I **Replication code uses "64" bits. Informational tokens are "DB2 v9.1.0", "n060704", "WR21350", and FixPak "1". DB2 instance path is "/home/inst1/sqlllib". Replication path is "/home/inst9/sqlllib".**

Explanation: This message appears after invoking the asnlevel program.

User response: This message is for your information only. No action is required.

ASN0609W *program_name : program_identifier : A license for WebSphere Event Publisher was not found.*

Explanation: The Q Capture program found a license for replication, but did not find a license for event publishing. The Q Capture program does not terminate because no active publications were found. Q Capture will not be able to activate new publications.

User response: If you plan to create and activate publications, install a license for WebSphere Event Publisher. Otherwise, no action is required.

ASN0610W *program_name : program_identifier : A license for WebSphere Replication Server was not found.*

Explanation: The Q Capture program found a license for event publishing, but did not find a license for

replication. The Q Capture program does not terminate because no active Q subscriptions were found. Q Capture will not be able to activate new Q subscriptions.

User response: If you plan to create and activate Q subscriptions, install a license for WebSphere Replication Server.

ASN0612E *program_name : program_identifier : The publication publication_name was not activated because a license for WebSphere Event Publisher was not found.*

Explanation: The Q Capture program does not activate an publication unless it finds a license for event publishing. The Q Capture program does not terminate. This situation can happen for one of the following reasons:

- A event publishing license was never installed.
- The program was migrated from Version 8 to Version 9 and a Version 9 event publishing license was not installed.
- A try-and-buy license expired.

User response: Install a license for WebSphere Event Publisher.

ASN0613E *program_name : program_identifier : The Q subscription Q_subscription_name was not activated because a license for InfoSphere Data Replication was not found.*

Explanation: The Q Capture program does not activate a Q subscription unless it finds a license for replication. The Q Capture program does not terminate. This situation can happen for one of the following reasons:

- A replication license was never installed.
- The program was migrated from Version 8 to Version 9 or Version 10 and a Version 9 or Version 10 replication license was not installed.
- A try-and-buy license expired.

User response: Install a license for InfoSphere Data Replication.

ASN0614E *program_name : program_identifier : A license for InfoSphere Data Replication was not found.*

Explanation: The Q Apply program terminated because a valid license was not found. This situation can happen for one of the following reasons:

- The license was never installed.
- The program was migrated from Version 8 to Version 9 or Version 10 and a Version 9 or Version 10 license was not installed.

- A try-and-buy license expired.

User response: Install a license for InfoSphere Data Replication.

ASN0615E *program_name : program_identifier : The program could not call dynamic function function_name because the function could not be dynamically loaded from any of these libraries: libraries.*

Explanation: The program tried to call a dynamic link library (DLL) function but could not because either the program could not load a library with that function or the function address could not be resolved from the library that was loaded.

User response: Make sure that the libraries exist and are specified in STEPLIB on the z/OS operating system or the library path environment variable specific to your Linux, UNIX, or Windows system.

ASN0616E *program_name : program_identifier : The program cannot load any of these dynamic link libraries (DLLs): libraries. The operating system error number for all of the DLLs that the program tried to load is error_number.*

Explanation: The program tried to load the dynamic link libraries (DLLs) that are listed but received the operating system error that is listed.

User response: Correct the operating system error and start the program.

ASN0617E *program_name : program_identifier : Function function_name was not exported from dynamic link library (DLL) library_name. The OSSe (Operating System Service Everywhere) return code is return_code. The OSSe return code description is description.*

Explanation: The program could not resolve the address of the function in the listed DLL because the DLL did not export the function.

User response: Remove the listed DLL from the STEPLIB so that the program can try a DLL with a different version.

ASN0618W *program_name : program_identifier : The program cannot use Unicode Conversion Services (UCS) for conversion from coded character set identifier (CCSID) identifier to CCSID identifier. The UCS return code is return_code. The UCS reason code is reason_code.*

Explanation: The program could not use UCS to convert data between the listed CCSIDs. The program

will try to use the International Components for Unicode (ICU) character set for the conversion.

User response: If neither UCS nor ICU can be used for conversion, take the necessary steps to enable either UCS or ICU for the listed CCSIDs.

ASN0619W *program_name : program_identifier : The program cannot use Unicode Conversion Services (UCS) or International Components for Unicode (ICU) to convert data from coded character set identifier (CCSID) identifier to CCSID identifier.*

Explanation: The program tried to use both UCS and ICU to convert data between the listed CCSIDs, but was unable to use either method for conversion.

User response: Take the necessary steps to enable either UCS or ICU for the listed CCSIDs.

ASN0620E *program_name : program_identifier : The program log read failed with reason code reason_code. The log record that could not be read has log sequence number (LSN) lsn and is for table table_name.*

Explanation: The replication capture program received an error from the database log read API while getting a log record for the table.

User response: Use the reason code to determine the database log read error. Look for messages that were issued by the database, and look for subsequent capture messages that may give you more information about the error.

ASN0621E *program_name : program_identifier : The program could not read a log record for table table_name because the edit routine (EDITPROC) that was used to encode the contents of the row could not be called by DB2 to decode the row contents.*

Explanation: Replication supports tables with EDITPROC definitions. The database log read API invokes the EDITPROC to decode the row contents before returning them to the replication capture program. The API could not use the EDITPROC to decode the row contents. Depending on the error action option specified, capture may stop or deactivate the subscription or registration.

User response: Determine why the log read API could not use the EDITPROC to decode the row contents. Look for messages from the database that may give you more information about the error.

ASN0622E *program_name : program_identifier : The program is not connected to database database_name. The current server is current_server.*

Explanation: The program explicitly connected to the database server, but the database server was not the current server after the connection executed.

User response: Look for database-issued messages on the system console that may give you more information about the error.

ASN0623E *program_name : program_identifier : The thread_name thread is unable to lock a mutex for the list_name list. The error number is error_number. The reason is reason. The lock holder,lock_holder, has locked the mutex number times.*

Explanation: The program thread is unable to lock a mutex for a list.

User response: This message is for your information only.

ASN0624E *program_name : program_identifier : The thread_name thread is unable to unlock a mutex for the list_name list. The error number is error_number. The reason is reason. The lock holder,lock_holder, has locked the mutex number times.*

Explanation: The program thread is unable to unlock a mutex for a list.

User response: This message is for your information only.

ASN0625W *program_name : program_identifier : ODBC function function_name completed with return code return_code, SQL state SQL_state, and native database completion code completion_code.*

Explanation: Replication supports use of ODBC functions to access certain databases. An ODBC function has completed with an unsuccessful return code that may or may not result in an error condition to replication. A diagnostic message follows with message text from the native database.

User response: Check later messages to determine if a subsequent error condition occurred in replication. The information in these messages might help determine the cause of the error. Correct the problem as appropriate.

ASN0626E *program_name : program_identifier : The Q Apply program encountered an error while it was creating a file on the target server to temporarily store Classic replication source data before loading a target table (routine_name, function, dataset_name, dataset_type, error).*

Explanation: During the target table loading process for Classic replication, the Q Apply program creates a temporary intermediate file on the target server to place data that it selects from the source table. The Q Apply program then reads from the file before inserting the data into the target. The user ID that runs the Q Apply program must have the authority to create this file.

User response: Make sure the user ID that is running the Q Apply program has the authority to create the file and that there is enough disk space on the target server to create the file. Use the error information in the message to determine other reasons that the file could not be created.

ASN0627E *program_name : program_identifier : The program stopped because column column_name is missing from the control table table_name.*

Explanation: A required column is missing from a control table. This message only reports the first missing column; other columns might also be missing. The missing columns are likely to be columns that were added to the control tables for a new version, which indicates that the control tables need to be migrated.

User response: If you installed a new version but did not migrate the control tables, run the necessary migration scripts and start the program.

ASN0628E *program_name : program_identifier : The column option option for source column column_name, which is part of Q subscription or publication name, has an invalid value value.*

Explanation: The specified value is not valid for the column option.

User response: Ensure that the column option is specified correctly in the publication or Q subscription. Use the replication administration tools to perform the following steps:

1. Drop the Q subscription or publication and recreate it with a valid value for the column option.
2. Start the Q subscription or publication.

ASN0629E *program_name : program_identifier : The program has encountered an unsupported data type data_type while trying to activate a registration, Q subscription, or publication.*

Explanation: The program found one or more columns in the source table that are defined with unsupported data types. Replication of the XML data type for Q Replication requires a COMPATIBILITY value in the IBMQREP_CAPPARMS table of 0905 or higher.

User response: Redefine the registration, Q subscription, or publication without the columns that contain unsupported data types. For source tables with the XML data type, ensure that the Q Capture program is at Version 9.5 or higher and that the COMPATIBILITY value is 0905 or higher.

ASN0631E *program_name : program_identifier : The program could not load field procedure procedure_name for source table table_owner.table_name, column column_name. z/OS returned system completion code completion_code and reason code reason_code. The program stopped.*

Explanation: The Capture or Q Capture program tried to load the user-specified field procedure to decode data in the source column because the column is defined with this field procedure. An error occurred that prevented the program from calling the procedure.

User response: Correct the error that is indicated by the system completion code and reason code and warm start Capture or Q Capture. The system completion code and reason code are explained in MVS System Codes.

ASN0632E *program_name : program_identifier : The field procedure for source table table_owner.table_name, column column_name detected an error. The field procedure returned return_code and reason code reason_code. The program stopped.*

Explanation: The user-specified field procedure detected an error while decoding data in the column. An error message from the field procedure might follow this message.

User response: Correct the error that is indicated by the return and reason codes and warm start Capture or Q Capture. The meanings of the return and reason codes are defined by the field procedure.

ASN0633W *program_name : program_identifier :* **The Q Capture or Capture program has detected either a long-running transaction or possible problem. The program read a log record for the start of transaction *transaction_ID* at *timestamp*. But the program has not seen a commit or rollback log record for the transaction for *number* seconds. The authorization ID is *auth_ID*. The correlation ID is *corellid*. The starting log sequence number for the transaction is *LSN*. Q Capture or Capture has captured up to *LSN LSN*.**

Explanation: The program captured the start of the transaction but has not seen its commit or rollback log record for more than an hour. This might or might not be a problem.

User response: If this is a long-running transaction, ignore this message. On z/OS, you can use the DSN1LOGP utility to format the contents of the recovery log and determine whether it is a long-running transaction.

ASN0634W *program_name : program_identifier :* **The transaction *transaction_ID* has not been committed for more than *timestamp* seconds. The initial log sequence number for the transaction is *LSN*. The authorization ID is *authid*. The correlation ID is *corellid*.**

Explanation: In addition to the oldest uncommitted transaction that was mentioned in ASN0633W, this message indicates a transaction for which the Q Capture or Capture program has not seen a commit or rollback log record for more than an hour.

User response: If this is a long-running transaction, ignore this message. On z/OS, you can use the DSN1LOGP utility to format the contents of the recovery log and determine whether it is a long-running transaction.

ASN0635I *program_name : program_identifier :* **The Q Capture or Capture program will ignore transactions that are identified by authorization ID *authid*, authorization token *authtoken*, and plan name *planname*.**

Explanation: The Q Capture or Capture program will ignore log records for transactions that are identified by the values that were inserted into the AUTHID, AUTHTOKEN, and PLANNAME columns of the IBMQREP_IGNTRAN table.

User response: This message is for your information only. No action is required.

ASN0636E *program_name : program_identifier :* **The program could not connect to database *database_name*. The reason code is *reason_code*.**

Explanation: An error occurred when the program issued one of the following functions:

- A CONNECT to DB2 for VSE and VM
- A CONNECT to DB2 Call Attachment Facility (CAF)
- An implicit connect to DB2

User response: See the DB2 codes in the messages and codes publication of the DB2 database manager on your operating system for the appropriate reason code. For DB2 for z/OS errors, see the section in the administration guide that describes the Call Attachment Facility.

ASN0637I *program_name : program_identifier :* **Spilling was stopped to avoid exceeding the file size limit for *file_name*.**

Explanation: The Q Capture or Capture program was spilling an in-memory transaction to a file to relieve memory. Spilling was stopped because the file size reached the internal limit of 1 GB. Spilling might be resumed to a new file if memory is depleted again. No data is lost.

User response: This message is for your information only. No action is required.

ASN0638W *program_name : program_identifier :* **One or both of the DB2 instance-level variables DB2_SKIPINSERTED=ON or DB2_SKIPDELETED=ON is set and could interfere with replication.**

Explanation: When DB2_SKIPINSERTED=ON or DB2_SKIPDELETED=ON is set for the DB2 instance, a select statement does not return rows that are inserted and deleted but not yet committed, even under the cursor stability isolation mode. The following results are possible:

- When you start a subscription that has a load phase, rows that are inserted or deleted might not be picked up by the EXPORT utility when the table is loaded at the target and might be missing from the replicated table at the target.
- The asntdiff utility might return inconsistent results.

User response: Suspend any applications that insert into or delete from the source table until a CAPSTART signal for the table has been inserted into the IBMQREP_SIGNAL table. Also, suspend applications when you run the asntdiff utility.

ASN0639W *program_name : program_identifier :* **One or both of the DB2 instance-level variables DB2_SKIPINSERTED=ON or DB2_SKIPDELETED=ON is set and could cause rows to be lost for source table *table_owner.table_name* at the target.**

Explanation: When DB2_SKIPINSERTED=ON or DB2_SKIPDELETED=ON is set for the DB2 instance, a select statement does not return rows that are inserted and deleted but not yet committed, even under the cursor stability isolation mode. When you start a subscription that has a load phase, rows that are inserted or deleted might not be picked up by the EXPORT utility when the table is loaded at the target and might be missing from the replicated table at the target.

User response: Suspend any applications that insert into or delete from the source table until the Q subscription for the table is started.

ASN0640I *program_name : program_identifier :* **The program is waiting for the WebSphere MQ queue manager to become available.**

Explanation: The Q Capture program is running with the term invocation parameter set to N and Q Capture cannot connect to the WebSphere MQ queue manager. With term=N, Q Capture waits until it can either connect to the queue manager or is stopped by a stop command. After Q Capture connects to the queue manager, it continues capturing changes.

User response: This message is for your information only. No action is required.

ASN0641E *program_name : program_identifier :* **The program is shutting down to avoid an inconsistent state even though the TERM=N parameter was specified. You will need to restart the program.**

Explanation: The TERM=N parameter instructs the program to keep running when the DB2 database manager is not available and when the queue manager is not available.

This message was returned because either the DB2 database manager was not available or the queue manager was not available, but when the program temporarily stopped some threads to go into a waiting state, one or more of the threads did not stop gracefully. To avoid an inconsistent state, the entire program stopped.

User response: Restart the program.

ASN0642E *program_name : program_identifier :* **Compatibility level *level* is not valid. The Q Capture program stops.**

Explanation: The replication programs use the Q Capture compatibility parameter to determine the level of messages that are sent. The compatibility level that was found in the Q Capture control tables is not supported.

User response: Change the compatibility level to a valid value and start the Q Capture program. For information about compatibility, see "Coexistence support in Version 9.7 Q Replication and event publishing."

ASN0643E *program_name : program_identifier :* **An invalid wildcard pattern *pattern* was specified in the *column_name* column of the IBMQREP_IGNTRAN table. The Q Capture or Capture program shuts down to avoid ignoring transactions that you might want to replicate.**

Explanation: A single percentage sign (%) was inserted into the IBMQREP_IGNTRAN table. Because this character is treated as a wildcard to represent any number of characters or none, inserting a single % into the table instructs Q Capture to ignore all transactions and is therefore invalid.

User response: Update the IBMQREP_IGNTRAN table with a valid authorization token or plan name (z/OS) or authorization ID (Linux, UNIX, Windows). If the percentage sign is part of the data but is not being used as a wildcard character, use a backslash to escape the wildcard character (\\%). Then start the Q Capture or Capture program.

ASN0644W *program_name : program_identifier :* **A non-NULL value was specified in the *column_name* column of the IBMQREP_IGNTRAN table. This value was ignored.**

Explanation: On Linux, UNIX, or Windows, you can insert transaction identifiers into only the AUTHID (authorization ID) column of the IBMQREP_IGNTRAN table.

User response:

1. Specify a transaction identifier in the AUTHID column.
2. Reinitialize the Q Capture or Capture program.

ASN0646E *program_name : program_identifier :* **The TZ environment variable is not defined. The program will stop.**

Explanation: The replication program requires that the TZ environment variable be set to define the time zone and the offset from Coordinated Universal Time (CUT).

User response: Define the TZ environment variable in one of the following locations:

- etc/profile
- The .profile file in the home directory of the user ID that executes the replication program
- A file that is specified by the LE _CEE_ENVFILE environment variable. The LE _CEE_ENVFILE environment variable enables a list of environment variables to be set from a specified file.

ASN0647E *program_name : program_identifier : The version version of the database database_name is not supported by this version version of the Q Capture or Capture program.*

Explanation: Q Capture or Capture do not support capture of changes from this version of database.

User response: See the DB2 Information Center for details about capture support for different database levels.

ASN0648E *program_name : program_identifier : System catalog table table_name must be set with DATA CAPTURE CHANGES.*

Explanation: The Q Capture or Capture program looks for updates to the system catalog tables to detect ALTER TABLE ADD COLUMN and ALTER TABLE ALTER COLUMN SET DATA TYPE changes to tables. DATA CAPTURE CHANGES must be enabled for this system catalog table.

User response: Enable DATA CAPTURE CHANGES for the system catalog table.

ASN0649I *program_name : program_identifier : A schema-level Q subscription was created, but a Q subscription name already exists for the source table table_owner.table_name that specifies the send or receive queue queue_name, replication queue map queue_map_name, and is in state state. The existing Q subscription will be overwritten.*

Explanation: A schema-level Q subscription was created that specifies that all tables with the same schema should be replicated. The program tried to create a Q subscription for the table, but an inactive Q subscription for the table already exists. The program will proceed to create the new Q subscription and overwrite the inactive Q subscription.

User response: This message is for your information only. No action is required.

ASN0650I *program_name : program_identifier : The Q subscription name that uses send queue or receive queue queue_name and replication or publishing queue map queue_map_name was deleted.*

Explanation: When a table with a matching schema-level subscription is dropped, the Q subscription for the table is deleted.

User response: This message is for your information only. No action is required.

ASN0651E *program_name : program_identifier : Column column_name of table table_owner.table_name was altered. The Q Capture or Capture program at this version cannot automatically replicate data type changes. The program stopped.*

Explanation: Q Capture or Capture cannot automatically replicate the results of an ALTER TABLE ALTER COLUMN SET DATA TYPE statement unless the program is at Version 10.1.

User response: Take the following steps:

1. Reorganize the source table.
2. Alter the matching column in the target table.
3. Redefine the registration or Q subscription to ensure that the source and target columns are correctly mapped.
4. Restart the Q Capture or Capture program.

ASN0652I *program_name : program_identifier : The NMI service initialized successfully. The server listens on socket socket_number.*

Explanation: The program defined an AF_UNIX socket and is listening for NMI client connection requests.

User response: This message is for your information only. No action is required.

ASN0653I *program_name : program_identifier : The NMI service was stopped on socket socket_number.*

Explanation: The program received an NMI TERM message and closed all NMI client connections.

User response: This message is for your information only. No action is required.

ASN0661E *program_name : program_identifier : The job was canceled.*

Explanation: The job was canceled explicitly by the user or as a result of an unexpected event, such as an ABEND.

User response: Investigate the cause and restart the program.

ASN0662I *program_name : program_identifier : The procedure ADMIN_REVALIDATE_DB_OBJECTS was successfully run on table table_owner.table_name.*

Explanation: The program detected that a replication source table was altered and the ADMIN_REVALIDATE_DB_OBJECTS procedure was run to revalidate the table.

User response: This message is for your information only. No action is required.

ASN0663E *program_name : program_identifier : The procedure ADMIN_REVALIDATE_DB_OBJECTS was not successful. The table table_owner.table_name remains in reorg pending state.*

Explanation: After a table is altered, the table might be moved to reorg pending state. The table must be revalidated if it is in reorg pending state before any INSERT, UPDATE, or DELETE operations may be performed. The ADMIN_REVALIDATE_DB_OBJECTS stored procedure performs this validation. The call to ADMIN_REVALIDATE_DB_OBJECTS failed and the table remains in reorg pending state.

User response: Manually reorganize the table or ensure that the Q Capture or Capture program has EXECUTE privilege on the ADMIN_REVALIDATE_DB_OBJECTS procedure.

ASN0664W *program_name : program_identifier : The Q Capture or Capture program detected a transaction that exceeded the size size MB that is specified by the warntxsz parameter. The transaction ID is transaction_ID, transaction size is transaction_size MB, authorization ID is auth_ID, authorization token is auth_token, and plan name is plan_name.*

Explanation: The Q Capture or Capture program was started with the **warntxsz** parameter, which prompts the program to issue this warning message when transactions that are about to be replicated exceed a specified size. The size of a transaction has exceeded the **warntxsz** value. This might or might not be a problem.

User response: If this is an unexpected transaction, you can stop the program and use the provided transaction ID with the **ignore_transid** or **transid** runtime parameter to skip the transaction upon restart.

ASN0665E *program_name : program_identifier : The program did not start because it must be run from an APF-authorized library to authenticate Network Management Interface (NMI) client connection requests.*

Explanation: The program was started with the **nmi_enable** parameter to enable it to operate as an NMI server. The program must be APF-authorized in this situation.

User response: APF-authorize the program load library and restart the program.

ASN0666E *program_name : program_identifier : The replication NMI server is unable to get the NMI client security identity values. The return code is return_code. The error number is error_number. The reason is reason.*

Explanation: The program ioctl() SECIGET command did not get the NMI client security identity values because of a system error. The NMI client connection request is denied because the replication NMI server is unable to authenticate the requester's authorization ID.

User response: See the explanation of system error for the ioctl() C function to determine the cause of the error.

ASN0667E *program_name : program_identifier : The NMI client authentication function received SAF return code return_code from RACROUTE REQUEST=AUTH. The RACF return code is return_code. The RACF reason code is reason_code.*

Explanation: The NMI client connection request is denied because the replication NMI server is unable to authenticate the requester's authorization ID.

User response: Refer to the zOS Security Server RACROUTE Macro Reference (SA22-7692-13) for descriptions of the RACROUTE return and reason codes.

ASN0668I *program_name : program_identifier : The send queue queue_name for replication queue map queue_map_name has the following status: status.*

Explanation: A Q Capture QSTATUS command was issued to display the status for this send queue.

User response: This message is for your information only. No action is required.

ASN0669I *program_name : program_identifier : The receive queue queue_name for replication queue map queue_map_name has the following status: status. The current queue depth is queue_depth.*

Explanation: A Q Apply QSTATUS command was issued to display the status for this receive queue.

User response: This message is for your information only. No action is required.

ASN0670E *program_name : program_identifier : The send queue queue_name does not exist. The qstatus command is ignored.*

Explanation: The qstatus command was used for a send queue that does not exist.

User response: Check that the queue name or replication queue map name is correct and reissue the qstatus command.

ASN0671E *program_name : program_identifier : The receive queue queue_name does not exist. The qstatus command is ignored.*

Explanation: The qstatus command was used for a receive queue that does not exist.

User response: Check that the queue name or replication queue map name is correct and reissue the qstatus command.

ASN0672E *program_name : program_identifier : Column column_name in table table_owner.table_name for Q subscription, publication, or registration name is of an unsupported data type and cannot be replicated. The data type is data_type.*

Explanation: The Q Capture or Capture program does not support the data type of the column, even though the data type is supported by the source database. This situation can occur when the Q Capture or Capture program is at a lower version level than the database (for example a Version 9.7 Q Capture program running on a DB2 10 for z/OS subsystem).

User response: Remove the column from the table and from the Q subscription, publication, or registration, or upgrade the Capture or Q Capture program to the same level as the database that it works with.

ASN0673I *program_name : program_identifier : The Q Capture or Capture program stopped spilling to avoid exceeding the file size limit for the file file_name (transaction identifier trans_ID, authorization ID authid, correlation ID correlid, plan name plan_name).*

Explanation: The Q Capture or Capture program was

spilling an in-memory transaction to a file to relieve memory. Spilling was stopped because the file size reached the internal limit of 1 GB. Spilling might be resumed to a new file if memory is depleted again. No data is lost.

User response: This message is for your information only. No action is required.

ASN0674I *program_name : program_identifier : The program is checking for database connectivity.*

Explanation: When the program is running, it periodically checks for database connectivity. If the database is not available, the program continues running and waits until the database is available if the term=n option is specified in the invocation parameters of the program. After the database is available, the program connects and starts processing

User response: This message is for your information only. No action is required.

ASN0675I *program_name : program_identifier : The program generated seed value value for file file_name.*

Explanation: The program generated a seed and saved the seed value in the file. The program will generate a key from the seed and use the key to create an IPC queue.

User response: This message is for your information. No action is required.

ASN0676I *program_name : program_identifier : The program could not generate a seed to create an inter-process communication (IPC) key for path path_name. OSSE reason is reason.*

Explanation: The program is not able to generate a seed. The program generates a key from the seed and uses the key to create an IPC queue. The program tried to generate the seed after the ftok() function returned a key for a message queue that already exists.

User response: Remove the IPC file and its associated message queue manually by following the steps in, "Troubleshooting problems with IPC message queues."

ASN0677E *program_name : program_identifier : The Network Management Interface (NMI) socket name socket_name is not valid. The program will stop.*

Explanation: The program was started with the nmi_enable parameter so that it can operate as an NMI server. A valid NMI socket name must start with a forward slash character (/) and not be longer than 64 characters.

User response: Correct the NMI socket name and restart the program.

ASN0678E *program_name : program_identifier : The program is terminating because the DB2 for z/OS subsystem subsystem_name is not active and the program was started with the term=y parameter.*

Explanation: The **term=y** parameter causes the program to terminate when DB2 is unavailable.

User response: Take one of the following actions:

- Start the DB2 for z/OS subsystem and then restart the program with **term=y**.
- Start the program with **term=n**. The program will start, and then will wait for the subsystem to become active before the program begins to replicate data.

ASN0679E *program_name : program_identifier : The program is terminating because the DB2 for z/OS subsystem subsystem_name was started in restricted access mode and the program was started with the term=y parameter*

Explanation: The **term=y** parameter causes the program to terminate when DB2 is unavailable. Because the DB2 for z/OS subsystem was started in restricted access mode, the program terminated.

User response: Take one of the following actions:

- Start the DB2 for z/OS subsystem in full access mode and then restart the program with **term=y**.
- Start the program with **term=n**. The program will start, and will wait for the subsystem to run in full access mode before the program begins to replicate data.

ASN0685W *program_name : program_identifier : The TCP/IP GetAddrInfo socket call failed for service location location and service name name. The errno is number and the reason is reason.*

Explanation: The program invoked the GetAddrInfo socket call to determine the IP address that is associated with the specified host name and port or service.

This message can be returned for two reasons:

- An incorrect host name was specified in the EIF_CONN1 and EIF_CONN2 columns in the IBMQREP_APPLYPARMS table
- The correct host name was specified in the IBMQREP_APPLYPARMS table, but TCP/IP was unable to resolve that host name to an IP address

User response: Verify the following details:

- The EIF_CONN1 and EIF_CONN2 columns in the IBMQREP_APPLYPARMS table specify valid host names, IP addresses, service names, or ports
- TCP/IP is able to resolve the host name that is specified in the IBMQREP_APPLYPARMS table to an IP address

If changes were recently made to the name server or local host table, refresh the resolver.

ASN0686W *program_name : program_identifier : The TCP/IP create socket call failed for service location location and service name name. The errno is number and the reason is reason.*

Explanation: The program invoked the create socket call to create an endpoint for communication. The call failed for the specified reason.

User response: Determine the meaning of the reason that was issued by z/OS UNIX System Services or z/OS Communications Server.

ASN0687W *program_name : program_identifier : The TCP/IP SetSocketOpt socket call failed for service location location and service name name. The errno is number and the reason is reason.*

Explanation: The program invoked the SetSocketOpt socket call to set socket options. The call failed for the specified reason.

User response: Determine the meaning of the reason that was issued by z/OS UNIX System Services or z/OS Communications Server.

ASN0688W *program_name : program_identifier : The TCP/IP connect socket call failed for service location location, service name name. The errno is number and the reason is reason.*

Explanation: The program invoked the SetSocketOpt socket call to establish a connection with the server socket. The call failed for the specified reason.

User response: Determine the meaning of the reason that was issued by z/OS UNIX System Services or z/OS Communications Server.

ASN0689W *program_name : program_identifier : The TCP/IP send socket call failed for service location location and service name name. The errno is number and the reason is reason.*

Explanation: The program invoked the SetSocketOpt socket call to send data on the socket. The call failed for the specified reason.

User response: Determine the meaning of the reason that was issued by z/OS UNIX System Services or z/OS Communications Server.

ASN0691E *program_name : program_identifier : Error decoding log records for table table_name. Version: version_number*

Explanation: The Q Capture or Capture program could not decode log records for this table. The table structure (schema) as recorded in the IBMQREP_TABVERSION and IBMQREP_COLVERSION control tables is not consistent with the format of the log records.

User response: Follow these steps:

1. Manually deactivate the registration or Q subscription if it is active to allow capture to progress past this point in the log.
2. Run the REORG TABLESPACE utility for the table space that contains this table.
3. Activate the registration or Q subscription.

ASN0777I *pgmname : program_qualifier : Additional information message_text, reason code(s): rc1, rc2, rc3.*

Explanation: The *Additional information* shown in this message refers to an informational text message. The reason codes provide supplemental return code information related to this message text. If an informational code field is not applicable, it contains "*" (an asterisk).

User response: This message is for your information only, and no action is required.

ASN0888E *pgmname : program_qualifier : EEE error condition message_text, error code(s): rc1, rc2, rc3.*

Explanation: The *EEE error condition* shown in this message is the description of an EEE-specific error that occurred in the specified program with the specified qualifier (if displayed). The error codes provide supplemental information related to this message text. If an error code field is not applicable, it contains "*" (an asterisk).

User response: Use the information from the *EEE error condition* and from the specified error codes to determine the cause of the error. Contact IBM Software Support if you cannot resolve the error.

ASN0999E *pgmname : program_qualifier : Error condition message_text, error code(s): rc1, rc2, rc3.*

Explanation: The *Error condition* shown in this message is the description of an error that occurred in the specified program with the specified qualifier (if

displayed). The error codes provide supplemental information related to this message text. If an error code field is not applicable, it contains "*" (an asterisk).

User response: Use the information from the *Error condition* and from the specified error codes to determine the cause of the error. Contact IBM Software Support if you cannot resolve the error.

Chapter 37. ASN1000 - ASN1499

ASN1001E **APPLY** *apply_qualifier*. The Apply program encountered an SQL error. The **ERRCODE** is *error_code*. The **SQLSTATE** is *sqlstate*. The **SQLCODE** is *sqlcode*. The **SQLERRM** is *sqlerrm*. The **SQLERRP** is *sqlerrp*. The **server name** is *server_name*. The **table name** is *table_name*.

Explanation: An error occurred during the execution of an SQL statement.

User response: Refer to your database message reference for an explanation of the SQL error code.

ASN1002E **APPLY** *apply_qualifier*. The *table_name* could not be locked. **ERRCODE** is *error_code*, **SQLSTATE** is *sqlstate*, **SQLCODE** is *sqlcode*, **SQLERRM** is *sqlerrm*, **SQLERRP** is *sqlerrp*, **server name** is *server_name*, **table name** is *table_name*

Explanation: The Apply program could not lock the table.

User response: Refer to your database message reference.

ASN1003E **APPLY** *apply_qualifier*. The Apply program could not connect to the server *server*.

Explanation: The Apply program attempted to connect to the database and received a failing return code. There are many possible reasons why the Apply program could not connect to the database. For example, the Apply program would receive a failing return code if the database was down or too many users were accessing it.

User response: Look up the **SQLCODE** (from the Apply diagnostic log) in the DB2 messages and codes manual to determine why the connection failed. See the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for information about storing replication user IDs and passwords.

Refer to your database message reference for an explanation of the SQL error code.

ASN1006E **APPLY** *apply_qualifier*. The product registration module has unexpected content.

Explanation: The content of the registration module (ASNAPR61) for replication is not as expected for this version of DB2. No further use of the product is

possible until you provide the correct registration module.

User response: Verify that DB2 was installed without errors. If errors occurred, correct them and try again.

If DB2 was installed without error and you are correctly accessing the feature-registration module (ASNAPR61), contact IBM Software Support for assistance.

ASN1008E **APPLY** *apply_qualifier*. The subscription set with Apply qualifier *qualifier* and set name *set_name* is not defined correctly. **ERRCODE** is *error_code*.

Explanation: The subscription set is not defined correctly.

User response: Make sure that the subscription set has at least one member in the `ASN.IBMSNAP_SUBS_MEMBR` table or one `BEFORE_OR_AFTER` SQL statement in the `ASN.IBMSNAP_SUBS_STMTS` table.

ASN1009E **APPLY** *apply_qualifier*. There is no subscription set defined for Apply qualifier *qualifier*.

Explanation: There is no subscription set defined for Apply qualifier *qualifier*.

User response: Define at least one subscription set for Apply qualifier *qualifier*.

ASN1010E **APPLY** *apply_qualifier*. The Apply program could not insert row *row* into the audit trail table due to the following error: *error_code*.

Explanation: This is an SQL return code indicating that the audit trail table was not set up with the same structure as the `IBMSNAP_APPLYTRAIL` table.

User response: See the SQL reference information for your database and "Table structures for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1011E **APPLY** *apply_qualifier*. The copy request has incompatible source and target attributes. The SQL code is *error_code*.

Explanation: This is an SQL code indicating that the attributes of the target table must be compatible with the attributes of the source table.

User response: Refer to the SOURCE_STRUCTURE column in the register table for the compatibility of the source and target attributes.

ASN1012E **APPLY** *apply_qualifier*. **The source table structure is not valid. The error code is** *error_code*.

Explanation: This is an SQL return code indicating that the source table structure in the register table was not set up according to the SOURCE_STRUCTURE column in the register table.

User response: See "Table structures for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for valid SOURCE_STRUCTURE column values used in the IBMSNAP_REGISTER table.

ASN1013E **APPLY** *apply_qualifier*. **The target table structure is not valid. The error code is** *error_code*.

Explanation: The target table structure in the subscription-targets-member table was not valid.

User response: See "Table structures for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for valid target table structures.

ASN1014E **APPLY** *apply_qualifier*. **The Apply program could not find a source for the copy request because it could not find the change-data (CD) table. The error code is** *error_code*.

Explanation: The CD table was not defined in the IBMSNAP_REGISTER table because either the Apply program did not find the change data table name in the register table or the source table was not registered correctly.

User response: See "Table structures for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center and verify that the CD table is correctly defined in the IBMSNAP_REGISTER table.

ASN1015I **APPLY** *apply_qualifier*. **The Apply program is waiting for the Capture program at server** *server_name* **to advance the global SYNCHTIME. Verify that the Capture program is running.**

Explanation: This message is for your information only.

User response: Verify that the Capture program is running.

ASN1016I **APPLY** *apply_qualifier*. **Refresh copying has been disabled. The error code is** *error_code*.

Explanation: While attempting to perform a full refresh, the Apply program encountered a DISABLE_REFRESH column in the register table which was set on.

User response: Either turn off the DISABLE_REFRESH column or bypass the Apply program and perform a manual refresh.

ASN1017E **APPLY** *apply_qualifier*. **The Apply program could not find any target column names. The error code is** *error_code*.

Explanation: The Apply program could not find any columns in the IBMSNAP_SUBS_COLS table.

User response: Redefine the subscription set and subscription-set members. See "Subscribing to sources for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1018I **APPLY** *apply_qualifier*. **The Apply program is processing subscription set** *set_name(whos_on_first)*. **(set_number of total_sets).**

Explanation: This message is for your information only.

User response: This message is for your information only, and no action is required.

ASN1019E **APPLY** *apply_qualifier*. **The target table does not have any key columns. The error code is** *error_code*.

Explanation: The Apply program cannot find key column names in one of the columns requiring a unique index or primary key.

User response: Redefine the subscription set and subscription-set members. See "Subscribing to sources for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1020E **APPLY** *apply_qualifier*. **The Apply program could not reserve a storage block. The error code is** *error_code*.

Explanation: The Apply program could not obtain the required (memory) storage.

User response: Contact IBM Software Support.

ASN1021E **APPLY : *apply_qualifier* : The Apply program cannot read the work file *filename* because of a system error with ERRNO *error_code*.**

Explanation: The Apply program cannot read the work file because of a system error.

User response: See the explanation of system error for the C function to determine the cause of the error. On many operating systems, descriptions of ERRNO can be found in C header file named *errno.h*. Determine if the problem is caused by lack of space, and contact your system administrator to obtain what is needed.

ASN1022E **APPLY : *apply_qualifier* : The Apply program cannot write to the work file *filename* because of a system error with ERRNO *errno*. The error code is *error_code*.**

Explanation: Either the user ID not have the proper access authority for one or all of the files, or insufficient space remains after the Apply program writes to the target file.

User response: See the explanation of system error for the C function to determine the cause of the error. On many operating systems, descriptions of ERRNO can be found in C header file named *errno.h*. Determine whether the problem is caused by a lack of access authority or a lack of space, and contact your system administrator to obtain what is needed.

ASN1023E **APPLY : *apply_qualifier* : The Apply program cannot open the work file *filename* because of a system error with ERRNO *errno*. The error code is *error_code*.**

Explanation: The Apply program cannot open the work file because of a system error.

User response: See the explanation of system error for the C function to determine the cause of the error. On many operating systems, descriptions of ERRNO can be found in C header file named *errno.h*.

ASN1024E **APPLY : *apply_qualifier* : The Apply program cannot close the work file *filename* because of a system error with ERRNO *error-number*. The error code is *error_code*.**

Explanation: The Apply program cannot close the work file because of a system error.

User response: See the explanation of system error for the C function to determine the cause of the error. On many operating systems, descriptions of ERRNO can be found in C header file named *errno.h*.

ASN1025I **APPLY *apply_qualifier*. The Apply program completed processing for subscription set *set_name(whos_on_first)*. The return code is *return_code*.**

Explanation: This message is for your information only.

User response: This message is for your information only, and no action is required.

ASN1026I **APPLY *apply_qualifier*. The Apply program encountered an error while trying to bind. SQLSTATE is *sqlstate*, SQLCODE is *sqlcode*.**

Explanation: An error occurred during the execution of bind.

User response: Refer to your database message reference.

ASN1027E **APPLY *apply_qualifier*. There are too many large object (LOB) columns specified. The error code is *error_code*.**

Explanation: Too many large object (BLOB, CLOB, or DBCLOB) columns are specified for a subscription set member. The maximum number of columns allowed is 10.

User response: Remove the excess large object columns from the subscription set member.

ASN1028E **APPLY *apply_qualifier*. The before-image column for a key column was not found. The error code is *error_code*.**

Explanation: If you allow users to update columns in the source table that are part of the target key, (IS_KEY=Y in the IBMSNAP_SUBS_COLS table), you must specify that the Apply program use before-image values from the source table when it updates target key columns. Using before-image values enables Apply to search the target table for the old key value, delete the row, and insert a new row with the new key value. Two steps are required:

- You must define the source registration to capture the before-image values of the columns that make up the target key.
- You must select the subscription-set member option for the Apply program to use before-image values for target key updates. Selecting this option adds before-image column information to the IBMSNAP_SUBS_COLS table.

User response: Follow these steps:

1. Deactivate the subscription set.
2. In the Replication Center, open the Member Properties dialog for the subscription-set member.

3. On the Target-Table Index page, click the **Let the Apply program use before-image values to update target-key columns** check box.
4. Activate the subscription set.

ASN1029E **APPLY** *apply_qualifier*. The SQL statement of the subscription set named *set_name* with a *whos_on_first* value of *whos_on_first* did not execute successfully. The statement failed with **SQLCODE** *sqlcode* and **SQLSTATE** *sqlstate*. The apply program internal error code is *error_code*.

Explanation: The user-specified SQL statement did not execute successfully.

User response: Refer to the corresponding information in the IBMSNAP_APPLYTRAIL table and to the SQL manual of your database for detailed information.

ASN1031E **APPLY** *apply_qualifier*. The SQL statement is empty. The error code is *error_code*.

Explanation: The SQL statement is an empty string.

User response: Specify the SQL statement to be executed.

ASN1032E **APPLY** *apply_qualifier*. The Apply program log file could not be opened. The error code is *error_code*, and the return code is *return_code*.

Explanation: The Apply program could not open the log file.

User response: For more information on the return code, refer to the manual that describes troubleshooting for your particular operating system.

ASN1033E **APPLY** *apply_qualifier*. The Apply program could not write to the Apply log file. The error code is *error_code*, and the return code is *return_code*.

Explanation: The Apply program could not write to the log file.

User response: For more information on the return code, refer to the manual that describes troubleshooting for your particular operating system.

ASN1034I **APPLY** *apply_qualifier*. The Apply program initialization is successful.

Explanation: This message is issued at successful initialization of the Apply process.

User response: This message is for your information only, and no action is required.

ASN1035E **APPLY** *apply_qualifier*. The Apply program could not access the subscription columns table. The error code is *error_code*. The **SQLSTATE** is *sqlstate*. The **SQLCODE** is *sqlcode*. The **SQLERRM** is *sqlerrm*. The **SQLERRP** is *sqlerrp*. The server name is *server_name*. The table name is *table_name*.

Explanation: An error occurred during the execution of an SQL statement.

User response: Refer to your database message reference for SQL.

ASN1036E **APPLY** *apply_qualifier*. The column type *col_type* for expression *expression* is not valid. The error code is *error_code*.

Explanation: The value for the COL_TYPE column in the subscription columns table is not valid.

User response: Change the value to A, B, C, D, F, L, or R.

ASN1038E **APPLY** *apply_qualifier*. No column names or expressions were specified in the IBMSNAP_SUBS_COLS table.

Explanation: Column names or expressions for a copy statement must be specified.

User response: See "Subscribing to sources for SQL Replication" in the DB2 Information Center for more information about requirements for subscription definitions.

ASN1039E **APPLY** *apply_qualifier*. The Apply program plan, *plan_name*, could not be opened. The error code is *error_code*. The return code is *return_code*. The reason code is *reason_code*.

Explanation: The Apply program plan could not be opened.

User response: Refer to the Apply for z/OS Program Directory.

ASN1040E **APPLY** *apply_qualifier*. The Apply program encountered an z/OS error. The error code is *error_code*, and the return code is *return_code*.

Explanation: Execution of a z/OS system operation failed.

User response: Refer to your z/OS system library information.

ASN1041I **APPLY** *apply_qualifier*. **The Apply program was started using subsystem name:** *subsystem*.

Explanation: This message informs you that the Apply program started using the specified subsystem name.

User response: This message is for your information only, and no action is required.

ASN1042W **APPLY** *apply_qualifier*. **There are too many invocation parameters.**

Explanation: The number of parameters you specified when you invoked the Apply program exceeds the maximum allowed.

User response: Refer to the Capture and Apply chapter for your operating system for information on the appropriate number of invocation parameters.

ASN1043E **APPLY** *apply_qualifier*. **There is already one Apply instance running with this Apply program qualifier *qualifier*. The error code is *error_code*, and the reason code is *reason_code*.**

Explanation: Verification attempt failed.

User response: Make sure that only one instance of the Apply program with the specified Apply qualifier is running under this user ID on this subsystem or database.

ASN1044I **APPLY** *apply_qualifier*. **The Apply program will become inactive for *number* minutes and *number* seconds.**

Explanation: The Apply program becomes inactive in the following scenarios:

- At the end of an active Apply cycle, the Apply program will sleep for an amount of time calculated according to the following formula:
[SLEEP_MINUTES] - [time spent in the cycle]

For example, if SLEEP_MINUTES is set to '1' and the Apply cycle lasts 10 seconds, the Apply program will be inactive for 50 seconds at the end of the cycle.

- When the Apply program does not find any subscription data to process, the Apply program will automatically sleep for 5 minutes.

You can configure whether this message is returned when the Apply program becomes inactive by specifying the inamsg parameter with the asnapply command.

User response: This message is for your information only, and no action is required.

ASN1045I **APPLY** *apply_qualifier* : **The Apply version *version_number* program was started using database *database_name*.**

Explanation: This message informs you from which database the Apply program is running.

User response: This message is for your information only. No action is required.

ASN1047I **APPLY** *apply_qualifier*. **There are too many columns specified. The error code is *error_code*.**

Explanation: There are too many columns specified for a member in the subscription.

User response: The user must reduce the number of columns specified for the member in the subscription. The maximum number of columns supported is determined by the total length of all the column names. More columns can be specified if the length of the column names is shorter.

ASN1048E **APPLY** *apply_qualifier*. **The execution of an Apply cycle failed. See the Apply trail table for full details: *text***

Explanation: An Apply cycle failed. In the message, *text* identifies the *target_server*, *target_owner*, *target_table*, *stmt_number*, and *cntl_server*.

User response: Check the APPERRM fields in the audit trail table to determine why the Apply cycle failed.

ASN1049E **APPLY** *apply_qualifier*. **The Apply program encountered a system error. The error code is *error_code*. The return code is *return_code*.**

Explanation: Execution of a system operation failed.

User response: Refer to the system library information for your operating system.

ASN1050E **APPLY** *apply_qualifier*. **The Apply program encountered an operation that is not valid while updating the target table. The error code is *error_code*. The operation to be applied is *operation*.**

Explanation: The operation field of a row fetched from the source table is not valid.

User response: Contact IBM Software Support.

ASN1051W **APPLY** *apply_qualifier* : **The Apply program detected a gap in changed data between the source table *table_owner.table_name* and the target table. The error code is *error_code*.**

Explanation: The Apply program detected that the Capture program skipped logged changes or deleted rows from CD tables before the Apply program could copy the data. For example, the Capture program was cold started or retention limit pruning occurred.

User response: Look for message ASN0100I to determine if the Capture program was started recently. If so, look for ASN0529I to determine if the value of the STARTMODE parameter was COLD. If so, the Apply program performs a full refresh of target tables unless full refresh has been disabled. If you find that the Capture program performed retention limit pruning of a CD table, the corresponding subscription may need a full refresh to synchronize the target table with the source.

ASN1052E **APPLY** *apply_qualifier*. **The Apply program could not find the ASNLOAD program.**

Explanation: The Apply program cannot find the ASNLOAD program in the current directory.

User response: Make sure that ASNLOAD is in the directory from which you are invoking the Apply program.

ASN1053E **APPLY** *apply_qualifier*. **The execution of the ASNLOAD exit routine failed. The return code is** *return_code*.

Explanation: The ASNLOAD exit routine detected an error and passed the error information back to the Apply program. The following values are valid return codes:

98

An unexpected error has occurred (The ASNLOAD exit routine has failed with an unexpected error. No processing will be performed.)

99

DB2 pwdfile keyword supplied - password file not found (The pwdfile parameter was passed, but no password file was found. This is an error, and no connections or other processing will be performed.)

100

The CONNECT statement in the exit routine specified a user ID and password that use values supplied in the encrypted Apply password file. The password file and a user ID and password combination for the DB2 server were found, but the connection failed.

101

DB2 connect without a user/using phrase failed - no pwdfile found (A connection

without a user/using phrase was made because no password file was provided. The connection failed.)

102

DB2 connect without a user/using phrase failed - pwdfile found, no entry (A connection without a user/using phrase was made, because no server entry was found in the pwdfile for the DB2 server. The connection failed.)

103

DB2 connect with a user/using phrase failed - uid/pwd from asnload.ini used (A connection with a user/using phrase was made using values supplied in the asnload.ini file. This file and a userid/password combination for the DB2 server were found, but the connection failed.)

104

DB2 connect without a user/using phrase failed - no asnload.ini found (A connection without a user/using phrase was made, because no asnload.ini file was found. The connection failed.)

105

DB2 connect without a user/using phrase failed - no uid/pwd found for server (A connection without a user/using phrase was made. The asnload.ini file was found, but no uid/pwd combination was provided. The connection failed.)

106

User specified LOADX_TYPE = 2, no user code provided (The value of LOADX_TYPE in the table ASN.IBMSNAP_SUBS_MEMBR was set by the user to the value of 2, indicating that the user was supplying custom code in the ASNLOAD exit routine. However this code was not found, and the ASNLOAD exit routine failed when the Apply program passed a LOADX_TYPE value of 2.)

107

DB2 import utility failed (The import utility failed to execute. The SQL code returned by the utility is passed as the reason code.)

108

DB2 export utility failed (The export utility failed to execute. The SQL code returned by the utility is passed as the reason code.)

109

| | | |
|---|---|--|
| 110 | DB2 load utility failed (The load utility failed to execute. The SQL code returned by the utility is passed as the reason code.) | ASN1054E <i>APPLY apply_qualifier. The Apply program could not find a row in the IBMSNAP_REGISTER or IBMSNAP_PRUNCNTL table that corresponds to the subscription set member with a set name set_name, for source owner src_ownr, source table src_tbl, and source view qualifier src_view_qual.</i> |
| 111 | DB2 load utility failed - invoked as crossload (The load utility failed to execute. The load utility was invoked with the load from cursor option. The SQL code returned by the utility is passed as the reason code.) | Explanation: The source table registration is incorrect or incomplete. User response: Drop and redefine the registration. |
| 112 | The LOADX_TYPE was set to an not valid value (The ASNLOAD exit routine was invoked with a LOADX_TYPE value that was set by the user. The LOADX_TYPE value is not valid for this environment, and the ASNLOAD exit routine failed.) | ASN1055E <i>APPLY apply_qualifier. The Apply program could not find the prune control information for source owner src_ownr, source table src_tbl, source view qualifier src_view_qual, target owner tgt_ownr, and target table tgt_tbl.</i> |
| 113 | LOADX_TYPE 3 requires a nickname for select (The ASNLOAD exit routine failed. The ASNLOAD exit routine was invoked with a LOADX_TYPE value that was set by the user. The LOADX_TYPE value is not valid for this environment unless a nickname is created for the remote DB2 source table and stored in the ASN.IBMSNAP_SUBS_MEMBR table.) | Explanation: The source table registration is incorrect. User response: Drop the subscription and redo it. |
| 114 | LOADX_TYPE 4 is incompatible with target table (The ASNLOAD exit routine failed. The ASNLOAD exit routine was invoked with a LOADX_TYPE set by the user. The LOADX_TYPE value is not valid for this environment, because the target table cannot be serviced by the DB2 for Linux, UNIX, and Windows load utility.) | ASN1056E <i>APPLY apply_qualifier. The Apply program could not connect to the server due to lack of user ID/password. The error code is error_code.</i> Explanation: The Apply program could not find the password and user ID to connect to the server. User response: Make sure that the Apply program password file exists. The Apply program password file resides in the same directory from which you start the Apply program. If you are using DB2 Satellite Edition, make sure that the password and user ID are defined on the client systems. |
| 115 | LOADX_TYPE 5 is incompatible with target table (The ASNLOAD exit routine failed . The ASNLOAD exit routine was invoked with a LOADX_TYPE set by the user. The LOADX_TYPE value is not valid for this environment, because the target table cannot be serviced by the DB2 import utility.) | ASN1057E <i>APPLY apply_qualifier. The Apply program could not read the password in the Apply password file. The error code is error_code.</i> Explanation: The Apply program found no password. User response: If you want to use the AUTHENTICATION=SERVER scheme, you must provide a password, as described in the Apply program section in the Capture and Apply chapter for your operating system. |
| | The ASNDLCOPY exit routine has failed (The ASNLOAD exit routine called the ASNDLCOPY exit routine, because there were DATALINK columns for the subscription-set member. The ASNDLCOPY exit routine failed; therefore, the process that loads this subscription-set member also failed.) | ASN1058E <i>APPLY apply_qualifier. The Apply program could not close the password file. The error code is error_code.</i> |
| User response: Check the return code and the corresponding explanation. Check for additional information in the ASNLOAD message file and in the message files generated by the DB2 utility, if applicable. | | Explanation: The Apply program could not close the password file. User response: Contact IBM Software Support. |

ASN1059E *APPLY apply_qualifier. The Apply program detects syntax that is not valid for line line in the password file. The error code is error_code.*

Explanation: The Apply program could not recognize a line in the password file.

User response: Correct the syntax error in the password file. See, "asnpwd: Creating and maintaining password files" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1060E *APPLY apply_qualifier. The dynamic allocation for the temporary work file failed. The error code is error_code.*

Explanation: A system error was encountered during dynamic allocation.

User response: Contact IBM Software Support.

ASN1061E *APPLY apply_qualifier. The specified keyword parameter is not valid. The error code is error_code.*

Explanation: An invocation parameter that is not valid has been specified and has been ignored by the Apply program.

User response: Correct the invocation parameter. See "asnapply: Starting Apply" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1062W *APPLY apply_qualifier. The Apply program must use SELECT and INSERT statements to perform a full refresh of this subscription-set member. The following information pertains to this subscription-set member: the set name is set_name, the source owner is source_owner, the source table is source_table, the source view qualifier is source_view_qual, the target owner is target_owner, and the target table is target_table.*

Explanation: The ASNLOAD exit routine cannot detect a user-specified LOADX_TYPE value, and no utilities are available to process this subscription-set member. Therefore, the ASNLOAD exit routine passes full refresh control back to the Apply program. The ASNLOAD exit routine does not currently support and is not able to process some target table types (such as the Sybase and MS SQL Server target tables).

User response: This message is for your information only, and no action is required. However, you can set the value of the LOADX_TYPE to 1 for these subscription-set members in order to avoid unnecessary processing by the ASNLOAD exit routine.

ASN1063E *APPLY apply_qualifier. A subscription set cannot have more than 200 members. The error code is error_code.*

Explanation: The number of subscriptions has exceeded the maximum allowed number of 200.

User response: Remove excess members from the subscription set.

ASN1064W *APPLY apply_qualifier. The Apply program cannot perform a full refresh for the subscription set named set_name, because the Capture program for this source has not yet been cold started.*

Explanation: The Apply program cannot attempt a full refresh for the subscription set, because the Capture program for this source has never been cold started and is not ready to process the CAPSTART signals are be inserted by the Apply program.

User response: Start the Capture program for this source.

ASN1065E *APPLY apply_qualifier. The Apply program cannot process data for the subscription set set_name because one or more of the registrations for the source table are stopped.*

Explanation: The Apply program cannot process data for this subscription set, because at least one of the registrations has STATE='S' in the IBMSNAP_REGISTER table.

The Capture program stops a registration if there is a problem with the registration that requires your intervention. Refer to the STATE_INFO column in the IBMSNAP_REGISTER table for error information. The integrity of the captured data for the registration might be compromised, and the Apply program must perform a full-refresh. This problem might occur if the registered source table was altered with data capture none.

User response: Fix the stopped registrations using the information from the error messages. Reactivate the registrations. When you reactivate the registrations, the Apply program performs a full refresh.

ASN1066E *APPLY apply_qualifier. An internal Apply program error occurred. The error code is error_code.*

Explanation: An internal Apply program error occurred.

User response: Contact IBM Software Support.

ASN1067E **APPLY** *apply_qualifier*. **The Apply program has detected update conflicts and compensated rejected transactions. See the unit-of-work table for details. The error code is** *error_code*.

Explanation: More than one application updated the same row in a table from different locations. Some transactions have been rejected and compensated.

User response: See the SQL Replication table structures documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1068E **APPLY** *apply_qualifier*. **The Apply program has deactivated the subscription due to a constraint violation. The error code is** *error_code*.

Explanation: A constraint violation was detected when copying data from the source table to a target table. The Apply program has terminated and the subscription has been deactivated.

User response: Correct the constraint error and reactivate the subscription.

ASN1070E **APPLY** *apply_qualifier*. **The Apply program could not lock the target table. The ERRCODE is** *error_code*. **The SQLSTATE is** *sqlstate*. **The SQLCODE is** *sqlcode*. **The SQLERRM is** *sqlerrm*. **The SQLERRP is** *sqlerrp*. **The server name is** *server_name*. **The table name is** *table_name*.

Explanation: The Apply program could not lock the target tables before it was to check update conflicts.

User response: Verify that all the target tables are available before restarting Apply.

ASN1071E **APPLY** *apply_qualifier*. **The Apply program could not reposition the work file. The error code is** *error_code*.

Explanation: The Apply program has detected an error while reading the temporary work file.

User response: Contact IBM Software Support.

ASN1072E **APPLY** *apply_qualifier*. **The Apply program could not find the ASNDONE program.**

Explanation: The Apply program could not find the user exit program, ASNDONE.

User response: Verify that the ASNDONE program is located in the correct directory.

ASN1073E **APPLY** *apply_qualifier*. **The execution of the ASNDONE program failed. The return code is** *return_code*.

Explanation: An error occurred while calling the user exit program, ASNDONE.

User response: Contact IBM Software Support.

ASN1074E **APPLY** *apply_qualifier*. **The Apply program could not find the ASNDLCOPY program.**

Explanation: The Apply program did not find the ASNDLCOPY program in the current search path.

User response: Add the ASNDLCOPY program to the search path and run the Apply program again.

ASN1075E **APPLY** *apply_qualifier*. **The ASNDLCOPY program failed. The return code is** *return_code*. **Additional information can be found in the ASNDL file**

Explanation: The ASNDLCOPY program detected an error and passed the error information back to the Apply program. The following values are valid return codes:

98

Unexpected error occurred.

99

The arguments passed to the ASNDLCOPY program are not valid.

100

Unable to allocate memory.

101

Unable to open the ASNDLSRVMAP configuration file.

102

The number of entries in the ASNDLSRVMAP configuration file exceeds the maximum limit.

103

An entry that is not valid has been found in the ASNDLSRVMAP configuration file.

104

No user login information was found in the ASNDLUSER configuration file for a given file server.

105

An entry that is not valid has been found in the ASNDLPARM configuration file.

106

Unable to open the ASNDLUSER configuration file.

107

An entry that is not valid has been found in the ASNDLUSER configuration file.

108

An I/O error occurred when reading from the input file.

109

An entry that is not valid has been found in the input file.

110

Unable to open the input file.

111

Unable to open the result file.

112

An I/O error occurred when writing to the result file.

113

An error occurred when initializing the control channel of the FTP protocol.

114

An error occurred when sending data through the control channel.

115

Unable to log on to the file server with the given user and password.

116

The copy daemon has not yet started.

117

An error occurred when initializing the data channel of the FTP protocol.

118

Unable to retrieve the file from the source file server.

119

Unable to store the file on the target file server.

120

An error occurred when transferring files in the passive mode.

121

Cannot find the path mapping for the given file reference.

122

An error occurred when executing the FTP BINARY command.

123

An error occurred when executing the FTP SIZE command.

124

An error occurred when executing the FTP MODTIME command.

125

An error occurred when executing the FTP SITE UMASK command.

126

An error occurred when executing the FTP SITE TOUCH command.

127

An error occurred when executing the FTP SITE CHMOD command.

User response: Check the return code and its corresponding meaning. The return code is based on the sample ASNDLCOPY program that is shipped with the product. Additional information is provided in the log file.

ASN1076E The Apply program cannot read the format of the result file that was generated by the ASNDLCOPY program.

Explanation: The result file that was generated by the ASNDLCOPY program is not in an expected format.

User response: If you modified the ASNDLCOPY program, check that your changes are not causing the invalid format. If your changes are not the cause of the problem, check that your machine has enough space for the result file.

ASN1077E APPLY *apply_qualifier*. The Apply program encountered an DATALINK column value that is not valid while updating the target table. The error code is *error_code*.

Explanation: The DATALINK column field of a row fetched from the source table is not valid.

User response: Contact IBM Software Support.

ASN1078E APPLY *apply_qualifier*. The ASNDLCOPY program was terminated by the signal *signal_number*. Additional information can be found in the *filename* file.

Explanation: The ASNDLCOPY program terminated abnormally by the given signal.

User response: Check the specified log file for the cause of the error. If you modified the ASNDLCOPY program and the signal is generated by the modified

code, fix the code and rerun. Otherwise, contact IBM Software Support.

ASN1079E **APPLY** *apply_qualifier*. **MEMBER_STATE** is invalid for set *set_name* for **WHOS_ON_FIRST** *whos_on_first*, **source owner** *source_owner*, **source table** *source_table*, **source view qualifier** *source_view_qual*, **target owner** *target_owner*, and **target table** *target_table*.

Explanation: There was an invalid update made to the MEMBER_STATE column in the IBMSNAP_SUBS_MEMBR table. Valid values for this column are: 'N', 'L', 'S', or 'D'.

User response: Update the MEMBER_STATE column with a valid value and start the Apply program again. Refer to the IBMSNAP_SUBS_MEMBR table for details about the values in the MEMBER_STATE column.

ASN1080E **APPLY** *apply_qualifier* : No columns are defined for target table *table_owner.table_name* in set *set_name* (**whos_on_first** *whos_on_first_value*). The error code is *error_code*.

Explanation: The Apply program's ASN.IBMSNAP_SUBS_COLS control table contains no information about the columns in the specified target table. Column information may have been manually removed during editing of the SQL script generated by the Replication Center or ASNCLP command-line program.

User response: Drop and recreate the subscription. If you know of a reason that the subscription cannot be dropped, you need to insert column data manually into the ASN.IBMSNAP_SUBS_COLS table.

ASN1097I **APPLY** *apply_qualifier*. The Apply program stopped.

Explanation: The error reported previously caused the Apply program to stop.

User response: Fix the error reported before this message.

ASN1207E **APPLY** *apply_qualifier*. The subscription for *subscription* was not activated.

Explanation: The selected subscription is inactive.

User response: Either activate the subscription or select another one.

ASN1210E **APPLY** *apply_qualifier*. An Apply qualifier must be specified following the keyword -q.

Explanation: You must specify an Apply qualifier following the keyword -q.

User response: Specify an Apply qualifier following the keyword -q.

ASN1212E **APPLY** *apply_qualifier*. A read-only set name *set_name* is found following the keyword *keyword*.

Explanation: A read-only set name was specified following the keyword U or D.

User response: Specify only replica for the keywords U and D.

ASN1221I **APPLY** *apply_qualifier*. Set *set_name* has been successfully refreshed with *number* rows at *time*.

Explanation: This message is for your information only.

User response: This message is for your information only, and no action is required.

ASN1242E **APPLY** *apply_qualifier*. An SQL error occurred. **ERRCODE** is *error_code*, **SQLSTATE** is *sqlstate*, **SQLCODE** is *sqlcode*, **SQLERRM** is *sqlerrm*, **SQLERRP** is *sqlerrp*, **table name** is *table_name*.

Explanation: This message is for your information only.

User response: This message is for your information only, and no action is required.

ASN1243E **APPLY** *apply_qualifier*. There is no eligible subscription in the ASN.IBMSNAP_SUBS_SET table.

Explanation: Either a subscription set has not been selected or the apply qualifier is not valid.

User response: Verify the subscription names and apply qualifier.

ASN1304E **APPLY** *apply_qualifier*. The ASNSAT program terminated due to a Capture error.

Explanation: The Capture program returned an error.

User response: Determine the error from the Capture log file.

ASN1305E **APPLY** *apply_qualifier*. The ASNSAT program terminated due to an Apply error.

Explanation: The Apply program returned an error.

User response: Determine the error from the Apply log file.

ASN1310E **APPLY** *apply_qualifier*. The ASNSAT program encountered a system error while attempting to invoke the Capture program. Return code is *return_code*.

Explanation: An operating system error occurred while calling ASNCAP.

User response: Make sure that the Capture program is in the execution path.

ASN1311E **APPLY** *apply_qualifier*. The ASNSAT program encountered a system error while attempting to invoke the Apply program. Return code is *return_code*.

Explanation: An operating system error occurred while calling ASNAPPLY.

User response: Make sure that the Apply program is in the execution path.

ASN1312E **APPLY** *apply_qualifier*. The environment variable that specifies the default target server, DB2DBDFT, is not set.

Explanation: The target server name was not specified, and the ASNSAT program could not determine the default database name from the DB2DBDFT variable.

User response: Specify the target server name following the -t keyword.

ASN1314E **APPLY** *apply_qualifier*. An SQL error occurred while ASNSAT was getting the default Apply qualifier. SQLSTATE is *sqlstate*, SQLCODE is *sqlcode*.

Explanation: The user did not specify the Apply qualifier. The ASNSAT program encountered an error while retrieving the USER special register.

User response: Specify the Apply qualifier following the -q keyword.

ASN1315E **APPLY** *apply_qualifier*. Cannot connect to database server. SQLSTATE is *sqlstate*, SQLCODE is *sqlcode*.

Explanation: An error occurred while attempting to connect to the target database.

User response: Refer to your database message reference.

ASN1316E **APPLY** *apply_qualifier*. ASNSAT encountered an error while trying to bind. The SQLSTATE is *sqlstate*, SQLCODE is *sqlcode*.

Explanation: An error occurred while attempting to auto bind.

User response: Make sure that the bind file exists in the sqllib\bnd directory.

ASN1317E **APPLY** *apply_qualifier*. An SQL error occurred while ASNSAT was getting the CD_TABLE value from ASN.IBMSNAP_REGISTER table. SQLSTATE is *sqlstate*, SQLCODE is *sqlcode*.

Explanation: An SQL error occurred while selecting from the register table.

User response: Refer to your database message reference.

ASN1318E **APPLY** *apply_qualifier*. An SQL error occurred while ASNSAT attempted to get the DB2 node type. SQLSTATE is *sqlstate*, SQLCODE is *sqlcode*.

Explanation: An error occurred while retrieving the node type configuration parameter.

User response: Refer to your database message reference.

ASN1319I **APPLY** : *apply_qualifier* : The program is rebinding the package that is required for full refresh with the option CONCURRENTACCESSRESOLUTION WAIT FOR OUTCOME.

Explanation: To ensure that the full refresh of the target table successfully selects all data from the source server, the Apply program must bind the SQL packages with this option. The option forces the SELECT operation to wait until all in-progress transactions that modify the source table are completed before beginning to retrieve the rows.

User response: This message is for your information only. No action is required.

ASN1320W **APPLY** : *apply_qualifier* : The program was unable to bind the packages that are required for full refresh. Bind failed on package *package_name* from the file *path_filename*. The SQLCODE *sqlcode* was returned.

Explanation: The program encountered an error when it attempted to bind a required package for full refresh.

User response: Interpret the SQLCODE of the failed

bind call. If you need to ensure that no data is lost during a full refresh, fix the problem and then restart the Apply program. You can manually bind the required packages by connecting to the source database and running the following command:

```
db2 bind @qapplycs.lst
      CONCURRENTACCESSRESOLUTION WAIT_FOR_OUTCOME
```

If you are only performing differential refresh, you can ignore this warning.

ASN1321I *APPLY* *apply_qualifier* : Number rows were committed to the target table *table_owner.table_name* during a full refresh operation.

Explanation: The Apply program committed the indicated number of rows to the target table during the most recent commit of the full refresh. You can control the number of rows in each commit by using the `nickname_commit_ct` parameter.

User response: This message is for your information only. No action is required.

ASN1322W *APPLY* : *apply_qualifier* : The Apply program parameter `MONITOR_ENABLED` is set to Y, but a control table that is required to use the monitoring functionality, `ASN.IBMSNAP_APPLYMON`, does not exist. No Apply status information will be saved.

Explanation: When `MONITOR_ENABLED` is set to Y, the Apply program inserts information about its status into the `ASN.IBMSNAP_APPLYMON` table. The table was not found and so no status information can be saved.

User response: Take one of the following actions:

- Create the `ASN.IBMSNAP_APPLYMON` table on the Apply control server. You can use the following SQL statements:

```
CREATE TABLE ASN.IBMSNAP_APPLYMON
  (MONITOR_TIME TIMESTAMP NOT NULL,
   APPLY_QUAL CHAR(18) NOT NULL,
   WHOS_ON_FIRST CHAR(1),
   STATE SMALLINT,
   CURRENT_SET_NAME CHAR(18),
   CURRENT_TABOWNER VARCHAR(128),
   CURRENT_TABNAME VARCHAR(128)
  );
```

```
CREATE INDEX IXIBMSNAP_APPLYMON ON
  ASN.IBMSNAP_APPLYMON(MONITOR_TIME,
                       APPLY_QUAL,
                       WHOS_ON_FIRST);
```

Restart the Apply program with `MONITOR_ENABLED=Y` after the table is created.

- Continue without Apply status information. Apply issues the warning message once and keeps processing subscriptions if the table does not exist. You can change the value of the `MONITOR_ENABLED` column in the `ASN.IBMSNAP_APPPARMS` table to N and Apply will use this value the next time it is restarted.

ASN1323E *APPLY* *apply_qualifier* : The value of the `ARCH_LEVEL` column in the `ASN.IBMSNAP_SUBS_SET` table on the Apply control server is not at the expected level of 1001 for the subscription set. The subscription set was deactivated. Apply control server: *control_alias*. Subscription-set name: *set_name*.

Explanation: When you migrate to SQL Replication Version 10 and set Capture compatibility to 10xx on the source server, you must set the value of the `ARCH_LEVEL` column in the `ASN.IBMSNAP_SUBS_SET` table to 1001. The variable `xx` identifies the release within Version 10.

User response: Follow the steps in the SQL Replication Version 10 migration information to migrate the Apply control server, and then activate the subscription set.

ASN1324E *APPLY* *apply_qualifier* : The Apply program cannot work with the Capture program on source server *source_alias* because the Capture control tables have a value of *value* in the `COMPATIBILITY` column of the `IBMSNAP_CAPPARMS` table for the subscription set *set_name*. The subscription set was deactivated.

Explanation: The Capture server is at Version 10, which requires the Apply program and Apply control server to also be at Version 10 to support new restart information and 16-byte columns for log sequence numbers.

User response: Follow these steps:

- Upgrade the Apply program and Apply control server to Version 10 or above.
- Migrate the Apply control tables to Version 10.
- Reactivate the subscription set.

ASN1325E *APPLY apply_qualifier* : The target CCD table *table_owner.table_name* in the subscription set *set_name* was not upgraded to include 16-byte columns to hold log sequence numbers (LSN). The Capture program is at Version 10 with a value of 1001 in the COMPATIBILITY column of the IBMSNAP_CAPPARMS table. The subscription set was deactivated.

Explanation: When the Capture program is at Version 10, target CCD tables require 16-byte LSN columns (IBMSNAP_COMMITSEQ, IBMSNAP_INTENTSEQ) to hold log sequence number details from the Capture program.

User response: Follow the steps in the SQL Replication Version 10 migration information to upgrade target CCD tables to 16-byte LSN columns and reactivate the subscription set.

ASN1326E *APPLY apply_qualifier* : The Apply program does not support the source CCD table *table_owner.table_name* on source server *source_alias* for subscription set *set_name* because the source server uses 16-byte columns for log sequence numbers (LSN). The subscription set was deactivated.

Explanation: Because the Apply program is older than Version 10, it does not support 16-byte LSN columns. The source CCD table was upgraded to 16-byte LSN columns, and so the Apply program cannot work with this table.

User response: Upgrade the Apply program and Apply control server to Version 10 or above to support 16-byte LSN columns and reactivate the subscription set.

ASN1327E *APPLY apply_qualifier* : For subscription set *set_name*, a value of 1001 is set for ARCH_LEVEL in the ASN.IBMSNAP_SUBS_SET table, but on source server *source_alias* either the COMPATIBILITY column does not exist in the IBMSNAP_CAPPARMS table or the value in the COMPATIBILITY column is 0801. The compatibility value must be updated to 1001. The subscription set was deactivated.

Explanation: When the ARCH_LEVEL value for a subscription set is 1001, the Capture compatibility value must be changed to 1001.

User response: Run the Version 10 Capture migration and compatibility script, and reactivate the subscription set.

ASN1328E *APPLY apply_qualifier* : The subscription set *set_name* was deactivated because the parameter table *table_name* on source server *source_alias* is empty or contains more than one row.

Explanation: The program parameter table must have exactly one row.

User response: Follow these steps:

1. Make sure that the parameter table is in the right format and contains exactly one row. You can find the correct DDL for the table in the migration scripts in the sqllib/samples/repl/mig10 directory.
2. Reactivate the subscription set.

Chapter 38. ASN1500 - ASN1999

ASN1500I The replication action *action_name* started at *timestamp* with architecture level *architecture_level*. The Capture server is *capture_serveralias* and the Capture schema is *capture_schema*.

Explanation: Valid values for *action_name* are *Create Capture server control tables* and *Drop Capture server control tables*.

User response: This message is for your information only, and no action is required.

ASN1501I The replication action *action_name* started at *timestamp* with architecture level *architecture_level*. The Capture server is *capture_serveralias*, the remote server is *remote_servername*, and the Capture schema is *capture_schema*.

Explanation: Valid values for *action_name* are *Create Capture server control tables* and *Drop Capture server control tables*.

User response: This message is for your information only, and no action is required.

ASN1502I The replication action *action_name* started at *timestamp* with architecture level *architecture_level*. The Apply control server is *apply_serveralias*.

Explanation: Valid values for *action_name* are *Create Apply server control tables* and *Drop Apply server control tables*.

User response: This message is for your information only, and no action is required.

ASN1503I The replication action *action_name* started at *timestamp*. The Capture server is *capture_serveralias*, the Capture schema is *capture_schema*, the source owner is *source_owner*, and the source table, view, or nickname is *source_table*.

Explanation: Valid values for *action_name* are *Create Registration*, *Drop Registration*, *Alter Registration*, *Add Registration*, and *Promote Registration*.

User response: This message is for your information only, and no action is required.

ASN1504I The replication action *action_name* started at *timestamp*. The Capture server is *capture_serveralias*, the remote server is *remote_server*, the Capture schema is *capture_schema*, the source owner is *source_owner*, and the source table, view, or nickname is *source_table*.

Explanation: Valid values for *action_name* are *Create Registration* and *Drop Registration*.

User response: This message is for your information only, and no action is required.

ASN1505I The replication action *action_name* started. The subscription set information follows: the Apply control server is *control_server*, the Apply qualifier is *apply_qualifier*, the set name is *set_name*, the target server is *target_server* for remote server *remote_servername*, the Capture server is *capture_server* for remote server *remote_servername*, and the Capture schema is *capture_schema*.

Explanation: Valid values for *action_name* are *Create Subscription Set*, *Drop Subscription Set*, *Alter Subscription Set*, and *Promote Subscription Set*.

User response: This message is for your information only, and no action is required.

ASN1506I The replication action *action_name* started at *timestamp*. The subscription set information follows: the Apply control server is *control_server*, the Apply qualifier is *apply_qualifier*, the set name is *set_name*, the target server is *target_server*, the Capture server is *capture_server* for remote server *remote_servername*, and the Capture schema is *capture_schema*.

Explanation: Valid values for *action_name* are *Create Subscription Set*, *Drop Subscription Set*, *Alter Subscription Set*, and *Promote Subscription Set*.

User response: This message is for your information only, and no action is required.

ASN1507I The replication action *action_name* started at *timestamp*. The subscription set information follows: the Apply control server is *control_server*, the Apply qualifier is *apply_qualifier*, the set name is *set_name*, the target server is *target_server* for remote server *remote_server*, the Capture server is *capture_server*, and the Capture schema is *capture_schema*.

Explanation: Valid values for *action name* are *Create Subscription Set*, *Drop Subscription Set*, *Alter Subscription Set*, and *Promote Subscription Set*.

User response: This message is for your information only, and no action is required.

ASN1508I The replication action *action_name* started at *timestamp*. The subscription set information follows: the Apply control server is *control_server*, the Apply qualifier is *apply_qualifier*, the set name is *set_name*, the target server is *target_server*, the Capture server is *capture_server*, and the Capture schema is *capture_schema*.

Explanation: Valid values for *action name* are *Create Subscription Set*, *Drop Subscription Set*, *Alter Subscription Set*, and *Promote Subscription Set*.

User response: This message is for your information only, and no action is required.

ASN1510I The replication action *action_name* ended successfully at *timestamp*.

Explanation: Valid values for *action name* are *Create Capture server control tables*, *Drop Capture server control tables*, *Create Apply control server control tables*, and *Drop Apply control server control tables*.

User response: This message is for your information only, and no action is required.

ASN1511I The replication action *action_name* ended successfully for source owner *source_owner* and source table, view, or nickname *source_table*.

Explanation: Valid values for *action name* are *Create Registration*, *Drop Registration*, *Alter Registration*, *Add Registration Column*, and *Promote Registration*.

User response: This message is for your information only, and no action is required.

ASN1512I The replication action *action_name* ended successfully for Apply qualifier *apply_qual*, set name *set_name*.

Explanation: Valid values for *action name* are *Create Subscription Set*, *Drop Subscription Set*, *Alter Subscription Set*, *Add Statements to Subscription Set*, *Drop Statements from Subscription Set*, and *Promote Subscription Set*.

User response: This message is for your information only, and no action is required.

ASN1513I The replication action *action_name* ended successfully for Apply qualifier *apply_qual*, set name *set_name*, WHOS_ON_FIRST *whos_on_first*, source owner *source_owner*, source table *source_table*, source view qualifier *source_view_qual*, target owner *target_owner*, and target table *target_table*.

Explanation: The following values are valid for *action name*:

- *Add Subscription Member*
- *Add Subscription Member Column*
- *Drop Subscription Member*

User response: This message is for your information only, and no action is required.

ASN1514I The replication action ended at *timestamp* with *number* successes, *number* errors, and *number* warnings.

Explanation: At least one of the tasks in the ASNCLP script failed. The number of errors depends partly on the following options that determine whether the ASNCLP continues to process commands after errors:

SET RUN SCRIPT NOW

- STOP ON SQL ERROR ON (default)
- STOP ON SQL ERROR OFF

SET RUN SCRIPT LATER

- GENERATE SQL FOR EXISTING NO (default)
- GENERATE SQL FOR EXISTING YES

User response: The ASNCLP log file has messages for any tasks that fail. The messages are prefaced by an ASN error code. For more details about the task that prompted an error, look for an informational message for that task, for example *Create Q Subscription - ASN2003I*. To learn more about setting error options for script processing, see "How the ASNCLP handles errors while generating and running scripts" in the Information Management Software for z/OS Solutions Information Center or DB2 Information Center.

ASN1550E The replication action *action_name* ended in error. The value for the input parameter *input_parameter* is missing.

Explanation: The input parameter is mandatory for this action and is missing.

User response: Provide the mandatory parameter and rerun the Replication action.

ASN1551E The replication action *action_name* ended in error. The value *value* for the input parameter *input_parameter* is incorrect. The reason code is *reason_code*.

Explanation: The value provided for the input parameter is not a valid value. The following values are valid for the reason code:

0

Blocking minutes value should be between 0-999.

1

Commit Count value should be between 0-999.

2

Server Type value should be Capture Server.

3

Table type value should be one of the following types:

- USERTABLE
- CCD TABLE
- POINT IN TIME
- BASE AGGREGATE
- CHANGE AGGREGATE
- REPLICA
- USERCOPY

4

Remote Server Name value should be NULL.

5

Server Type value should be one of the following types:

- Capture Server
- Control Server
- Capture and Control Server
- Capture, Control and Target Server

6

Internal CCD tables must be noncomplete.

7

The Apply qualifier exceeds the maximum length of 18 characters.

8

The set name exceeds the maximum length of 18 characters.

9

Event names must be 128 bytes or fewer in length.

10

The source Capture schema name exceeds the maximum length of 128 bytes.

11

The target Capture schema name exceeds the maximum length of 128 bytes.

12

The BEFORE_OR_AFTER statement value must be 'A', 'B', or 'S'.

13

The EI_OR_CALL value must be 'C' or 'E'.

14

SQLSTATES must be 50 digits or fewer in length.

15

SQLSTATES must be numeric

16

The CONFLICT_LEVEL must be zero (0) or NONE.

17

The CHGONLY value must be 'N'.

18

The external CCD table is noncondensed and contains LOB columns.

19

The CONFLICT_LEVEL must be between 0 and 2.

20

The CHGONLY value must be 'Y' or 'N'.

21

The RECAPTURE value must be 'Y' or 'N'.

22

The DISABLE_REFRESH value must be 0 or 1.

23

The CHG_UPD_TO_DEL_INS value must be 'Y' or 'N'.

24

The STOP_ON_ERROR value must be 'Y' or 'N'.

25

The BEFORE_IMG_PREFIX value must be only one character.

26

The corresponding table space does not have the *New Tablespace* flag set to true in any of the previous scenarios.

27

The table name is not a valid control table. See the SQL Replication table structures documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for a valid list of control tables.

28

A federated server was found but a corresponding federated schema name was not provided. When calling the API for a federated system, ensure that the federated schema name is provided.

29

The specified remote source database name on the OS/400 system does not match the remote source database name of the registration.

User response: Provide a valid value for the input parameter, and rerun the replication action.

ASN1552E The replication action *action_name* ended in error. The value *value1* for input parameter *input_parameter1* is incompatible with the value *value2* for input parameter *input_parameter2*.

Explanation: The value provided for the replication parameter conflicts with another parameter specification.

User response: Provide valid values for the input parameters and rerun the replication action. See the replication system commands documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1553E The value *value1* for input parameter *input_parameter1* is incompatible with the value *value2* for the existing subscription set *subscription_set*, Apply qualifier *apply_qual*, and WHOS_ON_FIRST *whos_on_first*.

Explanation: The value provided for the replication

parameter conflicts with one of the values for the existing subscription set.

User response: Provide a valid value for the input parameter or change the subscription set definition, and rerun the replication action. See the replication system commands documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details.

ASN1560E The replication action ended in error. An SQL error was encountered. SQL message: *sql_message*.

Explanation: An error occurred during the execution of an SQL statement.

User response: Refer to your database message reference for SQL.

ASN1561E Connection to the server *server_alias* cannot be established. An SQL error was encountered. SQL message: *sql_message*.

Explanation: The connection to the specified server could not be established.

User response: Refer to your database message reference for SQL. Verify that the userid and password information is correct.

ASN1562E The replication action ended in error. An unexpected error occurred. Reference Code *reference_code*.

Explanation: The specified action cannot be performed because of a run time error.

User response: Contact IBM Software Support.

ASN1563E The replication action *action_name* ended in error. The replication architecture level *arch_level* does not support server *server_alias*.

Explanation: The specified replication architecture level is not supported on the specified server operating system, version, or release.

User response: See "Table structures for SQL Replication" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center to check the required value for the ARCH_LEVEL column in the register table.

ASN1564E The replication action *action_name* ended in error. The Capture server architecture level *arch_level* for Capture schema *capture_schema* does not support this replication action.

Explanation: The replication architecture level found in the IBMSNAP_REGISTER table for this Capture schema does not allow the specified replication action.

User response: Migrate the Capture control tables to Version 8 architecture level before retrying this action.

ASN1565E **The replication action *action_name* ended in error. The Apply control server architecture level *arch_level* does not support this replication action.**

Explanation: The replication architecture level found in the ASN.IBMSNAP_SUBS_SET table does not allow the specified replication action.

User response: Migrate the Apply control tables to Version 8 architecture level before retrying the action.

ASN1567W **The table space container information for table space *tablespace_name* cannot be read, because the DB2 stored procedure *procedure_name* in the library *library_name* cannot be found.**

Explanation: The DB2 stored procedure READTSCINFOS cannot be found on the Capture server or target server. The stored procedure is required to retrieve DB2 table space container information for that server.

User response: Determine whether the stored procedure exists on the server: check if the file db2rtsc exists in the function directory of the sqllib directory. The file db2rtsc might not exist if the server is a pre-V8 server. If the stored procedure does not exist, then edit the table space container definition that is provided in the output script.

ASN1568E **The name length *length* for the database object, *objectname* exceeds the allowed limit of *allowed_limit*.**

Explanation: The database object type provided in the second parameter allows a length that is smaller than the length of the actual object provided in the third parameter. As in the Properties file, the following values are valid for the object: *Table*, *Index*, *Tablespace*, *Table owner*, *Nickname*.

User response: Refer to the SQL Reference for the appropriate database, and provide the correct name length.

ASN1569E **The name of the database object to be created is identical to the existing name *objectowner.objectname* of type *object_type*.**

Explanation: The database object cannot be created because there is already a database object of the same type with the same name. As in the properties file, the following values are valid for the object: *Table*, *Index*, *Nickname*, *Tablespace*, *Table owner*.

User response: Provide a name for that object that does not already exist in DB2, and reissue the replication task.

ASN1570E **The database object *object*, *objectowner.objectname* does not exist.**

Explanation: The database object does not exist in the DB2 catalog. This object must exist in order to be defined as a source or target of a subscription set, as per the replication action. This object might have been defined as part of an existing registration or subscription-set definition but is not found in the DB2 catalog. As in the properties file, the following values are valid for the object: *Table*, *Index*, *Nickname*, *Tablespace*, *Table owner*, *View*.

User response: Provide a name that already exists in DB2, and reissue the replication task. If the object was defined as part of an existing registration or subscription-set definition, verify that the object exists in the DB2 catalog.

ASN1571E **The database table *tableowner.tablename* cannot be created: the database definition is not valid for data type *datatype* and column *column_name*. The reason code is *reason_code*.**

Explanation: The following values are valid for the *reason code*:

- 0 The datatype is not supported on this platform.
- 1 The length of the column is not supported on this platform.
- 2 The precision or scale of the column is not supported on this platform.

User response: Refer to the SQL reference for the appropriate database.

ASN1572E **The row size *row_size* for the database object *objectowner.objectname* of type *object_type* exceeds its database buffer pool row size *bufferpool_rowsize*. The database object cannot be created.**

Explanation: The row size of a table cannot exceed the table space page size for that table. The table space page size is derived from the buffer pool page size to which it belongs. No script is generated.

User response: You might have to create the table in a different table space. Refer to your DB2 platform documentation.

ASN1573E The number of columns *number_columns* for the database object *objectowner.objectname* of type *object_type* exceeds the database limit *db2_limit*. The database object cannot be created.

Explanation: The number of columns that a database object (table or index) can contain depends on the DB2 platform but cannot exceed a predefined number. No script is generated. The following values are valid for object type: *table*, *index*.

User response: Redesign the DB2 object.

ASN1574E The DB2 page size *page_size* for table space *tablespace_name* is not valid.
Reason code *reason_code*.

Explanation: The page size must be valid for the table space to be created successfully. The following values are valid for reason code:

- 0 Page size is not equal to the page size of the given buffer pool.
- 1 Page size is not equal to one of the following: 4K, 8K, 16K, 32K.

User response: Refer to the DB2 *SQL Reference* for appropriate page size ranges or values.

ASN1575W The DB2 table *tableowner-tablename* will be created in the DB2 default table space.

Explanation: No table space name was specified indicating where to create the specified table, so the table will be created in the DB2 default table space. This might be a problem if the default table space specifications are not appropriate for the specified table.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the table to be in its own table space, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the table.

ASN1576W The DB2 index *index_name* will be created in the DB2 default index space or table space.

Explanation: A table space (for workstation operating systems) or an index space (for z/OS operating systems) was not provided into which the specified index might be created. Therefore, the index is created using the DB2 defaults. This might be a problem if the default specifications are not appropriate for the specified index.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the index to be in its own table space or index space, then reissue the replication task with the appropriate specifications. No action is

required if the default is appropriate for the index.

ASN1577W The DB2 table space *tablespace* will be created in the DB2 default database.

Explanation: For z/OS operating systems only, a database was not provided into which the specified table space might be created. Therefore, the table space is created using the DB2 defaults. This might be a problem if the default specifications are not appropriate for the specified table space.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the table space to be in its own database, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the table space.

ASN1578I The DB2 table space *tablespace* will be created in the DB2 default storage group.

Explanation: For workstation and z/OS operating systems only, a storage group was not provided into which the specified table space might be created. Therefore, the table space is created using the DB2 defaults. This might be a problem if the default specifications are not appropriate for the specified table space.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the table space to be in its own storage group, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the table space.

ASN1579I The DB2 index *index_name* will be created in the DB2 default storage group.

Explanation: For workstation and z/OS operating systems only, a storage group was not specified into which the DB2 index might be created. Therefore, DB2 created the index using the default specification. This might be a problem if the default specifications are not appropriate for the specified index.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the index to be in its own storage group, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the index.

ASN1580I The DB2 table space *tablespace* will be created in the DB2 default buffer pool.

Explanation: For workstation and z/OS operating systems only, a buffer pool was not provided into which the specified table space might be created. Therefore, the table space is created using the DB2 defaults. This might be a problem if the default

specifications are not appropriate for the specified table space.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the table space to be in its own buffer pool, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the table space.

ASN1581I **The DB2 index *index_name* will be created in the DB2 default buffer pool.**

Explanation: For workstation and z/OS operating systems only, a buffer pool was not provided into which the specified index might be created. Therefore, the index is created using the DB2 defaults. This might be a problem if the default specifications are not appropriate for the specified index.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the index to be in its own buffer pool, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the index.

ASN1582W **The table space *tablespace* will be created in buffer pool *buffer_pool* but the buffer pool does not exist or is not active.**

Explanation:

- For applications on a DB2 for z/OS database, the buffer pool is not active into which the table space might be created.
- For applications on a DB2 for Linux, UNIX, and Windows database, the buffer pool does not exist into which the specified table space might be created.

User response:

- For a DB2 for z/OS database, make sure the buffer pool is active at the time of running the script.
- For a DB2 for Linux, UNIX, and Window database, make sure that the buffer pool exists at the time of running the script.

ASN1583E **The PageSize *page_size* for Table space *tablespace* does not match the default buffer pool PageSize.**

Explanation: The given PageSize does not match the PageSize of the default buffer pool. The table space cannot be created.

User response: Change the PageSize or choose another buffer pool.

ASN1584E **The replication action *action_name* ended in error. The Capture server Replication architecture level *arch_level* for Capture schema *capture_schema* is not a valid architecture level.**

Explanation: The Replication architecture level found in the captureschema.IBMSNAP_REGISTER does not allow the specified replication action.

User response: Drop the control tables on the Capture control server manually because the architecture level is not supported. Create the control tables with a valid architecture level.

ASN1585E **The replication action *action_name* ended in error. The Apply control server Replication architecture level *arch_level* is not a valid architecture level.**

Explanation: The Replication architecture level found in the ASN.IBMSNAP_SUBS_SET does not allow the specified replication action.

User response: Drop the control table on the Apply control server manually because the architecture level is not supported. Create the control tables with a valid architecture level.

ASN1586W **The DB2 table *tableowner.tablename* will be created in the DB2 default database.**

Explanation: For z/OS operating systems only, a database was not provided into which the specified table might be created. Therefore, the table is created using the DB2 defaults. This might be a problem if the default specifications are not appropriate for the specified table.

User response: Refer to the *SQL Reference* for the DB2 defaults. If you require the table space to be in its own database, then reissue the replication task with the appropriate specifications. No action is required if the default is appropriate for the table.

ASN1587E **The value *value* for the parameter *parameter_name* of the database object *object_name*, which has a type of *type*, is not valid.**

Explanation: The provided value is not valid or conflicts with another parameter value.

User response: Refer to the SQL reference for valid values.

ASN1588E **The value *encoding_scheme* that was provided for the parameter encoding scheme is not valid for the DB2 server *server_name*.**

Explanation: The provided value for the encoding scheme is not valid for the DB2 version of the server. No script is generated.

User response: Refer to the SQL reference for a valid value of the encoding scheme for the DB2 version.

ASN1589W The calculation of the size of the table space container *container* of the table space *tspace* resulted in an incorrect container size. Therefore the container size has been changed to size *size* megabytes.

Explanation: The calculation of the table space container size has resulted in a value that is too low to be used in a valid table space container definition. To ensure that the definition will be accepted by DB2, a replication-specific minimum container size has been provided for the table space container definition.

User response: For the calculation based on a percentage of the current source table size, check whether the source table contains data and if the statistics of the source table are up to date (using the RUNSTATS utility) . For the calculation based on a number of rows, check whether the number of rows is realistic.

ASN1590E The DB2 table space *table_sp_name* is partitioned and in the DB2 *object_type* group. It should not be partitioned and it should be in the *object_type* IBMCATGROUP.

Explanation: The provided table space is a partitioned table space. Creation of the replication control tables in a partitioned table space is not supported. No script is generated.

User response: Specify a table space that is not partitioned.

ASN1600E The remote server *remote_server_name* cannot be found.

Explanation: The specified remote server name cannot be found in the federated catalog table SYSIBM.SYSSERVERS, for the SERVERNAME value that was provided. The non-DB2 relational server cannot be accessed.

User response: Verify the input provided for the remote server name and try the action again.

ASN1601E The REMOTE AUTHID information for the remote server *remote_servername* cannot be found.

Explanation: The remote authentication information cannot be found in the federated catalog table SYSIBM.SYSUSEROPTIONS, for the SERVERNAME value provided. The non-DB2 relational server cannot be accessed.

User response: Verify the input provided for the remote server name and try the action again.

ASN1602E The server *server_alias* does not support access to federated servers.

Explanation: The federated replication functions are only supported on DB2 for Linux, UNIX, and Windows Version 8 and higher.

User response: Make sure that the specified database server is at the correct level or do not issue the replication task against a server that does not support it.

ASN1603E The Apply control server cannot reside on a non-DB2 relational server.

Explanation: Non-DB2 relational servers can be Capture control servers or target servers, but they cannot be Apply control servers.

User response: Specify a DB2 server as the Apply control server.

ASN1604E The remote table *remoteowner.tablename* exists in the non-DB2 relational server, but the provided nickname *nicknameowner.nickname* cannot be found in the federated server.

Explanation: The specified remote table exists in the remote database but the corresponding nickname is not found in the federated database.

User response:

1. Refer to "Configuring data sources" in the DB2 Information Center for details on how to create a nickname.
2. Create the nickname in the federated database.
3. Issue the replication task again.

ASN1605E The nickname *nicknameowner.nickname* exists in the federated server but the remote table *remoteowner.remotetable* cannot be found in the non-DB2 relational server.

Explanation: The nickname for the specified remote table exists but the corresponding remote table does not exist in the remote database.

User response:

1. Drop the nickname.
2. Depending on the table type, perform the following actions:
 - If the table is a user table, create the remote table in the remote server.
 - If the table is a replication control table on the Capture control server, perform the following actions:
 - a. Copy the data from the existing control tables on the Capture control server.

- b. Drop the control tables on the Capture control server.
- c. Create the control tables on the Capture control server.
3. Create the nickname in the federated server.
4. Issue the replication task again.

ASN1606W The nickname *nickname_owner.nickname_name* exists in the Federated server but the remote table *table_owner.table_name* cannot be found in the non-IBM server.

Explanation: The nickname for the specified remote table exists but the corresponding remote table does not exist in the remote database. Although this is an orphan nickname, this inconsistent state is still tolerated when dropping Replication definitions. A script is generated.

User response: The source nickname is not dropped when dropping the replication definitions. To ensure a consistent catalog, drop the nickname.

ASN1607W It is strongly recommended to alter the nickname *nickname_owner.nickname_name* that is defined for the replication subscription target to alter the local data type of column *column_name* from *existing_local_datatype* to *recommended_local_datatype* and ensure the proper source-to-target mapping for column data types.

Explanation: A mismatch was found between a source column data type and its corresponding nickname target column data type. The mismatch does not violate DB2 compatibility rules, but it might cause a problem to native non-IBM end-user applications. The problem does not occur during replication of the column data. The problem does occur if end-user applications retrieve the data. For example, if the nickname data type is created using the default mappings from the non-DB2 relational data type to the DB2 data type, the column will hold the broadest range of data type values, which might clash with the end-user application requirement of a more restrictive data type. A script is generated.

User response: Check the target to ensure that the nickname data type you need at the target is indeed the source column data type. If it is, then issue an ALTER NICKNAME statement to change the local data type of the nickname column. When you alter the nickname local data type to be the same as the source column data type, you enforce that the end-user application on the non-DB2 relational server sees the same data type as the source column data type.

ASN1608I The nickname *source_nickname* for the source and the nickname *ccd_nickname* for the consistent-change data (CCD) table have a column data type that is altered. The local data type column *local_datatype* is set to *changed_datatype* because the remote data type is *remote_datatype*. Reason code *reason_code*.

Explanation: When creating the nickname for a CCD table, the nickname is altered based on the data type of the CCD table that is created in the non-DB2 relational server to ensure the proper data type setting. A script is generated that updates the definitions that were provided in the replication administration tool.

User response: No action is required if the replication updates are acceptable.

ASN1609E The nickname *nicknameowner.nickname* exists in the Federated server but the remote table *remoteowner.remotetable* does not contain all the necessary columns.

Explanation: The target table nickname exists, and contains only a subset of the columns requested in the subscription.

User response: Use another Nickname as the target table or change the subscription to match the columns in the existing nickname.

ASN1620E Both Capture control tables and Apply control tables already exist. Capture control tables exist with architecture level *capture_arch_level* and Capture schema *capture_schema*. Apply control tables exist with architecture level *apply_arch_level*.

Explanation: The IBMSNAP_REGISTER table for this Capture schema and the ASN.IBMSNAP_SUBS_SET table already exist at the given server.

User response:

- If the architecture level of the existing IBMSNAP_REGISTER table for this Capture schema is 0201:
 - If the IBMSNAP_REGISTER table is already populated with valid replication definitions, migrate the Capture control tables to the latest version supported by replication.
 - If the IBMSNAP_REGISTER table is empty, drop the older version of the Capture control tables and reissue the replication task again.
- If the architecture level of the existing Capture control tables is Version 8 or later and it is the architecture level you intend, consider creating Capture control tables with a different Capture schema name.

- If the architecture level of the existing ASN.IBMSNAP_SUBS_SET table is 0201:
 - If the ASN.IBMSNAP_SUBS_SET table is already populated with valid replication definitions, migrate the Apply control tables to the latest version supported by replication
 - If the ASN.IBMSNAP_SUBS_SET control table is empty, drop the older version of the Apply control tables and reissue the replication task again.
- If the architecture level of the existing Apply control tables is Version 8 or later and it is the architecture level you intend, consider creating Apply control tables on a different server.

ASN1621W At least one row was found in the control table *table_owner-table_name*. Dropping this control table will drop all replication definitions stored in the table.

Explanation: The control tables of the replication schema selected for the drop request are not empty. Replication control information will be deleted if the generated scripts are executed.

User response: Run the generated scripts only if:

- You understand the impact to existing dependent subscription sets of dropping the control tables from the Capture control server.
- You understand the impact to existing dependent subscription sets (for multi-tier scenarios) of dropping control tables from the Apply control server.
- You do not want replication to run the Capture or the Apply processes for these definitions anymore.

If the architecture level is 0201, migrate the Capture or Apply control tables to the latest architecture level before dropping the control tables.

ASN1622E The replication action *action_name* ended in error. The required control table *controlowner.controltable* could not be found.

Explanation: Replication definitions are stored in Replication control tables. These tables must exist before a registration or subscription definition can be created. The existence of the IBMSNAP_REGISTER table is used to check if the control tables for the Capture control server already exist for a particular Capture schema. The existence of the IBMSNAP_SUBS_SET table is used to check if the control tables on the Apply control server already exist. The existence of IBMSNAP_SUBS_MEMBR is checked at the time of checking for the existence of a subscription member.

User response: If the control table IBMSNAP_SUBS_MEMBR table does not exist, then

your environment is in an inconsistent state. You must drop all the control tables from the Apply control server and then create them before attempting the action.

Alternatively, if the control tables IBMSNAP_REGISTER or IBMSNAP_SUBS_SET do not exist, create them before adding registration or subscription definitions on a control server. Otherwise, you can do the following:

1. If you are doing a registration-related action, check if the appropriate Capture schema was provided; or if the appropriate Capture control server was provided as input.
2. If you are doing a subscription-related action, check if the appropriate Apply control server was provided as input.
3. If you are creating a subscription set that contains target tables that need to be auto-registered at the target server (CCD or replica), then check if the appropriate control tables for the Capture control server exist at the subscription target server.

ASN1623W The Replication control table, *controlowner.controltable* could not be found and is not dropped.

Explanation: The Drop Capture control tables or Drop Apply control server control tables action was issued and the control table was missing. The script will not generate the appropriate DROP statement for that control table.

User response: This message is for your information only, and no action is required.

ASN1624I The server *server_alias* is not a known Replication Capture server for *capture_schema*.

Explanation: The captureschema.IBMSNAP_REGISTER table could not be found. A server is defined as a Replication Capture server when the appropriate Capture server control tables (including the IBMSNAP_REGISTER table) exist on the server.

User response: Create the appropriate Capture server control tables, if needed.

ASN1625I The server *server_alias* is not a known Replication Apply control server.

Explanation: The ASN.IBMSNAP_SUBS_SET table could not be found. A server is defined as a Replication Apply control server when the appropriate Apply control server control tables (including the IBMSNAP_SUBS_SET table) exist on the server.

User response: Create the appropriate control tables on the Apply control server, if needed.

ASN1626E Capture server control tables already exist for architecture level *arch_level* with the same Capture schema.

Explanation: The table captureschema.IBMSNAP_REGISTER already exists at the given server.

User response:

- If the architecture level of the existing captureschema.IBMSNAP_REGISTER table is 0801 or 0805, consider the following options:
 - Running the command is not necessary because the tables already exist with the same Capture schema.
 - Run the command under a different Capture schema.
- If the architecture level of the existing captureschema.IBMSNAP_REGISTER control table is 0201:
 - Migrate the Capture control server control tables to the Version 8 architecture, if the existing captureschema.IBMSNAP_REGISTER is already populated with valid Replication definitions.
 - If the control table is empty, simply drop the pre-V8 Capture server control tables and issue the Replication task again.

Otherwise, the architecture level is not valid. You need to drop the tables manually before attempting to create the tables.

ASN1627E Some Capture server control tables already exist with the same Capture Schema but for which an architecture level cannot be determined.

Explanation: The table captureschema.ASN.IBMSNAP_REGISTER does not exist although other Capture server control tables were found at the given server. Capture server control tables cannot be created until the tables are dropped. The Replication definitions at the Capture server are in an inconsistent state.

User response: Drop the remaining Capture server control tables to clean up the Capture control server definitions, and reissue the Create control table task. Loss of data occurs, so look at the content of the remaining control tables before issuing the drop task.

ASN1628E The Capture server control tables are not at the architecture level requested.

Explanation: The table captureschema.IBMSNAP_REGISTER does not exist with the provided architecture level. No script is generated.

User response: Issue the replication task again at the appropriate architecture level for the appropriate

Capture control server and Capture schema.

ASN1629E No Capture server control tables were found for the provided Capture schema.

Explanation: No control tables exist on the Capture control server. No control tables are dropped, and no script is generated.

User response: Issue the replication task again at the appropriate architecture level for the appropriate Capture control server and Capture schema.

ASN1630W Some Capture server control tables already exist with Capture schema *capture_schema* but their architecture level cannot be determined. The replication action *action_name* for the provided architecture level *arch_level* and Capture schema will drop control tables that might not belong to the architecture level provided.

Explanation: The IBMSNAP_REGISTER table does not exist on the Capture server. The replication architecture level is unknown, and if you provide an incorrect architecture level, you might lose critical data. No checks occur to determine whether a particular architecture level for the Capture control tables can be inferred. The control table is dropped if it exists. A script is generated.

User response: Issue the task again with the appropriate architecture level for replication.

ASN1631E Apply control server control tables already exist for architecture level *arch_level*.

Explanation: The table ASN.IBMSNAP_SUBS_SET already exists at the given server. No script is generated.

User response: If the architecture level of the existing ASN.IBMSNAP_SUBS_SET control table is 0201:

- If the existing ASN.IBMSNAP_SUBS_SET is already populated with valid Replication definitions, migrate the Apply control server control tables to the Version 8 architecture,
- If the table is empty, simply drop the pre-V8 Apply control server control tables and reissue the Replication task again.

Otherwise, the architecture level is not valid. You need to drop the tables manually before attempting to create the tables.

ASN1632E Some Apply control server control tables already exist but for which an architecture level cannot be determined.

Explanation: The table ASN.IBMSNAP_SUBS_SET does not exist although other Apply control server control tables were found at the given server. Apply control server control tables cannot be created until the tables are dropped. The Replication definitions at the Apply control server are in an inconsistent state. No script is generated.

User response: Drop the remaining control tables on the Apply control server to clean up the Apply control server replication definitions. Reissue the *Create control table* task. Loss of data occurs, so look at the content of the remaining control tables before issuing the *Drop* task.

ASN1633E The Apply control server control tables are not at the architecture level requested.

Explanation: The table ASN.IBMSNAP_SUBS_SET does not exist with the provided architecture level. No script is generated.

User response: Issue the replication task again at the appropriate architecture level for the appropriate Apply control server.

ASN1634E No Apply control server control tables were found.

Explanation: There are no control tables to drop from the Apply control server. No script is generated.

User response: Issue the replication task again at the appropriate architecture level for the appropriate Apply control server.

ASN1635W Some Apply control tables already exist but their architecture level cannot be determined. The replication action *action_name* for the provided architecture level *arch_level* will drop control tables that might not belong to the architecture level provided.

Explanation: The table ASN.IBMSNAP_SUBS_SET does not exist on the Apply control server. The replication architecture level is unknown, and if you provide an incorrect architecture level, you might lose critical data. No checks occur to determine whether a particular Apply control server control table architecture level can be inferred. If the control table exists, it is dropped. A script is generated.

User response: Reissue the task with the appropriate architecture level for replication.

ASN1636E The Replication Action of Manual Full Refresh ended with an error for the Apply qualifier *apply_qual* and set name *set_name*. The synchpoint in the *capschema.IBMSNAP_PRUNCNTL* table for the source member *sourceowner.sourcetable* and the target member *targetowner.target_table* is not translated by the Capture program.

Explanation: The synchpoint is either less than 0 or equal to hex zeros.

User response: Make sure you run the before load script to translate the hex zeros and capture is running on the server.

ASN1637E The replication action 'Manual Full Refresh' ended in error for the Apply qualifier *apply_qualifier* and the set name *set_name*. The target structure of at least one of the target subscription-set members in the given subscription set is greater than eight. None of the subscription-set members is eligible for a manual full refresh.

Explanation: The target structure of at least one of the target subscription-set members in the given subscription set is greater than eight. A manual full refresh does not support target structures that are greater than eight.

User response: Make sure that the target structure of the subscription-set member is less than or equal to eight, and then reissue the replication task.

ASN1638W The subscription-set member with a target of *targetowner.targetname* and a source of *sourceowner.sourcenname* is not complete. This subscription-set member is not included in the manual full refresh.

Explanation: The manual full refresh supports complete targets only. The given subscription-set member is not complete and cannot be included.

User response: No action is required.

ASN1639E The replication action 'Manual Full Refresh' ended in error for the Apply qualifier *apply_qualifier* and the set name *set_name*. None of the target subscription-set members in the given subscription set is complete or eligible for a manual full refresh.

Explanation: The manual full refresh supports complete targets only, and none of the targets is complete.

User response: Make sure that at least one of the

subscription-set members in the subscription set is complete, and reissue the replication task.

ASN1640E The replication action ended in error for the Apply qualifier *apply_qualifier* and the set name *set_name*. There are no subscription-set members in the subscription set.

Explanation: The subscription set does not contain any subscription-set members.

User response: Add at least one subscription-set member to the subscription set, and reissue the replication task.

ASN1641E The replication action *action_name* ended in error. This action on an OS/400 system is supported only through OS/400 commands.

Explanation: Neither the replication center nor the command line supports the replication action on an OS/400 system. The possible actions might be: creating capture server control tables, creating apply server control tables, dropping capture server control tables, or dropping apply server control tables.

User response: Issue OS/400 commands to perform the replication action.

ASN1650I The replication action *action_name* started at *timestamp*. The monitor server is *server_name* and the Group_or_Contact name is *group_name_or_contact_name*.

Explanation: The replication action started at the specified monitor server.

User response: This message is for your information only, and no action is required.

ASN1653I The replication action *action_name* for *group_contact_or_condition_name* ended successfully at *timestamp*. The monitor server is *server_name*.

Explanation: The replication action ended successfully at the specified monitor server.

User response: This message is for your information only, and no action is required.

ASN1654E The replication action *action_name* ended in error. The length of the input parameter *parameter-name*, *parameter_length* exceeds the limit *maximum-limit*.

Explanation: The length of the specified input parameter is longer than the maximum allowable length. No script is generated.

User response: Verify the input parameter value, and re-enter the parameter value.

ASN1655E The replication action *action_name* ended in error. The value *input_value* of the input parameter *input_parameter* is incorrect.

Explanation: The value of the specified input parameter is not correct.

User response: Refer to your documentation for valid parameter values.

ASN1656E The replication action *action_name* ended in error. The value of the input parameter *input_parameter* is missing.

Explanation: A value for this specified input parameter is mandatory for this action. However, the value is missing. No script is generated.

User response: Enter a value for this mandatory input parameter, and rerun the replication action.

ASN1657E The replication action *action_name* ended in error. At least one optional parameter value must be specified.

Explanation: You must specify at least one optional parameter value when issuing a command in which each parameter value is optional. No script is generated.

User response: Issue the command again with the correct parameters.

ASN1658E The replication action *action_name* ended in error. The value *value1* of the input parameter *input_parameter1* must be different than the value *value2* of the input parameter *input_parameter2*.

Explanation: The value of one input parameter is the same as the value of another input parameter and will result in the creation of inconsistent definitions. No script is generated.

User response: Issue the command again with valid parameter values.

ASN1659E The replication action *action_name* ended in error. The contact *contact-name* already exists.

Explanation: The specified contact name already exists in one of the rows in the ASN.IBMSNAP_CONTACTS table. Contact names must be unique. No script is generated.

User response: Issue the command again with a different contact name.

ASN1660E The replication action *action_name* ended in error. The contact *contact-name* does not exist.

Explanation: The specified contact name does not exist in any of the rows in the ASN.IBMSNAP_CONTACTS table. The contact name must exist in the ASN.IBMSNAP_CONTACTS table before you can alter, substitute, delegate, or drop the name. No script is generated.

User response: Issue the command again with a different contact name.

ASN1661E The replication action *action_name* ended in error. The contact *contact-name* cannot be dropped, because dropping the contact empties each associated group.

Explanation: A group should have at least one associated contact. The specified contact is the last contact in each associated group, and the last contact cannot be dropped. No script is generated.

User response: Drop each associated group before attempting to drop the contact.

ASN1662E The replication action *action_name* ended in error. The contact *contact-name* cannot be dropped, because the contact is associated with one or more conditions.

Explanation: The contact name that you are attempting to drop is the only contact associated with conditions for either the Capture or Apply components. No script is generated.

User response: Use the SUBSTITUTE option in the DROP CONTACT command, or use the SUBSTITUTE command to change the contact name of the conditions. If you do not need the conditions, drop the conditions and then drop the contact.

ASN1663E The replication action *action_name* ended in error. The value *startdate_value* that is specified for the start date is greater than the value *enddate_value*, which is specified for the end date.

Explanation: You cannot enter a start date that is beyond the end date. No script is generated.

User response: Issue the command again with a valid combination of dates.

ASN1664E The replication action *action_name* ended in error. The group *group-name* already exists.

Explanation: The specified group name already exists in one of the rows in the ASN.IBMSNAP_GROUPS table. Group names must be unique.

User response: Change the group name, and issue the command again.

ASN1665E The replication action *action_name* ended in error. The group *group_name* does not exist.

Explanation: The specified group name does not exist in any of the rows in the ASN.IBMSNAP_GROUPS table. The group name must exist in the ASN.IBMSNAP_GROUPS table before you can alter or drop the group name. No script is generated.

User response: Verify the group name, and reissue the command.

ASN1666E The replication action *action_name* ended in error. The group *group_name* cannot be dropped because it is associated with one or more conditions.

Explanation: The group that you are attempting to drop is the only group associated with conditions for either the Capture or Apply components. No script is generated.

User response: In order to drop the group, alter the contacts of the associated conditions and then reissue the command.

ASN1667E The replication action *action_name* ended in error. The contact *contact-name* is not associated with the specified group *group_name*.

Explanation: The contact name that you are attempting to drop is not associated with the specified group.

User response: Verify the specified contact name and reissue the command.

ASN1668E The replication action *action_name* ended in error. The contact *contact-name* is already associated with the specified group *group_name*.

Explanation: The contact name that you specified is already associated with the specified group.

User response: No action is required.

ASN1671E The replication action *action_name* ended in error. The alert condition *condition-name* already exists for the monitor qualifier *mon-qual*, the server *server-name*, the schema or qualifier *schema-or-qualifier*, and the subscription-set name *set-name*.

Explanation: The alert condition that you are attempting to create already exists with the same

specified parameters on the monitor control server.

User response: Verify this alert condition and issue the command again.

ASN1672E **The replication action *action_name* ended in error. The alert condition *condition-name* does not exist for the monitor qualifier *mon-qual*, the server *server-name*, the schema or qualifier *schema-or-qualifier*, and the subscription-set name *set-name*.**

Explanation: The alert condition that you are attempting to drop or to alter does not exist on the monitor control server.

User response: Verify the alert name and issue the command again.

ASN1673W **The condition *condition_name* is valid only at the apply qualifier level.**

Explanation: The condition name is not valid with a subscription-set name value. The name of the subscription set will be ignored.

User response: Do not specify the subscription-set name value.

ASN1674W **The condition *condition_name* is valid only with update-anywhere subscription sets.**

Explanation: The condition name is valid only with update-anywhere subscription sets.

User response: Do not set this condition. This condition will be ignored.

ASN1675I **This is a test message from the Replication Center.**

Explanation: This message is used to send a test e-mail verifying the e-mail address entered in the contact.

User response: This message is for your information only, and no action is required.

ASN1677E **The replication action *action_name* ended in error. The apply qualifier *apply-qual* and the subscription-set name *set-name* do not exist on the server *server-name*.**

Explanation: The apply qualifier and the subscription-set name do not exist in the IBMSNAP_SUBS_SET table on the specified apply control server.

User response: Supply a valid apply qualifier and a valid subscription-set name.

ASN1678E **The replication action *action_name* ended in error. The capture schema *cap-schema* does not exist on the server *server-name*.**

Explanation: The capture schema does not exist in the ASN.IBMSNAP_CAPSCHEMAS table on the specified capture control server.

User response: Supply a valid capture schema.

ASN1679E **The replication action *action_name* ended in error. The contact *contact_name* that you attempted to substitute is not associated with a condition.**

Explanation: The contact name does not exist in the ASN.IBMSNAP_CONDITIONS table. A contact can be substituted only if it exists in the ASN.IBMSNAP_CONDITIONS table. No script is generated.

User response: Supply a valid contact name.

ASN1680I **The replication action *action_name* started at *time*. The Monitor server is *server_name*.**

Explanation: This message is for your information only.

User response: No action required.

ASN1681E **The request to create the monitor server failed. Monitor control tables of the architecture level *arch_level* already exist.**

Explanation: The monitor control tables already exist on the given server. No SQL script is generated.

User response: If the existing monitor control tables are of an older level:

- If the existing monitor control tables are already populated with valid replication definitions, migrate the monitor control tables to the latest version supported by replication.
- If the ASN.IBMSNAP_ALERTS table and the ASN.IBMSNAP_CONDITIONS table are empty, drop the older version monitor control tables and reissue the replication task again.

Otherwise, the architecture level is not valid. You need to drop the tables manually before attempting to create the tables.

If the architecture level of the existing monitor control tables stored in the ARCH_LEVEL column of the ASN.IBMSNAP_MONPARMS table is at the latest level supported by replication, there is no need to run the command since the control tables already exist.

ASN1682E The replication action ended in error.
No Monitor control tables were found.

Explanation: There are no Monitor control tables to drop. No script will be generated.

User response: Issue the replication task again for the appropriate server containing the Monitor control tables.

ASN1683E The replication action *action_name* ended in error. A row with the specified capture schema already exists in the ASN.IBMSNAP_CAPSCHEMAS table even though the Capture control tables do not exist under that Capture schema.

Explanation: The IBMSNAP_CAPSCHEMAS table already contains the Capture schema value that was specified as input for the "Create Capture Server Control tables" action. The Capture schema value must be unique.

User response: Make sure that you have specified the correct value for the Capture schema field and delete the row containing the specified Capture schema from the IBMSNAP_CAPSCHEMAS table. Reissue the task.

ASN1684W The nickname *nicknameowner_nickname* for the replication control table *controlowner.control_table* was not found on the federated database.

Explanation: You tried to drop a nickname of a replication control table that was already dropped or deleted; therefore, the generated script does not contain a DROP statement for that nickname.

User response: This message is for your information only, and no action is required.

ASN1685W The object *objectowner.object_name* of type *object_type* could not be found on the non-DB2 relational server.

Explanation: You tried to drop a replication object that was already dropped or deleted; therefore, the generated script does not contain a DROP statement for that replication object.

User response: This message is for your information only, and no action is required.

ASN1686E The name length *length* for the non-DB2 relational object *object_name* exceeds the allowed limit of *allowed_limit*.

Explanation: A non-DB2 object name was specified whose length is longer than the maximum length allowed for the actual object.

User response: Refer to the SQL reference for your database.

ASN1687E The replication action *action_name* ended in error. The table space *tablespace_name* does not belong to the IBMCATGROUP node group.

Explanation: The specified table space does not belong to the default IBMCATGROUP node group. The replication action does not support this node group.

User response: Verify the table space name and reissue the task.

ASN1688E The replication action *action_name* ended in error. The specified table space options *tableowner.tablename* are not valid.

Explanation: You can customize the table space when you create the control tables. You can use existing table space, a new table space, or a table space that has already been specified for another control table earlier in the same session. The values specified in these table space options are ambiguous. There are either no values, or more than one value is specified for these table space options.

User response: Verify the values of the table space options and reissue the task.

ASN1689E The replication action *action_name* ended in error. There are no alert conditions for the monitor qualifier *monitor_qualifier_name*, the server *capture_or_apply_server*, and the schema or qualifier *schema_or_qualifier*.

Explanation: No alert conditions are defined for the specified monitor qualifier, Capture or Apply control server, and schema or Apply qualifier.

User response: Define at least one alert condition for the monitor qualifier, Capture or Apply control server, and schema or Apply qualifier.

ASN1700E The column *tableowner.tablename.columnname* of data type *data_type* cannot be included in the registration. Reason code *reason_code*.

Explanation: The column cannot be supported by the replication capture program as defined. No script to register the specified is generated. The following values are valid for the reason code:

- | | |
|---|---|
| 0 | The data type is not supported. |
| 1 | The column is already registered. |
| 2 | z/OS fieldproc column. |
| 3 | This column does not qualify as a before-image column. |
| 4 | The data type is not supported through DB2 for federated. |

- 5 The column does not exist in the source object.
- 6 The maximum number of registered LOB columns was exceeded for that table.
- 7 The column name starts with the before-image prefix.
- 8 This column does not qualify as a before-image column or as an after-image column.
- 9 A mixed-case column name is not supported when the source table is on a non-DB2 server.
- 10 This column name is a duplicate of one already provided for this source.

User response: Check the reason code to determine why the column cannot be registered. Refer to the replication documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for additional explanations or restrictions.

ASN1701E The provided locksize value *lock_size* for the given table space *tablespace_name* is not valid.

Explanation: Locksize should be equal to the P(PAGE), R(ROW) or A(ANY), in the case of z/OS operating system.

User response: Provide the correct locksize and submit your action again.

ASN1702W Replication definitions for the registered column *objectowner.objectname.columnname* has been changed to support null values.

Explanation: before-image columns are required to support null values. If no before-image column value is present, an INSERT statement will fail. A script is generated to update user-provided definitions.

User response: This message is for your information only; no action is required.

ASN1703E The table *tableowner.tablename* cannot be registered for change-capture replication. Reason code *reason_code*.

Explanation: The table cannot be supported by the capture program as defined. No script is generated. The following values are valid for the reason code:

- 0 The table with a z/OS validproc.
- 1 Existing internal CCD table.
- 2 Existing CD table.
- 3 DB2 catalog table (Windows, UNIX, iSeries)
- 4 The table is already registered.

- 5 The source for an internal CCD table is not a registered source.
- 6 The source is a CD table and cannot be registered.
- 7 This source name is a duplicate for this session.
- 8 The source is a replication control table.
- 9 Not one of the source columns qualifies for registration.
- 10 The maximum number of registered LOB columns has been exceeded for this table.
- 11 Structured data types are not supported.
- 12 The before-image prefix can be only one character.
- 13 An internal error occurred.
- 14 A blank character is not a valid before-image prefix.
- 15 **iSeries:** The source table or view should not have blank spaces in it.
- 16 **iSeries:** A blank space character is not valid in the CD Table Owner/Name field.
- 17 **iSeries:** For a given source, you cannot register some before-image columns and some after-image only columns. Either all or none of the columns must have the before-image.
- 18 The CD name for this source is a duplicate; that CD name already exists for this session.
- 19 The source object type is not a valid object type for replication.

User response: Check the reason code to determine why the table cannot be registered for change-capture replication. Refer to the replication documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for additional explanations and restrictions.

ASN1704E The view *viewowner.viewname* cannot be registered. Reason code *reason_code*.

Explanation: The view cannot be supported by the Replication Capture mechanism, as defined. No script is generated. The following values are valid for the reason code:

- 0 None of the dependent tables for the view are registered.
- 1 The source-table columns on which the view is dependent are not registered.
- 2

| | | | |
|----|---|----------|--|
| 3 | The view is on an internal ccd. | 21 | The view contains the join of a CCD and a non-CCD table. |
| 4 | The view is already registered. | | The view defined at the CCD table must be complete and condensed. |
| 5 | The view has an 'OUTER JOIN' syntax. | 22 | The dependent table is a nickname. |
| 6 | The view includes more than one table or view column with a function, and no correlation is provided in the view definition for each table. | 23 | A federated registration expects a nickname to be registered as a source. |
| 7 | The view contains a reference to an aggregate function. | | User response: Check the reason code to determine why the view cannot be registered. If appropriate, fix the problem and register the view again. |
| 8 | The view contains a subselect/subquery. | <hr/> | |
| 9 | The view contains a reference to another view. | ASN1705E | The change data object, <i>objectowner.objectname</i> already exists in the server. |
| 10 | The view has an UNION. | | Explanation: The change data table or view cannot be used for the current source to be registered, because it already exists at the Capture server. No script is generated. |
| 11 | No correlation is provided for the column. | | User response: Provide a different name for the change data object. |
| 12 | The base table does not have the schema name. | <hr/> | |
| 13 | The base table does not exist. | ASN1706W | A column <i>column_name</i> is added to a registered source <i>sourceowner.sourcename</i>. The registered source maintains an Internal CCD table. The new column has to be first added to the CCD table subscription member before adding to any existing or not-yet existing subscription member. |
| 14 | The view contains Table Expression as Table. | | Explanation: If the new column is needed in dependent subscription sets, you first have to add the column to the internal CCD subscription member before adding the column to any subscription member needed. |
| 15 | The dependent table does not exist. | | User response: Provide a different name for the change data object. |
| 16 | A view on view cannot be registered. | <hr/> | |
| 17 | The given source object is not a view. | ASN1707W | The replication action Alter Registration for <i>sourceowner.sourcename</i> is not in effect until a Capture REINIT command is issued at the Capture server. |
| 18 | This source view is a duplicate for this session. | | Explanation: The registered source is successfully updated. However, the Capture program does not recognize the corresponding captureschema.IBMSNAP_REGISTER table updates until a REINIT command forces it to do so. A script is generated. A Capture command is required afterwards for the effect of the script to be in action. |
| 19 | The view definition cannot be supported. | | |
| 20 | The view has an asterisk (*) instead of a specific column name in the view definition. | | |

User response: To make the changes effective immediately:

1. Run the generated script.
2. Issue a REINIT of the appropriate Capture program, for the appropriate Capture schema.

ASN1708E **The table, view or nickname *objectowner.objectname* is not a Replication registered source.**

Explanation: The specified replication object is not defined in the replication control tables. No script is generated.

User response: Ensure that the object is correctly specified in the command and it exists.

ASN1709W **Associated subscription sets will not be valid after the registered source *sourceowner.sourcename* is dropped.**

Explanation: Subscription members rely on the underlying source registrations that define the source member. If you drop a registered source table, the dependent source members of a subscription set are no longer valid. The subscription sets that are associated with the specified registration source can be found in the captureserver.IBMSNAP_PRUNCNTL table at the Capture control server, where the SOURCE_OWNER and SOURCE_TABLE correspond to the registered source that is dropped. The appropriate Apply control server and subscription set names are columns in the IBMSNAP_PRUNCNTL table. The associated subscription sets fail if Apply is running. A script is generated.

User response: Deactivate or drop the dependent subscription sets before running the script, if the registered source has dependent subscription sets.

ASN1710W **Dependent view registered sources will not be valid after the registered source *sourceowner.sourcetable* is dropped.**

Explanation: View registrations rely on the underlying registration of the tables that make up the view definition. If you drop a registered source table, you invalidate any view registration that is based on the table. The views that might be affected can be found in the captureserver.IBMSNAP_REGISTER table at the Capture server, where the PHYS_CHANGE_OWNER and PHYS_CHANGE_TABLE are the same as the CD_OWNER and CD_TABLE of the registered source that is dropped. The associated subscription sets, which depend on the view registrations, fail, if Apply is running. A script is generated.

User response: Deactivate or drop the appropriate subscription sets, or view registrations, before running the script, if the registered source has dependent view registrations.

ASN1711W **The source *sourceowner.sourcename* is still active so dropping it will result in a Capture failure.**

Explanation: An active registration has a SYNCHPOINT value that is not null in its captureschema.IBMSNAP_REGISTER table. When the Capture program started, it expected all active registrations to always exist and be valid. So the Capture program needs to be signaled that a registered source was dropped because the drop action invalidates the registration information. Failure to provide that information to the Capture program causes the Capture program to fail. A script is generated, but is NOT ready to run.

User response:

1. Deactivate the appropriate registration (via the Replication Center GUI, or by issuing the STOP signal and a command type of CMD).
2. Wait for a SIGNAL_STATE of Complete in the captureschema.IBMSNAP_SIGNAL table.
3. Run the script that drops the registration.

ASN1712E **The table, view, or nickname *objectowner.objectname* is not a valid Replication registered source. Reason code *reason_code*.**

Explanation: Inconsistent information was found for this registered source in the Capture server control tables. No script is generated.

User response: Drop the registered source and create the registration again.

ASN1713E **The registered source *sourceowner.sourcename* cannot be deactivated. Reason code *reason_code*.**

Explanation: The following values are valid for reason code:

- | | |
|---|---|
| 0 | The source is registered as a FULL REFRESH and therefore cannot be deactivated. |
| 1 | The source is a CCD and CCD registrations cannot be deactivated. |
| 2 | The source is a view and view registrations cannot be deactivated. |

User response: This message is for your information only, and no action is required.

ASN1714E **The registered source *sourceowner.sourcename* cannot be altered. Reason code *reason_code*.**

Explanation: The following values are valid for the reason code:

- | | |
|---|---|
| 0 | The CD table for this source has RRN column |
|---|---|

(iSeries only). The RRN column must be the last column in the table, so the source cannot be altered.

- 1 The source is a view, and view registrations cannot be altered.
- 2 The source is registered for full refresh and cannot be altered.
- 3 The source table column does not match the column being altered.
- 4 The column is a LOB, DATALINK, or ROWID data type and does not qualify for a before-image value.
- 5 The before-image column value cannot be null or a blank character.
- 6 An after-image value has not been registered for the given column.
- 7 The before-image prefix cannot be updated if it is used with an existing registered source.
- 8 The use of the current before-image prefix makes one of the columns ambiguous in this registered source.
- 9 The before-image prefix can be only one character.
- 10 An internal error occurred.
- 11 The specified registration source name is a duplicate of a source included with that was changed, and the script for that registration was not yet executed.
- 12 The conflict level cannot be updated for the replica registration.

User response: Check the reason code to determine why the source cannot be altered, and refer to the replication documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details on how to correct these errors.

ASN1715E The replication action ended in error. The native OS/400 message is *as400native_message*.

Explanation: An error was encountered while issuing the appropriate command on the OS/400 operating system or iSeries servers. No script is generated.

User response: Refer to the OS/400 Console Log for more detailed error information.

ASN1716W The replication action ended with a warning. The native OS/400 message is *as400native_message*.

Explanation: A warning was encountered while issuing the appropriate command on the OS/400

operating system or iSeries server. A script is generated.

User response: Refer to the iSeries Console Log for more detailed warning information.

ASN1717I The replication action ended with an Informational Clause. The native OS/400 message is *as400native_message*.

Explanation: An informational message was encountered while issuing the appropriate command on the OS/400 operating system or iSeries server. A script is generated.

User response: This message is for your information only, and no action is required.

ASN1718E The nickname *nicknameowner.nickname* cannot be registered. Reason code *reasoncode*.

Explanation: The nickname is not supported by the Capture program as defined. No script is generated. The following values are valid for the reason code:

- 0 The internal CCD table (your CD table) already exists.
- 1 The nickname is on the native catalog table.
- 2 The nickname is already registered.
- 3 A federated registration expects a nickname as a source to be registered.
- 4 No columns are eligible for the Capture program.
- 5 The provided nickname is a duplicate from an earlier registration, but the corresponding script has not been executed.
- 6 A federated registration supports only user tables.
- 7 A federated registration supports only noncondensed and noncomplete CCD tables.
- 8 The CCD nickname provided is a duplicate of a CCD nickname from an earlier registration, but the script for that registration has not yet been executed.
- 9 Remote server information is not provided for the nickname registration.

User response: Check the reason code to determine why the nickname cannot be registered.

ASN1719W The non-IBM triggers that were defined for registered source *nicknameowner.nickname* will be dropped. Any additional logic later provided by users in these triggers will be lost.

Explanation: Dropping a registered source implies dropping all the objects that were created during the

source registration, regardless of later updates. A script is generated.

User response: Copy the trigger logic before dropping the registered source, if needed.

ASN1720E Change Data table information for the source nickname *nicknameowner.nickname* is not found in the *captureschema.IBMSNAP_REGISTER* table.

Explanation: A row is found in the *captureschema.IBMSNAP_REGISTER* table for the given source nickname but the CCD table information for that source is missing. The Change data table information is required to drop the replication definitions. A script is not generated.

User response: Please make sure the correct source name is given and call the action again.

ASN1722W The view *view_owner.viewname* will be registered as full refresh, because all the base tables of this view are registered as full refresh.

Explanation: The view must be registered as full refresh, because the base tables of this view are registered as full refresh only or are not registered replication sources.

User response: No action is required.

ASN1723W The view *viewowner.viewname* will be registered for change-capture replication, because one or more base tables from this view are registered for change-capture replication.

Explanation: The view must be registered for change-capture replication, because the base tables of this view are registered for change-capture replication.

User response: No action is required.

ASN1724E The name of the object that you are creating on the non-DB2 relational server is identical to the *objectowner.objectname* of type *objecttype*.

Explanation: The object that you specified cannot be created, because there is an existing object with the same type and same name on the non-DB2 relational server.

User response: Provide a unique name for the object, and reissue the replication task.

ASN1725W The trigger named *triggerowner.trigger_name* already exists on the remote table *remoteowner.remotetablename*. You must not run the generated script until you have determined how to merge the contents of the existing trigger with the generated trigger definition.

Explanation: An trigger with this name already exists on the remote table in the non-DB2 relational database. The database manager might not indicate a conflict and might subsequently overwrite your existing trigger if you run the CREATE TRIGGER statement in the generated script. The database manager also might return a SQL error indicating that the object already exists. Generated trigger names cannot be customized, because customized triggers cannot be dropped when the registration is dropped.

User response: First, determine how to merge the pre-existing triggers with the generated triggers. Then, either create your own script to merge your existing logic with the trigger logic that is generated by the replication tool, or update the script that is generated by the replication tool to include your existing trigger definitions.

ASN1726W The trigger named *triggerowner.trigname* does not exist in the remote table *owner.tablename* on the remote server *rmtservername*.

Explanation: The trigger does not exist on the remote database. The trigger might have been dropped.

User response: No action is required.

ASN1727I The registered source *registered_source* is deactivated.

Explanation: The specified registered source has already been deactivated.

User response: This message is for your information only, and no action is required.

ASN1728W The CCSID *Unicode_ASCII_EBCDIC* of the change data (CD) table *cdowner.cdname* for the source table *sourceowner.sourcetable* does not match the CCSID *Unicode_ASCII_EBCDIC* of the IBMSNAP_UOW table for the capture schema *capture_schema*.

Explanation: For the given capture schema, the Apply program will join the IBMSNAP_UOW table and the CD table of the given source if the column JOIN_UOW_CD in the ASN.IBMNSNAP_MEMBR table is set to Y. This column contains Y if the target type of the associated subscription-set member is not user copy, or if any columns of the IBMSNAP_UOW table are

used in the WHERE clause of the subscription-set member. If the Apply program joins tables with different encoding schemes, an error will occur. For more information about encoding schemes, see "UNICODE and ASCII encoding schemes" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center.

User response: For subscription members that will use this registration, define the target table with a type of user copy and do not use any IBMSNAP_UOW columns in the WHERE clause.

ASN1729E The registration for the nickname *nicknameowner.nickname* cannot be dropped. The reason code is *reasoncode*.

Explanation: The registration for this nickname cannot be dropped. No script is generated. The following value is valid for the reason code:

- 0 The specified nickname is a duplicate of a nickname included with a prior registration drop. However, the script for that registration drop has not yet been executed.

User response: Review the reason code explanation, and refer to the replication documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for additional explanations and restrictions.

ASN1730W The procedure *procedureowner.procedurename* already exists in the remote server *remote_server*. The content of the existing procedure must be merged with the generated procedure definition before the generated script is run.

Explanation: A procedure with this name already exists in the non-DB2 relational database. The database manager might not indicate a conflict and could subsequently overwrite your existing procedure if you run the CREATE PROCEDURE statement in the generated script. Or, an SQL error might be returned indicating that the object already exists. Generated procedure names cannot be customized because customized procedures cannot be dropped when the registration is dropped.

User response: Determine how to merge the existing procedures with the generated procedures. Then, either create your own script to merge your existing logic with the procedure logic that is generated by the replication tool, or update the script that is generated by the replication tool to include your existing procedure definitions.

ASN1731W The column *column_name* of data type *datatype* in the non-DB2 relational database table will be converted to a data type *datatype* in the nickname by the federated wrapper. The reason code is *reason_code*.

Explanation: The following value is valid for the reason code:

- 0 The non-DB2 relational database is an Oracle database and the Number column in that table will be converted to a Double data type in the nickname by the federated server.

User response: This message is for your information only, and no action is required.

ASN1732E The replication action *action_name* ended in error. The source table *srcowner.srctable* must be journaled with both before-images and after-images.

Explanation: The native OS/400 command requires the source table to be journaled with both the before-image and after-image columns before registration.

User response: Journal the source table with both the before-images and after-images before executing the generated script for the table registration.

ASN1733E The registration of the source *srcowner.srctable* cannot be dropped. The reason code is *reason_code*.

Explanation: The following value is valid for the reason code:

- 0 The provided source is a duplicate from a prior DROP registration, but the corresponding script was not executed.

User response: Check the name of the source provided and issue the task again.

ASN1734W Definitions for registered sources that have dependent views will not be updated with the columns that were for the registered source *srcowner.srctable*.

Explanation: You are attempting to alter or add one or more columns to an existing registered source. The source registration has dependent view registrations, but the altered columns will not be reflected in the view registration definitions.

User response: To update the view registration definition so that new column information is reflected in the view registration:

1. Drop the current view registration.
2. Re-create the view registration.

No action is required if you do not want the new column information to be reflected in the view registration.

ASN1735E The table, view, or nickname *owner.name* cannot be registered. There are subscription set members associated with the source table from a previous registration that was dropped. The definitions requested for the current registration conflict with the definitions stored in the replication control tables for the existing subscription set members. The reason code is *reason_code*.

Explanation: The table, view, or nickname was previously registered and dropped. However, the associated subscription set members were not dropped and the information for these members still exists in the replication control tables. To avoid any problems with the Capture and the Apply programs, the table or nickname is being re-registered and the information being requested needs to match the information that is stored in the control tables for the orphan members. The following values are valid for the reason code:

- 0 You are attempting to register the table or nickname for FULL REFRESH ONLY. The rows in the IBMSNAP_PRUNCNTL table indicate that this table or nickname was previously registered with the change-capture mechanism and the registration was dropped.
- 1 You are attempting to register the table or nickname with the change-capture mechanism. The rows in the IBMSNAP_PRUNCNTL table indicate that this table or nickname was previously registered as FULL REFRESH ONLY and the registration was dropped.
- 2 You are attempting to register the table or nickname with the change-capture mechanism. The source table was previously defined and was dropped. The values of the PHYS_CHANGE_OWNER and PHYS_CHANGE_TABLE columns in the IBMSNAP_PRUNCNTL table do not match with the input values specified for the CD/CCD owner or the CD/CCD table fields.

User response: Take the specific actions for the following reason codes:

- 0 Register the table or nickname with the change-capture mechanism.
- 1 Register the table or nickname as FULL REFRESH ONLY.
- 2 Register the table or nickname and make sure that the change data or the consistent change data owner and the table names match those that are in the IBMSNAP_PRUNCNTL field.

Do not customize these values so that the replication action will take the values stored in the control table.

ASN1736W The name of the before-image column for the column *column_name* of data type *datatype* will be truncated. The length of the column name plus the before-image prefix exceeds the allowed column name limit *allowed_limit* for the non-DB2 relational database.

Explanation: When the specified before-image prefix is appended to the column name, the name is longer than the maximum allowed column name limit on the non-DB2 relational server. The column name is truncated from the trailing end such that the length of the column name is equal to the maximum length.

User response: This message is for your information only, and no action is required.

ASN1737W The replication action *action_name* is unable to derive enough information to re-create the procedure or the trigger definition on the IBMSNAP_PRUNCNTL table.

Explanation: The procedure or trigger definition on the IBMSNAP_PRUNCNTL table prunes the CCD tables for all the registered nicknames. Each time that a nickname is registered or dropped from the registration, the procedure or trigger on the IBMSNAP_PRUNCNTL table is dropped and rebuilt using the new information for that particular nickname. The information specified for this registration does not contain the necessary information to define the procedure or the trigger on the IBMSNAP_PRUNCNTL table; therefore, the generated script does not contain a DROP or CREATE statement for the procedure or trigger.

User response: Drop the registration for the missing nickname or the missing table.

ASN1738W The procedure or trigger *name* on the IBMSNAP_PRUNCNTL table could not be found on the remote server.

Explanation: The procedure or trigger definition on the IBMSNAP_PRUNCNTL table recognizes the CCD tables for all the registered nicknames. Each time a registration is created, the procedure or trigger on the IBMSNAP_PRUNCNTL table is dropped and re-created with the existing and new registration information. The definitions stored in the control tables indicate that there are prior registrations on this server, but the procedure on the IBMSNAP_PRUNCNTL table could not be found. The procedure will still be created on the IBMSNAP_PRUNCNTL table with the definitions stored in the control tables.

User response: Ensure that the generated procedure or trigger statement on the IBMSNAP_PRUNCNTL table has all the registered sources.

ASN1739W The column name *column_name* of data type *datatype* in the user table will be created as a new data type *new_datatype* in the CCD table in the non-DB2 relational database. The reason code is *reason_code*.

Explanation: The following value is valid for the reason code:

0 You cannot insert or update a column of data type `TIMESTAMP` in a Sybase or a Microsoft SQL Server database. A CCD table is being created on one of the non-DB2 relational systems with a column using data type `TIMESTAMP` (IBMSNAP_SYBTMSTMP column or the IBMSNAP_MSTMSTMP column). A table in a Sybase or Microsoft SQL Server database cannot have more than one column with a data type of `TIMESTAMP`; therefore, the data type of the column in the user table will be converted to a `Binary(8)` data type.

User response: No action is necessary if the new data type is acceptable. If the new data type is unacceptable, do not run the generated script. Remove the specified column selection from the registration definition and reissue the task.

ASN1740W The CCD nickname *nicknameowner.nickname* exists in the federated server but the remote CCD table *tableowner.tablename* does not exist in the non-DB2 server. This table information will not be included in the re-creation of the PRUNCNTL procedure or trigger.

Explanation: The procedure or trigger definition on the IBMSNAP_PRUNCNTL table recognizes the CCD tables for all the registered nicknames. Each time a registration is created, the procedure or trigger on the IBMSNAP_PRUNCNTL table is dropped and re-created with the existing and new registration information. The remote CCD table does not exist on the non-DB2 server even though the nickname exists in the federated server. This table will not be included in the re-creation of the PRUNCNTL trigger or procedure.

User response: Determine why the CCD table was dropped. If the table was dropped by mistake, either drop the registration of that source or re-create the CCD table with the original definition.

ASN1741W The CCD nickname *nicknameowner.nickname* does not exist in the federated server. This table information will not be included in the re-creation of the PRUNCNTL procedure or trigger.

Explanation: The procedure or trigger definition on the IBMSNAP_PRUNCNTL table recognizes the CCD tables for all the registered nicknames. Each time a registration is created, the procedure or trigger on the IBMSNAP_PRUNCNTL table is dropped and re-created with the existing and new registration information. The CCD nickname does not exist in the federated server. This table will not be included in the re-creation of the PRUNCNTL trigger or procedure.

User response: Determine why the CCD nickname was dropped. If the nickname was dropped by mistake, either drop the registration of that source or re-create the CCD nickname.

ASN1742E The source nickname *nickname_owner.nickname* on the platform *platform* cannot be registered without first migrating registrations to the architecture level *arch_level*.

Explanation: There are existing registrations in the IBMSNAP_REGISTER table for non-DB2 relational sources that pertain to an earlier architecture level. You cannot register a new source until you migrate all registrations to the new architecture level.

User response: Migrate your existing registrations for non-DB2 relational sources to the current architecture level. For details, refer to Migrating to SQL Replication Version 8.

ASN1800E The subscription set *set_name* already exists for the Apply qualifier *apply_qual*, WHOS_ON_FIRST *whos_on_first*, at the Apply control server *server_alias*.

Explanation: There can only be one subscription set with the same name, for a given Apply qualifier and Apply control server. No script is generated.

User response: Create a new set name, or add new members to the existing set.

ASN1801E The statement number *statement_number* is associated with a statement string length *statement_stringlength* that exceeds the maximum statement length for the Apply qualifier *apply_qual*, the set name *set_name*, WHOS_ON_FIRST value *whos_on_first*, at the Apply control server *server_alias*.

Explanation: The length of the statement exceeds the allowed limit (1024 in V8). No script is generated.

User response: Rework the statement string so that its length is less than the allowed limits.

ASN1802W The Replication subscription source member is defined with RECAPTURE='N'. All changes to the replica target will not be propagated to the other replica targets.

Explanation: In an update-anywhere scenario, changes made in one target replica will not be recaptured at the source when RECAPTURE='N'. If there is more than one target replica that subscribes to the same source, then the changes made to one target replica will not be reflected in the other replica targets.

User response: If you want the changes propagated to the other replica targets, set RECAPTURE='Y'.

ASN1803I There exist *orphan_statements* Replication subscription set statements from a previously defined subscription set that was later dropped. These orphan statements are not dropped for the subscription set, for the provided Apply qualifier, at the provided Apply control server.

Explanation: A previous subscription set was dropped, without dropping all its appropriate statements. A script is generated for the new subscription set which shares the same name as the previous subscription set that was dropped. The previous subscription Statements are not dropped.

User response: Issue a Drop Subscription Statements to delete the orphan statements.

ASN1804I The Replication subscription set MAX_SYNCH_MINUTES *maxsynch_minutes* is not within the allowed range for the provided subscription set and Apply qualifier, at the provided Apply control server. The Replication default value is used instead.

Explanation: The valid range for this column is 0 to 999.

User response: No action is required if the default value of 30 minutes is acceptable.

ASN1805I The Replication subscription set COMMIT_COUNT *commitcount_value* is not within the allowed range for the provided subscription set and Apply qualifier, at the provided Apply control server. The Replication default value is used instead.

Explanation: The valid range for this column is 0 to 999.

User response: No action is required if the default value of 0 minutes is acceptable.

ASN1806E The replication action ended in error for the Apply qualifier *apply_qualifier*, subscription-set name *set_name*, WHOS_ON_FIRST value *whos_on_first*, source member *sourceowner.sourcetable*, source view qualifier *source_view_qual*, target member *targetowner.targettable*. The subscription-set member cannot be added to the provided subscription set. Reason code *reason_code*.

Explanation: The subscription set would be not be valid if the member were added. No script is generated. The following values are valid for the reason code:

- | | |
|----|---|
| 0 | The subscription set has reached its maximum limit for members. |
| 1 | The source member for the Capture schema is not the same as the subscription set for the Capture schema. |
| 2 | The iSeries source member is not the same as the subscription set journal. |
| 3 | The condensed table member structure is incompatible with the other member structures. |
| 4 | The source member does not support change-capture replication, but the target member relies on change-capture. The target structure is either a CCD or replica table, but the source has no CD table. |
| 5 | The source member is not a complete table. |
| 6 | The target member definition expects the existence of the target table, but the target table does not exist. |
| 7 | The target member definition asks for creation of the target table, but the target table already exists. |
| 8 | The set contains only full refresh supported target tables, but the new member supports change-capture replication. |
| 10 | The set contains only target tables supported by change-capture replication, but the new member supports full refresh only. |
| 11 | Replica rule: if target member is a replica, the source member can be either a replica or a user table. |
| 12 | The target structure is not supported for this operating system. |

- 13 The target structure is a CCD, which is set as a registration source (autoregistration), but the structure is not complete
- 14 The source member is not registered.
- 15 The source member columns have column definitions, but the target type is not an aggregate.
- 16 At least one of the excluded target columns from the subscription set is neither nullable nor NOT NULL with defaults.
- 17 The target member is a view that cannot be updated.
- 18 The subscription-set member already exists.
- 19 Unable to find a target column or expression with a valid mapping to the registered source.
- 20 Multiple effective sources have been found but have not been defined consistently.
- 21 The external CCD table is noncondensed and contains either DataLink or LOB columns.
- 22 The source member journal library or journal name does not match.
- 23 The remote journal name is not valid.
- 24 The journal name or library is not valid.
- 25 A replica table with remote journaling is not allowed.
- 26 An internal CCD table already exists for the specified registered source table.
- 27 The source and target servers must be the same for internal CCD tables.
- 28 The internal CCD table must be noncomplete.
- 29 The source table is remotely journalled and contains LOBs or DATALINK columns.
- 30 No related information exists in the IBMSNAP_PRUNCNTL table.
- 31 No related information exists in the IBMSNAP_PRUNE_SET table.
- 32 An internal CCD table with a view as a source is not allowed.
- 33 The target table is already shared by another existing member, and a conflict has been detected between the replication definitions with respect to the values of the NAME, IS_KEY, EXPRESSION, and TYPE fields stored in the IBMSNAP_SUBS_COLS table and the definitions requested for the new member.
- 34 The source is on a non-DB2 server and the requested target type is a replica. This scenario is not supported.
- 35 The specified source member does not have a

valid nickname at the target server. The error occurred for one of the following reasons:

- You did not provide a nickname when creating the source member.
- You provided a nickname that does not exist at the target server and one cannot be created for you.
- You provided an object that is not a nickname.
- You provided a nickname that does exist at the target server but it is not associated with the specified source member.

User response: For reason codes 1 through 34, either create the new member in a different subscription set or create a new subscription set for the new member.

For reason code 35, make sure that you provide a valid nickname that already exists at the target server and is associated with the source member.

ASN1807I **The replication subscription member is added to the provided subscription set and Apply qualifier, at the provided Apply control server with an informational clause. Reason code *reason_code*.**

Explanation: This message is for your information only, and no action is required. A script is generated. The following values are valid for reason code:

- 0** The new set results in a mixture of replica and read-only target members.
- 1** The subscription set supports transaction commit counts, but the target member does not qualify for transaction processing.
- 2** At least one member has a target member that is a CCD table, but not all members have a target member CCD table. Different generations of tables are contained in the same set.
- 3** The target member is a non-condensed, non-complete CCD table, without extra columns from IBMSNAP. This target table is essentially the same as the CD table.

User response: Review the reason codes in the explanation, and respond with the following options:

- 0** Consider keeping all the replica tables in one set, and the read-only tables in another.
- 1** Consider keeping all the target types that support transaction processing in the same set, and other tables in another set.
- 2** To maintain a consistent age of data across sets, consider keeping all the CCD target tables, which are part of the same generation, in the same set. Also, consider keeping all

non-CCD target tables in a different set. The second set depends on the data being in the first set, as, for example in the middle-tier stage.

- 3 Consider whether you need the CCD target table.

ASN1808E The replication action ended in error for Apply qualifier *apply_qualifier*, set name *set_name*, WHOS_ON_FIRST *whos_on_first*, source member *sourceowner.sourcetable*, source view qualifier *source_view_qual*, target member *targetowner.targettable*. The subscription target member expects its index key columns to be updated but at least one index key does not have its before-image column registered in the subscription source member.

Explanation: The target table allows its index keys to be updated (PRIMARY_KEY_CHG = 'Y'). To support this requirement, the Apply program needs to access the before-image columns of the index keys. So these before-image columns must exist in the Change Data table for the source member. If they do not exist, the Apply program fails. A script is not generated.

User response: For each column of the subscription target index:

- Check if the before-image column for that column is already registered in the source member, at the Capture server.
- If not, register the appropriate before-image column.

ASN1809W The replication action expects a subscription index key columns to be updated but the subscription member will be added to the subscription set without allowing updates to the target index key columns. Reason code *reason_code*.

Explanation: In the cases listed, the PRIMARY_KEY_CHG settings are meaningless. A script is generated that was updated with Replication definitions that override user-provided definitions. The following values are valid for reason code:

- 0 The target table type CCD:
PRIMARY_KEY_CHG is not valid.
- 1 The value of the CHG_UPD_TO_DEL_INS in the IBMSNAP_REGISTER table is set to 'Y'.
- 2 The target table is not condensed.

User response: This message is for your information only, and no action is required.

ASN1810W The subscription set member is defined on the source server with DB2 referential integrity constraints, but the target member is a replica that does not preserve these constraints.

Explanation: Referential integrity constraints at the target table are not enforced by DB2 at the replica site. This might not be the intended behavior at the replica site. A script is generated, and the script might not be ready to execute.

User response: Update the generated script to include the appropriate referential constraints at the target if needed.

ASN1811W The index definition for the target subscription member cannot guarantee proper uniqueness at the target. Reason code *reason_code*.

Explanation: The Apply program relies on the unique index definition to correctly update and delete rows in the target table, for some target types. If the provided index does not guarantee uniqueness, the Apply program will have some rework. Understand your application to ensure that this will not be the case. The following values are valid for reason code:

- 0 At least one column is generated by an SQL function, which does not guarantee the uniqueness of the index.
- 1 In a unique index, nullable columns are not generated by SQL functions.

User response: For each column of the subscription target index:

- Check if the column type in the Apply control server ASN.IBMSNAP_SUBS_COLS, COL_TYPE, is 'F'
- If so, redefine the index column expression not to include a SQL expression, or remove that column from the index key (ASN.IBMSNAP_SUBS_COLS, IS_KEY column is set to 'N').

ASN1812E The replication action ended in error for Apply qualifier *apply_qualifier*, set name *set_name*, WHOS_ON_FIRST value *whos_on_first*, source member *sourceowner.sourcetable*, source view qualifier *source_view_qual*, target member *targetowner.targettable*. The subscription target member cannot be added because the required target key is not valid. Reason code *reason_code*.

Explanation: Target members that require a unique index are target types of point-in-time, user copy, and replica tables, and condensed CCDs. If these targets do

not have a unique index, the Apply program fails. A script is not generated. The following values are valid for reason code:

- 0 The target table does not already exist but the target key information could not be derived from the source table.
- 1 The target key information cannot be found nor derived, and the RRN is not defined for the CD table (iSeries only).
- 2 The target table or view already exists but the required target key information is missing.
- 3 The target table or view already exists but the target key information is incompatible with the existing partitioning key information.
- 4 The target table of type replica already exists. The specified column is part of the target key information but the source key information does not contain this column. The keys on both the source and the target members must match when the target table is of type replica.
- 5 The specified target key column is not found in the target table definition.

User response: Define a valid target key. Take the specific actions for the following reason codes:

- 0 Create the appropriate unique index on the source table so that it can be used to derive the replication suggested index.
- 2 Provide the required target key information.
- 3 Refer to the SQL reference for the DB2 rules on partitioning indexes. For example, the key that you provided might not include the required partitioning key.
- 4 Modify the target or the source key information appropriately so that replication can derive the correct index.
- 5 Verify the column name.

ASN1813I **The Replication subscription source member is defined on the source server with some DB2 constraints, but the subscription target member does not preserve these constraints. Reason code**
reason_code.

Explanation: Constraints at the source table are not enforced by DB2 if they are not specified during the target member definition. This might not be the intended behavior at the replica table server. A script is generated that might not be ready to execute. Constraints are described in the following valid values for reason code:

- 0 At least one NOT NULL WITH DEFAULT clause in the target member.

- 1 Partitioned table space.

User response: Update the generated script to include the appropriate DB2 constraints at the target, if needed.

ASN1814E **The target column *column_name* of data type *datatype* cannot be added to subscription target member *tableowner.tablename*. Reason code**
reason_code.

Explanation: The subscription member fails the subscription column checks. A script is not generated. The following values are valid for reason code:

- 0 The column data type is not supported by replication. Data types that are not supported by DB2 prompt the replication programs to issue message ASN1648E.
- 1 The target data type is incompatible with the corresponding source data type.
- 2 The column is not found in the source table registration.
- 3 The column type is not supported for federated targets.
- 4 The target column is a LOB. The maximum number of LOB columns is exceeded for the target member.
- 5 The source column contains SQL column function, but the target member structure is neither base aggregate nor change aggregate.
- 6 The target table type is replica, and the source column is a LOB column.
- 7 The target table type is replica, and source column is DATALINK value. But the CONFLICT_LEVEL > 0.
- 8 A noncondensed CCD target table with LOB columns is not supported.
- 9 The column is not in the existing target table.
- 10 The target column is already defined in the subscription member definition.
- 11 The data-type definition of the specified column prevents either an insert or an update operation or both on the column.
- 12 The specified target column name does not match the mapped source column name. The target column name cannot be different from that of the source column name because the target table is a replica.
- 13 The target column is mapped from a before-image column in the source and the source column can contain null values; however, the target column is neither nullable nor not null with default.

- 14 The specified target column name does not match with the mapped CD column name. The target column name cannot be different from that of the CD because the target table is an internal CCD.

User response: Review the reason code in the explanation and respond as follows:

- 0 Change the data type to one that is supported.
- 1 Make sure the target data type matches the source data type.
- 2 Register the column of the source table.
- 3 Choose a valid data type that is supported for federated targets.
- 4 Make sure the number of LOB columns at the target member does not exceed the allowable limit.
- 5 Change either the source column expression or the target table structure.
- 6 Remove the LOB column for the replica target from the subscription member.
- 7 Remove the DATALINK column from the subscription member if the replica needs a conflict level greater than 0. Otherwise, change the replica conflict level.
- 8 Remove the LOB columns.
- 9 Verify the column name.
- 10 Verify the column name.
- 11 Remove the column from the subscription target member.
- 12 Make sure that the target column name matches the mapped column name in the source column name.
- 13 Modify the definition of the column in the target member to allow null or default values.
- 14 Make sure that the target column name matches the mapped column name in the CD table.

ASN1815E The replication action ended in error. The subscription set *set_name* for Apply qualifier *apply_qual*, WHOS_ON_FIRST *whos_on_first* is to be dropped if empty, but at least one member exists for this set. The subscription set cannot be dropped.

Explanation: The subscription set is not dropped because at least one member exists in the ASN.IBMSNAP_SUBS_MEMBR at the provided Apply control server, for the provided Apply qualifier of the particular subscription set. A script is not generated.

User response: Drop the subscription members that

still exist, and then drop the subscription set. Alternatively, issue the *Drop Subscription Set* task, with no requirement for the subscription set to be empty.

ASN1816W The Replication subscription set contains at least one member that will be dropped once the subscription set is dropped.

Explanation: When a subscription set is dropped successfully, all the set members are automatically dropped also.

User response: This message is for your information only, and no action is required.

ASN1818W The Replication subscription member is updated with new subscription Where Clause predicates. Previous predicates already exist for the subscription member. They will be overwritten by the new predicate information.

Explanation: The specified member already contains a predicate. The new predicate overwrites the old one. A script is generated.

User response: Provide the complete predicate clause for the replication task. You might need to update the predicate clause if does not contain all of the existing predicate clause.

ASN1819W The Replication subscription set is disabled successfully. Note however that disabling a subscription set has a direct impact to the Capture pruning logic of all source members for that subscription set.

Explanation: The Capture pruning logic does not prune any CD table until the dependent subscription members have been populated by the Apply program. A script is generated. It might need to be updated, if the disabling of the subscription set is not the ideal choice, and dropping the subscription set is a better option.

User response: If the subscription set is going to be disabled for a considerable amount of time that the pruning process of the CD tables will be impacted, or if the impact to the CD tables for the dependent registered sources will impact dramatically the Capture program and the Capture server CD tables, then consider dropping the subscription set and creating it again later, instead of simply disabling it. Alternatively, deactivate the appropriate registrations.

ASN1820E The replication string for subscription set *set_name*, apply qualifier *apply_qual*, WHOS_ON_FIRST *whos_on_first* contains DB2 syntax that is not valid. The string type is *string_type*, string text is *string_text* and SQL message is *sql_message*.

Explanation: The specified string is not a valid. A script is not generated.

User response: Please correct the appropriate object syntax and issue the replication task again.

ASN1821W Dependent subscription sets will no longer be valid after the existing subscription set is dropped, if this subscription set contains target members that are registered sources at its target server.

Explanation: The dependent subscriptions rely on their source member tables to exist. If these source members are maintained as replication targets, and these targets are dropped, then the Apply program fails when it processes the dependent subscription sets. Dependent subscription sets might be impacted if the captureserver.IBMSNAP_PRUNCNTL table at the target server contains SOURCE_OWNER or SOURCE_TABLE rows for which these values are the target tables being dropped. A script is generated.

User response: Deactivate or drop the dependent subscription sets before running the script, if required.

ASN1822E The replication action ended in error for Apply qualifier *apply_qual*, set name *set_name*, source member *sourceowner.sourcenname*, target member *targetowner.targetname*. The provided subscription member does not exist for the provided subscription set.

Explanation: The specified member cannot be found in the ASN.IBMSNAP_SUBS_MEMBR for the provided Apply qualifier at the provided Apply control server.

User response: Make sure the Apply qualifier, set name, member name, and control server provided are correct.

ASN1823E The subscription set *set_name* does not exist for the Apply qualifier *apply_qual*, WHOS_ON_FIRST *whos_on_first*, at the Apply control server *server_alias*.

Explanation: The specified subscription set cannot be found in the ASN.IBMSNAP_SUBS_SET for the provided Apply qualifier at the provided Apply control server.

User response: Make sure the Apply qualifier, set

name, member name, and control server provided are correct.

ASN1824W The Replication subscription set was updated with a COMMIT_COUNT of 0.

Explanation: The source is a view of multiple tables, and the commit count is null for the set. A Commit count of 0 is enforced for the set.

User response: This message is for your information only, and no action is required.

ASN1825W The replication action to drop a member did not drop the view.

Explanation: Even though the request was made to drop the view the action was not completed as requested.

User response: You need to manually drop the view.

ASN1826W The Capture schema *capture_schema* does not exist at the provided source server.

Explanation: Please make sure the Capture schema exists before adding any members to the subscription set.

User response: Create the Capture server control tables at the source server with the Capture schema specified.

ASN1827W The column *target_columnname* of the target member *target_member* does not preserve a DB2 column attribute of the corresponding column *source_columnname* of the source member *source_member*. Reason code *reason_code*.

Explanation: A DB2 column attribute of the source column differs from the corresponding target column. The following values are valid for reason code:

- | | |
|---|---|
| 1 | The source column is nullable and the target column is not nullable. |
| 2 | The source column is not nullable and the target column is nullable. |
| 3 | The source column has a default value and the target column has none. |
| 4 | The target column has a default value and the source column has none. |

User response: If the reason code is 1, then check whether there are null values in the source column that will be applied to the target column. If necessary, change the target column to NULLABLE. Alternatively, update the generated script to include the appropriate DB2 attributes at the target, if needed.

ASN1828E The replication action ended in error for Apply qualifier *apply_qual*, set name *set_name*, source member *sourceowner.sourcename*, target member *targetowner.targetname*. When the subscription target server is a non-IBM target server, either the action is not supported or is supported with restrictions. Reason Code is *reason_code*.

Explanation: These are current restrictions. No script is generated. The following values are valid for reason code:

- 0 Not supported.
- 1 Supported for the following target table structures: point-in-time, CCD, user copy.

User response: This message is for your information only, and no action is required.

ASN1829I A valid nickname *nicknameowner.nickname* is found for the subscription target table. Data type mapping rules for columns are enforced. The nickname is used as is.

Explanation: An existing target nickname was found in the federated database that is valid for this subscription (the data type mapping checks for columns are valid). However, there is no check to validate the existence of the target table on the non-DB2 relational database server.

User response: Make sure that the remote table exists for the provided nickname. Otherwise, the Apply program fails.

ASN1830E The replication action ended in error for the Apply qualifier *apply_qualifier*, subscription set *set_name*, whos on first value of *whos_on_first*, source member *sourceowner-sourcetable*, source view qualifier *source_view_qual*, target member *targetowner-targettable*, and predicate of *predicate*. The subscription-set member cannot be added to this subscription set. The reason code is *reason_code*.

Explanation: The subscription-set member is invalid, and no script is generated. A possible reason code is:

- 0 The predicate references columns from non-existing CD or UOW tables.

User response: Verify the accuracy of the specified predicate, and refer to the documentation on advanced change predicate features.

ASN1831E The replication action ended in error. No subscription statements exist for the subscription set *set_name* for the Apply qualifier *apply_qual*, with a whos on first value of *whos_on_first*, at the Apply control server *control_server*.

Explanation: No subscription statements exist for the specified subscription-set name with this Apply qualifier.

User response: Verify that the specified subscription-set name under this Apply qualifier contains subscription-set statements.

ASN1832W A column named *column_name* already exists in the ASN.IBMSNAP_SUBS_COLS control table.

Explanation: The specified column already exists in the ASN.IBMSNAP_SUBS_COLS table.

User response: No action is required.

ASN1833E The CCSID *Unicode_ASCII_EBCDIC* of the change data (CD) table *cdowner.cdname* for the source table *sourceowner.sourcetable* does not match the CCSID *Unicode_ASCII_EBCDIC* of the IBMSNAP_UOW table for the capture schema *capture_schema*. The provided subscription member definition would require a join of these two tables.

Explanation: For the given capture schema, the Apply program will join the IBMSNAP_UOW table and the CD table of the given source if either the target type of the associated subscription-set member is not user copy, or if any columns of the IBMSNAP_UOW table are used in the WHERE clause of the subscription-set member. If the Apply program processes such a subscription-set member defined with the given source table and capture schema by joining the CD table of the source table with the IBMSNAP_UOW table, an error will occur because of the different encoding schemes of the tables. For more information about encoding schemes, see "UNICODE and ASCII encoding schemes" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center.

User response: Take one of the following actions:

- Select a target type of user copy and do not use columns of the IBMSNAP_UOW table in the WHERE clause of the subscription member.
 - Register the source using a different capture schema and create the CD table in a table space with the same encoding scheme as those of the IBMSNAP_UOW table of the new capture schema.
-

ASN1834W The default target capture schema of 'ASN' will be used for the subscription set.

Explanation: The subscription set requires a target Capture schema, and the default value of 'ASN' is used.

User response: No action is required if the default is appropriate for the target Capture Schema column in this subscription set.

ASN1835W The target column *column_name* of data type *datatype* has been added to the subscription-set member target *tableowner.tablename* and source *sourceowner.sourcename*, but the corresponding source column *column_name* of data type *datatype* can contain data that is not applicable to the target column. Reason code *reason_code*.

Explanation: The source column definition does not exactly match with the target column definition. Therefore, if the data that is selected from the source by the Apply program is not appropriate for the target column, the Apply program might fail or might modify the source data (by truncating it). Note: If your application does not generate data that will make the Apply program fail, there is no problem with the definition mismatch.

The following values are valid for the reason code:

1

The target column length is less than the resolved source column expression.

2

The target column scale is less than the resolved source column expression.

3

The target column precision is less than the resolved source column expression.

4

The target and source column data types are compatible only for certain source values.

User response: If possible, change the definitions at the target site to be compatible with the source definitions. (This is usually driven by the applications that run at the target site.)

If you must keep the definition mismatch (because you have a specific reason why the target definitions must be different than the source definitions), review your applications to ensure that the definition mismatch will not cause a run-time problem.

ASN1836W The target table *owner.name* will not be dropped because it is registered as a source under the Capture schema *schemaname*.

Explanation: The target table is registered as a source under the specified Capture schema. If the table is dropped, the registration will no longer be valid.

User response: Drop the registration for the table, and then drop the table.

ASN1837W The DB2 target *tableowner.tablename* is not dropped.

Explanation: The target table is a replica or an external consistent-change data (CCD) table (CCD). The table is also registered at the target server so it may be the source for dependent targets. The table cannot be dropped automatically.

User response: Drop the registered source for the replica or external CCD table. Then, manually drop the DB2 target table or delete the replica or external CCD subscription member. The deletion of the member will drop the DB2 target table.

ASN1838E The replication action *action_name* ended in error. The OS/400 table *tableowner.tablename* cannot be used as a replica target subscription member. The reason code is *reason_code*.

Explanation: The following values are valid for the reason code:

0

The specified target table name does not exist or its schema is not created using either the "Create Collection" or "Create Schema" command; therefore, the specified table is not journaled. The journal information of the target table is required when registering the target table in the 'F' direction subscription set for the replica. The table needs to be created on an OS/400 system.

1

The specified target table exists but there is no journaling information for the table. The journal information of the target table is required at the time of registering the target table in the 'F' direction set for the replica.

User response: The action you must take depends on the reason code:

0

If the table target table does not exist: Create the target table, journal it with both before-images and after-images, then reissue the task. If the target table exists: Journal the table with both before-images and after-images, then reissue the task.

- 1 Journal the specified target table with both before-images and after-images, then reissue the task.

ASN1839I The base table *tableowner.tablename* of the source view *viewowner.viewname* has been chosen to be updated by the replica target.

Explanation: The source member is of type view in this replica scenario. Views cannot be updated; therefore, the registered base table of the view that has the largest number of columns will be used as the target member in the 'F' direction.

User response: Verify that your scenario is still valid.

ASN1840W The target database is on an OS/400 system.

Explanation: The target database is on an OS/400 system and the target table space definition will be ignored.

User response: This message is for your information only, and no action is required.

ASN1841E The expression *expression_name* specified in the target column *colname* for the subscription target member *targetowner.targetname* failed. The checks performed against the respective subscription source member *sourceowner.sourcename* returned an SQL error (*sqlcode* = *sqlcode*, *sqlstate* = *sqlstate*).

Explanation: A series of checks are performed on the given expression value for the target member against the source member to ensure data integrity rules. The specified expression failed in one or more of the checks performed.

User response: Refer to the SQL message returned and the SQL reference for any additional information.

ASN1842W The internal CCD table *ccdowner.ccdname* has fewer number of columns than the CD table *cdowner.cdtable* of the source object *sourceowner.sourcename*. For this source, existing members that were defined before the internal CCD was defined might become invalid.

Explanation: The internal CCD table takes precedence over the CD table when Apply tries to apply the changes to the target. An internal CCD is created for the specified source, which already has a CD table and subscription members defined. Some of the columns in the CD table do not exist in the internal CCD table and those members that were defined before the CCD table was created will become invalid.

User response: Make sure that your scenario is valid.

ASN1843E The replication action *action_name* ended in error for Apply qualifier *apply_qualifier*, set name *set_name*, WHOS_ON_FIRST *whos_on_first*, source member *source_owner.sourcetable*, source_view_qualifier *source_view_qualifier*, and target member *targetowner.targetname*.

Explanation: For OS/400 only. Duplicate rows are not allowed in the IBMSNAP_REG_EXT table for each combination of source_owner, source_table, and source_view_qualifier.

User response: Delete the row from the IBMSNAP_REG_EXT table that does not belong and reissue the task.

ASN1844E The replication action *action_name* ended in error for Apply qualifier *apply_qualifier*, set name *set_name*, WHOS_ON_FIRST *whos_on_first*, source member *source_owner.sourcetable*, source_view_qualifier *source_view_qualifier*, and target member *targetowner.targetname*. The specified source table and the source table registration reside on different servers.

Explanation: The source table and the source table registration can reside on different servers only if both the source database and the Capture database are located on iSeries machines. At least one of the servers is not on an iSeries system.

User response: Make sure the source table and the source table registration reside on the same server before reissuing the task.

ASN1845E The replication action ended in error. The table name length *length* of the target member *targetowner.targetname* exceeds the allowed limit of *max_limit* imposed by the z/OS replication control tables on the server *server*.

Explanation: The z/OS replication control tables are created with the maximum length of the table name as 30 characters. The error message was received because either the Capture control server or the Apply control server or both are on a z/OS system, and the specified target table name length is greater than 30. The subscription set member cannot be created.

User response: Ensure that the table name length of the target member is not longer than 30 characters. In some cases the source database restrictions limit the capabilities at the target, because the target table name must be stored at the Capture control server and at the Apply control server. In such cases, you can create a

view of the target table using a name that is less than or equal to 30 characters, and define the subscription set member using the target view name.

ASN1846E The set contains one or more invalid members. The original error message was *original_message*.

Explanation: An error has occurred while promoting the given set because the set has one or more invalid members.

User response: Verify the information that you are inputting is correct, and reissue the task.

ASN1847E The view *viewowner.viewname* contains an unsupported definition.

Explanation: You are unable to map a view column to the corresponding base table column because the view definition is unsupported.

User response: Verify that the view definition is correct, and reissue the task.

ASN1848W The index or constraint *index_constraint_name* could not be created on *targetowner.targetname* of type *object_name*. The view definition did not provide a column mapping for all key columns.

Explanation: The target object is a view. A target key will be created on the view. The unique index supporting this target key could not be created on the base table of the view because the view columns used for the target key could not be mapped to the base table columns.

User response: Create the index on the base table manually.

ASN1849E The replication subscription set *set_name* cannot be promoted for the specified Apply qualifier at the specified Apply control server. The reason code is *reason_code*.

Explanation: The following value is valid for the reason code:

0 Either the Capture server or the target server is a federated database. Promote subscription works only for subscription sets on two DB2 systems.

User response: This is a current restriction, and no action is required.

ASN1850W The value of the column *column_name* of data type decimal will be truncated when it is mapped to a column with data type of integer.

Explanation: The specified column can contain rational digits, but it is mapped to an integer column. Some of the digits in the original data might be truncated.

User response: This message is for your information only, and no action required.

ASN1851E The DB2 object *object_name* of type *type* cannot be added to the table *tableowner.tablename* because a conflict exists between the new properties and the partitioning key of that table.

Explanation: The specified object cannot be added to the table definition because the new definitions and the existing table properties cannot coexist according to the DB2 rules for the coexistence of table constraints for partitioned tables.

User response: See the *SQL Reference* for further details on table constraints.

ASN1852E The DB2 object *object_name* of type *type* cannot be added to the table *tableowner.tablename* because a conflict exists between the new properties and the existing partitioning key *partitioning_key*.

Explanation: The specified object cannot be added to the table definition because the new definitions and the existing table properties for the partitioning key cannot coexist according to the DB2 rules for the coexistence of table constraints for partitioned tables.

User response: See the *SQL Reference* for further details on table constraints.

ASN1853E The DB2 table *tableowner.tablename* definitions conflict with the index *indexowner.indexname* and the partitioning keys of the table.

Explanation: The specified index definition and the partitioning keys have conflicting definitions and cannot co-exist according to the DB2 rules for the coexistence of index properties and partitioning keys for partitioned tables.

User response: See the *SQL Reference* for further details on table constraints.

ASN1854E The definition of the table *tableowner.tablename* is not complete on a z/OS platform. The reason code is *reason_code*.

Explanation: The following value is valid for the reason code:

- 0 The table definition is missing the required properties such as unique indexes supporting a ROWID, a primary key, and unique constraints, or LOB column properties such as a ROWID, LOB table spaces, auxiliary tables, and auxiliary indexes. At least one of the these mentioned needs to be present in the table definition to make it complete.

User response: Redefine the table definition to include the required table properties.

ASN1855E The after-image column *column_name* used for the target key change was not found in the replica table *tableowner.table_name*.

Explanation: You selected the target key change option for the key column during subscription. The corresponding after-image column was not found in the replica source table.

User response: Make sure that the original after-image column is available for subscription in the replica source.

ASN1856E The replication action *action_name* ended in error. The check for the OS/400 collection name *collection_name* failed. The native OS/400 message is *OS/400_message*.

Explanation: The specified OS/400 collection is not valid. The script will not be generated

User response: Refer to the iSeries Console Log for more details on error information.

ASN1857E The replication action *action_name* ended in error. The source table *sourceowner.sourcename* has an internal CCD table that is defined as condensed, and the target *targetowner.targetname* is being requested as a non-condensed target.

Explanation: The definitions stored in the replication control tables indicate that the source member already exists in another subscription set with an internal CCD defined as condensed. The same source member is being added as a source in another subscription set, but the target member is non-condensed. The data at the source side is already condensed and cannot be copied to a non-condensed target. The script will not be generated

User response: Modify the target member definitions and reissue the task.

ASN1858E The table type *type* for table *tableowner.tablename* is not supported for this DB2 platform.

Explanation: The specified type of the table or view is not supported for the DB2 platform. You can include only certain types of tables as sources or targets in a subscription, depending on the DB2 platform.

User response: For more information about platform-specific support for replicated tables, see "Registering DB2 tables as sources" and "Subscribing to sources for SQL Replication" in the DB2 Information Center.

ASN1859E The replication action *action_name* ended in error for Apply qualifier *apply_qualifier*, set name *set_name*, WHOS_ON_FIRST *whos_on_first*, source member *source_owner.sourcetable*, source_view_qualifier *source_view_qualifier*, target member *targetowner.targetname*. The crossloader option is being used for this member, but the values for the nickname owner and the nickname table are missing.

Explanation: The value for the LOADX_TYPE column in the IBMSNAP_SUBS_MEMBR table is 3 for the specified member. This value indicates that the crossloader will be used for this member. A nickname for the source table must be created at the target server and specified in the LOAD_SRC_N_OWNER and LOAD_SRC_N_TABLE columns. One or both of these values are missing.

User response: Provide a value for either the LOAD_SRC_N_OWNER field or the LOAD_SRC_N_TABLE field, and reissue the task.

ASN1860W The target member *targetowner.targetname* does not exist, and a script to drop the member will not be generated.

Explanation: The target table, view or nickname does not exist on the target database and it will not be dropped.

User response: This message is for your information only, and no action is required.

ASN1861E The subscription set member cannot be created for target *targetowner.targetname* because the target is already a member of the set, and the existing member definition conflicts with the requested definition.

Explanation: The columns in the

IBMSNAP_SUBS_COLS control table already contain values for NAME, IS_KEY, EXPRESSION, and TYPE and those values differ from the ones requested for the new member. A script for creating the new member is not generated.

User response: To create the new subscription set member, use the same target properties as those that are stored in the IBMSNAP_SUBS_COLS table.

ASN1862E The replication action ended in error for the Apply qualifier *apply_qualifier*, subscription-set name *set name*, WHOS_ON_FIRST value *whos_on_first*, source member *sourceowner.sourcetable*, source view qualifier *source_view_qual*, target member *targetowner.targettable*. The subscription set member cannot be updated for the provided set.

Explanation: The specified source member does not have a valid nickname at the target server. The error occurred for one of the following reasons:

- You did not provide a nickname when creating the source member.
- You provided a nickname that does not exist at the target server and one cannot be created for you.
- You provided an object that is not a nickname.
- You provided a nickname that does exist at the target server but it is not associated with the specified source member.

User response: Make sure that you provide a valid nickname that already exists at the target server and is associated with the source member.

ASN1863E The replication action ended in error. The table owner length *length of the source member sourceowner.sourcename* exceeds the allowed limit of *max_limit* imposed by the replication control tables on the server *server*.

Explanation: The Capture server is on a Version 8 new-function mode DB2 subsystem and the Apply control server is either on a workstation or an earlier version of DB2 for z/OS. The maximum length allowed for a schema name on a Version 8 new-function mode DB2 subsystem is much longer than the length allowed on this server.

User response: Redefine your setup to use a Version 8 new-function mode DB2 subsystem as the Apply control server and issue the task.

ASN1864E The replication action ended in error. The table owner length *length of the target member targetowner.targetname* exceeds the allowed limit of *max_limit* imposed by the replication control tables on the server *server*.

Explanation: The target server is on a Version 8 new-function mode DB2 subsystem and either the Apply control server or the Capture control server is on a workstation or an earlier version of DB2 for z/OS. The maximum length allowed for a schema name on a Version 8 new-function mode DB2 subsystem is much longer than the length allowed on the control server.

User response: Redefine your setup to use a Version 8 new-function mode DB2 subsystem as either the Apply control server or the Capture control server.

ASN1865E The replication action ended in error. The architecture level *arch_level* on the Capture control server for the Capture schema *capture_schema* is not supported on the Version 8 new-function mode DB2 subsystem.

Explanation: The specified Capture control server is on a Version 8 new-function mode DB2 subsystem and the architecture level value obtained from the *capture_schema*.IBMSNAP_REGISTER table must be 0805. The architecture level obtained from the table is not 0805. You are running replication in Version 8 new-function mode; however, you did not migrate your control tables to support the new architecture level. No script is generated.

User response: Run the migration program AASNSAMP member (ASNM2V8) to alter the replication control table definitions. This migration program sets the architecture level to 0805 and alters some of the columns of the control tables to a longer length.

ASN1866E The replication action ended in error. The architecture level *arch_level* on the Apply control server is not compatible with the Version 8 new-function mode DB2 subsystem where it is defined.

Explanation: The specified Apply control server is on a Version 8 new-function mode DB2 subsystem and the architecture level obtained from the ASN.IBMSNAP_SUBS_SET table does not support Version 8 new-function mode DB2 subsystems. No script is generated.

User response: Migrate the Apply control tables to the latest architecture level and then retry the action.

ASN1867E The replication action ended in error for the Apply qualifier *apply_qualifier*, subscription-set name *set_name*, WHOS_ON_FIRST value *whos_on_first*, source member *sourceowner.sourcetable*, source view qualifier *source_view_qual*, target member *target_owner.target_table*. The subscription-set member cannot be added to the provided subscription set without first migrating subscription sets to the current architecture level. Reason code *reason_code*.

Explanation: The subscription set would not be valid if the member were added. No script is generated. The following values are valid for the reason code:

- 0 There are existing registrations in the schema.IBMSNAP_REGISTER table for Oracle sources that pertain to an earlier architecture level. You cannot add a new member to this set until you migrate all registrations to the new architecture level.
- 1 There are existing subscription-set members in the ASN.IBMSNAP_SUBS_SETS table for Oracle sources that pertain to an earlier architecture level. You cannot subscribe to a new source until you migrate all subscription-set members to the new architecture level.

User response: Migrate your existing registrations for non-DB2 relational sources to the current architecture level. For details, refer to Migrating to SQL Replication Version 8.

ASN1900E The table or view *objectowner.objectname* cannot be promoted to the new server. Reason code *reason_code*.

Explanation: The following values are valid for reason code:

- 0 The table type on this operating system is not supported for promote request.
- 1 The source server operating system required for the promote needs to match the target server operating system.
- 2 The table or view does not exist.

User response: Review the reason code in the explanation and respond as follows:

- 0 This message is for your information only, and no action is required.
- 1 Current restriction.
- 2 Verify that the table or view exists on the source server operating system.

ASN1901E The registered source *source_owner.source_name* cannot be promoted for Capture schema *capture_schema* at Capture server *capture_server*. Reason code *reason_code*.

Explanation: The following values are valid for the reason code:

- 0 The table or view is not a registered source.
- 1 The registered source is a replica table.
- 2 The registered source is on DB2 for iSeries but has a remote journal.
- 3 The table or view has already been promoted.
- 4 A view on a view is not supported by the replication promote registration function.

User response: Review the reason code in the explanation, and respond as follows:

- 0 The table or view name that you specified in the IBMSNAP_REGISTER table contained no entries for the specified Capture schema. The table or view registration cannot be promoted for this particular Capture schema.
- 1 The table you specified is of type replica (with a SOURCE_STRUCTURE column value of 7) in the captureschema.IBMSNAP_REGISTER table. The table cannot be promoted as a registered source. A replica can be promoted only within the context of a subscription set to ensure that proper definitions are maintained between the source user table and the replica target.
- 2 The registered source is maintained on DB2 for iSeries with a remote journal, which can be promoted only with SQL script.

ASN1902W Make sure that the schemas exist on the promoted Capture server before running the script. The Replication definitions will be incomplete if the object does not exist at the promoted Capture server.

Explanation: The promote tasks allows you to provide a new Capture server and new Capture server schemas. However, the promote tasks do not connect to the new Capture server to verify the names and existence of the Capture server and schemas. You must verify this information before running the script to ensure that the script executes successfully.

User response: To generate the required objects, run the appropriate SQL prior to running the script.

ASN1903W The object *objectowner.objectname* does not exist on the promoted Apply control server, yet some promoted objects depend on its existence. Failure to create this object will result in incomplete Replication definitions at promoted Apply control server. Reason code *reason_code*.

Explanation: Since the promote tasks allow you to provide a new Apply control server name, the tasks detect whether some required objects exist to ensure proper execution of the generated script. A script is generated, but is not ready to run. The following values are valid for reason code:

- 0 The Apply control server control tables do not exist.
- 1 The registration information for all source members of a promoted set.

User response: To generate the required objects, run the appropriate SQL prior to running the script.

ASN1904I The Replication subscription member is promoted successfully for the provided Apply qualifier, at the provided Apply control server. Reason code *reason_code*.

Explanation: This message is for your information only; no action is required. A script is generated that might need some updates before being executed. The following values are valid for reason code:

- 0 The source member structure is incompatible with the target member structure.
- 1 The target member is a replica (replica1) that is also the source member of another replica (replica2.) The RECAPTURE value for the registration row of replica2 does not allow updates from the user table to be replicated at replica2.
- 2 The source member is a user table that is also the source member of more than one replica (replica1 and replica2). The RECAPTURE value for the registration row of the user table does not allow updates at replica1 to be replicated at replica2, and vice-versa.

User response: Review the reason code in the explanation, and respond as follows:

- 0 Check the ASN.IBMSNAP_SUBS_MEMBR table, TARGET_STRUCTURE column. The value in the column should be compatible with the corresponding source member captureschema.IBMSNAP_REGISTER table, SOURCE_STRUCTURE column.
- 1, 2 Update the values, if needed.

ASN1905W The Capture server alias and the Capture schema name on both the host system and new system are the same. The generated replication definitions cannot work if run on the host system.

Explanation: The promote task detected that the Capture server alias and the Capture schema name are the same on both on the host and new system. The generated SQL script must be modified, or it will fail when executed.

User response: Take one of the following actions: 1) Run the same task with different Capture server alias and Capture schema name values for the host or new system. 2) Or, change the Capture server alias or Capture schema name in the generated script.

ASN1950E ASNCLP : An unexpected token *token_name* was found. Valid tokens include *list_of_tokens*.

Explanation: The command was entered with incorrect syntax.

User response: Check the documentation to verify your command syntax.

ASN1951E ASNCLP : The command was entered with profile *profile_name* that is not valid.

Explanation: A profile must exist before it can be used in a command.

User response: Issue the corresponding SET PROFILE command, and then re-enter the command that failed.

ASN1952E ASNCLP : The program encountered an internal error.

Explanation: The Replication command line processor encountered an unrecoverable error condition.

User response: Obtain the log file with the error, and contact IBM Software Support.

ASN1953I ASNCLP : Command completed.

Explanation: All commands of this ASNCLP session completed successfully. Please note that some individual commands in this session may have produced errors, warnings, or informational messages.

User response: Check the ASNCLP log file for any errors, warnings, or informational messages produced by the commands in this session.

ASN1954E ASNCLP : Command failed.

Explanation: At least one command in the ASNCLP session failed, and processing stopped.

User response: Look in the ASNCLP log file to diagnosis the error. Then fix the error, and try the command again.

ASN1955I ASNCLP : The program will use the following files: *capture_script_file_name* for the Capture SQL script, *control_script_file_name* for the control SQL script, *target_script_file_name* for the target SQL script, and *log_file_name* for the log file.

Explanation: The ASNCLP session generated information in the specified files.

User response: This message is for your information only, and no action is required.

ASN1956I ASNCLP : The program now generates the script for action: *action_name*.

Explanation: All input for this command has been successfully parsed, and the command that generates the script is now invoked.

User response: This message is for your information only, and no action is required.

ASN1957E ASNCLP : The value *value* for the input parameter *input_parameter* is incorrect. The reason code is *reason_code*.

Explanation: The value of the input parameter is incorrect. The following values are valid reason codes:

- 1
The input parameter is a character value but should be a numeric value.
- 2
The input parameter is a numeric value but should be a character value.
- 3
The command line processor cannot access the specified file.
- 4
The date must be in the format yyyy-mm-dd and the time must be in the format hh:mm:ss.ffffff. All components of date and time are mandatory except for the fraction part of time, denoted by fffffff. The hyphens and colons are also mandatory.
- 5

The value of RELEASE must be 91 (Version 9.1), 95 (Version 9.5), 97 (Version 9.7), or 97.3 (Version 9.7 Fix Pack 3).

User response: Check the reason code, and provide a valid input parameter value.

ASN1976E *pgmname* : *program_qualifier*. The specified database alias *db_alias_name* already exists in the password file *password_file_name*.

Explanation: The key that you specified already exists in the password file.

User response: Enter this command again using the MODIFY parameter instead of the ADD parameter.

ASN1977E *pgmname* : *program_qualifier*. The value of the input parameter *parameter_name* is missing.

Explanation: The indicated input parameter must be specified.

User response: Invoke the utility again using a valid input parameter value.

ASN1978E *pgmname* : *program_qualifier*. The value of the input parameter *parameter_name* is too long.

Explanation: The asnpwd command supports a maximum of eight characters for the length of the database alias and a maximum of 128 characters for the length of both the user ID and the password. The specific lengths of the user ID and the password are dependent upon the operating system that you are using.

User response: Invoke the API using an input parameter with a valid length.

ASN1979E *pgmname* : *program_qualifier*. The program encountered an unexpected token *token_name*. Expected tokens include *list_of_tokens*.

Explanation: The command was entered with incorrect syntax.

User response: Check the documentation to verify your command syntax.

ASN1980E *pgmname* : *program_qualifier*. The program did not complete successfully because *reason*.

Explanation: The asnpwd command encountered system problems as indicated in the message.

User response: Take action based on the information

in the message. Enter the command again after fixing the error.

ASN1981I *pgmname : program_qualifier. The program completed successfully using password file password_file_name.*

Explanation: The asnpwd command completed successfully.

User response: This message is for your information only, and no action is required.

ASN1982E *pgmname : program_qualifier. The specified database alias db_alias_name does not exist in the password file password_file_name.*

Explanation: The key that you specified with the MODIFY or DELETE parameter does not exist in the password file.

User response: Enter the command again using the ADD parameter.

ASN1983E *pgmname : program_qualifier. The program cannot find the password file password_file_name.*

Explanation: No password file was found.

User response: Verify that the password file exists in the specified path. If you are using the Password Management utility for the first time, use the INIT parameter.

ASN1984E *pgmname : program_qualifier. The program cannot be initialized because the password file password_file_name already exists.*

Explanation: The password file already exists in the specified path.

User response: Verify that the password file has been deleted. Then retry the command.

ASN1985E *pgmname : program_qualifier. The program encountered an internal error when using the password file password_file_name.*

Explanation: The operating system produced an unexpected error when trying to access the password file. No information about this error is available. However, this error might have occurred if the password file was manually edited causing the format of the file to change.

User response: Retry the command. If the problem persists, use the INIT parameter to create a new password file.

ASN1986E *pgmname : program_qualifier. The password file file_name contains encrypted information that cannot be listed.*

Explanation: Password files created using Encrypt All contain only encrypted data (alias, user ID, and password). Encrypted data cannot be listed. Data can be listed only from password files created using the Encrypt Password parameter.

User response: If you want to use the List option you must maintain a password file in which only the password is encrypted. You must delete the old password file using the asnpwd delete command and create a new password file using the asnpwd init encrypt password command.

ASN1987E *program_name : program_identifier : The option_name option that was specified for the ASNPWD command is not supported on this operating system.*

Explanation: On 64-bit Windows operating systems, the ADD, MODIFY, DELETE, and LIST options of the ASNPWD command are not supported for password files that were created using the ASNPWD command before Version 9.5 Fix Pack 2. The ASNPWD command began using a new encryption method starting with this fix pack.

User response: Take one of the following actions:

- Create a new password file using the ASNPWD INIT option. If you use this method, the password file can only be decrypted and read by replication programs that are at or newer than Version 9.5 Fix Pack 2. Do not use this method if an older replication program needs to read the password file.
- Use a different operating system to change the password file by running the ASNPWD command with the ADD, MODIFY, DELETE, or LIST options.

ASN1988E *program_name : program_identifier : The password file name contains encrypted information that cannot be read, listed, or updated.*

Explanation: Replication programs can read only password files that were created using the Advanced Encryption Standard (AES) algorithm. This algorithm is used by the asnpwd command when it creates encrypted password files.

User response: Create a new password file by using the asnpwd command and the INIT ENCRYPT option, for example:

```
asnpwd INIT ENCRYPT PASSWORD USING
mypwd.aut.
```

Chapter 39. ASN2000 - ASN2499

ASN2000I The action *action_name* started at *time_stamp*. The Q Capture server is *capture_server*, and the Q Capture schema is *capture_schema*.

Explanation: This message is for your information only.

User response: This message is for your information only. No action is required.

ASN2001I The action *action_name* started at *time_stamp*. The Q Apply server is *apply_server*, and the Q Apply schema is *apply_schema*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2002I The action *action_name* started at *time_stamp*. The publication name is *pub_name*, the Q Capture server is *capture_server*, the Q Capture schema is *capture_schema*, and the source table, view, or nickname is *table_name*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2003I The action *action_name* started at *time_stamp*. Q subscription name: *name*. Q Capture server: *capture_server*. Q Capture schema: *capture_schema*. Q Apply server: *apply_server*. Q Apply schema: *apply_schema*. The source table is *table_name*. The target table or stored procedure is *table_name*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2004I The action *action_name* started at *time_stamp*. The publishing queue map name is *queue_map_name*, the Q Capture server is *capture_server*, and the Q Capture schema is *capture_schema*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2005I The action *action_name* started at *time_stamp*. The replication queue map name is *queue_map_name*, the Q Capture server is *capture_server*, the Q Capture schema is *capture_schema*, the Q Apply server is *apply_server*, and the Q Apply schema is *apply_schema*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2006I The action *action_name* ended successfully at *time_stamp* for the Q Capture server *capture_server* and Q Capture schema *capture_schema*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2007I The action *action_name* ended successfully at *time_stamp* for the Q Apply server *apply_server* and Q Apply schema *apply_schema*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2008I The action *action_name* ended successfully at *time_stamp* for the publication name *pub_name*, Q Capture server *capture_server*, Q Capture schema *capture_schema*, and source table, view, or nickname *table_name*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2009I The action *action_name* ended successfully at *time_stamp* for the Q subscription name *name*. The Q Capture server is *capture_server*, the Q Capture schema is *capture_schema*, the Q Apply server is *apply_server*, the Q Apply schema is *apply_schema*, the source table is *table_name*, the target table or stored procedure is *table_name*.

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2010I **The action *action_name* ended successfully at *time_stamp* for the publishing queue map name *queue_map_name*, Q Capture server *capture_server*, and Q Capture schema *capture_schema*.**

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2011I **The action *action_name* ended successfully at *time_stamp* for the replication queue map name *queue_map_name*. The Q Capture server is *capture_server* and the Q Capture schema is *capture_schema*. The Q Apply Server is *apply_server* and the Q Apply schema is *apply_schema*.**

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN2018E **The value is missing for the input parameter *input_parameter*.**

Explanation: The value for the input parameter is mandatory but it is missing. A script is not generated.

User response: Provide a value for the mandatory parameter and rerun the action.

ASN2019E **The value *value* for input parameter *input_parameter* is incorrect. Reason code: *reason_code*.**

Explanation: The value provided for the input parameter is not a valid value. The following values are valid for the reason code:

0

In this context, the SUBTYPE column value must be 'U' or 'B'.

1

The SLEEP_INTERVAL column value must be greater than 0 and less than 2147483647.

2

The ALL_CHANGED_ROWS column value must be 'Y' or 'N'.

3

The BEFORE_VALUES column value must be 'Y' or 'N'.

4

The CHANGED_COLS_ONLY column value must be 'Y' or 'N'.

5

The HAS_LOADPHASE column value must be 'N' or 'E' for publications.

6

The HAS_LOADPHASE column value must be 'N', 'T', or 'E' for Q Subscription.

7

The SUPPRESS_DELETES column value must be 'Y' or 'N'.

8

The MESSAGE_FORMAT column value must be 'X' for publishing queue maps.

9

The MESSAGE_FORMAT column value must be 'C' for replication queue maps.

10

The MSG_CONTENT_TYPE column value must be 'T' or 'R' for publishing queue maps.

11

The MSG_CONTENT_TYPE column value must be 'T' for replication queue maps.

12

In the IBMQREP_SENDQUEUES table, the ERROR_ACTION value must be 'T' or 'S' for publishing queue maps or replication queue maps.

13

The ERROR_ACTION column value for Q subscriptions must be 'T', 'S', 'Q', or 'D'.

14

The CONFLICT_ACTION column value must be 'F', 'T', 'D', 'S', or 'Q' for Q subscriptions.

15

The LOADTYPE parameter must be '0', '1', '2', '3', '4', or '5'.

16

The SOURCENODE parameter must be greater than zero.

17

The TARGETNODE parameter must be greater than zero.

18

| | | | |
|----|--|----|--|
| 19 | The NUM_APPLY_AGENTS parameter must be greater than 0 and less than 100. | 34 | The STARTMODE column value must be 'COLD', 'WARMSI' or 'WARMNS'. |
| 20 | The MEMORY_LIMIT parameter must be greater than 0 and less than 100. | 35 | The COMMIT_INTERVAL column value must be greater than 100 and less than 600000. |
| 21 | The MSG_FORMAT column value for Q subscriptions must be 'C'. | 36 | The SIGNAL_LIMIT column value must be greater than 0 and less than 2147483647. |
| 22 | The valid values for the STATE parameter is 'A' only. | 37 | The length of the DBSPACE name is longer than 18 characters. |
| 23 | The TARGET_TYPE parameter must be '1' (user table) or '5' (stored procedure). | 38 | The specified option for the lock size is not valid. |
| 24 | The SUBNAME parameter cannot contain the semicolon character (;). | 39 | The number of pages specified cannot be less than 0. |
| 25 | The SENDQ or RECVQ column value cannot contain special characters or blanks. | 40 | The number of header pages cannot be less than 1 or greater than 8. |
| 26 | The MONITOR_INTERVAL column value must be greater than 0 and less than 2147483647. | 41 | The specified value for percent index, percent free space for data pages, or percent free space for index pages cannot be less than 0 or greater than 99. |
| 27 | The MONITOR_LIMIT column value must be greater than 0 and less than or equal to 10080. | 42 | The value for the DBSPACE name is not valid. |
| 28 | The TRACE_LIMIT column value must be greater than 0 and less than 2147483647. | 43 | The parameter loadphase must be 'N' for all target objects that are not tables. |
| 29 | The PRUNE_INTERVAL column value must be greater than 0 and less than 3600. | 44 | Targets that are not tables (such as nicknames and stored procedures) cannot be loaded. Therefore, the parameter HAS_LOADPHASE must be 'N'. |
| 30 | The AUTO_STOP column value must be 'Y' or 'N'. | | The value for the input parameter PARALLEL SENDQS cannot be greater than 4. |
| 31 | The LOG_REUSE column value must be 'Y' or 'N'. | | User response: Provide a valid value for the input parameter, and rerun the action. See the Q Replication documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for details. |
| 32 | The LOGSTDOUT parameter must be 'Y' or 'N'. | | |
| 33 | The TERM parameter must be 'Y' or 'N'. | | |

ASN2020E The value of the input parameter *input_parameter* is too long. The length of the value *value* exceeds the maximum allowed limit *maximum_limit*.

Explanation: See message text.

User response: Specify a valid value for the input parameter. See the documentation in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center for more information about the valid values for each parameter.

ASN2021E The action ended in error. An internal error occurred. Return code *return_code*.

Explanation: The specified action cannot be performed because of an internal or a run-time error.

User response: Contact IBM Software Support. Report the return code that was returned in the message. Provide an Analyzer report.

ASN2022E The action ended in error. An SQL error was encountered. SQL message is *sql_message*.

Explanation: An error occurred during the execution of an SQL statement.

User response: Refer to your database message reference for SQL.

ASN2023E The database platform *platform*, version *version* on server *server_alias* is not supported.

Explanation: Q Replication and event publishing are not supported on the specified platform, version, or release.

User response: See the system requirements page on the replication technical support Web site for details about supported platforms, versions, and releases.

ASN2050E The receive queue *queue_name* specified for the replication queue map *queue_map_name* is already used by another replication queue map at the Q Apply server *server_alias* and the Q Apply schema *apply_schema*.

Explanation: The receive queue can be used only once for one replication queue map and under a single Q Apply schema.

User response: Provide a different receive queue name that is not being used by any other replication queue map for this Q Apply schema, and rerun the action.

ASN2051E Control tables already exist at the given Q Capture server for Q capture schema *capture_schema*.

Explanation: At least one Q Capture control table already exists under the given schema at this server. No script is generated.

User response: Rerun the 'Create Capture Control Tables' action under a different Q Capture schema.

ASN2052E Control tables already exist at the given Q Apply server for Q Apply schema *apply_schema*.

Explanation: At least one Q Apply control table already exists under the given schema at the provided server. No script is generated.

User response: Rerun the 'Create Apply Control Tables' action under a different Q Apply schema.

ASN2053E No control tables were found for Q Capture schema *capture_schema*.

Explanation: No control tables exist on the Q Capture server for the specified Q Capture schema. If the action 'Drop Capture Control Tables' was issued, no script is generated to drop the control tables. If any other action, such as 'Create publication', was issued, then the control tables could not be retrieved from the database under the specified Q Capture schema and a script for the action is not generated.

User response: Issue the action again for the appropriate Q Capture server and Q Capture schema.

ASN2054E No control tables were found for Q Apply schema *apply_schema*.

Explanation: No control tables exist on the Q Apply server for the specified Q Apply schema. If the action 'Drop Apply Control Tables' was issued, no script is generated to drop the control tables. If any other action, like 'Create Q Subscription' was issued, then the control tables could not be retrieved from the database under the specified Q Apply schema and a script for the action is not generated.

User response: Issue the action again for the appropriate Q Apply server and Q Apply schema.

ASN2055I The replication control table *table_name* was not found and was not dropped.

Explanation: The 'Drop Q Capture control tables' or 'Drop Q Apply control tables' action was issued and the control table was missing. The script will not generate the appropriate DROP statement for that control table.

User response: This message is for your information only. No action is required.

ASN2056W At least one row was found in the control table *table_name*. If you choose to drop this control table, any replication definitions stored in the table will also be dropped.

Explanation: A script was generated to drop the specified control table, but that table is not empty. If you run the script, replication control information that is in the table will be deleted and your existing publications or Q subscriptions will be affected.

User response: Determine what would happen to your publications or Q subscriptions if you dropped the control table. Run the generated script only if the result is acceptable to you.

ASN2057W The table space *tablespace_name* already exists at the given server. The control tables will be created in this table space.

Explanation: The control tables are created in recommended table spaces by default if you do not override the default and specify other table spaces. In this situation, the recommended table space already exists, so the control tables will be created in an existing table space.

User response: Modify the script to specify a different table space if you do not want to create the tables in the existing table space. No action is required if you want the control tables to be created in the existing table spaces.

ASN2058E The Q Capture control table *table_name* was not created in table space *tablespace_name* in database *logicaldb_name* because the locksize is *locksize*. This control table must be created in table space with locksize *expected_locksize*.

Explanation: On the z/OS platform, some control tables require a table space with row locksize, while others require a table space with page locksize. The table space selected for the control table does not have the appropriate locksize.

User response: Issue the 'Create the Q Capture control tables' again using table spaces with the appropriate locksize.

ASN2059E The Q Apply control table *table_name* was not created in table space *tablespace_name* in database *logicaldb_name* because the locksize is *locksize*. This control table must be created in table space with locksize *expected_locksize*.

Explanation: On the z/OS platform, some control tables require a table space with row locksize, while others require a table space with page locksize. The

table space selected for the control table does not have the appropriate locksize.

User response: Issue the 'Create the Q Apply control tables' again using table spaces with the appropriate locksize.

ASN2075E The object *object_name* of type *object_type* exists at the specified server and under the specified Q Capture schema.

Explanation: Duplicate values for the specified type are not allowed in the control tables for the same Q Capture schema. No script is generated.

User response: Provide a different name for the object that is unique for the Q Capture schema, and rerun the action.

ASN2076E The send queue object *sendqueue_name* specified for publishing queue map *queue_map_name* is already used by another publishing queue map at the Q Capture server *server_alias* and the Q Capture schema *capture_schema*.

Explanation: The send queue can be used only once for one publishing queue map and under a single Q Capture schema.

User response: Provide a different send queue name that is not being used by any other publishing queue map for the Q Capture schema, and rerun the action.

ASN2077E The object *object_name* of type *object_type* does not exist at the server *capture_server* for the Q Capture schema *capture_schema*.

Explanation: The action was invoked on a Q subscription whose catalog definitions were not found in the Q Capture control tables.

User response: Drop the specified object from the Q Apply server, and if needed, recreate.

ASN2078E The object *object_name* of type *object_type* does not exist at the server *apply_server* for the Q Apply schema *apply_schema*.

Explanation: The action was invoked on a Q subscription whose catalog definitions were not found in the Q Apply control tables.

User response: Drop the specified object from the Q Capture server, and if needed, recreate.

ASN2081E The table *table_name* cannot be published. Reason code: *reason_code*.

Explanation: The table is not supported for publication. No script is generated. The following values are valid for the reason code:

- 0 The provided source table is a DB2 catalog table. Publishing of changes from system catalog tables is allowed only if the Q Capture program runs on z/OS.
- 1 The Q Capture server is on a z/OS platform and the source table is defined with an EDITPROC option, but DB2 for z/OS is not at the correct level.
- 2 The Q Capture server is on a z/OS platform and the source table is defined with a Valid proc option.
- 3 The source table has one or more LOB columns but there is no uniqueness defined.
- 4 No columns in the source table were selected for publication.
- 5 No key columns could be derived for the published columns at the source table.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 0 Provide a source table other than a DB2 catalog table and rerun the action.
- 1 DB2 for z/OS source tables that are defined with an edit routine (EDITPROC) to provide additional data security are supported. To use these tables as sources, the DB2 subsystem that contains the tables must be at Version 8 with APAR PK13542.
- 2 Provide a different source table that is not defined with the Valid proc option and rerun the action.
- 3 Create a unique database property on any one of the source columns that is subscribed other than a LOB column, and then rerun the action.
- 4 Select at least one column in the source table and rerun the action.
- 5 Specify the key columns and rerun the action.

ASN2082E The column *tablename-columnname* of data type *datatype* cannot be included in the publication. Reason code: *reason_code*.

Explanation: The provided column is not supported by publications. The following values are valid for the reason code:

- 0 The data type of the column is not supported for publication.
- 1 The maximum number of published LOB columns was exceeded for this table.
- 2

Columns of LOB data type cannot be part of the key definition of a table.

User response: Review the reason codes, provide a different column name in the publication, and rerun the action.

ASN2083E The Q subscription *name* cannot be created because it already exists under the specified Q Capture or Q Apply schema.

Explanation: See message text.

User response: Provide another Q subscription name that does not already exist at the server and rerun the action.

ASN2084E The object *objectowner.objectname* of type *object_type* exists at the specified server and under the specified Q Apply schema.

Explanation: A catalog definition already exists in the control tables for the provided object name and Q Apply instance. Duplicate values for this type are not allowed in the same Q Apply instance.

User response: Provide a different name for the object and rerun the action.

ASN2086E The stored procedure *storedproc_owner.storedproc_name* cannot be used by the associated Q subscription. Reason code: *reason_code*.

Explanation: The stored procedure is not valid. No script is generated. The following values are valid for reason code:

- 0 The stored procedure has no parameters.
- 1 The first parameter in the stored procedure must be 'OPERATION'.
- 2 The second parameter in the stored procedure must be 'SUPPRESSION_IND'.
- 3 The third parameter in the stored procedure must be 'SRC_COMMIT_LSN'.
- 4 The fourth parameter in the stored procedure must be 'SRC_TRANS_TIME'.
- 5 The 'INOUT' attribute value of the OPERATION parameter is not equal to 'INOUT'.
- 6 The 'INOUT' attribute value of the 'SUPPRESSION_IND', 'SRC_COMMIT_LSN', or 'SRC_TRANS_TIME' parameter is not equal to 'IN'.
- 7 The stored procedure must contain at least five parameters: four required parameters ('OPERATION', 'SUPPRESSION_IND',

'SRC_COMMIT_LSN', 'SRC_TRANS_TIME')
and at least one additional parameter.

User response: Review the reason codes in the explanation, provide a valid stored procedure and rerun the action.

ASN2087E The stored procedure parameter *owner.name.parameter* cannot be included in the Q subscription. Reason code: *reason_code*.

Explanation: The stored procedure parameter is not supported by the Q subscription. No script is generated. The following values are valid for the reason code:

- 0 The parameter's 'INOUT' attribute value is not equal to 'IN'.
- 1 The parameter is a special prerequisite parameter required by the Q Apply program and this parameter should not take part in the Q subscription.
- 2 The parameter is a before-value parameter, and before-value parameters cannot be mapped to a source column in the Q subscription.
- 3 The parameter is part of the key but the corresponding before-value parameter does not appear in front of this parameter in the stored procedure definition.

User response: Review the reason codes in the explanation. Either modify the parameter definition of the stored procedure or provide a stored procedure with supported parameters and rerun the action.

ASN2088E The column or stored procedure parameter *targetowner.targetname.name* cannot be added to the Q subscription. Reason code: *reason_code*.

Explanation: The column or the stored procedure parameter is not supported or not valid. No script is generated. The following values are valid for the reason code:

- 0 The target is a stored procedure and the parameter should exist in the procedure definition before adding the parameter to the Q subscription.
- 1 The column is already participating in the Q subscription.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 0 Modify the definition of the stored procedure to include this parameter in the stored procedure and rerun the action.
- 1 Verify the column name and provide a

different column name not already
participating in the Q subscription.

ASN2100E The length *length* of the string *string* for parameter *parameter* in the database object *object_name* of type *object_type* at the server *server_alias* is longer than the maximum limit *max_limit*.

Explanation: This type of database object requires a shorter string than the one that was supplied for the parameter.

User response: Refer to the SQL reference for the appropriate database to find the valid length of the string, and provide the correct name length.

ASN2101E The database object *object_name* of type *object_type* already exists at server *server_alias*.

Explanation: The database object cannot be created because there is already a database object of the same type with the same name.

User response: Provide a name for that object that does not already exist in the database and reissue the task.

ASN2102E The database object *object_name* of type *object_type* does not exist at the server *server_alias*.

Explanation: The database object does not exist in the database catalog. There is no information to retrieve for this object, and it cannot participate in the publication or Q subscription.

User response: Ensure that you specify the correct name of an existing database object and run the task again.

ASN2104E The row size *row_size* of all columns in the database object *object_name* of type *object_type* at server *server* exceeds the allowed maximum limit *max_limit*.

Explanation: The combined row size of all the columns in the specified object cannot exceed the limit. This object cannot be created or altered.

User response: Either assign a table space with a larger page size for this object, or specify fewer columns or shorter columns for this object. Refer to the database specific SQL reference for details on the row size limits of database objects.

ASN2105E The number of columns *num_columns* in the database object *object_name* of type *object_type* at server *server_alias* exceeds the database limit *max_limit*.

Explanation: The number of columns that a database object (table or index) can contain depends on the database platform, but it cannot exceed a predefined maximum number. No script is generated to create or alter the object.

User response: Redefine the database object with the appropriate number of columns. Refer to the SQL Reference for the appropriate database to find the valid number of columns for the database object.

ASN2106E The page size *page_size1* for the table space *tablespace_name* at server *server_alias* does not match the page size *page_size2* of the table space *bufferpool* *buffer_pool*. Reason code: *reason_code*.

Explanation: The page size for the table space must match the page size of the bufferpool that is assigned to this table space. The following values are valid for the reason code:

- 0 The bufferpool specified was assigned to the table space and the page size obtained from the bufferpool is not equal to the page size specified for the table.
- 1 Bufferpool information is not specified for the table space and so the default bufferpool is selected for this table space. But the page size of the default bufferpool does not match the page size of the table space.

User response: Check the input and take the following actions:

- If the bufferpool information is specified in the table space, modify the page size of the table space to match the page size of the bufferpool and rerun the action.
- If the default bufferpool was used for the table space, modify the page size of the table space to match the page size of the default bufferpool and rerun the action.

ASN2118E The database object *object_name* of type *object_type1* cannot be deleted because dependent objects of type *object_type2* still exist at server *server_alias*.

Explanation: The database object cannot be deleted because one or more objects exist at the specified server that are dependent on the existence of this object.

User response: Check the provided object name and issue the task again.

ASN2119E The parameter *parm_name* is missing for the database object *object_name* of type *object_type* at server *server_alias*.

Explanation: The parameter value is required for the specified object.

User response: Check the input and provide a valid value for the required parameter.

ASN2120E The parameter *parm_name* is not supported for the database object *object_name* of type *object_type* at server *server_alias*.

Explanation: See message text.

User response: Check the input and ensure that no value is specified for the parameter. Refer to the database specific SQL reference for further details.

ASN2121E The parameter *parm_name* cannot be used in an update operation for the database object *object_name* of type *object_type* at server *server_alias*.

Explanation: You cannot update the value of this parameter.

User response: Check the input and ensure that no value is specified for the parameter.

ASN2122E The value *value* for the parameter *parm_name* of the database object *object_name* of type *object_type* at server *server_alias* contains the invalid character *char*.

Explanation: The values for the parameter can only contain certain characters.

User response: Remove the invalid character from the provided value. Refer to the database specific SQL reference to find out which characters are valid for the provided parameter.

ASN2123E The value *value1* for the parameter *parm1* of the database object *object_name* of type *object_type* conflicts with the value *value2* for parameter *parm2*.

Explanation: The values provided for the specified parameters cannot coexist at the server.

User response: Modify the input to reflect a valid value combination. Refer to the database specific SQL reference to find out which value combinations are valid.

ASN2124E The value *value* for the parameter *parm* of the database object *object_name* of type *object_type* at server *server_alias* is higher than the maximum value allowed *max_value*.

Explanation: The provided value is too high. Only values lower than or equal to the specified maximum are supported.

User response: Check the input and provide a lower

number. Refer to the database specific SQL reference for the valid range of values.

ASN2125E The value *value* for the parameter *parm* of the database object *object_name* of type *object_type* at server *server_alias* is lower than the minimum value allowed *min_value*.

Explanation: The provided value is too small. Only values higher than or equal to the specified minimum value are supported.

User response: Check the input and provide a higher number. Refer to the database specific SQL reference for the valid range of values.

ASN2126E The database object *object_name1* of type *object_type1* at server *server_alias* contains duplicate objects of type *object_type2* of the name *object_name2*.

Explanation: The duplicate objects are not allowed.

User response: Check the input and make sure each object is provided only once.

ASN2127E The database object *object_name* of type *object_type1* at server *server_alias* contains *num_objects* objects of type *object_type2* but the maximum limit allowed is *max_limit*.

Explanation: Too many objects were specified for the provided parameter.

User response: Check the input and reduce the number of these database objects to the maximum.

ASN2128E The column *col_name* of data type *datatype* cannot be used in the database object *object_name* of type *object_type* at the server *server_alias*. Reason code: *reason_code*.

Explanation: The column is not supported in the specified database object. The following values are valid for the reason code:

- 0 The data type of the column is not valid.
- 1 The column cannot be used in the database object.
- 2 The column length is too long.
- 3 The data type of the referencing column does not match the data type of the referenced column in a foreign key constraint.

User response: Review the reason codes listed in the explanation. Choose another column or choose another database object for this column.

ASN2129E A conflict exists between property *prop_name* of type *prop_type* and another property in the database object *object_name* of type *object_type* at server *server_alias*.

Explanation: The object was not created or updated because conflicting properties were found and they cannot coexist. The following values are valid for the reason code:

- 0 You are trying to create the object and you specified a property that conflicts with another property specified in the same session.
- 1 You are trying to update the object and you specified a property that conflicts with another property specified in the same session.
- 2 You are trying to update the object and you specified a property that conflicts with an existing property in the object.

User response: If both the properties are new and are specified, then specify only one of the properties and issue the task. If one of the properties already exists in the object, then do not specify the new property in the object.

ASN2130E The database object *object_name* of type *object_type* cannot be created or updated because column *column_name* does not exist in the table *table_name* at server *server_alias*.

Explanation: See message text.

User response: Check the input and provide another name for the column that exists in the table.

ASN2131E The column *column_name* in the target table *table_name* cannot be used for Q Replication because the column is read-only.

Explanation: See message text.

User response: Choose another target column. Refer to the database specific SQL reference to determine which columns cannot be updated.

ASN2132E The Column *column_name* in the target table *table_name* is not mapped to a source column for replication. Because this column is not nullable and it does not have a default value, the Q subscription cannot be created or updated.

Explanation: Any columns in a target table that do not participate in the Q subscription must be either nullable or not null with default. At least one column in the target table that you did not select for the Q subscription does not meet the criteria.

User response: Take one of the following actions:

- Map the specified column in the target table to a column in the source table so that the column participates in the Q subscription.
- Enable the nullability property for the target column.
- Specify a default value for this column.

ASN2135E The data type *target_datatype* of the target column *target_colname* is not compatible with the data type *source_datatype* of the source column *source_colname*.

Explanation: The values of the source column cannot be applied to the target column because the data types of the columns are not compatible.

User response: Take one of the following actions:

- If the target exists and is a user table, change the column mapping to a different column that has a compatible data type.
- If the target exists and is a nickname, alter the data type of the nickname column to match the source column data type. See "Nickname data types required for federated Q Replication" in the DB2 Information Center for details.
- If the target table does not exist, modify the data type of the target column so that it is compatible with the source column.

ASN2136W The attribute of the target column *target_column* and the attribute of the source column *source_column* do not match. Reason code: *reason_code*.

Explanation: A difference between the values of the source and target column for the same attribute was detected. A script for creating or updating the Q subscription is generated. The following values are valid for the reason code:

- | | |
|---|--|
| 0 | The source column is nullable whereas the target column is not nullable. |
| 1 | The source column is not nullable whereas the target column is nullable. |
| 2 | The encoding schemes for the source column and the target column are different. |
| 3 | The source column has default values and the target column does not have default values. |
| 4 | The source column does not have default values and the target column has default values. |

User response: Change the column mapping or column properties if the difference can cause potential problems. If the difference is acceptable, no action is necessary.

ASN2137W The compatibility between the source column *source_column* and the target column *target_column* exists only for certain values of the source column. Reason code: *reason_code*.

Explanation: The values of the source column cannot be always applied to the target column. A script to create or update the Q subscription is still generated. The following values are valid for the reason code:

- | | |
|---|--|
| 0 | The source column can contain null values, but the target column does not allow null values. Null values in the source column cannot be applied to the target column. |
| 1 | The source column data type is either a CHAR or a VARCHAR, whereas the target column data type is a DATE, TIME, or a TIMESTAMP. Any string values in the source column that are not a valid representation of a time, date, or timestamp cannot be applied to the target column. |
| 2 | The source column contains numbers that are outside of the range of values that the target column supports. |
| 3 | The length of the VARCHAR data type of the source column is longer than the length of the CHAR or VARCHAR data type of the target column. |

User response: Review the reason codes in the explanation, and take the appropriate action:

- If the target table exists, change the column mapping.
- If the target table does not exist, change the column properties.
- Check the source and the target column. If the source column will only contain values that are allowed for the target column, no action is required.

ASN2138W The values of the source column *source_column* of data type *source_datatype* will potentially lose fractional digits when replicated to the target column *target_column* of data type *target_datatype*.

Explanation: There can be potential loss of data between the source and the target columns because the target column data type definition indicates that the precision allowed in the column is much lower than that of the source column.

User response: Take one of the following actions:

- If the target table exists, change the column mapping.
- If the target table does not exist, change the column properties.

- Check the source and the target column. If the values will never exist in the source column or if the truncation is acceptable in the target column, no action is necessary.

ASN2139E The value *value* for the parameter *parm* in the database object *object_name* of type *object_type* at server *server_alias* is not valid because it is not equal to one of the following values: *valid_values*.

Explanation: See message text.

User response: Check the input and provide a value from the list of values specified.

ASN2140E The value *value* for the parameter *parm1* in the database object *object_name* of type *object_type* at server *server_alias* is not a multiple of *factor*.

Explanation: The values for the specified parameter in the database object must be multiples of the provided factor.

User response: Check the input and provide a value that is a multiple of the specified factor.

ASN2141E The value *value* for the parameter *parm* of the database object *object_name* of type *object_type* at server *server_alias* is not valid.

Explanation: The value for the specified parameter is not valid on the database platform.

User response: Check the input and provide another value. Refer to the database specific SQL reference for valid values for this parameter.

ASN2142E The attribute *parm* is not valid for LOB table space *tablespace_name* on server *server_alias*.

Explanation: The following attributes are not valid for a LOB table space: PCTFREE, TRACKMOD, COMPRESS, SEGSIZE.

User response: Remove the attributes that are not valid for LOB table space (PCTFREE, TRACKMOD, COMPRESS, SEGSIZE), and rerun the action.

ASN2144E The table *table_name* at the server *server_alias* was not created or updated because the definition of the foreign key *foreign_key* is not valid. Reason code: *reason_code*.

Explanation: No script is generated. The following values are valid for the reason code:

- 0 The number of referencing columns is lower than the referenced columns.

- 1 The number of referencing columns is higher than the number of referenced columns.

User response: Check the input and ensure that the number of referenced columns matches the number of referencing columns.

ASN2145E The column name *column_name* for the table *table_name* at server *server_alias* was specified more than once in the same session.

Explanation: See message text.

User response: Check the input and remove the duplicate name.

ASN2146E The publication or Q subscription cannot be created or updated because required information is missing. Reason code: *reason_code*.

Explanation: No script is generated. The following values are valid for the reason code:

- 0 The source table information is missing.
- 1 The target table information is missing.
- 2 No source columns were selected.
- 3 No target columns were selected.

User response: Check the input and provide values for the specified parameters.

ASN2147E The publication or Q subscription cannot be updated or created because column *column_name* does not exist in table *targetowner-targetname* at server *server_alias*. Reason code: *reason_code*.

Explanation: No script is generated. The following values are valid for the reason code:

- 0 The column does not exist in the source table.
- 1 An existing target table was provided but the specified column does not exist in the target table.

User response: Check the input and select columns that exist in the source or target table.

ASN2148E The publication or Q subscription cannot be created or updated because column *column_name* in table *table_name* at server *server_alias* cannot be a replication key column. Reason code: *reason_code*.

Explanation: No script is generated. The following values are valid for the reason code:

0

The column that is selected for the key column is not part of the list of selected source columns.

1

The column that is selected for the key column is not part of the list of selected target columns.

User response: Check the input and choose key columns that are selected as the source or the target columns.

ASN2149E The target *table_name* of type *target_type* is read-only.

Explanation: The specified database object cannot be a target because it cannot be updated.

User response: Choose another database object as a replication target. Refer to the database specific SQL reference for information about which database objects can be updated.

ASN2150E The Q subscription cannot be created because the replication target key is missing. Reason code: *reason_code*.

Explanation: No script is generated. The following values are valid for the reason code:

0 User input is missing for the target key.

1 The target table does not exist and the source table does not contain any unique database properties for the selected source columns.

2 Both source and target tables exist, but they do not contain any unique database properties on the selected target and source columns.

User response: Check the input and make sure the appropriate columns are selected as part of the target key.

ASN2151E The source column *source_column* does not have a mapped target column.

Explanation: The action create or update Q Subscription cannot be completed because the specified source column does not have a target column to be mapped.

User response: Check the source column and verify that it is mapped to a valid target column and rerun the action.

ASN2152E The target column *target_column* does not have a mapped source column.

Explanation: The action create or update Q subscription cannot be completed because the specified target column does not have a source column to be mapped.

User response: Check the target column and verify that it is mapped to a valid source column and rerun the action.

ASN2153E The option *option* that was provided is not supported for the platform *platform* on server *server*.

Explanation: The ASNCLP option is not supported for the server platform.

User response: Check the syntax and provide the correct options.

ASN2154E The specified environmental value is not relevant in this particular context. The token *token* will be ignored.

Explanation: See message text.

User response: Check the syntax and provide the correct values.

ASN2155E The publication or replication queue map *queue_map_name* cannot be deleted because it is used by one or more publications or Q subscriptions, respectively. Reason code: *reason_code*.

Explanation: The following values are valid for the reason code:

0 You are trying to delete a publishing queue map but there exist one or more publications that depend on this queue map.

1 You are trying to delete a replication queue map but there exist one or more Q subscriptions that depend on this queue map.

User response: Review the reason codes in the explanation, and take the appropriate action:

0 If you want to delete the publications that use the publishing queue map, delete them first and then delete the publishing queue map.

1 If you want to delete the Q subscriptions that use the replication queue map, delete them first and then delete the replication queue map.

ASN2156W The uniqueness of the replication key values is not enforced at the source table, at the target table, or at both.
Reason code: *reason_code*.

Explanation: The following values are valid for the reason code:

- 0 The uniqueness of the replication key is enforced at the target table, but not at the source table.
- 1 The uniqueness of the replication key is enforced at the source table, but not at the target table.
- 2 The uniqueness of the replication key values is not enforced at both the source and target tables.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 0 Make sure that the application on the source table only generates unique values for the source columns that are mapped to the replication target key columns.
 - 1 Make sure that the application on the target table only generates unique values for the target columns that were selected as replication target key columns.
 - 2 Take the actions that are described for reason codes 0 and 1.
-

ASN2157I No script was generated because no action was required.

Explanation: The action is not required because the objects in the database already reflect the required state.

User response: This message is for your information only. No action is required.

ASN2158E No objects of type *object_type* were found for the provided query *query*. The command will be ignored.

Explanation: In the command line interface, you specified a query to retrieve multiple objects that all participate in an action (such as Create Publication or Create Q Subscription). The query did not return any objects. No script is generated for the create, update, or delete action.

User response: Redefine the query so that it can return at least one object of the specified type and reissue the command.

ASN2159E The value *value* provided for the option *option* cannot be used in conjunction with *token*.

Explanation: In the ASNCLP program, some parameters are not supported in certain scenarios. For example, if the subscription specifies to create a new target table, then the TYPE value cannot be STOREDPROC because there is no support for creating stored procedures.

No script is generated.

User response: Refer to the ASNCLP documentation for a list of valid combinations for the keywords.

ASN2160E The object *object* of type *type* cannot be created at server *server* because the logical database was not specified.

Explanation: The logical database must be specified in order to create the specified object type on a z/OS platform.

User response: Specify a value for the logical database and rerun the action.

ASN2161E The table *table_name* at server *server* cannot be created or updated because the column *column_name* is not valid.
Reason code: *reason_code*.

Explanation: The following values are valid for the reason code:

- 0 The GRAPHIC data type is not supported for ASCII or EBCDIC table spaces.
- 1 The column encoding scheme "FOR MIXED DATA" is not supported for ASCII or EBCDIC table spaces.
- 2 The GRAPHIC data type is not supported at the specified server.

User response: Specify a UNICODE table space for this table and rerun the action.

ASN2162W The uniqueness of the columns of the unique constraint or index *name* in the target table *table_name* at server *server* is not enforced in the source table *table_name* at server *server*.

Explanation: If uniqueness for one or more of the replicated columns is enforced at the target table but not at the source table, then the source column value combinations that violate this uniqueness cannot be replicated to the target. Data in the source table cannot be applied to the target table.

User response: If all unique database constraints or indexes at the target table cannot be violated by column value combinations from the source table, no

action is necessary. Otherwise, do not run the script that was generated. Instead, reconsider dropping the unique constraint or index at the target table and run the action again.

ASN2163E A script to create or update the Q subscription was not generated. The source table is *table_name*. The target table is *table_name*. Reason code: *reason_code*.

Explanation: The following values are valid for the reason code:

- 0 The target table does not exist and the source table does not have a unique key to be used for new target table.
- 1 The reason differs depending on the target type:

Target tables

The target table exists and neither the target table nor the source table has a unique index or primary key.

Stored procedure targets

The stored procedure does not have the correct number of parameters for the key columns. Two stored procedure parameters are required for each key column at the source: one for the before value of the key and one for the after value. The parameter for the before value must come before the parameter for the after value. The before-value parameter must begin with an X.

- 2 An automatic index could not be derived for the target table.

User response: Review the reason codes in the explanation. Modify the source table definition, the target table or stored procedure definition, or both. Reissue the action.

ASN2164E The column *tablename.columnname* of data type *datatype* cannot be included in the Q subscription. Reason code: *reason_code*.

Explanation: The following values are valid for the reason code:

- 0 The data type of the column is not supported for Q subscription.
- 1 The maximum number of subscribed LOB columns was exceeded for this table.
- 2 The target is a federated server and write access to a LOB data type for the federated data source is not permitted.

User response: Review the reason codes, provide a

different column name in the Q subscription, and rerun the action.

ASN2165E The properties of the queue map *queue_map_name* cannot be updated because its associated *source|target* database or subsystem *name* was not added to the Replication Center object tree.

Explanation: Before you can change the properties of a publishing or replication queue map, you must add both the source and target servers that are specified for the queue map to the Replication Center.

User response: Use the Add Q Capture Server or Add Q Apply Server wizards to add the missing server to the Q Capture Servers or Q Apply Servers folder in the object tree.

ASN2200E The Q Apply schema *apply_schema* or the Q Capture schema *capture_schema* for server *server* could not be used for multidirectional replication.

Explanation: On each server that participates in multidirectional replication, the Q Apply and Q Capture control tables must have the same schema.

User response: Make sure that the Q Apply control tables and the Q Capture control tables have the same schema name and that they exist on the same server that you are using for multidirectional replication.

ASN2201E In the IBMQREP_SUBS table for Q Capture schema *schema_name*, the values in the SOURCE_NODE column are not equal to each other.

Explanation: There are one or more Q subscriptions in the Q Capture control tables for the given schema that contain different values for the SOURCE_NODE field. In multidirectional replication, the rows in the SOURCE_NODE column must contain the same value.

The report log generated by ASNCLP contains details about the Q subscription that failed.

User response: Verify that all the SOURCE_NODE values for the Q subscriptions are the same in the IBMQREP_SUBS table.

ASN2202E One or more Q subscriptions (SUBTYPE *subtype*) cannot be created because one or more reported errors were not corrected.

Explanation: The Q subscriptions cannot be created for multidirectional replication because they are not valid. Refer to the error message given for each Q subscription.

User response: Fix the invalid Q subscriptions and try creating them again.

ASN2203E The Q subscription name *name* is not unique for the Q Capture schema *capture_schema* on server *server* for multidirectional replication.

Explanation: The same name must not be provided for multiple Q subscriptions when setting up a multidirectional replication configuration. For each instance of the Q Capture program, all Q subscription names must be unique. Because the name of the Q subscription is stored at both the source and target server, be sure that the name is compatible with the code pages for both the source and target servers.

User response: Make sure that each Q subscription name is unique.

ASN2204E The table *table_name* is on server *server*, but the Q Apply and Q Capture schemas *schema* were not found on that server.

Explanation: The Q Capture and Q Apply control tables with the given schema were not found on the same server as the table that you are replicating. In multidirectional replication, the Q Capture and Q Apply control tables must be on the same server as the replicated table and they must use the same schema name.

User response: Ensure that the table and the Q Apply and Q Capture control tables that it uses for multidirectional replication are on the same server.

ASN2205W Q subscription *name* for Q Apply schema *apply_schema* on Q Apply server *server* is not valid and was not created as part of the mass operation.

Explanation: When creating multiple Q subscriptions at one time for unidirectional replication, if one or more of the Q subscriptions are not valid they are not created; however, the rest of the Q subscriptions are created if they are valid.

User response: Correct the Q subscriptions that were not valid and then create them.

ASN2206E The connection information provided for Q Apply schema and Q Capture schema *source_schema* on server *server* to Q Apply schema and Q Capture schema *target_schema* on server *server* points to itself.

Explanation: In multidirectional replication, the Q Apply schema and the Q Capture schema at the Q Apply server (source) cannot also be used as the Q Capture server (target).

User response: Make sure that the schema name and server for the Q Apply and Q Capture tables at the Q Apply server are different from the schema name and server for the Q Apply and Q Capture tables at the Q Capture server.

ASN2207E The replication queue map *queue_map_name* is not valid for setting up a connection from the source Q Apply and Q Capture schema *source_schema* on the Q Capture server *server* to the target Q Apply and Q Capture schema *target_schema* on the Q Apply server *server*.

Explanation: This replication queue map does not match the specified schemas that are used for multidirectional replication.

User response: Ensure that the replication queue map maps from the specified Q Capture control tables to the specified Q Apply control tables.

ASN2208E Connection information was not provided for the SUBGROUP *subgroup*.

Explanation: The Q subscription group must have all required information defined before it can proceed with any CREATE, ALTER, or DELETE operations.

User response: Ensure that all of the replication queue maps are provided between the source and target tables.

ASN2209E The connection information exists for SUBGROUP *subgroup* and it cannot be overwritten.

Explanation: Once you define connection information for a Q subscription group, you cannot change it.

User response: Do not give connection information for a Q subscription group more than once.

ASN2210E A Q subscription with the SUBGROUP name *subgroup* does not exist on the Q Capture server *server* using Q Apply and Q Capture schema *schema*.

Explanation: The Q subscription group information could not be found.

User response: Make sure that there is at least one Q subscription on the Q Capture schema for the SUBGROUP name.

ASN2211E No Q subscriptions exist for the reference table *table_name* for the SUBGROUP *subgroup* using Q Apply and Q Capture schema *schema* on server *server*.

Explanation: See message text.

User response: Make sure that there is a Q subscription of SUBTYPE 'P' (peer-to-peer) or 'B' (bidirectional) on the Q Capture schema that uses the given reference table as a source.

ASN2212E **An internal inconsistency was detected.**
Reason code: *reason_code*

Explanation: The Q subscription group must be synchronized with the persistent information on the database before being used. The following values are valid for the reason code:

- 0 The Q subscription group already exists but it was not synchronized, or it is not in a synchronized state.
- 1 Connection information between the nodes was missing.

User response: Contact IBM Software Support. Provide an Analyzer report.

ASN2213E **Replication queue map *queue_map_name* was specified with duplicate source and target information: Q Apply schema *apply_schema* on Q Apply server *apply_server*, and Q Capture schema *capture_schema* on Q Capture server *capture_server*.**

Explanation: A replication queue map with duplicate source and target information was specified. Replication queue maps must be unique between the connections they form. For example, the first connection is arbitrary. The second connection must share one end with either end of the first connection. The third connection must share one end with one of the three points formed by the previous two connections.

User response: Make sure that the replication queue map specified is unique to the Q Apply server and the Q Capture server.

ASN2214E **Replication queue map *queue_map_name* between Q Apply schema *apply_schema* on Q Apply server *apply_server* and Q Capture schema *capture_schema* on Q Capture server *capture_server* does not connect to any of the replication queue maps previously specified in the connection information.**

Explanation: The replication queue maps must be connected to each other. Each replication queue map must share a Q Capture or Q Apply schema with any of the previously specified replication queue maps.

User response: Specify the replication queue map such that it shares a Q Capture or Q Apply schema

with any of the previously specified replication queue maps.

ASN2215E **The maximum number of Q Apply and Q Capture schema pairs allowed for a peer-to-peer Q subscription group was exceeded. The maximum number of pairs allowed is *max_limit*.**

Explanation: See message text.

User response: Reduce the number for Q Apply and Q Capture schema pairs used in this configuration for peer-to-peer replication.

ASN2216E **Q subscriptions of SUBTYPE *subtype* can be assigned only to a Q subscription group with other Q subscriptions of the same SUBTYPE.**

Explanation: Q subscriptions of SUBTYPE 'P' (peer-to-peer) cannot be added to a bidirectional replication configuration, and Q subscriptions of SUBTYPE 'B' (bidirectional) cannot be added to a peer-to-peer replication configuration.

User response: When adding Q subscriptions to a SUBGROUP, ensure that the Q subscriptions have the same SUBTYPE as others in the SUBGROUP.

ASN2217E **The pair of Q Apply and Q Capture schemas from which each Q subscription originates was not specified. The attributes for the two Q subscriptions cannot be updated for bidirectional replication.**

Explanation: When setting up bidirectional replication, you can specify different attributes for each of the two Q subscriptions. To do so, you must specify from which Q Apply and Q Capture schema each Q subscription originates.

User response: To assign attributes to a Q subscription when setting up bidirectional replication, specify the Q Apply and Q Capture schema pair from which the Q subscription originates.

ASN2218E **An internal inconsistency was detected.**
Reason code: *reason_code*.

Explanation: The Q subscriptions are not synchronized with the persistent information on the database to perform an UPDATE or DELETE operation. The following value is valid for the reason code:

- 0 The SUBGROUP already exists but was not synchronized, or it is not in a synchronized state.

User response: Contact IBM Software Support. Provide an Analyzer report.

ASN2219E The number of Q Apply and Q Capture schema pairs *schema* has exceeded the number allowed for SUBGROUP *subgroup* with Q subscriptions of SUBTYPE *subtype*.

Explanation: The Q subscription group configuration limits the number of Q Apply and Q Capture schema pairs. Bidirectional replication allows only two Q Apply and Q Capture schema pairs per SUBGROUP.

User response: Make sure that the number of Q Apply and Q Capture schema pairs is less than or equal to the maximum allowed for that SUBTYPE.

ASN2220E There are not enough Q Apply and Q Capture schema pairs *schema* for SUBGROUP *subgroup* with Q subscriptions of SUBTYPE *subtype*. The minimum number of pairs required is *number*.

Explanation: See message text.

User response: Make sure that the number of Q Apply and Q Capture schema pairs is equal to the minimum allowed for that type of Q subscription.

ASN2221E The connection information is not valid for this SUBGROUP *subgroup* for SUBTYPE *subtype*.

Explanation: Bidirectional and peer-to-peer replication require that all servers are connected to each other.

User response: Validate the connection information provided for this SUBGROUP and SUBTYPE and fix any problems that prevent all servers from connecting to each other. Then recreate the Q subscriptions.

ASN2222E The number *number* of attribute type *object_type* on table *table_name* on server *server* does not match the number *number* of object type *object_type* on table *table_name* on server *server*.

Explanation: The number of columns, unique indexes, primary keys, foreign keys, unique keys, and check constraints must be the same for the source and the target tables that participate in multidirectional replication configurations.

User response: Either change the definition of the specified attribute in the source or the target table so that the number of attributes match, or choose a different combination of source and target tables.

ASN2224E The object *object* of type *object_type* on table *table_name* on server *server* does not match an object on table *table_name* on server *server*.

Explanation: The attribute must be the same for the

source and the target tables that participate in multidirectional replication configurations.

User response: Either change the definition of the specified attribute in the source or the target table so that it matches the attribute on the other table, or choose a different combination of source and target tables.

ASN2225E ASNCLP: More than one SET SUBGROUP statement was found in the script for creating, updating, or deleting a SUBGROUP.

Explanation: Only one SUBGROUP can be specified per script. The ASNCLP program cannot work with more than one subgroup in a session. The SUBGROUP was not created, updated, or deleted.

User response: Ensure that you specify only one SET SUBGROUP statement per script and run the script again.

ASN2226E ASNCLP: The Q subscription SUBTYPE *subtype* does not match the SUBTYPE of other Q subscriptions in a Q subscription group.

Explanation: The information specified in a create subscription command was inconsistent with other create subscription commands belonging to this multidirectional replication configuration. Different types of Q subscriptions were detected for the Q subscription group. Only one SUBGROUP can be specified, and the SUBGROUP must have Q subscriptions of one SUBTYPE. No Q subscriptions were created.

User response: Ensure that you specify either SUBTYPE 'P' or 'B' under a specific SUBGROUP. Do not mix types of Q subscriptions in one Q subscription group. Rerun the create subscription command.

ASN2227E ASNCLP: The number of set connection statements is not sufficient for the number of set schema statements.

Explanation: In a multidirectional replication configuration, for every n nodes you need $n*(n-1)$ connections. For example, if there are 3 nodes, you need $3*(3-1)=6$ connections.

User response: Ensure that the set connection statement and the set schema statement conform to the rules.

ASN2228E ASNCLP: Q subscriptions for the SUBGROUP *subgroup* already exist on the specified servers.

Explanation: While trying to create more Q subscriptions, you provided SUBGROUP information

instead of providing referential information to that SUBGROUP.

User response: Provide a reference to the SUBGROUP rather than defining the SUBGROUP again.

ASN2229I The action *action_name* started at *time_stamp*. The Q Apply server is *apply_server*, the Q Apply schema is *apply_schema*, and the remote server name is *remote_servername*.

Explanation: The action started successfully at the Q Apply server.

User response: This message is for your information. No action is required.

ASN2230I The action of creating a Q subscription started at *time_stamp*. Q subscription name: *name*. Q Capture server: *capture_server*. Q Capture schema: *capture_schema*. Q Apply Server: *apply_server*. Remote server name: *remote_servername*. Q Apply Schema: *apply_schema*. The source table is *table_name*. The target table or stored procedure is *table_name*.

Explanation: The Q subscription is starting to be created.

User response: This message is for your information. No action is required.

ASN2231E The database object *object_name* of type *object_type* does not exist at the data source *datasource* that was defined through the remote server *remote_server*.

Explanation: The database object does not exist in the database catalog at the data source that you specified. Because no information can be retrieved for this object, this object cannot participate in the Q subscription.

User response: Ensure that you specify the correct name of an existing database object and run the task again.

ASN2232E The database object *object_name* of type *object_type* already exists at the data source *datasource*, which is defined through the remote server *remote_server*.

Explanation: The database object cannot be created because another object already exists that has the same name at the specified data source.

User response: Provide a name for that object that does not already exist in the database, and run the task again.

ASN2233E The data source *datasource_name* that was defined by using the remote server *remote_servername* is not supported as a Q Replication target.

Explanation: Only a subset of the data source platforms are supported as targets in Q Replication.

User response: See "Supported sources and targets" in the DB2 Information Center. Provide a remote server name that is defined on a supported target and run the task again.

ASN2234W The length *length1* of the source column *source_colname* is greater than the length *length2* of the target column *target_colname*. The source data applied to the target might get truncated.

Explanation: If the actual value of the source column cannot be stored in the target column completely, the data that is applied to the target might get truncated. As long as the actual values from the source are not larger as the target column definition, no data is modified.

User response: No action is needed if truncation is acceptable or if the source column values are always small enough to be applied to the target column. If truncation is not acceptable, then do not select this column to be part of the Q subscription.

ASN2235E The name of the target nickname *owner.name* conflicts with another database object of type *object_name*.

Explanation: For federated targets, the replication administration tools always create a new target nickname, which cannot have the same name as a table, view, or other nickname that exists on that system. The name that you provided already exists.

User response: Provide a name for the nickname that is not the same name as a table, view, or other nickname that already exists on that system.

ASN2236W The target table *schema.tablename* of the target nickname *schema.nickname* has a foreign key defined on it. The nickname does not have the foreign key defined on it but needs to.

Explanation: Any referential constraints that are defined on the source table are not automatically carried over to the nickname. You need to issue an alter statement to carry the constraints over to the nickname so that the information is in the DB2 catalog tables.

If both the parent and the child tables are participating in the Q subscription, the Q Apply program relies on the information that is stored in the DB2 catalog tables about the referential integrity constraints in order to apply changes to the parent and the child target tables

in the correct sequence. If the referential constraint information is not carried over to the nicknames, the Q Apply program might encounter an error while loading the nicknames if the data violates the referential integrity rules.

User response: Issue the ALTER NICKNAME statement if the parent data source tables are also participating in the Q subscriptions. See "Altering a nickname (DB2 command line)" in the DB2 Information Center for the correct syntax to alter a nickname.

ASN2237E The table *object* cannot be created at server *server* because you did not specify the table space.

Explanation: The replication administration tools require for all z/OS tables to be created using an explicit table space assignment.

User response: Assign a table space to the z/OS table.

ASN2238W Transaction mode processing is not supported when you are replicating data from a non-DB2 relational server. In this situation, NULL is the only valid value for the COMMIT_COUNT attribute of the subscription set. The COMMIT_COUNT *commitcount_value* that you specified will be ignored.

Explanation: When the source of a subscription set is a non-DB2 relational server, the Apply program uses table mode processing. For table mode processing, the Apply program fetches answer sets for the subscription-set members one member at a time, until all data has been processed. At the end of the data processing for the whole set, the Apply program issues a single commit. The COMMIT_COUNT of NULL is used to tell the Apply program to use table mode processing. The COMMIT_COUNT that you specified is not NULL. Therefore the COMMIT_COUNT is invalid and is ignored.

User response: This message is for your information only. No action is required.

ASN2239I The action *action_name* ended successfully at *time_stamp* for the Q Apply server *apply_server*, the Q Apply schema *apply_schema* and the remote server name *remote_servername*.

Explanation: The action ended successfully at the Q Apply server.

User response: This message is for your information only. No action is required.

ASN2240I The Q subscription was finished being created at *time_stamp* for the Q subscription name *name*. The Q Capture server is *capture_server*, the Q Capture schema is *capture_schema*, the Q Apply server is *apply_server*, the Q Apply schema is *apply_schema*, the remote server is *remote_servername*, the source table is *table_name* and the target table or stored procedure is *table_name*.

Explanation: The Q subscription was successfully created.

User response: This message is for your information. No action is required.

ASN2241I The Q subscription was finished being dropped at *time_stamp* for the Q subscription name *name*. The Q Capture server is *capture_server*, the Q Capture schema is *capture_schema*, the Q Apply server is *apply_server*, the Q Apply schema is *apply_schema*, the remote server is *remote_servername*, the source table is *table_name* and the target table or stored procedure is *table_name*.

Explanation: The Q subscription was successfully dropped.

User response: This message is for your information. No action is required.

ASN2251E The template *template_name* cannot be dropped because it is being used by one or more monitor suspensions.

Explanation: At least one row exists in the IBMSNAP_MONSUSPENDS table that refers to the provided template name.

User response: Drop all the suspensions that use this template and then drop the template.

ASN2252W The monitor program will be suspended for the entire period between the start date *start_date* and end date *end_date* because a template is not provided.

Explanation: Please see message text.

User response: If you intend to suspend the monitor program for the entire period then no action is necessary. However, if you intend to suspend the monitor program for a certain duration on certain days between the start date and end date then you need to provide a template that is defined using these attributes.

ASN2253E The Replication Alert Monitor control tables on server *server_name* are at a Version 8 architecture level. They must be migrated to a Version 9 architecture level before the suspension function can be used.

Explanation: Monitor suspension templates and suspensions can only be created if the Monitor control tables have been migrated to a level that supports this function. No script is generated.

User response: Run the migration program to upgrade the control table's architecture level and create the set of new control tables before invoking the action.

ASN2254E The template *template_name* cannot be created as another template with the same name already exists at the server *server_name*.

Explanation: A row already exists in the IBMSNAP_TEMPLATES table with the same value in the TEMPLATE_NAME column as the provided template name. No script is generated.

User response: Provide a valid name for the template and reissue the command.

ASN2255E The suspension *suspension_name* cannot be created because another suspension with the same name already exists at the server *server_name*.

Explanation: At least one row exists in the IBMSNAP_SUSPENDS table that has the same value in the SUSPENSION_NAME column as the provided suspension name. No script is generated.

User response: Provide a valid value for the suspension name and reissue the command.

ASN2256E The value for the end date *end_date* is not valid because it is earlier than the start date *start_date*.

Explanation: The end date value must be greater than the starting date value for the monitor program to suspend during that period.

User response: Provide a valid value for the end date and reissue the command.

ASN2257E The value for the start date *start_date* is not valid because it is in the past.

Explanation: The start date value must represent a date value in the future in order to indicate to the monitor program when to start suspending.

User response: Provide a valid value for the start date and reissue the command.

ASN2258E The value *value* provided for the input parameter START DATE or END DATE is out of range.

Explanation: The START DATE and END DATE input parameters require a value that is within the allowed range of the DB2 TIMESTAMP data type.

User response: Make sure that the specified value is within the allowed range. Refer to the DB2 SQL Reference for acceptable values for the TIMESTAMP data type.

ASN2259E The value *value* specified for the input parameter *parameter* is incorrect.
reason_text

Explanation: A value specified as an input parameter to the CREATE TEMPLATE or ALTER TEMPLATE command is either incorrect or outside of the required value range.

User response: Please specify a correct value for the input value based on the *reason_text*.

ASN2261E An operating system error occurred while executing the command *command_name* on the host *host_name* by the stored procedure *schema_name* in database *database_name*. The operating system error message is *message*.

Explanation: The operating system could not execute the command.

User response: Ensure that the fenced user of the DB2 instance of the specified database can execute the provided command.

ASN2262E The requested action failed because the name *name* that was provided for the WebSphere MQ queue manager is not valid or known at the host system *host_name*. The WebSphere MQ error code is *error_code*.

Explanation: The replication administration tool must be able to connect to the specified queue manager to perform the requested action. The queue manager name that was specified does not match any existing queue managers on the host system.

User response: Make sure that the queue manager exists on the host system and check the queue manager name. WebSphere MQ object names are case-sensitive. Read "2058 (080A) (RC2058): MQRC_Q_MGR_NAME_ERROR" in the WebSphere MQ information center for more details.

ASN2263E The requested action failed because the WebSphere MQ queue manager *host_name-queue_manager_name* is not available for connection. The WebSphere MQ error code is *error_code*.

Explanation: The replication administration tool must be able to connect to the specified queue manager to perform the requested action. The queue manager might not be running. Other possible causes can be found in the error code that was returned by WebSphere MQ.

User response: Start the queue manager if it is not running, for example by using the `strmqm` command. Otherwise, read "2059 (080B) (RC2059): MQRC_Q_MGR_NOT_AVAILABLE" in the WebSphere MQ information center for more details.

ASN2264E The requested action failed because the command server for the WebSphere MQ queue manager *host_name-queue_manager_name* is not running. The WebSphere MQ error code is *error_code*.

Explanation: The replication administration tool must be able to connect to the specified queue manager to perform the requested action. The command server must be running to issue commands to the queue manager and transmit responses back to the sender.

User response: Start the command server, for example by using the `strmqcsv` command.

ASN2265E The stored procedure *schema_name* in subsystem *subsystem_name* is not authorized to access the WebSphere MQ queue manager *host_name-queue_manager_name*.

Explanation: The effective user of the stored procedure does not have sufficient authority to connect to the queue manager.

User response: Ensure that the subsystem user, stored procedure definer, or the DB2 user has the authority to access WebSphere MQ. Read "2035 (07F3) (RC2035): MQRC_NOT_AUTHORIZED" in the WebSphere MQ information center for more details.

ASN2266E An error occurred while accessing the WebSphere MQ queue manager *host_name-queue_manager_name*. The WebSphere MQ reason code is *reason_code*.

Explanation: A problem occurred while accessing the queue manager.

User response: View the explanation for this reason code in the WebSphere MQ information center and correct the problem.

ASN2267E A WebSphere MQ queue with the specified name *queue_name* does not exist in the queue manager *host_name-queue_manager_name*. The WebSphere MQ error code is *error_code*.

Explanation: A queue with the specified name cannot be found at the specified queue manager. The queue name that was specified might be incorrect.

User response: Check the queue name for misspellings and ensure that the queue exists in the queue manager (for example using the `DISPLAY QUEUE` command). Read "2085 (0825) (RC2085): MQRC_UNKNOWN_OBJECT_NAME" in the WebSphere MQ information center more details.

ASN2268E The action *action* failed on WebSphere MQ queue *host_name-queue_manager_name-queue_name*. The WebSphere MQ reason code is *reason_code*.

Explanation: A problem occurred while putting a test message on a queue or getting a message from a queue.

User response: Look up the reason code in the WebSphere MQ information center for more details.

ASN2270E The stored procedure *schema_stored_procedure_name* in database *database_name* is not authorized to access the WebSphere MQ queue manager *queue_manager_name* because the operating system user ID *user_ID*, which is the DB2 fenced user of the instance that contains this database, is not a member of the operating system group for WebSphere MQ applications (usually `mqm`) at the host *host_name*.

Explanation: The fenced user of the database instance does not have the authority to access the WebSphere MQ environment or the specified queue manager.

User response: Ensure that the fenced user ID of the database instance is part of the operating system group assigned to WebSphere MQ (usually `mqm`). Read "2035 (07F3) (RC2035): MQRC_NOT_AUTHORIZED" in the WebSphere MQ information center for more details.

ASN2271W The WebSphere MQ queue manager *host_name-queue_manager_name* has the version *version_number* that is older than the minimum supported version *version_number*.

Explanation: Q Replication and event publishing require a minimum version level of WebSphere MQ.

User response: Check the installation for the correct

version, or install an WebSphere MQ server at the minimum version level.

ASN2272W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* has an invalid definition and cannot be opened. The WebSphere MQ error code is *error_code*.

Explanation: The specified queue is defined incorrectly.

User response: Look up the reason code in the WebSphere MQ information center for more details, fix the problem, and then retry the action.

ASN2273W The WebSphere MQ base queue (BASE_Q) *queue_name1* that is referenced by the alias queue *host_name-queue_manager_name-queue_name2* does not exist.

Explanation: The WebSphere MQ alias queue definition is inconsistent. Any alias queue used by Q Replication or event publishing must be defined with an existing base queue.

User response: Check the parameter BASE_Q of the alias queue for misspellings, and check if the specified queue exists. If necessary, update the BASE_Q parameter of the alias queue, or create a base queue with the specified name. Ensure that the base queue has the correct type and parameters as required by Q Replication.

ASN2274W The WebSphere MQ transmission queue *queue_name* for the remote queue *host_name-queue_manager_name-queue_name* does not exist.

Explanation: When a remote queue such as the Q Capture send queue or Q Apply administration queue is created, you use the XMITQ attribute to specify the transmission queue that is used. No transmission queue exists with the name that was specified with the attribute XMITQ for the specified remote queue in the same queue manager.

User response: Check the attribute XMITQ of the specified remote queue for misspellings, or create a transmission queue for the specified remote queue.

ASN2275W The maximum message size (MAXMSGL) *size1* of the WebSphere MQ queue *host_name-queue_manager_name-queue_name* is greater than the maximum message size (MAXMSGL) *size2* of its queue manager *host_name-queue_manager_name*.

Explanation: The maximum message size (MAXMSGL) of the specified queue must be less than

or equal to the maximum message size (MAXMSGL) of the queue manager.

User response: Reduce the maximum message size of the queue or increase the maximum message size of the queue manager.

ASN2276W The maximum message size (MAXMSGL) *size* of the WebSphere MQ queue *host_name-queue_manager_name-queue_name* that is used as the send queue is greater than the maximum message size (MAXMSGL) *size* of the receive queue *host_name-queue_manager_name-queue_name*.

Explanation: The maximum message size (MAXMSGL) of the send queue must be less than or equal to the maximum message size (MAXMSGL) of the receive queue.

User response: Reduce the maximum message size of the send queue or increase the maximum message size of the receive queue.

ASN2277W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* cannot be used as a Q_replication_queue_type because it is neither a local queue nor an alias queue that references a local queue.

Explanation: The Q Capture administration queue and restart queue must be local queues or alias queues that reference a local queue.

User response: Specify a local queue or an alias queue that references a local queue.

ASN2278W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* cannot be used as a send queue because it is not a local queue, a remote queue, or an alias queue that directly or indirectly references a local or remote queue.

Explanation: When you create or change a replication queue map, the send queue that you specify must be a local queue, remote queue, or an alias queue that references a local or remote queue.

User response: Specify a queue of a correct type.

ASN2279W The maximum message size (max_message_size) *size* of the queue map *queue_map_name* is greater than the maximum message size (MAXMSGL) *size* of the WebSphere MQ queue *host_name-queue_manager_name-queue_name* that is used as the send queue.

Explanation: The `max_message_size` attribute of the queue map, which specifies the maximum size of the buffer that is used for sending messages over a send queue, must not be larger than the WebSphere MQ maximum message length (MAXMSGL) attribute that is defined for the queue.

User response: Use the Replication Center or ASNCPLP command-line program to decrease the `max_message_size` value for the queue map or increase the MAXMSGL attribute of the queue.

ASN2280W The maximum message size (`max_message_size`) *size1* of the replication queue map *replication_queue_map_name* is greater than the maximum message size (MAXMSGL) *size2* of the WebSphere MQ queue *host_name-queue_manager_name-queue_name* that is used as the send queue.

Explanation: The maximum message size of the replication queue map must be less than or equal to the maximum message size of the specified send queue.

User response: Reduce the `max_message_size` value of the replication queue map or increase the MAXMSGL value of the queue.

ASN2281W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* cannot be used as a *Q_replication_queue_type*. Although the Q Capture and Q Apply programs use the same queue manager, the specified queue is neither a local queue nor an alias queue that directly or indirectly references a local queue.

Explanation: When the Q Capture and Q Apply programs use the same queue manager, the send queue, receive queue, and administration queue that you specify for a replication queue map must be local queues or alias queues that directly or indirectly reference local queues.

User response: Specify a queue of a correct type.

ASN2282W Although the Q Capture and Q Apply programs use the same queue manager *host_name-queue_manager_name*, the queue *queue1* that is used as a *Q_replication_queue_type1* and the queue *queue2* that is used as a *Q_replication_queue_type2* are not identical or are not alias queues that refer to the same local queue.

Explanation: When the Q Capture and Q Apply programs use the same queue manager, the send queue and receive queue must be either the same local queue

or alias queues that refer to the same local queue. The same holds true for the Q Capture and Q Apply administration queues.

User response: Choose one queue for both purposes.

ASN2283W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* cannot be used as a receive queue in the replication queue map *replication_queue_map_name*. Although the Q Capture and Q Apply programs use different queue managers, the specified queue is neither a local queue nor an alias queue referencing a local queue directly or indirectly.

Explanation: When Q Capture and Q Apply use different queue managers, the receive queue must be local queues or alias queues referencing a local queue directly or indirectly.

User response: Specify a queue of a correct type.

ASN2284W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* cannot be used as *Q_replication_queue_type* in the replication queue map *replication_queue_map_name*. Although the Q Capture and Q Apply programs use different queue managers, the specified queue is neither a remote queue nor an alias queue that directly or indirectly references a remote queue.

Explanation: When Q Capture and Q Apply use different queue managers, the send queue and Q Apply administration queue must be remote queues or alias queues that directly or indirectly reference remote queues.

User response: Specify a queue of a correct type.

ASN2285W The WebSphere MQ model queue *queue_name* does not exist in queue manager *host_name-queue_manager_name*, although the Q subscription specifies a load phase.

Explanation: For a Q subscription that uses a load phase, a model queue needs to exist in the queue manager that the Q Apply program uses so that the Q Apply program can create a spill queue dynamically. By default, Q Apply expects the model queue to be named IBMQREP.SPILL.MODELQ.

User response: Check the name of the model queue for misspellings, create a new model queue, or specify no load phase for the Q subscription.

The following attributes are mandatory for the model queue:

- Message delivery sequence (MSGDLVSQ): FIFO

- Maximum queue depth (MAXDEPTH): 500000 or higher
- Maximum message size (MAXMSGL): 100000 or higher
- Definition type (DEFTYPE): PERMDYN

ASN2286W The value *value1* of the parameter *parameter_name* of the WebSphere MQ model queue *host_name-queue_manager_name-queue_name* is incorrect or too low. The required value is *value2*.

Explanation: The model queue that is used to create spill queues for the Q Apply program needs to have the following parameter values:

- Message delivery sequence (MSGDLVSQ): FIFO
- Maximum queue depth (MAXDEPTH): 500000 or higher
- Maximum message size (MAXMSGL): 100000 or higher
- Definition type (DEFTYPE): PERMDYN

User response: Update the model queue parameter values, or specify a model queue that has the correct parameter values.

ASN2287W The WebSphere MQ queue *host_name-queue_manager_name-queue_name* cannot be used as *Q_replication_queue_type* because the value *value1* of the attribute *parameter_name* does not match the required value *value2*.

Explanation: The specified WebSphere queue attribute is not valid for the specified purpose.

User response: Choose a different queue or update the attribute to the required value.

ASN2288W The test message put on WebSphere MQ queue *host_name1-queue_manager_name1-queue_name1* did not arrive at the queue *host_name2-queue_manager_name2-queue_name2*.

Explanation: The message flow between the first and the second queue does not work properly.

User response: Test messages can only be delivered successfully if the Q Capture and Q Apply programs are not running. Ensure that the Q Capture and Q Apply programs are stopped. Check the WebSphere MQ setup for error messages and test the message flow by using the programs that are provided by WebSphere MQ (amqsput, amqsget, amqsbcbg, amqsbr).

ASN2289W The test message put on WebSphere MQ queue *host1_queue_manager1_queue1* was received at the queue *host2_queue_manager2_queue2*, but the content has been distorted.

Explanation: The message flow between the first and the second queue does not work properly. Messages can be transmitted, but their contents are being distorted.

User response: Check the code pages of the source and target queue managers, and use the WebSphere MQ tools to check for configuration faults that can lead to the distortion of messages between the two queues.

ASN2290I The test message put on WebSphere MQ queue *host_name1-queue_manager_name1-queue_name1* was received at the queue *host_name2-queue_manager_name2-queue_name2* used as *Q_replication_queue_type*.

Explanation: The test message confirms that the message flow between the two queues is functioning properly.

User response: This message is for your information only. No action is required.

ASN2291I *number* tests failed.

Explanation: The WebSphere MQ environment for the current object was validated by several checks, and problems were found.

User response: For any tests that failed, examine the corresponding messages and correct the problem.

ASN2293E The OPERATOR CONSOLE option is not valid for the Replication Alert Monitor with monitor qualifier *monitor_qualifier* and monitor server *server_name* because the monitoring server is on Linux, UNIX, or Windows and not on z/OS.

Explanation: The OPERATOR CONSOLE option is used to send alerts to the z/OS console. This option is not valid for Linux, UNIX, or Windows operating systems.

User response: Send alerts to a contact or contact group for this monitor.

ASN2294E The replication action *action_name* ended in error. The architecture level *arch_level* of the Q Capture server *server* for the Q Capture schema *Q_capture_schema* is not compatible with the Version 9 DB2 database for Linux, UNIX, or Windows where the Q Capture control tables are defined.

Explanation: In order to capture data and schema changes made to a Version 9 DB2 database on Linux, UNIX and Windows platforms, a Version 9 Q Capture server is required. Version 9 Q Capture server refers to both the architecture level of the control tables and the version of the Q Capture program instance running on the server. But based on the value of the architecture level obtained from the IBMQREP_CAPPARMS table, the control table structure is at a level prior to Version 9. This indicates there are inconsistencies in the data stored in the control tables or the control tables have not been migrated.

User response: Migrate your existing Q Capture control server to the current architecture level.

ASN2296E The table *schema_name* cannot be used as a CCD table because it does not have the correct column structure. The column *column_name* is missing or has an invalid definition.

Explanation: CCD tables must have the following four columns:

- IBMSNAP_INTENTSEQ
- IBMSNAP_OPERATION
- IBMSNAP_COMMITSEQ
- IBMSNAP_LOGMARKER

User response: Choose an existing table with the correct CCD columns, or use the Replication Center or ASNCLP program to create a new CCD table.

ASN2297E The table *schema_name* cannot be used as a user table because it has a CCD column structure. The column *column_name* indicates a CCD table structure.

Explanation: CCD tables cannot be used as user table targets for Q subscriptions.

User response: Choose a different existing table that does not contain any CCD columns, or use the Replication Center or ASNCLP program to create a new target table.

ASN2298E The operation *operation* cannot be completed because the architecture level *arch_level* of the schema *schema* for the program *program_name* is older than the required architecture level *arch_level*.

Explanation: The specified operation requires a higher architecture level.

User response: Check if the provided schema is correctly specified for the requested operation. If not, choose another, appropriate schema for the provided program.

ASN2299E The CCD table *schema.name* cannot be populated by Q Apply schema *Q_apply_schema* because the Q subscription *Q_subscription_name* in this schema already uses the CCD table as a target.

Explanation: CCD tables populated by Q Apply can be the target of only one Q subscription.

User response: Remove all but one Q subscription with the CCD table as a target.

ASN2300E The CCD option *option* cannot be specified together with the value *value* for the Q subscription attribute *attribute_name*. Valid attribute values are *value_list*.

Explanation: The command attempts to create a Q subscription with a target table type of CCD and some of the specified parameter values are not compatible with this type of target.

User response: Specify a valid value for the Q subscription attribute.

ASN2301E The Q subscription *receive_queue_Q_subscription_name* cannot be used to populate the SQL registration of the CCD table *schema_name* because the value *value1* of the attribute *attribute_name* of the Q subscription does not match the value *value2* of the SQL registration.

Explanation: The attributes for "complete" and "condensed" have to match between the Q subscription and the SQL registration.

User response: Update the attributes of the existing Q subscription or SQL registration.

ASN2302E The configuration of the Q Apply schema *Q_apply_schema* cannot be changed to populate registrations in the SQL Capture schema *capture_schema* because Q Apply is currently configured to populate registrations in the SQL Capture schema *capture_schema*.

Explanation: A Q Apply program (which is identified by a Q Apply schema) can only send changes to CCD tables that are registered in a single SQL Capture schema. The specified Q Apply schema is already being used to send changes to CCD tables in a SQL Capture schema. If you want to use the specified Q Apply schema for CCD tables in a different SQL Capture schema, change the SQL_CAP_SCHEMA value in the IBMQREP_APPLYPARMS table for the specified Q Apply program.

User response: Check the Q Apply schema for misspellings, check the configured SQL Capture schema in the Q Apply schema, or update the Q Apply schema. To update the schema, use the ALTER APPLY CONFIGURATION command in the ASNCLP command-line program or the Change Saved Parameters window for the Q Apply program in the Replication Center.

ASN2303W The configuration of the Q Apply schema *Q_apply_schema* is being changed from populating registrations in the SQL Capture schema *capture_schema* to populate registrations in the SQL Capture schema *capture_schema*. Existing, inactive Q subscriptions populating CCD target tables may become invalid.

Explanation: Q Apply has already been configured to use a different SQL capture schema. The update will overwrite the existing value and invalidate any existing Q Subscriptions with CCD targets.

User response: No action is necessary if the existing Q subscriptions can be invalidated. Otherwise, choose a different Q Apply schema.

ASN2304E The configuration of the Q Apply schema *Q_apply_schema* cannot be changed from populating registrations in the SQL Capture schema *capture_schema* to populate registrations in the SQL Capture schema *capture_schema*. There are active Q Subscriptions populating CCD target tables in this Q Apply schema.

Explanation: Q Apply has already been configured to use a different SQL capture schema. The update will overwrite the existing value and invalidate any existing Q Subscriptions with CCD targets.

User response: Deactivate the Q Subscriptions, or use

a different Q Apply schema.

ASN2305E The Q Apply schema *Q_apply_schema* does not contain a Q subscription that uses the provided CCD table *schema_name* as a target table.

Explanation: A Q subscription must exist for this CCD table before it can be registered as a source populated by Q Replication.

User response: Check the table name for misspellings, and verify that the provided table is a target table for a Q Subscription in the provided Q Apply schema.

ASN2306E The subscription member cannot be added to the subscription set because the existing members in this set are using registrations that are not populated by a Q Apply schema, and the new member is using a registration that is populated by the Q Apply schema *Q_apply_schema*.

Explanation: A subscription set can only contain member with registrations that are populated by a Q Apply schema, or members with registrations that are not populated by a Q Apply schema.

User response: Choose another subscription set that is empty or that contains compatible members.

ASN2307E The subscription member cannot be added to the subscription set because the existing members in this set are using registrations that are populated by the Q Apply schema *Q_apply_schema*, and the new member is using a registration that is not populated by a Q Apply schema.

Explanation: A subscription set can only contain member with registrations that are populated by a Q Apply schema, or members with registrations that are not populated by a Q Apply schema.

User response: Choose another subscription set that is empty or that contains compatible members.

ASN2308E The subscription member cannot be added to the subscription set because the existing members in this set are using registrations which are populated by the Q Apply schema *schema*, and the new member is using a registration that is populated by the Q Apply schema *Q_apply_schema*.

Explanation: A subscription set can only contain member with registrations that are populated by a Q Apply schema if the Q Apply schema and the receive queue are the same for all registrations of the members.

User response: Choose another subscription set that is empty or that contains compatible members.

ASN2309E The subscription member cannot be added to the subscription set because the existing members in this set are using registrations that are populated by the Q Apply schema *Q_apply_schema* using the receive queue *receive_queue_name*, and the new member is using a registration that is populated using the receive queue *receive_queue_name*.

Explanation: A subscription set can only contain member with registrations that are populated by a Q Apply schema if the Q Apply schema and the receive queue are the same for all registrations of the members.

User response: Choose another subscription set that is empty or that contains compatible members.

ASN2310W The SQL registration for the CCD table *schema_name* in the SQL Capture schema *capture_schema* will become unpopulated.

Explanation: If the Q subscription populating a CCD table registration in a SQL Capture schema is dropped, the registration will no longer receive updates from Q Apply.

User response: No action is required.

ASN2311I The table *table_owner.table_name* has an edit routine (EDITPROC) defined on it. Replication requires DB2 Version 8 APAR PK13542 or later to correctly handle the data in such tables.

Explanation: Replication supports tables with EDITPROC definitions, but the DB2 subsystem that contains the tables must be upgraded to the specified APAR. With the APAR, DB2 invokes the edit routine to transform the row back to its original state before giving the data to the replication log reader program.

User response: Upgrade the DB2 instance that contains the table to the specified APAR.

ASN2312E The target type of *type* does not support before-image columns.

Explanation: The only valid target type for before-image columns is CCD.

User response: Change the target type for the Q subscription to CCD.

ASN2313E The before-image column prefix *prefix* is invalid because its length *length* is not within the allowed range of 1 to 4 characters.

Explanation: The prefix that is used to identify values in a before-image column must be at least 1 and at most 4 characters long.

User response: Change the length of the prefix to a value that is within the required range.

ASN2314E The before-image column prefix *prefix* is invalid because it contains the character *character*, which is not supported.

Explanation: The before-image column prefix can only contain alphabetical characters that the database allows at the start of column names and within column names.

User response: Replace or remove invalid characters from the prefix.

ASN2315E The before-image column *column1* does not exist for the column *column2*.

Explanation: The Q subscription was set up so that the column from the source table has a corresponding before-image column in the target table. Before-image column names are generated by adding a before-image prefix to the name of the target column that maps to the source column. No column could be found in the target table with this generated name.

User response: Verify that the target table contains a column with this generated name. If not, change the Q subscription to update the before-image prefix value and regenerate the script. If the before-image is not needed, remove it from the target column specification.

ASN2316E The target column *column_name* cannot be selected for before-images because the data type *data_type* of the column does not support before images.

Explanation: Q Replication supports before-image columns for all data types except for large object (LOB) data types.

User response: Remove this column from the before-image column selection.

ASN2317E The column *column_name* cannot be selected for a before-image column because the length *length* of the combined column name and before-image prefix *combined_name* exceeds the limit of the database *database_alias*.

Explanation: The combined length of the column name and before-image prefix must be within the database limits for column-name length.

User response: Shorten the before-image prefix.

ASN2318E The column *column_name1* cannot be selected as a replication target column because it is used to hold the before-image values of column *column_name2*.

Explanation: Before-image columns cannot be target columns for replication.

User response: Either change the value of the before-image column prefix or select a different before-image column for the specified after-image column.

ASN2319W The before-image columns in the target table *table_owner.table_name* do not have the same before-image prefix. You will not be able to register this table with the SQL Replication Apply program for a three-tier distribution.

Explanation: Before-image columns in a CCD table that is used as the middle tier in three-tier replication must use the same single character for the before-image prefix. The before-image columns in the specified table do not have the same prefix. The table cannot be used in a three-tier architecture.

User response: Modify the before-image prefix if you intend to use this table in three-tier replication architecture.

ASN2320E The column *column_name1* requested to be a before-image column for the column *column_name2* does not exist in the target table.

Explanation: To be used as a before-image column, a column must first exist in the target table.

User response: Specify a column name that exists in the target table.

ASN2321E The column *column_name1* with data type *data_type* cannot be used as a before-image column for column *column_name2* with data type *data_type* because their data types are not compatible.

Explanation: The data types for an after-image column and its corresponding before-image column must be compatible.

User response: Make sure that you specify the correct column names.

ASN2322E The table *table_owner.table_name* cannot be registered as an external CCD table because the before-image columns do not have the same prefix.

Explanation: The before-image columns in a CCD table must share a common prefix in order to register the table as an external CCD. The before-image columns in the specified table do not share a common prefix.

User response: To use this CCD table as a SQL Replication source, redefine the Q subscription so that the before-image columns for the table share a common prefix, then perform the action again.

ASN2323E Before values cannot be sent for the column *column_name* in the table *table_owner.table_name* because the column is not part of the publication.

Explanation: You have requested that the Q Capture program send before values for the specified column, but this column was not included in the publication.

User response: Either select the column to be part of the publication or do not request before values for the column.

ASN2324I The default before-image prefix, X, resulted in conflicting column names in the target table *table_owner.table_name*. A different before-image prefix, *prefix*, will be used to generate unique before-image column names.

Explanation: The columns in the table should be unique. However, the default before-image prefix resulted in conflicting names. The specified prefix will be used instead.

User response: This message is for your information only. No action is required.

ASN2325E The Q Apply schema *schema* cannot be used to populate CCD tables that are registered for SQL Replication because the SQL Capture schema has not been specified.

Explanation: You must first configure the Q Apply program to use a specific SQL Capture schema for the CCD tables that it populates before you register the tables as sources for SQL Replication.

User response: Update the Q Apply schema using the ALTER CONFIGURATION APPLY command in the ASNCLP command-line program, or by using the Change Saved Parameters window for the Q Apply program in the Replication Center.

ASN2326E The SQL registration for the CCD target table *table_owner.table_name* of the Q subscription *receive_queue_name.q_subscription_name* cannot be created because this Q subscription uses the Q Capture schema *schema* and the Q Apply schema *schema*. These schemas are different from the Q Capture schema *schema* and the Q Apply schema *schema* that are used by other Q subscriptions that were selected.

Explanation: Only Q subscriptions that are replicated from the same Q Capture schema to the same Q Apply schema can be registered at the same time.

User response: Exclude the specified Q subscription from the group of Q subscriptions that you are registering.

ASN2327E The existing SQL registration of the CCD table *table_owner.table_name* in the SQL Capture schema *schema* cannot be modified to be used by the Q Apply program because the value *value1* for the property *property1* in this registration conflicts with the value *value2* for the property *property2* in the Q subscription *receive_queue_name.q_subscription_name*.

Explanation: The SQL registration and Q subscription need to have the same values for the CCD table properties CONDENSED and COMPLETE. The values can be either ON or OFF. The matching properties are required in order to change an existing registration to be used by the Q Apply program.

User response: Change the Q subscription so that the CCD table properties CONDENSED and COMPLETE match the properties of the SQL registration.

ASN2328E The Capture schema *schema* does not exist at the database *database_alias*.

Explanation: The specified Capture schema could not be found at this database.

User response: Check the name for misspellings and correct the name of the Capture schema.

ASN2329E The Capture schema *schema* has an architecture level of *level1* that is older than the required architecture level *level2*.

Explanation: The Capture schema is too old.

User response: Choose a different Capture schema, or migrate the Capture schema to the required architecture level.

ASN2330E The Capture schema *schema* is defined in the federated database *database_name* and cannot be used to register CCD tables that are targets for Q Replication.

Explanation: The Capture schema was set up for federated data sources and cannot be used for DB2 sources.

User response: Choose another Capture schema that is not set up for federated data sources.

ASN2331E The Capture schema *schema* that is managed by the Q Apply schema *schema* cannot be modified to the new schema *schema2* because an active Q subscription, *receive_queue_name.q_subscription_name*, is managing the SQL registration for the target CCD table *table_owner.table_name* in the existing schema.

Explanation: If the Capture schema were changed, it would invalidate the active Q subscription.

User response: Deactivate the Q subscription and retry the action.

ASN2332W The Capture schema *schema* that is managed by the Q Apply schema *schema* will be modified to the new schema *schema2*. The existing and inactive Q subscription *receive_queue_name.q_subscription_name* may become invalid.

Explanation: The requested update of the Capture schema could invalidate Q subscriptions.

User response: No user action is necessary.

ASN2333E The length *length1* of the value *schema* for the Capture schema is higher than the maximum supported length *length2*.

Explanation: The specified Capture schema has an invalid length.

User response: Reduce the length of the Capture schema and rerun the action.

ASN2334W The SQL registration for the CCD table *table_owner.table_name* will be modified so that it will not longer be populated by the Q subscription *receive_queue_name.q_subscription_name*.

Explanation: A SQL registration can receive updates only from one Q subscription. If another Q subscription is selected to send updates, the previous one will no longer be able to update the target as well.

User response: Choose one of the following options:

- If you would like the SQL registration to be populated by the specified Q subscription, then no action is necessary. You can execute the generated script.
- If you would like to continue using the original Q subscription to populate the SQL registration, then do not execute the generated script and cancel the action.

ASN2335E The SQL registration for the CCD table *table_owner.table_name* cannot be found in the SQL Capture schema *schema* at the database *database_alias*.

Explanation: The SQL Capture schema at this database does not contain the specified SQL registration.

User response: Check the name of the CCD table for which the registration was specified, and retry the action.

ASN2336E A SQL registration for the CCD table *table_owner.table_name* already exists in the SQL Capture schema *schema* at the database *database_alias*.

Explanation: A SQL registration was specified but another registration already exists with the same name in the SQL Capture schema on this database.

User response: Check the name of the CCD table for which the registration was specified, and retry the action.

ASN2337W The SQL registration for the CCD table *table_owner.table_name* will become unpopulated when the Q subscription is dropped.

Explanation: When the Q subscription that is populating this CCD table is dropped, the SQL registration of the CCD table will no longer receive data from the source.

User response: No user action necessary.

ASN2338W The SQL registration for the CCD table *table_owner.table_name* will be dropped. Ensure that the SQL subscription-set member for the target table *table_owner.table_name* at the SQL Apply control server *database_alias* in the SQL subscription set *set_name* that uses the Apply qualifier *apply_qualifier* has been deactivated.

Explanation: When this SQL registration is dropped, the subscription-set member that receives data from the registration should be inactive to avoid data loss.

User response: Check if the specified member has been deactivated.

ASN2339W The SQL registration for the CCD table *table_owner.table_name* will be updated to be populated by the Q subscription *receive_queue_name.q_subscription_name*. Ensure that the subscription-set member for the target table *table_owner.table_name* at the SQL Apply control server *database_alias* in the SQL subscription set *set_name* that uses the Apply qualifier *apply_qualifier* has been deactivated.

Explanation: When the SQL registration is updated, the subscription-set member that receives data from the registration should be inactive to avoid data loss.

User response: Check if the specified member has been deactivated.

ASN2340E The subscription-set member cannot be added to the subscription set because the existing members in this set are using source registrations that are not populated by a Q Apply program, and the new member is using a source registration that is populated by the Q Apply schema *schema*.

Explanation: A subscription set can only contain members with source registrations that are populated by one Q Apply program using the same receive queue, or that are not populated by Q Apply.

User response: Choose another subscription set that is empty or contains compatible members.

ASN2341E The subscription-set member cannot be added to the subscription set because the existing members in this set are using source registrations that are populated by Q Apply schema *schema* and the new member is using a source registrations this is not populated by a Q Apply program.

Explanation: A subscription set can only contain members with source registrations that are populated by one Q Apply program using the same receive queue, or that are not populated by Q Apply.

User response: Choose another subscription set that is empty or contains compatible members.

ASN2342E The subscription-set member cannot be added to the subscription set because the existing members in this set are using source registrations that are populated by Q Apply schema *schema* using the receive queue *receive_queue*, and the new member is using a source registration this is populated by Q Apply schema *schema2* using the receive queue *receive_queue*.

Explanation: A subscription set can only contain members with source registrations that are populated by one Q Apply program from the same receive queue.

User response: Choose another subscription set that is empty or contains compatible members.

ASN2343E The SQL registration for the target CCD *table_owner.table_name* cannot be modified to be managed by the Q Apply program because the value *value* of the property *property* is unsupported.

Explanation: Q Apply can only manage specific SQL registrations. For details about the restrictions, see "Creating a three-tier data distribution configuration (Q Replication to SQL Replication)" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center.

User response: Update the SQL registrations, or choose a different SQL registration in a different SQL Capture schema.

ASN2344I The target CCD *table_owner.table_name* for this Q subscription has an existing SQL registration in the SQL Capture schema *schema*, and the Q Apply program is configured to manage SQL registrations in this schema. If this Q subscription becomes active, the Q Apply program will manage the target CCD table as a source for SQL Replication.

Explanation: The target CCD is already registered and Q Apply has been configured correctly, so the target CCD will automatically become a source for SQL Replication.

User response: This message is for your information only. No action is required.

ASN2345E The resulting data type *data_type1* of the expression *expression* is not compatible with the data type *data_type2* of the mapped target column *column_name*.

Explanation: In order to replicate an expression to a target column, the Q Apply program requires the resulting data type of the expression to match the data type of the target column. The data types of the

expression and the target column to which it is mapped do not match.

User response: Modify your expression so that its resulting data type matches the data type of the specified target column, or choose another column in the target that has the same data type as the expression.

ASN2346E The column *column_name* that is specified in the expression *expression* does not exist in the source table *table_owner.table_name*.

Explanation: A column must exist in the source table to be part of an expression, and the specified column could not be found.

User response: Specify a different column that exists in the source table or remove the column from the expression.

ASN2347E The expression *expression* cannot be mapped to the target column *column_name* because the expression can have a null value and the target column is defined with a NOT NULL attribute.

Explanation: The null characteristics of the expression and the target column must match. None of the source columns chosen in the expression is defined with a NOT NULL attribute, so the resulting value of the expression can be null. The target column is defined as NOT NULL.

User response: Include at least one source column that is defined as NOT NULL in the expression, or map the expression to another column in the target that can hold null values.

ASN2348E The column *column_name* in the target table *table_owner.table_name* cannot be directly mapped to a source column because the specified target column already maps to an expression *expression*.

Explanation: A target column cannot be mapped to both a source column and an expression.

User response: Choose a different target column for the mapping.

ASN2349E The column *column_name* in the target table *table_owner.table_name* cannot be mapped to an expression *expression* because the target column is already mapped directly to a source column *column_name*.

Explanation: A target column cannot be mapped to both an expression and a source column.

User response: Choose a different target column for the mapping.

ASN2350E **The Replication administration tool is running at a DB2 Version 8 level. The tool cannot be used to administer a DB2 Version 9 server for Q Replication.**

Explanation: You cannot use a pre-Version 9 Replication administration tool to administer a DB2 Version 9 server because of the changes in the structure of the replication control tables and the replication architecture level.

User response: Migrate the DB2 client where the replication administration tool is running to Version 9.

ASN2351E **The CCD table *schema.name* cannot be defined as condensed because the database object *name* of type *typename* enforces uniqueness of the columns *column_names*, which are not part of the replication key.**

Explanation: The Apply program will fail when trying to apply changes to the target table if this table contains a primary key, unique constraint, or unique index that enforces the uniqueness of columns that are not part of the replication key. This is usually the case if there is a second unique constraint or index defined on this table. **Note:** You can have secondary unique indexes on CCD columns, for example, IBMSNAP_INTENTSEQ and IBMSNAP_COMMITSEQ.

User response: Remove the unique constraint or index, or consider generating a new CCD target table.

ASN2353E **On z/OS, you cannot use the Replication Center to perform operational tasks on replication and event publishing programs such as starting, stopping, or reinitializing.**

Explanation: The Replication Center function to run commands that operate on replication or publishing programs on z/OS has been deprecated.

User response: To start the programs, use JCL or a started task. To issue commands to running replication or publishing programs such as stop or reinitialize, use the MODIFY command.

ASN2354E **You cannot use the Replication Center on z/OS to perform operational tasks such as starting, stopping, or reinitializing replication and event publishing programs.**

Explanation: The Replication Center function to run commands that operate on replication or publishing programs on z/OS has been deprecated.

User response: To start the programs, use JCL or a

started task. To issue commands to running replication or publishing programs such as stop or reinitialize, use the MODIFY command.

ASN2355E **The command *command* cannot be used with a Classic replication source.**

Explanation: A Classic replication server was set as the capture server in a prior SET SERVER CAPTURE command. The specified command is not valid for Classic replication sources. The following commands are valid for a Classic source:

- CREATE REPLQMAP
- ALTER REPLQMAP
- DROP REPLQMAP
- CREATE QSUB
- ALTER QSUB
- DROP QSUB
- LOADDONE
- START QSUB
- STOP QSUB
- LIST
- ALTER CAPPARMS

User response: Take one of the following actions:

- Issue the environment command SET SERVER CAPTURE to set the Capture server to another server where the command is supported.
- Use one of the supported commands for Classic sources.

ASN2356W **A default schema of ASN is assigned because the capture server is a Classic replication server. The schema that was provided is not used.**

Explanation: A Classic replication server was set as the capture server in a prior SET SERVER CAPTURE command. Only one set of Classic capture control tables exists on the server. These tables have a schema of ASN and the ASNCLP program always uses a default schema of ASN.

User response: This message is for your information only. No action is necessary.

ASN2357E **The option *option* cannot be used with a Classic replication source.**

Explanation: A Classic replication server was set as the capture server in a prior SET SERVER CAPTURE command. The specified option is not valid for Classic sources. The following options are not supported:

- SUPPRESS DELETES
- SEARCH CONDITION
- TRGCOLS EXCLUDE (for new targets)
- LOAD TYPE with values 1, 2, or 3

User response: Take one of the following actions:

- Issue the environment command SET SERVER CAPTURE to set the Capture server to another server where the option is supported.
- Specify a different option for Classic sources. For LOAD TYPE, the supported options are 0 (no load) or 4 (load from Classic source).

ASN2358E The configuration file *file_name* does not exist at the specified location. Please check the directory name, file name, and permissions on the file.

Explanation: The ASNCLP program did not find the specified configuration file by using the path and file name that was provided in the *file_name* parameter.

User response: Ensure that the path and file name are correct. Also, check whether the permissions on the file are set correctly to allow the ASNCLP program to read the file.

ASN2359E The command to create Q subscriptions did not complete successfully. These *number* Q subscriptions were defined on the Classic server but not on the Q Apply server: *Q_subscription_list1*. These *number* Q subscriptions were not defined on either server: *Q_subscription_list2*. These *number* Q subscriptions were defined on both servers: *Q_subscription_list3*.

Explanation: An error occurred during processing of the CREATE QSUB command. Some of the Q subscriptions were defined on the Classic server only, and some were not defined on either server. The Q subscriptions that are defined on only one server are incomplete and must be deleted.

User response: Follow these steps:

- Use the DROP QSUB command in the ASNCLP program to delete the subscriptions that are defined on the Classic server only.
- Use the CREATE QSUB command to create the Q subscriptions that were not defined on either server.

The Q subscriptions that were defined on both servers are complete and no further action is required.

ASN2360E The command to delete one or more Q subscriptions did not complete successfully. These *number* Q subscriptions were deleted on the Classic server but not on the Q Apply server: *Q_subscription_list1*. These *number* Q subscriptions were not deleted from either server: *Q_subscription_list2*. These *number* Q subscriptions were deleted from both servers: *Q_subscription_list3*.

Explanation: An error occurred during processing of the DROP QSUB command. Some of the Q subscriptions were deleted only on the Classic server and some were not deleted at either server. The Q subscriptions that were deleted only on the Classic server still exist on the Q Apply server and must be dropped manually.

User response: Use the DROP QSUB command in the ASNCLP program to delete the Q subscriptions that were only deleted on the Classic server and the Q subscriptions that were not deleted on either server.

ASN2361E The Q subscription cannot be created because only a subset of the columns from the source table *table_owner.table_name* were selected to take part in replication and the source server is a Classic replication server.

Explanation: For Classic replication, all columns in the source table must be selected for replication.

User response: Specify all columns in the source table to take part in the Q subscription and issue the CREATE QSUB command again.

ASN2362E The action *action_name* ended in error because the architecture level of the Q Apply control tables under the schema *schema* is at pre-Version 9 and the source is a Classic replication server.

Explanation: The Q Apply control tables need to be at a Version 9 architecture level in order to use a Classic replication server as a source. The specified Q Apply schema is at pre-V9 architecture level.

User response: Take one of the following actions:

- Migrate the Q Apply control tables to Version 9.
- Create a new set of Q Apply control tables with Version 9 architecture.

ASN2363E The action *action_name* ended in error because capture control tables were not found on the Classic replication server.

Explanation: The control tables need to exist prior to performing the specified action but were not found. The control tables are created when the Classic replication product is installed.

User response: See the IBM InfoSphere Classic Information Center for details on how to configure a Classic replication server, and complete the steps before defining or using any replication objects.

ASN2364E The action *action_name* ended in error because the IBMQREP_CAPPARMS table is empty.

Explanation: The ASNCLP program requires a row in the IBMQREP_CAPPARMS control table before it can define objects for replication. No row was found in the table.

User response: Please run the ASNCLP command ALTER CAPPARMS to populate the table with a row prior to performing this action.

ASN2365I The Q Apply schema has Q subscriptions or replication queue maps for which the source is a Classic replication server. These definitions are not displayed in the list shown here.

Explanation: Q subscriptions or replication queue maps for which the source is a Classic replication server are not displayed by the Replication Center.

User response: This message is for your information only. No action is required.

ASN2366E The SET command *command_name* cannot be used with an Oracle source server.

Explanation: The specified command is not supported for Oracle sources.

User response: See the ASNCLP documentation in the DB2 Information Center for supported commands for Oracle sources.

ASN2367E The command could not be processed. Oracle servers are not supported as replication targets unless an intermediate federated server is installed and configured.

Explanation: To replicate data to an Oracle target, you must configure a Q Apply program to run within a DB2 federated instance, which requires InfoSphere Federation Server.

User response: Install and configure InfoSphere Federation Server and then use the SET SERVER command to specify a Q Apply server that is on a federated DB2 instance.

ASN2368E The CREATE QSUB command does not support the values B or P for the SUBTYPE keyword if the source is Oracle.

Explanation: Bidirectional (type B) and peer-to-peer (type P) replication are not supported for native Oracle sources. The only allowed Q subscription type is unidirectional (type U).

User response: For the SUBTYPE keyword, specify the value U and reissue the command.

ASN2369E The load type option *option* cannot be used with an Oracle source.

Explanation: An Oracle server was specified as the Q Capture server in a prior SET SERVER command. For Oracle servers, the values 1, 2, 3, and 4 are not supported for the LOAD TYPE keyword.

User response: Take one of the following actions:

- Issue the SET SERVER command and specify another Q Capture server where this option is supported.
- For an Oracle source, specify the values 0 (no load) or 5 (Oracle load) for the LOAD TYPE keyword.

ASN2370E The action *action* ended in error because the architecture level of the Q Apply control tables is *level* and the source is an Oracle server.

Explanation: The Q Apply control tables need to be at a architecture level of 0905 (Version 9.5) or above to use an Oracle server as a source. The specified Q Apply schema is at a previous architecture level.

User response: Take one of the following actions:

- Migrate the Q Apply control tables to the required architecture level.
- Create a new set of Q Apply control tables under the required architecture level and use that schema for the specified action.

ASN2371E The table *table_name* cannot be used as a Q Replication source because table-level supplemental logging is not set.

Explanation: Table-level supplemental logging should be set to ON so that Oracle maintains the required log data for replication.

User response: Issue the ALTER command to alter the source table to enable table-level supplemental logging.

ASN2372E The database *database_name* cannot be used as a Q Replication source because database-wide supplemental logging is not enabled.

Explanation: The Q Capture program uses Oracle LogMiner, which requires minimal supplemental logging to be enabled. This function logs the minimal amount of information needed for Oracle LogMiner to identify, group, and merge the redo operations from DML changes. Minimal supplemental logging ensures that Oracle LogMiner and Q Capture have sufficient information to support chained rows and various storage arrangements, such as cluster tables.

User response: Issue the ALTER command to alter the

database to enable database-wide minimal supplemental logging.

ASN2373E The database that was specified as a Q Capture server cannot be used as a source. Q Replication supports only Oracle version 10g and above.

Explanation: To replicate data from an Oracle source, you must specify a source database that is at Oracle version 10g or later as a Q Capture server.

User response: Specify a database that is at a supported level and retry the command.

ASN2375E The XML keyword was provided in the CREATE PUB command, but the publishing queue map that was specified for the publication uses a DELIMITED message format.

Explanation: The message format must match between the publication and the publishing queue map that it uses.

User response: Take one of the following actions:

- If you want a publication with XML message format, specify a publishing queue map that uses XML format.
- If you want a publication with DELIMITED message format, change the message format for the publication to DELIMITED.

ASN2376E The publishing queue map *queue_map_name* cannot be promoted because it does not exist within the specified Q Capture schema on the specified Q Capture server.

Explanation: Before you can promote an object, the object must exist within the schema that you specify. A publishing queue map with the specified name does not exist within this schema.

User response: Make sure that the correct publishing queue map, schema, and server are specified and issue the command again.

ASN2377E The publication *publication_name* cannot be promoted because it does not exist within the specified Q Capture schema on the specified Q Capture server.

Explanation: Before you can promote an object, the object must exist within the schema that you specify. A publication with the specified name does not exist within this schema.

User response: Make sure that the correct publication, schema, and server are specified and issue the command again.

ASN2378E The replication queue map *queue_map_name* cannot be promoted because it does not exist within the Q Capture schema *Q_Capture_schema* on the specified Q Capture server or within the Q Apply schema *Q_Apply_schema* on the specified Q Apply server.

Explanation: Before you can promote a replication queue map, the queue map must exist within the Q Capture and Q Apply schemas that you specify. A replication queue map with the specified name does not exist within these schemas.

User response: Make sure that the correct replication queue map, schema, and server are specified and issue the command again.

ASN2379E The command is not supported because the target database is accessed through a federated server.

Explanation: The specified command is not supported for federated targets.

User response: Specify a server where the command is supported.

ASN2380I The operation *operation* for program *program_name* has started for schema or qualifier *name* on database or subsystem *name* at host *host_name* at timestamp.

Explanation: The requested operation was initiated.

User response: This message is for your information only. No action is required.

ASN2381I The operation *operation* ended successfully at timestamp *timestamp*.

Explanation: A replication or event publishing program successfully carried out the described operation.

User response: This message is for your information only. No action is required.

ASN2383I The program status is *status*. Additional information: *information*.

Explanation: This message describes the status of a replication or event publishing program.

User response: This message is for your information only. No action is required.

ASN2384E The operation *operation* ended in error because the program *program_name* is not active.

Explanation: The program that was prompted to carry out the described operation was not running, so the

operation could not be executed.

User response: Start the program, and then request the operation again.

ASN2385I **The detailed status of the active program is *detailed_status*.**

Explanation: The requested status information describes the current state of the program in detail.

User response: This message is for your information only. No action is required.

ASN2386I **The program parameter *parameter* has a value of *value* and was defined by *origin*.**

Explanation: The message reflects the parameter value that the program is currently using. This value might differ from the value that is saved in the control tables. When a program stops and restarts, it uses saved values unless you override the saved values at startup or while the program is running.

User response: This message is for your information only. No action is required.

ASN2387E **The operation *operation* cannot be executed because the access method *method* is not available on the database or subsystem *name* at host *host*. Reason: *reason*. Suggested action: *action*. Perform the suggested action to enable this access method for future use.**

Explanation: The requested operation requires the stated access method to be available.

User response: Perform the suggested action and retry the operation.

ASN2388E **The operation was terminated at *timestamp* because the timeout of *number* seconds was exceeded.**

Explanation: A replication or event publishing program could not carry out the requested operation within the specified time.

User response: Check the status of the program and retry the operation.

ASN2389E **The operation *operation* cannot be executed because a prerequisite setup step is missing on the database or subsystem *name* at host *host*. Reason: *reason*. Suggested action: *action*. Perform the suggested action to enable this operation for future use.**

Explanation: The requested operation requires a prerequisite setup step to be completed.

User response: Perform the suggested action to complete the prerequisite setup step and retry the operation.

ASN2390E **An error occurred during communication with the active program *program_name*. Error message: *message*. Response: *response*.**

Explanation: An internal error occurred during communication with the replication program.

User response: Check the DB2 Information Center for details about any required changes to your replication setup, and retry the action.

ASN2391E **The operation *operation* for the program *program_name* is not supported for database_type databases.**

Explanation: Some operations by replication programs are not supported on some database types.

User response: Check the replication and event publishing documentation in the DB2 Information Center to see which operations are supported.

ASN2392W **The DB2 instance could not be determined. The DB2 instance name of *instance_name* is used as a default to perform the requested operation on the database *database_name*. To support other DB2 instance names for this operation, perform the suggested action *action*.**

Explanation: The DB2 instance name is required to perform the operation, but the instance name cannot be determined. A default instance name is used.

User response: If the name of the DB2 instance is the same as the default, no action is necessary. Otherwise perform the suggested action.

ASN2394W **The recommended access method *method* is not available to perform the operation *operation* on the database or subsystem *name* at host *host_name*. Reason: *reason*. Suggested action: *action*. Perform the suggested action in order to make the recommended access method available for future use.**

Explanation: A recommended method of accessing the database or subsystem was not available. An alternative method will be tried to perform the requested operation.

User response: Perform the suggested action and retry the operation.

ASN2395E The program *program_name* could not be started because an instance of this program with the same schema or qualifier is already running on the database or subsystem *name*.

Explanation: Only one program instance with a given schema or qualifier can run within a database.

User response: Either stop the running program that has the same schema or qualifier, or drop and recreate the schema or qualifier for the second program.

ASN2396E The operation *operation* ended in error because the program returned one or more error messages: *error_messages*.

Explanation: The program received the request to perform an operation, but returned an error.

User response: Check the error message, correct the problem, and retry the operation.

ASN2397E The replication queue map cannot be created because the MAXAGENTS_CORRELID value (*number_agents*) is greater than the NUM_APPLY_AGENTS value (*total_agents*).

Explanation: The total number of apply agents is specified by the NUM_APPLY_AGENTS value in the IBMQREP_RECVQUEUES table for a receive queue and replication queue map. The MAXAGENTS_CORRELID value specifies the number of transactions to apply in parallel from the same job name. This setting allows the Q Apply program to limit parallelism during batch workloads to prevent lock contention.

User response: Set the MAXAGENTS_CORRELID value to be lower than the NUM_APPLY_AGENTS value.

ASN2398E The replication queue map could not be updated because the MAXAGENTS_CORRELID value (*number_agents*) is greater than the NUM_APPLY_AGENTS value (*total_agents*).

Explanation: The total number of apply agents is specified by the NUM_APPLY_AGENTS value in the IBMQREP_RECVQUEUES table for a receive queue and replication queue map. The MAXAGENTS_CORRELID value specifies the number of transactions to apply in parallel from the same job name. This setting allows the Q Apply program to limit parallelism during batch workloads to prevent lock contention.

User response: Set the MAXAGENTS_CORRELID value lower than the NUM_APPLY_AGENTS value.

ASN2399E The table *capture_schema.IBMQREP_IGNTRAN* table does not have a row. Your configuration requires that a row exists and that the row contains the authorization ID for the Q Apply program.

Explanation: You are using the peer-to-peer configuration with value-based conflict detection. This configuration requires that the IBMQREP_IGNTRAN table be populated with the authorization ID of the Q Apply program so that the Q Capture program can ignore the transactions and avoid recursion.

User response: Specify the ALTER SERVER ADD IGNORE TRANSACTIONS command.

ASN2400I A row with the specified information already exists in the *capture_schema.IBMQREP_IGNTRAN* table. An SQL script to insert the ignore-transactions specification is not generated.

Explanation: The command did not generate a SQL script because a row already exists.

User response: This message is for your information only. No action is required.

ASN2401I A row with the specified information does not exist in the *capture_schema.IBMQREP_IGNTRAN* table. An SQL script to delete the ignore-transactions specification is not generated.

Explanation: The command did not generate a SQL script because a row does not exist in the IBMQREP_IGNTRAN table.

User response: This message is for your information only. No action is required.

ASN2402E The command *command* is only supported for Q Replication and event publishing.

Explanation: The specified command is not applicable to SQL Replication.

User response: Specify the environment by using the ASNCLP SESSION SET TO Q REPLICATION command before using this command.

ASN2403E The specified source table *owner.table_name* does not have any subscriptions or publications in the capture schema *capture_schema*. The ALTER ADD COLUMN command did not add a signal into the

IBMQREP_SIGNAL table.

Explanation: The command failed because subscriptions or publications for the specified source table do not exist.

User response: Check the source table name and reissue the command.

ASN2404E The column *column_name* is already subscribed for the subscription or publication *sub_name* for the provided source table *owner.table_name*.

Explanation: Because the column is already subscribed for the subscription or publication, the column cannot be added.

User response: If the wrong table or column name was specified, correct the statement and reissue the command. Otherwise, no action is required.

ASN2405W The referential integrity constraints *constraint_name* do not match for the source table *owner.src_table_name* and target table *owner.tgt_table_name*.

Explanation: The referential integrity constraints for the source and target tables for bidirectional and peer-to-peer subscriptions must exactly match. Because the constraints do not match for the source and target tables, you might have experience problems with the Q Apply program.

User response: If the mismatched constraints will not cause problems while running the Q Apply program, no action is necessary. Otherwise, ensure that the referential integrity constraints on the source and target tables match before running the replication programs.

ASN2406W The check constraints *constraint_name* do not match for the source table *owner.src_table_name* and target table *owner.tgt_table_name*.

Explanation: The check constraints for the source and target tables for bidirectional and peer-to-peer subscriptions must exactly match. Because the constraints do not match for the source and target tables, you might experience problems with the Q Apply program.

User response: If the mismatched constraints will not cause problems while running the Q Apply program, no action is necessary. Otherwise, ensure that the check constraints on the source and target tables match before running the replication programs. Scripts were generated because you set enforcement of constraints matching to 'no.'

ASN2407W The unique constraints *constraint_name* do not match for the source table *owner.src_table_name* and target table *owner.tgt_table_name*.

Explanation: The unique constraints for the source and target tables for bidirectional and peer-to-peer subscriptions must exactly match. Because the constraints do not match for the source and target tables, you might have experience problems with the Q Apply program.

User response: If the mismatched constraints will not cause problems while running the Q Apply program, no action is necessary. Otherwise, ensure that the unique constraints on the source and target tables match before running the replication programs.

ASN2408W The target member *target_owner-target_table* in the subscription set *subs_set_name* is an external consistent-change data (CCD) table. This table type does not require a join of the changed-data (CD) and unit-of-work (UOW) tables to populate the LOGMARKER column. You specified a non-zero commit count value for the subscription set. If you use this CCD table as a source in a three-tier distribution and the block factor is non-zero, data loss problems might result while running the Apply program.

Explanation: Because the CCD target type is '9,' the Apply program will not perform a join of the CD and UOW tables to get a value for the LOGMARKER column. You specified a non-zero value for the commit count when you created the subscription set. If this CCD table is used as a source in the second tier, you should specify zero for the blocking factor or MAX_SYNCH_MINUTES column in the IBMSNAP_SUBS_SET table for the subscription set.

User response: Specify zero for the blocking factor or the MAX_SYNCH_MINUTES column in the IBMSNAP_SUBS_SET table.

ASN2410E The keyword GENERATE SQL FOR EXISTING is not supported in SQL Replication.

Explanation: The GENERATE SQL FOR EXISTING option causes the SET RUN SCRIPT command to ignore certain errors when generating SQL script. This option is supported only in Q Replication. However, the current session is SQL Replication.

User response: Remove the GENERATE SQL FOR EXISTING keyword from the input file.

ASN2411W The object *object_name* of type *object_type* already exists under the specified Q Capture schema. A script to create the object has been generated regardless, because the option to ignore errors while generating the script was chosen.

Explanation: Data already exists in the Q Capture control tables for the publication or publication queue map with this name. The SQL for creating the object has been generated regardless, because the GENERATE SQL FOR EXISTING YES option was set in the SET RUN SCRIPT LATER command. Running this generated SQL against the given control tables will cause SQL errors.

User response: No action is needed.

ASN2412W The object *object_name* of type *object_type* already exists under the specified Q Capture schema and Q Apply schema. A script to create the object has been generated regardless, because the option to ignore errors while generating the script was chosen.

Explanation: Data already exists in the control tables for Q Capture or Q Apply for the subscription or replication queue map with this name. The SQL for creating the object has been generated regardless, because the GENERATE SQL FOR EXISTING YES option was set in the SET RUN SCRIPT LATER command. Running this generated SQL against the given control tables will cause SQL errors.

User response: No action is needed.

ASN2413W The database object *object_name* of type *object_type* already exists at the server *server_alias*. A script to create the object has been generated regardless, because the option to ignore errors while generating the script was chosen.

Explanation: The table, tablespace, or index already exists in the database. The SQL to create this object has been generated regardless, because the GENERATE SQL FOR EXISTING YES option was set in the SET RUN SCRIPT LATER command. Running this generated SQL against the given database will cause SQL errors.

User response: No action is needed.

ASN2414W The value of the *parameter_name* parameter of the WebSphere MQ model queue *host_name-queue_manager_name-queue_name* is too low. The value is set to *value1* but the required value must be at least *value2*.

Explanation: The model queue for the spill queue for

the Q Apply program needs to have the following parameter values:

- Maximum queue depth (MAXDEPTH): 500000 or higher
- Maximum message size (MAXMSGL): 100000 or higher

User response: Update the model queue parameter values, or specify a model queue that has the correct parameter values.

ASN2415E The load type *load_type* is not valid for the Q subscription *name* because the same target table *table_owner.table_name* is participating in another Q subscription *name* that has a conflicting load type of *load_type*.

Explanation: Either the wrong target table name was specified for the Q subscription, or the specified load type is incorrect. Target tables that are populated from multiple source tables (this process is often called data consolidation) require certain load types. If the target table name is correct, then at least one existing Q subscription has a load type that replaces all contents in the table (type 4 or 5). If two or more Q subscriptions were to replace data in the target table, a loss of data could occur.

User response: Take one of the following actions:

- Correct the target table name if the wrong one was specified.
 - Change the selection of the load type to use select and load with insert (104) or select and import with insert (105).
-

ASN2416E The load type *load_type* is supported only for Q subscriptions where the Q Capture server is a Classic server. The command failed.

Explanation: Load types 4 (Select and load with replace), 104 (Select and load with insert), 5 (Select and import with replace), and 105 (Select and import with insert) are only supported where the source is a Classic server.

User response: Choose a different load option and rerun the command.

ASN2417E Use the replication administration tools to start the Q subscription *name*.

Explanation: The target table that was specified for the Q subscription is the same as the target table for one or more other Q subscriptions. The Q Apply program does not start these Q subscriptions automatically because one of them has a load type (4 or 5) that replaces all data in the target table. The Q subscription must be started before the other Q

subscriptions that specify the same target table to ensure no loss of data.

The Q subscription is using one of the following load types:

104

Select and load with insert

5

Select and import with replace

105

Select and import with insert

The Q subscriptions need to be started in a particular order so there is no loss of data. The option to manually start the Q subscription was specified for this Q subscription.

User response: Use the Q Replication Dashboard, ASNCLP command-line program, or Replication Center to start the Q subscription. For details see, "Starting Q subscriptions."

ASN2418E The load type *load_type* is not valid when HAS LOAD PHASE is set to *value*.

Explanation: The HAS LOAD PHASE value was specified as N, indicating that the target table is not loaded. However, a load type value greater than 0 was specified.

User response: Take one of the following actions:

- If you want Q Apply to load the target table, pick a HAS LOAD PHASE value of I.
- If you do not want Q Apply to load the target table, keep the HAS LOAD PHASE value as N and remove the LOAD TYPE option from the CREATE QSUB command.

ASN2419E An unexpected value *value* was found in column *column_name* of replication control table *table_name*.

Explanation: The specified control table contains at least one value that prevents processing of the current action. The table might have been updated with SQL that was not generated by the Replication Center or ASNCLP.

User response: Use one of the replication administration tools to drop and recreate the object with the unexpected value. For example, if the object is a Q subscription, drop and recreate it.

ASN2420I The replication key was defined to include all of the selected target columns because no primary key, unique constraint, or index was found on both the source and the target table within the range of the replicated columns.

Explanation: If no primary key, unique constraint, or index are found on the source or target, the administration tool automatically specifies all valid replicated columns as key columns for replication. Some subscribed columns, such as LOB columns, cannot be used as keys.

User response: This message is for your information only. No action is required.

ASN2421W The column *column_name* is automatically excluded from the Q subscription.

Explanation: Columns defined as ROWID and GENERATED ALWAYS are excluded automatically from bidirectional and peer-to-peer Q subscriptions.

User response: This message is for your information only. No action is required.

ASN2422E The *source|target* database or subsystem that was provided in the SET SERVER command for the Q subscription does not match the *source|target* database or subsystem that is associated with the replication queue map *queue_map_name* for the Q subscription.

Explanation: The source and target servers that are specified for a Q subscription must be the same as the source and target servers for the queue map that used to define the Q subscription.

User response: To determine the name of the server that was specified for the queue map, use the Replication Center or view the CAPTURE_SERVER column in the IBMQREP_RECVQUEUES table or the APPLY_SERVER column in the IBMQREP_SENDQUEUES table. Use this value in the SET SERVER command.

ASN2423E The CREATE QSUB command failed because the source table *table_owner.table_name* that was specified has a primary key but does not have a unique index on the key columns.

Explanation: DB2 for z/OS tables with a primary key must also have a unique index on the primary key columns. Otherwise, the table definition is incomplete and the ASNCLP cannot create a Q subscription for the table.

User response: Create a unique index on the

appropriate primary key columns in the source table, and retry the command.

ASN2424E The option to specify the release level of the replication control tables on server *server_name* is not available because the server is on z/OS or an Oracle source database. The script for creating control tables was not generated.

Explanation: The ability to specify the release, or architecture level, of the Q Capture or Q Apply control tables is only supported on DB2 for Linux, UNIX, and Windows.

User response: Redefine the control tables, making one of the following changes:

ASNCLP command-line program

Omit the RELEASE keyword in the CREATE CONTROL TABLES FOR command.

Replication Center

Do not change the Q Capture release in the Create Q Capture Control Tables wizard or Q Apply release in the Create Q Apply Control Tables wizard.

ASN2425E Generating the script for creating Q Capture or Q Apply control tables failed because the specified release level is older than the release level of the DB2 database. Release level specified: *specified-release-level*. Server name: *server-name*. DB2 database release level: *DB2-release-level*.

Explanation: The release, or architecture level, that is specified for creating Q Capture or Q Apply control tables must match or be newer than the release of the DB2 instance where the control tables are being created.

User response: Redefine the control tables, specifying a release that is the same or newer than the DB2 database.

ASN2426I ASNCLP : The input clause *input_parameter* is ignored on native z/OS platforms.

Explanation: Some input parameters are not needed when the ASNCLP program runs on z/OS platforms with JCL:

- The SET OUTPUT MULTIDIR command for bidirectional and peer-to-peer replication is superseded by DD statements that reference locations where the appropriate source and target SQL statements are written.
- The LOAD MULTIDIR REPL SCRIPT command is superseded by DD statements that reference the

location of the ASNCLP input script for setting up bidirectional or peer-to-peer replication.

- The PASSWORD keyword is not used because user authentication is handled through the communications database (CDB).

User response: This message is for your information only. No action is required.

ASN2427E ASNCLP : The -EXE keyword is not supported on native z/OS platforms.

Explanation: When the ASNCLP program runs on z/OS platforms with JCL, the execute-immediately mode that is indicated by the -EXE keyword is not supported.

User response: Provide the commands in an input file that is executed with JCL.

ASN2428E The database alias *alias* that was retrieved from the CAPTURE_ALIAS column of the IBMQREP_CAPPARMS table does not match the alias *alias* that was specified to connect to the database.

Explanation: The CAPTURE_ALIAS column stores the alias information that was used to connect to the database when the Replication Center or ASNCLP command-line program was used to create Q Capture control tables. However, the alias that is stored in the column does not match the alias that was used for the current database connection. This situation can happen if the database was cataloged with one alias on the system where the tools were used, and later cataloged on a different system by using a different alias.

User response: Catalog the database again and use the same alias that was used when control tables were created.

ASN2429I ASNCLP : The DB2 subsystem *subsystem_ID* is used for the default connection.

Explanation: The stored procedure SYSPROC.ADMIN_INFO_SSID returned the name of the communications database (CDB) to which the ASNCLP first connects. Use this CDB to set up connectivity information for replication source and target servers to which the ASNCLP needs to connect.

User response: This message is for your information only. No action is required.

ASN2430W You mapped a source table column with a CHAR data type to a target table column with the Informix Boolean data type. Ensure that only the following values are inserted or updated into the source column: t, f, and NULL. These are the only allowed values for the

Informix Boolean type.

Explanation: The Informix Boolean column in the target table only accepts the values t, f, and NULL. The t and f are not case sensitive, so T and F are also allowed. Any other values that are inserted or updated into the source column that is mapped to the Boolean column will cause an error at the target.

User response: If you expect to insert or update any values other than t, f, or NULL into the source column, map the column to a target column that does not use the Informix Boolean data type.

ASN2431E The Q subscription *Q_subscription_name* cannot be created because the Q Capture program's compatibility level *compatibility_level* is higher than the Q Apply program's architecture level *arch_level*.

Explanation: An older version of the Q Apply program can only work with a newer version of the Q Capture program if the value of the COMPATIBILITY column in the IBMQREP_CAPPARMS table matches the Q Apply version, which is saved in the ARCH_LEVEL column in the IBMQREP_APPLYPARMS table.

User response: Change the Q Capture COMPATIBILITY value to match the version of the Q Apply program, or upgrade the Q Apply program to the same version as the Q Capture program.

ASN2432E The Q Capture program's architecture level *arch_level* at server *capture_server* and schema *capture_schema* does not match the Q Apply architecture level *arch_level* at server *apply_server* and schema *apply_schema*. The Q subscription cannot be created or started.

Explanation: Either the Q Capture program or Q Apply program is at Version 9.7, and the other program is at an older version. In a multidirectional replication configuration, all of the programs must be at the same version. The ARCH_LEVEL and COMPATIBILITY columns in the IBMQREP_CAPPARMS table must match the value of the ARCH_LEVEL column in the IBMQREP_APPLYPARMS table.

User response: Upgrade the down-level Q Capture or Q Apply program to Version 9.7 or to the equivalent PTF on z/OS, and then create or start the Q subscription.

ASN2433E Q Capture program's architecture level *arch_level* at server *capture_server* and schema *capture_schema* is identical to the Q Apply architecture level *arch_level* at server *apply_server* and schema *apply_schema*. However, the compatibility level *compatibility_level* at the Q Capture server is at a lower level. The Q subscription cannot be created or started.

Explanation: In a multidirectional replication configuration, the values of the ARCH_LEVEL column in the IBMQREP_APPLYPARMS table and the values of the ARCH_LEVEL and COMPATIBILITY columns in the IBMQREP_CAPPARMS table must all match.

User response: Update the value of the COMPATIBILITY column to match the Q Apply and Q Capture ARCH_LEVEL columns, then create or start the Q subscription.

ASN2434I ASNCLP : ASNCLP is unable to retrieve the subsystem identifier (SSID) of the DB2 subsystem that is used for the default connection. SQLCODE *sql_code*.

Explanation: On native z/OS platforms, the ASNCLP calls the SYSPROC.ADMIN_INFO_SSID stored procedure to return the name of the communications database (CDB) that the program uses for connectivity to source and target servers. The stored procedure might not be installed, or the ASNCLP does not have the EXECUTE privilege on the stored procedure.

User response: Ensure that SYSPROC.ADMIN_INFO_SSID is installed and that the user ID that runs the ASNCLP has the correct privileges to execute the procedure.

ASN2435E ASNCLP : The input clause *input_parameter* is not allowed on native z/OS platforms.

Explanation: On native z/OS platforms, the ASNCLP uses the SET BIDI NODE SERVER or SET PEER NODE SERVER commands to specify the paired Q Capture and Q Apply programs at each server for bidirectional or peer-to-peer replication. The SET SERVER MULTIDIR command is not used on native z/OS.

User response: Use the SET BIDI NODE SERVER or SET PEER NODE SERVER command.

ASN2436E The database alias *alias* that was retrieved from the APPLY_ALIAS column of the IBMQREP_APPLYPARMS table does not match the alias *alias* that was specified to connect to the database.

Explanation: The APPLY_ALIAS column stores the alias information that was used to connect to the

database when the Replication Center or ASNCPLP command-line program was used to create Q Apply control tables. However, the alias that is stored in the column does not match the alias that was used for the current database connection. This situation can happen if the database was cataloged with one alias on the system where the tools were used, and later cataloged on a different system by using a different alias.

User response: Catalog the database again and use the same alias that was used when control tables were created.

ASN2437E A Q subscription cannot be created because the Q Capture control tables are at architecture level *arch_level*, which is not supported by the Q Apply architecture level *arch_level* on Linux, UNIX, and Windows, or by Q Apply on z/OS at any architecture level.

Explanation: When the Q Capture program is at Version 9.8 (a value of 0908 in the ARCH_LEVEL column of the IBMQREP_CAPPARMS table), replication is only supported when the Q Apply control tables are in a DB2 for Linux, UNIX, or Windows database that is at Version 9.7 Fix Pack 2 or higher.

User response: Take one of the following actions depending on the location of the Q Apply control tables:

DB2 for Linux, UNIX, and Windows

Upgrade DB2 to V9.7 Fix Pack 2 or later.

DB2 for z/OS

Use a pre-V9.8 Q Capture program in the configuration.

ASN2438E The Q subscription cannot be created because the Q Capture program is at Version *version* and the value of the COMPATIBILITY column in the IBMQREP_CAPPARMS table is *current_value*. The expected value for COMPATIBILITY is *expected_value*.

Explanation: When the Q Capture program is at the specified architecture level (as recorded in the ARCH_LEVEL column of the IBMQREP_CAPPARMS table), the value of the COMPATIBILITY column must be the expected value that is shown in the message text.

User response: Use the Q Replication Dashboard, Replication Center, or SQL to update the value of the COMPATIBILITY column.

ASN2439E The same queue manager name *name* was provided for queue managers that are on two different systems. Queue manager names must be unique across the network. The command failed.

Explanation: WebSphere MQ requires that queue managers that are connected within a network have unique names even when they are on different systems.

User response: Provide a unique queue manager name or let the CREATE MQ SCRIPT command pick a unique queue manager name by not specifying the queue manager name in the QMANAGER keyword.

ASN2440E The Replication Center cannot obtain the DB2 for z/OS subsystem ID because the SYSPROC.ADMIN_INFO_SSID stored procedure is not installed.

Explanation: The Replication Center uses the subsystem ID to connect to DB2 for z/OS to create control tables or perform other tasks. The stored procedure for providing the subsystem ID is SYSPROC.ADMIN_INFO_SSID. On DB2 for z/OS Version 10, SYSPROC.ADMIN_INFO_SSID is installed by default. On DB2 for z/OS version 8 and Version 9, the job to install the SYSPROC.ADMIN_INFO_SSID stored procedure is DSNITJSG.

User response: Install the SYSPROC.ADMIN_INFO_SSID stored procedure at the source or target subsystem, and then retry the action.

ASN2441E The CREATE QSUB command does not support the option REPLICATE ADD COLUMN if the source or target DB2 subsystem is earlier than Version 10.

Explanation: The REPLICATE ADD COLUMN function is valid only if both the source and target are at DB2 for z/OS Version 10 or later.

User response: Either specify both a source and target subsystem that are at Version 10 or later, or do not specify the REPLICATE ADD COLUMN keywords.

ASN2442E The action *action_name* cannot be completed because the required information for the bidirectional or peer-to-peer node *node_number* was not provided in an earlier command.

Explanation: The CREATE CONTROL TABLES or CREATE REPLQMAP commands contain a reference to a node number. The node number is used to identify the servers that are involved in bidirectional or peer-to-peer replication. However, the node information was not set in a prior SET BIDI NODE or SET PEER NODE command.

User response: Issue the SET BIDI NODE or SET PEER NODE command to identify the server with a

node number before specifying the node number in the CREATE CONTROL TABLES and CREATE REPLQMAP commands.

ASN2443W The target table *table_owner.table_name* contains a BUSINESS_TIME WITHOUT OVERLAPS (BTWO) clause in its primary key or unique constraint specification, but the source table *table_owner.table_name* does not have the same option. The Q Apply program might encounter errors.

Explanation: The target is defined with a BUSINESS_TIME period clause that uses the WITHOUT OVERLAPS option. With this option, the BUSINESS_TIME period value for a matching key cannot be overlapped. If the source table contains overlapped BUSINESS_TIME values, the Q Apply program will receive SQL errors when it applies these changes to the target table.

User response: Drop the primary key or unique constraint on the target table and redefine it without the WITHOUT OVERLAPS option.

ASN2444E The source table *table_owner.table_name* is not a system period temporal table or bitemporal table, nor is versioning enabled for the table. A Q subscription or subscription-set member will not be created for the history table.

Explanation: The attempt to create a Q subscription or subscription-set member for the history table failed because the source table is not a system period temporary table, bitemporal table, or a table with versioning enabled. Therefore, the source table does not have a history table.

User response: Take one of the following actions:

- Alter the source table by specifying the ADD PERIOD SYSTEM_TIME and ADD VERSIONING USE HISTORY TABLE clauses, and then reissue the CREATE QSUB or CREATE MEMBER command.
- Do not specify INCLUDE HISTORY TABLE when creating the Q subscription or subscription-set member.

ASN2445E The attempt to create a Q subscription or subscription-set member for the source table *table_owner.table_name* failed. The source table has a SYSTEM_TIME or BUSINESS_TIME period defined, but the target table *table_owner.table_name* does not have a SYSTEM_TIME or BUSINESS_TIME period defined.

Explanation: A Q subscription or subscription-set member cannot be created unless both the source and target table are correctly defined. In order to define a

temporal table on the target, either a SYSTEM_TIME or BUSINESS_TIME period must be defined.

User response: Define a SYSTEM_TIME or BUSINESS_TIME period on the target table by using the ALTER TABLE statement, and then reissue the CREATE QSUB or CREATE MEMBER command.

ASN2446W No period *period_name* exists on the source table *table_owner.table_name*. The PERIOD keyword that was used in the CREATE QSUB or CREATE MEMBER command is ignored.

Explanation: You can only use the PERIOD keyword when you are subscribing a table that has a SYSTEM_TIME or BUSINESS_TIME period defined.

User response: Do not specify the PERIOD clause when creating a subscription, or define a period on the source table and then reissue the command.

ASN2447W The target temporal table *table_owner.table_name* has versioning enabled even though the source history table *table_owner.table_name* is part of a Q subscription or subscription-set member. SQL conflicts might occur when rows are applied to the target table.

Explanation: The versioning attribute on the temporal table at the target causes DB2 to update the table in addition to updates that are made by the Q Apply or Apply program. This can cause SQL conflicts at the table.

User response: Either remove the subscription on the history table or perform an ALTER TABLE to remove versioning on the target temporal table.

ASN2448E The ASNCLP program could not create a history table for the target table *table_owner.table_name* because no name was specified for the history table in the CREATE QSUB or CREATE MEMBER command.

Explanation: To create a Q subscription or subscription-set member for a temporal table, the target table must either already have a history table, or you must provide a name for the history table so the ASNCLP program can create the history table.

User response: Specify a name that can be used to create the target history table or create a history table for the target table. Then reissue the command.

ASN2449E An error occurred while trying to generate the DataStage definition (.dsx) files that create jobs to read from a consistent-change data (CCD) table. The files were not created. Reason code: *reason_code*.

Explanation: One of the following problems, identified by a reason code, caused the generation of the .dsx files to fail:

1

The IBMSNAP_FEEDETL control table does not exist at the Apply control server. This table is required to identify SQL Replication subscription sets with at least one CCD table member that is read by IBM InfoSphere DataStage.

2

The subscription set does not include any member CCD tables that are noncondensed and noncomplete.

User response: Review the reason codes in the explanation, and take the appropriate action:

1

Create the IBMSNAP_FEEDETL control table manually by using the sample in the SQLLIB/samples/repl/sql directory and regenerate the .dsx file.

2

Ensure that the subscription set has at least one member CCD table that is noncondensed and noncomplete.

ASN2450E Stored procedure target are not supported when either the Q Capture control tables or Q Apply control tables are at *arch_level* architecture level.

Explanation: When the version of the control tables at the source or target is 0908 (Version 9.8), replication to stored procedure targets is not supported.

User response: Select a different target type.

ASN2451E The ERROR ACTION 'B' is not supported when the target table is on Linux, UNIX, or Windows.

Explanation: The value 'B' for the ERROR ACTION keywords in the CREATE QSUB command prompts the Q Apply program to spill change messages for the Q subscription to a temporary spill queue when an error occurs. ERROR ACTION 'B' is only supported for z/OS target tables.

User response: Choose a different value for ERROR ACTION and rerun the CREATE QSUB command.

ASN2458E The value B for the ERROR ACTION keyword in the CREATE QSUB command is not supported when the source or target table has referential integrity constraints.

Explanation: You can specify B for ERROR ACTION to prompt the Q Apply program to put change messages for the Q subscription in a temporary spill queue when an SQL error occurs and until the error is resolved. However, the B value is not supported when the source or target table for the Q subscription is involved in a referential integrity relationship with other tables.

User response: Specify a different value for ERROR ACTION and resubmit the CREATE QSUB command.

ASN2459E The target nickname *nickname_owner.nickname_name* references a table on the remote server *remote_server* but the Q Apply control table nicknames reference control tables on a different remote server *remote_server*. The target nickname and Q Apply control table nicknames must point to the same remote server.

Explanation: The target nickname and the Q Apply control table nicknames must reference tables on the same remote server because the Q Apply program updates the nickname target and some control tables in the same commit scope and the federated wrapper does not support two-phase commit.

User response: Take one of the following actions and then retry the CREATE QSUB command:

- Create a target nickname that points to the same remote server as the Q Apply control tables.
- Let the ASNCPL program create the nickname by not specifying the EXIST clause.

ASN2460E A Q subscription could not be created for source table *table_owner.table_name* because the LOAD from CURSOR option was specified but no nickname was provided. The LOAD from CURSOR option that uses a cataloged DB2 alias instead of a nickname is not available because the Q subscription includes one or more XML columns.

Explanation: The Q subscription specifies that the Q Apply program should load the target table (automatic load) by using the LOAD from CURSOR utility. However, this utility cannot be used.

This message is returned when the target database is at Version 9.7 Fix Pack 4 or later and no nickname is provided.

In DB2 pureScale environments only:

Because the DB2 pureScale Feature is not supported in a federated environment, the LOAD from CURSOR utility can use an alias instead of a nickname in a DB2 pureScale environment. However, the LOAD from CURSOR utility cannot use an alias if the Q subscription includes one or more XML columns. This message is returned in a DB2 pureScale environment when no nickname is provided and there are one or more XML columns in the Q subscription.

User response: Recreate the Q subscription, specifying a nickname for the LOAD FROM CURSOR option, or select a different load option such as "Best available: Let the Q Apply program choose a load option" in the Replication Center (LOAD TYPE 0 in the ASNCLP command-line program) or the EXPORT and IMPORT combination of utilities.

In DB2 pureScale environments only:

Recreate the Q subscription, specifying a different load option such as "Best available: Let the Q Apply program choose a load option" in the Replication Center (LOAD TYPE 0 in the ASNCLP command-line program) or the EXPORT and IMPORT combination of utilities.

ASN2461E **LOAD from CURSOR was chosen for the Q subscription (LOAD TYPE 1) but the keywords to specify a nickname for the source table were not provided in the CREATE QSUB command. The source server *server_name* is at version *version* and this version requires a source nickname for LOAD from CURSOR.**

Explanation: At the specified server level, the LOAD from CURSOR option requires a nickname for the source table. Nicknames are not required for Version 9.7 Fix Pack 4 and later.

User response: Use the NICKNAME or NEW NICKNAME RMT SERVERNAME keywords to specify an existing or new nickname and reissue the CREATE QSUB command.

ASN2462E **The CREATE QSUB command specifies to create multiple Q subscriptions that use the LOAD from CURSOR option with a nickname for the source table. However, only a single nickname *nickname_owner.nickname_name* was specified for the source table.**

Explanation: When you create multiple Q subscriptions, it is recommended that you use the NAMING PREFIX option with the NEW NICKNAME RMT SERVERNAME keywords so that the ASNCLP program can generate a unique nickname for each source table.

User response: Modify the CREATE QSUB command to use the NAMING PREFIX option and reissue the command.

ASN2463E **The *command* command is not supported because one or more of the servers in your configuration is a DB2 for z/OS subsystem or pre-Version 10 DB2 for Linux, UNIX, and Windows database.**

Explanation: The provided command is only supported if all the databases that are involved in the configuration are on DB2 for Linux, UNIX, and Windows Version 10 or higher.

User response: Reissue the command on a DB2 for Linux, UNIX, and Windows database at Version 10 or higher.

ASN2464E **A schema-level Q subscription cannot be created for schema *schema_name* because a schema-level Q subscription already exists for another schema *existing_schema* that uses the same queue map *queue_map_name*, and these two schemas overlap.**

Explanation: The expression that was used to select schemas for the schema-level Q subscription overlaps or conflicts with the schema of an existing schema-level Q subscription. Overlapping schemas are not supported. For example, if a schema-level Q subscription is created for schema pattern "ANU%".T%, you cannot create another schema-level Q subscription for pattern "ANU1".T%" that uses the same replication queue map.

User response: Provide a different expression that does not overlap with existing schema-level Q subscriptions or drop the existing schema-level Q subscription and reuse this expression.

ASN2465E **A schema-level Q subscription of configuration type *configuration_type* cannot be created for schema *schema* because table-level Q subscriptions for some tables that belong to this schema already exist under a different configuration type *configuration_type* and use the same replication queue map *queue_map_name*.**

Explanation: Q subscriptions that belong to a schema-level Q subscription and use the same replication queue map should all be of the same configuration type (unidirectional, bidirectional, or peer-to-peer).

User response: Drop the existing Q subscription before creating the schema-level Q subscription.

ASN2466E A schema-level Q subscription of configuration type *configuration_type* cannot be created for schema *schema* because table-level Q subscriptions for some tables that belong to this schema have the same configuration type and use the same replication queue map *queue_map_name*, but have different values for properties such as conflict action and conflict rule.

Explanation: Table-level Q subscriptions that belong to the same schema-level Q subscription, have the same configuration type (unidirectional, bidirectional, or peer-to-peer), and use the same replication queue map must have the same conflict action and conflict rule.

User response: Drop the existing Q subscriptions before creating the schema-level Q subscription.

ASN2467E The name *name* that was provided to identify the list of Q subscription options for Q subscription type *type* cannot be used because another list of Q subscription options with the same name already exists for Q subscription type *type* at server *server_name* under Q Capture schema *schema*.

Explanation: The CREATE SUBSCRIPTION OPTIONS command sets options for all Q subscriptions of a given configuration type (unidirectional, bidirectional, or peer-to-peer) within a given schema. The name for the list of options must be unique for a Q Capture schema.

User response: Provide a different name for the list of options and reissue the command.

ASN2468W The table *table_owner.table_name* matches the provided schema pattern *schema_pattern* and source table pattern *table_pattern* but is already part of a Q subscription. The SCHEMA_SUBNAME column for this Q subscription will be updated to *schema_subname*.

Explanation: A table-level Q subscription cannot be created for the provided table because one already exists. But the SCHEMA_SUBNAME column in the IBMQREP_SUBS and IBMQREP_TARGETS tables will be updated to the provided value to make this Q subscription part of the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN2469W The option TARGET EXISTS VALIDATION NO was chosen and so the ASNCLP program will assume that the source and target tables are identical and will not perform any checks on source and target tables.

Explanation: When the option TARGET EXISTS VALIDATION NO is chosen, the ASNCLP program does not perform any checks such as whether the target table exists, whether the column mappings are compatible between the source and target tables, and whether constraints match between source and target tables. The ASNCLP program will only generate the insert statements into the control tables to create the Q subscription.

User response: To avoid problems with the Q Capture and Q Apply programs, make sure that the target table exists and that the source and target tables are identical before using this option.

ASN2470E The options name *options_name* that was specified in the OPTIONS clause in the CREATE SCHEMASUB or CREATE QSUB command does not exist.

Explanation: The options name identifies a set of options for all table-level Q subscriptions within a schema-level Q subscription. A row that is identified by the options name should exist in the IBMQREP_SUBS_PROF table before table-level Q subscriptions can be created using the specified attributes.

User response: Issue a CREATE SUBSCRIPTION OPTIONS command to first create the options before using the options name in the CREATE SCHEMASUB or CREATE QSUB command. If the CREATE SUBSCRIPTION OPTIONS command is already specified in the input file, make sure that the SET RUN SCRIPT NOW command is used to run the generated script immediately.

ASN2471E The type *type* of the Q subscription that was specified in the CREATE SUBSCRIPTION OPTIONS command with the options name *options_name* does not match the Q subscription type *type* that was specified in the CREATE SCHEMASUB or CREATE QSUB command.

Explanation: The Q subscription types (unidirectional, bidirectional, or peer-to-peer) should match between the CREATE SUBSCRIPTION OPTIONS command and the CREATE SCHEMASUB or CREATE QSUB commands that use the options. Otherwise, the Q subscription properties might be incompatible.

User response: Provide another options name that has the same Q subscription type that was specified in the CREATE SCHEMASUB or CREATE QSUB command. You can also omit the OPTIONS clause and the ASNCLP program will pick the product defaults for that Q subscription type.

ASN2472I The START SCHEMASUB command was processed for schema-level Q subscription *name*. A CAPSTART signal will be inserted into the *schema.IBMQREP_SIGNAL* table for *number* Q subscriptions that are part of the schema-level Q subscription.

Explanation: The START SCHEMASUB command was issued with the ALL option. All Q subscriptions that belong to the schema-level Q subscription in new and inactive state are activated.

User response: This message is for your information only. No action is required.

ASN2473I The START SCHEMASUB command was processed for schema-level Q subscription *name*. A CAPSTART signal was inserted into the *schema.IBMQREP_SIGNAL* table for the schema-level Q subscription.

Explanation: A CAPSTART signal was inserted to activate the specified schema-level Q subscription. The Q Capture program will start capturing changes for source tables that are part of the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN2474I The START SCHEMASUB command was processed for schema-level Q subscription *name*. A CAPSTART signal was not inserted into the *schema.IBMQREP_SIGNAL* table for the schema-level Q subscription because the Q Capture program is already capturing changes for source tables that are part of the schema-level Q subscription.

Explanation: The schema-level Q subscription is already in active state. No script is generated to start the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN2475I The STOP SCHEMASUB command was processed for schema-level Q subscription *name*. A CAPSTOP signal was inserted into the *schema.IBMQREP_SIGNAL* table for *number* table-level Q subscriptions that belong to the schema-level Q subscription.

Explanation: The STOP SCHEMASUB command was issued with the ALL option. All active Q subscriptions that belong to the schema-level Q subscription are deactivated.

User response: This message is for your information only. No action is required.

ASN2476I The STOP SCHEMASUB command was processed for schema-level Q subscription *name*. A CAPSTOP signal was inserted into the *schema.IBMQREP_SIGNAL* table for the schema-level Q subscription.

Explanation: In response to the STOP SCHEMASUB command, the ASNCLP program generated a SQL script to insert a CAPSTOP signal. The signal prompts the Q Capture program to stop capturing changes for all source tables within the schema.

User response: This message is for your information only. No action is required.

ASN2477I The STOP SCHEMASUB command was processed for schema-level Q subscription *name*. However, a CAPSTOP signal was not inserted into the *schema.IBMQREP_SIGNAL* table because the schema-level Q subscription is already in inactive state.

Explanation: Because the schema-level Q subscription is already inactive, the ASNCLP program did not generate a script to stop the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN2478E A schema-level Q subscription with name *name* does not exist at the server *server* under schema *schema*. No script is generated.

Explanation: The schema-level Q subscription that was specified in the DROP SCHEMASUB command does not exist in the IBMQREP_SCHEMASUBS table and cannot be deleted.

User response: Provide an existing schema-level Q subscription name and reissue the command.

ASN2479I *Number* table-level Q subscriptions that belong to the schema-level Q subscription *name* will be deleted.

Explanation: The DROP SCHEMASUB command was issued with the ALL option. All table-level Q subscriptions that belong to the schema-level Q subscription will be deleted.

User response: This message is for your information only. No action is required.

ASN2480E The provided list of Q subscription options *options_list* does not exist at server *server_name* under Q Capture schema *schema* and cannot be deleted.

Explanation: The IBMQREP_SUBS_PROF table stores the user-defined options for all Q subscriptions of a given type within a schema-level Q subscription. A row with the provided options name cannot be found in the IBMQREP_SUBS_PROF table. No SQL script is generated to delete the options list.

User response: Provide a different name for the options list and rerun the DROP SUBSCRIPTION OPTIONS command.

ASN2481E The set of Q subscription options *options_name* cannot be deleted because one or more schema-level Q subscriptions are using the options.

Explanation: The Q Capture and Q Apply programs refer to the saved list of Q subscription options to create new Q subscriptions for schema-level Q subscriptions that use this options name. If the options list is deleted, errors will occur for the running replication programs.

User response: Provide the name of a different set of Q subscription options to delete, or delete any schema-level Q subscriptions that are using the options and then delete the options.

ASN2482I The MQDEFAULTS option to pick default values for WebSphere MQ objects was chosen in the CREATE CONTROL TABLES command. The ASNCLP program will assign the following defaults: queue manager: *queue_manager*; administration queue: *admin_queue*; restart queue: *restart_queue*.

Explanation: When the MQDEFAULTS option is specified, the ASNCLP program creates control tables with the same defaults for WebSphere MQ objects that are generated by the CREATE MQ SCRIPT command.

User response: This message is for your information only. No action is required.

ASN2483I The MQDEFAULTS option to pick default values for WebSphere MQ objects was used in the CREATE REPLQMAP command. The ASNCLP program will assign the following defaults: administration queue: *admin_queue*; receive queue: *recv_queue*; send queue: *send_queue*.

Explanation: When the MQDEFAULTS option is specified, the ASNCLP program creates a replication queue map with the same defaults for WebSphere MQ

objects as the defaults that are generated by the CREATE MQ SCRIPT command.

User response: This message is for your information only. No action is required.

ASN2484W The option to drop control tables on the server that is specified by node number *number* was chosen, but the control table *table_name* does not exist at server *server_name* under the specified schema *schema*.

Explanation: The command specifies that the ASNCLP program should drop both Q Capture and Q Apply control tables at the server that is identified by the node number. However, the action cannot be carried out because both Q Capture and Q Apply control tables do not exist.

User response: Provide a different schema that has both Q Capture and Q Apply control tables or use the regular DROP CONTROL TABLES command to individually drop either the Q Capture or Q Apply control tables at the server.

ASN2485E The keyword *keyword* is not supported if either the source server or target server is not on DB2 for Linux, UNIX, and Windows Version 10 or higher.

Explanation: The option is only supported on DB2 for Linux, UNIX, and Windows V10 or higher.

User response: Remove the option and reissue the command.

ASN2486E The action *action_name* cannot be completed because the required information for node *node_number* was not provided.

Explanation: Node numbers are used to identify the servers in bidirectional replication. A reference to a node number was made in the CREATE CONTROL TABLES FOR or CREATE REPLQMAP command but the node information was not set in a prior SET BIDI NODE command.

User response: Issue the SET BIDI NODE command before you use the node number in a CREATE CONTROL TABLES FOR or CREATE REPLQMAP command.

ASN2487I The queue manager for MQSERVER *server_number* at host *hostname* and port *port_number* is local to the system where the ASNCLP program runs. The ASNCLP program will run the WebSphere MQ setup script for this server.

Explanation: If you use the RUN NOW option with

the CREATE MQ SCRIPT command and the ASNCLP program is running on the same system as the queue manager, the ASNCLP generates the script and then runs the script. If the queue manager is on a different system, the ASNCLP does not run the script even if RUN NOW is specified.

User response: This message is for your information only. No action is required.

ASN2488W The queue manager for MQSERVER *server_number* at host *hostname* and port *port_number* is remote from the system where the ASNCLP program will not run the WebSphere MQ setup script for this server even though the RUN NOW option is specified.

Explanation: If you use the RUN NOW option with the CREATE MQ SCRIPT command and the ASNCLP program is running on a different system than the queue manager, the ASNCLP generates the script but does not run the script. If the queue manager is on the same system, the ASNCLP runs the script when RUN NOW is specified.

User response: Run the generated WebSphere MQ script by using the following command:

```
runmqsc <queue_manager_name> <
    <filepath>/<script_file_name>
```

ASN2489E The option WITH LOGMARKER is specified in the CREATE QSUB command and the target table or nickname *owner.name* exists but does not have the IBMSNAP_LOGMARKER column.

Explanation: The target table or nickname must have the IBMSNAP_LOGMARKER column to use the WITH LOGMARKER option for the Q subscription.

User response: Perform one of the following actions and rerun the command:

- Alter the target table to add the IBMSNAP_LOGMARKER column with a data type of TIMESTAMP and NOT NULL WITH DEFAULT.
- Use another target table that has the IBMSNAP_LOGMARKER column.

ASN2490E The target table or nickname *owner.name* contains the column IBMSNAP_LOGMARKER but the data type of the column *data_type* does not match the required data type *data_type*.

Explanation: The data type of the column IBMSNAP_LOGMARKER must match the Q Replication requirements or the Q Apply program does not activate the Q subscription.

User response: Perform one of the following actions and rerun the command:

- Use a different target with a valid definition of the IBMSNAP_LOGMARKER column.
- Use a new target table.

ASN2491E The option WITH LOGMARKER is not supported because either the Q Apply server is on DB2 for z/OS and the control tables are not at Version 10 or higher (ARCH_LEVEL of 100Z) or on DB2 for Linux, UNIX, and Windows and not at Version 9.7 Fix Pack 4 or higher.

Explanation: The WITH LOGMARKER option of the CREATE QSUB command adds the timestamp of when the source row was changed to the target table or nickname. This option is only supported in the following cases:

- The Q Apply server on Linux, UNIX, and Windows at Version 9.7 FP4 or higher.
- The Q Apply server is on z/OS and the architecture level is 100Z or higher, which corresponds to Version 10, and the PTF that corresponds to Version 9.7 Fix Pack 4 is installed.

User response: Upgrade the Q Apply server to the version and PTF or fix pack that supports this option, and resubmit the CREATE QSUB command.

ASN2494E A schema-level Q subscription with name *name* already exists at the server *server_name* under schema *schema*.

Explanation: The name that was provided with the CREATE SCHEMASUB command already exists in the IBMQREP_SCHEMASUBS table. A duplicate schema-level Q subscription cannot be created.

User response: Provide a different name and reissue the CREATE SCHEMASUB command.

ASN2496E Target consistent-change data (CCD) tables have not been upgraded to 16-byte log sequence number (LSN) columns that are required when the target server is at DB2 version *version*.

Explanation: When the source server is at DB2 Version 10, target CCD tables require 16-byte LSN columns (IBMSNAP_COMMITSEQ, IBMSNAP_INTENTSEQ) to hold log sequence numbers.

User response: Follow the steps in the Version 10 migration information to upgrade target CCD tables to 16-byte LSN columns and and restart the Q Apply program.

ASN2497E **The WITH LOGMARKER option cannot be used because the source table *table.owner_table.name* has an IBMSNAP_LOGMARKER column.**

Explanation: You can populate an IBMSNAP_LOGMARKER column in a target table in two different but mutually exclusive ways:

- By specifying the WITH LOGMARKER keywords in the CREATE QSUB command. With this method, the IBMSNAP_LOGMARKER column is added to the target table and is populated with timestamp values taken directly from the source recovery log.
- By including an IBMSNAP_LOGMARKER column in the source table and mapping the column to a matching IBMSNAP_LOGMARKER column in the target table. With this method, the target column is populated with values from the source column. This method can be used in a three-tier configuration where the IBMSNAP_LOGMARKER column at Tier 2 is populated from the recovery log at Tier 1, and in turn its values are replicated to populate the column at Tier 3.

If you have a source column named IBMSNAP_LOGMARKER, you cannot use the WITH LOGMARKER keywords.

User response: Reissue the CREATE QSUB command, omitting the WITH LOGMARKER option. Map the IBMSNAP_LOGMARKER column in the source table to the matching target column. This method ensures that values in the IBMSNAP_LOGMARKER column are from the source table and not from the log.

ASN2498E **The action *action* failed. The architecture level *arch_level* of the Apply control tables is not compatible with the Capture compatibility value *compatibility*.**

Explanation: The source server is at DB2 Version 10, which uses 16-byte log sequence numbers (LSN). To replicate data from Version 10, the Apply program must work with control tables that are migrated to Version 10 and 16-byte LSN values.

User response: Migrate the Apply control server to Version 10.

ASN2499E **The action *action* failed. The source and target consistent-change data (CCD) tables are incompatible.**

Explanation: Both source and target CCD tables must have either 10-byte or 16-byte columns to hold log sequence numbers (LSN).

User response: Upgrade the target CCD table to 16-byte LSN columns if the source CCD table uses 16-byte LSN columns, or upgrade the source CCD table to 16 byte-LSN columns. The source CCD table can be either an external CCD table that you create or a CCD

table that is maintained by an Apply program or Q Apply program, in which case migration of these servers is also required.

Chapter 40. ASN2500 - ASN2999

ASN2500E **The TYPE CCD option cannot be used because the source table *table.owner_table.name* has the following IBMSNAP_% columns: *columns*.**

Explanation: You can populate the IBMSNAP_% columns in a target table in two different but mutually exclusive ways:

- By specifying the TYPE CCD keyword in the CREATE QSUB command. With this method, the target type is set to 2 and the IBMSNAP_% columns in the target table are populated with values taken directly from the source recovery log.
- By including the IBMSNAP_% columns in the source table and mapping the column to a matching IBMSNAP_% column in the target table. With this method, the target column is populated with values from the source column. This method can be used in a three-tier configuration where the IBMSNAP_% column at Tier 2 is populated from the recovery log at Tier 1, and in turn its values are replicated to populate the column at Tier 3.

If you have a source column named one of the IBMSNAP_% columns, you cannot use the TYPE CCD keyword.

User response: Reissue the CREATE QSUB command, omitting the TYPE CCD option to indicate that the target is a regular user table. Map the IBMSNAP_% columns in the source table to the matching target column. This method ensures that the IBMSNAP_% column is from the source table and not from the log.

ASN2503E **The TARGET keyword is not supported in the ALTER ADD COLUMN command because of the DB2 version, architecture level, or compatibility settings of the source or target server. Reason code *reason_code*.**

Explanation: The option to specify target column names with the ALTER ADD COLUMN command is allowed only in the following cases:

Linux, UNIX, Windows

Source or target servers are on DB2 Version 9.7 Fix Pack 5 or higher and the value of the COMPATIBILITY column in the IBMQREP_CAPPARMS table is 0907 or higher.

z/OS

The architecture level of the Q Capture or Q Apply control tables (ARCH_LEVEL column in the IBMQREP_CAPPARMS or

IBMQREP_APPLYPARMS table) is 0907 or higher and the value of the COMPATIBILITY column in the IBMQREP_CAPPARMS table is 0907 or higher.

The following values are valid for the reason code:

0

The Q Capture server is on DB2 for Linux, UNIX, and Windows V9.7 FP5 or higher but COMPATIBILITY is lower than 0907.

1

The Q Capture server is on DB2 for Linux, UNIX, and Windows V9.7 FP4 or lower.

2

The Q Capture server is on DB2 for z/OS and ARCH_LEVEL is lower than 0907.

3

The Q Capture server is on DB2 for z/OS but COMPATIBILITY is lower than 0907.

4

The Q Apply server is on DB2 for Linux, UNIX, and Windows V9.7 FP4 or lower.

5

The Q Apply server is on DB2 for z/OS and ARCH_LEVEL is lower than 0907.

User response: Follow these steps:

1. Upgrade the Q Capture server, Q Apply server or both to either DB2 for Linux, UNIX, and Windows V9.7 FP5 or higher with the required compatibility setting or DB2 for z/OS with the required compatibility setting and architecture level.
2. Rerun the command.

ASN2506E **A schema-level Q subscription *subscription_name* that uses replication queue map *queue_map_name* cannot be created because the expression *expression* that was provided for the schemas or tables to include or exclude from the Q subscription is invalid.**

Explanation: You can use the percentage sign (%) as a wild card to specify schema-level Q subscriptions and the Q Capture program automatically creates a Q subscription for all tables within the schemas with names that match the wild-card expression. However the wild card can only be specified as a suffix when you use the following keywords:

- OWNER LIKE
- NAME LIKE
- EXCLUDE NAME

Wild cards that are used as a prefix (for example "%AN" or "%AN%") are not supported, and no wild cards are allowed with the EXCLUDE OWNER keywords.

User response: Change to an allowed wild-card expression and reissue the CREATE SCHEMASUB command.

ASN2508I **The REINIT SCHEMASUB command was processed for the schema-level Q subscription *name*.**

Explanation: The REINIT SCHEMASUB command prompts the ASNCPL program to generate a SQL script to insert a REINIT_SCHEMASUB signal into the IBMQREP_SIGNAL table. The signal causes the Q Capture program to reread the saved options for the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN2509I **The REINIT SCHEMASUB command was processed for schema-level Q subscription *name* but a script to insert a REINIT_SCHEMASUB signal into the *schema*.IBMQREP_SIGNAL table was not generated because the schema-level Q subscription is inactive.**

Explanation: Schema-level Q subscriptions must be in active (A) state for a REINIT_SCHEMASUB signal to be processed.

User response: Issue the START SCHEMASUB command and then issue the REINIT SCHEMASUB command.

ASN2510E **A schema-level Q subscription with name *name* does not exist at the server *server* under schema *schema* and so could not be deleted.**

Explanation: The attempt to delete a schema-level Q subscription failed because the specified Q subscription name does not exist in the IBMQREP_SCHEMASUBS table.

User response: Provide an existing schema-level Q subscription name and rerun the action.

ASN2511E **The column *column_name* of data type *data_type* in target table *table_owner.table_name* cannot be populated with the timestamp of when the source row was changed because the column data type does not match the required data type *required_data_type*.**

Explanation: The target table must have a timestamp column in order for the table to serve as a point-in-time table. This type of table stores the time that each row change was committed at the source database.

User response: Redo the action to create the Q subscription after making one of these changes:

- Use a different column in the target table that is defined with the TIMESTAMP data type.
- Use a different target table that includes a timestamp column.

ASN2512E **The column *column_name* in target table *table_owner.table_name* cannot be populated with the timestamp of when the source row was changed because the target column is already mapped to a source column *column_name* in table *table_owner.table_name*.**

Explanation: The specified column in the target table is already mapped to a source column, so you cannot specify this column to be populated with the timestamp of when the source row was changed. The timestamp is taken from the recovery log at the source database.

User response: Redo the action to create the Q subscription after making one of these changes:

- Use a different eligible column in the target table to map to the source recovery log.
- Remove the source column mapping:
 - Use the TRGCOLS EXCLUDE option to exclude the target column from any direct mapping to a source column.
 - If you are already using the TRGCOLS INCLUDE option with the column, then remove the column from the list of target columns to include.

ASN2513E **The target table or nickname *owner.name* cannot be used as a point-in-time table or nickname because it does not contain any timestamp columns.**

Explanation: A point-in-time table or nickname includes an additional column that is populated with the timestamp of when the source row was changed. To use this option, the target table or nickname must have at least one column that is defined with the TIMESTAMP data type.

User response: Perform one of the following actions:

- In the Replication Center, deselect the **Point-in-time** option on the Source and Target page of the Q Subscription Properties notebook for this table or nickname.
- Alter the target table to include a timestamp column and then redo the action.
- Do not create a Q subscription that includes this target table or nickname.

ASN2514E The target table or nickname *owner.name* does not have any eligible column to store the timestamp of when the source row changed. All of the timestamp columns in the table or nickname are already mapped to source table columns.

Explanation: A point-in-time table or nickname includes an additional column that is populated with the timestamp of when the source row was changed. To use this option, the target table or nickname must have at least one timestamp column that is not mapped to a source column.

User response: Redo the action to create the Q subscription after making one of these changes:

- In the Replication Center, clear the **Point-in-time** check box on the Source and Target tab of the Q Subscription Properties notebook.
- Make one of the timestamp columns in the target table columns eligible for the point-in-time mapping by following these steps:
 1. On the Column Mapping page, remove the source-target mapping for a timestamp column in the target table.
 2. Select the target column and click **Map to LOGMARKER** to map the column to the value from the source recovery log.

ASN2515W An existing nickname could not be found on the federated database *database_name* for the remote table *table_owner.table_name*. A new nickname *owner.name* will be created.

Explanation: When you create multiple Q subscriptions for federated targets, the Replication Center uses the target object profile that you created to determine whether to use an existing nickname or create a new nickname. The option to use an existing nickname was selected in the profile, but the Replication Center did not find an existing nickname that points to the table on the non-IBM server. A new nickname will be created for the remote table.

User response: This message is for your information only. No action is required.

ASN2516E The nickname *owner.name* exists on the federated database *database_name*, but the nickname does not point to the remote table *table_owner.table_name* that was specified as the target table for the Q subscription.

Explanation: When you create multiple Q subscriptions for federated targets, the Replication Center uses the target object profile that you created to determine whether to use an existing nickname or create a new nickname. The option to use an existing nickname was selected in the profile and a nickname exists on the federated database that matches the profile. However, this nickname does not point to the remote target table.

User response: Take one of the following actions:

- On the Target Tables page of the Create Q Subscriptions wizard, click **Change**, and then on the Target Object Profiles notebook, clear the **Use existing nicknames** checkbox. The Replication Center then creates a nickname that points to the remote target table.
- Create a nickname that points to the remote target table, and use this existing nickname for the Q subscription.

ASN2517E More than one nickname points to the remote table *table_owner.table_name* and so the Replication Center cannot determine which nickname to use.

Explanation: When you create multiple Q subscriptions for federated targets, the Replication Center uses the target object profile that you created to determine whether to use an existing nickname or create a new nickname. The option to use an existing nickname was chosen in the profile, but you must use the Q Subscription Properties notebook to specify which nickname to use.

User response: Follow these steps:

1. On the Review Q Subscriptions page of the wizard, select the Q subscription that specifies this remote table and click **Properties**.
2. On the Source and Target tab of the Q Subscription Properties notebook, select a nickname from the **Existing nicknames** list.

ASN2521E The compatibility value *value* is not supported because the Q Capture control tables are at Version 10.1.

Explanation: The Q Capture compatibility parameter specifies the level of messages that are sent to the Q Apply program. When the Q Capture program is at Version 10.1 on Linux, UNIX, and Windows, the supported values for the compatibility parameter are 1001 and 0907.

User response: Use the Q Replication Dashboard or SQL to set the value in the COMPATABILITY column of the IBMQREP_CAPPARMS table to 1001 if the Q Apply program is at Version 10.1 on Linux, UNIX, or Windows, or 0907 if the Q Apply program is at Version 9.7. Then reissue the ASNCLP command.

ASN2523E You cannot use the Replication Center on Oracle source databases to perform operational tasks such as starting or stopping replication and event publishing programs or managing send or receive queues.

Explanation: The Replication Center function to run commands that operate on replication or publishing programs is not supported on Oracle source databases.

User response: To start the Q Capture program, open a console on the Oracle server and use the asnoqcap command. To issue commands to a running Q Capture program, use the asnoqcmd command.

ASN2524E One or more values are missing for the required parameter or parameters *parameter_names*. You must specify values for both MQLOG and MQDATA when you specify the MULTIINST option with the CREATE MQ SCRIPT command. No script is generated.

Explanation: In a multi-instance queue manager configuration, you must specify the MQLOG and MQDATA keywords with CREATE MQ SCRIPT. MQLOG specifies the directory where the queue manager writes its log, and MQDATA specifies the directory that is used to hold the queue manager data files. All instances of the multi-instance queue manager use the specified directories.

User response: Reissue the CREATE MQ SCRIPT command, specifying values for MQLOG and MQDATA.

ASN2525E The PARALLEL SENDQS option is not supported because one or more of the databases in your configuration is DB2 for Linux, UNIX, and Windows. No script is generated.

Explanation: The PARALLEL SENDQS option is only supported if all of the databases that are involved in the configuration are on DB2 for z/OS.

User response: Specify DB2 for z/OS subsystems for the Q Capture and Q Apply servers and rerun the command.

ASN2526E The PARALLEL SENDQS option is not supported because the Q Capture control tables, Q Apply control tables, or both are at an architecture level that is lower than 1001. No script is generated.

Explanation: The PARALLEL SENDQS option is supported only if both Q Capture and Q Apply control tables are at architecture level of 1001 or higher.

User response: Migrate the control tables to the 1001 architecture level and rerun the command.

ASN2527I A value *value* was specified for the PARALLEL SENDQS option, which indicates that the send queue *queue_name* will run in multiplexing mode. All of the underlying parallel send queues for this send queue must exist before you start replication. A SQL script to create or alter the replication queue map was generated.

Explanation: Because the send queue will run in multiplexing mode, the Q Capture program expects all underlying parallel send queues to exist before replication is started. The number of send queues and channels that you create should match the value that was specified for PARALLEL SENDQS.

User response: Make sure that the appropriate number of queues are created before you start replication.

ASN2528I The number of parallel send queues for replication queue map *queue_map_name* was changed from *old_value* to *new_value* while the queue map is active. A script to alter the queue map was generated. After the script is run, make sure to reinitialize or stop and start the Q Capture program so that it begins to use the new value.

Explanation: The new value is not used until the Q Capture program rereads the properties of queue map, which requires that you reinitialize or stop and start Q Capture

User response: Reinitialize or stop and start the Q Capture program.

ASN2529E Using a multiplexed send queue in two different queue maps is not supported. The send queue *queue_name* that was specified for replication queue map *queue_map_name* is also being used by replication queue map *queue_map_name*. The second queue map is defined to use send queue multiplexing. It uses *num_parallel_sendqs* send queues and also

specifies the send queue *queue_name*.

Explanation: The send queue that was specified for the queue map is being used by another queue map that is defined with multiplexing. Although there is no entry in the IBMQREP_SENDQUEUES table with this send queue name, it can be inferred that the send queue is already being used because of the base send queue name.

User response: Provide a different send queue name for the queue map that you are creating or altering and rerun the command.

ASN2530I **The replication queue map *queue_map_name* is defined with number parallel send queues. The validation checks are performed on all of the send queues.**

Explanation: The ASNCLP program will perform validation checks on all of the parallel send queues that are specified for the queue map, and will also check all of the WebSphere MQ channels that are associated with the send queues.

User response: This message is for your information only. No action is required.

ASN2531E **The PARALLEL SENDQS option was specified with the CREATE MQ SCRIPT command, but the required keywords PLATFORM ZOS were not specified. No script is generated.**

Explanation: Because send queue multiplexing is supported only on z/OS, if you specify the keywords PARALLEL SENDQS you must also specify PLATFORM ZOS.

User response: Specify the PLATFORM ZOS keywords and rerun the command.

ASN2539E **The EXIST option was specified under the INCLUDE HISTORY clause to indicate that a history table exists on the target. However, no history table was found for the target table *table_owner.table_name*. No script was generated.**

Explanation: The EXIST option specifies that a history table is already created for the target temporal table. In this situation the ASNCLP program creates a Q subscription for both the temporal table and the history table. The ASNCLP program was unable to find a history table, so no Q subscription was created.

User response: Take one of the following actions and rerun the CREATE QSUB command:

- Remove the EXIST clause and use the HIST_TARGET NAME keywords to specify a name for the history table so that the ASNCLP program can create the history table.
- Create the history table, alter the target table to enable versioning, and use the EXIST option.

ASN2540E **The target table *table_owner.table_name* has the SYSTEM_TIME period column defined but the ASNCLP program did not find a history table. No script was generated.**

Explanation: To create a Q subscription for a target temporal table, the table must have an associated history table. Versioning also must be enabled on the table to make it a complete system temporal table.

User response: Create the history table, alter the target table to add versioning, and rerun the CREATE QSUB command.

ASN2541E **A primary key or unique index was detected on the IBMQREPVERTIME column or IBMQREPVERNODE column in the source table *table_name* on server *server_name*. Q subscriptions for peer-to-peer replication do not allow primary keys or unique indexes on these columns.**

Explanation: The IBMQREPVERTIME and IBMQREPVERNODE columns are used to track different versions of row changes in peer-to-peer replication. These columns are reserved for use by Q Replication and cannot have primary keys or unique indexes.

User response: Remove the primary key or unique index that is defined on the IBMQREPVERTIME column, IBMQREPVERNODE column, or both, and then reissue the action to create a Q subscription.

Chapter 41. ASN3500 - ASN3999

ASN3750I The exception-reporting tool successfully started.

Explanation: In response to an asnqexp command, the replication exception-reporting tool successfully started.

User response: This message is for your information only. No action is required.

ASN3751I The exception-reporting tool will use the most recent start time of the Q Capture program as the start time for generating an exceptions report.

Explanation: By default, the exception-reporting tool uses the most recent Q Capture start time as the start time for the exceptions report. The tool uses the timestamp when all stranded changes have been delivered and processed at the failover site as the end time of the report. You can change the start and end times by using the exception_report_start and exception_report_end parameters of the asnqexp command.

User response: This message is for your information only. No action is required.

ASN3752I The exception-reporting tool will use the timestamp when all stranded changes have been delivered and processed at the failover site as the end time of the exceptions report.

Explanation: By default, the exception-reporting tool uses the most recent Q Capture start time as the start time for the exceptions report, and the timestamp when all stranded changes have been delivered and processed at the fail-over site as the end time of the report. You can change the start and end times by using the exception_report_start and exception_report_end parameters of the asnqexp command.

User response: This message is for your information only. No action is required.

ASN3753I *exception_number* exceptions were detected during the *duration* time between the start of the reporting period *exception_report_start* (Q_Apply_time_zone) and the end of the reporting period *exception_report_end* (Q_Apply_time_zone).

Explanation: The exception reporting tool ran to completion and generated a report in the directory from which the asnqexp command was invoked or in

a directory that was specified by the file parameter.

User response: This message is for your information only. Review the generated exception report.

ASN3754I No exceptions were found during the *duration* time between the start of the report period from *exception_report_start* (Q_Apply_time_zone) to the end of the report period *exception_report_end* (Q_Apply_time_zone).

Explanation: The exception reporting tool ran to completion and generated a report in the directory from which the asnqexp command was invoked or in a directory that was specified by the file parameter. No exceptions were found during the reporting period.

User response: This message is for your information only. No action is required.

ASN3755E An invalid timestamp *timestamp* was used with the exception_report_start parameter when the exception-reporting tool started. The tool will stop.

Explanation: The exception_report_start parameter requires a timestamp that uses any of the following formats: YYYY-MM-DD-HH.MM.SS.ssssss, YYYY-MM-DD-HH.MM.SS, YYYY-MM-DD-HH.MM, YYYY-MM-DD-HH, and YYYY-MM-DD.

User response: Restart the tool and specify a timestamp in a valid format.

ASN3756E An invalid timestamp *timestamp* was used with the exception_report_end parameter when the exception-reporting tool was started. The tool will stop.

Explanation: The exception_report_end parameter requires a timestamp that uses any of the following formats: YYYY-MM-DD-HH.MM.SS.ssssss, YYYY-MM-DD-HH.MM.SS, YYYY-MM-DD-HH.MM, YYYY-MM-DD-HH, and YYYY-MM-DD.

User response: Restart the tool and specify a timestamp in a valid format.

ASN3757E An invalid timestamp *timestamp* was used with the exception_report_end parameter when the exception-reporting tool was started. The end time should be larger than the sum of the start time *timestamp* plus the monitor interval *interval_value*. The tool will stop.

Explanation: When you specify an end point for the exceptions report, the timestamp that you provide for the `exception_report_end` parameter must be at least as long as one Q Apply monitor interval (specified by the `monitor_interval` parameter). The default Q Apply monitor interval is 60000 milliseconds or 1 minute on z/OS and 30000 milliseconds or 30 seconds on Linux, UNIX, and Windows. Typically the exception-reporting period would be significantly longer than one monitor interval.

User response: Restart the tool and specify a timestamp with an appropriate value.

ASN3758E **The exception-reporting tool could not determine the timestamp that denotes when Q Capture was started to publish stranded transactions. A report was generated but no exceptions are shown in the report.**

Explanation: The `asnxqexp` command was issued, but no value was specified for the `exception_report_start` or `exception_report_end` parameters. When these parameters are not provided, the exception-reporting tool assumes that you want a report about exceptions that occurred during the delivery of stranded transactions after the most recent Q Capture restart. However, the tool must know when Q Capture was started to provide this type of report. The ASN0572I message that contains the timestamp of the most recent Q Capture start could not be found in the `DESCRIPTION` column of the `IBMQREP_CAPTRACE` table at the site of the outage. In this situation, the tool issues a report that does not contain any information about exceptions, but provides timestamps for recent Q Capture starts that you can potentially use as values for the `exception_report_start` or `exception_report_end` parameters.

User response: Take one of the following actions, depending on what type of exceptions report you want:

- To generate a general exceptions report (a report that shows the exceptions whose causes are not limited to stranded transactions after a failover situation), reissue the `asnxqexp` command and specify the `exception_report_start` and `exception_report_end` parameters. These parameters specify the start time and end time of the exception-reporting window.
- To generate a report of the exceptions that were caused by the delivery of stranded transactions, reissue the `asnxqexp` command and specify the `exception_report_start` parameter with the restart time of Q Capture after the failover. If Q Capture was stopped and started multiple times before all stranded transactions were republished, specify the first time that Q Capture was started.

All timestamps should be in the time zone of the Q Apply server.

ASN3759E **The exception-reporting tool did not find any records in the IBMQREP_APPLYMON table from the time period *timestamp to timestamp*, or could not determine the monitor interval of the Q Apply program. No exception report was generated.**

Explanation: The exception-reporting tool could not find any row in the `IBMQREP_APPLYMON` table in this exception-reporting window and therefore did not generate a report. The following are possible reasons:

1. Q Apply was not started before the `asnxqexp` command was issued.
2. The value of the Q Apply `monitor_interval` parameter is too large.
3. Q Apply had not performed an insert into `IBMQREP_APPLYMON` before the `asnxqexp` command was issued and after the exception-reporting start time.

User response: Take one of the following actions before reissuing the `asnxqexp` command, depending on the reason for the problem:

1. Start the Q Apply program.
2. Reduce the `monitor_interval` to a value that is less than 60000 milliseconds.
3. Wait for three monitor intervals.

ASN3760W **The exception report that was generated by the `asnxqexp` command might not be accurate and complete because the Q Apply program was stopped during the exception-reporting window from *timestamp to timestamp*. This stoppage might have occurred before Q Apply finished writing all monitor statistics to the `IBMQREP_APPLYMON` table.**

Explanation: The exceptions reporting tool detected transactions that might have been applied but were not reported in the `IBMQREP_APPLYMON` table. Two things could have occurred:

- Q Apply was stopped before transactions or stranded transactions were applied.
- All transactions or stranded transactions were applied but the statistics in the `IBMQREP_APPLYMON` table do not reflect all of the transactions.

User response: If you believe that Q Apply was stopped before all transactions or stranded transactions were applied, restart Q Apply and let it run until the value in the `ROWS_PUBLISHED` column of the `IBMQREP_APPLYMON` table is 0 for three consecutive monitor intervals. To avoid this problem in the future, take one of the following actions:

- Shorten the value of the Q Apply `monitor_interval` parameter so that it is less than 60000 milliseconds.

- If Q Apply must be stopped, let it run until the value in the ROWS_PUBLISHED column of the IBMQREP_APPLYMON table is 0 for three consecutive monitor intervals.

ASN3761W The exception-reporting tool detected that the Q Apply program is still processing data. The exception report might not be accurate and complete.

Explanation: The exception-reporting tool determined that Q Apply is still applying changes. Some of these changes might be stranded transactions that might cause exceptions. Thus the tool cannot provide a full report.

User response: Take one of the following actions:

- Reissue the asnqexp command with a timestamp for the exception_report_end parameter.
- Reissue the asnqexp command until this warning message is not returned. By doing this, you can verify that Q Apply has finished processing all transactions that were stranded by the outage.

ASN3762W The exception-reporting tool could not determine when the Q Apply program finished processing stranded transactions. The value of the exception-reporting window end time will appear as "N/A" in the report.

Explanation: To generate a report of all exceptions that occurred during the processing of stranded transactions, the tool needs information from the Q Apply program. The following problems might have occurred:

- There is no row in the IBMQREP_APPLYMON table at the failover site where the value of ROWS_NOT_APPLIED or ROWS_APPLIED is greater than zero and MONITOR_TIME is later than the start time of the exception-reporting window.
- The Q Apply program is still running.

User response: Take one of the following actions, depending on the reason for the problem:

- Reissue the asnqexp command and specify an approximate timestamp for the exception_report_start parameter.
- Wait for three Q Apply monitor intervals and then reissue the asnqexp command.
- Reissue the asnqexp command with the exception_report_start and exception_report_end parameters to specify both the start time and end time of the exception-reporting window.

Chapter 42. ASN4000 - ASN4499

ASN4003E *program_name : program_identifier : The data type or the length of the source column source_col and of the target column target_col are not compatible.*

Explanation: See message text.

User response: Rerun the command ensuring that you compare columns of equal type, precision, and length.

ASN4004E *program_name : program_identifier : The program encountered an SQL error at line line_number. The function name is function_name. The SQLCODE is sqlcode. The SQLSTATE is sqlstate.*

Explanation: A negative SQLCODE was returned when the program executed an EXEC SQL statement or CLI call. This message sometimes is followed by a second message that provides more information about the action being executed by the program when the SQLCODE was encountered.

User response: See the messages and codes documentation of the DB2 database manager on your operating system for an explanation of this SQLCODE and for information about corrective actions that might need to be taken. If the program issued another message following this one, see the explanation and user response for that message. The line and function information are for IBM Software Support only.

ASN4005E *program_name : program_identifier : The program encountered an SQL warning at line line_number. The function name is function_name. The SQLCODE is sqlcode. The SQLSTATE is sqlstate.*

Explanation: A warning SQLCODE was returned when the program executed an EXEC SQL statement or CLI call. This message sometimes is followed by a second message that provides more information about the action being executed by the program when the SQLCODE was encountered.

User response: See the messages and codes documentation of the DB2 database manager on your operating system for an explanation of this SQLCODE and for information about corrective actions that might need to be taken. If the program issued another message following this one, see the explanation and user response for that message. The line and function information are for IBM Software Support only.

ASN4006I *program_name : program_identifier : Between the source table and the target table, there are common_rows common rows, source_rows rows that are unique to the source table, and target_rows rows that are unique to the target table.*

Explanation: See message text.

User response: Review details in the difference table. If necessary, run the asntrep command to synchronize the tables.

ASN4007E *program_name : program_identifier : The program found an incorrect subscription definition. Reason code: reason_code.*

Explanation: Possible reason codes are:

1. The target table did not have key columns specified in the IS_KEY column of the IBMSNAP_SUBS_COLS table or the IBMQREP_TRG_COLS table.
2. No columns for the table were found for the target table in the IBMSNAP_SUBS_COLS table or the IBMQREP_TRG_COLS table.

User response: Run the Analyzer for details to help you fix the subscription definition. Use the replication administration tools to fix the subscription definition. Rerun the asntdiff or asntrep command.

ASN4008E *program_name : program_identifier : The program stopped because the WHERE parameter where_clause caused a DB2 PREPARE statement to fail.*

Explanation: The program stopped because the specified WHERE parameter caused an SQL PREPARE statement to fail. This statement is necessary to obtain the name of the target table by selecting rows from the IBMSNAP_SUBS_MEMBR table for SQL Replication or from the IBMQREP_TARGETS table for Q Replication.

User response: Use an interactive command line processor to debug the WHERE clause for the appropriate control table: IBMSNAP_SUBS_MEMBR or IBMQREP_TARGETS. Rerun the command using the valid WHERE clause in the WHERE parameter.

ASN4009E *program_name : program_identifier : The rows that were fetched from the source table were not in the order of the dynamic collation. The program stopped.*

Explanation: The cursor that was used to fetch rows

from the source table fetched the rows out of order. This problem occurs because the actual collation of the key columns in the source table did not match the estimated collating sequence that was determined by using the difference table. The difference table may have been created in a table space that has different attributes than the table space that contains the source table. Alternatively, the date or timestamp format may be different between the source and target.

User response: For information about creating the difference table in the correct table space, see the Technote, "Code page and table space considerations for running asntdiff."

If your table contains date or timestamp formats, see the Technote, "Why I am getting asn4009e or value of 'U 2' in DIFF column for every row from asntdiff when source and target tables are identical?"

ASN4010I *program_name : program_identifier :* **Number of differences found between the source and target table: number. The details can be found in database source_DB, difference table difference_table.**

Explanation: See message text.

User response: Review the difference table and, if necessary, run the asntrep command to fix the differences.

ASN4011I *program_name : program_identifier :* **No differences were found between the source and target tables.**

Explanation: See message text.

User response: This message is for your information only. No action is required

ASN4012I *program_name : program_identifier :* **The program is comparing tables using the list of parameters following this message.**

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN4013E *program_name : program_identifier :* **The program stopped because the WHERE clause was not valid. It returned number_of_rows subscription members or Q subscriptions.**

Explanation: The clause provided in the WHERE parameter returned no row or more than one row from either the IBMSNAP_SUBS_MEMBR or the IBMQREP_TARGETS table.

User response: Use an interactive command line processor to debug the WHERE clause for the appropriate control table: IBMSNAP_SUBS_MEMBR or IBMQREP_TARGETS. Rerun the command using the valid WHERE clause in the WHERE parameter.

ASN4014E *program_name : program_identifier :* **The source table structure source_structure is not supported.**

Explanation: The following source table structures are supported for SQL Replication: user, replica, user copy, and point-in-time. Only user tables are supported for Q Replication.

User response: Modify the WHERE clause to select a supported source structure and rerun the asntdiff or asntrep command.

ASN4015E *program_name : program_identifier :* **The target table structure target_structure is not supported.**

Explanation: The following target table structures are supported for SQL Replication:

- User tables
- Replica tables
- User copy tables
- Point-in-time tables
- Consistent-change-data (CCD) tables that are condensed and complete

The following target table structures are supported for Q Replication:

- User tables
- Consistent-change-data (CCD) tables that are condensed and complete

User response: Modify the WHERE clause to select a supported target structure and rerun the asntdiff or asntrep command.

ASN4018E *program_name : program_identifier :* **The difference table table_owner_table_name already exists in database database_name and its columns are not compatible with the key columns for the source table.**

Explanation: The difference table already existed and the asntdiff program is trying to reuse the table to write the differences. The key column information in the table does not match the key columns of the current source table for which the asntdiff program is run.

User response: Delete the difference table or invoke the asntdiff program with parameter DIFF_DROP=Y.

ASN4019I *program_name : program_identifier : The program applied the rows from the difference table to the target table as follows: number rows were inserted, number rows were updated, and number rows were deleted.*

Explanation: The message summarizes the differences that were applied to the target table to synchronize it with the source table.

User response: This message is for your information. No action is required.

ASN4020I *program_name : program_identifier : The value in the SUPPRESS_DELETES column is set to Y for this subscription, so number deletes were ignored.*

Explanation: When the SUPPRESS_DELETES value for a subscription is Y, the asntdiff program ignores the rows that are unique to the target and reports no differences. SUPPRESS_DELETES=Y prompts Q Capture and SQL Capture to not process deletes for a subscription. This may result in rows that exist in the target table but not in the source table.

User response: This message is for your information. No action is required.

ASN4021E *program_name : program_identifier : The difference table table_owner.table_name does not exist at the Classic replication server server_name. The program stopped.*

Explanation: The difference table that is required to run the asntdiff utility does not exist at the Classic replication server. The utility does not automatically create the difference table. The table must be created manually.

User response: Create the difference table at the Classic replication server and issue the asntdiff command.

ASN4022E *program_name : program_identifier : The file system that contains spill key file spill_filename ran out of space after number keys were spilled to the file. The program stopped.*

Explanation: The asntdiff utility temporarily spills keys that it reads from the source table to disk when a memory limit is reached. The file system that contains this spill file ran out of space while the utility was writing keys to it. Unless all keys can be written to the spill file, the asntdiff utility cannot continue. By default, the utility will create the spill file in the temporary directory.

User response: Either free space on the file system temporary directory that is being used to spill keys, or

invoke the asntdiff utility with the diff_path parameter to specify a directory with more space. To view the path of the temporary file, run the asntdiff command with the DEBUG=Y option.

ASN4023E *program_name : program_identifier : The program encountered an SQL error at line line_number. The function name is function_name. The SQLCODE is sqlcode. The SQLSTATE is sqlstate.*

Explanation: An SQLCODE was returned when the program executed an ODBC/CLI statement against a Classic server. This message sometimes is followed by a second message that provides more information about the action that the program was executing when the error occurred.

User response: See the Classic system messages in the IBM WebSphere Classic information center for an explanation of this SQLCODE and for information about corrective actions. If the program issued another message after this one, see the explanation and user response for that message. The line number and function name are for IBM Software Support only.

ASN4024E *program_name : program_identifier : The repair table table_owner.table_name already exists in database database_name and its columns are not compatible with the columns of the difference table table_owner.table_name in database database_name.*

Explanation: A repair table at the target database is a copy of the difference table at the source database that contains only the keys to be deleted. In this instance, the repair table already exists from a previous asntrep command invocation, but the table contains incompatible columns.

User response: Delete the repair table and then reissue the asntrep command, or issue the command with the parameter DIFF_DROP=Y.

ASN4026E *program_name : program_identifier : An error occurred while parsing the RANGECOL option. The asntdiff program stopped.*

Explanation: The RANGECOL option was specified but the syntax was incorrect.

User response: Specify the correct RANGECOL syntax.

ASN4027E *program_name : program_identifier : The SQL data type of the source column that was specified for the RANGECOL option is invalid. Valid data types are DATE, TIME, or TIMESTAMP.*

Explanation: Either the wrong column was specified for the RANGECOL option, or the specified column has an incorrect data type.

User response: Specify a source column of DATE, TIME, or TIMESTAMP data type and rerun the program.

ASN4028E *program_name : program_identifier :* **An error was detected while accessing the IBMTDIFF_REFRESH table. The sqlcode is sql_code**

Explanation: DB2 returned an error while the ASNTDIFF program was trying to access the IBMTDIFF_REFRESH table to process the RANGECOL option.

ASN4029E *program_name : program_identifier :* **The UPTO_VALUE column in the IBMTDIFF_REFRESH table was not set.**

Explanation: To use the REFRESHONLY option, the UPTO_VALUE column in the IBMTDIFF_REFRESH table must have been set initially.

User response: The UPTO_VALUE can be initialized when the ASNTDIFF program is run using the FROM:, TO:, or FROM: and TO: RANGECOL option, or you can insert a value for the UPTO_VALUE column in the IBMTDIFF_REFRESH table and rerun asntdiff with the REFRESHONLY option.

ASN4030I *program_name : program_identifier :* **The asntdiff program that was invoked with the RANGECOL option processed rows from timestamp to timestamp. The schema is schema. The subscription name is subscription_name, and the number of rows that were processed is number.**

Explanation: This message reports the result of running the asntdiff program with the RANGECOL option. The number of rows that were processed is set to -1 if the run was not successful.

User response: If the number of rows processed is -1, check the asntdiff.log file for error messages that explain why the run did not succeed.

ASN4031E *program_name : program_identifier :* **The associated target column name for the source column column_name that was specified for the RANGECOL parameter was not found in the target table.**

Explanation: The asntdiff program could not find the target column in the IBMQREP_TRG_COLS table (Q Replication) or in the IBMSNAP_SUBS_COLS table (SQL Replication).

User response: Use the replication administration tools to redefine the Q subscription or subscription set

member so that the correct target column is mapped to the source column.

ASN4034W *program_name : program_identifier :* **The asntdiff program compared TIMESTAMP columns of different precision. In its comparison, the program treats column column_name in database database_name as if it had a truncated value of data_length so the precision values match.**

Explanation: When comparing TIMESTAMP columns that have different precision, asntdiff treats the column with the larger precision as if it has a matching precision and then compares the resulting values. For example, when comparing columns that are defined as TIMESTAMP(12) and TIMESTAMP(6), asntdiff would treat both as TIMESTAMP(6) and then report any differences in the column values.

User response: Take the data truncation into account when you interpret the results of the asntdiff command. You can use the asntdiff -f option to explicitly cast a TIMESTAMP column to a given precision for the comparison.

ASN4035W *program_name : program_identifier :* **Database database_name has VARCHAR2 support enabled. The asntdiff program cannot differentiate between NULL values and empty strings at this database.**

Explanation: With VARCHAR2 support enabled, a database treats NULL values and empty character strings as equivalent. When comparing character data with another database that does not have VARCHAR2 support enabled, the asntdiff program treats NULL values and empty strings as different values and flags a difference. For example, empty strings at two databases would be treated as different values because the database with VARCHAR2 support enabled converts the empty string to a NULL value.

User response: When you interpret the results of the asntdiff command, take these differences into account.

ASN4038W *program_name : program_identifier :* **A change condition was detected for Q subscription name. Differences that are reported may be a result of this predicate.**

Explanation: Change conditions allow you to use log record variables such as \$OPERATION and \$AUTHID to filter information about how a row was created when the Q Capture program is replicating or publishing changes. Unlike search conditions, change conditions do not reference information from the source table, and therefore asntdiff might show some rows as different between source and target.

User response: Do not use the asntrep utility to repair the table if you suspect that a large number of the operations that it would make are due to the effect of the change condition that is being used by Q Capture. Repairing the target table in this manner would undo the effect of filtering at the source.

ASN4039W *program_name : program_identifier : The asntdiff utility was unable to verify the existence of the difference table table_name. A select query against the catalog table table_name at the source server server was rejected. A select query against the difference table also was not allowed. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc.*

Explanation: The difference table is created in the source database or subsystem to store differences that the asntdiff utility finds. If you specify DIFF_DROP=N, you can create the difference table yourself; otherwise the utility creates the table. The utility was unable to verify the existence of a difference table at the specified server. The SQLCODE and SQLSTATE explain the reasons that the queries that were used for verification were rejected.

User response: Check the SQLCODE and SQLSTATE information. Look for additional error messages such as ASN4040W or ASN4041E that contain information about how to fix the problem.

ASN4040E *program_name : program_identifier : The asntdiff utility was unable to create a difference table named table_name at the server server. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc.*

Explanation: The asntdiff command creates a difference table in the source database or subsystem to store differences that it finds. The utility was unable to create a difference table at the specified server. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE and fix any problems that are preventing asntdiff from creating a difference table. Then reissue the asntdiff command.

ASN4041E *program_name : program_identifier : The difference table table_name was not created even though the value of N for the DIFF_DROP parameter specifies that the table will be created manually rather than automatically by asntdiff. The program will stop.*

Explanation: With DIFF_DROP=N, the asntdiff utility expects to find a user-created difference table at the source database or subsystem. The utility looks for a

table with the default schema and name of ASN.ASNTDIFF unless you use the DIFF and DIFF_SCHEMA parameters to specify a different schema and name. The utility was unable to find the specified table.

User response: Set the DIFF_DROP parameter to Y if you want the asntdiff utility to create the difference table automatically. Otherwise, create the difference table manually before you run asntdiff again.

ASN4044W *program_name : program_identifier : The asntdiff utility tried to change the SQL authorization ID to id, but the attempt to issue a SET CURRENT SQLID query against the server server failed. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc. The utility will continue the work using the current SQL authorization ID.*

Explanation: You can specify the SOURCE_SQLID or TARGET_SQLID parameters to change the SQL authorization ID that is used by the asntdiff utility to perform database operations. The utility uses SET CURRENT SQLID to perform this change, but the operation failed. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE for explanations of why the SET CURRENT SQLID operation failed. Fix any problems and then reissue the asntdiff command.

ASN4045E *program_name : program_identifier : The asntdiff utility was unable to create a global temporary table named table_name at the server server. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc. The program will stop.*

Explanation: The asntdiff utility was unable to create a global temporary table at the specified server. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE for explanations of why asntdiff could not create a global temporary table at the server. Fix any problems and then reissue the asntdiff command.

ASN4046E *program_name : program_identifier : The asntdiff utility was unable to insert data into the table table_name at the server server. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc. The program will stop.*

Explanation: As part of its processing, the asntdiff utility creates a global temporary table at the specified server and performs inserts. The utility was unable to insert the records into the specified table. The

SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE for explanations of why asntdiff could not insert records into the global temporary table. Fix any problems and then reissue the asntdiff command.

ASN4047E *program_name : program_identifier : The asntdiff utility was unable to select data from the table table_name at the server server. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc. The program will stop.*

Explanation: The asntdiff utility was unable to issue a SELECT statement against the specified table. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE for explanations of why asntdiff could not issue a SELECT statement against the table. Fix any problems and then reissue the asntdiff command.

ASN4050W *program_name : program_identifier : The asntdiff utility was unable to determine if the SELECT statement statement uses expression_or_alias for key columns.*

Explanation: To automatically determine the value of the NUMBLOCKS parameter, the asntdiff utility must determine if the statements that were specified for the SOURCE_SELECT and TARGET_SELECT parameters use expressions or aliases for key columns. The utility could not determine whether expressions or aliases are used.

User response: If the expressions or alias are not used for the key columns, this warning message can be ignored. Otherwise, to achieve a better performance, remove the key columns from the SELECT statements or change the ORDER BY clauses in SOURCE_SELECT or TARGET_SELECT parameters. If neither of these changes is possible, set the value of the NUMBLOCKS parameter to 1 before the next run.

ASN4051W *program_name : program_identifier : The asntdiff utility was unable to determine if at least one index can be used to fetch the data from the table table_name by using key columns. The key columns are specified in the ORDER BY clauses clauses that were specified for the SOURCE_SELECT and TARGET_SELECT parameters.*

Explanation: The asntdiff utility was unable to access the catalog tables SYSIBM.SYSKEYS and SYSIBM.SYSINDEXES. By accessing the catalog tables, the asntdiff utility can determine if one index exists that can be used for fetching the result sets from the SOURCE_SELECT and TARGET_SELECT queries. The utility requires this information to automatically

determine the value for the NUMBLOCKS parameter.

User response: If an eligible index exists for key columns in the ORDER BY clauses, you can ignore this message. Otherwise, for better performance set the value of the NUMBLOCKS parameter to 1 before the next run.

ASN4052E *program_name : program_identifier : The asntdiff utility was unable to create more parallel threads because of a system resource limit. The program will stop.*

Explanation: An operating system thread limit was reached for the user ID that runs the asntdiff utility. The asntdiff utility creates at least six threads for parallel table comparison when the PARALLEL parameter is specified. The number of threads can be higher when you use the NUMTHREADS parameter.

User response: Take one of these actions:

- Increase the resource limit for the BPXPRMxx member of the PARMLIB: MAXTHREADS and MAXTHREADTASKS. For running the asntdiff utility in the parallel option, the minimum number of threads is six. To achieve the best performance, you should set the system thread limit to 21 or larger.
- Lower the value of the NUMTHREADS parameter.

ASN4053E *program_name : program_identifier : The asntdiff utility ran out of memory space. The program will stop.*

Explanation: The asntdiff utility is unable to compare the tables because of insufficient memory allocation.

User response: Take one of these actions:

- On z/OS, increase the region size.
- When using the parallel option, reduce the value of the NUMTHREADS parameter.

ASN4055E *program_name : program_identifier : An error occurred while the asntdiff utility was terminating an existing thread. The program will stop.*

Explanation: Some threads of the asntdiff utility are waiting for returns from the stored procedure calls from DB2.

User response: Kill the threads manually and set the value of parameter NUMBLOCKS to zero before the next run. You can use the -DISPLAY THREAD command to get status information about the threads.

ASN4056I *program_name : program_identifier : The asntdiff utility received a user signal to stop the compare operation. The program will stop.*

Explanation: The utility received an operating system signal to cease operations that was initiated by user input.

User response: This message is for your information only. No action is required.

ASN4058E *program_name : program_identifier : The asntdiff utility, running in parallel mode, is unable to compare the two specified tables because their code pages (CCSID) are different. The program will stop.*

Explanation: The result sets of the queries that were used for the SOURCE_SELECT and TARGET_SELECT parameters have different CCSID values. This situation can occur when the DB2 subsystems that contain the two tables that are being compared use different code pages. The parallel option requires that both result sets have the same CCSID value.

User response: Change the value of the PARALLEL parameter from Y to N and rerun the utility.

ASN4059E *program_name : program_identifier : The asntdiff utility in parallel mode cannot be run in this environment. The program will stop.*

Explanation: With this release, the asntdiff utility in parallel mode cannot be run in this environment.

User response: Change the value of the PARALLEL parameter from Y to N and rerun the utility.

ASN4060I *program_name : program_identifier : The asntdiff utility changed the value of the NUMBLOCKS parameter to value. The reason code is reason_code.*

Explanation: The value for NUMBLOCKS was reduced. The following values are valid for the reason code:

- 1 Each block contains too few rows.
- 2 The source table contains too few rows.
- 3 An alias was used for key columns.
- 4 An expression was used as a key column for the SQL statement that was used with the SOURCE_SELECT or TARGET_SELECT parameter.
- 5

An eligible index cannot be found based on the provided ORDER BY clauses.

6

At least one instance of the unsupported keyword DESC is used with the SOURCE_SELECT or TARGET_SELECT parameters.

User response: To improve performance, review the reason codes in the explanation and take the appropriate action:

1

Reduce the value of the NUMBLOCKS parameter.

2

Change the value of the PARALLEL parameter from Y to N.

3

Remove the alias from the SELECT statements in the SOURCE_SELECT or TARGET_SELECT parameters.

4

Remove the expression from the SELECT statements or change the ORDER BY clauses in the SOURCE_SELECT or TARGET_SELECT parameters.

5

Respond to reason code code 5 in one of the following ways:

- Create an eligible index for the key columns.
- Change the ORDER BY clauses in the SOURCE_SELECT or TARGET_SELECT parameters.

6

Remove the DESC keyword from any SQL query that is used for the SOURCE_SELECT parameter or the TARGET_SELECT parameter

ASN4061E *program_name : program_identifier : The number of columns (number) in the SQL statement that was used with the SOURCE_SELECT parameter does not equal the number of columns (number) that were used with TARGET_SELECT. The two statements must have the same number of columns for the tables to be compared. The program will stop.*

Explanation: The number of columns that are used with SOURCE_SELECT must match the number of columns that are used with TARGET_SELECT.

User response: Change the SQL statements that are

used with SOURCE_SELECT and TARGET_SELECT and rerun the asntdiff utility.

ASN4062E *program_name : program_identifier :*
Column column_name with data type data_type is not supported as a non-key column for the asntdiff utility when it runs in parallel mode. The program will stop.

Explanation: When you specify the PARALLEL parameter to run the asntdiff utility in parallel mode, comparison of columns with the specified data types is not supported.

User response: Remove references to columns with unsupported data types from the queries that are used with the SOURCE_SELECT and TARGET_SELECT parameters, then rerun the asntdiff utility. If the column references cannot be removed, rerun the utility with PARALLEL=N.

ASN4063E *program_name : program_identifier :* **The asntdiff utility, running in parallel mode, is unable to process the SQL statement statement that was used with the SOURCE_SELECT or TARGET_SELECT parameter. The program will stop.**

Explanation: The query statement that was used with the SOURCE_SELECT or TARGET_SELECT parameter does not use the expected format.

User response: Rewrite the statement using the following format:

```
SELECT xxx FROM yyy (WHERE zzz) ORDER BY aaa
```

Then rerun the asntdiff utility. If the statement cannot be rewritten in this format, rerun the asntdiff utility with PARALLEL=N.

ASN4064E *program_name : program_identifier :* **The provided value value for the NUMTHREADS parameter is too small. The value should be between 6 and 21. The program will stop.**

Explanation: The asntdiff utility, running in parallel mode, must create at least six threads. For optimal performance NUMTHREADS should be set to 21.

User response: Increase the system resource limit for application threads if necessary and set the value of NUMTHREADS to 21. Then rerun the asntdiff utility. If the NUMTHREADS value cannot be increased because of a system resource limit, rerun the utility with PARALLEL=N.

ASN4065W *program_name : program_identifier :* **The parameter parameter is only used for environment. The value will be ignored in this run.**

Explanation: The specified parameter can only be used in some environments. The parameter is ignored in this run.

User response: Remove the parameter in future runs.

ASN4066I *program_name : program_identifier :* **Table partitioning was unsuccessful. The comparison will be performed in one pass against the entire results set.**

Explanation: The asntdiff utility was unable to partition the table due to an unsuccessful statement rewrite or a failure in data fetching from the source table.

User response: Check the SQLCODE and SQLSTATE in ASN0552E, if issued.

ASN4067E *program_name : program_identifier :* **Table partitioning was unsuccessful. The program will stop.**

Explanation: The asntdiff utility was unable to partition the table because of an unsuccessful statement rewrite or a failure in fetching data from the source table.

User response: Set the value of parameter NUMBLOCKS to one before the next run.

ASN4069E *program_name : program_identifier :* **The stored procedure procedure_name that is used by the asntdiff utility in parallel mode was unable to return the expected results. Reason: reason. The program will stop.**

Explanation: The specified stored procedure failed. The reason message from the stored procedure explains the problem.

User response: Check the reason message. At the server, verify that the stored procedure is correctly defined and that the configuration of the Workload Manager environment is valid. Then rerun the utility.

ASN4070E *program_name : program_identifier :* **The stored procedure procedure_name was not bound at the server server. SQLERRMC sqlerrmc. The program will stop.**

Explanation: The asntdiff utility was able to run the specified stored procedure that is required for parallel mode. But the embedded SQL statements in the stored procedure cannot be executed. The associated DBRM was not found by DB2. See the text for the SQLCODE -805.

User response: Bind the DBRM of the stored procedure before the next run. For information on binding asntdiff, see the ASNQBNBDT sample data set.

ASN4072E *program_name : program_identifier : The asntdiff utility was unable to issue the DB2 stored procedure procedure_name at the server server. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc. The program will stop.*

Explanation: The asntdiff utility was unable to run the specified stored procedure. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE. At the server, verify that the definition of the stored procedure is correctly specified and the configuration of the Workload Manager environment is valid.

ASN4073E *program_name : program_identifier : The ORDER BY clauses that were used with the asntdiff input parameters SOURCE_SELECT and TARGET_SELECT are different. The command failed.*

Explanation: When you use the asntdiff -f (file) option, you specify the SOURCE_SELECT and TARGET_SELECT parameters to provide SQL SELECT statements that specify the rows that you want to compare in the source and target tables. An ORDER BY clause is required in each of these SQL statements, and the ORDER BY clauses must be the same for both parameters.

User response: Change the SQL statements in the input file so that both ORDER BY clauses match, and reissue the asntdiff command.

ASN4074E *program_name : program_identifier : The column or columns that are specified in the ORDER BY clauses for the SOURCE_SELECT and TARGET_SELECT parameters of the asntdiff command do not reference a unique key or unique composite key. The command failed.*

Explanation: When you use the asntdiff -f (file) option, you specify the SOURCE_SELECT and TARGET_SELECT parameters to provide SQL SELECT statements that specify the rows that you want to compare in the source and target tables. The SQL statements must have an ORDER BY clause, and the clauses must reference a unique key or unique composite key.

User response: Rewrite the ORDER BY clauses in the input file and reissue the asntdiff command.

ASN4075W *program_name : program_identifier : The asntdiff utility was unable to drop the table table_name at the server server. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sqlerrmc.*

Explanation: The asntdiff utility was unable to issue a drop statement against the specified difference table or the global temporary table that it created. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE. Manually drop the specified table and rerun the utility.

ASN4076E *program_name : program_identifier : The asntdiff utility stopped because the file system that contains spill key file spill_filename is unable to store all of the key values of the source result set. In each comparison, the total size of key columns multiplied by the number of rows in the source table should be less than 2 GB.*

Explanation: When the asntdiff utility reaches a memory limit, it temporarily spills key column values that it reads from the source table to disk. The maximum allowable size of the spill file is 2 GB. Unless it can write all key values to the spill file, the asntdiff utility stops.

User response: Split the compare operation into smaller tasks:

- For the -f (input file) command option, partition the original comparison by specifying ranges of rows to compare. You can use any combination of columns and choose the appropriate partition column values as the boundary values in the search conditions. Columns that you choose for range values should be indexed for better performance. The corresponding search conditions should be added into the WHERE clauses of both the SOURCE_SELECT and TARGET_SELECT parameters.
- If you are not using the -f option, partition the comparison by specifying the RANGECOL parameter to indicate ranges of rows to compare. The selected columns must be DATE, TIME, or TIMESTAMP. If columns with these data types are not available, try to run the asntdiff utility with the -f command option.

Chapter 43. ASN5000 - ASN5499

ASN5101I **MONITOR** *monitor_qualifier*. The Replication Alert Monitor program started successfully.

Explanation: This message appears after a successful start of the Replication Alert Monitor program.

User response: This message is for your information only, and no action is required.

ASN5102I **MONITOR** *monitor_qualifier* : The Replication Alert Monitor version *version_number* program initialized successfully and is monitoring *number* alert conditions.

Explanation: The Replication Alert Monitor program successfully started.

User response: This message is for your information only, and no action is required.

ASN5103I **MONITOR** *monitor_qualifier*. The Replication Alert Monitor program re-initialized successfully and is monitoring *number-of-alert-conditions* alert conditions.

Explanation: The Replication Alert Monitor program successfully re-initialized.

User response: This message is for your information only, and no action is required.

ASN5104W **MONITOR** *monitor_qualifier*. *number-of-alert-conditions* alert conditions were ignored.

Explanation: The Replication Alert Monitor program initialized or re-initialized. Some alert conditions that are not valid might be excluded as noted in previously issued messages.

User response: Check the IBMSNAP_CAPTRACE table for messages about excluded alert conditions.

ASN5107I **MONITOR** *monitor_qualifier*. The Replication Alert Monitor program stopped.

Explanation: The Replication Alert Monitor program terminated.

User response: This message is for your information only, and no action is required.

ASN5108W **MONITOR** *monitor_qualifier*. The EMAIL_SERVER parameter was not set so e-mail notification cannot be sent.

Explanation: You started the Replication Alert Monitor without specifying a value for the EMAIL_SERVER parameter, and the EMAIL_SERVER column in the IBMSNAP_MONPARMS table is empty. The Replication Alert Monitor will store any alerts in the IBMSNAP_ALERTS table.

User response: If you want to receive e-mail notification for alerts, restart the Replication Alert Monitor with the EMAIL_SERVER value set.

ASN5109W **MONITOR** *monitor_qualifier*. The MONITOR_ERRORS parameter was ignored because the EMAIL_SERVER parameter was not set.

Explanation: The Replication Alert Monitor started but the MONITOR_ERRORS parameter was ignored because there was no value specified for the EMAIL_SERVER parameter. The MONITOR_ERRORS parameter requires the EMAIL_SERVER parameter. The Replication Alert Monitor will store any alerts in the IBMSNAP_ALERTS table.

User response: If you want to be notified by e-mail when an alert occurs, restart the Replication Alert Monitor with the MONITOR_ERRORS parameter and the EMAIL_SERVER parameter set.

ASN5111I **MONITOR** *monitor_qualifier*. *number-of-rows* rows were pruned from the table *schema.table-name* at timestamp *timestamp*.

Explanation: The Replication Alert Monitor program pruned rows from the IBMSNAP_ALERTS or the IBMSNAP_MONTRACE table.

User response: This message is for your information only, and no action is required.

ASN5117E **MONITOR** *monitor_qualifier*. There are no valid Alert Conditions for this monitor qualifier on Monitor Server *monitor_server*.

Explanation: The Replication Alert Monitor program cannot find any alert conditions for this monitor qualifier.

User response: Verify that the monitor qualifier name used with the monitor_qual parameter is correct, and check that the alert conditions are enabled. Also, check

any previously issued messages.

ASN5118E **MONITOR** *monitor_qualifier*. **The program cannot connect to the monitor control server** *server_name*. **The SQLCODE is** *sqlcode*, **and the SQLSTATE is** *sqlstate*.

Explanation: The Monitor program tried to connect to the monitor control server and failed with the corresponding SQLCODE.

User response: Read the corresponding action for this SQLCODE, and correct the error.

ASN5119E **MONITOR** *monitor_qualifier*. **The program cannot connect to the server** *server_name*. **The SQLCODE is** *sqlcode*, **and the SQLSTATE is** *sqlstate*.

Explanation: An SQL CONNECT statement failed when the Monitor program tried to connect to the monitored Capture or Apply control server.

User response: Read the corresponding action regarding this SQLCODE, and correct the error.

ASN5121E **MONITOR** *monitor_qualifier*. **A contact does not exist for the alert condition with a component of** *component*, **a server of** *server*, **a schema or qualifier of** *schema_or_qualifier*, **and a condition name of** *condition*.

Explanation: The specified contact does not exist in the IBMSNAP_CONTACTS table for this alert condition.

User response: Verify the contact information, and correct the alert condition.

ASN5122E **MONITOR** *monitor_qualifier*. **The contact group** *group-name* **does not exist or is empty. The component is** *component*, **the server is** *server*, **the schema or qualifier is** *schema_or_qualifier*, **and the condition name is** *condition*.

Explanation: The contact group specified in an alert condition does not have corresponding contacts in the IBMSNAP_CONTACTGRP table or does not exist in the IBMSNAP_CONTACTGRP table. A contact group cannot be empty.

User response: Verify the contacts for this group, and correct the alert condition.

ASN5123E **MONITOR** *monitor_qualifier*. **The table** *table_name* **was not found for program** *program_name*. **The server is** *server_name*. **The schema is** *schema*. **The condition name is** *condition_name*.

Explanation: The Monitor program could not find a table while attempting to monitor a condition on the server used by the program.

User response: Verify that the table exists on the server, or correct the alert condition.

ASN5124E **MONITOR** *monitor_qualifier*. **The table** *table-name* **is not found. The Apply control server is** *apply-control-server*. **The Apply qualifier is** *apply-qualifier*. **The subscription-set name is** *set-name*. **The condition name is** *condition-name*.

Explanation: The Replication Alert Monitor program cannot find a table when attempting to monitor a condition on the Apply control server.

User response: Verify that the table exists on the Apply control server, or correct the alert condition.

ASN5125E **MONITOR** *monitor_qualifier*. **The Apply qualifier** *apply-qualifier* **or the subscription set** *set-name* **is not found.**

Explanation: The Replication Alert Monitor program cannot find the Apply qualifier or the subscription set when attempting to monitor a condition on the Apply control server.

User response: Verify that the Apply qualifier and the subscription set exist on the Apply control server, or correct the alert condition.

ASN5126E **MONITOR** *monitor_qualifier*. **There is an error sending a notification. The return code is** *rc*.

Explanation: When an alert notification was attempted, the ASNMAIL exit returned an error. The reasons for the error may include:

- | | |
|----|---|
| 1 | SMTP protocol failed. Verify the address of your e-mail server with your administrator. |
| 2 | SMTP socket failed. Verify the configuration of your e-mail server or client. |
| 3 | The e-mail address is not valid. Verify the e-mail address. |
| 4 | Software error. |
| 99 | ASNMAIL exit not found. |

User response: For return codes 1 and 2, verify the configuration of your e-mail server and client. For

return code 3, check whether the e-mail address is correct. For return code 99, verify that the ASNMAIL exit is properly installed.

ASN5127E **MONITOR** *monitor_qualifier*. **An invalid value *value* exists in the column *column-name* of the table *table-name*.**

Explanation: This message indicates that the Replication Alert Monitor program found a column with a value that is not valid during program initialization.

User response: Verify the values of the column definitions in the specified table.

ASN5128W **MONITOR** *monitor_qualifier*. **An alert condition was not processed because a row is missing from the table *table_name*. The condition name is *condition_name*. The server is *server_name*. The schema or qualifier is *schema_or_qualifier*.**

Explanation: While processing an alert condition, the Monitor program did not find a row that was required in the specified table. These are the row requirements for the alert conditions:

- CAPTURE_LASTCOMMIT alert condition needs at least one row in the IBMSNAP_RESTART table.
- CAPTURE_LATENCY alert condition needs to read the global row (GLOBAL_RECORD = Y) in the IBMSNAP_REGISTER table.
- APPLY_SUBSDELAY and APPLY_SUBSINACT alert conditions need at least one row for the Apply qualifier specified in the IBMSNAP_SUBS_SET table.
- QCAPTURE_SUBSINACT alert condition needs a row for the subscription name in the IBMQREP_SUBS table.
- QAPPLY_QDEPTH and QAPPLY_SPILLQDEPTH alert conditions need a row in the IBMQREP_APPLYPARMS table.

User response: Verify that the replication program (Capture, Apply, Q Capture, or Q Apply) is running with the appropriate settings. Also, verify that the alert conditions are specified correctly. Perhaps the schema or server name is wrong.

ASN5129I **MONITOR** *monitor_qualifier*. **The Replication Alert Monitor on server *server-name* reports an e-mail alert.**

Explanation: The Replication Alert Monitor program sent an e-mail alert.

User response: This message is for your information only, and no action is required.

ASN5130I **MONITOR** *monitor_qualifier*. *capture_message*. **The Capture control server is *capture-server*. The schema is *schema*. The monitor control server is *monitor-server*.**

Explanation: The Replication Alert Monitor program retrieved a Capture program message from the IBMSNAP_CAPTRACE table when processing CAPTURE_ERRORS or CAPTURE_WARNINGS conditions.

User response: Read the Capture program message, and take appropriate action. Also, check any errors or warnings from the Capture control server.

ASN5131I **MONITOR** *monitor_qualifier*. *apply_message*. **The Apply control server is *apply-server*. The Apply qualifier is *apply-qualifier*. The monitor control server is *monitor-server*.**

Explanation: The Replication Alert Monitor program retrieved an Apply program message from the IBMSNAP_APPLYTRAIL or the IBMSNAP_APPLYTRACE table when processing APPLY_SUBSFALING, APPLY_ERRORS, or APPLY_WARNINGS alert conditions.

User response: Read the Apply program message, and take appropriate action. Also, check any errors or warnings from the Apply control server.

ASN5133I **MONITOR** *monitor_qualifier*. **The following alert *message_number* has occurred *number_of_times* times in the last *number_of_minutes* minutes. The notification for this alert will be suspended.**

Explanation: This message is issued after an alert has been detected the number of times specified in the MAX_NOTIFICATIONS_PER_ALERT parameter (the default is 3) for the number of minutes specified in the MAX_NOTIFICATIONS_MINUTES parameter (the default is 60 minutes).

User response: This message is for your information only, and no action is required.

ASN5134I **MONITOR** *monitor_qualifier*. **Alerts issued.**

Explanation: The text of this message appears in the subject line of the e-mail alerts sent by the Replication Alert Monitor program.

User response: This message is for your information only, and no action is required.

ASN5135W **MONITOR** *monitor_qualifier*. **Too many alerts** *number_of_alerts* **for server** *server_name* **between** *lower_bound_time* **and** *upper_bound_time*, **schema or qualifier** *schema_qual_name*, **condition name** *condition_name*.

Explanation: The Alert Monitor reached the maximum of alerts allowed for a monitor cycle (1024), or memory can not be allocated for them. The Alert Monitor will send the notifications and will update the Monitor control server and will reconnect to the server starting in the next condition. Some alerts for the specified alert condition might not be sent and inserted in the Monitor control server.

User response: You should verify the specified alert condition directly in the server to check if alerts were lost.

ASN5136W **MONITOR** *monitor_qualifier*. **There has been an error calling DAS component. The return code is** *rc* **for server** *server_name* **for schema or qualifier** *schema_qual_name* **and condition name** *condition_name*.

Explanation: While processing the condition name for the given server, the DAS component returned an error.

User response: Verify that DAS is running properly in both the client and the remote server.

ASN5137E **MONITOR** *monitor_qualifier*. **An error occurred while calling the WebSphere MQ API. The condition name is** *condition_name*. **The server is** *server_name*. **The schema is** *schema*. **The queue manager is** *queue_manager*. **The queue name is** *queue_name*. **The MQI API is** *mqi_api_name*. **Reason code:** *reason_code*.

Explanation: An unsuccessful reason code was generated by the WebSphere MQ API. Possible reasons are:

- WebSphere MQ is down at the specified server.
- The queue manager is not started.
- The queue name does not exist.

The Monitor program could not complete the action.

User response: Make sure that WebSphere MQ is running at the specified server and that the necessary objects exist, such as the queue manager and queue names.

ASN5150W **MONITOR** *monitor_qualifier*. **The program** *program_name* **is not running. The server is** *server_name*, **and the schema is** *schema*.

Explanation: The alert condition STATUS indicates

that the program specified in this message is not running.

User response: Verify the status of the program on the specified server.

ASN5151W **MONITOR** *monitor_qualifier*. **The elapsed time since the last commit of the Capture program exceeds the threshold value. The Capture control server is** *capture-server*. **The schema is** *schema*. **The last commit time is** *time*. **The threshold is** *seconds* **seconds.**

Explanation: The CAPTURE_LASTCOMMIT alert condition detects that the difference between the current timestamp value and the value of the MAX_COMMIT_TIME column in the IBMSNAP_RESTART table is greater than the threshold value of this alert condition as specified by the PARM_INT column value in the IBMSNAP_CONDITIONS table.

User response: Check the Capture control server, and determine the reason for the commit delay.

ASN5152W **MONITOR** *monitor_qualifier*. **The current Capture latency exceeds the threshold value. The Capture control server is** *capture_server*. **The schema is** *schema*. **The Capture latency is** *latency* **seconds. The threshold is** *threshold* **seconds.**

Explanation: The CAPTURE_CLATENCY alert condition detects that the difference between the CURR_COMMIT_TIME and the MAX_COMMIT_TIME column values in the IBMSNAP_RESTART table is greater than the threshold value of this alert condition as specified by the PARM_INT column value in the IBMSNAP_CONDITIONS table.

User response: Check the Capture control server, and determine the reason for the Capture latency.

ASN5153W **MONITOR** *monitor_qualifier*. **The latency exceeds the threshold value for program** *program_name*. **The server is** *server_name*. **The schema is** *schema*. **The latency is** *latency* **seconds. The threshold is** *threshold* **seconds.**

Explanation: One of the following has occurred:

- While processing the CAPTURE_HLATENCY alert condition, the Monitor program detected that the difference between the MONITOR_TIME and the SYNCTIME column values in the IBMSNAP_CAPMON table is greater than the threshold value. The threshold value of the alert condition is specified in the PARM_INT column value in the IBMSNAP_CONDITIONS table.
- While processing the QCAPTURE_LATENCY alert condition, the Monitor program detected that the

difference between the MONITOR_TIME and the CURRENT_LOG_TIME column values in the IBMQREP_CAPMON table is greater than the threshold value. The threshold value of the alert condition is specified in the PARM_INT column value in the IBMSNAP_CONDITIONS table.

User response: Check the server to determine why the latency threshold is being exceeded.

ASN5154W MONITOR *monitor_qualifier*. The memory used by the program *program_name* exceeds the threshold value. The server is *server_name*. The schema is *schema*. The amount of memory used is *memory* bytes. The threshold is *threshold* megabytes.

Explanation: While processing the CAPTURE_MEMORY or the QCAPTURE_MEMORY alert condition, the Monitor program detected that the value of the CURRENT_MEMORY column in the IBMSNAP_CAPMON table or in the IBMQREP_CAPMON table is greater than the threshold. The threshold value of the alert condition is specified in the PARM_INT column in the IBMSNAP_CONDITIONS table.

User response: Check the server to determine the reason for the excessive memory usage. Increase the memory_limit parameter of the program, if necessary.

ASN5155W MONITOR *monitor_qualifier*. The memory used by the Q Apply program exceeds the threshold value. The Q Apply server is *qapply_server*. The Q Apply schema is *schema*. The receive queue is *queue_name*. The amount of memory used is *memory* bytes. The threshold is *threshold* megabytes.

Explanation: While processing the QAPPLY_MEMORY alert condition, the Monitor program detected that the value of the CURRENT_MEMORY column in the IBMQREP_APPLYMON table is greater than the threshold value. The threshold value of the alert condition is specified in the PARM_INT column in the IBMSNAP_CONDITIONS table.

User response: Check the Q Apply server to determine the reason for the excessive memory usage. If necessary, increase the value of the memory_limit parameter for the replication queue map that uses the specified receive queue.

ASN5156W MONITOR *monitor_qualifier*. The size of a transaction exceeded the threshold for the program *program_name*. The server is *server_name*. The schema is *schema*. The transaction size is *transaction_size* bytes. The threshold is *threshold* megabytes.

Explanation: While processing the CAPTURE_TRANSIZE or the QCAPTURE_TRANSIZE alert condition, the Monitor program detected that the value of the column in the IBMSNAP_CAPMON or IBMQREP_CAPMON table is greater than the threshold value. The threshold value for the alert condition is specified in the PARM_INT column in the IBMSNAP_CONDITIONS table.

User response: Check the server to determine the reason for the excessive transaction size.

ASN5157W MONITOR *monitor_qualifier*. The Q subscription *subscription_name* is inactive. The server is *server_name*. The schema is *schema*. State information: *stateinfo*.

Explanation: While processing the QCAPTURE_SUBSINACT alert condition, the Monitor program detected an inactive Q subscription in the IBMQREP_SUBS table.

User response: Check the IBMQREP_SUBS table on the server to determine why the Q subscription was put in the INACTIVE state.

ASN5160W MONITOR *monitor_qualifier*. The Apply program is not running. The Apply control server is *apply_server*, and the Apply qualifier is *apply_qualifier*.

Explanation: The APPLY_STATUS alert condition detects that the Apply program is not running.

User response: Verify the status of the Apply program at the specified Apply control server.

ASN5161W MONITOR *monitor_qualifier*. The subscription set is inactive and is in an error state. The Apply control server is *apply-control-server*. The Apply qualifier is *apply_qualifier*. The name of the subscription set is *set-name*. The WHOS_ON_FIRST value is *wof*.

Explanation: The APPLY_SUBSINACT alert condition detects that the subscription set is inactive and that the status is not equal to zero (0).

User response: If the subscription set should be active, check this subscription set on the Apply control server.

ASN5162W MONITOR *monitor_qualifier*. A full refresh occurred. The Apply control server is *apply-control-server*. The Apply qualifier is *apply_qualifier*. The name of the subscription set is *set-name*. The WHOS_ON_FIRST value is *wof*.

Explanation: The APPLY_FULLREFRESH alert condition detects that the target table was refreshed during the past monitor cycle.

User response: If the full refresh was in error, verify the cause of the full refresh for this specified target table.

ASN5163W MONITOR *monitor_qualifier*. The subscription is delayed beyond the threshold. The Apply control server is *apply-control-server*. The Apply qualifier is *apply-qualifier*. The name of the subscription set is *set-name*. The WHOS_ON_FIRST value is *wof*. The time delayed is *time*, and the threshold is *threshold* seconds.

Explanation: The APPLY_SUBSDELAYED alert condition detected a subscription set that complies with the following conditions: CURRENT_TIMESTAMP minus LASTRUN is greater than the threshold.

User response: Check any previous messages to see if this subscription set has an error and to verify that the Apply program is running.

ASN5164W MONITOR *monitor_qualifier*. The number of rows that were reworked in a subscription set exceeded the threshold. The Apply control server is *apply-control-server*. The Apply qualifier is *apply-qualifier*. The name of the subscription set is *set-name*. The WHOS_ON_FIRST value is *whos-on-first*. The number of reworked rows is *rows*, and the threshold is *threshold* rows.

Explanation: The APPLY_REWORKED alert condition detects a subscription set with a value in the SET_REWORKED column of the IBMSNAP_APPLYTRAIL table that exceeds the specified threshold. The Apply program reworks changes under the following conditions:

- If an insert fails because the row already exists in the target table, the Apply program converts the insert to an update of the existing row.
- If the update fails because the row does not exist in the target table, the Apply program converts the update to an insert.

User response: Look for other messages to determine the reason that this many rows were reworked.

ASN5165W MONITOR *monitor_qualifier*. Transactions have been rejected in the subscription set. The Apply control server is *apply-control-server*. The Apply qualifier is *apply-qualifier*. The name of the subscription set is *set-name*. The WHOS_ON_FIRST value is *wof*. The number of rejected transactions is *transactions*.

Explanation: The APPLY_TRANSREJECT alert

condition detects rejected transactions for this subscription set.

User response: Verify the reason why these transactions were rejected.

ASN5166W MONITOR *monitor_qualifier*. A manual full refresh is required. The Apply control server is *apply-control-server*. The Apply qualifier is *apply-qualifier*. The name of the subscription set is *set-name*.

Explanation: A full refresh is needed for the specified subscription set.

User response: Verify the reason why a full refresh is required.

ASN5167W MONITOR *monitor_qualifier*. End-to-end latency exceeded the threshold. The Apply control server is *apply-control-server*. The Apply qualifier is *apply-qualifier*. The name of the subscription set is *set-name*. End-to-End latency is *latency* seconds, and the threshold is *threshold* seconds.

Explanation: The APPLY_LATENCY alert condition detects that the end-to-end latency of this subscription set is greater than the threshold value of this alert condition as specified by the PARM_INT column value in the IBMSNAP_CONDITIONS table.

User response: Check the Apply control server to determine the reason for this excessive end-to-end latency value.

ASN5168W MONITOR *apply_control_server*. The processing of a subscription set finished in error. The Apply control server is *control_server_name*, and the Apply qualifier is *apply_qualifier*.

Explanation: The Replication Alert Monitor detected a failed subscription set that has the following information in the IBMSNAP_APPLYTRAIL table:

- STATUS column is not 0 or 2
- APPERRM column is NULL

User response: Examine the data in the IBMSNAP_APPLYTRAIL table at the Apply control server. If the row for the specified Apply qualifier does not contain information that indicates the source of the problem, contact IBM Software Support.

ASN5171W MONITOR *monitor_qualifier*. The Q Apply latency exceeds the threshold value. The server is *server_name*. The Q Apply schema is *schema*. The latency is *latency* milliseconds. The threshold is *threshold* seconds.

Explanation: While processing the QAPPLY_LATENCY alert condition, the Monitor program detected that the value of the APPLY_LATENCY column in the IBMQREP_APPLYMON table is greater than the threshold value. The threshold value of the alert condition is specified in the PARM_INT column in the IBMSNAP_CONDITIONS table.

User response: Check the Q Apply server to determine why the latency threshold is being exceeded.

ASN5172W **MONITOR** *monitor_qualifier*. The end-to-end latency exceeds the threshold value for the Q Apply program. The server is *server_name*. The schema is *schema*. The end-to-end latency is *eelatenency* milliseconds. The threshold is *threshold* seconds.

Explanation: While processing the QAPPLY_EELATENCY alert condition, the Monitor program detected that the value of the END2END_LATENCY column in the IBMQREP_APPLYMON table is greater than the threshold value of this alert condition. The threshold value is specified in the PARM_INT column value in the IBMSNAP_CONDITIONS table.

User response: Check the Q Apply server to determine why the end-to-end latency is exceeding the threshold.

ASN5173W **MONITOR** *monitor_qualifier*. An exception was found for the Q Apply program. The server is *server_name*. The schema is *schema*. The receive queue is *queue_name*. The reason code is *reason_code*. The SQLCODE is *sqlcode*. The SQLSTATE is *sqlstate*. The SQLERRMC is *sql_tokens*.

Explanation: While processing the QAPPLY_EXCEPTIONS alert condition, the Monitor program detected a row in the IBMQREP_EXCEPTIONS table. The SQLCODE or SQLSTATE explains the reason for the exception.

User response: Check the QAPPLY_EXCEPTIONS table on the server to determine the reason for the exception.

ASN5174W **MONITOR** *monitor_qualifier*. The depth of the spill queue *spill_queue_name* for Q subscription *subscription_name* and receive queue *receive_queue_name* exceeds the threshold value. The server is *server_name*. The schema is *schema*. The current queue depth is *current_depth%*. The threshold is *threshold%*.

Explanation: While processing the

QAPPLY_SPILLQDEPTH alert condition, the Monitor program detected that the depth of the spill queue for the associated receive queue exceeds the threshold value. The threshold value of the alert condition is in the PARM_INT column in the IBMSNAP_CONDITIONS table.

User response: Check the IBMQREP_APPLYTRACE table for additional information. Also, check the data in the IBMQREP_SPILLQS table, if available.

ASN5175W **MONITOR** *monitor_qualifier*. The depth of the queue *queue_name* exceeds the threshold value for the program *program_name*. The server is *server_name*. The schema is *schema*. The current queue depth is *current_depth%*. The threshold is *threshold%*.

Explanation: While processing the QAPPLY_QDEPTH alert condition, the Monitor program detected that the depth of the queue exceeded the threshold value. The threshold value, specified as a percent, is in the PARM_INT column in the IBMSNAP_CONDITIONS table. The value is obtained by dividing the current depth by the maximum depth attribute of the queue, and then multiplying it by 100 .

User response: Verify that the program is processing the transactions for the queue as expected. Also, verify that the queue is large enough for the volume of transactions being processed.

ASN5176I **MONITOR** : *monitor_qualifier* : Monitoring of server *monitored_server* was suspended at *timestamp*. The monitor suspension is *suspension_name*, and the suspension template is *template_name*.

Explanation: All monitoring activities for the specified server are suspended because the monitor cycle is within the time, day, and duration for the specified suspension and template.

User response: This message is for your information. No action is required.

ASN5177I **MONITOR** : *monitor_qualifier* : Monitoring activities for server *monitored_server* were resumed at *timestamp*.

Explanation: All previously suspended monitoring activities for the specified server are resumed because the monitor cycle is later than the time, day or duration for any suspensions defined for the specified server.

User response: This message is for your information. No action is required.

ASN5178W Monitor *monitor_qualifier* : **The send queue** *send_queue_name* **is inactive. The server is** *server_name*. **The Q Capture schema is** *schema_name*. **State information:** *information*.

Explanation: While processing the QCAPTURE_SENDQSINACT alert condition, the monitor program detected an inactive send queue in the STATE column of the IBMQREP_SENDQUEUES table.

User response: Use the replication administration tools or the state information that is provided in the message text to determine why the send queue was put in the INACTIVE state by Q Capture.

ASN5179W Monitor *monitor_qualifier* : **The receive queue** *receive_queue_name* **is inactive. The server is** *server_name*. **The Q Apply schema is** *schema_name*. **State information:** *information*. **State Time:** *timestamp*.

Explanation: While processing the QAPPLY_RECVQSINACT alert condition, the monitor program detected an inactive receive queue in the STATE column of the IBMQREP_RECVQUEUES table.

User response: Use the replication administration tools or the state information that is provided in the message text to determine why the receive queue was put in the INACTIVE state by Q Apply.

ASN5180W Monitor *monitor_qualifier* : **The number of deadlock retries for transactions on Q Apply receive queue** *receive_queue_name* **exceeded the threshold value. The server is** *server_name*. **The Q Apply schema is** *schema_name*. **The current number of deadlock retries is** *number*. **The alert threshold for deadlock retries is** *threshold*.

Explanation: While processing the QAPPLY_DEADLOCKRETRIES alert condition, the Monitor program detected that the number of deadlock retries for transactions on the receive queue exceeded the alert threshold. For details on how the threshold value is determined, see "Alert conditions for the Replication Alert Monitor" in the IBM Information Management Software for z/OS Solutions Information Center or DB2 Information Center.

User response: Check the environment. If a z/OS batch job is running, update the MAXAGENT_CORRELID value for the replication queue map that contains the receive queue to half the value of NUM_APPLY_AGENTS for the queue map. Both of these values are specified in the IBMQREP_RECVQUEUES table. For other OLTP

transactions, check the environment to reduce the number of deadlocks.

ASN5181W Monitor *monitor_qualifier* : **The Q Capture program did not send data for large object (LOB) columns because the transaction size exceeds the maximum message size for the send queue that was specified in the replication queue map or publishing queue map. The server is** *server_name*. **The schema is** *schema_name*. **The send queue is** *queue_name*. **The number of rejected LOB columns that could not be applied because they were too big to fit in the maximum message size is** *number*.

Explanation: The Q Capture program processed the row that contains LOB values but rejected the LOB data because it exceeded the value that was specified for max_message_size for the replication queue map or publishing queue map that contains the send queue. The data was too large for the transaction message.

The QCAPTURE_LOBSTOOBIG alert condition detects the number of LOB columns that were not sent because they did not fit within the max_message_size value.

User response: Use the replication administration tools to increase the max_message_size value for the replication queue map or publishing queue map. You might also need to increase the size of the MAXMSGL parameter for the WebSphere MQ send queue. Ensure that the value for the max_message_size parameter is at least 4KB smaller than value for the MAXMSGL parameter. If you cannot increase the max_message_size parameter, you can set the LOB_SEND_OPTION value in the IBMQREP_CAPPARMS table to S to send the LOB values in a LOB message after the transaction message. You need to issue the reinit command in either case. To synchronize the source and target tables, run the asntdiff and asntrep utilities.

ASN5182W Monitor *monitor_qualifier* : **The Q Capture program did not send data for XML columns because the transaction size exceeds the maximum message size of the send queue that was specified in the replication queue map or publishing queue map. The server is** *server_name*. **The schema is** *schema_name*. **The send queue is** *queue_name*. **The number of rejected XML columns that could not be applied because they were too big to fit in the maximum message size is** *number*.

Explanation: The Q Capture program processed the row that contained XML values but rejected the XML data because it exceeded the value that was specified for max_message_size for the replication queue map or publishing queue map that contains the send queue.

The data was too large for the transaction message.

The QCAPTURE_XMLDOCSTOOBIG alert condition detects the number of XML columns that were not sent because they did not fit within the max_message_size value.

User response: Use the replication administration tools to increase the max_message_size value for the replication queue map or publishing queue map. You might also need to increase the size of the MAXMSGL parameter for the WebSphere MQ send queue. Ensure that the value for the max_message_size parameter is at least 4KB smaller than value for the MAXMSGL parameter. Run the reinitq command to prompt the Q Capture program to read the changes to the queue map. To synchronize the source and target tables, run the asntdiff and asntrep utilities.

ASN5183W Monitor *monitor_qualifier* : The Q Capture program exceeded the threshold for the number of retries to put messages on the send queue. The Q Capture server is *server_name*. The schema is *schema_name*. The WebSphere MQ send queue is *queue_name*. The number of retries is *full_queue_retries* and the monitor alert threshold for retries is *threshold*.

Explanation: The Q Capture program could not put a message on the send queue using the MQPUT command. This condition can occur for several reasons:

- The Q Apply program is not processing messages on the receive queue.
- The Q Apply program is stopped.
- The Q Capture program is putting messages on the queue faster than WebSphere MQ can process them.
- Many changes accumulated at the source table while the Q Capture program was stopped.

This error is more likely to happen in a queue-sharing environment.

The QCAPTURE_QFULLNUMRETRIES alert condition detects a send queue with a QFULL_NUM_RETRIES column value in the IBMSNAP_SENDQUEUES table that exceeds the specified threshold.

User response: Check for the possible causes listed in the "Explanation" section of this message. Also, check the WebSphere MQ reason code that the Q Capture program issued in its diagnostic log and take the appropriate action.

ASN5184E MONITOR *monitor_qualifier* : An SQL error occurred when the monitor program was querying a monitor control table to determine status for program *program_name*. The monitor table name is *table_name*, the server name is *server_name*, the SQLCODE is *sqlcode*, the SQLSTATE is *sqlstate*.

Explanation: An error occurred while the monitor program was trying to access the specified program's monitor table to check program status. The monitor could not determine the status.

User response: Use the SQLCODE and SQLSTATE information to correct the problem. The monitor will continue running.

ASN5185E MONITOR *monitor_qualifier* : An SQL error occurred when the monitor program queried the *table_name* table to determine status for the Apply program with Apply qualifier *apply_qualifier*. The server name is *server_name*, the SQLCODE is *sqlcode*, and the SQLSTATE is *sqlstate*. The monitor could not determine the program status.

Explanation: An SQL error occurred that prevented the monitor program from accessing the control table to check Apply program status. The monitor continues running.

User response: Use the SQLCODE and SQLSTATE information to correct the problem. After the problem is corrected, the monitor program will issue an alert if the program status meets the specified alert condition.

ASN5190E MONITOR *monitor_qualifier*. An error occurred running the stored procedure *stored_proc_name*. The condition name is *condition_name*. The server is *server_name*. The schema is *schema*. The reason code is *reason_code*. The SQLCODE is *sqlcode*. The SQLSTATE is *sqlstate*. The SQLERRMC is *sql_tokens*.

Explanation: The Monitor program was not able to run the specified stored procedure. The SQLCODE and SQLSTATE explain the reason.

User response: Check the SQLCODE and SQLSTATE. At the server, verify that the definition of the stored procedure is correctly specified.

ASN5191W MONITOR *monitor_qualifier*. message

Explanation: This message appears when a user-defined alert condition issues a warning message.

User response: Read the issued message, and take appropriate action.

ASN5192E **MONITOR** *monitor_qualifier. message*

Explanation: This message appears when a user-defined alert condition issues an error message.

User response: Read the issued message, and take appropriate action.

ASN5193E **MONITOR** *monitor_qualifier. The stored procedure stored_proc_name was not found at server server_name. The condition name is condition_name. The schema is schema.*

Explanation: The stored procedure that is required by the Monitor program was not found at the server while processing the alert condition.

User response: Verify that the CREATE PROCEDURE statement was issued at the correct server for the stored procedure name specified in the message.

ASN5194I **MONITOR** *monitor_qualifier : The monitor program will monitor alert conditions within the following time range time_range for monitored server server_name.*

Explanation: The monitor program looks for alert conditions that fall within the specified time range.

User response: This message is for your information only. No action is required.

ASN5200E **ASNSCRT:** The replication process type is a required parameter and must be specified when invoking the asnsrct command.

Explanation: The asnsrct command was invoked without a specified replication process type.

User response: Enter the command again with a replication process type of -C, -A, or -M.

ASN5201E **ASNSCRT:** The database instance is a required parameter and must be specified when invoking the asnsrct command.

Explanation: The asnsrct command was invoked without a specified database instance.

User response: Enter the command again with a database instance name.

ASN5202E **ASNSCRT:** A replication process path is a required parameter and must be specified when invoking the asnsrct command.

Explanation: The asnsrct command was invoked without a specified replication process path.

User response: Enter the command again with a path to the asncap, asnaply, or asnmon command.

ASN5203E **ASNSCRT:** The Capture server is a required parameter and must be specified when invoking this asnsrct command.

Explanation: The asnsrct command was invoked without a specified Capture control server.

User response: Enter the command again with a Capture control server name.

ASN5204E **ASNSCRT:** The Apply control server is a required parameter and must be specified when invoking this asnsrct command.

Explanation: The asnsrct command was invoked without a specified Apply control server.

User response: Enter the command again with an Apply control server name.

ASN5205E **ASNSCRT:** The Apply qualifier is a required parameter and must be specified when invoking this asnsrct command.

Explanation: The asnsrct command was invoked without a specified Apply qualifier.

User response: Enter the command again with an Apply qualifier.

ASN5206E **ASNSCRT:** The monitor control server is a required parameter and must be specified when invoking this asnsrct command.

Explanation: The asnsrct command was invoked without a specified monitor control server.

User response: Enter the command again with a monitor control server name.

ASN5207E **ASNSCRT:** The monitor qualifier is a required parameter and must be specified when invoking this asnsrct command.

Explanation: The asnsrct command was invoked without a specified monitor qualifier.

User response: Enter the command again with a monitor qualifier.

ASN5208I **ASNSCRT: The replication service**
service_name **was created successfully.**

Explanation: The asnsCRT command successfully created the specified service.

User response: This message is for your information only, and no action is required.

ASN5209I **ASNSCRT: The replication service**
service_name **started successfully.**

Explanation: The asnsCRT command successfully started the specified service.

User response: This message is for your information only, and no action is required.

ASN5210E **ASNSCRT: The replication service**
service_name **was not created, because the display name already exists (either as a service name or as another display name) in the service control manager database.**

Explanation: The asnsCRT command cannot create the specified service, because the display name already exists as another service name or display name in the service control manager database.

User response: Go to service control manager database, and remove the service with the duplicate service or display name. Then re-enter the command.

ASN5211E **ASNSCRT: The replication service**
service_name **was not created, because the specified service name is not valid.**

Explanation: The asnsCRT command cannot create the specified service, because the system API returned an error code indicating that the service name is incorrect. The specified service might contain special characters in the instance name, database name, or schema name. Special characters are not allowed in the service name.

User response: Change the instance name, database name, or schema name if possible. Then re-enter the command.

ASN5212E **ASNSCRT: The replication service**
service_name **was not created, because the specified service name already exists.**

Explanation: The asnsCRT command cannot create the specified service, because another service with the same service name already exists in the service control manager.

User response: Remove the existing service with the same service name. Then re-enter the command.

ASN5213E **ASNSCRT: The replication service**
service_name **was not started, because the service binary file could not be found.**

Explanation: The asnsCRT command cannot start the specified service, because the corresponding asncap, asnappl, or asnmon command cannot be invoked using the system path specified by the PATH environment variable. If the fully qualified path is provided, the asnsCRT command cannot find the asncap, asnappl or asnmon command in that path.

User response: Make sure the specified path is correct. Then re-enter the command.

ASN5214E **ASNSCRT: The replication service**
service_name **did not start, because an instance of the service is already running.**

Explanation: The asnsCRT command cannot start the specified service, because the service is already running.

User response: This message is for your information only, and no action is required.

ASN5215E **ASNSCRT: The replication service**
service_name **did not start, because the service depends on a DB2 instance service that does not exist or has been marked for deletion.**

Explanation: The ASNSCRT command cannot start the specified service, because the corresponding DB2 instance service does not exist or has been deleted.

User response: Verify that the corresponding DB2 instance service exists in the service control manager. Then reissue the command.

ASN5216E **ASNSCRT: The replication service**
service_name **did not start, because this service depends on another service that failed to start.**

Explanation: The ASNSCRT command cannot start the specified service, because the corresponding DB2 instance service failed to start.

User response: Verify that the corresponding DB2 instance service started in the service control manager. Then reissue the command.

ASN5217E **ASNSCRT: The replication service**
service_name **did not start, because the service is disabled.**

Explanation: The asnsCRT command cannot start the specified service, because the service has been disabled in the service control manager.

User response: Verify that the service startup type is

set to either automatic or manual in the service control manager. Then re-enter the command.

ASN5218E **ASNSCRT: The replication service *service_name* did not start, because the service cannot log on. This error occurs if the service starts from an account that does not have the proper "Log on as a service" access right.**

Explanation: The **ASNSCRT** command cannot start the specified service, because the corresponding DB2 instance service cannot log on.

User response: Go to service control manager, and locate the specified service. Verify that the provided account name and passwords are correct. Then reissue the command.

ASN5219E **ASNSCRT: The replication service *service_name* was not created, because the service is marked for deletion.**

Explanation: The **asnsCRT** command cannot create the specified service, because the service was deleted.

User response: Close the service control manager window. Then re-enter the command.

ASN5220E **ASNSDROP: The service name is a required parameter and must be specified when invoking the **asnsdrop** command.**

Explanation: The **asnsdrop** command was invoked without a specified service name.

User response: Re-enter the command with a service name.

ASN5221I **ASNSDROP: The service *service_name* has been successfully removed.**

Explanation: The **asnsdrop** command was invoked with a specified service name.

User response: This message is for your information only, and no action is required.

ASN5222E **ASNSDROP: The replication service *service_name* cannot be removed, because the requested access is denied.**

Explanation: The **ASNSDROP** command cannot remove the specified service name, because the user does not have the appropriate permission to remove it.

User response: Verify that the current user has permission to log on to the corresponding DB2 instance. Then reissue the command.

ASN5223E **ASNSDROP: The replication service *service_name* cannot be removed, because the specified service name is not valid.**

Explanation: The **asnsdrop** command cannot remove the specified service name, because the service name contains illegal special characters.

User response: Go to service control manager, and locate the specified service. Verify that the service name is valid, and re-enter the command.

ASN5224E **ASNSDROP: The replication service *service_name* cannot be removed, because the specified service does not exist.**

Explanation: The **asnsdrop** command cannot remove the specified service name, because the service name does not exist in the service control manager.

User response: Go to service control manager, and locate the specified service. Verify that the service name is correct, and re-enter the command.

ASN5225E **ASNSDROP: The replication service *service_name* cannot be stopped, because other running services are dependent on it. The *service_name* is not removed.**

Explanation: The **asnsdrop** command cannot remove the specified service, because other services that are dependent on this specified service are currently running.

User response: Go to service control manager, and stop all services that are dependent on this specified service. Then re-enter the command.

ASN5226E **ASNSDROP: The replication service *service_name* cannot be removed, because the system is shutting down.**

Explanation: The **asnsdrop** command cannot remove the specified service, because the operating system is shutting down.

User response: Enter the command again after system restarts.

ASN5227I **ASNSDROP: The replication service *service_name* cannot be removed, because it has already been marked for deletion.**

Explanation: The **asnsdrop** command cannot remove the specified service, because the specified service has already been deleted.

User response: This message is for your information only, and no action is required.

ASN5228E *pgmname* : **The command cannot**
command_action **the replication service**
service_name, **because the system call**
API_func_name **returned an unexpected**
error code *error_code*.

Explanation: The asnsrct and asndrop commands make system calls in order to work with the services. The specified system call returned an unexpected error code that prevents the given command from completing the requested action.

User response: Enter the command again. This error code might indicate only a temporary system condition. For further information about the error code, check your operating system documentation.

ASN5229E **ASNSCRT: The account is a required parameter and must be specified when invoking the asnsrct command.**

Explanation: The **ASNSCRT** command was invoked without a specified account name for the corresponding DB2 instance.

User response: Reissue the command with an account name for the corresponding DB2 instance.

ASN5230E **ASNSCRT: The password is a required parameter and must be specified when invoking the asnsrct command.**

Explanation: The **ASNSCRT** command was invoked without a specified password for the corresponding DB2 instance.

User response: Reissue the command with the password for the corresponding DB2 instance.

ASN5231E **ASNSCRT: The replication service**
service_name **was not created, because the**
account name specified on the account
parameter does not exist.

Explanation: The **ASNSCRT** command was invoked with an unknown account name for the corresponding DB2 instance.

User response: Verify that the provided account name and passwords are correct. Then reissue the command.

ASN5232E **ASNSCRT: The required parameter** *path*
was not specified.

Explanation: When started as a service, a replication command must contain a path specified by the path keyword that is specific to the command (*capture_path* for asncap, *apply_path* for asnappl, and *monitor_path* for asnmon). If the path keyword is specified, the service is registered if no errors occur.

If the path keyword is not specified, the asnsrct command retrieves the DB2 global registry profile

variable DB2PATH. If this variable contains a non-null value, the asnsrct command adds the appropriate path keyword to the command using the value of DB2PATH. If this variable is not set, the asnsrct command cannot register the service.

User response: Enter the command again after you specify the appropriate path keyword or after you define the DB2 global registry profile variable DB2PATH.

ASN5233E **ASNSCRT: The replication service**
service_name **was not created, because the**
specified database instance
database_instance **does not exist.**

Explanation: The **asnsrct** command cannot create the specified service because the specified database instance does not exist in the system.

User response: Verify that the specified database instance exists in the system. Then reissue the command.

ASN5234E **ASNSCRT: The capture_server**
parameter is not specified.

Explanation: The asnsrct command was entered without a Q Capture server name. The *capture_server* parameter is required.

User response: Enter the asnsrct command again, with the *capture_server* parameter set to the Q Capture server name.

ASN5235E **ASNSCRT: The apply_server parameter**
is not specified.

Explanation: The asnsrct command was entered without a Q Apply server name. The *apply_server* parameter is required.

User response: Enter the asnsrct command again, with the *apply_server* parameter set to the Q Apply server name.

Chapter 44. ASN6000 - ASN6499

ASN6000W **QCAPLATENCY** *max-latency* **EXCEEDS**
threshold FOR schema

Explanation: The average Q Capture latency during the last monitor interval exceeded the alert threshold. Q Capture latency measures the average time between transactions being committed at the source database and the Q Capture program sending the transaction to be replicated. The threshold was set for the QCAPTURE_LATENCY alert condition in the Replication Alert Monitor program.

User response: To reduce Q Capture latency, use the Performance Advisor in the Q Replication Dashboard or follow the instructions in "Actions to take when Q Replication latency is too high."

ASN6001W **TRANSIZE** *max-transize* **EXCEEDS**
threshold FOR schema.

Explanation: One or more replicated transactions exceeded the threshold for memory usage during the last monitor interval. The threshold was set for the QCAPTURE_TRANSIZE alert condition in the Replication Alert Monitor program.

User response: If the amount of memory use is not acceptable, you can increase the value of the Q Capture memory_limit parameter.

ASN6002W **QSUBINACT** *subname* **FOR schema**

Explanation: The Q Capture program deactivated the Q subscription, which resulted in the QCAPTURE_SUBSINACT alert condition in the Replication Alert Monitor.

User response: Look for messages from the Q Capture program in the Q Replication Dashboard, Q Capture log, or IBMQREP_CAPTRACE table that state why the Q subscription was deactivated.

ASN6003W **QCAPMEM** *max-memsize* **EXCEEDS**
threshold FOR schema

Explanation: Total memory usage by the Q Capture program during the last monitor interval exceeded the threshold that was set for the QCAPTURE_MEMORY alert condition in the Replication Alert Monitor.

User response: [Optional] If the amount of memory use that is reported in the *max-memsize* runtime token in this message is acceptable, increase the value of the Q Capture memory_limit parameter to a value that is larger than the amount that is reported in the *max-memsize* runtime token.

ASN6004W **QAPPLATENCY** *max-latency* **EXCEEDS**
threshold FOR schema recvq

Explanation: The average Q Apply latency during the last monitor interval exceeded the alert threshold. Q Apply latency measures the time it takes for a transaction to be applied to a target table after the Q Apply program gets the transaction from a receive queue. The threshold was set for the QAPPLY_LATENCY alert condition in the Replication Alert Monitor program.

User response: To reduce Q Apply latency, use the Performance Advisor in the Q Replication Dashboard or follow the instructions in "Actions to take when Q Replication latency is too high."

ASN6005W **AVGE2ELATENCY** *max-latency*
EXCEEDS *threshold FOR schema recvq*

Explanation: The average end-to-end latency during the last monitor interval exceeded the alert threshold. End-to-end latency measures the total time that replication requires to capture changes and apply those changes to a target database. The threshold was set for the QAPPLY_EELATENCY alert condition in the Replication Alert Monitor program.

User response: To reduce Q Apply latency, use the Performance Advisor in the Q Replication Dashboard or follow the instructions in "Actions to take when Q Replication latency is too high."

ASN6006W **QAPPMEM** *max-mem* **EXCEEDS** *threshold*
FOR schema recvq

Explanation: The peak memory usage by the Q Apply program for processing transactions from this receive queue during the last monitor interval exceeded an alert threshold. The threshold was set for the QAPPLY_MEMORY alert condition in the Replication Alert Monitor program.

User response: To allocate more memory for processing transactions from this receive queue, increase the value of the memory_limit parameter for the replication queue map that includes the receive queue.

ASN6007W **QAPPEXCPT** *sqlcode*
SQLSTATE *sqlstate* **FOR schema recvq**

Explanation: The Q Apply program inserted a row in the IBMQREP_EXCEPTIONS table because of a conflict or SQL error at a target. The Replication Alert Monitor reported the insertion because the QAPPLY_EXCEPTIONS alert condition is enabled.

User response: Use the SQLCODE and SQLSTATE values to determine the cause of the exception. You can use the Q Replication Dashboard to get more detail about the exception, or look at the IBMQREP_EXCEPTIONS table.

ASN6008W RECVQINACT FOR *schema recvq*

Explanation: The Q Apply program stopped processing messages on the specified receive queue. This action is also known as deactivating or stopping the receive queue. The Replication Alert Monitor reported the action because the QAPPLY_RECVQINACT alert condition is enabled.

User response: Look for messages from the Q Apply program that explain why it stopped processing messages on the receive queue.

ASN6009W QDEPTH *max-qdepth* EXCEEDS *threshold* FOR *schema qname*

Explanation: The fullness of the specified receive queue or spill queue exceeded an alert threshold. The Replication Alert Monitor reported the action because one or both of the alert conditions QAPPLY_QDEPTH or QAPPLY_SPILLQDEPTH are enabled.

User response: Take one of the following actions to reduce the fullness of the receive queue or spill queue:

Receive queue

Follow the instructions in "Checking for and preventing a full receive queue."

Spill queue

If possible, stop the Q Capture program. If necessary, increase the value for the MAXDEPTH attribute of the spill queue.

ASN6010W PGMDOWN *pgmtype* FOR *schema*

Explanation: The program stopped. The Replication Alert Monitor reported the action because one or both of the alert conditions QCAPTURE_STATUS or QAPPLY_STATUS are enabled.

User response: If the stoppage was unexpected, look for messages from the program to indicate why it stopped.

ASN6011W AVG END-TO-END LATENCY LATENCY_VALUE IS ABOVE THRESHOLD FOR SCHEMA_NAME RECVQ_NAME IN WORKLOAD_NAME.

Explanation: An average end-to-end latency event was defined in the IBMQREP_APPEVTDEFS table. The average latency during the event interval exceeded the maximum threshold.

User response: Check the Q Apply server to

determine why the latency threshold was exceeded. Further actions might be needed to reduce latency.

ASN6012I AVG END-TO-END LATENCY LATENCY_VALUE IS BELOW THRESHOLD FOR SCHEMA_NAME RECVQ_NAME IN WORKLOAD_NAME.

Explanation: An average end-to-end latency event was defined in the IBMQREP_APPEVTDEFS table. The average latency had exceeded the maximum threshold but now is back within the required specification.

User response: This message is for your information only. No action is required.

ASN6013W PIT CONSISTENCY LATENCY LATENCY_VALUE IS ABOVE THRESHOLD FOR SCHEMA_NAME RECVQ_NAME IN WORKLOAD_NAME.

Explanation: A latency event for point-in-time (PIT) consistency was defined in the IBMQREP_APPEVTDEFS table. The value of PIT-consistency latency exceeds the maximum threshold.

User response: Check the Q Apply server to determine why the maximum threshold was exceeded. Further actions might be needed to reduce the latency.

ASN6014W PIT CONSISTENCY LATENCY LATENCY_VALUE IS BELOW THRESHOLD FOR SCHEMA_NAME RECVQ_NAME IN WORKLOAD_NAME.

Explanation: A latency event for point-in-time (PIT) consistency was defined in the IBMQREP_APPEVTDEFS table. The value of PIT-consistency had exceeded the maximum threshold and is now is back within the required specification.

User response: This message is for your information only. No action is required.

Chapter 45. ASN7000 - ASN7499

ASN7000I *program_name : program_identifier : number*
subscriptions are active. number
subscriptions are inactive. number
subscriptions that were new and were
successfully activated. number
subscriptions that were new could not
be activated and are now inactive.

Explanation: This message describes the number of publications or Q subscriptions that are in various states.

User response: This message is for your information only. No action is required.

ASN7001I *program_name : program_identifier : The*
command_name **command was received.**

Explanation: The command that was received might be processed asynchronously. The program will issue a message when the command is processed and store that message in the trace table (IBMQREP_CAPTRACE).

User response: This message is for your information only. No action is required.

ASN7002E *program_name : program_identifier : The*
send queue queue_name is full. This error
occurred while processing publication or
Q subscription name (publishing or
replication queue map queue_map_name).

Explanation: The number of messages reached the number that is set for the MAXDEPTH attribute on the send queue. The ERROR_ACTION value determines how the program behaves due to this error. The value for the Q Capture program is stored in the IBMQREP_SENDQUEUES table.

User response: Do not clear the queue or all of the messages containing your changes will be removed.

Make sure that Q Apply or a subscribing application is running.

If necessary, increase the value for the MAXDEPTH attribute of the send queue. After you increase the value, take all of the appropriate actions that are necessary:

- If the Q Capture program terminated because of this error, warm start the Q Capture program.
- If the publications or Q subscriptions were deactivated, activate them.

ASN7004E *program_name : program_identifier : The*
source column source_column subscribed
in publication or Q subscription name
does not exist in the source table
table_name.

Explanation: The column shown in the IBMQREP_SRC_COLS table does not exist in the source table. This error occurred due to one of the following problems:

- The column name was specified incorrectly in the publication or Q subscription.
- The source table changed since the publication or Q subscription was created.
- Extra columns were specified in the publication or Q subscription.

Therefore, the publication or Q subscription cannot be activated.

User response: Ensure that the column name is specified correctly in the publication or Q subscription.

If the publication or Q subscription is not correct (for example, because it contains more columns than are available at the source table), use the replication administration tools to perform the following steps:

1. Delete and recreate the publication or Q subscription.
2. Activate the publication or Q subscription.

ASN7006E *program_name : program_identifier : The*
publication or Q subscription name was
deactivated due to an error.

Explanation: The specified publication or Q subscription was deactivated due to an error that was specified in a previous message.

User response: Look in the Q Capture diagnostic log or in the IBMQREP_CAPTRACE table to find the error message that precedes this one. Take the appropriate action to correct the problem (follow the instructions in the user response of the previous message). Activate the publication or Q subscription.

ASN7007E *program_name : program_identifier : The Q*
Capture program cannot warm start
because it detected database partitions
for which it does not have restart
information. Those database partitions
are node_id,...,node_id.

Explanation: This error message occurred due to one of the following reasons:

- One or more new database partitions were added to the database. The Q Capture program could not process the log records for the new partitions because it tried to warm start with the add_partition parameter set to 'N'. This parameter determines whether the Q Capture program starts reading the log file for the partitions that were added since the last time the Q Capture program was started. By default, this parameter is set to 'N'. The add_partition parameter must be set to 'Y' before the Q Capture program can read the log files.
- The restart message in the restart Q was corrupted.

User response: If new database partitions were added, warm start the Q Capture program using the add_partition='Y' option.

If the restart message in the restart queue was corrupted, cold start the Q Capture program.

ASN7008I *program_name : program_identifier : The program was successfully reinitialized. number subscriptions are active. number subscriptions are inactive. number subscriptions that were new were successfully activated. number subscriptions that were new could not be activated and are now inactive.*

Explanation: The REINIT command was processed successfully. As a result, the Q Capture program recognizes any changed attributes for publications, Q subscriptions, and queue maps. Any new publications or Q subscriptions were automatically activated. Refer to the message text for the number of publications or Q subscriptions that are in various states.

User response: This message is for your information only. No action is required.

ASN7010I *program_name : program_identifier : The program successfully activated publication or Q subscription name (send queue queue_name, publishing or replication queue map queue_map_name) for source table table_name.*

Explanation: The Q Capture program is replicating changes for the active publication or Q subscription.

When the Q Capture program starts, it activates any new publications or Q subscriptions. While the Q Capture program runs, it performs the following actions:

- Activates any new publications or Q subscriptions when it reinitializes.
- Activates any new or inactive publications or Q subscriptions when it receives a capstart signal.

User response: This message is for your information only. No action is required.

ASN7011E *program_name : program_identifier : Source table table_name of publication or Q subscription name (send queue queue_name, publishing or replication queue map queue_map_name) was not found.*

Explanation: The source table specified in the publication or Q subscription does not exist. Most likely the wrong table name was specified when attempting to activate the publication or Q subscription. Otherwise, the source table might have been dropped intentionally since the publication or Q subscription was created.

User response: Take one of the following actions:

- Ensure that the source table is specified correctly in the control tables and that it exists in the database. Activate the publication or Q subscription again.
- If you dropped the source table intentionally, delete the publication or the Q subscription.

ASN7012I *program_name : program_identifier : The publication or Q subscription name was successfully reinitialized.*

Explanation: The Q Capture program successfully reinitialized the publication and it will process transactions according to the new attributes.

User response: This message is for your information only. No action is required.

ASN7013I *program_name : program_identifier : The publication or Q subscription name was deactivated.*

Explanation: The publication or Q subscription was deactivated because the CAPSTOP signal was received.

User response: This message is for your information only. No action is required.

ASN7015E *program_name : program_identifier : This is additional information for message ASN0575E. The program detected an unrecoverable WebSphere MQ error for send queue send_queue_name of publishing or replication queue map queue_map_name. The error action specified for the queue map is error_action.*

Explanation: An unrecoverable WebSphere MQ error occurred on this send queue. The Q Capture program stops. The error action I to deactivate the publications or Q subscriptions on the send queue has been deprecated.

User response: Look in the Q Capture diagnostic log file or the IBMQREP_CAPTRACE table to find the error message that precedes this one, typically ASN0575E.

Take the appropriate action to correct the WebSphere MQ problem (follow the instructions in the user response of the previous message). After you fix the error, start the Q Capture program in warm mode.

ASN7016E *program_name : program_identifier : The LOADDONE signal failed because the publication or Q subscription name has no load phase.*

Explanation: The publication or Q subscription does not have a load phase (HAS_LOADPHASE=N). This error could be caused by an incorrect publication or Q subscription name specified in the LOADDONE signal or in the load done received message.

User response:

If the publication or Q subscription has a load phase and you inserted the LOADDONE signal

Make sure that the publication name or Q subscription name in the signal is correct and re-insert the LOADDONE signal.

If the publication or Q subscription has a load phase and you sent the load done received message

Make sure that the publication name or Q subscription name in the message is correct and resend the load done received message.

ASN7017I *program_name : program_identifier : The target table table_name is ready to be loaded from source table table_name for publication or Q subscription name.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7018E *program_name : program_identifier : The source table table_name for publication or Q subscription name does not have the DATA CAPTURE CHANGES attribute set.*

Explanation: The source table is incorrectly defined. No data will be captured for this publication or Q subscription.

User response: Alter the source table so that the DATA CAPTURE CHANGES attribute is set and restart the publication or Q subscription. For example:

```
alter table src_owner.src_table
  data capture changes
```

ASN7019I *program_name : program_identifier : signal_name signal was received and will be processed.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7020I *program_name : program_identifier : The program reached the end of the active log and terminated because the AUTOSTOP option is specified.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7021E *program_name : program_identifier : The publication name did not start because an internal load phase is specified.*

Explanation: For publications, only two load options are supported. The HAS_LOADPHASE value in the IBMQREP_SUBS table must be set to one of the following values:

N The target will not be loaded.

E An application other than the Q Apply program loads the target table.

An internal load option (HAS_LOADPHASE=I) is not supported for publications.

User response: Specify a supported load option for the publication (ensure that the correct HAS_LOADPHASE value is in the IBMQREP_SUBS table) and start the publication.

ASN7022W *program_name : program_identifier : The publication or Q subscription name does not exist in the IBMQREP_SUBS control table. The signal signal was ignored.*

Explanation: The program cannot find this publication or Q subscription in the IBMQREP_SUBS control table.

User response: Ensure that the publication or Q subscription name is accurate for the signal.

ASN7023W *program_name : program_identifier : Source table table_name does not have a primary key, unique constraint, or unique index.*

Explanation: Since the source table does not have a primary key, unique constraint, or unique index, all columns except LOB columns were selected as the replication key. This might impact performance because of the amount of data sent. The related publications or Q subscriptions are still activated.

User response: Create a primary key, a unique constraint, or a unique index on the source table and reinitialize all related publications or Q subscriptions if necessary.

ASN7024E *program_name : program_identifier : The program cannot process the restart message from the restart queue restart_queue.*

Explanation: The program detected an internal inconsistency. This could happen if you are falling back from a higher-level PTF or version of Q Capture to a previous PTF level or previous version. Q Capture saves its restart information using a new format. Q Capture can warm start by using restart information from a prior level of Q Capture, but prior levels of Q Capture will not be able to use the new restart information format to warm restart.

User response: Check that the correct version of the Q Capture program is running and that the restart queue was not improperly altered. If you need to warm restart a prior level of Q Capture, you must start Q Capture from a known point in the log by specifying the LSN and MAXCMTSEQ parameters. Look for message ASN7109I in the Q Capture diagnostic log. The value of LSN is described as "the lowest log sequence number of a transaction still to be committed." The value of MAXCMTSEQ is described as "the highest log sequence number of a successfully processed transaction."

ASN7025E *program_name : program_identifier : The program cannot access the restart queue restart_queue.*

Explanation: See message text.

User response: Make sure the restart queue is correctly specified and accessible to application programs (GET enabled).

ASN7026E *program_name : program_identifier : Invalid signal type signal_type and subtype subtype.*

Explanation: The signal type or signal subtype is not supported.

User response: Check the SIGNAL_TYPE and SIGNAL_SUBTYPE of the row inserted in the IBMQREP_SIGNAL control table. Ensure that the signal type or signal subtype is correct and try inserting the signal again.

ASN7027E *program_name : program_identifier : The restart queue restart_queue is empty.*

Explanation: The Q Capture program could not warm start. See message text.

User response: Cold start the Q Capture program.

ASN7028E *program_name : program_identifier : The MAX_MESSAGE_SIZE specified for send queue queue_name for publishing or replication queue map queue_map_name in the IBMQREP_SENDQUEUES table is larger than the MAXMSGL queue attribute value of number bytes defined in WebSphere MQ.*

Explanation: The MAX_MESSAGE_SIZE column value of this send queue is not valid because it exceeds the MAXMSGL attribute of the WebSphere MQ queue manager. The send queue is invalidated.

User response: Decrease the MAX_MESSAGE_SIZE column value for this send queue. If necessary, increase the MAXMSGL attribute:

- For remote send queues, increase the MAXMSGL attribute of the transmission queue.
- For local send queues, increase the MAXMSGL attribute.

Restart any publication or Q subscription on the send queue.

If the send queue is remote, the MAX_MESSAGE_SIZE value must be at least 4 KB smaller than the MAXMSGL attribute of both the transmission queue and the queue manager. This 4 KB difference accounts for the extra information that is carried in the message header while the message is on the transmission queue.

ASN7029E *program_name : program_identifier : The restart message exceeded the maximum message size of the restart queue restart_queue.*

Explanation: See message text.

User response: Increase the maximum message size of the restart queue and warm start the Q Capture program.

ASN7030E *program_name : program_identifier : The restart queue queue_name is full. The Q Capture program will terminate.*

Explanation: The number of messages reached the number that is set for the MAXDEPTH attribute on the restart queue.

User response: If necessary, increase the value for the MAXDEPTH attribute of the restart queue. After you increase the value, cold start the Q Capture program.

ASN7033E *program_name : program_identifier : An invalid search condition was detected for publication or Q subscription name. Column name: column_name. Reason code: reason_code.*

Explanation: The following values are valid for the reason code:

0

An ordinary identifier for a column name must start with a letter followed by letters, numbers, or underscores.

1

The column name is not followed by a double quotation mark.

4

The search condition is missing the keyword WHERE.

5

The search condition contains a column that is not in the source table

6

The search condition does not contain any source table column names.

7

The combined search condition and change condition does not contain any source table column names or log record variables.

User response: Ensure that the search condition or change condition is in the proper format and restart the publication or Q subscription.

ASN7034W *program_name : program_identifier : The program cannot process a message from the administration queue queue_name. Message type: message_type. Message content: message_content.*

Explanation: See message text.

User response: Make sure that the format and the attributes of the control message are correct and put the message on the administration queue again.

ASN7035W *program_name : program_identifier : The program received a WebSphere MQ exception report message on the administration queue queue_name that contains an unsupported message. Message type: message_type. Message content: message_content.*

Explanation: See message text.

User response: Make sure that your WebSphere MQ configuration is correct.

ASN7036E *program_name : program_identifier : The program cannot process a message from the administration queue queue_name. Message type: message_type. Message content: message_content.*

Explanation: The program detected an internal inconsistency.

User response: Contact IBM Software Support. Provide the exact message that you received, which identifies the Q Capture schema. Provide the diagnostic log file. Also, provide a trace dump (ASNTRC DMP) file, if available.

ASN7037E *program_name : program_identifier : Initialization of the XML parser failed. The expected XML parser version is version_number. This error message was received: error_message.*

Explanation: An internal inconsistency was detected.

User response: Contact IBM Software Support. Provide the diagnostic log file. Also, provide a trace dump (ASNTRC DMP) file, if available.

ASN7038E *program_name : program_identifier : The XML control message does not satisfy the validity constraints of the XML schema for control messages as specified in the mqsub.xsd file. The message will be deleted from the queue and ignored. XML control message: xml_message.*

Explanation: See message text.

User response: Correct the XML control message so that it satisfies the validity constraints of the XML schema for control messages as specified in the mqsub.xsd file, and put it on the administration queue again.

ASN7039W *program_name : program_identifier : An error occurred while parsing an XML control message. The message was ignored and the program continues to run. Error message: error_message. XML control message: xml_message.*

Explanation: The XML control message is not valid and was ignored. The Q Capture program continues to run.

User response: Fix the XML control message and resend it.

ASN7040E *program_name : program_identifier : A DOM error occurred while parsing an XML control message. DOM exception code: exception_code. XML control message: xml_message.*

Explanation: The XML control message is ignored and the Q Capture program continues to run.

User response: Contact IBM Software Support. Provide the diagnostic log file. Also, provide a trace dump (ASNTRC DMP) file, if available.

ASN7043W *program_name : program_identifier : A message_type message was received in the administration queue for the nonexistent or inactive publication or Q subscription name. The control message was ignored.*

Explanation: The request that was received on the administration queue could not be performed because the publication or Q subscription does not exist or is inactive.

User response: Ensure that the name of the publication or Q subscription is correct and that it is in the active state. Put the control message on the administration queue again.

ASN7045E *program_name : program_identifier : The send queue queue_name of publishing or replication queue map queue_map_name is not active. The REINITQ command was ignored.*

Explanation: The REINITQ command must be issued for an active send queue.

User response: Check that the queue name is correct and reissue the REINITQ command or reinitialize the queue from the Replication Center.

ASN7046I *program_name : program_identifier : Send queue queue_name of publishing or replication queue map queue_map_name was successfully reinitialized. The following attributes were refreshed: ERROR_ACTION is error_action, HEARTBEAT_INTERVAL is heartbeat_interval, MAX_MESSAGE_SIZE is max_message_size.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7051W *program_name : program_identifier : The publication or Q subscription name was deactivated because no key column is specified in the IBMQREP_SRC_COLS table.*

Explanation: At least one column in the source and target tables must be specified as a key column for the publication or Q subscription. The name of this column is specified in the IBMQREP_SRC_COLS table.

User response: Ensure that at least one column in the source and target tables is specified as a key column for the publication or Q subscription, then activate the publication or Q subscription.

ASN7052E *program_name : program_identifier : The source table table_name for publication or Q subscription name contains LONG VARCHAR or LONGVARGRAPHIC columns but it does not have the DATA CAPTURE CHANGES attribute set to INCLUDE LONGVAR COLUMNS.*

Explanation: The source table is incorrectly defined. No data will be captured for this publication or Q subscription.

User response: Alter the source table so that the DATA CAPTURE CHANGES attribute is set to INCLUDE LONGVAR COLUMNS and restart the publication or Q subscription. For example:

```
alter table owner.table data capture changes
include longvar columns
```

ASN7057I *program_name : program_identifier : The SUB_ID of publication or Q subscription name in the control table IBMQREP_SUBS is not unique. A new SUB_ID is generated.*

Explanation: A duplicate SUB_ID was found in the control table IBMQREP_SUBS for publication or Q subscription name. The value in this column must be unique.

User response: This message is for your information only. No action is required.

ASN7059E *program_name : program_identifier : Q subscription name was not activated because the GROUP_MEMBERS column in the IBMQREP_SUBS table is not null.*

Explanation: The GROUP_MEMBERS column in the IBMQREP_SUBS table must be null for an inactive Q subscription in multidirectional replication.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7060E *program_name : program_identifier : The Q subscription name was not activated. The Q subscription corresponding to the Q subscription name for multidirectional replication either is in an invalid state or it does not exist in the same Q subscription group in the IBMQREP_TARGETS table.*

Explanation: To activate a Q subscription for

multidirectional replication, all the corresponding Q subscriptions in the same Q subscription group in the IBMQREP_TARGETS table must be in the active or inactive state. This error might occur if any Q subscriptions in the same Q subscription group are in the process of being initialized or deactivated.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7061E *program_name : program_identifier : Q subscription name was not activated because another Q subscription name that shares the same Q subscription group group_name and table table_name in the IBMQREP_SUBS table has a different SOURCE_NODE source_node.*

Explanation: In multidirectional replication, all Q subscriptions that share the same Q subscription group and table must have the same SOURCE_NODE in the IBMQREP_SUBS table.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7062E *program_name : program_identifier : Q subscription name was not activated because Q subscription name is in an invalid state state.*

Explanation: To activate a Q subscription for peer-to-peer replication, all the other Q subscriptions in the same Q subscription group must be in the active or inactive state in the IBMQREP_SUBS table. This error might occur if any Q subscriptions in the same Q subscription group are in the process of being initialized or deactivated.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7063E *program_name : program_identifier : Q subscription name was not activated because another Q subscription name, which shares the same Q subscription group, is in the process of being activated.*

Explanation: In peer-to-peer replication, multiple Q subscriptions in a Q subscription group cannot start at the same time.

User response: Check the IBMQREP_SUBS table to ensure that the other Q subscription is in the active ('A') state. Then activate the Q subscription that was placed in the inactive state.

ASN7065E *program_name : program_identifier : Q subscription name has a source or target node number that is outside of the supported range: low_limit to high_limit.*

Explanation: In multidirectional replication, the source node and target node in a Q subscription must be in the supported range.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7067E *program_name : program_identifier : Unable to process signal signal_name because of invalid SIGNAL_INPUT_IN signal_input_in, expected format: signal_input_format.*

Explanation: The program requires a parameter of the signal to be specified in SIGNAL_INPUT_IN. This error will be issued if some of the parameters are missing or in an unexpected format.

User response: Ensure that the SIGNAL_INPUT_IN is specified correctly and re-insert the signal.

ASN7073E *program_name : program_identifier : The REINIT_SUB command failed because the publication or Q subscription name is not active.*

Explanation: The REINIT_SUB command requires that a publication or Q subscription is active.

User response: Check that the publication or Q subscription was specified correctly in the command parameters and reissue the command.

ASN7074E *program_name : program_identifier : The REINIT_SUB command or the REINIT_SUB signal is not valid for Q subscription name. The Q subscription was not reinitialized.*

Explanation: The REINIT_SUB command or the REINIT_SUB signal applies only to publications or to Q subscriptions of type U (unidirectional) that are in the ACTIVE state. Unidirectional Q subscriptions or publications in a non-ACTIVE state, or Q subscriptions of type B (bidirectional) or P (peer-to-peer) cannot be reinitialized.

User response: For a publication, ensure that you specify the correct name and that the publication is in ACTIVE state. For a Q subscription, ensure that you specify the correct name, that the SUBTYPE in the IBMQREP_SUBS table is U, and that the Q subscription is in ACTIVE state. Reissue the command or reinsert the signal.

ASN7079W *program_name : program_identifier : Row operation messages cannot be sent in compact message format on send queue `queue_name` of queue map `queue_map_name`. Transaction messages are sent instead.*

Explanation: Both row operations and transaction messages are valid in XML message format, but only transaction messages are valid in compact format.

User response: If you want to publish row operations, the MESSAGE_FORMAT in the IBMQREP_SENDQUEUES control table must be set to 'X'.

ASN7080E *program_name : program_identifier : A WebSphere MQ exception report message was received in the administration queue with the WebSphere MQ reason code `reason_code`.*

Explanation: This error is issued when an MQ exception report message is generated by the receiving MCA (message channel agent) in response to a message delivery problem to the target WebSphere MQ queue manager.

User response: See the reason code documentation in the WebSphere MQ application programming reference for an explanation of the reason code and for information about actions that you might need to take for the target WebSphere MQ queue manager.

ASN7081W *program_name : program_identifier : The send queue for publication or Q subscription `name` that was specified in a WebSphere MQ exception report message no longer exists in the IBMQREP_SUBS control table.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7083E *program_name : program_identifier : The publication or Q subscription `name` was not activated because one or more columns of LOB data type were specified as key columns for the publication or Q subscription.*

Explanation: Columns of LOB data type must not be specified as key columns for a publication or Q subscription.

User response: Make sure that all columns in the publication or Q subscription that are specified as key columns are of a data type other than LOB. Activate your publication or Q subscription.

ASN7084E *program_name : program_identifier : Q subscription `name` was not activated because some columns of the source table are not included.*

Explanation: All source columns and version columns must be included in the Q subscription for peer-to-peer replication.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7085E *program_name : program_identifier : Q subscription `name` was not activated because the version column `name` is not defined in the logical table `table_name`.*

Explanation: Logical tables for Q subscriptions in peer-to-peer replication must contain version columns.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7086E *program_name : program_identifier : Q subscription `name` was not activated because version column `column_name` in logical table `table_name` was not defined with data type `datatype`, or the version column was not defined in the source table.*

Explanation: The data type of the version column was not defined correctly, or the version column was not defined in the logical table for the Q subscription in peer-to-peer replication.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7087E *program_name : program_identifier : Q subscription `name` (send queue `queue_name` of replication queue map `queue_map_name`) was not activated because its send queue does not use compact message format.*

Explanation: In multidirectional replication, the send queue of a Q subscription must use the compact message format.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7088E *program_name : program_identifier : For Q subscription `name`, the program cannot find information in the IBMQREP_SUBS table for the SUBGROUP, SOURCE_NODE, or*

TARGET_NODE columns.

Explanation: Q subscriptions in multidirectional replication must have a Q subscription group, a source node, and a target node.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7089E *program_name : program_identifier : The publication or Q subscription name was not activated because it contains LOB columns and no primary key, unique constraint, or unique index exists on source table table_name.*

Explanation: The source table for a publication or Q subscription that contains LOB columns must have one or more key columns. All columns that are specified as key columns must be of a data type other than LOB data type

User response: Use the replication administration tools to do take the following steps:

1. Ensure that the source table has a primary key, unique constraint, or unique index. Or, remove the LOB data type columns from the publication or Q subscription.
2. Activate the publication or Q subscription.

ASN7090E *program_name : program_identifier : Q subscription name was not activated because it contains a search condition.*

Explanation: A Q subscription in peer-to-peer replication cannot contain a search condition.

User response: If this Q subscription is for unidirectional or bidirectional replication, change the subscription type to 'U' or 'B' respectively. If this Q subscription is for peer-to-peer replication, remove the search condition. In any case, redefine the Q subscription using one of the replication administration tools and activate the Q subscription.

ASN7091W *program_name : program_identifier : The program cannot find the corresponding Q subscription in the IBMQREP_SUBS table for Q subscription name.*

Explanation: See message text.

User response: Redefine the Q subscription using one of the replication administration tools.

ASN7093E *program_name : program_identifier : Q subscription name_or_subid was not activated because it is not the only Q subscription in Q subscription group subgroup for bidirectional replication in the IBMQREP_SUBS table.*

Explanation: Only one Q subscription is allowed in a Q subscription group for bidirectional replication.

User response: Redefine the Q subscription using one of the replication administration tools, and then activate the Q subscription.

ASN7094E *program_name : program_identifier : The signal_name signal failed because the publication or Q subscription name has an invalid subtype subtype or an invalid state state.*

Explanation: The program received an unexpected or invalid signal in regards to a subtype or state of a particular publication or Q subscription.

User response: Make sure that the publication or Q subscription is expecting this signal. For example, a CAPSTART signal should only be targeted at an inactive Q subscription. Re-insert the signal if needed.

ASN7095E *program_name : program_identifier : The column SUB_ID in the IBMQREP_SUBS control table is null for the publication or Q subscription name that is either in active or loading state.*

Explanation: The SUB_ID in the table IBMQREP_SUBS cannot be null for a previously active publication or Q subscription.

User response: Use the ASNCLP command-line program or Replication Center to stop and then start the publication or Q subscription.

ASN7096E *program_name : program_identifier : The publication or Q subscription name contains a semicolon.*

Explanation: The publication or Q subscription name in the IBMQREP_SUBS control table may not contain a semicolon.

User response: Make sure that the publication or Q subscription name does not contain a semicolon and then start the publication or Q subscription.

ASN7097E *program_name : program_identifier : The Q subscription name did not start because there is a Q subscription for the same table in another Q subscription group for bidirectional or peer-to-peer replication.*

Explanation: See message text.

User response: Restart the Q subscription after ensuring that there are no other Q subscriptions for the same table in another Q subscription group.

ASN7098W *program_name : program_identifier : The ADDCOL signal failed because column `column_name` is already part of the publication or Q subscription name that was specified in the signal insert.*

Explanation: The source table column that was specified in the ADDCOL signal insert is already part of the Q subscription or publication definition.

User response: Ensure that the column is correctly specified and then insert the ADDCOL signal again if necessary.

ASN7100I *program_name : program_identifier : Column `column_name` has been added to the publication or Q subscription name.*

Explanation: Changes to the column are now being captured based on the Q subscription or publication definition. For Q subscriptions, the new column will be added to the target table if the column does not already exist in the target table.

User response: This message is for your information. No action is required.

ASN7101E *program_name : program_identifier : The column `column_name` that was specified in the ADDCOL signal for publication or Q subscription name cannot be added because is not nullable and does not have a default value.*

Explanation: Columns that are being added to Q subscriptions or publications with the ADDCOL signal must be nullable or defined as NOT NULL WITH DEFAULT.

User response: Ensure that the column that is being added is correctly defined, and then insert the ADDCOL signal again.

ASN7102E *program_name : program_identifier : The Q Capture program detected restart information that does not pertain to that program. The restart information came from the restart queue `queue_name` and it pertains to Q Capture server name and to Q Capture schema `schema`.*

Explanation: The Q Capture program is unable to warm start because the restart information in the restart queue is invalid.

User response: Ensure that you specified the restart queue correctly, or empty the restart queue and cold start the Q Capture program.

ASN7103W *program_name : program_identifier : The Q Capture program cannot find partitions that were previously known to DB2. Those database partitions are `partition_name`.*

Explanation: The Q Capture program is validating the database partition information that is stored in the restart information. The program found that some database partitions that it knew about from a prior invocation have been removed. The Q Capture program will update its database partition information and continue normally.

User response: Validate that the listed database partitions were actually removed. If so, no further action is required. If the database partitions have not been removed, stop the Q Capture program because it will not capture transactions on the missing database partitions. Start the Q Capture program with a start mode of cold.

ASN7104E *program_name : program_identifier : The publication name was not activated because it contains a topic topic that differs from active publications on send queue `queue_name` of publishing queue map `queue_map_name`.*

Explanation: All of the publications that use the same send queue must have the same topic. The publication was not activated because it did not have the same topic.

User response: Ensure that all of the publications that use the same send queue have the same topic.

ASN7105I *program_name : program_identifier : Restart information exists for send queue `queue_name`, However, this send queue of replication queue map `queue_map_name` does not exist in the IBMQREP_SENDQUEUES table.*

Explanation: The send queue was deleted from the IBMQREP_SENDQUEUES since the last run.

User response: This message is for your information. No action required.

ASN7106I *program_name : program_identifier : During warmstart, the Q subscription name is in a state other than inactive, but the restart information for the corresponding send queue `queue_name` of queue map `queue_map_name` is not found.*

Explanation: The restart information for the send queue is not present in the restart message. The restart message might be corrupted. The Q Capture program will fix its restart information.

User response: message is for your information. No action required.

ASN7107E *program_name : program_identifier :* **The number of ADDCOL signals that were inserted within this WebSphere MQ commit interval has exceeded the maximum of 20.**

Explanation: The Q Capture program can process a maximum of 20 ADDCOL signals within one WebSphere MQ commit interval. This interval specifies how often a Q Capture program commits messages to WebSphere MQ and is set by using the Q Capture program's `commit_interval` parameter.

User response: Insert the ADDCOL signals that failed so that the Q Capture program can commit them within a new interval.

ASN7108I *program_name : program_identifier :* **At program initialization, the lowest log sequence number of a transaction still to be committed (LSN) is `restart_lsn` and the highest log sequence number of a successfully processed transaction (MAXCMTSEQ) is `maxcmtseq`.**

Explanation: The Q Capture program starts and records the values of the restart message. Q Capture will start reading the log from the restart log sequence number (the lowest LSN of a transaction still to be committed) and will ignore any transactions whose commit is before the maximum commit sequence number (the highest LSN of a successful transaction). In a partitioned database, these log indicators come from the partition where the Q Capture program is running.

You can use the LSN and MAXCMTSEQ values to start Q Capture from a known point in the recovery log.

User response: This message is for your information only. No action is required.

ASN7109I *program_name : program_identifier :* **At program termination, the lowest log sequence number of a transaction still to be committed (LSN) is `restart_lsn` and the highest log sequence number of a successfully processed transaction (MAXCMTSEQ) is `maxcmtseq`.**

Explanation: The Q Capture program terminates and records the values of the restart message. During the next run, Q Capture will start reading the log from the restart log sequence number (the lowest LSN of a transaction still to be committed) and will ignore any transactions whose commit is before the maximum commit sequence number (the highest LSN of a successful transaction). In a partitioned database, these log indicators come from the partition where the Q

Capture program is running.

You can use the LSN and MAXCMTSEQ values to start Q Capture from a known point in the recovery log.

User response: This message is for your information only. No action is required.

ASN7110E *program_name : program_identifier :* **The LSN parameter must be specified with the MAXCMTSEQ parameter.**

Explanation: If you want to start the Q Capture program from a specific point in the recovery log, you must specify both the LSN parameter and the MAXCMTSEQ parameter. The LSN parameter specifies the lowest LSN of a transaction still to be committed. The MAXCMTSEQ parameter specifies the highest log sequence number of a successful transaction.

User response: Restart the Q Capture program and specify both the LSN parameter and the MAXCMTSEQ parameter.

ASN7111E *program_name : program_identifier :* **The LSN parameter and the MAXCMTSEQ parameter cannot be specified in cold start mode.**

Explanation: The Q Capture program was stopped because the LSN parameter and the MAXCMTSEQ parameter must be specified in warm start mode.

User response: Restart the Q Capture program in warm start mode with the LSN parameter and the MAXCMTSEQ parameter.

ASN7112I **Server name**

ASN7113I **Schema name**

ASN7114I **Program status**

ASN7115I **Time since program started**

ASN7116I **Log file location**

ASN7117I **Number of active Q subscriptions**

ASN7118I **Current memory**

ASN7119I **Log reader currency**

| | |
|----------|---|
| ASN7120I | Last committed transaction published (LSN) |
| ASN7121I | Current application memory |
| ASN7122I | Q Capture program status |
| ASN7123I | Up |
| ASN7124I | Down |
| ASN7125I | Path to DB2 log files |
| ASN7126I | Oldest DB2 log file needed for Q Capture restart |
| ASN7127I | Current DB2 log file captured |
| ASN7128I | <i>program_name : program_identifier : The db2flsn utility could not find the log file that contains LSN lsn. The return code is return_code. The reason is reason. The Q Capture program continues with normal processing.</i> |

Explanation: Q Capture executes the db2flsn (find log sequence number) command when Q Capture starts to display the name of the log file that includes the oldest DB2 log sequence number that is needed for a restart. This number is shown in the ASN7155I message. Q Capture executes db2flsn when Q Capture stops to display the name of the log file that includes the last DB2 log sequence number that was captured. This number is shown in the ASN7156I message. Q Capture also executes db2flsn when it receives the status show details command to display the oldest DB2 log file that is needed for a restart and the current DB2 log file that was captured. If the user ID that runs the Q Capture program does not have read privilege for the SQLOGCTL.LFH.1 and SQLOGCTL.LFH.2 log control files in the database home directory for the Q Capture server, the db2flsn program generates return code -101.

User response: If the utility return code is -101, give the user ID that runs the Q Capture program read privilege on the SQLOGCTL.LFH log control files.

| | |
|----------|--|
| ASN7129I | <i>program_name : program_identifier : The program could not open a pipe to execute the db2flsn utility. The return code is return_code. The reason is reason.</i> |
|----------|--|

Explanation: The Q Capture program was unable to open a pipe to execute the db2flsn utility.

User response: This message is for your information. No action is required.

ASN7130E *program_name : program_identifier : The Q Capture program is unable to decode log records for the source table table_owner.table_name. The program stopped.*

Explanation: The Q Capture program encountered a different version of the table because the table was altered. The Q Capture program was not able to determine the correct version by reading from the IBMQREP_COLVERSION and IBMQREP_TABVERSION tables. The tables might have been dropped.

User response: Ensure that the IBMQREP_COLVERSION and IBMQREP_TABVERSION tables have not been dropped. If they are missing, restore these tables and restart the Q Capture program.

ASN7131I *program_name :program_identifier : The Q Capture program will switch to cold start mode because warm start information was not found.*

Explanation: The Q Capture program switched to cold start because restart information is missing and WARMSI start mode was specified.

User response: This message is for your information. No action is required.

ASN7132E *program_name :program_identifier : Warm start information was not found and there are active Q subscriptions. The Q Capture program will terminate.*

Explanation: The Q Capture program stopped because information that is needed to restart is missing and Q subscriptions are active.

User response: Restart the Q Capture program in cold start mode if needed. You can also use the asncap command with the lsn and maxcmtseq parameters to specify a restart point.

ASN7133I *program_name :program_identifier : The publication or Q subscription name was stopped.*

Explanation: The publication or Q subscription was stopped because of a signal or an error.

User response: If the deactivation was due to an error, look in the Q Capture diagnostic log file or IBMQREP_CAPTRACE table for messages that might indicate the problem. Fix the error and start the publication or Q subscription.

ASN7134E *program_name : program_identifier :* **The publication or Q subscription name failed to activate during the initialization of the Q Capture program because of an error.**

Explanation: The publication or Q subscription was not started because of an error.

User response: Look in the Q Capture diagnostic log file or IBMQREP_CAPTRACE table for messages that might indicate the problem. Fix the error and start the publication or Q subscription.

ASN7135E *program_name : program_identifier :* **The publication or Q subscription name was not activated because one or more columns of LONG data type were specified as key columns for the publication or Q subscription.**

Explanation: Columns of LONG type are not valid for the key that is used for publishing or replication (IS_KEY > 0 in the IBMQREP_SRC_COLS or IBMQREP_TRG_COLS control tables).

User response: Use the Replication administration tools to redefine the columns that are used as a replication or publishing key, and then start the publication or Q subscription.

ASN7136E *program_name : program_identifier :* **The max_message_size value that is specified for remote send queue queue_name for publishing or replication queue map queue_map_name in the IBMQREP_SENDQUEUES table must be at least 4000 bytes less than the MAXMSGL attribute value of number bytes that is defined for the transmission queue queue_name in WebSphere MQ.**

Explanation: The max_message_size value of the queue map that uses this send queue is too large. If the send queue is remote, the max_message_size value must be at least 4 KB smaller than the MAXMSGL attribute of both the transmission queue and the queue manager. This 4 KB difference, which is specified by MQ_MSG_HEADER_LENGTH, accounts for the extra information that is carried in the message header while the message is on the transmission queue.

User response: Use the ASNCLP or Replication Center to reduce the MAX_MESSAGE_SIZE value for the queue map. If necessary, increase the MAXMSGL attribute of the transmission queue, send queue, remote queue, the queue channels, and the client-server connection channels (WebSphere MQ client). Restart any publications or Q subscriptions that use the send queue.

ASN7137W *program_name : program_identifier :* **The specified MAX_MESSAGE_SIZE for the publication or Q subscription name might be too small for the send queue queue_name, publishing queue map or replication queue map queue_map_name. A maximum of size bytes might be required.**

Explanation: The MAX_MESSAGE_SIZE value of the publishing queue map or replication queue map that specifies this send queue might too small to contain a row in the WebSphere MQ message. If this problem occurs when the Q Capture program is running, Q Capture stops.

User response: Increase the MAX_MESSAGE_SIZE value for the queue map. Reinitialize the queue map by using the reinitq subcommand of the MODIFY or asnqcmd command.

ASN7138W *program_name : program_identifier :* **The Q Capture program is converting the source columns from code page code_page to code page code_page for delimited publication publication_name (send queue send_queue_name).**

Explanation: The source database columns are encoded using a code page that is different from the code page that is specified in the MESSAGE_CODEPAGE column of the IBMQREP_SENDQUEUES table. The Q Capture program will convert all data in the messages that it puts on the send queue.

User response: No action is necessary if data conversion is acceptable. If you wish to avoid data conversion, stop the Q Capture program and update the MESSAGE_CODEPAGE value to match the source column code page. Then start Q Capture.

ASN7139W *program_name : program_identifier :* **The value of HEARTBEAT_INTERVAL in the IBMQREP_SENDQUEUES table for send queue send_queue_name (publishing queue map queue_map_name) must be 0 for delimited message format.**

Explanation: If you specify that a publishing queue map is used for delimited messages, the value for the interval that the Q Capture program uses to send heartbeat messages must be 0 (no heartbeat messages are sent).

User response: No action is necessary. The Q Capture program overrides the HEARTBEAT_INTERVAL value. If you do not wish to see this message, use the ASNCLP command-line program or Replication Center to change the heartbeat interval for the queue map to 0. Then use the reinitq command or replication administration tools to reinitialize the send queue.

ASN7140E *program_name : program_identifier : The delimited publication publication_name (send queue send_queue_name) did not start because the BEFORE_VALUES attribute of the publication is not Y.*

Explanation: When publishing delimited messages, the Q Capture program always sends before and after values for an update operation. If you specify that a publication uses delimited message format, you must enable before values.

User response: Use the ASNCLP or Replication Center to specify before values for the publication and then start the publication.

ASN7141E *program_name : program_identifier : The delimited publication publication_name (send queue send_queue_name) did not start because the CHANGED_COLS_ONLY attribute of the publication is not N.*

Explanation: When publishing delimited messages, the Q Capture program always sends all columns for a row even if they are not changed. If you specify that a publication will use delimited message format, the CHANGED_COLS_ONLY attribute must be N.

User response: Use the ASNCLP or Replication Center to specify CHANGED_COLS_ONLY=N for the publication.

ASN7142I *program_name : program_identifier : The Q Capture program received a WebSphere MQ error with reason code reason_code when it put a message on queue queue_name using the MQPUT command. The program will retry putting the message up to number times at number second intervals before stopping.*

Explanation: The Q Capture program could not put a message on the queue and received one of the following WebSphere MQ reason codes from the MQPUT operation:

- 2051: MQRC_PUT_INHIBITED
- 2053: MQRC_Q_FULL
- 2192: MQRC_STORAGE_MEDIUM_FULL
- 2346: MQRC_CF_STRUC_IN_USE
- 2373: MQRC_CF_STRUC_FAILED

When this error is received, the Q Capture program retries sending the message based on the WebSphere MQ qfull_num_retries and qfull_retry_delay parameters. This error might occur if the Q Apply program is not processing the messages on the receive queue. The Q Apply program might be stopped or the Q Capture program might be putting messages on the queue faster than WebSphere MQ can process them.

The error might also occur if many changes accumulated at the source table while the Q Capture program was stopped. This error is more likely to happen in a queue-sharing environment.

User response: This message is for your information only. No action is required.

ASN7143I *program_name : program_identifier : The WebSphere MQ MQPUT command on queue queue_name was successful after number retries.*

Explanation: The Q Capture program needed to retry the MQPUT operation multiple times because the queue was full.

User response: This message is for your information only. No action is required.

ASN7145E *program_name : program_identifier : Oracle LogMiner found a dictionary mismatch while reading from V\$LOGMNR_CONTENTS for table table_name. LogMiner status is status. Info is informational-message. All Q subscriptions for the table were stopped.*

Explanation: Oracle LogMiner found a dictionary mismatch and could not translate the log record. Possible causes include:

- Database-wide supplemental logging is not enabled.
- Table-level supplemental logging is not enabled for the subscribed table.
- A subscribed table was altered and Q Capture was restarted at a system change number (SCN) location that is earlier than the DDL change.

All of these situations can cause a mismatch between the online dictionary and the format of the log records.

User response: Enable database-wide supplemental logging on the Oracle source database. Additionally, enable table-level supplemental logging for all subscribed tables. Start all Q subscriptions for the table.

ASN7146E *program_name : program_identifier : Oracle LogMiner found a missing system change number (SCN) range in the list of redo log files that were being processed. The Oracle error message is 'oracle_message'. The Q Capture program cannot proceed because of missing log records. The program will stop.*

Explanation: Log records that are needed by the Q Capture program are missing. The program stops so that you can locate missing redo log files or archive log files.

User response: Use the Oracle error message to

identify which redo or archived logs contain these SCN values and make these files available to Oracle, then start Q Capture. If these missing log files are no longer available, a cold start is required.

ASN7147E *program_name : program_identifier : An error occurred while attempting to start an Oracle LogMiner session. The Oracle error is error.*

Explanation: The Q Capture program uses Oracle LogMiner to read log records for the Oracle source table. LogMiner returned an error while trying to start.

User response: See the Oracle messages and codes documentation for the cause and action related to this error and perform the recommended action.

ASN7148E *program_name : program_identifier : The program encountered an Oracle Call Interface (OCI) error while performing the operation operation. The Oracle error is error.*

Explanation: An error occurred when the Q Capture, Q Apply, or Replication Alert Monitor program issued an OCI call. The error might be caused by an Oracle problem that needs to be investigated, such as an out-of-space condition, or Oracle might be unavailable for use by applications. This message is sometimes followed by a second message that provides information about what the replication program was doing when the error occurred.

User response: See the Oracle messages and codes documentation for an explanation of this error code and for information about corrective actions that you might need to take in Oracle. If a replication program issued another message immediately after this one, see the explanation and user response for that message.

ASN7149E *program_name : program_identifier : The program encountered an Oracle C++ Call Interface (OCCI) error while performing an SQL operation. The operation is operation. The server name is server_name. The SQL request is sql_request. The table name is table_name. The Oracle error code is error_code. The error message is error_message.*

Explanation: An error occurred when the Q Capture, Q Apply, or Replication Alert Monitor program issued an OCCI call. The error might be caused by an Oracle problem that needs to be investigated, such as an out-of-space condition, or Oracle might be unavailable for use by applications. This message is sometimes followed by a second message that provides information about what the replication program was doing when the error occurred.

User response: See the Oracle messages and codes

documentation for an explanation of this error code and for information about corrective actions that you might need to take in Oracle. If a replication program issued another message immediately after this one, see the explanation and user response for that message.

ASN7150E *program_name : program_identifier : The program encountered an Oracle C++ Call Interface (OCCI) error while performing the operation operation. The Oracle error code is error_code. The error message is error_message.*

Explanation: An error occurred when the Q Capture, Q Apply, or Replication Alert Monitor program issued an OCCI call. The error might be caused by an Oracle problem that needs to be investigated, such as an out-of-space condition, or Oracle might be unavailable for use by applications. This message is sometimes followed by a second message that provides information about what the replication program was doing when the error occurred.

User response: See the Oracle messages and codes documentation for an explanation of this error code and for information about corrective actions that you might need to take in Oracle. If a replication program issued another message immediately after this one, see the explanation and user response for that message.

ASN7151E *program_name : program_identifier : The program encountered an unsupported data type during internal processing. The column name is column_name. The table name is table_name.*

Explanation: An internal error occurred during processing of an Oracle log record.

User response: Redefine the Q subscription for the table if this data type is not supported. Otherwise, contact IBM Software Support.

ASN7153W *program_name : program_identifier : A TRUNCATE TABLE log record was detected and ignored for table table_name.*

Explanation: TRUNCATE TABLE statements are not replicated by the Q Capture program. Any rows that were deleted from the source table by the statement were not deleted from the target table.

User response: This message is for your information only, and no action is required.

ASN7154E *program_name : program_identifier : The source table **table_name** for publication or Q subscription **name** does not have table-level supplemental logging enabled for all columns. No data will be captured for this publication or Q subscription.*

Explanation: The source table is incorrectly defined. The SUPPLEMENTAL LOG DATA attribute must be defined for all columns.

User response: Alter the source table so that SUPPLEMENTAL LOG DATA (ALL) COLUMNS is enabled and restart the publication or Q subscription. For example:

```
alter table <src_owner>.<src_table>
  add supplemental log data (all) columns;
```

ASN7155I *program_name : program_identifier : The Q Capture program started processing log records in log file **log_file_name**.*

Explanation: When the Q Capture program started, this log file contained the first DB2 log sequence number (LSN) or Oracle system change number (SCN) that the program read.

User response: This message is for your information only. No action is required.

ASN7156I *program_name : program_identifier : The Q Capture program stopped processing log records in log file **log_file_name**.*

Explanation: When the Q Capture program stopped, this log file contained the last DB2 log sequence number (LSN) or Oracle system change number (SCN) that the program read.

User response: This message is for your information only. No action is required.

ASN7157E *program_name : program_identifier : The value **value** in the COMPATIBILITY column of the IBMQREP_CAPPARMS control table is not allowed for Q Capture for Oracle sources. Only values of '0905' and above are allowed. The Q Capture program stopped.*

Explanation: At startup, the Q Capture program found an unsupported value in the COMPATIBILITY column. Q Capture for Oracle sources can only work with Q Apply programs at Version 9.5 or later. The COMPATIBILITY column indicates the version of messages that Q Capture sends to Q Apply.

User response: Update the COMPATIBILITY column in IBMQREP_CAPPARMS to '0905' or above and restart the Q Capture program.

ASN7158E *program_name : program_identifier : The value **value** in the LOB_SEND_OPTION column of the IBMQREP_CAPPARMS control table is not supported for Q Capture for Oracle sources. The only supported value is 'I'. Large-object (LOB) data is sent inline.*

Explanation: At startup, the Q Capture program found an unsupported value in the LOB_SEND_OPTION column. Q Capture for Oracle sources only supports the inline send option for LOBs.

User response: Update the LOB_SEND_OPTION column in IBMQREP_CAPPARMS to "I".

ASN7160W *program_name : program_identifier : The average time that it took the Q Capture program to receive a batch of log records from the Oracle LogMiner utility during the most recent monitor interval exceeded the Q Capture commit interval. The minimum, maximum, and average fetch durations (in milliseconds) were **minimum_fetch_time**, **maximum_fetch_time**, **average_fetch_time**.*

Explanation: During the last period when Q Capture wrote performance statistics to its control tables, the average amount of time that Q Capture took to receive a batch of log records from LogMiner was greater than the value of the Q Capture **commit_interval** parameter. This parameter sets the length of time that Q Capture waits to commit messages that are on the send queue to WebSphere MQ. An average fetch time that is larger than the commit interval slows Q Capture performance.

User response: Update the **commit_interval** parameter to a value in milliseconds that is greater than or equal to the average fetch duration that is shown in this error message.

ASN7166W *program_name : program_identifier : The data for LOB or XML columns has not been sent because the transaction size **size_bytes** exceeds the maximum message size **size_bytes** of the send queue. The subscription or publication is **name**. The LOB or XML column **column_name** is in the source table **table_name**. The row has key columns **key_columns** with key values **key_values**. The error action of the send queue will be performed.*

Explanation: The LOB or XML data was too large for the transaction message.

User response: Increase the MAX_MESSAGE_SIZE value in the IBMQREP_SENDQUEUES table. You might also need to increase the size of the MAXMSGL parameter for the WebSphere MQ send queue. Ensure that the value for the MAX_MESSAGE_SIZE parameter

is at least 4KB smaller than value for the MAXMSGL parameter.

If you cannot increase the MAX_MESSAGE_SIZE parameter, you can set the LOB_SEND_OPTIONS value in the IBMQREP_CAPPARMS table to 'S' to send the LOB values in a LOB message following the transaction message.

ASN7167E *program_name : program_identifier : The Q Capture program could not activate a Q subscription or publication for table table_owner.table_name. The table was altered. The table version is version.*

Explanation: The Q Capture program cannot decode log records if a table is altered. After you create a Q subscription or publication that specifies a source table, you must not alter the table before activating the Q subscription or publication. The table version must be 0.

User response: Reorganize the table and activate the Q subscription or publication.

ASN7168E *program_name : program_identifier : The characters that were chosen for a column delimiter, record delimiter, string delimiter, and decimal point are not unique for publishing queue map queue_map_name that specifies send queue queue_name.*

Explanation: None of the four types of delimiters that are used for delimited messages in event publishing can use the same character.

User response: Use the Q Replication Dashboard or Replication Center to change the publishing queue map so that all of the delimiters have unique values. Then restart the Q Capture program.

ASN7169E *program_name : program_identifier : The characters that were chosen for a column delimiter, record delimiter, string delimiter, and decimal point are invalid for the publishing queue map queue_map_name that specifies send queue queue_name.*

Explanation: The delimiters that are used for delimited messages in event publishing cannot be alphanumeric characters (0-9, aA-zZ).

User response: Use the Q Replication Dashboard or Replication Center to change the publishing queue map so that none of the delimiters use alphanumeric characters. Then restart the Q Capture program.

ASN7170E *program_name : program_identifier : The value of the max_message_size parameter for replication or publishing queue map queue_map_name that specifies send queue queue_name was exceeded.*

Explanation: A WebSphere MQ message was put on the send queue that is larger than the maximum message size limit that was defined for the queue map that uses this queue.

User response: Increase the value of max_message_size for the queue map and reinitialize the Q Capture program.

ASN7171E *program_name : program_identifier : The program stopped because the data for a LOB or XML column was too large and the error action of the replication or publishing queue map queue_map_name that includes send queue queue_name was to stop the program.*

Explanation: Data in LOB or XML columns is sent within the transaction message. If the LOB or XML data is large, you might need to increase the maximum message size for WebSphere MQ messages. Two parameters govern the allowable size of messages: max_message_size for queue maps, and MAXMSGL (maximum message length) for the WebSphere MQ queue.

User response: Increase the value of max_message_size for the replication or publishing queue map. You might also need to increase the size of MAXMSGL for the WebSphere MQ send queue that is specified in the queue map. Ensure that the value for max_message_size is at least 4 KB smaller than the value for MAXMSGL.

ASN7172E *program_name : program_identifier : The publication publication_name (send queue queue_name, publishing queue map queue_map_name) that is set up for delimited message format (message_format=D) cannot be activated because the Q Capture lob_send_options parameter is set to S (send LOB values in a separate message).*

Explanation: The delimited message format does not support sending LOB values in a separate message. The delimited message format only supports sending a LOB value within the transaction message (lob_send_options=I).

User response: Use the replication administration tools to set lob_send_options=I. Then reinitialize the Q Capture program and start the publication.

ASN7173W *program_name : program_identifier : The send queue queue_name was deactivated because of an MQPUT error error_number. The Q Capture program continues putting messages on other send queues based on the Q (stop queue) error action in the IBMQREP_SENDQUEUES table. You can use the startq command to resume replication or publishing on the queue after you correct the problem.*

Explanation: The Q error action prompts Q Capture to continue putting messages on active send queues even when one send queue is deactivated. Q Capture saves restart information for Q subscriptions or publications that use the deactivated send queue, and can recapture changes for this queue from the log and bring it up to the same restart point as the other queues.

User response: Use the WebSphere MQ error number to determine the problem that caused the error on the queue, and then use the startq command to resume replication or publishing on the queue.

ASN7174I *program_name : program_identifier : The send queue queue_name for replication queue map queue_map_name was activated because of a startq command.*

Explanation: Q Capture recaptures changes for this queue from the log and brings it up to the same restart point as the other queues.

User response: This message is for your information only. No action is required.

ASN7175E *program_name : program_identifier : The send queue queue_name does not exist. The startq command is ignored.*

Explanation: The startq command was issued for a send queue that does not exist.

User response: Check that the queue name is correct and reissue the startq command.

ASN7176I *program_name : program_identifier : The send queue queue_name for replication queue map queue_map_name was deactivated because of a stopq command. Q Capture has captured up to log sequence number lsn for this queue. The maximum commit sequence for this queue is maxcmtseq.*

Explanation: Q Capture continues to put messages on active send queues even when one send queue is deactivated. Q Capture saves restart information for Q subscriptions or publications that use the deactivated send queue, and can recapture changes for this queue

from the log and bring it up to the same restart point as the other queues.

User response: Use the startq command to resume replication or publishing on the queue.

ASN7177I *program_name : program_identifier : The send queue queue_name does not exist or is not known to the Q Capture program. The stopq command is ignored.*

Explanation: The stopq command was issued for a send queue that does not exist or is not known to the Q Capture program.

User response: Check that the queue name is correct and reissue the stopq command. If the queue was added while replication is active, start the Q subscriptions that use the queue so that the Q Capture program becomes aware of the queue.

ASN7178I *program_name : program_identifier : All send queues are in an inactive (I) state. You can use the STARTQ command to resume replication or publishing on inactive queues.*

Explanation: While all queues are inactive, the Q Capture program continues reading the log for signals such as CAPSTART, continues inserting into its monitor tables, and waits for commands.

User response: This message is for your information only. No action is required.

ASN7179W *program_name : program_identifier : The Q Capture program is unable to activate the Q subscription name for source table table-owner_table.name because the send queue queue_name that is used by the Q subscription is in an inactive (I) state and other active Q subscriptions use the send queue. The CAPSTART signal to activate the Q subscription is ignored by Q Capture.*

Explanation: When a send queue is inactive either because it was stopped by the stopq command or because of an error, the Q Capture program cannot activate Q subscriptions that specify the queue unless no other active Q subscriptions use the queue, or the queue is started.

User response: Take one of the following actions:

- Use the **startq** command to resume replication or publishing on the queue.
- Stop all other active Q subscriptions that use the queue and then start the Q subscription that is specified in this message.
- Stopping and starting the other active Q subscriptions prompts a new load (full refresh) of the target tables for all of the Q subscriptions. If you

want to prompt a new load for just the table that is specified in the message text, use the **startq** command.

ASN7180E *program_name : program_identifier :* **The Q Capture program log read failed because the DB2 compression dictionary that was used to create the compressed log record no longer exists. The log record that could not be read was for the source table *table_owner.table_name*. The reason code is *reason_code*.**

Explanation: The Q Capture program received an error from the DB2 log read interface. The reason code is a DB2 for z/OS reason code or on Linux, UNIX, or Windows a SQLCODE. The error indicates that the data on a log record cannot be processed because the compression dictionary for the corresponding DB2 table or table space is permanently gone. The compressed table or table space that contains this source table was probably reorganized by the REORG utility that ran without the KEEPDICTIONARY option. The Q Capture program cannot read the remaining compressed log records from source changes that occurred before the reorganization, and so it deactivated the Q subscription. Q Capture then activates the Q subscription and prompts a new load of the target table if a load is specified.

User response: If the Q subscription specifies an automatic load of the target table, no action is required. If the Q subscription specifies a manual load, you must take actions to reload the target and issue a LOADDONE signal to Q Capture. Q Apply will spill changes to the spill queue until it sees the LOADDONE message from Q Capture. If the Q subscription specifies no load, you must take actions to reload the target and reactivate the Q subscription.

ASN7181E *program_name : program_identifier :* **The Q Capture program log read failed because the DB2 compression dictionary that was used to create the compressed log record is temporarily unavailable. The log record that could not be read was for the source table *table_owner.table_name*. The reason code is *reason_code*.**

Explanation: The Q Capture program received an error from the DB2 log read. The error indicates that the data on a log record cannot be processed because the compression dictionary for the corresponding DB2 table or table space is temporarily unavailable. For z/OS, the reason code is a z/OS diagnostic code. For Linux, UNIX, and Windows, the reason code is an SQL code. One cause for the problem could be that the compressed table space is in the STOPPED state when the DB2 Log Read Interface attempts to read the compression dictionary. DB2 takes a latch on the source

compressed table space to access the dictionary and the latch does not work if the table space is stopped. This error prompts the Q Capture program to deactivate the Q subscription for the source table.

User response: The problem will be resolved when the compression dictionary becomes available. Restart the Q Capture program.

ASN7182W *program_name : program_identifier :* **The signal *signal_type* failed because the send queue *queue_name* that is used by the Q subscription *name* is in an inactive (I) state. The Q Capture program ignored the signal.**

Explanation: When a send queue is inactive either because it was stopped by a **stopq** command or because of an error, the Q Capture program cannot process signals for Q subscriptions that use the queue until the queue is started.

User response: Use the startq command to resume replication or publishing on the queue and then reinsert the signal or repeat the action that resulted in the signal insert.

ASN7183I *program_name : program_identifier :* **Data partition *data_partition_number* was added to the source table *table_owner.table_name*.**

Explanation: The source partitioned table was altered to add a partition. The Q Capture or Capture program does not replicate this alteration, and so the new partition is not automatically added to the target table. Subsequent inserts, updates, and deletes to this newly added partition are captured. If the target is an identically partitioned table and the new source partition does not exist at the target, errors occur when the Q Apply or Apply program tries to apply rows to the target.

User response: If you require the new partition at the target you must manually add it.

ASN7184W *program_name : program_identifier :* **Data partition *data_partition_number* was attached to the source table *table_owner.table_name*.**

Explanation: The source partitioned table was altered to attach a partition. The Q Capture or Capture program does not support replication of this alteration, and so the new partition is not automatically attached at the target and the existing data in the new partition is not replicated to the target. Subsequent inserts, updates, and deletes to this newly attached partition are replicated.

User response: If you require the new partition at the target you must manually add it. If you require the

attached data at the target, you must manually load it into the target.

ASN7185W *program_name : program_identifier : Data partition data_partition_number was detached from the source table table_owner.table_name.*

Explanation: The source partitioned table was altered to detach a partition. The Q Capture or Capture program does not support replication of this alteration, and so the detached partition is not detached at the target and the DELETE operations are not replicated to the target.

User response: If you no longer want the partition or its data at the target, detach the partition or delete the data.

ASN7186W *program_name : program_identifier : An operation by utility utility_name was detected on table table_name for Q subscription name (receive queue queue_name, replication queue map queue_map_name). The handling of rows that are added or changed depends on the type of DB2 utility that performed the data changes and whether the Q subscription specifies that load operations at the source table should be replicated (value of CAPTURE_LOAD column in the IBMQREP_SUBS control table).*

Explanation: The Q Capture program detected data changes at the source table by one of the following utilities:

DB2 for z/OS

- LOAD SHRLEVEL NONE RESUME YES
- LOAD SHRLEVEL NONE REPLACE
- REORG TABLESPACE DISCARD
- CHECK DATA DELETE YES LOG NO
- RECOVER PIT
- EXCHANGE DATA ON CLONE TABLESPACE

DB2 for Linux, UNIX, and Windows

DB2 LOAD

For the following DB2 utilities, Q Capture acts based on the CAPTURE_LOAD value:

- LOAD SHRLEVEL NONE RESUME YES
- LOAD SHRLEVEL NONE REPLACE
- REORG TABLESPACE DISCARD
- CHECK DATA DELETE YES LOG NO
- RECOVER PIT
- EXCHANGE DATA ON CLONE TABLESPACE

- DB2 LOAD (Linux, UNIX, Windows)

The CAPTURE_LOAD option supports the following actions when a data change operation is detected:

R (restart)

The Q subscription is restarted and the target table is loaded with the data from the source table. The type of load is determined by the LOAD_TYPE value in the IBMQREP_TARGETS control table.

W (warning)

The Q subscription is not restarted, and this warning message is issued.

User response: If CAPTURE_LOAD is set to W or the DB2 utility is not supported by CAPTURE_LOAD, take action to ensure that the source and target tables are synchronized.

ASN7187E *program_name : program_identifier : The value of the CAPTURE_LOAD column in the IBMQREP_SUBS control table is set to R for Q subscription name. This value is only supported in unidirectional, bidirectional, and peer-to-peer Q subscriptions with a maximum of two servers. However, a load operation was detected on the source table. The Q Capture program stopped.*

Explanation: CAPTURE_LOAD is set to R but the Q subscription type is peer-to-peer with more than two servers. The R value is only supported for unidirectional, bidirectional, and peer-to-peer replication with two servers.

User response: Change the value of CAPTURE_LOAD to W for all Q subscriptions that participate in the peer-to-peer group. Because the source table was loaded, the source and target tables are not synchronized. Take the following actions:

1. Manually synchronize the tables by using the asntdiff and asntrep utilities.
2. Set CAPTURE_LOAD to W.
3. Start Q Capture in warm mode. Q Capture issues a warning message when it detects the source table load operation in the log.

ASN7188W *program_name : program_identifier:* The value of the CAPTURE_LOAD option in the IBMQREP_SUBS control table is set to R. The value of the HAS_LOADPHASE option in the IBMQREP_SUBS control table is set to N. If the Q Capture program detects a load operation, the Q subscription *name* (send queue *queue_name*, replication queue map *queue_map_name*) will be restarted by Q Capture without a load phase.

Explanation: If the Q Capture program detects a load operation and the CAPTURE_LOAD value is set to R, Q Capture restarts the Q subscription. During the restart process the Q Apply program uses the HAS_LOADPHASE value to determine the type of load to use. If the HAS_LOADPHASE value is set to N the Q Apply program does not load the source table. The source and target tables become out of synch because the target table is not loaded with the latest data from the source table. If the source and target tables become out of synch, you must synchronize the tables.

User response: In the IBMQREP SUBS control table, review the settings for the CAPTURE_LOAD option and the HAS_LOADPHASE option for the affected Q subscription.

ASN7189I *program_name : program_identifier :* The Q subscription *name* (send queue *queue_name*, replication queue map *queue_map_name*) for table *table_owner.table_name* was restarted because a *utility_name* operation was detected in the database recovery log.

Explanation: The Q Capture program detected a utility operation that changed data in the source table and the CAPTURE_LOAD option is set to R. The Q subscription was restarted. The restart prompts a new load of the target table, which includes, if applicable, any newly loaded data.

User response: This message is for your information only. No action is required.

ASN7190W *program_name : program_identifier :* The value of the CAPTURE_LOAD column in the IBMQREP_SUBS control table is set to R for Q subscription *name* (send queue *queue_name*, replication queue map *queue_map_name*). However, the compatibility level is *compatibility_level*. The Q Capture program will report but not replicate load operations that it detects at the source table.

Explanation: The CAPTURE_LOAD option of R is only valid with a compatibility level of 0907 or higher, which requires both the Q Capture and Q Apply

program to be at Version 9.7 or higher. When CAPTURE_LOAD is set to R for Version 9.7 and higher, load operations at the source table are replicated by using the load type that was chosen for the Q subscription.

User response: To enable the replicate load feature, upgrade the Q Apply program to Version 9.7 or higher and set the COMPATIBILITY value in the IBMQREP_CAPPARMS table to 0907 or higher.

ASN7191E *program_name : program_identifier :* The database *database_name* is not in ARCHIVELOG mode. The Q Capture program stopped.

Explanation: The Q Capture program tried to perform a cold or warm start and the source database was not configured correctly for the Q Capture program to work with the Oracle LogMiner utility. The LOG_MODE column in the V\$DATABASE dynamic view must be set to ARCHIVELOG.

User response: Change the configuration of the database to enable ARCHIVELOG mode. For more information on enabling archival logging, see the "Managing Archived Redo Logs" chapter in the Oracle Database Administrator's Guide.

ASN7192E *program_name : program_identifier :* The database *database_name* is not in minimal supplemental logging mode. The Q Capture program stopped.

Explanation: The Q Capture program tried to perform a cold or warm start and the source database was not configured correctly for the Q Capture program to work with the Oracle LogMiner utility. The SUPPLEMENTAL_LOG_DATA_MIN column in the V\$DATABASE dynamic view must be set to YES or IMPLICIT.

User response: Enable minimal supplemental logging for the database. For more information, see "Supplemental Logging" in the Oracle Database Utilities guide.

ASN7193I Timed out waiting for reply

ASN7194W *program_name : program_identifier :* The Q Capture program detected that a partition was attached to table *table_owner.table_name*. Q Capture was behind in the log when the first Q subscription for this table was started. Q Capture might have incorrectly captured rows for the table before it was attached as a partition.

Explanation: When the Q Capture program starts the first Q subscription for a table, it gathers table information from the system catalogs. If a partition was

attached to a table before Q Capture starts the first Q subscription for the table, the system catalog information might not correctly reflect where Q Capture is in the log. For a newly attached partition, Q Capture might incorrectly capture rows from the table before it was attached as a partition

User response: Stop and then start the Q subscription.

ASN7195W *program_name : program_identifier :* **The Q Capture program detected that a partition was detached from table *table_owner.table_name*. Q Capture was behind in the log when the first Q subscription for this table was started. For a newly detached partition, Q Capture might incorrectly not capture rows from the partition before it was detached.**

Explanation: When the Q Capture program starts the first Q subscription for a table, it gathers table information from the system catalogs. If a partition was detached from the table before Q Capture starts the first Q subscription for a table, the system catalog information might not correctly reflect where Q Capture is in the log. For a newly detached partition, Q Capture might incorrectly miss rows from the partition before it was detached.

User response: Stop and then start the Q subscription.

ASN7198E *program_name : program_identifier :* **The *trans_batch_sz* parameter is set to a value greater than 1, but the Q Capture *lob_send_options* parameter is set to S rather than the required setting of I. The Q Capture program stops.**

Explanation: Batching of transactions can only be used when large object (LOB) data is sent within a single transaction message. If the *trans_batch_sz* parameter is set to greater than 1, the *lob_send_options* parameter must be set to I (inline) rather than S (separate).

User response: Set the *lob_send_option* parameter to I and restart the Q Capture program.

ASN7199E *program_name : program_identifier :* **The log record variable *variable_name* that was specified in the *CHANGE_CONDITION* column of the *IBMQREP_SUBS* table for Q subscription or publication *name* is not valid for this database. The Q subscription or publication is deactivated.**

Explanation: The \$AUTHOKEN and \$PLANNAME variables cannot be used for change conditions on Linux, UNIX, and Windows. These variables are only available in the database log record on z/OS operating

systems. On Linux, UNIX, and Windows, you can use the \$OPERATION and \$AUTHID variables.

User response: Ensure that the search condition and change condition are in the proper format and restart the publication or Q subscription. See "Log record variables to filter rows (unidirectional replication)" in the DB2 Information Center for details about the correct way to specify a search condition or change condition.

ASN7200E *program_name : program_identifier :* **The variable *variable_name* is not valid in the *SEARCH_CONDITION* column of the *IBMQREP_SUBS* table for Q subscription or publication *name*. Log record variables should be specified in the *CHANGE_CONDITION* column only. The Q subscription or publication is deactivated.**

Explanation: The *SEARCH_CONDITION* column must only contain source table column names or constants. Specify log record variables in the *CHANGE_CONDITION* column.

User response: Ensure that the search condition and change condition are in the proper format and restart the publication or Q subscription. See "Log record variables to filter rows (unidirectional replication)" in the DB2 Information Center for details about the correct way to specify a search condition or change condition.

ASN7201E *program_name : program_identifier :* **The restart information in queue *queue_name* is incompatible with the version *version* of Q Capture.**

Explanation: The restart message cannot be processed by this version of Q Capture, possibly because Q Capture was recently upgraded.

User response: Start Q Capture using the *asncap* command, prompting the program to read from the beginning of the log by specifying the parameter "migrate=y". Use this option only the first time that you start Q Capture and specify *startmode=warmns*.

ASN7202E *program_name : program_identifier :* **The Q Capture program could not open the restart file *file_name* for reason *reason*. The program did not start.**

Explanation: The Q Capture program was started with the *override_restartq* parameter set to Y (yes). This setting prompts Q Capture to look in the restart file instead of the restart queue for restart information. However, Q Capture was unable to open the restart file. The file might have been inadvertently deleted, or the user ID that is associated with the Q Capture

process does not have sufficient authority to open the file.

User response: Verify that the user ID that is associated with the Q Capture process has sufficient authority. If the file was inadvertently deleted, create a new restart file and restart Q Capture. Q Capture writes the restart information to standard output (stdout) after it starts reading the log and after its last WebSphere MQ commit before terminating (Q Capture also writes the restart information to the restart data set or file at this time). Copy the restart information from the standard output or job log into the restart message.

ASN7203E *program_name : program_identifier : The send queue queue_name does not exist in the restart information. The program did not start.*

Explanation: The Q Capture program was started with the override_restartq parameter set to Y (yes). This setting prompts Q Capture to get its restart information from a file instead of the restart queue, and enables you to specify individual restart points for one or more send queues. However, Q Capture was unable to find a queue that is used by any active Q subscriptions in the restart file.

User response: In the restart file, specify a queue that is used by an active Q subscription and restart Q Capture.

ASN7205E *program_name : program_identifier : The Q Capture program cannot be started in warm mode by using the override_restartq=y parameter because the lsn and maxcmtseq parameters were also specified. Either specify lsn and maxcmtseq or specify override_restartq.*

Explanation: The Q Capture program can be started with global restart information by using the lsn and maxcmtseq parameters together during a warm start. Or, Q Capture can be started from different restart points for each send queue by using the override_restartq=y parameter. You cannot specify both lsn/maxcmtseq and override_restartq.

User response: Start Q Capture with global warm start information (lsn and maxcmtseq) or start Q Capture with warm start information from the restart file (override_restartq).

ASN7206E *program_name : program_identifier : The Q Capture program cannot write to the restart queue override file file_name because of a system error error_code. Q Capture will continue running.*

Explanation: The Q Capture program writes restart information for each send queue to a file with the data set name or file name capture_path/qrestart, where

capture_path is the value of the capture_path parameter. Either the user ID that started the Q Capture program does not have access authority for the file, or the file is too small.

User response: Ensure that the user ID that started the Q Capture program has correct access and security permissions for the file. Also, ensure that sufficient space has been allocated for the file.

ASN7207I *program_name : program_identifier : The contents of the restart queue override file are as follows: <send_queue_name>, <restart_lsn>, <maxcmtseq>, <partition> (only for partitioned databases):*

Explanation: The Q Capture program records restart information in a data set or file in addition to the restart message. Q Capture starts reading the log from the restart log sequence number (the lowest LSN of a transaction still to be committed) and ignores any transactions whose commit is before the maximum commit sequence number (the highest LSN of a successful transaction). In a partitioned database, these log indicators come from the partition where the Q Capture program is running.

User response: This message is for your information only. No action is required.

ASN7208I *program_name : program_identifier : An ALTER TABLE ALTER COLUMN statement was detected for column column_name of table table_name with a new data type of data_type. The column was automatically altered for Q subscription name.*

Explanation: The Q Capture program detected that a column that is part of a Q subscription was altered. A message is automatically sent to the Q Apply program to make a matching alteration of the column at the target.

User response: This message is for your information only. No action is required.

ASN7209I *program_name : program_identifier : Column column_name of table table_name was automatically added to publication or Q subscription name.*

Explanation: The Q Capture program detected that a column was added to a subscribed table, and the value of the repl_addcol parameter is Y (yes). The parameter specifies that newly added columns should automatically be added to the Q subscription for the table, and so the column was added. The column will be added to the target table if it does not already exist.

User response: This message is for your information only. No action is required.

ASN7210I *program_name : program_identifier : Q subscription name that corresponds to schema-level subscription name was successfully created for the source table table_owner.table_name that uses send queue queue_name and replication queue map queue_map_name.*

Explanation: The schema-level subscription specifies that the replication programs should automatically create Q subscriptions for all tables with the specified table owner. The table-level Q subscription was successfully created.

User response: This message is for your information only. No action is required.

ASN7211E *program_name : program_identifier : A Q subscription name already exists for the source table table_owner.table_name, and the Q subscription cannot be overwritten because it is in a state other than inactive (I). The Q subscription specifies the send queue queue_name and queue map queue_map_name. The program will stop.*

Explanation: A schema-level Q subscription exists that specifies that the replication programs should automatically create Q subscriptions for all tables with the same table owner. The Q Capture program tried to create a Q subscription for the specified table, but a Q subscription for the table already exists. If the Q subscription were inactive, a new Q subscription would have been created that overwrote the existing Q subscription.

User response: Follow these steps:

1. Use the replication administration tools to delete the existing Q subscription or set its state to inactive (I) so that it can be overwritten.
2. Start the Q Capture program.

ASN7212E *program_name : program_identifier : The Q subscription name could not be started because it was in the process of being stopped.*

Explanation: The Q Capture program had previously received a CAPSTOP signal for this Q subscription and was in the process of stopping it.

User response: After the Q subscription state changes to I (inactive), use the replication administration tools or insert a CAPSTART signal to start the Q subscription again.

ASN7213I *program_name : program_identifier : The program will publish changes up to database commit timestamp timestamp and then stop.*

Explanation: A stop command was issued with the captureupto option that specifies that Q Capture stop after publishing database transactions that are committed up to and including the provided timestamp. The timestamp is treated as being in the database time zone.

User response: This message is for your information only. No action is required.

ASN7214I *program_name : program_identifier : The program will publish changes up to database commit timestamp timestamp, wait for changes to be sent or applied, and then stop.*

Explanation: You can specify when a Q Capture program stops by using the parameter named captureupto. You can cause a Q Capture program to stop after the program publishes all of the database transactions that are committed up to and including a certain timestamp by specifying that timestamp with the captureupto parameter. This message is returned when the Q Capture program was stopped, and the parameter named captureupto was specified. The timestamp is treated as being in the database time zone.

The stopafter option was also specified in the stop command. The program will wait to stop until one of the following conditions is met, depending on which condition was specified:

- The transmission queue is empty (or Q Apply has consumed all messages if the queue is local).
- All changes up to the stopping point are applied at the target.

User response: This message is for your information only. No action is required.

ASN7215I *program_name : program_identifier : The program will publish transactions up to database commit timestamp timestamp for send queue queue_name and then deactivate the queue.*

Explanation: A stopq command was issued with the captureupto option that specified that Q Capture stop publishing changes to the given queue after publishing database transactions that are committed up to and including the specified timestamp. The timestamp is treated as being in the database time zone.

User response: This message is for your information only. No action is required.

ASN7216I *program_name : program_identifier : The program will publish transactions up to database commit timestamp timestamp for send queue queue_name, deactivate the queue, and then wait for changes to be sent or applied.*

Explanation: A stopq command was issued with the captureupto option that specified that Q Capture stop publishing changes to the given queue after publishing database transactions that are committed up to and including the specified timestamp. The timestamp is treated as being in the database time zone.

The stopafter option was also specified in the stopq command. The program will wait to stop until one of the following conditions is met, depending on which condition was specified:

- The transmission queue is empty (or Q Apply has consumed all messages if the queue is local).
- All changes up to the stopping point are applied at the target.

User response: This message is for your information only. No action is required.

ASN7217I *program_name : program_identifier : The program will publish changes up to the end of the active log and then stop.*

Explanation: A stop command was issued with the captureupto option eol (end of log), which specifies that Q Capture stop after publishing database transactions up to the end of the active database log.

User response: This message is for your information only. No action is required.

ASN7218I *program_name : program_identifier : The program will publish changes up to the end of the active log, wait for changes to be sent or applied, and then stop.*

Explanation: A stop command was issued with the captureupto option eol (end of log), which specifies that Q Capture stop after publishing database transactions up to the end of the active database log.

The stopafter option was also specified in the stop command. The program will wait to stop until one of the following conditions is met, depending on which condition was specified:

- The transmission queue is empty (or Q Apply has consumed all messages if the queue is local).
- All changes up to the stopping point are applied at the target.

User response: This message is for your information only. No action is required.

ASN7219I *program_name : program_identifier : The program will publish transactions up to the end of the active log for send queue queue_name, and then deactivate the queue.*

Explanation: A stopq command was issued with the captureupto option eol (end of log), which specifies that Q Capture stop publishing database transactions to the given queue after it reaches the end of the active database log.

User response: This message is for your information only. No action is required.

ASN7220I *program_name : program_identifier : The program will publish transactions up to the end of the active log for send queue queue_name, deactivate the queue, and then wait for changes to be sent or applied.*

Explanation: A stopq command was issued with the captureupto option eol (end of log), which specifies that Q Capture stop publishing database transactions to the given queue after it reaches the end of the active database log.

The stopafter option was also specified in the stopq command. After the queue is deactivated, the program will monitor and issue a message when one of the following conditions is met:

- The transmission queue is empty (or Q Apply has consumed all messages if the queue is local).
- All changes up to the stopping point are applied at the target.

User response: This message is for your information only. No action is required.

ASN7221E *program_name : program_identifier : The program found an incorrect value value for the captureupto parameter. The command or signal is ignored.*

Explanation: A stop or stopq command contains a captureupto parameter with an invalid timestamp. The timestamp must be specified in the time zone of the Q Capture server, in a full or partial timestamp format. The full timestamp uses the following format: YYYY-MM-DD-HH.MM.SS.mmmmmm

For examples of acceptable formats for partial timestamps, see, "Stopping the Q Capture program at a specified point."

As an alternative, you can specify the keyword CURRENT_TIMESTAMP, and the Q Capture program substitutes the current time, causing it to publish transactions that are committed up to that current time.

User response: Correct the timestamp format and reissue the command or signal.

ASN7222I *program_name : program_identifier : In response to a stop or stopq command with the stopafter option, the program will monitor send queue queue_name until all published messages have been delivered to the receive queue. The current queue depth of the local queue or transmission queue queue_name is queue_depth.*

Explanation: A stop or stopq command was specified with stopafter=data_sent. Q Capture will monitor the given send queue until all messages have been delivered to the receive queue (that is, until the queue depth becomes 0). For a stop command, Q Capture will terminate after all active queues are empty. For a stopq command, Q Capture will issue a message once the queue is empty. If Q Capture and Q Apply share a local queue, Q Capture will wait until Q Apply has removed all messages from the queue. In configurations where multiple remote queues share a transmission queue, Q Capture might wait longer than expected because of non-Q Capture messages being sent on the transmission queue.

User response: This message is for your information only. No action is required.

ASN7223I *program_name : program_identifier : In response to a stop or stopq command with the stopafter option, the program will wait for all transactions that are published to send queue queue_name to be applied at the target.*

Explanation: A stop or stopq command was specified with option stopafter=data_applied. Q Apply will notify Q Capture when all messages are applied for the given queue. For a stop command, Q Capture will stop after Q Apply sends a response for all active queues. For a stopq command, Q Capture will issue a message when it is notified by Q Apply. If Q Apply is not running, Q Capture will wait indefinitely.

User response: This message is for your information only. No action is required.

ASN7224I *program_name : program_identifier : The Q Apply browser thread for receive queue queue_name (replication queue map queue_map_name) notified the Q Capture program that all transactions that were published to send queue queue_name up to commit sequence number commit_seq and commit timestamp timestamp were applied. The browser thread continues to process messages.*

Explanation: A stop or stopq command was issued for Q Capture with option stopafter=data_applied. Q Capture requested that Q Apply send a notification when Q Apply has applied all transactions up to the

time that the command was issued. Q Apply has notified Q Capture, but continues to process the receive queue.

User response: This message is for your information only. No action is required.

ASN7225I *program_name : program_identifier : All published messages have been drained from send queue queue_name.*

Explanation: A stop or stopq command was issued for Q Capture with option stopafter=data_sent. Q Capture detected that all messages have been delivered to the remote receive queue (or consumed by Q Apply if it shares a local queue with Q Capture), and that the queue depth is now 0.

User response: This message is for your information only. No action is required.

ASN7226E *program_name : program_identifier : The Q Capture program could not read the restart file file_name for reason reason. The program did not start.*

Explanation: The Q Capture program was started with the override_restartq parameter set to Y (yes). This setting prompts Q Capture to look in the restart file instead of the restart queue for restart information. However, Q Capture was unable to read the restart file.

User response: Verify that the user ID that started Q Capture has read authority on the restart file, which has the data set name or file name capture_path/qrestart, where capture_path is the value of the capture_path parameter.

ASN7227I *program_name : program_identifier : All transactions that were published to send queue queue_name (replication queue map queue_map_name) have been applied at the target.*

Explanation: A stop or stopq command was specified with option stopafter=data_applied. Q Capture detected that all transactions that were published to the given send queue were applied to the target.

User response: This message is for your information only. No action is required.

ASN7228E *program_name : program_identifier : The program could not wait for published messages to be drained from send queue queue_name. The reason code is reason_code. The local queue or transmission queue is queue_name.*

Explanation: Q Capture could not reliably determine when all messages have been delivered to the remote receive queue, or drained from the local queue if Q Capture and Q Apply are sharing a local queue. Q

Capture no longer monitors the queue depth. The following values are valid for the reason code:

1

The transmission queue name could not be derived from the remote queue definition.

2

The transmission queue appears to be shared between Q Capture and another application.

User response: Review the reason codes in the explanation, and take the appropriate action:

1

Check that the remote queue definition contains a valid transmission queue name.

2

Q Capture cannot wait for messages to drain when other applications are sharing the transmission queue. Use a dedicated transmission queue for the send queue to avoid this problem.

ASN7229E *program_name : program_identifier : The program could not wait for published transactions on send queue queue_name to be applied at the target. The reason code is reason_code.*

Explanation: Q Capture could not determine when all transactions that were published to the given send queue were applied at the target. Q Capture is no longer waiting for the transactions to be applied. The following values are valid for the reason code:

1

Q Apply is not currently processing the associated receive queue.

2

One of the following conditions exists:

- Q Apply does not appear to be running, is not processing any receive queues that are receiving messages from this Q Capture, or the response message is not arriving on the Q Capture administration queue.
- A response was not received in a reasonable amount of time.

User response: Review the reason codes in the explanation, and take the appropriate action:

1

Issue a startq command for the corresponding receive queue at the target to allow the transactions to be applied.

2

Check that Q Apply is running, all relevant receive queues are active, and there is connectivity for messages on the administration queue.

ASN7230I *program_name : program_identifier : The program stopped successfully after the captureupto criteria, stopafter criteria, or both for a stop command or signal completed.*

Explanation: A stop command or signal was issued with captureupto criteria, stopafter criteria, or both that caused Q Capture to wait for a certain condition to become true before stopping.

User response: This message is for your information only. No action is required.

ASN7232E *program_name : program_identifier : The send queue queue_name for Q subscription name is not defined in the IBMQREP_SENDQUEUES table.*

Explanation: The program tried to load the Q subscription into memory, but the send queue name that was specified for the Q subscription could not be found in the IBMQREP_SENDQUEUES table. The error might have occurred because the Q subscription was updated with SQL other than that generated by the Replication Center or ASNCPLP program, or the SQL generated by these tools was modified.

User response: Use one of the administration tools to drop and recreate the Q subscription.

ASN7236W *program_name : program_identifier : The Q subscription for source table table_owner.table_name was successfully created. When the Q Capture program created the entry for this new Q subscription in control table table_owner.table_name based on the schema-level subscription name that uses send queue queue_name, Q Capture deleted an existing Q subscription entry for the same source table.*

Explanation: After a CREATE TABLE operation for a table with a matching schema-level subscription is detected, the Q Capture program automatically creates the Q subscription entry for this table in the IBMQREP_SUBS and IBMQREP_SRC_COLS control tables. If the Q subscription entry already exists in one or both of these control tables, Q Capture deletes the prior Q subscription entry and creates the new entry.

User response: You might want to determine why the Q subscription entry for the same source table already existed in this control table. If Q Capture was restarted from a prior point in the log, it might have encountered a log record for the same CREATE TABLE operation.

ASN7237W *program_name : program_identifier : In response to a CREATE TABLE operation at the source database or subsystem, the Q Capture program tried to create a Q subscription for source table `table_owner.table_name` based on the schema-level Q subscription `name` that uses send queue `queue_name`. However, the individual Q subscription could not be created because the insertion of Q subscription information into the IBMQREP_SUBS control table failed. The SQL code or internal replication reason code is `sqlcode_or_reason_code`. Q Capture will continue processing changes for other Q subscriptions that use the send queue.*

Explanation: After a CREATE TABLE operation for a table with a matching schema-level Q subscription is detected, the Q Capture program automatically creates the Q subscription entry for this table in the IBMQREP_SUBS control table. If the insertion of the Q subscription entry fails, the Q subscription is not created and replication does not start for the source table.

User response: If you want to replicate this source table, use the replication administration tools to create a Q subscription for the table. Make sure that the target table exists or let the replication tools create the target table. Check the SQL code or internal reason code to determine why the insert into IBMQREP_SUBS failed.

ASN7238W *program_name : program_identifier : The program tried to create a new Q subscription entry in the IBMQREP_SRC_COLS control table for source table `table_owner.table_name` based on the schema-level Q subscription `name` that uses send queue `queue_name`. However, the insertion of Q subscription information into the control table failed. The SQL code or internal replication reason code is `sqlcode_or_reason_code`. The Q subscription `name` for this source table will be deactivated. Q Capture will continue processing changes for other Q subscriptions on the send queue.*

Explanation: After a CREATE TABLE operation for a table with a matching schema-level Q subscription is detected, the Q Capture program automatically adds Q subscription information about the columns in the source table to the IBMQREP_SRC_COLS control table. However, the insertion of this information into the control table failed. The Q subscription will be deactivated in the IBMQREP_SUBS table.

User response: Follow these steps:

1. Check the SQL code or internal reason code to determine why the insertion into IBMQREP_SRC_COLS failed.
2. Use the replication administration tools to delete and recreate the Q subscription for the table. Make sure that the target table exists or let the replication tools create the target table.

ASN7240E *program_name : program_identifier : The command to start the Q Capture program contained both the `autostop=Y` and `term=N` parameters. The Q Capture program will stop.*

Explanation: You cannot specify `autostop=Y` when the value of the term parameter is N (No). You also cannot specify `term=N` when the saved value of `autostop` in the IBMQREP_CAPPARMS table is Y.

User response: Restart the Q Capture program while specifying either the `autostop` parameter or the `term=N` parameter. Update the value of `autostop` in the IBMQREP_CAPPARMS table if necessary.

ASN7241E *program_name : program_identifier : An invalid expression `expression_list` was detected in the SCHEMA_NAME column of the IBMQREP_SCHEMASUBS table for the schema-level Q subscription `name` that uses replication queue map `queue_map_name`. The Q Capture program will stop.*

Explanation: You can use the percentage sign (%) as a wild card to specify schema-level Q subscriptions and the Q Capture program automatically creates a Q subscription for all tables within the schemas that match the wild card expression. However, you cannot specify the wild card before a schema name, for example "%AN" or "%AN%".

User response: Follow these steps:

1. Change the expression in the schema name in the IBMQREP_SCHEMASUBS table to an allowed wild card expression.
2. Start the Q Capture program.
3. Insert a START_SCHEMASUB signal in the IBMQREP_SIGNAL table.

ASN7242E *program_name : program_identifier : An invalid expression `expression_list` was detected in the OBJECT_NAME column of the IBMQREP_SCHEMASUBS table for the schema-level Q subscription `name` that uses replication queue map `queue_map_name`. The Q Capture program will stop.*

Explanation: You can use the percentage sign (%) as a wild card to specify schema-level Q subscriptions and

the Q Capture program automatically creates a Q subscription for all tables within the schemas that match the wild card expression. However, you cannot specify the wild card before an object name, for example "%AN" or "%AN%".

User response: Follow these steps:

1. Change the expression in the object name in the IBMQREP_SCHEMASUBS table to an allowed wild card expression.
2. Start the Q Capture program.
3. Insert a START_SCHEMASUB signal in the IBMQREP_SIGNAL table.

ASN7243E *program_name : program_identifier :* **An overlapping expression was detected either in the SCHEMA_NAME or OBJECT_NAME column of the IBMQREP_SCHEMASUBS table for the schema-level Q subscription name that uses replication queue map queue_map_name. The Q Capture program will stop.**

Explanation: Overlapping expressions in the schema name are not allowed. For example, expressions such as "ANU%" and "AN%" cannot be used at the same time.

User response: Follow these steps:

1. Change the expression in the schema name or object name in the IBMQREP_SCHEMASUBS table.
2. Start the Q Capture program.
3. Insert a START_SCHEMASUB signal in the IBMQREP_SIGNAL table.

ASN7244E *program_name : program_identifier :* **An expression was detected in the SCHEMA_NAME column of the IBMQREP_EXCLSCHEMA table. The Q Capture program will stop.**

Explanation: You can use an expression to exclude objects like tables from schema-level Q subscriptions. The expressions are stored in the IBMQREP_EXCLSCHEMA table. However, you cannot use expressions in the SCHEMA_NAME column of the IBMQREP_EXCLSCHEMA table.

User response: Remove the expression in the SCHEMA_NAME column of the IBMQREP_EXCLSCHEMA table and start the Q Capture program.

ASN7247I *program_name : program_identifier :* **The Q Capture program successfully loaded the schema-level subscription name and its corresponding profile. The Q subscription uses replication queue map queue_map_name and specifies that Q subscriptions should automatically be created in schema schema_name for object_name objects.**

Explanation: As a result of processing the log record that corresponds to a REINIT_SCHEMASUBS signal, the Q Capture program successfully loaded the schema-level subscription that is defined in the IBMQREP_SCHEMASUBS table and the corresponding profile that is defined in the IBMQREP_SUBS_PROF table.

User response: This message is for your information only. No action is required.

ASN7248W *program_name : program_identifier :* **The Q Capture program detected a table table_owner.table_name that matches the schema-level subscription name but is not part of a table-level Q subscription.**

Explanation: Q Capture program found an unsubscribed table for the schema-level subscription that is defined in the IBMQREP_SCHEMASUBS table. This situation is unexpected because the replication administration tools create Q subscriptions for all existing tables within a schema when a schema-level Q subscription is created. However, this situation can occur if new tables are created between the time that the table-level Q subscriptions are created and the schema-level subscription is loaded. For example, a user could generate the script to create Q subscriptions for all existing tables, but wait days before running this script, during which time new tables are created.

User response: Use the replication administration tools to create a Q subscription for this table if desired.

ASN7249E *program_name : program_identifier :* **An invalid expression expression_list was detected in the OBJECT_NAME column of the IBMQREP_EXCLSCHEMA table.**

Explanation: You can use the percentage sign (%) as a wild card in the OBJECT_NAME column of the IBMQREP_EXCLSCHEMA table to specify the names of objects such as tables that should be excluded from a schema-level subscription. Use a single % to specify that all objects for the given schema be excluded. However, you cannot specify the wild card before an object name, for example "%AN" or "%AN%".

User response: Change the expression in the OBJECT_NAME column of the IBMQREP_EXCLSCHEMA table and start the Q Capture program.

ASN7250E *program_name : program_identifier : The Q subscription name for source table table_owner.table_name that is based on the schema-level Q subscription name (receive queue queue_name, replication queue map queue_map_name) could not be activated. Replication should be manually set up for this table.*

Explanation: After a CREATE TABLE operation is detected for a table with a matching schema-level Q subscription, the Q Capture program informs the Q Apply program so that it can automatically create and activate a table-level Q subscription. If the Q subscription activation fails, the Q subscription will be deactivated and future changes will be ignored for this source table.

User response: Use the replication administration tools to create a new Q subscription for the table or to activate the existing Q subscription. Make sure that the source and target tables exist or let the replication tools create the target table. Also, verify that the table will be loaded either by the Q Replication programs or externally.

ASN7251I *program_name : program_identifier : The Q Capture program successfully deactivated the schema-level Q subscription name that uses replication queue map queue_map_name, schema schema_name, and object object_name.*

Explanation: In response to a STOP_SCHEMASUB signal that was inserted into the IBMQREP_SIGNAL table, the Q Capture program deactivated the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN7252I *program_name : program_identifier : The Q Capture program successfully reinitialized the schema-level Q subscription name that uses replication queue map queue_map_name, schema schema_name, and object object_name.*

Explanation: In response to a REINIT_SCHEMASUB signal that was inserted into the IBMQREP_SIGNAL table, the Q Capture program reinitialized the schema-level Q subscription.

User response: This message is for your information only. No action is required.

ASN7253E *program_name : program_identifier : An ALTER TABLE ALTER COLUMN SET DATA TYPE operation was detected for source table table_owner.table_name. The Q Capture program does not support replication of this operation when its compatibility level is level. All Q subscriptions for this table will be deactivated.*

Explanation: The Q Capture program can replicate some DDL operations at source tables. However, replication of the ALTER TABLE ALTER COLUMN SET DATA TYPE operation is not supported if the value of the Q Capture compatibility parameter is 0907 or lower.

User response: Use the replication administration tools to update the value of the compatibility parameter to 1001 or newer and start the Q subscriptions.

ASN7254W *program_name : program_identifier : The DB2 database configuration parameter LOG_DDL_STMTS must be set to YES before DDL statements can be replicated.*

Explanation: At least one schema-level subscription is active, but the database is not currently logging DDL statements. To replicate CREATE TABLE and other DDL statements, the database configuration must be updated so that DB2 logs these statements.

User response: Update the database configuration parameter LOG_DDL_STMTS to YES. The change is reflected in the database immediately.

ASN7255E *program_name : program_identifier : An operation by utility utility_name was performed on table table_owner.table_name while Q subscription name (send queue queue_name, replication queue map queue_map_name) was still in loading state. The Q subscription is being deactivated.*

Explanation: A utility operation was detected in the log while the Q subscription was still in loading state. The Q subscription was set up to replicate source table loads (CAPTURE_LOAD R). Invoking a utility while the Q subscription is in loading state disrupts the load protocol and might lead to SQL errors while changes are being applied from the spill queue to target tables.

User response: Wait for the Q subscription to stop and become inactive. Then start the Q subscription. If CAPTURE_LOAD is set to R, always wait for the Q subscription to become active before invoking a utility that triggers a reload of the target table.

ASN7256W *program_name : program_identifier : A signal_type signal was received for source table table_owner.table_name, Q subscription name. The Q Capture program cannot process this signal when its compatibility level is level. The Q Capture program ignores the signal.*

Explanation: The signal that was received cannot be processed because the resulting message that would be sent to Q Apply is restricted by the current compatibility level.

User response: If Q Apply is at the same level as Q Capture, set the Q Capture compatibility level to match the Q Capture version that is installed and reissue the signal.

ASN7257W *program_name : program_identifier : Replication of DDL operations is not available when the Q Capture compatibility level is level. No DDL operations will be replicated for tables that match schema-level subscription patterns.*

Explanation: The current compatibility level prevents sending of messages to Q Apply that are required to automatically create Q subscriptions and replicate DDL operations. Replication of DDL operations requires a compatibility level of 1001 or higher.

User response: If Q Apply is at the same level as Q Capture, set the compatibility level in Q Capture to match the Q Capture version that is installed.

ASN7263W *program_name : program_identifier : An ALTER TABLE RENAME COLUMN operation was detected for a source table. The Q subscription was stopped or the receive queue that is used for this Q subscription was stopped, depending on what error action was specified for the Q subscription. Source table: table_name. Q subscription: subname. Queue name: queue_name. Replication queue map: queue_map_name. New column name: column_name.*

Explanation: The Q Capture program detected an ALTER TABLE RENAME COLUMN operation for the source table. The replication programs do not automatically handle the renaming of a column that is part of a Q subscription. The programs took action based on the value in the ERROR_ACTION column of the IBMQREP_TARGETS table.

User response: Follow these steps:

1. Change the name of the corresponding column at the target table to match the new source column name.

2. Update the entries for this column in the SRC_COLNAME column of the IBMQREP_SRC_COLS table and the SOURCE_COLNAME and TARGET_COLNAME columns of the IBMQREP_TRG_COLS table.
3. Check the value of ERROR_ACTION in the IBMQREP_TARGETS table and take one of the following actions:
 - If the Q subscription was stopped (ERROR_ACTION of D), start the Q subscription.
 - If the receive queue was stopped (ERROR_ACTION of Q), start the queue.
 - If the Q Apply program was stopped (ERROR_ACTION of S), start the Q Apply program.

ASN7307W *program_name : program_identifier : The table table_owner.table_name that is part of the schema-level Q subscription name that uses send queue queue_name was not enabled for replication because the table was explicitly created with DATA CAPTURE NONE.*

Explanation: When a CREATE TABLE operation for a table with a matching schema-level Q subscription is detected, the table must have the DATA CAPTURE CHANGES attribute set. If the CREATE TABLE command explicitly specifies DATA CAPTURE NONE, it is treated as an indication that the user does not want this table replicated, and this takes precedence over any matching schema-level Q subscription.

User response: If you want to replicate this table, follow these steps:

1. Alter the table with DATA CAPTURE CHANGES.
2. Use the replication administration tools to create a table-level Q subscription for the table.

Chapter 46. ASN7500 - ASN7999

ASN7504E *program_name : program_identifier : The Q Apply program cannot obtain number bytes of storage in routine routine_name. The program stops.*

Explanation: The Q Apply program did not have enough physical memory that was requested from the operating system to allocate an internal data structure.

User response: Ensure that enough memory is available from the operating system before starting the Q Apply program.

ASN7505E *program_name: program_identifier: The attributes of the target columns are not set up correctly for the Q subscription name on the receive queue queue_name, for the replication queue map queue_map_name. The reason code is reason_code.*

Explanation: An error in the setup of the columns of the IBMQREP_TRG_COLS table. The following values are valid for the reason code:

0

The number of columns specified in the IBMQREP_TRG_COLS table does not match the number of columns in the target table. There are more columns defined in the IBMQREP_TRG_COLS table than exist in the target.

1

The value of the MSG_COL_NUMBER column in the IBMQREP_TRG_COLS table is not unique.

2

The key columns are not the leading columns with respect to the MSG_COL_NUMBER in the IBMQREP_TRG_COLS table.

3

The value in the MSG_COL_NUMBER column is out of range. The value is either less than zero or greater than the maximum columns that are defined by the IBMQREP_TRG_COLS table.

4

No columns are defined that can uniquely identify a row. Q Replication requires that at least one column is unique.

5

6

One or more of the following columns are NULL in the IBMQREP_TRG_COLS table: MSG_COL_NUMBER, MSG_COL_TYPE, MSG_COL_CODEPAGE, or MSG_COL_LENGTH.

No unique index corresponds to the declared key columns. The DB2 database must have a constraint that one row from the source table corresponds to exactly one row in the target table. The Q Apply program uses unique index information to schedule parallelism of transactions. Therefore a unique index must be declared for each target table.

7

The target column has SQL Expressions in a key column but the source columns are not part of replication IS_KEY. All source columns used in a key expression must be defined as key columns in the IBMQREP_SRC_COLS control table at the capture server.

8

The Q subscription specifies an expression on both XML columns and key columns. You cannot replicate XML expressions on both XML columns and key columns.

User response: Review the reason codes in the explanation, and take the appropriate action:

Reason Codes 0, 1, 2, 3, and 5 require the same user response. For each of these reason codes take the following user response action.

Redefine your Q subscription by using the replication administration tools and activate the Q subscription.

4

Redefine the Q subscription and identify the unique columns by using one of the replication administration tools. Then activate the Q subscription.

6

1. Choose a different set of columns that have a unique index defined as the key columns for this Q subscription,
2. or, create a unique index or a unique constraint on the target table.
3. Then restart the corresponding Q subscription.

7

Set the value of IS_KEY to 1 for the affected source columns in the IBMQREP_SRC_COLS control table.

8

Redefine your Q subscription and remove the expression on either the key columns or the XML columns. Then restart the corresponding Q subscription.

If the problem persists, contact IBM Software Support. Provide an Analyzer report.

ASN7506E *program_name : program_identifier : The target target_name for Q subscription name (receive queue queue_name, replication queue map queue_map_name) does not exist. The Q subscription is not loaded and changes cannot be replicated to this target.*

Explanation: The target for the Q subscription as defined in the IBMQREP_TARGETS table does not exist. Either the target was specified incorrectly or it does not exist.

User response: Make sure that the target exists in the target database. Verify that the values are correct for this target in the TARGET_OWNER, TARGET_NAME columns in the IBMQREP_TARGETS table. Also, make sure that the value in the TARGET_TYPE column is appropriate for the target. For example, if the target is a stored procedure make sure that the target type in the IBMQREP_TARGETS table is '5'. If the Q subscription is not valid, redefine it using one of the replication administration tools.

ASN7510E *program_name : program_identifier : The Q Apply program received an ADD COLUMN message for Q subscription name (receive queue queue_name, replication queue map queue_map_name), but the Q subscription is inactive.*

Explanation: The Q subscription must be active before you can add a new column to it.

User response: Start the Q subscription.

ASN7512E *program_name : program_identifier : The Q Apply program could not activate the Q subscription name (receive queue queue_name, replication queue map queue_map_name). Reason code: reason_code.*

Explanation: The Q subscription definition is incorrect. The error may have occurred because the Q subscription was updated with SQL other than that generated by the Replication Center or ASNCLP program, or the SQL generated by these administrative

tools was modified. The following values are valid for the reason code:

0

The Q subscription does not exist in the IBMQREP_TARGETS control table.

1

The value of the STATE column in the IBMQREP_TARGETS control table is not I.

2

The type of Q subscription on the Q Capture and Q Apply servers is different.

3

The DESCRIBE statement failed for the target.

4

The value in the TARGET_COLNAME column in the IBMQREP_TRG_COLS table does not match any of the target columns in the target.

5

A value in the TARGET_COLNAME, MSG_COL_TYPE, or MSG_COL_LENGTH columns in the IBMQREP_TRG_COLS table does not match the name, type, or length of the corresponding column in the target table or stored procedure.

6

The SOURCE_COLNAME column in the IBMQREP_TRG_COLS table does not match the value of the SRC_COLNAME in the IBMQREP_SRC_COLS table.

7

The Q subscription ID is not unique for the receive queue.

8

The Q subscription is defined as bidirectional and, according to the conflict rules, the Q Capture sending options are incorrect. The CONFLICT_RULE column in the IBMQREP_TARGETS table is set to A or C and the sending option BEFORE_VALUES column in the IBMQREP_SUBS table is not set to Y.

9

The values of the SOURCE_OWNER, SOURCE_NAME columns in the IBMQREP_TARGETS table do not match the values in the IBMQREP_SUBS table. Also the value of the SOURCE_SERVER column in the IBMQREP_TARGETS table is not the server on which the Q Capture program is running.

10

| | | | |
|----|---|----|--|
| | The value of the IS_KEY column in the IBMQREP_TRG_COLS table does not match the value of the IS_KEY column in the IBMQREP_SRC_COLS table. | | |
| 11 | | 20 | The Q subscription is type 'U' (unidirectional), but the CONFLICT_RULE column in the IBMQREP_TARGETS table is not set to K. |
| | The code page of the source column cannot be converted to the Q Apply program code page. | | The target type for this Q subscription is a CCD table. CCD tables must have the following four columns: |
| 12 | | | <ul style="list-style-type: none"> • IBMSNAP_INTENTSEQ • IBMSNAP_OPERATION • IBMSNAP_COMMITSEQ • IBMSNAP_LOGMARKER |
| | The CONFLICT_ACTION column in the IBMQREP_TARGETS table is set to F, but the Q Capture sending options are incorrect. For the CONFLICT_ACTION F, the message needs to include all the columns not just the changed columns. The CHANGED_COLS_ONLY should be set to N in the IBMQREP_SUBS table. | 21 | |
| 13 | | | The target type for this Q subscription is a CCD table and the subscription type is either for bidirectional or peer-to-peer replication. CCD tables are only supported in unidirectional replication. This error should not occur unless one of the following is true: |
| | The CONFLICT_RULE in the IBMQREP_TARGETS is set to K, but the Q Capture sending options are incorrect. The BEFORE_VALUES column should be set to N in the IBMQREP_SUBS table. | | <ul style="list-style-type: none"> • The Q subscription has been updated with SQL other than that generated by the Replication Center or ASNCPL command-line program. • The SQL script generated by the Replication Center or ASNCPL was modified before it was run. |
| 14 | | 22 | |
| | The Q Capture sending options specifies CHANGED_COLS_ONLY=N in the IBMQREP_SUBS table, but the value should be Y because of one or both of the following settings in the IBMQREP_TARGETS table: | | The target is a noncondensed CCD table. The only valid value for CONFLICT_ACTION in the IBMQREP_TARGETS table is F (force). Because all incoming changes are always inserted, the only valid Q Capture sending options in the IBMQREP_SUBS table are BEFORE_VALUES = Y and CHANGED_COLS_ONLY = N. |
| | <ul style="list-style-type: none"> • CONFLICT_ACTION is not F • CONFLICT_RULE is not A | | |
| 15 | | 23 | |
| | There are extra columns in the target that do not exist in the IBMQREP_TRG_COLS table and they are NOT NULL and are not defined with default. | | |
| 16 | | | The target type for the Q subscription is a condensed and complete CCD table. The value for CONFLICT_ACTION in the IBMQREP_TARGETS table must be F (force) or I (ignore). For F (force) the only valid Q Capture sending options in the IBMQREP_SUBS table are BEFORE_VALUES = Y and CHANGED_COLS_ONLY = N. For I (ignore) the only valid Q Capture sending options are BEFORE_VALUES = N and CHANGED_COLS_ONLY = Y. |
| | The value of the SEARCH_CONDITION column in the IBMQREP_SUBS table is not specified correctly. | | |
| 17 | | 24 | |
| | The values of the SOURCE_NODE and TARGET_NODE columns in the IBMQREP_SUBS and the IBMQREP_TARGETS table do not match. | | |
| 18 | | | The target is a condensed and noncomplete CCD table. The CONFLICT_ACTION in the IBMQREP_TARGETS table is F (force). The only valid Q Capture sending options in the IBMQREP_SUBS table are BEFORE_VALUES = Y and CHANGED_COLS_ONLY = N. |
| | The stored procedure parameters which represent the columns do not match the SRC_COLNAME fields of the IBMQREP_SRC_COLS table. | | |
| 19 | | | |

ASN7512E

| | | | |
|----|--|----|---|
| 25 | The target is a CCD table. The value in one or both of the columns CCD_CONDENSED and CCD_COMPLETE in the IBMQREP_TARGETS table are null. | 2 | IBMQREP_TARGETS control table to I. On the Q Capture server, deactivate and activate the Q subscription. |
| 26 | The target type for this Q subscription is a noncomplete CCD table, but the HAS_LOADPHASE column in the IBMQREP_SUBS table has a value other than N. | 3 | Check that the value of the SUBTYPE column in the IBMQREP_SUBS table matches the value of the SUBTYPE column in the IBMQREP_TARGETS tables. Redefine the Q subscription by using one of the replication administration tools. |
| 27 | The target is a CCD table. The Q Apply program read a message from the Q Capture program that has a version lower than Version 9 (for example, Version 8.2). The Q Apply program can process such down-level messages, but Version 9 function such as CCD targets cannot be supported in this configuration. | 4 | Refer to the Q Apply diagnostic log file or the IBMQREP_TRACE table for the message ASN0552E, which contains the SQL return code information. |
| 28 | The target column information in IBMQREP_TRG_COLS table is not set up correctly. | 5 | Verify that the Q subscription is set up correctly. Do a describe on the target table and the values stored in the IBMQREP_TRG_COLS table for the target table to identify differences. Redefine the Q subscription using one of the replication administration tools. |
| 29 | The target type for this Q subscription is a stored procedure and the Q subscription is for replication from or to a database with the DB2 pureScale Feature. Stored procedure targets are only supported in unidirectional replication and are not supported with the DB2 pureScale Feature. | 6 | Verify that the Q subscription is set up correctly. Do a describe on the target table or stored procedure and check the values stored in the IBMQREP_TRG_COLS table to identify differences. Redefine the Q subscription using one of the replication administration tools. |
| 30 | The target table is a system period temporal table or a bitemporal table on a DB2 whose version is older than V10.1 Fix Pack 2 on Linux, UNIX, and Windows or APAR PM31315 on z/OS. The Q Apply program does not support replication to temporal tables before these versions. | 7 | Redefine the Q subscription using one of the replication administration tools. |
| | User response: Review the reason codes in the explanation, and take the appropriate action: | 8 | Redefine the Q subscription using one of the replication administration tools. |
| 0 | Verify the definition of the Q subscription, the SUBNAME column in the IBMQREP_SUBS table and the IBMQREP_TARGETS table. Redefine the Q subscription using the replication administration tools and activate the Q subscription. | 9 | Redefine the Q subscription using one of the replication administration tool. If the problem persists, contact IBM Software Support. Provide an Analyzer report. |
| 1 | Deactivate the Q subscription by setting the value of the STATE column of the | 10 | Redefine the Q subscription using one of the replication administration tool. If the problem persists, contact IBM Software Support. Provide an Analyzer report. |
| | | 11 | Redefine the Q subscription using one of the replication administration tool. If the problem persists, contact IBM Software Support. Provide an Analyzer report. |

| | | | |
|----|---|----|--|
| | Refer to the Q Apply diagnostic log file for the message ASN0568E, which specifies the CCSID that was not converted. Recommendation: Change the Q Apply code page to be the same as the source database code page. | | Do one of the following things: |
| 12 | | | <ul style="list-style-type: none"> Correct any SQL that was generated by the Replication Center or ASNCLP program and then modified. Drop the Q subscription and recreate it by using one of the administration tools. |
| | Redefine Q subscription using one of the replication administration tools. If the problem persists, contact IBM Software Support. Provide an Analyzer report. | 22 | |
| 13 | | | Change the CONFLICT_ACTION column in the IBMQREP_TARGETS table and the BEFORE_VALUES and CHANGED_COLS_ONLY columns in the IBMQREP_SUBS table to the values that are specified in the explanation. |
| | Redefine Q subscription using one of the replication administration tools. If the problem persists, contact IBM Software Support. Provide an Analyzer report. | 23 | |
| 14 | | | Change the CONFLICT_ACTION column in the IBMQREP_TARGETS table and the BEFORE_VALUES and CHANGED_COLS_ONLY columns in the IBMQREP_SUBS table to the values that are specified in the explanation. |
| | Redefine Q subscription using one of the replication administration tools. If the problem persists, contact IBM Software Support. Provide an Analyzer report. | 24 | |
| 15 | | | Change the CONFLICT_ACTION column in the IBMQREP_TARGETS table and the BEFORE_VALUES and CHANGED_COLS_ONLY columns in the IBMQREP_SUBS table to the values that are specified in the explanation. |
| | Modify the target attributes or subscribe to those extra columns. | 25 | |
| 16 | | | In the IBMQREP_TARGETS table, set the CCD_CONDENSED column to Y for condensed CCDs or N for noncondensed CCDs; set the CCD_COMPLETE column to Y for complete CCDs or N for noncomplete CCDs. |
| | Check the text in the SEARCH_CONDITION column of the IBMQREP_SUBS table and ensure that only the column names of the table being replicated are enclosed in square brackets. Redefine the Q subscription using one of the replication administration tools. | 26 | |
| 17 | | | If the target table needs to be loaded, the CCD table must be complete. If you need a noncomplete CCD table, the value for HAS_LOADPHASE must be N. |
| | Redefine the Q subscription using one of the replication administration tools. If the problem persists, contact IBM Software Support. Provide an Analyzer report. | 27 | |
| 18 | | | Upgrade the Q Capture program at the source system to Version 9 and then start the Q subscription. |
| | Redefine the Q subscription using one of the replication administration tools. If the problem persists, contact IBM Software Support. Provide an Analyzer report. | 28 | |
| 19 | | | Look in the Q Apply diagnostic log file or IBMQREP_APPLYTRACE table for messages that might indicate the problem with the column definition. Redefine the Q subscription using one of the replication administration tools, and then start the Q subscription. |
| 20 | | 29 | |
| | Use the Replication Center or ASNCLP program to add the four mandatory columns to the CCD table, and then activate the Q subscription. | | |
| 21 | | | |

Use a source database, target database, or both that does not have the DB2 pureScale Feature.

30

Upgrade DB2 on the target database to V10.1 Fix Pack 2 on Linux, UNIX, and Windows or APAR PM31315 on z/OS.

ASN7513W *program_name : program_identifier : A row change for SUB_ID subid was received, but there is no active Q subscription (receive queue queue_name, replication queue map queue_map_name). The Q Apply program cannot apply the change. Reason code: reason_code.*

Explanation: A row in the transaction does not belong to an active Q subscription due to one of the following reasons:

- 0 The Q subscription is inactive due to ERROR_ACTION or CONFLICT_ACTION but the Q Capture program did not stop sending changes yet. Refer to the IBMQREP_EXCEPTIONS table for the reason why the Q subscription was inactivated.
- 1 The Q Apply program never activated the Q subscription because the Q subscription setup is not valid.
- 2 The Q subscription does not exist in the IBMQREP_TARGETS table.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 1. Check for a Q subscription that corresponds to the SUB_ID in the IBMQREP_SUBS and IBMQREP_TARGETS control tables.
- 2. If the Q subscription exists, look at the STATE_INFO column in the IBMQREP_TARGETS table to determine if the Q subscription was deactivated due to the CONFLICT_ACTION or ERROR_ACTION.
 - a. If the Q subscription was deactivated, the Q Capture program will eventually stop sending changes for this Q subscription. No action is required on your part.
 - b. If the Q subscription was not deactivated, then it was never activated to begin with. Refer to message ASN7512E in the Q Apply diagnostic log for the exact reason code and the required response.

ASN7514W *program_name : program_identifier : Administration queue queue_name is full. Q subscription: name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The number of messages reached the number that is set for the MAXDEPTH attribute on the

administration queue. The administration queue is full and the Q Apply browsers cannot write to it. Perhaps the Q Capture program is not reading from that administration queue.

User response: Ensure that the Q Capture program is running. If it is not, restart it.

If necessary, increase the value for the MAXDEPTH attribute of the administration queue.

ASN7515E *program_name : program_identifier : The Q Apply program cannot process message type message_type of version message_version.*

Explanation: This message version is not supported by this version of the Q Apply program. The Q Capture program is not compatible with the Q Apply program version.

User response: Set the value of the COMPATIBILITY column in the IBMQREP_CAPPARMS table to match the version of the Q Apply program. The Q Apply version is recorded in the ARCH_LEVEL column of the IBMQREP_APPLYPARMS table. If Q Capture works with multiple Q Apply programs, set COMPATIBILITY to match the version of the oldest Q Apply program. You can use the Replication Center or Q Replication Dashboard to change the Q Capture COMPATIBILITY value.

ASN7516E *program_name : program_identifier : The SQL statement for operation operation on target table table_owner.table_name is too long.*

Explanation: The combined size of the target table schema (owner name) and table name cannot exceed 256 bytes.

User response: Specify a target table for the Q subscription that is within the length limit.

ASN7517E *program_name : program_identifier : The load done received message for the Q subscription name (receive queue queue_name, replication queue map queue_map_name) cannot be processed. Reason code: reason_code.*

Explanation: The Q Apply program received the load done receive message from Q Capture, but cannot process the message. The following values are valid for the reason code:

- 0 The Q subscription information that is loaded in memory is not correct.
- 1 The Q subscription state in the STATE column of the IBMQREP_TARGETS table is not correct.
- 2 There is no spill queue defined for this Q

subscription. Either the information in the SPILLQ column in the IBMQREP_TARGETS is wrong or the physical queue does not exist.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 0 Check if the Q subscription is in inactive (I) state in the STATE column in the IBMQREP_TARGETS table and refer to the Q Apply diagnostic log file for the reason. Fix the problem, and activate the Q subscription.
- 1 The Q subscription value in STATE column in the IBMQREP_TARGETS table is not as expected. Deactivate and then activate the Q subscription.
- 2 Deactivate and then activate the Q subscription.

ASN7519E *program_name : program_identifier : The Q subscription name (receive queue queue_name, replication queue map queue_map_name) cannot be loaded because the value in the SUB_ID column is NULL.*

Explanation: The SUB_ID can be NULL only if the Q subscription state is 'I' (Inactive). The Q subscription definition was altered.

User response: Deactivate the Q subscription, redefine it, and activate it.

ASN7522E *program_name : program_identifier : The Q Apply program stopped because it encountered an error for Q subscription name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The Q Apply program encountered an error or conflict while trying to apply a transaction for this target. The error or action for this Q subscription is 'S' (stop the Q Apply program). Refer to the Q Apply diagnostic log file and the IBMQREP_EXCEPTIONS table for details as to why the error occurred.

User response: Address the problem that was identified in the diagnostic log file or the IBMQREP_EXCEPTIONS table and restart the Q Apply program. No changes will be lost. If the problem is only related to the target for this Q subscription, deactivate the Q subscription and restart the Q Apply program. The Q Apply program will be able to apply changes to other targets.

ASN7523E *program_name : program_identifier : The Q Apply program encountered an error or conflict for Q subscription name (receive queue queue_name, replication queue map queue_map_name). The Q Apply program will stop reading from the receive queue.*

Explanation: The Q Apply program encountered an error or conflict while trying to apply a transaction for this target. The error or action for this Q subscription is 'Q' (stop reading from the queue). Refer to the Q Apply diagnostic log file and the IBMQREP_EXCEPTIONS table for details as to why the error occurred.

User response: Address the problem that was identified in the diagnostic log file or the IBMQREP_EXCEPTIONS table and resume reading from the queue using the startq command. No changes will be lost. If the problem is only related to the target for this Q subscription, deactivate the Q subscription and resume reading from the queue. The Q Apply program will be able to apply changes to other targets.

ASN7524E *program_name : program_identifier : The Q Apply program encountered an error or conflict for Q subscription name (receive queue queue_name, replication queue map queue_map_name). It stopped reading from the receive queue. The error occurred while applying a row from the transaction corresponding to the LSN LSN.*

Explanation: The Q Apply program encountered an error or conflict while trying to apply a transaction for this target. The error or action for this Q subscription is 'Q' (stop reading from the queue).

User response: Address the problem that was identified in the diagnostic log file or the IBMQREP_EXCEPTIONS table and resume reading from the queue using the startq command. No changes will be lost. If the problem is only related to the target for this Q subscription, deactivate the Q subscription and resume reading from the queue. The Q Apply program will be able to apply changes to other targets.

ASN7525W *program_name : program_identifier : The receive queue queue_name (replication queue map queue_map_name) is not in active state and is not being processed by the Q Apply program. If the skiptrans parameter was specified, it is ignored because the queue is inactive.*

Explanation: The receive queue is inactive. The STATE column in the IBMQREP_RECVQUEUES table is marked as I (inactive) for this receive queue. If the skiptrans parameter was specified on Q Apply program invocation, the parameter is ignored.

User response: This message is for your information only. No action is required. If you want the Q Apply program to process the queue, use one of the following methods:

- Issue the startq command without restarting Q Apply.
- Change the STATE column in the IBMQREP_RECVQUEUES table to A and restart Q Apply.

ASN7526I *program_name : program_identifier : The Q Apply program has started processing the receive queue queue_name for replication queue map queue_map_name.*

Explanation: The Q Apply program has started reading from the receive queue.

User response: This message is for information only. No action is required.

ASN7527I *program_name : program_identifier : The Q Apply browser for receive queue queue_name for replication queue map queue_map_name is activated as a result of the startq command.*

Explanation: The startq command was issued for the receive queue, and the Q Apply program has started the Q Apply browser for the queue.

User response: This message is for your information only. No action is required.

ASN7528I *program_name : program_identifier : The Q Apply program for the Q subscription name (receive queue queue_name, replication queue map queue_map_name) will use the utility_name utility to load table table_name.*

Explanation: The loadphase for this Q subscription is specified as internal. Q Apply has chosen this utility to perform the load.

User response: This message is for your information only. No action is required.

ASN7529I *program_name : program_identifier : The utility_type utility for table target_name completed successfully for the Q subscription name (receive queue queue_name, replication queue map queue_map_name). The message from the utility is text.*

Explanation: The load utility specified finished successfully. Look at the detail message from the utility for the statistics.

User response: This message is for your information only. No action is required.

ASN7530E *program_name : program_identifier : The load utility utility_name for table target_name failed for Q subscription name (receive queue queue_name, replication queue map queue_map_name). Detailed message from the load utility is text.*

Explanation: The load utility encountered an error for the Q subscription. The Q subscription cannot be activated.

User response: Refer to the detailed message from the load utility for the return code and the SQL return code. Fix the problem and then activate the Q subscription. If the problem persists, select a different load utility.

ASN7531I *program_name : program_identifier : The load utility utility_name was selected by the Q Apply program for Q subscription name (receive queue queue_name, replication queue map queue_map_name). Reason code: reason_code.*

Explanation: The Q subscription specifies that the Q Apply program will load the target table. Q Apply chose this load utility based on the replication environment. The following values are valid for the reason code:

0

The target is DB2 for Linux, UNIX, and Windows Version 8 or later and one of the following conditions is true:

- A nickname is defined for the source table or the source table is local to the target system, and the server option CONCURRENT_ACCESS_RESOLUTION=W is supported. In this case LOAD from CURSOR uses the nickname to load the target table.
- No nickname is defined for the source table and the source table has no XML columns. In this case, LOAD from CURSOR is performed by using a cataloged DB2 alias.

1

The target is DB2 for Linux, UNIX, and Windows Version 8 or later and the source does not have a nickname and the source is remote to the target.

2

The target is DB2 for z/OS Version 7 or later.

3

The target is DB2 for Linux, UNIX, and Windows Version 8 or later and either the source nickname is defined or the source is

| | | | |
|----------|---|---|---|
| | local to the target, or the target is DB2 for z/OS Version 7 or later. | | <ul style="list-style-type: none"> The source table is remote from the target database and an XML column is part of the Q subscription. |
| 4 | The target is DB2 for Linux, UNIX, and Windows Version 8 or later. | 3 | The target is not distributed. |
| 5 | The target is DB2 for Linux, UNIX, and Windows Version 8 or later and one of the following conditions is true: <ul style="list-style-type: none"> The target is a nickname. The target is remote and the source table contains large object (LOB) columns that are part of a unidirectional Q subscription. | 4 | The load type is not valid. |
| 6 | The target is Version 8 or later of DB2 for Linux, UNIX, and Windows and the source nickname is defined but does not use the server option <code>CONCURRENT_ACCESS_RESOLUTION=W</code> . This option is required to use LOAD from CURSOR to ensure that the utility waits until all in-progress transactions that modify the source table are completed before beginning the load. Q Apply selected the next best alternative utility for the load. | 5 | The target is a nickname and the load type specified is not valid. For nicknames, the valid load type is EXPORT/IMPORT. However, importing into a nickname is not supported for federated targets when the Q Apply server uses the DB2 pureScale Feature. |
| | User response: This message is for your information only. No action is required. | 6 | The Q subscription is bidirectional or peer-to-peer and the EXPORT/IMPORT method can cause rows to be recaptured at the target. Valid load types are LOAD FROM CURSOR or EXPORT/LOAD. |
| ASN7532E | <i>program_name : program_identifier :</i> The Q Apply program could not select the utility_type load utility for the Q subscription name (receive queue queue_name, replication queue map queue_map_name). Reason code: reason_code. | 7 | The target is a remote database and the source table contains large object (LOB) or XML columns that are part of a Q subscription. The EXPORT/LOAD option does not support loading remote targets with LOB or XML data. The only valid load type is LOAD FROM CURSOR. |
| | Explanation: The following values are valid for the reason code: | 8 | To specify an automatic load, you must use LOAD from CURSOR because the target is remote and the source table contains LOB columns that are part of a bidirectional or peer-to-peer Q subscription. On a remote target, LOB or XML data cannot be loaded using EXPORT/LOAD. For a bidirectional or peer-to-peer Q subscription the IMPORT utility is not a valid option because loaded data might be recaptured at the target. To use LOAD from CURSOR, the target must be DB2 for Linux, UNIX, or Windows Version 8 or later. If the source is remote from the target, you can use LOAD from CURSOR over a source nickname or a cataloged DB2 alias. |
| 0 | The target is distributed but it is not DB2 for Linux, UNIX, or Windows Version 8 or later. | 9 | The target is Version 8 or later of DB2 for Linux, UNIX, and Windows and the source does not use the server option <code>CONCURRENT_ACCESS_RESOLUTION=W</code> . This option is required to use LOAD from CURSOR to ensure that the utility waits until all in-progress transactions that modify the source table are completed before beginning |
| 1 | The target is not DB2 for z/OS Version 7 or later. | | |
| 2 | The target is DB2 for Linux, UNIX, or Windows Version 8 or later and one of the following conditions is true: <ul style="list-style-type: none"> A nickname is not defined for the source table and the source table is not local to the target database. | | |

the load. If you need to use this utility, you must update the federated server options by using the following command:

```
db2 alter server <name>
options (add
CONCURRENT_ACCESS_RESOLUTION 'W')
```

10

Moving data from XML to LOB columns (BLOB, CLOB, DBCLOB) is not supported by the DB2 load utilities (LOAD/IMPORT). No automatic load type could be chosen.

11

LOAD from CURSOR over a cataloged DB2 alias cannot be selected because the Q subscription includes XML columns.

User response: Refer to the detailed message from the utility for the return code and SQL return code. Select a different load type for this Q subscription. Activate the Q subscription. For reason code 8, if you cannot set up LOAD FROM CURSOR, specify a manual load or no load for HAS_LOADPHASE in the IBMQREP_TARGETS control table. For reason code 10, you must specify a manual load or no load.

ASN7533E *program_name : program_identifier :*
Column *column_name* in the target table is not part of the Q subscription name (receive queue *queue_name*, replication queue map *queue_map_name*) and it is not nullable or it has no default values.

Explanation: Any column in the target table that is not part of the Q subscription must be nullable or have a default value.

User response: Modify the attributes of the target or add the column to the Q subscription. Deactivate and then activate the Q subscription.

ASN7534E *program_name : program_identifier :* **The stored procedure is not valid for the following reason *reason_code*.**

Explanation: A stored procedure that is used as a target must define four mandatory parameters plus parameters that can be mapped back to source columns.

The following example shows a valid stored procedure.

```
CREATE TABLE mySource (Parm1 INT NOT NULL,
Parm2 VARCHAR(20) NOT NULL, Parm3 TIMESTAMP,
Parm4 DATE, Parm5 CHAR(2),
PRIMARY KEY(Parm1, Parm2));
CREATE statement for stored procedure:
CREATE PROCEDURE storedprocedure_name(
INOUT operation integer,
IN suppression_ind VARCHAR(size),
IN src_commit_lsn CHAR(10) FOR BIT DATA,
IN src_trans_time TIMESTAMP,
IN XParm1 INT NOT NULL,
```

```
IN Parm1 INT NOT NULL,
IN XParm2 VARCHAR(20)
NOT NULL IN Parm2 VARCHAR(20)
NOT NULL,
IN Parm3 TIMESTAMP,
IN Parm4 DATE,
IN Parm5 CHAR(2) )
```

0

The stored procedure must contain at least five parameters: four required parameters (OPERATION, SUPPRESSION_IND, SRC_COMMIT_LSN, and SRC_TRANS_TIME) and at least one parameter each source column.

1

The first parameter in the stored procedure must be OPERATION.

2

The second parameter in the stored procedure must be SUPPRESSION_IND.

3

The third parameter in the stored procedure must be SRC_COMMIT_LSN.

4

The fourth parameter in the stored procedure must be SRC_TRANS_TIME.

5

The INOUT attribute value of the OPERATION parameter is not equal to INOUT.

6

The INOUT attribute value of the SUPPRESSION_IND, SRC_COMMIT_LSN, or SRC_TRANS_TIME parameter is not equal to IN.

7

The first parameter must have the parameter mode INOUT. All other parameters must have parameter mode IN.

8

No parameters were found that map to the before values of key columns. Before values for key columns are required for key updates. Before values of key columns have the name of the source column prefixed by the letter X. For example, if the key parameter is named Col3, then the before value for that key parameter must be named XCol3.

9

A key column does not have a parameter in the stored procedure that maps to the key

| | | | |
|-------|---|---------------------|--|
| | column's before value. Before values for key columns are required for key updates. Before values of key columns have the name of the source column prefixed by the letter X. For example, if the key parameter is named C013, then the before value for that key parameter must be named XC013. | | Declare the parameters SUPPRESSION_IND, SRC_COMMIT_LSN, and SRC_TRANS_TIME to be IN parameters. |
| 10 | One of the four mandatory parameters was not of the expected data type. Expected data types are: | 7 | Ensure that the first parameter has the parameter mode INOUT and that all other parameters have parameter mode IN. |
| | 1. OPERATION : INTEGER | 8 | Ensure that the stored procedure has a parameter that maps to the before value for each key column. Add any missing parameters to handle before values for all key columns. |
| | 2. SUPPRESSION_IND : VARCHAR(x) | 9 | Ensure that the stored procedure has a parameter that maps to the before value for each key column. Add parameters to handle before values for all key columns. |
| | 3. SRC_COMMIT_LSN : CHAR(10) FOR BIT DATA. If the COMPATIBILITY column in the IBMQREP_CAPPARMS table is 1001 or higher, the data type must be VARCHAR(16) FOR BIT DATA. | 10 | Ensure that the four mandatory parameters have the following data types: |
| | 4. SRC_TRANS_TIME : TIMESTAMP | | 1. OPERATION : INTEGER |
| | User response: Modify the stored procedure by performing the action that corresponds to the reason code. Register the stored procedure again and restart the Q subscription. | | 2. SUPPRESSION_IND : VARCHAR(x) |
| 0 | Add the following mandatory parameters as the first parameters: OPERATION, SUPPRESSION_IND, SRC_COMMIT_LSN, SRC_TRANS_TIME). Next add one parameter that matches the data type for each primary key column and one identical parameter for the before image of updated key columns. Then add one parameter for each non-key column with data types that match each source column. | | 3. SRC_COMMIT_LSN : CHAR(10) FOR BIT DATA or VARCHAR(16) FOR BIT DATA |
| 1 | The first parameter in the stored procedure must be OPERATION. | | 4. SRC_TRANS_TIME : TIMESTAMP |
| 2 | The second parameter in the stored procedure must be SUPPRESSION_IND. | | |
| 3 | The third parameter in the stored procedure must be SRC_COMMIT_LSN. | | |
| 4 | The fourth parameter in the stored procedure must be SRC_TRANS_TIME. | | |
| 5 | Declare the OPERATION parameter to be INOUT. | | |
| 6 | | | |
| <hr/> | | | |
| | | ASN7535E | <i>program_name</i> : <i>program_identifier</i> : In multidirectional replication, Q subscription name (receive queue queue_name, replication queue map queue_map_name) is not valid in the IBMQREP_TARGETS table. Reason code: reason_code. |
| | | Explanation: | The following values are valid for reason codes: |
| | | 0 | The Q subscription does not exist in the IBMQREP_TARGETS table. |
| | | 1 | The SUBTYPE value is set to 'P' (peer-to-peer), but either the CONFLICT_RULE value is not set to 'V' (check version) or the CONFLICT_ACTION value is not set to 'F' (force the change). |
| | | 2 | There is no Q subscription in the IBMQREP_SUBS table for some members of the Q subscription group. |
| | | 3 | The SUBGROUP column is NULL. |
| | | 4 | In the IBMQREP_TARGETS table, the SOURCE_NODE, the TARGET_NODE, or both, do not match the values in the IBMQREP_SUBS table. |
| | | 5 | The Q subscription definition in the |

IBMQREP_SUBS table does not match the Q subscription definition in the IBMQREP_TARGETS table for a particular SUBGROUP column.

- 6 The IBMQREP_SUBS table and the IBMQREP_TARGETS table have the same number of Q subscriptions, but the SUBGROUP values for those Q subscriptions do not match.
- 7 There is no Q subscription in the IBMQREP_TARGETS table for some members of the Q subscription group.
- 8 One or more of the following conditions exist:
 - The target table does not exist.
 - The version columns do not exist: "ibmqrepVERTIME", "ibmqrepVERNODE".
 - The version columns exist but they have the wrong data types or defaults.
- 9 In the IBMQREP_TARGETS table, the CONFLICT_RULE value 'V' (check version) is not allowed when SUBTYPE is 'B' (bidirectional). This conflict rule is valid only for SUBTYPE 'P' (peer-to-peer). For Q subscriptions in bidirectional replication, the valid conflict rule values are 'K', 'C', or 'A'.
- 10 There is more than one Q subscription defined in the IBMQREP_SUBS table for a given SUBGROUP and TARGET_NAME. In bidirectional replication (SUBTYPE='B'), there must be only one Q subscription in the IBMQREP_SUBS table and one in the IBMQREP_TARGETS table for a given SUBGROUP.

User response: Review the reason codes in the explanation, and take the appropriate action:

- For reason code 8: Verify that the target table exists and has the proper columns. If it does not have the proper columns, use one of the replication administration tools to create the table with the proper columns or manually create the proper columns. If the target table does not exist, redefine the Q subscription to and from the node using one of the replication administration tools.
- For all other reason codes: Redefine the Q subscriptions to and from this node using one of the replication administration tools.

ASN7536E *program_name : program_identifier : Q subscription name (send queue queue_name, replication queue map queue_map_name) is not properly defined in the IBMQREP_SUBS table. The error was detected while trying to add Q subscription name (receive queue queue_name, replication queue map queue_map_name) to the SUBGROUP.*

Reason code: *reason_code.*

Explanation: The error was detected while trying to add a new node to the Q subscription group. The Q subscription going from the new node to an active member is not defined correctly in the IBMQREP_SUBS table. The following values are valid for the reason code:

- 0 The Q subscriptions in the Q subscription group do not have the same SUBTYPE value. All the Q subscriptions for the subscription group must have the same SUBTYPE value. The SUBTYPE value must be either 'P' (peer-to-peer) or 'B' (bidirectional).
- 1 The SOURCE_NODE in the IBMQREP_SUBS table is not the same as the TARGET_NODE in the IBMQREP_TARGETS table.
- 2 The GROUP_MEMBERS column is not NULL in the IBMQREP_SUBS table.
- 3 The STATE column must be 'I' (inactive) in the IBMQREP_SUBS table.
- 4 More than one of the Q subscriptions has the same value for the TARGET_NODE column in the IBMQREP_SUBS table.

User response: Redefine the Q subscription using one of the replication administration tools.

ASN7537E *program_name : program_identifier : Q subscription name (receive queue queue_name, replication queue map queue_map_name) is not valid in the IBMQREP_TARGETS table for multidirectional replication. Reason code: reason_code.*

Explanation: The valid values for the reason code are:

- 0 The Q subscriptions do not have the same SUBTYPE. All of the Q subscriptions for the group must have the same SUBTYPE. The SUBTYPE must be either 'P' (peer-to-peer) or 'B' (bidirectional).
- 1 The TARGET_NODE in the IBMQREP_TARGETS table is not the node of the new member Q subscription.
- 2 The state for this Q subscription is not 'I' (inactive). The STATE column should be 'I' (inactive) in the IBMQREP_TARGETS table of a Q subscription that is a member.
- 4 More than one of the Q subscriptions in the SUBGROUP has the same value for the SOURCE_NODE in the IBMQREP_TARGETS table.

User response: Redefine the Q subscriptions to and from this node using one of the replication administration tools.

ASN7538E *program_name : program_identifier : For multidirectional replication, there is either no row in the IBMQREP_TARGETS table for the Q subscription name (receive queue queue_name, replication queue map queue_map_name), or there is no match for this Q subscription in the IBMQREP_SUBS table at this server.*

Explanation: The Q subscription is not correctly defined for peer-to-peer or bidirectional replication.

User response: Redefine the Q subscriptions to and from the table at this server using one of the replication administration tools.

ASN7539E *program_name : program_identifier : During deactivation of all the Q subscriptions for the same table, the Q subscription name (receive queue queue_name, replication queue map queue_map_name) was not in the inactive state and it does not have a corresponding Q subscription entry in the IBMQREP_SUBS table at this server.*

Explanation: Some Q subscriptions cannot be found. Some Q subscriptions might have been deleted, or were not created to begin with. All Q subscriptions going to and from this node cannot be deactivated.

User response: The Q subscriptions were not properly defined because they are not fully connecting all physical tables from this logical table. Therefore the deactivation protocol cannot automatically deactivate all of the Q subscriptions to and from this node.

To make sure that changes are not replicating to and from this table, take the following steps:

1. At this server, insert a CAPSTOP signal for the Q subscriptions that are in the IBMQREP_SUBS table for this logical table. Do this step only for those Q subscriptions that are not already inactive (for example, they might be in the active state or loading state). Insert the signal for the Q subscriptions one at a time until all of the Q subscriptions are deactivated.
2. Repeat the previous step for all of the Q subscriptions that are at all other servers and are replicating changes to the first server.

ASN7540I *program_name : program_identifier : The RI constraint constraint_name on target target_name was dropped for Q subscription name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The Q subscription is specified with a loadphase, either internal or external. During the load, the referential constraints are dropped from the target

table and saved in the IBMQREP_SAVERI table.

User response: This message is for your information only. No action is required.

ASN7541I *program_name : program_identifier : The RI constraint constraint_name was added for the target target_name, for Q subscription name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The Q Apply program has finished loading the Q subscription, and has added the referential constraints back to the target. The referential constraints were removed from the IBMQREP_SAVERI table.

User response: This message is for your information only. No action is required.

ASN7542E *program_name : program_identifier : There was no Q subscription found in the IBMQREP_SUBS table that matches Q subscription name (receive queue queue_name, replication queue map queue_map_name) during the initialization of this Q subscription.*

Explanation: The Q subscription definition in the IBMQREP_SUBS table was dropped during the initialization of a new member.

User response: Redefine the Q subscriptions to and from the node using one of the replication administration tools.

ASN7543W *program_name : program_identifier : The REINITQ command was issued to change the memory limit for the receive queue queue_name, replication queue map queue_map_name. However, the memory limit memory_limit MB was not changed.*

Explanation: A REINITQ command was issued to change the memory limit for the receive queue, but the memory limit value specified was the same as the existing value.

User response: If you want to change the memory limit from the current value, reissue the REINITQ command for the receive queue, specifying a new memory limit in MB.

ASN7544W *program_name : program_identifier : The REINITQ command is issued to change the number of agents for the receive queue queue_name, replication queue map queue_map_name. However, the number of agents number was not changed.*

Explanation: A REINITQ command was issued to

change the number of apply agents for the receive queue, but the number of agents specified was the same as the existing number.

User response: If you want to change the number of agents for this receive queue, reissue a REINITQ command for the receive queue that specifies a new number of agents.

ASN7545W *program_name : program_identifier : The REINITQ command for receive queue queue_name, replication queue map queue_map_name was not processed as it would reduce the number of agents to zero.*

Explanation: REINITQ command was issued to change the number of agents; however, after the change there would be 0 agents left. Because of this fact, the REINITQ command was not processed.

User response: Verify that the correct number of agents was specified for this receive queue.

ASN7546W *program_name : program_identifier : The REINITQ command cannot be processed because too many agent threads were specified for the replication queue map. The maximum number of agent threads that are allowed is number. The original number of agents, number, stays the same.*

Explanation: The Q Apply program allows a maximum number of agent threads for each replication queue map. If the number is exceeded, the REINITQ command to refresh the attributes of the queue map fails.

User response: Reduce the number of agent threads that you specify for the queue map and then reissue the command.

ASN7547I *program_name : program_identifier : The REINITQ command was processed successfully. The number of agents decreased from number to number for receive queue queue_name, replication queue map queue_map_name.*

Explanation: The REINITQ command that you issued to decrease the number of agents for the receive queue was processed successfully.

User response: This message is for your information only. No action is required.

ASN7548I *program_name : program_identifier : The REINITQ command was processed successfully. The number of agents increased from number to number for receive queue queue_name, replication*

queue map queue_map_name.

Explanation: The REINITQ command that was issued to increase the number of agents was processed successfully.

User response: This message is for your information only. No action is required.

ASN7549I *program_name : program_identifier : The REINITQ command was processed successfully. The memory limit was set from memory_limit MB to memory_limit MB for the receive queue queue_name, replication queue map queue_map_name.*

Explanation: The REINITQ command was issued successfully to increase the memory limit for this receive queue.

User response: This message is for your information only. No action is required.

ASN7550E *program_name : program_identifier : The Q Apply program was expecting a message of type msgA but found message type msgB as the last message of the transaction (receive queue queue_name, replication queue map queue_map_name).*

Explanation: An internal inconsistency was detected.

User response: Contact IBM Software Support. Provide an Analyzer report.

ASN7551E *program_name : program_identifier : The Q Apply program detected a gap in message numbers on receive queue queue_name (replication queue map queue_map_name). It read message ID message_ID, but expected to find message ID message_ID. The Q Apply program cannot process any messages until it finds the expected message.*

Explanation: To make it easier for Q Apply programs or subscribing applications to detect missing or out-of-sequence messages, Q Replication and event publishing programs use a dense numbering system. A dense numbering system is one where each message is assigned a positive integer with no gaps between numbers. The Q Capture program assigns a unique identifier to each message that it puts on the send queue. These identifiers are necessary because the Q Apply program processes messages in a strict order. When a message does not arrive in the receive queue in its expected order, the Q Apply program continues to look for the message. The Q Apply program will apply all changes that belong to messages that it already received.

User response: Look for the message with the expected message ID on all the dead letter queues of all

the WebSphere MQ queue managers that are used to transmit messages between the Q Capture and Q Apply programs. If you recover the message, put it on the receive queue, preserving the WebSphere MQ message header information (especially the message ID). If the message cannot be recovered, follow these steps:

1. Use the stopq command to stop the Q Apply program from reading from the receive queue.
2. Deactivate all of the Q subscriptions for this replication queue map.
3. Empty the send queue and the receive queue.
4. Use the startq command so that the Q Apply program resumes reading from the receive queue.
5. Activate all the Q subscriptions for this replication queue map.

ASN7552W *program_name : program_identifier : The Q Apply program is polling receive queue queue_name, replication queue map queue_map_name for message ID message_ID.*

Explanation: The Q Apply program detected a gap in the message ID and cannot proceed until that message ID is found. This message will continue to be issued periodically until the message with the expected message ID is put on the send or receive queue.

User response: Look for the message with the expected message ID on all the Dead Letter Queues of all the WebSphere MQ queue managers that are used to transmit messages between the Q Capture and Q Apply programs. If you recover the message, put it on the receive queue, preserving the WebSphere MQ message header information (especially the message ID). If the message cannot be recovered, follow these steps:

1. Use the stopq command to stop the Q Apply program from reading from the receive queue.
2. Deactivate all the Q subscriptions for this replication queue map.
3. Empty the send queue and the receive queue.
4. Use the startq command so that the Q Apply program resumes reading from the receive queue.
5. Activate all the Q subscriptions for this replication queue map.

ASN7553E *program_name : program_identifier : The Q Apply program stopped reading from receive queue queue_name, replication queue map queue_map_name, because it found a message that is older (timestamp message_seq) than the last message that it read (timestamp message_seq).*

Explanation: The message ID contains a timestamp (in integer format). The Q Apply program cannot continue reading from the queue if it encounters a message that is older than the one it read. This error likely occurred

for one of the following reasons:

- There are two Q Capture programs writing messages to the same receive queue. This setup is not supported. The Q Apply program is expecting messages from one Q Capture program on any given receive queue.
- The system clock where the Q Capture program runs was set back in time and either the Q Capture program was cold started or it was warm started but there were no active Q subscriptions for this replication map queue.

User response: If more than one Q Capture program is writing to the same receive queue, follow these steps:

1. Deactivate all the Q subscriptions that feed the receive queue.
2. Redefine the set up so that the messages from each Q Capture program go to separate receive queues.
3. Delete all message from all the send queues and the receive queue that were involved in the invalid setup.
4. Issue the startq command so that the Q Apply program starts reading from the receive queue again.
5. Activate all the Q subscriptions.

If the system clock was set back in time at the source, follow these steps:

1. Stop the Q Capture program.
2. Either reset the clock to the original time (or later), or wait until enough time passes so that the current time reaches the original time.
3. Deactivate all the Q subscriptions that feed the receive queue.
4. Delete all message from all the send queues and the receive queue that were involved in the invalid setup.
5. Issue the startq command so that the Q Apply program starts reading from the receive queue again.
6. Activate all the Q subscriptions.

ASN7554W *program_name : program_identifier : The RI constraint constraint_name could not be added for the target target_name for Q subscription name (receive queue queue_name, replication queue map queue_map_name) because a dependent Q subscription name is in inactive (I) state. The constraint was saved in the IBMQREP_SAVERI table. The program will continue to apply changes for the Q subscription without the RI constraint.*

Explanation: See message text.

User response: Determine why the dependent Q subscription is in the inactive state. If it has not been

activated, activate it. If it is in inactive state due to an error, fix the error by looking at the Q Apply diagnostic log file. After the dependent Q subscription is active, the Q Apply program will add the RI constraints that were saved in the IBMQREP_SAVERI table.

ASN7555W *program_name : program_identifier : The RI constraint constraint_name could not be added for the target table_name for Q subscription name (receive queue queue_name, replication queue map queue_map_name) because it has an RI constraint with a table that is unknown to the Q Apply program.*

Explanation: The Q Apply program encountered SQLCODE 667 when trying to add RI constraints for the Q subscription. The Q Apply program might not know about the dependent table because there is no Q subscription defined on it, or because a Q subscription is defined on it but a CAPSTART signal was never issued. After the dependent Q subscription is active, the Q Apply program will add the RI constraints that were saved in the IBMQREP_SAVERI table.

User response: Find the dependent table and define a Q subscription on it. If a Q subscription is defined, issue a CAPSTART signal on it.

ASN7557W *program_name : program_identifier : The memory limit for receive queue queue_name for replication queue map queue_map_name is too low. It is memory_limit MB, but it should be equal to or greater than memory_limit MB.*

Explanation: See message text.

User response: Increase the MEMORY_LIMIT value in the IBMQREP_RECVQUEUES table for the specified replication queue map.

ASN7558E *program_name : program_identifier : The spill queue queue_name is full. Q subscription name. Replication queue map queue_map_name.*

Explanation: The number of messages on the spill queue has reached the number set by the MAXDEPTH attribute for that spill queue. The spill queue is full and the Q Apply program cannot write to it. Perhaps the Q Apply program is still loading the target table, but the corresponding source table is being heavily updated.

User response: If possible, stop the Q Capture program. If necessary, increase the value for the MAXDEPTH attribute of the spill queue.

ASN7559W *program_name : program_identifier : The memory limit for receive queue queue_name for replication queue map queue_map_name is too high. It is memory_limit MB, but it should be less than memory_limit MB.*

Explanation: See message text.

User response: Lower the MEMORY_LIMIT value in the IBMQREP_RECVQUEUES table for the specified replication queue map.

ASN7583W *program_name : program_identifier : An error occurred while retrieving LOB data. The Q Apply program could not find a LOB message having LOBId lobid.*

Explanation: See message text.

User response: Drain all messages from the queue and activate the Q subscription.

ASN7584E *program_name : program_identifier : The Q subscription name (receive queue queue_name, replication queue map queue_map_name) is defined as type stored procedure. The stored procedure proc_name could not be found.*

Explanation: The Q subscription cannot be activated as the stored procedure name as defined in the IBMQREP_TARGETS could not be found or the stored procedure is not registered.

User response: Check for the stored procedures schema and name in the TARGET_OWNER,TARGET_NAME column of the IBMQREP_TARGETS table. Register the stored procedure using the "CREATE PROCEDURE" statement.

ASN7586E *program_name : program_identifier : The memory limit for receive queue queue_name for replication queue map queue_map_name is too low memory_limit MB. It should be at least 3 times the incoming message size message_size.*

Explanation: The memory_limit is the amount of memory used by the Q Apply program to store transactions to be applied. It needs to be big enough to contain at least 1 incoming message.

User response: Check the maximum message size of the incoming send queue in the IBMQREP_SENDQUEUES table. The memory limit needs to be able to contain many messages else performance would not be acceptable. So you need to increase the memory_limit value in the IBMQREP_RECVQUEUES table and issue the REINITQ command or stop and start the Q Apply process.

ASN7587E *program_name : program_identifier : The memory limit for receive queue queue_name for replication queue map queue_map_name is too high memory_limit MB. It should be less than memory_limit MB.*

Explanation: The memory limit for the receive queue is too high.

User response: Lower the MEMORY_LIMIT value in the IBMQREP_RECVQUEUES table for the specified replication queue map.

ASN7588E *program_name : program_identifier : The startq command for receive queue queue_name, replication queue map queue_map_name cannot be completed. Reason code: reason_code*

Explanation: The following values are valid for the reason code:

- 0 The queue is already being processed.
- 1 There is no information about this queue in the IBMQREP_TARGETS table.
- 2 The program is still in the initialization phase.
- 3 The value in the MAXAGENTS_CORRELID column in the IBMQREP_RECVQUEUES table is incorrect.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 0 No action required.
- 1 Check the IBMQREP_RECVQUEUES table for the information of the queue. Redefine the Q Subscription using one of the replication administration tool.
- 2 Reissue the command later.
- 3 Change the value of MAXAGENTS_CORRELID column in the IBMQREP_RECVQUEUES table and reissue the command. This value should be less than the NUM_APPLY_AGENTS value.

ASN7589E *program_name : program_identifier : The receive queue queue_name for replication queue map queue_map_name cannot be found in table IBMQREP_RECVQUEUES.*

Explanation: The Q Apply browser while trying to update this receive queue state, could not find an entry corresponding to the receive queue name, in the IBMQREP_RECVQUEUES table.

User response: Check the RECVQ column in the IBMQREP_RECVQUEUES table. Redefine the Q subscription using one of the replication administration

tool, and activate the Q subscription again.

ASN7590I *program_name : program_identifier : The Q Apply program stopped reading from the queue queue_name for replication queue map queue_map_name. Reason code: reason_code.*

Explanation: The following values are valid for the reason code:

- 0 The Q Apply program is running with the AUTOSTOP parameter option and it encountered an empty receive queue.
- 1 The Q Apply browser received the STOPQ command.

User response: This message is for your information only. No action is required.

ASN7591I *program_name : program_identifier : The Q Apply program started an apply agent agent_name for receive queue queue_name, replication queue map queue_map_name.*

Explanation: The Q Apply agents started for this receive queue.

User response: This message is for your information only. No action is required.

ASN7592E *program_name : program_identifier : The Q Apply program (Q subscription name, queue queue_name, replication queue map queue_map_name) encountered an error while converting a code page of the field field_name of the message message_type. The original value is field_value.*

Explanation: A character field inside an internal message that was sent by the Q Capture program and processed by the Q Apply program cannot be converted to Q Apply program code page (the code page code is called CCSID). It is highly recommended that you run the source database, the Q Capture program, the Q Apply program, and the target database in the same code page to avoid unnecessary conversions. In this case the ICU library used to convert one code page to another is not able to convert the Q Capture program code page to the Q Apply program code page. For DB2, you can set the application code page with the command db2set DB2CODEPAGE=[CCSID].

User response: Issue the command to set the Q Apply program code page to the Q Capture program code page, or set it to a code page that can be converted by the ICU library.

ASN7593E *program_name : program_identifier : An unsupported data type exists in column column_name for Q subscription name (receive queue queue_name, replication queue map queue_map_name) and it cannot be replicated. The data type code is data_type.*

Explanation: Some data types are not supported in Q Replication and Event Publishing and some data types can be used only under certain circumstances.

User response: Check the data type of the column, and ensure that DB2 supports the data type on the target. Also check the Q Replication documentation for information about data types that are not supported.

ASN7594W *program_name : program_identifier : The Q Apply program removed a message from the receive queue queue_name, replication queue map queue_map_name because the message format was incorrect. Reason code: reason_code.*

Explanation: The following values are valid for the reason code:

- 0 The message that was sent by the Q Capture program is an XML message intended for event publishing. It is not supported by the Q Apply program.
- 1 The Q Apply program could not identify that the message is sent from the Q Capture program.

User response: Review the reason codes in the explanation, and take the appropriate action:

- 0 Change the message format to 'C' in the IBMQREP_SENDQUEUES table for the send queue for this replication queue map.
- 1 If a program other than the Q Capture program is putting messages on the specified queue, stop that program. The queue must be used exclusively by one Q Capture program. If no other program is using the queue, contact IBM Software Support. Provide an Analyzer report.

ASN7595W *program_name : program_identifier : The Q Apply program received a Q subscription deactivate message, but the SUB_ID SUB_ID (receive queue queue_name, replication queue map queue_map_name) cannot be found or the Q subscription is in the inactive state.*

Explanation: The Q Apply program is not replicating or loading the Q subscription corresponding to the incoming message.

User response: If the Q subscription is already inactive (state 'I' in the IBMQREP_TARGETS table) nothing needs to be done. If the Q subscription SUB_ID cannot be found in the IBMQREP_TARGETS table, the Q subscription is defined only for the Q Capture program. Define the Q subscription so that the target information is part of the Q subscription.

ASN7597E *program_name : program_identifier : The Q subscription name (receive queue queue_name, replication queue map queue_map_name) is about to be disabled because of the conflict action or error action. Future row changes to the source table that is specified for this Q subscription will be skipped until the Q subscription is restarted.*

Explanation: The Q Apply program encountered an error or conflict while applying a change from a transaction, and the Q subscription specifies a conflict action, error action, or both to deactivate the Q subscription when conflicts or errors occur. The row that caused the conflict or error is saved to the IBMQREP_EXCEPTIONS table. All further changes from the source table are ignored until the Q subscription is restarted, which will trigger a new load of the target table.

User response: Check the IBMQREP_EXCEPTIONS table for details about the conflict or error. Fix the problem, and then start the Q subscription again. If no row exists in the IBMQREP_EXCEPTIONS table, then either the row was deleted or the error occurred during the initial load process. During the load process, no further information is stored in the IBMQREP_EXCEPTIONS table. Instead, the information is stored in the Q Apply log file (either the present log file or a previous log file).

ASN7598E *program_name : program_identifier : The transaction message for the Q subscription name (receive queue queue_name, replication queue map queue_map_name) does not contain before values of the column.*

Explanation: The Q subscription probably specifies N in the BEFORE_VALUES column of the IBMQREP_SUBS table but the CONFLICT_RULE column in the IBMQREP_TARGETS table specifies C or A. Such a configuration is not valid. This configuration might have resulted if the IBMQREP_SUBS table was modified without first stopping the Q subscription.

User response: Ensure that the Q subscription is defined correctly by following these steps:

1. Stop the Q subscription.
2. Use the replication administration tools to correctly define the Q subscription. If the conflict rule for the Q subscription is C (check changed columns) or A

(check all columns), then the BEFORE_VALUES column in the IBMQREP_SUBS table must be set to Y.

3. Start the Q subscription

ASN7605I *program_name : program_identifier : Q subscription name (receive queue queue_name, replication queue map queue_map_name) is active, but the RI constraints were not added back to the target target_name because there is a dependent Q subscription in loadphase.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7606I *program_name : program_identifier : Q subscription name (receive queue queue_name, replication queue map queue_map_name) is active.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7607I *program_name : program_identifier : Q subscription name (receive queue queue_name, replication queue map queue_map_name) finished loading the target. The changes will be applied from the spill queue queue_name.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7608I *program_name : program_identifier : The program finished loading Q subscription name (receive queue queue_name, replication queue map queue_map_name), and notified the Q Capture program.*

Explanation: See message text.

User response: This message is for your information only. No action is required.

ASN7609W *program_name : program_identifier : Due to codepage differences between the source and target, the converted data for column name of Q subscription name for receive queue queue_name is longer than the target column length (converted data length is number_of_bytes, source data length is number_of_bytes).*

Explanation: The source and target tables do not use

the same codepage. While the Q Apply program was performing the codepage conversion, the data of one column expanded. Because the expanded data does not fit in the target column, a database error will occur, which will trigger the Q Apply error action for this Q subscription. The Q Apply error action will be logged in the Q Apply log file and the IBMQREP_APPLYTRACE table.

User response: Alter the target table column to increase the column's length. For example, if the column is declared as VARCHAR(10), you can extend it to be VARCHAR(30). Look in Q Apply log file or the IBMQREP_APPLYTRACE table for the message that will be issued after the Q Apply error action is triggered.

ASN7610E *program_name : program_identifier : The Q Apply program received a message to add a column to the Q subscription, but the SUB_ID subscription_identifier (receive queue queue_name, replication queue map queue_map_name) is not active.*

Explanation: An AddCol signal was issued to the Q Capture program, which resulted in a message being sent to the Q Apply program. However, the Q subscription is not known by the Q Apply program either because the Q subscription was never started, it was stopped because of an error, or it was dropped.

User response: Look in the Q Apply log file or the IBMQREP_APPLYTRACE table for related messages. For example, if you see the message ASN7512E, then the Q subscription was never started, and if you see the message ASN7597E, then the Q subscription was stopped. If the Q subscription exists, restart it. If the Q subscription does not exist, create and start it.

ASN7611I *program_name : program_identifier : The column name that was added for the Q subscription name (receive queue queue_name, replication queue map queue_map_name) already exists in the target table schema_name.table_name.*

Explanation: The Q Apply program tried to alter the target table to add a new column, but the new column already existed. The Q Apply program will add this column to the Q subscription and start replicating changes.

User response: This message is for your information. No action is required.

ASN7612I *program_name : program_identifier : Column column_name, has been added to Q subscription name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The Q Apply program processed the

AddCol signal successfully. Changes for the newly added column are now being replicated to the target table.

User response: This message is for your information. No action is required.

ASN7613I *program_name : program_identifier : Start processing queue (receive queue queue_name, replication queue map queue_map_name), application single byte codepage CCSID, double byte codepage CCSID, source codepage CCSID, endian conversion required needed_or_not, float conversion encoding.*

Explanation: This message contains the codepage information for the replication queue map.

User response: Use this information to know which codepages (CCSID) are used by the Q Apply program and possibly to debug issues that are related to codepage conversion.

ASN7614W *program_name : program_identifier : The local clock is at least number_of_seconds seconds behind the clock of the system that is sending messages on receive queue queue_name for replication queue map queue_map_name. This time difference causes delays in processing rows that are in peer-to-peer configurations.*

Explanation: The Q Apply program detected that the timestamp of a replicated row in a peer-to-peer configuration is ahead of the local time. To avoid inconsistencies the apply agent that is processing this change will not apply the change until the local time is equal to or more than the change's own time.

User response: Adjust the time of the local clock. The local operating system and software packages offer commands to adjust the local time. The local time on the system that is running the Q Apply program should be moved ahead to match the time of the machine where the Q Capture program is running. It is not advisable to move a clock back in time.

ASN7615W *program_name : program_identifier : The program is unable to prune the IBMQREP_DELTOMB table because not all queues in the peer-to-peer configuration are started, or the queues are inactive and are not sending heartbeat messages.*

Explanation: The Q Apply program periodically deletes rows from the IBMQREP_DELTOMB table if there are any active Q subscriptions in a peer-to-peer configuration. A row is not deleted until the Q Apply program sees a transaction message or a heartbeat

message that has a later timestamp than the current row. A queue is either not being processed, no messages are going through it, or no messages are being put on it by the Q Capture program.

User response: Restart an unprocessed queue using the command:

```
asnqacmd APPLY_SERVER=
database_or_subsystem_name
startq=receive_queue_name
```

Also, make sure that all of the send queues that send messages to this Q Apply program have a heartbeat value that is a maximum of 3600 seconds (1 hour).

ASN7616E *program_name : program_identifier : The column name with a LOB data type in the add column message for Q subscription name (receive queue queue_name, replication queue map queue_map_name) does not exist in the target table schema_name. The table is table_name.*

Explanation: On DB2 for z/OS, any LOB column (BLOB, CLOB, DBCLOB) that is added to a Q subscription must be defined on the target table before you issue the ADDCOL signal. The column that was in the ADDCOL signal does not exist at the target.

User response: Create the LOB column on the target table and reissue the ADDCOL signal.

ASN7617E *program_name : program_identifier : The target table definition of schema_name. The table table_name is incomplete for the column name with a LOB data type that is in the add column message to be added for Q subscription name (receive queue queue_name, replication queue map queue_map_name). Table status is status.*

Explanation: An AddCol signal was issued to Q Capture program to add a LOB column (BLOB, CLOB, DBCLOB) to the Q subscription. The column was found, but its definition is not complete for one of the following reasons:

- L** An auxiliary table or auxiliary index has not been defined for a LOB column.
- P** The table lacks a primary index.
- R** The table lacks an index that is required on a row ID.
- U** The table lacks an index that is required on a unique key.

User response: Check the table status and complete the definition to add the LOB column to the target table by adding the corresponding index. Then reissue the AddCol signal.

ASN7618I *program_name : program_identifier :*
Transaction LSN was rolled back (error message SQL0911) with reason code code. Apply agent agent_id was applying a operation for Q subscription name (receive queue queue_name, replication queue map queue_map_name). This transaction will be retried until the deadlock retry limit is reached.

Explanation: This informational message (that corresponds to SQL0911) indicates that a transaction is retried because of a lock timeout or deadlock. Reason code:

2

Transaction got rolled back

68

Lock timeout

The transaction log sequence number (LSN) can be used to identify whether the same transaction gets retried multiple times. A transaction is retried multiple times until the number in DEADLOCK_RETRIES from the IBMQREP_APPLYPARMS table is reached. If this limit is reached, the Q Apply program stops processing changes for this replication queue map.

User response: For peer-to-peer replication only, run the runstats program on the target tables and the Q Apply internal tables: IBMQREP_DONEMSG, IBMQREP_SPILLEDROWS, IBMQREP_DELTOMB.

In a peer-to-peer environment where many conflicts are expected, many deadlocks should be expected as well. It is also advised to increase the DEADLOCK_RETRIES limit to a value that is greater than the default. On z/OS, the target table and the Q Apply internal tables (IBMQREP_DONEMSG, IBMQREP_SPILLEDROWS, IBMQREP_DELTOMB) should use row level locking. The database or subsystem generates deadlock event files that you can use to find the exact participant of a deadlock.

It is also possible to lower the number of apply agents for this replication queue map to have less concurrency and avoid chances of deadlocks.

ASN7619W *program_name : program_identifier :* **The Q Apply program approximated a floating point conversion. The IEEE floating point number from the source is value. The z/OS floating point number that was applied to the target table is value.**

Explanation: Changes are replicated between two float columns. The source is a non-z/OS server and the target is a z/OS server. A non-z/OS server uses the IEEE floating point number representation standard, while a z/OS server has its own format. Conversion is possible between the two formats. However, every

value cannot be represented, and some values must be approximated. The replication configuration is not altered if this warning is issued, and the Q Apply error action will not be performed. Because of this approximation, the target table will not necessarily match the source table exactly.

Float columns should be avoided for primary key columns when you are replicating between a z/OS and a non-z/OS server to avoid breaking the one-to-one mapping between the source and target rows if there is a chance that floating point values get approximated.

User response: This message is for your information. No user response is required.

ASN7620W *program_name : program_identifier :* **The program could not gather monitoring information for one or more Q Apply agents or spill agents for a receive queue. Any data that the program did gather in this monitor interval was stored to be reported after the next monitor interval. Number of agents for which the program could not gather monitoring information: number. Receive queue name: queue_name.**

Explanation: The Q Apply program's monitoring thread was not able to gather monitoring information for at least one Q Apply agent or spill agent. The monitoring data for the monitoring interval is not reliable, so reporting will be skipped for this monitoring interval. However, the data will be reported in the next monitoring interval, so it will not be lost. If a similar message continues to be reported, then the Q Apply agent or spill agent stopped or unexpected errors were detected in structures that are used in memory to collect performance statistics.

User response: Respond to the message in one of the following ways:

- If this message appears only once, no action is required.
- If multiple monitoring cycles are skipped, perform the following troubleshooting steps:
 1. Look for additional error messages that might identify the problem, and try to correct the problem.
 2. If you cannot identify or correct the problem, stop and start the Q Apply program.

Determine the best time to start and stop the Q Apply program so that you minimize the effect on data replication, which should be continuing because replication is not affected by this problem.

ASN7621E *program_name : program_identifier : Source column name does not map to target column name because of reason code code for Q subscription name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: While activating the Q subscription, the Q Apply program detected that the source column is not compatible with the target column. As a result, the Q subscription will not be started.

- 1 Data type mismatch between source and target.
- 2 Length mismatch between source and target. The target column must be equal to or longer than the source for character data types. If you are replicating from GRAPHIC to CHAR (or from VARGRAPHIC to VARCHAR), then the target column must be at least twice as long as the source column (because GRAPHIC and VARGRAPHIC data types are double byte).
- 3 Replicating from NULL to NOT NULL is not allowed.
- 4 Replicating from LONG VARCHAR to VARCHAR requires the VARCHAR column to have a length of the maximum size.
- 5 The Q subscription is in a bidirectional or peer-to-peer configuration and either the data type or column length does not match between source and target.

User response: Modify the target or source columns (or both) so that they are compatible. Recreate and restart the Q subscription.

ASN7622W *program_name : program_identifier : Source column name mapped to target column name is longer than the target. Some truncation might occur. Q subscription name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The maximum length of a VARCHAR column is less than the length of a LONG VARCHAR column. Some incoming data might be truncated while it is replicated.

User response: If data truncation is not acceptable, create the target column as LONG VARCHAR. Then redefine and restart the Q subscription. If truncation is acceptable, no action is required.

ASN7623E *program_name : program_identifier : The Q Apply program cannot process the AddCol signal for Q subscription name (receive queue queue_name, replication queue map queue_map_name). The target is a nickname that does not support the alter statement, and the column that is being added name is not part of the nickname definition.*

Explanation: An AddCol signal was issued to the Q Capture program for a Q subscription that has a nickname as the target. Because the Q Apply program cannot alter the nickname directly, the AddCol signal cannot be processed.

User response: Redefine the nickname to include the new column and reissue the AddCol signal.

ASN7624I *program_name : program_identifier : Found number secondary unique index(es) for Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The Q Apply program loads unique index information for all target tables so that it can correctly order transactions and apply them in parallel. This message specifies how many secondary unique indexes were found. The number does not include the unique index that is used to be the replication key.

User response: This message is for your information. No action is required.

ASN7625E *program_name : program_identifier : The target type for Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name) is a nickname but the Q Capture sending options are not set up correctly.*

Explanation: When the target type of a Q subscription is a nickname and this nickname has more than one unique index, the Q Apply program requires the Q Capture program to send the before values of columns that have changed, and the values of columns that have not changed. These sending options allow Q Apply to apply transactions in the correct order without violating unique constraints on the target. These options are not required when the target type is a user table because DB2 provides the necessary index identifier that is used to order transactions. However, this identifying information is not always available from non-DB2 databases.

User response: At the source system change the BEFORE_VALUES attribute for this Q subscription to be Y and CHANGED_COLS_ONLY to be N in the IBMQREP_SUBS table.

ASN7626E *program_name : program_identifier : The target table referenced by nickname owner.nickname should be empty before the IMPORT utility begins. The Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name) will be deactivated.*

Explanation: When you use the IMPORT utility to load data into federated targets, the target table should be empty. The IMPORT utility does not support the REPLACE option.

User response: Delete the contents of the target table and activate the Q subscription.

ASN7627E *program_name : program_identifier : The Q Apply browser browser_name (receive queue queue_name, replication queue map queue_map_name) terminated with error code error_code.*

Explanation: A previous error or condition caused the Q Apply browser to terminate. The reason code could be an SQLCODE, a WebSphere MQ reason code, or a replication diagnostic code.

User response: Check the Q Apply log file for any diagnostic or error messages prior to this message.

ASN7628E *program_name : program_identifier : The Q Apply browser browser_name (receive queue queue_name, replication queue map queue_map_name) terminated because it had pending RI constraints that got a violation and there are no in-flight transactions to resolve the constraint violation.*

Explanation: Q Apply browser tried to add the RI constraints to the tables after the table was loaded. However, the operation failed because of a constraint violation. Since there are no more transactions that can resolve this violation, the Q Apply browser stopped.

User response: Compare the source and target tables and fix the rows that cause the constraint violation.

Tip: You can use the asntdiff program to compare the source and target tables.

ASN7629E *program_name : program_identifier : The spill queue name spill_queue_name for Q subscription Q_subscription_name on receive queue receive_queue_name using the model queue name model_queue_name exceeded the maximum of 48 characters for a spill queue name by number.*

Explanation: The spill queue name is generated in Q Apply by concatenating the name of the model queue with the Q subscription information. The maximum

length allowed is 48 characters.

User response: Update the MODELQ column in the IBMQREP_TARGETS table with a model queue name that is smaller in length, so the spill queue name will not exceed 48 characters.

ASN7630I Time period used to calculate average

ASN7631I All transactions applied as of (time)

ASN7632I All transactions applied as of (LSN)

ASN7633I Oldest in-progress transaction

ASN7634I Average end-to-end latency

ASN7635I Average Q Capture latency

ASN7636I Average WSMQ latency

ASN7637I Average Q Apply latency

ASN7638I Current memory

ASN7639I Current queue depth

ASN7640I Log reader currency

ASN7641I Q Apply program status

ASN7642I Receive queue

ASN7643I *program_name : program_identifier : The target nickname_owner.nickname_name is a nickname for the Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The target of the Q subscription is a nickname that refers to a table in the non-DB2 database

User response: This message is for your information. No action is required.

ASN7644W *program_name : program_identifier : Q Capture MQPUT timestamp occurs number milliseconds after Q Apply MQGET on receive queue queue_name. Possible clock skew between source and target servers.*

Explanation: The clocks at the source and target server machines may not be synchronized.

User response: Adjust the clocks at the source or target server so that they match. This will avoid the clock skew.

ASN7645E *program_name : program_identifier : Target table table_name for Q subscription q_subscription_name (receive queue receive_queue_name, replication queue map queue_map_name) does not exist.*

Explanation: The Q Apply program tried to find the target table when activating the Q subscription, but the table does not exist in the target database.

User response: Redefine the Q subscription with the replication administration tools, and either specify a table that does exist or let the administration tools create a new target table. Then activate the Q subscription.

ASN7646E *program_name : program_identifier : The Q Apply program encountered an error for the Q subscription Q_subscription_name on the receive queue queue_name for the replication queue map queue_map_name while Q Apply was writing to the IBMQREP_EXCEPTIONS table. The original row had a source commit LSN of LSN, a source transaction ID of transaction_ID, an SQLDETAILS of SQLDETAILS, and a ROWDETAILS of ROWDETAILS.*

Explanation: The Q Apply program was unable to insert a row into the IBMQREP_EXCEPTIONS control table. It unsuccessfully tried to apply a row to the replication target. The information in this message identifies the row that the agent tried to apply.

User response: Use the SQLCODE, SQLERRMC, OPERATION, and REASON in the message to determine the problem that prevented the Q Apply program from applying the row at the target. Also, ensure that the IBMQREP_EXCEPTIONS table for this Q Apply schema exists and is available. For details see the DB2 DESCRIBE command in the IBM DB2 Information Center.

ASN7647E *program_name : program_identifier : The CCD column column_name for Q subscription subname has an incorrect data type of data_type, incorrect data length of data_length, or both. The expected data type is data_type and the expected data length is data_length.*

Explanation: The specified CCD column for the Q subscription has an unexpected data type, data length, or both.

User response: Alter the CCD table so that the

specified column has the correct data type and data length.

ASN7648E *program_name : program_identifier : SQL expressions are not supported for data_type columns. MAPPING_TYPE is not set up correctly in IBMQREP_TRG_COLS table for Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name).*

Explanation: The MAPPING_TYPE value of E in IBMQREP_TRG_COLS table for the Q subscriptions specifies that a SQL expression is used to transform data. However, expressions are not supported for the specified data type.

User response: Take one of the following actions:

- If you want to specify a SQL expression for the Q subscription, change the Q subscription so that any columns with unsupported data types are not replicated.
- Use the ASNCLP command-line program or Replication Center to change the Q subscription so that expressions are not specified and the MAPPING_TYPE value in the IBMQREP_TRG_COLS table is R (regular).

Start the Q subscription after you make either change.

ASN7649E *program_name : program_identifier : Column column_name of Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name) is a constant, derived constant, or CCD audit column. This column cannot have a before-image column defined at the target.*

Explanation: SQL expressions with constants (such as "IBM") or derived constants (such as CURRENT TIMESTAMP) or CCD audit columns (such as IBMSNAP_LOGMARKER) cannot have before-image columns defined at the target.

User response: Use the ASNCLP command-line program or Replication Center to change the Q subscription so that before-values are not specified for any ineligible expressions or columns, then start the Q subscription.

ASN7650I *program_name : program_identifier : The Q Apply browser thread (receive queue queue_name, replication queue map queue_map_name) will stop after it applies all source transactions that were committed up to timestamp timestamp. The user input was input.*

Explanation: A command was issued to specify that

the Q Apply browser thread stop at the specified timestamp. The browser thread will stop after it processes all transactions with a source commit time that is less than or equal to the specified timestamp.

User response: This message is for your information only. No action is required.

ASN7651I *program_name : program_identifier :* **The Q Apply browser thread (receive queue *queue_name*, replication queue map *queue_map_name*) processed a source transaction with a commit timestamp that was greater than the user-specified stopping point *timestamp*. The browser thread will stop.**

Explanation: The Q Apply program was started with the applyupto parameter, which specified that the browser thread stop at a specified timestamp. The browser thread detected a transaction with a source commit time that was greater than or equal to the timestamp, and so the browser thread will stop.

User response: This message is for your information only. No action is required.

ASN7652E *program_name : program_identifier :* **The Q Apply program found an invalid timestamp for the applyupto parameter when it started. The field *field* of the timestamp *timestamp* is not correct. The Q Apply program will stop.**

Explanation: The timestamp that was provided for the applyupto parameter in the command to start the Q Apply program has an invalid format. The expected format is one of the following: YYYY-MM-DD-HH.MM.SS.MMMMMM, YYYY-MM-DD-HH.MM.SS, YYYY-MM-DD-HH.MM, YYYY-MM-DD-HH, HH.MM, or HH.

User response: Restart the Q Apply program and specify a valid timestamp for the applyupto parameter.

ASN7653E *program_name : program_identifier :* **The command to start the Q Apply program contained both the autostop and applyupto parameters. The Q Apply program will stop.**

Explanation: The applyupto parameter cannot be specified when the value of the autostop parameter is Y (yes). So you cannot specify both applyupto and autostop =Y when you start the Q Apply program. You also cannot specify applyupto when the saved value of autostop in the IBMQREP_APPLYPARMS table is Y.

User response: Restart the Q Apply program while specifying either the autostop parameter or the applyupto parameter. Update the value of autostop in the IBMQREP_APPLYPARMS table if necessary.

ASN7654E *program_name : program_identifier :* **The command to change the Q Apply operating parameters contained the autostop parameter when the applyupto parameter was already specified. The command was not processed.**

Explanation: When you use the asnqacmd command to operate the Q Apply program, the autostop and applyupto parameters cannot be enabled at the same time.

User response: Reissue the asnqacmd command without specifying the autostop parameter.

ASN7655E *program_name : program_identifier :* **The value of the MAXAGENTS_CORRELID is greater than the value of NUM_APPLY_AGENTS in the IBMQREP_RECVQUEUES table for receive queue *recv_queue_name* and replication queue map *repl_qmap_name*. The Q Apply program will stop the browser thread for this receive queue.**

Explanation: The MAXAGENTS_CORRELID value cannot be greater than the NUM_APPLY_AGENTS value.

User response: Set the MAXAGENTS_CORRELID value lower than the NUM_APPLY_AGENTS value in the IBMQREP_RECVQUEUES table.

ASN7656W *program_name : program_identifier :* **The MAXAGENTS_CORRELID value is equal to the NUM_APPLY_AGENTS value in the IBMQREP_RECVQUEUES table for receive queue *recv_queue_name* and replication queue map *repl_qmap_name*. The browser thread (replication queue map *repl_qmap_name*) will not limit parallelism based on the correlation ID.**

Explanation: Because the MAXAGENTS_CORRELID value is the same as the NUM_APPLY_AGENTS value, the Q Apply browser will not serialize transactions based on the correlation ID.

User response: If lock contention is a problem, you can serialize transactions by setting the MAXAGENTS_CORRELID value smaller than the NUM_APPLY_AGENTS value. If you set the MAXAGENTS_CORRELID value to zero, you will not receive this warning and the browser will continue applying transactions in parallel.

ASN7657E *program_name : program_identifier* **The Q Apply program failed to process the spillsub command for the specified Q subscription *subscription_name* and specified receive queue *queue_name* (replication queue map *queue_map_name*). The subscription state is *subscription_state*. Reason code: *reason_code*.**

Explanation: The following values are valid for the reason code:

- 0
The program is in the initialization phase.
- 1
The specified receive queue was not found.
- 2
The specified Q subscription was not found.
- 3
The spillsub command can only be issued to subscriptions in A (active) state. The Q subscription was in a state other than active.
- 4
The target table for the specified Q subscription has a referential integrity relationship with another table.
- 5
The temporary spill queue could not be created.

User response: Perform the action that corresponds to the reason code:

- 0
Reissue the command after the program has finished initializing.
- 1
Check that the name of the receive queue is correct and reissue the command.
- 2
Check that the name of the Q subscription is correct and reissue the command.
- 3
If the state of the Q subscription is I or T, an error might have occurred with the Q subscription before the spillsub command was issued. Check the previous messages from the Q Apply program to determine the cause of the error. For other states, you must wait for the Q subscription to become active.
- 4

The spillsub command is not supported for target tables that have referential integrity relationships. To stop applying replicated transactions to target tables with referential integrity relationships, you must stop the entire receive queue by using the stopq command.

5

Check the log file for the Q Apply program for previous diagnostic or error messages to determine the cause of the failure.

ASN7658E *program_name : program_identifier* **The Q Apply agent thread encountered an error while spilling a row message for the Q subscription *Q_subscription_name* to the spill queue *spill_queue_name*. The browser thread for receive queue *queue_name* (replication queue map *queue_map_name*) will stop.**

Explanation: The Q Apply program was following the specified error action for the Q subscription because of an error during the attempt to apply a row. The specified action prompts Q Apply to place messages in a temporary spill queue until the error is corrected. However, Q Apply was unable to place the Q subscription in spill mode and so Q Apply will stop the browser thread.

User response: Look in the Q Apply diagnostic log file or the IBMQREP_APPLYTRACE table for messages that might indicate the cause of the failure. Make sure that the WebSphere MQ model queue that is used to create spill queues for the Q subscription is created, and that the queue is correctly specified in the IBMQREP_TARGETS table.

Note: If the DFTMODELQ parameter was used when Q Apply was started, Q Apply uses the model queue that is specified with this parameter.

After the problem is corrected, use the startq command to restart the receive queue.

ASN7659I *program_name : program_identifier* : **The Q Apply browser thread (receive queue *queue_name*, replication queue map *queue_map_name*) has placed the Q subscription *Q_subscription_name* in spill mode based on the spillsub command. The spill queue is *queue_name*.**

Explanation: The Q Apply program is placing messages for the Q subscription in a temporary spill queue as directed by the spillsub command.

User response: This message is for your information only. No action is required.

ASN7660I *program_name : program_identifier* **The Q Apply browser thread (receive queue *queue_name*, replication queue map *queue_map_name*) has resumed the Q subscription *Q_subscription_name* based on the resumesub command.**

Explanation: The Q subscription that was placed in spill mode will be returned to normal processing as directed by resumesub command. Spilled messages will be applied. New messages will continue to be spilled and referential integrity constraints are still dropped on the target table. When the spill queue is empty, the constraints will be added back and the Q subscription will be placed in active (A) state.

User response: This message is for your information only. No action is required.

ASN7661I *program_name : program_identifier* **The REINITQ command was processed successfully. The Q Apply program will limit parallelism to *number* agents for receive queue *recv_queue_name* and replication queue map *queue_map_name*.**

Explanation: The MAXAGENTS_CORRELID parameter specifies the maximum number of agent threads that can process transactions with the same z/OS correlation ID or job name. This parameter allows you to reduce lock contention by limiting parallelism for transactions from the same batch job on a z/OS Q Capture server. Reducing lock contention can lower latency.

By using this setting, the Q Apply program will limit parallelism during batch workloads to prevent lock contention.

User response: This message is for your information only. No action is required.

ASN7662E *program_name : program_identifier* : **The Q Apply program failed to process the resumesub command for the specified Q subscription *subscription_name* and receive queue *queue_name* (replication queue map *queue_map_name*). The Q subscription state is *state*. Reason code: *reason_code*.**

Explanation: The following values are valid for the reason code:

- 0 The program is in the initialization phase.
- 1 The specified receive queue was not found.
- 2 The specified Q subscription was not found.

3

The resumesub command can only be issued to Q subscriptions in S (spilling) or P (pending) state. The Q subscription was in a state other than spilling or pending.

User response: Perform the action that corresponds to the reason code:

0

Reissue the command after the program finishes initializing.

1

Check that the name of the receive queue is correct and reissue the command.

2

Check that the name of the Q subscription is correct and reissue the command.

3

Take one of the following actions, depending on the Q subscription state:

- If the state is I or T, an error might have occurred with the Q subscription before the resumesub command was issued. Check the previous messages from the Q Apply program to determine the cause of the error. After the error is fixed, reissue the resumesub command.
- If the state is A, R, or F, the Q subscription already resumed and you do not need to issue the resumesub command.

ASN7663E *program_name : program_identifier* **The Q Apply program failed to process the loadonesub command for the specified Q subscription *subscription_name* and specified receive queue *queue_name* (replication queue map *queue_map_name*). The subscription state is *subscription_state*. Reason code: *reason_code*.**

Explanation: The following values are valid for the reason code:

0

The program is in the initialization phase.

1

The specified receive queue was not found.

2

The specified Q subscription was not found.

3

The loadonesub command can only be issued for Q subscriptions in the E state.

4

The loadonesub command should be issued only for the peer-to-peer Q subscription that was used for the CAPSTART signal. This subscription is the same subscription that has the source table that was used as the load source for the manual load.

User response: Perform the action that corresponds to the reason code:

0

Reissue the command after the program has finished initializing.

1

Check that the name of the receive queue is correct and reissue the command.

2

Check that the name of the Q subscription is correct and reissue the command.

3

If the state of the Q subscription is I or T, an error might have occurred with the Q subscription before the loadonesub command was issued. Check the previous messages from the Q Apply program to determine the cause of the error. For all other states, verify that you specified the correct Q subscription and verify that the state of the Q subscription is 'E.' Reissue the command.

4

Check that you specified the correct peer-to-peer Q subscription for the loadonesub command. Use the Q subscription that specifies the source table that you used to perform the initial load. Reissue the command for the correct Q subscription.

ASN7664W *program_name : program_identifier* **The value of the MAXAGENTS_CORRELID column is greater than the value of the NUM_APPLY_AGENTS column in the IBMQREP_RECVQUEUES table. The browser thread (receive queue *recv_queue_name* and replication queue *map repl_qmap_name*) will use the previous value for the MAXAGENTS_CORRELID.**

Explanation: The MAXAGENTS_CORRELID value cannot be greater than the NUM_APPLY_AGENTS value.

User response: Set the MAXAGENTS_CORRELID value lower than the NUM_APPLY_AGENTS value in the IBMQREP_RECVQUEUES table.

ASN7665E *program_name : program_identifier* : **The Q Apply program (receive queue *receive_queue_name*, replication queue *map queue_map_name*) cannot apply segmented large object (LOB) data to an XML column in the target table for Q subscription *Q_subscription_name*.**

Explanation: The Q Apply program encountered an error while applying a row that contained LOB data to a target table with an XML column. If the error occurred during spill queue processing, the Q subscription is deactivated. Otherwise the browser thread terminates.

User response: To replicate LOB data to an XML column at the target, the LOB data must be sent within the transaction message, also known as inline, rather than in segments. Follow one of these procedures:

If the Q subscription was deactivated

1. Stop the Q Capture program.
2. Update the LOB_SEND_OPTION column in the IBMQREP_CAPPARMS table with a value of I.
3. Start the Q Capture program in warm mode.
4. Use the replication administration tools to activate the Q subscription.

If the browser thread stopped

1. Change the value of the STATE column for the Q subscription to I in the IBMQREP_TARGETS table.
2. Use the asnqacmd startq command to start the browser for the receive queue.
3. Wait until the STATE column in the IBMQREP_SUBS table changes to a value of I.
4. Stop the Q Capture program.
5. Update the LOB_SEND_OPTION column in the IBMQREP_CAPPARMS table with a value of I.
6. Start the Q Capture program in warm mode.
7. Use the replication administration tools to activate the Q subscription.

ASN7666E *program_name : program_identifier* : **The INSERT_BIDI_SIGNAL parameter was set to N, but the IBMQREP_IGNTRAN table does not have a valid entry for Q Apply's authorization ID *auth_ID* or plan name *plan_name*. The Q Apply program stopped. Specify a valid entry in the IBMQREP_IGNTRAN table or set INSERT_BIDI_SIGNAL to Y.**

Explanation: You can use one of two methods to

prevent recapture of transactions in a bidirectional replication environment:

- Accept the default behavior in which the Q Capture and Q Apply programs use signal inserts (insert_bidi_signal=Y).
- Use the Q Capture program's ability to ignore specified transactions (insert_bidi_signal=N and an entry to the IBMQREP_IGNTRAN table with Q Apply's authorization ID or plan name).

User response: If you want to use the ignore transaction method, insert a valid identifier for the Q Apply program into the AUTHID or PLANNAME column of the IBMQREP_IGNTRAN table at the server that is shared by the Q Apply program and Q Capture program in the bidirectional configuration. Then restart the Q Apply program.

ASN7667E *program_name : program_identifier :* **The INSERT_BIDI_SIGNAL parameter was set to N, but the IBMQREP_IGNTRAN table does not exist. The Q Apply program stopped. Create the IBMQREP_IGNTRAN table and specify Q Apply's authorization ID or plan name (z/OS). Alternatively, set INSERT_BIDI_SIGNAL to Y.**

Explanation: You can use one of two methods to prevent recapture of transactions in a bidirectional replication environment:

- Accept the default behavior in which the Q Capture and Q Apply programs use signal inserts (insert_bidi_signal=Y).
- Use the Q Capture program's ability to ignore specified transactions (insert_bidi_signal=N and an entry to the IBMQREP_IGNTRAN table with Q Apply's authorization ID or plan name).

User response: If you want to use the ignore transaction method, take these actions:

1. Upgrade the Q Capture control tables to Version 9.1 Fix Pack 2 or newer to add the IBMQREP_IGNTRAN table.
2. Insert a valid identifier for the Q Apply program into the AUTHID or PLANNAME column of the IBMQREP_IGNTRAN table at the server that is shared by the Q Apply program and Q Capture program in the bidirectional configuration.
3. Restart the Q Apply program.

ASN7668E *program_name : program_identifier :* **The unique index index_name on the target table table_owner.table_name for the Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name) has more key columns than the maximum number of limit key columns.**

Explanation: The Q Apply program detected that the number of key columns for the unique index exceeds the maximum number that DB2 allows. The limit differs by platform and version. For more information, see the DB2 Information Center.

User response: Drop and recreate the unique index so that it is within the maximum number of key columns.

ASN7669W *program_name : program_identifier :* **The Q Apply browser thread on receive queue queue_name (replication queue map queue_map_name) will ignore the value of MAXAGENTS_CORRELID in the IBMQREP_RECVQUEUES table and not limit parallelism based on correlation ID. The reason code is reason_code.**

Explanation: The Q Apply program uses the MAXAGENT_CORRELID setting for the receive queue only if the Q Capture server is on z/OS and at Version 9.1 with the PTF for APAR PK49430 or later. The following reason codes apply:

0

The Q Capture program is an older version that does not support sending correlation ID information.

1

The Q Capture server is not on z/OS. Correlation ID information is not valid for non-z/OS sources.

User response: Upgrade the Q Capture server, or if the Q Capture server is on Linux, UNIX, or Windows, set the MAXAGENT_CORRELID column in the IBMQREP_RECVQUEUES table to NULL for this receive queue.

ASN7670I *program_name : program_identifier :* **Transaction transaction_ID was skipped successfully from receive queue queue_name (replication queue map queue_map_name).**

Explanation: The transaction ID that was specified with the skiptrans parameter on program invocation or with the startq command was successfully skipped from the receive queue.

User response: This message is for your information only. No action is required.

ASN7671I *program_name : program_identifier :* **Transaction IDs in the range beginning_ID to ending_ID were successfully skipped from receive queue queue_name (replication queue map queue_map_name). Number of transactions skipped: number.**

Explanation: The transaction IDs in the range that was specified in the skiptrans parameter on program invocation or with the startq command were read from the receive queue but not applied. The messages were deleted from the receive queue.

User response: This message is for your information only. No action is required.

ASN7673E *program_name : program_identifier :*
Transaction transaction_id could not be skipped from receive queue queue_name (replication queue map queue_map_name).
Reason code: reason_code. The program stops reading from the queue.

Explanation: The transaction ID that was specified with the skiptrans parameter on program invocation or with the startq command has an invalid format or was not found on the receive queue. The following values are valid for the reason code:

0

The transaction ID to skip was not found on the receive queue.

1

The transaction ID to skip has an invalid format

User response: Review the reason codes in the explanation, and take the appropriate action:

0

Use the asnqmfmt tool or asnqxfmt tool on z/OS to check if the transaction ID exists on the receive queue. If it exists and you want to skip it, start the program again while specifying this transaction ID with the skiptrans parameter.

1

The following characters can be used to specify a transaction ID with the skiptrans parameter:

- Uppercase A-F
- Lowercase a-f
- Numerics 0-9
- Special characters: "-" and ":" only

For example, the following transaction IDs are allowed:

0000:0000:0000:51a1:0000

00000000000051a10000

0000:0000:0000:0000:51a1-
0000:0000:0000:0000:51a8

Note:

- Leading zeros cannot be omitted in a transaction ID. The transaction ID can be specified with or without colons. For example, 0000:0000:0000:51a1:0001 is treated the same as 00000000000051a10001.
- The length of each unit between colons in the input transaction ID should be equal to 4. For example, 0000:1090:1234:5671:001 is not allowed because the length of the last unit is less than 4.
- Wild card characters such as (*) are not allowed. For example, a range clause such as "*" -1123:0000:01fa:bbc2:0001" is not allowed.
- No spaces are allowed between the delimiters (semi-colon, hyphen, comma) that are used in the syntax for skiptrans parameter.
- The range of transaction IDs to skip includes beginning and ending transaction IDs.

ASN7674E *program_name : program_identifier :* **The Q Apply program encountered an invalid Q subscription state while processing a schema message for Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name). Q Apply expected the state to be W but found a state of state. The Q subscription was not activated.**

Explanation: The Q Apply program received a schema message for the Q subscription that indicates that the target table was loaded. The Q subscription state in the IBMQREP_TARGETS table should be W but is a different value.

User response: Review the setup and restart the Q subscription.

ASN7675I *program_name : program_identifier :* **The Q Apply program successfully processed a schema message for Q subscription Q_subscription_name (receive queue queue_name, replication queue map queue_map_name). Q Apply changed the state of the Q subscription to E (being loaded by an external application).**

Explanation: The Q Apply program received a schema message for the Q subscription and successfully changed its state.

User response: This message is for your information only. No action is required.

ASN7676I *program_name : program_identifier :* **The Q Apply program successfully processed a schema message for Q subscription *Q_subscription_name* (receive queue *queue_name*, replication queue map *queue_map_name*). Q Apply changed the state of the Q subscription to L.**

Explanation: The Q Apply program received a schema message for the Q subscription and successfully changed its state.

User response: This message is for your information only. No action is required.

ASN7677I **Restart point for Q Capture (MAXCMTSEQ)**

ASN7678I **Agents processing transaction**

ASN7679I **Agents waiting for transaction**

ASN7680I **Agents processing internal messages**

ASN7682I **Agents in initializing state**

ASN7684E *program_name : program_identifier :* **The Q Apply agent thread encountered an error while spilling a row message for the Q subscription *name* to the spill queue *queue_name*. The browser thread for receive queue *queue_name* (replication queue map *queue_map_name*) will stop.**

Explanation: The Q Apply program had an unrecoverable failure while trying to put a row message for the Q subscription on a temporary spill queue. Q Apply stopped the browser thread.

User response: Check the Q Apply log file or any diagnostic or error messages before this message to determine the cause of the failure. For problems with the spill queue, the diagnostic message lists the WebSphere MQ error code. The spill queue might be full. Fix the cause of the underlying failure and use the startq command to restart message processing on the receive queue.

ASN7687E *program_name : program_identifier :* **The Q Apply program encountered an SQL error while it was running in buffered insert mode for Q subscription *name* (receive queue *queue_name*, replication queue map *queue_map_name*). The program stops reading from the queue.**

Explanation: The Q Apply program encountered an error or conflict while trying to apply a transaction for this Q subscription in buffered insert mode. When it

runs in this mode, Q Apply ignores regular conflict or error handling options, and any SQL error that is encountered by a browser causes Q Apply to stop reading from the affected receive queue. No entry is added to the IBMQREP_EXCEPTIONS table. Refer to the Q Apply diagnostic log file for details about why the error occurred.

User response: Address the problem that was identified in the diagnostic log or restart the Q Apply program with buffered_inserts=n. Then activate the receive queue by using the startq command to move past the exception. No changes are lost. The Q Apply program continues to apply changes to other targets.

ASN7690I **Current queue percentage full**

ASN7694I *program_name : program_identifier :* **Before it loads the target table *table_owner.table_name*, the program will bind the packages that are required for the DB2 EXPORT utility with the option CONCURRENTACCESSRESOLUTION WAIT FOR OUTCOME in COLLECTION ASN.**

Explanation: To ensure that the DB2 EXPORT utility successfully exports all data from the source server, Q Apply must bind the SQL packages for this utility (in @db2ubind.lst) with this option. This option forces DB2 EXPORT to wait until all in-progress transactions that modify the source table are completed before beginning the export.

User response: This message is for your information only. No action is required.

ASN7695E *program_name : program_identifier :* **The program was unable to bind the packages that are required for the DB2 EXPORT utility. The Q subscription *name* (receive queue *queue_name*, replication queue map *queue_map_name*) will be deactivated. Bind failed on package *pkg_name* from the file *path_filename*. The SQLSTATE *sqlstate* was returned.**

Explanation: The program encountered an error while attempting to bind a required package.

User response: Interpret the SQLCODE of the failed bind call. Fix the problem and then activate the Q subscription. You can manually bind the required packages by connecting to the source database and running the following command:

```
db2 bind @db2ubind.lst CONCURRENTACCESSRESOLUTION
WAIT FOR OUTCOME COLLECTION ASN.
```

Rebinding the utilities package in collection ASN does not affect the use of those utilities by other

applications. If you cannot rebind the utilities packages, you might also consider a different load type for the Q subscription. See "IBMQREP_TARGETS table" in the DB2 Information Center for a list of the available load types.

ASN7697E *program_name : program_identifier : The Q Apply browser thread detected a message that has either an incorrect format or was put on the receive queue queue_name, replication queue map queue_map_name by an application other than the Q Capture program. The browser thread stopped reading from the queue but left the message on the queue so that the problem can be diagnosed. The WebSphere MQ message ID of the message is message_ID. The queue was placed in inactive (I) state.*

Explanation: A valid replication message comes from the Q Capture program and has C (compact) format. The Q Apply browser thread read a message that is not a valid replication message. The message was left on the queue for diagnostic purposes.

User response: Use the asnqmfmt command to format the first 50 messages on the queue and save the formatted message output. Contact IBM Software Support with the formatted message output and Q Apply diagnostic log file. Use the following syntax for asnqmfmt:

```
asnqmfmt receive_queue_name queue_manager_name
-mqmd -hex -l 50 -o filepath
```

Where filepath is the location on your system where you want asnqmfmt to save the text file with the formatted messages output, and the name of the file.

Inspect the output to see if the "MsgId" field contains the specified message ID. If not, continue getting the asnqmfmt output of more messages until the output contains the ID in question.

ASN7699I *program_name : program_identifier : The Q Apply browser thread for receive queue queue_name (replication queue map queue_map_name) applied all transactions up to the transaction commit point that is identified by OLDEST_COMMIT_LSN_value if the source table is on z/OS or OLDEST_COMMIT_SEQ_value if the source table is on Linux, UNIX, or Windows.*

Explanation: The Q Apply program's caught-up point with respect to the source is reported in the message text. You can use the value that is appropriate for your source operating system for the Q Capture maxcmtseq

parameter to restart Q Capture without the need for a full refresh of target tables in situations where the Q Capture restart information or WebSphere MQ messages are lost.

User response: This message is for your information only. No action is required.

ASN7700I *program_name : program_identifier : The Q Apply program received a deactivate subscription message for the Q subscription name (receive queue queue_name, replication queue map queue_map_name). The Q subscription will be deactivated.*

Explanation: The program received a message from the Q Capture program to deactivate the Q subscription. A user might have deliberately deactivated the Q subscription, or the Q Capture program might have deactivated it because of an error during processing.

User response: This message is informational only. If the deactivation of this Q subscription was unexpected, check the associated Q Capture log file for reasons for the deactivation.

ASN7701W *program_name : program_identifier : The load utility utility_name will not wait for uncommitted transactions to complete at the source database before it begins to select source data. Reason code: reason_code.*

Explanation: The program determines whether the source database uses CURRENTLY COMMITTED locking semantics by default. If the load utility does not disable this behavior, data loss will result if in-flight transactions are in progress at the time a load phase is initiated. See the reason code for possible workarounds. For all situations, the default workaround to prevent potential data loss is to suspend applications from the time the Q subscription is started (when a CAPSTART signal is inserted into the IBMQREP_SIGNAL table) until at least the start of the load phase (when the Q subscription state changes to L or A in the IBMQREP_SUBS table). Possible reason codes are as follows:

0

The DSNUTILS package (z/OS only) that enables LOAD from CURSOR cannot be bound at the source database with WAIT FOR OUTCOME semantics. An application suspend is required.

1

LOAD from CURSOR (Linux, UNIX, and Windows only) has no option to set WAIT

FOR OUTCOME behavior. Either perform an application suspend or perform the following steps:

1. Bind the db2cli packages at the source database with the following command:

```
db2 bind @db2cli.lst
      CONCURRENTACCESSRESOLUTION
      WAIT_FOR_OUTCOME COLLECTION ASN
```
2. Add the following name-value pair to the db2cli.ini file at the federated database, under the stanza that declares the options for the server definition to which the nickname belongs:

```
[data_source_name]
```

```
CURRENTPACKAGESET=ASN
```

Where data_source_name is the source database that the db2cli.bnd packages were bound against.

2

The LOAD utility from a local database (Linux, UNIX, and Windows only) has no option to set WAIT FOR OUTCOME behavior. An application suspend is required.

User response: Refer to the details in the "Explanation" section that are associated with the reason code in the message text.

ASN7704W *program_name : program_identifier : The Q Apply program with schema schema_name is running against a federated target. One or more nickname options are not correctly set for optimal performance of queries that are issued by the Q Apply program against the nicknames that refer to the tables in the non-DB2 target database. See the message explanation for more detail.*

Explanation: To improve performance for federated targets, use the following settings:

Oracle targets

Set VARCHAR_NO_TRAILING_BLANKS = Y for the RECVQ column of the IBMQREP_DONEMSG nickname and for the SPILLQ column of the IBMQREP_SPILLEDROW and IBMQREP_SPILLQS nicknames.

Microsoft SQL Server targets

Set BINARY_REP = Y for the MQMSGID column of the IBMQREP_DONEMSG nickname.

User response: Change the settings to the recommended ones if you are observing slow performance from the Q Apply program. The following

SQL can be used to add the settings:

Oracle targets

```
alter nickname
  apply_schema.IBMQREP_DONEMSG
alter COLUMN RECVQ OPTIONS
  (ADD VARCHAR_NO_TRAILING_BLANKS 'Y')
alter nickname
  apply_schema.IBMQREP_SPILLEDROW
alter COLUMN SPILLQ OPTIONS
  (ADD VARCHAR_NO_TRAILING_BLANKS 'Y')
alter nickname
  apply_schema.IBMQREP_SPILLQS
alter COLUMN SPILLQ OPTIONS
  (ADD VARCHAR_NO_TRAILING_BLANKS 'Y')
```

Microsoft SQL Server targets

```
alter nickname
  apply_schema.IBMQREP_DONEMSG
alter COLUMN MQMSGID
  OPTIONS (ADD BINARY_REP 'Y')
```

ASN7705I *program_name : program_identifier : The Q Apply browser thread for receive queue queue_name (replication queue map queue_map_name) encountered an empty receive queue while running with the NOWAIT qualifier set in the applyupto parameter. The browser thread stops.*

Explanation: The Q Apply program was started with the applyupto parameter with a specified timestamp and the NOWAIT optional qualifier. NOWAIT specifies that the browser thread should stop when it encounters an empty receive queue even if the specified timestamp has not been exceeded by any transactions processed so far.

User response: This message is for your information only. No action is required.

ASN7706W *program_name : program_identifier : The browser thread for receive queue queue_name (replication queue map queue_map_name) did not process a WebSphere MQ message with size zero. The message did not have any content. The WebSphere MQ message ID was msgid.*

Explanation: Messages with a size of zero are not expected. The Q Apply browser continues normal processing after not processing this message. Because this message did not have any content, Q Apply could not determine if this was a transaction message or a control message such as a heartbeat message.

User response: Check the Q Capture log for other errors that might have occurred while Q Capture was replicating or publishing messages to WebSphere MQ.

ASN7707E *program_name : program_identifier :* **A row change was received from a source database with the DB2 pureScale Feature, and the Q subscription *subscription_name* (receive queue *queue_name*, replication queue map *queue_map_name*) specifies a CCD target. CCD tables are supported only in unidirectional replication in an environment without DB2 pureScale. The Q Apply browser thread will stop.**

Explanation: CCD tables are not supported in this configuration.

User response: Migrate the target to a version of DB2 database with DB2 pureScale that supports CCD targets.

ASN7708W *program_name : program_identifier :* **The Q Apply browser thread for receive queue *queue_name* (replication queue map *queue_map_name*) placed the Q subscription *subscription_name* in pending state. An error occurred and the error action for the Q subscription specifies that messages be placed in a temporary spill queue until the error is resolved. The spill queue *spill_queue_name* will be used.**

Explanation: The error action that was specified for the Q subscription (value of B in the ERROR_ACTION column of the IBMQREP_TARGETS table) prompts the Q Apply program to direct data for the Q subscription to a spill queue when SQL errors are detected.

User response: Take the following actions:

1. Look for messages in the IBMQREP_APPLYTRACE table or Q Apply log file that indicate which SQL error occurred.
2. Fix the error.
3. Issue the resumesub command to prompt Q Apply to apply data from the spill queue and begin normal processing for the Q subscription.

ASN7709E *program_name : program_identifier :* **The Q Apply program encountered an error for Q subscription *subscription_name* (receive queue *queue_name*, replication queue map *queue_map_name*) and the error action that was specified for the Q subscription is B (break transaction and send changes to a spill queue). However, error action B is not allowed for tables with referential integrity constraints.**

Explanation: The table that is specified for this Q subscription has a referential integrity relationship with one or more other tables. Error action B is not allowed

for tables with referential integrity constraints.

User response: Take the following actions:

1. Use the replication administration tools to choose a different error action for the Q subscription.
2. If the Q subscription is disabled, restart the Q subscription.
3. If the browser thread for the receive queue was stopped, use the startq command to restart message processing on the receive queue.

ASN7710E *program_name : program_identifier :* **A Q subscription *subscription_name* already exists for table *table_owner.table_name* that specifies the receive queue *queue_name* (replication queue map *queue_map_name*). However, the Q subscription is in a state other than inactive (I). The program will stop reading from the queue.**

Explanation: The Q Apply program tried to create a Q subscription for the table in response to a schema-level Q subscription at the source, but a Q subscription for the table already exists. To create a new Q subscription, the existing Q subscription must be in inactive state so that it can be deleted. The Q Apply program will stop reading from the receive queue.

User response: Follow these steps:

1. Deactivate the existing Q subscription.
2. Use the replication administration tools to delete the existing Q subscription.
3. Issue a startq command to tell Q Apply to start reading from the receive queue.

ASN7711I *program_name : program_identifier :* **Q subscription *subscription_name* was successfully created for the target table *table_owner.table_name*. That table uses receive queue *queue_name*, and that table uses replication queue map *queue_map_name*. The Q subscription was automatically created because a schema-level subscription *subscription_name* was defined at the source.**

Explanation: The Q Apply program was prompted to automatically create a Q subscription by the Q Capture program because a schema-level subscription was defined at the source for all tables with the same owner. The Q subscription was successfully created.

User response: This message is for your information only. No action is required.

ASN7712I *program_name : program_identifier : The DDL operation operation was successfully applied at the target. Receive queue name: queue_name. Replication queue map name: queue_map_name. SQL statement: SQL_statement.*

Explanation: The Q Apply program successfully replicated the specified DDL operation. The SQL statement might be truncated if it is too long to be printed.

User response: This message is for your information only. No action is required.

ASN7713E *program_name : program_identifier : The DDL operation operation that was sent on receive queue queue_name (replication queue map queue_map_name) could not be successfully applied. The SQL statement is SQL_statement. The program will stop reading from the queue.*

Explanation: The Q Apply program failed to execute the DDL operation that was replicated from the source. An SQL error is likely to have occurred.

User response: Follow these steps:

1. Look for prior messages that mention SQLSTATE and SQLCODE information.
2. Fix the error.
3. Issue the startq command on the receive queue.

ASN7719W *program_name : program_identifier : The alteration of column column_name in table table_owner.table_name (Q subscription name, receive queue queue_name, replication queue map queue_map_name) placed the index of the altered column in a REBUILD-pending state. DB2 issued completion code SQL610W with sqlerrmc sqlerrmc_out during execution.*

Explanation: The Q Apply program performed an ALTER TABLE ALTER COLUMN SQL statement that changed the status of an index for the altered column to REBUILD pending. REBUILD-pending state indicates that the affected index or index partition is broken and must be rebuilt from the data. All subsequent DML SQL statements on the affected table will fail until the index is removed from REBUILD-pending state.

User response: You can reset the REBUILD-pending status for an index with any of these operations:

- REBUILD INDEX
- REORG TABLESPACE SORTDATA
- REPAIR SET INDEX with NORBDPEND
- START DATABASE command with ACCESS FORCE

ASN7721I *program_name : program_identifier : The Q Apply browser thread for receive queue queue_name (replication queue map queue_map_name) is already actively processing the queue. The startq command will be ignored.*

Explanation: A command to start the browser thread for the receive queue was received, but the browser thread is already active.

User response: This message is for your information only. No action is required.

ASN7722I *program_name : program_identifier : The Q Apply browser thread for receive queue queue_name (replication queue map queue_map_name) is already inactive. The stopq command will be ignored.*

Explanation: A command to stop the browser thread for the receive queue was received, but the browser thread is already inactive.

User response: This message is for your information only. No action is required.

ASN7723E *program_name : program_identifier : The stopq command for receive queue queue_name (replication queue map queue_map_name) cannot be completed. Reason code: reason_code.*

Explanation: The following values are valid for the reason code:

1

The IBMQREP_RECVQUEUES table contains no information about this queue.

2

The program is still in the initialization phase.

User response: Review the reason codes in the explanation, and take the appropriate action:

1

Check the IBMQREP_RECVQUEUES table for the information about the queue. Redefine the Q subscription by using one of the replication administration tools.

2

Reissue the command later.

ASN7727W *program_name : program_identifier : The column **column_name** that was added to the Q subscription **name** (receive queue **queue_name**, replication queue map **queue_map_name**) already exists at the target table **table_owner.table_name**. The default value of the target column does not match the default value of the source column.*

Explanation: The Q Apply program tried to alter the target table to add a new column, but the new column already exists with a default value that is different from the default value at the source. Q Apply added the column to the Q subscription and started replicating changes.

User response: Check the definitions of the source and target tables to determine the default column values. If the mismatch is not intended, alter the source or target table.

ASN7728W *program_name : program_identifier : The DDL operation operation that was sent on the receive queue **queue_name** (replication queue map **queue_map_name**) will not be applied because the referenced Q subscription **name** with SUBID **subid** was not found or is inactive.*

Explanation: The Q Apply program did not process a replicated DDL operation because the related Q subscription was either not found or inactive.

User response: No action is required if you deleted the Q subscription or expected it to be inactive. Otherwise, follow these steps:

1. Manually alter the target table with the DDL operation that could not be replicated.
2. Look for previous error messages to determine what caused the Q subscription to become inactive, and fix the error.
3. Use the replication administration tools or a CAPSTART signal to activate the Q subscription.

ASN7729W *program_name : program_identifier : The Q subscription **name** was successfully created for source table **table_owner.table_name**. When the Q Apply program created the entry for this new Q subscription in control table **table_owner.table_name** based on the schema-level Q subscription **name** that uses receive queue **queue_name** (replication queue map **queue_map_name**), it deleted an existing entry for the same source table.*

Explanation: After a CREATE TABLE operation is detected for a table with a matching schema-level Q

subscription, the Q Capture program informs the Q Apply program so that it can automatically create the Q subscription entry in the IBMQREP_TARGETS and IBMQREP_TRG_COLS control tables. If the Q subscription entry already exists in one or both of these control tables, Q Apply deletes the prior Q subscription entry and creates the new entry.

User response: You might want to determine why the Q subscription entry for the same target table already existed in one or both of these control tables. If Q Capture was restarted from a prior point in the log, it might have encountered a log record for the same CREATE TABLE operation and replicated details of this operation to the target.

ASN7730W *program_name : program_identifier : In response to a CREATE TABLE operation at the source database or subsystem, the Q Apply program tried to create a Q subscription for source table **table_owner.table_name** based on the schema-level Q subscription **name** that uses replication queue map **queue_map_name** on receive queue **queue_name**. However, the individual Q subscription could not be created because the insertion of Q subscription information into the IBMQREP_TARGETS control table failed. The SQL code or internal replication reason code is **sqlcode_or_reason_code**. Q Apply will continue processing changes for other Q subscriptions on this receive queue.*

Explanation: After a CREATE TABLE operation for a table with a matching schema-level Q subscription is detected, the Q Capture program informs the Q Apply program so that it can automatically create the Q subscription entry in the IBMQREP_TARGETS table. If the insertion of the Q subscription entry fails, the Q subscription is not created and replication does not start for the source table.

User response: If you want to replicate this source table, use the replication administration tools to create a Q subscription for the table. Make sure that the source and target tables exist or let the replication tools create the target table. Check the SQL code or internal reason code to determine why the insert operation into IBMQREP_TARGETS failed.

ASN7731W *program_name : program_identifier :* **The program tried to create a new Q subscription entry in the IBMQREP_TRG_COLS control table for source table *table_owner.table_name* based on the schema-level Q subscription name that uses replication queue map *queue_map_name* and receive queue *queue_name*. However, the insertion of Q subscription information into the control table failed. The SQL code or internal replication reason code is *sqlcode_or_reason_code*. The Q subscription for this source table will be deactivated. Q Apply will continue processing changes for other Q subscriptions that use this receive queue.**

Explanation: After a CREATE TABLE operation for a table with a matching schema-level Q subscription is detected, the Q Apply program automatically adds Q subscription information about the columns in the target table to the IBMQREP_TRG_COLS control table. However, the insertion of this information into the control table failed. The Q subscription will be deactivated in the IBMQREP_TARGETS table.

User response: Follow these steps:

1. Check the SQL code or internal reason code to determine why the insert into IBMQREP_TRG_COLS failed.
2. Use the replication administration tools to delete and recreate the Q subscription for the table. Make sure that the target table exists or let the replication tools create the target table.

ASN7732E *program_name : program_identifier:* **The Q Apply program could not find the spill queue *spill_queue_name* for Q subscription name that uses receive queue *queue_name* (replication queue map *qmap_name*). Q subscription state is *state*. The Q subscription will be deactivated.**

Explanation: When a Q subscription specifies that the target table should be loaded with data from the source, either by the Q Apply program (automatic load) or outside of replication (manual load), a spill queue must exist to hold changes from the source table during the loading process. The Q Apply program could not find a spill queue with the name that is specified for the Q subscription, and so it deactivated the Q subscription. This situation can occur if the same model queue name is used for multiple Q Apply programs that share the same queue manager and one Q Apply program deletes the spill queue after using it.

Model queue names can be specified in two ways:

- At the Q subscription level (the names are stored in the MODELQ column of each Q Apply program's IBMQREP_TARGETS table).
- At the Q Apply instance level by using the *dftmodelq* initialization parameter.

When multiple Q Apply programs use the same queue manager, all model queue names that are specified through either of these methods must be unique.

User response: Check to see if the specified spill queue was deleted. If multiple Q Apply programs share the same queue manager, ensure that all model queue names are unique. After you change the spill queue name, activate the Q subscription.

ASN7733W *program_name : program_identifier :* **The DDL operation operation that was sent on the receive queue *queue_name* (replication queue map *queue_map_name*) will not be applied because the object that the DDL would create at the target database or subsystem already exists. The DDL statement is *SQL_statement*.**

Explanation: The Q Apply program did not process a replicated DDL operation because the object that was created at the source and for which the DDL was replicated already exists at the target.

User response: No action is required unless the object at the target does not match the object definition at the source. In that case, manually alter the object to match the source definition.

ASN7734W *program_name : program_identifier :* **The DDL operation operation that was sent on receive queue *queue_name* (replication queue map *queue_map_name*) will not be applied because the object that was requested to be dropped at the target was not found. The SQL statement is *SQL_statement*.**

Explanation: The Q Apply program did not process a replicated DDL operation because the object that was to be dropped was not found.

User response: No action is required if the object was intentionally dropped. Otherwise, look for prior messages to determine what caused the object to be missing.

ASN7735E *program_name : program_identifier :* **The command to start the Q Apply program contained both the *autostop=Y* and *term=N* parameters. The Q Apply program will stop.**

Explanation: You cannot specify *autostop=Y* when the value of the *term* parameter is N (No). You also cannot specify *term=N* when the saved value of *autostop* in

the IBMQREP_APPLYPARMS table is Y.

User response: Restart the Q Apply program while specifying either the autostop parameter or the term=N parameter. Update the value of autostop in the IBMQREP_APPLYPARMS table if necessary.

ASN7737E *program_name : program_identifier :* **The Q subscription name (receive queue queue_name, replication queue map queue_map_name) does not have a valid model queue defined. The Q subscription will be deactivated. Reason code: reason_code**

Explanation: The model queue that is associated with the Q subscription is not valid or was defined incorrectly on WebSphere MQ. The Q Apply program could not activate the Q subscription. Model queues can be specified in the MODELQ column of the IBMQREP_TARGETS table (Q subscription level) or with the dftmodelq initialization parameter (Q Apply instance level) The following values are valid for the reason code:

0

The model queue name that was specified does not exist in WebSphere MQ.

1

The model queue was not defined on WebSphere MQ as a permanent dynamic type of queue.

User response: Review the reason codes in the explanation, and take the appropriate action:

0

Define the model queue on WebSphere MQ and activate the Q subscription.

1

Change the model queue definition on WebSphere MQ to include permanent dynamic type by using the DEFTYPE (PERMDYN) keywords, then activate the Q subscription.

ASN7740E *program_name : program_identifier :* **The Q Apply program could not activate the Q subscription name (receive queue queue_name, replication queue map queue_map_name). Reason code: reason_code.**

Explanation: The specified table-level Q subscription is based on a schema-level Q subscription that is defined at the source database or subsystem. The Q Apply program encountered errors while trying to activate this table-level Q subscription. Further changes will be ignored for this Q subscription. The following values are valid for the reason code:

0

The value of the STATE column in the IBMQREP_TARGETS control table is not I or U.

1

The target is a nickname. Nickname targets are not supported for Q subscriptions that are created based on schema-level Q subscriptions.

2

The target table does not have a unique index or primary key that corresponds to the replication key that was chosen by Q Capture and Q Apply and the attempt to create such a unique index failed.

User response: For reason codes 0 and 1, if you want to replicate this source table, use the replication administration tools to create a new Q subscription for the table. Make sure that the source and target tables exist or let the replication tools create the target table. For reason code 2, look at the Q Apply log file or IBMQREP_APPLYTRACE table for messages that might explain the reason that the unique index could not be created, try to correct the problem, and then use the replication administration tools to create a new Q subscription for the table.

ASN7741I *program_name : program_identifier :* **The Q Apply program successfully created a unique index index_name on the target table table_owner.table_name. The index is for Q subscription name, and that Q subscription is based on schema-level subscription name (receive queue queue_name, replication queue map queue_map_name). The index was created because there was no primary key or unique index on the target table that matched the replication key of the Q subscription.**

Explanation: Q Apply successfully created a unique index on the target table that is based on the replication key of the Q subscription. For a Q subscription that is based on a schema-level subscription, if the target table is missing a unique index or primary key that matches the replication key that was chosen by Q Capture, then Q Apply creates a unique index for the columns of the replication key before it activates the Q subscription.

User response: This message is for your information only. No action is required.

ASN7742E *program_name : program_identifier : A Q Apply browser thread has stopped because a row change was received for the Q subscription named name (receive queue queue_name, replication queue map queue_map_name) to a stored procedure target that uses an invalid data type.*

Explanation: Databases with the DB2 pureScale Feature use 16-byte log sequence numbers (LSN) to record the number of each pureScale member. If the source database uses the DB2 pureScale Feature, you can continue to use stored procedure targets if the following conditions are met:

- You change the stored procedure target to use the VARCHAR(16) data type for the SRC_COMMIT_LSN parameter.
- The Q Capture compatibility parameter is set to 0907. In this case, Q Capture continues to use 10-byte LSNs and the CHAR(10) data type is supported for the SRC_COMMIT_LSN parameter.

User response: Change the stored procedure target to use the VARCHAR(16) data type for the SRC_COMMIT_LSN parameter, or change the value for compatibility to 0907.

ASN7748W *program_name : program_identifier : The DDL operation operation that was sent on receive queue queue_name (replication queue map queue_map_name) was applied, but the operation caused an SQL warning. The SQLCODE is sqlcode. The SQLSTATE is sqlstate. The SQLERRMC is sql_tokens. The SQL statement is SQL_statement. The program continues to process the receive queue.*

Explanation: The Q Apply program performed the DDL operation that was replicated from the source, but a SQL warning was reported. The program continues to process the receive queue.

User response: Review the reported SQLSTATE and SQLCODE information. To address the SQL warning, perform any corrective steps that are required or suggested in the message or DB2 documentation.

ASN7749W *program_name : program_identifier : The program is unable to perform an insert operation on the event table schema.IBMQREP_APPEVENTS. The SQL code is SQLCODE. The events that were generated for this interval will not be added to the table.*

Explanation: The event thread failed with an unexpected SQL code. The generated events for this interval are not inserted into the IBMQREP_APPEVENTS table. The program tries again

at the next interval. This error does not cause the program to terminate.

User response: If this SQL code indicates a temporary error, then no action is required. Otherwise, take action as indicated for the SQL error.

ASN7750W *program_name : program_identifier : The program is unable to send EIF messages to the event server (host:port). The subsequent EIF messages for this interval will not be sent to the server.*

Explanation: The TCP/IP connection to the event server failed. The generated EIF messages for this interval are not sent, and the program tries again at the next interval. This error does not cause the program to terminate.

User response: Review the most recently issued ASN0686W, ASN0687W, ASN0688W, or ASN0689W messages to determine the cause of the error.

ASN7753I *program_name : program_identifier : The Q Apply program is now ready to generate events. Events will be sent to the following destinations: destinations.*

Explanation: The Q Apply program created a separate thread for generating events that are defined in the IBMQREP_APPEVTDEFS table. You can direct events to the console, IBMQREP_APPEVENTS table, and Event Interface Facility (EIF) servers. The EIF servers can be specified by using two different mechanisms:

- The eif_conn1 and eif_conn2 parameters of the asnqapp command
- The eif_conn1 and eif_conn2 columns of the IBMQREP_APPLYPARMS table

User response: This message is for your information only. No action is required.

ASN7754E *program_name : program_identifier : The program did not recognize the event type event_type that was inserted into the schema.IBMQREP_APPEVTDEFS table.*

Explanation: The IBMQREP_APPEVTDEFS table stores definition for events that are generated by Q Apply. The value that was provided for the event definition is not valid. Valid values include AVERAGE_LATENCY, LONG_COMMITTED_TRANS, HEARTBEAT, QUEUE_UP, QUEUE_DOWN, and PIT_CONSISTENCY_LATENCY.

User response: Insert a valid event type and reissue the task.

ASN7755E *program_name* : *program_identifier* : The value *input_value* that was inserted into the *schema*.IBMQREP_APPEVTDEFS table is invalid for the row whose primary key is (WORKLOADNAME *workload_name*, QMAPNAME *queue_map_name*, TYPE type). Reason code: *reason_code*.

Explanation: The following values are valid for the reason code:

1

When TYPE is either AVERAGE_LATENCY or PIT_CONSISTENCY_LATENCY, the value of RESET (the reset average latency threshold) must be less than 80 percent of MAX (maximum average latency threshold).

2

When TYPE is AVERAGE_LATENCY, the value of MAX (the maximum average latency threshold) must be greater than or equal to 250 milliseconds.

3

When TYPE is AVERAGE_LATENCY, the value of RESET (the reset average latency threshold) must be greater than or equal to 100 milliseconds.

4

When TYPE is LONG_COMMITTED_TRANS, the value of MAX (the maximum threshold for long-committed transactions) must be greater than or equal to 251 milliseconds.

5

When TYPE is PIT_CONSISTENCY_LATENCY, the value of MAX (the maximum threshold for point-in-time consistency latency) must be greater than or equal to 1500 milliseconds.

6

When TYPE is PIT_CONSISTENCY_LATENCY, the value of RESET (the reset threshold for point-in-time consistency latency) must be greater than or equal to 800 milliseconds.

7

For a specific replication group, the average maximum latency threshold (the value of MAX when TYPE is AVERAGE_LATENCY) must be greater than or equal to the long committed transaction latency threshold (the value of MAX when TYPE is LONG_COMMITTED_TRANS). All of the

event types for a replication group have identical WORKLOADNAME and QMAPNAME.

8

The valid values of EIF_WORKLOADTYPE are Active/Standby and Active/Query.

9

When EIF is Y in the IBMQREP_APPEVTDEFS table, a value must be provided for EIF_CONN1 or EIF_CONN2 in the IBMQREP_APPLYPARMS table.

10

The value of the EIF column in the IBMQREP_APPEVTDEFS table is Y, but the specified values for the parameters EIF_CONN1 and EIF_CONN2 in the IBMQREP_APPLYPARMS table are invalid. The value for EIF_CONN1 and EIF_CONN2 must be either the host name or the IPv4 address of the EIF server with the port number in parentheses. For example, 9.123.59.123(22008) or stlmvs1.svl.ibm.com(25009).

11

The valid values of the EIF, TABLE and CONSOLE column are N and Y.

12

The value of the CONSOLE column must be N when the value of the TYPE column is neither AVERAGE_LATENCY nor PIT_CONSISTENCY_LATENCY.

13

The value of the EIF column must be N when the value of the TYPE column is PIT_CONSISTENCY_LATENCY.

14

The value of the EIF column must be N when the value of the TYPE column is HEARTBEAT.

15

When TYPE is LONG_COMMITTED_TRANS, the column value of RESET must be NULL.

16

When TYPE is either AVERAGE_LATENCY or PIT_CONSISTENCY_LATENCY, the value of the RESET column must be less than 80 percent of the value of the MAX column.

17

When the EIF_WORKLOADTYPE is Active/Standby and MCG is enabled, the value of column WORKLOADNAME must be

equivalent to the corresponding MCGNAME, which is defined in the IBMQREP_RECVQUEUES table.

18

When the EIF_WORKLOADTYPE is Active/Query and MCG is enabled, the value of column WORKLOADNAME must be not identical to the corresponding MCGNAME, which is defined in the IBMQREP_RECVQUEUES table.

User response: Verify that the information that you entered is correct, and reissue the task.

ASN7756I *program_name : program_identifier :* **The Q Apply event thread is waiting for all browser threads to stop before it issues the QUEUE_DOWN event.**

Explanation: The QUEUE_DOWN event specifies that Q Apply has stopped processing messages on a receive queue. Before the event is generated, the event thread waits until all browser threads have stopped.

User response: This message is for your information only. No action is required.

ASN7757E *program_name : program_identifier :* **An initial attempt to connect to the primary or backup event server (host:port) failed. The Q Apply program stopped. Reason code: reason_code.**

Explanation: If the value of the EIF column in the IBMQREP_APPLYEVTDEFS table is Y, connection information must be specified for either a primary or backup event server. These servers are the destination for generated EIF messages. When Q Apply started, it tried to connect to an event server but was unsuccessful. The following values are valid for the reason code:

1

A TCP/IP GetAddrInfo socket call failed.

2

A TCP/IP create socket call failed.

3

A TCP/IP SetSockOpt socket call failed.

4

A TCP/IP connect socket call failed.

5

A TCP/IP send socket call failed.

User response: If you do not need to generate EIF events, change the value of EIF to N and restart the Q Apply program. Otherwise, take one of the following

actions depending on the reason code and then restart Q Apply:

1

Make sure that the EIF_CONN1 and EIF_CONN2 columns in the IBMQREP_APPLYPARMS table specify valid host names or IP addresses, service names, or ports. If the host name is correct, then TCP/IP is unable to resolve the host name to an IP address. Make changes that are needed to allow for proper resolution of host name to IP address. If the changes have been made to the name server or local host table, refresh the resolver.

2

Review the most recently issued ASN0686W message to determine the cause of the error.

3

Review the most recently issued ASN0687W message to determine the cause of the error.

4

Review the most recently issued ASN0688W message to determine the cause of the error.

5

Review the most recently issued ASN0689W message to determine the cause of the error.

ASN7758W *program_name : program_identifier :* **The connection to the primary or backup event server (host:port) was lost.**

Explanation: If the value of the EIF column in the IBMQREP_APPLYEVTDEFS table is Y, connection information must be specified for either a primary or backup event server. These servers are the destination for generated EIF messages. The existing connection to the event servers was lost. Q Apply will try to reconnect.

User response: Make sure that the TCP/IP protocol is working on the specified event servers.

ASN7759I *program_name : program_identifier :* **The program connected to the primary or backup event server (host:port).**

Explanation: Q Apply connects to a primary or backup event server to send EIF messages.

User response: This message is for your information only. No action is required.

ASN7761E *program_name : program_identifier :* **The Q Apply program detects on receive queue `queue_name` (replication queue map `queue_map_name`) that the Q Capture program is using parallel send queues. The `parallel_sendqs` option for the queue map must be set to y (yes) in this configuration. Q Apply stopped message processing on the receive queue.**

Explanation: Send queues are said to be parallel when multiple send queues that each use a different transmission queue all route to the same receive queue. When a Q Capture program uses parallel send queues, the replication queue map option `parallel_sendqs` must be set to y (yes) to handle messages that do not arrive in the receive queue in dense message order.

User response: Use one of the replication administration tools to change the value of the replication queue map option `parallel_sendqs` to y and then restart message processing on the receive queue.

ASN7762E *program_name : program_identifier :* **The setting for the Q Apply `prune_method` parameter of `setting` is not compatible with the `parallel_sendqs=y` option for replication queue map `queue_map_name` (receive queue `queue_name`).**

Explanation: When a receive queue is running with the queue map option `parallel_sendqs=y`, the Q Apply instance-level parameter `prune_method` must be set to 2.

User response: Restart the Q Apply program with `prune_method=2`, and restart message processing on the receive queue.

ASN7763I *program_name : program_identifier :* **The `reinitq` command was processed successfully. The Q Apply program will begin using the event definition that was inserted into the `schema.IBMQREP_APPEVTDEFS` table for the event type `type` in the workload `workload_name` of replication queue map `queue_map_name`.**

Explanation: The event definition was updated for a particular event type in the workload of the queue map. The `reinitq` command prompts Q Apply to begin using the updated definition.

User response: This message is for your information only. No action is required.

ASN7764W *program_name : program_identifier :* **The value `input_value` that was inserted into the `schema.IBMQREP_APPEVTDEFS` table is invalid for the row whose primary key is {`WORKLOADNAME workload_name`, `QMAPNAME queue_map_name`, `TYPE type`}. Reason code: `reason_code`.**

Explanation: The following values are valid for the reason code:

- 1
When TYPE is either AVERAGE_LATENCY or PIT_CONSISTENCY_LATENCY, the value of RESET (the reset average latency threshold) must be less than 80 percent of MAX (maximum average latency threshold).
- 2
When TYPE is AVERAGE_LATENCY, the value of MAX (the maximum average latency threshold) must be greater than or equal to 250 milliseconds.
- 3
When TYPE is AVERAGE_LATENCY, the value of RESET (the reset average latency threshold) must be greater than or equal to 100 milliseconds.
- 4
When TYPE is LONG_COMMITTED_TRANS, the value of MAX (the maximum threshold for long-committed transactions) must be greater than or equal to 251 milliseconds.
- 5
When TYPE is PIT_CONSISTENCY_LATENCY, the value of MAX (the maximum threshold for point-in-time consistency latency) must be greater than or equal to 1500 milliseconds.
- 6
When TYPE is PIT_CONSISTENCY_LATENCY, the value of RESET (the reset threshold for point-in-time consistency latency) must be greater than or equal to 800 milliseconds.
- 7
For a specific replication group, the average maximum latency threshold (the value of MAX when TYPE is AVERAGE_LATENCY) must be greater than or equal to the long committed transaction latency threshold (the value of MAX when TYPE is LONG_COMMITTED_TRANS). All of the

| | | |
|----|--|---|
| 8 | event types for a replication group have identical WORKLOADNAME and QMAPNAME. | The valid TYPE values include AVERAGE_LATENCY, LONG_COMMITTED_TRANS, HEARTBEAT, QUEUE_UP, QUEUE_DOWN, and PIT_CONSISTENCY_LATENCY. |
| 9 | The valid values of EIF_WORKLOADTYPE are Active/Standby and Active/Query. | 18 |
| 10 | When EIF is Y in the IBMQREP_APPEVTDEFS table, a value must be provided for EIF_CONN1 or EIF_CONN2 in the IBMQREP_APPLYPARMS table. | 19 |
| 11 | The value of the EIF column in the IBMQREP_APPEVTDEFS table is Y, but the specified values for the parameters EIF_CONN1 and EIF_CONN2 in the IBMQREP_APPLYPARMS table are invalid. The value for EIF_CONN1 and EIF_CONN2 must be either the host name or the IPv4 address of the EIF server with the port number in parentheses. For example, 9.123.59.123(22008) or stlmvs1.svl.ibm.com(25009). | When the EIF_WORKLOADTYPE is Active/Query and MCG is enabled, the value of column WORKLOADNAME must be not identical to the corresponding MCGNAME, which is defined in the IBMQREP_RECVQUEUES table. |
| 12 | The valid values of the EIF, TABLE, and CONSOLE columns are N and Y. | User response: 1. Verify that the information that you entered is correct. If not redo the update of the IBMQREP_APPEVTDEFS table. 2. Issue the reinitq command for the receive queue that is specified in the queue map to prompt Q Apply to begin using the updated definitions. |
| 13 | The value of the CONSOLE column must be N when the value of the TYPE column is neither AVERAGE_LATENCY nor PIT_CONSISTENCY_LATENCY. | <hr/> ASN7765W <i>program_name : program_identifier : The reinitq command was issued after an update to the event definitions for the receive queue queue_name, replication queue map queue_map_name. However, none of the event definitions in the IBMQREP_APPEVTDEFS table were changed.</i> |
| 14 | The value of the EIF column must be N when the value of the TYPE column is PIT_CONSISTENCY_LATENCY. | Explanation: A reinitq command was issued to prompt the Q Apply program to begin using updated event definitions for the receive queue, but the definitions that were updated in the IBMQREP_APPEVTDEFS table were the same as the existing values. |
| 15 | The values of the EIF column must be N when the value of the TYPE column is HEARTBEAT. | User response: If you want to change the related event definitions from the current values, follow these steps: 1. Redo the update of the IBMQREP_APPEVTDEFS table, making sure the new values are included. 2. Issue the reinitq command for the receive queue that is specified in the queue map to prompt Q Apply to begin using the updated definitions. |
| 16 | When the value of the TYPE column is LONG_COMMITTED_TRANS, the value of the RESET column must be NULL. | |
| 17 | When the value of the TYPE column is either AVERAGE_LATENCY or PIT_CONSISTENCY_LATENCY, the value of the RESET column must be less than 80 percent of the value of the MAX column. | |

Part 5. AUD Messages

Chapter 47. AUD0000 - AUD0499

| | |
|----------|----------------------|
| AUD0000I | Operation succeeded. |
|----------|----------------------|

| | |
|----------|-------------------|
| AUD0001N | Operation failed. |
|----------|-------------------|

| | |
|----------|----------------------|
| AUD0002N | Syntax error. Usage: |
|----------|----------------------|

| | |
|----------|--|
| AUD0003N | I/O error on accessing <i>directory/filename</i> , make sure the directory/file exists and has the right permission. |
|----------|--|

| | |
|----------|--|
| AUD0004N | File system full, can't write to <i>filename</i> . |
|----------|--|

| | |
|----------|------------------------|
| AUD0016N | Invalid database name. |
|----------|------------------------|

| | |
|----------|--|
| AUD0017N | Can not read audit configuration file. |
|----------|--|

| | |
|----------|--|
| AUD0018N | An error occurred when updating the audit settings for the instance <i>sqlcode</i> . Check the <i>sqlcode</i> 's meaning for more information. |
|----------|--|

| | |
|----------|--|
| AUD0019N | System error occurred. Retry the command, if still failing, take a trace and then contact IBM support. |
|----------|--|

| | |
|----------|---|
| AUD0020N | Error occurred during pruning. <i>db2audit.log</i> has been truncated to size 0, the contents originally in <i>db2audit.log</i> can be found in <i>db2audit.cpy</i> |
|----------|---|

| | |
|----------|---|
| AUD0021N | <i>filename</i> already exists. It could be generated by previous <i>db2audit</i> operation. Make sure that it does not contain any information you need, remove it and then retry the command. |
|----------|---|

| | |
|----------|----------------------------------|
| AUD0022N | The Audit Log File is corrupted. |
|----------|----------------------------------|

| | |
|----------|---|
| AUD0023N | An error occurred while auditing the requested event, <i>sqlcode</i> <i>sqlcode</i> . Check the <i>sqlcode</i> 's meaning for more information. |
|----------|---|

| | |
|----------|--|
| AUD0026I | A request to start the DB2 audit facility has been processed. Note that audit may have already been started on the instance. |
|----------|--|

| | |
|----------|---|
| AUD0027I | A request to stop the DB2 audit facility has been processed. Note that audit may have already been stopped on the instance. |
|----------|---|

| | |
|----------|---|
| AUD0028N | An invalid ASCII delimiter was specified for the extract. This delimiter must either be a single character, or a 4 byte string representing a valid hexadecimal character value such as 0xff. |
|----------|---|

| | |
|----------|--|
| AUD0029N | The NODE option was not specified on the archive option, but the database server is not running. Either indicate which node the archive is to take place on, or issue a <i>db2start</i> so the archive will take place on all nodes. |
|----------|--|

| | |
|----------|---|
| AUD0035N | The active audit log <i>directory-or-filename</i> has already been archived, and new events cannot be logged to it. |
|----------|---|

| | |
|----------|--|
| AUD0036N | Extract can not be performed on the active audit log file <i>directory-or-filename</i> . |
|----------|--|

Part 6. CCA Messages

This section contains the Configuration Assistant Interface (CCA) messages. The messages are listed in numeric sequence.

Chapter 48. CCA1000 - CCA1499

| | |
|----------|--|
| CCA1001I | Use alphanumeric characters only. The first character cannot be numeric. |
|----------|--|

| | |
|----------|------------------------------------|
| CCA1002I | Use decimal numbers, 0 to 9, only. |
|----------|------------------------------------|

| | |
|----------|---|
| CCA1003I | No DB2 system was found on the network. |
|----------|---|

| | |
|----------|--|
| CCA1004I | No DB2 instance was found on the selected system(s). |
|----------|--|

| | |
|----------|--|
| CCA1005I | No DB2 database was found on the selected instance(s). |
|----------|--|

Chapter 49. CCA2000 - CCA2499

CCA2001W No files were found matching the file pattern specified.

Explanation: A request was made to read from files matching the specified file pattern. No files were found matching the pattern.

User response: Correct the file pattern specified and retry the operation.

CCA2002W An error was encountered while updating the TCP/IP services file.

Explanation: An attempt to add a service name and port number to the TCP/IP services file failed or Network Information Services is being used and only the local services file was updated. The port number was used to catalog the node instead of the service name.

User response: To use the service name instead of the port number in the node directory entry, the node must be manually uncataloged and then recataloged using the service name. The services file must also be updated manually. If Network Information Services (NIS) is being used, the local services file may have been updated but the NIS server must be updated manually. In this case, the node was also cataloged using the port number.

CCA2003W The Discover request did not return data for one or more DB2 systems.

Explanation: The discover request did not return data for one or more of the requested DB2 systems. One of the following may have occurred:

- The Administration Server to which the discover request was sent has not been started.
- An error occurred attempting to perform the discover request.
- The DB2 system to which the discover request was sent is not configured for discovery.

User response: Verify that the DB2 system to which the discover request was sent is enabled for discovery. If enabled for discovery, ensure that the Administration Server on the DB2 system is running.

CCA2004W The nname value specified is not unique.

Explanation: The nname value specified is already being used by another NetBIOS application on the network.

User response: Select 'YES' to use the specified nname

or 'NO' to cancel the request. If 'YES' is selected, any application using the existing nname will be affected.

CCA2005W The socket number specified is not unique.

Explanation: The socket number specified is being used by another DB2 instance on the workstation.

User response: Select 'YES' to use the specified socket or 'NO' to cancel the request. If 'YES' is selected, any application using the existing socket number will be affected.

CCA2006W An entry already exists in the TCP/IP services file for the service name and port number specified.

Explanation: An entry already exists in the TCP/IP services file for the service name and port number specified. Another application may be using the entry.

User response: Select 'YES' to use the existing entry or 'NO' to cancel the request. If 'YES' is selected, any application using the existing entry will be affected.

CCA2007W The port number specified is being used with a different service name.

Explanation: The TCP/IP services file contains an entry that uses the port number specified but the associated service name does not match the service name specified.

User response: Select 'YES' to use the specified service name and port number or 'NO' to cancel the request. If 'YES' is selected, a new entry will be added to the services file. Any applications using the entry with the existing port number may be affected.

CCA2008W The service name specified is being used with a different port number.

Explanation: The TCP/IP services file contains an entry that uses the service name specified but the associated port number does not match the port number specified.

User response: Select 'YES' to use the specified service name and port number or 'NO' to cancel the request. If 'YES' is selected, the existing entry in the services file, using the service name, will be updated to use the port number specified. This may affect any applications using the existing entry.

CCA2009W The request was canceled by the user.

Explanation: The request was canceled by the user.

User response: None.

CCA2010W An attempt to update the APPC stack failed.

Explanation: An attempt to add the Transaction Program name to the APPC stack failed.

User response: The APPC stack must be manually updated. If the Transaction Program name is not added to the stack, remote connections to the server will not be possible.

CCA2011W An attempt to add the service name and port number to the TCP/IP services file failed.

Explanation: An attempt to add the service name and port number to the TCP/IP services file failed. The database manager configuration file has been updated with the service name specified.

User response: The service name and port number must be manually added to the TCP/IP service file. If the entry is not added to the services file, remote connections to the server will not be possible.

CCA2012W The Discover request did not find any DB2 systems.

Explanation: A request to search the network for DB2 systems completed successfully but no DB2 systems were found. Following is a list of possible reasons why no DB2 systems were found:

- Search discovery was not enabled on any DB2 systems (that is, in the DBM configuration file of the administration server on the DB2 system, DISCOVER = SEARCH was not specified).
- The DB2 system was not set up with the appropriate discovery protocol for the client to find it (that is, DISCOVER_COMM on the administration server does not contain a protocol that matches one specified for DISCOVER_COMM on the client).
- The DB2 system is on the other side of a router or bridge, and the routers and bridges on the network were configured such that the discovery packet was filtered out and not passed on.

User response: Following is a list of possible actions to take in order for discovery to find DB2 systems:

- Set DISCOVER = SEARCH in the DBM configuration file of the administration server on all DB2 systems that you wish to be found.
- Set DISCOVER_COMM for the administration server to include the protocol that the client will use to issue a discover request (that is, set DISCOVER_COMM to include at least one of the protocols specified in DISCOVER_COMM on the client).

- Have the network administrator reconfigure the router or bridge to allow discovery packets (for the specified protocol) to be passed through.

CCA2013W Remote database cataloged using APPC but the stack was not configured.

Explanation: A request to catalog a database resulted in a node being cataloged that uses the APPC protocol. The node was cataloged using the symbolic destination name which was retrieved from the profile specified. The APPC stack was not configured because there was insufficient information in the profile to configure the stack or APPC was not detected on the DB2 system. It was not possible to use a different protocol because no other matching protocol was detected on the client.

User response: If APPC is not installed on the client, uncatalog the database and recatalog the database manually using a protocol that is available on both the client and server. If APPC is installed, configure the stack if it has not already been configured.

CCA2014W The transaction program name specified is not unique or has already been configured.

Explanation: The transaction program name specified is already being used by another DB2 instance or a non DB2 application on this server.

User response: Select 'YES' to use the specified transaction program name or 'NO' to cancel the request. If "YES" is selected, for all applications using the transaction program name concurrently, APPC will only be operational for the first one started. If any new APPC parameter values have been specified, the APPC stack will be updated with these values.

CCA2015W The service name and port number specified are being used in different entries in the services file.

Explanation: The TCP/IP services file contains entries that use the service name and port number specified but they are not being used in the same entry.

User response: Select 'YES' to use the specified service name and port number or 'NO' to cancel the request. If 'YES' is selected, the existing entry in the services file, using the service name, will be updated to use the port number specified. This may affect any applications using the existing entries.

CCA2016W The password will be saved as clear text.

Explanation: The password will be saved as clear text in db2cli.ini file.

User response: If password security is a concern, deselect the 'Save password' check box.

Chapter 50. CCA3000 - CCA3499

CCA3000C **An internal error has occurred. Reason code** *reason-code*.

Explanation: An unexpected internal error has occurred.

User response: Turn trace on and retry the steps that caused the error. If the problem reoccurs, save the trace information to a file and contact IBM Support with the following information:

- Problem description
- Message number
- Reason code
- Trace file

CCA3001N **The specified service name and port number conflicts with existing values in the TCP/IP services file.**

Explanation: The service name and port number entered by the user conflicts with existing values in the TCP/IP services file. The service name may already be used with a different port number, the port number may already be used with a different service name or both.

User response: Specify a service name and port number that does not conflict with existing entries in the services file.

CCA3002N **An I/O error occurred.**

Explanation: An error was encountered while attempting to open, read, change the file position or close a file.

User response: If a file name was specified, verify that the file name is valid and that the user has permission to access the file. Also check for any disk and operating system errors.

CCA3003N **The format of the file is not valid.**

Explanation: An error was encountered while reading from a file. The format of the file is not valid. Possible errors include:

- The file contains invalid data.
- The file does not contain expected data.
- The order of the data in the file is incorrect.

User response: If a file name was specified, and the file has been modified by the user, regenerate the file and retry the operation. If the problem persists, and the file was not modified by the user, or the problem occurred during a Discover request, turn trace on and

retry the steps that caused the error. If the problem reoccurs, save the trace information to a file and contact IBM Support with the following information:

- Problem description
- Message number
- Trace file
- File which is causing the error if a file name was specified

CCA3004N **An attempt to allocate memory failed.**

Explanation: An error was detected while attempting to allocate memory.

User response: Terminate other applications running on the system that may be using large amounts of memory. If the problem persists, turn trace on and retry the operation. If the problem reoccurs, save the trace information to a file and contact IBM Support with the following information:

- Problem description
- Message number
- Trace file

CCA3005N **An error was encountered while writing to a file.**

Explanation: An error was detected while writing to a profile. The error could also be encountered when updating a host system password and errors are being recorded in the file db2pem.log.

User response: Verify that the file system on which the file resides is not full and is not damaged. Also check for any operating system errors.

CCA3006N **No matching communication protocol was detected.**

Explanation: The database cannot be cataloged because none of the protocols available on the client match any of the protocols available at the server.

User response: Ensure that the client and server have at least one matching communication protocol that can be detected on both. If a matching protocol is installed on both the client and the server, the protocol could not be detected. In this case, catalog the database and node manually.

CCA3007N **The database alias name specified is not valid.**

Explanation: The length of the database alias specified is not a valid or the alias contains invalid characters.

User response: Correct the alias name and resubmit the request.

CCA3009N The application requestor name specified is not valid.

Explanation: The length of the application requestor name specified is not valid or the name contains invalid characters.

User response: Correct the application requestor and resubmit the request.

CCA3010N The length of the parameter value specified is not valid.

Explanation: The length of the parameter value specified for the application requestor is not valid.

User response: Correct the parameter value and resubmit the request.

CCA3011N The target database name specified is not valid.

Explanation: The length of the target database name specified is not valid or the name contains invalid characters.

User response: Correct the target database name and resubmit the request.

CCA3012N Add ODBC data source failed.

Explanation: A request to add an ODBC data source failed. The error could be caused by an out of memory error, a disk full condition or a disk failure.

User response: Verify that the disk on which the ODBC.INI and DB2CLI.INI files reside is not full and that the disk is not damaged. In addition, if other applications are using large amounts of memory, terminate the applications and retry the operation.

CCA3013N Remove ODBC data source failed.

Explanation: A request to remove an ODBC data source failed. The error could be caused by an out of memory condition or a disk failure.

User response: If other applications are using large amounts of memory, terminate the applications and retry the operation. Also verify that the disk on which the ODBC.INI and DB2CLI.INI files reside is not damaged.

CCA3014N The bind request cannot be processed.

Explanation: The bind request cannot be processed because another bind operation is already in progress.

User response: Complete or terminate the bind in progress and resubmit the bind request.

CCA3015N The adapter specified is not valid.

Explanation: The adapter specified was not detected on the DB2 system.

User response: Specify an adapter that is available and resubmit the request.

CCA3016N The nname value specified is not unique.

Explanation: The nname value specified is already being used by another NetBIOS application on the network.

User response: Specify a unique nname and retry the operation.

CCA3017N The path specified for the file is not valid.

Explanation: An attempt was made to open the specified file, but the path specified is invalid or does not exist.

User response: Ensure that the path specified is valid and the path for the file exists.

CCA3018N The user does not have sufficient authority to access the file.

Explanation: An attempt was made to access the requested file, but the user does not have the required authority to access the file.

User response: Ensure that the user has the required authority to access the file.

CCA3019N The file name specified is a directory.

Explanation: An attempt to access the file specified failed because the name specified is a directory and not a file.

User response: Specify a valid file name and retry the operation.

CCA3020N An attempt to access the specified file failed because of a share violation.

Explanation: An attempt to access the file specified failed because of a share violation. Another process may have the file opened in exclusive mode.

User response: The file is currently being accessed by another process in exclusive mode. Ensure that no other process is accessing the file and retry the operation or specify another file name.

CCA3021N An attempt to retrieve, add or remove variable *variable-name* from the DB2 Profile Registry failed with Return Code *return-code*.

Explanation: An attempt to retrieve, add or remove the indicated variable from the DB2 Profile Registry failed. The return code indicates the cause of the problem. The possible return codes are as follows:

- -2 The specified parameter is invalid
- -3 Insufficient memory to process the request
- -4 Variable not found in the registry
- -7 DB2 Profile Registry is not found on this DB2 system
- -8 Profile not found for the given instance
- -9 Profile not found for the given node
- -10 UNIX registry file lock time-out

User response: For return code:

- -2 Ensure that the parameter has been specified correctly.
- -3 Terminate other applications using large amounts of memory and retry the operation.
- -4 Ensure that the variable is set in the DB2 Profile Registry.
- -7 Ensure that the DB2 Profile Registry has been created.
- -8 Ensure that the profile has been created for the instance.
- -9 Ensure that the profile has been created for the node.
- -10 Ensure that the registry file is not locked by another process.

CCA3022C An attempt to retrieve the address of the function *procedure-name* from library *library-name* failed with Return code *return-code*.

Explanation: An attempt to retrieve the address of a function from the indicated library failed.

User response: Verify that the correct version of the library is being used. If the incorrect version is being used, install the correct version. If the problem persists, turn trace on and retry the steps that caused the error. If the problem reoccurs, save the trace information to a file and contact IBM Support with the following information:

- Problem description
- Message number
- Return code
- Trace file

CCA3023C An attempt to load library *library-name* failed with Return code *return-code*.

Explanation: An attempt to load the indicated library failed.

User response: Verify that the path in which the library resides is included in the library path. Also ensure that there is enough memory available to load the library. If the problem persists, turn trace on and retry the steps that caused the error. If the problem reoccurs, save the trace information to a file and contact IBM Support with the following information:

- Problem description
- Message number
- Return code
- Trace file

CCA3024C An attempt to unload library *library-name* failed with Return code *return-code*.

Explanation: An attempt to unload the indicated library failed.

User response: Turn trace on and retry the steps that caused the internal error. If the problem reoccurs, save the trace information to a file and contact IBM Support with the following information:

- Problem description
- Message number
- Return code
- Trace file

CCA3025N One or more of the IPX/SPX parameters specified is not valid.

Explanation: One or more of the input parameters is not valid. Following is a list of possible errors:

- One or more of the fileserver, objectname and ipx_socket parameters is NULL.
- Only the fileserver parameter or objectname parameter is set to "*".
- The fileserver and/or objectname parameters are not set to "*" for Windows and Solaris.
- The objectname value specified is not unique.
- The ipx_socket value specified is not unique.
- The ipx_socket value specified is not in the valid range.
- The DB2 system failed to attach to the fileserver specified.

User response: Verify the following:

- The fileserver, objectname and ipx_socket parameters are not NULL.
- If the value specified for fileserver is "*", the value for objectname must also be "*".

- On Windows and Solaris, both fileserver and objectname must be "*".
- The value specified for objectname, if not "*", is unique for all DB2 instances and IPX/SPX applications registered at the fileserver.
- The value specified for ipx_socket is unique across all DB2 instances on the DB2 system.
- The value specified for ipx_socket is within the valid range.
- The fileserver specified exists and that it is up and running.

Correct all errors and retry the operation.

CCA3026N No available NetBIOS adapters detected.

Explanation: A NetBIOS adapter was not detected on the DB2 system. The database cannot be cataloged.

User response: Catalog the database and node manually if an adapter is available on the DB2 system.

CCA3027N The port number specified is out of range.

Explanation: The TCP/IP port number specified is out of range. The maximum value that can be specified for the port number is 65534.

User response: Specify a port number that does not exceed the maximum value and retry the operation.

CCA3028N The DB2INSTANCE variable is not valid.

Explanation: The DB2INSTANCE environment variable is not set or is set to the Administration Server instance. The Configuration Assistant cannot run under the Administration Server Instance.

User response: Set the DB2INSTANCE variable to an instance other than the Administration Server instance.

CCA3029N Updating of the ODBC data source settings failed.

Explanation: A request to update the settings of an ODBC data source failed. The error could be caused by an out of memory error, a disk full condition or a disk failure.

User response: Verify that the disk on which the DB2CLI.INI file resides is not full and that the disk is not damaged. In addition, if other applications are using large amounts of memory, terminate the applications and retry the operation.

CCA3030N Values missing for configuring APPC.

Explanation: A request was made to catalog a database using APPC or to configure a server instance for APPC. The request could not be completed because one or more parameters were not specified.

User response: Ensure that all required parameters have been specified and retry the operation.

CCA3031N The APPC stack has not been configured for the database selected.

Explanation: The database selected is using APPC for the database connection. However, the APPC stack has not been configured for the connection.

User response: Configure the APPC stack for the database selected.

CCA3051N A protocol protocol interface failure has occurred with Return code *return-code*.

Explanation: The failure occurred while attempting to access the protocol interface.

User response: Verify that the protocol is operational.

CCA3052N The specified item *item-name* was not found.

Explanation: The specified item name could not be found in the configuration data.

User response: Verify that you have specified the item name correctly.

CCA3053N The ODBC DSN specified *DSN-name* is invalid.

Explanation: The ODBC DSN specified is an invalid name.

User response: Ensure that you are using valid characters for the ODBC DSN name.

CCA3054N The ODBC DSN could not be registered.

Explanation: The ODBC DSN registration attempt failed.

User response: Verify that ODBC is installed properly and is functional.

CCA3055N The specified item *item-name* already exists.

Explanation: The specified item name already exists in the configuration data.

User response: Verify that you have specified the item name correctly. Use a different item name or delete the old item and resubmit the request.

CCA3056N The hostname *host-name* was not found.

Explanation: The specified hostname could not be resolved on the network.

User response: Ensure that the hostname, as specified, is correct and is a valid hostname on your network.

CCA3057N The service name *service-name* was not found.

Explanation: The specified service name was not found in the local services file.

User response: Ensure that the service name, as specified, is correct and that there is a valid entry for that service name in your local services file.

CCA3058N The local system object is not allowed to be removed.

Explanation: The local system object appears if this is a server installation, and does not appear if this is a client installation. You have no direct control over the removal of this object because it has special properties that are required by the server installation type.

User response: No action is required.

CCA3059N The local system object is not allowed to be changed.

Explanation: The local system object appears if this is a server installation, and does not appear if this is a client installation. You have no direct control over the changing of this object because it has special properties that are required by the server installation type.

User response: No action is required.

CCA3060N The selected system object is not allowed to be changed.

Explanation: The selected system is using communication protocol that is not supported by the DB2 Administration Tools.

User response: No action is required.

CCA3061N Incomplete server configuration.

Explanation: The server configuration information contained in the server profile is missing data necessary to complete the requested operation. Refer to the db2diag log file for details.

User response: Contact your system administrator to verify that the server configuration is correct.

CCA3062N "Common" is a reserved data source name.

Explanation: "Common" is a reserved data source name by DB2 CLI.

User response: Re-enter another data source name.

CCA3063N The specified service name and port number conflicts with existing values in the TCP/IP services file. Do you wish to overwrite the existing values in the services file?

Explanation: The service name and port number entered by the user conflicts with existing values in the TCP/IP services file. The service name may already be used with a different port number, the port number may already be used with a different service name or both.

User response: Click Yes to overwrite the existing values in the services file with the new values.

Click No to cancel the action and keep the existing values in the services file.

CCA3064N The new data source name specified already exists.

Explanation: A data source entry already exists with the same name and its content does not match the specification of the new entry. Therefore, it cannot be reused.

User response: Use a different data source name.

CCA3065N The requested operation is not available in offline (OFFLINE) mode.

Explanation: An operation or task has been requested which is not valid or not applicable while the mode is offline (OFFLINE). The operation cannot proceed.

User response: Change the mode from offline (OFFLINE) mode before retrying this operation.

CCA3066N The requested operation is not available in remote (REMOTE) mode.

Explanation: An operation or task has been requested which is not valid or not applicable while the mode is remote (REMOTE). The operation cannot proceed.

User response: Change the mode from remote (REMOTE) mode before retrying this operation.

Chapter 51. CCA5000 - CCA5499

CCA5000N The user ID specified is not valid.

Explanation: The user ID specified does not exist.

User response: Enter the correct user ID and resubmit the request.

CCA5001N The password specified is incorrect.

Explanation: The password specified for the user ID is incorrect.

User response: Enter the correct password for the user ID and resubmit the request.

CCA5002N The password for the user ID has expired.

Explanation: The password for the user ID has expired and cannot be updated.

User response: Contact your system administrator to have the password reset.

CCA5003N The new password specified is not valid.

Explanation: The new password specified is not valid.

User response: Enter a valid password and resubmit the request.

CCA5004N An unexpected error occurred.

Explanation: An unexpected error occurred while attempting to update the password for the user ID specified. Additional information may have been written to the file db2pem.log in the instance directory.

User response: Contact your system administrator for further assistance and provide the information from the file db2pem.log.

CCA5005N The new password does not match the verify password.

Explanation: The new password does not match the verify password.

User response: Type the new password in both text boxes again.

CCA5006N The port number specified is not valid.

Explanation: The port number specified is out of range. It should be greater than zero and less than 65535.

User response: Type the new port number and retry the operation.

CCA5007N The parameter value specified is not valid.

Explanation: The parameter value specified is out of range.

User response: Type the new parameter value and retry the operation.

CCA5008N The adapter number specified is not valid.

Explanation: The adapter number specified is out of range. It should be between 0 and 255.

User response: Type the new adapter number and retry the operation.

Part 7. CIE Messages

This section contains the DB2 Text Search messages. The messages are listed in numeric sequence.

Chapter 52. CIE0000 - CIE0499

CIE0001I **Operation completed successfully.**

Explanation: No errors were encountered during the execution of this command.

User response: No action required.

CIE0002I **The DB2 Text Search catalog and text indexes are at the latest release level for the database. The system has not been upgraded.**

Explanation: The database version value in the SYSIBMTS.TSDEFAULTS administrative view is set to the latest version. The database is up to date.

User response: No action required.

CIE0003I **Index update started.**

Explanation: Update processing has started for a text search index, because either a UPDATE INDEX command was executed, or an automatic update was performed according to the update frequency specified for the index. Depending on the number of documents, processing may take some time, but the index is still available for searches.

User response: No response is required.

CIE0004I **Index update ended.**

Explanation: Update processing has ended for a text search index.

User response: No response is required.

CIE0005I **DB2 Text Search Index update statistics: Total number of documents processed: *number*. Number of documents inserted: *number*. Number of documents updated: *number*. Number of documents deleted: *number*. Number of empty documents skipped successfully: *number*. Number of documents not indexed because of internal document errors: *number*.**

Explanation: Update is complete. The statistics indicate the number of documents that were processed by DB2 Text Search.

User response: No action required.

CIE0006I **DB2 Text Search Index update statistics: Number of documents processed in this commit cycle: *number*. Number of documents inserted: *number*. Number of documents updated: *number*. Number of documents deleted: *number*. Number of empty documents skipped successfully: *number*. Number of documents not indexed because of internal document errors: *number*. Number of documents processed in total: *number*.**

Explanation: One commit cycle is complete. The statistics indicate the number of documents processed in this commit cycle.

User response: No action required.

CIE0007I **DB2 Text Search Index update statistics: Number of documents processed in total is *total-number-of-docs*. Number of documents inserted is *number-of-inserted-docs*. Number of documents updated is *number-of-updated-docs*. Number of documents deleted is *number-of-deleted-docs*. Number empty documents skipped successfully is *number-of-empty-docs*. Number of documents resubmitted for indexing due to internal document errors is *number-of-docs-resubmitted*. Updated statistics for resubmitted documents: Number of documents inserted is *number-of-resubmitted-docs-inserted*. Number of documents updated is *number-of-resubmitted-docs-updated*. Number of documents not indexed because of document errors is *number-of-resubmitted-docs-not-indexed*.**

Explanation: Update processing is complete. The statistics indicate the number of documents that were processed by DB2 Text Search.

User response: No action required.

CIE0008I **Index reorganization started.**

Explanation: DB2 Text Search index reorganization has started.

The phase of reading and parsing documents has ended and a Text Search index reorganization is necessary.

User response: No response is required.

CIE0009I Index reorganization ended.

Explanation: The DB2 Text Search index reorganization has ended.

User response: No response is required.

CIE0010I Text Search Index update statistics: Number of documents resubmitted for indexing due to internal document errors is *number-of-docs-resubmitted*. Updated statistics for resubmitted documents: Number of documents inserted is *number-of-resubmitted-docs-inserted*. Number of documents updated is *number-of-resubmitted-docs-updated*. Number of documents not indexed because of document errors is *number-of-resubmitted-docs-not-indexed*.

Explanation: Update processing is complete for the resubmitted documents. The statistics indicate the number of documents that were processed by DB2 Text Search.

User response: No action required.

CIE0011I Index build started.

Explanation: The DB2 Text Search index update build has started.

User response: No response is required.

CIE0012I Index build ended.

Explanation: The DB2 Text Search index build operation has ended.

User response: No response is required.

CIE0013I No dependent text-maintained table was found in pending mode.

Explanation: The RESET PENDING command executes a SET INTEGRITY statement for any dependent tables that manage full-text search. No such table was found.

User response: No action required.

CIE0016I The text search server for partition *partition-number* on host *host-name* and port *port-number* is in started state.

Explanation: You can determine the status of the text search server by using the db2ts START FOR TEXT command with the STATUS option. This message is returned when a connection to the text search server can be established successfully.

You can perform text search queries and text index administration operations in the database.

User response: No action is required.

CIE0017I The text search server for partition *partition-number* on host *host-name* and port *port-number* is in stopped state.

Explanation: You can determine the status of the text search server by using the db2ts START FOR TEXT command with the STATUS option. This message is returned when a connection to the text search server cannot be established successfully.

You cannot perform text search queries or text index administration operations in the database.

User response: No action required.

CIE0018I Search services are not active.

Explanation: DB2 Text Search services were found to be inactive when the db2ts START FOR TEXT command verified the status of the services.

User response: No action required.

CIE0019I Search services are active.

Explanation: DB2 Text Search services were found to be active when the db2ts START FOR TEXT command verified the status of the services.

User response: No action required.

CIE0020I DB2 Text Search index update statistics during a *time-in-minutes* minute interval: Total number of documents processed is *total-number-of-docs*. The percentage completed is *percentage-processed*%.

Explanation: The statistics indicate the number of documents that were processed in this time interval.

User response: No action required.

CIE0021I DB2 Text Search services on host *host-name* and port *port-number* are running successfully.

Explanation: A stand-alone text search server has been detected. For a stand-alone text search server setup, the initialization of DB2 Text Search server is done on a separate host or environment.

In this case, the db2ts START FOR TEXT command does not attempt to start a server but instead checks the status of the stand-alone text search server to report success.

User response: No user response is required.

CIE0090W A directory named *directory-name* already exists. richtextTool will create a symbolic link named *link-name* to that directory in order to reuse existing Outside In Technology library files.

Explanation: richtextTool will attempt to create a directory with the specified name. If that directory already exists, richtextTool will then create a symbolic link to that directory instead so that it can reuse existing Outside In Technology library files.

User response: No response is required.

CIE0091I The rich text feature has been successfully configured for use on DB2 Text Search Instance named *instance-name*.

Explanation: richtextTool has been run and has successfully enabled the rich text feature on the named DB2 Text Search instance.

User response: No response required.

CIE0092I DB2 Text Search instance server *instance-server* has been set to the default. Rich text is not supported.

Explanation: richtextTool has been run and has successfully disabled rich text support on the named DB2 Text Search instance.

DB2 Text Search is now running in default mode for the instance.

User response: No response is required.

CIE0093I Rich text support configuration has started for DB2 Text Search instance named *instance-name*.

Explanation: richtextTool has been called. Configuration for rich text support has started on the named DB2 Text Search instance.

User response: No response is required.

CIE0094I Default support configuration, no rich text, has started for DB2 Text Search instance *instance-name*.

Explanation: richtextTool has been called. Configuration for default support, no rich text, has started on the named DB2 Text Search instance.

User response: No response is required.

CIE0095I The Outside In Technology product library files have been setup for DB2 Text Search.

Explanation: The installation of the Outside In Technology product library files needs to be done prior to running richtextTool.

This only needs to be done once for each DB2 installation.

User response: For each instance that requires rich text support, login as the instance owner and run the 'richtextTool enable' command to enable rich text support for that instance.

CIE0201E Text search index *schema-name.index-name* already exists.

Explanation: The text search index that you specified already exists in this database.

User response: Check the name specified and the database you are using. Use the SYSIBMTS.TSINDEXES view to see the existing text search indexes.

CIE0202W Instance services already active.

Explanation: A db2ts start command was issued but the instance services are already running.

User response: No further action required.

CIE0203W Some instance services already active.

Explanation: The services are already active, so you do not need to start them.

User response: No response is required.

CIE0204E Instance services already started.

Explanation: An attempt was made to start the text search instance services, but the text search instance services are already running.

User response: If you are simply trying to start the text search instance services, no further action is necessary.

If there were errors using DB2 Text Search that indicated that the text search services were not running, then attempt to stop and restart using the 'db2ts start for text' command.

Check the current system status and configuration. If you continue to get errors while starting the text search services, the following steps might be necessary:

- On UNIX, use the ipcs command to check the resources. Stop all applications, such as DB2 and DB2 Text Search. If further resources are listed, clean them up using ipcrm.

- In the case of multiple DB2 instances running on the system, ensure that the DB2 Text Search for each instance has been configured for a unique communication port.

CIE0205E Instance services already stopped.

Explanation: An attempt was made to stop the text search instance services, but they are not running.

User response: If you are simply trying to stop the text search instance services, no further action is necessary.

If there were errors using DB2 Text Search that indicated that the text search services should be stopped and restarted, the following steps might be necessary:

- On UNIX, use the `ipcs` command to check the resources. Stop all applications, such as DB2 and DB2 Text Search. If further resources are listed, clean them up using `ipcrm`.
- In the case of multiple DB2 instances running on the system, ensure that the DB2 Text Search for each instance has been configured for a unique communication port.

CIE0206E Environment variable *variable-name* not set.

Explanation: A required environment variable is not set.

User response: Check your environment, specify the required variable, and reissue the command.

CIE0207E No help information available.

Explanation: No help information is available for the specified message identifier or command.

User response: Check that the message identifier was specified correctly. If the identifier was correct then there is no help information available using the 'db2ts help' or 'db2ts ?' commands. The DB2 Text Search documentation can be used for further help.

**CIE0208E DB2 Text Search Index update completed with *number* document errors.
Event view: *schema-name.event-name*.**

Explanation: Document specific errors occurred in a DB2 Text Search Index update operation. Only documents without specific errors were indexed successfully. All document errors are written to the event table.

User response: Check the event view for more information about the document errors. Correct the errors and update the affected documents. Retry the DB2 Text Search Index update operation. Clean up the event log after the problems have been fixed. In some

cases, errors may be written to the `db2diag.log`.

CIE0209I A DB2 Text Search index with identifier *text-search-id* does not exist.

Explanation: The CLEANUP FOR TEXT operation has detected and removed inconsistent DB2 Text Search Index information.

The CREATE INDEX ... FOR TEXT command adds DB2 Text Search Index information to the database catalog and creates a text index properties file in the collection directory.

DB2 Text Search Index information is removed from both the database catalog and the collection directory with the DROP INDEX ... FOR TEXT or DISABLE DATABASE FOR TEXT command.

In some scenarios, like dropping and re-creating a database, without first dropping the text indexes, the catalog information is lost while the information in the collection directory persists. The CLEANUP command removes inconsistent DB2 Text Search Index information after such scenarios have occurred.

User response: No response is required. This message is for informational purposes.

CIE0210E DB2 Text Search index cleanup completed with errors.

Explanation: Errors occurred in a DB2 Text Search index cleanup operation. Only indexes without specific errors were cleaned up successfully. All cleanup errors are written to the `db2diag.log` file.

User response: Check the `db2diag.log` for more information about the cleanup errors. Correct the errors and retry the DB2 Text Search cleanup operation.

CIE0211W The CREATE INDEX FOR TEXT command with BACKUP option is not supported for DB2 Text Search indexes on tables in a partitioned database setup.

Explanation: The index without backup options was not created successfully.

User response: Errors generated while creating the index have been written to the `db2diag.log` file. Correct the errors and retry the CREATE INDEX FOR TEXT operation.

CIE0212W Incomplete enablement of the Text Search server. Reason code: *reason-code*.

Explanation: The database is enabled for text search but some additional configuration steps need to be completed.

The reason code indicates which configuration step needs to be completed:

1

Text server configuration is unavailable for the database.

2

Review default collection directory.

User response: Respond to this error according to the reason code:

1

Insert a text server configuration into the SYSIBMTS.TSSERVERS administrative view.

2

Verify that a default collection directory is configured for the text search server.

CIE0302E Failed to set integrity for dependent tables.

Explanation: The RESET PENDING command executes a SET INTEGRITY statement for dependent tables used to manage full-text search.

The command failed to complete successfully.

User response: Run the SET INTEGRITY command with the IMMEDIATE CHECKED option enabled for any the base table's dependent tables that are still in pending mode.

CIE0304E System function failed. Errno: *error-number*.

Explanation: A system error occurred that does not allow further processing.

User response: Additional information can be found using the ERRNO information in the error message. On many operating systems, descriptions of ERRNO can be found in the C header file named errno.h.

CIE0305E Feature not supported.

Explanation: This functionality is not supported by DB2 Text Search.

User response: Examine the operation to determine the unsupported operation. Check if the operation might have violated any documented restrictions with text search indexes.

CIE0306E Table *schema-name.table-name* has no primary key.

Explanation: An attempt was made to create an index on a table that does not have a primary key.

User response: Call "db2 alter table" to ensure the existence of a primary key. Then try to create the index again.

CIE0307E Directory *directory-name* does not exist.

Explanation: The directory which was specified does not exist on the file system.

User response: Create the directory. Ensure that the instance owner is able access it.

Then attempt the operation again.

CIE0308E The size of an object exceeds the maximum allowed size. Two-part object name: *schema-name.object-name*. Maximum allowed size: *maximum-size*. Size of object: *key-column-size*.

Explanation: The internal representation of the key columns exceeds the maximum size.

User response: Change the layout of the table before creating the index again. Try using smaller key columns, which will also improve performance.

CIE0309E Object *schema-name.object-name* number key columns, exceeds max of *maximum-number*.

Explanation: The number of key columns you specified is greater than the maximum number of key columns supported.

User response: Change the layout of the table before creating the index again.

CIE0310E File *file-name* is not readable.

Explanation: The file specified cannot be read.

User response: Check the access rights for the file. Take into account that the stored procedure runs as a fenced user ID, which may also require rights to work on this file.

CIE0311E File *file-name* cannot be opened.

Explanation: The file specified cannot be opened.

User response: Check that the file exists. Also check the access rights for the file. Take into account that the administration procedure runs as a fenced user ID, which may also require rights to work on this file.

CIE0312E The unsupported type *type* was returned by *schema-name.object-name*.

Explanation: In a CREATE INDEX command, a column type transformation was specified that returns an unsupported datatype.

User response: Choose a different column type transformation function.

CIE0313E Column type *type* not supported.

Explanation: You have specified a column type that is not in the list of supported types.

User response: Ensure that the CREATE INDEX command uses only supported column types (for text column, attribute columns, user exit columns, and primary key columns). Make the appropriate changes to the command and try again. If the error persists, start a trace and also check db2diag.log. Report the error to your IBM service representative.

CIE0314E UDF *schema-name.function-name* does not exist.

Explanation: The specified user-defined function does not exist in this database.

User response: Check the name specified for this user-defined function, or register the user-defined function in the database you are using.

CIE0315E A NULL value has been passed to an internal user-defined function.

Explanation: DB2 has passed a NULL value to an internal user-defined function.

User response: First make sure the specified base table has a primary key. Change your select statement to avoid this problem. Switch on the trace function and pass the returned information to your IBM service representative.

CIE0316E Text search index *schema-name.index-name* does not exist.

Explanation: The specified text search index does not exist in this database.

User response: Check the name specified and the database you are using. Use the SYSIBMTS.TSINDEXES view to see the existing text search indexes.

CIE0317E The codepage conversion from *source-codepage* to *target-codepage* failed.

Explanation: The codepage conversion could not be performed because no adequate converter exists.

User response: Report the error to your IBM service representative.

CIE0318E Object *schema-name.object-name* does not exist.

Explanation: The specified table does not exist in this database.

User response: Ensure the schema and table name exists in the database.

CIE0319E Column *column-name* does not exist in *schema-name.table-name*.

Explanation: The specified column does not exist.

User response: Check the column name that you specified.

Check the table, view, or database you are using.

CIE0320E Table space *table-space-name* does not exist.

Explanation: The specified table space does not exist in this database.

User response: Ensure the table space exists within the database.

CIE0321E Table space *table-space-name* is not of type regular or large.

Explanation: The specified table space is not of type regular or large. An operation attempting to create a table in a system temporary or user temporary table space will fail.

User response: Reissue this command, specifying a table space that is regular or large type.

CIE0322W Specified or default database already enabled for text.

Explanation: The database you specified is already enabled for DB2 Text Search.

User response: Check the database name that you specified.

Also check to see if the DB2DBDFT variable (which implies an implicit connection) has been set.

CIE0323E Specified or default database not enabled for text.

Explanation: The database you specified is not enabled for DB2 Text Search.

User response: Check the database name you specified and the setting of the DB2DBDFT variable. If the database name is correct, use the command ENABLE DATABASE FOR TEXT.

**CIE0324E The command failed because the user does not have CONTROL privilege.
Table name: *schema-name.object-name*.
User: *user-id*.**

Explanation: You do not have the authority to use this command because it requires CONTROL authority. Only the owner of this table can use this command or grant access to use this command.

User response: Have the owner of the table provide

you with the required authorization.

CIE0325E **The operation failed because the user does not have DBADM authority.**

Explanation: You do not have the required authority to use this command because it requires database administration (DBADM) authority.

Only an administrator of this database can use this command.

User response: Obtain the required authorization from the security administrator.

CIE0326E **Disabling the database failed because text search indexes are active in the database.**

Explanation: A database cannot be disabled until all text search indexes are dropped.

At least one text search index exists in the specified or default database.

User response: See the SYSIBMTS.TSINDEXES view for the existing text search indexes. Drop the existing indexes using the DROP INDEX command, or specify the FORCE option with the DISABLE DATABASE command.

CIE0327E **The operation failed because the specified code page is not supported. Specified code page: *code-page*.**

Explanation: The code page that you specified is not supported.

User response: Specify a valid code page.

CIE0328E **The operation failed because the specified language is not supported. Specified language: *language*.**

Explanation: The specified language is not supported.

User response: Specify a valid language.

CIE0329E **The operation failed because the specified format is not supported. Specified format: *format*.**

Explanation: The specified format is not supported.

User response: Specify a valid format.

CIE0330E **The invalid value *value* was specified for parameter *parameter-name*.**

Explanation: The specified value for the parameter is invalid.

User response: Specify a valid value.

CIE0331E **The operation failed because the format of the specified update frequency is invalid. Too many terms beginning with *character* were specified in the value for the update frequency.**

Explanation: The syntax for the update frequency is not correct because too many terms beginning with *character* are specified for the index update frequency.

User response: Ensure that the DAY, HOUR, and MINUTE parameters are only specified once.

CIE0332E **Incorrect command syntax.**

Explanation: Unexpected end of command.

User response: Check the command syntax. Verify that you specified the required parameters.

CIE0333E **The command failed because of an unexpected token. Unexpected token: *token-name*.**

Explanation: The command syntax is incorrect.

User response: Check the command syntax and issue your command again.

CIE0334E **The command failed because the following token is too long: *token*.**

Explanation: The token is too long.

User response: Check the command syntax and verify that the token is reduced to the maximum size allowed.

CIE0335E **The command failed because the following token occurs twice in the update frequency: *token*.**

Explanation: You specified an incorrect syntax for the update frequency.

User response: Ensure that the DAY, HOUR, and MINUTE parameters are only specified once.

CIE0336E **Invalid range *range* for parameter-*name*. Valid range: *start* - *finish*.**

Explanation: The value should be in the allowed range.

User response: Update your command. Change the value to match those in the allowed range.

CIE0337E **Library *library* cannot be loaded.**

Explanation: A library could not be found.

User response: Check that the library is located in the library path and available. Start and stop DB2 to ensure that the current settings are used.

CIE0338E **Cannot load function** *function-name* **from library** *library-name*.

Explanation: A library entry point cannot be loaded.

User response: The library access appears to be invalid. Check that the library is only specified once.

CIE0339E **Cannot find executable program** *program-name*.

Explanation: The program file cannot be located or accessed.

User response: Check if the program file is located in the bin or adm directory of the DB2 server. The installation is corrupt if the file cannot be found.

CIE0340E **Cannot start executable program** *program-name*.

Explanation: The program cannot be started.

User response: Check if the program is located in the bin or adm directory of the DB2 server and that the appropriate libraries are installed. For further information, call the program manually on the server.

CIE0341E **Update index errors. Event view:** *schema-name.event-name*.

Explanation: Errors occurred in an update index operation. Any document errors are written to the event table.

User response: Check the event view for more information about the document errors. Clean up the event log after the problems have been fixed. In some cases, errors might be logged to db2diag.log.

CIE0342E **Text index creation failed because data type** *data-type* **of column** *column-name* **is not supported.**

Explanation: To use text search functionality with data in a given column of a table, you must create a text index on that column. Generally, you can create a text index on a column directly using the CREATE INDEX FOR TEXT command.

For some data types, such as user-defined types, it is not supported to create a text index directly against a column of one of those data types. You can still include data in columns of those unsupported data types in text searches, but some additional steps are required to create the text index. This message is returned when an attempt is made to create a text index directly against one of these unsupported data types.

User response: Respond to this error in one of the following ways:

- Alter the data type of the specified column:

1. Alter the table using the ALTER TABLE statement with the ALTER COLUMN SET DATA TYPE clause so that the data type of the specified column is one of the supported data types.

2. Run the CREATE INDEX FOR TEXT command again.

- If the data type of the specified column is binary, specify the code page of the binary data in the column in the CREATE INDEX FOR TEXT command.

- If the data type of the specified column is one of the unsupported data types, perform the following procedure:

1. Create a conversion function that casts data from the unsupported data type to one of the supported data types.

2. Specify the conversion function in the CREATE INDEX FOR TEXT command.

CIE0343E **The path** *path* **not an absolute path.**

Explanation: An absolute path should be specified.

User response: Check the path and use an absolute path in the command.

CIE0344E **No corresponding text search index found.**

Explanation: An attempt was made to search on a column without a text search index.

User response: Check the column you are searching on, or create a text search index on the column.

CIE0345N **No text search index on column** *column-name*, **table** *table-name*.

Explanation: A text search query on a table column failed because it has no valid and active text search index.

User response: Check the column you are searching on, or create a text search index on the column. If the table column has an invalid text search index, drop the text index and create a new one.

For example, a text search index gets invalidated if the underlying base table is dropped and recreated.

CIE0346E **The ALTER INDEX command failed because no command options were specified.**

Explanation: The ALTER INDEX command changes the characteristics of an index, such as the update frequency option. None of the characteristics that can be changed was specified.

User response: Specify at least one command option. Refer to the command syntax for all possible options.

CIE0347E Conflict with existing text search index on same column.

Explanation: A text search index defined on the same column already exists. Creating more than one text search index on a column is not supported.

User response: If it is no longer needed, drop the existing text search index on the specified column using the DROP INDEX command. Then try reissuing this command.

CIE0348E Attribute name missing.

Explanation: Whenever a column expression is used in the attribute expression, an attribute name must be supplied. For example:

(C1+C2 AS myname)

User response: Add "AS <attribute-name>" to the attribute expression.

CIE0349E The CREATE INDEX command failed because a user exit column does not exist.

Explanation: One or more of the user exit columns specified does not exist in the table the index is created on.

User response: Correct the CREATE INDEX command.

CIE0350E Invalid ATTRIBUTE expressions.

Explanation: The column list in the attribute expression is not valid.

User response: Correct the attribute column list in the CREATE INDEX command. Make sure the columns exist in the specified table. If a function is applied on a column, verify that it is used correctly.

CIE0351E Attribute *attribute* of column *column-name* is not supported.

Explanation: For attribute columns, the only supported data type is DOUBLE.

User response: Make sure the attribute columns of the table for the text column to be indexed are of type DOUBLE.

It is possible to use cast operators in attribute column expressions. Refer to the SQL Reference for which data types can be cast to double.

CIE0352E The invalid value *value* was specified for index configuration parameter *parameter-name*. A valid value is *value*.

Explanation: The specified value for the configuration parameter is incorrect. For valid values of the

parameters, refer to the command syntax.

User response: Correct the index configuration parameter value in the CREATE INDEX command.

CIE0353E The parameter *parameter-name* not a valid index configuration parameter.

Explanation: The index configuration option is not known.

User response: Check the CREATE INDEX command syntax. An example of a valid index configuration option is: COMMENT.

For example:

index configuration(COMMENT 'my comment')

CIE0354E Function *function-name* failed. Error *error*.

Explanation: A Windows function failed with the specified error code which does not allow further processing.

User response: Use the specified Windows system error code to get detailed error information.

CIE0355E Incorrectly specified update frequency *frequency*.

Explanation: The syntax for the update frequency statement is not correct.

User response: Correct the update frequency statement according to the syntax specification.

CIE0356E Table space *table-space-name* has more than one partition, which is not allowed for this text search index.

Explanation: For an index on a nickname, the table space of the administration tables must consist of only one partition.

User response: Specify a table space in a single-partition database partition group for the administration tables.

CIE0357E Conflicting administration operation in progress.

Explanation: Another administration command is still running or terminated abnormally without releasing the command lock.

User response: If there is no other administration command running (for the text search index), clear the lock manually using CLEAR COMMAND LOCKS. Be aware that someone else may be running the administration command holding the lock. Retry this command later.

CIE0358E Conflicting administration operation without FORCE option.

Explanation: Another administration command is still running or terminated abnormally without releasing the command lock.

User response: If there is no other administration command running, clear the lock manually using CLEAR COMMAND LOCKS. Be aware that someone else may be running the administration command holding the lock. Retry this command later. For a DISABLE DATABASE command, you may specify the FORCE option which stops all other commands on that database.

CIE0359E Empty command text with procedure *schema-name.procedure-name*.

Explanation: The indicated administration procedure was called with an invalid command text parameter.

User response: Correct the command text parameter value of the administration procedure. For valid commands, refer to the DB2 Text Search documentation.

CIE0360E Executable program *program-name* terminated abnormally.

Explanation: When executing a DB2 Text Search command, the executable was called, but terminated abnormally.

User response: Verify that the executable was not terminated explicitly by user interaction (a signal). If it was not, start a trace, rerun the command and report the error to your IBM service representative.

CIE0361E Parameter *parameter-name* missing.

Explanation: Internal error - when executing a DB2 Text Search command, an administration executable program was called with a missing parameter.

User response: Try to change DB2 Text Search command parameters to avoid the problem. If the error persists, switch on the trace function and report the error to your IBM service representative.

CIE0362E Query failed. Use simpler attribute expressions.

Explanation: DB2 Text Search creates a query based on the expressions in your command to select data for indexing from the database. The query failed because it was too complex.

User response: Reduce the complexity of attribute expressions.

CIE0363E Cannot specify UPDATE MINIMUM with *schema-name.index-name* text search index.

Explanation: The text search index was created with the RECREATE INDEX ON UPDATE option. In this context, the UPDATE MINIMUM may not be specified. Update minimum is only effective if the index is updated incrementally.

User response: If you want to recreate the index each time an update is performed, remove the UPDATE MINIMUM setting. If you want to use UPDATE MINIMUM, do not specify RECREATE INDEX ON UPDATE.

CIE0364E Inconsistent objects in specified or default database.

Explanation: At least one DB2 Text Search object is missing or corrupted. This can occur when DB2 Text Search database objects under schema SYSIBMTS are changed manually.

User response: Issue a DISABLE DATABASE command using the FORCE option. You can then enable the database for text again by using the ENABLE DATABASE command. In this case, all text search indexes are lost.

CIE0365E The user exit function (*calltype=call-type*) returned with *code=error-code*, *reasoncode=reason-code*.

Explanation: The user exit function indicated a problem. Depending on the returned code, the index update was either terminated (in case of SQL_ERROR) or continued.

User response: Check the document that was passed to the user exit function and the user exit function itself.

CIE0366E Cannot load user exit library *library-name*. *Reasoncode=reason-code*.

Explanation: The user exit library could not be loaded.

User response: Make sure that the library can be dynamically loaded.

CIE0367E User exit function *function-name* not found in library *library-name*. *Reasoncode=reason-code*.

Explanation: The user exit library could be loaded, but it does not contain the specified function.

User response: Make sure that the user exit library contains the specified function. For C++ implementations, check if the function is defined as extern "C".

CIE0368E **User exit function returned invalid data for *name*. Invalid value=*value*, context information:*context-info-1*, *context-info-2*.**

Explanation: The user exit function was called for a document and returned invalid data in the output structure.

User response: Ensure that valid data is return, if the return code is SQL_SUCCESS. Check the user exit function implementation.

CIE0369E **User exit column *column-name* not in the table the index is created on.**

Explanation: The user exit columns specified in the CREATE INDEX command must be columns of the table the index is created on. Expressions are not allowed.

User response: Rerun the corrected CREATE INDEX command.

CIE0370E **User exit function (*calltype=call-type*) threw an exception.**

Explanation: The user exit function has thrown a C++ exception. The index update was stopped.

User response: Correct the implementation of the user exit function.

CIE0371E **No row found in *name*. "IBMSNAP_REGISTER" for source table *schema-name.table-name* and capture change table *schema-name.table-name*.**

Explanation: No valid entry was found in the IBMSNAP_REGISTER table for the replication capture table characteristics specified in the CREATE INDEX command. A valid entry must contain the specified source table for the index in columns SOURCE_OWNER and SOURCE_NAME, with SOURCE_VIEW_QUAL=0 and the specified replication capture table in columns PHYS_CHANGE_OWNER and PHYS_CHANGE_TABLE.

Possible causes:

The specified source table was not registered as a replication source for the replication capture table.

User response: Register the source table correctly for DB2 Replication, or specify a correct replication capture table for the source table.

CIE0372E **Invalid value *value* for replication-setting in *name*. "IBMSNAP_REGISTER" for source table *schema-name.table-name* and capture change table *schema-name.table-name*.**

Explanation: A replication setting found in the

IBMSNAP_REGISTER table is not allowed.

Possible causes:

1. The column CHG_UPD_TO_DEL_INS does not contain the value 'Y'.
2. The column CCD_CONDENSED contains the value 'Y'.

User response: When registering the source table for DB2 Replication, ensure that update operations are transformed into pairs of delete and insert operations. In addition, ensure that no condensed replication capture tables are used.

CIE0373E **Source table *schema-name.table-name* and capture change table *schema-name.table-name* are on different servers (*source-table-server* and *capture-change-table-server*).**

Explanation: The specified source table and replication capture table must reside on the same server.

User response: This operation is not currently supported.

CIE0374E **Specified format not allowed for an XML column.**

Explanation: For columns of type XML, only the 'XML' format is allowed.

User response: Specify format 'XML' or none.

CIE0375E **Alias *schema-name.alias-name* not allowed in replication clause.**

Explanation: You are not allowed to specify an alias for a nickname in a replication clause.

User response: Specify the nickname instead of the alias, or create a new nickname for the remote table.

CIE0376E **Table in replication clause cannot be a view.**

Explanation: A view cannot be specified in the replication clause.

User response: This operation is not currently supported.

CIE0377E **Only instance owner *user-id* can use that command.**

Explanation: You do not have the required authority to use this command.

On UNIX systems: The DB2 instance owner is the authorization name specified in the db2icrt command.

On Windows: It is the authorization name running the DB2 instance service. It need not be the instance name.

User response: On UNIX, verify that the DB2 instance environment variable DB2INSTANCE is set correctly, or log in as the instance owner.

On Windows, check that the DB2 instance service, DB2 - <install-copy> - <instance-name>, associated with the current instance is using your account and password.

CIE0378E **Option *option* not allowed for current connection.**

Explanation: The usage of this option is allowed only when connecting to a DB2 z/OS system.

User response: Check that the connection options relate to a DB2 z/OS system, or reissue the current command without specifying the option.

CIE0379E **The ALTER INDEX SET ACTIVE operation failed because an active index already exists on the specified column.**

Explanation: You can activate a text search index using the db2ts ALTER INDEX command or the SYSPROC.SYSTS_ALTER procedure with the SET ACTIVE clause. Multiple text search indexes can exist at the same time for a given column. However, only one text search index can be active on a column at any time.

User response:

- To switch the activation status between text indexes, use the ALTER INDEX command with SET ACTIVE enabled, but without specifying the UNILATERAL option.
- Issue the DROP INDEX command to remove other active text indexes if they are no longer required. Then reissue your command.

CIE0380E **Invalid codepage *codepage* for non-binary text column.**

Explanation: For text column data types that are not binary (that is, not BLOB and not strings-types FOR BIT DATA), DB2 always stores the data in the database codepage.

User response: Omit the codepage clause.

CIE0381E **Minimum update frequency *minimum-frequency* is less than allowed-frequency minute limit.**

Explanation: The update frequency defines a minimum time interval between two automatic index updates that is lower than the allowed value.

User response: Change the update frequency specification to keep to the indicated limit.

Note that the time interval between the latest entry on day 6 and the earliest entry on day 0 must also be considered.

CIE0382E **NumberOfMatches not supported.**

Explanation: This function is not supported by DB2 Text Search.

User response: Use either DB2 Net Search Extender or change your SQL query so that it can run without the NumberOfMatches function.

CIE0383E **Operation took longer than limit of *limit* seconds.**

Explanation: An administration operation did not succeed within the time limit. Possible reasons are workload problems while using text search instance services.

User response: Ensure that text search instance services are up and running.

Check db2diag.log for further information.

CIE0384E **Text search index not created on *schema-name.table-name* because the table is not a base table.**

Explanation: Text search indexes can be created only on a base table. Creating a text search index on non-base tables such as views or MQTs is not supported.

User response: Check whether the column belongs to a base table or not. Create the text search index on a base table column.

CIE0385E **The text search index could not be created on the specified table because the table is a partitioned table. Table: *schema-name.table-name*.**

Explanation: Text search indexes cannot be created on a partitioned table.

User response: Check if the table is partitioned. Create the text search index only on a table that is not partitioned.

CIE0386E **Database already enabled for DB2 Net Search Extender.**

Explanation: The specified or default database is already enabled for a different type of text search index. A database may be associated with only one text search component.

User response: If you do not want to use DB2 Net Search Extender, disable the database using the DB2TEXT DISABLE DATABASE command. After disabling the database for DB2 Net Search Extender, reissue the DB2 Text Search command.

CIE0387E Cannot use NULL values for required parameters.

Explanation: A DB2 Text Search administration procedure argument that is not optional received NULL as the specified value. NULLS can only be used for optional arguments.

User response: Correct the arguments specified for the procedure call.

CIE0388E Undefined client error.

Explanation: A processing error occurred when calling an internal function.

No details are available for this error.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0389E Internal function processing error.

Explanation: The method could not complete because it is not implemented or because the parameters specified make it impossible to execute.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0390E Invalid object referenced.

Explanation: A processing error occurred when calling an internal function. An invalid or uninitialized object was used or referenced.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0391E Code page *source-codepage* to *target-codepage* convert error.

Explanation: A processing error occurred when calling an internal function. A passed data could not be converted from the source code page to the target codepage.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0392E The following string needs data type *data-type*. String: *string*.

Explanation: A processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0393E Required property *property* absent.

Explanation: A processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0394E Property=*property* value=*value* failed as type=*type*.

Explanation: A processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0395E No db2ts start command was issued.

Explanation: A command was called which requires text search instance services.

User response: Start the text search instance services with the db2ts start command.

CIE0396E DB2 Text Search Instance services not configured to support rich text.

Explanation: A command which requires rich text support for DB2 Text Search was called.

Rich text has not yet been configured for the DB2 Text Search Instance.

User response: Use richtextTool to enable and configure DB2 Text Search for rich text support.

CIE0397E Table column *schema-name.table-name* has no valid text search index.

Explanation: A text search query on a table column failed because it has no valid and active text search index.

User response: Ensure that the table column has a valid and active text search index before performing a text search query. If the table column has an invalid text search index, drop the text index and create a new one.

For example, a text search index gets invalidated if the underlying base table is dropped and recreated.

CIE0398N Text search index *schema-name.index-name* is not valid and cannot be used.

Explanation: Text index is invalid and cannot be updated or used.

User response: Drop and recreate the invalid text search index.

CIE0399E **Failed to set integrity for the following table:** *schema-name.table-name*.

Explanation: The RESET PENDING command executes a SET INTEGRITY statement for dependent tables that are used to manage full text search.

The command failed to complete successfully.

User response: Run the SET INTEGRITY command for the failed table, with the IMMEDIATE CHECKED option enabled.

CIE0401E **A search system operation failed.**
Message id: *message-id*. **Exception**
es::exception-part1: *exception-part2* **was**
thrown by file *file-name* **line** *line-number*.

User response:

CIE0402E **Failed to set integrity for dependent tables.**

Explanation: The RESET PENDING command executes a set integrity statement for dependent tables used to manage full-text search. The command failed to complete successfully.

User response: Issue the SET INTEGRITY command with the IMMEDIATE CHECKED enabled for all of the base table's dependent tables that are still in pending mode.

CIE0403E **The data type of column** *column-name* **is not supported for text search indexes in this database.**

Explanation: To use text search functionality with data in a given column of a table, you must create a text index on that column. Generally, you can create a text index on a column using the CREATE INDEX FOR TEXT command.

The behavior of the text search functionality is affected by the value of the DB2_COMPATIBILITY_VECTOR registry variable. Specifically, creating a text index on a column of type DATE or TIMESTAMP(0) is not supported when the DB2_COMPATIBILITY_VECTOR registry variable is set to 40, where it enables the use of the DATE data type as TIMESTAMP(0) as a combined date and time value.

This message is returned when an attempt is made to create a text index for a column of type DATE or TIMESTAMP(0) while the DB2_COMPATIBILITY_VECTOR registry variable is set to 40.

User response: Respond to this error in one of the following ways:

- Alter the data type of the specified column:

1. If possible, alter the table so that the data type of the specified column is not DATE or TIMESTAMP(0).
 2. Run the CREATE INDEX FOR TEXT command again.
- Disable the compatibility features:
 1. Disable the compatibility features that are preventing the index creation by setting the DB2_COMPATIBILITY_VECTOR registry variable to NULL.
 2. Run the CREATE INDEX FOR TEXT command again.

CIE0404E **Format** *format-type* **is not allowed for column type** *column-type*.

Explanation: The specified index format type should be supported by the indicated column type. For example, INSO format is only supported by BLOB column and XML format is only supported by XML column.

User response: Check whether the specified index format type is supported by the indicated column type.

CIE0405E **CJKSegmentation for language**
language-name **is not supported.**

Explanation: CJKSEGMENTATION is supported only for Chinese, Japanese and Korean character sets.

User response: Remove CJKSEGMENTATION as part of the CREATE INDEX command

CIE0406E **Token CJKSEGMENTATION occurs twice in index configuration.**

Explanation: You specified an incorrect syntax for the CJKSEGMENTATION.

User response: Ensure that the CJKSEGMENTATION parameter is only specified once.

CIE0408E **User** *user-name* **does not have DATAACCESS authority to perform the DB2 Text Search upgrade.**

Explanation: The text search user needs to have DATAACCESS authority to perform the upgrade.

User response: Grant DATAACCESS privilege to the user.

CIE0409E **The DB2 Text Search catalog has not been upgraded to the current version.**

Explanation: Before upgrading DB2 Text Search indexes, all catalog tables or views associated with text search should be upgraded to the latest version.

User response:

1. Connect to the database.

2. Call the upgrade procedures in the following order:
 - a. SYSPROC.SYSTS_UPGRADE_CATALOG
 - b. SYSPROC.SYSTS_UPGRADE_INDEX.

CIE0410E **Incorrect usage. Parameters are missing or incorrect parameters are given.**

CIE0411C **Cannot find the db2level command to detect if the instance is 32-bit or 64-bit.**

Explanation: DB2 Text Search was not able to locate the db2level command. DB2 Text Search uses this command to determine if the DB2 instance is 32-bit or 64-bit.

User response: The db2level command binary is missing from the DB2 installation or cannot be found via search path for executable files. Resolve any DB2 installation or search path for executable files issues before retrying this operation.

CIE0412E **DB2INSTANCE environment variable is not set.**

Explanation: In order to configure rich text support, DB2 Text Search requires that the DB2INSTANCE environment variable is set.

User response: Set the DB2INSTANCE variable and retry richtextTool command.

CIE0413E **The DB2 Text Search directory named *directory-name* does not exist.**

Explanation: DB2 Text Search has not been configured correctly.

User response: Ensure that DB2 Text Search has been configured correctly.

CIE0414E **The user named *user-name* does not have write permission to the directory named *directory-name*.**

Explanation: richtextTool cannot create a directory to store unzipped files.

User response: Refer to documentation on user permissions needed to execute this tool.

CIE0415C **DB2 Text Search cannot locate the DB2 installation.**

Explanation: richtextTool cannot locate the DB2 installation.

User response: Resolve DB2 installation issues before retrying this operation.

CIE0416E **Rich text support cannot be configured because DB2 Text Search is running.**

Explanation: The rich text support cannot update configuration files while the DB2 Text Search server is running.

User response: Stop the DB2 Text Search server by issuing the command, db2ts stop for text.

Then retry configuring rich text support.

CIE0417E **Cannot access directory named *directory-name* to locate Outside In Technology product zip files.**

Explanation: richtextTool cannot access the directory given by the user.

User response: Check that the directory exists and you have read permission to that directory.

CIE0418E **richtextTool failed to extract files from the zipfile named *compressed-file-name*.**

Explanation: richtextTool cannot extract files to complete configuring rich text support.

User response: Examine and correct the error messages returned from the file unzipping operation.

Then retry running richtextTool.

CIE0419E **File extraction failed. Cannot locate the file named *file-name* in the Outside In Technology product directory.**

Explanation: richtextTool cannot find a required file in the Outside In Technology product directory.

User response: Review the output of the file extraction operation to correct any issues that might have caused unzip failure.

If rich text has not been configured for this copy of DB2, remove the directory named stellent under the directory named db2tss and retry running richtextTool.

CIE0420E **Files are missing from a directory named *directory-name*.**

Explanation: Files are missing from the named directory. A failure may have occurred during the run of the configuration tool. System configuration may be in an inconsistent state.

User response: Manually remove the named directory. Then rerun richtextTool to configure rich text support.

CIE0421E **Cannot locate version *version* Outside In Technology product Html Export zip file named *compressed-file-name*.**

Explanation: Cannot locate the proper version of the Outside In Technology product zip file.

User response: Refer to the error message to see which file is missing.

Obtain the missing file then retry running richtextTool.

CIE0422E **Cannot locate version *version* Outside In Technology product Search Export zip file *compressed-file-name*.**

Explanation: Cannot locate the proper version of the Outside In Technology product zip file.

User response: Refer to the error message to see which file is missing.

Obtain the missing file then retry running richtextTool.

CIE0423E **richtextTool cannot access the file named *file-name*.**

Explanation: richtextTool cannot access the file referenced in the error.

User response: Ensure the file exists and correct any access permission problems to that file.

CIE0424E **The DB2 Text Search configuration directory named *directory-name* for the DB2 instance named *instance-name* does not exist.**

Explanation: DB2 Text Search has not been correctly configured or is damaged.

User response: Retry configuring rich text support by running richtextTool.

CIE0425E **DB2 Text Search cannot set the required access permission for files located in the directory named *directory-name*.**

Explanation: Correct file permissions are needed so the rich text support program can launch properly.

richtextTool cannot set the correct read and execute permissions to extracted files on its own.

User response: Examine the error messages returned from the chmod command and correct the problems.

If rich text has not been configured before for this copy of DB2, remove the directory named stellent under the directory named db2tss before retrying richtextTool.

CIE0426E **DB2 Text Search cannot detect the DB2COPY name for the instance that is being configured.**

Explanation: richtextTool cannot derive the DB2COPY name from DB2 instance name. DB2 may not have been setup correctly.

User response: Correct any DB2 setup problems before configuring rich text support.

CIE0427E **richtextTool cannot restore a configuration file named *file-name* from an existing copy.**

Explanation: richtextTool tried to restore an existing copy of a DB2 Text Search configuration file and hit an error.

User response: System configuration maybe in an inconsistent state.

Check if system is running out of disk space. Correct the problem, then rerun richtextTool.

CIE0428E **The DB2 command line environment is not set.**

Explanation: richtextTool requires the DB2 command line environment to detect the DB2 environment and to execute DB2 utilities when configuring rich text support.

richtextTool cannot configure DB2 Text Search rich text support without the DB2 command line environment setup.

User response: Run richtextTool in a DB2 command window.

CIE0429E **The home directory of the DB2 instance named *instance-name* does not exist. The home directory is named *directory-name*.**

Explanation: DB2 instance setup is incorrect.

User response: Correct any DB2 instance setup issues before running richtextTool.

CIE0430E **richtextTool cannot be run because the contents of the file named *file-name* are incorrect.**

Explanation: The contents of the referenced file are incorrect. System configuration might be in an inconsistent state.

User response: Check system resources. Check if the system is running out of disk space.

Correct the problem and rerun richtextTool.

CIE0431E **A text search index can not be created because the sequence of primary key columns are not the same as base table.**

Explanation: The sequence of primary key columns should be kept the same as base table, otherwise a text search index can not be created.

User response: Ensure that the column sequence in a compound primary key matches the column sequence in the base table.

CIE0432E **The database directory cannot be found on the local file system.**

Explanation: The text search administration operation failed to find the local database directory.

User response: Issue the 'db2 list database directory' command to confirm that the local database directory exists. Then try the operation again.

CIE0440E **User *user-id* does not have the authority to create a text search index on the table *schema-name.table-name*.**

Explanation: "CREATE INDEX ... FOR TEXT" requires one of the following:

- DBADM authority
- CONTROL privilege on the table
- INDEX privilege on the table with either IMPLICIT_SCHEMA privilege on the database or CREATEIN privilege on the index schema.

For more information about DB2 Text Search create index, see the topic called "DB2 Text Search CREATE INDEX command" in the DB2 Information Center.

User response: Ensure that the user has the required authority to perform the create text index operation.

CIE0441E **User *user-id* does not have authority to enable the database *database-name* for DB2 Text Search.**

Explanation: Enabling text search for a database requires DBADM authority.

For more information about enabling DB2 Text Search, see the topic called "Enable Database for Text text search command" in the DB2 Information Center.

User response: Ensure that the user has the required authorization to perform the ENABLE operation.

CIE0442E **The instance owner does not have the authority to perform the text index command on database *database-name*.**

Explanation: The user attempted to run a text index command without the instance owner holding the appropriate authority.

Different DB2 Text Search index commands require varying levels of database authority. For a list of required authorities for DB2 Text Search index commands, see the topic called "DB2 Text Search index command authority requirements" in the DB2 Information Center.

User response: Ensure that the instance owner has DBADM with DATAACCESS authority in order for the user to perform text index operations.

CIE0443E **User *user-id* does not have authority to update the DB2 Text Search index named *index-schema-name.index-name*.**

Explanation: Update text index requires one of the following:

- DATAACCESS authority
- CONTROL privilege on the table.

For more information about DB2 Text Search update, see the topic called "DB2 Text Search UPDATE INDEX command" in the DB2 Information Center.

User response: Ensure that the user has the required authorization to perform the update text index operation.

CIE0444E **User *user-id* does not have DATAACCESS authority to perform the DB2 Text Search index command.**

Explanation: You need to have DATAACCESS authority to perform the DB2 Text Search index command.

Different DB2 Text Search index commands require varying levels of database authority. For a list of required authorities for DB2 Text Search index commands, see the topic called "DB2 Text Search index command authority requirements" in the DB2 Information Center.

User response: Ensure that the SECADM grants DATAACCESS authority to the user.

CIE0445N **The requested operation cannot be executed. Run the REBIND command for packages *package-list*.**

Explanation: The list of packages are invalid. You must run the REBIND command.

User response: Run the REBIND command on the packages before running this text index command.

CIE0446E **The document exceeds the maximum size limit of *max-size* bytes.**

Explanation: DB2 Text Search can index documents up to the maximum size limit. The current document exceeds this limit and is not indexed.

User response:

1. Reduce the document size
2. Submit the modified document

CIE0447N **Failed to retrieve the text-maintained dependent tables for**
schema-name.object-name.

Explanation: The command executes a query to retrieve information from the text index administration tables. Querying the text index information failed.

User response: Ensure that the text index administration tables are accessible. For more information see the db2diag.log.

CIE0448N **The backup configuration is inconsistent.**

Explanation: The backup infrastructure was enabled for this text index, but the backup configuration is unavailable or inconsistent.

User response: Review and update the backup configuration according to the documentation for Text Search index backup options.

CIE0449N **Text search index *schema-name.index-name* not enabled for Backup.**

Explanation: Text search indexes can only be backed up when a backup mode is specified.

User response: To enable backup for a text search index, use the ALTER INDEX FOR TEXT operation to set the backup mode for the text search index to IMMEDIATE or DEFERRED.

CIE0450N **The backup configuration for the text search index is invalid.**

Explanation: The backup configuration refers to backup modes in the Text Search server. The configuration does not match with the supported backup modes.

User response: Update the backup configuration to match with a supported backup mode for the Text Search server. Review the Text Search server documentation for supported modes.

CIE0451E **Lock cannot be cleared for text search index *schema-name.index-name*.**

Explanation: The text index is locked for an operation on the database level.

User response: Repeat the CLEAR COMMANDLOCKS operation without specifying a text search index name.

CIE0452E **User *user-name* does not have the authority to alter text search index *schema-name.index-name*.**

Explanation: To alter a text search index, the authorization ID required one of the following:

- DBADM role on the database
- CONTROL role on the table
- ALTERIN privilege on the index

User response: Ensure that the SECADM assigns appropriate privileges.

CIE0453E **User *user-name* does not have the authority to drop the text search index *schema-name.index-name*.**

Explanation: To drop a text search index, the authorization ID requires one of the following roles or privilege:

- DBADM authority on the database
- CONTROL privilege on the table
- DROPIN privilege on the schema

User response: Ensure that the SECADM assigns appropriate privileges.

CIE0454E **User *user-name* needs role *role-name* for this operation.**

Explanation: Operations on text search indexes require the appropriate SYSTS_ADM, SYSTS_MGR or SYSTS_USR roles.

User response: Ensure that the SECADM assigns the appropriate role.

CIE0456E **Start operation for a stand-alone text search server setup has to be executed separately.**

Explanation: When using a stand-alone text search server setup, an operation like starting a server has to be executed in the stand-alone text search server environment using the startup script. The db2ts utility from the DB2 instance environment cannot be used to start the stand-alone text search server.

You can use the db2ts utility to check the text search server status. Ensure that the text search server for the DB2 instance is registered by issuing SYSTS_CONFIGURE for a text search enabled database.

User response: Run the following script to successfully start the text search server:

- /bin/startup.sh (Linux or Unix platforms)
 - \bin\startup.bat (Windows platform)
-

CIE0457E Expected values for DB2 Text Search defaults are not found.

Explanation: Text search defaults are setup in the SYSIBMTS.TSDEFAULTS view. These values are typically added when enabling a database for text search and subsequent configuration steps such as during execution of the SYSTS_CONFIGURE procedure.

User response: Ensure that database level defaults for text search are setup correctly by successfully issuing the ENABLE operation and the SYSTS_CONFIGURE procedure for a text search enabled database.

CIE0458E The specified administration option option cannot be used in this setup.

Explanation: Some options for administrative operations are supported only in certain product environments.

User response: Ensure that only the supported options for the instance type or database are specified in the command.

Check the DB2 instance type and ensure that the used options apply to it.

Retry the operation after removing or changing the option so that it applies to your setup.

CIE0459E Data preparation constructor constructor-name is invalid

Explanation: There is no constructor configured with this name, or there is no value assigned.

User response: Ensure that the configuration for data preparation is valid in SYSIBMTS.TSDEFAULTS. Use the DATAPREP: prefix to define the constructor.

CIE0460E Parallel index update errors across multiple database partitions.

Explanation: When updating text search indexes in a partitioned database environment, the index update operation runs on each relevant database partition in parallel. Errors occurred on one or more database partitions.

User response:

1. Ensure the database has been activated after the DB2 instance started.
2. Check the event view and db2diag.log file for more information about the errors. You might need to cleanup any remaining locks for the text search index using the CLEAR COMMAND LOCKS command.
3. Correct the errors and retry the DB2 Text Search update index operation.

CIE0462I The text search catalog has been upgraded to the current version.

Explanation: In order to make text search work, the text search index also needs to be updated.

User response: Update the text search index by calling "SYSPROC.SYSTS_UPGRADE_INDEX".

CIE0463E The DB2 Text Search catalog has not been upgraded to the current version.

Explanation: The text search catalog must be upgraded to the current version.

User response: Update the text search catalog to the current version.

Chapter 53. CIE0500 - CIE0999

CIE0701E **Internal error. Location:** *location-1*,
location-2.

Explanation: An internal processing error that does not allow further processing occurred.

User response: Try starting and stopping the text search instance services, as well as DB2. If the error persists, start a trace and also check db2diag.log.

CIE0702E **Memory allocation error.**

Explanation: The system has run out of memory.

User response: Increase the available memory size for the user or stop other processes that are running in parallel.

CIE0703E **Error using shared system resources.**

Explanation: A request to share or access system resources, such as shared memory or semaphores, cannot be fulfilled.

User response: Check the current system status and configuration. On UNIX, use the `ipcs` command to check the resources. Stop all applications, such as DB2 and DB2 Text Search. If further resources are listed, clean them up using `ipcrm`.

CIE0705E **Instance services error.**

Explanation: An instance services error occurred.

User response: Check db2diag.log for further information, or clean up your shared resources. Also see CIE00703.

CIE0706E **Instance services could not be stopped.**

Explanation: The `db2ts stop` command has not stopped the instance services and there are still processes running.

User response: To recover from this error state, reissue the stop command after allowing any text search administration operations running in the background to complete.

In the case of administration commands that have terminated abnormally (or that you want to stop), clear the lock manually using `CLEAR COMMAND LOCKS`.

If there are no signs of other administration operations running, or if the stop command still does not work, try restarting the instance services and then issue a stop command. As a last resort, stop DB2 and all applications, and then clean up system resources.

CIE0707E **Contacting instance services timed out.**

Explanation: A timeout occurred when contacting the text search services. Possible reasons are:

1. The DB2 Text Search instance is not configured properly.
2. Rich text feature was not disabled before the upgrade.
3. Network issues.

User response:

1. Verify that DB2 Text Search is configured properly .
2. Verify that rich text support is configured properly .
3. Ensure that the network communication is working.

CIE0708E **An error occurred while opening the following Windows service: *service-name*. Error that was returned: *error*.**

Explanation: The specified service cannot be found on the Windows system.

User response: Check if the specified service is installed on the Windows system. Use the specified Windows system error code to get detailed error information.

CIE0709E **Cannot find DB2 instance profile path.**

Explanation: A function to obtain the DB2 instance profile path failed.

User response: Check the DB2 instance configuration.

On Windows, try to create a DB2 instance without specifying the instance profile path information and retry the command.

CIE0710E **Administration table space *text-index-table-space* is not defined in the same database partition group as the table space of *schema-name.object-name*. Database partition group: *database-partition-group*.**

Explanation: The table space of the administration tables must be distributed over different partitions in exactly the same way as the table containing the text column to be indexed.

The system checks whether the specified table space is defined in the same database partition group.

User response: Specify a table space that is defined in the same partition group as the table containing the text column to be indexed.

CIE0711E **Invalid operation on log table**
schema-name.table-name. **Operation:**
operation.

Explanation: The log table keeps track of operations executed on the table containing the indexed text column. This table contains an invalid value due to manual changes made to this table.

User response: Check the log table, and delete or correct the invalid entry.

CIE0712E **The table named** *schema-name.table-name* **contains an incorrect syntax expression in the column named** *column-name*.

Explanation: There is an error in the expression list in the specified text column.

User response: Check the delimiter Begin and End pairs.

CIE0713E **Length of index properties** *length* **exceeds max.** *maximum-allowed-length*.

Explanation: The total length of index properties exceeds the maximum size.

User response: Report the error to your IBM service representative.

CIE0714E **Setting environment variable**
environment-variable **failed.**

Explanation: Setting the specified environment variable failed.

User response: Report the error to your IBM service representative.

CIE0715E **Call to** *function-name*: **rc=***return-code*,
SQLCODE=*sqlcode*.

Explanation: An internal processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0716E **Error creating shared memory.**

Explanation: The shared memory resource could not be created due to a previous error or a permission problem.

User response: Check db2diag.log for further information, or clean up your shared resources. Also see error CIE00703.

CIE0717E **Value for parameter** *parameter-name* **too long.**

Explanation: The value exceeds the allowable maximum size.

User response: Check the maximum size. Correct the parameter to be within limits.

CIE0718E **Log table** *schema-name.table-name* **will be modified in next UPDATE.**

Explanation: When starting an incremental index update, a time stamp is created. This serves as a threshold for processing change records. Changes occurring concurrently to the incremental update are processed during the next update.

In certain situations, changes in transactions may be uncommitted at the time the update starts, but are committed while the index update is performed, which can lead to inconsistencies.

To avoid inconsistent situations of this kind, the change records prior to the threshold time stamp are not deleted from the log table, even if they have been partially processed.

During the next incremental update, the changes will be re-applied to the index. In case of delete operations, this can lead to warning messages in the log table indicating that the document was already deleted.

User response: If CIE00718 errors frequently occur, consider dropping and re-creating the index with a modified time stamp threshold for incremental index update.

For example:

```
CREATE INDEX ...
      INDEX CONFIGURATION(UPDATEDELAY 30)
```

This means that processing during an incremental update only changes records older than 30 seconds, and avoids interference with concurrent change transactions of less than 30 seconds.

CIE0719E **Target database system**
database-system-name **not supported.**

Explanation: You tried to execute a command with a connection to a database system that is not supported by DB2 Text Search.

User response: Check your environment and ensure that the DB2 Text Search is used against a supported database.

CIE0720E **Cannot find type or version info for server** *server-name*.

Explanation: The type and version information for the server could not be found in the DB2 catalog view 'SERVERS'.

User response: Make sure that the DB2 federated environment is set up correctly.

CIE0721E Signal *signal* caught.

Explanation: The program received a signal.

User response: If you did not interrupt the program, contact your IBM service representative.

CIE0722E Corrupted instance services input file *file-name*.

Explanation: The file containing index information is corrupted.

User response: For obvious problems caused by manual editing of the file, use your system editor and try to correct the problem. You may have truncated an entry, or deleted the end-of-line character.

If this does not restore the file content, try the following:

- Call command 'db2ts stop for text' to stop the instance services
- Delete or rename the corrupted file
- Call command 'db2ts start for text' to re-start the services
- Use command 'db2ts cleanup for text' to recreate the file based on the catalog information of all databases for a DB2 instance.

CIE0723E Windows Exception *exception* caught.

Explanation: The program received a Windows exception.

User response: If you did not interrupt the program, contact your IBM service representative.

CIE0724E Exception *exception*, address=*address*, flags=*flag*.

Explanation: The program received a Windows exception.

User response: If you did not interrupt the program, contact your IBM service representative.

CIE0725E DB2 Text Search is not supported in a multiple partition database environment.

Explanation: A text search index can be created only on DB2 configurations with a single partition.

User response: Ensure that the DB2 instance is not configured for with multiple database partitions.

CIE0726E JDK_PATH setting is incorrect.

Explanation: DB2 Text Search requires use of the Java Runtime Environment (JRE) that is part of the DB2 installation.

An error has occurred in finding the JRE due to a problem with the JDK_PATH value.

User response: Ensure that DB2 was installed successfully.

Check the value of JDK_PATH, ensure it is pointing to the directory where the JRE was installed and that the JRE is accessible.

The command 'db2 get dbm cfg' may be used to verify the JDK_PATH setting.

CIE0727E Java classpath not found.

Explanation: A directory in the Java classpath for supporting DB2 Text Search was not found.

User response: Check the DB2 installation. Ensure that the DB2 Text Search option was selected and that the installation completed successfully.

CIE0728E Cannot close file *file-name*.

Explanation: The file specified cannot be closed.

User response: Check if the file can be accessed.

Also check if there is enough free space on the system.

CIE0729E Cannot copy file *file-name-1* to *file-name-2*.

Explanation: The first file cannot be copied to the second file.

User response: Check if the second file already exists and is writable.

Also check if there is enough free space on the system.

CIE0730E Cannot remove file *file-name*.

Explanation: The specified file cannot be removed from the system.

User response: Check if the file can be accessed.

CIE0731E Failed to write on file *file-name*.

Explanation: The file specified is not writable.

User response: Check if the file can be accessed.

Also check if there is enough free space on the system.

CIE0732E **Line *line* daemon index entry parsing error.**

Explanation: The daemon file ciedem.dat is corrupted.
The erroneous line is shown.

User response:

1. Rename the file (Ensure that there is no file ciedem.dat in the instance directory.)
2. Stop the instance services with the following command:
db2ts stop for text
3. Start the instance services with the following command:
db2ts start for text
4. Re-create the daemon file: db2ts cleanup for text
(This may take some time, because all databases are searched for index information.) In case of problems, contact your IBM service representative.

CIE0733E **Error in ciedem.dat: database-partition *database-partition* invalid.**

Explanation: The daemon file ciedem.dat is corrupted.
The value for a DB2 partition is not numeric.

User response:

1. Rename the file (Ensure that there is no file ciedem.dat in the instance directory.)
 2. Stop the instance services with the following command:
db2ts stop for text
 3. Start the instance services with the following command:
db2ts start for text
 4. Re-create the daemon file with the following command:
db2ts cleanup for text
- (This may take some time, because all databases are searched for index information.) In case of problems, contact your IBM service representative.

CIE0734E ***number* index cleanup errors.**

Explanation: The specified number of errors occurred during cleanup processing.

User response:

1. Inspect the db2diag.log to determine what kind of errors occurred.
2. Reissue the command after correcting the cause of the error.

CIE0735E ***number of number* collections could not be deleted during this operation.**

Explanation: During the DROP or DISABLE command, some collections of the text search engine could not be dropped.

User response: Ensure that the text search instance services are up and running.

If you working in a partitioned database, you need to check all servers. If they are running, try to stop and start those servers again. Check db2diag.log for details.

CIE0736E **Cannot delete collection *collection-name* from text search index *schema-name.index-name*.**

CIE0737E **Cannot open file *file-name*. Reason code: *reason-code*.**

Explanation: DB2 Text Search could not open the specified file.

Reason codes:

11

The file could not be opened for read access.

12

The file could not be opened for write access.
It is expected to exist.

22

The file could not be opened for write access.
It must not exist.

32

The file could not be opened for write access
or could not be deleted.

User response: Verify that the path of the specified file exists.

Verify that the path and file permissions allow the file to be opened as required for the reason code specified.

CIE0738E **Cannot read nodes configuration file *file-name*.**

Explanation: The specified file cannot be read. It is probably corrupted.

User response: Verify that the file contents comply with the format of the nodes configuration file. If this is not the case, correct the file accordingly.

CIE0739E **Error in nodes config file *file-name*. Reason code: *reason-code*.**

Explanation: The content of the specified file is not consistent.

Reason codes are interpreted as follows:

1

The required file entry specifying the node on which the instance services are running is missing.

2

The file contains multiple file entries specifying the node on which the instance services are running. Only one is required.

User response: Check the file contents and correct the file according to the reason code.

CIE0740E **Nodes config file** *file-name* **needs entry for node** *database-partition*.

Explanation: The file specified is not consistent. There is at least one node entry missing.

User response: Adapt the nodes configuration file and verify that it is in sync with the db2nodes.cfg file.

CIE0741E **Unsupported command** *command-name* **with** *schema-name.procedure-name* **procedure**.

Explanation: The command specified is not supported in conjunction with the administration procedure.

User response: Use the db2ts executable to execute the command.

CIE0742E **No nodegroup information for table space** *table-space*, **table schema** *schema-name*, **table name** *table-name*, **node group name:** *database-partition-group-name*.

CIE0743E **DB2 Installation copy for instance** *instance* **was not found**.

Explanation: The DB2 installation copy to which the given instance belongs could not be found.

User response: Check the DB2 Installation. Please ensure that the following:

1. The DB2 Text Search option was selected and the installation completed successfully.
2. Verify that the instance exists by opening a DB2 CLP Window for each DB2 installation copy and executing the 'DB2ILIST.EXE' command. If the instance exists then it will be listed in one of the DB2 CLP Windows.

CIE0744E **Unable to shut down text search instance services**.

Explanation: You attempted to stop the text search instance services, but they could not be stopped.

User response: If there were errors using DB2 Text

Search that indicated that the text search services should be stopped and restarted, the following steps might be necessary:

- On UNIX, use the ipcs command to check the resources. Stop all applications, such as DB2 and DB2 Text Search. If further resources are listed, clean them up using ipcrm.
- In the case of multiple DB2 instances running on the system, ensure that the DB2 Text Search for each instance has been configured for a unique communication port.

CIE0745E **Cannot find file path** *path*.

Explanation: A file in the DB2 Text Search installation path was not found.

User response: Check the DB2 installation. Ensure that the DB2 Text Search option was selected and that the installation completed successfully.

CIE0746E **Authentication error. Reason code:** *reason-code*.

Explanation: Authentication failure occurred during a text search administration operation or a text search query. The reason code provides more information on the cause for the failure:

Reason Code:

1

Invalid token

2

File not found or unexpected file permissions

3

A system call for executing a process failed

4

A file system call failure

5

A file system call failure

6

Failure spawning a process

7

Memory allocation failure

User response: Check the DB2 installation. Ensure that the DB2 Text Search option was selected and that the installation completed successfully.

A DB2 instance is associated with specific text search instance services that cannot be accessed from a different DB2 instance. Ensure that the correct text search instance services are being accessed.

Stop and restart DB2 Text Search instance services.

Consider running the tool for generating the authentication token as the authorized user.

CIE0747E Timeout starting the text search server.

Explanation: The startup of text search instance services took much longer than normal.

User response: If there were errors using DB2 Text Search that indicated that the text search services should be stopped and restarted, the following steps might be necessary:

- On UNIX, use the `ipcs` command to check the resources. Stop all applications, such as DB2 and DB2 Text Search. If further resources are listed, clean them up using `ipcrm`.
- Check the configuration of DB2 Text Search.

CIE0748E Shared memory not available.

Explanation: The shared memory resource could not be accessed due to a previous error or permission problem.

User response: Ensure that the text search instance services were started successfully. Check `db2diag.log` for further information, or clean up your shared resources.

Stop and restart text search instance services.

CIE0749E Could not delete collection directory.

Explanation: During the execution of the `DISABLE` command, the collection directory for text search indexes could not be removed.

User response: By default, the name of the subdirectory is system generated and is under the `db2collections` subdirectory in the database path, `DBPATH`, used by the database. The `db2collections` directory is created during the `ENABLE` command and deleted when the database is disabled for text search.

CIE0750E Command failed. Database concurrently associated with NSE.

Explanation: Only one type of text search index may be associated with a database at any given time. A database can use either DB2 Text Search or DB2 Net Search Extender for text search indexes, but not both concurrently.

User response: If you do not want to use DB2 Net Search Extender, disable the database using the `DB2TEXT DISABLE DATABASE` command. If you would rather use DB2 Net Search Extender, try this command again after disabling the database for DB2 Text Search.

CIE0751E Unexpected condition detected.

Explanation: A processing error occurred when calling an internal function. An unexpected condition was detected. Operation could not complete.

User response: If the error persists, start a trace and check the `db2diag.log`. Report the error to your IBM service representative.

CIE0752E Mutex error. Details: *details*.

Explanation: An error occurred when a mutex function was called.

User response: If the error is persistent, turn on the trace for further error determination.

CIE0753E Conditional variable error. Details: *details*.

Explanation: An error occurred when a conditional variable was called.

User response: If the error is persistent, turn on the trace for further error determination.

CIE0754E Thread error. Details: *details*.

Explanation: An error occurred when a thread function was called.

User response: If the error is persistent, turn on the trace for further error determination.

CIE0755E Communication failure.

Explanation: A processing error occurred when calling an internal function. The communication layer reported a failure without throwing an exception.

User response: If the error persists, start a trace and check the `db2diag.log`. Report the error to your IBM service representative.

CIE0756E Connect failed on host *host-name* and port *port-number*.

Explanation: The connection to the Text Search server can fail for various reasons, including the following:

- The DB2 Text Search server is not started.
- Mismatch in the DB2 Text Search configuration for the DB2 instance and the database in which the operation is executed.
- The DB2 Text Search server for the DB2 instance is either not configured or configured incorrectly.
- Port conflict or network communication issues.

User response: Perform the following troubleshooting steps:

1. Start the DB2 Text Search instance services.

2. Ensure that the text search configuration for the text search enabled database matches the configuration for the DB2 instance. For a physical database partitioned environment, the SYSIBMTS.TSSERVERS administrative view must always include the actual host name.
 - a. Update the SYSIBMTS.TSSERVERS view with the Text Search server configuration data.
 - b. Call the SYSPROC.SYSTS_CONFIGURE procedure.
 - c. Issue the FLUSH PACKAGE CACHE DYNAMIC command to remove old data from the database cache.
3. Use the configTool to check that the integrated or stand-alone text search server is correctly configured for the DB2 instance.
4. Ensure that the correct communication port is specified, and that if multiple DB2 instances are running on the same system, the Text Search server for each instance has been configured to use a unique communication port. Verify proper function of the network connection.

CIE0757E Receive failed on host named *host-name* and port *port-number*.

Explanation: DB2 Text Search instance services have detected one or more communication errors.

A receive error on the named host and port occurred while responding to a request.

User response: Check the DB2 Text Search configuration.

Ensure that the communication port specified is correct.

If the error persists, start a DB2 trace and check the db2diag.log. Contact IBM software support for further assistance.

CIE0758E Send failed on host named *host-machine-name* and port *port-number*.

Explanation: DB2 Text Search instance services have detected one or more communication errors.

A send error occurred while sending a search request to the named host and port.

User response: Check the DB2 Text Search configuration.

Ensure that the communication port specified is correct.

If the error persists, start a DB2 trace and check the db2diag.log. Contact IBM software support for further assistance.

CIE0759E Communications timeout on host named *host-machine-name* and port *port-number*.

Explanation: DB2 Text Search instance services have detected one or more communication errors.

A request on the named host and port timed out.

User response: Check the DB2 Text Search configuration.

Ensure that the communication port specified is correct.

If the error persists, start a DB2 trace and check the db2diag.log. Contact IBM software support for further assistance.

CIE0760E Protocol error.

Explanation: A processing error occurred when calling an internal function. A protocol error was encountered while processing a text search administration operation or a text search query.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0761E Authentication failure.

Explanation: Authentication failed during a text search administration operation or a text search query.

User response: Ensure correct configuration and that the operation is being requested using the text search instance services associated with the DB2.

If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0762E Application error.

Explanation: A processing error occurred when calling an internal function. Server application error occurred while processing a text search administration operation or a text search query.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0763E Parse error at host *host-name* on port *port-number*.

Explanation: A processing error occurred when calling an internal function. A parse error occurred processing a response to a request on the named host and port.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0765E Unknown communication error.

Explanation: A processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0766E Invalid search client state.

Explanation: A processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0767E Invalid data *data*.

Explanation: A processing error occurred when calling an internal function. A Base64-encoded string contains invalid data.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0768E String *string* not aligned.

Explanation: The Base64-encoded string is either not aligned at a 4 character boundary, or not padded properly.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0769E Buffer size *size* too small.

Explanation: A processing error occurred when calling an internal function due to a small buffer size.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0770E *method-name* memory allocation failed.

Explanation: A processing error occurred when calling an internal function. Unable to allocate memory in the named method.

User response: If the error persists, start a trace and check the db2diag.log. Report the error to your IBM service representative.

CIE0771E The DB2 Text Search services might not be active.

Explanation: The DB2 Text Search services are not accessible. The most probable reason is the services have not been started yet.

User response: To start the DB2 Text Search services, execute the following command:

db2ts start for text

CIE0772E The DB2 Text Search services are in an inconsistent state and need to be restarted.

Explanation: The DB2 Text Search services are in an inconsistent state, it is necessary to restart the services.

User response: To bring DB2 Text Search services into a consistent state, issue the following commands:

db2ts stop for text
db2ts start for text

CIE0773E The text search authentication token is not valid.

Explanation: The DB2 Text Search authentication token is invalid and must be regenerated.

User response: To generate a valid authentication token, issue the following command:

configTool generateToken

CIE0774E The DB2 Text Search configuration could not be found. More diagnostic information is available in the db2diag log files.

Explanation: A failure occurred while attempting to access the DB2 Text Search configuration parameters.

User response: Respond to this error in one or more of the following ways:

- Refer to the db2diag log files for more details about this error.
- To view the DB2 Text Search configuration parameters, issue the following command:
configTool printAll

CIE0775E The DB2 Text Search configuration parameters could not be read from the configuration file. More diagnostic information is available in the db2diag log files.

Explanation: A failure occurred while attempting to read the DB2 Text Search configuration parameters.

User response: Respond to this error in one or more of the following ways:

- Refer to the db2diag log files for more details about this error.
 - Make sure the tmp subdirectory in the DB2 instance directory is available and has read and write permissions for the instance owner.
-

CIE0776E Restore failed for text search index
schema-name.index-name.

Explanation: A failure occurred while attempting to restore the DB2 Text Search index.

The most common reasons are:

1. The text search index is not configured for backup.
2. The restore settings in the text search backup configuration do not match with a supported restore mode in the Text Search Server.

User response: See the db2diag.log for more details about the error.

1. If the text search index is not configured for backup, enable a backup mode through the ALTER INDEX FOR TEXT operation to enable subsequent restores.
2. If the restore settings in the text search backup configuration are invalid, update the configuration to match with a supported restore mode in the Text Search server.

CIE0777E Failed to schedule the text search task.
Ensure the DB2 administrative task scheduler is properly setup.

Explanation: A failure occurred while attempting to add the Text Search task in the DB2 scheduler.

User response: Use the following instructions to verify the setup of the DB2 scheduler. Ensure:

- the DB2_ATS_ENABLE registry variable is enabled
- the SYSTOOLSPACE table space is available in the database
- the database is activated

Repeat the text search index operation to add the task in the scheduler.

CIE0778E Schedule with taskname *task-name*
already exists.

Explanation: A scheduler task with this taskname already exists. Schedule tasknames for text indexes are generated using the unique text index identifier, however, it is possible that an independently created manual schedule taskname causes a conflict.

User response: Repeat the operation, this results in a new index identifier and thus in a new schedule taskname. If the error persists, make sure that the TSSCHED_ prefix is reserved for the system-generated text index schedule tasknames.

CIE0779E Error processing DB2 Text Search server configuration.

Explanation: A failure occurred while attempting to process the DB2 Text Search configuration parameters.

User response: Check the validity of the information

added during or after the enablement of the database for text search in the SYSIBMTS.TSSERVERS administrative view like host, port or token.

Update the parameters that appear invalid by issuing the SYSTS_CONFIGURE procedure.

CIE0780E Cannot create collection *collection-name*
for text search index
schema-name.index-name.

Explanation: An error occurred while creating a collection for the text search index. Creating an index in a partitioned database environment might create multiple collections.

It is essential that all collections be created successfully for the CREATE INDEX operation to be successful.

User response: Check read and write permissions for the directories where the collections are being created. Verify the following:

- The defaultDataDirectory parameter is printed by configTool. Ensure that it is accessible to the user starting the text search server.
- Make sure there is sufficient disk space for the index being created.

See the db2diag.log for more details about this error.

CIE0781E DB2 Text Search is not supported in a
DB2 pureScale environment.

Explanation: DB2 Text Search is not supported in a DB2 pureScale environment.

User response: Reconfigure DB2 Text Search.

CIE0782E Incorrect number of active DB2 Text
Search servers.

Explanation: You are receiving this error due to the following inconsistency in the SYSIBMTS.TSSERVERS view:

- There are no rows in SYSIBMTS.TSSERVERS view.
- More than one row has SERVERSTATUS set to indicate an active server.

User response: Ensure that there is only one entry for an active text search server in the SYSIBMTS.TSSERVERS view.

Part 8. CLI Messages

This section contains the Call Level Interface (CLI) messages. The messages are listed in numeric sequence.

Chapter 54. CLI0000 - CLI0499

CLI0001W Disconnect error.

Explanation: An error occurred during the disconnect. However, the disconnect succeeded.

User response: Validate that the communication between the client and the server are still active.

CLI0002W Data truncated.

Explanation: An output buffer specified is not large enough to contain the data.

User response: Increase the size of the output buffer.

CLI0003W Privilege not revoked.

Explanation: The statement was a REVOKE statement and the user did not have the specified privilege.

User response: No action needed.

CLI0004W Invalid connection string attribute.

Explanation: An invalid or unsupported connection string attribute was specified in the connection string but the driver was able to connect to the data source anyway.

User response: No action needed.

CLI0005W Option value changed.

Explanation: The driver did not support the specified option value and substituted a similar value.

User response: No action needed.

CLI0006W SQLCancel treated like a close.

Explanation: The SQLCancel call was treated like a SQLFreeStmt call with the SQL_CLOSE option.

User response: No action needed.

CLI0008I Capture mode terminated.

Explanation: If the message is received at connect time when running in either capture or match modes, SQL statements are not captured into or matched from the capture file. This can be due to incorrectly configured mandatory db2cli.ini keywords. If the message is received at disconnect time, this is considered normal processing.

User response: Ensure that the mandatory static capture or static match db2cli.ini keywords (StaticMode, StaticPackage, StaticCapFile) are properly

configured as specified in the documentation.

CLI0100E Wrong number of parameters.

Explanation: The number of parameters specified in SQLSetParam or SQLBindParameter was less than number of parameters in the SQL statement.

User response: Respecify the SQL statement or provide more input parameters through SQLSetParam or SQLBindParameter.

CLI0101E The statement did not return a result set.

Explanation: The previous statement does not result in a result set.

User response: Respecify the SQL statement.

CLI0102E Invalid conversion.

Explanation: The conversion between the application data type and SQL data type is not supported by the driver.

User response: Respecify a data conversion which is supported by the driver.

CLI0103E Too many columns.

Explanation: The number of columns specified in SQLBindCol is greater than the number of columns in the current result set.

User response: Respecify the SQL statement or reset the bound column information in SQLBindCol.

CLI0104E Unable to connect to data source.

Explanation: The driver was unable to establish a connection with the data source.

User response: Ensure the server is started and that the communications between the client and server are correct.

CLI0105E Connection in use.

Explanation: The specified connection handle is already being used and the connection is still opened.

User response: Allocate a new connection using SQLAllocConnect and retry the connection or terminate the existing connection.

CLI0106E Connection is closed.

Explanation: The connection specified by the connection handle is no longer active.

User response: Establish a new connection.

CLI0107E Connection failure during transaction.

Explanation: The connection failed during the execution of the function and it cannot be determined whether the COMMIT or ROLLBACK occurred before the failure.

User response: Establish a new connection.

CLI0108E Communication link failure.

Explanation: The connection between the driver and the data source failed during execution of this function.

User response: Establish a new connection.

CLI0109E String data right truncation.

Explanation: The data specified through SQLSetParam or SQLBindParameter is larger than the maximum size allowed for the corresponding use of a parameter marker.

User response: Respecify the parameter using SQLSetParam or SQLBindParameter.

CLI0110E Invalid output or indicator buffer specified.

Explanation: The returned data was NULL but the output or indicator buffer specified was a NULL buffer.

User response: Respecify the output or indicator buffer supplying a non-NULL buffer and retry the operation.

CLI0111E Numeric value out of range.

Explanation: Returning the numeric data would have caused the whole part of the number to be truncated.

SQLPutData was called more than once for a parameter and the input data was not of type character or binary.

User response: Respecify the output bindings either through SQLBindCol or SQLGetData to avoid creating a numeric data truncation.

Do not call SQLPutData for a parameter if the application data type specified for that parameter through SQLSetParam or SQLBindParameter is not SQL_C_CHAR or SQL_C_BINARY.

CLI0112E Error in assignment.

Explanation: The data sent for a parameter or column was incompatible with the data type of the associated table column.

User response: Respecify the output binding through

SQLBindCol or SQLGetData or the input binding through SQLSetParam or SQLBindParameter.

CLI0113E Invalid date/time format.

Explanation: The data sent for a date/time field was invalid. An invalid datetime format was detected; that is, an invalid string representation or value was specified.

User response: Respecify the date data.

CLI0114E Datetime field overflow.

Explanation: The data sent for a date, time, or timestamp parameter, or column, was invalid.

User response: Respecify the date, time, or timestamp data.

CLI0115E Invalid cursor state.

Explanation: The statement is not positioned on a row.

User response: Position the statement on a row by calling SQLFetch or SQLExtendedFetch and retry the operation.

CLI0116E Invalid transaction state.

Explanation: There was a transaction in progress when SQLDisconnect was called.

User response: Call SQLTransact before calling SQLDisconnect.

CLI0117E Invalid cursor name.

Explanation: An invalid or duplicate cursor name was specified in SQLSetCursorName.

User response: Respecify a valid cursor name through SQLSetCursorName.

CLI0118E Invalid SQL syntax.

Explanation: An invalid or incorrect SQL statement was specified.

User response: Respecify a valid SQL statement.

CLI0119E Unexpected system failure.

Explanation: An unexpected system failure occurred while processing the function.

User response: Restart the application and try again.

CLI0120E Memory allocation failure.

Explanation: The driver was unable to allocate memory required to support execution or completion of the function.

User response: Verify that your system has enough memory to complete the desired operation.

CLI0121E Invalid column number.

Explanation: The value specified in iCol is less than zero, greater than the number of columns in the result set, or greater than the maximum number of columns allowed in a result set.

User response: Respecify a valid value for iCol.

CLI0122E Program type out of range.

Explanation: The value specified for fCType is not valid.

User response: Respecify a valid value for fCType.

CLI0123E SQL data type out of range.

Explanation: The value specified for fSQLType is not valid.

User response: Respecify a valid value for fSQLType.

CLI0124E Invalid argument value.

Explanation: The value specified for an argument was invalid. Possible reasons may be a null pointer, invalid length, invalid option, etc.

User response: Re-examine the argument passed in to the function and determine which argument is invalid.

CLI0125E Function sequence error.

Explanation: This function was called in an invalid sequence.

User response: Correct the sequence in your application and retry the operation.

CLI0126E Operation invalid at this time.

Explanation: The operation the system is trying to perform is invalid at this time.

User response: Correct the sequence of operation and retry the operation.

CLI0127E Invalid transaction code.

Explanation: The transaction option specified in SQLTransact was not SQL_COMMIT or SQL_ROLLBACK.

User response: Specify either SQL_COMMIT or SQL_ROLLBACK and retry the operation.

CLI0128E Unexpected memory handling error.

Explanation: Memory handling error.

User response: An unexpected error occurred in the driver while handling some internal memory buffer. Restart your application.

CLI0129E An attempt to allocate a handle failed because there are no more handles to allocate.

Explanation: A CLI handle is a variable that refers to a data object allocated and managed by DB2 CLI. There are four types of handles in CLI:

- Environment handle
- Connection handle
- Statement handle
- Descriptor handle

This message is returned when all handles that can be allocated have been allocated, and an attempt is made to allocate a handle using one of the following functions:

- SQLAllocEnv
- SQLAllocConnect
- SQLAllocStmt
- SQLAllocHandle
- SQLExecute
- SQLExecDirect

User response: Free handles that are no longer being used in the application by calling SQLFreeEnv, SQLFreeConnect, or SQLFreeStmt, or SQLFreeHandle.

To reduce the incidence of this error take one or both of the following actions:

- Modify the application to release handles more frequently.
- Increase the number of statement handles that can be allocated by increasing the CLIPkg CLI/ODBC configuration parameter using the BIND command.

CLI0130E No cursor name available.

Explanation: There was no open cursor on the statement and no cursor had been set with SQLSetCursorName.

User response: Specify a cursor name through SQLSetCursorName.

CLI0131E Invalid string or buffer length.

Explanation: The length of the buffer specified is invalid.

User response: Specify a valid buffer length.

CLI0132E Descriptor type out of range.

Explanation: The description type specified is not valid.

User response: Specify a valid descriptor type.

CLI0133E Option type out of range.

Explanation: The option type specified is not valid.

User response: Specify a valid option type.

CLI0134E Invalid parameter number.

Explanation: The number specified for the parameter number is less than zero or greater than the maximum parameter supported by the data source.

User response: Specify a valid parameter number.

CLI0135E Invalid scale value.

Explanation: The scale value specified is invalid.

User response: Specify a valid scale value.

CLI0136E Function type out of range.

Explanation: The function type is invalid.

User response: Specify a valid function type value.

CLI0137E Information type out of range.

Explanation: The information type is invalid.

User response: Specify a valid information type value.

CLI0138E Column type out of range.

Explanation: The column type is invalid.

User response: Specify a valid column type value.

CLI0139E Scope type out of range.

Explanation: The scope type is invalid.

User response: Specify a valid scope type value.

CLI0140E Uniqueness option type out of range.

Explanation: The uniqueness option type is invalid.

User response: Specify a valid uniqueness option type value.

CLI0141E Accuracy option type out of range.

Explanation: The accuracy option type is invalid.

User response: Specify a valid accuracy option type value.

CLI0142E Direction option out of range.

Explanation: The direction option is invalid.

User response: Specify a valid direction option value.

CLI0143E Invalid precision value.

Explanation: The precision value is invalid.

User response: Specify a valid precision value.

CLI0144E Invalid parameter type.

Explanation: The parameter type is invalid.

User response: Specify a valid parameter type value.

CLI0145E Fetch type out of range.

Explanation: The fetch type is invalid.

User response: Specify a valid fetch type value.

CLI0146E Row value out of range.

Explanation: The row value is invalid.

User response: Specify a valid row value.

CLI0147E Concurrency option out of range.

Explanation: The concurrency option is invalid.

User response: Specify a valid concurrency option value.

CLI0148E Invalid cursor position.

Explanation: The cursor position is invalid.

User response: Specify a valid cursor position value.

CLI0149E Invalid driver completion.

Explanation: The driver completion is invalid.

User response: Specify a valid driver completion value.

CLI0150E Driver not capable.

Explanation: The operation is valid but not supported by either the driver or the data source.

User response: Specify a valid operation.

CLI0151E No data pending.

Explanation: SQLParamData or SQLPutData was called but there is no data at execute pending on this statement.

User response: Respecify the parameters through SQLSetParam or SQLBindParameter.

CLI0152E Not a string value.

Explanation: The function expected a string argument to be supplied.

User response: Respecify the argument to the function.

CLI0153E Invalid file name length.

Explanation: The length specified for the file name is invalid.

User response: Specify a valid file name length.

CLI0154E Invalid connection state.

Explanation: An attempt to change the connection type was made but the connection or another connection (in the case of Distributed Unit of Work) is already opened.

User response: Do not attempt to change the connection type once the connection is opened.

CLI0155E File name length is too long.

Explanation: The length supplied for the file name is greater than supported.

User response: Specify a valid file name length.

CLI0156E Error closing a file.

Explanation: An unexpected condition occurred while closing a file.

User response: Retry the operation.

CLI0157E Error opening a file.

Explanation: An unexpected condition occurred while opening a file.

User response: Retry the operation.

CLI0158E Error writing to a file.

Explanation: An unexpected condition occurred while writing to a file.

User response: Retry the operation.

CLI0159E Error deleting a file.

Explanation: An unexpected condition occurred while deleting a file.

User response: Retry the operation.

CLI0164E Nullable type out of range.

Explanation: An invalid nullable value was specified.

User response: Respecify the nullable value.

CLI0165E Error in row.

Explanation: An error occurred while fetching one or more rows. (Function returns SQL_SUCCESS_WITH_INFO.)

User response: Retry the operation.

CLI0166E PARMLIST syntax error.

Explanation: The PARMLIST value in the stored procedures catalog table contains a syntax error.

User response: Respecify the row for this stored procedure.

CLI0167E Operation was cancelled.

Explanation: SQLCancel was called on the statement.

User response: Resubmit the operation.

CLI0171E The database was unable to connect.

Reason: *reason-text*.

Explanation: The database connection failed. The 'Reason:' will explain why the database was unable to connect.

User response: Fix the problem, and try the connect again.

CLI0172E The database was unable to connect. Would you like to continue with the configuration? Reason: *reason-text*.

Explanation: The database connection failed. The 'Reason:' will explain why the database was unable to connect.

User response: Fix the problem, and try the connect again or continue with the configuration.

CLI0175E No ODBC Driver Manager has been found. Cannot open ODBCINST.INI file.

Explanation: Neither one of the Visigenic or the Intersolv ODBC Driver Managers have been installed. An ODBC Driver Manager must be present before the DB2 ODBC Driver can be registered into the associated ODBCINST.INI file.

User response: Install an ODBC Driver Manager, and try running this command again.

CLI0176E Unable to write to the ODBCINST.INI file.

Explanation: An error has occurred while writing to an ODBCINST.INI file to be used with an ODBC Driver Manager. This may be due to invalid format of the original ODBCINST.INI file.

User response: Rename the existing ODBCINST.INI file to something else and try running this command again. If this error persists, contact your technical support.

CLI0177E Cannot determine where the IBM Data Server Client is installed.

Explanation: An error has occurred while trying to locate where the IBM Data Server Client is installed.

User response: Check that the IBM Data Server Client has been correctly installed.

CLI0178I The DB2 ODBC Driver has been successfully registered.

Explanation: The DB2 ODBC Driver has been registered with the installed ODBC Driver Manager(s).

User response: The user can now use the appropriate ODBC Administrator tool from the ODBC Driver Manager vendor to configure the ODBC data sources.

CLI0179E Unable to display dialog box.

Explanation: The DB2 ODBC Driver cannot open a dialog box with the window handle provided by the application.

User response: The application must pass a valid window handle.

CLI0180E Invalid bookmark value.

Explanation: The argument fFetchOrientation was SQL_FETCH_BOOKMARK and the bookmark pointed to by the value SQL_ATTR_FETCH_BOOKMARK_PTR statement attribute was not valid.

User response: Re-specify a valid bookmark value.

CLI0181E Invalid descriptor index.

Explanation: The value specified for the column number argument is invalid.

User response: Re-specify a valid column number.

CLI0182W Fractional truncation.

Explanation: The data for one of the columns was truncated.

User response: None needed.

CLI0183E Associated statement is not prepared.

Explanation: The descriptor handle was associated with an IRD, and the associated statement handle was not in a prepared state.

User response: Prepare the statement associated with the descriptor.

CLI0184E Cannot modify an implementation row descriptor.

Explanation: The descriptor handle was associated with an IRD, and IRD record fields cannot be updated.

User response: Specify a valid descriptor and field.

CLI0185E Invalid use of an automatically allocated descriptor handle.

Explanation: Only explicit descriptors allocated using the SQLAllocHandle function can be used in this manner.

User response: Use an explicitly allocated descriptor.

CLI0186E Server declined cancel request.

Explanation: Communication errors caused the server to decline the cancel request.

User response: None needed.

CLI0187E Non-character and non-binary data sent in pieces.

Explanation: SQLPutData was called more than once for a column that was not a character or binary column.

User response: Only call SQLPutData once for non-character and non-binary data.

CLI0188E Attempt to concatenate a null value.

Explanation: A previous call to SQLPutData for this parameter specified an input buffer length of SQL_NULL_DATA

User response: Call SQLPutData with a valid input buffer length or ensure SQLPutData is called the correct number of times.

CLI0189E Inconsistent descriptor information.

Explanation: The associated descriptor failed a consistency check.

User response: Ensure that descriptor fields are valid, and all interdependent fields are set correctly.

CLI0190E Descriptor invalid on indirect reference.

Explanation: The descriptor handle was invalid, or the TYPE, OCTECT_LENGTH, DATA_PTR, INDICATOR_PTR, or OCTECT_LENGTH_PTR field of the ARD contained an invalid value.

User response: Ensure the indirect descriptor handle is valid and that the deferred fields are set correctly.

CLI0191E Invalid attribute value.

Explanation: The value given was not correct for the specified attribute.

User response: Use a correct value for the specified attribute.

CLI0192E Invalid database alias specified.

Explanation: You have specified an invalid or non-existent database alias name through the -d parameter of the DB2CAP command.

User response: Specify the name of an existing database alias. To find out which databases are currently catalogued, use the LIST DATABASE DIRECTORY command.

CLI0193E Capture file *capture-file-name* not found or in use.

Explanation: DB2CAP is unable to open the file <capture-file-name> that you specified as a BIND parameter value. Either the file does not exist or some other process is using it exclusively.

User response: Specify the name of an existing capture file or wait until the file is released by the other process.

CLI0194E Command syntax error. The correct syntax is: db2cap [-h | -?] bind capture-file -d db-alias [-u userid [-p password]]

Explanation: The parameters of the DB2CAP command, as well as the command name itself, must be typed exactly as shown in the syntax diagram provided. They are case-sensitive on UNIX platforms and case-insensitive on Intel platforms. There must be at least one space between the parameter and its value. For example: -u userid is correct, but -userid is incorrect.

In the syntax diagram, square brackets ([]) denote optional parameters. If you specify the userid but omit the password, you will be prompted to enter the password, in which case it won't be visible when you type it.

If you specify the -h or -? parameter, the command will

display the syntax help and no other parameter will be processed.

User response: Make sure the syntax of the command is as described in the diagram.

CLI0195E Invalid capture file: does not contain any valid statement groups.

Explanation: A valid capture file must contain at least one valid statement group.

User response: Re-capture the application using the same capture file and make sure that at least one SQL statement was captured, or edit the capture file and add statement groups manually. When adding a statement group manually, make sure you specify all the required keywords for the type of SQL statement you are adding. Then submit your request again.

CLI0196E Switch user operation invalid at this time.

Explanation: The application cannot switch the user within a unit of work.

User response: Correct the sequence of operation. The application must switch the user before beginning a unit of work.

CLI0197E A trusted context is not enabled on this connection. Invalid attribute value.

Explanation: The attribute SQL_ATTR_TRUSTED_CONTEXT_USERID or SQL_ATTR_TRUSTED_CONTEXT_PASSWORD cannot be set to reuse a trusted connection before an initial trusted connection has been established.

User response: Ensure that the SQL_ATTR_USE_TRUSTED_CONTEXT has been set to enable the trusted context and that a trusted connection has been established.

CLI0198E Missing trusted context userid.

Explanation: The attribute, SQL_ATTR_TRUSTED_CONTEXT_PASSWORD was set without providing the required attribute, SQL_ATTR_TRUSTED_CONTEXT_USERID.

User response: Ensure that when setting the attribute, SQL_ATTR_TRUSTED_CONTEXT_PASSWORD that the SQL_ATTR_TRUSTED_CONTEXT_USERID attribute is also set.

CLI0199E Invalid connection string attribute.

Explanation: An invalid or unsupported connection string attribute was specified in the connection string. The driver was unable to establish a connection with the data source.

User response: Specify a valid connection string attribute and retry.

CLI0200E Invalid DSN value.

Explanation: The value attempting to be used for the DSN in the connection string or connection API is invalid.

The driver did not attempt a connection with the data source.

User response: Specify a valid DSN and retry.

CLI0201E Invalid database value.

Explanation: The database name or database alias in the connection string or connection API is invalid.

The driver did not attempt a connection with the data source.

User response: Specify a valid database name and retry.

CLI0202E Error opening the CLI Client Optimization feature capture file.

Explanation: CLI Client Optimization feature is unable to open the file specified in the pureQueryXML CLI parameter value. Either the file does not have the correct permissions or some other process is using it exclusively.

The CLI Client Optimization feature is enabled when either "captureMode" or "executionMode" CLI keywords are specified.

The pureQueryXML CLI keyword specifies the capture file location using either an absolute path or a relative path to the directory location where the CLI application is running.

During the capture phase of the Client Optimization feature, the capture file can either exist or not exist in the location specified. If the file exists, the Client Optimization feature will require read and write permission to the capture file. If the files does not exist, the Client Optimization feature will require write permission to the directory to create the file.

During the match phase of the Client Optimization feature, the capture file must exist and the Client Optimization feature will require read permission to the specified file.

CLI will return an error if any of these permissions are not available.

User response: Specify the name of a capture file with appropriate permissions for the pureQueryXML CLI keyword or disable the Client Optimization feature.

CLI0203E Client Optimization feature is not supported when static profiling is enabled in the DB2 CLI driver.

Explanation: The DB2 CLI driver static profiling function cannot be used with the Client Optimization feature. CLI does not allow enabling both Client Optimization and the static profiling function simultaneously for a given database connection.

User response: Use the Client Optimization feature or the DB2 CLI driver static profiling function. Client Optimization can be disabled by removing the captureMode and executionMode CLI keywords. DB2 CLI driver static profiling can be disabled by removing the StaticMode CLI keyword.

CLI0204E Dynamic execution of SQL statement is not allowed.

Explanation: The Client Optimization feature failed to find a matching SQL statement in the pureQueryXml capture file when executionMode CLI keyword is set to STATIC and the allowDynamicSQL CLI keyword value is FALSE. Dynamic execution is not allowed with this configuration.

If a CLI application is expected to run any SQL statement dynamically, specifying a value of TRUE for allowDynamicSQL will enable an SQL statement to run dynamically if no matching statement found in the pureQueryXml capture file.

Alternatively, an incremental capture phase may be used to refresh the pureQueryXml capture file to include any missing SQL statement. This capture phase must be followed by a Configure phase and a StaticBinder phase of the Client Optimization feature before the application can be run with an executionMode value of STATIC.

User response: Specify a value of TRUE for allowDynamicSQL keyword or refresh the pureQueryXml capture file to include the SQL statement.

CLI0205E Cannot find library db2clxml4c.

Explanation: db2clxml4c cannot be found in the library path.

The library path is defined by the environment variable for the following operating systems:

AIX (Java 1.1)

LD_LIBRARY_PATH

AIX (Java 1.2 or later)

LIBPATH

HP-UX

SHLIB_PATH

Linux

LD_LIBRARY_PATH

Silicon Graphics IRIX

LD_LIBRARY_PATH

Solaris Operating Environment

LD_LIBRARY_PATH

Windows operating systems

PATH

The db2clxml4c library is dynamically loaded if a CLI application has requested to use the Client Optimization feature. CLI attempts to locate this library from the library path. The DB2 CLI drivers include this library in the default library path when installed. If the db2clxml4c library cannot be located, the Client Optimization feature may be disabled by removing the captureMode and executionMode CLI keywords.

User response: Ensure your application environment is configured correctly to use DB2. On UNIX platforms, ensure that the sqllib/db2profile script has been run to set your environment variables.

CLI0206E Unable to read an element of the pureQueryXml capture file.

Explanation: The Client Optimization feature reads the pureQueryXml capture file. If one or more entries in this file are corrupted or invalid, the Client Optimization feature cannot proceed.

Any manual changes made to pureQueryXml capture file, may cause this error. If the file is not modified and it was completely generated by the CLI driver or the Configure tool, please contact IBM technical support for further help.

User response: Ensure that the pureQueryXml capture file is correct or recapture the pureQueryXml capture file.

CLI0207E Invalid file extension for pureQueryXml capture file.

Explanation: The pureQueryXml capture file must have a file extension of ".xml" or ".pdqxml".

User response: Specify a file with a supported file extension for the pureQueryXml CLI keyword.

CLI0208E The codepage could not be set for this connection because the connection is part of an XA environment.

Explanation: Setting the codepage at the connection level is not supported in an XA environment.

This message is returned when an application attempts to set the connection attribute

SQL_ATTR_CLIENT_CODEPAGE after a call to xa_open has been made.

The codepage was not changed.

User response: Respond to this message in one of the following ways: To open the connection without changing the codepage, reopen the connection without specifying the SQL_ATTR_CLIENT_CODEPAGE connection attribute. To change the codepage, use the DB2CODEPAGE environment variable.

CLI0209E The application failed to execute the query because the capturedOnly keyword is set to TRUE but client optimization found no matching statements in the pureQueryXML capture file.

Explanation: You can improve SQL and XQuery performance by collecting SQL and XQuery statements into pureQueryXML capture files.

To run only the matching statements in the pureQueryXML files, you can set the capturedOnly keyword to TRUE. SQL statements that are not in the pureQueryXML file will not be executed when the capturedOnly keyword to TRUE.

This message is returned when the capturedOnly keyword is set to TRUE, but there are no matching statements in the pureQueryXML file.

User response: To use client optimization with a pureQueryXML file, perform the following steps:

1. Capture SQL and XQuery statements into a pureQueryXML file by performing the following steps:
 - a. Begin capturing mode by setting the captureMode property to ON.
 - b. Run all logical paths that contain the SQL or XQuery statements that you want to capture.
 - c. End capturing mode by setting the captureMode property to OFF.
2. Re-run the application.

CLI0210E The statement could not be executed because the code page of the parameter markers or result set columns does not match the code page that is currently specified in the SQL_ATTR_OVERRIDE_CHARACTER_CODEPAGE connection attribute.

Explanation: In general, when the code page of a client application differs from the code page of the database server to which the application is connected, character conversion happens automatically when data is inserted or retrieved.

You can override this automatic behavior for parameterized insert and update operations and for

fetching query results using the connection attribute `SQL_ATTR_OVERRIDE_CHARACTER_CODEPAGE`. When the code page of a statement's parameter markers or result set columns matches the value of `SQL_ATTR_OVERRIDE_CHARACTER_CODEPAGE`, no character conversion happens when that statement is executed.

This message is returned when an attempt is made to execute a statement for which the describe information for parameter markers or result set columns does not match the code page to which the connection attribute `SQL_ATTR_OVERRIDE_CHARACTER_CODEPAGE` is set.

User response: To insert or retrieve the data without character conversion, perform the following steps:

1. Determine the code page indicated by the describe information for the parameter markers or result set columns of the current statement.
2. Free any statement handles associated with the current connection.
3. Set the connection attribute `SQL_ATTR_OVERRIDE_CHARACTER_CODEPAGE` to the same code page as the code page indicated by the describe information for the parameter markers or result set columns of the current statement.
4. Reallocate the current statement handle.
5. Execute the statement again.

CLI0211E Configuring the specified property failed because the specified value is incompatible with existing property settings.

Explanation: You can improve the performance of your database application using pureQuery client optimization. For example, you can use pureQuery client optimization to substitute an SQL statement that executes very quickly for an SQL statement that executes slowly without rewriting your application.

The behavior of the pureQuery client optimization feature is controlled by a collection of properties. For example, to enable the substitution of SQL statements, you can set the `enableDynamicSQLReplacement` property to TRUE. This message is returned when an attempt is made to set one of these properties to a value that is not compatible with other, existing property values.

This message is returned when an attempt is made to configure pureQuery client optimization in one of the following incompatible combinations:

- An attempt was made to enable both dynamic SQL replacement and capture mode.
Dynamic SQL replacement is not supported when pureQuery is in capture mode. Specifically, `enableDynamicSQLReplacement` cannot be set to TRUE when `captureMode` is set to ON, and

`captureMode` cannot be set to ON when `enableDynamicSQLReplacement` is set to TRUE.

- An attempt was made to enable both static execution mode and capture mode.

Static execution mode is not supported when pureQuery is in capture mode. Specifically, `executionMode` cannot be set to STATIC when `captureMode` is set to ON, and `captureMode` cannot be set to ON when `executionMode` is set to STATIC.

User response: Respond to this message in one of the following ways:

- If you do not need to use the pureQuery client optimization feature, disable the feature by setting the following configuration:
`captureMode=OFF` [or keep it unset]
`executionMode=DYNAMIC` [or keep it unset]
- To use the pureQuery client optimization feature, use the appropriate compatible combination of pureQuery client optimization feature properties:
 1. Enable either dynamic SQL replacement or capture mode, but not both:
 - Enable dynamic SQL replacement with the following property settings:
`enableDynamicSQLReplacement=TRUE`
`captureMode=OFF` [or keep it unset]
 - Enable capture mode with the following property settings:
`captureMode=ON`
`enableDynamicSQLReplacement=FALSE`
[or `enableDynamicSQLReplacement` unset]
 2. Enable capture mode or use static execution mode, but not both:
 - Enable capture mode with the following property settings:
`captureMode=ON`
`executionMode=DYNAMIC` [or keep it unset]
 - Use static execution mode with the following property settings:
`captureMode=OFF` [or keep it unset]
`executionMode=STATIC`

CLI0212W A seamless failover occurred during an execute request.

Explanation: A seamless failover occurred during an execute request from member to member or group to group. However, the execute request was completed successfully.

User response: No action required.

CLI0213E Retrieving the pureQueryXML file from the repository failed.

Explanation: pureQuery runtime uses the runtime group ID and the connection information to determine

the repository from where to retrieve the pureQueryXML file.

The reasons why the pureQueryXML file that you indicated by using the pureQueryXmlRepository or propertiesGroupId property cannot be retrieved from the repository can be:

- The repository indicated by the pureQueryXmlRepository property is not valid or is not accessible.
- The runtime group id specified by the propertiesGroupId property is not a valid runtime group id or is not activated.

User response: Perform all the actions that apply to your case:

- Ensure that you have access to the repository specified in the pureQueryXmlRepository property.
- Ensure that the runtime group id specified in the propertiesGroupId property is activated.
- If necessary, adjust the values of the pureQueryXmlRepository or propertiesGroupId property so that the pureQueryXML file can be accessed from the repository.

CLI0214E The propertiesGroupId property was not specified in the configuration file.

Explanation: The propertiesGroupId property was not set to the runtime group ID defined in the repository for the pureQueryXML file.

pureQuery Runtime uses the runtime group ID and the connection information to determine the repository from where to retrieve the pureQueryXML file. If the propertiesGroupId is not specified, the pureQueryXML file cannot be retrieved from the repository.

User response: Set the propertiesGroupId property to the runtime group id defined in the repository where the pureQueryXML file is stored.

CLI0215E The connection failed because pureQuery client optimization attempted to access the pureQueryXML capture file, but a valid file name was not specified with the pureQueryXML property.

Explanation: You can improve the performance of your database application using pureQuery client optimization. For example, you can capture information about SQL statements, as the statements run, into a pureQueryXML capture file. That capture file (also known as the pureQueryXML file) can later be used to execute the same SQL statements with improved performance. The name of the pureQueryXML capture file is specified by using the pureQueryXML property.

Client optimization can be configured and controlled by using a collection of pureQuery-related properties, including the following examples:

- enableDynamicSQLReplacement - Setting enableDynamicSQLReplacement to "TRUE" causes alternative SQL statements that are in the pureQueryXML file to be executed.
- executionMode - Setting executionMode to "STATIC" causes matching SQL statements that are in the pureQueryXML file to be executed statically.

This message is returned when an attempt is made to use pureQuery client optimization functionality that requires the pureQueryXML file but the name of the pureQueryXML file has not been specified.

User response: Specify the name of the pureQueryXML file with the pureQueryXML property, and then run the application again.

CLI0217E The pureQuery client optimization feature was unable to use the specified pureQueryXML capture file because the version of the capture file is not supported by CLI.

Explanation: You can improve SQL and XQuery performance by collecting information about SQL and XQuery statements into pureQueryXML capture files. The version of the pureQueryXML capture file is determined by the method used to generate the capture file. Here are examples of ways to generate a pureQueryXML capture file:

- Using the db2cap command
- Using the pureQuery utility called GeneratePureQueryXml

DB2 for Linux, UNIX, and Windows Call Level Interface (CLI) does not support pureQueryXML capture files that are generated by the GeneratePureQueryXml utility.

This message is returned when a DB2 CLI application attempts to use pureQuery client optimization with a pureQueryXML capture file that was generated by the GeneratePureQueryXml utility.

User response: Regenerate a pureQueryXML capture file using a method other than the GeneratePureQueryXml utility.

CLI0219E Certificate-based authentication failed because the DB2 client or data server driver was configured to use certificate-based authentication, but a password was also specified with the connection.

Explanation: You can use SSL client authentication, based on just a user ID, with DB2 database servers using certificate-based authentication. Certificate-based authentication allows you to use SSL client authentication without the need to store and maintain database passwords on the database client.

You can use certificate-based authentication by specifying configuration parameters such as SSLClientKeystash or SSLClientKeystoreDBPassword in your CLI configuration file, db2cli.ini, or in your data server driver configuration file, db2dsdriver.cfg.

When certificate-based authentication is configured to supply authentication information, a password cannot be specified in other ways (such in the db2dsdriver.cfg configuration file, in the db2cli.ini configuration file, or in the connection string.) This message is returned when certificate-based authentication is configured to supply authentication information, and a password is also specified in another way.

User response: To authenticate using certificate-based authentication, perform the following steps:

1. Specify either SSLClientKeystash or SSLClientKeystoreDBPassword in the CLI configuration file or the data server driver configuration file.
2. Verify that no password is specified in the db2dsdriver.cfg configuration file, in the db2cli.ini configuration file, or in the connection string.

CLI0220E **Certificate-based authentication failed because two mutually exclusive configuration parameters were both specified.**

Explanation: You can use certificate-based authentication by specifying either the SSLClientKeystash configuration parameter or the SSLClientKeystoreDBPassword configuration parameter in your CLI configuration file, db2cli.ini, or in your data server driver configuration file, db2dsdriver.cfg or in the connection string.

The configuration parameters SSLClientKeystash and SSLClientKeystoreDBPassword are mutually exclusive. This message is returned when the SSLClientKeystash configuration parameter and the SSLClientKeystoreDBPassword configuration parameter are both specified in either the CLI configuration file or the data server driver configuration file.

User response: To authenticate using certificate-based authentication, specify either the SSLClientKeystash configuration parameter or the SSLClientKeystoreDBPassword configuration parameter in the CLI configuration file, db2cli.ini, or in your data server driver configuration file, db2dsdriver.cfg or in the connection string.

CLI0221E **Certificate-based authentication failed because the DB2 client or data server driver was configured to use certificate-based authentication, but the SSLClientLabel parameter was not specified.**

Explanation: You can use certificate-based

authentication by specifying the SSLClientLabel parameter in your CLI configuration file, db2cli.ini, or in your data server driver configuration file, db2dsdriver.cfg or in the connection string.

When certificate-based authentication is configured to supply authentication information, the SSLClientLabel parameter must be specified in the db2dsdriver.cfg configuration file or in the db2cli.ini configuration file or in the connection string. This message is returned when certificate-based authentication is configured to supply authentication information, but the SSLClientLabel parameter is missing or invalid.

User response: To authenticate using certificate-based authentication, specify the SSLClientLabel parameter in the CLI configuration file or in the data server driver configuration file or in the connection string.

CLI0222E **Authentication failed because the SSLClientLabel parameter was specified but the DB2 client or data server driver was not configured to use certificate-based authentication.**

Explanation: You can use certificate-based authentication by setting the authentication keyword to "CERTIFICATE" in your connection string, in your CLI configuration file, db2cli.ini, or in your data server driver configuration file, db2dsdriver.cfg.

When certificate-based authentication is not configured to supply authentication information, the SSLClientLabel parameter should not be specified. This message is returned when certificate-based authentication is not configured to supply authentication information, but the SSLClientLabel parameter is set in the connection string or in the db2cli.ini configuration file or db2dsdriver.cfg configuration file.

User response: To authenticate using certificate-based authentication, perform the following steps:

1. Set the authentication parameter to CERTIFICATE.
2. Specify SSLClientLabel in the CLI configuration file or in the data server driver configuration file.

To use an authentication method other than certificate-based authentication, remove the SSLClientLabel parameter from the connection string, db2cli.ini configuration file, and the db2dsdriver.cfg configuration file.

CLI0223E **The SQLReloadConfig function call was not processed because another execution of the SQLReloadConfig function or the SQLSetConnectAttr function with SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute is already in progress.**

Explanation: You can reload a configuration property

from the IBM data server driver configuration file by calling the SQLReloadConfig function.

Only one instance of the SQLReloadConfig function process can run at one time. This message is returned when an attempt is made to execute the SQLReloadConfig function or call the SQLSetConnectAttr function with SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute while another instance of the SQLReloadConfig function is running.

User response: Wait for the instance of the SQLReloadConfig function or the SQLSetConnectAttr function with SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute that is already running to stop, and then call the SQLReloadConfig function again.

CLI0224E The SQLReloadConfig function call failed because an invalid value was specified for the ConfigProperty argument.

Explanation: You can reload configuration properties from a specified section of the IBM data server driver configuration file by calling the SQLReloadConfig function. You can identify which section to reload by specifying the ConfigProperty argument.

This message is returned when a value other than DSD_ACR_AFFINITY is specified for the ConfigProperty argument.

User response: Call the SQLReloadConfig function again, specifying valid values for all arguments.

CLI0225E The SQLReloadConfig function call failed because the SQLReloadConfig function cannot access the IBM data server driver configuration file.

Explanation: You can reload a configuration property from the IBM data server driver configuration file by calling the SQLReloadConfig function.

The SQLReloadConfig function attempts to locate the IBM data server driver configuration file in several places:

- The directory specified in the DB2DS DRIVER_CFG_PATH registry variable
- The default IBM data server driver configuration file location for the driver type and platform

This message is returned when either the SQLReloadConfig function is unable to locate the IBM data server driver configuration file, or when the SQLReloadConfig function is unable to read the IBM data server driver configuration file.

User response: Perform the following troubleshooting steps:

1. If the DB2DS DRIVER_CFG_PATH registry variable is set, verify that the IBM data server driver configuration file is located in the directory specified in the DB2DS DRIVER_CFG_PATH registry variable.
2. If the DB2DS DRIVER_CFG_PATH registry variable is not set, verify that the IBM data server driver configuration file is located in the default location for the driver type and platform.
3. Verify that the user id that was used to execute the application has read access for the IBM data server driver configuration file.

CLI0226E The SQLReloadConfig function failed because sections in the IBM data server driver configuration file other than the section that was specified in the ConfigProperty argument have been updated.

Explanation: You can reload configuration properties from a specified section of the IBM data server driver configuration file by calling the SQLReloadConfig function. You can identify which section to reload by specifying the ConfigProperty argument.

This message is returned when the SQLReloadConfig function detects that sections other than the section that was specified with the ConfigProperty argument have been updated.

User response:

1. Return the properties in the IBM data server driver configuration file to values that match the currently active configuration.
2. Update only those sections in the IBM data server driver configuration file that are supported with the SQLReloadConfig function.
3. Execute the SQLReloadConfig function again, specifying the updated sections in ConfigProperty argument.

CLI0227E The SQLReloadConfig function failed because the CLI subsystem is not initialized.

Explanation: You can reload a configuration property from the IBM data server driver configuration file by calling the SQLReloadConfig function.

An environment handle must be allocated before the SQLReloadConfig function can be called.

This message is returned when the SQLReloadConfig function is called before an environment handle has been allocated.

User response: Allocate an environment handle by calling the SQLAllocHandle function before calling the SQLReloadConfig function

CLI0228E The SQLReloadConfig function failed because the IBM data server driver configuration file is missing one or more of the required sections. Database: *database-name*. Host name: *host-name*. Port number: *port-number*.

Explanation: You can reload a configuration property from the IBM data server driver configuration client configuration file by calling the SQLReloadConfig function.

The SQLReloadConfig function requires a subset of sections to be configured in the IBM data server driver configuration file. The list of required sections includes the following sections: <client>, <affinitylist>, and <alternateserverlist>.

This message is returned when an attempt is made to call the SQLReloadConfig function when the required sections in the IBM data server driver configuration file are not configured.

User response: Configure all of the sections in the IBM data server driver configuration file that are required by the SQLReloadConfig function, and then call SQLReloadConfig again.

CLI0229E The SQLReloadConfig function failed because the IBM driver or client cannot establish a connection to any of the servers that are specified in the <alternateserverlist> section of the IBM data server driver configuration file. Database name: *database-name*. Host name: *host-name*. Port number: *port-number*.

Explanation: You can reload a configuration property from the IBM data server driver configuration file by calling the SQLReloadConfig function.

The SQLReloadConfig function requires that the IBM client or driver can connect to the servers that are listed in the <alternateserverlist> section of the IBM data server driver configuration file.

This message is returned when an attempt is made to call the SQLReloadConfig function when the IBM client or driver cannot connect to any of the servers that are listed in the <alternateserverlist> section of the IBM data server driver configuration file.

User response:

1. For any server to which the IBM client or driver cannot connect, resolve the problems that are preventing the IBM client or driver from connecting.
2. Update the <alternateserverlist> section of the IBM data server driver configuration file to specify only server to which the IBM client or driver can connect.
3. Call SQLReloadConfig again.

CLI0230E The SQLReloadConfig function failed because the SQLReloadConfig function failed to find an appropriate port number for a service name. Database name: *database-name*. Host name: *host-name*. Port number: *port-number*. Service name: *service-name*.

Explanation: You can reload a configuration property from the IBM data server driver configuration file by calling the SQLReloadConfig function.

The SQLReloadConfig function attempts to look up the port number for specified service names.

This message is returned when the SQLReloadConfig function cannot look up the port number for a specified service name.

User response: Specify only service names that are available for lookup.

CLI0231W The SQLReloadConfig function succeeded. However, the IBM client or driver could not access one or more of the servers that are specified in the <alternateserverlist> section of the IBM data server driver configuration file. Database name: *database-name*. Host name: *host-name*. Port number: *port-number*. List of unreachable servers, by host name and port number: *host-name-and-port-number-pairs*.

Explanation: You can reload a configuration property from the IBM data server driver configuration file by calling the SQLReloadConfig function.

The SQLReloadConfig function requires that the IBM client or driver can connect to every server listed in the <alternateserverlist> section of the IBM data server driver configuration file.

This message is returned when the SQLReloadConfig function successfully reloads configuration information, but the IBM client or driver cannot connect to one or more of the servers that are listed in the <alternateserverlist> section of the IBM data server driver configuration file.

User response: Optional:

1. Determine to which servers the IBM client or driver is unable to connect.
2. For any server to which the IBM client or driver cannot connect, resolve the problems that are preventing the IBM client or driver from connecting.
3. Update the <alternateserverlist> section of the IBM data server driver configuration file to specify only server to which the IBM client or driver can connect.
4. Call SQLReloadConfig again.

CLI0232E Attempt to set new primary member for the client affinities configuration has failed as client affinities feature has not been enabled.

Explanation: The client affinities element entries are not found in the IBM data server driver configuration file. The client affinities feature must be configured in the IBM data server driver configuration file before you can override the primary member with the SQLSetConnectAttr() API and the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute.

User response: Create the required client affinities element entries in the IBM data server driver configuration file and rerun the application.

CLI0233E The CLI driver is unable to establish a connection to the new primary member.

Explanation: The new primary member that is set with the SQLSetConnectAttr() API and the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute is not reachable.

User response: Specify another member in the <clientaffinitydefined> subsection of the IBM data server driver configuration file and call the SQLSetConnectAttr() API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute.

CLI0234E New unknown primary member cannot be specified for the client affinities configuration.

Explanation: The new primary member that is specified through the SQLSetConenctAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute was not present in the <alternateserverlist> subsection of the IBM data server driver configuration file.

User response: Specify existing member in the <alternateserverlist> subsection of the IBM data server driver configuration file and call the SQLSetConnectAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute or the SQLReloadConfig API.

CLI0235E A concurrent attempt to set new primary member for the client affinities configuration failed.

Explanation: An application might have attempted to set new primary member by calling the SQLSetConnectAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute in different threads. Also, an application might have attempted to set new primary member by calling SQLSetConnectAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute

at the same time as calling the SQLReloadConfig API.

User response: Ensure that only one attempt to set the new primary member is made at any one time. The new primary member can be set by calling the SQLSetConnectAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute or calling the SQLReloadConfig API after modifying the IBM data server driver configuration file.

CLI0236W The NULL value that was specified for the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY connection attribute has been ignored.

Explanation: You can use the client affinities feature to specify an ordered list of members to which the automatic client reroute feature will redirect application connections. You can configure the client affinities feature in the IBM data server driver configuration file.

The first member in the list of client affinities is called the primary member. You can temporarily override the primary member that is specified in the IBM data server driver configuration file by calling the SQLSetConnectAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY connection attribute.

To return to using the primary member that is specified in the IBM data server driver configuration file, you must specify a value of NULL for the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY connection attribute.

This message is returned when a value of NULL is specified for the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY connection attribute, but the primary member has not been overridden and the automatic client reroute feature is already using the primary member that is specified in the IBM data server driver configuration file.

User response: No action is needed.

CLI0237W The specified primary member is already set.

Explanation: You can use the client affinities feature to specify an ordered list of members to which the automatic client reroute feature will redirect application connections. You can configure the client affinities feature in the IBM data server driver configuration file.

The first member in the list of client affinities is called the primary member. You can temporarily override the primary member that is specified in the IBM data server driver configuration file by calling the SQLSetConnectAttr API with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute.

The new primary member that is specified in the

SQLSetConnectAttr API call with the SQL_ATTR_OVERRIDE_PRIMARY_AFFINITY attribute is same as the current primary member in the client affinities list.

User response: No action is needed.

CLI0238W Unretrieved OUT parameter data from the stored procedure call was lost.

Explanation: The application called the SQLFetch API before all OUT parameter data was retrieved with the SQLGetData API.

User response: If complete OUT parameter data is required, call the stored procedure again and complete the retrieval of OUT parameter with the SQLGetData API before fetching the result-set.

CLI0239E The batch CALL statement cannot be processed without the absolute parameter length value.

Explanation: You specified the SQL_DATA_AT_EXEC value for the length of the parameter in the SQLBindParameter API call. The SQL_DATA_AT_EXEC value is not supported by the batch CALL statement.

User response: Specify an absolute value for length of the parameter in the SQLBindParameter API call.

CLI0240E The batch CALL statement failed because the SQL_ATTR_ALLOW_INTERLEAVED_GETDATA attribute is specified.

Explanation: The SQL_ATTR_ALLOW_INTERLEAVED_GETDATA attribute is not supported by the batch CALL statement.

User response: Disable the SQL_ATTR_ALLOW_INTERLEAVED_GETDATA attribute.

CLI0241E The batch CALL statement failed because the SQL_ATTR_INTERLEAVED_PUTDATA or SQL_ATTR_INTERLEAVED_STREAM_PUTDATA statement attribute is specified.

Explanation: The SQL_ATTR_INTERLEAVED_PUTDATA and SQL_ATTR_INTERLEAVED_STREAM_PUTDATA attributes are not supported by batch CALL statements. The SQL_ATTR_INTERLEAVED_PUTDATA and SQL_ATTR_INTERLEAVED_STREAM_PUTDATA attributes are enabled through the SQLSetStmtAttr API.

User response: Disable the SQL_ATTR_INTERLEAVED_PUTDATA or

SQL_ATTR_INTERLEAVED_STREAM_PUTDATA attribute.

CLI0242E The batch CALL statement failed because the array input chain feature was specified.

Explanation: The array input chain feature is not supported by the batch CALL statement. The array input chain feature is enabled through the following methods:

- Enabling the ArrayInputChain keyword in the db2cli.ini file or the IBM data server driver configuration file.
- Setting the SQL_ATTR_CHAINING_BEGIN or the SQL_ATTR_CHAINING_END attribute in the SQLSetStmtAttr API.

User response: Disable the array input chain feature.

CLI0243E The batch CALL statement failed because the stored procedure contains the array data type argument.

Explanation: The stored procedure with array data type arguments is not supported by the batch CALL statement.

User response: Call the stored procedure without the use of the batch CALL statement.

CLI0244E Attempt to retrieve the latest connected member name has failed as client affinities feature has not been enabled.

Explanation: The client affinities element entries are not found in the IBM data server driver configuration file. The client affinities feature must be configured and the <alternateserverlist> subsection must be present in the IBM data server driver configuration file before you obtain the member name with the SQLGetConnectAttr API and the SQL_ATTR_GET_LATEST_MEMBER_NAME attribute.

User response: Create the required client affinities element entries in the IBM data server driver configuration file and rerun the application.

CLI0245E Attempt to set the SQL_ATTR_DATE_SEP attribute has failed because required SQL_ATTR_DATE_FMT attribute has not been set.

Explanation: The SQL_ATTR_DATE_FMT attribute must be set with the SQLSetConnect API or the SQLSetEnvAttr API before the SQL_ATTR_DATE_SEP attribute can be set.

User response: Set the SQL_ATTR_DATE_FMT attribute with the SQLSetConnect API or the SQLSetEnvAttr API before you set the

SQL_ATTR_DATE_SEP attribute.

CLI0246E **Attempt to set the
SQL_ATTR_TIME_SEP attribute has
failed because required
SQL_ATTR_TIME_FMT attribute has
not been set.**

Explanation: The SQL_ATTR_TIME_FMT attribute must be set with the SQLSetConnect API or the SQLSetEnvAttr API before the SQL_ATTR_TIME_SEP attribute can be set.

User response: Set the SQL_ATTR_TIME_FMT attribute with the SQLSetConnect API or the SQLSetEnvAttr API before you set the SQL_ATTR_TIME_SEP attribute.

Chapter 55. CLI0500 - CLI0999

CLI0600E Invalid connection handle or connection is closed.

Explanation: Connection had been closed prior to the operation.

User response: Ensure connection close is not called prior to the operation.

CLI0601E Invalid statement handle or statement is closed.

Explanation: Statement had been closed prior to the operation.

User response: Ensure statement close and statement connection close are not called prior to the operation.

CLI0602E Memory allocation error on server.

Explanation: Cannot allocate memory on server.

User response: Ask your database administrator to check the DB2 JDBC log file for details. Rerun your program.

CLI0603E CallableStatement get* method was called without registerOutParameter.**

Explanation: Get*** method was called on a parameter that was not registered using registerOutParameter.

User response: Add registerOutParameter call for the parameter.

CLI0604E CallableStatement get* method was called without calling execute.**

Explanation: The CallableStatement get*** method was called prior to calling CallableStatement execute.

User response: Make sure that CallableStatement execute is called prior to any CallableStatement get*** method.

CLI0605E CallableStatement get* method did not match the type used in registerOutParameter.**

Explanation: The CallableStatement get*** method was not the one that corresponds to the type used in registerOutParameter for this parameter.

User response: Change to the get*** method that matches to the type used in registerOutParameter for this parameter. (Refer to JDBC specification)

CLI0606E The returned value from a column is incompatible with the data type corresponding to the get* method.**

Explanation: The value in a CHAR/VARCHAR/LONGVARCHAR column is not a valid numeric value.

User response: Use the appropriate get methods other than the ones that return numeric values.

CLI0607E Invalid datetime format.

Explanation: The value in CHAR, VARCHAR, LONGVARCHAR, GRAPHIC, or VARGRAPHIC column is invalid date/time/ timestamp value.

User response: Use appropriate get*** method other than getDate/getTime/ getTimestamp.

CLI0608E Invalid conversion.

Explanation: The get*** method is invalid on this column type.

User response: Use the valid get*** method (see JDBC specification) to retrieve data from this column.

CLI0609E Numeric value out of range.

Explanation: The value of this column is too large/small for a short or int value, conversion will cause data loss.

User response: Use the get method that can accommodate the value.

CLI0610E Invalid column number.

Explanation: The column number is less than 1 or greater than the total number of columns in the ResultSet.

User response: Make sure the column number is no less than 1 and no greater than the total number of columns.

CLI0611E Invalid column name.

Explanation: The given column name cannot be found in the ResultSet.

User response: Ensure the column name is correct.

CLI0612E Invalid parameter number.

Explanation: The given parameter number is less than 1 or greater than the total number of parameters.

User response: Ensure the parameter number is no less than 1 and no greater than the total number of parameters.

CLI0613E Program type out of range.

Explanation: The object given in the PreparedStatement/CallableStatement setObject is not a valid object type.

User response: Make sure the object to be set is one of the object types allowed in setObject (see JDBC specification).

CLI0614E Error sending to the socket, server is not responding.

Explanation: Error happened while sending data to server, server may be down.

User response: Ensure the server is up, and rerun your program.

CLI0615E Error receiving from socket, server is not responding.

Explanation: Error happened while receiving from server, server may be down.

User response: Ensure JDBC server is up, and rerun your program.

CLI0616E Error opening socket.

Explanation: Cannot open the socket to server, server may be down.

User response: Ensure JDBC server is up, and rerun your program.

CLI0617E Error closing socket.

Explanation: Cannot close the socket to server. Server may be down.

User response: Ensure JDBC server is up, and rerun your program.

CLI0618E Userid and/or password invalid.

Explanation: The given userid/password is not valid.

User response: Make sure the userid/password is correct, and rerun your program.

CLI0619E Invalid UTF8 data format.

Explanation: When getUnicodeStream is called on columns that are not DBCS columns, the decoding from UTF8 format is done, but the data cannot be decoded properly.

User response: Use getString, getBytes, getAsciiStream, or getBinaryStream instead.

CLI0620E IOException, error reading from input stream.

Explanation: Error happened when reading data to be put from input stream.

User response: Make sure the file exists, and file length is correctly given.

CLI0621E Unsupported JDBC Server configuration.

Explanation: The target JDBC Server configuration is not supported. If you are running the Control Center, the target JDBC Server must be a standalone db2jd process (started via db2jstrt) and not a 2-tier native JDBC Server.

User response: Reconfigure the JDBC Server using db2jstrt on the port targeted by the Control Center.

CLI0622E Error accessing JDBC administration service extensions.

Explanation: The Control Center depends upon administration support services that run with the JDBC server. The Control Center was not able to locate or access these administration services.

User response: Ensure that the Control Center's administration services are installed with the JDBC server.

CLI0623E Code page conversion tables do not exist.

Explanation: Cannot find the codepage conversion tables.

User response: Ensure the conversion tables (from the local codepage to Unicode and from Unicode to the local codepage) are installed.

CLI0624E Code page conversion tables are not loaded.

Explanation: Cannot load the codepage conversion tables.

User response: Ensure your JDBC application has access to the code page tables and the tables are not corrupted.

CLI0625E You specified JDBC 1.22 behavior; cannot use JDBC 2.0 functions.

Explanation: The function you are trying to call is a new function defined in JDBC 2.0, but you specified you want JDBC 1.22 behavior.

User response: Do not set the JDBCVERSION keyword to "122" if you want to use JDBC 2.0 functions. It should remain unset or be set to "200".

CLI0626E *function-name* is not supported in this version of DB2 JDBC 2.0 driver.

Explanation: This feature is not supported by this version of the DB2 JDBC 2.0 driver.

User response: Do not use this feature.

CLI0627E The result set is not scrollable.

Explanation: The result set is not scrollable.

User response: Ensure you set the statement/result set attribute to SCROLLABLE.

CLI0628E No. *parameter-number* parameter marker in parameter set no. *set-number* is not set.

Explanation: The set<data-type> method has not been called for this input parameter.

User response: Call the set<data-type> method to specify the data type of and input value for this input parameter.

CLI0629E *function-name* is not supported for this column.

Explanation: This operation cannot be called for this column.

User response: Change to other method.

CLI0630E Unknown result set type/concurrency number.

Explanation: An invalid value for result set type or concurrency has been specified.

User response: Change the value to the proper value given in the specification.

CLI0631E No random access for mixed char/clob column.

Explanation: Random access is not supported for this mixed char column.

User response: Retrieve the data sequentially.

CLI0632E Invalid fetch size; it must be between 0 and maxRows, inclusively.

Explanation: The fetch size specified must be greater than or equal to 0, and less than or equal to maxRows.

User response: Modify the value.

CLI0633E Cannot call relative() when there is no current row.

Explanation: The relative method was called when the cursor is not at a valid row.

User response: First move the cursor to a valid row (using absolute, next, etc), then call relative.

CLI0634E Error allocating CLI environment handle.

Explanation: SQLAllocEnv failed during initialization.

User response: Ensure the DB2INSTANCE environment variable is set properly.

CLI0635E *function-name* is not supported in an applet.

Explanation: <function-name> is not supported in an applet.

User response: Avoid using <function-name> in an applet.

CLI0636E Property *property-name* is not specified for the Context object.

Explanation: Property *property-name* is not specified for the Context object in your Java application.

User response: Ensure that the property *property-name* is specified for the Context object in your Java application.

CLI0637E *object* cannot be found.

Explanation: <object> does not exist.

User response: Ensure that the <object> exists.

CLI0638E *object* already exists.

Explanation: <object> cannot be created as it already exists.

User response: Ensure that the operation is correct.

CLI0639E Empty string.

Explanation: An empty string is not allowed.

User response: Correct the string specified.

CLI0640E *object* cannot be listed.

Explanation: <object> cannot be listed.

User response: Ensure that the <object> can be listed.

CLI0641E There is a SELECT statement in the batch.

Explanation: A SELECT statement is not allowed in the batch.

User response: Remove the SELECT statement from the batch.

CLI0642E Invalid fetch direction.

Explanation: The fetch direction specified is not supported.

User response: Correct the fetch direction.

CLI0643E There is no statement in the batch.

Explanation: The batch does not have any statement.

User response: Add a statement to the batch.

CLI0644E Invalid row value to the absolute() call.

Explanation: The row value specified to absolute() is invalid.

User response: Correct the row value.

**CLI0645E Error registering driver *class-name*.
Message: *message*. SQLSTATE: *sqlstate*.
SQLCODE: *sqlcode*.**

Explanation: The DriverManager is unable to register the DB2 JDBC driver.

User response: Correct the problem as indicated by the returned message, SQLSTATE and SQLCODE, and run your program again.

CLI0646E Cannot find library *library-name*.

Explanation: *library-name* cannot be found in the library path. The library path is defined by the environment variable for the following operating systems:

AIX (Java 1.1)
LD_LIBRARY_PATH

AIX (Java 1.2 or later)
LIBPATH

HP-UX SHLIB_PATH

Linux LD_LIBRARY_PATH

Silicon Graphics IRIX
LD_LIBRARY_PATH

Solaris Operating Environment
LD_LIBRARY_PATH

Windows operating systems
PATH

User response: Ensure your application environment is configured correctly to use DB2. On UNIX platforms, ensure that the sqllib/db2profile script has been run to set your environment variables.

CLI0647E Error allocating DB2 environment handle. Return code = *return-code*.

Explanation: The DB2 CLI environment could not be established.

User response: Ensure your application environment is configured correctly to use DB2. On UNIX platforms, ensure that the sqllib/db2profile script has been run to set your environment variables. To interpret the return code, refer to the *CLI Guide and Reference* section on SQLAllocHandle().

CLI0648N The ResultSet is closed.

Explanation: The ResultSet was closed prior to the operation.

User response: Ensure ResultSet.close() was not called prior to the operation. If you are processing multiple ResultSets using getMoreResults(), please ensure you call getResultSet() prior to accessing the new ResultSet.

CLI0649N executeQuery is only allowed for statements that return a ResultSet.

Explanation: executeQuery is only allowed for statements that return a ResultSet. Such SQL statements include SELECT, VALUES, WITH, and CALL statements.

User response: Use executeUpdate() or execute().

CLI0650N A query is not allowed in an executeUpdate().

Explanation: You cannot issue a query with executeUpdate().

User response: Use executeQuery() or execute().

CLI0651N The stream contains more or less data than is specified.

Explanation: The number of bytes or characters in the stream is not equal to the given length.

User response: Specify the correct length.

CLI0652W The execution of the query will not be atomic because the database server does not support atomic operations for non-insert queries.

Explanation: You can use the SQL_ATTR_PARAMOPT_ATOMIC statement attribute to specify whether the processing of multiple parameter markers should be done altogether as one operation (atomic) or as multiple operations (non-atomic.) For example, specifying a value of SQL_ATTR_PARAMOPT_ATOMIC_YES causes the underlying processing of the query to be performed as an atomic operation.

This message is returned when the following conditions are all true for a given query:

- The SQL_ATTR_PARAMOPT_ATOMIC attribute is set to SQL_ATOMIC_YES for the query.
- The query is not an insert operation.
- The query is executed against a DB2 for z/OS database server, which does not support atomic processing of non-insert queries.

User response: You do not need to respond to this warning. There is no method for forcing a non-insert query to be executed atomically against a DB2 for z/OS database server.

Part 9. Net Search Extender Messages

Note that the SQL states returned from the search function are 38600 plus the CTE error number.

Chapter 56. CTE0000 - CTE0499

CTE0015W No dependent text-maintained table found requiring reset.

Explanation: The RESET PENDING command executes a set integrity statement for dependent tables used to manage full-text search. No dependent table was found that matches these criteria for the given table.

CTE0016W Failed to set integrity for *schema-name.table-name*.

Explanation: The RESET PENDING command executes a set integrity statement for dependent tables used to manage full-text search. The command failed to complete successfully.

For more details see the topic called "RESET PENDING" in the DB2 Information Center.

User response: Run a "Set Integrity for <schema>.<table> IMMEDIATE CHECKED" for the failed table.

CTE0017W Failed to set integrity for dependent tables.

Explanation: The RESET PENDING command executes a set integrity statement for dependent tables used to manage full-text search. The command failed to complete successfully.

For more details, see the topic called "RESET PENDING" in the DB2 Information Center.

User response: Run a "Set integrity for <dependent-table> immediate checked" for any dependent tables for the base table that are still in pending mode.

CTE0100E A DB2 operation failed. DB2 information: %2 %4.

Explanation: A DB2 error occurred that does not allow further processing.

User response: For more detailed information on this DB2 error, use the following command: db2 ? SQLxxx.

CTE0101E A search engine operation failed.
Reason code: %2, %3, %4, %5, %6.

Explanation: A Search Engine error occurred that does not allow further processing.

User response: For more detailed information, see the Search Engine reason code descriptions.

CTE0102E A general system function failed. Error: %2.

Explanation: A system error occurred that does not allow further processing.

User response: Additional information can be found on UNIX in the errno.h header file.

CTE0103E An internal error occurred. Location: %1, %2.

Explanation: An internal processing error that does not allow further processing. Try to start and stop the update and locking services, as well as DB2.

User response: If the error persists, start a trace and also check the db2diag.log.

CTE0104E Memory allocation error (search engine).

Explanation: The system has run out of memory.

User response: Increase the available memory size for the instance owner, or stop other processes running parallel.

CTE0105E Memory allocation error.

Explanation: The system has run out of memory.

User response: Increase the available memory size for the user, or stop other processes running parallel.

CTE0106E Table %1.%2 has no primary key.

Explanation: You tried to create an index on a table that does not have a primary key.

User response: Call the db2 alter table to ensure the existence of a primary key. Then try to create the index again.

CTE0107E Directory %1 does not exist.

Explanation: You specified a directory which does not exist.

User response: Create the directory, ensure accessibility to the instance owner. Then try to specify the directory again. Note that in a distributed DB2 environment, this directory has to exist on every physical node.

CTE0108E The internal size %4 of the key columns on object %1.%2 is larger than maximum allowed size of %3.

Explanation: The internal representation of the key columns exceeds the maximum size.

User response: Change the layout of the table before creating the index again. Use smaller key columns, which also benefit performance.

CTE0109E The number of key columns %3 on object %1.%2 is larger than the allowed maximum of %4.

Explanation: A maximum number of 14 key columns is supported.

User response: Change the layout of the table before creating the index again.

CTE0110E The primary key on object %1.%2 is larger than allowed.

Explanation: The primary key length exceeds the maximum size. The primary key length limit is based on the page size of the table space used by the table:

| Max Key Length | Page size |
|----------------|-----------|
| 1007 | 4K |
| 2031 | 8K |
| 4079 | 16K |
| 4094 | 32K |

Note that if the primary key consists of more than one column, the mentioned limits must be reduced by 2 bytes for each additional column.

User response: Change the layout of the table before creating the index again. Use smaller key columns, which also benefit performance.

CTE0111E The file %1 is not readable.

Explanation: The file specified cannot be read.

User response: Check the access rights for the file. Take into account that the Stored Procedure runs as a fenced user ID, which may also require rights to work on this file.

CTE0112E The file %1 cannot be opened.

Explanation: The file specified could not be opened.

User response: Verify that the file is correctly specified.

CTE0113E Error converting model file %1 to UTF-8 encoding.

Explanation: The specified CCSID or the default database CCSID does not match the model file CCSID.

User response: Ensure correct specification of the model file CCSID.

CTE0114E Unable to register document model %1 in file %2.

Explanation: The model file could not be used.

User response: Check that the model file syntax is correctly specified.

CTE0115E A locking problem occurred. Lock Manager information: %1 %2.

Explanation: An internal locking problem occurred.

User response: Check the current locks using the db2text control command. Using the same command, clean up the pending locks. If this does not help, stop and restart the locking and update services.

CTE0116E Operation conflicts with existing lock.

Explanation: You have tried to use a command that is currently not allowed when other commands are running on the index.

User response: Check the locks held on this index to see which commands are currently running. Wait until the other commands have finished. If the operation is no longer running but the lock is still active, clean up the locks for the index and try again.

CTE0117E All available lock space for databases is used. Change the configuration.

Explanation: You tried to work on more databases than are configured in your lock file.

User response: Change the number of databases you want to work in parallel with in your lock configuration db2extlm.cfg. Restart the update and locking services using the db2text stop and db2text start commands.

CTE0118E All available lock space for indexes on a databases is used. Change the configuration.

Explanation: You tried to work on more indexes for one database than are configured in your lock file.

User response: Change the number of indexes you want to work in parallel with in your lock configuration file db2extlm.cfg. Restart the update and locking services using the db2text stop and db2text start commands.

CTE0119E All available space for locks on an index is used.

Explanation: The operations you are running require more locks for one index than are configured in your lock configuration file.

User response: Change the number of locks you want to work in parallel with in your lock configuration db2extlm.cfg. Restart the update and locking services using the db2text stop and db2text start commands.

CTE0120E Update and locking services configuration file error.

Explanation: The configuration file db2extlm.cfg is in error.

User response: Check the db2extlm.cfg file and correct the error. Restart the update and locking services using the command db2text start.

CTE0121E The update and locking services configuration file cannot be opened.

Explanation: The file db2extlm.cfg could not be opened.

User response: Check if the file exists and that it can be accessed. If the file cannot be accessed, try to update your db2 instance using db2iupdt.

CTE0122E A syntax error was found in the update and locking services configuration file.

Explanation: A syntax error was found in the update and locking services configuration file.

User response: Check the update and locking services configuration file for errors.

CTE0126E The update and locking service input file %1 is corrupted.

Explanation: A required file for update and locking services is corrupted.

User response: Check if the file exists and if it can be accessed. If you can access the file, rename the file and restart the update and locking services. The file should be created again. However, this action removes all of the specified frequency updates for create index.

CTE0127E An update and locking service error has occurred. Reason code: %1.

Explanation: An internal error has occurred in the update and locking service area.

User response: Stop DB2 and Net Search Extender and then clean up your shared resources. Try to start both again. If this does not work, report the problem to your IBM representative.

CTE0129E NULL values are not allowed to be passed as parameters.

Explanation: DB2 has passed a NULL value to an internal user-defined function.

User response: First make sure the specified base table has a primary key. Change your select statement to avoid this problem. Switch on the trace function and pass the returned information on to IBM Services.

CTE0130E The specified search argument exceeds the maximum length. The current search argument length is %1 and the maximum supported length is %2.

Explanation: The length of the specified search argument is %1. The maximum length must not exceed %2.

User response: Reduce the length of your search argument to %2.

CTE0131E The user-defined function %1.%2 does not exist.

Explanation: The specified user-defined function does not exist in this database.

User response: Check the name specified for this user-defined function, or register the user-defined function in the database you are using.

CTE0132E The text index %1.%2 does not exist.

Explanation: The specified text index does not exist in this database.

User response: Check the name specified and the database you are using. Use the db2ext.text columns view to see the existing text indexes.

CTE0133E The text index %1.%2 already exists.

Explanation: The text index that you specified already exists in this database.

User response: Check the name specified and the database you are using. Use the db2ext.text columns view to see the existing text indexes.

CTE0135E The object %1.%2 does not exist.

Explanation: The specified object name does not exist in this database.

User response: Check the object name specified and the database you are using.

CTE0136E The column %1 does not exist in %2.%3.

Explanation: The specified column does not exist.

What to do

Check the column name that you specified. Check the table, view, or database you are using.

CTE0137E The table space %1 does not exist.

Explanation: The specified table space does not exist in this database.

User response: Check the name specified and the database you are using.

CTE0138E The table space %1 is not regular.

Explanation: The specified table space is not regular. The event table can only be created in a regular table space.

User response: Use this command again with a regular table space.

CTE0139E The environment variable %1 is not set.

Explanation: A required environment variable is not set.

User response: Check your environment, specify the required variable, and use the command again.

CTE0140E The database %1 is already enabled for text.

Explanation: The database you specified is already enabled for text.

User response: Check the name that you specified. Also check the DB2DBDFT variable that implies an implicit connection.

CTE0141E The database %1 is not enabled for text.

Explanation: The database you specified is not enabled for text.

User response: Check the database name you specified and the DB2DBDFT variable. If the database name is correct, use the command db2text enable database for text.

CTE0142E The command requires control authority on %1.%2 granted to user %3.

Explanation: You do not have the authority to use this command.

User response: Only the owner of this table can use this command or provide you with the required authorization.

CTE0143E The command requires database administration authority for user %1.

Explanation: You do not have the required authority to use this command.

User response: Only the owner of the database can use this command or provide you with the required authorization.

CTE0144E There is at least one text index active in database %1.

Explanation: You cannot disable your database until all text indexes are dropped.

User response: See the db2ext.text columns view for the existing indexes. Drop the existing indexes using the DROP INDEX command or specify the FORCE option with the DISABLE DATABASE command.

CTE0145E The CCSID %1 is not supported.

Explanation: The CCSID that you specified is not supported.

User response: Specify a valid CCSID.

CTE0146E The language %1 is not supported.

Explanation: The specified language is not supported.

User response: Specify a valid language.

CTE0147E The format %1 is not supported.

Explanation: The specified format is not supported.

User response: Specify a valid format.

CTE0148E The specified format %1 does not accept a model file.

Explanation: The format %1 does not support model files.

User response: Use a format that accepts a model file, or remove the model file from your command.

CTE0149E Too many terms (beginning with %1) are specified for the index update frequency.

Explanation: The syntax for the update frequency is not correct.

User response: Ensure that the DAY, HOUR, and MINUTE parameters are only specified once.

CTE0150E Unexpected end of command. Check the command syntax.

Explanation: The command syntax is not correct, or you are creating an index on a partitioned table and you omitted the ADMINISTRATION TABLES IN clause.

User response: Check the command syntax. Verify that you specified the required parameters. If you are creating an index on a partitioned table, specify the ADMINISTRATION TABLES IN clause.

CTE0151E Token %1 is unexpected. Check the command syntax.

Explanation: The syntax of the command is not correct.

User response: Check the command syntax and verify that the token you are using is allowed in the specific command.

CTE0152E Token %1 is too long.

Explanation: The token is too long.

User response: Check the command syntax and verify that the token is reduced to the maximum size allowed.

CTE0153E Token %1 occurs twice in the update frequency.

Explanation: You specified an incorrect syntax for the update frequency.

User response: Ensure that the DAY, HOUR, and MINUTE parameters are only specified once.

CTE0154E The value %1 for %2 is out of range. The valid range is %3 - %4.

Explanation: You specified an incorrect value. The value should be in the allowed range.

User response: Update your command. Change the value to match those in the allowed range.

CTE0155E The search string is empty.

Explanation: You specified an empty search string.

User response: Check that the search string includes valid alphanumeric characters.

CTE0157E Syntax error near %1.

Explanation: You specified an incorrect search syntax.

User response: Check the syntax near %1. Correct and try again.

CTE0158E The freetext search string is missing.

Explanation: Specify a freetext string.

User response: Check that the search string after "is about" includes valid alphanumeric characters.

CTE0159E Search string exceeds the allowed length of %1.

Explanation: The search string is too long.

User response: Reduce the size of the search string and try again.

CTE0160E No section name has been specified in the search string.

Explanation: You need to specify a valid section name.

User response: Add a valid section name and try again.

CTE0162E The escape command could not be processed.

Explanation: Your search string includes too many special characters that can be used as masking characters.

User response: Reduce the number of special characters in your search term, or avoid the escape command. The following special characters can be used: ! * + , _ . : ; { } ~ | ? [] ` = \

CTE0163E No thesaurus name specified in thesaurus clause.

Explanation: A thesaurus search is requested without a thesaurus name.

User response: Specify a thesaurus name in your search argument.

CTE0164E Syntax error in thesaurus relation %1.

Explanation: The specified syntax for the thesaurus relation is not correct.

User response: Update the thesaurus relation according to the syntax specification.

CTE0166E Freetext must be the last statement in search query.

Explanation: It is not allowed to have further operators after the "is about" token.

User response: Rewrite the query string. The last operator must be "is about".

CTE0167E **Syntax error in free text query %1.**

Explanation: The syntax for the free text string is not correct.

User response: Update the free text string according to the syntax specification.

CTE0168E **A left parenthesis in a section statement is missing.**

Explanation: The syntax for the section statement is not correct.

User response: Update the section statement according to the syntax specification.

CTE0169E **A comma or right parenthesis is missing in a section statement.**

Explanation: The syntax for the section statement is not correct.

User response: Update the section statement according to the syntax specification.

CTE0170E **A closing double quote is missing.**

Explanation: The specified syntax for the search term is not correct.

User response: Update the search term according to the syntax specification.

CTE0171E **An open double quote for a section name is missing.**

Explanation: The syntax for the section statement is not correct.

User response: Update the section statement according to the syntax specification.

CTE0172E **The closing double quote for the section name is missing.**

Explanation: The syntax for the section statement is not correct.

User response: Update the section statement according to the syntax specification.

CTE0173E **One escape character must be defined in an escape clause.**

Explanation: There can be no more than one character in an escape clause.

User response: Remove the additional characters in the escape clause.

CTE0174E **A blank character is not allowed as an escape character.**

Explanation: It is not allowed to have a blank character in an escape clause.

User response: Change the escape clause to a clause with a valid character.

CTE0175E **An escape clause is defined but no mask character is found in the search phrase.**

Explanation: An escape clause is specified without using a mask character.

User response: Remove the escape clause.

CTE0176E **The succeeding character of an escape character in the phrase is neither the same character nor a mask character.**

Explanation: The character after the escape character must be either a masking character or the escape character itself.

User response: Change the search string to correctly use the escape character.

CTE0177E **The number value %1 is invalid.**

Explanation: The specified number in the search argument is not valid.

User response: Check the documentation about the valid range. Update the value in the search argument.

CTE0178E **Mask characters in fuzzy phrase must be preceded by an escape character.**

Explanation: Masking together with fuzzy search is not allowed.

User response: Update the search string with an escape character.

CTE0179E **Thesaurus name %1 exceeds allowed length of %2.**

Explanation: Primary keys longer than 60 bytes are not supported.

User response: Change the layout of the table before creating the index again.

CTE0180E **Thesaurus %1 can not be found.**

Explanation: The thesaurus specified cannot be found.

User response: Check that the thesaurus files are located in the thesaurus directory or fully qualified.

CTE0181E Library %1 cannot be loaded.

Explanation: A library cannot be found.

User response: Check that the library is located in the library path and available. Start and stop DB2 to ensure that the current settings are used.

CTE0182E Function %1 cannot be loaded from library %2.

Explanation: A library entry point cannot be loaded.

User response: The library accessed seems to be invalid. Check that the library is specified only once.

CTE0183E Error occurred using shared system resources.

Explanation: A request to share system resources like shared memory or semaphores cannot be fulfilled.

User response: Check the current system status and configuration. On UNIX, use the `ipcs` command to check the resources. Stop all applications, such as DB2 and Net Search Extender. If further resources are listed, clean them up using `ipcrm`.

CTE0184N No DB2TEXT START command was issued.

Explanation: A Net Search Extender (NSE) command was called which requires the NSE instance services to be started. You might get this messages if the service is not started.

In a multiple partitioned database environment:

- NSE instance services must be started on all database partitions.
- For more information about starting NSE instance services in a multiple partitioned database environment, see the topic called "Starting Net Search Extender (NSE) instance services in a multiple partitioned database environment".

User response: Use the DB2TEXT START command to start NSE instance services.

CTE0185E The Net Search Extender (NSE) instance services are already active.

Explanation: This message can be returned when the `db2text start` command is issued when the update and locking services are already running.

User response: No further action is required.

CTE0186E Update and locking service error occurred, check the db2diag log file for details.

Explanation: An update and locking service error occurred.

What to do

Check the `db2diag` log file for further information, or clean up your shared resources. See also CTE0183E.

CTE0187E Update and locking services are still active, use FORCE option to stop the services.

Explanation: The `db2text stop` command has not stopped the locking services, there are still processes running.

User response: Check with `db2text control` which processes are running and wait for those to finish. If you need to stop them, use the `FORCE` option.

CTE0188E There is a temporary problem using update and locking services. Please try again.

Explanation: The `db2text stop` command has not stopped the locking services. Programs are still running or an inconsistent situation is found.

User response: Check with `db2text control` which processes are running and wait for those to finish. To stop them, use the `FORCE` option.

CTE0189E The executable program %1 cannot be found.

Explanation: The program file cannot be located or accessed.

User response: Check if the program file is located in the `bin` or `adm` directory of the DB2 server. The installation is corrupt if the file cannot be found.

CTE0190E The executable program %1 cannot be started.

Explanation: The program cannot be started.

User response: Check if the program is located in the `bin` or `adm` directory of the DB2 server and that the appropriate libraries are installed. For further information, call the program manually on the server.

CTE0191E The drop index operation is incomplete. Check the db2diag log files for details.

Explanation: The drop index operation is incomplete, possibly caused by the `FORCE` option.

User response: Using the `FORCE` option drops everything regardless of any errors. Check the index directory for pending files and remove these manually.

CTE0192E Errors occurred in an update index operation. Check event table %1. %2 and db2diag log files for details.

Explanation: During the index update process, any document errors are written to the event table.

User response: Check the event table for more information about the document errors. Clean up the event log after the problems have been fixed.

CTE0194E The type %1 of column %2 is not supported.

Explanation: You used a column that is not in the list of the supported ones.

User response: Check CREATE INDEX for a list of valid columns for keys and indexing. Make the appropriate changes to the command and try again.

CTE0195E %1 is not an absolute path.

Explanation: An absolute path on the server is required.

User response: Check the path and write an absolute path in the command.

CTE0198E No corresponding text index.

Explanation: There is no text index on the column.

User response: Check if the text index still exists.

CTE0199E There is no text index corresponding to column *column-number* of table *table-name*.

Explanation: A text search query on a table column failed because it has no valid and active text search index.

User response: Check the column you are searching on, or create a text index on the column. If the table column has an invalid text search index, drop the text index and create a new one.

For example, a text search index gets invalidated if the underlying base table is dropped and recreated.

CTE0200E At least one command option must be specified.

Explanation: The ALTER INDEX command changes the characteristics of an index, such as the update and storage options. None of the characteristics to be changed was specified.

User response: Specify at least one command option. Refer to the command syntax for all possible options.

CTE0201E There is a conflict with an existing text index on the same column.

Explanation: A text index defined on the same column was created with different parameters from this create index command.

User response: Correct the parameter values in the create index command. Make sure that following parameters have the same value for the existing index and the index to be created: ccsid, language, format, document model, index configuration, column function, and attributes.

CTE0202E The object %1.%2 must be a view when key columns are specified.

Explanation: The specified object is not a view. The KEY COLUMNS FOR INDEX ON VIEW clause is only allowed when indexing a column of a view.

User response: Remove the KEY COLUMNS FOR INDEX ON VIEW(SQL-columnname-list) clause.

CTE0203E The text index %1.%2 was not created with the CACHE TABLE option. This is required for command execution.

Explanation: This command can only be executed if the specified index was created with the CACHE TABLE option.

User response: Create an index with CACHE TABLE option. Refer to the documentation for the command syntax.

CTE0204E An attribute name is missing. Add "AS <attribute name>" to the attribute expression.

Explanation: Whenever a column expression is used in the attribute expression, an attribute name must be supplied. For example: (C1+C2 AS myname).

User response: Add "AS <attribute name>" to the attribute expression.

CTE0205E CACHE TABLE expressions are not valid.

Explanation: The column list in the cache table expression is not valid.

User response: Correct the cache table column list in the create index command. Make sure the columns exist in the specified table. If a function is applied on a column, verify that it is used correctly.

CTE0206E ATTRIBUTE expressions are not valid.

Explanation: The column list in the attribute expression is not valid.

User response: Correct the attribute column list in the create index command. Make sure the columns exist in the specified table. If a function is applied on a column, verify that it is used correctly.

CTE0207E KEY COLUMNS FOR INDEX ON VIEW not specified for index on view %1.%2.

Explanation: If indexes on views are created, the KEY COLUMNS FOR INDEX ON VIEW(SQL-columnname-list) clause must be specified. The list of column names specifies the columns that UNIQUELY identify a row in the view.

User response: Include the KEY COLUMNS FOR INDEX ON VIEW(SQL-columnname-list) clause in the create index command.

CTE0208E INITIAL SEARCH RESULT ORDER columns are not valid.

Explanation: The column list in the INITIAL SEARCH RESULT ORDER(SQL-order-by list) expression is not valid.

User response: Correct the order by column list in the create index command. Check if the syntax is correct and the columns exist in the specified table. If a function is applied on a column, verify that it is used correctly.

CTE0209E The type %1 of attribute column %2 is not supported, type DOUBLE is required.

Explanation: For attribute columns, the only supported data type is DOUBLE.

User response: Make sure the attribute columns of the table with the text column to be indexed are of type DOUBLE. It may be possible to use cast operators in attribute column expressions. Refer to the SQL Reference for data types which can be cast to double.

CTE0210E The value %1 for index configuration parameter %2 is not valid. A valid value is %3.

Explanation: The specified value for the configuration parameter is incorrect. For valid values of the parameters refer to the command syntax.

User response: Correct the index configuration parameter value in the create index command.

CTE0211E %1 is not a valid index configuration parameter.

Explanation: The index configuration option is not known.

User response: Check the create index command syntax. Valid index configuration options are TreatNumbersAsWords and IndexStopWords. These have to be comma separated: index configuration(treatnumberaswords 1, indexstopwords 1).

CTE0212E Internal index configuration file %1 could not be saved.

Explanation: The internal configuration file for the index could not be saved.

User response: Make sure the instance owner has write permissions to the directory the file should be saved in. If a file with the same name already exists, make sure that it is writable for the instance owner.

CTE0213E Internal index configuration file template %1 could not be loaded.

Explanation: The internal index configuration file template could not be read.

User response: Make sure the file exists in the correct location and is readable.

CTE0214E Internal error when setting new entry [%1],%2=%3 for index configuration file.

Explanation: Internal error while writing an internal configuration file for the index.

User response: If the file exists, check if it is readable and writable for the instance owner. Check that there is enough space on the device where the file is located.

CTE0215E Index creation on alias %1.%2 is not supported. Use base table %3.%4 instead.

Explanation: The index cannot be created on the alias.

User response: Type in the create index command with the base table.

CTE0217E The schedule service is already active.

Explanation: The service is already active, you do not need to start it.

User response: No action required.

CTE0218E **The function *function-name* failed with error code *error-code*.**

Explanation: A Windows operating system function failed with the specified error code which does not allow further processing.

User response: Use the specified Windows system error code to get detailed error information.

CTE0219E **The service named *service-name* could not be opened. Error code *error-code*.**

Explanation: The specified service cannot be found on the Windows operating system.

This message can be returned when the database manager attempted to start the Net Search Extender (NSE) services in response to the DB2TEXT START command, but could not start the services on one or more of the database partitions.

For example, if Net Search Extender (NSE) is not installed on one or more computers where the database partitions are located, then the database manager will not be able to successfully start Net Search Extender (NSE) instance services.

User response: Check if the specified service is installed on the Windows system. Use the specified Windows system error code to get detailed error information.

CTE0220E **The DB2 instance profile path could not be found.**

Explanation: Internal DB2 function to obtain the DB2 instance profile path failed.

User response: Create a DB2 instance without specifying the instance profile path information and retry the command.

CTE0221E **UpdateFrequency *%1* is incorrectly specified.**

Explanation: The syntax for the update frequency statement is not correct.

User response: Correct the update frequency statement according to the syntax specification.

CTE0222E **The schedule service input file *%1* is corrupted.**

Explanation: The scheduler file containing index update information is corrupted.

User response: Use your system editor and try to correct the problem. Maybe an entry has been truncated, or the ending line character has been deleted. If this does not restore the file content, try the following: Call command db2text stop to stop the scheduler. Delete the scheduler service file. Call

command db2text start to start the scheduler. Use command db2text alter index ... to recreate the update frequency entries for all concerned indexes.

CTE0223E **File *%1* could not be closed.**

Explanation: The file specified cannot be closed.

User response: Verify that the file is correctly specified.

CTE0224E **File *%1* could not be copied to *%2*.**

Explanation: The first file cannot be copied to the second file.

User response: Verify that the files are correctly specified. Check if the second file already exists and is read only. Also check if there is enough free space on the system.

CTE0225E **File *%1* could not be removed.**

Explanation: The file specified cannot be removed from the system.

User response: Verify that the file is specified correctly and check the file access rights.

CTE0227E **A write operation on file *%1* failed.**

Explanation: The file specified is not writable.

User response: Verify that the file is correctly specified and check the file access rights. Also check if there is enough free space on the system.

CTE0228E **The user has insufficient access rights at the operating system level.**

Explanation: The command requires administrator rights at the operating system level.

User response: Ensure that you have operating system administrator rights. Check if you are a member of the administrator group.

CTE0231E ***%1* is not defined in same nodegroup (*%4*) as the table space of *%2.%3*.**

Explanation: The table space of the administration tables is required to be distributed over different nodes in exactly the same way as the table containing the text column to be indexed. To enforce this, it is checked whether the specified tablespace is defined in the same nodegroup.

User response: Specify a table space that is defined in the same nodegroup as the table containing the text column to be indexed.

CTE0232E The specified or default table space %1 is not single-noded. This is necessary for an index on a view, or when the CACHE TABLE option is specified.

Explanation: An index on a view or with the CACHE TABLE option enabled is only supported for tables on a single node.

User response: Put the table in a single-noded table space if the default table space caused this error. Alternatively, specify another single-noded table space, if you specified a multi-noded table space.

CTE0233E There is a conflicting administration command running. Please retry this command later.

Explanation: Another administration command is still running or terminated abnormally without releasing the command lock.

User response: Check with CONTROL LIST which locks are still active. If there is an active lock but no command running, clear the lock manually using the CONTROL CLEAR command. Be aware that someone else may be running the administration command holding the lock.

CTE0234E There is a conflicting administration command running on a text index. Please retry this command later, or specify the FORCE option of a DISABLE DATABASE command.

Explanation: Another administration command is still running or terminated abnormally without releasing the command lock.

User response: Check with CONTROL LIST which locks are still active. If there is an active lock but no command running, clear the lock manually using the CONTROL CLEAR command. Be aware that someone else may be running the administration command holding the lock. For a DISABLE DATABASE command you may specify the FORCE option which stops all other commands on that database.

CTE0235E No valid license found for Net Search Extender.

Explanation: There was no valid license found for Net Search Extender.

User response: Check whether the license was correctly installed with db2lic. Make sure existing instances are updated after the product install.

CTE0236E Only Node0 is supported on MPP instances.

Explanation: Text Indexes can only be created on MPP instances, if the table with the text column to be indexed resides on Node0.

User response: Check the node group of the tablespace in which the table is defined.

CTE0237E Internal error: log table %1.%2 contains an invalid operation %3.

Explanation: The log table keeps track of operations executed on the table containing the indexed text column. This table might be corrupted, as it contains an entry not written by Net Search Extender.

User response: Check the log table and delete the corrupted entry.

CTE0238E Internal error: table %1.%2 contains an incorrect syntax expression in column %3.

Explanation: There is an error in the expression list in the specified text column.

User response: Check the delimiter Begin and End pairs.

CTE0239E Internal error: total length of index properties %1 exceeds maximum %2.

Explanation: The maximum size of the index properties (1016 bytes) is exceeded. The properties contain the instance, index, and work directory as well as other information.

User response: Make sure these path names are not too long.

CTE0240E Internal error: setting environment variable %1 failed.

Explanation: Setting the specified environment variable failed. There may be a problem with the environment setup.

User response: Check your OS specific guidelines.

CTE0242E Value *value* for parameter *parameter-name* is invalid.

Explanation: The search stored procedure or the table valued function DB2EXT.TEXTSEARCH was called with invalid parameters.

User response: Correct the parameter values of the search stored procedure or SQL table-valued function. For valid parameters refer to the documentation.

CTE0243E The cache for text index %1.%2 has not been activated.

Explanation: A Net Search Extender operation requires an activated cache. The cache is currently not activated. These are the possible reasons: The cache has never been activated after the last DB2TEXT START command. The cache has been explicitly deactivated with the DB2TEXT DEACTIVATE CACHE command.

User response: Perform a DB2TEXT ACTIVATE CACHE command for the index and rerun the Net Search Extender operation.

CTE0244E Internal error: call to %1 returns rc=%2, SQLCODE=%3.

Explanation: An internal processing error occurred when calling an internal function.

User response: If the error persists, start a trace and check the db2diag.log. Report the error.

CTE0245E The requested cache size exceeds the available cache size. Increase the maximum cache size to a value > %1 or decrease the pctfree value.

Explanation: The cache size necessary to load all data exceeds the MAXIMUM CACHE SIZE value for an index. This can be detected during activation of the cache (the DB2TEXT ACTIVATE command), or by an index update operation while the cache is activated.

User response: If the error was reported in a DB2EXT ACTIVATE command, recalculate the maximum cache size using the DB2EXT.MAXIMUM_CACHE_SIZE function and alter the MAXIMUM CACHE SIZE setting for the index. Eventually decrease the PCTFREE value. If the maximum number of documents is exceeded during incremental update, rebuild the cache with the commands db2 deactivate cache and db2text activate cache recreate.

CTE0246E File %1 is empty.

Explanation: A DB2TEXT CREATE INDEX command failed because the document model file specified in the command is empty.

User response: Specify a valid document model file in the command.

CTE0247E A Net Search Extender stored procedure could not be created.

Explanation: A DB2TEXT ENABLE DATABASE command failed to create the internal stored procedure DB2EXT.CTESRVSP.

User response: Check the additional DB2 error message associated with a CREATE PROCEDURE statement for details. If the error cannot be corrected by

removing an existing stored procedure with an identical name, start a trace and report the error.

CTE0248E The generated search string is too long. Reduce the complexity of search query.

Explanation: A Net Search Extender query is too long or too complex to be processed by the base search engine. The complexity is affected by thesaurus expansions, FUZZY FORM OF expressions, and masking characters.

User response: Reduce complexity or length of the query.

CTE0249E The executable program named *program-name* terminated abnormally.

Explanation: While executing a Net Search Extender (NSE) command, the database manager ran the named executable program. That program terminated abnormally.

User response:

1. Verify, that the executable was not terminated explicitly by user interaction.
2. If the program was not terminated by a user, rerun the command that failed.
3. If the problem persists:
 - Turn on the DB2 trace utility to collect diagnostic information.
 - Rerun the command that failed
 - If the command fails again, contact DB2 support with the collected diagnostic information.

CTE0250E The return type %1 of column type transformation function %2.%3 is not supported.

Explanation: In a DB2TEXT CREATE INDEX command a column type transformation was specified that returns an unsupported datatype. Supported datatypes are: CHARACTER, VARCHAR, LONG VARCHAR (deprecated), CLOB, GRAPHIC, VARGRAPHIC, LONG VARGRAPHIC (deprecated), DBCLOB, BLOB, and XML.

User response: Choose a different column type transformation function.

CTE0251E Internal error: the column type %1 is not supported.

Explanation: A column type is used that is not in the list of supported types.

User response: Check create index for a list of valid columns for Keys and Indexing. Make the appropriate changes to the command and try again. If the error persists, start a trace and also check the db2diag.log. Report the error to IBM Services.

CTE0252E The parameter %1 is missing.

Explanation: Internal error - when executing a Net Search Extender command, an administration executable program was called with a missing parameter "%1".

User response: Try to change Net Search Extender parameter commands to avoid the problem. If the error persists, switch on the trace function and report the error to IBM Services.

CTE0253E The document listed in the log view was not found.

Explanation: The contents of a text document that is listed in the log view has changed and could not be accessed.

User response: Check that the document exists and the read/access permissions of the text documents to be included in the index.

CTE0254E The cache for index %1 is already activated.

Explanation: The index has already been activated with the ACTIVATE CACHE command.

User response: Check the specified index name and the database that you are using.

CTE0255E A column name for a cache result column expression is missing. Add "AS <cache column name>" to the expression.

Explanation: A cache result column expression must be named. For example: 'C1+C2 AS myresult'.

User response: Add "AS <cache column name>" to the expression.

CTE0256E The query necessary to select data for indexing failed. Reduce the complexity of the attribute, cache table, or the initial search result order expressions.

Explanation: Net Search Extender creates a query from the expressions in your command to select data for indexing from the database. The query failed because it was too complex.

User response: Reduce the complexity of attribute, cache table, or initial search result order expressions.

CTE0257E : Error creating shared memory.

Explanation: The shared memory resource could not be created due to a previous error or permission problem.

User response: Check db2diag.log for further

information, or clean up your shared resources. See also error CTE0183E.

CTE0258E Shared memory version error.

Explanation: The shared memory resource could not be accessed because it is corrupted or there is a version conflict.

User response: Check db2diag.log for further information. Disable and re-enable the database and then try again.

CTE0259E Cannot insert entry in global shared memory. Entry already exists.

Explanation: An entry to be inserted in global shared memory already exists because of a previous error.

User response: Check db2diag.log for further information. Restart the update and locking services using the commands db2text stop and db2text start.

CTE0260E Cannot access entry in global shared memory. Entry not found.

Explanation: An entry to be removed from global shared memory does not exist because of a previous error.

User response: Check db2diag.log for further information. Try to restart the update and locking services using the commands db2text stop and db2text start.

CTE0261E There is at least one cache activated for a text index in this instance. Deactivate the cache for any activated index using the DEACTIVATE CACHE command, or use the FORCE option to stop.

Explanation: The db2text stop command can only be used if you run a DEACTIVATE CACHE command for all text indexes that have been activated with the ACTIVATE CACHE command.

User response: Deactivate the cache for any activated index using the DEACTIVATE CACHE command or use the FORCE option to stop.

CTE0262E The value for parameter %1 is too long.

Explanation: The value exceeds the maximum allowable size.

User response: Check the maximum size.

CTE0263E The text index %1.%2 was created with the RECREATE INDEX ON UPDATE option. In this context, the UPDATE MINIMUM or COMMITCOUNT FOR UPDATE may not be specified.

Explanation: Update minimum and commitcount for update are only effective if the index is updated incrementally.

User response: If you want to recreate the index each time an update is performed, remove the UPDATE MINIMUM and COMMITCOUNT FOR UPDATE settings. If you want to use UPDATE MINIMUM and COMMITCOUNT FOR UPDATE, do not specify RECREATE INDEX ON UPDATE.

CTE0264E Errors occurred in an activate index operation. Check event view %1.%2 and the db2diag.log for details.

Explanation: During the index activate process, errors are written to the event table and the db2diag.log file.

User response: Check the event table for more information about the document errors. Clean up the event log after the problems have been fixed.

CTE0265E The table space of a user table or administration table space (%1) is not only defined on node 0.

Explanation: If text indexes are created on MPP instances, the table space of the user table must only reside on Node0.

User response: Use a table where the table space resides on Node0.

CTE0266E ValueFrom %1 must be smaller than ValueTo %2.

Explanation: The values specified in the attribute search are not valid. If the search syntax is 'BETWEEN ValueFrom AND ValueTo', the lower boundary (ValueFrom) must be smaller than upper boundary(ValueTo).

User response: Change the boundaries in the 'BETWEEN ValueFrom AND ValueTo' clause.

CTE0267E The Net Search Extender database objects in the database %1 are in an inconsistent state.

Explanation: At least one Net Search Extender object is missing or corrupted. Either the database has not been migrated after installation of a new Net Search Extender product version, or a database user has changed or dropped Net Search Extender internal object(s). In this case, all text indexes are lost and the database has to be disabled for text.

User response: For a database migration to the current version please follow the migration description found in the Net Search Extender documentation. Alternatively, issue a DB2TEXT DISABLE DATABASE command using the FORCE option. You can then enable the database for text again by using the

DB2TEXT ENABLE DATABASE command.

CTE0270E Logtable %1.%2 could not be modified after incremental update. Entries are to be processed during the next UPDATE.

Explanation: When starting an incremental index update, a timestamp is created. This serves as a threshold for change records to be processed. Changes occurring concurrently to the incremental update are then processed later, during the next update. In certain situations, there can be changes in transactions that are uncommitted at the time the update starts, but are committed while the index update is being performed. This may potentially lead to inconsistencies.

To avoid such an inconsistent situation, the change records prior to the threshold timestamp are not deleted from the logtable, although they have been partially processed. On the next incremental update the changes will be re-applied to the index.

User response: On the next index update the changes are re-applied to the index. In case of delete operations, this can lead to the following error: CTE0101E: ItlEnReasonCode_Docmap_ docid_not_found.

Note that this error can be ignored, as the document was already deleted. If CTE0270E errors frequently occur, consider dropping and re-creating the index with a modified timestamp threshold for incremental index update. For example: db2text "CREATE INDEX ... INDEX CONFIGURATION(UPDATEDELAY 30)"

This means that processing during an incremental update run only changes records older than 30 seconds and avoids interference with concurrent change transactions of less than 30 seconds.

CTE0271E Cache not usable, DEACTIVATE and ACTIVATE RECREATE required.

Explanation: The cache is in an inconsistent state because the maximum cache size has been reached.

User response: Check that the maximum cache size is still sufficient. Then call the following db2text commands: DEACTIVATE CACHE and ACTIVATE CACHE RECREATE.

CTE0272E Insufficient cache size. Increase the PCTFREE value or use DEACTIVATE and ACTIVATE [RECREATE] to recreate the cache.

Explanation: All reserved memory for the cache has been used.

User response: Rebuild the cache by using the following sequence of db2text commands: DEACTIVATE CACHE, ALTER INDEX MAXIMUM CACHE SIZE, and ACTIVATE CACHE RECREATE.

CTE0273E **The cache for index %1, %2 is already activated.**

Explanation: The index has already been activated with the ACTIVATE CACHE command.

User response: Check the specified index name and the database that you are using.

CTE0274E **The target database system %1 for the connection is not supported.**

Explanation: You tried to execute a DB2TEXT command with a connection to a database system that is not supported by Net Search Extender.

CTE0275E **The type and version information for server %2 could not be found.**

Explanation: The type and version information for the server could not be found in the DB2 catalog view 'SERVERS'.

User response: Make sure that the DB2 federated environment is set up correctly.

CTE0277E **A cache memory segment could not be attached.**

Explanation: The system cannot allocate enough memory to load a large cache segment, or the cache segment cannot be opened because it has been previously deleted.

User response: Check your system settings and increase the amount of paging space and free memory. For large cache sizes you may need to prepare your system. Refer to the Net Search Extender documentation. Use the DEACTIVATE and ACTIVATE [RECREATE] commands to recreate the cache. If the problem persists, check db2diag.log for additional information.

CTE0278E **On an AIX 32-bit system, change the MAXDATA setting before activating a large cache.**

Explanation: When you use the search stored procedure on an AIX 32-bit system, you may need to change the MAXDATA setting for the db2fmp executable.

User response: Refer to the Net Search Extender documentation for details about changing the MAXDATA setting.

CTE0279E **The size of the cached data has reached a system limit.**

Explanation: By decreasing the PCTFREE value, you can increase the maximum data size during cache

activation. This enables the system to reserve less freespace in the cache.

User response: Use a lower PCTFREE value or reduce your amount of text data to be cached. Use the DEACTIVATE and ACTIVATE [RECREATE] commands to recreate the cache.

CTE0280E **There is not enough disk space to write persistent cache files.**

Explanation: The system can not write a large enough file for persistent cache in the cache directory.

User response: Change the persistent cache directory to an empty file system by using the ALTER INDEX command. Alternatively, reduce the cache size by decreasing the PCTFREE or MAXIMUM CACHE SIZE values or by using a temporary cache.

CTE0281E **Deletion of persistent cache file %1 has failed.**

Explanation: The file does not exist or cannot be accessed.

User response: Check if this file still exists and delete it manually.

CTE0282E **The number of documents in the cache has reached a system limit.**

Explanation: By decreasing the PCTFREE value, you can increase the maximum number of document entries to be cached during cache activation. This enables the system to reserve less freespace in the cache.

User response: Use a lower PCTFREE value or reduce the amount of document entries in the cache. Use the DEACTIVATE and ACTIVATE [RECREATE] commands to recreate the cache.

CTE0283E **A cache memory segment could not be created.**

Explanation: The system cannot allocate enough memory for loading a large cache segment into memory. By decreasing the PCTFREE value, you achieve a smaller cache segment size.

User response: Check your system settings and increase the amount of paging space and free memory. You can also decrease the cache size by using a lower PCTREE value. For large cache sizes, you may need to prepare your system. Refer to the Net Search Extender documentation. Use the DEACTIVATE and ACTIVATE [RECREATE] commands to recreate the cache. If the problem persists, check db2diag.log for additional information.

CTE0284E The text index is located on node %1, but the search function was called on node %2.

Explanation: The search stored procedure or table valued function DB2EXT.TEXTSEARCH was not called on the node where the index is located. The search function will not automatically be distributed to the correct node.

User response: Set the DB2NODE environment variable to the node where the index is connected before connecting to the database.

CTE0285E Search function is not allowed for a text index which is distributed to multiple nodes.

Explanation: The table valued function DB2EXT.TEXTSEARCH must not be called with indexes that are distributed to multiple nodes, since it will not be automatically distributed to the correct nodes, but executed on the coordinator node.

User response: Use the CONTAINS, SCORE or NUMBEROFMATCHES function in a multiple node environment.

CTE0286E No row found in %1."IBMSNAP_REGISTER" for source table %2.%3 and capture change table %4.%5.

Explanation: No valid entry was found in the IBMSNAP_REGISTER table for the replication capture table characteristics specified in the DB2TEXT CREATE INDEX command. A valid entry must contain the specified source table for the index incolumns SOURCE_OWNER and SOURCE_NAME, with SOURCE_VIEW_QUAL=0 and the specified replication capture table in columns PHYS_CHANGE_OWNER and PHYS_CHANGE_TABLE.

The specified source table was not registered as a replication source for the replication capture table.

User response: Register the source table correctly for DB2 Replication, or specify a correct replication capture table for the source table.

CTE0287E Invalid value %1 for %2 in "%3"."IBMSNAP_REGISTER" for source table %4.%5 and capture change table %6.%7.

Explanation: A replication setting found in the IBMSNAP_REGISTER table is not allowed. Possible causes: 1.The column CHG_UPD_TO_DEL_INS does not contain the value 'Y'. 2.The column CCD_CONDENSED contains the value 'Y'.

User response: When registering the source table for DB2 Replication, ensure that update operations are

transformed into pairs of delete and insert operations. In addition, ensure that no condensed replication capture tables are used.

CTE0288E Source table %1.%2 and capture change table %3.%4 are on different servers (%5 and %6).

Explanation: The specified source table and replication capture table must reside on the same server.

CTE0289E The wrapper %1 is not supported.

Explanation: The wrapper is not supported. Refer to the Net Search Extender documentation for a list of supported wrappers.

CTE0290E The alias %1.%2 is not allowed in the replication clause.

Explanation: You are not allowed to specify an alias for a nickname in a replication clause.

User response: Specify the nickname instead of the alias, or create a new nickname for the remote table.

CTE0291E The specified format is not allowed for a column of type XML.

Explanation: For columns of type XML only the 'XML' format is allowed.

User response: Specify format 'XML' or none.

CTE0292E Windows Exception %1 was caught, address=%2, flags=%3.

Explanation: A windows exception occurred. Exception name, address and flags are provided.

CTE0293E Windows exception %1 was caught.

Explanation: A Windows exception occurred.

CTE0294E Search argument processing problem.

Explanation: A search argument processing error occurred due to an incorrect environment setup.

User response: Check that the locale charmap value matches the DB2 code page and is available on the system.

CTE0295E Invalid CCSID %1 specified for non-binary text column.

Explanation: For non-binary text column data types, DB2 always stores the data in the database CCSID. Only the database CCSID is valid for non-binary text columns.

User response: Omit the CCSID-clause or specify a valid CCSID.

CTE0296E The library %1 cannot be found on %2. Please check the Net Search Extender installation.

Explanation: You tried to execute a Net Search Extender command by using the DB2 Control Center. Net Search Extender is not installed correctly on the target system.

User response: Check if Net Search Extender is installed correctly on the target system.

CTE0297E The database %1 is associated with DB2 Text Search

Explanation: You are not allowed to associate a database with more than one text search component.

User response: If you want to continue using DB2 Text Search with the database, then no action is required. If you would rather use DB2 Net Search Extender, disable the database from DB2 Text Search and try this command again.

CTE0298E The database %1 is already associated with DB2 Text Search. This command cannot be executed.

Explanation: You are not allowed to associate a database with more than one text search component.

User response: If you do not want to use DB2 Net Search Extender, disable the database using the DB2TEXT DISABLE DATABASE command. If you would rather use DB2 Net Search Extender, try this command again after disabling the database with DB2 Text Search.

CTE0300N User *user-id* does not have authority to create text index on the table *schema-name.table-name*.

Explanation: "CREATE INDEX .. FOR TEXT" requires one of the following:

- DBADM authority
- CONTROL privilege on the table
- INDEX privilege on the table with either IMPLICIT_SCHEMA privilege on the database or CREATEIN privilege on the index schema.

For more information about DB2 Net Search Extender create index, see the topic called "Net Search Extender CREATE INDEX command" in the DB2 Information Center.

User response: Ensure that the user has authority to perform the create index operation.

CTE0301E The instance owner does not have the authority to perform the Net Search Extender command.

Explanation: The user attempted to run a text index command without the instance owner id holding the appropriate authority

Different DB2 Text Search index commands require varying levels of database authority. For a list of required authorities for DB2 Net Search Extender commands, see the topic called "DB2 Net Search Extender command authority requirements" in the DB2 Information Center.

User response: Ensure that the instance owner has DBADM with DATAACCESS authority in order for the user to perform Net Search Extender operations.

CTE0302E User *user-id* does not have authority to update the DB2 Net Search Extender index named *index-schema-name.index-name*.

Explanation: Update text index requires one of the following:

- DATAACCESS authority
- CONTROL privilege on the table.

For more information about DB2 Net Search Extender update, see the topic called "Net Search Extender UPDATE INDEX command" in the DB2 Information Center.

User response: Ensure that the user has authorization to perform the update index operation.

CTE00303N Table column *column-name.table-name* has no valid and active text index.

Explanation: A text search query on a table column failed because it has no valid and active text index.

User response: Ensure that the table column has a valid and active text index before performing a text search query.

If the table column has an invalid text index, drop the text index and create a new one.

For example, a text index gets invalidated if the underlying base table is dropped and recreated.

CTE0304N Index *schema-name.index-name* is not valid and cannot be used.

Explanation: The text index is invalid and cannot be updated or used.

User response: Drop and recreate the invalid text index.

CTE0305E Failed to retrieve the text-maintained dependent tables for *schema-name.table-name*.

Explanation: The command executes a query to retrieve information from the text index administration tables. Querying the text index information failed.

User response: Ensure that the text index administration tables are accessible. For more information see the db2diag.log.

CTE0306N The specification of a COMMITCOUNT value is not allowed for a text index created with the REPLICATION CAPTURE TABLE option.

Explanation: If the REPLICATION CAPTURE TABLE option is specified the COMMITCOUNT value cannot be specified for the index.

User response: Retry the command without the COMMITCOUNT option for the index since REPLICATION CAPTURE TABLE option is being used.

CTE0307N The table specified in the replication clause must not be a view.

Explanation: The viewname has been specified in the replication clause instead of the table name.

User response: Retry command with the view name replaced with the table name for the replication clause.

CTE0308N A text index on a view cannot be created using the REPLICATION CAPTURE TABLE option.

Explanation: The viewname has been specified in the replication clause instead of the table name.

User response: Retry command with the view name replaced with the table name for the replication clause.

CTE0309W Search result truncated by mask limit.

Explanation: The result set for the search was truncated due to the mask resolution limit.

User response: To get the complete set of results for the search query, increase the mask resolution limit and re execute the search query.

CTE0310N Net Search Extender index update failed due to insufficient disk space for index *schema-name.index-name* in database *database-name*. Required disk space: *required-disk-space* KB. Available disk space: *available-disk-space* KB.

Explanation: The available space for index update is less than the estimated disk space required.

User response: Check the available free space on the file system that holds the index and work directories for the index.

Increase the available space on the file system and try to update the index again.

CTE0311E The CREATE INDEX command is missing the "ADMINISTRATION TABLES IN" clause for a partitioned table.

Explanation: The command syntax is incorrect.

User response: Check the command syntax. Verify that you have specified the ADMINISTRATION TABLE IN clause in the CREATE INDEX command for partitioned tables.

CTE0312E A common secondary group does not exist for a fenced user and an instance owner.

Explanation: A common secondary group is required for a fenced user and an instance owner.

User response: Add a common secondary group for fenced user and instance owner.

CTE0360E Index update is not allowed due to insufficient disk space.

Explanation: Do not update the index as the available space for index update is less than the estimated disk space required.

User response: Check the available free space for the file system that contains the index and work directory for the index. Increase the available space on the file system and try updating index again.

CTE0451E The specified document format %1 is not supported by the highlighting UDF.

Explanation: The document format %1 does not support highlighting.

User response: Use a document format that is supported by the highlighting UDF.

CTE0452E Syntax error near option %1 in the highlighting UDF.

Explanation: You specified an incorrect syntax near the specified option.

User response: Check the syntax near option %1. Correct and try again.

CTE0453E **The return size of the highlighting UDF is too small.**

Explanation: The requested parts of the highlighted document does not fit into the return parameter of the highlighting UDF.

User response: Decrease the window number, the window size and/or the number of sections from which hits should be displayed. This will reduce the document parts returned to the user.

CTE0454E **Error converting the parameters of the highlighting UDF from codepage %1 to codepage UTF8.**

Explanation: The parameters of the highlighting UDF in the specified CCSID (which may be the default database CCSID), cannot be converted to UTF8.

User response: Ensure correct specification of the CCSID.

CTE0455E **The database codepage %1 is not supported in the highlighting UDF.**

Explanation: The database has a codepage which is not supported by the highlighting UDF.

CTE0456E **The highlighting UDF only supports documents in codepage UTF8.**

Explanation: Only documents in codepage UTF8 support the highlighting UDF.

CTE0457E **The value %1 for parameter %2 is not valid in the highlighting UDF.**

Explanation: A value for a highlighting parameter is not valid.

User response: Check the parameter value and ensure that the value is allowed in the data range.

CTE0458E **Usage: db2exthl <new size in kilo bytes>.**

Explanation: The parameter for the db2exthl utility is not correct.

User response: Provide a value between 1 and 1048576.

Chapter 57. CTE0500 - CTE0999

CTE0841E Missing command option %1.

Explanation: A required command option was not specified.

User response: Check the specified parameters and add the missing parameter.

CTE0842E No value is specified for the command option %1.

Explanation: A required value for a command option was not specified.

User response: Check the specified parameters and add the missing option.

CTE0843E No numeric value is specified for the command option %1.

Explanation: A string instead of a number has been specified.

User response: Check the specified parameters and change the string to the correct number.

CTE0844E The definition file path %1 is too long.

Explanation: The specified path is too long and could not be processed.

User response: Use a shorter path and try again.

CTE0845E No definition file is specified.

Explanation: The definition file needs to be specified.

User response: Add a valid definition file and try the call again.

CTE0846E The definition file name %1 is too long.

Explanation: The specified definition file name is too long.

User response: Reduce the length of the definition file name to the size allowed.

CTE0847E The definition file %1 does not exist.

Explanation: The specified definition file could not be found.

User response: Check that the definition file is in the correct path and can be accessed by the current user.

CTE0849E The dictionary file %1 could not be locked.

Explanation: The process was not able to lock the dictionary file. Either you do not have write access, or another process has opened the file for writing.

User response: Check the running processes to ensure that no process is locking the dictionary file and check your access rights.

CTE0850E Output file %1 already exists.

Explanation: The specified output file could not be overwritten.

User response: Check that you are able to create the thesaurus in the specified directory.

CTE0851E The integrity of the dictionary file %1 is lost.

Explanation: The thesaurus dictionary files are corrupted.

User response: Clean up the directory and compile your definition file again.

CTE0852E Dictionary file %1 version error.

Explanation: Your dictionary file was generated with an older version of the thesaurus compiler.

User response: Compile your definition file again with the current version of the thesaurus compiler.

CTE0853E The existing dictionary %1 cannot be overwritten.

Explanation: An existing dictionary cannot be overwritten.

User response: Check your write access right on the dictionary file, its directory location and subdirectory location.

CTE0855E A thesaurus term is incorrectly specified.

Explanation: There is a syntax error in your definition file.

User response: Check your Net Search Extender documentation for information on creating a thesaurus definition file and thesaurus support.

CTE0856E The definition file %1 is empty.

Explanation: An empty definition file is not allowed.

User response: Check your Net Search Extender

documentation for information on creating a thesaurus definition file and thesaurus support.

CTE0857E **No block starting line found in file**
file-name **at line** *line-number*.

Explanation: There is a syntax error in your thesaurus definition file.

User response: A block has to start with 'WORDS'. Check your Net Search Extender documentation for information on thesaurus concepts.

CTE0858E **An invalid relationship is specified in**
file %1 at line %2.

Explanation: There is a syntax error in your definition error.

User response: You have to examine your 'associated-term-definition'. Check your Net Search Extender documentation for information on creating a thesaurus definition file.

CTE0859E **The relationship number is out of range**
in file %1 at line %2.

Explanation: The user-defined relations are all based on the associative type. They are identified by unique numbers between 1 and 128.

User response: Verify your relationship numbers.

CTE0860E **No terms are defined in file %1 at line**
%2.

Explanation: Required terms are not specified.

User response: Check your Net Search Extender documentation for information on creating a thesaurus definition file.

CTE0861E **The thesaurus term in file %1 at line %2**
is too long.

Explanation: The length of the thesaurus term is restricted to 64 bytes.

User response: Alter the size of your thesaurus term and try again.

CTE0862E **Strength is incorrectly specified in file**
%1 at line %2.

Explanation: There is a syntax error in your definition file.

User response: Check your Net Search Extender documentation for information on creating a thesaurus definition file and thesaurus support.

CTE0863E **Strength is out of range in file %1 at**
line %2.

Explanation: The strength value should be specified between 1 and 100.

User response: Change the strength value so that it is a numerical value from 1 to 100.

CTE0864E **Internal error: Thesaurus compiler failed**
with reason code %1.

Explanation: An internal processing error occurred that does not allow further processing. Try to start and stop the update and locking services, as well as DB2.

User response: If the error persists, start a trace and also check the db2diag.log.

CTE0865E **The directory %1 could not be created.**

Explanation: The specified directory could not be created.

User response: Check if the directory already exists and the permissions on the directory.

CTE0866E **The directory %1 could not be removed.**

Explanation: The directory could not be removed.

User response: Check that you have write permissions on the specified directory.

Chapter 58. CTE5000 - CTE5499

CTE5251E The checknseindex command failed because the following invalid parameter or value was specified: *invalid-input*.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

This message is returned when the checknseindex command is invoked with an invalid parameter. There are multiple ways this error could occur, including the following examples:

- An erroneous parameter was specified.
- A parameter was specified twice.
- A parameter was specified in upper case. (The parameters for the checknseindex command must be specified in lower case.)
- A mandatory parameter was not specified.

User response:

1. Review the checknseindex command syntax by issuing the following command:
`checknseindex -h`
2. Issue the checknseindex command again, specifying valid parameters.

CTE5252E The checknseindex command failed because a required parameter was not specified. Missing parameter: *missing-parameter*.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

This message is returned when the checknseindex command is invoked and one or more required parameters were not specified.

User response:

1. Review the checknseindex command syntax by issuing the following command:

`checknseindex -h`

2. Issue the checknseindex command again, specifying valid parameters.

CTE5253E The checknseindex command failed because no value was specified for the following parameter: *parameter*.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

This message is returned when the checknseindex command is invoked and no value was specified for a parameter that requires a value.

User response:

1. Review the checknseindex command syntax by issuing the following command:
`checknseindex -h`
2. Issue the checknseindex command again, specifying values for all parameters that require a value to be specified.

CTE5254E The NSE index validation utility failed to validate the specified index because the utility could not access the specified path. Specified path: *specified-path*.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

This message is returned when the checknseindex utility cannot access the NSE index directory path that was specified with the -p parameter. There are multiple ways this error could occur, including the following examples:

- The specified path does not exist.
- The checknseindex utility does not have required access permissions for the specified path. (The

checknseindex utility runs with the same privileges as the user who calls the checknseindex command.)

User response:

1. List information about current NSE indexes, including index directories, by using the db2ext.textindexes view.
2. Verify that the specified NSE index directory path exists.
3. Verify that the user who is running the checknseindex command has read permission for the specified NSE index directory path.
4. Issue the checknseindex command again, specifying a valid NSE index directory path.

CTE5255E **The NSE index validation utility failed because the utility could not find the specified NSE index in the specified path. Specified NSE index: *specified-index*. Specified path: *specified-path*.**

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

This message is returned when the checknseindex utility cannot find the NSE index directory path that was specified with the -p parameter.

User response:

1. List information about current NSE indexes, including index directories, by using the db2ext.textindexes view.
2. Issue the checknseindex command again, specifying a valid NSE index and index directory path.

CTE5256E **The NSE index validation utility failed to validate the specified index because the utility was unable to validate the NSE mapping index. Reason code: *reason-code*. Diagnostic data: *data*.**

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

The checknseindex utility validates a specified NSE

index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the checknseindex utility encounters an internal, unhandled error while attempting to investigate the NSE mapping index.

There are multiple ways this error could occur. The reason code indicates more specifically why the checknseindex utility failed. Possible reason codes include the following:

1

The checknseindex utility was unable to allocate memory for an internal structure.

2

The checknseindex utility was unable to open an internal file that is related to the NSE mapping index.

The diagnostic details in the *data* runtime token are useful for IBM support personnel only. The runtime token *data* might sometimes be empty.

User response: Run the checknseindex command again.

If this error continues, contact IBM support for assistance.

CTE5257E **The NSE index validation utility found the NSE mapping index to be invalid. Reason code: *reason-code*. Diagnostic data: *data*.**

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

The checknseindex utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the checknseindex utility determines that the NSE mapping index is not valid.

There are multiple ways this error could occur. The reason code indicates more specifically why the NSE mapping index is not valid:

1

NSE mapping index-related files are damaged or contain invalid contents.

2

Data control records are invalid.

3

Index control records are invalid.

4

The size of the reverse mapping file does not match the expected file size.

The diagnostic details in the *data* runtime token are useful for IBM support personnel only. The runtime token *data* might sometimes be empty.

User response: If this error continues, either recreate the index or contact IBM support for assistance.

CTE5258I The NSE index validation utility found the NSE mapping index to be valid.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

The checknseindex utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the checknseindex utility determines that the NSE mapping index is valid.

User response: No user response is required.

CTE5259E The NSE index validation utility failed to validate the specified index because the utility was unable to validate the NSE internal index.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

The checknseindex utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the checknseindex utility encounters an internal, unhandled error while attempting to investigate the NSE internal index.

User response: Run the checknseindex command again.

If this error continues, contact IBM support for assistance.

CTE5260E The NSE index validation utility found the NSE internal index to be invalid.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

The checknseindex utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the checknseindex utility determines that the NSE internal index is not valid.

User response: If this error continues, either recreate the index or contact IBM support for assistance.

CTE5261I The NSE index validation utility found the NSE internal index to be valid.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

The checknseindex utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the checknseindex utility determines that the NSE internal index is valid.

User response: No user response is required.

CTE5262E The NSE index validation utility failed to validate the specified index because the utility was unable to confirm whether the NSE mapping index and the NSE internal index are consistent.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of

significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the `checknseindex` utility.

The `checknseindex` utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the `checknseindex` utility encounters an internal, unhandled error while attempting to compare the NSE mapping index with the NSE internal index.

User response: Run the `checknseindex` command again.

If this error continues, contact IBM support for assistance.

CTE5263E The NSE index validation utility found the NSE mapping index and the NSE internal index to be inconsistent. Reason code: *reason-code*. Diagnostic data: *data*.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the `checknseindex` utility.

The `checknseindex` utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the `checknseindex` utility determines that the NSE mapping index and the NSE internal index are not consistent with each other.

There are multiple ways this error could occur. The reason code indicates more specifically why the NSE mapping index and the NSE internal index are not consistent with each other:

1

The total number of documents in the NSE mapping index does not match the total number of documents in the NSE internal index.

2

The highest document number in the NSE mapping index does not match the highest document number in the NSE internal index. This message is returned with reason code 2 when a mismatch in the highest document number is found after the initial index update.

3

The highest document number in the NSE mapping index does not match the highest document number in the NSE internal index. This message is returned with reason code 3 when a mismatch in the highest document number is found after an incremental index update.

The diagnostic details in the *data* runtime token are useful for IBM support personnel only. The runtime token *data* might sometimes be empty.

User response: If this error continues, either recreate the index or contact IBM support for assistance.

CTE5264I The NSE index validation utility found the NSE mapping index and the NSE internal index to be consistent.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the `checknseindex` utility.

The `checknseindex` utility validates a specified NSE index by investigating several internal structures of the NSE index, including the NSE mapping index and the NSE internal index. This message is returned when the `checknseindex` utility determines that the NSE mapping index and the NSE internal index are consistent with each other.

User response: No user response is required.

CTE5265I The NSE index validation utility found the specified index, *specified-index*, to be valid.

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the `checknseindex` utility.

This message is returned when the `checknseindex` utility determines that the NSE index is valid.

User response: No user response is required.

CTE5266E **The NSE index validation utility found the specified index, *specified-index*, to be invalid.**

Explanation: To enable Net Search Extender (NSE) to search more efficiently through documents that are stored in a DB2 database, you can create an NSE text index on the documents. An NSE index consists of significant terms that are extracted from those documents.

To troubleshoot problems with NSE indexes, you can validate an NSE index by using the checknseindex utility.

This message is returned when the checknseindex utility determines that the NSE index is not valid.

User response: If this error continues, either recreate the index or contact IBM support for assistance.

CTE5267E **The NSE index validation utility failed to validate the index due to an invalid command syntax.**

Explanation: You can detect corruption or other abnormalities in Net Search Extender text indexes by using the checknseindex utility.

This message is returned when the checknseindex command is run with invalid parameters. For example, this message can be returned when two mutually exclusive parameter are specified.

User response: Issue the checknseindex command again, specifying only valid parameters.

CTE5268I **Multiple NSE indexes exist on the same column *column_name*.**

Explanation: More than one NSE index has been created on the specified column for the same schema, database and table.

The checknseindex utility can only validate one NSE index at a time.

User response:

1. Select an NSE index.
 2. Run the checknseindex utility again for that index.
-

CTE5269W **The internal key list is full for index ID *index-id*, which might cause index reorganization to fail.**

Explanation: Net Search Extender indexes maintain an internal key list structure, index reorganization will fail if the key list structure get full.

This message is returned when the key list for the given index has reached its threshold. The index is still functional and will continue to index documents. However, this condition can cause index reorganization to fail.

User response: Create another NSE index for this column to avoid index reorganization failure.

CTE5270W **NSE internal index key list is nearly full for index ID *index-id*. Create another NSE index for the column to avoid possible index reorganization failure.**

Explanation: The NSE internal index key list has almost reached its threshold. The index is still functional and will continue to index documents. However, this condition can eventually cause index reorganization to fail.

User response: Create another NSE index for this column to avoid index reorganization failure.

Part 10. DB2 Messages

This section contains the messages generated by the command line processor. The command line processor returns DB2 and SQL messages. The messages are listed in numeric sequence.

Chapter 59. DB20000 - DB20499

DB20000I The *command* command completed successfully.

Explanation: No errors were encountered during the execution of this command.

User response: No action required.

Chapter 60. DB21000 - DB21499

DB21001E The option *option-letter* specified after the 'db2' command or in the DB2OPTIONS variable is incorrect.

Explanation: The option specified is not supported. The supported options are:

Option Description

-a Display SQLCA
-c Auto-commit
-e Display SQLCODE/SQLSTATE
-f Read from input file
-l Log commands in history file
-n Remove new line character
-o Display output
-p Display interactive prompt
-r Save output report to file
-s Stop execution on cmd error
-t Set stmt termination character
-v Echo current command
-w Display FETCH/SELECT warnings
-x Suppress printing of column headings
-z Save all output to file

The command cannot be processed.

User response: Resubmit the command with a valid option.

DB21002E The parameter for option *option-letter* specified after the 'db2' command or in the DB2OPTIONS variable is missing or incorrect.

Explanation: The following lists the options with parameters:

| Option | Description |
|--------------|--|
| ----- | ----- |
| -ec | Display SQLCODE |
| -es | Display SQLSTATE |
| -f<filename> | Read from input file <filename> |
| -l<filename> | Log commands in history file <filename> |
| -r<filename> | Save output report to file <filename> |
| -td<x> | Set termination char to 'x' |
| -z<filename> | Save all output to file <filename> |

User response: Resubmit the command with the valid option and parameter.

DB21003E The value *value* in *environment-variable* is not valid.

Explanation: The value for DB2BQTRY must be between 0 and 4294967295. The value for DB2BQTIME, DB2RQTIME or DB2IQTIME must be between 1 and 4294967295.

User response: Set the environment variable with the correct value and resubmit the command.

DB21004E You cannot specify both an input file and a command when invoking the Command Line Processor.

Explanation: You cannot specify both the -f option and a command line command when invoking the command line processor.

User response: Correct the error and issue the command again.

DB21005E An error occurred while accessing the file *filename*.

Explanation: The following could have caused the error:

- file permissions do not allow file access
- the file does not exist

User response: Correct the error and try again.

DB21006E The input command is too long. The maximum length is *length*.

Explanation: The input command cannot exceed the length specified.

User response: Correct the error and resubmit the command.

DB21007E End of file reached while reading the command.

Explanation: The last command was not executed because end of file was reached. Terminate the last command with a ';' (or your defined termination character) if the -t option was used. Remove the '\n' from the last line in the command if the +t option was used.

User response: Correct the error and resubmit the command.

DB21008E Commands can only be entered from the command line processor interactive mode or file input mode.

Explanation: The user attempted to enter a command line processor command from the DOS prompt.

User response: Use command line processor interactive mode or file input mode.

DB21009E **This command must be launched from a command window running with full administrative privileges.**

Explanation: You can not run this command from a command window running with a reduced set of privileges. The DB2 installation provides you with the "Command Window - Administrator" shortcut that has the appropriate privileges to run this command.

User response: Launch the "Command Window - Administrator" shortcut and re-run the command.

DB21010I **Help given for *help-command-phrase*.**

Explanation: This message only appears in the history file specified with the -I option.

User response: No action required.

DB21011I **In a partitioned database server environment, only the table spaces on the current node are listed.**

Explanation: Only the table spaces on the current node are visible to the LIST TABLESPACES command.

User response: To list the table spaces on another node, you must issue the LIST TABLESPACES command on that node.

DB21015E **The Command Line Processor backend process request queue or input queue was not created within the timeout period.**

Explanation: Either the values for the DB2BQTRY and DB2BQTIME environment variables need to be increased or the command line processor back-end program "db2bp" cannot be started. The "db2bp" program must reside in the correct database manager install path and users must have execute permission on the file.

On Linux and UNIX platforms, ensure that the file system has enough file blocks and inodes.

User response: Correct the error and resubmit the command.

DB21016E **The Command Line Processor encountered a system error while sending the command to the backend process.**

Explanation: One of the following could have happened:

- The backend process was abnormally terminated.
- A system error occurred when reading from or writing to the back-end process queues.
- A system error occurred when reading from the front-end process output queue.

User response: Resubmit the command. If the error occurs again, get help from your system administrator.

DB21017E **The Command Line Processor encountered a system error with the front-end process output queue. Reason code = *reason-code*.**

Explanation: A system error occurred while creating or reading from the front-end process output queue.

If the reason code is -2499, the command line processor output queue conflicts with an existing queue.

User response: Resubmit the command. If the error occurs again, record the message number and reason code, and get help from your system administrator.

DB21018E **A system error occurred. The command line processor could not continue processing.**

Explanation: One of the following caused the system error:

- Too much data is being output to the screen. Pipe the output to a file that you can view once processing is complete.
- Command line processor did not successfully install its interrupt signal handler.
- Command line processor did not successfully open the back-end process queues.
- Command line processor did not successfully start up the back-end process.
- The back-end process was abnormally terminated.
- The front-end process did not successfully allocate or free memory.
- The front-end and/or back-end process did not successfully dynamically load a library.
- Command line processor received one of the following program termination signals:
 - SIGILL
 - SIGTRAP
 - SIGEMT
 - SIGBUS
 - SIGSEGV
 - SIGSYS

User response: Retry the command. If the problem persists, record the DB2 message number. If the trace was active, save the trace information and contact your technical support with the following information:

- Problem description

- DB2 message number
- SQLCA if possible
- Trace file if possible

DB21019E An error occurred while accessing the directory *directory*.

Explanation: The following could have caused the error:

- directory permissions do not allow access
- the directory does not exist

User response: Correct the error and resubmit the command.

DB21020E Unable to create the default message file *file*.

Explanation: To properly process this command, CLP needs a file where the messages issued during processing are saved before they are displayed on the console. It attempted to create such a file in a directory normally used for that purpose (such as /tmp on UNIX platforms), but the attempt failed because the directory does not exist.

User response: Correct the error and resubmit the command.

DB21021E No Admin Server instance is defined. The command failed.

Explanation: You issued a command that needs to use an Administration Server instance, but no such instance is defined.

User response: Define an Administration Server instance and resubmit the command.

DB21022E Unable to switch to Administration Server instance *instance-name*.

Explanation: You issued a command that needs to use an Administration Server instance. Command line processor tried to switch to Administration Server instance *instance-name* but failed. The possible reasons may be:

- The Administration Server instance is not set correctly.
- Command line processor was already attached to a DB2 instance.
- Command line processor was connected to a database.

User response: Check that a valid Administration Server instance is set up before using this command. Also, you may need to issue the DETACH, or CONNECT RESET, or TERMINATE command before trying your request again.

DB21023E The command is not valid when executed from Administration Server.

Explanation: You issued one of the following commands: GET DBM CONFIGURATION, RESET DBM CONFIGURATION, or UPDATE DBM CONFIGURATION, but these commands cannot be executed from Administration Server.

User response: Issue one of the following commands that are valid on Administration Server: GET ADMIN CONFIGURATION, RESET ADMIN CONFIGURATION, or UPDATE ADMIN CONFIGURATION.

DB21024I This command is asynchronous and may not be effective immediately.

Explanation: This message is displayed after executing either a FORCE command with the ASYNC clause or a REORG TABLE command with the INPLACE clause.

User response: No action required.

DB21025I One or more of the parameters submitted for immediate modification were not changed dynamically. Client changes will not be effective until the next time the application is started or the TERMINATE command has been issued. Server changes will not be effective until the next DB2START command.

Explanation: Some changes to the database manager configuration could not be applied immediately. For these parameters, the changes become effective only after DB2 is started. This typically occurs after a DB2START on a server and after the application restarts on a client.

User response: To see which parameter changes took effect dynamically and which ones did not, retrieve the database manager configuration parameters and display details using the following command:

```
DB2 GET DBM CFG SHOW DETAIL
```

Changes to database manager configuration parameters can only take effect dynamically if you are attached to the instance. Not all configuration parameters support dynamic update. To see which parameters can be changed dynamically, refer to the Administration Guide.

If several parameters were submitted in a group, submit the parameters individually. In those cases where a configuration parameter could not change dynamically, do one or more of the following:

- for user applications: stop and start the application
- for CLP: TERMINATE and reconnect
- for a server: issue DB2STOP and DB2START

DB21026I One or more of the parameters submitted for immediate modification were not changed dynamically. For these configuration parameters, all applications must disconnect from this database before the changes become effective.

Explanation: The database configuration command has been successfully processed. However, not all changes were processed right away. After the applications disconnect from the database, the first connection to the database will cause the changes to become effective.

User response: To see which parameter changes took effect dynamically and which ones did not, retrieve the database configuration parameters and display the parameter details using the following command:

```
DB2 GET DB CFG FOR
    <database-alias> SHOW DETAIL
```

Changes to database configuration parameters can only take effect dynamically if you are connected to the database. Not all configuration parameters support dynamic update. To see which parameters can be changed dynamically, refer to the Administration Guide.

If several parameters were submitted in a group, submit the parameters individually. In those cases where a configuration parameter could not change dynamically, do one or more of the following:

- Ensure that all applications disconnect from the database and issue a db2 connect to command
- Rebind your packages after the new configuration parameters take effect so that the new values will be used during the bind.
- Use the FLUSH PACKAGE CACHE command to invalidate dynamic statements in the SQL cache.

DB21027E Isolation level may not be changed while connected to a database.

Explanation: The user attempted to change the isolation level while connected to a database. The command is ignored.

User response: If a change in isolation level is necessary, disconnect from the current database, then set the isolation level and re-connect to the database.

DB21028E The cursor *cursor-name* has not been declared.

Explanation: The specified cursor needs to be declared before issuing an OPEN, FETCH, or CLOSE SQL statement.

User response: Declare the cursor and resubmit the command.

DB21029E The cursor *cursor-name* has already been declared and opened.

Explanation: The user attempted to declare a cursor that is open.

User response: Close the cursor and resubmit the open command.

DB21030E The cursor *cursor-name* has not been opened.

Explanation: The specified cursor needs to be opened.

User response: Open the cursor and resubmit the command.

DB21031E The SQL statement using the cursor *cursor-name* (internal-cursor) returned:

Explanation: This message shows the internal cursor name for the user defined cursor. Some SQL error messages may show the internal cursor name. This message is displayed before the SQL message.

User response: Correct the SQL error and resubmit the command.

DB21032E You have already declared the maximum number of cursors.

Explanation: Command line processor supports 100 cursors declared with the WITH HOLD attribute and 100 cursors declared without the WITH HOLD attribute.

User response: Re-declare one of the existing cursors and resubmit the command.

DB21033E The command is not valid for this DB2 database server.

Explanation: The following commands are supported only when accessing a DB2 for Linux, UNIX, or Windows database server:

- LIST TABLES
- LIST PACKAGES
- REORGCHK
- INGEST

User response: Do not issue the command for this database.

DB21034E The command was processed as an SQL statement because it was not a valid Command Line Processor command. During SQL processing it returned:

Explanation: You can execute database utilities, SQL statements, and online help using the DB2 command line processor (CLP). You can use the CLP in interactive mode, in command mode (where each command must

be prefixed by "db2"), or in batch mode.

This message is displayed in the following scenarios:

1. Text is passed to the CLP and it is not a valid CLP command.
2. An error occurs when the CLP passes the text to the database manager to process as an SQL statement. When this scenario happens, this message is returned, followed by the error message that was returned by the database manager.

This error can be caused by the following kinds of problems with the specified command or SQL statement:

- Text is passed to the CLP in command mode or in batch mode and there are special characters, such as quotes, in the text that are not identified with an escape character.
- There is a syntax error.

User response: Respond to this error in one of the following ways:

- If there are special characters, such as quotes, in the command, use an escape character, such as the backslash character, to cause the operating system shell to ignore those special characters.
- Correct the syntax error by performing the following steps:
 1. Review the full text of the error returned by the database manager to determine the cause of that error.
 2. Modify the command to correct the cause of the problem.
 3. Resubmit the command.

DB21035E **The maximum number of items in the *list-name* list has been exceeded. The maximum number is *number*.**

Explanation: The number of items in the list cannot exceed the maximum specified. This error may be caused by an invalid range specification.

User response: Correct the error and resubmit the command.

DB21036E **The *command* command failed.**

User response: Retry the command with the trace active. If the problem persists, save the trace information and contact your technical support with the following information:

- Problem description
- DB2 message number
- Trace file

DB21037W **No data sources found.**

User response: There were no ODBC data sources found of the type (USER or SYSTEM) that you specified. Retry the command by specifying the other type (SYSTEM or USER).

DB21040E ***number* is not a valid in-doubt transaction number.**

Explanation: *number* is not one of the listed in-doubt transaction numbers.

User response: Select one of the listed transaction numbers and resubmit the command.

DB21041W ***number* indoubt transaction(s) not shown.**

Explanation: Command line processor could not display all the in-doubt transactions. The number of transactions not displayed is *number*.

User response: Finish processing the current list of indoubt transactions in order to reduce the total number of indoubt transactions in the system. Next, reissue the list indoubt transactions command.

DB21042E **You must specify a transaction number.**

Explanation: A transaction number must be specified with the commit (c), rollback (r), or forget (f) in-doubt transaction subcommands.

User response: Reissue the command with the appropriate transaction number.

DB21043E ***subcommand* is not a valid request.**

Explanation: The in-doubt transaction subcommand specified is not valid. Valid subcommands are:

Subcommand Description

| | |
|------------|--|
| c <number> | Heuristically commit the in-doubt transaction <number>. |
| r <number> | Heuristically roll-back the in-doubt transaction <number>. |
| f <number> | Heuristically forget the in-doubt transaction <number>. |
| l <number> | List all in-doubt transactions or the in-doubt transaction <number>. |
| q | Exit LIST INDOUBT TRANSACTION prompting. |

NOTE: The *number* command is not available in DB2 Extended Enterprise Edition.

User response: Correct the error and resubmit the command.

DB21044E **Cannot COMMIT transaction number *number*.**

Explanation: The transaction must be in a prepared state (i) when you issue the commit (c) subcommand.

User response: Issue another command.

DB21045E **Cannot ROLLBACK transaction number *number*.**

Explanation: The transaction must be in a prepared (i) or idle (e) state. when you issue the rollback (r) subcommand.

User response: Issue another command.

DB21046E **Cannot FORGET transaction number *number*.**

Explanation: The transaction must be in a heuristically committed (c) or heuristically rolled-back (r) state when you issue the forget (f) command.

User response: Issue another command.

DB21050E *state* is not a valid SQLSTATE.

Explanation: The sqlstate specified is not valid or could not be found. Valid states are numerical and two or five digits in length.

User response: Resubmit the command using a different state.

DB21051E **The command is not supported for this environment.**

Explanation: The command requested is currently not supported by the command line processor in the environment being used.

User response: Resubmit the command on a different platform or within a different environment.

DB21052I *command* has been submitted for execution.

Explanation: The command line processor is attempting to execute the indicated command. It will be unable to determine whether the command has completed successfully.

User response: If the command appears unsuccessful, submit it from outside the command line processor.

DB21053W **Automatic escalation will occur when you connect to a database that does not support *isolation-level*.**

Explanation: Isolation levels are database dependent. Some, like NC, are only supported by specific databases. If you select an isolation level that is not

supported by the database you are connecting to, it will automatically escalate to a supported level.

User response: Connect to a database that supports the isolation level you have selected, or select a different isolation level.

DB21054E **The command line processor encountered a system error with the *command* command. Reason code = *reason-code*.**

Explanation: While processing the portion of the command indicated by the reason code, the command line processor was unable to get enough memory.

| Reason code | Description | Syntax |
|-------------|--|--|
| 1 | Data file | LOAD FROM file/pipe/dev... |
| 2 | Lob data | LOBS FROM lob-path... |
| 3 | Directory | USING directory... |
| 4 | Source | FROM dir/dev... |
| 5 | Target | TO dir/dev... |
| 6 | Tablespace | TABLESPACE tblspace-name... |
| 7 | Tblspace-def | MANAGED BY... |
| 8 | Container data | USING (PATH..., PATH...) or USING ({FILE DEVICE}..., {FILE DEVICE}...) |
| 9 | Log path | log-directory ..., log-directory ... |
| 10 | Node list | node-number ..., node-number ... |
| 11 | Partitioned load option | [PARTITIONED DB CONFIG] partitioned-db-option [{partitioned-db-option}...] |
| 12 | Storage path | ON drive/path, drive/path... |
| 17 | Ingest field definition | (\$field-name field-type, ...) |
| 18 | SQL statement on the INGEST command | INSERT ..., UPDATE ..., MERGE..., or DELETE ... |
| 19 | Column name in the SQL statement on the INGEST command | any column name in the statement |
| 20 | Field name in the SQL statement on the INGEST command | any field name in the statement |

User response: You might be able to avoid the error by one of the following:

- Specify a larger value for one of the database manager configuration parameters that control heap or memory size.
- Modify the portion of the command indicated by the reason code to contain fewer elements. For example,

if the reason code is 17, remove some field definitions from the INGEST command.

DB21055W Command not valid for type 2 connections.

Explanation: The information returned by the GET CONNECTION STATE command does not apply to type 2 connections.

User response: Issue QUERY CLIENT to verify CONNECT = 1.

DB21056W Directory changes may not be effective until the directory cache is refreshed.

Explanation: If directory caching (DBM CFG dir_cache) is enabled, database, node, and DCS directory files are cached in memory. Directory changes may not become effective until the directory cache is refreshed. Refer to the dir_cache configuration parameter in the ADMIN guide for a description on directory caching.

User response: To refresh CLP's directory cache, issue a db2 TERMINATE. To refresh the directory information for another application, stop and restart that application. To refresh the directory information for the database, stop (db2stop) and restart (db2start) the database.

DB21057E Invalid tape device specified.

Explanation: The tape device passed to the operating system was not accepted. On Windows NT it must be of the form "\\.\TAPEx" where x represents the drive number (0 is the first).

User response: Reissue the command specifying a valid tape device.

DB21058E Invalid tape position specified.

Explanation: The tape mark position specified is not valid. On Windows NT the first tape position to which a backup is record is 1. Subsequent backup images begin at tape mark 2 and so on.

User response: Reissue the command specifying a valid tape position.

DB21059E Invalid tape blocksize specified.

Explanation: The tape mark blocksize specified is not in the range supported by the tape device. In addition, for backup/restore to work it must also be a factor of or multiple of 4096.

User response: Reissue the command specifying a valid tape size.

DB21060E General tape failure.

Explanation: An unexpected return code was returned from the tape operation.

User response: Resubmit the command. If the error occurs again, get help from your system administrator.

DB21061E Command line environment not initialized.

Explanation: You have tried to invoke the command line processor from a command window that was not started by db2cmd.exe.

User response: Issue DB2CMD to start a command window that has the command line processor environment initialized.

DB21070W One or more of the configuration parameters were set to AUTOMATIC in a case where the parameter does not support AUTOMATIC.

User response: If the parameter changes were submitted as a group, resubmit the changes individually to see which parameter changes were successful.

If only one parameter was submitted then this message indicates that the value AUTOMATIC is not supported for this parameter.

To find out which configuration parameters support the AUTOMATIC value, refer to the Administration Guide.

DB21071W Configuration parameter values were changed on all nodes, however not all nodes were able to change the values dynamically.

User response: If parameter changes were submitted in a group, issue the update command individually for more detailed information.

At the next restart of the unsuccessful nodes, the new values will take effect.

Changes to database manager configuration parameters can only take effect dynamically if you are attached to the instance. Not all configuration parameters support dynamic update. To see which parameters can be changed dynamically, refer to the Administration Guide.

To identify the nodes that failed in their dynamic application, attach to the instance on each node and issue the following command:

```
DB2 GET DBM CFG SHOW DETAIL
```

DB21080E **No previous RESTORE DATABASE command with REDIRECT option was issued for this database alias, or the information about that command is lost.**

Explanation: You have tried to run the RESTORE DATABASE command with the CONTINUE or ABORT option. However, either you didn't previously issue a RESTORE DATABASE command with the REDIRECT option, or you did issue such a command but for a different database alias than the one specified on the failing command. Another possible cause for this message could be that the information about a correctly issued previous RESTORE DATABASE ... REDIRECT command was lost. This can happen if the CLP back-end process terminates abnormally, or if you issue the TERMINATE command.

User response: Restart the whole redirected restore process by issuing the RESTORE DATABASE ... REDIRECT command and SET TABLESPACE CONTAINERS commands again. Then issue the RESTORE DATABASE ... CONTINUE command.

DB21081E **No sections were found in the db2cli.ini file.**

Explanation: You have tried to list CLI parameters in the db2cli.ini file by using the GET CLI CONFIGURATION command, but the file is empty. There are no sections in that file.

User response: To update the db2cli.ini file, use the UPDATE CLI CONFIGURATION command.

DB21082E **You don't have enough authority to run the UPDATE CLI CONFIGURATION command.**

Explanation: You need to have SYSADM authority to run this command.

User response: Obtain the necessary authority from your database administrator, then issue the command again.

DB21083E **Section *section* not found.**

Explanation: Section <section> was not found in the db2cli.ini file.

User response: Specify an existing section and issue the command again.

DB21084E **NEW and CONFIRM passwords not the same.**

Explanation: You were using the ATTACH or CONNECT command and specified that you wanted to change your password. The new password needs to be specified twice, by using the NEW and CONFIRM clauses, or by responding to the prompts. The two

passwords you specified as your new password were different.

User response: Specify the same password twice.

DB21085I **This instance or install (instance name, where applicable: *instance-name*) uses 32-or-64 bits and DB2 code release *release* with level identifier *level*.**

Explanation: This message is the output of the db2level command and provides detailed information about the code level of a given DB2 product. The information might be requested by DB2 service personnel to aid in resolving problems.

When the db2level command is run with a DB2 product that does not have a database manager instance, such as the IBM Data Server Driver Package or the IBM Data Server Driver for ODBC and CLI, the value of the first run-time token, *instance-name*, will be "*" .

User response: Record all displayed information to provide to DB2 service personnel.

The db2level executable should not be copied from one machine to another. Only DB2 installation and service maintenance programs should manipulate this file.

The db2level executable may not display all information about private test fixes that the customer has received from DB2 service personnel and installed on top of an officially supported service level.

DB21086I **The incremental RESTORE operation of this backup image completed successfully, but there are additional backup images that must be restored to complete the entire incremental RESTORE operation.**

Explanation: To complete an incremental RESTORE operation, every incremental backup image in its restore chain must be restored. The current operation succeeded, but there are additional backup images that must be restored before the overall RESTORE operation is complete.

User response: Restore the next backup image.

DB21089I **DB2 Administration Server *server-name* uses DB2 code release *ver-rel-mod* with level identifier *level-id*. Informational tokens are *build-id1*, *build-id2*, and Fix Pack *FixPak-number*. Product is installed at *install-path*.**

Explanation: This message is the output of the db2daslevel command and provides detailed information about the code level of a given DB2 Administration Server. The information may be

requested by DB2 service personnel to aid in resolving problems.

User response: Record all displayed information to provide to DB2 service personnel.

The db2daslevel executable should not be copied from one machine to another. Only DB2 installation and service maintenance programs should manipulate this file.

The db2daslevel executable may not display all information about private test fixes that the customer has received from DB2 service personnel and installed on top of an officially supported service level.

DB21100E **Stored procedure *procedure-name* exists in multiple schemas.**

Explanation: The specified procedure name was found in more than one schema.

User response: Re-issue the CALL command with a fully qualified procedure name (*schema.procedure-name*).

DB21101E **Too few parameters were specified for stored procedure "*procedure-name*" (expected *quantity*).**

Explanation: The definition of this stored procedure contains more parameters than those specified in the CALL command.

User response: Verify the number of parameters in the stored procedure and re-issue the command.

DB21102E **Too many parameters were specified for stored procedure "*procedure-name*" (expected *quantity*).**

Explanation: The definition of this stored procedure contains fewer parameters than those specified in the CALL command.

User response: Verify the number of parameters in the stored procedure.

DB21103E **The data type of parameter *parameter-number* is not supported by the CALL command in CLP.**

Explanation: Currently in CLP, one cannot call stored procedures with a parameter whose data type is that of parameter *parameter-number*.

User response: Do not call this stored procedure from CLP.

DB21104E **Parameter *parameter-number* for stored procedure "*procedure-name*" should be an INPUT parameter.**

Explanation: The parameter *parameter-number* is defined as an INPUT or INPUT/OUTPUT parameter.

However a "?" was specified for this parameter.

User response: Replace the "?" in parameter *parameter-number* by the input value of this parameter and re-issue the CALL command.

DB21105E **Parameter *parameter-number* for stored procedure "*procedure-name*" should be an OUTPUT parameter.**

Explanation: The parameter *parameter-number* is defined as an OUTPUT parameter. However a input value was specified for this parameter.

User response: Replace the input value in parameter *parameter-number* by a "?" and re-issue the CALL command.

DB21106E **Stored procedure "*procedure-name*" is undefined.**

Explanation: The stored procedure has not been defined or it has not been cataloged in the system catalogs.

User response: Verify that the stored procedure exists in the system catalogs. Try dropping and re-creating the stored procedure. Then re-issue the CALL command.

DB21107E **An error has occurred when accessing the file *file-name*. Reason code: *reason-code*.**

Explanation: For CLP commands REGISTER XMLSCHEMA, REGISTER XSROBJECT, ADD XMLSCHEMA DOCUMENT, COMPLETE XMLSCHEMA, or DECOMPOSE XML DOCUMENT, some possible reason codes include:

- 1 The file size is zero.
- 2 The file cannot be found.
- 3 Access to the file was denied. The user does not have permission to open the file.
- 4 Unexpected end of file encountered while reading from the file.
- 5 Access to the file was denied. The file is in use.
- 6 A media error was encountered while accessing the file.
- 7

The size of the document exceeds the maximum size supported by the command.

User response: Fix any errors and rerun the application.

DB21108E **The request cannot be executed because the DB2 server version does not support this functionality.**

Explanation: Some new functionality is not supported against older DB2 server versions.

User response: Execute the request against a DB2 server where the latest DB2 server version has been installed, or upgrade the server to the latest DB2 server version.

DB21109E **The syntax is incorrect for the *directive* directive. Reason Code: *reason-code*.**

Explanation: The new terminator cannot be set, because of a problem with the --#SET TERMINATOR syntax fragment in the directive. The terminator has not been changed. One of the following reason codes may apply:

Reason codes:

- 1 An invalid terminator was specified. A valid terminator contains one or two characters, but cannot contain a space or a tab.
- 2 The directive does not include a new terminator.
- 3 There must be a space after the words --#SET TERMINATOR.

User response: Revise the syntax of your directive, and then reissue the directive.

Chapter 61. DB21500 - DB21999

DB21500I **The DB2MSCS command completed successfully.**

Explanation: The user request was successfully processed.

User response: No action required.

DB21501E **The db2mscs command failed because the following invalid option was specified:** *option-name*

Explanation: You can create the infrastructure for DB2 failover support with Microsoft Cluster Server (MSCS) using the db2mscs utility.

Valid arguments for the db2mscs command are:

- -f InputFileName Specifies the DB2MSCS.CFG input file to be used by the MSCS utility. If this parameter is not specified, the DB2MSCS utility reads the DB2MSCS.CFG file that is in the current directory.
- -d TraceFileName Turns on debug tracing and specifies the name of the trace output file.
- -u InstanceName Undo the DB2MSCS operation for the instance.
- -l Username Specifies the user name of the domain account for the DB2 service (specified as domain\user).
- -p Password Specifies the password of the domain account for the DB2 service.

User response: Enter the command again with valid parameters.

DB21502E **Cannot open the configuration file** *file-name*.

Explanation: A configuration file could not be opened. Depending on the filename in the message text, this error can be explained as follow:

- If the filename in the message text was the name of the input file specified for the DB2MSCS command, then input file can not be found.
- If the filename was "db2system", then the database manager configuration file for target instance is missing.
- If the filename was "db2mscs.bak", then the backup configuration file could not be created in the instance directory.
- If the filename was "db2mscs.bak" and an undo operation was performed, then the backup configuration file from the instance directory could not be opened when performing the undo operation.

User response: Depending on the file that was in

error, the problem may be corrected as follow:

- If the filename in the message text was the name of the input file specified for the DB2MSCS command, then ensure that the file exists in the current directory, or that a fully qualified filename is specified for the command.
- If the database manager configuration file for target instance is missing, then drop and recreate the instance.
- If the backup configuration file could not be created in the instance directory, then ensure that the instance profile directory exists and that the current logon account has write access to the directory.
- If the backup configuration file from the instance directory could not be opened when performing the undo operation, then ensure that the MSCS disk that contains the instance profile directory is online on the current machine and retry the operation.

DB21503E **Not enough memory is available to process this command.**

Explanation: There was not enough memory to continue processing the command.

User response: Ensure that the system has sufficient real and virtual memory. Close all applications that are not in use to free up additional memory for the system.

DB21504E **The value *parameter-value* specified for parameter *parameter-name* exceeded the maximum length for that parameter. The maximum length of the parameter is *length* bytes.**

Explanation: The value *parameter-value* specified for the keyword *parameter-name* exceeded the maximum allowable limit for that parameter.

User response: Specify a value that satisfies length restrictions.

DB21505E **Parameter *parameter-name1* must be specified before parameter *parameter-name2* in the DB2MSCS configuration file *file-name*.**

Explanation: The sequence of parameters specified in the DB2MSCS configuration file is not valid. The group name must be specified before any other resource parameter can be specified. For each resource, the resource name parameter must be specified before any resource parameter can be specified.

User response: Modify the DB2MSCS configuration file so that the sequence of parameters is correct.

DB21506E The cluster *cluster-name* cannot be accessed. Ensure that the cluster name is correct and the cluster service on the current current machine has been started.

Explanation: The DB2MSCS utility could not open the cluster because either the cluster name was incorrect or the cluster service on the current machine has not been started.

User response: If the cluster service has not been started on the current machine, then start the cluster service by running the command "net start clussvc" or by starting the "Cluster" service from the Services dialog. If the cluster name was specified incorrectly in the DB2MSCS configuration file, then modify the cluster name and resubmit the command.

DB21507E The instance name *instance-name* is not valid.

Explanation: The instance name specified in the DB2MSCS configuration file is not valid, or the DB2INSTANCE environment variable was not set to a valid instance name.

User response: If the instance name was specified in the DB2MSCS configuration file, then verify that the instance name is valid and resubmit the command. If the instance name was not specified in the configuration file, then ensure the DB2INSTANCE environment variable is set to the name of a valid DB2 instance.

DB21509E The keyword *keyword* specified in configuration file *file-name* is only valid for the partitioned database instance.

Explanation: The keyword specified is only valid if the target instance is a partitioned database instance. For example, the DB2_NODE keyword should only be specified for the partitioned database instance.

User response: Comment out the invalid keyword in the configuration file and resubmit the command.

DB21510E Internet address *internet-address* specified in configuration file *file-name* is not valid.

Explanation: The value specified for either the IP address or the subnet mask does not conform with the internet address format. A valid internet address format has the form: "nnn.nnn.nnn.nnn", where nnn is a number from 0 to 255.

User response: Correct the invalid address in the configuration file and resubmit the command.

DB21511E Node *node-number* does not exist.

Explanation: The node number specified in the DB2_NODE keyword does not correspond to a valid database partition number.

User response: Correct the DB2_NODE parameter to specify an existing node number.

DB21512E The keyword *keyword* is not a valid DB2MSCS keyword.

Explanation: The keyword specified was not a valid DB2MSCS keyword.

User response: Use a valid DB2MSCS keyword. For more information about the keywords, please refer to the "Administration Guide".

DB21513E The DB2MSCS utility failed to create MSCS group *group-name* because of system error: *error-msg*

User response: Please refer to the Windows system error message for further information.

DB21514E The DB2MSCS utility failed to complete because of system error: *error-msg*

Explanation: The DB2MSCS utility failed to complete because of a Windows system error.

User response: Please refer to the Windows system error message for further information.

DB21515E The required resource property specified by keyword *keyword-name* is missing for the resource *resource-name*.

Explanation: A resource could not be created because one of its required parameters was not specified. For example, for IP Address resource, the IP Address and the subnet mask must be specified. For the Network Name resource, the network name must be specified.

User response: Ensure that the required parameter is specified and resubmit the command.

DB21516E DB2MSCS cannot bring resource *resource-name* online. Ensure that the properties of the resource are set correctly.

Explanation: After a resource is created, the DB2MSCS utility validates the resource by attempting to bring the resource online. Failing to bring a resource online indicates that either the resource property was not specified correctly or that the cluster network was not functioning properly.

User response:

- If a disk resource was in error, then ensure that the disk subsystem and the disk device driver is

functioning properly. The Event Viewer should be used to examine if any disk device driver problem was recorded in the Event logs.

- If an IP Address resource was in error, then ensure that the IP parameters are correct and that the network(s) where the IP address resides on is functioning properly. Also, the IP address that is specified for DB2 must NOT be used by any other machine in the network. If you are not sure what parameters to use for the IP address, consult with your network administrator.
- If a Network Name resource was in error, then ensure that the network is functioning properly, and that the value specified for the Netname parameter has not been used by any machine in the network. Note that the Network Name parameter is not required. As a work around, you may want to comment out the Network Name parameter and proceed.
- If a DB2 resource was in error, then the db2diag.log should be examined for any DB2 errors.

DB21517E **MSCS Network** *network-name* **is not active.**

Explanation: The network parameter specified for the IP address is not active.

User response: From the Cluster Administration view, activate or enable the target network and resubmit the command.

DB21518E **There is no active MSCS network.**

Explanation: The network parameter was not specified for the IP address and there was no network available to be used.

User response: A valid MSCS network must be configured. Please refer to your cluster documentation for how to add and configure an MSCS network.

DB21519E **DB2MSCS cannot bring the resource**
resource-name **offline. Ensure that the**
properties of the resource are set
correctly.

Explanation: The DB2MSCS could not bring a resource offline. The resource may be in use by the cluster software.

User response: Retry the operation. If the problem persists, run with the trace option and contact your IBM Service Representative for further assistance.

DB21520E **The DB2PATH profile variable is not defined.**

Explanation: The DB2PATH registry profile variable is not defined for the current machine. The DB2PATH must be set to the path where DB2 is installed.

User response: Set the DB2PATH to the directory where DB2 is installed using the db2set command. For example, db2set -g DB2PATH=D:\SQLLIB.

DB21521E **DB2MSCS cannot read from file**
file-name.

Explanation: The DB2MSCS utility can not read data from the indicated file.

User response: Ensure that the file is not locked and that the current logon user have sufficient authority to read the file.

DB21522E **Cannot open machine registry for**
machine *machine-name*. **Ensure that the**
machine is active and that the current
logon account has Local Administrator
authority.

Explanation: The DB2MSCS could not open the registry for the remote machine for read and write access. By default, only users that belong to the Local Administrator groups on that machine have read and write access to the machine registry. This error is also returned if the remote machine is not active.

User response: Ensure that the target machine is active, then logon to a domain account that belongs to the Local Administrator group on the target machine and resubmit the command.

DB21523E **Cannot close machine registry for**
machine *machine-name*. **Ensure that the**
machine is active and that the current
logon account has Local Administrator
authority.

Explanation: After opening the remote registry, the DB2MSCS utility failed to close the handle to the remote registry because of an internal error.

User response: Run with the trace option and contact your IBM Service Representative for further assistance.

DB21524E **Failed to create the resource**
resource-name. **System error:** *error-msg*

Explanation: The command failed to create the target resource because of a Windows system error.

User response: Refer to the Windows system error message for additional information.

DB21525E **Failed to add dependency to the**
resource *resource-name*. **System error:**
error-msg

Explanation: The command failed to add dependency for the target resource because of a Windows system error.

User response: Refer to the Windows system error

message for additional information.

DB21526E **Failed to move resource** *resource-name*.
System error: *error-msg*

Explanation: The command failed to add move resource because of a Windows system error.

User response: Refer to the Windows system error message for additional information.

DB21527E **No disk resource is specified for the group** *group-name*.

Explanation: At least one disk resource must be specified for each group.

User response: Assign one or more disk resource to the group indicated in the error message.

DB21528E **The value** *keyword-value* **specified for the INSTPROF_DISK keyword does not match any disk in the same group.**

Explanation: The INSTPROF_DISK keyword was used to specify the location where the content of instance profile directory will be copied to. The value for the INSTPROF_DISK keyword did not match one the name of a disk resource in the same group.

User response: Set the INSTPROF_DISK to the name of one of the disk resources in the same group.

DB21529E **The DB2MSCS utility cannot access the registry of machine** *machine-name*. **Ensure that the machine is active and that the current logon account has Local Administrator authority.**

Explanation: The DB2MSCS utility cannot access the registry of the target machine.

User response: Logon to a domain account that belongs to the Local Administrator group on the target machine and resubmit the command.

DB21530E **The DB2MSCS utility cannot access the cluster registry for the cluster** *cluster-name*. **Ensure that the cluster is active and the current logon account has Local Administrator authority.**

Explanation: To administer a cluster, users must have either administrative permissions on both nodes or specific permissions to administer the cluster. By default, the local Administrators group on both nodes has permissions to administer the cluster.

User response: Logon to an account that has sufficient access to the Cluster. To give a user permissions to administer a cluster without giving the user Administrative permissions on both nodes: Run the Cluster Administration GUI Right-click the cluster

name, and then click Properties Click Security (or Permissions) Specify which users and groups that may administer the cluster.

DB21531E **Cannot obtain property for MSCS disk.**
System error: *error-msg*

Explanation: The DB2MSCS utility can not obtain the drive letter from the MSCS disk resource. This problem usually occurs when the disk resource specified by the INSTPROF_DISK keyword is an IBM Netfinity disk resource, "IPSHA Disk".

User response: Do not use the INSTPROF_DISK keyword. Instead, use the INSTPROF_PATH keyword to explicitly specify the target location where the instance profile directory will be copied to.

DB21532E **A internal error occurred. File:** *file-name*,
Line *line-number*. **Please contact your IBM Service Representative.**

Explanation: The DB2MSCS failed because of an internal error.

User response: Run with the trace option and contact your IBM Service Representative for further assistance.

DB21533E **An error occurred during the migration of the DB2 instance, rc =** *error-code*.

Explanation: After all the required MSCS resources had been created, the DB2MSCS utility failed to migrate the DB2 instance to run in a clustered environment because of an internal error. During an instance migration, the utility performs the following steps:

- Copy the instance directory to the location specified by the INSTPROF_DISK or INSTPROF_PATH keyword.
- Move the DB2 registry profile variables from the machine registry to the cluster registry.
- Set the DB2INSTPROF registry variable to point to the new instance profile location.
- Set the DB2CLUSTERLIST to the name of the current machine.

User response: Before running the DB2MSCS utility, ensure that:

- On the current machine, the instance can be started and stopped successfully from the command line.
- On other cluster node(s), the same instance must be stopped and optionally dropped.
- All the disk resources are active on the current machine and can be moved back and forth successfully between the cluster nodes.
- The current logon user has sufficient access to the local machine registry and the cluster registry.

- If the problem persists, contact your IBM Service Representative and provide both the DB2MSCS traces and DB2 traces.

DB21534E **An error occurred during addition of MSCS node to the DB2 instance, rc = *error-code*.**

Explanation: The utility failed to add the other MSCS node to the DB2 instance. During this operation, the utility will do the followings:

- Update the DB2 cluster machine list by adding the name of the target machine to the DB2CLUSTERLIST registry variable.
- Create the DB2 service and the registry instance profile for the current DB2 instance on the target node.

User response: Before running the DB2MSCS utility, ensure that:

- On the current machine, the instance can be started and stopped successfully from the command line.
- On other cluster node(s), the same instance must be stopped and optionally dropped.
- All the disk resources are active on the current machine and can be moved back and forth successfully between the cluster nodes.
- The current logon user has sufficient access to the target machine registry and the cluster registry.
- If the problem persists, contact your IBM Service Representative and provide both the DB2MSCS traces and DB2 traces.

DB21535E **The instance-owning database partition server is not on the current machine.**

Explanation: When migrating a partitioned database instance, the DB2MSCS utility must be run on the instance owning machine.

User response: Run the DB2MSCS utility from the instance owning machine.

DB21536E **The username *userid* is not valid.**

Explanation: The username specified is not valid.

User response: Specify a valid username.

DB21537E **The password *password* is not valid.**

Explanation: The password specified is not valid.

User response: Specify the correct password.

DB21538E **The password for the account *account-name* has expired.**

Explanation: The password for the target account has expired.

User response: Reset the password and resubmit the command.

DB21540E **Group *group-name* requires at least one network name resource.**

Explanation: When migrating a partitioned database instance, a network name resource must be created for the group that contains the instance owning node.

User response: Specify to create a network name resource in the group indicated.

DB21541E **An error occurred when removing the MSCS node from the DB2 instance, rc = *error-code*.**

Explanation: During an "undo" operation, the utility failed to remove an MSCS node from the DB2 instance because of an internal error.

User response: Manual clean up is required. To manually clean up the instance do the following:

- Stop and drop the DB2 instance.
- Remove all DB2 resources and their dependent resources from the Cluster Administrator window.

DB21542E **An error occurred while attempting to remove failover support for the instance. Failover support is still active for this instance, rc = *error-code*.**

Explanation: During an "undo" operation, the utility failed to uncluster the DB2 instance because of an internal error.

User response: Manual clean up is required. To manually clean up the instance do the followings: Stop and drop the instance. Remove all DB2 resources and their dependent resources from the Cluster Administrator window.

DB21543E **The resource name *resource-name* does not match any IP resource in the same group.**

Explanation: A network name resource must be configured to depend on an IP Address resource in the same resource group.

User response: Specify the name of an IP Address resource that resides in the same group as a dependency for the Network Name resource.

DB21544E **The MSCS resource *resource-name* already exists.**

Explanation: The resource name specified already exists in the cluster.

User response: Specify a different resource name.

DB21545E **The module *file-name* was loaded, but the function *function-name* is not found. Ensure that you are running on the version of DB2 that supports MSCS.**

Explanation: The utility failed to obtain the address of a required function because the version of DB2 is not compatible with the version of the DB2MSCS utility.

User response: Use the version of the DB2MSCS utility that is shipped with the DB2 product.

DB21546E **The module *file-name* could not be loaded.**

Explanation: The utility failed to load the required DLL.

User response: Reinstall the DB2 product.

DB21547E **Error occurred while moving group *group-name* to node *node-number*. System error: *error-msg***

Explanation: The utility failed to move the group to the target node because one or more resources can not be moved.

User response: Ensure that all cluster nodes are active and that all disk resources can be moved back and forth between cluster nodes. If the problem persists, contact your IBM Service Representative.

DB21548E **A logon account for the DB2 service must be specified for a partitioned database system. Specify a valid logon account using the DB2_LOGON_USERNAME and DB2_LOGON_PASSWORD keywords.**

Explanation: The DB2 service for a partitioned database system must be configured to be run under a valid domain account.

User response: Specify a valid domain account using the DB2_LOGON_USERNAME and DB2_LOGON_PASSWORD keywords.

DB21549N **Network name *network-name* specified for node *node-number* is not valid.**

Explanation: The network name value specified could not be resolved.

User response: Ensure that the network name is registered in the DNS server or in a local etc/hosts file.

DB21600N **The *command* command did not complete successfully.**

Explanation: Errors were encountered during the execution of this command.

User response: Contact your technical service representative. Information in the db2diag.log file will enable your service representative to determine the cause of the failure.

Chapter 62. DB22000 - DB22499

DB22000E **Error: The syntax of the DB2LSWTCH command is incorrect.**

Explanation: The local switcher promotes or demotes the DB2 copy it is installed under:

- db2lswtch.exe -db2 -promote
- db2lswtch.exe -client -promote
- db2lswtch.exe -all -promote
- db2lswtch.exe -db2 -demote
- db2lswtch.exe -client -demote
- db2lswtch.exe -all -demote

The command options are:

- -db2 -promote to promote the current copy to the default DB2 copy.
- -client -promote to promote the current copy to the default IBM database client interface copy.
- -all -promote to promote the current copy to the default DB2 and IBM database client interface copy.
- -db2 -demote to demote the current default DB2 copy.
- -client -demote to demote the current default IBM database client interface copy.
- -all -demote to demote the current default DB2 and IBM database client interface copy.

User response: Enter the command again using valid parameters.

DB22001E **An error occurred during the initialization of the local switcher environment. No change was done to any copy on the system.**

Explanation: An error occurred when the local switcher tried to look up information in the Windows system registry.

User response: Contact your IBM service representative.

DB22002E **An error occurred during the promotion of the copy *name*. The promotion will continue and this copy will be assigned as the default on the system.**

Explanation: An error occurred during one of the tasks that the local switcher performs to promote the copy it is installed under to be the default one on the system. The local switcher will attempt to complete the promotion and will continue with the remaining tasks. The copy will still be assigned as the default on the system.

User response: Manual intervention is needed to complete the task or tasks that failed during the promotion of the copy.

DB22003E **An error occurred during the demotion of the DB2 copy *name*. The demotion will be aborted and this copy will remain as the default on the system.**

Explanation: An error occurred during one of the tasks that the local switcher performs to demote the copy it is installed under to no longer be the default on the system. The demotion will be aborted and the local switcher will be called to promote the same copy so that any tasks that have been completed are rolled back.

User response: Try to manually correct the problem that resulted in this error and attempt to run the same operation again. If the problem persists, contact your IBM service representative.

DB22004I **The copy *name* has been successfully demoted and is no longer the default DB2 copy on the system.**

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

DB22005I **The copy *name* has been successfully promoted to be the default DB2 copy on the system.**

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

DB22006E **An error occurred during the process of updating the global system environment variables to include the environment variables for the copy *name*.**

Explanation: During promotion of the DB2 copy it is installed under, the local switcher modifies the global system environment variables such as PATH, INCLUDE, LIB and CLASSPATH to add values pointing to the path of its copy to them. This operation has encountered a problem.

User response: Inspect the global system environment and ensure that any environment variables that need to be modified correctly include values pointing to the path of the copy that has been promoted.

DB22007E **An error occurred during the process of updating the global system environment variables to remove the environment variables for the copy *name*.**

Explanation: During demotion of the DB2 copy it is installed under, the local switcher removes values pointing to the path of the copy from global system environment variables such as PATH, INCLUDE, LIB and CLASSPATH. This operation has encountered a problem.

User response: Inspect the global system environment and ensure that any environment variables that need to be modified no longer include values pointing to the path of the copy that has been demoted.

DB22008E **An error occurred when registering the IBM ODBC driver for the copy *name* as the default driver on the system.**

Explanation: During promotion of its copy, the local switcher registers the default IBM DB2 ODBC driver and ensures that this driver points to the path of the default IBM database client interface copy. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22009E **An error occurred when unregistering the IBM ODBC driver for the copy *name* as the default driver on the system.**

Explanation: During demotion of the copy it is installed under, the local switcher unregisters the default IBM DB2 ODBC driver and ensures that this driver no longer points to the path of the demoted copy. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22010E **An error occurred when registering the IBM .NET data provider for the copy *name* in the global assembly cache.**

Explanation: During the promotion of the copy it is installed under, the local switcher copies the IBM .NET provider for this copy to the Global Assembly Cache of the Windows operating system. This is done using special Windows APIs. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22011E **An error occurred when unregistering the IBM .NET data provider for the copy *name* in the global assembly cache.**

Explanation: During the demotion of the copy it is installed under, the local switcher removes the IBM

.NET provider of this copy from the Global Assembly Cache of the Windows operating system. This is done using special Windows APIs.

This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22012E **An error occurred when registering the IBM OLEDB provider for the copy *name* as the default provider on the system.**

Explanation: During the promotion of the DB2 copy it is installed under, the local switcher registers the default IBM OLEDB provider and ensures that the default IBM OLEDB provider points to the provider of this copy. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22013E **An error occurred when unregistering the IBM OLEDB provider of the copy *name* as the default provider on the system.**

Explanation: During the demotion of the DB2 copy it is installed under, the local switcher unregisters the default IBM OLEDB provider and ensures that the default IBM OLEDB provider no longer points to the provider of this copy. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22014E **An error occurred during the update of the ODBC system DSN *name*.**

Explanation: During the demotion of the copy it is installed under, the local switcher modifies any system Data Source Names to ensure they still function. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22015E **An error occurred during the operating system call *func_name*. The operating system message is *msg* and the return code is *return_code*.**

Explanation: A call to an operating system function has failed. The operating system error message and return code contain further information on the reason for the failure.

User response: Contact your IBM service representative.

DB22016E **An error occurred when registering the IBM Add-in of the copy *name* with VisualStudio 2003.**

Explanation: During the promotion of its copy, the local switcher registers the IBM Visual Studio 2003 Add-in with the installed copy of Visual Studio. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22017E **An error occurred when unregistering the IBM Add-in of the copy *name* with VisualStudio 2003.**

Explanation: During the demotion DB2 copy it is installed under, the local switcher unregisters the IBM Visual Studio 2003 Add-in with the installed copy of Visual Studio. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22018E **An error occurred while trying to start the DB2 Management Service of the copy *name*.**

Explanation: During the promotion of the DB2 copy it is installed under, the local switcher starts the DB2 Management Service belonging to this copy. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22019E **An error occurred while trying to stop the DB2 Management Service of the copy *name*.**

Explanation: During the demotion of the DB2 copy it is installed under, the local switcher stops the DB2 Management Service belonging to this copy. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22020E **The copy *name* is not the default DB2 copy and therefore cannot be demoted.**

Explanation: The local switcher has been run with a -db2 -demote switch from an alternate copy. The local switcher can only demote the default DB2 copy on the system.

User response: If you need to demote the default DB2 copy, run the local switcher from the installation path of the default DB2 copy.

DB22021E **Another copy *name* is defined as the default DB2 copy on the system. Promotion of another copy cannot proceed unless the current default copy is demoted.**

Explanation: The local switcher has been called to promote an alternate copy of DB2 while another copy is already assigned as the default DB2 copy on the system.

User response: Demote the current default DB2 copy before promoting an alternate copy.

DB22022E **An error occurred while setting the "default" value under the HKLM\Software\IBM\InstalledCopied registry key to the name of the copy *name*.**

Explanation: The "default" value under the HKLM\Software\IBM\InstalledCopied key contains the name of the current default copy. The local switcher has encountered a problem setting this value to the name of the copy that has been promoted.

User response: Attempt the operation again. If the problem persists, contact your IBM service representative.

DB22023E **An error occurred while updating the "default" value under the HKLM\Software\IBM\InstalledCopied registry key to remove the name of the copy *name*.**

Explanation: The "default" value under the HKLM\Software\IBM\InstalledCopied key contains the name of the current default copy. The local switcher has encountered a problem removing the name of the copy that has been demoted.

User response: Attempt the operation again. If the problem persists, contact your IBM service representative.

DB22024E **An error occurred while registering the COM server *name*.**

Explanation: During the promotion of the DB2 copy it is installed under, the local switcher registers the COM servers that are part of this copy with the operating system. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22025E **An error occurred while unregistering the COM server *name*.**

Explanation: During the demotion of the DB2 copy it is installed under, the local switcher unregisters the COM servers that are part of this copy with the

operating system. This operation has encountered a problem.

User response: Contact your IBM service representative.

DB22026E Error: The syntax of the db2swtch command is incorrect.

Explanation: The db2swtch utility can display the DB2 copies or IBM data server driver copies installed on the system, and set the default DB2 copy or default IBM database client interface copy. This utility can also be used to redirect the references in machine.config between common IDS .NET data provider and SQLI IDS .NET data provider:

```
db2swtch [-db2|-client]
          [-d {name of DB2 copy or IBM data
            server driver copy}]
          [-l]
          [-IDS common|SQLI]
          [-h|-?]
```

The command options are:

(no arguments)

launches the utility in graphical mode

-d {name of DB2 copy}

sets the default DB2 and default IBM database client interface copy to the name specified

-db2 -d {name of DB2 copy}

sets the default DB2 copy to the name specified

-client -d {name of DB2 copy or IBM data server driver copy}

sets the default database client interface copy to the name specified

-l displays the list of DB2 copies and IBM data server driver copies on the system

-IDS common

redirects the IDS .NET data provider reference in machine.config to common IDS .NET data provider

-IDS SQLI

redirects the IDS .NET data provider reference in machine.config to SQLI IDS .NET data provider

-h|-? displays help

User response: Enter the command again using valid parameters.

DB22027E An error occurred when registering the performance counters for the copy *name*.

Explanation: During the promotion of the DB2 copy it is installed under, the local switcher registers the DB2 performance counters for this copy. This operation has encountered a problem. This is not usually a severe problem unless you intend to use the DB2 performance counters.

User response: Contact your IBM service representative.

DB22028E An error occurred when unregistering the performance counters for the copy *name*.

Explanation: During the demotion of the DB2 copy it is installed under, the local switcher unregisters the DB2 performance counters for this copy. This operation has encountered a problem. This is not usually a severe problem unless you intend to use the DB2 performance counters.

User response: Contact your IBM service representative.

DB22029E An error occurred when registering the DB2 system tray for the default instance of the copy *name*.

Explanation: During the promotion of the DB2 copy it is installed under, the local switcher registers the DB2 system tray for the default instance of this copy. This operation has encountered a problem. This is not usually a severe problem.

User response: Contact your IBM service representative.

DB22030E An error occurred when unregistering the DB2 system tray for the copy *name*.

Explanation: During the demotion of the DB2 copy it is installed under, the local switcher unregisters the DB2 system tray for this copy. This operation has encountered a problem. This is not usually a severe problem.

User response: Contact your IBM service representative.

DB22031E An error occurred when modifying the display name of the shortcut group for the copy *name*.

Explanation: During the promotion or demotion of the DB2 copy it is installed under, the local switcher modifies the display name of the shortcut group for this copy to indicate if it is the default DB2 copy or not. This operation has encountered a problem. This is not usually a severe problem.

User response: Contact your IBM service representative.

DB22032I **The copy *name* has been successfully demoted and is no longer the default IBM database client interface copy on the system.**

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

DB22033I **The copy *name* has been successfully demoted and is no longer the default DB2 and IBM database client interface copy on the system.**

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

DB22034I **The copy *name* has been successfully promoted to be the default IBM database client interface copy on the system.**

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

DB22035I **The copy *name* has been successfully promoted to be the default DB2 and IBM database client interface copy on the system.**

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

DB22036E **The copy *name* is not the default IBM database client interface copy and therefore cannot be demoted.**

Explanation: The local switcher has been run with a -client -demote switch from an alternate copy. The local switcher can only demote the default IBM database client interface copy on the system.

User response: If you need to demote the default database client interface copy, run the local switcher from the installation path of the default client interface copy.

DB22037E **The copy *name* is not default DB2 and IBM database client interface copy and therefore cannot be demoted.**

Explanation: The local switcher has been run with a -all -demote switch from an alternate copy. The local

switcher can only demote the copy which is both the default DB2 copy and the default IBM database client interface copy on the system.

User response: If you need to demote the default DB2 and IBM database client interface copy, run the local switcher from the installation path of the copy which is both the default DB2 copy and the default IBM database client interface copy.

DB22038E **Another copy *name* is defined as default IBM database client interface copy on the system. Promotion of another copy cannot proceed unless the current default database client interface copy is demoted.**

Explanation: The local switcher has been called to promote an alternate copy of database client interface while another copy is already assigned as the default IBM database client interface copy on the system.

User response: Demote the current default IBM database client interface copy before promoting an alternate copy.

DB22039E **An error occurred while setting the "defaultClientInterfaceCopy" value under the HKLM\Software\IBM\InstalledCopied registry key to the name of the copy *name*.**

Explanation: The "defaultClientInterfaceCopy" value under the HKLM\Software\IBM\InstalledCopied key contains the name of the current default copy. The local switcher has encountered a problem setting this value to the name of the copy that has been promoted.

User response: Attempt the operation again. If the problem persists, contact your IBM service representative.

DB22040E **An error occurred while updating the "defaultClientInterfaceCopy" value under the HKLM\Software\IBM\InstalledCopied registry key to remove the name of the copy *name*.**

Explanation: The "defaultClientInterfaceCopy" value under the HKLM\Software\IBM\InstalledCopied key contains the name of the current default database client interface copy. The local switcher has encountered a problem removing the name of the copy that has been demoted.

User response: Attempt the operation again. If the problem persists, contact your IBM service representative.

DB22041E The DB2 copy *name* does not exist.

Explanation: The specified DB2 copy does not exist.

User response: Specify an existing DB2 copy and reissue the command.

DB22042E An error occurred when setting up the DB2 instance configuration for the IBM Tivoli Monitoring for Databases: DB2 Agent in the copy *name*.

Explanation: During the promotion of the DB2 copy the local switcher attempted to set up the DB2 instance configuration for the Monitoring Agent for DB2 in this copy.

User response: Contact your IBM service representative.

DB22043E An error occurred when removing the DB2 instance configuration for the IBM Tivoli Monitoring for Databases: DB2 Agent in the copy *name*.

Explanation: During the demotion of the DB2 copy the local switcher attempted to remove the DB2 instance configuration for the Monitoring Agent for DB2 in this copy.

User response: Contact your IBM service representative.

DB22044E An error occurred when creating the configuration file for the IBM Tivoli Monitoring for Databases: DB2 Agent in the default copy installation directory.

Explanation: An internal error occurred when creating the configuration file for the Monitoring Agent for DB2 in the default copy installation directory.

User response: Ensure that you have write access to the default copy installation directory.

DB22045E An error occurred when removing the configuration file for the IBM Tivoli Monitoring for Databases: DB2 Agent in the default copy installation directory.

Explanation: An internal error occurred when removing the configuration file for the Monitoring Agent for DB2 in the default copy installation directory.

User response: Ensure that you have write access to the default copy installation directory.

Chapter 63. DB29000 - DB29499

DB29320W Output has been truncated.

Explanation: The full query result cannot be fetched as supplied.

User response: The CLP query can be rewritten to fetch a shorter string. Use of another interface to DB2 can also help in overcoming the CLP's limitation.

Chapter 64. DB29500 - DB29999

DB29501E DB2 has detected an inconsistent environment. Please investigate the following : *error*

Explanation: An error in the operating environment has been detected.

User response: This error may arise due to an inconsistent DB2 or operating system environment. Attempt to correct the problem identified in the error message and re-issue the command.

DB29502E An incompatible Java Runtime Environment has been detected. Please verify that the required level of the Java Runtime Environment has been installed. The Java application may be launched using the Java runtime installed with DB2.

Explanation: The installed Java runtime is not supported by the DB2 Java application.

User response: The installed DB2 Java runtime will be used to launch the DB2 Java application. Other Java applications using the installed Java runtime should be terminated prior to executing the DB2 Java application to ensure correct behaviour.

DB29503E Connections to multiple databases is not supported.

Explanation: You cannot connect to more than one database.

User response: Connect to only one database.

DB29504E Connections to multiple users is not supported.

Explanation: You cannot connect to more than one user.

User response: Connect to only one user.

DB29523W Are you sure you want to CANCEL the query?

DB29524E LogSQLException() Failed

DB29525E LogSQLException() *param-1 param-2*

DB29526E You cannot drop this table; it does not exist. Please select a job that has results, and try again, or select a recurring job

that is scheduled and try again.

Explanation: A request was made to drop a result table for a job which did not have an associated result table.

User response: Select a job that has a results or select a recurring job that is scheduled and try again.

DB29527W Are you sure you want to cancel job #*param-1*?

DB29528W Are you sure you do NOT want to display job #*param-1*?
WARNING: If you select this option this job will never be displayed again.

DB29530W This option will remove confirmation messages from actions like 'cancel' and 'hide'. Is this something you really want to do?

DB29537E Unknown Type: *param-1*

DB29542E Connection SQL error: *param-1*
Error Number: *param-2*
Error Message: *param-3*

DB29543E Resource DLL rqsres.dll not found.

Explanation: Resource DLL rqsres.dll could not be loaded successfully.

User response: During initialization of DB2 Query Patroller Recurring Query Scheduler the resource DLL rqsres.dll could not be loaded as it was either damaged or deleted. Please reinstall DB2 Query Patroller QueryEnabler.

DB29544E Resource DLL qeres.dll not found.

Explanation: Resource DLL qeres.dll could not be loaded successfully.

User response: During initialization of DB2 Query Patroller QueryEnabler the resource DLL qeres.dll could not be loaded as it was either damaged or deleted. Please reinstall DB2 Query Patroller QueryEnabler.

DB29545E Name invalid: Query Name can only be 20 characters or less in length. Please change the name accordingly and retry.

Explanation: A name for a query was specified which

DB29546E

was longer than 20 characters.

User response: Specify a query name that has at most 20 characters.

DB29546E **Name invalid: Query Name can only contain alphanumeric characters and spaces...("a..z", "A..Z", "0..9"). Please change the name accordingly and retry.**

Explanation: A name for a query was specified which contained non-alphanumeric characters.

User response: Specify a query name composed of only alphanumeric characters.

Chapter 65. DB210000 - DB210499

DB210200I All applications must disconnect from the database before the changes become effective.

Explanation: The ADD or DROP DATALINKS MANAGER command has been successfully processed. However, the change cannot be made effective until all applications have disconnected from the database. After the applications disconnect from the database, the first connection to the database will cause the changes to become effective.

User response: Ensure that all applications disconnect from the database and issue a CONNECT statement.

DB210201I Reconcile utility should be run on database tables that contain links to files on the DB2 Data Links Manager. The DB2 Data Links Manager does not perform any unlink processing for these files.

Explanation: The DROP DATALINKS MANAGER command has been successfully processed. Confirm that no database tables contained links to files on the DB2 Data Links Manager before dropping the DB2 Data Links Manager. If such links exist, they should be removed from the database tables by using the reconcile utility. Note that the files themselves will remain in linked state on the file system. For additional details, see the usage notes for the DROP DATALINKS MANAGER command in the *Command Reference*.

DB210202E Less than *number* MB of memory is dedicated to your server. No recommendation has been made. The current values match the suggested values.

Explanation: Because the amount of memory dedicated to the server is so small, the Performance Configuration wizard cannot make a recommendation. Configuration parameters will not be changed.

User response: If you can dedicate more memory to the server, re-run the command with a larger value specified for the option MEM_PERCENT.

DB210203I AUTOCONFIGURE completed successfully. Database manager or database configuration values may have been changed if you chose to apply changes. The instance must be restarted before any such applied changes come into effect. You may also want to rebind your packages after the new configuration parameters take effect so that the new values will be used.

Explanation: The Configuration Advisor has generated its recommendations, and has updated the configuration parameters and buffer pool sizes as requested by the user.

The changes will not take effect until the instance is restarted on the server.

Packages were bound with the old configuration parameters. They need to be rebound, after the new parameters take effect, to exploit these new parameters.

User response: When you are ready to use the new configuration parameters, stop the instance, and then restart the instance. Rebind your packages if necessary

DB210204E An error occurred when the Configuration Advisor tried to get system information from your server.

Explanation: An unexpected error occurred. The Configuration Advisor cannot continue.

User response: Contact your system administrator.

DB210205W The Configuration Advisor was unable to increase the sizes of the buffer pools due to other memory requirements determined from your responses. The buffer pool sizes are left unchanged. The use of the suggested set of configuration values may cause paging on the server.

Explanation: This is a warning that there may not be enough memory dedicated to your server to run the workload that you specified.

User response: Review your input for this command to check that the workload description is appropriate or add more memory to your server.

DB210206W The transactions rate you entered is more than ten times the average number of connected applications. If you are using a transaction manager, ignore this message. If not, consider changing the rate.

Explanation: The Configuration Advisor verifies that the transactions rate is reasonable. The transaction rate may be too high if you are not using a transaction manager.

User response: If you are using a transaction manager, ignore this message. Otherwise, use a lower number of transactions per minute, or increase the average number of connected applications.

DB210207E Unable to autoconfigure *database-name*. Creation of database *database-name* failed.

Explanation: An unexpected error occurred. The Configuration Advisor cannot continue, and the database you were trying to create failed.

User response: Try running CREATE DATABASE without the AUTOCONFIGURE option, and then running AUTOCONFIGURE in order to configure your database.

DB210208E Cannot autoconfigure databases in EEE system. Creation of database *database-name* failed.

Explanation: The AUTOCONFIGURE option is not available for EEE. The database was not created.

User response: Try running CREATE DATABASE without the AUTOCONFIGURE option, and then running AUTOCONFIGURE in order to configure your database.

DB210209I The database was created successfully. Please restart the instance so configuration changes take effect.

Explanation: The Configuration Advisor has updated the database manager configuration parameters, database configuration parameters, and buffer pool sizes in the SYSBUFFERPOOLS catalog.

The changes will not take effect until the instance is stopped on the server.

Packages were bound with the old configuration parameters. They need to be rebound, after the new parameters take effect, to exploit these new parameters.

User response: When you are ready to use the new configuration parameters, stop the instance, and then start the instance. Rebind your packages if necessary.

DB210210E The parameter for APPLY is not supported by the Configuration Advisor for this server release.

Explanation: Only APPLY DB OR DBM or APPLY NONE are valid parameters for the Configuration Advisor for this server release.

User response: Reissue the command using APPLY DB OR DBM or APPLY NONE.

DB210211W AUTOCONFIGURE completed successfully, but has only calculated recommendations for the current database partition. You may wish to run AUTOCONFIGURE on all database partitions or propagate these recommendations to all database partitions on your system.

Explanation: The Configuration Advisor has generated its recommendations, and has updated the configuration parameters and buffer pool sizes as requested by the user for only the current database partition.

The changes will not take effect until the instance is restarted on the server, but should be propagated to other database partitions on this system by running AUTOCONFIGURE for each database partition, or through separate user action.

Packages were bound with the old configuration parameters. They need to be rebound, after the new parameters take effect, to exploit these new parameters.

User response: When you are ready to use the new configuration parameters, stop the instance, and then start the instance. Rebind your packages if necessary.

DB210212W The recommendations made by the Configuration Advisor require a greater percentage of memory than that specified. No recommendation has been made. The current values match the suggested values.

Explanation: There is not enough memory available to be allocated for the recommendations made by the Configuration Advisor.

User response: Resubmit the command, specifying a greater percentage of memory.

DB210213W The recommended number of disks for the specified inputs exceeds the number available. No recommendation has been made. The current values match the suggested values.

Explanation: There are not enough disks available for the recommendations made by the Performance Configuration Wizard.

User response: Resubmit the command once there are more disks available.

DB210214W The recommended amount of memory for the specified inputs exceeds the amount available. No recommendation has been made. The current values match the suggested values.

Explanation: There is not enough memory available to be allocated for the recommendations made by the Configuration Advisor.

User response: Resubmit the command either with different input values or once a greater amount of memory is available.

DB210215W The Configuration Advisor was unable to assign a minimum amount of memory to the buffer pools due to other memory requirements determined from your responses.

Explanation: The Configuration Advisor is unable to provide a set of recommendations based on the specified database requirement with the specified memory resource.

User response: Increase the percentage of physical memory allocated, if more memory resource is available. Otherwise, increase the amount of physical memory on the server.

DB210220E The command number specified is not valid.

Explanation: An invalid command number was specified for the EDIT or RUNCMD command. This command number does not exist in the current CLP interactive mode command history.

User response: Run the HISTORY command to see a list of valid command numbers and resubmit the EDIT or RUNCMD command specifying a valid command number.

DB210221E The editor *editor* could not be launched.

Explanation: An invalid editor was specified for the EDIT command. This editor either does not exist or it is not contained in the PATH.

User response: The editor used in the EDIT command is decided in the following order:

- by using the value of the EDITOR parameter in the EDIT command, if it was specified
- by using the value of the DB2_CLP_EDITOR registry variable, if it was set
- by using the value of the VISUAL environmental variable, if it was set

- by using the value of the EDITOR environmental variable, if it was set.

Specify a valid editor contained in the PATH in the EDITOR parameter of the EDIT command or in the appropriate registry/environmental variable.

DB210222E The value specified for *parameter* is not in the valid range of *number* and *number*.

Explanation: An invalid value was specified for *parameter*. This value does not lie within the valid range of *number* and *number*.

User response: Refer to the appropriate documentation to find out the valid range for *parameter* and resubmit the command specifying a valid value.

DB210223E The command *command* can only be run in the CLP's interactive mode.

Explanation: An attempt was made to run a DB2 command in either CLP's command or batch mode, but this command can only be run in CLP's interactive mode.

User response: Reissue the same command within the CLP's interactive mode.

Chapter 66. DB216000 - DB216499

DB216001I The DECOMPOSE XML DOCUMENTS command successfully decomposed all documents. The number of documents requested was *number-requested*.

Explanation: No errors were encountered during the execution of this command.

User response: No user action is required.

DB216002W One or more documents could not be decomposed. The number of documents successfully decomposed was *number-successful*. The number of documents attempted was *number-attempted*.

Explanation: Error conditions prevented one or more documents from being decomposed. For more details on the status of each document, refer to messages file specified by the MESSAGES parameter.

User response: To decompose the XML documents that were not successfully decomposed:

1. Review the error message information in the file specified by the MESSAGES parameter. If necessary, see the db2diag log file for more details about each error. The pertinent entries in the db2diag log file are identified by the document ID.
2. Correct the errors described in the error messages file.
3. Issue the DECOMPOSE XML DOCUMENTS command again, as follows:
 - If you specified a positive COMMITCOUNT value, reissue the command only for the documents that were not successfully decomposed.
 - If you specified a COMMITCOUNT value of 0, use one of the following options:
 - Perform a rollback, and then reissue the command for all the documents.
 - Commit the changes, and then reissue the command only for the documents that were not successfully decomposed.

Chapter 67. DB250000 - DB250499

DB250000I The command completed successfully.

Explanation: The Command Line Processor Plus command completed successfully.

User response: No action is required.

DB250001I CLPPlus has successfully read the configuration file named *configuration-file-name*.

Explanation: When the DB2DS DRIVER_CFG_PATH environment variable is set, or when the path to the configuration file is found automatically by Command Line Processor Plus (CLPPlus), CLPPlus will read the configuration file during a connect attempt. This message is only displayed when verbose is set to on using the CLPPLUS command.

User response: No user response is required.

DB250002E CLPPlus failed to read the configuration file named *configuration-file-name* in the directory named *directory*.

Explanation: During a connect attempt, Command Line Processor Plus (CLPPlus) failed to read the configuration file specified by the DB2DS DRIVER_CFG_PATH environment variable.

User response: Verify that the information in the configuration file is correct by performing the following steps:

- Verify that the path to the configuration file matches what is specified by the DB2DS DRIVER_CFG_PATH environment variable.
- Verify that the name of the configuration file matches what is specified by the DB2DS DRIVER_CFG_PATH environment variable.

DB250003E Could not find a DSN alias named *dsn-alias-name* in the configuration file named *configuration-file-name*. *dsn-alias-name* will be used as the database name in the subsequent interactive connect attempt.

Explanation: When a DSN alias name specified during a connect attempt is not found in the configuration file read by CLPPlus, the DSN alias name is used as a database name in the subsequent interactive CLPPlus connect attempt. All other values for the connect must be specified by the user.

User response: Continue with the interactive connect. You can also check the contents of the configuration file

and use an existing DSN alias or create one for future use.

DB250004E The connection failed because an empty string was specified for the value for the database name.

Explanation: Command Line Processor Plus (CLPPlus) requires you provide a valid database name when prompted. Entering an empty string for a database name is not allowed.

User response: When prompted to retry the connection attempt, specify a valid database name.

DB250005E The command failed because the command is applicable only in a spawned console window.

Explanation: A specific command, meant only for use in the spawned console window, was issued in the current console window.

User response: Use only supported Command Line Processor Plus (CLPPlus) commands in the current console window.

DB250006E Command Line Processor Plus (CLPPlus) failed to start.

Explanation: The default invocation of CLPPlus on Linux and Unix requires XServer running on your desktop.

User response: Verify that the DISPLAY environment variable is set properly. Or, you can start a command line version of CLPPlus using the CLPPLUS command with the *-nw* option.

DB250007I CLPPlus canceled processing the current SQL statement due to a user initiated interrupt.

Explanation: The current SQL statement being processed by CLPPlus has been canceled. A user interrupt initiated the cancel request.

User response: If this was intentional then no response is required. If this was unintentional, reissue the canceled SQL statement.

DB250008E The CLPPlus command failed because the shell or environment variable with name *variable-name* could not be found.

Explanation: A CLPPlus command can include embedded shell or environment variables. If the

variable embedded in your CLPPlus command is not defined in the operating system the command fails.

User response: You can:

- Remove the reference to the variable and issue the command.
- Define the variable and issue the command.

DB250009E **The IMPORT CLPPlus command failed because the SKIPCOUNT value *n* is greater than the total number of rows in the import file.**

Explanation: The IMPORT CLPPlus command was run with the SKIPCOUNT option. The value specified for SKIPCOUNT exceeds the number of rows in the import file read. No rows were imported.

User response: You can:

- reset the SKIPCOUNT to a value less than the number of rows in the import file by using the SET CLPPlus command and retry IMPORT.
- remove the SKIPCOUNT option and retry IMPORT.

DB250010E **The explain facility failed due to an error on the database server.**

Explanation: When AUTOTRACE is set to ON in CLPPlus, and a select statement is run on an Informix database server, CLPPlus implicitly tries to set explain to on. SET EXPLAIN ON failed on the Informix database server and the statement failed.

User response: Check with your Informix database administrator for the reason why SET EXPLAIN ON failed. Retry after resolving the error.

DB250011E **Explain failed to generate data for the issued statement.**

Explanation: When AUTOTRACE is set to ON in CLPPlus, and an SQL statement is run on an Informix database server, the explain plan is only generated for SELECT statements. A statement other than a SELECT statement was issued.

User response: Retry with valid SELECT statement.

DB250012E **CLPPlus failed to process explain data.**

Explanation: When AUTOTRACE is set to ON in CLPPlus, and an SQL statement is run on an Informix database server, CLPPlus is unable to parse the explain plan data generated in XML format.

User response: None

DB250013E **A CLPPlus connection failed because security mechanism *security-mechanism* specified for DSN alias *dsn-alias* in the db2dsdriver.cfg file is not supported by CLPPlus.**

Explanation: The security mechanism specified for a DSN alias in the db2dsdriver.cfg file is not supported in CLPPlus. Any connection attempt using this DSN alias fails and is terminated.

User response: You can:

- specify a valid CLPPlus security mechanism in the db2dsdriver.cfg file for the applicable DSN alias and retry.
- remove the security mechanism parameter for the applicable DSN Alias from the db2dsdriver.cfg file and retry.

DB250014I **DSN alias *dsn-alias* is not found in the configuration file named *config-file*.**

Explanation: A DSN alias specified during a connect attempt is not found in the dsdriver configuration file read by CLPPlus. CLPPlus continues to search the LDAP Directory server configured by the user for that DSN alias.

User response: None required. CLPPlus continues to search in the LDAP directory server for an entry named *dsn-alias*.

DB250015I **CLPPlus successfully established a connection with LDAP directory server *LDAP-server*.**

Explanation: Connection information for an LDAP directory server was found in the dsdriver configuration file and a successful connection was established with that LDAP server during a CLPPlus connect attempt.

User response: None

DB250016E **DSN alias *dsn_alias* was not found in LDAP directory server *LDAP-server*. *dsn_alias* is used as the database name in the subsequent interactive CLPPlus connect attempt.**

Explanation: The DSN Alias name specified during a CLPPlus connect attempt is not found in LDAP directory server. The DSN alias name is treated as database name in the subsequent interactive connect attempt. All other required values for the connection must be specified interactively.

User response: You can:

- continue with the interactive connect attempt.
- use an existing DSN alias name and retry the connect attempt.

- create the DSN alias that was not found and retry the connect attempt.

DB250017E CLPPlus failed to establish a connection with the following LDAP directory server: *LDAP-server*. The following alias was used as the database name in an interactive connect attempt: *dsn-alias*.

Explanation: CLPPlus failed to establish a connection with an LDAP directory server. The DSN Alias name used during the initial connect attempt will be treated as the database name in a subsequent interactive connect attempt. All other required values for the connection are specified interactively.

User response: Respond to this message in one of the following ways:

- Continue with interactive connect.
- Verify and correct the LDAP directory server configuration in the db2dsdriver.cfg file, and retry the connection.
- Check with the LDAP administration to ensure the LDAP directory server is running and retry the connection.

DB250018E A connection attempt during the CLPPlus environment start sequence failed. The CLPPlus environment failed to start.

Explanation: During the CLPPlus environment start sequence, an option to connect to a database is available. If the database connection attempt fails, the CLPPlus start sequence stops and exits.

The reason for the database connection failure is captured by an accompanying JCC error.

User response: Check the accompanying JCC error code and rerun the CLPPLUS command with valid database connection string parameters.

DB250100E The command *text* is not a valid CLPPlus command.

Explanation: The text entered is not a valid Command Line Processor Plus (CLPPlus) command.

User response: Specify a valid command. Refer to the documentation if required for a list of valid commands.

DB250101E The command syntax is invalid. An unexpected token *token* was found following *text*. Expected values include: *token-list*.

Explanation: A syntax error in the command string was detected at the specified token following the text. The text field indicates the 20 characters of the input command string that preceded the token that is not valid.

As an aid, when possible, a partial list of valid tokens is provided in the *token-list*. This list assumes the statement is correct to that point. If the valid values would be unknown user inputs, *token-list* might be blank.

The command cannot be processed.

User response: Refer to the documentation of the command to identify the valid command syntax.

Examine and correct the command in the area of the specified token. Issue the command again.

DB250102E The command syntax is invalid. The value of command parameter *parameter-name* of type *type* is invalid. Valid values include: *token-list*.

Explanation: The command syntax is invalid. A command parameter *parameter-name* is of an invalid data type.

The command cannot be processed.

User response: Examine and correct the parameter value the area of the specified token. Issue the command again.

DB250103E The command syntax is invalid. The value of command parameter *parameter-name* is out of range. The valid range is *lower-bound* to *upper-bound*, inclusive.

Explanation: This message is returned when a value for a parameter is specified that is out of range for that parameter. The token *parameter-name* indicates the parameter for which an out of range value was specified.

User response: Issue the command again, specifying a value for the named parameter that is in the given range.

DB250104E The command is not supported for the *database-server* data server.

Explanation: The command is not supported for specified data server. The command failed.

The command cannot be processed.

User response: Refer to the documentation for the supported commands for the current data server.

DB250105W The CLPPlus command SET LOCALE failed because the locale *locale* is not valid or not supported in the CLPPlus environment.

Explanation: The CLPPlus environment, by default, sets the locale with the value set in the operating system. You can also set or change the locale used by

the CLPPlus environment with the SET LOCALE CLPPlus command.

The SET LOCALE command specified a locale that is not valid or not supported in the CLPPlus environment. The SET LOCALE command failed and the CLPPlus environment continues to use the last valid locale setting.

User response: Specify a valid and supported locale and run the command again.

DB250200E **The feature *feature-name* can not be enabled because the feature does not exist in the current installation.**

Explanation: The feature name *feature-name* is not valid. For valid feature names, search the DB2 documentation.

User response: Specify a valid feature name, then try the request again.

DB250201E **No database connection exists, but a database connection is required.**

Explanation: A database connection is required to issue the command. There is no database connection.

User response: Establish a database connection using the clplus command and try again.

DB250202E **Connections to databases of type *database-type* are not supported by this feature.**

Explanation: The database type is not valid. For valid database types, refer to the documentation.

User response: Specify the name of a database of a valid database type, then try the request again.

DB250203E **Connection failed to be established.**

Explanation: An attempt to establish a database connection failed. No database connection exists.

Possible reasons include:

- Invalid user-name
- Invalid password
- Invalid hostname
- Invalid database name
- Invalid port number

User response: If explicitly connecting to a database, verify and correct the command syntax and try again. Otherwise try again and if the connection still cannot be established, contact your database administrator.

DB250204E **An attempt to locate a file *filename* failed. The command cannot be processed.**

Explanation: A required file *filename* could not be located. Possible reasons include:

- Invalid file name
- Invalid file path
- File not in the file path
- Permissions on file or path do not allow reading of files

User response:

- Verify that the file name is correct and that the file exists in the specified location.
- Verify that the permissions on the file path and file allow read access of files.

DB250205E **Attempt to read file *file-name* failed.**

Explanation: An attempt to read the file parameter *file-name* failed. Possible reasons include:

- Invalid permissions on the file system.

The command cannot be processed.

User response: If the file is a user-defined file, verify that the file permissions allow the file to be read.

DB250206E ***env-variable-name* is not a defined environment variable.**

Explanation: The name *env-variable-name* is not a defined Command Line Processor Plus (CLPPlus) environment variable name or the variable is not set to a required value.

User response: Verify and correct the command syntax, specify a valid environment variable name, and reissue the command or verify and set, if needed, the following CLPPlus environment variables set to required values:

On Windows operating systems:

Variable: CLPPLUS_HELP
Value: sqllib/bin/clplus.bat

On UNIX operating systems:

Variable: CLPPLUS_HELP
Value: sqllib/bin/clplus.sh

DB250207E **Attempt to create file *filename* failed.**

Explanation: The required file *filename* could not be created. Possible reasons include:

- File already exists
- Invalid filename
- Invalid file path
- Invalid permissions on the file system

- Insufficient memory in the file system

User response: If the file is a user-defined file, verify and correct the filename or filepath, verify that the path is valid, verify that files can be written to the path, and that there is sufficient memory. If the file is a system-generated file, re-issue the command - if the command still fails, contact your database administrator.

DB250208E Invalid command sequence. A prerequisite command must be issued before this command. A possible prerequisite command is *prereq-command*.

Explanation: For some CLPPlus commands, there is a prerequisite that one or more other CLPPlus commands be issued first. If the prerequisite commands are not issued, the current command cannot be issued.

The command has not been processed.

User response: Issue the recommended prerequisite command, then reissue the current command. For more information about the command, see the CLPPlus documentation.

DB250209I A connection has been established. Database type: *product-name*. Database version: *version*. Host name: *hostname*. Port: *port-number*. Database name: *database-name*. User: *auth-id*.

Explanation: A database connection has been established to a database of type: *product-name* of version: *version*. The host name of the database server is *hostname*. The port number is *port-number*. The name of the database is *database-name*. The authorization ID of the user is *auth-id*.

Database commands can now be issued for this database.

Queries and database operations can now be executed for this database.

Command Line Processor Plus (CLPPlus) commands that require a connection can be executed until the connection is closed or reset.

User response: Proceed with using CLPPlus. Issue CLPPlus commands, database queries and operations.

DB250210I The connection to database *database-name* has been closed.

Explanation: The connection to database *database-name* has been closed. No database communications can be made at this time.

User response: To execute database operations or to issue commands that require a database connection, re-establish a database connection by issuing the CLPPLUS command.

DB250211E The database object named *object-name* was not found in the database.

Explanation: You can list packages or tables that were created by the current user in the database by issuing the LIST PACKAGES command or the LIST TABLES command.

You can retrieve meta-data information about database objects such as tables, views, functions, and others by issuing the DESCRIBE command, and specifying the database object by name.

This message is returned when either the LIST command or DESCRIBE command is issued for a specified database object, and that object cannot be found in the database.

If this message was returned when the LIST command was issued, it is possible that the database object being searched for was created by a user other than the current user. By default, the LIST command lists only those database objects that were created by the current user.

User response: Verify that the database object exists, and was created by the current user.

DB250212E The database object named *object-name* with type *type* was not found in the database.

Explanation: You can retrieve meta-data information about database objects such as tables, views, functions, and others by issuing the DESCRIBE command, and specifying the database object by name.

This message is returned when the DESCRIBE command is issued for a specified database object, and that object cannot be found in the database.

User response: Verify that the database object exists, and was created by the current user.

DB250300E An unexpected SQL error occurred while disconnecting from a database.

Explanation: An SQL error occurred while committing data or closing a database connection for which no further information is available.

User response: If required, re-establish a connection to the database before performing further database operations.

DB250301E An unexpected SQL error occurred while processing a Command Line Processor Plus command.

Explanation: An unexpected SQL related error occurred on the data server which has interrupted the processing of the Command Line Processor Plus command. The command execution has not completed.

User response: Reissue the command.

DB250302E An asynchronous notification *notification-name* received from the data server process with process ID (PID) *process-id*.

Explanation: An asynchronous notification from the data server occurred with name *notification-name* and process ID *process-id* might require your immediate attention.

User response: Using the process ID value, look up the process and related notification. Respond to any urgent requirements if required. Command Line Processor Plus is still active.

DB250303E The Command line processor plus line buffer size is less than the size of column *column-name*.

Explanation: Column values cannot be displayed, because the Command line processor plus line buffer size is less than the size of column *column-name*. The column values would be truncated.

User response: Specify a line buffer width that is greater than the column width and then issue the command or execute the SQL buffer again.

DB250400E The buffer line number is invalid. Valid values can be between *min-value* and *max-value*.

Explanation: The command attempts to operate on a line in the buffer and the buffer line number provided is invalid. The line number could be wrong because there are no lines in the buffer or because the line number is out of range for the buffer. It is either less than "min-value" or greater than "max-value".

User response: Re-submit the command with a valid line number.

DB250401E Spooling cannot be turned off because it is already off.

Explanation: The SPOOLING OFF command has already been processed. The command cannot be processed.

User response: Because the command is already processed, no further action is required.

DB250402E The command or statement terminator value is invalid.

Explanation: The SQL terminator must be exactly one character long.

User response: Set the terminator character to a single character value.

DB250403E The input string *string* specified in the CHANGE command was not found in the SQL buffer.

Explanation: To successfully issue the CHANGE command the input string must be found in the SQL buffer and the input string *string* was not found.

User response: Re-submit the command with valid argument(s).

DB250405E The DESCRIBE command does not support objects of type *object-type*. Valid types include *token-list*.

Explanation: The DESCRIBE command can be used to access the definitions of a sub-set of database object types in a database. The object type *object-type* is not supported. Valid objects include those in *token-list*.

User response: Refer to the DESCRIBE command documentation for a list of valid database object types. Specify the name of a supported database object and try again.

DB250406E The database object *name* of type *object* does not contain columns.

Explanation: The issued command requires the specification of a database object that contains columns. Valid objects include:

- tables
- views
- catalog tables
- catalog views

The specified object did not contain any columns.

User response: Specify the name of a database object that contains columns.

DB250407E The separator character *character* is invalid.

Explanation: The character *character* specified as a separator character is invalid. Its use is not compatible with the use of the quotation mark character within the command syntax.

User response: Do the following:

1. Specify a different separator character or terminator character.
2. Verify and correct the command syntax.
3. Reissue the command.

DB250408E No user substitution variables are defined.

Explanation: A user substitution variable was referenced, but no user substitution variables are defined.

The command cannot be processed.

User response: Define a user substitution variable and try again.

DB250409E The command cannot be executed when the SQL buffer is empty.

Explanation: The command cannot be issued when there is no data in the SQL buffer.

User response: Add statements to the SQL buffer and re-submit the command.

DB250410I The Command Line Processor Plus is currently paused with the pause value set to *pause-text*.

Explanation: The PAUSE command was previously issued with the pause text *pause-text*. To exit the paused state initiated by the PAUSE command, the enter key must be pressed.

User response: Press enter to exit the paused state.

DB250411E The data type, *type*, of the column named *column-name* is not supported for this computation function.

Explanation: The data type of the named column is not supported for the computation function specified in the COMPUTE command.

Valid numeric data type include the following:

- SMALLINT
- LONG
- DOUBLE
- FLOAT
- DECIMAL
- INTEGER

User response: Specify a column of a valid data type as an argument to the function.

DB250412E No defined computations are available to be displayed.

Explanation: No computations have been defined using the COMPUTE command.

User response: If you want to calculate and display compute information, issue the COMPUTE command at least once and then issue the SHOW COMPUTE command.

DB250413E A compute with label *label* is already defined.

Explanation: A compute with the name *label* is already defined for the column.

User response: Define the computation using a

different name or modify the computation so that it specifies a different column.

DB250414I There are no columns defined.

Explanation: You can list the columns that have been defined by issuing the SHOW COLUMN command or by issuing the COLUMN command with no parameters.

This message is returned when the SHOW COLUMN command is issued, or when the COLUMN command is issued with no parameters, and there are no columns to display.

User response: You do not need to respond to this message.

DB250415W No breakpoints were cleared, because there are no breakpoints defined for the specified column.

Explanation: You can clear any breakpoints that have been defined for a column by issuing the CLEAR BREAK command, specifying the column name.

This message is returned when the CLEAR BREAK command has been issued for a specified column, and no breakpoints were defined for that column.

User response: You do not need to respond to this message.

DB250416W There were no compute actions defined for the column named *column-name* to clear.

Explanation: You can clear all the compute actions that have been defined for a given column by issuing the CLEAR COMPUTES command, specifying the column from which to clear the compute actions.

This message is returned when no compute actions have been defined for the specified column.

User response: You do not need to respond to this message.

DB250417E The following function, specified in the COMPUTE command, is not a supported function for a compute action: *function*.

Explanation: You can use the COMPUTE command to define a compute action to be performed for a given column when a breakpoint for that column is reached.

The syntax of the COMPUTE command is:

```
COMPUTE <function-1> LABEL <label-1>
      OF <column-1>
      [{, <function-2> LABEL <label-2>
        OF <column-2>} ...]
      ON <column-n>
```

where <function> can be one of:

[SUM | MAX | MIN | AVG | COUNT | NUM]

This message is returned when an invalid value for <function> is specified.

User response: Issue the COMPUTE command again, specifying a valid function.

DB250418W The specified compute action was previously defined.

Explanation: The COMPUTE command is used with the BREAK command. You can define a breakpoint for a specified column by issuing the BREAK command. You can use the COMPUTE command to define a compute action to apply to the specified column when a breakpoint for that column is reached during the execution of a SQL statement. A compute action is a function such as SUM, COUNT, or MAX that is applied to all elements of the column.

This message is returned when the COMPUTE command is issued to define a compute action for a column that is identical to a compute action that was previously defined for that column.

User response: You do not need to respond to this message.

DB250419W The SQL buffer is empty.

Explanation: As you issue CLPPlus commands, such as COMPUTE, and BREAK, the resulting objects and actions are stored in the SQL buffer.

Here is a list of just some of the actions that you can perform on a SQL buffer:

- You can edit the contents of the SQL buffer by issuing the EDIT command.
- You can append to the contents of the SQL buffer by issuing the APPEND command.
- You can clear the contents of the SQL buffer by issuing the CLEAR command.

When a SQL statement is executed, the contents of that SQL buffer affect the execution. For example, breakpoints and compute actions have been defined for a column, they will take effect when the SQL statement is run.

This message is returned when an attempt is made to act on the SQL buffer, but the SQL buffer is empty.

User response: You do not need to respond to this message.

DB250420E Creation of a new bind variable failed because a bind variable already exists with the specified name.

Explanation: You can create bind variables to contain values, and use those bind variables in SQL statements.

The scope of bind variables is the database session. After you connect to a database, you can create and use bind variables for the duration of that database session. When you disconnect from the database, those bind variables cease to exist.

You can create bind variables using the following syntax:

VARIABLE <bind-variable-name> <data-type>

This message is returned when an attempt is made to create a new bind variable, specifying the same name for that new variable as the name of a bind variable that already exists.

User response: To create a new bind variable, issue the VARIABLE command again, specifying a name that is not the same as the name of any existing bind variable.

DB250421I There are no breakpoints defined.

Explanation: You can list breakpoints that have been defined by issuing the BREAK command with no parameters.

This message is returned when the BREAK command is issued with no parameters, and there are no breakpoints to display.

User response: You do not need to respond to this message.

DB250422W No columns were cleared, because there are no columns defined.

Explanation: You can clear any columns that have been defined by issuing the CLEAR COLUMN command.

This message is returned when the CLEAR COLUMN command has been issued, and no columns were defined.

User response: You do not need to respond to this message.

Chapter 68. DB250500 - DB250999

DB250500W An error occurred during the initialization of the CLPPlus terminal. However, processing can continue.

Explanation: This message is returned when an internal error occurs while the CLPPlus terminal is being initialized. The initialization did not fail, and you can proceed with CLPPlus commands.

User response: You do not have to respond to this message.

DB250501E Column with name *column-name* is not defined.

Explanation: The column *column-name* specified is not defined.

User response: Define the column and retry the command, or retry the command with a defined column.

DB250502E The name *name* is too long. The maximum length is *maximum-length*.

Explanation: The length of the name specified is too long.

User response: Retry the command with a name shorter than the maximum length specified.

DB250503E No format and display attributes found for column *column-name*.

Explanation: The specified column does not have any format or display attributes set.

User response: Using the COLUMN command, define any required format and display attributes for the column and retry the command.

DB250504E The ACCEPT command failed to accept the value for the substitution variable *substitution-variable* after three consecutive input value format errors.

Explanation: The value of the substitution variable does not conform to the format specified in the ACCEPT command. Users are prompted to retry the command with a suitable substitution variable until three consecutive attempts have failed. The ACCEPT command has failed.

User response: Ensure the substitution variable has a value that conforms to the format specified in the ACCEPT command and retry the ACCEPT command.

DB250505E Error fetching server output messages.

Explanation: Server side procedures can have output messages. When you ask to print these messages on the client console, using the 'set serveroutput on' command, CLPPlus attempts to retrieve the server output messages from the server after the procedure has run.

An error has occurred when retrieving these server output messages.

User response: Check the server-side db2diag.log for any associated errors.

DB250506E A column with alias name *column-alias* is already defined.

Explanation: An alias name *column-alias* already exists for the column.

User response

Retry the command with a unique alias name.

DB250507I No errors found.

Explanation: The Command Line Processor Plus command (CLPPlus) SHOW ERRORS returned successfully and found no errors.

User response: No user response is required.

DB250508I No errors found for *procedure-name*.

Explanation: The Command Line Processor Plus command (CLPPlus) SHOW ERROR returned no errors for the procedure *procedure-name*.

User response: No user response is required.

DB250601E The SQL statement could not be executed because the following bind variable is used in the SQL statement, but does not exist: *bind-variable-name*.

Explanation: You can create bind variables to contain values, and use those bind variables in SQL statements.

The scope of bind variables is the database session. After you connect to a database, you can create and use bind variables for the duration of that database session. When you disconnect from the database, those bind variables cease to exist.

You can create bind variables using the following syntax:

VARIABLE <bind-variable-name> <data-type>

DB250601E

You can use bind variables in SQL statements like this:

```
DELETE employee WHERE empno = :id;
```

where ":id" is a bind variable.

This message is returned when an attempt is made to execute a SQL statement that contains a bind variable that has not been defined.

User response: Respond to this message in one of the following ways:

- Define the variable that was used in the SQL statement, and then attempt to execute the SQL statement again.
- Attempt to execute the SQL statement again, specifying a bind variable that is defined.

Chapter 69. DB255000 - DB255499

DB255000I **The RDF utilities were unable to sort the RDF triples.**

Explanation:

- You need to have access to the SORT command on Linux and UNIX operating systems.
- You need to have the Cygwin program installed on the Windows operating system.

User response: Ensure you have the proper permission or application installed for your operating system. Then attempt the operation again.

DB255001E **The SPARQL query failed because an unexpected database error occurred while the query was being processed. Query: *query-details*. SQLCODE: *sqlcode*. SQLSTATE: *sqlstate*.**

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as `createRdfStore`. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned when an attempt is made to query RDF data that is stored in a DB2 database, but an error was encountered while the query was being processed.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the `db2diag` diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the `db2diag` log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.
4. Execute the SPARQL query again.

DB255002E **Creation of an RDF store failed because an unexpected database error occurred while the underlying database objects were being created. SQLCODE: *sqlcode*. SQLSTATE: *sqlstate*.**

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as

`createrrdfstore`. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned when an attempt is made to create a new RDF data store, but an error was encountered while the store was being created.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the `db2diag` diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the `db2diag` log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.
4. Run the RDF command again to create the new RDF store.

DB255003E **Dropping an RDF store failed because an unexpected database error occurred while the underlying database objects were being removed. SQLCODE: *sqlcode*. SQLSTATE: *sqlstate*.**

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as `createrrdfstore`. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned when an attempt is made to drop an existing RDF store from a DB2 database, but an error was encountered while the underlying tables and other database objects were being dropped and cleaned up.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the `db2diag` diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the `db2diag` log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.

3. Take corrective actions to fix the error.
4. Call the `droprdfstore` command to drop the RDF store.

DB255004E **Updating the statistics for an RDF store failed because an unexpected database error occurred during the operation.**
SQLCODE: *sqlcode*. **SQLSTATE:** *sqlstate*.

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as `createrrdfstore`. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned when an attempt is made to update the statistics for an RDF store in a DB2 database, but an error was encountered while the statistics for the underlying database objects were being updated.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the `db2diag` diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the `db2diag` log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.
4. Run the statistics updates RDF operation again.

DB255005E **Reorganizing an RDF store failed because an unexpected database error occurred during the operation.**
SQLCODE: *sqlcode*. **SQLSTATE:** *sqlstate*.
Reason code: *reason-code*.

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as `createrrdfstore`. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

To create tables optimized for an RDF store, you can reorganize the RDF store. Reorganizing an RDF store involves three steps, and this message can be returned when the reorganization operation fails in one of these three steps. The runtime token *reason-code* indicates which steps the reorganization operation failed:

1

Identifying whether the RDF store needs to be reorganized using the `reorgcheckrdfstore` command.

2

Creating tables and scripts to prepare to reorganize the RDF data using the `reorgrdfstore` command.

3

Transferring the RDF data into the new, optimized tables using the `reorgcomplete` command.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the `db2diag` diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the `db2diag` log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.
4. Execute the reorganization operation again.

DB255006E **Updating data in an RDF store failed because an unexpected database error occurred during the operation.**
SQLCODE: *sqlcode*. **SQLSTATE:** *sqlstate*.

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as `createrrdfstore`. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned when an attempt is made to update data in an RDF store in a DB2 database, but an error was encountered while the operation was being processed.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the `db2diag` diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the `db2diag` log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.
4. Execute the RDF command again.

DB255007E Creation of an RDF store failed because the current user id does not have the required authority or privileges to create tables-or-external-functions in the specified schema. Schema name: *schema-name*.

Explanation: The following privileges are required to create an RDF store:

- CREATETAB permission to create tables in a schema.
- CREATE_EXTERNAL_ROUTINE permission to create external functions.

User response: Assign the proper privilege to the user. Then issue the command again.

DB255008E Creation of an RDF store failed because the page size of the table spaces in the specified schema is too small.

Explanation: An RDF store requires a table space with a page size of 32 KB to successfully create all the columns.

User response: Create a table-space with 32 KB page size. Then use it to create the RDF store.

DB255009E Creation of an RDF store failed because an RDF store with the specified name already exists in the specified schema. Specified RDF store name: *store-name*. Schema name: *schema-name*.

Explanation: There is a naming conflict for the RDF store you are attempting to create.

User response: Use a different name to create the RDF store.

DB255010E The RDF operation failed because the current user id does not have SELECT privileges on a required RDF store table. RDF store name: *store-name*. Table name: *table-name*.

Explanation: The necessary read permission for the RDF store metadata table is missing.

User response: Ensure user has CONNECT and READ authority for the RDF store.

DB255011E Creation of an RDF store failed. The RDF utility was unable to create a required table for the new RDF store because a table with the same name already exists in the schema. Table name: *table-name*. Schema name: *schema-name*.

Explanation: One or more tables failed to create successfully due to a naming conflict.

User response:

1. Rename existing tables which are causing the conflict.
2. Edit objectnames file and provide a different physical name for the tables which are causing the conflict.

DB255012E The RDF command failed because the specified RDF store does not exist in the specified schema. Specified RDF store name: *store-name*. Specified schema: *schema-name*.

Explanation: The RDF store you are attempting to connect to does not exist.

User response:

1. Check the name of the RDF store that you are connecting to and issue the command again.
2. Create the RDF store and then connect to it.

DB255013E The droprdfstore command failed because the current user id does not have required authority or privileges to drop a table in the RDF store. Table name: *table-name*. Schema name: *schema-name*.

Explanation: The user does not have the proper authority to drop the RDF store table.

User response: Ensure the user has the DROP TABLE privilege for this database.

Then issue the command again.

DB255014E The RDF operation failed because the current user id does not have SELECT privileges on the data tables for the RDF store. RDF store name: *store-name*. Schema name: *schema-name*.

Explanation: The user does not have the proper authority to read the data tables accessed using RDF commands or APIs.

User response: Ensure the user has been granted READ permission for these tables.

Then issue the command or API again.

DB255015E The RDF command failed because a mandatory parameter was not specified. Parameter name: *parameter-name*. Command name: *command-name*.

Explanation: This is a mandatory parameter for this command line tool.

User response: Check your command to ensure it contains the mandatory parameter.

Then issue the command again.

DB255016E The RDF command failed because the specified RDF store name is too long.
Specified RDF store name: *store-name*.

Explanation: The maximum length of a store name is 128.

User response: Ensure the store name is less than the maximum number of characters. Then issue your command again.

DB255017E Gathering statistics for an RDF store failed because the current user id does not have the required authority or privileges on a table in the RDF store.
Table name: *table-name*. Schema name: *schema-name*.

Explanation: The user does not have the proper authority to gather the latest statistics.

User response:

- Ensure the user has READ privilege on all tables for this database.
- Ensure the user has UPDATE privilege on all the statistics tables.

Then issue the command again.

DB255018E Creation of a new RDF graph failed because the specified graph name, "DEF", is invalid.

Explanation: DEF is reserved for IBM use. It can not be used as the name of any user created graph.

User response: Rename the graph. Avoid using DEF as the name.

DB255019E The RDF operation failed because the specified RDF quad contains a URI that is too long. Graph: *graph*. Subject: *subject*. Predicate: *predicate*. Object: *object*.

Explanation: An URI in the quad has exceeded the maximum length of 2048 characters.

User response: Examine the quad to reduce the length of the offending URI to less than 2048 characters.

DB255020E Parsing of a SPARQL query failed.
Specified SPARQL query: *query*. Line number of the problem in the query: *line-number*. Column number of the problem in the query: *column-number*.

Explanation: Check your SPARQL query at the location specified.

User response: Rectify the error and issue the query again.

DB255021E The SPARQL feature *feature-name* is not supported.

Explanation: IBM APIs do not support the feature.

User response: Edit your query to refrain from using the unsupported feature.

DB255022E The RDF command failed because the following parameter is invalid:
parameter-name.

Explanation:

1. Parameters can not be abbreviated.
2. Parameters can be issued in lower case only.

User response:

1. Check the spelling of the parameter.
2. Ensure the parameter is in lower case.

Issue the command again.

DB255023E The `createrrdfstoreandloader` command failed because the RDF utility could not create the file specified with the `-storeloadfile` parameter or the utility could not create the file specified with the `-storeschemafile` parameter.

Explanation: The user specified directory can not be created on the file system.

User response: Ensure the directory structure already exists on the file system. Then Issue the command again.

DB255024E The RDF command failed because the current user id does not have INSERT and UPDATE privileges on the data tables for the RDF store. RDF store name: *store-name*. Schema name: *schema-name*.

Explanation: The user does not have the proper authority to write to the data tables accessed using RDF commands or APIs.

User response: Ensure the user has been granted WRITE permission for these tables.

Then issue the command or API again.

DB255025E The RDF command failed because the specified RDF store name does not satisfy DB2 database table naming conventions. Specified RDF store name: *store-name*.

Explanation: The name must conform to naming conventions for table objects. Specifically, the name must not contain any space characters.

User response: Ensure the store name does not contain any invalid characters. Then issue your command again.

DB255026E The RDF command failed because the specified schema name does not satisfy DB2 database schema naming conventions. Specified schema name: *schema-name*.

Explanation: The *schemaName* does not follow the rules of DB2 schema names.

User response: Ensure the *schemaName* rules are not being violated. Then issue your command again.

DB255027E The RDF operation failed because the specified RDF quad contains a literal string that is too long. Graph: *graph*. Subject: *subject*. Predicate: *predicate*. Object: *object*. Number of characters in the specified string literal: *num-chars*.

Explanation: A literal in the quad has exceeded the maximum character length as configured for the store limit.

User response:

- Examine the quad to reduce the length of the offending item to less than the configured length for the store.
- Increase the configured length of literals for the store.

Then issue your query again.

DB255028E The SPARQL query was not processed because the query form does not match the method used to execute the query. Expected query form: *query-form*.

Explanation: SPARQL supports four types of queries:

- Describe
- Ask
- Select
- Construct

The query type does not match the user operation. For example, *execDescribe*, *execSelect*, *execConstruct*, etc.

User response: Ensure the query matches one of the supported query types.

DB255029E The RDF command failed because not all required table names were specified. RDF command name: *command-name*. Names of tables that were not specified: *table-name-list*.

Explanation: If you choose to provide table names when issuing either the *createrdfstore* or

createrdfstoreandloader commands, you need to specify names for all the tables.

User response: Provide all the table names. Then issue your command again.

DB255030E An unexpected database error occurred during the operation. Operation name: *operation-name*. Query: *query-details*. SQLCODE: *sqlcode*. SQLSTATE: *sqlstate*.

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as *createRdfStore*. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned in two scenarios:

- A DB2 database RDF utility command encounter a database error while performing an RDF utility operation.
- An attempt is made to query RDF data that is stored in a DB2 database, but an error was encountered while the query was being processed.

When this message is returned because a query fails, the details of the query will appear in the *query-details* runtime token. When this message is returned because an RDF utility command fails, the *query-details* runtime token will contain the following place holder: *"*"*.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the *db2diag* diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the *db2diag* log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.
4. Run the RDF operation again.

DB255031E The RDF operation failed because the RDF utilities could not find the Cygwin library.

Explanation: The Cygwin application needs to be on the path if you are running your program on the Windows operating system.

User response: Ensure the Cygwin application has been added to the path on Windows.

DB255032E The reorgswitchrdfstore command failed because reorganized tables for the RDF store are unavailable.

Explanation: The reorgswitchrdfstore command transfers the RDF store to the new tables. No new reorganized tables are currently available for this store.

User response: Issue the reorgrdfstore command first. After it completes successfully, then issue the reorgswitchrdfstore command.

DB255033E The following RDF command failed because the command is being issued remotely. Command name: *command-name*.

Explanation: This command can only be run locally on the DB2 database Server.

User response: Command *command-name* has to be run locally on the DB2 database server.

DB255034I RDF data reorganization has completed.

Explanation: Table data has been successfully reorganized into the new tables.

User response: Begin using the reorganized tables by issuing the reorgswitchrdfstore command.

DB255035E Connecting to an RDF store failed because an unexpected database error occurred during the operation. **SQLCODE:** *sqlcode*. **SQLSTATE:** *sqlstate*.

Explanation: You can take advantage of the Resource Description Framework (RDF) with DB2 databases using the DB2 database RDF utility commands such as createrdfstore. You can query RDF data that is stored in DB2 databases using the SPARQL query language.

This message is returned when an attempt is made to connect to an RDF store in a DB2 database, but an error was encountered while the connection was being processed.

To assist with troubleshooting, the underlying database error code is included in the runtime token *sqlcode*. More diagnostic information can be found in the db2diag diagnostic log files. The runtime token *sqlstate* might sometimes be empty because not all SQL errors have an associated *sqlstate*.

User response: Respond to this error by performing the following troubleshooting steps:

1. Search in the db2diag log files for the *sqlcode* and *sqlstate* given in the runtime tokens *sqlcode* and *sqlstate*.
2. Examine the underlying database error message returned by the DB2 database.
3. Take corrective actions to fix the error.

4. Execute the RDF operation again.

DB255036I The reorgrdfstore command completed successfully.

Explanation: You can create tables with optimal number and length of columns for an RDF store using the reorgrdfstore command.

This message is returned when the command completes successfully.

User response: To switch the RDF store to use the newly reorganized tables, use the reorgswitchrdfstore command.

DB255037I None of the RDF store tables need to be reorganized.

Explanation: When significant amounts of data have been inserted into a RDF store, the number and length of columns for various tables in the store might no longer be optimal for your data. This might impact query and insert performance for the RDF store. You can determine whether any or all of the RDF store tables need to be reorganized using the reorgcheckrdfstore command.

This message is returned when the reorgcheckrdfstore utility determines that no RDF store tables need to be reorganized.

User response: No user response is required.

DB255038E The reorgrdfstore command failed because an invalid value was specified for the -table parameter. Specified table name: *specified-table-name*.

Explanation: You can create tables with optimal number and length of columns for an RDF store using the reorgrdfstore command.

You can specify which RDF store table needs to be reorganized using the -table parameter with the reorgrdfstore command. This message is returned when the value specified with the -table parameter is not a valid name of an RDF store table.

User response: Call the reorgrdfstore command again, specifying a valid RDF store table name with the -table parameter.

DB255039E RDF store creation failed because the value of the database LOGFILSIZ parameter is not large enough.

Explanation: The minimum required DB2 LOGFILSIZ parameter value is 20000 or more. If you get this error even with a value of 20000, you need to increase it further.

User response:

1. Update the LOGFILSIZ database configuration parameter to a larger value.
2. Issue the store creation command again.

DB255040E RDF store creation failed because the table space specified for a RDF Store table in the -objectnames parameter does not exist. Table space name: *table-space-name*.

Explanation: The table space specified in the properties file of the -objectnames parameter does not exist.

User response: Create the table space or check the spelling of the table space. Then issue the command again.

DB255041E The queryrdfstore command failed because reopt ONCE or ALWAYS option was specified but the corresponding REOPT package is not bound.

Explanation:

- You need to bind the REOPT NULLIDR1 package if you are using the reopt ONCE option.
- You need to bind the REOPT NULLIDRA package if you are using the reopt ALWAYS option.

For more information, see: <http://www.ibm.com/developerworks/data/library/techarticle/dm-1011reopt/>

User response: Bind the respective NULLIDR1 or NULLIDRA packages. Then issue the queryrdfstore command again.

DB255042E The loadrdfstore command failed because the input load file is not specified as the last argument to the loadrdfstore command.

Explanation: The loadrdfstore command requires an input load file. This load file is used to list the commands to load data into an RDF store.

User response: Ensure the load file is specified. Then issue the loadrdfstore command again.

DB255043E The loadrdfstore command failed because the format of the input load file is not supported.

Explanation: The loadrdfstore command requires an input load file in a supported file format.

The following load file formats are supported:

- nquad
- ntriple
- rdxml
- trig

- turtle
- n3

User response: Ensure the load file is formatted properly. Then issue the loadrdfstore command again.

DB255044E The queryrdfstore command failed because the SPARQL query is missing.

Explanation: The queryrdfstore command accepts the SPARQL query as an input parameter in one of the following two ways:

- as a file using the file parameter
- as an inline string specified as the last argument to the queryrdfstore command

The command is missing the required query string.

User response: Ensure that a SPARQL query is available to the queryrdfstore command as a parameter. Then issue the command again.

Part 11. DBA Messages

This section contains the messages generated by the Database Administration (DBA) tools. The messages are listed in numeric sequence.

Chapter 70. DBA0000 - DBA0499

DBA0000I **The function completed successfully.**

Explanation: This message is for your information only.

User response: No action is required.

DBA0001E **An internal error occurred when allocating memory.**

Explanation: An internal error occurred when allocating memory. The application is terminated.

User response: Try closing some applications to release memory. If the problem persists, contact IBM Service.

DBA0002E **An internal error occurred. Unexpected input was received from the request.**

Explanation: An unexpected input value was received from the request. The application is terminated.

User response: Contact IBM Service.

DBA0003E **An inconsistency in the table definition was detected when processing the alter request.**

Explanation: When processing an alter table request, an inconsistency was detected between the last known table definition and the table definition fetched from the database. The table definition might have changed outside of the context of the administration tool. The table cannot be altered. The request is terminated.

User response: Select the "Refresh" action on the "Tables" pop-up menu to get an updated list of tables from the database. Try altering the table again.

DBA0004E **An error occurred when validating a connection to a database.**

Explanation: An invalid connection was detected. A problem can occur if the server is brought down, the currently held connections are no longer valid, and a new connection cannot be obtained. The request is terminated.

User response: Ensure that the server is started. Select the "Connect" action in the pop-up menu of the database that you want to connect to. Try the action again.

If the action fails again, ensure that your network is operating for the target system and that the database is up and running on that system.

If the network is operational, a system internal error

might have occurred. Contact IBM Service.

DBA0005E **An error occurred when validating an attach to an instance.**

Explanation: An invalid connection was detected. The database might have stopped or the connection might no longer exist. The request is terminated.

User response: Ensure that the server is started. Select the "Connect" action in the pop-up menu of the database that you want to connect to. Try the action again.

If the action fails again, ensure that your network is operating for the target system and that the database is up and running on that system.

If the network is operational, a system internal error might have occurred. Contact IBM Service.

DBA0006E **An error occurred when validating a handle for a persistent database connection or instance attach.**

Explanation: An invalid connection was detected. The database might have stopped or the connection might no longer exist. The request is terminated.

User response: Ensure that the server is started. Select the "Connect" action in the pop-up menu of the database that you want to connect to. Try the action again.

If the action fails again, ensure your network is operating for the target system and that the database is up and running on that system.

If the network is operational, a system internal error might have occurred. Contact IBM Service.

DBA0007E **An internal error occurred when setting the context type for connections.**

Explanation: An error occurred when setting the context type for connections established by the application. The application is terminated.

User response: Contact IBM Service.

DBA0008E **An internal error occurred. An unexpected list type was requested.**

Explanation: The type of object to be listed was not recognized as a valid type by the application. The application is terminated.

User response: Contact IBM Service.

DBA0009E **An internal error occurred. An unexpected request type was received.**

Explanation: The type of request to be performed was not recognized as a valid type by the application. The application is terminated.

User response: Contact IBM Service.

DBA0010E **An internal error occurred. An unexpected object type was detected.**

Explanation: The type of object to be processed was not recognized as a valid type by the application. The application is terminated.

User response: Contact IBM Service.

DBA0011E **There are too many open connections.**

Explanation: The maximum number of open database connections was reached. The request is terminated.

User response: Either:

- Disconnect from other databases and connect to the database you want. Select the "Disconnect" action in the pop-up menu of the databases from which you want to disconnect. Select the "Connect" action in the pop-up menu of the database to which you want to connect.
 - Update the configuration parameter to increase the number of connections allowed.
-

DBA0012E **Unable to allocate a connection handle.**

Explanation: An error occurred when attempting to allocate a connection handle. The application is terminated.

User response: Contact IBM Service.

DBA0013W **The maximum number of objects that can be listed was reached.**

Explanation: The maximum number of objects that can be listed by the application was reached. The maximum number is 20,000.

User response: None

DBA0014E **An internal error occurred. An unexpected drop type was requested.**

Explanation: The type of object to be dropped was not recognized as a valid type by the application. The application is terminated.

User response: Contact IBM Service.

DBA0015E **An internal error occurred. A latch request failed.**

Explanation: A latch request failed. The application is terminated.

User response: Contact IBM Service.

DBA0016E **An internal error occurred. An unlatch request failed.**

Explanation: An unlatch request failed. The application is terminated.

User response: Contact IBM Service.

DBA0017E **No error information is available. See the administration tools log.**

Explanation: No error information can be displayed.

User response: See the administration tools log for error information.

DBA0018E **The Administration Server is not started. Start the Administration Server and try the action again.**

Explanation: The Administration Server is not started.

User response: Start the Administration Server using the DB2ADMIN START command on the target system and try the action again.

DBA0019E **The object *object-name* is in use. Try the action again at a later time.**

Explanation: The object on which the action was requested or a related object is already being used in another action. Modifications to the object or a related object might be in progress.

The requested action cannot be performed at the same time as the current action.

Examples:

- If a table is being altered, a request to rename the same table will not be allowed until the interaction with the Alter window is complete. However, a request to sample the contents of that table is allowed.
- A request to drop an object (such as a table) will not be allowed if any other actions are pending on related objects (such as the database, views, or triggers).

The action is not performed.

User response: Try the action again at a later time when the object is available.

DBA0020E **An error occurred opening the administration tools trace file.**

Explanation: An error occurred opening the administration tools trace file.

User response: Contact IBM Service.

DBA0021E **An error occurred writing to the administration tools trace file.**

Explanation: An error occurred writing to the administration tools trace file.

User response: Contact IBM Service.

DBA0022W **An error occurred closing the administration tools trace file.**

Explanation: An error occurred when attempting to close the administration tools trace file. Processing continues.

User response: Ensure that the administration tools log file was specified correctly.

DBA0023W **An error occurred opening the administration tools log file.**

Explanation: An error occurred when attempting to open the administration tools log file. Processing continues.

User response: Ensure that the administration tools log file was specified correctly.

DBA0024W **An error occurred writing to the administration tools log file.**

Explanation: An error occurred when attempting to write to the administration tools log file. Processing continues.

User response: Ensure that the administration tools log file was specified correctly and can be written to.

DBA0025W **An error occurred closing the administration tools log file.**

Explanation: An error occurred when attempting to close the administration tools log file. Processing continues.

User response: Ensure that the administration tools log file was specified correctly.

DBA0026E **An internal administration tools error occurred.**

Explanation: An unrecoverable error occurred.

User response: Contact IBM Service.

DBA0027E **An internal error occurred. A lock contention in the administration tools was detected.**

Explanation: Lock contention in the administration tools was detected.

User response: Contact IBM Service.

DBA0028I **The current administration tools session will be terminated due to an unrecoverable error.**

Explanation: A severe internal error occurred.

User response: Contact IBM Service.

DBA0029C **The application programming interface program did not complete successfully for object *name* with reason code *name*. Try again or contact your local support.**

Explanation: An application programming interface call did not complete.

User response: Try the operation again or contact your system administrator.

DBA0030E **An error occurred while attempting to connect to database *name*.**

Explanation: The requested operation could not complete because it requires a connection to database *name*. The connection was not successful.

User response: Make sure the database is accessible. Try explicitly connecting to the database using the connect action on the database pop-up menu. A common error on connection is an invalid userid and password. Make sure you have supplied the proper userid and password. Try the operation again or contact your system administrator.

DBA0031E **An error occurred while attempting to attach to instance *name*.**

Explanation: The requested operation could not complete because it requires an attachment to the instance *name*. The attachment was not successful.

User response: Make sure the instance is accessible. Try explicitly connecting to the instance using the connect action on the instance pop-up menu. A common error on connection is an invalid userid and password. Make sure you have supplied the proper userid and password. Try the operation again or contact your system administrator.

DBA0032C An error occurred while attempting to load DLL *name*.

Explanation: Either the dynamic link library *name* was not found or an error occurred while reading the file.

User response: Verify that the DLL is installed and undamaged, and its location is part of the LIBPATH parameter in the CONFIG.SYS file.

DBA0033C *name* cannot be dropped because it is currently in use by the administration tools. Close some windows and retry the command.

Explanation: Either the object *name* you selected to drop, or an item that is owned by the object you selected to drop, is still being used by another administration tool. For example, you might have an ALTER table window open for a database that you are trying to drop. In this case, you cannot drop the database until you close its ALTER table window.

User response: Close all the windows using the object you selected, or close all the windows using items owned by that object, and then retry the command.

DBA0034C *name* cannot be removed because it is currently in use by the administration tools. Close some windows and retry the command.

Explanation: Either the object *name* you selected to remove, or an item that is owned by the object you selected to remove, is still being used by another administration tool. For example, you might have an ALTER table window open for a database that you are trying to remove. In this case, you cannot remove the database until you close its ALTER table window.

User response: Close all the windows using the object you selected, or close all the windows using items owned by that object, and then retry the command.

DBA0035C The operation cannot be performed because the object *name* is currently in use by the administration tools. Close some windows and retry the command.

Explanation: This operation requires *name* to be in a *name* state.

User response: Make sure all other pending operations on the selected object are complete, and then retry the command.

DBA0036I The database configuration was successfully updated. All applications must disconnect from database *name* before the changes come into effect. If a backup has already been taken it is recommended to take another backup of the database in order to pick up the new configuration values.

Explanation: The database configuration was successfully updated, but the active database configuration cannot be changed until all applications are disconnected from the database. Once all applications are disconnected, the changes will come into effect with the first connection to the database.

User response: Ensure that all applications are disconnected from the database, and then reconnect to the database.

DBA0037I The instance configuration was successfully updated. The instance *name* must be stopped before most changes come into effect.

Explanation: The instance configuration was successfully updated.

Changes to the configuration parameter "dftdbpath" take effect immediately.

For the remaining configuration parameters, the changes do not take effect until all applications have disconnected from the database, and the instance is successfully stopped and started again.

User response: To make all changes take effect, ensure that all applications have disconnected from the database, stop the instance, and then start it again.

You can stop the instance by issuing the db2stop command. You can then start the instance by issuing the db2start command.

DBA0039W There are currently database objects locked or database connections in use. Continue shutdown?

Explanation: The database objects in the Control Center are locked when actions, such as Alter Table, are taken against them. Other actions will maintain connections to the database that are required to complete the task.

User response: If you choose to continue with the shutdown, then all locks and connections are forced and the tools will shut down. It can be dangerous to do so, however, because forcing some actions (like Restore Database) halfway through the task can leave the database in a corrupted state.

It is safest to first go back to the DB2 tools and make sure there are no database operations outstanding.

DBA0040W Shutting down will stop the Snapshot monitor. Continue shutdown?

Explanation: Snapshot monitoring has been started on one or more database objects and to continue the monitoring, you must have the DB2 tools running.

User response: Choose to continue the shutdown only if you no longer require the Snapshot monitor.

DBA0041I The requested operation could not be performed because the DB2 Administration Tools is an untrusted applet.

Explanation: Applets loaded over the network are usually considered to be untrusted. Different Web browsers and applet viewers may impose different restrictions on untrusted applets including most local system operations (For example, read files, write files, and print files).

User response: Consult your system administrator and see if these restrictions can be relaxed or customized.

DBA0042I Do you want to clear the setting for the userid and password that was used to connect to this database.

Explanation: If you choose to clear this setting, you will prevent unauthorized access to this database from this workstation. However, you will be prompted to enter the correct userid and password the next time you connect to this database.

User response: Click on Yes to clear the userid and password setting. Click on No to save the userid and password setting until you exit the Control Center.

DBA0043I The command has not been executed.

Explanation: The operation did not occur and has been aborted.

User response: Resubmit the command.

DBA0044N The Java Virtual Machine heap size is not large enough to handle your request.

Explanation: If your request is accessing a large amount of data, this is probably an expected occurrence.

Note that if your request is not accessing a large amount of data, this occurrence may indicate a memory problem.

User response: Increase the size of the Java Virtual Machine heap and restart the application. To increase the Java Virtual Machine heap size include the -Xmx<size> option when invoking the application,

where <size> is the maximum size in bytes. Optionally specify a different unit of measurement by appending the letter k or K to indicate kilobytes, or m or M to indicate megabytes.

For example, to start the Control Center with a Java Virtual Machine heap size of 128 megabytes, enter: db2cc -Xmx128m.

DBA0045N At least one required field in the Register XSR dialog is incomplete.

Explanation: All required fields must be filled in before the XSR object registration process can be completed by the Register XSR dialog.

User response: Verify that all required fields have been completed in the Register XSR dialog and try again.

DBA0046N A file required for XSR object registration cannot be found.

Explanation: When you add XSR objects through the Register XSR dialog, the files you reference must be available during the registration process. At least one file cannot be found.

User response: Verify that all files you reference during the XSR object registration process are available and try again.

DBA0047N The requested native XML data store functionality is not available on this database.

Explanation: This message is expected behavior when you attempt to use the XML functionality supported by a native XML data store administration tool against a database which does not support native XML data store. Note that XML Extender is not part of the native XML data store.

User response: If the DB2 instance does not support native XML data store functionality, no user response is necessary. You cannot use the XML functionality supported by your native XML data store administration tool against this database.

For further information on how to use the supported native XML data store functionality, visit the DB2 Information Center at <http://publib.boulder.ibm.com/infocenter/db2luw/v9index.jsp> and search for "native XML data store" and "overview".

DBA0048N Automatic storage cannot be added to this database.

Explanation: An attempt has been made to access the 'Add Automatic Storage' window for a non-automatic storage database.

User response: To add storage for a non-automatic

storage database, select a table space listed in the Table Spaces folder. Right click on the table space name and select "Alter..." . The Alter Table Space notebook opens. Select the Containers tab to add storage for your database.

DBA0049N **Registration with the XML schema repository (XSR) failed because there was an error reading or processing the selected file.**

Explanation: The Register XSR notebook only supports certain file extensions when you are adding objects to the XSR. If you are using a supported extension, then the error occurred during the processing of the file.

User response: The supported extensions are .dtd (XML DTDs), .mod (XML DTD Modules), .ent (External Entities), and .xsd (XML Schemas). If you are using a supported extension, try validating the file contents using an XML editor. If there are no errors, try the registering the object again. If registration still fails, contact IBM Service for additional help.

DBA0099N **The Java Virtual Machine heap size is not large enough to service your request.**

Explanation: If your request is accessing a large amount of data, this is probably a normal response. However, if your request is not accessing a large amount of data, this may indicate a memory problem.

User response: Modify the maximum size of the Java Virtual Machine heap with the -Xmx<size> option when invoking the program. Specify the maximum size, in bytes. Append the letter k or K to indicate kilobytes, or m or M to indicate megabytes.

For example, db2cc -Xmx128m

If the problem persists, contact IBM Service.

DBA0100I **Control Center is updating node directory information for instance - instance. Incorrect information is being corrected. The current request cannot be performed.**

Explanation: The Control Center detected that the nodetype changed at the instance and is recataloging the instance based on the new information. Depending on the nodetype value for the instance, actions from the Control Center may be different.

User response: If available, try the request again.

DBA0101W **Entries specified in the exception list will not be used. Do you want to continue?**

Explanation: The Nodes Contain Exceptions check box is deselected, but exception containers are specified in the List Exception dialog.

User response: Click on Yes to clear the exception containers list, or click on No or Cancel to reselect the Nodes Contain Exceptions check box.

DBA0102W **Unable to detect nodetype for instance - instance. Reason Code: reason-code.**

Explanation: Nodetype is a database manager configuration parameter that identifies whether the instance consists of a single or multiple database partition servers.

Discovery tries to resolve the nodetype value for you. This distinction between instances is required because actions in the Control Center can be different between single and multiple database partition servers environments.

If the reason code is -1, discovery could not map the catalogued instance to a valid DB2 instance.

All other reason codes map to a valid SQL message. Check the help for the corresponding SQL message.

User response: Discovery has the following requirements:

1. DB2 Administration Server must be set up and running at the remote instance you want to catalog or access.
2. Ensure the instance is listed in the registry. Issue db2set -l from the host you are cataloging.
3. The following global DB2 Registry values must be set:
 - DB2SYSTEM
 - DB2ADMINSERVER
4. The following instance DB2 Registry values must be set:
 - DB2COMM
5. The following administration server configuration parameters must be set:
 - DISCOVER = KNOWN and DISCOVER_COMM = null

or

 - DISCOVER = SEARCH and DISCOVER_COMM = protocol (for example, TCPIP)

Verify the setting of these registry values by entering db2set -all.

DBA0103W Changes made to the configuration parameter values will not be used. Do you want to continue?

Explanation: Another database partition was selected, but the changes to the configuration parameters apply to the previously selected database partition.

User response: Click on Yes to obtain the configuration parameter values for the next database partition and discard the changes that were specified for the previously selected database partition.

DBA0113E Exception containers are specified, but not all nodes in the database partition group are included. No containers are specified for the following nodes: *node-list*.

Explanation: The specified exception containers have no common container. Because no common container is specified, the exception container must be defined at every node in the database partition group.

User response: Click on OK to add containers to the specified nodes.

DBA0114W Control Center detected incorrect node directory information for instance - *instance*. Incorrect information is being corrected. Please exit and restart the Control Center. The current request cannot be performed.

Explanation: The Control Center detected that the nodetype changed at the instance.

User response: Exit the Control Center and restart it.

DBA0115I Node directory information has been corrected for this instance. Instance *instance* should be refreshed.

Explanation: Because the cataloged information for the instance is updated, the displayed objects and actions may not be correct.

User response: Select the "Refresh" action from the instance pop-up menu to update the Control Center, then try the request again.

DBA0116I The operation cannot be performed by the JDBC server at this time. Retry this operation later.

Explanation: The JDBC server is currently busy with an existing operation and cannot perform the user operation.

User response: Wait until the JDBC server completes the existing operation and retry the operation.

DBA0117W Database partition information could not be retrieved. Reason code = *reason-code*, object = *object*.

Explanation: Function which relies on successful retrieval of database partition information will be disabled or removed. This may include menu items as well as actual dialog function.

User response: Make sure that the DAS is started at the server where the object exists. If the DAS is started, then treat the return code as a DAS error and continue troubleshooting.

DBA0200E An attempt to open or read from file *filename* failed. Error code = *error-code*.

Explanation: An attempt was made to either open or read from a file using the Database Administration Server (DAS). The attempt failed.

User response: Verify that the DAS is running, and that the file exists and has read permission for the userid associated with the DAS.

If the problem persists, contact your IBM service representative with the corresponding file name and error code.

DBA0201E An attempt to open or write to file *filename* failed. Error code = *error-code*.

Explanation: An attempt was made to either open or write to a file using the Database Administration Server (DAS). The attempt failed.

User response: Verify that the DAS is running, and that the file system is not full and has write permission for the userid associated with the DAS.

If the problem persists, contact your IBM service representative with the corresponding file name and error code.

DBA0202E An attempt to open or read from file *filename* failed. Error code = *error-code*.

Explanation: An attempt was made to either open or read from a file. The attempt failed.

User response: Verify that the file exists and has appropriate read permission.

If the problem persists, contact your IBM service representative with the corresponding file name and error code.

DBA0203N An attempt to open or write to file *filename* failed. Error code = *error-code*.

Explanation: An attempt was made to either open or write to a file. The attempt failed.

User response: Verify that the file system is not full

and that it has appropriate write permission.

If the problem persists, contact your IBM service representative with the corresponding file name and error code.

Chapter 71. DBA0500 - DBA0999

DBA0900N The path *path-name* already exists.

Explanation: The path *path-name* was specified in the context of a non-existing path, but this path already exists.

User response: Specify a path that does not exist.

DBA0901N The path *path-name* does not exist.

Explanation: The path *path-name* was specified in the context of a existing path, but this path either does not exist or it is not accessible.

User response: Specify a existing path which is accessible.

DBA0902E The file or directory specified by path *path-name* is read-only.

Explanation: An attempt was made to modify the file or directory specified by path *path-name*, which is currently read-only.

User response: Specify a file or directory which is not read-only.

DBA0903E The file or directory specified by path *path-name* is currently in use by another application.

Explanation: An attempt was made to modify or delete the file or directory specified by *path-name*, which is currently in use by another application.

User response: Close all applications accessing the file or directory and retry the operation.

DBA0904E The file or directory specified by *path-name* is not accessible.

Explanation: The file or directory specified by *path-name* is not accessible to the current user.

User response: Verify the permissions set on the file or directory allow access to the current user and retry the operation.

DBA0905E The path or device *name* is not available.

Explanation: The path or device *name* is not available at this time.

User response: Verify the path or device is still available and retry the operation.

DBA0906E The directory *directory-name* is not empty.

Explanation: The directory *directory-name* was specified in the context of an empty directory, but this directory is not empty.

User response: Verify the directory is empty and retry the operation.

DBA0907E The path *path-name* does not refer to a directory.

Explanation: The path *path-name* was specified in the context of a directory, but this path does not refer to a directory.

User response: Verify the path specifies a valid directory and retry the operation.

DBA0908E The path *path-name* is not valid.

Explanation: The path *path-name* does not refer to a valid path.

User response: Specify a valid path and retry the operation.

DBA0909E The path *path-name* is too long.

Explanation: The length of the path *path-name* exceeds the maximum value permitted by the operating system.

User response: Specify a path whose length is within the maximum value allowed by the operating system and retry the operation.

DBA0910E There is no more space available on the file-system.

Explanation: There is no more space available on the file-system to permit the operation to complete successfully.

User response: Make sure there is enough space on the file-system to complete the operation and then retry the operation.

DBA0911E The maximum number of files that can be open has been reached.

Explanation: The maximum number of files that is permitted to be open by the operating system has been reached.

User response: Close one or more open files and retry the operation.

DBA0912E The end of the file *file-name* has been reached.

DBA0913E

Explanation: An attempt was made to read or seek beyond the end of the file *file-name*.

User response: Verify the operation is not attempting to read or seek beyond the end of the file and retry the operation.

DBA0913E A physical I/O error occurred.

Explanation: An undetermined physical I/O error while accessing the file-system.

User response: Retry the operation. If the problem persists, contact IBM Service.

Chapter 72. DBA1000 - DBA1499

DBA1005W This action will force both Local Databases and Gateway applications. Do you want to continue?

Explanation: You requested to force all applications on the instance.

User response: Click Yes to process the request or No to cancel it.

DBA1006E An invalid plugin extension has been detected by the Control Center.

Explanation: The file "db2plug.zip" has been corrupted or not properly setup.

User response: The "db2plug.zip" file should be included in the tools directory under the sqllib directory.

Refer to the documentation for instructions on defining the "db2plug.zip" file and recreate the "db2plug.zip" file.

If the problem persists, contact the system administrator for assistance or use the Control Center trace command to determine if the plugin classes are being loaded. The command "db2cc -tf <filename>" will place the Control Center trace information in the specified filename. When specifying a filename, you must provide the absolute path to the file. To determine if the plugin classes are being loaded, search through the file for lines containing the text "PluginLoader".

DBA1007W An invalid object name or parameter was entered. The Control Center will start with the Systems folder as the root of the navigator tree.

Explanation: An invalid object name means that the system, instance, subsystem, or database name entered cannot be found, either because it does not exist or because the combination of names entered is invalid. For example, the database name entered may exist but not within the specified system name. An invalid parameter means that something other than -h, -i, -sub, or -d was entered.

User response: Ensure that you are using valid object names and parameters, and try again.

DBA1100W Less than *number* MB of memory is dedicated to your server. No recommendation has been made. The current values on the Results Page match the suggested values.

Explanation: Because the amount of memory

dedicated to the server is so small, the Configuration Advisor cannot make a recommendation.

The suggested values in the "Results Page" are the same as the current values.

User response: If you can dedicate more memory to the server, turn to the "Server Page", increase the amount of memory, and try again. Otherwise, click on Cancel to close the Configuration Advisor.

DBA1101I Stripe set information about the table space container could not be obtained. Stripe sets will not be shown.

Explanation: The connected user might not have sufficient authority to obtain stripe set information.

User response: SYSADM authority is required to obtain stripe set information.

DBA1102E An error occurred when the Configuration Advisor tried to save the suggested values in a system file.

Explanation: Either the Configuration Advisor could not create a CFG subdirectory under your instance directory or it could not save a system file containing recommendations into the CFG subdirectory of your instance directory.

User response: Make sure the disk containing the instance directory is not full and you have write access to this disk. Afterwards, click on Done to try again.

DBA1103W The Configuration Advisor was unable to assign a minimum amount of memory to the buffer pools due to other memory requirements determined from your responses.

Explanation: The Configuration Advisor is unable to provide a set of recommendations based on the specified database requirement with the specified memory resource.

User response: Increase the percentage of physical memory allocated, if more memory resource is available. Otherwise, increase the amount of physical memory on the server.

DBA1104I The instance and database configuration parameters were successfully updated. The instance *name* must be stopped before most changes come into effect. You may want to rebind your packages after the new configuration parameters take effect so that the new values will be used during the bind.

Explanation: The Configuration Advisor has updated the database manager configuration parameters, database configuration parameters, and buffer pool sizes in the SYSBUFFERPOOLS catalog.

The changes will not take effect until the instance is stopped on the server.

Packages were bound with the old configuration parameters. They need to be rebound, after the new parameters take effect, to exploit these new parameters.

User response: When you are ready to use the new configuration parameters, stop the instance, and then start the instance.

Rebind your packages if necessary.

DBA1107E An error occurred when the Configuration Advisor tried to get system information from your server.

Explanation: An unexpected error occurred. The Configuration Advisor cannot continue.

User response: Contact your system administrator.

DBA1108W The Configuration Advisor was unable to increase the sizes of the buffer pools due to other memory requirements determined from your responses. The buffer pool sizes are left unchanged. The use of the suggested set of configuration values may cause paging on the server.

Explanation: This is a warning that there may not be enough memory dedicated to your server to run the workload that you specified.

User response: Review your selections on previous pages of the Configuration Advisor to check that the workload description is appropriate or add more memory to your server.

DBA1109W The transactions rate you entered is more than ten times the average number of connected applications. If you are using a transaction manager, ignore this message. If not, consider changing the rate.

Explanation: The Configuration Advisor verifies that the transactions rate is reasonable. The transaction rate

may be too high if you are not using a transaction manager.

User response: If you are using a transaction manager, ignore this message. Otherwise, use a lower number of transactions per minute, or increase the average number of connected applications.

DBA1110I The system database directory has no entries.

Explanation: There are currently no entries found in this database directory.

User response: No action is required.

DBA1111E The database alias *name* already exists for another database.

Explanation: You requested that a backup be recovered into a new database, but the database name you specified is already used as an alias for an existing database.

When a database is created, the database is cataloged in the system database directory using the database name as the alias, and this alias must be unique.

User response: Specify a database name that is not being used as an alias, or request that the backup be recovered into the existing database.

DBA1112E The database alias *name* was not found in the system database directory.

Explanation: You requested that a backup be recovered into an existing database, but a database no longer exists with this alias.

User response: Select an existing database, or request that the backup be recovered into a new database.

DBA1113E Restart failed because indoubt transactions exist against the database and the connection to the database was dropped.

Explanation: The restart operation encountered indoubt transactions. This left the database in an inconsistent state. The DBA Utility dropped the connection to the database.

User response: Resolve the indoubt transactions. If they are not resolved, applications must restart the database whenever it is required.

If you are in an XA/DTP environment, and the transaction manager that was using the database is available, instruct it to resolve the indoubt transactions.

Otherwise, use the Command Line Processor to manually complete each indoubt transaction.

DBA1114E Database *name* is in backup pending state. A full backup must be done before the database can be used.

Explanation: This state will prevent the database from being updated. The database must be backed up before any updates can happen.

User response: Back up the database.

DBA1115E Enter a value for *name*.

Explanation: A value is required.

User response: Provide the value and retry the request.

DBA1116E The value for *parameter* must be between *minimum* and *maximum*.

Explanation: The value must be within the specified range.

User response: Enter a value in the specified range and retry the request.

DBA1117E The value for *name* is not valid.

Explanation: The value entered is not valid.

User response: Correct the value and retry the request.

DBA1118E The value for *name* must not exceed *maximum*.

Explanation: The value entered is not valid.

User response: Correct the value and retry the request.

DBA1119E The value for *name* must not be smaller than *minimum*.

Explanation: The value entered is not valid.

User response: Correct the value and retry the request.

DBA1121I A job was created for the request *description* on system *system*. Its job number is *number*. To view the status and output of the job, use the Jobs page on the Journal.

Explanation: The job was successfully started.

User response: To view the status and output of the job, use the Jobs page on the Journal.

DBA1122I The job has completed for the request *description* on system *system*. For the status and output of the job, view job number in the Jobs page on the Journal.

Explanation: The job has completed.

User response: To view the status and output of the job, use the Jobs page on the Journal.

DBA1123I Job *description* is terminated.

Explanation: The job was canceled in response to a user request.

User response: No action is required.

DBA1126I The database has no table spaces.

Explanation: The database has no table spaces.

User response: No action is required.

DBA1127I The table space has no containers.

Explanation: The table space has no containers.

User response: No action is required.

DBA1128E No containers have been specified. Select "Add" in order to create containers.

Explanation: The current list does not contain any new containers.

User response: Select "Add" and fill in the necessary fields to add a container to the list.

DBA1129E Provide a value for *name*.

Explanation: The operation you requested can be performed only if you enter a value for *name*.

User response: Enter a value, and then click on OK.

DBA1130E The *name* entries must not exceed *maximum* characters.

Explanation: You entered a value that exceeded the maximum number of characters.

User response: Enter a valid value, and then click on OK.

DBA1131E *device* is not a valid tape device.

Explanation: The operating system did not accept the specified tape device. Local tape devices are of the form "\\.\TAPEn", where n represents the drive number (0 is the first).

User response: Enter a valid name for the tape device, and then click on OK.

DBA1132E Tape device *device* does not support a tape mark blocksize of *blocksize*.

Explanation: The Control center uses a tape mark blocksize that is not supported by the specified tape device.

User response: Use the Command Line Processor to submit the request. Specify a supported blocksize in the DB2 INITIALIZE TAPE command.

DBA1133E An error occurred operating tape device *device*.

Explanation: An unexpected error occurred while operating the specified tape device.

User response: Resubmit the request. If the problem persists, contact your system administrator.

DBA1134I A job was created for the request *description*. However, no job-related entries can be created in the Journal because the database administrative server (DAS) instance could not be accessed.

Explanation: The job started successfully, but the status and output of the job will not be logged in the Journal. This can happen if the DAS instance is not started.

User response: To have the status and output of the job logged in the Journal, start the DAS instance first before starting a backup operation.

DBA1135I The job has completed for the request *description*.

Explanation: The job has completed.

User response: No action is required.

DBA1136E A backup at the table space level cannot be performed because the database *db* is not enabled for rollforward recovery.

Explanation: You can only back up a database at the table space level if you have first enabled that database for rollforward recovery.

User response: To enable the database for rollforward recovery, you can either:

- Use the Database Logging Wizard to change your logging type to ARCHIVE logging.
 - Set either the `logarchmeth1` or `logarchmeth2` database configuration parameter to a value other than OFF, disconnect all existing applications from the database, and perform an offline backup of the database.
-

DBA1137E The value for *name* must not be smaller than *minimum units*.

Explanation: The value entered is smaller than the minimum *minimum units*.

User response: Enter a valid value and retry the request.

DBA1138E The value for *name* must not exceed *maximum units*.

Explanation: The value entered exceeded the maximum of *maximum of units*.

User response: Enter a valid value and retry the request.

DBA1139E You must allocate at least *pages* for the system catalog.

Explanation: There cannot be more than *pages* in each table space container. Each page of data requires 4 KB. So 1 MB of storage is equal to 256 of these 4 KB pages.

User response: Enter a valid value and retry the request.

DBA1140E You cannot allocate more than *pages* pages in a single container.

Explanation: Each page of data requires 4 KB. So 1 MB of storage is equal to 256 of these 4 KB pages.

User response: Enter a valid value and retry the request.

DBA1141E You must allocate at least *pages* pages in each container.

Explanation: The minimum number of pages in a table space container is related to the extent size of the table space. It is 5 times the extent size plus one page. So for a table space with an extent size of 32 pages (the default value), the minimum size of a single container is 161 pages.

User response: Enter a valid value and retry the request.

DBA1142W The LOGARCHIVE database configuration parameter has currently been set to ON for database *name*. As soon as all applications have disconnected from the database, the database will be in backup pending state. This state prevents the database from being updated. The database must be backed up before any additional updates can occur.

Explanation: Before the database manager can start logging each transaction in your database, it must have

a full offline database backup to use as a starting point.

User response: Make a full offline database backup as soon as possible. If this is not possible, set the LOGARCHIVE parameter to OFF.

DBA1143E You cannot perform an offline backup right now because at least one application is still connected to the database. You can either try again later, or disconnect all applications connected to the database immediately by selecting force from the instance object's pop-up menu.

Explanation: An offline backup cannot be performed until all applications connected to the selected database have stopped. You can either wait for the applications to disconnect themselves, or you can force the connected applications right now.

User response: Either try the action later, or, to force all applications connected to the selected database immediately, select "Force" from the instance object's pop-up menu.

DBA1144E The container *container-name* already exists.

Explanation: Each table space container must be unique.

User response: Specify a file name or directory that does not currently exist on the system.

DBA1145E The container *container-name* has been specified more than once.

Explanation: Each table space container must be unique.

User response: Specify a file name or directory that is not currently in your list of new containers.

DBA1146E The container *container-name* is included in both the user and temporary storage spaces.

Explanation: Each table space container must be unique.

User response: Change the file name or directory of the container in either the user or temporary table space.

DBA1147E The container *container-name* is included in both the catalog and temporary storage spaces.

Explanation: Each table space container must be unique.

User response: Change the file name or directory of

the container in either the catalog or temporary table space.

DBA1148E The container *container-name* is included in both user and catalog storage spaces.

Explanation: Each table space container must be unique.

User response: Change the file name or directory of the container in either the user or catalog table space.

DBA1149E *param-1* does not have enough free space for the containers you specified.

Explanation: When a DMS container is created it consumes the full amount of space you allocated to it.

User response: Change the location of the container to an area that has more storage, or increase the available storage in the current location of the container.

DBA1150E A column with the same name already exists.

Explanation: The column was not added because a column with the same name is already specified or exists in the table being created or altered.

User response: Specify a different column name.

DBA1151W The requested operation was not performed against the column or constraint. Reason code = *reason-code*.

Explanation: An attempt was made to add, alter, or remove a column or constraint during table creation or alteration. The reason code given indicates the actual rule violation, as follows:

1

A column or constraint with the same name already exists.

2

The column participates in the primary or distribution key and cannot be removed.

3

The column participates in a unique key and cannot be removed.

4

The column participates in a foreign key and cannot be removed.

5

The column participates in a table dimension and cannot be removed.

6

- 7 The column may participate in a table check constraint. The operation was allowed.
- 8 Primary and unique keys cannot be defined as a subset of all dimension columns.
- 9 Dimensions cannot be defined such that any primary or unique key becomes a subset of all dimension columns.
- 10 Primary and unique keys must be defined as supersets of the distribution key.
- 11 The distribution key must be defined as a subset of all existing primary and unique keys.
- 12 An object with that column set already exists.
- 13 The column participates as a partitioning column and cannot be removed.
- 14 The column participates as a partitioning column and its data type cannot be changed.
- Columns of XML data type are not supported in partitioned tables.

User response: Correct the condition that is causing the operation to fail and retry the operation.

DBA1152E A constraint with the same name already exists.

Explanation: The constraint was not added because a constraint with the same name is already specified or exists in the table being created or altered.

User response: Specify a different constraint name.

DBA1153E The object is already in the list.

Explanation: The specified object already exists in the list. The action is not performed.

User response: Specify another object or close the window.

DBA1154E The specified system, instance, and database do not identify a recognized database.

Explanation: The system, instance, and database specifications do not identify a database that is known

to the application. The action is not performed.

User response: Either ensure that the system, instance, and database names are correct; or add the system, instance, and database that you want to access to the application.

DBA1155E The object *object-name* no longer exists in the database. The request cannot be performed.

Explanation: The specified object no longer exists in the database and cannot be operated on. The request is terminated.

User response: Select the "Refresh" action in the pop-up menu of the object folder to get an updated list of objects.

DBA1156W The request completed with warnings. See the administration tools log for details.

Explanation: The requested action completed, but one or more warning messages were issued.

User response: See the administration tools log for more information.

DBA1157E No authorities or privileges were granted or revoked on *object-name* for *user-name*. See the administration tools log for details.

Explanation: The requested action completed. No action was taken against the authorities or privileges on the specified object for the specified user or group.

User response: None

DBA1158I Number of records read during the load = *count*. Number of records skipped before the load begins = *count*. Number of rows loaded into the target table = *count*. Number of records that could not be loaded = *count*. Number of duplicate rows deleted = *count*. Number of records loaded successfully and committed to the database = *count*.

Explanation: The action completed with the specified results.

User response: None

DBA1159I Export completed. *item-description* = *count*.

Explanation: The export action completed successfully with the specified results.

User response: None

DBA1160I Import completed. Number of records read during the import = *count*. Number of records skipped before the import begins = *count*. Number of rows imported into the target table = *count*. Number of records that could not be imported = *count*. Number of duplicate rows deleted = *count*. Number of records imported successfully and committed to the database = *count*.

Explanation: The import action completed successfully with the specified results.

User response: None

DBA1161I *item-description* = *count*.

Explanation: The import or export completed with the specified results.

User response: None

DBA1162I Load completed. Number of records read during the load = *count*. Number of records skipped before the load begins = *count*. Number of rows loaded into the target table = *count*. Number of records that could not be loaded = *count*. Number of duplicate rows deleted = *count*. Number of records loaded successfully and committed to the database = *count*.

Explanation: The load action completed successfully with the specified results.

User response: None

DBA1163E The table was not copied. See the administration tools log for details.

Explanation: One or more warnings or errors occurred when copying the table. The table was not copied.

User response: See the administration tools log for more information.

DBA1164E The specified target table already exists. The source table was not copied.

Explanation: The copy table action failed because the target table exists.

User response: Either specify a new target table or delete the existing target table and try the action again.

DBA1165E No privileges were granted or revoked for *user-name*.

Explanation: Privileges were not granted to or revoked from the specified user or group since no privileges have been changed.

User response: Resubmit the command after making changes to the current privileges.

DBA1166E The identifier *identifier-name* contains unbalanced quotes. Correct the identifier and try the action again.

Explanation: The specified identifier contains unbalanced single or double quotes. The naming rules require balanced quotes. The action is not performed.

User response: Correct the identifier and try the action again.

DBA1167E The identifier *identifier-name* contains embedded blanks. Correct the identifier and try the action again.

Explanation: The specified identifier contains blank characters, which violates the naming rules for ordinary identifiers. The action is not performed.

User response: Remove the blanks or delimit the identifier with double quotes and try the action again.

DBA1168E The first character in the identifier *identifier-name* is not valid. Correct the identifier and try the action again.

Explanation: The first character of the specified identifier is invalid. It must follow the naming rules for ordinary identifiers. The action is not performed.

User response: Replace the first character with a valid one or delimit the identifier with double quotes then try the action again.

DBA1169I The command script that can update the instance and database configuration parameters was created successfully. Use the Script Center to run or schedule the script.

Explanation: A new script was saved successfully in the Script Center. You may open the Script Center to view, run or schedule the new command script.

After the script is run, the changes will not take effect until the instance is stopped on the server.

Packages were bound with the old configuration parameters. They need to be rebound, after the new parameters take effect, to exploit these new parameters.

User response: To view, run or schedule the newly saved script, open the Script Center by selecting the Script Center icon from the icon bar.

After you have run the script and when you are ready to use the new configuration parameters, stop the instance, and then start the instance.

Rebind your packages if necessary.

DBA1170E The page size of the table space and the page size of the buffer pool *buffer-pool* must be equal.

Explanation: You must choose an existing buffer pool whose page size matches that of the table space you wish to create. The default buffer pool has the default page size of 4K. If any other page size is desired for the table space, a buffer pool with a matching page size must exist.

User response: From the Advanced Table Space window, change the page size of the table space to match that of the buffer pool or change the buffer pool to one whose page size matches that of the table space. If there are no buffer pools of the desired page size, you can create one from the Create Buffer Pool window.

DBA1171N The temporary table space *tablespace* specified cannot be used for this utility operation.

Explanation: The backup and restore commands cannot be performed on a temporary table space.

User response: Resubmit the command using only valid non-temporary table spaces.

DBA1172W Since the database is in backup pending state, you must perform an offline database backup. This has been preselected for you. Other options that are no longer available having been disabled.

Explanation: The only valid operation for a database in Backup Pending state is to perform an offline database backup. The wizard has ensured that these options are selected. It has also disabled the other options in the wizard that are no longer available.

User response: Complete the wizard, and run the command. It is recommended that you perform a full database backup.

DBA1173N A restore at the table space level cannot be performed because the database *database* is not enabled for rollforward recovery.

Explanation: You can only restore a database at the table space level if you have enabled the database for rollforward recovery.

User response: To enable the database for rollforward recovery, set either the **logarchmeth1** or **logarchmeth2** database configuration parameter to a value other than OFF, disconnect all existing applications from the database, and perform an offline backup of the database.

DBA1174I You have chosen more than one image for the same database partition. In most cases, you will not wish to do this.

Explanation: In most scenarios, choosing more than one backup image to restore is not useful, and will actually slow down the total time to complete the restore.

User response: Ensure that only one image for each database partition is selected.

DBA1175W This functionality requires an image for each database partition.

Explanation: Rolling forward to a point in time requires an image for each database partition. You have not chosen all database partitions on the Available Images page, and so cannot use this option.

User response: Return to the Available Images page and select one image for each database partition in your database.

DBA1176N The Control Center failed to automatically catalog the Tools Catalog Database *database-name* for system *system-name*.

Explanation: Possible explanations are:

- The database is specified incorrectly at the DB2 Administration Server for system *system-name*.
- The database manager that contains the database is not configured for TCP/IP communication.

User response: Possible solutions include:

- Verify that the DB2 Administration Server configuration parameters are correctly specified.
- Verify that the database manager is configured for TCP/IP communications.

If the problem persists, contact IBM Support.

DBA1177N A database connection to the Tools Catalog Database *database-name* cannot be made. **SQLCODE** *sqlcode*

Explanation: A connection to the Tools Catalog Database *database-name* as defined in the DB2 Administration Server configuration parameters cannot be made.

This message will appear under one of the following conditions:

1. The Tools Catalog Database has not been created.
2. The Tools Catalog Database has been created but the DB2 Administration Server configuration parameters have not been updated.
3. The Tools Catalog Database is configured properly but the connection failed.

User response: Suggestions corresponding to the three conditions are:

1. Create a Tools Catalog Database by issuing the CREATE TOOLS CATALOG command from the CLP.
2. Update the DB2 Administration Server configuration parameters to define the Tools Catalog Database. Four parameters need to be updated with the UPDATE ADMIN CONFIGURATION command:
 - SCHED_ENABLE
 - TOOLSCAT_INST
 - TOOLSCAT_DB
 - TOOLSCAT_SCHEMA
3. Refer to the user response for the SQLCODE returned from the connect operation.

DBA1178W Task *task-name* was created with a more recent version of the Control Center.

Explanation: The selected task was created with a more recent version of the Control Center. Some functionality for the task are not supported by the Control Center on the local machine.

User response: You can choose to proceed and edit the task. When saving the task, you can:

- Save the changes to the existing task. This will cause the loss of functionality that was specific to the more recent version of the Control Center.
- Save changes by creating a new task. This will not alter the original task.

DBA1179W Task *task-name* created with a more recent version of the Control Center will be overwritten.

Explanation: A task originally created with a more recent version of the Control Center will be overwritten. Proceeding with this operation will cause the loss of functionality that was specific to the more recent version of the Control Center.

User response: You can choose to:

- Save the changes to the existing task. This will cause the loss of functionality specific to the more recent version of the Control Center.
- Save changes by creating a new task. This will not alter the original task.

DBA1180W No SQL available to display.

Explanation: Show SQL only displays the SQL statements that change the current privileges. If no privileges have been changed, no SQL statements are generated.

User response: Resubmit the command after making changes to the current privileges.

DBA1181W The catalog partition cannot be backed up at the same time as other database partitions in an offline backup.

Explanation: In an offline backup, the catalog partition cannot be backed up at the same time as other database partitions. If you continue, the backup will fail.

User response: Ensure that the catalog partition has not been grouped with other database partitions on the "Partitions" page of this wizard.

DBA1183N Unable to edit task *task-name*

Explanation: The selected task was created with a more recent version of the Control Center. There is no editor associated with this type of task in this version of the Control Center.

User response: Use a newer version of the Control Center to edit this task.

DBA1184W The database will immediately be taken offline and a full database backup will be performed.

Explanation: As part of the execution of this wizard an offline full database backup will now be performed. This operation will remove all current users from the system and make the database unavailable for the duration of the backup. This operation may run for an extended period.

User response: Click OK to continue or Cancel to return back to the wizard.

DBA1185W Automatic maintenance is enabled for this database.

Explanation: This database has been configured for automatic maintenance, and is performing automatic database backups. Instead of performing a manual database backup through the Backup wizard, you may want to modify your automatic backup settings. This is performed using the Configure Automatic Maintenance wizard.

User response: Launch the Configure Automatic Maintenance wizard to configure your automatic database backup settings or continue to manually backup using the Backup wizard.

DBA1186W The operation has completed successfully, except for the required database backup. Your database is currently in Backup Pending state.

Explanation: The only valid operation for a database in Backup Pending state is to perform an offline database backup. You can perform an offline backup in the Backup wizard.

User response: Launch the Backup wizard or press Cancel to return.

DBA1187N The days of the month entered are invalid.

Explanation: As part of the specification for the maintenance window you have selected to have it apply only on specific days of the month. The days and day ranges entered contain invalid characters, and/or invalid ranges of days.

User response: Enter a valid list of days and day ranges and press OK.

DBA1188E The maintenance settings could not be updated or retrieved.

Explanation: The automatic maintenance settings are stored in tables in the database. Either the tables could not be accessed or the settings stored within the table contain invalid data.

User response: If the problem persists, contact IBM Support.

DBA1189I There are unrecognized settings in the automatic maintenance settings configuration for this database.

Explanation: The unrecognized settings are ignored.

User response: No action required.

DBA1350I The Design Advisor was not able to recommend any indexes. Set a longer time limit on the Calculate page and try calculating recommended objects again.

Explanation: The Design Advisor was not able to recommend any objects because the maximum search time was too short.

User response:

1. On the Calculate page in the Design Advisor, set a longer time limit or no limit at all.
2. Try running the Design Advisor again.

DBA1351I The Design Advisor was not able to recommend any objects. Set a larger size limit on the Limits page and try calculating recommended objects again.

Explanation: The Design Advisor was not able to recommend any object because the maximum logical size of all objects was too small.

User response:

1. On the Limits page in the Design Advisor, set a larger maximum logical size or no maximum at all.
2. Try running the Design Advisor again.

DBA1352I The Design Advisor was not able to recommend any objects that will improve the performance of your workload.

DBA1353W You specified only one SQL statement in your workload. Your database might have additional activity. Do not drop any objects unless you are sure they are not needed for other reasons.

Explanation: Since you only specified one SQL statement in your workload, it is unlikely that it represents all of the activity against your database.

The Drop page of the Design Advisor lists all existing objects that were not recommended for this one SQL statement.

User response: Do not drop any objects unless you are sure they are not needed for other reasons.

DBA1354E An error occurred while compiling the SQL statements in the specified workload.

Explanation: The Design advisor encountered a problem while compiling the SQL statements in the specified workload. One or more statements may contain a SQL syntax error. These statements are excluded from the advisor's analysis.

User response: Open the Workload Details Dialog to view the statements which contain an error. From the Workload page, make the necessary corrections by editing the statement in the workload. Once corrected, select Recommendation page to start analysis.

DBA1355E The ADVISE and the EXPLAIN tables do not match.

Explanation: The ADVISE and the EXPLAIN tables do not correspond with the current version of IBM DB2.

User response: If the problem persists, contact IBM Support.

DBA1356E The Recommendation Advisor encountered a non-severe error while attempting to retrieve recommendations for alert.

Explanation: An error occurred while the Recommendation advisor was retrieving information to resolve the current Health Indicator alert.

User response: If the problem persists, contact IBM Support.

DBA1357E **An error occurred while attempting to launch the *GUI-tool***

Explanation: An unexpected error occurred while attempting to launch the specified tool

User response: Attempt the launch the tool using an alternate method. If the problem persists, contact IBM Support.

DBA1358E **The Design Advisor encountered an unexpected error. Return code = *return-code*.**

User response: If the problem persists, contact IBM Support.

Chapter 73. DBA1500 - DBA1999

DBA1500E The local system object is not allowed to be removed.

Explanation: The local system object appears in the tree if this is a server installation, and does not appear if this is a client installation. You have no direct control over the removal of this object because it has special properties that are required by the server installation type.

User response: No action is required.

DBA1501E The local system object is not allowed to be changed.

Explanation: The local system object appears in the tree if this is a server installation, and does not appear if this is a client installation. You have no direct control over the changing of this object because it has special properties that are required by the server installation type.

User response: No action is required.

DBA1502E The unknown system object is not allowed to be removed.

Explanation: The unknown system object appears in the tree only if orphaned instances that are not recognized as belonging to any existing system are found. You have no direct control over the removal of this object because it is shown automatically only when needed.

User response: No action is required. However, if you want to remove the need to show this object in the tree, then do the following steps:

- Expand the unknown system object to show all orphaned instances.
- If necessary, add new system objects to the control center to contain the orphaned instances.
- If necessary, use the change action to alter the system name on any remaining orphaned instances.
- If the unknown instance object also appears, then you also have orphaned databases that are not recognized as belonging to any existing instance. Remove the need to show these objects by doing steps similar to these, adding instances or altering databases until no orphans are left to be shown.

DBA1503E The unknown system object is not allowed to be changed.

Explanation: The unknown system object appears in the tree only if orphaned instances that are not recognized as belonging to any existing system are

found. You have no direct control over the changing of this object because it is shown automatically only when needed and has only special properties that cannot be changed.

User response: No action is required.

DBA1510E Specify a system name.

Explanation: You must specify a system name before you can add this new system object.

User response: Specify a system name in the appropriate entry field, and then try the action again.

DBA1511E The specified system name is already in use. You must specify a system name that is unique.

Explanation: You specified a system name that matches the name used for another system object in the tree. System names must be unique within the control center, and are case-sensitive. You cannot use the names 'Local' or 'Unknown' because these are system names reserved for use by the local and unknown system objects. To avoid confusion, you should also try not to use the names of existing instances or databases as system names, although the tool does not prevent you from doing this.

User response: Specify a unique system name, and then try the action again.

DBA1520W The operating system type for this system object has been detected to be *newos*, but the local information shows it to be *oldos*. Do you want to update the local information now to reflect the correct operating system type?

Explanation: When you add a new system object, you specify an operating system type, which is stored locally for that remote system. Later, when actual connections are made to the remote system, the true operating system type is reported back. It is possible for the original operating system type information to have been set incorrectly so that it does not match the true settings on the remote system. In this case, it is best to update the local information to match what settings truly exist on the remote system.

User response: Select Yes to update the local information now, or select No to leave the local information as it is.

DBA1521W The server version type for this system object has been detected to be *newtyp*, but the local information shows it to be *oldtyp*. Do you want to update the local information now to reflect the correct server version type?

Explanation: When you add a new system object, the server version type is automatically set to be either DB2 V5 or DRDA, depending on the operating system type you choose. Later, when actual connections are made to the remote system, the true server version type, which may be a previous release such as DB2 V2, is reported back. In this case, it is best to update the local information to match what settings truly exist on the remote system.

User response: Select Yes to update the local information now, or select No to leave the local information as it is.

DBA1522W The operating system type and server version type for this system object have been detected to be *newos* and *newtyp* respectively, but the local information shows these to be *oldos* and *oldtyp* respectively. Do you want to update the local information now to reflect the correct settings for these values?

Explanation: When you add a new system object, you specify an operating system type, which is stored locally for that remote system. Also, the server version type is automatically set to be either DB2 V5 or DRDA, depending on the operating system type you choose. Later, when actual connections are made to the remote system, the true operating system type and server version type are reported back. The original settings for this information may have been incorrect so that they do not match the true settings on the remote system. In this case, it is best to update the local information to match what settings truly exist on the remote system.

User response: Select Yes to update the local information now, or select No to leave the local information as it is.

DBA1530E The specified system has not been configured to support any protocols that your local system uses.

Explanation: The specified system has been detected on the network, however the Administration Server for this remote system has not been configured to support any protocols that your local system uses.

User response: You must configure communications on the remote system Administration Server, including the DB2COMM parameter, before you can successfully access it.

DBA1533E Unable to export the server profile.

Explanation: The target file system might not have enough disk space to complete the export operation.

User response: Free some disk space on the target file system and retry the operation. If the problem persists, contact your DB2 system administrator.

DBA1534W A valid mail server has not been configured.

Explanation: The current value of the SMTP_SERVER DB2 Administration Server configuration parameter is not valid.

User response: Use the Troubleshoot Health Alert Notification wizard to identify a valid mail server.

DBA1540E The active local instance object is not allowed to be removed.

Explanation: The active local instance object appears in the tree always as the first instance under the local system object. This is a special instance object used to represent the current DB2INSTANCE environment variable setting. You have no direct control over the removal of this object because it has special properties that are required by the local system.

User response: No action is required.

DBA1541E The active local instance object is not allowed to be changed.

Explanation: The active local instance object appears in the tree always as the first instance under the local system object. This is a special instance object used to represent the current DB2INSTANCE environment variable setting. You have no direct control over the changing of this object because it has special properties that are required by the local system.

User response: No action is required.

DBA1550E Specify an instance name.

Explanation: You must specify an instance name before you can add this new instance object.

User response: Specify an instance name in the appropriate entry field, and then try the action again.

DBA1551E The specified instance name is already in use. You must specify an instance name that is unique, otherwise do not specify the instance name and one will be generated for you automatically.

Explanation: You specified an instance name that either matches the name used for another instance object in the tree, or is being used by a system object in the tree. Instance names must be unique within the

control center, and are always capitalized. System objects also use an automatically generated instance name to store their own protocol information, with the instance name usually being a short form of the chosen system name. You cannot use the name 'LOCAL' or the current DB2INSTANCE environment variable setting for the instance name.

User response: Specify a different instance name, and then try the action again.

DBA1552E Specify a remote instance.

Explanation: You must specify a remote instance before you can add this new instance object.

User response: Specify a remote instance in the appropriate entry field, and then try the action again.

DBA1560E Specify a destination name.

Explanation: You must specify a destination name for the selected protocol before you can correctly add this new object.

User response: Specify a destination name in the appropriate entry field, and then try the action again.

DBA1561E Specify a file server.

Explanation: You must specify a file server for the selected protocol before you can correctly add this new object.

User response: Specify a file server in the appropriate entry field, and then try the action again.

DBA1562E Specify a workstation name.

Explanation: You must specify a workstation name for the selected protocol before you can correctly add this new object.

User response: Specify a workstation name in the appropriate entry field, and then try the action again.

DBA1563E Specify a host name.

Explanation: You must specify a host name for the selected protocol before you can correctly add this new object.

User response: Specify a host name in the appropriate entry field, and then try the action again.

DBA1564E Specify a service name.

Explanation: You must specify a service name for the selected protocol before you can correctly add this new object.

User response: Specify a service name in the appropriate entry field, and then try the action again.

DBA1565E Specify a computer name.

Explanation: You must specify a computer name for the selected protocol before you can correctly add this new object.

User response: Specify a computer name in the appropriate entry field, and then try the action again.

DBA1566E Specify an instance name.

Explanation: You must specify an instance name for the selected protocol before you can correctly add this new object.

User response: Specify an instance name in the appropriate entry field, and then try the action again.

DBA1567E This action or function is not available for this instance type or the associated system.

Explanation: The selected action, or associated function, is currently not supported for the type of instance against which the action or function was initiated. Some instance types, for example Satellite Edition, necessarily do not support full functionality.

User response: Refer to the documentation describing available function for this instance type.

DBA1568W Database connection was established but did not establish a syncpoint two-phase connection.

Explanation: An attempt was made to connect to a database using syncpoint two-phase connection type. A connection was established but failed to make a syncpoint two-phase connection. Multisite updates involving this database connection will not be successful.

User response: Check that the multisite update scenario has been configured properly for your environment, ensuring that the coordinating instance is restarted after configuration is complete. If DB2's Syncpoint Manager is being used, check in the db2diag log file to see if it was started successfully.

DBA1569E The action cannot be completed on this instance because its remote instance name is not known.

Explanation: The remote instance name field is an optional parameter and was not specified. To perform this action on a remote instance, the remote instance name must be known.

User response: Invoke the Change Instance window and specify the name of the remote instance.

DBA1570E The DB2 Administration Server for this system has not been defined.

Explanation: The current operation requires an attachment to the system's administration server but its instance name has not been specified.

User response: Invoke the change action on the target system and specify its instance name, operating system and any other protocol parameters.

DBA1571W The name of the local system does not match the DB2SYSTEM environment variable.

Explanation: The system name of the local admin node is different from the value of the DB2SYSTEM environment variable.

User response: There are three options:

- Change the value of the DB2SYSTEM environment variable to match the system name of the local admin node.
 - Uncatalog the local admin node and recatalog it with a system name that matches the value of the DB2SYSTEM environment variable.
-

DBA1572E The instance name *instance-name* specified is not valid.

Explanation: The instance name specified is not valid. The instance name may be 1 to 8 characters and all of the characters must be from the database manager base character set. Also, the instance name must not match that of any other catalogued instance.

User response: Specify a different value for the instance name.

DBA1573E This action is not available for the selected system.

Explanation: Some Control Center actions require that an Administration Server be available on the target system. The system that you selected does not have an Administration Server. The requested action is not available for this system.

User response: No action is required.

DBA1580E New database objects are not allowed to be added to the unknown instance.

Explanation: The unknown instance object appears in the tree only if orphaned databases that are not recognized as belonging to any existing instance are found. The tool tries to prevent you from creating new orphans whenever possible, so you are not allowed to add new ones to the unknown instance.

User response: No action is required.

DBA1581E Specify a database name.

Explanation: You must specify a database name before you can add this new database object.

User response: Specify a database name in the appropriate entry field, and then try the action again.

DBA1582E Specify an alias.

Explanation: You must specify an alias before you can add this new database object.

User response: Specify an alias in the appropriate entry field, and then try the action again.

DBA1583E The specified alias is already in use.

Explanation: You specified an alias that matches the alias used for another database object in the tree. Database aliases must be unique within the control center, and are always capitalized.

User response: Specify a unique alias, and then try the action again.

DBA1590W The *dbase* database was created using an earlier version of DB2 than DB2 V5. Only limited function will be allowed for this database until it is migrated.

Explanation: All databases that are found under DB2 V5 type system objects are first assumed to be DB2 V5 databases. This database has now been detected to actually be a back-level database that has not yet been migrated to the DB2 V5 level. You can still access this database, but only limited function is available for back-level databases until you actually migrate them.

User response: No action is required.

Chapter 74. DBA2000 - DBA2499

DBA2000E You have not specified a job description.

Explanation: A description for the job is required information.

User response: Enter a job description in the entry field, and then click on OK.

DBA2001E You have not specified the number of hours for your selection.

Explanation: The number of hours for your selection is required information.

User response: Enter a number in the hours entry field.

DBA2002E You must select at least 1 day of the week for your selection.

Explanation: To schedule a recurring job, you must specify at least one day of the week for the job.

User response: Select at least one day of the week when scheduling a recurring job.

DBA2003E You specified a date that is not valid.

Explanation: The date you specified is not valid. Perhaps the date you have specified is in the past.

User response: Enter a valid date.

DBA2004E You specified a date or time or combination of date and time that is not valid.

Explanation: The next execution date and time cannot be calculated, so no jobs can be scheduled. One possible error is when the time specified for a job to be run once is in the past. Remember that the date and time combination must be valid on the managed host where the schedule job is to be run.

User response: Enter a valid date and time combination.

DBA2005E You have not specified a script name.

Explanation: A name for the script is required information.

User response: Enter a name for the script file, and then click on OK.

DBA2006E You have not specified a comment.

Explanation: You indicated that you want to record a

comment when the job is completed, but the comment entry field is empty.

User response: Enter a comment, or clear the 'Record a comment' check box.

DBA2007E You specified a script name that is not valid.

Explanation: The command script is saved as a file in the file system. The script path is at most 255 bytes long. The file name part of the script name can be at most 8 bytes long. If an extension is specified, it cannot be longer than 3 bytes. The script name must not contain any blanks. For example, valid script names can be: j:\script\new.cmd c:\data\test\crtddb.bat

User response: Enter a valid script name, and then click on OK.

DBA2008I Job *job-id* was created successfully.

Explanation: A new job was scheduled successfully. You may turn to the journal job summary to view the newly created job.

User response: To view the newly created job, turn to the journal job summary by selecting the Journal icon from the icon bar.

DBA2009E A system error occurred. A Scheduler InfoBase call was unsuccessful. RC = *return-code*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2010E The Scheduler service is not up and running. Restart the Admin. Server, and then retry the command.

Explanation: To perform any action involving the Job Scheduler or the Script Center, the Scheduler service must be up and running.

User response: Restart the Admin. Server, and then retry the command.

DBA2011E A system error occurred. Scheduling type (DBA or PGM) is not valid. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2012E A system error occurred. OS error was detected. RC = *return-code*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2013E The job id was not found. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2014E A system error occurred. Unable to change the status of the job. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2015E A system error occurred. Memory allocation was unsuccessful in client. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2016E A system error occurred. Scan handle is not valid. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2017I Are you sure you want to remove the selected jobs?

Explanation: You specified to remove the selected jobs. You have the chance now to double-check the jobs

that you selected to remove, and either to continue or to cancel the request.

User response: Select YES to remove the selected jobs. Select NO to cancel the remove request.

DBA2018I Are you sure you want to remove the selected saved scripts?

Explanation: You specified to remove the selected scripts that are currently saved. You have the chance now to double-check the scripts that you selected to remove, and either to continue or to cancel the request.

User response: Select YES to remove the selected scripts. Select NO to cancel the remove request.

DBA2019E You have not specified the number of days for your selection.

Explanation: The number of days for your selection is required information.

User response: Enter a number in the days entry field.

DBA2020E A system error occurred. Unable to determine script type. Command script was not executed.

Explanation:

User response:

DBA2021I Are you sure you want to discard your changes?

Explanation: You specified to have your changes not saved in the Script Center. You have the chance now either to continue discarding your changes, or to cancel the request and resume editing.

User response: Select YES to discard the changes. Select NO to resume editing the command script.

DBA2022I Command script *script-id* was created successfully.

Explanation: A new script was saved successfully in the Script Center. You may turn to the Script Center to view the new command script.

User response: To view the newly saved script, turn to the Script Center by selecting the Script Center icon from the icon bar.

DBA2023E Instance *instance-name* does not exist. Command script was not executed.

Explanation: The instance name associated with the script file does not exist.

User response: You may select "Edit" from the Script Center to update the instance name to an existing instance.

DBA2024E Command script *script-id* does not exist.

Explanation: The scheduled job cannot be executed because the script file associated with the job does not exist.

User response:

DBA2025E A system error occurred. This problem may also be related to memory allocation on the client. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2026I Are you sure you want to remove the selected job?

Explanation: You specified to remove the selected job. You have the chance now to double-check the job that you selected to remove, and either to continue or to cancel the request.

User response: Select YES to remove the selected job. Select NO to cancel the remove request.

DBA2027I Are you sure you want to remove the selected saved script?

Explanation: You specified to remove the selected script that is currently saved. You have the chance now to double-check the script that you selected to remove, and either to continue or to cancel the request.

User response: Select YES to remove the selected script. Select NO to cancel the remove request.

DBA2028E You have not specified the number of weeks for your selection.

Explanation: The number of weeks for your selection is required information.

User response: Enter a number in the weeks entry field.

DBA2029E A system error occurred. Scheduler key was not found. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2030E You did not specify a system name. Select one from the list.

Explanation: The system name is required information to process the action.

User response: Select a system name from the drop-down list.

DBA2031I Job *job-id* was rescheduled successfully.

Explanation: A job that you requested to be rescheduled was rescheduled successfully.

User response:

DBA2032E You have not specified a script name.

Explanation: A name for the script is required information.

User response: Enter the script name in the entry field, and then click on OK.

DBA2033E You have not entered the command script.

Explanation: The command script is empty.

User response: Enter the command script content, and then click on OK.

DBA2034I Command script *script-id* was updated successfully.

Explanation:

User response:

DBA2035E A system error occurred. Unable to retrieve the command script *script-id*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2036E A system error occurred. Unable to create the command script. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2037E A system error occurred. Unable to replace the command script *script-id*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2038E A system error occurred. Unable to copy the command script *script-id*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2039E You have not specified the number of months for your selection.

Explanation: The number of months for your selection is required information.

User response: Enter a number in the months entry field.

DBA2040E You must select at least 1 day of the month for your selection.

Explanation: To schedule a recurring job, you must specify at least one date per month for the job.

User response: Select at least one date per month when scheduling a recurring job.

DBA2041E A system error occurred. Unable to remove the command script *script-id*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2042E A system error occurred. Unable to run the command script *script-id*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2043I Job *job-id* was created successfully for running command script *script-id*.

Explanation:

User response:

DBA2044E You have not specified a script description.

Explanation: A description for the script file is required information.

User response: Enter a description for the script file, and then click on OK.

DBA2046E Command script *script-id* is system-generated. Content cannot be edited.

Explanation: The script associated with a backup or reorganize table action is not in readable format.

User response: Select a different command script file to edit.

DBA2047E A system error occurred. The script type is unknown. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2048E The file name you specified already exists. Specify a different script file name.

Explanation: The name of each script file must be unique. You specified a file name that already exists.

User response: Specify an unique script file name.

DBA2049E A system error occurred. File access was denied.

Explanation: The application was unable to access the file in the file system.

User response: Retry the command. If the problem persists, contact your system administrator.

DBA2050E A system error occurred. The disk is full.

Explanation: There is no more room on the disk. Processing cannot continue.

User response: Delete unwanted files from the file system. If the problem persists, contact your system administrator.

DBA2051E A system error occurred. A hard disk error occurred.

Explanation: An internal error relating to the hard disk occurred. The application was unable to access a file.

User response: Retry the command. If the problem persists, contact your system administrator.

DBA2052E Unable to execute the job due to logon failure. Userid = *user-id*.

Explanation: The password was invalid for the given userid or the userid is invalid.

User response: Enter a valid userid and password when prompted.

DBA2053E A system error occurred. Unable to remove job history entry.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2054E A system error occurred. Unable to run job *job-id*.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2055E Command script *script-id* has been removed from the Script Center.

Explanation:

User response:

DBA2056I Command script *script-id* was not changed.

Explanation:

User response:

DBA2057E A system error has occurred. The path was not valid.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2058E A line in the file is greater than *number* bytes.

Explanation: The maximum number of bytes in a line is exceeded.

User response: Rewrite the long line into separate lines without exceeding the limit.

DBA2059E A system error occurred. There was a sharing violation.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2060E A system error occurred. Unable to copy script *script-id*.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2061I Script *script-name* was created successfully.

Explanation: A new script file with the name *script name* was created successfully.

User response: To view the newly created script, turn to the Script Center by selecting the Script Center icon from the icon bar.

DBA2062E A system error occurred. Unable to create a new script. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2063E A system error occurred. Unable to retrieve the results of job *job-id*. Restart the Admin. Server, and then retry the command.

Explanation: An internal error occurred.

User response: Restart the Admin. Server, and then retry the command. If the problem persists, contact your system administrator.

DBA2064E A Script Center record already exists for this script file name. Specify a different script file name.

Explanation: The name of each script file must be unique within the Script Center. You specified a file name that already exists.

User response: Specify an unique script file name.

DBA2065E A system error occurred. The script file does not exist in the file system.

Explanation: The script file was not found in the file system. The file may have been erased.

User response: Remove the script entry and recreate it.

DBA2067I The REORG TABLE command completed successfully.

Explanation:

User response:

DBA2068E You have not entered the instance name.

Explanation: The name of the instance is required information.

User response: Enter the instance name in the entry field, and then click on OK.

DBA2069W The maximum number of *number* jobs has been created. Please remove obsolete jobs.

Explanation:

User response: Remove jobs that are no longer needed before creating another new job.

DBA2070W The maximum number of *number* scripts has been created. Please remove obsolete scripts.

Explanation:

User response: Remove scripts that are no longer needed before creating another new script.

DBA2071W The maximum number of *number* bytes has been reached. Data shown is truncated. The complete file *file-name* can be found at the server.

Explanation: The editor cannot handle a file as large as the current one. Data shown had to be truncated.

User response: The entire file can be viewed at the server in the indicated location.

DBA2072E The script associated with job *job-id* has been removed from the Script Center.

Explanation: The script file does not exist anymore.

User response:

DBA2073E You have not specified a userid and password.

Explanation: A userid and password for running the job are required.

User response: Enter a valid userid and password for running the job, and then click on OK.

DBA2074E A system error occurred. The working directory path was not valid.

Explanation: The specified working directory does not exist so it cannot be used to run the script.

User response: Use the Browse... button to select a valid working directory path.

DBA2075I Job *job-id* has been submitted. Use the Jobs page on the Journal to view the results.

Explanation:

User response:

DBA2076E The job terminated abnormally.

Explanation:

User response:

DBA2077E File name *file-name* already exists in the file system.

Explanation: File name must be unique.

User response: Specify another file name.

DBA2078E An unexpected error occurred. There is no job output.

Explanation: There is no job output from running the job.

User response: A possible error could be an invalid file extension is used on the operating system. Recreate the script with a valid extension.

DBA2081W The file *file-name* exceeds the maximum number of bytes (*maximum-bytes*) for the Script Center. You must use an external editor to view or edit the file.

Explanation: The file is larger than the maximum size that the Script Center can display.

User response: Use an external editor to view or edit the file.

DBA2082W **The file permissions for the script do not allow Script Center to read the script with the current user ID.**

Explanation: The file permissions for the script do not allow Script Center to read the script with the current user ID.

User response: Ensure that you have set appropriate file permissions for the script.

DBA2083I **You are about to edit a dialog based task and in doing so will remove the ability to edit it with its dialog. Do you wish to continue?**

Explanation: The task you are about to edit was created by a DBA Tools dialog or a Wizard. It is recommended to use the 'Edit with dialog' action on this task instead of editing it manually. If you choose to edit this task manually, you will lose the ability to 'Edit with dialog' in the future. Other capabilities may also be lost, such as detailed 'Show Progress'.

User response: Select 'No' to keep the task in a dialog based task format. Select 'Yes' to convert the task to a simple task and continue editing it manually.

DBA2151E **The script cannot be scheduled because it was not saved.**

Explanation: The script cannot be scheduled because you canceled the save action and therefore, the script was not saved in the Script Center.

User response: Save the script and then try scheduling again.

DBA2152I **The script has to be saved to the Script Center before it can be scheduled.**

Explanation: Only scripts that are saved in the Script Center can be scheduled and run.

User response: Select OK to save the script to the Script Center, or Cancel if you do not want to schedule the script.

DBA2153I **File *file-name* was opened successfully.**

Explanation:

User response:

DBA2154I **Script *script-name* was opened successfully.**

Explanation:

User response:

DBA2155I **Do you want to overwrite file *file-name*?**

Explanation: The file *file name* currently exists in the file system and the save operation will overwrite its contents.

User response: Select OK if you want to overwrite the contents of the file, or Cancel if you do not want to alter its contents.

DBA2156I **Do you want to overwrite script *script-name*?**

Explanation: The script *script name* currently exists in the Script Center and the save operation will overwrite its contents.

User response: Select OK if you want to overwrite the contents of the script, or Cancel if you do not want to alter its contents.

DBA2157I **File *file-name* was updated successfully.**

Explanation:

User response:

DBA2158I **File *file-name* was created successfully.**

Explanation: A new file with the name *file name* was created successfully in the file system.

User response:

DBA2160I **The script has to be saved to the Task Center before it can be scheduled.**

Explanation: Only scripts that are saved in the Task Center can be scheduled and run.

User response: Select OK to save the script to the Task Center, or Cancel if you do not want to schedule the script.

DBA2161I **Do you want to overwrite script *script-name*?**

Explanation: The script *script name* currently exists in the Task Center and the save operation will overwrite its contents.

User response: Select OK if you want to overwrite the contents of the script, or Cancel if you do not want to alter its contents.

DBA2171I **Do you want to discard all your changes to the script, and exit from the Command Center?**

Explanation: You specified to exit from the Command Center, but the changes to the current script have not yet been saved. Your changes will be lost unless you cancel your exit request.

User response: Select Yes to discard your changes and exit from the Command Center, or select No to cancel the exit request.

DBA2172I Do you want to discard your changes to the current script, and create a new script?

User response: Select Yes to discard your changes and to create a new script, or select No to cancel the request.

DBA2173I The Command Center is still processing a command. Try again later.

Explanation: You tried to run a script or issue a set of commands, but the Command Center cannot process your command until it has finished processing the current one.

User response: Please wait a few moments and try the command again. To stop the current command, click on the turning gears icon on the tool bar or select "Terminate" from the Results menu on the Results page.

DBA2174E The file *file-name* specified on the command line could not be opened. Make sure that the file exists and that you have read access to the file, and then try again.

Explanation: You specified the name of a file containing a script that you want to import, but the Command Center could not access that file.

User response: Make sure that you have entered the file name correctly on the command line, that the file exists, and that your user ID has read access to that file.

DBA2175E There is not enough memory to create the Command Center window. The application cannot continue.

Explanation: The system encountered a memory allocation problem when trying to create the Command Center window.

User response: Make sure there is enough memory on your system to run the Command Center, and then try again.

DBA2176E The system environment did not initialize. The application cannot continue. Restart the database manager and try again, or contact your local support personnel.

Explanation: The command processor environment did not initialize properly.

User response: Make sure you have enough disk space and then try again.

DBA2177E The system could not create an access plan because a problem relating to Visual Explain has occurred.

Explanation: Command Center has encountered a problem when trying to create an access plan.

User response: Please refer to the help from the previous messages.

DBA2178E You need a database connection in order to create an access plan. Use the "Connect" statement to connect to a database, and then try the command again.

User response: Establish a database connection first using the "Connect" statement, and then create an access plan for that database. To see a list of available databases, issue the command "list database directory".

DBA2179E An access plan cannot be created because the statement you used is not in the correct format or is not a query.

Explanation: The statement you used to try to create an access plan is either not in the correct format for a query or is not a query.

User response: Check your statement for errors in query format, make necessary changes, and retry the statement. Refer to Visual Explain documentation for guidance on forming valid queries.

DBA2180I The program name "db2" before the statement will be ignored because it is not required.

Explanation: You don't need to type "db2" before a statement in the Command Center. It is only required at the operating system command prompt.

User response: No action is required.

DBA2181I The Command Center is not available at this time. Retry this operation later.

Explanation: The JDBC server does not support multiple WEBCC Command Center sessions. The JDBC server is already busy with an existing Command Center session.

User response: Wait until the current Command Center session is finished and retry the operation.

DBA2182N A script name following the pattern *script-name-pattern* cannot be saved because the Script Center already contains too many scripts with names that follow the same pattern. Remove old scripts that follow this pattern from the Script Center then invoke the *tool-name* again.

Explanation: There are too many automatically generated scripts with names like *script_name_pattern*.

User response: In the Script Center, remove at least one script with a name like *script_name_pattern*, then invoke the *tool_name* again.

DBA2192E Database connection failed.

Explanation: A JDBC connection to the database failed. SQL Assist requires a JDBC connection to function.

Explanation: Verify that the JDBC driver is running. Reconnect to the database and try launching SQL Assist again.

If the JDBC driver is running and the problem persists, contact IBM Support.

DBA2193W *window-name* cannot run without *program-name*. Ensure that the Database Tools subcomponent is installed.

Explanation: To invoke *window-name*, the *program-name* application must be installed on the server. By default, the Database Tools subcomponent of the Administration and Configuration Tools component is selected during installation. The Database Tools subcomponent must be selected during installation in order for the *program-name* application to be installed.

User response: Ensure the Database Tools subcomponent is installed. If Database Tools was not selected during the installation of the database server, you will need to run DB2 Setup again.

- 1 Stop all DB2 services.
- 2 Run DB2 Setup.
- 3 Select Custom install.
- 4 Uncheck all components except Administration and Configuration Tools.
- 5 In the details for Administration and Configuration Tools, select Database Tools.
- 6 Complete the install.
- 7 Restart all DB2 services.
- 8 Use the *window-name*.

DBA2194E The XQuery Assist component could not be opened.

Explanation: The XQuery Assist component is part of the IBM Data Studio, which might not be properly installed.

User response: Reinstall IBM Data Studio.

Chapter 75. DBA3000 - DBA3499

DBA3007W Are you sure you want to delete the explained statement?

Explanation: You are about to delete the explained statement. This will delete the rows in the explain tables.

User response: Ensure that you want to delete the explained statement.

DBA3008W Are you sure you want to delete the explained statements?

Explanation: You are about to delete the explained statements. This will delete the rows in the explain tables.

User response: Ensure that you want to delete the explained statements.

DBA3009E A system error occurred. The Visual Explain tool could not continue processing.

Explanation: An unexpected error occurred causing the operation to stop.

User response: Possible solutions include:

- Verify that your system has sufficient memory.
- Where appropriate, stop other programs using the system.

If the problem persists, contact your system administrator.

DBA3010E The explain snapshot has been deleted. Your request cannot be completed.

Explanation: The explain snapshot that is requested has been deleted from the explain tables.

User response: Regenerate the snapshot.

DBA3011E This explain snapshot was created using a different version of the database manager from the Visual Explain Tool. Your request cannot be completed.

Explanation: The explain snapshot was created using a newer version or an older version of the database manager from the Visual Explain Tool. The access plan graph cannot be constructed by the Tool.

User response: Use the same version of the database manager as the Visual Explain Tool to graph the explain snapshot.

DBA3012I String not found.

Explanation: The search string is not found in the text being displayed.

User response: None.

DBA3013E Text is missing. You must specify the text string to be found.

Explanation: You did not specify a search string.

User response: Specify a search string.

DBA3014E No selection is made in the list. You must select at least one entry in the list.

Explanation: You did not make a selection in the list.

User response: Select at least one entry in the list.

DBA3015I There are no referenced columns associated with this access plan.

Explanation: There were no columns referenced in the query for the referenced table.

User response: None.

DBA3016I There are no page fetch pairs statistics associated with this index.

Explanation: There are no page fetch pairs statistics stored in the catalog table for this index. Refer to the System Catalog Statistics section in the *Administration Guide* for more details.

User response: None.

DBA3017I There are no referenced functions associated with this access plan.

Explanation: This access plan did not require the use of any functions.

User response: None.

DBA3018I There are no column distribution statistics associated with this column.

Explanation: There are no column distribution statistics stored in the catalog table for this column. Refer to the System Catalog Statistics section in the *Administration Guide* for more details.

User response: None.

DBA3019I There are no indexes associated with this table.

Explanation: There were no indexes defined on the referenced table at the time of explain and there are no indexes currently defined in the table.

User response: None.

DBA3020E An error occurred while a COMMIT against the database was being attempted and no SQLCODE can be returned.

Explanation: An unexpected error occurred.

User response: Contact your system administrator.

DBA3021E An error occurred while a ROLLBACK against the database was being attempted and no SQLCODE can be returned.

Explanation: An unexpected error occurred.

User response: Contact your system administrator.

DBA3022E The Get Database Manager Configuration API passed back a non-zero return code.

Explanation: An unexpected error occurred causing the operation to stop.

User response: Contact your system administrator.

DBA3023E The Get Database Configuration API passed back a non-zero return code.

Explanation: An unexpected error occurred causing the operation to stop.

User response: Contact your system administrator.

DBA3024E An error occurred while saving to the file *file*.

Explanation: Possible problems include:

- The file cannot be opened for write.
- The file name does not conform to the file system convention.

User response: Possible solutions include:

- Change the file attribute to read write.
 - Specify the correct file name.
-

DBA3025E An error occurred while retrieving from the file *file*.

Explanation: Possible problems include:

- The file is not readable.
 - The file does not exist.
-

User response: Specify the correct file name.

DBA3026E No printers are installed.

Explanation: There are no printers installed on your machine.

User response: Install a printer on your machine.

DBA3033I There are no referenced table spaces associated with this access plan.

Explanation: This access plan did not reference any database objects contained in a table space.

User response: None.

DBA3034W The snapshot was created using code page *codepage* but is not converted to the code page of the window running the Visual Explain Tool.

Explanation: The snapshot of the access plan graph displayed was created on a code page different from the code page of the window running the Visual Explain Tool. Code page conversion is not successful because the snapshot is too large. The access plan graph is displayed without code page conversion. You may not be able to display the current statistics when you open the statistics windows.

User response: Change the code page of the window running the Visual Explain Tool to the same one used to create the snapshot and display the access plan graph again.

DBA3035W The snapshot was created using code page *codepage* but is not converted to the code page of the window running the Visual Explain Tool. The value of the database manager configuration parameter **ASLHEAPSZ** is too small.

Explanation: The snapshot of the access plan graph displayed was created on a code page different from the code page of the window running the Visual Explain Tool. Code page conversion is not successful because the user defined function used to convert between the code pages do not have enough memory to run. The access plan graph is displayed without code page conversion. You may not be able to display the current statistics when you open the statistics windows.

User response: Possible solutions include:

- Change the database manager configuration parameter **ASLHEAPSZ** of your server to 512 and display the access plan graph again. Note that the configuration parameter may not take effect until the next time you do a **db2start** on your server.
-

- Change the code page of the window running the Visual Explain Tool to the same one used to create the snapshot and display the access plan graph again.

DBA3036I Explain tables have been created to process your dynamic explain request.

Explanation: One or more explain tables have been created under the current user ID. These tables will be used to store the information needed by your dynamic explain request.

User response: None.

DBA3037E The explain snapshot has been corrupted. Your request cannot be completed.

Explanation: The explain snapshot that is requested has been corrupted in the explain tables.

User response: Regenerate the snapshot.

DBA3038I There are no referenced buffer pools associated with this access plan.

Explanation: This access plan did not reference any database objects contained in a buffer pool.

User response: None.

DBA3039E Explain tables cannot be found.

Explanation: Explain tables do not exist.

User response: Create explain tables using one of the following methods:

- Run the EXPLAIN.DDL sample command file. This file is located in the sqllib\misc directory. To run the command file, go to this directory and issue the **db2 -tf EXPLAIN.DDL** command.
- Call the SYSPROC.SYSINSTALLOBJECTS procedure.

DBA3040E The command is specified incorrectly.

Explanation: The command has invalid syntax.

User response: Issue the command with the -h option to see the correct syntax.

DBA3041E The name of the database must be between 1 and *maximum* characters long.

Explanation: The specified database name is either empty or too long.

User response: Reissue the command with a valid database name.

DBA3042E The explainable statement text must be between 1 and *maximum* characters long.

Explanation: The specified explainable statement text is either empty or too long.

User response: Reissue the command with valid explainable statement text.

DBA3043E The query tag cannot have more than *maximum* characters.

Explanation: The specified query tag is too long.

User response: Reissue the command with a valid query tag.

DBA3044E The user ID cannot have more than *maximum* characters.

Explanation: The specified userid is too long.

User response: Reissue the command with a valid userid.

DBA3045E The valid range for query number is from 0 to *maximum*.

Explanation: The specified query number is either too small or too large.

User response: Reissue the command with a valid query number.

DBA3046E The valid values for optimization class are 0, 1, 2, 3, 5, 7 and 9.

Explanation: The specified optimization class is not one of the valid values.

User response: Reissue the command with a valid optimization class.

DBA3047E The database *name* cannot be found.

Explanation: The database name specified in the command does not exist in the system database directory.

User response: Reissue the command with an existing database name or catalog the database in the system database directory.

DBA3059I There is no details information associated with the selected operator *operator*.

Explanation: The details information of the selected operator is not available for the current release of the DB2 Administration Tools.

User response: None.

DBA3060I **There is no help information associated with the selected operator *operator*.**

Explanation: The help information of the selected operator is not available for the current release of the DB2 Administration Tools.

User response: None.

DBA3061I **There are no column groups associated with this table.**

Explanation: There were no column groups defined on the referenced table at the time of explain and there are no column groups currently defined in the table.

User response: No response required.

DBA3062N **Generate DDL cannot be performed on system tables.**

Explanation: Tables with schema SYSIBM, SYSCAT, or SYSSTAT are system tables.

User response: Perform generate DDL on non-system tables.

DBA3063N **One or more system tables have been selected for Generating DDL. DDL will be generate for selected user tables only. Do you want to continue?**

User response: Tables with schema SYSIBM, SYSCAT, and SYSSTAT are system tables.

User response: Click Yes to generate DDL for non-system tables only. Click No to abort Generate DDL operation.

DBA3064N **Only the first *number* tables will be processed by db2look.**

Explanation: The db2look -t option will only process *number* tables.

User response: Move the list of selected tables to the list of available tables.

DBA3065E **Snapshot cannot be processed due to a CLI error.**

Explanation: A CLI error has occurred during snapshot processing. An access plan cannot be displayed.

User response: Check the CLI configuration by issuing the following command:
db2 get cli cfg
for section <db-name>

If LONGDATACOMPAT is set to 1, try cataloging the database with a different alias:
catalog db <db-name> as <db-alias-name>

Set LONGDATACOMPAT=0 for the database alias:
db2 update cli cfg for section <db-alias-name> using longdatacompat 0

Submit the query for EXPLAIN in the alias database.

If LONGDATACOMPAT is not set to 1, or the parameter is not being set in the CLI configuration, contact IBM Service.

DBA3066E **Statement text for a snapshot contains an empty string.**

Explanation: The statement text of an explain record contains an empty string. Explained statement history cannot be retrieved successfully.

User response: Check the CLI configuration by issuing the following command:
db2 get cli cfg
for section <db-name>

If LONGDATACOMPAT is set to 1, try cataloging the database with a different alias:
catalog db <db-name>
as <db-alias-name>

Set LONGDATACOMPAT=0 for the database alias:
db2 update cli cfg
for section <db-alias-name> using longdatacompat 0

Submit the query for EXPLAIN in the alias database.

If LONGDATACOMPAT is not set to 1, or the parameter is not being set in the CLI configuration, contact IBM Service.

Chapter 76. DBA4000 - DBA4499

DBA4000W Are you sure you want to stop monitoring?

Explanation: You have requested to stop monitoring all database objects that are currently monitored and to exit the Performance Monitor.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4001I There is no monitor summary for your selection.

Explanation: You requested to show monitor activity on a level of object in which nothing is being monitored. There is no summary activity to be shown.

User response: Retry the action when a monitor is started on a object that corresponds to the level you have selected.

DBA4002W Are you sure you want to stop monitoring *name*?

Explanation: You have requested to stop monitoring an object that is currently being monitored.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4003I Only the default level is changed for IBM supplied monitors.

Explanation: For an IBM-supplied monitor, the only attribute that you can change is the level for which this monitor is designated as the default.

User response:

DBA4004E Multiple selection is only allowed for the Remove action.

Explanation: You selected more than one monitor from the list of monitors. The only action you can perform on multiple objects is Remove.

User response: Deselect all monitors except one from the list and retry the action.

DBA4005W Are you sure you want to remove the selected monitors?

Explanation: You requested to remove the selected monitors from the list.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4006W Are you sure you want to reset the counters?

Explanation: You requested to reset the database counters.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4007W Are you sure you want to stop monitoring all objects in this level?

Explanation: You have requested to stop monitoring all objects in this level that are currently being monitored.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4008W Are you sure you want to stop monitoring all database partitions in this object?

Explanation: You have requested to stop monitoring all database partitions in this object that are currently being monitored.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4009I The IBM supplied monitors are being created.

Explanation: The IBM supplied monitors are being created. This may take a while.

User response:

DBA4010E The Performance Monitor encountered a problem and cannot continue. Retry the action. If the problem persists, take a trace at the client and the server, and contact support personnel.

Explanation: Unexpected error has occurred.

User response: Take a trace at the client and the server and contact support.

DBA4011E The program *name* is not a valid name. The program cannot be run. Check the name and retry the action.

Explanation: You have specified an invalid program name.

User response: Specify a valid program name and retry the action.

DBA4012E The program *name* cannot be run. The currently executing thread is not allowed to execute the specified program.

Explanation: A security exception is thrown when executing the program specified. Run command works in application mode only.

User response: Retry the action when you run Control Center as an application.

DBA4013I Database manager node *name* is down.

Explanation: From the Tools settings, node status page, you have selected to be informed about database manager node status when it is down.

User response: Restart the database manager node to continue processing.

DBA4014I The state of database manager node *name* is unknown.

Explanation: From the Tools settings, node status page, you have selected to be informed about database manager node status when the state is unknown.

User response:

DBA4015E *name* contains invalid input. Please re-enter the entry.

Explanation: Your entry contains invalid character.

User response: Retry your action with new entry.

DBA4016E New monitor cannot be created. The maximum number of monitors is *number*.

Explanation: The create or save as or copy request fails. The maximum number of monitors has been reached.

User response: Remove unused monitors and retry your action.

DBA4017E The monitoring connection cannot be established. The maximum number of monitoring connections is *number*.

Explanation: The maximum number of monitoring connections has been reached.

User response: Retry your action later or stop any running monitoring session on your machine.

DBA4018W Your entry in *name* is too long. The maximum number of characters is *number*. Your entry is truncated.

Explanation: Your entry has exceeded the limit.

User response:

DBA4019E Your entry in *name* is too long. The maximum number of characters is *number*. Retry your action with valid input.

Explanation: Your entry has exceeded the limit.

User response: Retry the action with valid input.

DBA4020W You made changes to the performance monitor *name*. Do you want to save your changes?

Explanation: You made one or more changes to the performance monitor that are not yet saved, and you requested to exit from the Show Monitor window or set the running monitor a different monitor. Your changes will be lost unless you click Yes.

User response: Click "Yes" to process the request or "No" to cancel it.

DBA4021W You made changes to the performance variable settings for *name*. Do you want to apply your changes to the next snapshot?

Explanation: You made one or more changes to the performance variable settings that are not yet saved, and you selected a different performance variable to work with. Your changes will be lost unless you click Yes.

User response: Click "Yes" to apply the changes to the next snapshot of "No" cancel it.

DBA4023E You entered a combination of alarm and warning threshold values that is not valid for performance variable *name*. Change one or more threshold values and try again.

Explanation: The upper alarm threshold value must be greater than the upper warning threshold value and so on; that is, upper alarm threshold value > upper warning threshold value > lower warning threshold value > lower alarm threshold value. The threshold value is a float no longer than 9 digits.

User response: Verify the threshold values and try the operation again.

DBA4024W The Summary page cannot show more than *number* columns. Turn to the Details page to see all the performance variables.

Explanation: The Summary page currently displays as many columns as it can accommodate, but some performance variables are not shown.

User response: Turn to the Details page to see all the performance variables.

DBA4025W The *type* view cannot show more than *number* performance variables.

Explanation: The maximum number of performance variables are shown.

User response: Remove one or more performance variables from the *type* view before adding more.

DBA4026E A non-applicable monitor has been set as default for this level.

Explanation: The monitor must include at least one performance variable for this level.

User response: Modify the monitor to include at least one performance variable for this level.

DBA4027I Monitoring has not been started as there is nothing to monitor.

Explanation: Currently there are no databases or gateway connections to be monitored for this instance.

User response: Catalog a database or gateway connection to start monitoring.

DBA4060E The Performance Monitor is currently active with another monitor for this instance.

Explanation: Only one monitor can be active at a time per instance. For example, if your default monitor for databases is different from your default monitor for connections, you will not be able to monitor databases and monitor (or list) connections at the same time.

User response: In the List Monitors window, you can:

1. Find out which monitor is running and stop it so you can specify a different one.
 2. Find out which monitor is running and specify it for your next monitoring action.
 3. Create a monitor that contains performance variables for all levels, and then select it as your default monitor for each level.
-

DBA4065E The threshold value cannot be blank or contain more than 9 characters. Enter a valid threshold value.

Explanation: You entered a threshold value which is not valid or your value has been modified internally. For example, 123456789 will be modified to 123,456,789.0 which exceeds the number of characters allowed in the threshold.

User response: Re-enter the threshold value and retry your action.

DBA4070E The JDBC server encountered a problem. The list of performance monitors could not be processed.

Explanation: An unexpected problem is encountered.

User response: Verify that the JDBC server is up and running. If it is, take a trace at the JDBC server and then contact support personnel.

DBA4071W Some monitors are currently in use locally. The refreshed list shows local monitors only.

Explanation: The list of monitors is refreshed with local monitors only.

User response: To show the complete list, stop the active monitors and retry the operation.

DBA4072E This monitor name already exists. Specify a unique name.

Explanation: Performance monitor names in a JDBC server must be unique regardless of the monitor's creator.

User response: Specify a unique monitor name.

DBA4074E The JDBC server encountered a problem. The operation could not be performed.

Explanation: An unexpected problem is encountered.

User response: Verify that the JDBC server is up and running. If it is, take a trace at the JDBC server and then contact support personnel.

DBA4075E This performance monitor is empty so it cannot be saved. Add at least one performance variable to the monitor and retry the operation.

Explanation: Monitors need to contain at least one performance variable in order to be valid.

User response: Add at least one performance variable to the monitor and retry the operation.

DBA4076E The action cannot be performed because you are not the creator of the monitor.

Explanation: You must own the monitor to be able to change, rename or remove it.

User response: If you were trying to change a monitor, copy it under a different name and then change it.

DBA4079E This monitor is currently in use locally. Stop the monitor, and then remove or rename it.

Explanation: You cannot remove or rename an active monitor.

User response: Stop the monitor before removing or renaming it.

DBA4080E This monitor is currently in use remotely. Wait till the monitor is stopped, and then remove or rename it.

Explanation: You cannot remove or rename an active monitor.

User response: Wait till the monitor is stopped before removing or renaming it.

DBA4083E This monitor could not be found. Another user might have removed it. Stop all monitoring and then refresh your list of monitors.

Explanation: Another user might have removed the monitor that you are using.

User response: Stop all monitoring and refresh your list of monitors.

DBA4090E Connection to the JDBC server has ended abnormally. Close the Control Center and retry the operation. If the problem persists, take a trace at the JDBC server and contact support personnel.

Explanation: An unexpected error has occurred.

User response: Close the Control Center and retry the operation. If the problem persists, take a trace at the JDBC server and contact support personnel.

DBA4095I This function has not been implemented on this object.

Explanation: This function has been planned for a future release but is not currently implemented on this object.

User response: None.

DBA4220I Full database restore cannot be done from a table space backup image.

Explanation: Full database restore cannot be done from a table space backup image.

User response: None.

DBA4221I Table space selection is not allowed for the media type.

Explanation: Table space selection is not allowed for the media type.

User response: None.

DBA4222I Table space selection is not allowed since the database is in backup pending state.

Explanation: Table space selection is not allowed since the database is in backup pending state.

User response: None.

DBA4223I Full database required when backup image information is entered manually.

Explanation: Full database required when backup image information is entered manually.

User response: None.

Chapter 77. DBA4500 - DBA4999

DBA4730E Performance monitoring encountered a problem while accessing the monitor server. Monitoring cannot continue. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Contact your technical service representative with the smcode.

DBA4731E A resource problem occurred on the Monitor Server. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Free up some server resource and retry the operation. If the problem persists, contact your technical service representative with the smcode.

DBA4732E A problem occurred when trying to attach to the database instance. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Retry the operation. If the problem persists, contact your technical service representative with the smcode.

DBA4733W The monitored database instance is stopped. Restart the instance. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Restart the instance. If the problem persists, contact your technical service representative with the smcode.

DBA4734E Performance monitoring encountered a problem while accessing the monitor server. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Retry the operation. If the problem persists, contact your technical service representative with the smcode.

DBA4735E A problem occurred when trying to attach to the database instance. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Contact your technical service representative with the smcode.

DBA4736E Unable to determine the current territory code or current code page in use on the client. **smcode:** *Error-code*

Explanation: The client's territory code and code page must be sent to the server so that the data is returned in the correct code page. The Monitor Server was unable to determine the active territory code or code page.

User response: Refer to your operating system documentation to determine how to set the territory code and active code page for your system. Retry the operation after setting the territory code and code page.

DBA4737E Unable to set the current client territory code or current code page on the server. **smcode:** *Error-code*

Explanation: The client's territory code and code page must be sent to the server so that the data is returned in the correct code page. The Monitor Server was unable to set the client's active territory code or code page on the server.

User response: Refer to your operating system documentation to determine how to activate the client's territory code and active code page on your server's system. Retry the operation after activating the territory code and code page.

DBA4738E The Administration Server on system *name* is not started. **smcode:** *Error-code*

Explanation: The listed Administration Server on system *name* is not started.

User response: Start the Administration Server and retry the operation.

DBA4739E The client system's code page could not be determined. The client system might not be set up correctly. **smcode:** *Error-code*

Explanation: This is a client system setup problem.

User response: Contact IBM Service.

DBA4740E The client code page *name* is not available on instance *name*. **smcode:** *Error-code*

Explanation: The listed code page has not been installed on instance *name*.

User response: Contact your system administrator. Your administrator should install code page *name* on instance *name*.

DBA4741E The client code page *name* is not available on instance *name*, at node *name*. **smcode:** *Error-code*

Explanation: The listed code page has not been installed on instance *name*, at node *name*.

User response: Contact your system administrator. Your administrator should install code page *name* on instance *name* at node *name*.

DBA4742E The Administration Server *name* cannot be found. **smcode:** *Error-code*

Explanation: Possible reasons for this problem include: 1) Remote instance name is not valid. 2) Hostname is not valid.

User response: Contact your system administrator.

DBA4743E The Administration Server *name* detected a problem. **sqlcode:** *Error-code*. **smcode:** *Error-code*

Explanation: The Administration Server *name* had an internal error.

User response: Contact your system administrator.

DBA4744E The Control Center could not communicate with the parallel tools on instance *name* at node *name*. Try starting communications. **smcode:** *Error-code*

Explanation: Communications may not have been started.

User response: Contact your system administrator.

DBA4745E The TCP/IP port configuration between the Control Center and its listener (db2cclst) on instance *name* at node *name* is not correct. **smcode:** *Error-code*

Explanation: The /etc/services file at node *name* might contain incorrect information.

User response: Contact your system administrator. Your administrator should check that the port name of db2ccmsrv is entered correctly in the /etc/services file on node *name*.

DBA4746E Instance node host name *name* is not valid on instance *name*. **smcode:** *Error-code*

Explanation: The db2nodes.cfg file for instance *name* might contain incorrect information.

User response: Contact your system administrator. Your administrator should validate the host name found in the db2nodes.cfg file for instance *name*.

DBA4747E The Performance Monitor did not receive a response from node *name* within the snapshot capture interval. **smcode:** *Error-code*

Explanation: A performance or communications problem on the listed node is preventing a response from being sent, or the snapshot capture interval is not long enough to allow each node time to respond to the request.

User response: Contact your system administrator. Your administrator can 1) Validate the performance on the listed node. 2) Increase the snapshot capture interval to allow more time for the node to respond to the snapshot request.

DBA4748E Logical node *name* on instance *name* is not valid. **smcode:** *Error-code*

Explanation: The db2nodes.cfg file for instance *name* might contain incorrect information.

User response: Contact your system administrator.

DBA4749E The Monitor Server on instance *name* is unable to establish communications with node *name*. **smcode:** *Error-code*

Explanation: The communications subsystem may not have been started.

User response: Start the db2cclst process on the specified node. If problem still persists, contact your system administrator.

DBA4750E The instance being switched to does not exist. **smcode:** *Error-code*

Explanation: See the Administration Server error logs for more information.

User response: Contact your technical service representative with the smcode.

Chapter 78. DBA5000 - DBA5499

DBA5006E *evname* is an Event Monitor that writes to a named pipe. Event Analyzer does not support PIPE Event Monitors.

Explanation: The Event Analyzer only handles traces produced by FILE Event Monitors.

User response: Use a FILE Event Monitor.

DBA5007E Event Monitor *event-monitor* does not exist.

Explanation: The Event Monitor names specified with the -evm option could not be found in the catalogs for the database specified with the -db option. The Event Monitor may have been dropped or you may be trying to connect to the wrong database.

User response: Make sure that the Database alias specified with the -db is properly cataloged and that the Event Monitor has not been dropped. If the latter is the case, re-create the event monitor.

DBA5250I The number of records retrieved exceeded the limits of the window. A partial list is displayed.

Explanation: The window height required to display the retrieved objects exceeds a system limit.

User response: You can reduce the number of objects displayed in the window by doing the following:

- If the "Include" choice is available on the "View" menu, you can specify a subset of the objects to display in the window. When you select "Include", a window appears and prompts you for a sub set of the criteria.

DBA5300E Event Analyzer invocation error.
Usage:
db2eva [- db database-alias
-evm evmon-name]

Explanation: The command parameters for the db2eva command are:

-db database-alias

specifies the database for which the Event Monitor is defined, as cataloged on the machine where the trace is analyzed.

-evm evmon-name

the Event Monitor trace tables you want to analyze. The event monitor must be defined in the database specified in the -db parameter.

If you specify the database and event monitor names, db2eva connects to the database, and issues a 'select target from sysibm.syseventmonitors' to locate the table where the Event Monitor writes its trace. The Event Monitor and related tables must not have been dropped when using this mode.

If you do not specify any parameter, a dialog box will be opened to prompt for the database alias and the event monitor name you want to analyze.

User response: Re-issue the command with valid arguments.

DBA5311E Event Monitor invocation error.
Usage: db2emcrt database-alias

Explanation: The command parameter for the db2emcrt command is:

database-alias

specifies the database for which the event monitors are to be created or analyzed, as cataloged on the machine the db2emcrt command is invoked. database-alias overrides the database name specified in the trace.

User response: Re-issue the command with valid arguments.

DBA5350E An error occurred while opening file *file-name*.

Explanation: The specified file could not be opened in read-only mode for an unknown cause.

User response: Check the spelling of the directory specified by the -path option, its access authority, and that it contains a non-empty, readable trace file called 00000000.evt.

NOTE: The first trace file for an Event Monitor is always named 00000000.evt and cannot be renamed.

DBA5351E File *file-name* cannot be found.

Explanation: There is no 00000000.evt file in the directory specified by -path.

User response: Check the spelling of the directory specified by the -path option, its access authority, and that it contains the trace file 00000000.evt.

NOTE: The first trace file for an Event Monitor is always named 00000000.evt and cannot be renamed.

DBA5352E Path *path-name* cannot be found.

Explanation: The -path option specifies a directory that does not exist.

User response: Check the spelling of the directory specified by the -path option, its access authority, and that it contains the trace file 00000000.evt.

NOTE: The first trace file for an Event Monitor is always named 00000000.evt and cannot be renamed.

DBA5353E Access to *param-1* was denied.

Explanation: The specified file could not be opened in read-only mode. The -path option may be specifying a directory to which you do not have sufficient access authority.

User response:

- Check the access authority to the directory specified by the -path option.
- Verify that no one has locked the specified file in exclusive mode.

DBA5354E The data cannot be read from *file-name*.

Explanation: The Event Monitor trace file contains unexpected data, or cannot be accessed.

The trace file was likely corrupted during transmission or has been removed.

User response: Transmit your trace files from your server again. When transmitting trace files from a remote server, ensure that the transmission is done in binary mode.

DBA5355E An Event Monitor log header could not be found in *file-name*.

Explanation: The first file written to by an Event Monitor is named 00000000.evt and contains a data structure that identifies the trace characteristics. This data structure could not be read. Possible causes:

- The trace file was corrupted.
- The trace file is empty. This can happen if the Event Monitor is active but, has not flushed its buffers yet.

User response:

- If the trace was transmitted from a remote server, retransmit, ensuring that the transmission is done in binary mode.
- If the trace file is empty, turn off the Event Monitor to force it to flush its buffer.

DBA5356E Byte order not valid in *file-name*.

Explanation: The Event Monitor log header (the first record written to a new trace) specifies whether the trace contains little-endian or big-endian (e.g., AIX) data. The value found in the trace file specified neither of these two supported types of data.

The trace file was likely corrupted during transmission.

User response: Transmit your trace files from your server again. When transmitting trace files from a remote server, ensure that the transmission is done in binary mode.

DBA5357E Unexpected data was found at offset *offset-value* in file *file-name*.

Explanation: The Event Monitor trace file contains unexpected data.

The trace file was likely corrupted during transmission.

User response: Retransmit your trace files from the server, ensuring that the transmission is done in binary mode.

DBA5358I Cannot browse local file system without an active local Administration Server.

Explanation: The file browser requires a Database Administration Server (DAS) instance to browse file systems. If this is a client installation, the DAS instance is not present and local file browsing is disabled.

User response: For server installations, ensure the DAS instance is started and retry the operation.

Chapter 79. DBA5500 - DBA5999

DBA5500E A system error has occurred. The Event Analyzer tool could not continue processing.

Explanation: For an unknown reason a system error has occurred when trying to initialize the application environment.

User response: Re-issue the command.

DBA5501W Do you want to delete the event files associated with the event monitor *evname* that you are removing?

Explanation:

User response: To delete the event files for the selected event monitor, click on OK. If you do not want to delete the event files, but still want the event monitor removed, click on NO. If you do not want to remove the event monitor, or delete the event files click on CANCEL.

DBA5502E No selection was made from the list of Event types. You must select one or more types from the list.

Explanation: When creating an event monitor, at least one Event type must be selected.

User response: Select one or more of the Event type check boxes from the Create Event Monitor window.

DBA5503I An Event monitor *evname* has been successfully created.

Explanation: The event monitor has been successfully created and has been added to the database list of event monitors.

User response: Click on OK to remove the message.

DBA5504W Some event monitors were not removed because their event files could not be deleted. See the messages page from the journal notebook for the list of event monitors not removed.

Explanation: Some of the selected event monitors may have been successfully removed, but others were not, because their associated event trace files could not be deleted.

This occurs when either event files are read-only, or they are being used by another active process.

User response: See the message page of the journal notebook for the list of event monitors that were not

removed. To delete the files, end the process using them and check the access authority for the files. Then select the event monitor from the list of event monitors and choose remove from the menu.

DBA5505E Maximum file size was not specified.

Explanation: If maximum file size in 4K pages is selected on the create event monitor Options window, then the maximum size of the event files must be specified.

User response: Enter a maximum file size of between 1 and 2,147,483,647 4K pages, or you can select the No Maximum option. If the No Maximum option is selected, the maximum number of event files is 1.

DBA5506E Maximum number of files was not specified.

Explanation: If Maximum number of files is selected on the create event monitor Options window, then the maximum number of event files must be specified.

User response: Enter a maximum number of event files of between 1 and 2,147,483,647, or you can select No Maximum option.

DBA5507E Buffer size was not specified.

Explanation: An event monitor Buffer size in 4K pages must be specified.

User response: You must enter an event monitor buffer size of between 1 and 2,147,483,647 4K pages.

DBA5508W Do you want to delete the event files associated with the event monitors that you are removing?

User response: To delete the event files for the selected event monitors, click on OK. If you do not want to delete the event files, but still want the event monitors removed, click on NO. If you do not want to remove the selected event monitors, or delete the event files click on CANCEL.

DBA5509E The name entered for the event monitor is not valid.

Explanation: The name entered in the event monitor Name field contains one or more characters that are not allowed. An event monitor can contain alphanumeric characters and must begin with a letter.

User response: Remove from the name any characters

that are not alphanumeric and make sure the name begins with a letter.

DBA5510E When trying to activate the monitor, the directory for its data files was not found.

Explanation: When activating an event monitor, the directory that was specified when the event monitor was created must exist. In order to flush its buffers, an event monitor needs an existing file and path.

User response: Create the directory that was specified when creating the event monitor. Try to activate the event monitor again.

DBA5511E Entry *name* is not valid because the value is out of the acceptable range.

Explanation: The values for an event monitor's Maximum file size, Maximum number of files, and Buffer size must be between 1 and 2,147,483,647.

User response: Enter a value in the entry fields that is within the given range.

DBA5512E The character *character* entered for the event monitor name is not allowed.

Explanation: A character entered in the event monitor name entry field is not a valid character for an event monitor name.

User response: Remove from the event monitor name any characters that are not alphanumeric and make sure the name begins with a letter.

DBA5513E The dynamic link library "DB2EVA.DLL" could not be loaded successfully.

Explanation: During the initialization of the event analyzer the dll could not be loaded because the dll has been deleted or removed.

User response: Reinstall the event analyzer by selecting the performance monitor during the installation process.

DBA5514I The event files for event monitor *evname* could not be deleted.

Explanation: The event files for the event monitor could not be deleted because the files are being used by an active process or are read-only.

User response: To delete the files, end the process that is using them and check the access authority for the files. Select the event monitor again and choose remove from the menu.

DBA5515E The database alias name or database name *dbname* could not be found.

Explanation: The alias name that was entered on the command line has not been cataloged on this machine.

User response: Either catalog the alias on this machine and reenter the command or re-issue the command with a valid alias name.

DBA5516W Viewing the trace for this event monitor may take some time. Do you wish to continue processing?

Explanation: The user has requested to view a large trace for an event monitor. Reading the files could take some time to complete.

User response: If you want to continue processing the trace click on OK. If you do not wish to continue processing the trace file then click on NO.

DBA5517E The event monitor could not be started due to a reason code of *reason-code*.

Explanation: The event monitor could not be started for the reason specified by *reason-code*:

1. The event monitor path is invalid.
2. The device, where the event monitor is to be started, is not ready.
3. Attachment to the admin server failed.

User response: To start the event monitor, make sure that the event monitor path is valid, the admin server started, and the proper authorities exist on the device. Click on OK.

If the problem persists, manually create the event monitor directory on the device and click on OK.

To create the event monitor without starting it, deselect the Start now check box and click on OK.

DBA5518E The event monitor could not be started because its output directory could not be created.

Explanation: The directory could not be created for one of the following reasons.

1. Proper authorities do not exist for creating directories on the requested file system.
2. Create event monitor does not support directory creation on non-Intel platforms.
3. Cannot create a directory of greater than 8 characters on a FAT file system.

User response: Make sure the proper authorities exist for directory creation.

On non-Intel platforms create the event monitor output directory manually.

DBA5519I **Event monitor table cannot be found in the event source database.**

Explanation: The event monitor table that holds the event monitor records cannot be found in the source database specified.

User response: Ensure that the event monitor table exists in the source database and the table name is in the SYSIBM.SYSEVENTTABLES.

DBA5520E **Event monitor control table cannot be found in the event source database.**

Explanation: The event monitor control table cannot be opened.

User response: Ensure that the event monitor control table exists in the source database and the table name is in the SYSIBM.SYSEVENTTABLES.

DBA5521I **Connection to the database is not established. The event monitor list cannot be retrieved.**

Explanation: A connection to the database is not established.

User response: Please make sure the database manager command has been issued. If the problem persists, contact IBM Service.

Chapter 80. DBA6000 - DBA6499

DBA6001I The replication subscription set, *subscription-set*, is in use. Please try again later.

Explanation: The Apply program is currently processing the replication subscription and this action cannot be deactivated. Wait until the Apply program has completed its processing and retry your command.

User response: No action is required.

DBA6002I The replication source, *source-name-1.source-name-2*, was defined successfully. To begin capturing changes to the replication source, start the Capture program from the command line.

Explanation: The Capture program requires a defined replication source in order to run.

User response: Make sure the Capture program is running.

DBA6003I The replication subscription was defined successfully. To begin replicating the subscription set, ensure that the Apply program has been started from the command line.

Explanation: The Apply program requires a defined replication source in order to run.

User response: Make sure the Apply program is running.

DBA6004I The replication subscription was successfully changed.

Explanation: This is for your information only.

User response: No action is required.

DBA6005I The replication sources were successfully removed.

Explanation: This is for your information only.

User response: No action is required.

DBA6006I The replication source, *source-name*, could not be removed.

Explanation: This is for your information only.

User response: No action is required.

DBA6007I The replication subscription, *subscription-name*, could not be removed.

Explanation: This is for your information only.

User response: No action is required.

DBA6008I The replication subscriptions were successfully removed.

Explanation: This is for your information only.

User response: No action is required.

DBA6009I *file-name* is locked by another action.

Explanation: The file is currently being used by another action.

User response: Wait until the action is completed and retry your command.

DBA6010I The join was successfully defined.

Explanation: This message is for your information only.

User response: No action is required.

DBA6012I The replication source was defined successfully. To begin capturing changes to the replication source, ensure that the Capture program has been started from the command line.

Explanation: The Capture program requires a defined replication source in order to run.

User response: Make sure the Capture program is running.

DBA6013I The replication subscription was successfully updated.

Explanation: This message is for your information only.

User response: No action is required.

DBA6014I The replication subscription was successfully cloned.

Explanation: This message is for your information only.

User response: No action is required.

DBA6015I The SQL statements for this action have been saved in a file for future editing and execution.

Explanation: This message is for your information only.

User response: No action is required.

DBA6016I The SQL script file has been run successfully.

Explanation: This message is for your information only.

User response: No action is required.

DBA6100I The Apply program will only perform full refresh copies to this set.

Explanation: This message is for your information only.

User response: No action is required.

DBA6101E The table, *table-name-1.table-name-2*, cannot be found.

Explanation: The control table specified could not be located in the database or was accidentally dropped.

User response: Create the control table. See the Tables Chapter for more details.

DBA6102I This action may take some time to complete. Please wait...

Explanation: This message is for your information only.

User response: No action is required.

DBA6103E The SQL file, *file-name*, cannot be run.

Explanation: The user may have edited the SQL statement file for a replication action, such as defining a replication source, and there is an error in the SQL statement file.

User response: Correct the error in the SQL statement file and run the file again.

DBA6104W No unique prefix character was found. The prefix character will be set to NULL.

Explanation: A prefix character is an alphabetical or numerical character used to distinguish before-image columns in the source table. If a prefix character cannot be assigned because all possible prefix characters are already in use, you will receive this warning. The change data table cannot be created.

User response: Uncheck the before-image columns and rerun the action.

DBA6105E The field, *field-name*, requires your input.

Explanation: The user has not filled in a required field.

User response: Specify a parameter and select OK.

DBA6106E The replication action does not currently support this database system.

Explanation: The replication action must support your database system in order to run.

User response: Contact your IBM Service representative.

DBA6107E Please enter a different value for the field, *field-name*.

Explanation: When cloning a replication subscription, the **target server** and **Apply qualifier** fields must be different from the original **target server** and **Apply qualifier** fields.

User response: Enter a different target server name or Apply qualifier and select OK.

DBA6108E The action cannot be completed due to a syntax error.

Explanation: The SQL statement contains a syntax error.

User response: Check the SQL reference for the proper syntax. For subscription actions, check the where clause, the rows page, or the CREATE column. For define-join actions, check the CREATE VIEW statement.

DBA6109E SHOW COLUMNS action cannot be completed.

Explanation: The SHOW COLUMNS action requires that a table be selected.

User response: Select a table from the Selected Table list box and then press the SHOW COLUMNS key.

DBA6110E The SQL file, *file-name*, cannot be opened.

Explanation: The SQL file either does not exist, is in the wrong subdirectory, or is engaged in another process.

User response: Make sure you enter the correct file name and that the file is in the correct subdirectory. If the file is engaged in another process, terminate the process and reenter the file name.

DBA6111E The file, *file-name*, cannot be opened.

Explanation: The file either does not exist or is in the wrong subdirectory.

User response: Make sure you enter the correct file name on the Replication page of the Tool Settings notebook.

DBA6112I You have chosen to substitute user-defined values to define replication objects in the Tool Settings notebook. Please supply these values in the file *file-name*.

Explanation: This message is for your information only.

User response: No action is required if you have already supplied user-defined values in the file.

DBA6113E *keyword* is an invalid keyword in file *file-name*.

Explanation: The specified keyword is invalid and may be a typographical error.

User response: Use the keywords as given in the default file "DPREPL.DFT".

DBA6114E The value for keyword *keyword* cannot be found in file *file-name*.

Explanation: The value of this keyword is required for this action to complete.

User response: Specify a value for this keyword in the file.

DBA6115E None of the items in the column list were selected. This action cannot be completed.

Explanation: At least one item in the column list must be selected for this action to complete.

User response: Select one or more items from the column list and press OK.

DBA6116E The value for keyword *keyword* is invalid in file *file-name*.

Explanation: The value of this keyword is either too long, an incorrect data type, or in an incorrect format.

User response: Correct the value and invoke the action again.

DBA6117W The before-image column *column-list* was truncated to *number* bytes because the length of the original column was at the database limit before being appended to the before-image prefix *before-image-prefix*.

Explanation: The length of the column from the original source table was at the database limit. When appended to the before-image prefix, some characters at the end of the column name were truncated to fit the database limit. This new before-image column name may no longer be unique because the truncated characters may be needed to guarantee a unique name.

User response: Uncheck the before-image column when the truncated name is not unique and rerun the action.

DBA6118W Datatypes *datatype-list* are not supported on the database *database*. The datatypes will be converted to *datatype-list* in the SQL script.

Explanation: Some datatypes from the source are not supported on the named database. However, equivalent datatypes could be found on the named database. The generated SQL statements will use the converted datatypes in place of the unsupported datatypes.

User response: No action is necessary if the converted datatypes are acceptable. If they are not acceptable, uncheck the target columns and rerun the action.

DBA6119W Datatypes *datatype-list* are not supported on the database *database*. The columns containing these datatypes will be excluded from the SQL script.

Explanation: Some datatypes from the source are not supported on the named database. No equivalent datatypes could be found on the named database. The generated SQL statements will exclude any columns containing the unsupported datatypes.

User response: Contact IBM service if you would like to recommend a specific datatype be supported on the named database in the future.

DBA6120W The source table *table-name* was not defined with a primary key. Specify a primary key for the target table in the Target Columns page.

Explanation: The define subscription action could not find a qualifying primary key for the source table, therefore the define subscription could not determine an appropriate primary key for the target table.

User response: Specify a primary key for the target table in the Target Columns page of the Advanced notebook.

DBA6121W An incompatible Java Runtime Environment has been detected. Please verify that the required level of the Java Runtime Environment has been installed.

Explanation: The installed Java runtime is not supported by the DB2 Java application.

User response: Refer to the Control Center readme.htm for information on prerequisite Java Runtime Environment levels.

DBA6123E The requested action cannot be performed because the schema names of the objects exceed *number* bytes in length.

Explanation: The data replication components support schema names up to 128 bytes in length. You tried to perform a replication action against objects that have longer schema names. The action cannot be performed.

User response: Select objects whose schema names are 128 bytes or less, and perform the action again.

DBA6124W The requested action cannot be performed on some of the objects you selected because the schema names of the objects exceed *number* bytes in length.

Explanation: The data replication components support schema names up to 128 bytes in length. You tried to perform a replication action against some objects that have longer schema names. The action will be performed only on those objects whose schema names are 128 bytes or less.

User response: No action is required.

DBA6125W The field *field*, contains a schema name exceeding *number* bytes in length. The schema name will be truncated to *number* bytes.

Explanation: The data replication components support schema names up to 128 bytes in length. The schema name in the field exceeds 128 bytes and it will be truncated to 128 bytes.

User response: No action is required. However, you can type a shorter schema name in the field, if you do not want the existing schema name truncated.

Chapter 81. DBA7000 - DBA7499

DBA7000E **The license for *product-name* could not be removed. RC = *return-code*.**

Explanation: The license could not be removed. Either the license could not be found in the nodelock file or the nodelock file could not be updated.

User response: Check the permissions of the nodelock file and that the license exists in it.

The location of the file is platform specific:

- AIX - /var/ibm/nodelock
- Windows operating systems - %DB2PATH%/license/nodelock
- All other operating systems - /var/lum/nodelock

DBA7001E **The license could not be added to the nodelock file. RC = *return-code*.**

Explanation: There was an error transmitting the license file or installing the license.

User response: Please check the permissions of the nodelock file. The location of the file is platform specific:

- AIX - /var/ibm/nodelock
- Windows operating systems - %DB2PATH%/license/nodelock
- All other operating systems - /var/lum/nodelock

Alternatively you can also add the license using the command line utility db2licm.

If the problem persists please contact IBM Service.

DBA7002E **The license certificate is invalid. Please try the command again with a valid license file. RC = *return-code*.**

Explanation: The license certificate contained in the file specified does not contain a valid db2 license or is in an unrecognizable format.

User response: Please obtain a valid DB2 license certificate from your IBM representative or authorized dealer.

DBA7003W **The list of current users could not be retrieved. RC = *return-code*.**

Explanation: The daemon or service that tracks the users is not functioning correctly or has not been started. You must have either the concurrent user policy or registered user policy activated.

User response: Turn on either the concurrent user

policy or registered user policy and restart all the instances.

If the problem persists contact IBM Service.

DBA7004E **The registered user could not be removed from the registered user list. RC = *return-code*.**

Explanation: The registered user list on the server could not be updated with the changes you have requested.

User response: Check to make sure the user exists and that this is a valid action for this product. If the problem persists contact IBM Service.

DBA7005E **The registered user could not be added. RC=*return-code*.**

Explanation: The registered user list on the server could not be updated with the changes you have requested.

User response: Check to make sure the username is valid and ensure that the registered user policy is valid for this product. If the problem persists contact IBM Service.

DBA7006E **The license policy could not be updated. RC = *return-code*.**

Explanation: The attempt to update the license policy for this product did not succeed.

User response: Restart the instance and the License Server using the command db2licd -xxx. On a Windows operating system, you can also restart the instance and the License Server using the Services Control Manager.

If the problem persists contact IBM Service.

DBA7007E **The statistics regarding license usage could not be retrieved. RC = *return-code*.**

Explanation: Either the statistics are corrupt, there was a communication error, or no data is available. Data is generated on connects and disconnects after a policy is updated and the instances are restarted.

User response: If the statistics are corrupt you can remove the statistics database. This will remove all previously collected statistics. The location is platform specific and depends on where the product is installed.

- UNIX - DB2 installation directory/misc/db2licst
- Windows operating systems -DB2 installation directory\license\db2licst.dat

DBA7008I The license policy will not take effect fully until the instances are restarted.

Explanation: When the instances are restarted it rereads the license configuration.

User response: Restart the instances on the server.

DBA7009E This product *product* does not support the specified license policy. RC = *return-code*.

Explanation: The license policy specified is not available for use with this product.

User response: Specify a valid license policy.

DBA7010E The License Center was unable to retrieve a list of products from the specified server. RC = *return-code*.

Explanation: The list of products on this server could not be obtained.

User response: Please restart the administration instance and the Control Center and try again. If the problem persists contact IBM Service.

DBA7011E You do not have proper authority to perform the specified action.

Explanation: The user does not have permission to perform this operation.

User response: Please login with a userid that has the proper permissions to perform this command.

DBA7012E No statistics are available for the specified time period.

Explanation: No statistics are available for the time period that you specified.

User response: Please enter a valid date range where statistics exist.

DBA7013E The license certificate *filename* could not be found. Please try the command again with a valid license file. RC = *return-code*.

Explanation: The file specified does not exist or cannot be read.

User response: Please obtain a valid DB2 license certificate from your IBM representative or authorized dealer.

DBA7014E Restart the Control Center and the administration instance and try the command again.

Explanation: Restart the Control Center and the administration instance and try the command again.

User response: If the problem persists please call IBM Service.

DBA7015E The license for DB2 OLAP Server cannot be updated. The DB2 OLAP processes are currently active.

Explanation: DB2 cannot update the license for DB2 OLAP Server while DB2 OLAP Server is running.

User response: To update your OLAP license, please stop all OLAP processes and reinstall this DB2 license.

DBA7016E The selected system does not have any instances specified.

Explanation: A valid DB2 instance for the selected system was not found. Verify that you have properly cataloged at least one valid DB2 instance residing on the selected DB2 system.

User response: Catalog a valid DB2 instance for the selected system, restart Control Center and try the command again.

DBA7017E No valid DB2 instances were found on the selected system.

Explanation: A valid DB2 instance for the selected system was not found. Verify that the selected systems has at least one working DB2 instance.

User response: Create a valid DB2 instance on the selected system and try the command again.

DBA7018E The number of processors on the server exceeds the number of processor entitlements for this product.

Explanation: You are exceeding the number of processors allowed by your license agreement.

User response: Purchase additional processor based entitlements from your IBM representative or authorized dealer and update your license using the db2licm command.

DBA7019E No supported DB2 instances were found on the selected system.

Explanation: A supported DB2 instance for the selected system was not found. Verify that the selected system has at least one DB2 instance supported by this version of DB2 License Center.

User response: Create a supported DB2 instance on

the selected system and try the command again.

DBA7100W The file *file-name* could not be opened. The page *page-name* in the Information Center, which corresponds to that file, will not be shown.

Explanation: The data for the Information Center is stored in a number of NDX files. The file *file-name* might be missing from the correct directory.

User response: Ensure that the NDX files are installed on the file system.

Pages in the Information Center and their corresponding NDX files:

- Concepts - db2booksuc.ndx
- Tasks - db2tasksuc.ndx
- Reference - db2refsuc.ndx
- Troubleshooting - db2msgsuc.ndx
- Sample Programs - db2sampsuc.ndx
- Tutorials - db2webuc.ndx

DBA7101E The Information Center cannot be displayed because none of the Information Center's NDX files can be opened.

Explanation: The Information Center must read at least one NDX file to display any links to DB2 information.

The data for the Information Center is stored in a number of NDX files. The files might be missing from the file system.

User response: Ensure that the NDX files are located in the file system.

Pages in the Information Center and their corresponding NDX files:

- Concepts - db2booksuc.ndx
- Tasks - db2tasksuc.ndx
- Reference - db2refsuc.ndx
- Troubleshooting - db2msgsuc.ndx
- Sample Programs - db2sampsuc.ndx
- Tutorials - db2webuc.ndx

DBA7102W The file *file-name* is not formatted properly. The following lines were ignored: *line-numbers*. Ensure that these lines are formatted correctly.

Explanation: Each line in an NDX file must be in the form of a comma-separated list. Here is an example:

```
"IBM",1,"IBM Home Page",
http://www.ibm.com
```

The format of the list is as follows:

- 1 The name of the entry as it appears in the tree or list of items in the Information Center
- 2 The number representing the category of this entry
- 3 The description that appears on the status line when this item is selected
- 4 The full Web address indicating the location of the entry

User response: If you have edited this file, ensure that each line is formatted properly as follows:

- The name and description must start and end with double quotation marks.
- The name and description cannot contain any double quotation marks.
- Use commas as separators.
- The first entry in the file is a number corresponding to the number of entries in the file. You must update this number manually if you add or remove entries.

DBA7200E A maximum of 10 columns can be selected as input to a geocoder.

Explanation: More than 10 columns were selected as input to a geocoder.

User response: Move column names from the **Selected columns** box to the **Available columns** box until the **Selected Columns** box lists 10 names or fewer.

DBA7201E The database is not enabled to perform Spatial Extender operations.

Explanation: The database must be enabled for Spatial Extender before you can perform Spatial Extender operations.

User response: Right-click the database and select **Spatial Extender->Enable** from the menu.

DBA7300W Database *database-name* cannot be displayed because it does not have a valid node cataloged.

Explanation: There is no entry in the node directory corresponding to the node name provided for the specified database in the database directory.

User response: Ensure that there is a node catalog entry corresponding to the node at which the database was cataloged.

DBA7301W There are no instances to display in the Health Center.

Explanation: There are no entries in the node catalog and DB2INSTANCE is not set.

User response: Catalog a node or specify the default DB2 instance on the client.

DBA7302E Alert information cannot be retrieved for instance *instance-name* or any of its objects because the instance catalog partition is down.

Explanation: The health monitor cannot be queried for alert information because the instance partition and monitor are down.

User response: Restart the database partition and try the action again.

DBA7303E Alert information cannot be retrieved for instance *instance-name* or any of its objects because the instance is down.

Explanation: The health monitor cannot be queried for alert information because the instance and monitor are down.

User response: Restart the instance and try the action again.

DBA7304W The object that was selected *object-name* is no longer present under the current toggle filter.

Explanation: The selected object changed state on the most recent refresh of the data. Therefore, the object is no longer viewable in the navigation view on the current toggle filter selection.

User response: Select another instance or database from the current toggle filter selection or select a different toggle button to view other states.

DBA7305E The alert that was selected no longer exists.

Explanation: The selected alert no longer exists in the health monitor.

User response: If the NOTIFYLEVEL setting is set to include the alert type, review the administration notification log to see details on the alert.

DBA7306E The alert no longer exists. Do you want to close the Details dialog?

Explanation: The current details dialog is showing information on an alert that no longer exists.

User response: Click on Yes to close the dialog, or click on No to leave the details dialog open.

DBA7307W The database manager configuration parameter value has changed since the dialog was last refreshed. Do you want to continue?

Explanation: The value for the configuration parameter was changed since this details dialog was last refreshed.

User response: Click on Yes to apply the change and overwrite the current setting, or click on No to return to the details dialog. If you click on No, refresh the contents of the dialog to see the current setting of the configuration parameter.

DBA7308W The database configuration parameter value has changed since the dialog was last refreshed. Do you want to overwrite the current value?

Explanation: The value for the configuration parameter was changed since this details dialog was last refreshed.

User response: Click on Yes to apply the change and overwrite the current setting, or click on No to return to the details dialog. If you click on No, refresh the contents of the dialog to see the current setting of the configuration parameter.

DBA7309W One or both threshold setting values have changed since the dialog was last refreshed. Do you want to overwrite the current values?

Explanation: The value for the warning and/or alarm thresholds have changed since the details dialog was last refreshed.

User response: Click on Yes to apply the change and overwrite the current setting, or click on No to return to the details dialog. If you click on No, refresh the contents of the dialog to see the current setting of the threshold values.

DBA7310I The threshold setting configuration updates have been applied.

Explanation: The threshold setting configuration updates have been applied.

User response: No action is required.

DBA7311I The configuration parameter update has been applied.

Explanation: The configuration parameter update has been applied. It can take some time before the changes take effect.

User response: No action is required.

DBA7312I The configuration settings have been reset to the original install default settings.

Explanation: The configuration settings have been reset to the original install default settings. These health indicator settings were used at product install time as the initial settings for the health monitor.

User response: No action is required.

DBA7313I The configuration settings for object *object-name* have been reset to the current object default health indicator settings.

Explanation: The configuration settings for the specified object have been reset to the current object default health indicator settings. These settings are based on the global health indicator default settings for the object type of the specified object.

User response: No action is required.

DBA7314W The updated configuration settings can be applied to existing objects. Do you want to propagate all changes to existing objects and overwrite current settings on those objects? If you select No, only the default settings will be updated.

Explanation: The new global default health indicator settings can be propagated to all existing objects that are affected based on the object type for the changed health indicators. The global settings can be updated without propagating changes to existing objects.

User response: Click on Yes to apply the global default health indicator updates and these changes to existing objects, or click on No to apply only the global default health indicator setting updates.

DBA7315W You must provide a script location that is accessible by all database partitions for the action to succeed on all database partitions. Do you want to continue creating the script?

Explanation: The script being defined must be accessible to all database partitions for the alert action to succeed on all database partitions.

User response: Click on Yes to continue creating the script, or click on No to return to the Script Details dialog.

DBA7316E You have chosen a start record value greater than the highest record number in the notification log. No records can be returned.

Explanation: There are fewer records in the notification log than specified as the argument for start

record. No records can be displayed.

User response: Decrease the start record number in the filter condition.

DBA7317W You have chosen an end record value greater than the highest record number in the notification log. Do you want to continue?

Explanation: The notification log has fewer records that are specified as the end record count. Log records can only be retrieved to the end of the file.

User response: Click on Yes to retrieve all the log records up to the end of the file, or click on No to change the filter condition.

DBA7318I There are currently *number* alerts in the Health Center. For more details, launch the Health Center from the toolbar or from the beacon on the status line.

Explanation: There are current alerts to be viewed in the Health Center.

User response: No action is required.

DBA7319E The selected instance *instance-name* has been dropped since the last refresh.

Explanation: The selected object no longer exists so alert information cannot be retrieved.

User response: Recreate the object and re-try the action.

DBA7320E The selected database *database-name* has been dropped since the last refresh.

Explanation: The selected object no longer exists so alert information cannot be retrieved.

User response: Recreate the object and re-try the action.

DBA7321E The object *object-name* for which this alert was generated has been dropped. The requested action cannot be completed.

Explanation: The selected object no longer exists so the action cannot be completed.

User response: Recreate the object and re-try the action.

DBA7323I Removing the contact from the contact list does not change any configured notification settings in the Task Center or Health Monitor.

Explanation: Any notification actions using the

removed contact will not be updated. These notification actions will fail.

User response: Remove any notification actions involving the removed contact.

DBA7324I Contact *contact-name* selected for notification no longer exists. The contact has been removed from the selected list.

Explanation: Invalid contacts have been removed from the list selected for health notification.

User response: No action required.

DBA7325W There are unsaved changes in the dialog. Do you want to discard the changes and refresh the details?

Explanation: You have chosen the refresh or reset option on this dialog. All unapplied changes will be discarded.

User response: Select No, if you want to retain the changes. Otherwise, select Yes to discard any unapplied changes.

DBA7326E No scheduler systems could be found. Task actions cannot be specified.

Explanation: No scheduler systems are cataloged in the admin node directory. Therefore, no tasks can be located for the Select Tasks dialog.

User response: Ensure the DB2 Administration Server for the scheduler system is cataloged in the admin node directory.

DBA7327I The configuration parameter updates have been applied. These changes to the global object type defaults do not affect existing settings for specific database objects.

Explanation: The updated configuration settings apply only to the global default health indicator settings. If you have previously specified settings for a health indicator on a particular object, those settings will not be affected by these changes.

User response: No action is required.

DBA7328I Evaluation has been disabled on the health indicator for the selected alert. The health indicator will not be evaluated on the Health Monitor refresh.

Explanation: The alert configuration for the health indicator of the selected alert has been updated to disable evaluating thresholds or states. This change will not be reflected until the next time the health monitor refreshes.

User response: No action is required.

DBA7329W Task *task-name* cannot be found in the task metadata. Task name is shown as *task-id-number.task-suffix-number*.

Explanation: The Select Actions dialog retrieves the task name for task actions from the task metadata. There was no task matching the task id number and suffix number in the task metadata so a task name cannot be displayed in the dialog.

User response: Check that the task exists in the metadata. If it has been deleted, update the action for the health indicator with the correct task id number and suffix number.

DBA7330E The action cannot be completed on this instance because its remote instance name is not known.

Explanation: The remote instance name field is an optional parameter and was not specified. To perform this action on a remote instance, the remote instance name must be known.

User response: Drop and recatalog the instance specifying the REMOTE_INSTANCE parameter.

DBA7331W There are one or more orphaned contacts in the health notification contact list for the instance *instance-name*.

Explanation: An orphaned contact is a contact that exists in the health notification contact list but is not defined in the contact list stored on the system specified by the CONTACT_HOST configuration parameter of the DB2 administration server.

Currently, there is at least one orphaned contact in the health notification contact list for the instance *instance-name*.

User response: If you do not need the orphaned contacts any more, delete them. Otherwise, redefine them so that they will receive health alert notifications. You can use the Troubleshoot Health Alert Notification wizard to perform these actions.

DBA7332W The filter for the monitoring task *monitoring-task-name* is too complex to be displayed.

Explanation: The filter for the monitoring task *monitoring-task-name* is too complex to be displayed in the Set Up Activity Monitor wizard. As a result, you will not be able to view or modify the filter using the Set Up Activity Monitor wizard.

User response: To view the filter, query the view SYSTOOLS.AM_TASKS. To modify the filter, call the stored procedure SYSPROC.AM_SAVE_TASK.

Chapter 82. DBA7500 - DBA7999

DBA7500N The threshold value specified is not within the threshold range.

Explanation: The threshold value specified is not valid. It has been reset to its previous value.

User response: Specify a value within the valid range.

DBA7501N The configuration parameter value specified is not within the valid range.

Explanation: The configuration parameter value specified is not valid. It has been reset to its previous value.

User response: Specify a value within the valid range.

DBA7502N The specified configuration parameter value is not valid.

Explanation: The configuration parameter value specified is not valid. It has been reset to its previous value.

User response: Specify a valid configuration parameter value.

DBA7503N The specified threshold value is not valid.

Explanation: The threshold value specified is not valid. It has been reset to its previous value.

User response: Specify a valid number.

DBA7504N An instance attachment is required to invoke Memory Visualizer.

Explanation: Memory Visualizer requires an instance attachment in order to be able to display the memory utilization and allocation.

User response: Select or specify a version 8.1 instance or higher.

DBA7510W The action is not currently available with this version of the DB2 server. For this action, the DB2 server must be at level *db2-version*.

Explanation: The Control Center client and the DB2 Server are at different levels. The requested action is not available for the current level of the DB2 server. The DB2 server must be at level *db2-version*.

User response: Upgrade the DB2 server to the required level.

DBA7511W The action is not currently available with this version of the database. For this action, the database must be at level *db-version*.

Explanation: The requested action is not available for the current level of the database. The database must be at level *db-version*.

User response: Upgrade the database to the required level.

DBA7512W The action is not currently available because the required license is not found.

Explanation: The requested action is not available because the DB2 Administration Server is not started or because the license is not installed.

User response: Ensure the DB2 Administration Server is started and that the required license is installed.

DBA7513W The action is not currently available.

Explanation: The action may not be available for the following reasons:

1. The correct license is not installed
2. The correct product is not installed.
3. The client and the DB2 server are at different levels.
4. The DB2 Administration Server is not started.

User response:

1. Ensure the proper license is installed.
2. Ensure the product is installed.
3. Ensure the client and the DB2 server are at the same level.
4. Ensure the DB2 Administration Server is started.

DBA7514W A satellite control database could not be found for the instance *instance*.

Explanation: A satellite control database does not exist in instance *instance*, or it is not cataloged in the database directory.

User response: To work with a satellite control database in instance *instance*, ensure that it is created and cataloged in the database directory. Or you can work with satellite control databases in other instances (if any).

DBA7515W Since the database is a down level server, new options that are not available in previous versions have been

removed or disabled.

Explanation: Since you are connected to a down level server, the new options that are not available on your server version have been disabled or are no longer available.

User response: Complete the dialog, and run the command.

DBA7516N Memory Visualizer could not find the specified data file.

Explanation: The specified file does not exist or it is an invalid Memory Visualizer data file.

User response: Verify that the file name specified is correct.

DBA7517N Memory Visualizer could not write to the specified data file.

Explanation: Memory Visualizer could not save the data file to the specified path and file name.

User response: Verify that the specified path exists and that the file permissions are correct.

DBA7604N No storage management snapshot for the current database partition group can be found.

Explanation: To take a storage management snapshot, from the Control Center, expand the object tree until you find the database partition group object. Right-click the database partition group object and select Manage Storage from the pop-up menu.

User response: Choose a different option on the Redistribute Strategy page to proceed with the redistribute wizard, or close the wizard and take a storage snapshot first before continuing with the redistribute operation.

DBA7608N An error occurred when attempting to create the registry for stepwise redistribute settings for *database-name.database-partition-group-name*.

Explanation: Without the registry records the stepwise redistribute tasks will fail.

User response: Verify your file permissions and recreate the redistribute settings registry records using the Wizard.

DBA7609W The database partition group is not partitioned.

Explanation: The current database partition group contains only one database partition.

User response: There is no need for running redistribute utility for it.

DBA7610W The table is not distributed.

Explanation: The table is defined under a database partition group which contains only one database partition.

User response: No data distribution file is created for the given table.

Altering the distribution key can only be done to a table whose table space is associated with a single-partition database partition group.

You can add or drop distribution keys, using the ALTER TABLE statement.

You can also use the Design Advisor to migrate from a single-partition to a multiple-partition database. Search the Information Center for a topic called: "Using the Design Advisor to migrate from a single-partition to a multiple-partition database".

DBA7611N The stepwise redistribute property type is not valid.

Explanation: The requested stepwise redistribute property entry cannot be retrieved or updated.

User response: Enter a valid property type.

DBA7612N An error has occurred when accessing file.

Explanation: The possible reasons include:

- File format error
- Communication error
- Memory allocation error when reading from file

User response: Fix any format errors in the file. Check the db2diag log file for more information to help resolve the communication or memory allocation errors. Rerun the application.

DBA7613W Affected tables within the database partition group may be locked and unavailable for long periods of time.

Explanation: Redistribution can be a log-intensive activity.

User response: It is recommended to run redistribute on a test database before running it on your production system to determine how long the redistribution will take.

To minimize the risk of running out of active log space during a redistribute, it is recommended that you perform some or all of the following actions:

- Reduce or eliminate other log-intensive database activities
- Use infinite logging
- Use stepwise redistribution as recommended by the Wizard

DBA7614W Affected tables within the database partition group may be locked and unavailable for long periods of time when redistributing multiple database partition groups.

Explanation: Redistribution can be a log-intensive activity.

User response: It is recommended to run redistribute on a test database before running it on your production system to determine how long the redistribution will take.

To minimize the risk of running out of active log space during a redistribute, it is recommended that you perform some or all of the following actions:

- Reduce or eliminate other log-intensive database activities
- Use infinite logging
- Schedule the redistribution of each database partition group to happen at a different time to minimize contention for available log space
- Use stepwise redistribution as recommended by the Wizard

DBA7615N One or more tables required by the application are not defined with the expected definition.

Explanation: The tables the application is trying to access do not match the expected definition. This may be caused by one or more of the following problems:

- Table does not exist
- Table column definition is not the same as expected

User response: The tables may have to be dropped and recreated with the appropriate definition. If this error is being encountered during a drop, then you can force the tables to be dropped by specifying the FORCE option in the drop command.

DBA7616N Invalid partition weight for stepwise redistribute operation.

Explanation: The value of the partition weight must be between 0 to 32767.

User response: Enter an integer within the valid range and try again.

DBA7617N One or more input parameters do not contain valid values expected by the routine.

Explanation: This may be caused by one or more of the following problems:

- Numerical parameter is out of the expected range
- Object referenced by a parameter may no longer exist

User response: Verify your input parameters with the

routine specification, and try again.

DBA7618W Verify that all databases at instance *instance-name* are cataloged.

Explanation: When adding database partitions to an instance, all temporary table spaces for all databases must be altered with new containers on the new database partitions. A database must be cataloged at that instance for the Add Database Partition wizard to detect its temporary table spaces.

User response: Review the list of databases on the first page of the Add Database Partition wizard.

If there are any databases missing:

1. Close the wizard.
2. Close the Add Database Partitions launchpad.
3. Add the missing databases to the databases folder in the Control Center.

If there are databases listed that do not actually exist at the instance, remove them from the databases folder in the Control Center.

DBA7619W Verify that all databases at instance *instance-name* are cataloged.

Explanation: When dropping database partitions from an instance, all data in all databases needs to be moved off the affected database partitions. A database must be cataloged at that instance for the Add Database Partition wizard to detect its database partition groups.

User response: Review the list of databases in the Drop Database Partition from Database Partition Groups window.

If there are any databases missing:

1. Close the window.
2. Close the Drop Database Partitions launchpad.
3. Add the missing databases to the databases folder in the Control Center.

If there are databases listed that do not actually exist at the instance, remove them from the databases folder in the Control Center.

DBA7620N There is no distribution key defined for the given table.

Explanation: A distribution key is a column (or group of columns) that is used to determine the database partition in which a particular row of data is stored. Tables without a distribution key are only allowed in single-partition database partition groups. Without a distribution key, a data distribution file cannot be generated, and data distribution analysis cannot be performed.

User response: Altering the distribution key can only

be done to a table whose table space is associated with a single-partition database partition group.

You can add or drop distribution keys, using the ALTER TABLE statement.

You can also use the Design Advisor to migrate from a single-partition to a multiple-partition database. Search the Information Center for a topic called: "Using the Design Advisor to migrate from a single-partition to a multiple-partition database".

DBA7621W Tables require distribution keys before they can be redistributed from one to many database partitions. Review the database design before moving from one to many database partitions.

Explanation: A distribution key is required to partition table data across a set of database partitions.

Tables without a distribution key are only allowed in single-partition database partition groups.

User response: Use the Alter Table notebook to add distribution keys to all tables that will be partitioned.

Consider creating different database partition groups. Usually, smaller tables can stay on a small number of database partitions, even just one, while larger tables should be distributed across all or most of the database partitions.

Altering the distribution key can only be done to a table whose table space is associated with a single-partition database partition group.

You can add or drop distribution keys, using the ALTER TABLE statement.

You can also use the Design Advisor to migrate from a single-partition to a multiple-partition database. Search the Information Center for a topic called: "Using the Design Advisor to migrate from a single-partition to a multiple-partition database".

DBA7627N Unable to load snapshot data.

Explanation: The selected snapshot no longer exists.

User response: Select a different snapshot from the Current Snapshot list on the Storage Management toolbar.

If no other snapshot exists, you can capture a new snapshot by right clicking on a database, database partition group, or table space, and selecting the Capture Snapshot menu option.

DBA7628N Drop Database Partition task must start after the last data redistribution is complete. The last redistribution task is scheduled to start at *start-date start-time*.

Explanation: All data redistribution tasks must be

completed before the Drop Database Partition task can run. Redistribution moves the data off the database partitions that will be dropped.

User response: Choose a new start time for the Drop Database Partition task or just save it to the Task Center and run it manually once all the redistribution tasks are complete.

DBA7630W The following hosts could not be found: *host-names*.

Explanation: Add Database Partition will fail if these host names cannot be found on the network. Host names specified must refer to existing systems at the time that this Add Database Partition task is scheduled to execute.

User response: Verify that all of the listed host names are correct. Ensure that the specified host names exist when this task is scheduled to execute.

DBA7631W The instance *instance-name* will be stopped and restarted when this task is run.

Explanation: The script to be executed contains the following commands:

```
db2stop force
db2start
```

This will stop the database instance, refusing access to all connected applications. All databases in this instance will not be available while the instance restarts.

User response: Run this task, which will restart the instance at the scheduled time, or schedule the task to run at a different time, or save the task to the Task Center and run it later.

DBA7632N The database partitions *database-partitions* cannot be removed from the following database partition groups *partition-groups*.

Explanation: All database partition groups must have at least one database partition.

User response: When altering a database partition group, you must keep at least one database partition in the database partition group. If you do not want any database partitions in the database partition group, you must drop that database partition group and lose any existing data.

DBA7633N The database partitions *partition-names* cannot be dropped because they are catalog partitions for one or more databases.

Explanation: A catalog partition cannot be dropped without first dropping the corresponding database. On

Windows systems, you cannot drop database partition 0.

User response: Select one or more other database partitions to drop from the instance, or drop the databases for the catalog partition you wish to drop before trying to drop the database partition.

DBA7634N The container name *container-name* is invalid.

Explanation: The container name must be a valid path if the container type is directory or file, or a valid raw device if the container type is raw device.

User response: Redefine the container name to a valid path or raw device then attempt to change it, or add a new container and then remove the invalid one.

DBA7666W Configuration changes have been made for the previously selected database partition. Do you want to save the changes?

Explanation: Another database partition was selected, but the configuration changes to the previously selected database partition have not been saved.

User response: Click Yes to obtain the configuration parameter values for the next database partition and save the changes that were specified for the previously selected database partition.

Click No to obtain the configuration parameter values for the next database partition and discard the changes that were specified for the previously selected database partition.

Click Cancel to go back to the configuration parameter values for the previously selected database partition.

DBA7900N The SYSPROC.ALTOBJ stored procedure input parameter *parameter* is invalid.

Explanation: An invalid parameter value was specified.

User response: Reissue the stored procedure call with valid parameter values:

- parameter 1 (IN), *execModeName*, choose from one of the following six values: 'GENERATE', 'VALIDATE', 'APPLY_CONTINUE_ON_ERROR', 'APPLY_STOP_ON_ERROR', 'UNDO' and 'FINISH'
- parameter 2 (IN), *sqlStmt*, valid CREATE TABLE DDL is expected
- parameter 3 (IN/OUT), *alterId*, use -1 if no alter plan and ID has been previously generated, use the previously generated integer id, if a existing plan is to be followed
- parameter 4 (OUT), *msg*, use ? in the call statement

DBA7901W There has been a request to transform column *column-name* from source type *source-type* to target type *target-type*. However, there is no default casting function available.

Explanation: The stored procedure attempts to use built-in scalar column functions to transform the existing data in the table from its source type into its target type. There is no built-in column function that can perform the data type transformation requested. Ensure that a UDF is created to transform the data, otherwise the existing data may not be transferred to the new table.

User response: If you have a column function that can transform the data to its target type, update the SYSTOOL.ALTOBJ_INFO table with the following statement:

```
UPDATE SYSTOOLS.ALTOBJ_INFO_V
SET SQL_STMT='edited SELECT statement'
WHERE ALTER_ID='<alterId>'
AND SQL_OPERATION='SELECT'
AND EXEC_MODE LIKE '_1____'
```

Alternatively, change the target column type specified in the third parameter, so that there is a built-in column function that can be used for data transformation. Then call the stored procedure again with the new input.

DBA7902N The table to be altered contains at least one column data type that is not supported by the SYSPROC.ALTOBJ stored procedure.

Explanation: DATALINK type, structured type UDTs, and reference type UDTs are not supported. The SYSPROC.ALTOBJ stored procedure cannot be used to alter a table containing non-supported column data types.

User response: Do not attempt to alter a table containing non-supported column data types.

DBA7903N The table type is not supported by the SYSPROC.ALTOBJ stored procedure.

Explanation: The following types of tables cannot be altered using SYSPROC.ALTOBJ:

- Materialized Query Tables.
- Type tables or a tables that are the scope of any existing reference column.
- Remote tables referenced by a nickname.
- Tables with activated row or column level access control.
- Tables protected with row level LBAC.

User response: Do not attempt to alter the table using the SYSPROC.ALTOBJ stored procedure.

DBA7904N The user does not have *authority-name* authority, which is required to run the SYSPROC.ALTOBJ stored procedure.

Explanation: DBADM authority is required to run the SYSPROC.ALTOBJ stored procedure.

User response: Obtain the necessary authority before running the stored procedure.

DBA7905N The drop column action cannot be performed at the same time as adding or renaming columns.

Explanation: When dropping existing columns, column names are used to match the existing columns to their new definitions. When adding or renaming columns, the column index is used to match the existing and the new columns.

User response: The drop column operation must be attempted separately from adding or renaming column operations, however you can add and rename columns at the same time. The sequence of the existing columns cannot be changed using alter.

If you need to perform more than one type of alter operation, the SYSPROC.ALTOBJ must be called multiple times, one for each compatible alter operation group using different DDLs.

DBA7906N The number of SQL statement generated for the alter process exceeds the expected limit.

Explanation: This may occur when a large number of related objects for the table being altered, such as triggers, aliases, views, SQL stored procedures, and materialized query tables.

User response: The SYSPROC.ALTOBJ stored procedure cannot be used to alter the given table.

DBA7907W The alter process completed with errors.

Explanation: This warning may occur when there are related objects of the table that are dependent on the columns that have been changed, or the existing data cannot be transformed into the their target data type, or conform to the new constraints.

User response: Use the SYSPROC.ALTOBJ stored procedure with the following parameter inputs so that the related objects can be recreated.

```
CALL SYSPROC.ALTOBJ('UNDO', CAST
  (NULL AS VARCHAR(2)), alterID, ?)
```

Or examine the errors in the SYSTOOLS.ALTOBJ_INFO_V, update the SQL statements that caused runtime error, and run the stored procedure again in APPLY mode.

DBA7908E Some of the functions are not available because a set of stored procedures has not been installed for *subsystem-name*. The following stored procedures may be missing: *stored-procedures*.

Explanation: The Control Center could not perform any of the following:

- Call SYSPROC.DSNUTILS to determine which utilities are installed.
- Call SYSPROC.DSNWZP to retrieve the subsystem parameters.
- Locate at least one of the OS/390 Enablement stored procedures.

User response: Ensure that the OS/390 Enablement stored procedures and DB2 stored procedures DSNWZP and DSNUTILS are installed and that the activation steps were performed. See the DB2 management Clients Program Directory for more information.

Chapter 83. DBA8000 - DBA8499

DBA8000E **Unable to generate a new policy ID because the maximum number of allowable suffix IDs has been reached.**

Explanation: Each policy has an internally-generated policy ID, which consists of a prefix and a suffix. The prefix is the subsystem identifier and the suffix is a number between 000 and 999. The Object Maintenance wizard fails to generate a new policy ID because the maximum number of allowable suffix IDs has been reached.

User response: Open the List Policies dialog to remove the policies that are no longer required, then attempt to create the new policy again.

DBA8001E **An error was encountered during the initialization of the wizard.**

Explanation: Some of the pages of the wizard were not initialized because an error was encountered while executing a DB2 Administration Server request to read the data set.

User response: Refer to the DB2 Administration Server's First Failure Data Capture Log for additional information.

DBA8002E **You have selected to modify the selection criterion for this work set. This change will cause all alert objects associated with this workset, identified using the current selection criterion, to be deleted from the health monitor maintenance tables.**

Explanation: Modifying the selection criterion will change the set of objects to be evaluated by the health monitor. This invalidates the alert objects in the maintenance tables that were identified by the last health monitor evaluation for this work set, and they will be deleted from the health monitor maintenance tables.

User response: Click Yes to change the selection criterion for this work set, and to delete all alert objects currently associated with this work set from the health monitor maintenance tables.

Click No to keep the current selection criterion for this work set.

DBA8010N **The tree view of the XML document cannot be built.**

Explanation: Explanation: The document may be corrupted.

User response: Ensure that the document is well-formed and that the character encoding is set correctly.

DBA8011N **The source view of the XML document cannot be built due to an internal error.**

User response: Contact IBM.

DBA8012I **The XML document is empty.**

User response: No action is required.

DBA8013I **The instance attach completed successfully.**

User response: No action is required.

DBA8015N **An I/O error occurred.**

Explanation: An error occurred while attempting to open or close a file.

User response: Verify that the file name is valid and that the user has permission to access the file. Also check for any disk and operating system errors.

DBA8016N **An I/O error occurred.**

Explanation: An error occurred while attempting to read the result set.

User response: Ensure that the data source is available. If the data source is on a network, please ensure the network is working correctly.

Part 12. DBI Messages

This section contains the error messages which may be generated during installation and configuration of DB2 products. The messages are listed in numeric sequence.

Chapter 84. DBI1000 - DBI1499

DBI1001I **An invalid argument was entered for the db2icrt command.**

Explanation: You can create a DB2 database manager instance, including a DB2 pureScale instance by using the db2icrt command. You can also create an initial DB2 member and cluster caching facility (CF) as part of the creation of the DB2 pureScale instance by using the db2icrt command.

This message is returned when the db2icrt command is invoked with incorrect or missing parameters.

User response: Enter the command again using valid parameters.

DBI1002I **Usage:**
db2uit [-d] [-Q] [-D]
 [-q <field>]
 [-i <id>]
 [-a <action>]
 [-r <runlevels>]
 [-p <process>]

Explanation: An invalid argument was entered for the db2uit command. Valid arguments for this command are:

-d turn debug mode on

-q query a field given the search parameters. Without -q, this becomes an update (if -i ID is already present) or insert. Valid fields: i, r, a, p

-i ID of the entry

-r run levels

-a action to perform

-p process to launch

-Q Quiet: no output

-D Delete record

FileName file to modify

RETURN CODES:

-q If no field matches, non-zero, zero otherwise

update/insert If file could not be updated, non-zero, zero otherwise

User response: Rerun the command with valid arguments.

DBI1003I **Usage:**
dlfmcr [-b BackupDir]
 -p PortName
 -g DLFMGid
 DLFMAdmin

Explanation: An invalid argument was entered for the dlfmcr command. Valid arguments for this command are:

-h | -? display the usage information.

-p PortName is the port name or port number to be used by this instance.

-b BackupDir is the directory to be used to create the backup of the dlfm database in.

-g DLFMGid is the DLFM Admin group (this group must be a secondary group of DLFMAdmin)

DLFMAdmin is the name of the DB2 Data Links Manager Administrator.

User response: Enter the command again as follows:
dlfmcr -p PortNum -g DLFMGid DLFMAdmin

DBI1004W **libdb2-link detected. This DB2 copy may encounter problems loading its libraries.**

Explanation: The db2ln command was run from another DB2 copy which creates symlinks in /usr/lib. These symlinks interfere with multiple DB2 copies running on the same system, whether those copies are the same version at different fixpacks, or different versions.

User response: From the installation path where the db2ln command was run, run the db2rmln command to remove these links.

DBI1006I **An invalid argument was entered for the db2idrop command.**

Explanation: The syntax for the db2idrop command is:

```
db2idrop [-h|-?]
          [-d]
          [-g]
          [-f]
          InstName
```

Valid arguments for this command are:

-h|-?

Display the usage information.

-d

Turn debug mode on.

-g

This parameter is required when db2idrop is used with a DB2 pureScale instance. Specifies that you want to drop the DB2 pureScale instance on all hosts. This parameter requires all DB2 members and all cluster caching facilities (CFs) are stopped on all the hosts in the DB2 pureScale instance.

-f

This parameter is deprecated. This parameter flags applications. If this flag is specified all applications using the instance will be forced to terminate. This parameter is not supported on a DB2 pureScale environment.

InstName

Specifies the name of the instance.

User response: Enter the command again using valid parameters.

DBI1007I Usage:
dlfmdrop DLFMAdmin

Explanation: An invalid argument was entered for the dlfmdrop command. Valid arguments for this command are:

-h|-? display the usage information.

DLFMAdmin

is the name of the DB2 data Links Manager Administrator.

User response: Enter the command again as follows:
dlfmdrop DLFMAdmin

DBI1008E Install media and AIX version mismatch.

Explanation: Installing DB2 64-bit on AIX requires different CDs for AIX version 4 and AIX version 5.

db2setup has detected that your system is running AIX version 5, but you are installing from the DB2 for AIX version 4 CD.

User response: Insert the DB2 for AIX version 5 CD

and try the installation again.

DBI1009E Install media and AIX version mismatch.

Explanation: Installing DB2 64-bit on AIX requires different CDs for AIX version 4 and AIX version 5.

db2setup has detected that your system is running AIX version 4, but you are installing from the DB2 for AIX version 5 CD.

User response: Insert the DB2 for AIX version 4 CD and try the installation again.

DBI1010E The host *host-name* is not offline.
Command: *command-name*

Explanation: The host is still active in the GPFS cluster. Prior to updating the GPFS level, the host must be offline.

User response: Shutdown the GPFS cluster on the specified host by executing the specified command. For example, /usr/lpp/mmfs/bin/mmshutdown

DBI1011I An invalid argument was entered for the db2iupdt command.

Explanation: You can use the db2iupdt command to update a DB2 instance to a higher level within a release, convert a DB2 instance to a DB2 pureScale instance, or manage DB2 members or cluster caching facilities (CFs) of a DB2 pureScale instance.

This message is returned when the db2iupdt command is invoked with incorrect or missing parameters.

User response: Enter the command again using valid parameters.

DBI1013I Usage:
dlfmupdt [-h|-?] DLFMAdmin

Explanation: An invalid argument was entered for the dlfmupdt command. Valid arguments for this command are:

-h|-? display the usage information.

DLFMAdmin

is the name of the DB2 Data Links Manager Administrator.

User response: Enter the command again as follows:
dlfmupdt [-h|-?] DLFMAdmin

DBI1014E Instance upgrade to 32-bit is not supported.

Explanation: Upgrading from 64-bit to 32-bit is not supported in this version.

User response: No actions required.

DBI1015I **The db2iupgrade command failed because an invalid argument was specified.**

Explanation: You can upgrade a DB2 database manager instance to a new release by using the db2iupgrade command.

This message is returned when the db2iupgrade command is called and invalid parameters or arguments are specified.

User response: Run the command again, specifying only valid parameters and arguments.

DBI1016I **Program *program-name* is performing uninstallation. Please wait.**

DBI1017I **installFixPack is updating the DB2 product(s) installed in location *location*.**

DBI1019E **DB2 Administration Server *DAS* cannot be updated.**

Explanation: An attempt to update the DB2 Administration Server *DAS* has failed because:

- This version of the dasupdt command cannot be used to update the DB2 Administration Server.
- To use this level of code for the DB2 Administration Server, an upgrade is required (instead of an update).

User response: To move the DB2 Administration Server up to this code level, issue the dasmgr command from the <current DB2 installation path>/instance directory to upgrade the DB2 Administration Server.

DBI1020I **An invalid argument was entered for the db2setup command.**

Explanation: The syntax for the db2setup command is:

```
db2setup [-h|-?]
          [-t <TrcFile>]
          [-r <RspFile>]
          [-c]
          [-l <LogFile>]
          [-i <Lang>]
          [-f <nobackup>]
```

Valid arguments for this command are:

-h|-?

Displays help information.

-t

Turns on the debug mode. The debug information is written to the file name specified.

-r

Performs a response file install using the file name specified. This parameter is mandatory when -c is specified.

-c

Validates the contents of the response file specified in the -r parameter. When specified, you must also specify the -r parameter.

-l

Writes the log to the file name specified. For root installations, the default log file is /tmp/db2setup.log. For non-root installations, the default log file is /tmp/db2setup_<userID>.log, where <userID> represents the user ID that owns the non-root installation. If the IBM Tivoli System Automation for Multiplatforms (SA MP) is being installed with db2setup, the install log file for the SA MP will be located in the same directory as DB2 log files.

-i

Displays the db2setup utility in the language specified.

-f

This applies to the non-root upgrade only. Force db2setup to not backup installation files when the components are updated. If you choose not to backup the files, the space requirement of the installation directory is reduced. However, choosing not to backup the files also means that if any errors occur, the DB2 installer will not be able to perform a rollback operation. In this case, you will need to manually clean up the files and reinstall the product.

User response: Enter the command again using valid parameters.

DBI1021I **Usage:**
db2imigrev [-h|-?] [-m Version]
InstName

Explanation: An incorrect argument was entered. Valid arguments for this command are:

-h|-? display the usage information.

-d turn debug mode on.

-m specify DB2 version.

Version

is the version the instance will be migrated backward to.

InstName

is the name of the instance which will be migrated backward from version 8 to the previous version of DB2.

User response: Enter the command again as follows:
db2imigrev [-h|-?] InstName

DBI1022I Usage:
db2imigrev [-h|-?] [-f] InstName

Explanation: An incorrect argument was entered.
Valid arguments for this command are:

- h|-?** display the usage information
- d** turn debug mode on.
- f** is the force applications flag. If this flag is specified then all applications using the instance will be forced to terminate.

InstName
is the name of the instance which will be migrated backward from version 8 to a previous version of DB2.

User response: Enter the command again as follows:
db2imigrev [-h|-?] [-f] InstName

DBI1023I Usage:
db2ginfo [-h|-?] [-y] [-c InstName] [-a AuthType] [-u FencedID] OutputDir

Explanation: An invalid argument was entered for the db2ginfo command. Valid arguments for this command are:

- h|-?** display usage information.
- y** executes script (only shows warning otherwise).
- c InstName** specify instance to test.
- a AuthType** specify authentication type (SERVER, CLIENT, or SERVER_ENCRYPT).
- u FencedID** specify fenced user ID.

OutputDir
is the directory where the output file, dbginfo.txt, will go.

User response: Enter the command again as follows:
db2ginfo [-h|-?] [-y] [-c InstName] [-a AuthType] [-u FencedID] OutputDir

DBI1024I Usage:
db2iauto [-h|-?] -on|-off InstName

Explanation: An incorrect argument was entered.
Valid arguments for this command are:

- h|-?** display the usage information
- on|-off** enables or disables autostarting of an instance.

InstName
is the name of the instance.

User response: Enter the command again as follows:
db2iauto [-h|-?] -on|-off InstName

DBI1025I Usage:
dascrt [-d] -u ASUser

Explanation: An invalid argument was entered for the dascrt command. Valid arguments for this command are:

- d** enters debug mode, for use with DB2 Service
- u ASUser** is the user the DAS will run under

User response: Enter the command again as follows:
dascrt -u ASUser

Confirm that user IDs and group names used to create the DB2 instance are valid.

DBI1026I Usage:
dlfmfsmd [-h|-?] [-j] dlfsMntPt

Explanation: An invalid argument was entered for the dlfmfsmd command. Valid arguments for this command are:

- h|-?** display the usage information.
- j** modify the named file system from dlfs to jfs on AIX or ufs on Solaris Operating Environment.

dlfsMntPt
Mount Point for Data Links Manager Filter file system (including the slash: /).

User response: Enter the command again as follows:
dlfmfsmd dlfsMntPt

DBI1027I Usage:
command-name [-d]

Explanation: An invalid argument was entered for the command. Valid arguments for this command are:

- d** enters debug mode, for use with DB2 Service

User response: Review the command syntax and enter the command again.

DBI1028I Usage:
dasupdt [-d] [-D] [-h|-?]

Explanation: An invalid argument was entered for the dasupdt command. Valid arguments for this command are:

- d**
enters debug mode, for use with DB2 Service
- D**
Moves the DAS from a higher code level on one path to a lower code level installed on another path.
- h|-?**
Displays usage information.

User response: Enter the command again as follows:
dasupdt

DBI1029I Usage:
db2nqadm [start|stop]

Explanation: An invalid argument was entered for the db2nqadm command. Valid arguments for this command are:

start start all NetQ services and daemon
stop stop all NetQ services and daemon

You must specify one of the arguments.

User response: Enter the command again with the proper argument.

DBI1030E Kernel parameter *name* must be set to *value*.

Explanation: DB2 requires certain kernel parameters to be updated.

User response:

- Update all necessary kernel parameters.
- Reboot the system
- Try this command again

DBI1031E The installation operation failed because the selected product requires DB2 Server Edition to be installed first.

Explanation: There are multiple methods for installing DB2 database software, including the DB2 Setup wizard and the db2_install command.

This message is returned when a component or Feature is selected to be installed, but the selected component or Feature cannot be installed until after DB2 Server Edition is installed.

User response: Install DB2 Server Edition.

DBI1032E The product you selected can not be installed on AIX with a version less than 4.3

Explanation: The DB2 product you selected requires an AIX version that is 4.3 or greater.

DBI1033E The installation operation failed because the selected product requires DB2 Server Edition to be installed first.

Explanation: There are multiple methods for installing DB2 database software, including the DB2 Setup wizard and the db2_install command.

This message is returned when a component or Feature is selected to be installed, but the selected component or Feature cannot be installed until after DB2 Server Edition is installed.

User response: Install DB2 Server Edition.

DBI1034W Missing prerequisite package.

Explanation: One of the following packages has been selected for installation: "<pkg-name>". Either the installed "<pkg-name>" file set level is below "<name>" or it cannot be detected.

User response: Ensure the "<pkg-name>" file set is version "<name>" or higher.

Your "<pkg-name>" level will not affect the installation of DB2.

DBI1035E Failed to mount file system *File-System*.

Explanation: An attempt to mount the specified file system has failed.

User response: Verify that the file system is defined. Correct any errors from the mount command and try again.

DBI1036E Failed to unmount file system *File-System*.

Explanation: An attempt to unmount the specified file system has failed.

User response: Correct any errors from the umount command and try again.

DBI1037E File system *File-System* does not exist.

Explanation: The file system specified is not defined on this workstation.

User response: Define the file system on the workstation and try again.

DBI1038E File system *File-System* cannot be modified.

User response: Ensure that you have write permission to the system file that contains file system definition, and try again.

DBI1039W Missing prerequisite package.

Explanation: One of the following packages has been selected for installation: "<name>" However, the installed "<name>" file set level is below "<name>"

User response: Install the “<name>” file set version “<name>” which is provided on the CD.

DBI1040I The translated version of following message is not available on this system.

User response: The translated version of following message is not available on this system. Please refer to the "Installation.Notes" file in the PRODDIR/Readme/LOCALE directory, where PRODDIR is the product directory and LOCALE is the locale name. For example, the product directory is:

/usr/opt/db2_08_xx, where xx is 01 or FPN,
where n is the FixPak number
or /opt/IBM/db2/V8.x, where x is 1 or FPN,
where n is the FixPak number

DBI1041E File system *inputMntPt* cannot be converted to dlfs.

Explanation: The File-systems `/`, `/var`, `/usr`, and `/tmp` cannot be converted to a `dlfs` File-system. This has been done to avoid booting problems that may arise by making these `dlfs` File-systems.

User response: Run the command again specifying a different File-system.

DBI1042E Base File-system *BaseFS* cannot be changed to dlfs.

Explanation: On Aix, a File-system can be changed to dlfs only if its vfs is fsm or jfs. Dlfs is not supported for any other vfs.

User response: Run the command again on a jfs or fsm file system.

DBI1043E The product you selected requires DB2 Administration Client installed first.

User response:

- Install DB2 Administration Client.
- Try to install the product again.

DBI1044E The installation of the DB2 National Language Package cannot begin because there is no prerequisite DB2 product installed at this install location:
install-path.

Explanation: The DB2 National Language Package can only be used to add national language support to installed DB2 products that include:

DBI1045I An invalid argument was entered for the db2 install command.

Explanation: The syntax for the db2_install command is:

```
db2_install [-b <installpath>]
             [-p <db2producttobeinstalled>]
             [-c <imagelocation>]
             [-l <logfile>]
             [-f NOTSAMP]
             [-f PURESCALE]
             [-f nobackup]
             [-f ignoreType1]
             [-t <trcFile>]
             [-n]
             [-m]
             [-L <language>]
             [-h|-?]
```

User response: Enter the command again using valid parameters.

```

DBI1046I      Usage:
               doce_install [-b <installpath>] [-
               p <db2producttobeinstalled>]
                           [-c <imagelocation>] [-
               l <logfile>] [-t <trcFile>]
                           [-n] [-L <language>] [-h | -
               ?]

```

Explanation:

- b** Specifies the path where the DB2 product will be installed. Mandatory when **-n** is specified. The length of the path is limited to 128 characters and is a full path name.
- p** Specifies the DB2 product to be installed. Mandatory when **-n** is specified.
- c** Specifies the location of the related DB2 National Language Pack (NLPACK). This parameter is mandatory when **-n** is specified.

The DB2 NLPACK locations needs to be provided explicitly if the all of the following conditions are all met:
 - The **-n** option is specified
 - The current installation requires needs to install National Language (non-English) language supports

- The DB2 NLPACK is neither on the DB2 DVD nor in the same subdirectory as the DB2 product being installed
- l** Specifies the log file. The default log file is /tmp/doce_install.log\$\$, where \$\$ is the process id.
- t** Turns on the debug mode. The debug information is written to the file name specified.
- n** Specifies non-interactive mode. When specified, you must also specify -b, -p, and -c.
- L** Specifies national language support. To install multiple languages at the same time, this parameter can be specified multiple times.
- For example, to install both English and German, specify -L EN -L DE.
- h|-?** Displays help information.

User response: Reenter the command.

DBI1047I An invalid argument was entered for the db2_deinstall command.

Explanation: The syntax for the db2_deinstall command is:

```
db2_deinstall -F <featureName> | -a | -r <RspFile>
               [-l <logFile>]
               [-b <installPath>]
               [-t <trcFile>]
               [-s GPFS]
               [-s TSAMP]
               [-f sqllib]
               [-h|-?]
```

For details about the parameters of the db2_install command, refer to the DB2 documentation for Linux, UNIX, and Windows.

User response: Enter the command again using valid parameters.

DBI1048I Usage:
doce_deinstall -a
 [-l <logFile>]
 [-b <installPath>]
 [-t <trcFile>]
 [-r <response_file>]
 [-h|-?]

Explanation:

- a** Removes all installed DB2 products in the current location.
- b <installPath>** This option is valid if the command is run from the DB2 media. It specifies the absolute path where the DB2 product was installed and

will be uninstalled. The command will prompt for the path if the option is not specified.

-l <logFile>

Specifies the log file. The default log file is /tmp/doce_deinstall.log\$\$, where \$\$ is the process id.

-t <trcFile>

Turns on the debug mode. The debug information is written to the file name specified.

-r <response_file>

Removes the Information Center using a response file. For example, doce_deinstall -r db2un.rsp. Cannot be combined with the -a parameter.

-h|-?

Displays help information.

User response: Reenter the command.

DBI1049I Usage:
db2ls [-q] [-b <baseInstallpathOfDB2>]
 [-c] [-f <feature rsp file ID>]
 [-l <logfile>]
 [-p] [-a]

Explanation:

-q

Specifies that this is a query of the local components. By default, only the visible components (features) are displayed unless -a is also specified. Mandatory when -p is specified.

-b

Mandatory when -q is specified and you are running the global db2ls command.

-c

Prints the output as a colon-separated list of entries rather than column-based. Allows you to use shell scripts to accurately work with this information.

For example:
#FEATURE:NAME:VRMF:FIXPACK:
SPECIALINSTALL:PREREQS.

-f

Queries to check if the specific feature is installed. If the feature is installed, the return code is zero. If the feature is not installed, the return code is non-zero.

-l

Specifies the log file. The default log file is /tmp/db2ls.log

-P

Lists the DB2 products installed. Cannot be combined with -f or -a.

-a

Lists all hidden components and visible features. The default is to only list visible features.

User response: Reenter the command.

DBI1050E Unsupported operating system - OS-name, version OS-ver.

Explanation: Your current operating system is not supported or this version of the operating system is not supported.

DBI1051E You cannot execute this program as root.

Explanation: This program can only be run under a non-root user ID.

User response: Login as an user ID other than root, and issue the command again.

DBI1052E The command failed because the current user does not have system root user authority.

Explanation: Special privileges are required to execute some DB2 database commands.

This message is returned when a non-root user ID issues a DB2 database command that requires root user authority.

User response: Login as root and then issue the command again.

DBI1053E File system of type *fsys-type* is not supported.

Explanation: Creating the instance's home directory on this type of file system is not supported. Instance's home directory must exist either on a local file system or on a NFS mounted file system.

User response: Change the home directory to a supported file system and recreate the instance.

DBI1054I The installFixPack command failed because an invalid argument was specified.

Explanation: You can update the installed DB2 database products in a given location, on all UNIX and Linux platforms by using the installFixPack command.

This message is returned when an invalid parameter

argument is specified with the installFixPack command.

User response: Enter the command again using valid parameters.

DBI1055E The message file *file-name* cannot be found.

Explanation: The message file required by the instance commands is missing from the system; it may have been deleted or the database products may not be installed properly.

User response: Verify that the product option containing the message file is installed correctly. If there are verification errors, reinstall the product option.

DBI1056I Usage:
db2chgpath [-d] [-f <relativePath/
FileName>]

Explanation:

-d Turns debug mode on.

-f Specifies a specific file name to update the runtime path. The specified file name should have the path name relative to the base of the current DB2 install location. Root authority is required, and this command must be run directly from DB2DIR/install directory.

User response: Reenter the command.

DBI1057E Required locale *locale* is not available.

Explanation: The required locale directory, /usr/lib/locale/LANG/LC_MESSAGES, does not exist, where LANG is the name of your current locale.

User response: Verify that the specified locale has been installed correctly. Also check for the read and execute permissions are set correctly on that directory.

DBI1058E gunzip command not found.

Explanation: The gunzip utility must be available on your system to install or update DB2 products.

User response: Ensure the gunzip utility is installed in the PATH environment variable. Reenter the command.

DBI1059E There still exists DB2 instance(s) related to the current installation directory where DB2 Suite filesets are going to be completely uninstalled. The deinstallation process can't continue because of this.

User response:

1. If the deinstallation is for moving up or down DB2 level for the current DB2 installation, you need to use installFixPack in the DB2 images of the desired DB2 Level to update the current DB2 installation.

2. If the purpose is to uninstall then reinstall DB2 in the same location, use installFixpack with the -f option.
3. If the purpose is to just uninstall DB2 in the current location, you need to drop the DB2 instances related to the installation directory if they are not needed any more. Or you need to update the DB2 instances to other DB2 installation location for the same version of DB2 and then restart db2_deinstall.

DBI1060E Invalid package name *pkg-name*.

Explanation: An incorrect name has been entered. The package either does not exist or the name has been entered incorrectly.

User response: Check to see if the name of the given package exists on the distribution media. If so, examine the name to see if it has been misspelled.

DBI1061E A file set or package name *pkg-name* is missing.

Explanation: An error occurred when verifying software pre-requisites. The indicated file set or package must be installed before using this command.

User response: Check to see if the name of the given package has been installed on your system.

DBI1062E A file set or package name *pkg-name* is missing.

Explanation: The indicated file set or package must be installed before using this command.

User response: Install the file set or package and then reissue this command.

DBI1063E PTF or patch *patch-name* is missing.

Explanation: The indicated PTF or patch must be installed before using this command. This PTF or patch is required to complete this command successfully.

User response: Install the required PTF or patch and and enter this command again.

DBI1064E The user *user* is being used by the DB2 instance and thus can not be used to create the DB2 Administration Server.

Explanation: The user that is being used to create the DB2 Administration Server is being used by the DB2 instance. You should not create the DB2 Administration Server with this user.

User response: Use a different user name that is not being used by the DB2 instance to create the DB2 Administration Server.

DBI1065E Program *program-name* terminated prematurely.

Explanation: An error has occurred during the execution and that has caused this program to terminate prematurely.

User response: Correct the problem and try the command again.

DBI1066E Program *program-name* terminated by request.

Explanation: This program has been terminated by the user.

User response: Issue the same command to restart the program.

DBI1067E DB2 Product Library for locale *locale-name* is not installed.

Explanation: No compressed-tar HTML files can be found in the DB2DIR/doc/"<locale-name>"/html directory, where

DB2DIR =
 /usr/opt/db2_08_xx, where xx is 01 or FPn,
 where n is the FixPak number
 or /opt/IBM/db2/V8.x, where x is 1 or FPn,
 where n is the FixPak number

User response: Install the DB2 Product Library for the required locale and issue this command again.

DBI1068E There is not enough space to uncompress and un-tar all the HTML file in the directory *dir-name*.

Explanation: After uncompress and un-tar, all the HTML files cannot be saved in the indicated directory because the file system is full.

User response: Increase the size of the file system or free up sufficient disk space in the file system. Issue the command again.

DBI1069E Unexpected error. Function = *fname*, Return code = *return-code*.

Explanation: An unexpected error occurred during the execution of this program.

User response: If the problem persists, contact IBM Support with the following information:

- Message number
 - Function name
 - Return code
 - Problem description
-

DBI1070I Program *program-name* completed successfully.

DBI1072I db2licm failed to update the nodelock file. Refer to the log file *log-name* for more information.

Explanation: db2licm failed to add the license to the nodelock file, so DB2 will run with a Try-and-Buy license until a product license is installed.

User response: Try the command again and if it continues to fail, manually add the license key to the nodelock file.

DBI1073E The parameter *parameter_name* is missing. This parameter is required for the *installer_name* command.

DBI1074E Administration server does not exist.

Explanation: An attempt was made to update or upgrade an administration server. An administration server does not exist.

User response: Create an administration server using the dasprt command.

DBI1075E Administration server cannot be migrated.

Explanation: The administration server cannot be migrated. The administration server is running at a level which is not a supported DB2 migration level.

User response:

- If the administration server is running the same version as DB2, use dasupdt to update the administration server.
 - If the administration server is running an unsupported migration level, drop the administration server using the dasdrop command, and re-create the administration server using the dasprt command at the current DB2 level.
-

DBI1076E The installation of DB2 pureScale failed because no supported RDMA capable network adapters were found.

Explanation: The DB2 pureScale Feature requires the use of Remote Direct Memory Access (RDMA) capable network adapters to provide a low latency cluster interconnect. A supported RDMA adapter was not found.

User response: Consult the documentation about installing DB2 pureScale for information on network considerations in the installation prerequisites.

DBI1077E The user *user* is being used by the DB2 Administration Server and thus can not be used to create the DB2 instance.

Explanation: The user that is being used to create the DB2 instance has been used by the DB2 Administration Server. You should not create the DB2 instance with this user.

User response: Use a different user name that is not being used by the DB2 Administration Server to create the DB2 instance.

DBI1078E Updating the IBM Tivoli System Automation for Multiplatforms (SA MP) failed. See the log file *log-file-name* for details. (TSAMP_RC=TSAMP_RC).

Explanation: The DB2 installer uses the installSAM utility to update the SA MP. The installSAM utility returned errors. The installSAM log file contains more complete information.

TSAMP_RC is the return code from the installSAM utility.

Note that you must have root authority to use the DB2 installer to update SA MP.

SA MP was not updated.

User response: See the mentioned log file for details.

To update SA MP manually, use the installSAM command.

For more information about the installSAM command, see the SA MP Base Component documentation.

DBI1079I Output is saved in the log file *log-name*.

Explanation: All processed and failed operations have been saved into this log file.

User response: Do not modify this file in any way. This file is for IBM Technical Support reference.

DBI1080E Disk full. Minimum space required in *dir-name* is KB KB, whereas only KB KB space is available.

Explanation: There is not enough free space available in the file system or directory.

User response: Free more disk space and retry the command.

DBI1081E The file or directory *file-name* is missing.

Explanation: A file or directory that is required to process this command is missing.

User response: If the file belongs to one of the database products then verify that the product is installed correctly and reinstall the product if necessary.

If the file belongs to an instance then the instance may have already been removed or it may be corrupted.

Check to see if the name of the given directory exists on your file system. If so, examine the name to see if it has been misspelled. Note that all file and directory names are case-sensitive on Unix.

DBI1082E The file or directory *file-name* already exists.

Explanation: A file or directory that the command needs to create already exists.

User response: Examine the specified file or directory. If the file or directory exists as a result of a previous successful completion of the command, no action is required. Otherwise, you need to either rename or remove the specified file or directory before trying the command again.

DBI1083E An attempt to remove *file-name* failed.

Explanation: An error occurred when removing the given file or directory. This may be due to an inappropriate ownership on this given file or directory.

User response: Adjust the file permissions or ownership of the specified file or directory and then issue the command again.

DBI1084E An attempt to create *file-name* failed.

Explanation: An error occurred when creating the given file or directory.

User response: Check that there is enough space in the file system and that you have write permission for the required directory. Suggested directory permissions are u=rwx,go=rx.

DBI1085E An attempt to move the file or directory *file-name* to *file-name* failed.

Explanation: An error occurred when attempting to relocate the given file or directory. The command could not be processed successfully.

User response: Determine why the specified file or directory could not be moved and then try the command again.

DBI1086E An attempt to copy the file or directory *file-name* to *file-name* failed.

Explanation: An error occurred when attempting to duplicate the given file or directory. The command could not be processed successfully.

User response: Determine why the specified file or directory could not be copied and then try the command again.

DBI1087E An attempt to create the link *filename* failed.

Explanation: The given link cannot be created. It could be that the write permission on its parent directory has been turned off, or that a file or directory with the same name already exists.

User response: Check permissions for the parent directory and restart the install process. Suggested directory permissions are u=rwx,go=rx.

DBI1088E Invalid access permission detected for directory *directory*.

Explanation: This message can be returned for multiple reasons, including the following reasons:

- The given name is not a directory, or the directory does not have read and execute permission.
- The given directory can be created but cannot be accessed.
- The permissions on any of the parent directories of the given directory do not meet requirements for DB2 database instance and database directories.

If the home directory of the DB2 database manager instance is located on a mounted drive, this message can be returned if the DB2 database install utility fails to validate the permissions of the mount point.

User response: Check permissions for the given directory and restart the install process. Suggested directory permissions are u=rwx,go=rx.

If the home directory of the DB2 database manager instance is located on a mounted drive, perform the following troubleshooting steps:

1. Verify the access permissions of the home directory of the instance owner by using the command "su".

Example:

```
/usr/bin/su <instance_owner> -c /bin/pwd
```

2. If the permissions do not include "r" and "x", unmount the drive, change the permissions, and remount the drive.

Example:

```
umount /<directory>
chown <instance_owner> /<directory>
chmod 755 /<directory>
mount /<directory>
```

DBI1089E Error in saving current instance related information in the *backup-dir* directory.

Explanation: There may be several reasons for this error. It is likely that one of the following situations has occurred:

- The given directory does not have sufficient permissions or does not have the write permission.
- There is no space left on the file system.

User response: Take an appropriate corrective action and issue the command again.

DBI1090E **An attempt to update *parameter* in *file-name* failed.**

Explanation: An error occurred while updating the given file.

User response: Verify that all occurrences of “<parameter>” are set to “<value>” in “<file-name>”.

DBI1091E **An attempt to change permissions for *directory* failed.**

Explanation: An error occurred when changing permissions on the given file or directory. This may be due to an inappropriate ownership on this file or directory.

User response: Check permissions and ownership for the given directory and restart the install process. Suggested directory permissions are u=rwx,go=rx.

DBI1092E **An attempt to change ownership for *directory* failed.**

Explanation: An error occurred when changing ownership on the given file or directory. This may be due to an inappropriate ownership on this given file or directory.

User response: Check ownership for the given directory and restart the install process.

DBI1093E **An attempt to change group ownership for *directory* failed.**

Explanation: An error occurred when changing group ownership on the given file or directory. This may be due to an inappropriate ownership on this file or directory.

User response: Check ownership for the given directory and restart the install process.

DBI1094E **Directory access error.**

Explanation: The tool encountered an access error during an I/O operation with the directory server.

User response: Ensure the directory server is on-line and accessible over the LAN.

DBI1095W **The file or directory *name* cannot be found.**

User response: Verify that the file set/package providing the file/directory has been installed.

DBI1096E **The file or directory *file-name* already exists.**

Explanation: A file or directory that the command needs to create already exists.

User response: Examine the specified file or directory. If you no longer need this file or directory, remove it and then run the command again. Otherwise, you will need to either rename or remove the specified file or directory before trying the command again.

DBI1097E **The file or directory *file-name* does not exist.**

Explanation: A file or directory that the command requires does not exist.

User response: If the file or directory does not exist, you may have to install the software that provides the file/directory. If the file or directory does exist, check to see that the PATH variable has been set properly. Once the file or directory is there in your PATH, try the command again.

DBI1098I **Usage:**
 db2cptsa [-h | -?] [-c] [-f] [-r]

Explanation: Valid arguments for this command are:

-h | -?

Displays help information.

-c

Verifies that the DB2 High Availability (HA) scripts exist in /usr/sbin/rsct/sapolicies/db2, and that they are at the proper level.

-f

Forces a reinstall of the DB2 HA scripts in /usr/sbin/rsct/sapolicies/db2. Without this argument, if the version of the DB2 HA scripts that are already installed is the same as or higher than the version of the scripts being installed, then the installed scripts are not overwritten.

-r

Removes the directory /usr/sbin/rsct/sapolicies/db2. This directory is where the DB2 HA scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) are located. These scripts and this directory will only be removed if the SA MP is not installed.

This utility installs or updates the DB2 HA scripts in /usr/sbin/rsct/sapolicies/db2. You need these DB2

HA scripts to use the IBM Tivoli System Automation for Multiplatforms (SA MP) with the DB2 HA feature.

By default, this utility installs the DB2 HA scripts in /usr/sbin/rsct/sapolicies/db2 if they aren't already installed there, or if the version of the scripts already installed is older than the version of the scripts being installed. This utility installs or updates the DB2 HA scripts if and only if the SA MP is already installed.

This command must be run with root authority.

This command can be found on the DB2 install media in the directory:

- db2/<plat>/tsamp

where <plat> is:

- aix for DB2 for AIX
- linux for DB2 for Linux on 32-bit AMD and Intel systems (x86)
- linuxamd64 for DB2 for Linux on AMD64 and Intel EM64T systems (x86-64)
- linuxppc for DB2 for Linux on POWER (iSeries(TM) and pSeries) systems
- linux390 for DB2 for Linux on System z9 and zSeries

The command is also available at <DB2DIR>/install/tsamp directory where <DB2DIR> is the installation path of DB2 database.

User response: Reenter the command.

DBI1099I The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) were successfully installed in /usr/sbin/rsct/sapolicies/db2.

Explanation: You need DB2 HA scripts to use the IBM Tivoli System Automation for Multiplatforms with the DB2 HA feature.

These DB2 HA scripts are located at /usr/sbin/rsct/sapolicies/db2. The DB2 installer detects whether these DB2 HA scripts need to be installed or updated.

The DB2 installer successfully installed the DB2 HA scripts.

User response: No action is required.

DBI1100E An instance lock is found. Another instance management command is running.

Explanation: The command failed because another instance management command (eg. db2icrt, db2idrop, or db2iupdt) is running. The instance management commands acquire an exclusive lock to avoid

conflicting simultaneous commands against the same instance.

User response: Wait for all other occurrences of the instance commands to finish executing before trying the command again. If no other instance commands are running then delete the file "instance.lock" in one of the following directories:

 /usr/opt/db2_08_xx, where xx is 01 or FPn,
 where n is the FixPak number
 or /opt/IBM/db2/V8.x, where x is 1 or FPn,
 where n is the FixPak number

DBI1101E The user name *user-name* is invalid.

Explanation: The given user name must be the login name of an existing user and cannot:

1. be more than 8 characters long.
2. start with "sql", "ibm" or "sys".
3. start with a numeral or contain characters other than a-z, _, or 0-9.

User response: Please check the "User, user ID and group naming rules" in DB2 Administration Guide for details.

Try the command again with a valid user name.

DBI1102E The DAS name is invalid.

Explanation: The DAS name must be the login name of an existing user and cannot

- be more than 8 characters long,
- start with "sql", "ibm", or "sys",
- or start with a numeral or contain characters other than a-z, \$, #, @, _, or 0-9.

User response: Issue the command with a valid DAS name.

DBI1103E The Administration Server already exists.

Explanation: The system detects that the Administration Server has already been created on your system. You can only create one Administration Server on each system.

User response: If you want to re-create the Administration Server, you have to drop the Administration Server first before re-creating it.

DBI1104E The Administration Server cannot be dropped.

Explanation: An attempt to drop the Administration Server failed. The system detects that the Administration Server does not exist.

User response: No action required.

DBI1105E The DB2 Data Links Manager Administrator already exists.

Explanation: The system detects that the DB2 Data Links Manager Administrator has already been created on your system. You can only create one DB2 Data Links Manager Administrator on each system.

User response: If you want to re-create the DB2 Data Links Manager Administrator, you have to drop the DB2 Data Links Manager Administrator first before re-creating it.

DBI1106E The DB2 Data Links Manager Administrator cannot be dropped.

Explanation: An attempt to drop the DB2 Data Links Manager Administrator failed. The system detects that the DB2 Data Links Manager Administrator does not exist or has been created under a different user ID.

User response: Verify that the name of the DB2 Data Links Manager Administrator as reported by the “dlfmist” command is correct and try again.

DBI1108E Valid group name or number *group-name* not specified for the DLFM user *user*.

Explanation: While creating a dlfm instance, or migrating a dlfm instance from DB2 V7 or earlier, the -g DLFMGid must be specified.

DLFMGid can be a numeric group id or a group name. The named group must be in the secondary group list of the named DLFM user.

If migrating a V8 or later dlfm instance, if -g DLFMGid is specified which is not the same as the DLFMGid already set up for the instance, the specified parameter will be ignored and the old Gid will be used.

User response: Set the group in the list of secondary groups for the DLFM user and reissue the command.

DBI1109E The kernel parameters must be updated before using this command.

Explanation: Some of the kernel parameters are not configured properly.

User response:

- Update all necessary kernel parameters.
- Reboot the system
- Try this command again

DBI1110I The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) were successfully updated in */usr/sbin/rsct/sapolicies/db2*.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

These DB2 HA scripts are located at */usr/sbin/rsct/sapolicies/db2*. The DB2 installer detects whether these DB2 HA scripts need to be installed or updated.

The DB2 installer successfully updated the DB2 HA scripts.

User response: No action is required.

DBI1111E The FencedID parameter must be specified using -u flag with this command.

Explanation: The FencedID parameter must be specified with this command. The FencedID parameter is set to the name of the user under which fenced User Defined Functions (UDFs) and fenced Stored Procedures will be run.

User response: Enter the command again adding the arguments “-u FencedID” where FencedID is the name of the user that you want your fenced UDFs or fenced Stored Procedures to run under.

For security reasons it is recommended that you do not use the instance name as the FencedID. However, if you are not planning to use fenced User Defined Functions (UDFs) or fenced Stored Procedures then setting the FencedID to the instance name may save you the step of having to create another user for the FencedID.

DBI1112E The FencedID parameter *fenced-id* is not valid.

Explanation: The FencedID parameter which specifies the user name under which fenced User Defined Functions and fenced Stored Procedures will be run:

- Must be the login name of an existing user.
- Cannot be set to the root or the bin user.

User response: Try the command again with a valid FencedID parameter.

DBI1113W Previous value of FencedID parameter *old-value* is different. New value *new-value* is ignored.

Explanation: The FencedID parameter entered on the command line is different from the one used before for this userid. New value for this parameter is ignored.

DBI1114W The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) were not found in */usr/sbin/rsct/sapolicies/db2*.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

The db2cpts utility expects to find these scripts, as

well as a file named "spec", in /usr/sbin/rsct/sapolicies/db2.

The DB2 installer could not find the DB2 HA scripts, or the spec file.

User response: Run the db2cpts utility with root authority to manually install the DB2 HA scripts.

DBI1115E Invalid AuthType parameter *auth-type* specified with -a flag.

Explanation: The AuthType parameter which specifies the authentication type used for the instance is not valid. Valid authentication types are:

- SERVER
- CLIENT
- SERVER_ENCRYPT

User response: Try the command again with a valid AuthType parameter.

DBI1116W The version of the DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) found in /usr/sbin/rsct/sapolicies/db2 is a lower version than the version of the same scripts on the current DB2 install media.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

SA MP and the DB2 HA feature being installed from the DB2 install media require DB2 HA scripts with a version the same as or higher than the version of the scripts also on the DB2 install media.

The version of the DB2 HA scripts currently installed is lower than the version of the scripts on the DB2 install media.

User response: Run the db2cpts utility with root authority to manually update the DB2 HA scripts.

DBI1117I The version of the DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) found in /usr/sbin/rsct/sapolicies/db2 is a higher version than the version of the same scripts on the current DB2 install media.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

SA MP and the DB2 HA feature being installed from the DB2 install media require DB2 HA scripts with a version the same as or higher than the version of the scripts also on the DB2 install media.

The version of the DB2 HA scripts currently installed is higher than the version of the scripts on the DB2 install media.

User response: No action is required.

DBI1118W DB2 has created a .profile file under the home directory of user *name* (instance owner or DAS) that contains the entries to set up a default environment (instance or DAS).

Explanation: The user ID for the instance or DAS *name* is using the Korn Shell but does not have a .profile file under its home directory. The DB2 installation script has created the .profile file for the user ID and has populated the file with entries necessary for setting the default instance or DAS environment. This is done so that when the user logs in, the user will automatically acquire the instance or DAS environment.

User response: If you are using the Korn Shell and want to manually set up the environment of the instance or DAS *name*, you can remove the .profile created by the DB2 script and manually set up the instance or DAS environment.

DBI1119I The version of the DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) found in /usr/sbin/rsct/sapolicies/db2 is the same version as the version of the scripts on the current DB2 install media.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

SA MP and the DB2 HA feature being installed from the DB2 install media require DB2 HA scripts with a version the same as or higher than the version of the scripts also on the DB2 install media.

The version of the DB2 HA scripts currently installed the same as the version of the scripts on the DB2 install media.

User response: No action is required.

DBI1120E Instance *inst-name* cannot be dropped.

Explanation: An attempt was made to drop an instance which may still be active.

DBI1121E Cannot drop the Administration Server using "db2idrop".

Explanation: An attempt was made to drop the Administration Server using "db2idrop". This operation is not allowed since it will cause the registry to be in an inconsistent state.

User response: Use the "dasidrop" command to remove the Administration Server.

DBI1122E Instance *inst-name* cannot be updated or upgraded.

Explanation: An attempt was made to either update or upgrade an instance. This instance cannot be updated or upgraded because:

- One of the following commands cannot be used to either update or upgrade this instance: db2iupdt, db2iupgrade, or db2nrupdt.
- The instance is still active.

User response: Ensure that you are using the correct version of one of the following commands: db2iupdt, db2iupgrade, or db2nrupdt. Also, ensure that there are no DB2 processes running at the instance. Retry the command.

DBI1123E Cannot update the Administration Server using “db2iupdt”.

Explanation: An attempt was made to update the Administration Server using “db2iupdt”. This operation is not allowed since it will cause the registry to be in an inconsistent state.

User response: Use the “dasiupdt” command to update the Administration Server.

DBI1124E Instance *inst-name* cannot be upgraded.

Explanation: An attempt was made to upgrade an instance. This instance cannot be upgraded because:

- The instance is still active.
- Upgrading this instance is not supported.
- This version of the db2iupgrade command cannot be used to upgrade this instance.

User response: Ensure that the instance is ready for upgrading and you are using the correct version of the db2iupgrade command. For more information on upgrading an instance, refer to the DB2 Information Center.

DBI1125E Cannot drop the DB2 Data Links Manager Administrator using “db2idrop”.

Explanation: An attempt was made to drop the DB2 Data Links Manager Administrator using “db2idrop”. This operation is not allowed since it will cause the registry to be in an inconsistent state.

User response: Use the “dlfmdrop” command to remove the DB2 Data Links Manager Administrator.

DBI1126W Instance upgrade failed because one or more views defined in the database might impact database upgrade.

Explanation: The db2iupgrade command calls the db2ckupgrade command. The instance upgrade fails if

the db2ckupgrade command returns any errors.

The db2ckupgrade command has identified one or more views that are dependent on database entities that have changed. This might cause the UPGRADE DATABASE command to fail.

The identified views are listed in the db2ckupgrade log file.

User response: Resolve the issue with the identified views or drop the views and then re-issue the re-issue the db2iupgrade command. After database upgrade, re-create the views that you dropped.

Alternatively, you can upgrade a test database that contains the identified views in a testing environment to determine if the UPGRADE DATABASE command will fail. If the database upgrade is successful, re-issue the db2iupgrade command with the -F parameter to force the upgrade.

DBI1127E The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) found in /usr/sbin/rsct/sapolicies/db2 were not removed because the SA MP was still installed.

Explanation: You need DB2 HA scripts to use the SA MP with the DB2 HA feature.

You cannot uninstall these scripts while the SA MP is still installed. This is to prevent unexpected behaviour resulting from using the SA MP with the DB2 HA feature without the scripts being installed.

The DB2 HA scripts were not uninstalled.

User response:

- If you intend to manually uninstall the SA MP as well as the DB2 HA scripts, uninstall the SA MP first using the uninstallSAM utility.
- Alternatively, use db2_deinstall with the -a and -F TSAMP arguments to uninstall the SA MP and the DB2 HA scripts, because db2_deinstall will uninstall these in the correct order.

DBI1128E Cannot drop the Administration Server using “dlfmdrop”.

Explanation: An attempt was made to drop the Administration Server using “dlfmdrop”. This operation is not allowed since it will cause the registry to be in an inconsistent state.

User response: Use the “dasidrop” command to remove the Administration Server.

DBI1129E Cannot update the Administration Server using “dlfmupdt”.

Explanation: An attempt was made to update the Administration Server using “dlfmupdt”. This operation is not allowed since it will cause the registry to be in an inconsistent state.

User response: Use the “dasiupdt” command to update the Administration Server.

DBI1130E The IBM Tivoli System Automation for Multiplatforms (SA MP) could not be installed or updated because system prerequisites were not met. See the log file *log-file-name* for details.

Explanation: There are system prerequisites for installing or updating SA MP. These prerequisites have not been met.

The mentioned log file has the details for the unsatisfied prerequisites.

To find out more about the system prerequisites for the SA MP Base Component, see the SA MP Base Component documentation's Installation and Configuration Guide.

SA MP cannot be installed or updated.

User response: Modify your system to meet the prerequisites; and use the installSAM utility to manually install SA MP.

DBI1131E The user ID *user-id* is invalid.

Explanation: An attempt to access the given user ID failed. One of the following situations has occurred:

- This user ID does not exist on the system.
- The home directory of this user is not set up properly.
- One of the user attributes needed by DB2 is unset.
- The UID of this user is 0

User response: Make sure a valid user ID with valid home directory, shell, primary group and secondary group has been used. Create a new user if necessary.

DBI1132E The DB2 installer could not find the correct license to install or update the IBM Tivoli System Automation for Multiplatforms (SA MP). The DB2 installer expected the license file to be found at *directory-name* on the DB2 install media.

Explanation: The regular license for SA MP is named sam31.lic, and the try-and-buy license is named sam31tb.lic. One, but not both, of these files must exist in the specified directory on the DB2 install media for the DB2 installer to install or update SA MP.

SA MP cannot be installed without the proper license.

User response: Ensure the appropriate license file is in the specified directory and reissue the command.

DBI1133E Command failed. The current host is part of a GPFS cluster.

Explanation: The current host is part of an IBM General Parallel File System (GPFS) cluster. You cannot remove the GPFS cluster from the host on which the command was run. Also, the uninstallGPFS command is not a supported command.

User response: To remove the GPFS file system and cluster on the remaining host, run the db2cluster command. For details about the steps to manually removing a clustered file system, see the DB2 Information Center.

DBI1134W The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) were not installed because the SA MP was not installed.

Explanation: You need DB2 HA scripts to use the SA MP with the DB2 HA feature. You do not need these scripts if you are not using the SA MP.

The DB2 HA scripts were not installed.

User response: No action is required.

DBI1135E The primary group *users* of user ID *test-user* is invalid.

Explanation: The primary group name of the instance user ID:

1. Cannot be one of "guests", "admins", "users" or "local".
2. Cannot start with "sql" or "ibm".
3. Cannot start with a numeral number.
4. Can contain only lowercase letters (a-z), numbers (0-9), and the underscore character (_).

User response: Ask a user with root privilege to change the primary group of the instance user ID to a valid group as described in this message's Explanation and then try the command again.

DBI1136W OLAP Start Kit not installed.

Explanation: The instance you are migrating has OLAP functionality. OLAP Starter Kit is not available in this version of DB2. The instance will be migrated only if the -F parameter was specified.

User response: If program aborts, rerun the same db2imigr command with the -F parameter to force migration. OLAP functionality will not be present.

DBI1137W Server Product not installed.

Explanation: The instance you are upgrading is a Server instance. A DB2 Server Product has not been detected. The instance will be upgraded only if the -F parameter is specified.

User response: If the db2iupgrade command fails, reissue this command with the -F parameter to force the upgrade. The instance will not have the DB2 Server functionality. To re-enable DB2 Server functionality in the instance, install a DB2 Server Product and issue the db2iupdt command to update the instance.

DBI1138W Query Patroller not installed.

Explanation: The instance you are upgrading has Query Patroller functionality. Query Patroller was not detected in the installed DB2 product.

User response: If the db2iupgrade command fails, reissue the command to force the upgrade. The instance will not have the Query Patroller functionality. To re-enable Query Patroller functionality in the instance, install Query Patroller and issue the db2iupdt command to update the instance.

DBI1139E Data-Links File Manager is not installed.

Explanation: The instance you are upgrading has Data-Links functionality. Data-Links File Manager is not supported in the version of the DB2 product you are upgrading to.

DBI1140W Command is being Forced. You have specified the -F flag to force the issued command past the point of failures mentioned in the previous warnings.

User response:

DBI1141E Visual Warehouse 3.1 instance detected. Visual Warehouse 3.1 instance migration to V8 is not supported. Only Visual Warehouse 5.2 and above can be migrated to V8.**DBI1142W Visual Warehouse 5.2 instance detected. Instance will be migrated but warehouse metadata will not be migrated.****DBI1143W Relational Connect is not installed.**

Explanation: Relational Connect was not detected in the version you are migrating to, but a setting file for DB2 and/or non-DB2 datasources was detected for the instance you are migrating. This configuration will be migrated but Relational Connect for DB2 has to be installed to make use of these configuration parameters.

User response: Install Relational Connect for DB2 to make use of the configuration parameters.

DBI1144E Relational Connect is not installed.

Explanation: Relational Connect was not detected in the version you are migrating to but environment variables have been found that indicate that you have non-DB2 data sources. This instance cannot be migrated if Relational Connect is not installed.

User response: Please install Relational Connect and try the command again.

DBI1146E DB2 Instance Setup wizard is not installed.

Explanation: The db2isetup script can perform configuration and instance setup tasks using a response file, or it can launch the DB2 Instance Setup wizard to perform these tasks using a graphical user interface. Currently, only the response file support mode is installed. To use the graphical mode, the DB2 Instance Setup wizard feature must be installed.

User response:

- Install the DB2 Instance Setup wizard by launching the DB2 Setup wizard from the CD (not available as part of the Runtime Client).
- Create a response file and reissue the command as db2isetup -r <responsefile>.

DBI1147E db2iupdt failed to update the user ID *userid*.

Explanation: The user ID specified is not a DB2 instance.

DBI1148E dasupdt failed to update the user ID *userid*.

Explanation: The user ID specified is not an Administration Server.

DBI1149E To execute this program, you must be the owner of the installation copy.

Explanation: The current DB2 copy was not installed by the user who is running the program.

User response: Log in as the user who installed the current copy of DB2 and rerun the command.

DBI1150W db2iupdt has updated this instance to a 64-bit non-server instance.

Explanation: You cannot create or upgrade to a 64-bit Express server instance. If you want to run a 64-bit server instance, you must have the DB2 Server Edition installed.

User response: If you do not require a server instance,

you can continue to use the created or upgraded instance. If you want to use a 64-bit server instance, you should install DB2 Server Edition, and run the following command:

```
db2iupdt -w 64 <instance-name>
```

DBI1151E The db2icrt command or the db2iupgrade command has created a 64-bit non-server instance.

Explanation: You cannot create or upgrade to a 64-bit Express server instance. If you want to run a 64-bit server instance, you must have the DB2 Server Edition installed.

User response: If you do not require a server instance, you can continue to use the created or upgraded instance. If you want to use a 64-bit server instance, you should install DB2 Server Edition, and run the following command:

```
db2iupdt -w 64 <instance-name>
```

DBI1152E To run this program, you can not be root user.

Explanation: *tool-name* can not be run by users with root authority.

User response: Rerun the command as a non-root user.

DBI1153E AIO is not enabled.

Explanation: AIO needs to be enabled on the system to run DB2.

User response: Ask a user with root privileges to enable AIO and try the command again.

DBI1154E *instance-name* is an instance that belongs to a DB2 copy installed by a non-root user. The instance cannot be updated from a DB2 copy installed by a root user. If you want to update the instance, run db2nrupdt from the DB2 copy installed by the non-root user.

DBI1155E The command failed for the DB2 non-root instance *instance-name*.

Explanation: The command failed because you tried one of the following actions, which are not supported:

1. Update the DB2 non-root instance to become a DB2 root instance
2. Upgrade the DB2 non-root instance to become a DB2 root instance
3. Drop the DB2 non-root instance by running the db2idrop command from a DB2 root installation location

User response:

1. To update the DB2 non-root instance, run the db2nrupdt command in the \$HOME/sqllib/instance/ directory, where \$HOME is the home directory of the user ID that installed the DB2 non-root instance.
2. Upgrading to a DB2 root instance is not supported for DB2 non-root instances. Refer to the DB2 Information Center for details on how to upgrade databases from DB2 root instances to non-root instances.
3. To drop the DB2 non-root instance, uninstall the DB2 non-root installation by running the db2_deinstall command in the command \$HOME/sqllib/install/ directory, where \$HOME is the home directory of the user ID that installed the DB2 product.

DBI1156I System prerequisites cannot be checked completely when installing as a non-root user.

Explanation: A complete system check cannot be performed because root authority is required.

User response: A non-root user must check system prerequisites from the DB2 documentation. A system administrator must make sure the system meets all the requirements to install and use the current DB2 level.

DBI1157E Invalid value.

Explanation: The *parameter-or-value* passed to the installer command *cmd* is not valid.

User response: Correct the error and run the command again.

DBI1158I DB2 will be installed to *installpath*.

DBI1159E db2iupdt can not update an instance from using a non-root installation to using a root installation or vise versa.

DBI1160I Non-root install is being performed.

DBI1161I Usage:
db2rfe [-h | -?] [-d] [-l *logfile_name*] [-q] -f <db2rfe_configuration_file>

Explanation:

- h | -? displays the usage information.
- d turns debug mode on.
- l logfile_name is a user specified log file.
- q query which installed root features are enabled

or not enabled or which configured service ports are reserved or not reserved. The -q option should not be used with the -f option.

-f is mandatory option. db2rfe_configuration_file is the full path name of db2rfe's configuration file.

Note: This command must be run from the \$DB2DIR/instance directory. If the db2rfe command detects any root features enabled or any port numbers reserved and used by the DB2 instance, a summary report shows the following information:

- The installed root features that have been or have not been enabled
- The DB2 service ports used by the DB2 instance that have been or have not been reserved in the system's services file.

DBI1162I Usage:

Explanation: You can configure non-root installations by using the db2nrcfg utility.

Syntax:

```
db2nrcfg [-h|-?]
          [-d]
          [-a AuthType]
          [-p PortName]
          [-s InstType]
```

User response: Call the db2nrcfg command, specifying valid parameters.

DBI1163I An invalid argument was entered for the db2nrupdt command.

Explanation: The syntax for the db2nrupdt command is:

```
db2nrupdt [-a AuthType]
           [-d]
           [-k ]
           [-j "TEXT_SEARCH" |
            -j "TEXT_SEARCH,portnumber"]
           [-h|-?]
```

For details about the parameters of the db2nrupdt command, refer to the DB2 documentation for Linux, UNIX, and Windows.

This command must be run from the \$HOME/sqlllib/instance directory where the \$HOME is the home directory of the non-root DB2 instance ID. This command can be run only as the non-root DB2 instance who owns the non-root DB2 installation at \$HOME/sqlllib.

User response: Enter the command again using valid parameters or run the command from the correct directory.

DBI1164E A value is needed for entry-name.

Explanation: To enable the remote connection, please provide a value for *entry-name* in the configuration file.

DBI1165E Instance *instance-name* is not accessible on the following database partitions: *database-partitions-list*

DBI1166E The version string for the file set contains invalid characters. Version string: *version*. File set: *fileset*.

Explanation: The file name of the specified file set has been modified or corrupted.

User response: Download the fix pack installation package again and re-run the command.

DBI1167I An invalid argument was entered for the db2ckgpfs command.

Explanation: The syntax for the db2ckgpfs command is:

```
db2ckgpfs [-d]
           [-l <logFile>]
           [-v install | media]
           [-s install | media]
           [-c]
           [-n <cluster-name>]
           [-h|-?]
```

For details about the parameters of the db2ckgpfs command, refer to the DB2 Information Center.

User response: Enter the command again using valid parameters.

DBI1168W While attempting to create or update a 64-bit instance, it was detected that the installed level of bos.rte.libc is lower than the minimum requirement for 64-bit instances.

Explanation: The installed level of bos.rte.libc is lower than 5.1.0.28, the minimum required level for supporting 64-bit instances.

User response: Before proceeding, download APAR IY32466 and update your system.

Information on downloading APARs is available from <http://www.ibm.com/aix>.

You can reissue the command with a -F parameter to force execution past the point of warning.

DBI1169E The DB2 HTML Documentation is not installed.

Explanation: The DB2 HTML Documentation is not installed or it cannot be found.

To update the DB2 HTML Documentation to this FixPak, the DB2 HTML Documentation must be installed.

User response: Install the DB2 HTML Documentation and try the command again.

DBI1170I An invalid argument was entered for the db2cktsa command.

Explanation: The syntax for the db2cktsa command is:

```
db2cktsa [-d]
          [-l <logFile>]
          [-v install | media]
          [-i install | media]
          [-c]
          [-n <cluster-name>]
          [-h|-?]
```

For details about the parameters of the db2cktsa command, refer to the DB2 Information Center.

User response: Enter the command again using valid parameters.

DBI1171E DB2 64-bit support is not installed.

Explanation: 64-bit instances are not supported or implicit on this platform.

User response:

- For Linux IA64, reissue the command without the -w option to create a 64-bit instance.
- For other platforms, reissue the command without the -w 64 option to create a 32-bit instance.

DBI1172W While attempting to create or update a 64-bit instance, it was detected that the installed level of one or both of bos.rte.libc and bos.adt.prof is lower than the minimum requirement for 64-bit instances.

Explanation: The installed level of one or both of bos.rte.libc and bos.adt.prof is lower than 4.3.3.50, the minimum required level for supporting 64-bit instances.

User response: Before proceeding, download one or both of the following PTFs and update your system.

- For bos.rte.libc, the PTF is U473728.bff
- For bos.adt.prof, the PTF is U473581.bff

Information on downloading PTFs is available from <http://www.ibm.com/aix>.

You can reissue the command with a -F parameter to force execution past the point of warning.

DBI1173W db2iupdt has updated this instance to a 64-bit non-server instance.

Explanation: You can only update a 32-bit Workgroup server instance to a 64-bit server instance if you have DB2 Server Edition installed.

User response: If you do not wish to update to a server instance, you can continue to use the updated instance. If you wish to change this instance back to a 32-bit Workgroup server instance, run the following command:

```
db2iupdt -w 32 <instance-name>
```

If you wish to use a 64-bit server instance, install DB2 Server Edition and run the following command:

```
db2iupdt -w 64 <instance-name>
```

DBI1174E The GPFS version installed is too low. Installed version: *installed-version*. Media version: *media-version*.

Explanation: The IBM General Parallel File System (GPFS) version level installed on your system is lower than the GPFS version level on the installation media. GPFS cannot be automatically upgraded. Also, the installGPFS command is not a supported command.

User response: Manually remove the GPFS file system and cluster, uninstall GPFS, then install the fix pack with the installFixPack command. For details, see the DB2 Information Center.

DBI1175W While attempting to create or update a 64-bit instance, it was detected that the installed level of the Solaris Operating Environment is lower than the minimum requirement for 64-bit.

Explanation: The installed level of the Solaris Operating Environment is lower than 5.7, the minimum required level for supporting 64-bit instances.

User response: Before proceeding, update your system to the minimum required level of the Solaris Operating Environment.

You can reissue the command with a -F parameter to force execution past the point of warning.

**DBI1176I Usage:
installAltFixPak [-h] [-s] [-y]**

Explanation:

- h display the usage information
- s install the exact same file sets/packages to the alternate path (for supported products) as those installed in /usr/opt/db2_08_01 or /opt/IBM/db2/V8.1.

If installAltFixPak detects that DB2 is not installed in /usr/opt/db2_08_01 or /opt/IBM/db2/V8.1, the -s option will not be used. In that case, it will call the db2_install utility.

- y You have reviewed and accepted the license terms and conditions located in the db2/license directory for all the DB2 Version 8 products on your system.

Please confirm that you have reviewed the license terms and conditions located in the db2/license directory for the installed products. If you accept these license terms and conditions, restart the FixPak installation and specify the '-y' option.

By specifying the '-y' option, you indicate your agreement to the license terms and conditions provided with this product.

DBI1177W Updating an instance to a code level that is lower than the current level used by the instance.

Explanation: An attempt has been made to update an instance to a code level that is lower than the one being used by the instance currently.

User response: To update an instance to a code level that is lower than the current level, issue the db2iupdt command with the -f level option as follows:

```
db2iupdt -f level
        <instance-name>
```

DBI1178W Updating a DAS to a code level that is lower than the current level used by the DAS.

Explanation: An attempt has been made to update a DAS to a code level that is lower than the one being used by the DAS currently.

User response: To update a DAS to a code level that is lower than the current level, issue the dasupdt command with the -D option as follows:

```
dasupdt -D
```

DBI1179E Client instances cannot be created using the db2icrt command when the PortName parameter is specified.

Explanation: The PortName parameter is used for inbound TCP/IP connections and is not applicable to client instances.

User response: Reissue the db2icrt command without specifying the PortName parameter.

DBI1180E 32-bit server instances are not supported.

Explanation: The current platform does not support 32-bit server instances.

User response: You cannot specify the bit size for the instance when you create or upgrade an instance. The bit size for new instances is determined by the operating system where DB2 database is installed.

DBI1181E DB2 64-bit support is not installed.

Explanation: 64-bit instances are not supported on this platform.

User response: Before proceeding, either

- update your system to the minimum required level to run 64-bit DB2, and reissue the command again, or
- drop the instance you are upgrading or updating, and re-create it as a 32-bit instance.

DBI1182W Either DB2 Warehouse Manager, or DB2 Warehouse Manager Connectors, or both are not supported for upgrade in this release.

Explanation: The instance you are trying to upgrade has Warehouse Manager or Warehouse Manager Connectors functionality. Warehouse Manager and Warehouse Manager Connectors are not supported in the product release to which you are trying to upgrade. A successful upgrade means that the upgraded instance will not have these capabilities. If this is the expected outcome, you can re-issue the db2iupgrade command.

User response: Re-issue the db2iupgrade command to force the upgrade. Either the Warehouse Manager, or the Warehouse Manager Connectors, or both will not be available once the upgrade is completed.

DBI1183W Spatial Extender not installed.

Explanation: The instance you are upgrading has Spatial Extender functionality. Spatial Extender was not detected in the installed DB2 product.

User response: If the db2iupgrade command fails, reissue the command to force the upgrade. The instance will not have the Spatial Extender functionality. To re-enable Spatial Extender functionality in the instance, install Spatial Extender and issue the db2iupdt command to update the instance.

DBI1184W Life Sciences Data Connect not installed.

Explanation: The instance you are migrating has Life Sciences Data Connect functionality. Life Sciences Data Connect was not detected in the installed DB2 product.

User response: Rerun the db2iupgrade command to force migration. Life Sciences Data Connect will not be present. You must install Life Sciences Data Connect and manually run db2iupdt on the instance you are migrating in order to re-enable Life Sciences Data Connect functionality.

DBI1185I **Server protocol *protocol* is no longer supported. It has been removed from DB2COMM.**

DBI1186I **Usage:**
db2cdbcr [-d] -n CDBName

Explanation: An incorrect argument was entered for the db2cdbcr command. Valid arguments for this command are:

-h | -? display the usage information

-d turn debug mode on.

-n CDBName

CDBName is the name of the Warehouse Control Database you want to create.

User response: Enter the command again as follows:
db2cdbcr -n CDBName

DBI1187E **Unable to complete DB2 Warehouse Control Database setup.**

Explanation: An attempt to create the DB2 Warehouse Control Database has failed.

User response: A database partition groups (created on database partition 0) and a table space (both called FLG32K) needs to be created as well. See DB2 DataWarehouse documentation for details.

DBI1188E **Migration failed. The group *group-name* is not in the secondary group list of the DLFM user *user-name*.**

User response: Set the group *group-name* in the list of secondary groups for the DLFM user *user-name* and reissue the command.

DBI1189E **There has been an attempt to use db2_deinstall on an image for a platform that does not match the current platform *platform* on which it is being run.**

Explanation: Possible causes include:

- This DB2 install image is not valid for the current platform.
- The current platform is not supported by DB2

User response: Install DB2 using the DB2 install image that corresponds with the current platform *platform*

DBI1190I **db2setup is preparing the DB2 Setup wizard which will guide you through the program setup process. Please wait.**

DBI1191I **db2setup is installing and configuring DB2 according to the response file provided. Please wait.**

DBI1192I **DB2 installation completed successfully. The install logs db2setup.log and db2setup.err are location in *location*.**

DBI1193W **DB2 installation completed with warnings. A minor error occurred while installing this computer. Some features may not function correctly. The install logs db2setup.log and db2setup.err are located in *location*.**

DBI1194E **An error occurred while installing DB2 on this computer. The installation cannot continue. The install logs db2setup.log and db2setup.err are located in *location*.**

DBI1195E **db2nrupt failed to update user *user-name*.**

Explanation: The instance can not be updated because it is still active.

User response: Make sure that the instance is stopped. Then rerun the command.

DBI1196E **The DB2 installer cannot detect the operating system on host *host-name*. System command used: *system-command*.**

Explanation: The DB2 installer cannot continue. Either the system command was run on an unsupported operating system, or, there is a problem with the specified system command.

User response: Ensure you are running on a supported operating system. For a list of supported operating systems, see to the DB2 Information Center.

If you are running the specified command on a supported operating system, ensure the command is functioning correctly. Contact your system administrator.

DBI1197E **An error occurred acquiring the processor type information on host *host-name*. System command used: *system-command*.**

Explanation: The DB2 installer cannot continue. Either the system command was run on an unsupported

processor type, or, there is a problem with the specified system command.

User response: Ensure you are running on a supported processor type. For a list of supported processor types, see the DB2 Information Center.

If you are running the specified command on a supported operating system, ensure the command is functioning correctly. Contact your system administrator.

DBI1198E An error occurred acquiring the Linux distribution on host *host-name*.

Explanation: The DB2 installer could not determine the Linux distribution on the specified host. This can happen if there is a problem with the Linux release files.

User response: Ensure that either the /etc/issue/release files, or the /etc/*release files, or both, are installed correctly. For a list of supported Linux distributions, see to the DB2 Information Center.

DBI1199I DB2IDROP completed successfully. The DB2INSTDEF registry variable was set to the dropped instance; it has been cleared. This registry variable should be explicitly set to another valid instance name.

DBI1200E The directory *directory-name* is not a version *version-number* instance.

Explanation: The command cannot proceed because the specified directory does not point to the required version of the database products.

User response: The command cannot be run with the specified instance.

DBI1201E The specified hosts are not running on the same type of operating systems. Host: *host-name1*. Operating system: *operating-system1*. Installation-initiating host: *host-name2*. Operating system: *operating_system2*.

Explanation: All hosts in the same cluster must be running on the same type of operating system.

User response: Ensure all hosts are created on the same operating system. For a list of supported operating systems, see to the DB2 Information Center.

DBI1202E Upgrading of instance *inst-name* is not supported.

Explanation: The instance cannot be upgraded because:

1. Upgrading from this version of the instance is not supported.
2. The instance is already using the current version of the product and upgrade is not required.

User response: The db2level command shows the current version and service level of the installed DB2 product. The db2ls command lists where DB2 products are installed on your system and list the DB2 product level.

Make sure that the instance is valid for upgrading and try the command again with a valid instance name.

DBI1203I The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) were successfully uninstalled.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

These scripts have been removed.

User response: No action is required.

DBI1204E The keyword *keyword* has an invalid value. Sample configuration file name: *file_name*

Explanation: The sample configuration file has examples of valid values for this keyword.

User response: Correct the error and rerun the command.

DBI1205E One or more local databases cannot be upgraded to the version from which you are running this command. Check the log file *logfile-name* for the list of errors.

Explanation: The following list provides the reasons why a database cannot be upgraded and the corrective action to take:

- Database is in Backup pending state. Perform backup of the database.
- Database is in Roll-forward pending state. Perform a roll-forward database to end of logs and stop.
- Database is in Restore pending state. Perform restore of the database.
- Database is inconsistent. Restart database to return it to a consistent state.
- Database has one or more table space not in normal state. Perform a roll-forward database.
- Database contains database objects that have a schema name of SYSCAT, SYSFUN, SYSIBM, SYSPUBLIC, or SYSSTAT. Drop the object and re-create the object with a correct schema name (qualifier). If the object is a table you should first

export its data, drop the table, re-create the table with a correct schema name, and then import or load the data to the new table.

- Database contains database objects that have a dependency on the SYSFUN.DIFFERENCES function. Possible violated objects and the associated action to correct the violation are:
 - constraint - alter table to drop the constraint
 - function - drop the function
 - trigger - drop the trigger
 - view - drop the view
- Database contains user-defined data types that are reserved in the new version. Rename the data types.
- Database contains orphan rows in system catalog tables. Contact IBM technical service representative.
- Cataloged database does not exist. Create cataloged database.
- SYSCATSPACE does not have at least 50% free pages if SYSCATSPACE is a DMS table space and AUTORESIZE is not enabled. Increase available space for SYSCATSPACE table space.
- Database does not allow successful connection as an HADR primary database. Stop HADR on primary database.
- A database has HADR standby role. Stop HADR on the standby database and drop the standby database. Upgrade the HADR primary database. Reinitialize the standby database from the copy of upgraded primary database via restore or db2inidb.

User response: Determine the reason why the database cannot be upgraded and take the corresponding action specified in the explanation.

DBI1206E **Enablement of root feature *feature-name* failed. See log file *log-file-name* for details.**

DBI1207E **Reserving the service for *feature-name* failed.**

Explanation: Can not reserve the service in the services file. Possible reasons are:

- The service name and port number are used in services file.
- The service name or port number are not valid values.

User response: Correct the error and rerun the command.

DBI1208E **Reserving the service for *feature-name* failed.**

DBI1209E ***keyword* is missing in the configuration file.**

Explanation: *keyword* is needed in the configuration file to enable the root feature *feature-name*.

User response: Provide a valid value for *keyword* and rerun the command.

DBI1210E **Unknown system error.**

DBI1211E **Upgrading of the *directory-name* failed.**

Explanation: A system error such as running out of disk space or memory has been encountered during the directory upgrade. Instance upgrade failed.

User response: Make sure there is sufficient disk space and memory before you reissue this command again.

DBI1212W **Upgrade of the local database directory on *path* failed.**

Explanation: The upgrade process of the local database directory is not complete because of an unexpected error occurred. Without the local database directory, databases that are cataloged in the local database directory will no longer be accessible.

User response: Perform one of the following actions:

1. Remove the local database directory and re-catalog the databases.
2. Keep the db2mgdbd.err error file and the sqlldbidr/sqllddir.bak local database directory backup file that are created in the *path* directory and contact IBM service for instruction of possible ways to repair your directory.

DBI1213I **Root feature *feature-name* has been enabled successfully.**

DBI1214I **Service for *feature-name* is successfully reserved.**

DBI1215I **Service for *feature-name* has been reserved successfully.**

DBI1216E **Uninstalling the IBM Tivoli System Automation for Multiplatforms (SA MP) failed because the command *uninstallSAM* was not found in the expected directory: *directory*.**

Explanation: db2_deinstall uses the uninstallSAM utility to uninstall the SA MP. db2_deinstall expects to find the uninstallSAM utility in the specified directory, but did not find it there in this case. This might be because the current version of the SA MP is earlier than version 2.2. The option to uninstall the SA MP with the

db2_deinstall command is only supported for SA MP Version 2.2 or higher.

The SA MP was not uninstalled.

User response: To uninstall the SA MP manually, use the uninstallSAM command.

For more information about the uninstallSAM command, see the SA MP Base Component documentation.

DBI1217E Valid installation identifier not detected.

Explanation: All valid identifiers from the image have been removed. The installer cannot determine which valid product(s) are available on the image and cannot display the installation choices.

User response: The product image has been modified and does not have a valid identifier for the installer to display the available options. Install using the original DB2 product image.

**DBI1218I The license certificate
license-certificate-file-name for the IBM
Tivoli System Automation for
Multiplatforms (SA MP) was
successfully installed.**

Explanation: SA MP requires a valid license certificate to work with the DB2 High Availability (HA) feature. This license certificate was successfully installed or updated.

User response: No response is required.

**DBI1219E The license certificate
license-certificate-file-name for the IBM
Tivoli System Automation for
Multiplatforms (SA MP) was not
successfully installed.**

Explanation: SA MP requires a valid license certificate to work with the DB2 High Availability (HA) feature. This license certificate was not successfully installed or updated.

If you used the DB2 installer to install or update this license certificate, you can find more detailed information about why the install or update failed in the DB2 install log file.

User response: To manually install or update this license certificate for the SA MP Base Component, issue the command:

- **samlcsm** -i *license-certificate-file-name*

For more information about the **samlcsm** command, see the SA MP Base Component documentation.

**DBI1220W The upgrade of the node directory on
path failed.**

Explanation: The process of upgrading a node directory did not complete because an unexpected error occurred.

User response: Perform one of the following actions:

1. Remove the node directory and re-catalog the node entries.
2. Keep the db2ugndd.err error file and the sqlnkdir/sqlnddir.bak node directory backup file that are created in the path directory and contact IBM service for instruction of possible ways to repair your directory.

**DBI1221W The node directory is corrupted and
cannot be upgraded.**

Explanation: Either the node directory files are not the same or the node directory file is corrupted.

User response: Remove node directory and re-catalog node entry.

**DBI1222W The system database directory is
corrupted and cannot be upgraded.**

Explanation: Either the primary and backup system database directory files are not the same or system database directory is corrupted.

User response: Remove system database directory and recatalog all database entries.

**DBI1223W The local database directory is corrupted
and cannot be upgraded.**

Explanation: Either the primary and backup database directory files are not the same or the local database directory file is corrupted.

User response: An error file db2mgdbd.err and a local database directory backup sqlldbdir/sqllddir.bak are created in the indicated directory path.

Keep these two files and contact IBM service for instruction of possible ways to repair your directory. Without the local database directory, databases that are cataloged in the local database directory will no longer be accessible.

**DBI1224E The target host is not running on the
same processor type as the source host.
Source host: source-host, processor type:
source-processor-type. Target host:
target-host, processor type:
target-processor-type.**

Explanation: The source host and target host must run on the same processor type.

User response: Ensure all hosts have the same

processor type. For a list of supported processors, see to the DB2 Information Center.

DBI1225W **The authentication type of one or more cataloged local databases have been changed.**

Explanation: One or more cataloged database entries were detected with an authentication type different than the authentication type for the instance that owns these databases. If no action is taken, all cataloged local database entries will inherit the authentication type from the instance.

User response: Check the log file for the db2ckupgrade command for a list of the cataloged local database entries that do not have the same authentication type as the instance. If you want the database to maintain its previous authentication type, you can either change the authentication type of the instance, or, you can move the database to another instance that has the desired authentication type. However, before changing the authentication type of the instance, you should make sure that you want all the cataloged local database entries to have the new authentication type.

DBI1226E **The target host is not running on the same Linux distribution as the source host. Source host: *source-host*, distribution: *source-host-distribution*. Target host: *target-host*, distribution: *target-host-distribution*.**

Explanation: The hosts must be running on the same Linux distribution. The target host must have the same Linux distribution as the source host.

User response: Ensure all hosts are created on the same Linux distribution. For a list of supported Linux distributions, see the DB2 Information Center.

DBI1227E *product-name* **is not a supported DB2 product for non-root installation.**

DBI1228E **The instance *inst-name* is not a server instance.**

Explanation: Some configuration tasks can be performed against DB2 server instances only. Such configuration tasks include remote connection and DB2 text search.

User response: If the specified instance is a client instance and a DB2 server product is installed, run db2iupdt or db2nrupdt to update the client instance to a server instance. Then try the configuration task again.

DBI1229E **DB2 copy upgrade on Windows is not supported for the IBM Data Server Runtime Client, DB2 Thin Client, or IBM Data Server Driver Package installed at the location *installation-location*.**

Explanation: One of the following conditions is true:

- The selected DB2 copy for upgrade is an installation of previous releases of DB2 Run-Time, DB2 Run-Time Client Lite, DB2 Runtime Client, or Data Server Runtime Client.
- The DB2 product being installed is IBM Data Server Runtime or IBM Data Server Driver Package.

The DB2 copy upgrade is not supported from previous releases of DB2 Run-Time, DB2 Run-Time Client Lite, DB2 Runtime Client, or Data Server Runtime Client. Also, DB2 copy upgrade is not supported to IBM Data Server Runtime or IBM Data Server Driver Package.

User response: Choose "Install New" to install the IBM Data Server Runtime Client, or IBM Data Server Driver Package in a new location. For IBM Data Server Runtime Client, upgrade the existing client instances to the new DB2 copy with the db2iupgrade command.

DBI1230E **DB2 copy upgrade is not supported for the selected DB2 copy installed at *installation-location*.**

Explanation: DB2 copy upgrade is not supported for the selected DB2 copy due to one or both of the following reasons:

1. One or more DB2 instances have an instance type that is not supported by the DB2 database product that you want to install. On Windows operating systems, upgrading to a DB2 instance type of a lower level is not supported.
2. The DB2 DAS is related to the DB2 installation location from which the DB2 copy upgrade was requested, but the DB2 product to be installed does not support the DB2 DAS.

User response: Choose the DB2 product which supports all the instances or DAS related to the DB2 copy to be upgraded.

DBI1231E **Upgrading a DB2 instance is not supported for the instance *instance-name*. (The original instance type is *original-instance-type*. The instance type to be upgraded is *upgraded-instance-type*.)**

Explanation: The specified DB2 instance cannot be upgraded to the current DB2 copy. On Windows operating systems, upgrading to a DB2 instance type of a lower level is not supported.

User response: Install the DB2 product which

supports the instance type and rerun the db2iupgrade command.

DBI1232E Cannot continue with the uninstallation.

Explanation: Possible reasons:

1. The database manager for the instance *inst-name* is still active.
2. The DB2 Text Search instance service is still active.

User response: Stop the database manager and the DB2 Text Search instance service then rerun the command.

DBI1233E Instance *inst-name* cannot be dropped.

Explanation: An instance cannot be dropped if the DB2 Text Search instance service is active.

User response: Stop the DB2 Text Search instance service then rerun the command.

DBI1234E The DB2 instance update is not supported for the instance *instance-name*. (The original instance type is *original-instance-type*. The instance type to be updated is *updated-instance-type*.)

Explanation: The specified DB2 instance cannot be updated. Downgrading the DB2 instance type is not supported.

User response: Install the DB2 product which supports the instance type and rerun the db2iupdt command.

DBI1235E Only one product can be installed at a time.

Explanation: The DB2 installer only supports the installation of one product at a time.

User response: Specify only one product and rerun the command.

DBI1236E The multi-partition instance *instance-name* cannot be upgraded or updated to a non multi-partition instance.

Explanation: The instance being upgraded or updated is a multi-partition instance. The current DB2 copy to which the instance is being upgraded to or updated does not support multi-partition instances.

User response: If the DB2 instance does not have any partitioned local databases and you want to use non multi-partition DB2 instance, update the current db2nodes.cfg to include only one partition and rerun the command.

If a partitioned local DB2 database with a DB2 instance exists, you should upgrade or update it to a DB2 copy

that supports multi-partition instances. DB2 Enterprise Server Edition supports multi-partition instances.

DBI1237E The specified component(s) *component-name* cannot be uninstalled since the DB2 database manager has not been stopped on all the instances related to the current DB2 installation.

Explanation: The components requested to be uninstalled are related to DB2 database manager. The DB2 database manager has to be stopped on all the instances related to the current DB2 copy before the specified components can be dropped.

User response: Stop the DB2 database manager on all the DB2 instances related to the current installation of DB2 and rerun the command.

DBI1238E The db2cluster_prepare command failed because the following invalid parameter was specified: *invalid-parameter*.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

This message is returned when the db2cluster_prepare command is called with an invalid parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying only valid parameters.

DBI1239I The DB2 Information Center is listening for requests on port *port-number*

DBI1240E The database manager could not be started.

Explanation: The database manager could not be started when attempting to check if cataloged local databases can be upgraded.

User response: Resolve why the database manager could not be started and try the command again.

DBI1241I The Information Center daemon is already active.

Explanation: The command to start the DB2 Information Center is already processed.

The command cannot be processed.

User response: The application is already listening.

DBI1242I The STOP command to stop the DB2 Information Center completed successfully.

DBI1243I The START command to start the DB2 Information Center completed successfully.

DBI1244I Directory for non-root installation of DB2 - *directory*

DBI1245E One or more space characters were found in the specified DB2 installation path *install-path*, which is not supported.

Explanation: The specified installation path of DB2 contains one or more space characters, which is not supported for DB2 installation on the Unix and Linux platforms.

User response: Specify an installation path without space characters and rerun the command.

DBI1246I The db2_deinstall command can only be run from a DB2 installation directory.

Explanation: Starting from DB2 9, the db2_deinstall command can only be run from a DB2 installation path. It can be found under DB2DIR/install, where DB2DIR represents the directory where the DB2 product is installed.

User response: Reenter the command from DB2DIR/install directory.

DBI1247W Root feature *feature-name* can not be enabled because the feature does not exist in the current installation.

DBI1248E You can only use this version of the db2_deinstall command to uninstall a DB2 version *version* product.

Explanation: To uninstall a DB2 product, you must use the db2_deinstall command that is at the same version level.

User response: Find the appropriate DB2 media and rerun the command or run the db2_deinstall command from the <DB2DIR>/install directory, where <DB2DIR> represents the installation path of the DB2 product you want to remove.

DBI1249E You can only use this version of the doce_deinstall command to uninstall the DB2 Information Center version *version*.

Explanation: To uninstall the DB2 Information Center, you must use the doce_deinstall command that is at the same version level.

User response: Find the appropriate DB2 Information Center media and rerun the command or run the doce_deinstall command from the <DB2DIR>/install directory, where <DB2DIR> represents the installation path of the DB2 Information Center you want to remove.

DBI1250E Applications are still using instance *instance-name*.

Explanation: There are applications that are still running that are using the specified instance. All applications using this instance must be terminated before the command can be completed successfully. You can get a list of the applications that are currently using the instance by issuing the command:

db2 list applications

User response: You can either wait for the applications to end by themselves, or you can explicitly force the applications to end. You can logon as the instance owner and run the command

db2 force application all

Note that some applications may have unexpected behavior when terminated using the indicated command. After all the applications have stopped, stop the database manager using "db2stop" command.

DBI1251N An error occurred while creating the cluster resource for this instance.

Explanation: The db2iupgrade command is unable to initialize the clustering service or is unable to find the existing DB2 resource to upgrade.

User response: Ensure that the current logon user account has sufficient authority to access the clustering service. Also ensure that the cluster service is active on the system, and that any resources that the DB2 resource you are upgrading depends on are online on the current node.

If the problem persists, contact your IBM service representative for assistance.

DBI1252N An error occurred while registering the DB2 server resource type in the clustering service.

Explanation: DB2 requires a resource type to be available in the clustering service but the command to register the resource type failed.

User response: Ensure that the current logon user account has sufficient authority to access the clustering service.

Register the resource type manually using the utility "db2wolfi.exe i".

If the problem persists, contact your IBM service representative for assistance.

DBI1253N **An error occurred while creating the cluster resource for this instance. The system cannot revert to its original settings.**

Explanation: An error occurred while creating the cluster resource required for this instance. The existing cluster resource for this instance was removed and could not be restored.

User response: Ensure that the current logon user account has sufficient authority to access the clustering service. Also ensure that that cluster service is active on the system, and that any resources that the DB2 resource you are upgrading depends on are online on the current node.

Recreate the resource for the instance using the "Cluster Administrator" available under the administrative tools control panel applet.

If the problem persists, contact your IBM service representative for assistance.

DBI1254N **An error occurred while upgrading the cluster resource for this instance because the existing resource cannot be removed.**

Explanation: In order to upgrade the instance, the cluster resource must be removed and a new cluster resource must be created using the DB2 server type. The command was not able to remove the existing cluster resource and therefore cannot upgrade this instance.

User response: Ensure that the current logon user account has sufficient authority to access the clustering service. Also ensure that the cluster service is active on the system, and that any resources that the DB2 resource you are upgrading depends on are online on the current node.

If the problem persists, contact your IBM service representative for assistance.

DBI1255E **The configuration file *filename* needed to create the specified instance type cannot be found.**

Explanation: The required configuration file to create the specified instance type can not be found.

User response: Specify the instance type supported by the current installation of DB2 and rerun the command.

DBI1256E **The db2cluster_prepare command failed because the -instance_shared_dev parameter was not specified.**

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

You must specify the absolute path to the shared device

on which the GPFS file system will be created using the required -instance_shared_dev parameter.

This message is returned when the db2cluster_prepare command is called without the -instance_shared_dev parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying a valid value for the -instance_shared_dev parameter.

DBI1257I **The DB2 Information Center is not listening for requests.**

DBI1258E **Unable to start DB2 Information Center.**

Explanation: The command to start the DB2 Information Center was unsuccessful.

User response: The messages are written in the log file doc/eclipse/workspace/.metadata/.log from the DB2 Information Center install path.

DBI1259I **Enabled root feature:**

DBI1260E **Cannot read from response file.**

Explanation: An attempt was made to read from the given response file failed. Ensure that the response file has read access permission and the given path to the response file is correct.

User response: Correct the access permission and the location of the response file and try again.

DBI1261E **The value is not valid for the keyword.**

Explanation: The value specified in the response file is not valid for its corresponding keyword. Refer to the sample response file for a list of valid entries.

User response: Correct the problem in the following line of the response file and try again.

DBI1262E **Unknown keyword detected in the response file.**

Explanation: The given keyword in the response file is not valid. Refer to the sample response file for a list of valid keywords.

User response: Correct the problem in the following line of the response file and try again.

DBI1263I **Not enabled root feature:**

DBI1264E This program failed. Errors encountered during execution were written to the installation log file. Program name: *program-name*. Log file name: *log-name*.

Explanation: This message is returned when some processes and operations have failed. Detailed information about the error was written to the log file.

User response: Contact IBM support to get assistance in resolving this issue. Keep the log file intact as this file is an important reference for IBM support.

DBI1265I Reserved service entry:

DBI1266I Refer to the log file *log-name* for more information.

Explanation: All processed and failed operations have been saved into this log file.

User response: Do not modify this file in any way. This file is for IBM Technical Support reference.

DBI1267I Not reserved service entry:

DBI1268E The file set *fileset* is not available on the installation media.

Explanation: One of the file sets to be installed is not found on the installation media. If the file set is not located on the installation media then it cannot be installed.

User response: Ensure that the file set is located on the installation media. If the file set is not located on the installation media de-select it and try again.

DBI1269I Summary:

DBI1270E Error detected in the response file.

Explanation: An error has been detected when processing the response file. Install cannot be continued unless the problem has been corrected.

User response: Correct the problem in the following line of the response file and try again.

DBI1271E The configuration file *configuration-file* is invalid.

DBI1272I To start using the DB2 instance *instance-name*, you must set up the DB2 instance environment by sourcing **db2profile** or **db2cshrc** in the **sqllib** directory, or you can open a new login window of the DB2 instance user.

Explanation: The DB2 instance cannot be used before

db2profile (for Bourne or Korn shell users) or **db2cshrc** (for C shell users) is sourced.

User response: To set up the DB2 instance environment, you can open a new login window under the ID that owns the DB2 instance, or source the DB2 instance environment by running the appropriate following command under the ID that owns the DB2 instance:

```
. $HOME/sqllib/db2profile
```

```
source $HOME/sqllib/db2cshrc
```

where \$HOME represents the home directory of the user ID that owns the DB2 instance.

DBI1273W The DB2 Text Search instance service could not be configured for the instance. Reason code: *reason-code*.

Explanation: An attempt to configure the instance for DB2 Text Search failed. Possible reasons are:

1

DB2 Text Search is not installed. Install DB2 Text Search from the installation media then try again.

2

The service name or port number specified is not valid. A valid service name must not exceed 14 characters in length, and a valid port number must be in the range 1024-65535.

3

The port number specified is being used by another application. Choose a port number that is available on the system.

4

The specified service name or port number conflict with existing values in the TCP/IP services file. Provide a service name and a port number that is available on the system.

5

The TCP/IP services file could not be accessed. Check that you have permission to read and write to the services file. Also, ensure that the content of the file is valid and no duplicate entries exist.

6

The default service name is reserved to a port number that is different from the one provided. Provide a service name and port number that do not conflict with existing entries in the TCP/IP services file.

7

The default service name is reserved to a port number that is being used by another application. Choose a service name and port number that is available on the system.

8

The service name provided is reserved to a port number that is being used by another application. Choose a service name and port number that is available on the system.

9

An attempt to create or copy DB2 Text Search directories or files in the instance directory failed.

10

An attempt to read or write to DB2 Text Search configuration files in the instance directory failed.

11

An attempt to create the DB2 Text Search instance service Windows service failed.

12

An unexpected internal error occurred.

User response: Contact your IBM service representative if the suggested action associated with the reason code did not solve the problem.

DBI1274N The DB2 Text Search instance service for the instance could not be stopped.

Explanation: An attempt to stop the DB2 Text Search instance service failed. The Windows service could not be stopped.

User response: Stop the DB2 Text Search server for the instance and try your command again.

DBI1275E The DB2 Text Search instance service for the instance could not be stopped.

Explanation: An attempt to stop the DB2 Text Search instance service failed.

User response: Stop the DB2 Text Search instance service for the instance and try your command again.

DBI1276E The DB2 installation at *install-path* must be updated while logged in under the user ID that performed the installation.

Explanation: You can only update the DB2 installation by running the installFixPack command under the user ID that performed the DB2 installation.

You have to run the installFixPack command as a user with root privilege to update a DB2 root installation.

User response: Rerun the command with an appropriate user ID.

DBI1277E Install path and runtime path do not match

Explanation: When setting up the runtime path with db2chgpath, a path was given that does not match the current path that DB2 is running from.

User response: Perform one of the following actions:

- Move the install path to match the runtime path.
- Rerun db2chgpath with the correct runtime path, to match the current install path.

DBI1278W Warning issued from the response file.

Explanation: A warning has been raised when processing the response file. DB2 Installer will carry on without stopping.

User response: If required, correct the problem in the following line of the response file and try again.

DBI1279I Notice issued from the response file.

Explanation: A notice has been issued when processing the response file. DB2 Installer will carry on without stopping.

User response: This is for your information only. No action is required.

DBI1280E The db2cluster_prepare command failed because one or more parameters were specified multiple times.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

This message is returned when the db2cluster_prepare command is called with a duplicate parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying each parameter only once.

DBI1281E The database manager configuration file could not be initialized.

Explanation: An error occurred when attempting to initialize the database manager configuration file. A DB2 instance could not be created or upgraded.

User response: Refer to the log file for troubleshooting. Resolve the problem and try the command again. If the problem persists, contact your IBM service representative.

DBI1282W The database manager configuration files could not be merged. The original configuration file was saved as *dbm-cfg-file-name*. (The original instance type is *original-instance-type*. The instance type to be upgraded or updated is *instance-type-upgraded*.)

Explanation: An error was detected when attempting to merge two database manager configuration files. The previous database manager configuration file could not be merged with the new database manager configuration file. This error could happen if the instance is being upgraded or updated from a higher instance type to a lower instance type.

If the instance update or upgrade was successful, a new instance is created using the default configuration settings for the instance-type-upgrade.

This error could happen if the instance is being upgraded or updated from a higher instance type to a lower instance type. The new instance uses the default configuration settings for the new instance type if update or upgrade was successful.

User response: Check the values in the new database manager configuration file and update the parameters as required using the *dbm-cfg-file-name* original configuration file.

DBI1283E Unable to update the communication information for *instance*.

Explanation: The SVCENAME parameter and/or the DB2COMM registry values could not be updated during instance creation.

User response: Make sure that the SVCENAME parameter is set to "<profile-name>" in the database manager configuration file and update using
db2 update dbm cfg using
SVCENAME "<profile-name>"

Also update the DB2COMM variable to tcpip by running
db2set DB2COMM=tcpip

DBI1284I *feature-name* has been set successfully.

DBI1285E Setting *feature-name* failed. See log file *log-file-name* for details.

DBI1286E The utility *utility-name* was not found.

Explanation: The utility must be available on your system to install or update DB2 products.

User response: Ensure that the utility is installed and its location is in the PATH environment variable.

Re-enter the command.

DBI1287E There has been an attempt to use db2_deinstall on an image for a platform that does not match the current platform on which it is being run.

Explanation: Possible causes include:

- This DB2 install image is not valid for the current platform.
- The current platform is not supported by DB2.

User response: Remove DB2 using the DB2 install image that corresponds with the current platform or directly run the command db2_deinstall from the directory 'DB2DIR/install', where DB2DIR is the DB2 installation path.

DBI1288E The execution of the program *program-name* failed. This program failed because you do not have write permission on the directory or file *directory-or-file-name*.

Explanation: Reasons for this error include:

- The given directory or file is on a local file system of the current machine, but the user id running the program does not have write permission for that directory or file.
- The given directory or file is on an NFS mounted file system mounted on the current machine, but the current user id does not have write permission for that directory or file.
- The current user id does not have write permission for the given directory or file located at the target partition. This can happen when the current machine's operating system is IBM AIX Version 6.1 or later and a Workload Partition root user id is attempting to write to a Global Partition.
- The current user id does not have write permission for the directory or file located at the target zone. This can happen when the current machine's operating system is Sun Solaris 10 or later and a local zone root user id is attempting to write to a global zone.
- The current machine is running the Sun Solaris Operating System in a local zone which does not have the write permission for the directory or file located on the global zone.

User response: There are several ways to respond to this message:

- Run the program with a user id that has write permission for the given directory or file.
- Modify the permissions of the directory or file to allow the user id to write to that directory or file, and rerun the program.
- Modify the authority of the user id to allow that user id to write to the given directory or file, and rerun the program.

DBI1289W **The DB2 install or instance utility *utility-name* failed to configure file permission and ownership properties for files related to DB2 Advanced Copy Services (ACS). See the log file *log-file-name* for more information.**

Explanation: The DB2 install or instance utility attempted to configure permission and ownership for DB2 ACS files that are located in the sqllib/acs directory under the DB2 instance home directory. The given log file contains information about why the DB2 install or instance utility failed.

DB2 ACS might not be usable because the DB2 install or instance utility failed to configure permission and ownership for these files.

User response: If you will not be using DB2 ACS, you do not need to reinstall IBM Data Server, or recreate the DB2 instance.

If you will be using DB2 ACS, perform the following steps:

1. See the given log file for more information about the failures.
2. Correct the cause of the failures and take the action recommended in the log file.

For more information about configuring DB2 ACS, see the topic called "Configuring DB2 Advanced Copy Services" in the DB2 Information Center.

DBI1290E **An attempt to create the profile registry *profile-name* failed.**

Explanation: An error occurred when creating the profile registry. One of the following situations occurred:

- Incorrect access permission on the top directory of the current DB2 installation.
- There is not enough space on your file system.

User response: Check the directory permissions for the top directory of the current DB2 installation.

DBI1291E **The instance *instance-name* was not found in the instance list.**

Explanation: The specified instance was not found in the instance list.

User response: Verify that the list of instances as reported by the "db2ilist" command is correct. Retry the command with a valid instance name.

DBI1292E **The instance *instance-name* is already in the instance list.**

Explanation: The instance that is being created or upgraded is already in the instance list.

User response: Use a different instance name that is not part of the instance list as reported by the db2ilist command.

DBI1293E **There is no DB2 Query Patroller Server set up on this system.**

User response: If DB2 Query Patroller was not installed via the DB2 Installer, please use the Installer to properly set up an instance for use with the DB2 Query Patroller Server. If one has been set up, then verify that it is recorded in the default.env file, and that the file is world readable. default.env can be found in /usr/opt/db2_08_xx, where xx is 01 or FPn, where n is the FixPak number or /opt/IBM/db2/V8.x, where x is 1 or FPn, where n is the FixPak number

DBI1294W **The DB2 install or instance utility *utility-name* failed to start DB2 Advanced Copy Services (ACS.) See the log file *log-file-name* for more information.**

Explanation: Before you can use DB2 ACS, the services must be started. The DB2 install or instance utility attempted to start the services, but failed.

You can use the DB2 instance, but you cannot use DB2 ACS until the services are started.

User response: To start DB2 ACS, perform the following steps:

1. See the given log file for more information about the failures.
2. Correct the cause of the failures.
3. Start DB2 ACS manually.

For more information about manually starting DB2 ACS, see the topic called "Starting DB2 Advanced Copy Services" in the DB2 Information Center.

DBI1295E **The instance list could not be updated.**

Explanation: The instance could not be added to or removed from the instance list. An error has occurred when adding or removing the instance from the instance list. One of the following situations occurred:

- Incorrect access permission on the registry profile.
- Profile registry is not set up properly.
- There is not enough space on your file system.

User response: Check the file permissions on the instance list profile under the top directory of the current DB2 installation.

DBI1296E The DB2 install or instance utility *utility-name* failed to stop DB2 Advanced Copy Services (ACS.) See the log file *log-file-name* for more information.

Explanation: Before you can update or upgrade a DB2 instance, DB2 ACS must be stopped. The DB2 install or instance utility attempted to stop the services, but failed.

You cannot update or upgrade the DB2 instance until DB2 ACS is stopped.

User response: To stop DB2 ACS, perform the following steps:

1. See the given log file for more information about the failures.
2. Correct the cause of the failures.
3. Stop DB2 ACS manually.

For more information about manually stopping DB2 Advanced Copy Services (ACS), see the DB2 Information Center.

| | |
|----------|--|
| DBI1297E | The instance profile <i>profile-name</i> could not be updated. |
|----------|--|

Explanation: An error occurred when updating the instance profile registry. One of the following situations occurred:

- Incorrect access permissions on the instance profile registry.
- Profile registry is not set up properly.
- There is not enough space on your file system.

User response: Check the file permissions on the instance list profile under the top directory of the current DB2 installation.

DBI1298E The DB2 install or instance utility *utility-name* failed to disable DB2 Advanced Copy Services (ACS.) See the log file *log-file-name* for more information.

Explanation: Before you can drop a DB2 instance, or remove a non-root install of IBM Data Server, DB2 ACS must be stopped. The DB2 install or instance utility attempted to disable the services, but failed.

You cannot drop a DB2 instance, or remove a non-root install of IBM Data Server until DB2 ACS is disabled.

User response: To disable DB2 ACS, perform the following steps:

1. See the given log file for more information about the failures.
2. Correct the cause of the failures.
3. Disable DB2 ACS manually.

For more information about manually disabling DB2 ACS, see the topic called "Disabling DB2 Advanced Copy Services" in the DB2 Information Center.

| | |
|----------|---|
| DBI1299E | DB2 profile registry variable <i>variable-name</i> was not nulled because it is a member of an aggregate set. |
|----------|---|

Explanation: Several DB2 profile registry variables can be grouped together and set as one aggregate DB2 profile registry variable. When you configure an aggregate DB2 profile registry variable, the DB2 profile registry variables that are members of that set are automatically configured to their predefined values.

DB2 profile registry variables that are members of an aggregate set, can not be nulled using the following db2set command syntax:

db2set *variable-name*=

The variable was not nulled.

User response: To null a DB2 profile registry variable that is a member of an aggregate set, use the following db2set command syntax to set the variable to NULL:

db2set -null *variable-name*

```
DBI1300N      db2set displays, sets, or removes DB2
                profile variables.
db2set [variable=[value]]
        [-g|-i instance [member-
number]]
        [-all]
        [-null]
        [-r [instance] [member-number]]
        [-im|-immediate]
        [-info]
        [-n DAS node[-u user[-
p password]]]
        [-l|-lr]
        [-v]
        [-ul|-ur]
        [-?|-h]
```

Explanation: The command options are:

-99-

Access the global profile variables.

-i

Specifies the instance profile to use instead of the current or default.

-n

Specifies the remote DB2 Administration Server node name.

-u

Specifies the user ID to use for the Administration Server attachment.

| | |
|-------------------------|--|
| -ul | Access the user profile variables |
| -ur | Refreshes user profile variables |
| -p | Specifies the password to use for the admin server attachment. |
| -r | Resets the profile registry for the given instance. The default/current instance will be used if none is provided. |
| -im -immediate | Specifies that the update takes effect the next time an SQL statement is compiled for registry variables that support this feature. |
| -info | Returns the properties of the specified variable. The properties state whether an immediate change is supported by the variable and whether the change is immediate by default |
| -l | Lists all instance profiles. |
| -lr | Lists all supported registry variables. |
| -v | Verbose mode. |
| -? | Displays the command help message. |
| -h | Same as -? option. |
| -all | Displays all occurrences of the local environment variables as defined in: <ul style="list-style-type: none"> • The environment, denoted by [e] • The user level registry, denoted by [u] • The node level registry, denoted by [n] • The instance level registry, denoted by [i] and • The global level registry, denoted by [g] |
| -null | Sets the variables value to null at the specified registry level to prevent looking up the value in the next registry level as defined in the variable value search order. |

Notes:

- db2set with no variable name displays all defined variables.
- db2set <variable> displays <variable>'s value.
- db2set <variable>= (nothing) deletes the <variable>.
- db2set <variable>=<value> modifies the <variable>'s value.
- db2set <variable> -null sets <variable>'s value to NULL.
- db2set <variable> -all displays all defined <variable>'s values.
- db2set -ur refreshes the current user profile.
- db2set <variable> -ul displays defined <variables> at the user level.
- db2set -all displays all defined variables in all registry levels.

User response:

DBI1301E The registry variable was not updated because an invalid value was specified.

Explanation: You can configure many aspects of DB2 database functionality by setting DB2 profile registry variables. You can set DB2 profile registry variables by using the db2set command.

This message is returned when an attempt is made to use the db2set command to set a DB2 profile registry variable to a value that is invalid.

User response:

1. Determine the range of valid values for the registry variable.
2. Call the db2set command again, specifying a valid value for the specified registry variable.

DBI1302E Invalid parameter detected.

Explanation: An invalid parameter was used.

User response: Use the -? option for the usage help message.

DBI1303W Variable not set.

Explanation: The variable was not set in the profile registry.

User response: No further action is required.

DBI1304E Unexpected error.

Explanation: The tool encountered an unexpected system error.

User response: Contact your DB2 service representative.

DBI1305E The profile registry was not found.

Explanation: The target machine does not have a profile registry setup.

User response: Create the registry on the target machine by installing DB2.

DBI1306N The instance profile is not defined.

Explanation: The instance is not defined in the target machine registry.

User response: Specify an existing instance name or create the required instance.

DBI1307N The instance node profile is not defined.

Explanation: The instance node is not defined in the target machine registry.

User response: Create the registry by installing the required DB2 product parts.

DBI1308E Out of memory condition encountered.

Explanation: The tool encountered an “out of memory resource” error.

User response: System is low on memory resources. Terminate non-essential applications or try again later.

DBI1309E System error.

Explanation: The tool encountered an operating system error.

User response: A system error was encountered during registry access. Ensure that there is enough space on the file system where the registry is located, and that there is a valid LAN connection if the registry is remote.

DBI1310E Remote registry access is not supported.

Explanation: The tool does not support the remote registry option.

User response: Consult the Command Reference on how to access the registry remotely.

DBI1311I Listing registry instance profiles...

DBI1312I Listing defined global variables...

DBI1313I Listing defined instance variables...

DBI1314I Listing all supported registry variables...

DBI1315E The specified variable is not an aggregate type registry variable.

Explanation: To display the group definition, the registry variable must be an aggregate type registry variable.

User response: Do not specify the -gd option or specify a valid aggregate variable name when using the -gd option.

DBI1316E Instance is not MPP.

Explanation: The target instance is not a DB2 MPP instance.

User response: Choose an MPP instance.

DBI1317E Instance node already exists.

Explanation: A duplicate instance node already exists.

User response: Choose a different instance node.

DBI1318E The value specified for the registry variable is too long.

Explanation: The value specified for the registry variable exceeds the maximum limit. The maximum length for a registry variable value is 255 bytes.

User response: Specify a shorter value for the registry variable.

DBI1319W The variable *variable-name* has been explicitly set and will not be affected by the configuration of the aggregate variable *aggregate-var-name*.

Explanation: An aggregate registry variable has been configured to a value whose group definition contains settings for a variable which has been explicitly configured. The explicitly configured value will be retained.

User response: If desired, explicitly configure the variable.

DBI1320W WARNING:

Explanation: This script is used to gather information useful to IBM Support to debug the problems you may be experiencing. This information may be of a sensitive nature. You may wish to edit the output file, db2ginfo.txt, before sending it to IBM Support.

User response: To signal your awareness and acceptance of the warning, run this script with the -y flag specified.

DBI1321W The ulimit setting for the DB2 instance *instance-name* does not comply with the recommended values for DB2 on the current platform.

Explanation: The ulimit setting for the current non-root DB2 instance has values for 'data' and 'nofiles' smaller than the values recommended by DB2 on the current platform. Set the 'data' to 'unlimited' and 'nofiles' to '65536' or the maximum value allowed on the system. If setting 'data' to 'unlimited' cannot be done, the following formula can be used as reference to decide the estimate value in kilobytes of 'data' required by DB2 for the current platform:

```
default_data_ulimit + ( ( 20 *
max_concurrent_active_databases ) +
estimated_max_connections ) * max_querydegree * 8192
```

where:

1. default_data_ulimit is the default value of 'data' of ulimit for a user on the system. For example, the default value is 262144 on IBM AIX system.
2. max_concurrent_active_databases is the maximum number of concurrent active databases for the DB2 instance.
3. estimated_max_connections is the estimated maximum number of database connections for the DB2 instance.
4. max_querydegree is the value of the parameter "MAX_QUERYDEGREE" for DB2 database manager configuration when the INTRA_PARALLEL is enabled in the DB2 database manager configuration. If INTRA_PARALLEL is not enabled, the max_querydegree's value is 1.

User response: Ask the system administrator with root privilege to update the ulimit setting for the current DB2 instance.

DBI1322I Usage:
db2iprune -r input_file_path
-o destination_dir_path | -c
[-
t trace_file]
[-l log_filename]
[-h | -?]

Explanation: An incorrect argument was entered for the db2iprune command. Valid arguments for this command are:

-r input_file_path

Specifies the full path to the input file that is to be used. The input file, or .prn file, contains a full list of removable components and is used to indicate which products, components, and languages you would like removed from the installation image.

-o destination_directory_path

Specifies the full path to where the new DB2 pruned image is copied. Make sure that you have write access to this directory.

-c

Specifies that you want to prune the source installation image directly. Ensure that the source installation image directory is writable.

-t trace_file

(On Linux and UNIX operating systems only.) Turns on the debug mode. The debug information is written to the file name specified.

-l log_filename

Enables error logging. On Linux and UNIX operating systems, if the -l option is not specified, the default log file name is tmpdir/db2iprune_username.log. On Windows operating systems, the log file db2iprune.log is written to the destination directory.

-h | -?

Displays the usage information.

User response: Enter the command again using the valid parameters. For details, see the DB2 Information Center for examples and usage notes for the db2iprune command.

DBI1323I Usage:
db2rspgn -d output_directory
[-i instance]
[-t trace-file]
[-h | -?]

Explanation: An incorrect argument was entered for the db2rspgn command. Valid arguments for this command are:

-d

Specifies the full path to the output directory for generated files. If the output directory specified is an existing directory, the directory must be empty and writable. If the output directory specified does not exist, the new directory is created if the location is writable. This parameter is mandatory.

-i instance

Generates the specified instance configuration and saves this information in the generated response file and instance configuration profile. This parameter is optional. By default, all instances are selected. To specify multiple instances, specify this parameter multiple times. For example, -i db2inst1 -i db2inst3.

-t trace-file

Linux and UNIX operating systems only. Turns on the debug mode. The debug information is written to the file name specified as trace-file.

-h|-?

Displays help information.

User response: Enter the command again using the valid parameters.

DBI1324W **Support of the *command-name* command is deprecated. For more information, see the DB2 Information Center.**

DBI1325E **The *command-name* command has failed.**

Explanation: At least one instance is still running the DB2 Text Search instance service.

User response: Stop the DB2 text search instance service and re-run the command.

DBI1326W **The Text Search instance services failed to restart.**

Explanation: Before running the DB2 instance command, the Text Search instance service was not stopped. The DB2 process stopped the Text Search instance service. However, the DB2 process is unable to restart the Text Search instance service.

User response: Manually restart the Text Search instance service.

DBI1327E **The db2icrt command failed because the home directory of the instance owner is a subdirectory under the directory in which the DB2 product is installed.**

Explanation: The home directory of the instance owner cannot be a subdirectory under the directory in which the DB2 product is installed.

For example, if the DB2 product is installed in the directory named /opt/db2v97, the home directory of the user who will be the instance owner cannot be in any directory under /opt/db2v97.

This message is returned in the following scenario:

- On Linux and UNIX operating systems: the home directory of the instance owner is a subdirectory under the directory in which the DB2 product is installed.

The DB2 instance was not created.

User response:

1. Configure that the user environment and directory structure so that the home directory of the instance owner is not a subdirectory of the directory in which the DB2 product is installed.
2. Run the db2icrt command again.

DBI1328E **The db2cluster_prepare command failed because no value was specified for the -instance_shared_dev parameter.**

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

You must specify the absolute path to the shared device on which the GPFS file system will be created using the required -instance_shared_dev parameter. The db2cluster_prepare utility will create the GPFS file system in the specified location.

This message is returned when the db2cluster_prepare command is called without a value specified for the -instance_shared_dev parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying a value for the -instance_shared_dev parameter.

DBI1329I **Usage:**
db2val [-h|-?] [-o] [-i inst_name1] | [-a] [-b db_name] [-t trace_file] [-d] [-s] [-l log_file]

Explanation: An invalid argument was entered for the db2val command. Valid arguments for this command are:

-o

Specifies that only the installation files will be validated; validation of the instance, database, and extended security will not be performed. If this parameter is specified, the -i, -a, -b, and -s parameters are ignored.

-i inst_name

Specifies that the name of the instance to validate. To specify validation of multiple instances, specify this parameter multiple times. For example, -i inst1 -i inst2.

- On Windows operating systems, if this parameter is not specified, the current instance will be used as the default value.
- On Linux and UNIX operating systems, this parameter can only be used by root users in a root installation of a DB2 copy.

-a

Validates all instances in the DB2 copy. On Linux and UNIX operating systems, this

parameter can only be used by root users in a root installation of a DB2 copy. This option overrides the `-i` option.

-b db_name

Validates the local database creation and connection with the given database name. This parameter is ignored in the following scenarios:

- the instance to be validated is a client instance
- `-a` and `-i` are not specified when `db2val` is run by root users (on Linux and UNIX only).

-t trace_file

Specifies the full path and name of trace file specified by `trace_file`.

-d

This parameter is deprecated and might be removed in a future release. Use the `-t` parameter instead. Valid only on Linux and UNIX operating systems. Use this parameter only when instructed by DB2 Support. Turns the debug mode on.

-s

Starts the DB2 database manager for the specified instance that is part of a partitioned database environment. This parameter overrides parameter `-i`.

-l log_file

Writes the log to the file name specified. Unless the `-l` parameter is specified, the default log path is:

- On Linux and Unix, `/tmp/db2valxx.log`
- On Windows, `My Documents\DB2LOG\db2valxx.log`

where `xx` is a generated value.

-h | -?

Displays the usage information.

User response: Enter the command again using the valid parameters.

DBI1330W The action might affect other instances.

Explanation: The file `IWH.environment` contains global information for all instances. Updating this file might affect all instances. Do you wish to continue?

DBI1331E Installation file validation for the DB2 copy *copy_name* failed.

Explanation: Some features or components have missing files or are the wrong size.

User response: To repair the DB2 installation, from the Add or Remove Programs tool, select Change for the DB2 copy, and select the Repair option. Rerun the `db2val` command.

DBI1332E Missing template file for IWH.environment.

Explanation: The template file for `IWH.environment` is needed.

DBI1333I Installation file validation for the DB2 copy *copy_name* was successful.

DBI1334E Installation file validation failed for the DB2 copy installed at *installation-location*. Reason code = *reason-code*.

Explanation: You can verify the basic functions of a DB2 copy, checking the state of installation files, instance setup, and local database connections, by using the `db2val` utility.

This message is returned when the `db2val` utility detected problems with the DB2 copy that is installed at the given location.

The reason code indicates more specifically what the problem is:

1

Some features or components have missing files or files with the wrong size.

2

The installation image inside the install path has missing files.

3

The embedded runtime path for some DB2 library or executable files is not set properly.

4

The installation path is not accessible to the database manager.

5

The `/etc/services` file is not accessible to the database manager.

User response: Respond to this message according to the reason code:

1

Respond to reason code 1 in one of the following ways:

- If your DB2 copy has the DB2 pureScale Feature installed:
 1. Reinstall the DB2 copy to a new path by running the installFixPack command from the installation media, specifying the -p <new-path> option and the -f nobackup option.
 2. Update the instance to the new copy using the db2iupdt command.
 3. Run the db2val command.
- If your DB2 copy does not have DB2 pureScale installed:
 1. Reinstall the DB2 copy to the same path by running the installFixPack command, specifying the "-f nobackup" option.
 2. Run the db2val command.

2

Copy the installation image into <DB2-installation-path>/sd directory.

3

Set the embedded runtime path by using one of the following methods:

- For root installations, run <DB2-installation-path>/install/db2chgpath as root.
- For non-root installations, run <DB2-installation-path>/install/db2chgpath as the copy owner.

Then run the db2val command.

4

Modify the access permissions for the installation path so that the DB2 database manager can access the installation path.

5

Modify the access permissions for the /etc/services file so that the DB2 database manager can access that file.

3. The start and end ports for Fast Communication Manager (FCM) are not reserved properly on all database partition servers.
4. The instance profile directory is not accessible from all database partition servers.
5. The instance service is running under a non-domain account. This will prevent the DB2 instance from starting.
6. The installation directory for the instance is not accessible from all database partition servers.
7. The code level on each database partition server is not consistent.
8. On Windows operating systems, if the db2val command is run to validate a multiple partition instance, domain user account authority is required.
9. An unexpected internal error occurred.

User response:

1. Ensure the format of db2nodes.cfg follows the DB2 standard.
2. Check the .rhosts file is configured on all the partitions. If the DB2RSHCMD registry variable is set to ssh, ensure the ssh remote shell utility is available.
3. Ensure the start and end ports for FCM are free and reserved on each computer participating in the partitioned database environment. The start port should be the same on each participating computer.
4. Ensure the instance profile directory on the instance-owning machine is shared among all the database partition servers. The instance profile directory should be in the UNC format.
5. Use a domain account for the instance service on all database partition servers.
6. Ensure the installation directory is accessible from all database partition servers.
7. Ensure the same code level is installed on all computers participating in the partitioned database environment.
8. Log on as a user with the domain user account authority and retry the command.
9. Rerun the db2val command. If this error persists, contact your IBM service representative.

DBI1335I **Installation file validation for the DB2 copy installed at *installation_location* was successful.**

DBI1336E **The partitioned database environment validation for the instance *instance_name* failed. Reason code = *reason_code*.**

Explanation:

1. The format of the DB2 node configuration file (db2nodes.cfg) is incorrect.
2. The other database partition servers are not accessible.

DBI1337I **The partitioned database environment validation for the instance *instance_name* was successful.**

DBI1338E **The validation of instance *instance_name* failed. Reason code=*reason_code*.**

Explanation:

1. The default instance is not set.
2. Some files in the instance directory are damaged. The symbolic links of the files are not pointing to

the current DB2 copy installation path, and the file permission and ownership might have been changed.

3. The DB2 Extended Security setup does not allow the current user to start an instance. If Extended Security is enabled, the user needs to be in the DB2ADMNS group.
4. To start the instance, read permission is required to the `/etc/services` file.
5. The appropriate authority is required to validate this instance.
6. The partitioned database environment setting is not correct.
7. The instance failed to start because of system errors.
8. An unexpected internal error occurred.

User response:

1. Set the DB2INSTDEF profile registry variable to an instance in the DB2 copy.
2. Run the `db2iupdt` or `db2nrupdt` command with the `-k` option for the instance.
3. Add the current user to the DB2ADMNS group and rerun the command.
4. Ensure you have read permission to the `/etc/services` file.
5. For authorization details, see the `db2val` command in the DB2 Information Center. Then log on as a user with the appropriate authority, and retry the command.
6. Check the partitioned database environment validation log file for details.
7. Check the log file for details.
8. Rerun the `db2val` command. If this error persists, contact your IBM service representative.

DBI1339I **The instance validation for the instance *instance_name* was successful.**

DBI1340I **Database validation for instance *instance_name* was successful.**

DBI1341E ***ID* does not have the authority to perform the requested command.**

Explanation: The appropriate authority is required to run the command. For authorization details, see the `db2val` command in the DB2 Information Center.

User response: Log on as a user with the appropriate authority, and retry the command.

DBI1342E **Instance *instance_name* does not exist in the current DB2 copy.**

Explanation: The specified instance does not exist in the current DB2 copy and will not be validated.

User response: Rerun the `db2val` command with an instance that is listed by the `db2ilist` command for the current DB2 copy.

DBI1343I **The `db2val` command completed successfully. For details, see the log file *log_path*.**

DBI1344E **The validation tasks of the `db2val` command failed. For details, see the log file *log_path*.**

DBI1345W **Read permission to the `/etc/services` file is not granted to every user. This lack of permission can cause the instance creation to fail.**

Explanation: Creating an instance requires the users who manage the instance to have read permission to the `/etc/services` file.

User response: Ensure users have read permission to the `/etc/services` file.

DBI1346W **Path *path* defined in the DFTDBPATH database manager configuration variable is not accessible on the current database partition server. This can cause the database creation to fail.**

DBI1347W **The `db2val` command completed with warnings. For details, see the log file *log_path*.**

DBI1348W **Read and execute permissions to the installation directory are not granted to everyone. This lack of permission might cause problems when using your DB2 product.**

DBI1349W **Installation file validation completed with warnings for the DB2 copy installed at *installation_location*. For details, see the log file *log_path*.**

DBI1350E **The database validation failed for the instance *instance_name*. Reason *code=reason_code*.**

Explanation: Database validation failed because of one of the following reasons:

1. Database creation failed because of system errors.
2. Database connection failed because of system errors.
3. An unexpected internal error occurred.

User response: Check the log file for details, or rerun the `db2val` command. If this error persists, contact your IBM service representative.

DBI1351E **You must be the instance owner to run this command.**

Explanation: To run this command, you must be the DB2 Query Patroller Server instance owner.

User response: Verify that the DB2 Query Patroller Server instance is correctly recorded in the file default.env. Log in as the DB2 Query Patroller Server instance owner and issue the command again. default.env can be found under the top directory of the current DB2 installation.

DBI1352E **The instance *instance-name* is not an ESE instance.**

Explanation: DB2 Query Patroller Server/Agent must be created on an ESE instance.

User response: Run this command again using a valid ESE instance name, or you may create and set up a new ESE instance first before running this command again.

DBI1353E **DB2 Query Patroller Server has been set up on an instance already.**

Explanation: DB2 Query Patroller Server can be set up on one DB2 instance only.

User response: Run the command dqplist to find out the named of the instance being used as the DB2 Query Patroller Server. If the instance name is correct then there is no need to run dqpcrt to set up a server. For setting up an agent you need to use the correct instance name. If this instance is no longer being used for that purpose, run the command

dqpdrop inst_name

to remove it. Then re-run the dqpcrt command again.

DBI1354E **The instance *instance-name* is not an DB2 Query Patroller Server instance.**

User response: Use dqplist to find out the proper server instance name and re-run the command again.

DBI1355I **Usage:**
dqpcrt [-h|-?] -s|-a
-p PortName InstName

Explanation: An incorrect argument was entered. Valid arguments for this command are:

PortName

Port name to be used with the DB2 Query Patroller Server/Agent

InstName

name of an instance that is to be designated as an DB2 Query Patroller Server instance

-s Creating a DB2 Query Patroller Server on the named instance

-a Creating a DB2 Query Patroller Agent on the named instance.

-h|-? display the usage information

The -a option is only valid on a multiple partitioned database instance.

User response: Re-enter the command with the proper syntax.

DBI1356I **Usage:**
dqplist [-h|-?]

Explanation: An incorrect argument was entered. Valid arguments for this command are:

-h|-? display the usage information

User response: Re-enter the command with the proper syntax.

DBI1357I **Usage:**
dqpdrop [-h|-?] InstName

Explanation: An incorrect argument was entered. Valid arguments for this command are:

InstName

name of the instance that you wish to remove the DB2 Query Patroller Server from

-h|-? display the usage information

This command can only be issued on the node where the DB2 Query Patroller Server was created.

User response: Re-enter the command with the proper syntax.

DBI1358W **Failed to modify the DB2 Query Patroller profile files.**

Explanation: An attempt to update the dqpprofile and/or dqpcshrc files have failed. Possible causes can be

- These files do not exist in the sqllib directory under the instance home directory.
- There is no write permission on these files.
- Failure to create a temporary file in /tmp directory.

User response: Check the existence and permission on these files. Verify that you can write to /tmp. Issue the command again.

DBI1359E **Unable to remove the DB2 Query Patroller Server from instance *instance-name*.**

Explanation: An attempt to remove the DB2 Query Patroller Server from the named instance has failed.

User response: Manually remove the file sqllib/cfg/dqplevel from the instance owner's home

directory. Also edit the file `default.env` and remove the `DQPSEVER` line from this file. `default.env` can be found under the top directory of the current DB2 installation.

DBI1360E Unable to complete DB2 Query Patroller set up.

Explanation: An attempt to set up the DB2 Query Patroller Server has failed. Possible cause include:

- Failed to create a database for the DB2 Query Patroller Server
- Failed to create a database partition group on the specified node
- Failed to create a table space in the nodegroup

User response: If the database or the database partition group creation failed, try to create them manually and re-run the command.

If it failed at creating the table space, check to see that you have write permission to the path you have designated. Try to create the table space manually.

DBI1361E DB2 Query Patroller Server code has not been installed.

Explanation: The DB2 Query Patroller Server code has not been installed, and is required for the current operation to complete.

User response: Install the DB2 Query Patroller Server code, and try the command again.

DBI1362W IWM schema objects have been detected.

Explanation: DB2 has detected IWM schema objects, possibly from a previous install of IWM.

User response: If you wish to keep the existing IWM schema objects, there is nothing to do. If you wish to remove them and re-create new ones, re-issue the `dqpsetup` command with the `-o` flag.

DBI1363E Failed to create IWM schema objects.

Explanation: An attempt to create the IWM schema objects has failed.

User response: Verify that the files `db2_qp_schema` and `iwm_schema.sql` are both in the directory `/usr/opt/db2_08_xx/bin`, where `xx` is 01 or FPn, where `n` is the FixPak number.

If both files exist, then the output of the IWM schema creation procedure can be found in `/tmp/iwmschcr.log`. Correct any error and re-run the `dqpsetup` command again.

DBI1364E No previous DB2 Query Patroller schema objects exist.

Explanation: No previous DB2 Query Patroller schema objects exist to be migrated.

User response: Run `dqpsetup` again with the correct syntax. Use `'dqpsetup -h'` to see the syntax.

DBI1365E Node number is not defined.

Explanation: The node number is not defined in the `db2nodes.cfg`.

User response: Examine the `db2nodes.cfg` file and make sure the entry you want is there. Enter the command again using the correct node number exactly as recorded in the `db2nodes.cfg` file.

DBI1366E The container *sms-path* is already in use.

Explanation: The directory named as a container for the table space already exists.

User response: Run the command again using a different path for the container for the table space.

DBI1367E Failed to bind the files listed in *bind-list*.

Explanation: There was a problem binding the bind files listed in named file(s).

User response: Perform the binding manually.

DBI1368E Failed to connect to DQP database *db2dbdft*.

Explanation: The database likely does not exist.

User response: Specify that you would like to setup a new DQP instance using all the required parameters to `dqpsetup`. This instance cannot be upgraded.

DBI1369W The partitioned database environment validation for the instance *instance_name* completed with warnings.

DBI1370W Unable to modify the *.profile* or the *.login* file of user *inst-name*.

Explanation: DB2 has failed to modify the `.profile` file or the `.login` file of this user. These files may not exist, or you may not have write permission to them. Without the modification you need to set your environment manually each time when you log in as this user to use DB2 Query Patroller.

User response: Set up the application development environment.

DBI1371E Missing arguments for setting up a Query Patroller Server.

Explanation: A Query Patroller instance must be set up as server when DB2 Query Patroller Server is installed. Therefore the following arguments must be supplied:

1. Node group.
2. Node number.
3. Database name.
4. Table space.
5. Result table space.
6. Table space path.

And optionally:
DMS Table space size.

User response: Specify the required and/or optional arguments in the response file.

DBI1372E Cannot set up a Query Patroller server instance.

Explanation: To set up a Query Patroller server instance, you must have DB2 Query Patroller Server installed.

User response:

- Install DB2 Query Patroller Server.
- If you wish to set up a Query Patroller agent, take out the arguments for node group, node number, database name, table space, result table space, table space path and DMS table space size.

DBI1373E Cannot set up Query Patroller server instance when creating a new ESE instance.

User response: Create the ESE instance and set up the instance as a Query Patroller server in separate steps.

DBI1374W The validation of instance *instance_name* completed with warnings.**DBI1375E Cannot set up a Query Patroller instance without ESE installed.**

User response: Install ESE, and DB2 Query Patroller Agent/Server.

DBI1376E To set up a Query Patroller instance, you need to provide the service name and port number of the IWM user.

User response:

- If you wish to set up a Query Patroller instance, enter the service name and port number of the IWM user.

- If you do not wish to set up a Query Patroller instance, take out the specification of node group, node number, database name, table space, result table space, table space path and DMS table space size from the response file.

DBI1377N db2isetup was unable to locate a suitable Java Runtime Environment on your computer. If a Java Runtime Environment *jre-version* is present, set your JAVA_HOME environment variable and run the command again. Otherwise, refer to the DB2 Installation Requirements for information on the recommended Java environment for your operating system. If you need to use your own JRE, set DB2USELOCALJRE=true.

DBI1378N db2setup was unable to locate a suitable Java Runtime Environment on your computer. If a Java Runtime Environment *jre-version* is present, set your JAVA_HOME environment variable and run the command again. Otherwise, refer to the DB2 Installation Notes for information on the recommended Java environment for your operating system. If you are unable to obtain a suitable Java Runtime Environment you can install this product from the command using the doc_install script. Run this script without any parameters for information on how to use it. If you need to use your own JRE, set DB2USELOCALJRE=true.

DBI1379I The db2val command is running. This can take several minutes.

DBI1380W The database *db_name* cannot be removed.

Explanation: The database might be in use by other applications.

User response: To manually drop the database, run the following command: db2 drop db *db_name*

DBI1381W Database validation for instance *instance_name* completed with warnings.

DBI1382W The database manager of the instance *instance_name* cannot be stopped.

Explanation: The database might still be in use.

User response: To stop the database manager, run the db2stop command.

DBI1383I You must update the DB2 registry for the current instance manually using the db2iset command, including the full path of the db2iset command. For example: /opt/ibm/db2/<release>/instance/db2iset -d *instance-name*.

User response:

DBI1384E The instance *instance-name* cannot be created.

Explanation: DB2 Text Search configuration can only be performed on server instances.

User response: Configure the DB2 Text Search feature on a server instance.

DBI1385W Partitioned instances are not supported with the DB2 Text Search feature.

Explanation: The DB2 Text Search feature cannot be configured for this instance.

User response: Configure a single-partition instance to use the DB2 Text Search feature

DBI1386N The account *user-name* is locked.

Explanation: The user's account is locked by the operating system.

User response: Contact the system administrator to unlock this user's account.

DBI1387I An invalid argument was entered for the db2isetup command.

Explanation: The syntax for the db2isetup command is:

```
db2isetup [-h|-?]
          [-i <language_code>]
          [-r <response_file>]
          [-t <trace_file>]
          [-l <log_file>]
```

Valid arguments for this command are:

-i language-code

Specifies two letter code for the preferred language in which to run the install. If unspecified, this parameter will default to the locale of the current user. For a list of language identifiers, see the DB2 Information Center.

-r response_file

Specifies the full path and file name of the response file to use .

-t tracefile

Specifies the full path and name of trace file specified by tracefile.

-l logfile

Writes to the log to the file name specified. For root installations, the path and filename defaults to /tmp/db2isetup.log. For non-root installations, the default log file is /tmp/db2isetup_userID.log, where userID represents the user ID that owns the non-root installation.

-h|-?

Displays the usage information.

User response: Enter the command again using valid parameters.

DBI1388I An invalid argument was entered for the db2nrupgrade command.

Explanation: The syntax for the db2nrupgrade command is:

```
db2nrupgrade [-d]
              [-a auth_type]
              [-b backup_dir]
              [-j "TEXT_SEARCH" |
               -j "TEXT_SEARCH,portnumber"]
              [-h|-?]
```

For details about the parameters of the db2nrupgrade command, refer to the DB2 documentation for Linux, UNIX, and Windows.

User response: Enter the command again using valid parameters.

DBI1389I Usage:
db2ckupgrade (<database_name> | -
e)
-l <logfile>
-u <userid> -
p <password>
[-h|-?]

Explanation:

database_name

Specifies the name of a database to be scanned.

-e

Specifies that all local cataloged databases are to be scanned.

-l logfile

Mandatory parameter. Specifies a log file to keep a list of errors and warnings generated for the scanned database.

-u userid

Specifies the user ID of the instance owner.

-p password

Specifies the password of the instance owner.

-h | ?

Displays the usage information.

User response: Reenter the command with the proper syntax.

DBI1390E The current instance cannot be upgraded.

Explanation: An attempt was made to upgrade the instance. This instance cannot be upgraded because either:

- The DB2 version running the db2nrupgrade command is not a supported version to upgrade the current instance.
- The instance is still active

User response: See the DB2 Information Center for supported versions. Also ensure that there are no DB2 processes running at the instance. Fix the error, and to upgrade the instance to the new DB2 Copy version, rerun the db2nrupgrade command. For example, db2nrupgrade -b *backupDir*

DBI1391E *directory-name* is not a valid path.

Explanation: The directory cannot be accessed.

User response: Enter a valid path and rerun the command.

DBI1392W The *deprecated-command-name* command is deprecated. The deprecated command is replaced by new command *new-command-name*.

Explanation: The deprecated command will be removed from a future release.

User response: Replace the deprecated command with the new command in any scripts or applications.

DBI1393W The -j "TEXT_SEARCH" option is supported only on server instances ese, wse and standalone.

Explanation: The current instance is of type client or DB2 pureScale. The command completed, but the -j "TEXT_SEARCH" option was ignored.

User response: To use DB2 Text Search specify a server instance.

DBI1394E The db2cluster_prepare command failed because the db2cluster_prepare utility cannot write to the following trace file path: *trace-file-path*.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

You can cause the db2cluster_prepare utility to generate a trace file by specifying an absolute path for the trace file using the -t parameter. The db2cluster_prepare utility will create the trace file in the specified location. You can use this trace file to investigate any errors that occur while the GPFS cluster is being created.

This message is returned when the db2cluster_prepare command is called with a trace file location specified with the -t parameter and the db2cluster_prepare utility cannot access the specified trace path.

The db2cluster_prepare utility runs with the same authority as the user who runs the db2cluster_prepare command. This means that the user who runs the db2cluster_prepare command must have access to the path specified with the parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying a trace file path that you can access.

DBI1395E The db2cluster_prepare command failed because the db2cluster_prepare utility cannot write to the following log file path: *log-file-path*.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

The db2cluster_prepare utility generates a log file that you can use to investigate any errors that occur while the GPFS cluster is being created. You can set the location for the log file by specifying an absolute path for the log file using the -log_file_path parameter. The db2cluster_prepare utility will create the log file in the specified location.

This message is returned when the db2cluster_prepare command is called with a log file location specified with the -log_file_path parameter and the db2cluster_prepare utility cannot access the specified log path.

The db2cluster_prepare utility runs with the same authority as the user who runs the db2cluster_prepare command. This means that the user who runs the db2cluster_prepare command must have access to the path specified with the -log_file_path parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying a log file path that you can access.

DBI1396E The db2cluster_prepare command failed because the utility was unable to create necessary temporary files in the following path: *path*.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

The db2cluster_prepare utility generates some temporary files for internal use. This message is returned when the db2cluster_prepare command is unable to create these internal, temporary files.

The db2cluster_prepare utility runs with the same authority as the user who runs the db2cluster_prepare command. This means that to run the db2cluster_prepare command, you must have access to the named path.

A GPFS file system was not created.

User response:

1. Ensure you have permission to access the named path.
2. Ensure that there is enough disk space available to create the temporary files.
3. Run the db2cluster_prepare command again.

DBI1397W The db2set command failed to set the database manager configuration parameter DB2_ATS_ENABLE value to YES.

Explanation: The Administrative Task Scheduler (ATS) could not be enabled. Enabling ATS is a prerequisite for using DB2 Text Search.

User response: To use DB2 Text Search, you must enable the Administrative Task Scheduler.

DBI1398E The db2cluster_prepare command failed because the db2cluster_prepare utility was unable to access the following directory: /tmp.

Explanation: You can create a IBM General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

The db2cluster_prepare utility generates a log file that you can use to investigate any errors that occur while the GPFS cluster is being created. You can set the location for the log file by specifying an absolute path with the -log_file_path parameter. If you do not use the -log_file_path parameter, the db2cluster_prepare utility will attempt to create a log file named db2prepare_cluster.log in the default directory location /tmp.

This message is returned when the db2cluster_prepare command is called without a log file location being specified with the -log_file_path parameter the

db2cluster_prepare utility cannot access the directory named /tmp.

The db2cluster_prepare utility runs with the same authority as the user who runs the db2cluster_prepare command. This means that you must be able to access the directory named /tmp to be able to call the db2cluster_prepare command without specifying the -log_file_path parameter.

A GPFS file system was not created.

User response: Respond to this error in one of the following ways:

- Run the db2cluster_prepare command again, specifying a log file path with the -log_file_path parameter that you can access.
- Allow the db2cluster_prepare utility to create the log file in the default location by performing the following two steps:
 1. Configure the /tmp directory permissions so that you have access to that directory.
 2. Run the db2cluster_prepare command again without using the -log_file_path parameter.

DBI1399I You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

Explanation: In a DB2 pureScale environment, all data and logs must be on a General Parallel File System. Existing data and logs that are not on GPFS must be moved to GPFS before upgrading to a DB2 pureScale environment. You can set up a DB2 managed GPFS cluster and file system, so that the data and logs can be moved to it in preparation for the upgrade to a DB2 pureScale environment using the db2cluster_prepare utility.

User response: No response is necessary. This message is for your information only.

DBI1400N The syntax of the db2licm command is incorrect.

Explanation: The db2licm tool performs basic license functions. It adds, removes, lists, and modifies licenses installed on the local system. Execute db2licm tool with -l parameter to find out the product identifier for your product:

```
db2licm [-a filename]
        [-e product-identifier HARD | SOFT]
        [-p product-identifier
        REGISTERED | CONCURRENT | OFF]
        [-r product-identifier]
        [-u product-identifier num-users]
        [-c product-identifier num-connectors]
        [-l]
        [-v]
        [-?]
```

The command options are:

-a

Adds a license for a product. Specify a file name containing valid license information. This can be obtained from your licensed product CD or contact your IBM representative or authorized dealer.

-e

Updates the enforcement policy on the system. Valid values are: HARD and SOFT. HARD specifies that unlicensed requests will not be allowed. SOFT specifies that unlicensed requests will be logged but not restricted.

-p

Updates the license policy type to use on the system. The keywords CONCURRENT, REGISTERED, or CONCURRENT REGISTERED can be specified. Specify OFF to turn off all policies.

-r

Removes the license for a product. After the license is removed, the product functions in "Try & Buy" mode. To get the password for a specific product, invoke the command with the -l option.

-u

Updates the number of user entitlements that have been purchased. Specify the password of the product for which the entitlements were purchased and the number of users.

-c

Updates the number of connector entitlements that have been purchased. Specify the password of the product for which the entitlements were purchased and the number of connectors.

-l

Lists all the products with available license information, including the product identifier.

-v

Displays version information.

-?

Displays help information. When this option is specified, all other options are ignored, and only the help information is displayed.

User response: Enter the command again using the valid parameters.

DBI1401I Command line DB2 License Manager.

Explanation: The db2licm tool performs basic license functions. It adds, removes, lists, and modifies licenses installed on the local system.

To determine the product identifier for your product issue the db2licm -l command.

```
db2licm [-a filename]
        [-e product-identifier HARD | SOFT]
        [-p product-identifier
REGISTERED | CONCURRENT | OFF]
        [-r product-identifier]
        [-u product-identifier num-users]
        [-c product-identifier num-connectors]
        [-l]
        [-v]
        [-?]
```

The command options are:

-a

Adds a license for a product. Specify a file name containing valid license information. This can be obtained from your licensed product CD or contact your IBM representative or authorized dealer.

-e

Updates the enforcement policy on the system. Valid values are: HARD and SOFT. HARD specifies that unlicensed requests will not be allowed. SOFT specifies that unlicensed requests will be logged but not restricted.

-p

Updates the license policy type to use on the system. The keywords CONCURRENT, REGISTERED, or CONCURRENT REGISTERED can be specified. Specify OFF to turn off all policies.

-r

Removes the license for a product. After the license is removed, the product functions in "Try & Buy" mode.

-u

Updates the number of user entitlements that have been purchased. Specify the password of the product for which the entitlements were purchased and the number of users.

-c

Updates the number of connector entitlements that have been purchased. Specify the password of the product for which the entitlements were purchased and the number of connectors.

-l

Lists all the products with available license information, including the product identifier.

-v

Displays version information.

-?

Displays help information. When this option is specified, all other options are ignored, and only the help information is displayed.

User response:

DBI1402I License added successfully.

DBI1403I License removed successfully.

DBI1404N Product identifier not found.

Explanation: The given identifier is either invalid, or a license for this product was not found in the nodelock file.

User response: Issue this command with -l option to check that the identifier entered is the correct product identifier for the product that you want to perform this action on. If you are using nodelock passwords, check that the license key for this product is installed in the nodelock file.

DBI1405I License policy type updated successfully.

DBI1406N Invalid license policy type.

Explanation: The license policy type that was entered was not valid for the product specified.

User response: Please enter a valid license policy. Options are:

- CONCURRENT
 - REGISTERED
 - CONCURRENT REGISTERED
 - OFF
-

DBI1407N Invalid license certificate file.

Explanation: License certificate file is not in the correct format.

User response: Enter the name of a file with the correct license certificate format.

DBI1408N The file *file-name* could not be opened.

Explanation: The file is not found or access to the file is denied.

User response: Enter the name of a file that exists and can be opened and try the command again.

DBI1409N Invalid enforcement policy type.

Explanation: The enforcement policy type specified is not valid for this product.

User response: Please enter a valid enforcement policy type that is supported by the specified product.

DBI1410I Concurrent entitlements updated successfully.

DBI1411I Enforcement policy type updated successfully.

DBI1412W A hard stop enforcement policy stops the use of the product by unlicensed users when licensing infractions have been exceeded.

DBI1413W A soft stop enforcement policy logs license infractions, but allows unlicensed users to still use the product.

DBI1414I The db2ls command is preparing and checking the installed DB2 copies on the system.

DBI1415I Checking active DB2 instances on AIX system workload partitions (WPARs) *workload-partition-list* for the related DB2 copy.

DBI1416N The license could not be added to the nodelock file automatically.

Explanation: The return code is "<return-code>".

User response: Please ensure the license certificate is readable. You may also enter the license into the nodelock file manually. Please see the license file for instructions.

DBI1417N The license specified could not be removed from the nodelock file.

Explanation: The return code is "<return-code>".

User response: Ensure that the license for this product exists in the nodelock file.

DBI1418I The number of licensed processors on this system has been updated successfully.

DBI1419N There was an error updating the number of licensed processors.

Explanation: The return code is "<return-code>".

DBI1420N This product does not support this type of license policy.

Explanation: The license policy specified does not apply to this product or is not supported.

User response: Please enter a valid license policy.

DBI1421N This product specified is not installed on this system.

Explanation: You can not configure a license policy for a product until the product is installed.

User response: Install the product before running this command or specify the correct product identifier. To list the products install on the system issue db2licm -l command.

DBI1422N The number of concurrent entitlements was not updated.

Explanation: The return code is "<return-code>".

User response: Please ensure the concurrent policy is enabled for this product.

DBI1423N This option requires the creation of an instance.

Explanation: Features that are required to perform this action are only accessible once an instance has been created.

User response: Please create an instance and issue this command again.

DBI1424N An unexpected error occurred while accessing processor information.

Explanation: The return code is "<return-code>".

User response: None.

DBI1425E The license for DB2 OLAP Server cannot be updated. The DB2 OLAP processes are currently active.

Explanation: DB2 cannot update the license for DB2 OLAP Server while DB2 OLAP Server is running.

User response: To update your OLAP license, please stop all OLAP processes and reinstall this DB2 license.

DBI1426I This product is now licensed for use as specified in the License Acceptance and License Information documents pertaining to the licensed copy of this product. USE OF THE PRODUCT CONSTITUTES ACCEPTANCE OF THE TERMS OF THE IBM LICENSE ACCEPTANCE AND LICENSE INFORMATION DOCUMENTS, LOCATED IN THE FOLLOWING DIRECTORY: *dir-name*

DBI1427I This product is now licensed for use as specified in the License Acceptance and License Information documents pertaining to the evaluation ("Try and Buy") copy of this product. USE OF THE PRODUCT CONSTITUTES ACCEPTANCE OF THE TERMS OF THE IBM LICENSE ACCEPTANCE AND LICENSE INFORMATION DOCUMENTS, LOCATED IN THE FOLLOWING DIRECTORY: *dir-name*

DBI1428N There was an error updating the number of licensed processors.

Explanation: The number of licensed processors entered exceeds the number of maximum licensed processors allowed for this product.

User response: Please enter number of licensed processors that does not exceed the defined maximum. If the number of processors on your system exceeds the maximum number of processors allowed for this product, please contact your IBM representative or authorized dealer.

DBI1429N This product does not support this combination of license policies.

User response: Please enter a valid combination of license policies. For example, you can specify "CONCURRENT REGISTERED" as a valid combination.

DBI1430N The license could not be added to the nodelock file because the license date is greater than operating system date.

User response: Please check your certificate file to ensure that the license start date precedes the current date (the date set on the operating system).

DBI1431N This user does not have sufficient authority to perform the specified action.

Explanation: This program can be run only under the root user ID or a user ID with SYSADM authority.

User response: Login with a user ID that has

permission to run this command.

DBI1432N The license could not be added to the nodelock file.

Explanation: This product has used the maximum number of evaluation licenses. The maximum number of evaluation licenses is *lic-number*.

User response: Run this command again with a permanent license key.

DBI1433N The number of license entitlements was not updated.

Explanation: The specified number of license entitlements is not in the valid range.

User response: Run this command again using a valid number of license entitlements.

DBI1434N DB2 has added the license entry to the nodelock file, however, this license entry is not active.

Explanation: DB2 failed to activate this license entry, therefore DB2 will run with the previous license configuration until this license is activated.

User response: Try the command again and if it continues to fail, edit the nodelock file manually or contact IBM support.

If you edit the nodelock file manually, move the new license entry to the top of the license entries list.

The location of the nodelock file is platform specific:

AIX /var/ibm/nodelock

Windows

\$DB2PATH/license/nodelock

All other operating systems:

/var/lum/nodelock

Refer to the DB2 Information Center for more information on licensing.

DBI1435E Error Opening DB2 First Steps. Cannot find supported Web browser.

Explanation: DB2 First Steps requires a supported Web browser.

User response: Install one of the following Web browsers:

- Internet Explorer 6.0 and up
- Mozilla 1.7 and up
- Firefox 2.0 and up

When using Windows operating systems, set the default browser to your choice from the previous list.

DBI1437E The db2cluster_prepare command failed because no value was specified for the -t parameter.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

You can cause the db2cluster_prepare utility to generate a trace file by specifying an absolute path for the trace file using the -t parameter. The db2cluster_prepare utility will create the trace file in the specified location. You can use this trace file to investigate any errors that occur while the GPFS cluster is being created.

This message is returned when the db2cluster_prepare command is called without a value specified for the -t parameter.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying a value for the -t parameter.

DBI1438E The db2cluster_prepare command failed because no value was specified for the parameter named -l.

Explanation: You can create a DB2 managed General Parallel File System (GPFS) for a DB2 pureScale environment using the db2cluster_prepare utility.

The db2cluster_prepare utility generates a log file that you can use to investigate any errors that might occur while the GPFS cluster is being created. You can set the location for the log file by specifying an absolute path with the parameter named -l. If you use the parameter named -l, the db2cluster_prepare utility will create the log file in the specified location.

This message is returned when the db2cluster_prepare command is called without a value specified for the parameter named -l.

A GPFS file system was not created.

User response: Run the db2cluster_prepare command again, specifying a value for the parameter named -l.

DBI1439W The option -immediate was used when running db2set, but the variable specified cannot be updated immediately

Explanation: The registry was successfully updated with the new value, however, it is never possible for the change to the variable specified by the db2set -immediate call to take effect immediately. The -immediate option has been ignored.

User response: Do not supply the -immediate option when updating this specific variable.

DBI1440W A db2set -immediate call was made, the offline registry was successfully updated, but db2set could not attach to the instance

Explanation: Although the offline registry was updated successfully, the value cannot be applied immediately by one or more members because db2set could not attach to the instance.

User response: Try attaching to the instance explicitly and re-run the command and verify that all members of the instance are healthy.

DBI1441W A db2set -immediate call was made but the dynamic update failed on one or more hosts or members

Explanation: The offline registry was successfully updated, and the registry update has been made immediately on all of the hosts or members db2set could access.

On the hosts or members that db2set could not access, however, the setting for the variable specified could not be applied immediately.

User response: Make sure that all members can be accessed from the host where the command was issued and re-run the command or try attaching to each host or member explicitly and re-run the command. Alternatively, all members will see the changes the next time they restart.

DBI1442W Some variables governed by the aggregate variable *aggregate_name* were updated dynamically, although some variables it governs do not support an immediate change.

Explanation: The aggregate registry variable was successfully updated immediately. Some of the variables governed by the aggregate variable were changed dynamically, but other variables governed by the aggregate variable require the instance to be restarted to take effect.

User response: Restart the instance for all of the variables governed by this aggregate registry variable to take effect.

DBI1443W One or more registry variables meant to take effect for the whole instance have different values on different hosts.

Explanation: On the Windows operating system, registry variables are stored within the operating system registry of each host. Currently at least one registry variable has been set to a number of different values across hosts. It is highly recommended to keep each registry variable set consistently on all hosts.

User response: Make sure that all DB2 registry

variable settings are consistent on all hosts and re-run the command.

DBI1444E The db2ls command cannot be found in the expected directory *directory-name*.

Explanation: The DB2 installation path passed to the db2ls command with the -b option must be the location of a copy with DB2 Version 9 or later installed.

User response: Specify a DB2 installation path of a copy with DB2 Version 9 or later installed.

DBI1445E The -b install-path option is mandatory when the -q option is specified. For a list of DB2 installation locations to query, run the db2ls command without the -b and -q options.

DBI1446I The %1 command is running.

DBI1447E A Java Runtime Environment is needed to run the program *program-name*.

Explanation: The execution of the command requires a Java Runtime Environment (JRE). If IBM Software Development Kit (SDK) for Java is installed with the related DB2 copy, the JRE is found in <DB2DIR>/java/jdk64/jre/bin for 64-bit Linux systems, or <DB2DIR>/java/jdk32/jre/bin for 32-bit Linux systems, where <DB2DIR> represents the installation path of the related DB2 copy.

If a JRE is not found in either path, the program looks for a JRE in either:

- the current user's \${JAVA_HOME}/jre/bin, or
- the default PATH environment.

In this case, the JRE must meet the current DB2 version requirements. For details, see "Java software support for DB2 products" in the DB2 Information Center.

User response: Install the IBM Software Development Kit (SDK) for Java and rerun the program.

DBI1448E An error occurred during the creation or removal of a DB2 entry on the main menu. For details, check the log file *log-file-name*.

Explanation: On Linux operating systems, main menu entries for DB2 tools can be created and removed manually or automatically. Main menu entries for DB2 tools, such as db2fs, can be created manually by running the db2addicons command, or created automatically when specific DB2 commands are run. These main menu entries can be removed manually by running the db2rmicons command, or removed automatically when specific DB2 commands are run. Errors occurred during either the creation or removal of

these main menu entries. For a list of the specific DB2 commands, see "Main menu entries for DB2 tools" in the DB2 Information Center.

User response: Check the error log file for details and rerun the program.

DBI1449I Usage: db2addicons [-h]

Explanation: The db2addicons command creates main menu entries for DB2 tools, such as db2fs, for the current user. The main menu entries for DB2 tools are created by manually running the db2addicons command, or are created automatically when specific DB2 commands are run. For a list of the specific DB2 commands, see "Main menu entries for DB2 tools" in the DB2 Information Center.

The only argument for this command is:

-h

displays the usage information.

User response: Enter the db2addicons command again without any argument.

DBI1450I Usage: db2rmicons [-h]

Explanation: The db2rmicons command removes main menu entries for DB2 tools, such as db2fs, for the current user. The main menu entries for DB2 tools are removed by manually running the db2rmicons command, or are removed automatically when specific DB2 commands are run. For a list of the specific DB2 commands, see "Main menu entries for DB2 tools" in the DB2 Information Center.

The only argument for this command is:

-h

displays the usage information.

User response: Enter the db2rmicons command again without any argument.

DBI1451E A DB2 instance environment is not set to run the program *program-name*.

Explanation: On Linux operating systems, the DB2 instance environment needs to be set for the current user before the command can be run. DB2 instance environment can be setup with:

- <Instance_HOME>/sqlib/db2profile (for Bourne shell and Korn shell users), or
- <Instance_HOME>/sqlib/db2chsrc (for C shell users)

where <Instance_HOME> is the instance owner's home directory.

User response: Set up the DB2 instance environment and rerun the program.

DBI1452E You do not have the authority to install the network version of the DB2 Information Center. Installing the product named *product-name* requires root authority. You can find the workstation version of the DB2 Information Center in the subdirectory called "workstation" of the folder located at *folder-location*, at IBM Passport Advantage, or at the following URL: www.ibm.com/support/docview.wss?rs=71&uid=swg27009474. You do not need to have root authority to install the workstation version of the DB2 Information Center.

DBI1453E Invalid component name *component-name*

Explanation: Specified component name is either incorrect or misspelled. If the component name is correct, this component name might not be valid on this operating system or platform.

User response: Check to see if the component name is a valid DB2 component for the current operating system or platform. Refer to the DB2 Information Center or to the sample uninstallation response DB2DIR/db2un.rsp for a list of valid component keys.

DBI1454E *instance-name* is a DB2 pureScale instance. Validation on this type of instance is not supported in the current DB2 release.

Explanation: The db2val command does not support validation on a DB2 pureScale instance.

User response: No action required.

DBI1455W The information on instance *instance_name* will not be added to the generated response file.

Explanation: The specified instance is a DB2 pureScale instance. When the db2rspgn command is run against a DB2 pureScale instance the information on the specified instance is not added to the generated response file and the instance profile is not created.

User response: No action required

DBI1456W The *command-name* command failed because it was executed on a member that is not the coordinating member of the instance *instance-name*.

Explanation: The command is only supported on the coordinating member of the instance.

User response: Run the command on the coordinating member of the instance.

DBI1457E **The DB2 product or feature could not be uninstalled. To uninstall, the following list of instances must be dropped:**
instance-list.

Explanation: The product or feature must not be associated with any instances when it is uninstalled. The db2ilist command is used to list installed instances.

User response: Drop the offending instances using the db2idrop command and re-run the db2_deinstall command.

DBI1458E **The db2iupdt or db2iupgrade commands are not supported on a non instance-owning member.**

Explanation: The db2iupdt command can be used to update an instance to a higher level within a release. The db2iupgrade command upgrades an instance to a DB2 copy of the current release from a DB2 copy of a previous release. On the Windows platform the DPF instance update and upgrade are only supported on the instance-owning member.

User response: Run the db2iupdt or db2iupgrade command on the instance-owning member.

DBI1459E **The db2_deinstall command failed because some DB2 database files are in use.**

Explanation: You can uninstall DB2 database products, features, or languages using the db2_deinstall command.

Before you can uninstall any part of a non-root installed DB2 database manager instance, all operating system processes and resources for that instance must be release. This message is returned when an attempt is made to uninstall a non-root installed instance while some DB2 database-related system resources have not been freed. For example, this message can be returned for the following reasons:

- DB2 database-related operating system processes are still running
- DB2 database-related libraries are still loaded in memory

User response:

1. Stop the DB2 database manager using the db2stop force command
2. Run the db2_deinstall command again

DBI1460E **Command parameter combination is invalid.**

Explanation: When using the db2setup command, if the -c parameter is specified to validate the contents of a response file, the -r parameter indicating the response file must also be specified.

User response: Rerun the db2setup command specifying both the -c and -r parameters.

DBI1461E **The distributed fix pack installation failed because the -H parameter was specified with the installFixPack command, but the DB2 installer detected no DB2 pureScale components in the base copy of the DB2 database product.**

Explanation: You can update the installed DB2 database products to the same level by using the installFixPack command.

In DB2 pureScale environments only, you can update multiple hosts in a DB2 cluster by specifying the -H parameter with the installFixPack command.

This message is returned when the -H parameter is specified with the installFixPack command in an environment that is not a DB2 pureScale environment.

User response: Rerun the installFixPack command without the -H parameter.

DBI1462I **The db2setup command is validating the response file provided. This can take several minutes.**

Explanation: Validation process in progress.

User response: No action required.

DBI1463E **The response file validation failed. Log file: *path_to_logfile***

Explanation: Invalid entries were detected in the response file. For more details, refer to the log file.

User response: Correct the response file errors and rerun the command.

DBI1464E **Conflicting installFixPack command parameters specified. Parameters: *conflicting-parameter-list.***

Explanation: There are specific parameters related to specific actions in the installFixPack command. Some parameters cannot be specified for some actions.

This message is returned when two parameters of the installFixPack command cannot be specified together because the syntax does not support such usage.

User response: Review the syntax of the installFixPack command and re-issue the command with the correct parameters.

DBI1465E **An error occurred in the host list file.**
Host file: *host-file-name*.

Explanation: An error was detected in the specified host list file. The error can occur for any of the following reasons: the file does not exist, the file is not readable, the file is empty or contains an invalid character, the path name is not specified, or the HOST keyword is not valid.

User response: Check the contents of the specified host list file, and rerun the command.

DBI1466E **Instance validation failed. Instance name:** *instance_name*. **Reason code:** *reason_code*. **Current host name:** *host_name*.

Explanation: The specified reason code indicates one of the following errors occurred:

- 1
The IBM Reliable Scalable Cluster Technology (RSCT) peer domain is not online
- 2
The IBM General Parallel File System (GPFS) cluster is not online
- 3
At least one db2 cluster services alert is indicated

User response: Depending on the specified reason code, take the appropriate action:

- 1
Bring the RSCT peer domain online by entering the following command: `db2cluster -cm -start -domain domain_name`
 - 2
Start the GPFS cluster by entering the following command: `db2cluster -cfs -start -host host_name`
 - 3
Clear all alerts by running the following command: `db2cluster -list -alert`
-

DBI1467N **The db2ndrop command was not able to drop the multiple partition instance member 0.**

Explanation: The db2ndrop command is used to drop a database partition server from an instance that has no databases.

The db2ndrop command cannot drop the instance owning database partition server (0). To drop a database partition server that is assigned to the logical port 0 from a machine that is running multiple logical

database partition servers, all other database partition servers assigned to the other logical ports must be dropped first. Each database partition server must have a database partition server assigned to logical port 0.

User response: Use the db2idrop command to drop the multiple partition instance member 0, after the rest of members have been dropped using the db2ndrop command.

DBI1468E **Instance validation failed. Instance:** *instance_name*. **Log file:** *path_to_logfile*.
Hosts: *host_name*.

Explanation: The validation failed on the specified hosts.

User response: Check the log file for more details.

DBI1469W **The DB2 copies of the hosts participating in the instance do not have the same code level or install path.**
Instance: *instance_name*.

Explanation: The instance might be in the rolling update stage.

User response: Update your instance to the same level and the same path on all participating hosts in the instance using the db2iupdt command.

DBI1470E **The instance validation failed on host** *host_name*. **Reason code =** *reason_code*.

Explanation: The specified reason code indicates one of the following errors occurred:

- 1
Validation on the specified host did not finish before time-out.
- 2
The SSH communication for the instance owner between the specified host and other hosts failed.
- 3
Instance shared directory cannot be accessed from the given host or is not shared between the given host and the local host.
- 4
Some files in the instance directory are damaged. The symbolic links of the files are not pointing to the instance shared directory.
- 5
The specified host is offline in the RSCT peer domain.
- 6

The specified host is stopped in the GPFS cluster.

7

Installation file validation failed for the DB2 copy on the specified host.

User response: Depending on the specified reason code, take the appropriate action:

1

Rerun the db2val command on the specified remote host.

2

Fix the SSH communication for these hosts. This requires you setup a passwordless SSH access for the instance owner. For details, see the DB2 Information Center.

3

The instance shared directory must be accessible from the specified host. The directory path must be accessible from all the hosts, and must have read/write permission.

4

To fix the symbolic links, run db2iupdt <instance_name> as root.

5

Bring the host online by running the following command: db2cluster -cm -start -host <host-name>

6

Bring the host online by running the following command: db2cluster -cfs -start -host <host-name>

7

Check the log file for more details.

DBI1471I db2setup validated the response file successfully. Log file: *path_to_logfile*.

Explanation: The validation process completed successfully.

User response: Check the log file for details.

DBI1472W The db2val -d parameter is deprecated. The deprecated parameter is replaced by the -t *trace_file* parameter.

Explanation: The deprecated parameter will be removed from a future release.

User response: Replace the deprecated parameter with the new parameter in any scripts or applications.

DBI1473W Parameter *deprecated-param-name* is deprecated. The deprecated parameter is replaced by new parameter *new-param-name*. Command: *command-name*.

Explanation: The deprecated parameter might be removed from a future release.

User response: Replace the deprecated parameter with the new parameter in any scripts or applications. For command syntax and parameter details, see the DB2 Information Center.

DBI1474E The instance shared mount point path is invalid. Command: *command-name*.

Explanation: In the specified command, the path specified for the instance shared mount point must be a new path and must not be a sub-directory of an existing GPFS file system.

User response: Rerun the command specifying a new path for the instance shared mount point.

DBI1475E The installation failed because only the root user is permitted to install DB2 pureScale.

Explanation: To install the DB2 pureScale Feature you must be logged in as the root user. Installation is restricted to the root user because components required by DB2 pureScale affect shared resources on the system.

User response: Review the prerequisites for DB2 pureScale before attempting another installation.

DBI1476E The db2_install command failed because the following mutually exclusive options were specified: "-f PURESCALE" and "-f NOSTAMP".

Explanation: The IBM Tivoli Storage Automation for MultiPlatform (SA MP) is a required component for the DB2 pureScale Feature. Running the db2_install command with the "-f PURESCALE" option and the option to exclude Tivoli SA MP from the installation, "-f NOSTAMP", is not permitted.

User response: Respond to this message in one of the following ways:

- To include DB2 pureScale in the installation, rerun the db2_install command without specifying the "-f NOSTAMP" option.
- If you do not want to install DB2 pureScale, rerun the db2_install command without specifying the "-f PURESCALE" option.

DBI1477E The db2cluster_prepare command failed because of a syntax error. Reason code: *reason-number*.

Explanation: The syntax entered was invalid. See the corresponding reason code for an explanation of the cause of the error.

1. The -cfs_takeover and -instance_shared_dev option are mutually exclusive, when specifying one option for db2cluster_prepare the other cannot be used.
2. The db2cluster_prepare command failed because one or more parameters were specified multiple times.
3. The db2cluster_prepare command failed because no value was specified for the -instance_shared_dev parameter.
4. The db2cluster_prepare command failed because no value was specified for the -t parameter.
5. The db2cluster_prepare command failed because no value was specified for the parameter named -l.
6. The db2cluster_prepare command failed because no value was specified for the -instance_shared_dev parameter.

User response: Rerun the db2cluster_prepare command with the correct syntax. For the correct syntax, see the command reference or run the command with the "-?" option, by typing "db2cluster_prepare -?".

DBI1478E The GPFS binary update failed. Installation detected one or more GPFS efices that were not originally installed by the DB2 product. Installation cannot proceed until these efices are manually removed. efices: *efix_list*

Explanation: The GPFS efices detected by the DB2 installer were not installed by the DB2 product. The DB2 installer can only remove efices installed as part of a DB2 product installation. The specified efices must be removed manually.

User response: Manually remove the efices, and rerun the DB2 command.

DBI1479N The syntax of the DB2NDROP command is incorrect.

Explanation: The DB2NDROP utility drops a node from a partitioned system.

DB2NDROP /n:node
[/i:instance]

The command options are:

- /i Specify the instance if different from the default/current instance

User response: Issue the DB2NDROP command with one of the valid command options.

DBI1480N Node *node* not found for instance *instance*.

Explanation: DB2NDROP failed because the node does not exist.

User response: Ensure the node number is correct and reissue the command.

DBI1481W Node *node* for instance *instance* has been deleted.

Explanation: The DB2NDROP processing has completed successfully.

User response: No further action is required.

DBI1482E Syntax error, the db2_deinstall command options to skip the removal and remove the same feature are mutually exclusive.

Explanation: The '-s' option skips the removal of a feature, and the '-F feature-name' option specifies that feature feature-name, should be removed. These two options must not contain the same feature identifier.

User response: If you want to remove the feature, rerun the command without the '-s feature-id' option. If you want to skip the removal of the feature, rerun the command without the '-F feature-id' option.

DBI1483E Syntax error, the '-F GPFS' option is not valid.

Explanation: The GPFS feature cannot be removed with '-F GPFS' option of the db2_deinstall command. By default, the GPFS feature will be removed with the 'db2_deinstall -a' command when you remove the last DB2 copy and GPFS is not still in use, if GPFS was installed by the DB2 installer.

User response: Rerun the db2_deinstall command without the '-F GPFS' option. See the uninstall documentation for more information.

DBI1493W A Reliable Scalable Cluster Technology (RSCT) peer domain was detected, IBM Tivoli System Automation for Multiplatforms (SA MP) cannot be removed by the DB2 uninstall process.

Explanation: If an RSCT peer domain is present, the Tivoli SA MP component is not uninstalled with the DB2 copy to prevent adverse behavior of other copies or applications, which might depend on the peer domain or Tivoli SA MP.

User response: If there are no other DB2 copies or applications dependant on Tivoli SA MP component and the component is no longer required, remove it.

Chapter 85. DBI1500 - DBI1999

DBI1500I **Usage:**
db2inst [-f response-file]

Explanation: An incorrect argument was entered. The command accepts an optional flag -f which is followed by the path of the “DB2 Installer” response file.

User response: Re-enter the command with a valid argument.

DBI1501E **An internal error was encountered.**

Explanation: An error was detected while performing an internal operation.

User response: Verify that the file system on which the file resides is not damaged. If the problem persists, contact IBM Support with the following information:

1. Message number
2. Internal error message description
3. Problem description

DBI1502E **An error was encountered when opening or reading from file, file-name.**

Explanation: An error was encountered when opening or reading from a file. One of the following errors occurred:

- An I/O error occurred opening or reading from the file.
- An expected value was not found in the file.
- The order of the data in the file is incorrect.

User response: If the file was modified by the user, verify that the data in the file is valid. If the file was not modified by the user, contact IBM Support with the following information:

1. Message number
2. Message description
3. Problem description

DBI1503E **An error was encountered when opening or writing to file, file-name.**

Explanation: An error was encountered when writing data to a file. One of the following errors occurred:

- Incorrect access permissions on a given directory.
- There is not enough space on your file system.

User response: Check to see if the permissions on the file's parent directory have been set correctly. If the problem persists, contact IBM Support with the following information:

1. Message number

2. Message description
3. Problem description

DBI1504E **An error occurred opening message files.**

Explanation: An attempt to open and read from “DB2 Installer” message catalog files failed. Before DB2 Installer starts up, it looks for two resource files -- db2inst.rcx and db2inst.cat in the following locations:

- DB2 product directory, or
- on the distribution media, or
- environment variables DB2IRCX and DB2ICAT.

User response: Set the environment variables DB2IRCX and DB2ICAT to the absolute location of db2inst.rcx and db2inst.cat respectively.

DBI1505E **An attempt to allocate memory failed.**

Explanation: An error was detected when attempting to allocate memory.

User response: Terminate other applications running on the system that may be using large amounts of memory. If the problem persists, contact IBM Support with the following information:

1. Message number
2. Message description
3. Problem description

DBI1506E **Installation failed because IBM Tivoli System Automation for Multiplatforms (SA MP) could not be upgraded.**

Explanation: The Tivoli SA MP product must be at the version level included in the DB2 installation image. To avoid any potential problems with the existing peer domain, the DB2 installer does not attempt to upgrade Tivoli SA MP when a peer domain is in use. There is a Reliable Scalable Cluster Technology (RSCT) peer domain on the system preventing the installer from upgrading the Tivoli SA MP installation.

User response: Manually upgrade the Tivoli SA MP product before attempting to install the DB2 pureScale Feature.

DBI1507E **An instance of the DB2 Installer is already started.**

Explanation: An error was detected when attempting to start up the DB2 Installer. Another instance of the DB2 Installer is still running.

User response: Terminate all instances of the DB2 Installer and restart the install process. If the problem persists, remove the lock file /tmp/.db2inst.lck and restart the DB2 Installer.

DBI1508W IBM Tivoli System Automation for Multiplatforms (SA MP) was updated as part of the DB2 product installation or upgrade process.

Explanation: Only one SA MP copy can be installed on a system, and the DB2 installation required that the SA MP product installed be upgraded to the version of SA MP integrated with the DB2 installation media.

User response: If other software has dependencies on SA MP verify that the software is compatible with the updated SA MP version.

DBI1509W The DB2 installer failed to update to the required version of IBM Tivoli System Automation for Multiplatforms (SA MP). Installed version number of Tivoli SA MP: *installed-SA-MP-version*. Required SA MP version level: *required-SA-MP-version*.

Explanation: Failure to upgrade Tivoli SA MP can be caused if the Reliable Scalable Cluster Technology (RSCT) peer domain is offline or in maintenance mode.

User response: Manually upgrade the Tivoli SA MP product after the DB2 installation or update completes. See the Tivoli SA MP Base Component Installation and Configuration Guide for more information.

DBI1510E The DB2 installer failed to install a DB2 copy because non-root users are restricted to a single DB2 copy and a copy preexists for the user.

Explanation: The DB2 instance processes are run with the user ID for non-root DB2 installations. All DB2 instances are required to run with a unique ID, this restricts a non-root user to a single copy.

User response: Upgrade the installed non-root DB2 copy with by upgrading with a response file, specifying the keyword "UPGRADE_PRIOR_VERSION". Alternatively, you can run "db2_install -m" to upgrade the installed copy.

DBI1511E Installing the IBM Tivoli System Automation for Multiplatforms (SA MP) failed. See the log file *log-file-name* for details. (TSAMP_RC=TSAMP_RC).

Explanation: The DB2 installer uses the installSAM utility to install SA MP. The installSAM utility returned errors. The installSAM log file contains more complete information.

TSAMP_RC is the return code from the installSAM utility.

Note that you must have root authority to use the DB2 installer to install SA MP.

SA MP was not installed.

User response: See the mentioned log file for details.

To install SA MP manually, use the installSAM command.

For more information about the installSAM command, see the SA MP Base Component documentation.

DBI1512E The installation failed because the operating system does not support the installation of DB2 pureScale.

Explanation: The DB2 product is supported on more operating systems than the DB2 pureScale Feature. The detected operating system does not support DB2 pureScale.

User response: Review the prerequisites for DB2 pureScale before attempting another installation.

DBI1513E Uninstalling the IBM Tivoli System Automation for Multiplatforms (SA MP) failed. See the log file *log-file-name* for details. (TSAMP_RC=TSAMP_RC).

Explanation: db2_deinstall uses the uninstallSAM utility to uninstall SA MP. The uninstallSAM utility returned errors. The uninstallSAM log file contains more complete information.

TSAMP_RC is the return code from the uninstallSAM utility.

Note that you must have root authority to use db2_deinstall to uninstall SA MP.

SA MP was not uninstalled.

User response: See the mentioned log file for details.

To uninstall SA MP manually, use the uninstallSAM command.

For more information about the uninstallSAM utility, see <http://publib.boulder.ibm.com/tividd/td/IBMTivoliSystemAutomationforMultiplatforms2.2.html>.

DBI1514I The DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP) were successfully installed.

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

These DB2 HA scripts are located at /usr/sbin/rsct/sapolicies/db2. The DB2 installer detects whether these DB2 HA scripts need to be installed or updated.

User response: No action is required.

DBI1515E An attempt to allocate resource for a subprocess failed.

Explanation: An error was detected when attempting to startup a subprocess.

User response: Terminate other applications running on the system that may be using large amounts of resources. If the problem persists, contact IBM Support with the following information:

1. Message number
2. Problem description

DBI1516E An attempt to terminate a subprocess failed.

Explanation: An error was detected when attempting to terminate a subprocess.

User response: Terminate other applications running on the system that may be using large amounts of resources. If the problem persists, contact IBM Support with the following information:

1. Message number
2. Problem description

DBI1517E An attempt to execute a command in a subprocess failed.

Explanation: An error was detected when attempting to execute a command in a subprocess. One of the following problems occurred:

- The command does not exist.
- Incomplete command search path.
- Incorrect access permissions on the command.
- System resource problem.

User response: Correct the problem and try again. If the problem persists, contact IBM Support with the following information:

1. Message number
2. Problem description

DBI1518E The DB2 installer returned an error while installing or updating the DB2 High Availability (HA) scripts for the IBM Tivoli System Automation for Multiplatforms (SA MP).

Explanation: You need DB2 HA scripts to use SA MP with the DB2 HA feature.

These DB2 HA scripts are located at /usr/sbin/rsct/sapolicies/db2. The DB2 installer detects whether these DB2 HA scripts need to be installed or updated.

The DB2 installer tried to install or update the DB2 HA scripts, but was not successful.

User response: Use the command db2cpts a to

manually install or update these scripts.

DBI1519E The DB2 product cannot extend the instance. The set up of DB2 cluster services failed on the remote host named *host-name*.

Explanation: The DB2 product is not able to extend the DB2 pureScale instance to the host. The set up of DB2 cluster services failed on the remote host.

User response: Examine the log file db2iupdt.log in /tmp/db2iupdt/ on the host, correct the conditions, and rerun the db2iupdt command.

DBI1520E Minimum terminal display size is 24 x 80.

Explanation: Your current terminal or window size is not large enough. This program requires a minimum terminal display size of 24 lines by 80 columns.

User response: Check your terminal or window size and try again.

DBI1521E An attempt to read in the terminal capability information failed.

Explanation: An error was detected when attempting to initialize your terminal display. There are two conditions where this type of error would occur:

- The function cannot find the environment variable TERM or it is unable to find the correct terminfo database entry for the terminal.
- There is not enough memory space available when initializing your terminal display.

User response: Set your environment variable TERM to a correct terminal type. If the problem persists, try to terminate other applications running on the system that may be using large amounts of memory. If the problem still occurs, contact IBM Support with the following information:

1. Message number
2. Problem description

DBI1522E An attempt to restore the terminal capability information failed.

Explanation: An error was detected when attempting to restore your terminal display to the original state.

User response: Set your environment variable TERM to a correct terminal type. If the problem persists, try to terminate other applications running on the system that may be using large amounts of memory. If the problem still occurs, contact IBM Support with the following information:

1. Message number
2. Problem description

DBI1523E **The DB2 product cannot change the preferred primary cluster caching facility to the host *host-name*.**

Explanation: The host name is not currently a cluster caching facility. The host is a DB2 member.

User response: Rerun this command with a cluster caching facility name.

DBI1525E **To completely drop a DB2 pureScale instance, the -g option must be used with the db2idrop command**

Explanation: The db2idrop command was run without using the -g option for a DB2 pureScale instance. To completely drop a DB2 pureScale instance, the -g option is required.

User response: If the DB2 pureScale instance needs to be completely dropped on all the hosts, rerun the db2idrop command with the -g option. Otherwise, run the db2iupdt command with the -drop option. For command syntax details, see the DB2 Information Center.

DBI1527E **The installation failed because the response file used specified the installation of the IBM General Parallel File System (GPFS) but did not specify DB2 pureScale.**

Explanation: The GPFS license is provided for use with the DB2 pureScale Feature only. The DB2 pureScale must be selected in the installation response file to install GPFS.

User response: Edit the installation response file to select both components for installation.

DBI1528E **The installation failed because the product installation path contains a version of a DB2 product that is not compatible with the version of the product on the installation image. Installation path: *DB2DIR*.**

Explanation: To add a product to a DB2 copy, the installed product and the product being installed must be at the same version and fix pack level.

User response: You can update the installed copy to the version on the installation image, before installing the new product to the same path. Alternatively, you can install a new DB2 copy.

DBI1529W **Manual installation of the added languages, products, or both is required on the other hosts in the instance. Added languages: *added-languages*. Added products: *added-products*.**

Explanation: When the installation-initiating host

installs the DB2 pureScale Feature, the installation image copied to members does not include language packs or products that were not selected. Excluding these files from the installation image results in a smaller disk footprint and less network traffic during installation. When adding a language feature or product after the initial DB2 pureScale instance is created, they must be manually installed on each of the other members of the instance because they do not exist in the installation image on the other members. After adding additional language features or products, new members added to the DB2 pureScale instance will include all languages and products that are installed on the installation-initiating host.

User response: Add the language features and products to other members by manually installing the products added to this host on the others.

DBI1530E **Cannot unpack image file.**

Explanation: An error has occurred while unpacking the tar image from the distribution media. One of the following situations has occurred:

- The given path name is a non-existing directory.
- The file system does not have enough space.
- The access permission on the given path is incorrect.
- The tar command was not found in your current search path.

User response: Restart the install process and retry with a correct path name.

DBI1531E **The db2val command is not supported for instance type *instance-type*.**

Explanation: This instance type cannot be validated using the db2val command.

User response: No action required.

DBI1532W **IBM Tivoli System Automation for Multiplatforms (SA MP) will be updated as part of the DB2 product installation or upgrade process.**

Explanation: Only one Tivoli SA MP copy can be installed on a system. The DB2 installation requires that the installed Tivoli SA MP product be upgraded to the version of SA MP integrated with the DB2 installation media.

User response: If other software has dependencies on Tivoli SA MP verify that the software is compatible with the updated Tivoli SA MP version.

DBI1533E **The db2_deinstall command failed because the db2_deinstall utility was unable to uninstall the GPFS product. The db2_deinstall utility was unable to uninstall the GPFS product because the host machine on which the db2_deinstall command was run is in a GPFS cluster.**

Explanation: You can uninstall DB2 database products, features, or languages on Linux and UNIX operating systems by using the db2_deinstall command. If IBM General Parallel File System (GPFS) software was installed as part of a DB2 pureScale Feature installation, the db2_deinstall utility will uninstall the GPFS product by default.

In order for the db2_deinstall utility to uninstall the GPFS product, the GPFS cluster and file system that the DB2 product or feature being uninstalled is using must be manually removed before the db2_deinstall command is run. You can cause the db2_deinstall utility to skip uninstalling the GPFS product by specifying the "-s GPFS" option.

This message is returned when the db2_deinstall command is run on a host machine that is part of a GPFS cluster and the "-s GPFS" option was not specified with the db2_deinstall command.

User response: Respond to this error in one of the following ways:

- If there are no DB2 database manager instances using a GPFS cluster on the current host machine, perform the following steps:
 1. Remove the GPFS cluster and file system manually.
 2. Complete the process of uninstalling DB2 database products and Features as well as uninstalling the GPFS product by calling the db2_deinstall command again.
- If other instances are using a GPFS cluster on the current host machine, uninstall the DB2 product or Feature without uninstalling the GPFS product by calling the db2_deinstall command again, specifying the "-s GPFS" option.

DBI1534E **The db2_deinstall command failed because of IBM Reliable Scalable Cluster Technology (RSCT) peer domain membership.**

Explanation: The db2_deinstall command cannot continue because the current host is part of a RSCT peer domain.

User response: If IBM Tivoli System Automation for Multiplatforms (SA MP) is required on this host, remove any DB2 pureScale instances that are on this host using either the "db2iupdt -drop" or "db2idrop -g" command. If the peer domain is needed, by-pass the Tivoli SA MP binary removal by running "db2_deinstall

-a -s TSAMP". If there are no resources using the peer domain on this host, before re-issuing the command, manually remove Tivoli SA MP. For details, see the DB2 Information Center.

DBI1535W **The DB2 installer cannot update IBM Tivoli System Automation for Multiplatforms (SA MP) to the required version level. Installed version number of Tivoli SA MP: *installed-SA-MP-version*. Required Tivoli SA MP version level: *required-SA-MP-version*.**

Explanation: The pre-installation validation detected that it cannot upgrade Tivoli SA MP. If the Reliable Scalable Cluster Technology (RSCT) peer domain is offline or in maintenance mode it can cause the validation to fail.

User response: Manually upgrade the Tivoli SA MP product after the DB2 installation or update completes. See the Tivoli SA MP Base Component Installation and Configuration Guide for more information.

DBI1536E **The DB2 installer failed because of an unknown error, which occurred during the installation of IBM Tivoli System Automation for Multiplatforms (SA MP).**

Explanation: Tivoli SA MP is a required component of the DB2 pureScale environment. The Tivoli SA MP product was not installed or updated due to the unexpected error.

User response: Manually install the Tivoli SA MP product, then install the DB2 pureScale product. See the Tivoli SA MP Base Component Installation and Configuration Guide for more information.

DBI1537E **The offline fix pack update operation failed because one or more members or CFs in the following DB2 database manager instance are still active: *instance-name*. Host name: *host-name*.**

Explanation: A DB2 database manager instance cannot be updated while any member or cluster caching facility (CF) is active.

User response:

1. Identify active members or CFs in the instance by using the following command:
db2instance -list
2. Stop active members by using the following command:
db2stop member <member-id>
3. Stop active CFs by using the following command:
db2stop CF <CF-id>
4. Retry the fix pack update operation.

DBI1538E The command failed because the specified DB2 instance is not a DB2 pureScale instance but a parameter that was specified with the command is valid for only DB2 pureScale instances. **Command name:** *command-name*. **Specified parameter that is valid for only DB2 pureScale environments:** *parameter*.

Explanation: For many DB2 database manager commands, some parameter are valid in only certain conditions.

This message is returned when a command is called with a parameter specified that is valid for only DB2 pureScale instances.

User response: Respond to this message in one of the following ways:

- Re-issue the command without specifying parameters that are valid for only DB2 pureScale instances.
- Update the instance to a DB2 pureScale instance and then call the command again.

DBI1539E The online fix pack update operation failed because the code level to which the instance is being updated is lower than the current effective code level.

Explanation: The code level of a given DB2 product is comprised of product details such as: version, release, modification level, interim fix pack level, and special build tokens. The current effective code level (CECL) is the code level at which an instance is said to be operating.

The `curr_eff_code_level` configuration parameter shows the CECL for an instance.

Online fix pack updates of an instance towards a code level that is lower than the CECL of the instance is not supported.

User response: Update the instance by performing an offline update.

DBI1540E An invalid distribution media path is specified.

Explanation: An error was detected when attempting to locate the installable images on the distribution media. One of the following errors occurred:

- The file system associated with the given path is not properly mounted.
- An incorrect access permission has been assigned to the given path.

User response: Check the path and try the command again. Note that all file and directory names in Unix are case sensitive.

DBI1541E An attempt to open or read from the given status file failed.

Explanation: An error was detected when attempting to open or read from the status file. If the status file has not been modified manually, then one of the following errors occurred:

- An incorrect file path has been specified.
- An incorrect access permission has been assigned to the status file.
- The file system associated with the given path is not properly mounted.

User response: Check the path and try the command again. Note that all file and directory names in Unix are case sensitive. If the problem persists, contact IBM Support with the following information:

1. Message number
2. Problem description
3. Status file

DBI1544E The `db2icrt` command failed to create all the required data structures for the instance.

Explanation: The DB2 instance was created but it is missing some data structures. Some DB2 functionality might not be available or have unpredictable behavior.

User response: Drop and create the instance again. If the problem persists, contact IBM support.

DBI1546E The fix pack update operation failed because a fix pack update is not supported from the current version of the DB2 instance.

Explanation: You can update the version of DB2 instance to a higher level within the same release.

This message is returned when an attempt is made to perform a fix pack update of a DB2 instance that has a version of the DB2 database product than is not the same version specified in the fix pack update.

User response: Perform an instance upgrade to the version of the DB2 database product specified in the fix pack update.

DBI1547E The online fix pack update operation failed because an offline fix pack update is in progress.

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update.

This message is returned when an attempt is made to perform an online fix pack update while an offline fix pack update is in progress. An online fix pack update

cannot be performed until the offline fix pack update is committed or canceled.

User response: To respond to this message, perform the following steps:

1. Wait until the offline fix pack update that is in progress is committed or canceled.
2. Initiate the online fix pack update again.

DBI1548E The offline fix pack update operation failed because an online fix pack update is in progress.

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update.

This message is returned when an attempt is made to perform an offline fix pack update while an online update is in progress. An offline fix pack update cannot be performed until the online fix pack update is committed or canceled.

User response: To respond to this message, perform the following steps:

1. Wait until the online fix pack update that is in progress is committed or canceled.
2. Retry the offline fix pack update.

DBI1549E The fix pack update operation failed because the specified code level is different than the code level of the fix pack update performed on one or more DB2 members or cluster caching facilities (CFs).

Explanation: The code level of a given DB2 product is comprised of product details such as: version, release, modification level, and fix pack level. During a fix pack update operation, it is possible for some DB2 members or CFs in a DB2 cluster to be at one code level while other members or CFs in the same DB2 cluster are at a second code level.

Performing a fix pack update of a DB2 instance is only supported when the members and CFs in a DB2 cluster are at a maximum of two distinct code levels.

This message is returned when an attempt is made to perform a fix pack update to a new code level that is different than the code level of the fix pack update performed on several DB2 members and CFs in the same DB2 cluster.

User response: To perform the fix pack update to the specified code level, perform the following steps:

1. Update all members and CFs in the DB2 cluster to the same code level as the fix pack update that is in progress.
2. Retry the fix pack update to the specified code level.

DBI1550E There is not enough disk space on the file system.

Explanation: An error was encountered when trying to allocate more disk space on the file system.

User response: Retry the command after enough space has been freed up for the specified operation.

DBI1551E The online fix pack update operation failed because performing an online fix pack update from the current code level to the new code level is not supported.

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update.

This message is returned when an attempt is made to perform an online fix pack update between two code levels, and online fix pack updates is not supported between those two code levels.

User response: Perform an offline fix pack update.

DBI1552E The online fix pack update operation failed because the instance does not have the minimum current effective code level (CECL) required to install the specified code level.

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update.

This message is returned when an attempt is made to perform an online fix pack update from code level A to code level C, and the online fix pack update to code level C requires the instance to have a CECL that is higher than code level A.

User response: Respond to this message in one of the following ways:

- Perform an offline fix pack update to the specified code level.
- Perform the online fix pack update to the specified code level by following these steps:
 1. Determine the minimum committed code level from which you can perform an online fix pack update to the new code level by using the `installFixPack -show_level_info` command.
 2. Update and commit the instance to the minimum code level that you determine in the previous step.
 3. Perform the online fix pack update to the specified code level.

DBI1554E The online fix pack update operation failed because it involves a special build that does not support online fix pack updates to or from the current code level without a service password.

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update. This message is returned in two cases:

- An attempt is made to perform an online fix pack update to a special build, but performing an online fix pack update from the current code level to that special build is not supported unless a service password is specified.
- An attempt is made to perform an online fix pack update from a special build to another code level, but performing an online fix pack update from the special build to new code level is not supported unless a service password is specified.

User response: Respond to this message in one of the following ways:

- Update the instance by performing an offline fix pack update.
- Work with the assistance of IBM support to set the service password.

DBI1556E The fix pack update operation failed because of an error with an internal DB2 database manager configuration file. Reason code: *reason-code*.

Explanation: You can update the version of a DB2 database manager instance to a higher code level by performing an offline fix pack update. You can update the version of a DB2 pureScale instance to a higher level without having to bring your database offline by performing an online fix pack update.

This message is returned when a problem is encountered with an internal configuration file that is used for fix pack updates.

The reason code is useful for assisting only IBM support personnel troubleshoot the problem:

- | | |
|---|--|
| 1 | An invalid record was found. |
| 2 | A duplicate record was found. |
| 3 | A required record does not exist. |
| 4 | The configuration file content is invalid, or the fix pack update operation encountered an error reading the configuration file. |

5

The fix pack update operation cannot find the configuration file.

6

An error occurred while writing to the configuration file.

User response: For reason code 6, confirm that your user id has read and write access permission to the `sqllib_shared` directory as well as read and write access permission to the subdirectories of the `sqllib_shared` directory.

Perform the fix pack update operation again.

If this error persists, take note of this message identifier and reason code and contact IBM software support.

DBI1557E The install utility determined that the new architecture level is the same as the current effective code level (CECL). The install fix pack operation cannot be committed because the levels are the same.

Explanation: You can update the version of DB2 database to a higher code level by performing offline fix pack updates or online fix pack updates to keep databases online.

A fix pack update has several steps. The final step is to commit the fix pack update. Before committing the fix pack update, you can optionally verify that the fix pack update operation is ready for commit by using the `-check_commit` parameter with the `db2iupdt` command or the `installFixPack` command.

This message is returned when the pre-commit verification or commit of the fix pack update operation determines that the new level of code is the same as the CECL.

User response: No response is required.

DBI1558E The install utility determined that there are DB2 members or cluster caching facilities (CFs) in the current instance that are at two or more different code levels. The install fix pack operation cannot be committed if the members and CFs are not all at the same level.

Explanation: You can update the version of DB2 database to a higher code level by performing offline fix pack updates or online fix pack updates to keep databases online.

A fix pack update has several steps. The final step is to commit the fix pack update. Before committing the fix pack update, you can optionally verify that the fix pack update operation is ready for commit by using the `-check_commit` parameter with the `db2iupdt` command.

or the installFixPack command.

This message is returned when the pre-commit verification or commit of the online fix pack update operation determines that not all DB2 members or CFs in the DB2 pureScale instance are at the same level of code.

User response: Update all DB2 members and CFs in the DB2 pureScale instance to the same code level. Then perform the pre-commit verification or commit of the online fix pack update again.

DBI1559E **The pre-commit verification or commit of the offline fix pack update operation failed because the instance must be offline.**

Explanation: An offline fix pack update has several steps. The final step is to commit the fix pack update. Before committing the fix pack update, you can optionally verify that the fix pack update operation is ready for commit.

This message is returned when the pre-commit verification or commit of the offline fix pack update operation determines that the DB2 database manager instance is online. The instance must be offline.

User response: Respond to this message by performing the following steps:

- Stop the instance by issuing the db2stop command.
- Commit the offline fix pack update.

DBI1560E **The online fix pack update operation failed because an invalid service password was specified.**

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update.

A service password might be required to perform a online fix pack update with a special build or a fix pack with test fix that has been provided by IBM support.

This message is returned when an invalid service password was specified to perform an online fix pack update that requires a service password.

User response: Work with the assistance of IBM support to obtain the required service password and then perform the online fix pack update again.

DBI1562E **The commit of the fix pack update operation failed because the database manager configuration cannot be updated.**

Explanation: You can update the version of DB2 database to a higher code level within a release by

performing an offline fix pack update or an online fix pack update to keep the database online.

This message is returned when an online fix pack update fails because the database manager configuration cannot be updated.

User response: Perform the fix pack update operation again.

DBI1563W **The commit of the online fix pack update was successful. However, the database manager configuration for one or more DB2 members or cluster caching facilities (CFs) was not updated.**

Explanation: You can update the version of DB2 pureScale instance to a higher code level within a release without having to bring your database offline by using an online fix pack update.

An online fix pack update has several steps. The final step is to commit the online fix pack update. In this final step, the database manager configuration for each DB2 member and CFs in the DB2 cluster is updated.

This message is returned when the online fix pack update operation encounters errors while updating the database manager configuration for one or more DB2 members or CFs.

User response: Respond to this message by performing the following actions:

1. Determine which members or CFs need to be restarted by issuing the following command:
db2pd -rustatus
2. Restart the affected members and CFs by issuing the db2stop command followed by the db2start command.

DBI1564E **The fix pack update operation failed at task name *task-name* with the following error message: *message-id*. Log file names: *log-file-name-list*.**

Explanation: You can update the version of DB2 database to a higher code level by performing offline fix pack updates or online fix pack updates to keep databases online.

A fix pack update performs several tasks. This message is returned when an error is encountered in one of these tasks.

User response: Respond to this message by performing the following troubleshooting steps:

- Review the diagnostic information in the given log file or log files.
- Review the explanation and user response for the given message identifier.
- Resolve the cause of the underlying error.
- Complete the fix pack update operation.

DBI1565E **The command failed because a parameter was specified that requires a service password, but the service password is not set.**

Explanation: You can update the version of DB2 database to a higher code level within a release without having to bring your database offline by using an online fix pack update.

This message is returned when an attempt is made to perform an online fix pack update that requires a service password, but the service password environment variable has not been configured.

User response: Work with the assistance of IBM support, to set the service password environment variable.

DBI1567E **The installFixPack command failed because an empty path was specified for the -p parameter.**

Explanation: If you are performing a fix pack update on a DB2 pureScale instance (dsf type instance), you must specify a path to install the fix pack that is different than the existing installation. This path must be the same on all DB2 members and cluster catching facilities (CFs).

This message is returned when an empty path is specified with the -p parameter because this parameter is mandatory for DB2 pureScale instances.

User response: Re-issue the installFixPack command specifying a valid new installation path with the -p parameter.

DBI1568E **The installFixPack command failed because the path that was specified with the -p parameter is the same as the current DB2 database product installation path, but the -f parameter was not specified.**

Explanation: If you are performing a fix pack update on a DB2 pureScale instance (dsf type instance), you must specify a path to install the fix pack that is different than the existing installation. This path must be the same on all DB2 members and cluster catching facilities (CFs).

To replace the existing installation, you must also specify the -f parameter. However, if you specify the same path as the existing installation, you cannot cancel the fix pack update.

This message is returned when the path that is specified with the -p parameter is the same path as the current DB2 product install path, but the -f parameter was not specified with the installFixPack command.

User response: Respond to this message in one of the following ways:

- Re-issue the installFixPack specifying an install path with the -p parameter that is not the same as the current DB2 product install path.
- >Re-issue the installFixPack command specifying the same install path with the -p parameter and specify the -f parameter.

DBI1569E **The instance creation or update operation failed because too many interconnect netnames were specified for a DB2 member or CF.**

Explanation: You can create a new DB2 database manager instance by using the db2icrt command. You can update an existing instance by using the db2iupdt command. For both the db2icrt command and the db2iupdt command, you can specify a list of interconnect netnames for a DB2 member by using the -mnet parameter, and you can specify a list of netnames for a cluster caching facility (CF) by using the -cfnet parameter. The maximum number of netnames that can be specified for a DB2 member or CF is four.

This message is returned when an attempt is made to specify more than four netnames for a DB2 member or CF while using the db2icrt command or the db2iupdt command.

User response: Call the command again, specifying no more than four netnames for any DB2 members or CFs.

DBI1570I **Usage:**
db2olset InstName

Explanation: An incorrect argument was entered for the db2olset command. Valid arguments for this command are:

-h|-? display the usage information

-d turn debug mode on.

InstName
is the name of the instance you wish to set up to use with the OLAP Starter Kit

User response: Enter the command again as follows:
db2olset InstName

DBI1571E **The command failed because the instance is a DB2 pureScale instance but a parameter was specified with the command that is not valid in DB2 pureScale environments. Command name: *command-name*. Specified parameter: *parameter*.**

Explanation: For many DB2 database manager utility commands, some parameters are valid in only certain conditions. This message is returned when a command is called in a DB2 pureScale environment with a parameter that is not valid for DB2 pureScale environments.

User response: Call the command again without specifying parameters that are not valid for DB2 pureScale environments.

DBI1572E The fix pack update operation failed because the DB2 member is involved in a restart light process.

Explanation: When a failed DB2 member cannot be restarted on its original host, or home host, DB2 cluster services restarts that member on one of the other available members in a DB2 pureScale instance. This process is known as restart light.

This message is returned when an attempt is made to perform a fix pack update on a DB2 member that is involved in a restart light process. Fix pack updates are not supported on DB2 members that are involved in a restart light process.

User response: Monitor the DB2 member and retry the fix pack update when the restart light process is resolved.

DBI1573W The DB2 installer did not update the RSCT software on this member or CF, because the currently installed version of RSCT is supported with the new DB2 fix pack level.

Explanation: DB2 cluster services includes three major components, including a cluster manager, Tivoli SA MP, which includes RSCT.

When you update your DB2 database product to a new DB2 fix pack level on a DB2 member or cluster caching facility (CF), the DB2 install utility will automatically update the Reliable Scalable Cluster Technology (RSCT) software on that member or CF, if needed. If the version of RSCT that is already installed on that member or CF is supported by the new fix pack level, the DB2 install utility will not update the RSCT software on that member or CF.

This message is returned when a fix pack is being updated on a DB2 member or CF, and the version of RSCT that is currently installed on the member or CF is supported with the new fix pack level.

User response: No action is required.

DBI1574W The fix pack update operation will update the version of Reliable Scalable Cluster Technology (RSCT) installed to a higher version on the local host.

Explanation: In a DB2 pureScale instance, the DB2 cluster services provide essential infrastructure for the instance to be highly available and to provide automatic failover and restart as soon as the instance has been created. RSCT is one of the three major components included by DB2 cluster services.

This message is returned when an attempt is made to

perform a fix pack update on DB2 member or CF that has an earlier version of RSCT installed. The fix pack update operation will update to a higher version of RSCT.

User response: If the fix pack update is cancelled, manual installation of the earlier version of RSCT is required.

DBI1575E The fix pack update operation failed because the instance on the DB2 member or cluster caching facility (CF) that is being updated cannot be stopped.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online. During the fix pack update, the instance is stopped to perform the update.

This message is returned when an attempt to stop the instance on the member or CF that is being updated fails.

User response: Stop the instance on the host by issuing the following command:

```
db2stop INSTANCE ON host-name FORCE
```

Once the instance is stopped, perform one of the following actions:

- Perform again the fix pack update.
- Cancel the fix pack update on the instance as follows:

```
db2start INSTANCE host-name
db2start member member-id or
db2start CF CF-identifier
```

DBI1576E The fix pack update operation failed because the DB2 member that is being updated cannot be stopped.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online. During the fix pack update, the member is stopped to perform the update.

This message is returned when the attempt to stop the DB2 member that is being updated fails.

User response: Stop the DB2 member by issuing the following command:

```
db2stop member member-id
```

Once the member is stopped, perform one of the following actions:

- Perform again the fix pack update.
- Cancel the fix pack update on the member by restarting the member as follows:

```
db2start member member-id
```

DBI1577E **The fix pack update operation failed because the cluster caching facility (CF) that is being updated cannot be stopped.**

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online. During the fix pack update, the CF is stopped to perform the update.

This message is returned when the attempt to stop the CF that is being updated fails.

User response: Stop CF by issuing the following command:

```
db2stop CF CF-identifier
```

Once the CF is stopped, perform one of the following actions:

- Perform again the fix pack update.
- Cancel the fix pack update on the CF by restarting the CF as follows:

```
db2start CF CF-identifier
```

DBI1578E **The fix pack update operation failed because the cluster manager cannot be put in maintenance mode to be updated.**

Explanation: To ensure that DB2 members or cluster caching facilities (CFs) are not restarted on the host when making updates to the cluster manager, you can put the cluster manager into maintenance mode.

This message is returned when an attempt is made to perform a fix pack update but the cluster manager cannot be put in maintenance mode to be updated.

User response: To respond to this message, manually put the cluster manager in maintenance mode. Then, to continue with the fix pack update, perform it again.

DBI1579E **The fix pack update operation failed because the IBM General Parallel File System (GPFS) cannot be put in maintenance mode to perform the update.**

Explanation: To ensure that DB2 members or cluster caching facilities (CFs) are not restarted on the host when making updates to GPFS, you can put GPFS into maintenance mode.

This message is returned when an attempt is made to perform a fix pack update but the update operation cannot put GPFS in maintenance mode to update it.

User response: To respond to this message, manually put GPFS in maintenance mode. Then, to continue with the fix pack update, perform it again.

DBI1580E **The fix pack update operation failed because the cluster manager cannot exit maintenance mode.**

Explanation: After making updates to the cluster manager, you can exit maintenance mode to remove the host from maintenance mode and put the domain online.

This message is returned when an attempt is made to perform a fix pack update but the cluster manager cannot exit maintenance mode.

User response: To respond to this message, manually exit the cluster manager from maintenance mode. Then, to continue with the fix pack update, perform it again.

DBI1581E **The fix pack update operation failed because IBM General Parallel File System (GPFS) cannot exit maintenance mode.**

Explanation: After making updates to GPFS, you can exit maintenance mode to remove a host from maintenance mode and put the domain online.

This message is returned when an attempt is made to perform a fix pack update but GPFS cannot exit maintenance mode.

User response: To respond to this message, manually exit GPFS from maintenance mode. Then, to continue with the fix pack update, perform it again.

DBI1582E **The online fix pack update operation failed because the instance being updated is not a DB2 pureScale instance.**

Explanation: You can update the version of DB2 pureScale instance to a higher code level within a release without having to bring your database offline by performing an online fix pack update.

This message is returned when an online fix pack update is performed on an instance with a type different than a DB2 pureScale instance.

User response: Perform an offline fix pack update on the specified instance.

DBI1583E **The installFixPack command failed because the list of instances specified includes DB2 pureScale instances and other types that are not pureScale instances.**

Explanation: You can specify a list of instances to update with the -I parameter of the installFixPack command. All the instances in the list have to be either pureScale instances or other types that are not pureScale instances. A list that includes pureScale instances and other types that are not pureScale instances is not supported.

This message is returned when -I parameter of the installFixPack command was specified with a list of instances that includes pureScale instances and other types that are not pureScale instances.

User response: To respond to this message perform one of the following actions or both:

- Issue the installFixPack command with one pureScale instance name.
- Issue the installFixPack command with a list of other type of instances that are not pureScale instances.

DBI1584E The installFixPack command failed because the list of instances specified have different installation paths.

Explanation: You can specify a list of instances to update with the -I parameter of the installFixPack command. All the instances in the list must have the same installation path. A list of instances that have different installation paths is not supported.

This message is returned when -I parameter of the installFixPack command was specified with a list of instances that have different installation paths.

User response: To respond to this message, issue the installFixPack command with a list instances that have the same installation path.

DBI1585E The commit of the fix pack update failed because the commit of the cluster manager update fails.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online.

A fix pack update has several steps. The commit is the final step that completes the fix pack update. The cluster manager update is committed during the commit.

This message is returned when the commit of the fix pack update cannot commit the cluster manager update.

User response: To respond to this message, manually commit the cluster manager update. Then, perform the commit for the fix pack update again.

DBI1586E The commit of the fix pack update failed because the commit of the IBM General Parallel File System (GPFS) update fails.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online.

An fix pack update has several steps. The commit is

the final step that completes the fix pack update. The GPFS update is committed during the commit.

This message is returned when the commit of the fix pack update cannot commit the GPFS update.

User response: To respond to this message, manually commit the GPFS update. Then, perform the commit of the fix pack update again.

DBI1587E The commit of the fix pack update failed because the commit of the current effective code level (CECL) or current effective architectural level (CEAL) update fails.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online.

A fix pack update has several steps. The commit is the final step that completes the fix pack update. The CECL and CEAL update is committed during the commit.

This message is returned when the commit of the fix pack update cannot commit the CECL and CEAL update.

User response: To respond to this message, manually commit the CECL and CEAL update. Then, perform the commit of the fix pack update again.

DBI1588E The online fix pack update operation failed because the instance on the DB2 member or cluster caching facility (CF) that is being updated cannot be restarted.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using an online fix pack update to keep databases online. During the online fix pack update, the instance is restarted to perform the update.

This message is returned when an attempt to restart the instance on the DB2 member or CF that is being updated fails.

User response: Start the instance on the DB2 member or CF by issuing the following command:
db2start INSTANCE ON host-name

DBI1589E The online fix pack update operation failed because the DB2 member or cluster caching facility (CF) that is being updated cannot be restarted.

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using an online fix pack update to keep databases online. During the online fix pack update, the DB2 member or CF is restarted to perform the update.

DBI1590E • DBI1593E

This message is returned when an attempt to restart the DB2 member or CF that is being updated fails.

User response: Manually restart the DB2 member or CF.

- To restart a DB2 member, issue the following command:
db2start member member-id
- To restart a CF, issue the following command:
db2start CF CF-id

DBI1590E **The fix pack update operation failed because the DB2 database software cannot be updated. Details about the error were written to the log file. Log file name: log-file-name.**

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online. During the fix pack update, the DB2 database software components is updated.

This message is returned when an attempt to update the DB2 software fails.

User response: To respond to this message, perform one of the following actions:

- Complete the fix pack update by performing the following steps:
 - Check error message in the log file and perform the corrective action.
 - Re-issue the fix pack update command.
- Cancel the fix pack update by restarting the instance.

DBI1591E **The fix pack update operation failed because either the SA MP software or the GPFS software cannot be updated. Details about the error were written to the log file. Log file name: log-file-name.**

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online. During the fix pack update, the following software components are updated:

- IBM Tivoli System Automation for Multiplatforms (SA MP)
- IBM General Parallel File System (GPFS)

This message is returned when an attempt to update the software for SA MP or GPFS fails.

User response: To respond to this message, perform one of the following actions:

- Complete the fix pack update by performing the following steps:
 - Check error message in the log file and perform the corrective action.

- Perform again the fix pack update operation.
- Cancel the fixpack update by issuing the installFixPack command with the -f level parameter from the install directory of the previous fix pack release.

DBI1592E **The fix pack update operation failed to update the instance configuration. Details about the error were written to the log file. Log file name: log-file-name.**

Explanation: You can update the version of DB2 pureScale instance to a higher code level by using offline fix pack updates or online fix pack updates to keep databases online. During the fix pack update, the existing instance configuration is updated.

This message is returned when attempt update the instance configuration fails.

User response: To complete the online fix pack update, perform the following actions:

- Check the log file and perform a corrective action.
- Start the instance on the local host by issuing the following command:
db2start INSTANCE ON host-name
- Start the DB2 member or CF by issuing one of the following commands:
db2start member member-id

or
db2start CF CF-id

To complete the offline fix pack update, perform the following actions:

- Check the log file and perform a corrective action.
- Perform again the fix pack update operation.

To cancel the fix pack update, issue the installFixPack command with -f level parameter from the install directory of the previous fix pack release.

DBI1593E **The DB2 command failed because the changes to the instance configuration cannot be committed.**

Explanation: During an instance creation or upgrade changes to the instance configuration are committed to complete the command.

This message is returned when committing the instance configuration changes during an instance creation or upgrade fails.

User response: Manually commit the instance configuration changes by issuing the following command:

db2iupdt -commit_new_level instance_name

DBI1594E The `setup_db2locssh` command failed because an invalid parameter was specified.

Explanation: You can install and configure a DB2 pureScale instance without having to enable passwordless SSH by running the `setup_db2locssh` command before performing the install or configuration operation.

This message is returned when the `setup_db2locssh` command is invoked with invalid or missing parameters.

User response: Enter the command again specifying only valid parameters.

DBI1595E The `setup_db2locssh` command failed because of an internal, runtime error.
Reason code: *reason-code*.

Explanation: You can install and configure a DB2 pureScale instance without having to enable passwordless SSH by running the `setup_db2locssh` command before performing the install or configuration operation.

This message is returned when the `setup_db2locssh` command fails with an unexpected error.

The reason code in the runtime token is an error number that can assist IBM support personnel to troubleshoot the cause of the error. The runtime token might sometimes be empty.

User response: Contact IBM support for assistance.

DBI1596W The user ID that was specified with the `-i` parameter was ignored because a non-root user ID already exists on the system. The command proceeded using the existing non-root user ID.

Explanation: In DB2 pureScale environments, the non-root user ID (sometimes referred to as "db2sshid") is used to establish a secure shell (SSH) network protocol between host machines.

You perform install, upgrade, or configuration tasks by using commands such as `db2icrt` or `db2iupdt`. When you call one of these commands you can configure a new non-root user ID by specifying the proposed user ID with the `-i` parameter.

This message is returned when a command such as `db2icrt` or `db2iupdt` is called and the `-i` parameter is specified with a proposed user ID, but the utility detects that a non-root user ID already exists on the system. Because there may exist only one non-root user ID on a given system at a time, the utility uses the existing non-root user ID instead of configuring a new one.

User response: No response is required.

DBI1597I Communication between host machines was performed using passwordless SSH because the ability to communicate without using passwordless SSH has not been configured.

Explanation: You can install and configure a DB2 pureScale instance without having to enable passwordless SSH by running the `setup_db2locssh` command before performing the install or configuration operation. Running the `setup_db2locssh` command enables DB2 database manager processes to communicate between host machines without using passwordless SSH.

Utilities such as `db2icrt` and `db2iupdt` will communicate without using passwordless SSH, if the capability to do so has been configured. This message is printed when a utility must use passwordless SSH because the capability to communicate without using passwordless SSH has not been configured.

User response: No response is required.

DBI1631E An error was encountered while creating a sample database.

Explanation: One of the following situations has occurred:

- There is not enough disk space on the file system.
- The DB2 product has not been installed properly, or some of the component files are missing.

User response: If this is not a disk space problem, try to re-install the product. If the problem persists, turn trace on and retry the steps to create a database through the CLP. Then, save the trace information to a file and contact IBM Support with the following information:

- Message number
- Problem description
- Trace file

DBI1632E An error was encountered while configuring to auto-start instance.

Explanation: One of the following situations has occurred:

- The feature to "auto-start" the instance is already enabled.
- An error was encountered while adding an entry to `/etc/inittab` file. This could be a file permission problem.
- The `db2uit` tool is missing.

User response: If any of these three situations apply, then correct the problem and retry the command. If the problem persists, contact IBM Support with the following information:

- Message number

- Problem description

DBI1633E **An error was encountered while removing the configuration to auto-start instance.**

Explanation: One of the following situations has occurred:

- The feature to "auto-start" the instance is not enabled.
- An error was encountered while removing an entry to "/etc/inittab" file. This could be a file permission problem.
- The db2uit tool is missing.

User response: If any of those three situations apply, then correct the problem and retry the command. If the problem persists, contact IBM Support with the following information:

- Message number
- Problem description

DBI1634W **Unable to update /etc/rc.db2v08 to auto-mount dlfs file system.**

Explanation: An attempt to update /etc/rc.db2v08 to enable the auto-mounting of the dlfs file system on system reboot has failed.

User response: Manually edit the file /etc/rc.db2v08, and add the following lines to it:

```
if [ -x /etc/rc.dlfs ]; then
    /etc/rc.dlfs
fi
```

DBI1635E **An error was encountered while adding to the Administration Server group list.**

Explanation: An error was detected when attempting to add a DB2 instance to the Administration Server group list.

User response: If you are running this command on an NIS client, try to append the group name of the DB2 instance to the secondary group of the Administration Server on your NIS server.

DBI1637W **By clicking OK, your previous settings for the instance will be lost.**

User response: Click OK if you wish to proceed. Otherwise click Cancel.

DBI1639E **Cannot set up a new instance.**

Explanation: The given instance name does not belong to an existing instance.

User response:

- Create a new instance.

- If you are running the interactive installer, from the instance creation panel choose to create a new instance.
- If you are running the response file installer, provide additional information for the user.
- Enter the name of an existing instance.

DBI1640W **The specified instance is not a server instance.**

Explanation: Some of the DB2 components can only be configured for DB2 server instances.

User response:

- If the specified instance is a client instance and DB2 server product is installed, run db2iupdt to update the client instance to a server instance.
- If the specified instance is a Data Link Administration Server, exit the installer and set up the instance from the command line, if you wish to do so.

DBI1651E **The specified UID is invalid.**

Explanation: An invalid UID has been entered. One of the following situations has occurred:

- The given UID is already assigned to an existing user on the system.
- The given UID either is too long or consists of invalid characters.
- The given UID less than or equal to 100.

User response: Retry with a different user ID.

DBI1652E **The specified user name is invalid.**

Explanation: An invalid user name has been entered. One of the following situations has occurred:

- The given user name already exists on the system.
- The given user name already exists on the system but under a different group name.
- The given user name either is too long or consists of invalid characters.

User response: Retry with a different user name.

DBI1653E **The specified group is invalid.**

Explanation: An invalid group ID or group name has been entered. One of the following situations has occurred:

- The given group ID already exists but under a different group name on the system, or the given group name already exists but under a different group ID.
- The given group ID or group name either is too long or consists of invalid characters.
- The given group ID is less than or equal to 100.

User response: Retry with a different group ID or group name.

DBI1654E The specified directory is invalid.

Explanation: An invalid directory has been entered. One of the following situations has occurred:

- The given user already exists but has a different home directory specified on the system.
- The given directory is either too long or consists of invalid characters.

User response: Retry with the user's home directory specified on the system or a different directory.

DBI1655E The specified password is invalid.

Explanation: An invalid password has been entered. One of the following situations has occurred:

- The given password has a length of zero.
- The given password does not match with the verify password.
- The given password contains invalid characters.

User response: Retry with a different password.

DBI1657E Instance name already in use.

Explanation: An instance with the same name already exists on your system.

User response: Verify with the list of instances as reported by the "db2ilist" command. Retry with a different user name.

DBI1701E One or more of the TCP/IP parameters specified is invalid.

Explanation: One of the following situations has occurred:

- The length of service name exceeds 14 characters.
- The port number specified is out of range. This number must be between 1024 and 65535.

User response: Correct the problem and try again.

DBI1702E The specified service name or port number conflicts with existing values in the TCP/IP services file.

Explanation: The service name or port number conflicts with existing values in the services file. The service name might already be used with a different port number, or the port number might already be used with a different service name.

User response: Specify a service name and port number that does not conflict with existing entries in the services file.

DBI1703E No valid service name or port number found.

Explanation: No valid service name or port number for this instance can be found in the TCP/IP services file. This information is either missing or is invalid.

If you installed "DB2 Extended Server Edition" product, a service name must be specified for the instance before using this command.

User response: Specify valid service name and port number for this instance and enter this command again.

DBI1704W TCP/IP is not running on your system.

Explanation: TCP/IP is not running on your system. TCP/IP services file will not be updated.

DBI1705E The port number specified *port-number* is being used by another application.

Explanation: The port number specified cannot be used because it is being used by another application that is currently running.

User response: Specify another port number that is not in use. The port number must be in the range 1024 to 65535.

DBI1709E An error was encountered when updating the TCP/IP services file.

Explanation: An attempt to add a service name and port number to the TCP/IP services file failed.

User response: If the user wishes to use the service name instead of the port number in the node directory entry, the node must be manually uncatalogued and then recatalogued using the service name. The services file must also be updated manually. If Network Information Services (NIS) is being used, the local services file may have been updated but the NIS server must be updated manually. In this case, the node was also catalogued using the port number.

DBI1711E One or more of the IPX/SPX parameters specified is invalid.

Explanation: One of the following situations has occurred:

- the file server, objectname, or ipx_socket parameter is null.
- the file server parameter is set to * but objectname is not *.
- the objectname or ipx_socket value specified is not unique.
- the ipx_socket value specified is not in the valid range.

User response: Correct the problem and try again.

DBI1715E **NetWare User ID or NetWare Password is invalid.**

Explanation: The NetWare User ID and NetWare Password that you specify must already exist and is used to register the DB2 Object Name at a NetWare file server and must have supervisory or equivalent authority.

User response: Correct the problem and try again.

DBI1720E **The specified User ID or Group ID cannot be added since NIS is running.**

Explanation: NIS is running on your system. A new User ID or Group ID cannot be created while you have NIS running on your system.

User response: Use a previously created User ID or Group ID and try again.

DBI1722E **The specified service cannot be added to the /etc/services file because NIS is running.**

Explanation: NIS is running on your system. A new service cannot be added to the /etc/services file while you have NIS running on your system.

User response: Use a previously created service name and port number from the /etc/services file and try again.

DBI1725W **Cannot create Agent Daemon and Logger services for Data Warehouse Agent.**

Explanation: TCP/IP is not running on the system. The install will proceed but you need to create Agent Daemon and Logger services manually.

User response: Activate TCP/IP on the system and run the installer.

DBI1740W **Security Risk.**

Explanation: Using the default instance user ID here allows this instance to have full access permission to other instances on your system. It is highly recommended to use a different user ID for security reasons.

User response: To avoid the possibility of a security risk, use a different user ID.

DBI1741W **Protocol not detected.**

Explanation: You have selected a protocol that is not detected. DB2 Installer cannot update all of the protocol's required settings. However, you can provide values for this protocol's settings.

User response: No actions required.

DBI1744W **The DB2 Data Links Manager Administrator is not created.**

Explanation: You have not created the DB2 Data Links Manager Administrator. Without the DB2 Data Links Manager Administrator, you cannot administer the DB2 Data Links Manager.

User response: No actions required.

DBI1745W **The Administration Server is not created.**

Explanation: You have not created a Administration Server. Without the Administration Server, you cannot provide services to support client tools that automate the configuration of connections to DB2 databases, and cannot administer DB2 from your server system or a remote client.

User response: No actions required.

DBI1746W **DB2 Instance is not created.**

Explanation: You have not created a DB2 Instance. A DB2 Instance is an environment where you store data and run applications.

User response: No actions required.

DBI1747W **The specified directory already exists.**

Explanation: The directory specified in the Home Directory field already exists. If you choose to use this directory you may run into permission problems.

User response: Choose a different directory if this may be a problem.

DBI1750W **Security Risk.**

Explanation: Using the default instance user ID here allows this instance to have full access permission to other instances on your system. It is highly recommended to use a different user ID for security reasons.

User response:

- Press OK to continue, or
 - Press Cancel to enter a different user ID.
-

DBI1751W **Protocol not detected.**

Explanation: You have selected a protocol that is not detected. DB2 Installer cannot update all of the protocol's required settings. However, you can provide values for this protocol's settings.

User response:

- Press OK to provide protocol settings, or
 - Press Cancel to ignore.
-

DBI1753W The Administration Server is not created.

Explanation: You have not created an Administration Server. Without the DB2 Administration Server, Data Warehouse Agent will not be fully set up.

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1754W The DB2 Data Links Manager Administrator is not created.

Explanation: You have not created the DB2 Data Links Manager Administrator. Without the DB2 Data Links Manager Administrator, you cannot administer the DB2 Data Links Manager.

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1755W The Administration Server is not created.

Explanation: You have not created a Administration Server. Without the Administration Server, you cannot provide services to support client tools that automate the configuration of connections to DB2 databases, and cannot administer DB2 from your server system or a remote client.

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1756W DB2 Instance is not created.

Explanation: You have not created a DB2 Instance. A DB2 Instance is an environment where you store data and run applications.

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1757W The specified directory already exists.

Explanation: The directory specified in the Home Directory field already exists. If you choose to use this directory you may run into permission problems.

User response:

- Press OK to use this directory, or
- Press Cancel to go back.

DBI1758W DB2 Instance or Administration Server detected.

Explanation: You have selected to removed the entire DB2 product. However, a DB2 Instance or Administration Server has been detected on your system. If you remove the DB2 product without dropping these instances first, your DB2 Instances may not be able to function properly later on.

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1759W DB2 Instance or Administration Server configurations detected.

Explanation: You do not have the Administration Server selected. This will cause all of your Administration Server and DB2 Instance configurations to be discarded.

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1760E To set up *pkg-name*, you will need to install the package.

Explanation: You have chosen to configure a DB2 component for an instance. You will not be able to do so unless the package is installed.

User response: Install the package and try again.

DBI1761W Software prerequisites violated on file set or package *pkg-name*.

Explanation: An error occurred when verifying software prerequisites on the indicated file set or package. Continuing may result in the product not functioning properly. Do you want to continue without installing its pre-requisites?

User response:

- Press OK to continue without installing its prerequisites, or
- Press Cancel to abort installation.

DBI1762W Software dependencies violated on file set or package *pkg-name*.

Explanation: An error occurred when verifying software dependencies on the indicated file set or package. Continuing may result in removing all of its software dependencies. Do you want to remove its software dependencies?

User response:

- Press OK to continue, or
- Press Cancel to go back.

DBI1763I English HTML documentation files will also be installed.

Explanation: The English HTML documentation files are required when you select any non-English HTML documentation files for installation. This is to ensure that users will be able to search the entire DB2 documentation library.

DBI1765W A non-existing user ID, *InstName*, has been detected from the registry.

Explanation: The registry is containing a non-existing user ID. This can be caused by just removing the instance user ID without first dropping the instance. This will not affect the normal operation of DB2, but it is recommended to remove this obsolete user ID from the instance list in the registry.

User response: To remove this user ID from the registry with the command: `db2iset -d "<InstName>"`

DBI1766W Cannot change the secondary group list of user-ID. Error code: *err-code*. Required group id: *group-ID*.

Explanation: An error code was returned when attempting to change the secondary group list of the given user ID. One of the following situations has occurred:

- NIS is running.
- One or more processes are currently being executed under the given user ID.

User response: Add the given group ID to the secondary group list of the user ID so that the Administration Server can function properly.

- If there happens to be any process run under the given user ID, terminate all of these processes and follow the instructions in this message to set up the secondary group list of this user ID.
- If you are running this command on an NIS client, try the instructions to set up the secondary group list of this user ID on your NIS server.

DBI1767W DB2 registry may contain corrupted information.

Explanation: The registry may contain invalid instance information. This is probably due to the fact that an instance or the Administration Server is not deleted properly before removing the DB2 product. If you are re-installing the product and want to keep the existing instance information in the registry, press Cancel. Otherwise, press OK to re-create the registry. Do you want to re-create the registry?

User response:

- Press OK to re-create the registry, or
- Press Cancel to keep the current registry information.

DBI1768W Unable to add *db2profile* to the .profile file or *db2cshrc* to the .login file of user *inst-name*.

Explanation: DB2 has failed to modify the .profile file or the .login file of this user. These files may not exist, or you may not have write permission to them. Without the modification you need to set your environment manually each time when you log in as this user to use DB2.

User response: Respond in one of the following ways:

- Add the profile to the .profile file
- Add the location of db2cshrc to the .profile file

DBI1769W Unable to remove DB2 profile entry from the .profile or the .login file of *inst-name*.

Explanation: DB2 has failed to comment out the DB2 environment setting line in the .profile file or the .login file of this user. The next time you log in as this user, you may get a message saying that the file db2profile (or db2cshrc) cannot be found. There can be different reasons:

- The two files (.profile and .login) may not exist.
- You may not have write permission on these files.

User response: If you receive a message concerning missing db2profile or db2cshrc, locate your .profile or .login file and comment out the following:

```
. sqllib/db2profile    from your .profile file
source sqllib/db2cshrc from your .login file
```

DBI1770E Unable to update the file /etc/vfs.

Explanation: The following DB2 Data Links Manager specific entry is required in the /etc/vfs file:

```
dlfs dlfs_num /sbin/helpers/dlfs_mnthelp none
```

where *dlfs_num* is to be set to 12, or some number between 8 and 15 if 12 is already assigned to another fs. The installation process has failed to insert this line into the /etc/vfs file. There may be different reasons:

- You do not have write permission to this file.
- The range 8-15 is all taken up, no number can be assigned to *dlfs*.

User response: Manually edit the /etc/vfs file, and add the entry for *dlfs* into this file.

DBI1771I File /etc/vfs has been copied to /tmp.db2.etcvfs.backup.

Explanation: For DB2 Data Links Manager to work, the following line has to be inserted into the /etc/vfs file:

```
dlfs dlfs_num /sbin/helpers/dlfs_mnthelp none
```

A backup copy of the original /etc/vfs file has been placed in /tmp/.db2.etcvfs.backup.

DBI1775W Modified an existing /etc/rc.dlfs file.

Explanation: DB2 has modified an existing /etc/rc.dlfs file with the new dlfs file system mount entry.

User response: Examine the file /etc/rc.dlfs to ensure that all entries are correct.

DBI1780W DB2 Data Links Manager is not supported on the version of your current operating system.

Explanation: The version of your current operating system is not supported to run DB2 Data Links Manager. However, you can still install and configure the product at your own risk. The following is a list of supported operating systems for DB2 Data Links Manager:

- "<OSlist>"

DBI1782E No language specified.

Explanation: You must specify at least one language from those available.

User response: Issue the command with no parameters for a complete list.

DBI1783E No topics specified.

Explanation: You must specify at least one topic from those available.

User response: Issue the command with no parameters for a complete list.

DBI1784E Language not available.

Explanation: You specified a language that is not available for installation.

User response: Issue the command with no parameters for a complete list.

DBI1785E Topic Package not available.

Explanation: You specified a documentation package that is not available for installation.

User response: Issue the command with no parameters for a complete list.

DBI1790E ODSSI utilities have not been installed, unable to create DB2 search index.

Explanation: ODSSI utilities have not been installed. These utilities (config_search, config_view and config_help) are required to create the search index for the DB2 online documentation.

User response: Install the ODSSI utilities, then run the following commands:

/opt/IBM/db2/V8.x/doc/db2insthtml lang_locale,
where x is 1 or FPn,
where n is the FixPak number

to create the search index for DB2 documentation in the language locale lang_locale. Note that whenever you install documentation in any language, English documentation is installed as well. So you should also run the following command to index the English files:
/opt/IBM/db2/V8.x/doc/db2insthtml en_US,
where x is 1 or FPn,
where n is the FixPak number

DBI1791W Possible error in creating/removing DB2 search index for html file.

Explanation: DB2 runs the system utilities config_view and config_help to create or remove DB2 search index to be used with SCOHELP. However, error messages may be returned if there exists other non-DB2 indexes on the system that have not been created/removed properly.

User response:

- If you were creating the DB2 index, you can start up SCOHELP and check to see if the DB2 entry has been added to the topics panel. Test searching for a string like "TCP/IP" and see if you get any results. If you are able to do the search, DB2 search index has been created successfully.
- If you were removing the DB2 index, you can start up SCOHELP and verify that the DB2 entry has been removed from the topics panel. If the entry is gone, then the DB2 index has been successfully removed.
- If you cannot verify that the DB2 index has been successfully created or removed, contact IBM support.

DBI1792I Creating search index for DB2 documentation files, please wait ...

Explanation: Search index for DB2 documentation files are being created. Depending on the language(s) you have selected, this may take some time.

DBI1793W The user ID userID is unknown.

Explanation: The specified user is unknown or cannot be found on the current system.

User response: Please verify that the user exists on this system and try again.

DBI1794E Attempt to put too many icons into one group.

Explanation: An attempt has been made to place more icons into the required folder than are allowed by the desktop manager.

User response: Please remove any unnecessary icons from your desktop and retry the command.

DBI1795E **There is insufficient free memory to create all the required icons.**

Explanation: The icon creation utility ran out of memory during icon generation.

User response: Close any unnecessary programs and try again.

DBI1796W **Attempt to create icons in a directory without proper permissions.**

Explanation: This user ID does not have the permissions required to create icons for the requested user.

User response: Log on as a user with root authority or with the permissions required to create files in the user directory and retry the command.

DBI1797I **Icons successfully created.**

DBI1900N **The Microsoft Cluster Server (MSCS) support is not accessible.**

Explanation: DB2 is unable to access the Microsoft Cluster Server (MSCS) cluster.

User response: Ensure the Microsoft Cluster Server (MSCS) support has been correctly installed and the Cluster Service has been started.

DBI1901N **An error occurred when attempting to update the user rights for the DB2 or Microsoft Cluster Server (MSCS) cluster service.**

Explanation: An error was encountered when attempting to grant the necessary user rights to the account specified for use with the DB2 service.

User response: Ensure a valid account has been specified and that you have the necessary privilege to grant user rights.

DBI1902N **An error occurred when attempting to update the DB2 or Microsoft Cluster Server (MSCS) cluster service configuration.**

Explanation: DB2 encountered an error when attempting to update the DB2 or Microsoft Cluster Server (MSCS) cluster service configuration.

User response: Ensure the Service Database has not been locked and is accessible. Also ensure the DB2 or Microsoft Cluster Server (MSCS) cluster service has been created.

DBI1903N **An error occurred when accessing the Cluster registry.**

Explanation: DB2 encountered an error when attempting to read or update the Cluster registry.

User response: Ensure the Microsoft Cluster Server (MSCS) support has been correctly installed and the Cluster Service has been started. Ensure that the current logon user account has sufficient authority to access the Cluster registry.

DBI1904N **The DB2 instance is not Clustered.**

Explanation: DB2 attempted to access an instance to perform cluster operations but the instance was not identified as a DB2 Cluster instance.

User response: Verify the instance was created with the Cluster options. The instance may be dropped and re-created using the Cluster option but the Node directory, Database Directory, and Database Manager Configuration file will be lost.

DBI1905N **A System Error occurred during a DB2 Cluster operation.**

Explanation: A system error was detected while DB2 was performing a Cluster operation.

User response: Check to make sure the Windows operating system is at the correct level for Cluster Support.

Also ensure the Microsoft Cluster Server (MSCS) support has been correctly installed and the Cluster Service has been started.

DBI1906N **A error occurred while attempting to add a node to a DB2 Cluster instance.**

Explanation: An error was encountered while attempting to update a node to support a DB2 Cluster instance.

User response: Ensure the remote system is on-line and accessible over the LAN. Also ensure you have the necessary privilege to update the Registry on the remote system.

DBI1907N **A error occurred while attempting to delete a node from a DB2 Cluster instance.**

Explanation: An error was encountered while attempting to delete a node from a DB2 Cluster instance.

User response: Ensure the remote system is on-line and accessible over the LAN. Also ensure you have the necessary privilege to update the Registry on the remote system.

A workstation cannot drop cluster support from itself.

The instance must be dropped to perform this task.

DBI1908N **A error occurred while attempting create the instance profile on a remote system.**

Explanation: A error was detected while attempting to create the instance profile on a remote system to enable Cluster Support.

User response: Ensure the remote system is on-line and accessible over the LAN. Also ensure you have the necessary privilege to update the Registry on the remote system.

DBI1909N **A error occurred while attempting delete the instance profile on a remote system.**

Explanation: A error was detected while attempting to delete the instance profile on a remote system to disable Cluster Support.

User response: Ensure the remote system is on-line and accessible over the LAN. Also ensure you have the necessary privilege to update the Registry on the remote system.

DBI1910N **Usage : DB2NCRT -I:Instance-Name -C:cluster-node**

Explanation: The user specified incorrect arguments to the DB2NCRT command.

User response: Re-issue the command specifying a clustered instance and a valid workstation name for the cluster node.

DBI1911N **Usage : DB2NLIST -I:Instance-Name**

Explanation: The user specified incorrect arguments to the DB2NLIST command.

User response: Re-issue the command specifying a clustered instance.

DBI1912I **The DB2 Cluster command was successful.**

Explanation: The user request was successfully processed.

User response: No action required.

DBI1913N **Unable to connect to the remote machine registry.**

Explanation: DB2 encountered an error when attempting to connect to the Windows registry on a remote machine.

User response: Ensure that the target machine is running and that the computer name specified is correct. Also ensure that the current logon user account

has sufficient authority to connect to the remote machine registry.

DBI1914N **Unable to connect to the Cluster registry.**

Explanation: DB2 encountered an error when attempting to connect to the Cluster registry.

User response: Ensure that the Microsoft Cluster Server (MSCS) support has been correctly installed and the Cluster Service has been started.

DBI1915N **The target machine does not belong to the Microsoft Cluster Server (MSCS) cluster.**

Explanation: The command failed for one of the following reasons:

(1)

The db2iclus utility attempted to remove a machine from a DB2 MSCS instance but the target machine had not been added to the DB2 MSCS instance, or

(2)

The db2iclus utility attempted to add a machine to a DB2 MSCS instance but the target machine does not have the required Microsoft Cluster Server (MSCS) support software installed.

User response: Follow the action corresponding to the cause of the problem:

(1)

Ensure that the machine is part of the DB2 MSCS instance by running the "db2iclus list" command, or

(2)

Install the Microsoft Cluster Server (MSCS) cluster software on the target machine and issue the command again.

DBI1916N **The instance name specified is not valid.**

Explanation: The instance specified by the instance name parameter does not exist for the local machine.

User response: To find a list of valid instances, run the db2ilist command. Issue the command again using a valid instance name.

DBI1917I **The instance upgrade or instance update completed successfully. However you might need to reconfigure any ODBC data sources for instances cataloged under this instance.**

Explanation: The instance upgrade or instance update is not able to determine which ODBC data sources are associated with the instance that was upgraded or removed. As a result, these data sources were not modified.

User response: Reconfigure the data sources using the DB2 CATALOG ODBC DATA SOURCE command, the Configuration Assistant or the Data Sources (ODBC) control panel applet so you can access these data sources.

DBI1918N An error occurred when accessing the Windows registry.

Explanation: DB2 encountered an error when attempting to read or update the Windows registry of the target machine.

User response: Ensure that the current logon user account has sufficient authority to access the Windows registry of the target machine. If you already have a clustered instance, ensure the cluster service is started. If the problem persists, contact your IBM service representative for assistance.

DBI1919N An error occurred when accessing the TCP/IP services file.

Explanation: DB2 encountered an error when reading or updating the TCP/IP services file.

User response: Ensure that the services file exists in the system and that the file can be accessed for read and write authority. Also, ensure that the contents of the file are valid and no duplicate of entries exist in the services file.

DBI1920N The DB2 instance profile path parameter is not valid.

Explanation: DB2 is unable to access the DB2 instance profile path specified because the path does not exist or write access has not been given to the current logon user account.

User response: Ensure that the path name for the DB2 instance profile points to a valid directory and that the current logon user account has write access to the directory.

DBI1921N The user account or password is not valid.

Explanation: The user account or the password is not valid.

User response: Issue the command again with the correct account.

DBI1922N The host name is not valid.

Explanation: The TCP/IP host name parameter is not valid or does not exist.

User response: Ensure that TCP/IP is operational on the system. If a domain name server is used, ensure that the domain name server machine is active. Issue the command again using the correct TCP/IP host name.

DBI1923N The TCP/IP port range parameter is not specified.

Explanation: The db2icrt utility requires the TCP/IP port range to be specified when creating a partitioned database instance if the entries for TCP/IP port range for the instance have not been added to the services file.

User response: Choose a range of TCP/IP ports that are available for the system and issue the db2icrt command again with the -r option.

DBI1924N The computer name is not valid.

Explanation: The computer name parameter is not valid

User response: Ensure that the target machine is running and that the computer name specified is correct. Also ensure that the current logon user account has sufficient authority to connect to the target machine registry.

DBI1925N Unable to query the status of the DB2 service.

Explanation: DB2 is unable to query the service status for the DB2 service.

User response: Ensure that the DB2 service exists on the target machine. Ensure that the current logon user account has sufficient authority to query the service status on the target machine.

DBI1926N The password is expired.

Explanation: The account password has expired.

User response: Change your password, then retry the request using the new password.

DBI1927N Usage:

```
db2iclus { ADD /u:Username,Password
           [/m:Machine name] |
           DROP [/m:Machine name] |
           MIGRATE /
           p:Instance profile path |
           UNMIGRATE }
           [/i:instance name]
           [/DAS DASname]
```

[/c:Cluster name]

Explanation: Valid arguments for this command are:

ADD to add an MSCS node to a DB2 MSCS instance

DROP to remove an MSCS node from a DB2 MSCS instance

MIGRATE
to migrate a non-MSCS instance into an MSCS instance

UNMIGRATE
to undo the MSCS migration

Valid options for this command are:

/DAS to specify the DAS instance name. This option is required when performing the cluster operation against the DB2 Administration Server.

/c to specify the MSCS cluster name if different from the default/current cluster

/p to specify the instance profile path. This option is required when migrating a non-MSCS instance into an MSCS instance

/u to specify the account name and password for the DB2 service. This option is required when adding another MSCS node to the DB2 MSCS partitioned database instance

/m to specify the remote computer name for adding or removing an MSCS node

/i to specify the instance name if different from the default/current instance

User response: Enter the command again using valid parameters

DBI1928N The user account is not specified.

Explanation: When creating a partitioned database instance or adding a node to a partitioned database instance, you must specify the user account parameter. Under the Windows environment, each database manager partition (or MPP node) is required to run under a valid Windows account to access the network share which contains the instance directory.

User response: Issue the command again with the -u option to specify the user account name and password.

DBI1929N The instance is already configured for MSCS support.

Explanation: An attempt to upgrade an instance into an MSCS instance failed because the instance is already configured for MSCS support.

User response: No action required.

DBI1930N The target machine is already part of the DB2 MSCS instance.

Explanation: An attempt to add an MSCS node to the DB2 MSCS instance failed because The target machine is already part of the DB2 MSCS instance.

User response: No action is necessary.

DBI1931N The database partition server (node) is active.

Explanation: The database partition server can not be dropped or modified because it is active.

User response: Ensure that the database partition server is stopped before attempting to drop or change it. To stop a database partition server, use the db2stop command as follows:

```
db2stop dbpartitionnum <db-partition-number>
```

DBI1932N Can not add the database partition server on a remote machine because the same instance on that machine already exists.

Explanation: The db2nprt command failed to add a new database partition server on a remote machine because there is an existing instance on the remote machine.

User response: If the instance on the remote machine is no longer used, remove the instance by running the db2idrop command on the remote machine.

DBI1933N The database partition server can not be moved to a new machine because there is existing database(s) in the system.

Explanation: When the db2nchg command is issued with the /m:machine option, the command will move the database partition server to a new machine. The db2nchg will fail if there is existing database(s) in the system.

User response: To move the database partition server to a new machine when there is database(s), use the db2start restart command. For more information on the db2start command, refer to the DB2 Command Reference.

DBI1934N An invalid argument was entered for the db2iupdt command.

Explanation: You can update a DB2 database manager instance to a higher level within a release, update an instance to a DB2 pureScale instance, or scale a DB2 pureScale instance by using the db2iupdt command.

This message is returned when the db2iupdt command is invoked with incorrect or missing parameters.

User response: Enter the command again using valid arguments.

DBI1935N The db2iupdt command failed to update the instance directory.

Explanation: The db2iupdt command failed to update the instance directory because of one of the following reasons:

- (1) the db2audit.cfg file is missing from the \SQLLIB\CFG directory.
- (2) the db2iupdt command does not have the required authority to create files or sub-directories within the instance directory.

User response: Ensure that the db2audit.cfg file exists in the \SQLLIB\CFG directory and that the current logon account has sufficient authority to create files and directories in the instance directory. The instance directory is located in \SQLLIB\InstName (where InstName is the name of the instance).

DBI1936N The db2iupdt command failed to update the database manager configuration file.

Explanation: The db2iupdt command failed to update the database manager configuration file because the previous database manager configuration file is corrupted or missing. The current instance is in inconsistent state and can not be used.

User response: Copy the default database manager configuration file to the instance directory then re-submit the command. The instance directory is located in \SQLLIB\InstName. The default database manager configuration file is named db2system and can be found in the \sqllib\cfg directory.

DBI1937W The db2nrcr command successfully added the node. The node is not active until all nodes are stopped and started again.

Explanation: The db2nodes.cfg file is not updated to include the new node until all nodes are simultaneously stopped by the STOP DATABASE MANAGER (db2stop) command. Until the file is updated, the existing nodes cannot communicate with the new node.

User response: Issue db2stop to stop all the nodes. When all nodes are successfully stopped, issue db2start to start all the nodes, including the new node.

DBI1940N The DB2 Administration Server is active.

Explanation: The DB2 Administration Server cannot be modified because it is active.

User response: Ensure that the DB2 Administration

Server is stopped before attempting to change it. To stop the DB2 Administration Server, use the db2admin command as follows:
db2admin stop

DBI1941W The DB2 Administration Server does not exist.

Explanation: The DB2 Administration Server was not found on the machine.

User response: Create the DB2 Administration Server on the machine.

DBI1942N The DB2 Administration Server already exists on the target machine.

Explanation: The operation could not be completed because the DB2 Administration Server already exists on the target machine.

User response: Remove the DB2 Administration Server from the target machine and retry the operation.

DBI1943N The DB2 Administration Server is not Clustered.

Explanation: DB2 attempted to access a DB2 Administration Server to perform cluster operations but the DB2 Administration server is not configured for MSCS support.

User response: Verify the DB2 Administration Server is configured for MSCS support.

DBI1944N The DB2 Administration Server is already configured for MSCS support.

Explanation: An attempt to upgrade a DB2 Administration Server failed because the DB2 Administration Server is already configured for MSCS support.

User response: No action required.

DBI1945N The DB2 Copy Name provided does not exist or the DB2 registry is corrupted.

Explanation: To update or upgrade an instance to a specific DB2 copy, the DB2 copy must exist. The DB2 Copy Name provided does not exist or the DB2 registry is corrupted.

User response: Check the DB2 registry table and provide a valid DB2 Copy Name. For a list of the DB2 copies installed, run the db2ls command.

DBI1946N An invalid argument was entered for the db2iupgrade command.

Explanation: You can upgrade a DB2 database manager instance from one release to a later release by using the db2iupgrade command.

This message is returned when the db2iupgrade command is invoked with an invalid argument.

User response: Enter the command again, specifying valid parameters.

DBI1947N The db2iupgrade command failed to update the database manager configuration file.

Explanation: The db2iupgrade command failed to update the database manager configuration file because the previous database manager configuration file is corrupted or missing. The current instance is in an inconsistent state and cannot be used.

User response: Copy the default database manager configuration file to the instance directory, then re-submit the command. The instance directory is located in \SQLLIB\instanceName. The default database manager configuration file is named db2system and can be found in the cfg directory under your installation path.

DBI1948N Files cannot be deleted.

Explanation: An error occurred when removing files or directories belonging to the specified instance.

User response: Ensure that you have write access to the location where the files are located.

DBI1949E Upgrading instance *inst-name* is not supported.

Explanation: The instance cannot be upgraded because either:

- Upgrading from this version of the instance is not supported.
- The instance is already using the current version of the product and upgrading is not required.

User response: Make sure that the instance version is supported for upgrade and try the command again with a supported instance name or a valid DB2 Copy Name. For details about which DB2 product versions are supported for upgrade, refer to the DB2 Information Center.

DBI1950W The instance *instance-name* is already in the instance list.

Explanation: The instance that is being created or upgraded is already in the instance list.

User response: Verify that the list of instances as reported by the db2ilist command is correct.

DBI1951W The instance *instance-name* was not found in the instance list.

Explanation: The specified instance was not found in the instance list.

User response: Verify that the list of instances as reported by the db2ilist command is correct.

DBI1952E The instance name *instance-name* is invalid.

Explanation: The instance name is invalid because of one of the following reasons:

1. The name is more than 8 characters long.
2. The name starts with sql, ibm, or sys.
3. The name contains characters other than a-z, \$, #, @, _, or 0-9.
4. The name does not exist.

User response: Try the command again with a valid instance name.

DBI1953E The instance drop or update operation failed because the instance is active.

Explanation: You can update a DB2 database manager instance to a higher level by using the db2iupdt command. You can remove an instance by using the db2idrop command.

Before you can update or drop an instance, all processes that are running for the instance must be stopped. This message is returned when an attempt is made to update or drop an instance that is active.

User response:

1. Wait for all applications that are using the instance to finish processing.
2. Stop the instance by using the db2stop command from the instance user ID.
3. Issue the db2idrop or db2iupdt command again.

DBI1954E The instance name is currently used by the DB2 Admin Server.

Explanation: The instance cannot be created because it is currently used by the DB2 Admin Server.

User response: Try the command again with a different instance name.

DBI1956I An invalid argument was entered for the db2ilist command.

Explanation: The db2ilist command lists the names of DB2 instances associated with the DB2 installation path where db2ilist is run.

The syntax for the db2ilist command is:

db2ilist [-h]

Valid arguments for this command are:

-h

Displays the usage information.

User response: Enter the command again without any parameters.

DBI1957E The syntax of the db2icrt command is incorrect.

Explanation: You can create a DB2 database manager instance by using the db2icrt command.

Syntax:

```
db2icrt InstName [-s {ese|wse|standalone|client}]
                  [-p instance profile path]
                  [-u username,password]
                  [-h hostname]
                  [-r baseport,endport]
                  [-j textSearchConfig]
                  [-?]
```

User response: Enter the command again using the valid parameters.

DBI1958N The instance profile cannot be added to the registry.

Explanation: An error occurred when adding the instance profile to the registry. The instance was not created.

User response: Contact your IBM service representative.

DBI1959N The instance directory cannot be created.

Explanation: An error occurred when creating the required files or directories for the new instance. The instance was not created.

User response: Ensure that you have write access to the location where the instance directory is created. The instance directory is created under the path where the product is installed. You can move the instance to another directory by using the /p option in db2iupgrade.

DBI1960N The DB2 service cannot be created.

Explanation: An error occurred when registering a DB2 service because of one of the follow reasons:

1. A service of the same name already exists or was marked as deleted but not cleaned up until the next system reboot.
2. You don't have sufficient access to create a Windows service

User response:

1. Ensure that the name of the instance does not match any of the existing Windows services. The list

of services in the system can be found in the Windows registry under HKEY_LOCAL_MACHINE\System\CurrentControlSet\Services. If the service was marked as deleted, you have to reboot the machine to have the service removed from the system.

2. Log on to a user account that belongs to the local Administrators group and retry the operation.

DBI1961N The node key for the new instance cannot be added to the registry.

Explanation: An error occurred when adding a node key for the new instance in the instance profile registry.

User response: Contact your IBM service representative.

DBI1962N The new node can not be added to the registry.

Explanation: An error occurred when adding the new node to the instance profile registry.

User response: Contact your IBM service representative.

DBI1963N An error occurred when granting the user rights to the account for the instance.

Explanation: When configuring the account for the DB2 Service, the account must be granted the user rights:

1. Act as part of the operating system
2. Create a token object
3. Increase quota
4. Logon as a service
5. Replace a process level token

An error occurred when granting the user rights to the account for the instance.

User response: If the machine belongs to a domain, ensure that the primary domain controller is active and accessible over the network. Otherwise, contact your IBM service representative.

DBI1964N The logon account cannot be assigned to the DB2 service.

Explanation: An error occurred when assigning a logon account to the DB2 service.

User response: Ensure user that the user name and the password for the logon account is valid.

DBI1965N **The node *node-number* was not found in the node list.**

Explanation: The specified node was not found in the node list.

User response: Verify that the node exists by displaying the list of nodes using the DB2NLIST command.

DBI1966N **The node cannot be deleted from the registry.**

Explanation: An error occurred when removing the specified node from the instance profile registry.

User response: Contact your IBM service representative.

DBI1967N **The DB2 service cannot be de-registered.**

Explanation: An error occurred when de-registering a DB2 service.

User response: Ensure that you have sufficient authority to delete a Windows service.

DBI1968N **An error occurred when changing the node configuration in the profile registry.**

Explanation: An error occurred when changing the node configuration in the profile registry.

User response: Contact your IBM service representative.

DBI1969N **An error occurred when creating a new file in the instance directory.**

Explanation: An internal error occurred when creating a new file in the instance directory.

User response: Ensure that you have write access to the instance directory.

DBI1970N **An error occurred when reading from or writing to a file in the instance directory.**

Explanation: An internal error occurred when reading from or writing to a file in the instance directory.

User response: Contact your IBM service representative.

DBI1971N **The instance profile cannot be removed from the registry.**

Explanation: An error occurred when removing the instance profile from the registry.

User response: Contact your IBM service representative.

DBI1972N **The instance directory cannot be removed.**

Explanation: An error occurred when removing the required files or directories belonging to the specified instance.

User response: Ensure that you have write access to the location where the instance directory is located.

DBI1973N **An attempt to configure the DB2 service to start automatically failed.**

Explanation: An error occurred when setting the DB2 service to start automatically.

User response: Reboot the machine and try the command again. If the problem persists, contact your IBM service representative.

DBI1974N **The profile variable DB2ADMINSERVER cannot be set in the profile registry.**

Explanation: An error occurred when setting the profile variable DB2ADMINSERVER in the profile registry.

User response: Contact your IBM service representative.

DBI1975N **The environment variable DB2ADMINSERVER cannot be removed from the profile registry.**

Explanation: An error occurred when removing the environment variable DB2ADMINSERVER in the profile registry.

User response: Contact your IBM service representative.

DBI1978E **Update of instance *inst-name* is not supported.**

Explanation: The DB2 copy of the instance does not have the same version number as the target DB2 copy. You can only update an instance from one copy to another if both copies have the same version. To move a DB2 instance between different versions, use the db2iupgrade command.

User response: Make sure that the instance is valid for update and then try the command again with a valid instance name.

DBI1980W

DBI1980W **Unable to create user ID *userID* in UPM.**

Explanation: The user ID that you provided during product installation could not be created in User Profile Management (UPM). This user ID is required to start the DB2 Administration Server.

User response: Manually create a user ID and password using UPM and then manually start the DB2 Administration Server. You may need to set the DB2 Administration Server to use this user ID and password combination using the “db2admin” command.

Chapter 86. DBI20000 - DBI20499

DBI20000E A DB2 instance could not be created as specified. A rollback of the instance creation will be started.

Explanation: The message indicates that the instance creation failed. Instance rollback happens on all involved hosts. Instance rollback is triggered if one of the following happens:

- Any operation (DB2 binary installation, GPFS cluster setting, RSCT peer domain setting, and instance file set creation, for example) fails on install-initiating host (IIH).
- Any operation fails on a CF host.
- Any operation fails on all member hosts.

Instance rollback does not roll back DB2 binaries.

User response: For details, see the installation log. By default, the installation log is located in the /tmp directory. Additional error messages and details are contained in the log file. After fixing the problems shown in the log file, re-run the instance creation command.

DBI20001E A DB2 instance could not be created as specified. A rollback of the instance creation has been performed.

Explanation: Instance creation rollback has been performed on all involved hosts. If DB2 binaries are successfully installed, they will not be rolled back. The DB2 pureScale instance that was partially created during the current installation has been removed.

User response: For details, see the installation log. By default, the installation log is located in the /tmp directory. After fixing the problems shown in the log file, re-run the instance creation command.

DBI20002E A DB2 instance has been partially created. DB2 members that could not be added successfully will not be included in the instance. A cleanup of the affected members will be started.

Explanation: The message occurs when the instance is partially created on some of the hosts. Instance rollback happens on the hosts where the instance creation did not complete. Some immediate messages in the installation log after this line are rollback related. Rollback of the partial instance creation is triggered if any operation (DB2 binary installation, GPFS cluster setting, RSCT peer domain setting, instance file set creation, for example) fails on a member host, then rollback happens on the corresponding host.

Instance rollback does not roll back DB2 binaries.

User response: To determine the members that were not successfully added into the instance, see the log file. Additional error messages and details are contained in the log file. After fixing the problems shown in this log file, re-run the instance creation command.

DBI20004E An error occurred while removing *registry_variable* global registry variable record on the following hosts:
host-name-list.

Explanation: This error occurs when removing the GPFS_CLUSTER global registry or the PEER_DOMAIN global registry from specific hosts.

User response: Contact IBM Software Support.

DBI20006E Installing the IBM Data Server Driver Package failed because the installer could not determine whether to install a new copy or to upgrade an existing copy because no copy name was specified.

Explanation: On Windows operating systems, the IBM Data Server Driver Package simplifies application deployment. This driver, which has a small footprint, is designed to be redistributed by independent software vendors (ISVs) and to be used for application distribution in mass deployment scenarios typical of large enterprises.

You can install the IBM Data Server Driver Package or perform a maintenance install on an existing copy of the driver package using either the DB2 Setup wizard or the SETUP command. When you install the IBM Data Server Driver Package, the behavior of the installer depends on whether there are copies of the driver package already installed and whether a driver package copy name is specified to the install utility.

This message is returned when an attempt is made to install the IBM Data Server Driver Package and the installer cannot determine whether to create a new copy of the driver package or to perform a maintenance install on an existing copy of the driver package.

User response: To install a new copy of the IBM Data Server Driver Package, use one of the following methods:

Using the GUI installer

- Install a new copy using the generated, default name by issuing the SETUP command with the /o option.

- Install a new copy using a name you choose by issuing the SETUP command with the /n <copy-name> option.

Using a silent install

Set the COPY_NAME keyword to a new copy name in the response file.

To perform a maintenance installation on an existing IBM Data Server Driver copy, specify which copy to upgrade using one of the following methods:

Using the GUI installer

Issue the SETUP command with the /n <copy-name> option to specify the copy you want to upgrade.

Using silent install

Set the COPY_NAME keyword to the existing copy you want to upgrade in the response file.

DBI20007E Upgrading the IBM Data Server Driver for ODBC and CLI failed because multiple copies of the driver exist on the system and the upgrade utility could not determine which copy to upgrade.

Explanation: You can install multiple copies of the IBM Data Server Driver for ODBC and CLI on the same machine. You might want to do this if you have two database applications on the same machine that require different versions of the driver. When multiple copies of the driver exist on this system, you can identify one of them as the "default" copy.

To upgrade a copy of the IBM Data Server Driver for ODBC and CLI when multiple copies of the driver are installed on the same system, the upgrade utility can determine which copy to upgrade in two ways: The name of the copy to upgrade can be specified with the upgrade command. If one of the copies is identified as the default and no name is specified with the upgrade command, the upgrade utility will upgrade the default copy.

This message is returned when an attempt is made to upgrade one of multiple copies of the IBM Data Server Driver for ODBC and CLI on a system and none of the existing copies is set as the default and no copy was specified by name with the upgrade command.

User response: To upgrade an existing IBM Data Server Driver copy, specify which copy to upgrade using one of the following methods:

Using the GUI installer

Issue the SETUP command with the /n <copy-name> option to specify the copy you want to upgrade.

Using silent install

Set the COPY_NAME keyword to the existing copy you want to upgrade in the response file.

DBI20008W Neither the Path environment variable nor the Classpath environment variable were updated during the installation of the IBM Data Server Driver Package because the following DB2 copy was previously installed: *copy-name*.

Explanation: When you install the IBM Data Server Driver Package, you can specify that the Data Server Driver Package should be the default database client interface.

When the Data Server Driver Package is set as the default database client interface during install, the Path and Classpath environment variables are automatically updated to include Data Server Driver Package details. However, if the Data Server Driver Package install detects that another IBM data server client or driver is already installed, the Path and Classpath environment variables will not be updated to include Data Server Driver Package details.

This message is returned when the Data Server Driver Package install detects that another IBM data server client or driver is already installed and as a result does not update the Path or Classpath environment variables.

User response: To configure the environment for the Data Server Driver Package, manually update the Path and Classpath environment variables.

DBI20009E Creation of Tivoli SA MP resources for the DB2 instance, *instance_name*, has failed.

Explanation: This error occurs because of a failure in the creation of Tivoli SA MP resources for the DB2 pureScale instance.

- This error only occurs during DB2 pureScale instance creation. The error will trigger the instance rollback.
- The command that the installer internally calls is db2cluster -cm -resources -create This command is run as the instance user. For more details, see the db2cluster command.
- The log file for this command is <instance_user>/sqllib/db2dump/db2diag.log. Since a rollback rolls back <instance_user>/sqllib, the file is backed up as /tmp/db2diag.log.process_id.
- There are multiple possibilities that may cause this error. A possible reason for this is that the db2nodes.cfg is incorrect or failed to be created. The file can be found at <instance_user>/sqllib/db2nodes.cfg. During rollback, the file is backed up as /tmp/db2nodes.cfg.process_id.

For more information, see the post-installation recommendations section of this log.

User response: When this error occurs and instance rollback is triggered, contact IBM service with all log and trace files. For details about trace and log file collection, see the DB2 Information Center.

DBI20010W The DB2 installer detected that a copy of Tivoli SA MP is already installed, and that the version of the currently installed Tivoli SA MP product is older than the version that is included in the DB2 product installation media. If the DB2 installation proceeds, Tivoli SA MP will be updated or upgraded to the following new version:
new-SA-MP-version.

Explanation: You can use the copy of the Tivoli System Automation for Multiplatforms (SA MP) Base Component that is integrated with your DB2 database product to manage the high availability of your DB2 database system. You can install IBM Tivoli SA MP by using either the DB2 installer or the installSAM installation script that is included in the DB2 product installation media.

The DB2 product is intended to use the version of Tivoli SA MP that is included in the DB2 product installation media. If an older version of Tivoli SA MP is already installed before the DB2 product is installed, Tivoli SA MP must be updated or upgraded to the version that is included in the DB2 product installation media..

This message is returned when the DB2 installer detects an older version of Tivoli SA MP is already installed.

User response: Optional: verify that other software on the system that is using the currently installed version of Tivoli SA MP is compatible with the new version of Tivoli SA MP.

DBI20011E Deletion of the RSCT peer domain, *domain_name*, failed. Manually remove the domain by running the following command: *command*.

Explanation: Deletion of the RSCT peer domain error can occur during a rollback or when running the db2idrop -g command.

User response: To manually remove the RSCT peer domain by performing the following steps:

1. Check if there is still resource attached to the peer domain by running lssam. If not, proceed to step 5.
2. Switch to the instance owner by su - *instance_owner*
3. Remove resource by db2cluster -cm -repair -resources
4. Switch back to root.
5. Remove RSCT peer domain by db2cluster -cm -delete -domain *domain_name*. For more

information, about manually cleaning up an IBM Reliable Scalable Cluster Technology peer domain, see the DB2 Information Center.

DBI20012W The DB2 installer successfully upgraded the SA MP Base Component. Previous version: *old-SA-MP-version*. New version: *new-SA-MP-version*. However, the license of the new SA MP product is a trial license, not a full license.

Explanation: You can upgrade IBM Tivoli System Automation for Multiplatforms (SA MP) using either the DB2 installer or the installSAM install script that is included in the IBM Data Server install media.

Depending on which DB2 product licenses you have purchased, there are conditions under which you can use IBM Tivoli System Automation for Multiplatforms (SA MP) that is integrated with IBM Data Server.

This message is returned when a fully licensed version of SA MP Base Component is upgrade to a version of SA MP that has only a trial license. It is possible that if the host machine on which the trial version of the SA MP product is installed is restarted, the SA MP product will not longer function.

User response: For full use of the SA MP product, update the trial license of the SA MP product by performing the following steps:

1. Acquire a full license key from either the IBM Passport Advantage web site or the activation CD in an installation media pack from IBM.
2. Register the license key by using the db2licm command.

DBI20013E The DB2 installer detected that the given device path *device_path* is a symbolic link on the host *host_name*.

Explanation: This message indicates that the given device path for tie-breaker device path, or instance shared device, is a symbolic link. The device path must be a storage device.

User response: Specify a valid storage device as the device path.

DBI20014E Removal of DB2 resources has failed for DB2 instance *instance_name*.

Explanation: Instance resource is the RSCT resource created for the DB2 pureScale instance during instance configuration. The instance is not in a usable state.

- This error might occur when a DB2 pureScale instance creation is rolled back or a DB2 pureScale instance is dropped.
- The command that the installer internally calls is db2cluster -cm -resources -delete running as the instance user. For more details, see the db2cluster command.

- During db2idrop -g, if instance resource deletion fails, no further instance dropping action will be performed. The log file for this command is <instance_user>/sqllib/db2dump/db2diag.log.
- Check if <instance_user>/sqllib/db2nodes.cfg is valid.
- During a rollback in instance creation, instance resource deletion will also be called. Since rollback rolls back <instance_user>/sqllib, the file is backed up as /tmp/db2diag.log.process_id.

User response: For more details, see the db2diag.log for error messages. This log file can be found in <instance_user>/sqllib/db2dump/db2diag.log. Correct the problem and rerun the db2idrop command.

DBI20015E **Shrinking of the RSCT peer domain, domain_name, failed on the following hosts: host-name-list. Manually remove the hosts by running the following command: command.**

Explanation: This error can occur during partial rollback or db2iupdt -drop. It occurs because shrinking of the RSCT peer domain has failed.

User response: Manually remove failed hosts from the RSCT peer domain by performing the following steps:

1. Check if there are still resources attached to the host by running lssam.
 - if a resource is still attached, proceed to step 2.
 - if a resource is not attached, run db2cluster -cm -remove -host host_name from IIH.
2. From the install-initiating host (IIH), switch to the instance owner by running su - instance_owner.
3. Remove resource by running db2cluster -cm -delete -resources.
4. Switch back to root.
5. Remove the failed hosts from the RSCT peer domain by running db2cluster -cm -remove -host host_name.
6. Switch to the instance owner again.
7. Recreate the resources by running db2cluster -cm -create -resources.
8. Switch back to root.

DBI20016E **The cluster interconnect netname net_name for host host_name is not on the same subnet as the installation-initiating host.**

Explanation: This error occurs when the cluster interconnect netname for a host is not on the same subnet as the installation-initiating host.

All cluster interconnects need to be on the same subnet for performance reasons (all hosts within the same subnet can usually be reached in one routing hop). For example if the cluster interconnect network is

configured with the network address 192.168.0.0/24, then all cluster interconnect netname addresses should start with 192.168.0 (for example, 192.168.0.1, 192.168.0.2, and so on).

User response: Check the network card configuration on the new host (for example run ifconfig -a) and check /etc/hosts if a name was used instead of an address. Reconfigure the network adapter or choose a different one.

DBI20017E **The cluster interconnect netname net_name is invalid.**

Explanation: The DB2 installer cannot identify the cluster interconnect netname.

User response: Check the following:

- Check the command for a typo error
- Check the machine's network configuration (for example check /etc/hosts) or
- Use the ping and nslookup tools to confirm the cluster interconnect netname maps to the same IP address on all hosts. Check that the IP address maps back to the same cluster interconnect netname on all hosts.

Enter the netname or IP address of a valid device.

DBI20018E **The cluster interconnect netname for host host-name failed to ping the installation-initiating host.**

Explanation: When adding a new host, validation checks to see if the new host can send a small data packet to the install-initiating host (IIH) and receive a reply. This send-reply test is commonly known as a ping. This message is the result of a failed ping.

User response: Check the network adapter and cable, or choose a different one. If there is a problem, you can verify the results by running this command line from the remote host's console: ping <installation-initiating host name>. After it is verified that the problem occurs outside the DB2 installer, there are various things you can check to find the source of the problem. For example:

- bad physical connections (for example a loose cable),
- a non-working network adapter driver,
- misconfigured network.

DBI20019E **The DB2 installer detected a conflict between the DEFAULT_INSTPROF record and the instance shared directory specified.**

Explanation: The DEFAULT_INSTPROF record in the Global Registry indicates the instance shared file system has already been set up by the DB2 installer. In

this case, the following options or keywords are not needed:

- For response file installations:
 - INSTANCE_SHARED_DEVICE_PATH
 - INSTANCE_SHARED_DIR
- For the db2icrt or db2iupdt command:
 - instance_shared_dev
 - instance_shared_dir

If the value for INSTANCE_SHARED_DIR (or instance_shared_dir) does not match with the existing instance shared file system mount point, the installation fails.

User response: Do not specify the instance shared directory.

DBI20020E The DB2 installer detected that GPFS cluster *cluster_name* was previously created on host *host_name*. However, the GPFS cluster is offline on that host.

Explanation: A GPFS cluster created by the DB2 software already exists. To create a DB2 pureScale instance by reusing this GPFS cluster, bring the cluster online on the specified host.

User response: To start the GPFS cluster, run the following command:

```
<installed_path>/bin/db2cluster -cfs -start -all
```

For details about the db2cluster command, see the DB2 Information Center.

DBI20021E The DB2 installer detected that the *reg-var* variable record is defined in the global registry. However, the value *dir-name* is not an existing directory.

Explanation: This message occurs if the global registry variable record specified indicates that a shared file system already exists. The DB2 product expects a shared file system exists for the instance shared file.

The directory name value is the mount point for the file system. The DB2 installer detects the mount point does not exist on the system. A possible reason is the specified global registry record is left over from a previous DB2 operation and the actual file system does not exist anymore, or, the filesystem was remounted somewhere else but the registry entry was not changed.

User response: Contact IBM Software Support.

DBI20022E The DB2 installer detected that the variable record *reg-var* is defined in the global registry. However, the GPFS cluster does not exist on host *host_name*.

Explanation: The global registry variable record GPFS_CLUSTER indicates that there is a DB2-created

GPFS cluster on the host. In this case, the installer detects there is no GPFS cluster on the host. The DB2 product expects an existing GPFS cluster on this host.

User response: Contact IBM Software Support.

DBI20023E The installation operation failed because a user-managed RSCT peer domain is active on the following host machine: *host-name*.

Explanation: This error occurs when the DB2 installer fails to create an IBM Reliable Scalable Cluster Technology (RSCT) peer domain on a host because another, user-managed RSCT peer domain is active.

The peer domain that was not created by the DB2 installer must be stopped or removed prior to creating an RSCT peer domain by the DB2 installer.

User response:

1. Stop the peer domain on the host specified by using the using db2cluster command or the stoprpdomain command.
2. Restart the installation process.

DBI20024E The DB2 installer detected that the given device path *device_path* is not a valid device on host *host_name*.

Explanation: This message indicates that the given device path for tie-breaker device path, or instance shared device, is not a storage device. The device path must be a storage device. A device can be invalid for the following reasons:

- there are existing files or directories
- the device does not exist
- the device is not a storage device, such as pipe, fifo device, character device
- root has no searchable access authority to the given path
- the path name is too long.

User response: Specify a valid storage device as the device path.

DBI20025E The DB2 installer detected that the IBM Reliable Scalable Cluster Technology (RSCT) peer domain *domain_name* was previously created on host *host_name* by DB2, but the peer domain is offline on that host.

Explanation: This error occurs when the DB2 installer detects that an RSCT peer domain was previously created on a host, but it is offline and the installation-initiating host is in a different RSCT peer domain. The remote host must be removed from the host list.

User response: To remove the host from a DB2 pureScale instance:

1. log on to a different host which belongs to the same DB2 pureScale instance
2. run db2iupdt [-d] -add -m lcf
<host_name>:<interconnect_name> -u <fenced_id>
<instance_owner>

For details about the db2iupdt command, see the DB2 Information Center.

DBI20026E The following hosts *host-name-list* are not included in the instance.

Explanation: This post installation message shows the hosts that have not been included in the instance due to an error. If DB2 binaries were successfully installed, they have not been removed. Look in the installation log for more information on the error that has caused this rollback.

User response: After fixing the problems shown in this log file, extend the instance to these hosts by running the instance extension command (db2isetup -add or db2iupdt -add).

DBI20027E The GPFS cluster creation failed on host *host_name*. Failed command: *command_name*.

Explanation: The GPFS cluster failed to be created during instance creation. This failure will trigger an instance creation rollback.

The db2cluster command log file can be found in /tmp/ibm.db2.cluster.*

User response: Refer to the db2cluster command log file. This log file can be found in /tmp/ibm.db2.cluster.*. Additional error messages and details are contained in the log file. After the problem is solved, rerun the command.

DBI20028E The GPFS cluster, *db2cluster_xxxx*, cannot be deleted using the *command_name* command.

Explanation: This message occurs when the GPFS cluster fails to be deleted during a rollback.

User response: Refer to the db2cluster command log file. This log file can be found in /tmp/ibm.db2.cluster.*. Additional error messages and details are contained in the log file. After the problem is solved, rerun the command to remove the host from GPFS cluster.

DBI20029E The GPFS cluster, *db2cluster_xxxx*, failed to create the file system *fs_name*. Failed command: *command_name*.

Explanation: This message occurs when the GPFS file system creation fails.

User response: Refer to the db2cluster command log file. This log file can be found in /tmp/ibm.db2.cluster.*. Additional error messages and details are contained in the log file. After the problem is solved, rerun the command.

DBI20030E The GPFS cluster, *db2cluster_xxxx*, failed to extend to host, *hostname*. Failed command: *command_name*.

Explanation: This message occurs when the GPFS cluster extension to the host fails.

User response: Refer to the db2cluster command log file. This log file can be found in /tmp/ibm.db2.cluster.*. Additional error messages and details are contained in the log file. After the problem is solved, rerun the command to extend the host.

DBI20031E The GPFS cluster, *db2cluster_xxxx*, failed to remove the file system *fs_name*. Failed command: *command_name*.

Explanation: This message occurs when the newly created GPFS file system fails to be deleted during a rollback.

User response: Refer to the db2cluster command log file. This log file can be found in /tmp/ibm.db2.cluster.*. Additional error messages and details are contained in the log file. After the problem is solved, rerun the command to remove the file system.

DBI20032E The host *host_name* did not finish validation within the time out period *time_period*. To change the time out time, set the new time out value (in minutes) using environment variable *environment-variable-name*.

Explanation: This error occurs because a validation process on the remote machine did not return within the required time out period.

User response: If your machine is slow, you can extend the time out time to give it more time to complete. To change the time out time, set the new time out value (in minutes) using environment variable "environmentvariable1". Rerun the installation command.

If the machine is hanging, collect all information files and contact IBM service.

DBI20033E The host *host_name* cannot be removed from the GPFS cluster *db2cluster_xxxx*.
Failed command: *command_name*.

Explanation: This message occurs when the GPFS cluster fails to remove the host.

User response: Refer to the db2cluster command log file. This log file can be found in /tmp/ibm.db2.cluster.*. Additional error messages and details are contained in the log file. After the problem is solved, rerun the command to remove the host from GPFS cluster.

DBI20034I No action is required, there is no user managed cluster to takeover.

Explanation: The db2cluster_prepare command did not detect a IBM General Parallel File System (GPFS). Running the db2cluster_prepare command is not required. When creating a DB2 pureScale instance, a GPFS file system will be created for you.

User response: None.

DBI20035E The PVID for device path *device_path* on the install-initiating host (IIH) *host_name* does not exist on remote host *remote_host_name*.

Explanation: This message occurs because there is no device path that can be found on the remote host with the same PVID as the given device path on the IIH. A matching PVID is required on the remote host.

User response: To check if the PVID is available on all hosts:

1. Get the PVID of the given device path in the IIH machine.
2. Make sure the same PVID exists on all hosts.

Configure the PVID on the remote hosts. For details on configuring the PVID, see the DB2 Information Center.

DBI20036E The remote host *host_name* belongs to an active IBM Reliable Scalable Cluster Technology (RSCT) peer domain which was previously created by the DB2 pureScale Feature.

Explanation: This error occurs when the DB2 installer detects that an RSCT peer domain was previously created by the DB2 pureScale Feature.

User response: To continue, perform one of the following:

- Remove the remote host from the host list, or,
- If the remote specified host belongs to a DB2 pureScale instance, remove it from that DB2 pureScale instance using the db2iupdt -drop command.

For details about the db2iupdt command, see the DB2 Information Center.

DBI20037E The remote host *host_name* belongs to an active IBM Reliable Scalable Cluster Technology (RSCT) Peer Domain which was previously created by the DB2 pureScale Feature, but the installation-initiating host is in a different RSCT peer domain.

Explanation: This error occurs when the DB2 installer detects that an RSCT peer domain was previously created by the DB2 pureScale Feature but the installation-initiating host is in a different RSCT peer domain.

User response: Perform one of the following actions:

- Remove the remote host from the host list.
- If the remote host belongs to a DB2 pureScale instance, remove that host from that DB2 pureScale instance by calling the db2iupdt command specifying the -drop parameter.

DBI20038E The remote host *host_name* belongs to a GPFS cluster.

Explanation: The DB2 product does not allow the joining of multiple GPFS clusters. The current host and the remote host indicated belongs to different GPFS clusters.

When the install-initiating host (IIH) is part of a GPFS cluster, the remote host should either not belong to a GPFS cluster, or, belong to the same GPFS cluster. When IIH is not part of a GPFS cluster, the remote host should not belong to a GPFS cluster.

User response: To resolve this error, perform one of the suggestions below and reissue the command:

- Remove the remote host from the host list.
- Prior to adding the host to the current DB2 pureScale setup, remove the host from the other GPFS cluster.
- For an instance creation, to reuse the GPFS cluster defined on the remote host, start the instance creation on the remote host.

Reissue the command.

DBI20039E The new copy, *new-copy-name*, must be installed to another path. The installation path specified, *new-copy-path*, contains an existing DB2 product with the DB2 copy name *existing-copy-name*.

Explanation: DB2 requires that new copies are installed to a path which does not contain an existing copy.

User response: To install the new copy, specify a new path. To work with the DB2 product installed at this

location, exit the installation and select the 'Work with Existing' option from the DB2 Setup Launchpad.

DBI20040E The path specified in the response file by the keyword **FILE**, *provided-path*, contains an installed DB2 copy with **DB2_COPY_NAME** *existing-copy-name* that does not match the copy name provided in the response file, *provided-copy-name*.

Explanation: The path and name pair provided in the response file by keywords **FILE** and **DB2_COPY_NAME** is invalid. The path specified contains an existing DB2 copy with a copy name that is different from the copy name provided. A path may only contain a single DB2 copy.

User response: Ensure the correct copy name and path are provided and run the command again.

DBI20041E You cannot drop host *host_name* from the instance *instance_name* because host *host_name* is the installation-initiating host (IIH). Host *host_name* can be dropped from another host in the instance.

Explanation: The `db2iupdt -drop` command will not drop the last cluster caching facility or the last member, in the DB2 pureScale instance. You must run the `db2iupdt -drop` command from a host that will still belong to the instance after you've dropped the cluster caching facility or member.

User response: The specified host can be dropped from another host in the instance using the `db2iupdt -drop` command.

DBI20042E The response file keyword **DB2_COPY_NAME** value, *provided-copy-name*, is invalid for the path, *provided-copy-path*, specified by the keyword **FILE**. That copy exists on path *existing-copy-path*.

Explanation: The **DB2_COPY_NAME** provided exists, but not on the path provided. When working with an existing copy of DB2 installed on Windows, the correct installation path must be specified.

User response: To perform a new installation, specify a unique value for **DB2_COPY_NAME**. To work with an existing copy, ensure the keyword **DB2_COPY_NAME** corresponds to the correct path specified by keyword **FILE**.

DBI20043E DB2 failed to install because of restrictive permission settings on the temporary directory, *tmp-directory*.

Explanation: The temporary directory is used by DB2 to store files necessary for creating and managing instances and database objects. To install or update DB2 the temporary directory must be enabled for read, write and execute access to all users.

User response: Enable read, write and execute permission on the temporary directory for all users. On Linux and UNIX this can be accomplished with the command `chmod 0777`. In a DB2 pureScale environment, be sure all hosts have the proper permission settings.

DBI20044E The parameter *parameter-name* is missing or invalid.

Explanation: The specified parameter is either missing, or invalid. Some parameters are mandatory in a response file for the specified command, however, certain combinations are not supported. For example, the **PREFERRED_PRIMARY_CF** parameter, or the **PREFERRED_SECONDARY_CF** parameter, can only be specified for a instance type of 'dsf'.

User response: Check the response file for the specified parameter. For more details on response files, and parameter combinations for the command, see the DB2 Information Center.

DBI20045E DB2 failed to create a DB2 pureScale instance because the detected operating system is not supported.

Explanation: The DB2 pureScale Feature for Linux is only supported on specific distributions.

User Action Consult the documentation for installing the DB2 pureScale Feature for information on software considerations in the installation prerequisites.

User response:

DBI20046E The DB2 pureScale instance could not be created because the uDAPL configuration file could not be accessed.

Explanation: The user Direct Access Programming Library (uDAPL) configuration file could not be accessed while attempting to create the DB2 pureScale instance. The DAPL configuration file is referenced by the DB2 pureScale Feature to verify the presence of a remote direct memory access (RDMA) capable network adapter. The file must be accessible on the path `/etc/dat.conf`.

User response: If a backup copy of the file exists, restore it to the original location.

DBI20047E The DB2 installation process failed on the following host or hosts: *host-list*.

Explanation: All software components of the DB2 pureScale Feature are installed and configured from a single host. If the installation of the feature fails on any one of the host machines, the DB2 installation process cannot complete.

There are several reasons why the installation process might have failed on a host machine:

- The host name cannot be found.
- The host is unreachable.
- Someone other than the root user is performing the installation.
- Your system does not meet all uDAPL or InfiniBand prerequisites.
- There is not enough disk space available in the /root or /tmp path to complete the installation.

User response: To help you identify and resolve the reason for the installation failure you can collect diagnostic data by using the db2support tool provided with the install image.

To collect diagnostic data from the affected host or hosts, run the 'db2support -install' command with the -host option. For more information about troubleshooting DB2 installation problems with the db2support -install command, see the DB2 Information Center.

DBI20048E The DB2 installation failed because an adapter could not be found to link to the cluster interconnect on one or more hosts. Hosts: *host-names*.

Explanation: All hosts in a DB2 pureScale environment must be connected to a low latency, high-speed interconnect.

User response: Ensure your environment meets the installation prerequisites listed in the DB2 Information Center and then rerun the installation.

DBI20049C The installFixPack command failed. The base DB2 copy contains DB2 pureScale instances, however the DB2 installer cannot detect the hosts in these instances. To install the fix pack on the local host only, rerun the installFixPack command with the -L parameter. To install the fix pack on all hosts, rerun the installFixPack command with the -H parameter specifying the host list file.

Explanation: The DB2 installer cannot detect which hosts require the fix pack to be applied.

User response: Rerun the command specifying to apply the fix pack to either the local host or to all hosts.

DBI20050I When the installFixPack command is run with the -p parameter, the instances are not updated automatically. You must manually update the instances using the db2iupdt command. If the fix pack is applied on multiple hosts, you must run the db2iupdt command on all hosts.

Explanation: The installFixPack -p parameter does not automatically update instances. This must be performed manually.

User response: Run the db2iupdt command as indicated.

DBI20051C The DB2 installation has failed on the remote hosts *host_name*. Review the logged errors then re-run the installFixPack on those machines.

Explanation: The fix pack was not installed successfully. The error log contains details.

User response: Review the error log for additional information.

DBI20052E The fix pack update failed because the DB2 pureScale component is not installed on this host.

Explanation: In environments that are not DB2 pureScale environments, a fix pack cannot be applied across hosts.

In DB2 pureScale environments, to apply fix pack update across multiple hosts, the DB2 copy must have the DB2 pureScale component installed on every host.

User response:

1. Remove the host that does not contain DB2 pureScale component from the host list file.
2. Rerun the installFixPack command.

DBI20053C The service level (version, release, modification level) on this host is different from the service level on the IIH. Current host level: *host_level*. IIH level: *iih_level*. To apply a fix pack update across multiple hosts, the base DB2 copy on all hosts must have the same service level. Apply the fix pack required to bring this host and the IIH to the same service level, and rerun the installFixPack command. Otherwise, rerun the installFixPack command with the -H parameter with the current host removed from the host list file.

Explanation: All hosts must have the same service level to proceed.

User response: Rerun the installFixPack as indicated.

DBI20054C The IHH host is not listed in the specified host list file. IHH host *host_name*.

Explanation: The IHH host must be listed in the host list file provided by the -H parameter.

User response: Add the IHH host to the host list file and rerun the installFixPack command.

DBI20055I An invalid value was entered for the port number when configuring the TCP/IP settings for a DB2 instance.
Value: *port-value*.

Explanation: A DB2 instance can be configured during product installation. TCP/IP port values are one of the settings that can be set during instance configuration. The TCP/IP port number is used by a DB2 instance to listen for incoming DB2 connections. This message is displayed when an invalid invalid port number is entered. This message also appears in the log file created during a silent install operation

User response: Enter a valid port value of between 1024 and 65535 and try again.

DBI20056I Install of *product-name* cannot start because Visual Studio 2008, Visual Studio 2010, or Visual Studio 2012 is not installed.

Explanation: The IBM Database Add-Ins for Visual Studio are included as a separately installable component with the DB2 Client and the DB2 servers. Once you are finished installing your DB2 product, you will be presented with an option to install the IBM Database Add-Ins for Visual Studio. If you choose this option and a supported version of Visual Studio is not installed, the installation fails because Visual Studio is a requirement to install the IBM Database Add-Ins for Visual Studio.

User response: Install Visual Studio 2008 or Visual Studio 2010 and then retry the product install.

DBI20057I *product-name* was only installed on Visual Studio 2008 and not on Visual Studio 2005 because Visual Studio 2005 Service Pack 1 was not installed.

Explanation: The IBM Database Add-Ins for Visual Studio are included as a separately installable component with the DB2 Client and the DB2 servers. Once you are finished installing your DB2 product, you will be presented with an option to install the IBM Database Add-Ins for Visual Studio.

If you have Visual Studio 2008 and Visual Studio 2005 but your Visual Studio 2005 is not at the Service Pack 1 level, then VSAI will only be installed on Visual Studio 2008.

User response: Install Visual Studio 2005 Service Pack 1 and retry the install of VSAI.

DBI20058I The user *user-name* could not be deleted during DB2 uninstall.

Explanation: DB2 can create Windows user accounts during the install process, for example a user account can be create as the instance owner. During uninstall you have the option to delete any Windows user accounts created by DB2 during the install process.

This message is displayed when, during a DB2 product uninstall, the uninstall utility tried to delete a Windows user account but failed.

User response: Delete the user account manually.

DBI20059I The user group *user-group-name* could not be deleted during DB2 uninstall.

Explanation: DB2 can create Windows group accounts during the install process. During uninstall you have the option to delete any Windows group accounts created by DB2 during the install process. This message is displayed when, during a DB2 product uninstall, the uninstall utility tried to delete a group account but failed.

User response: Delete the user group account manually.

DBI20060I The DB2 uninstall utility found users and groups on the system which were created by the DB2 install utility. Users: *user-list*. Groups: *group-list*.

Explanation: DB2 can create Windows user and group accounts during the install process. This message is displayed when, during a DB2 product uninstall, users and group accounts are found. These might be in use by another DB2 copy or by other non-DB2 applications. It is recommended keeping these accounts.

User response: One of the following:

- If you want to keep the users and groups then select Yes
- If you don't want to keep the users and groups then select No.

DBI20061I The DB2 uninstall utility found users on the system which have been created by the DB2 install utility. Users: *user-list*.

Explanation: DB2 can create Windows user accounts during the install process. This message is displayed when, during a DB2 product uninstall, user accounts were found. These might be in use by another DB2 copy or by other non-DB2 applications. It is recommended keeping these accounts.

User response: One of the following:

- If you want to keep the user accounts then select Yes
- If you don't want to keep the user accounts then select No.

DBI20062I **The DB2 uninstall utility found groups on the system which have been created by the DB2 install utility. Groups:** *group-list*.

Explanation: DB2 can create Windows user group accounts during the install process. This message is displayed when, during a DB2 product uninstall, user group accounts were found. These groups may be in use by another DB2 copy or by other non-DB2 applications. It is recommended keeping these accounts.

User response: One of the following:

- If you want to keep the group accounts then select Yes
- If you don't want to keep the group accounts then select No.

DBI20063E **The DB2 installer failed to configure and start the SSH server because GPFS is unavailable. DB2 database manager instance name:** *instance-name*. **Host machines:** *host-machine-name-list*.

Explanation: You can securely run commands from remote machines by using a secure shell (SSH) server product such as IBM Secure Shell Server for Windows or Open Secure Shell (OpenSSH) on Linux and AIX operating systems.

This message is returned when the DB2 installer is configuring a SSH server in an environment with IBM General Parallel File System (GPFS) and the installer is unable to complete that SSH configuration on one or more host machines because GPFS is offline.

User response:

1. On each of the host machines listed, restart GPFS manually.
2. Run the DB2 install operation again.

DBI20064E **DB2 copy upgrade is not supported from DB2 Version** *source-version* **to DB2 Version** *target-version*.

Explanation: A DB2 copy refers to one or more installations of DB2 database products in a particular location on the same computer. If you attempt to upgrade from an unsupported DB2 version, the DB2 copy upgrade will fail.

User response: Attempt the DB2 copy upgrade again specifying source and target DB2 versions that have supported upgrade paths.

DBI20065I **No action is required, the cluster is managed by DB2 cluster services.**

Explanation: Running the db2cluster_prepare command to takeover the cluster was unnecessary. DB2 cluster services are managing the cluster.

User response: None.

DBI20066E **The GPFS cluster does not meet the requirements for a DB2 pureScale instance. Cluster name:** *cluster*. **Failed command:** *command-name*.

Explanation: While running the command, the GPFS cluster failed a verification test of the requirements for the DB2 pureScale Feature.

User response: For information on setting up a GPFS cluster, see the DB2 Information Center.

DBI20067E **The operation failed because the DB2 global profile registry variable could not be updated on all hosts. Global registry variable:** *variable-name*. **Hosts not updated:** *comma-separated-host-list*.

Explanation: The DB2 global registry file on the identified hosts could not be modified during the operation.

User response: Make sure that the DB2 global profile registry can be accessed, then rerun the command.

DBI20068E **The command failed to update the DB2 global profile registry variable on one or more hosts. Registry variable:** *variable-name*. **Hosts not updated:** *comma-separated-host-list*. **Command to run on hosts:** *command-name*.

Explanation: The DB2 global registry file on the identified hosts could not be modified during the operation.

User response: Make sure that the DB2 global-profile registry can be accessed on the identified hosts, then rerun the command as specified on each host to update the DB2 global profile registry file with the variable record.

DBI20069E **Communication could not be established between the two hosts. Source host:** *host-name1*. **Target host:** *host-2*.

Explanation: All hosts in a DB2 pureScale instance require passwordless SSH access to all other hosts participating in the instance.

User response: Verify that the target machine is online and that passwordless SSH access from the source host has been configured.

DBI20070E An error occurred running the command because a required component is not installed. Missing component: *component-name*.

Explanation: The missing component is a prerequisite to running the command.

User response: Install the component on the host, then rerun the command.

DBI20071I The cluster takeover has completed successfully.

Explanation: The user-managed cluster is now managed by DB2 cluster services.

User response: None.

DBI20072W The DB2 cluster services tiebreaker disk validation failed. The tiebreaker specified has not been configured.

Explanation: The DB2 cluster services tiebreaker disk helps a subcluster obtain operational quorum in the event of a failure that leaves exactly half of the DB2 pureScale instance hosts able to communicate with each other. In this event, the subcluster will attempt to get an exclusive lock on the disk to achieve operational quorum. The DB2 pureScale instance can operate without a configured tiebreaker disk, however the instance is more resilient with one.

User response: For more information about DB2 cluster services tiebreaker support, see the shared storage considerations in the DB2 Information Center.

DBI20073E The specified command failed because the value specified for a keyword in the response file was invalid. Value: *keyword-value*. **Keyword:** *keyword-name*.

Explanation: All response file keywords must be applicable to the installation, and all keyword values must be valid. Some commands and keywords are only valid when run by the instance owner, or when the installation includes specific components, features, or configurations.

User response: Edit the response file to correct the reported error. The DB2 Setup wizard can be used to create a valid response file for future use. If you have an existing installation, the response file generator can be used to create a response file.

DBI20074E The command cannot be run due to insufficient authority to use the specified keyword. **Keyword:** *keyword-name*.

Explanation: Administrative or root authority is required to use the specified keyword when performing a response file installation.

User response: Retry installation with a user that has sufficient authority, or without the problematic keyword. Contact the system administrator if you require the response file installation to run with the keyword.

DBI20075E The DB2 installer failed to configure the IBM Secure Shell Server for Windows service to start automatically because the default port is in use. Default port number: *port_number*.

Explanation: The installation cannot configure IBM Secure Shell Server to automatically start on this port.

User response: Select a different port to manually install the service, stop the service currently running on the port before starting the IBM Secure Shell Server.

DBI20076E The directory specified for the IBM Secure Shell Server cannot be used.

Explanation: The IBM Secure Shell Server installation directory cannot be the same as a DB2 installation directory.

User response: Specify a directory different from a DB2 installation directory.

DBI20077E Updating the IBM Secure Shell Server for Windows failed during a fix pack installation.

Explanation: The DB2 installer detected an error with the version of the IBM Secure Shell Server binaries. The fix pack installation cannot continue.

User response: Ensure that the IBM Secure Shell Server files already installed in the system are not at a higher level than the fix pack image.

DBI20078W Because no value was specified for the response file keyword **AUTOSTART_SSH_SERVER**, the DB2 installer attempted to configure the IBM Secure Shell (SSH) Server for Windows service to start automatically on the default port. Default port number: *port-number*. However, that configuration failed because the default port is already in use. The service for IBM SSH Server will need to be started manually.

Explanation: An error occurred during the uninstallation of the IBM Secure Shell Server. IBM Secure Shell Server was not uninstalled.

User response: Login as an administrator user and manually remove the IBM Secure Shell Server.

DBI20079E **Removal of the IBM Secure Shell Server for Windows failed during the DB2 uninstall process.**

Explanation: An error occurred during the uninstallation of the IBM Secure Shell Server. IBM Secure Shell Server was not uninstalled.

User response: Login as an administrator user and manually remove the IBM Secure Shell Server.

DBI20080E **During uninstall, the DB2 installer detected that both a 32-bit version and a 64-bit version of the IBM Data Server Driver Package (DSDRIVER) are installed. 32-bit and 64-bit versions of the DSDRIVER product cannot coexist.**

Explanation: A 32-bit and 64-bit version of a DB2 product cannot coexist on a host machine.

User response: Uninstall the conflicting product before attempting to reinstall. You must manually uninstall one of the DSDRIVER products using the db2unins command. For db2unins command details, see the DB2 Information Center.

DBI20081E **There is not enough free disk space in *directory_name*. Free space detected in *directory_name* is *space_available*, while space needed is *space_needed*. Free up additional space and try again.**

Explanation: There is not enough free space available.

User response: Free more disk space and retry the command.

DBI20082E **The DB2 installer cannot create necessary database manager instance files because the DB2 installer does not have write permission for the specified shared disk partition device path. Specified shared disk partition device path: *shared-disk-path*.**

Explanation: In a pureScale environment, a shared disk path is required to hold shared instance and database files. When a DB2 database manager instance is created by using the db2icrt command, the shared disk path is specified by using the `-instance_shared_dev` parameter.

This message is returned when the DB2 installer does not have write permission for the specified shared disk path.

User response: Respond to this message in one of the following ways:

- Modify the permissions of the directory that is to be the specified shared disk path so that the DB2 database installer has write access to the path.

- Call the db2icrt command again, specifying a shared disk path with the `-instance_shared_dev` parameter for which the DB2 installer has write access.
-

DBI20083E **The DB2 pureScale instance could not be created because the DB2 installer found more than one uDAPL configuration file.**

Explanation: The user Direct Access Programming Library (uDAPL) configuration file, `dat.conf`, is referenced by the DB2 pureScale Feature to verify the presence of a remote direct memory access (RDMA) capable network adapter. Exactly one copy of the file must be accessible under the appropriate subdirectory of `/etc`.

This message is returned when the DB2 installer found more than one of these uDAPL configuration files. If multiple uDAPL configuration files exist, the DB2 installer cannot determine which set of configurations to use.

User response: Verify the uDAPL configuration manually, including verifying that only one copy of the `dat.conf` file exists on the system, and then run the instance creation operation again.

DBI20084E **The domain that was specified for the DB2 administrators group is not the same as the domain that was specified for the DB2 users group.**

Explanation: Two Windows groups are used by the DB2 database manager to enforce extended Windows security: the DB2 administrators group, which is named "DB2ADMINS" by default; and the DB2 users group, which is named "DB2USERS" by default. These groups are used by the DB2 database manager to provide protection at the operating system level.

The DB2 administrators group and the DB2 users group may both be local groups or they may both be domain groups. If the DB2 administrators group and the DB2 users group are domain groups, they must be in the same domain as each other.

User response: Specify the same domain for both the DB2 administrators group and the DB2 users group.

DBI20085E **The instance update operation for the instance named *instance-name* failed because some resources for other instances were not completely removed during previous instance drop operations.**

Explanation: You can update a DB2 database manager instance to a higher level within a release, update an instance other than a DB2 pureScale instance to a DB2 pureScale instance, or scale a DB2 pureScale instance by using the db2iupdt command.

This message is returned when an attempt is made to perform an instance update operation on a given instance, but the update operation cannot proceed because there is a conflict with resources that were not completely removed during a previous attempt to drop one or more instances.

User response:

1. Identify any entries in the global registry that were associated with instances that no longer exist.
2. Remove any entries in the global registry that were associated with instances that no longer exist by using the db2greg command.
3. Run the db2iupdt command again.

DBI20086E The update command failed because the specified CF or DB2 member is not present in the DB2 cluster. Specified CF or member: *id*.

Explanation: You can update a DB2 database manager instance to a higher level within a release, update an instance other than a DB2 pureScale instance to a DB2 pureScale instance, or scale a DB2 pureScale instance by using the db2iupdt command.

This message can be returned in two scenarios:

- The db2iupdt command is called with the -update parameter and a cluster caching facility (CF) that is not part of the DB2 cluster is specified with the -cf parameter
- The db2iupdt command is called with the -update parameter and a DB2 member that is not part of the DB2 cluster is specified with the -m parameter

User response: Respond to this error in one of the following ways:

- To update an existing CF in the DB2 cluster, call the db2iupdt command again specifying a CF that is present in the DB2 cluster.
- To add a new CF to the DB2 cluster, call the db2iupdt command again specifying the -add and -cf parameters.
- To update an existing member in the DB2 cluster, call the db2iupdt command again specifying a member that is present in the DB2 cluster.
- To add a new member to the DB2 cluster, call the db2iupdt command again specifying the -add and -m parameters.

DBI20087E The db2dsdriver.cfg file relocation failed because the file or folder is read-only.

Explanation: The db2dsdriver.cfg file was not moved to the destination directory. The following directory is not accessible:

- On Windows XP Professional and Windows Server 2003 the path is:

C:\Documents and Settings\
All Users\Application Data\IBM\DB2\
<driver-copy-name>\cfg

- On Windows Vista, Windows 7, and Windows Server 2008 the path is:

C:\ProgramData\IBM\DB2\<driver-copy-name>\cfg

User response: Modify the path so that the current user has read and write access to the file and folder. Then perform the silent installation again.

DBI20088E Installing or updating SSH utilities failed.

Explanation: The DB2 pureScale Feature uses General Parallel File System (GPFS) as the shared-disk file system to provide data access across all members in a DB2 pureScale environment. GPFS requires certain commands to run as root on all hosts that are part of the GPFS domain. There are two ways to configure your environment to meet this GPFS requirement:

- You can enable passwordless remote root login using standard SSH public key authentication.
- In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, passwordless remote root login is no longer required. You can use the non-root user ID (db2sshid) to use a secure shell (SSH) network protocol between hosts.

This message is returned when the DB2 installer fails to install the DB2 utilities that support using the db2sshid.

User response: Run the install operation again.

DBI20089E The db2dsdriver.cfg file relocation failed due to an unknown keyword in the response file. Keyword: *keyword*. Line: *line-number*.

Explanation: The response file used during silent installation has detected an unknown keyword.

User response:

1. Fix the error in the response file.
2. Attempt silent installation again.

DBI20090W The DB2 installer detected that the db2ssh directory is world writeable.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshid) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login. There are several DB2 SSH utilities located in the db2ssh directory to support using the db2sshid.

This message is returned when the DB2 installer detects that the db2ssh directory is w-writable. The

permissions of the db2ssh directory should allow only root users write access to the directory.

User response: Modify the access permissions for the db2ssh directory so that only root users may access those directories.

DBI20091W The DB2 installer failed to change the ownership of the db2locssh script to the db2sshd. The ownership of the db2locssh script has been set to the instance owner ID.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login.

One of the DB2 utilities that supports using the db2sshd is the db2locssh script. During install, the DB2 installer attempts to use operating system commands to set the owner of the db2locssh script to the db2sshd. However, if the installer is unable to set the owner of the db2locssh script to the db2sshd the installer will set the owner of the db2locssh script to the current instance owner ID.

This message is returned when the DB2 installer sets the owner of the db2locssh script to the current instance owner ID because the installer failed to set the owner of the db2locssh script to the db2sshd.

User response: Manually change the ownership of the db2locssh script to the db2sshd.

DBI20092W The DB2 installer detected a problem with the db2locssh utility that will prevent the use of the db2sshd.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login.

One of the DB2 utilities that supports using the db2sshd is the db2locssh script. During install, the DB2 installer verifies that the db2locssh script is functioning correctly.

This message is returned when the DB2 installer is unable to successfully verify the db2locssh functionality.

User response: Manually set up the db2sshd user ID.

DBI20093E The DB2 installer failed to set up the db2sshd because the GPFS file system is a user-managed file system.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login. Using the db2sshd is not supported with user-managed GPFS file systems.

This message is returned when an attempt is made to enable the db2sshd with a user-managed GPFS file system.

User response: Change the user-managed GPFS file system to be a DB2 managed GPFS file system.

DBI20094E The db2icrt command failed because the specified db2sshd does not exist on all hosts in the DB2 cluster. Specified db2sshd: *specified-id*.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login. To use the db2sshd, you need to create a user id on all hosts in the DB2 cluster, including the installation-initiating host, and then pass that user id to the db2icrt command with the -i parameter.

This message is returned when an attempt is made to specify a user id as a db2sshd but the specified user id does not exist on all hosts in the cluster.

User response: Call the db2icrt command again, specifying a user id that exists on all hosts in the DB2 cluster.

DBI20095E The distribution of the SSH public key failed because communication failed between two or more of the hosts in the DB2 cluster.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login. For the db2sshd to be used, a public key must be distributed to all hosts in the DB2 cluster.

This message is returned what the DB2 installer fails to distribute the public key to all of the hosts in the cluster during an upgrade or update operation.

User response: Manually distribute the public key of individual hosts to all hosts in the cluster.

DBI20096E The SSH configuration failed for the instance named *instance-name* on the following hosts in the cluster: *host-name-list*.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login. For the db2sshd to be used, a public key must be distributed to all hosts in the DB2 cluster.

This message is returned what the DB2 installer fails to create the public key during instance creation or or update operation.

User response: Manually distribute the public key of individual hosts to all hosts in the cluster.

DBI20097E The SSH configuration failed for the db2sshd *user-id* on the following hosts in the cluster: *host-name-list*.

Explanation: In DB2 Version 10.1 Fix Pack 2 and later fix packs, if you are using a DB2 managed GPFS file system for new DB2 pureScale installations, you can use the non-root user ID (db2sshd) to use a secure shell (SSH) network protocol between hosts instead of using passwordless remote root login. For the db2sshd to be used, a public key must be distributed to all hosts in the DB2 cluster.

This message is returned what the DB2 installer fails to create the public key during instance creation or or update operation.

User response: Manually distribute the public key of individual hosts to all hosts in the cluster.

DBI20098E The db2iupdt -add -m command failed to add a member because the CF is running on the host to be added.

Explanation: In an existing DB2 pureScale cluster, if the cluster caching facility (CF) is running, a member cannot be added to the CF host. Before a member can be added, the CF must be stopped using the db2stop command.

User response: On the host where you are adding the member, stop that CF using the db2stop command indicating the CF identifier.

DBI20099E The db2unins command failed because the command was run from an installed copy of the DB2 database product and the -f parameter, which can only be used when the command is run from the installation media, was specified.

Explanation: You can uninstall one or more DB2

database products, features, or languages by using the db2unins utility.

You can run the db2unins command either from the installed copy of DB2 database or from the DB2 database installation media:

- Some db2unins parameters are only supported when the command is run from an installed copy of DB2 database.
- Some db2unins command parameters are only supported when the command is run from the installation media.

You can perform a forced uninstallation of all DB2 database products on the system by specifying the -f parameter with the db2unins command. The -f parameter is one of the parameters that can only be used when the db2unins command is run from the installation media.

This message is returned when the command is run from an installed copy of DB2 database and the -f parameter is specified. The command fails because the -f parameter cannot be specified when the command is run from an installed copy of DB2 database.

User response: Respond to this error in one of the following ways:

- To perform a forced uninstallation of all DB2 database products on the system, run the db2unins command from the installation media, specifying the -f parameter.
- To use the db2unins command from the installed copy of DB2 database, call the db2unins command again without specifying the -f parameter.

DBI20100E The upgrade operation failed because upgrading the IBM Data Server Driver Package software from the current version to the new version is not supported. IBM Data Server Driver Package copy name: *copy-name*. Current version: *version-number*. New version: *version-number*.

Explanation: The IBM Data Server Driver Package software simplifies application deployment. This driver, which has a small footprint, is designed to be used by independent software vendors (ISVs) for mass application deployment scenarios that are typical of large enterprises.

This message is returned when an attempt is made to upgrade the version of the IBM Data Server Driver Package software, and updating from the current version to the new version is not supported.

User response: Respond to this error in one of the following ways:

- Install a new copy of the IBM Data Server Driver Package software.

- Upgrade to DB2 Version 9.7, and then upgrade from DB2 Version 9.7 to the latest release.
- Uninstall the current IBM Data Server Driver Package software and then install the new IBM Data Server Driver Package software.

DBI20101E The upgrade operation has halted because upgrading the IBM Database Add-Ins for Visual Studio software from the current version to the new version is not supported. Product name: *product-name*. Current version: *version-number*. New version: *version-number*.

Explanation: The IBM Database Add-Ins for Visual Studio software is a collection of features that integrate seamlessly into your Visual Studio development environment so that you can work with DB2 servers and develop DB2 procedures, functions, and objects.

This message is returned when an attempt is made to upgrade IBM Database Add-Ins for Visual Studio, and updating from the current version to the new version is not supported.

User response: Uninstall the current IBM Database Add-Ins for Visual Studio software and then install the new IBM Database Add-Ins for Visual Studio software.

DBI20102E The upgrade operation has halted and cannot proceed further because upgrading the IBM Data Server Driver Package software from the current version to the new version is not supported. IBM Data Server Driver Package copy name: *copy-name*. Current version: *version-number*. New version: *version-number*.

Explanation: The IBM Data Server Driver Package software simplifies application deployment. This driver, which has a small footprint, is designed to be used by independent software vendors (ISVs) for mass application deployment scenarios that are typical of large enterprises.

This message is returned when an attempt is made to upgrade the version of the IBM Data Server Driver Package software, and updating from the current version to the new version is not supported.

User response: Respond to this error in one of the following ways:

- Install a new copy of the IBM Data Server Driver Package software.
- Upgrade to DB2 Version 9.7, and then upgrade from DB2 Version 9.7 to the latest release.
- Uninstall the current IBM Data Server Driver Package software and then install the new IBM Data Server Driver Package software.

DBI20103E The upgrade operation has halted and cannot proceed further because upgrading the IBM Database Add-Ins for Visual Studio software from the current version to the new version is not supported. Product name: *product-name*. Current version: *version-number*. New version: *version-number*.

Explanation: The IBM Database Add-Ins for Visual Studio software is a collection of features that integrate seamlessly into your Visual Studio development environment so that you can work with DB2 servers and develop DB2 database objects.

This message is returned when an attempt is made to upgrade IBM Database Add-Ins for Visual Studio, and updating from the current version to the new version is not supported.

User response: Uninstall the current IBM Database Add-Ins for Visual Studio software and then install the new IBM Database Add-Ins for Visual Studio software.

DBI20104E The installFixPack command failed because invalid or conflicting parameters were specified.

Explanation: You can update installed DB2 database products to apply a fix pack by using the installFixPack command.

This message is returned when there is a problem with one or more of the specified command parameters, including the following example problems:

- A parameter was specified that is not supported with the installFixPack command
- A parameter was specified that is not valid for the operating environment or for the type of DB2 instance
- Conflicting parameters, parameters that are not supported together, were specified

User response:

1. Review the installFixPack command syntax in the documentation or by running the following command:
installFixPack -h
2. Call the installFixPack command again, specifying only valid parameters.

DBI20105E An error occurred while installing the following file set: *file-package-name*. Because these files were not successfully installed, functionality that depends on these files might not work as expected.

Explanation: The DB2 installer installs files related to DB2 database functionality as well as files for software that is bundled with the DB2 database software.

This message is returned when the DB2 installer

encounters a problem copying some files from the install media to the install location.

User response: Contact IBM Software Support for assistance.

DBI20106E Installation of the DB2 database product was not performed. The installer package DB2Server.msi cannot be invoked directly. The DB2 database product must be installed by using the setup command.

Explanation: On Windows operating systems, you can launch the DB2 Setup wizard to install DB2 database products. You can perform a response file installation by specifying the -u parameter with the setup command.

This message is returned when an attempt is made to use the DB2Server.msi installation package directly, which is not supported.

User response: Run the setup command.

DBI20107E The db2_deinstall command failed because the DB2 installer detected that the current host is a member of a GPFS cluster that includes other, remote host machines.

Explanation: You can remove DB2 database products or DB2 database components using the db2_deinstall command.

This message is returned when an attempt is made to uninstall DB2 database where the current host is part of an IBM General Parallel File System (GPFS) cluster that includes remote hosts.

User response: Manually remove the GPFS cluster and file system, and then run the db2_deinstall command again.

DBI20108W Validation of the RSCT peer domain failed. However, the validation failure is being ignored and the add member operation will proceed.

Explanation: When a new DB2 member is being added to a DB2 pureScale instance, it is DB2 install and configuration utilities that add the new member. While the DB2 install utility processes an add member request, the utility validates the IBM Reliable Scalable Cluster Technology (RSCT) peer domain.

This warning message is printed when the DB2 install utility detects a problem with the RSCT peer domain while the utility is processing an add member request. (The problem with the RSCT peer domain configuration could have been caused by an error during a previous add member operation, for example.)

This warning message is printed when the RSCT peer

domain validation failure does not cause the current add member operation to fail. This warning is printed for your information only.

User response: Unless the DB2 install utility encounters other errors while processing the add member request, no response is required to this warning message.

DBI20109W Validation of the instance home directory structure failed. However, the validation failure is being ignored and the add member operation will proceed.

Explanation: When a new DB2 member is being added to a DB2 pureScale instance, it is DB2 install and configuration utilities that add the new member. While the DB2 install utility processes an add member request, the utility validates several aspects of the environment and of the DB2 instance itself.

This warning message is printed when the DB2 install utility detects a problem with the instance home directory while the utility is processing an add member request. (The problem with the instance home directory could have been caused by an error during a previous add member operation, for example.)

This warning message is printed when the directory structure validation failure does not cause the current add member operation to fail. This warning is printed for your information only.

User response: Unless the DB2 install utility encounters other errors while processing the add member operation, no response is required to this warning message.

DBI20110W The DB2 install utility encountered an internal error while processing an add member operation. However, the validation failure is being ignored and the add member operation will be retried automatically.

Explanation: When a new DB2 member is being added to a DB2 pureScale instance, it is DB2 install and configuration utilities that add the new member. To accomplish the add member operation, the DB2 install utility performs many tasks.

This warning message is printed when the DB2 install utility encounters a recoverable error while processing an add member request. (The problem could have been caused by an error during a previous add member operation, for example.)

This warning message is printed when the failure does not cause the current add member operation to fail. This warning is printed for your information only.

User response: Unless the DB2 install utility encounters other errors while processing the add

member operation, no response is required to this warning message.

DBI20111E **The installation operation failed because one of the following invalid keywords was specified in the response file: ADVANCED_ENTERPRISE_SERVER_EDITION, ENTERPRISE_SERVER_EDITION, or WORKGROUP_SERVER_EDITION.**

Explanation: You can install DB2 database software without any user interaction by using a response file.

A response file is an English-only text file that contains setup and configuration information in the form of keywords and keyword-value pairs.

This message is returned when an attempt is made to perform an installation operation with a response file that includes an invalid keyword.

User response: In your response file, replace instances of the following keywords with the keyword DB2_SERVER_EDITION:

- ADVANCED_ENTERPRISE_SERVER_EDITION
- ENTERPRISE_SERVER_EDITION
- WORKGROUP_SERVER_EDITION

Part 13. DBT Messages

Chapter 87. DBT1000 - DBT1499

DBT1000I The tool completed successfully.

Explanation: The tool's processing completed without any errors.

User response: No further action is required.

DBT1001N The syntax of the DB2INIDB tool is incorrect.

Explanation: The DB2INIDB tool has the following syntax:

```
db2inidb <database_alias>
AS < SNAPSHOT | STANDBY | MIRROR >
[ RELOCATE USING config_file ]
```

User response: Resubmit the command using the correct syntax.

DBT1002N Database name *database-name* is invalid.

Explanation: The database name specified in the command is not valid. The database name must contain 1 to 8 characters and all of the characters must be from the database manager base character set.

User response: Resubmit the command using a valid database name.

DBT1003N Unable to find program *program-name*.

Explanation: An attempt to execute the specified program failed because it could not be found.

User response: Ensure that the specified program exists and verify that the program's path exists in the PATH environment variable. Correct the problem and resubmit the command.

DBT1004N Unable to execute program *program-name*.

Explanation: An error occurred while attempting to execute the specified program.

User response: Ensure that the specified program exists and that it contains the appropriate file permissions. Correct the problem and resubmit the command.

DBT1005N The file *file-name1* could not be copied to *file-name2*.

Explanation: An error occurred when attempting to duplicate the specified file. The command could not be processed successfully.

User response: Ensure that the specified source file

exists and that it contains the appropriate file permissions, and ensure that the specified target file does not already exist. Correct the problem and resubmit the command.

DBT1006N The *file-device-name* file or device could not be opened.

Explanation: An error occurred during an attempt to open the specified file or device.

User response:

1. Ensure that the file or device exists and that its permissions are correct.
2. Reissue the command.

DBT1007N An error occurred during an attempt to perform the *operation-name* operation on the *file-device-name* file or device.

Explanation: An attempt to perform the specified operation on the specified file or device was unsuccessful.

User response:

1. Correct the problem. Possible solutions include increasing the disk space or correcting the file permissions.
2. Reissue the command.
3. If the problem persists, contact your technical service representative.

DBT1008N Database *database-name* is not a split mirror image.

Explanation: An attempt was made to use the DB2INIDB tool on a database that is not a split mirror image.

A split mirror image is a mirrored copy of a database that is taken while I/O writes are suspended. I/O writes can be suspended using the following command while being connected to the database:

```
SET WRITE SUSPEND FOR DATABASE
```

Once the split mirror image has been taken, I/O writes on the primary database can be resumed with the following command:

```
SET WRITE RESUME FOR DATABASE
```

Before the split mirror image can be used, the DB2INIDB tool must be run in order to remove its suspended state and initialize it by either performing crash recovery or placing it in rollforward pending state.

User response: Resubmit the command using a split mirror image database.

DBT1009N Database *database-name* is not a recoverable database.

Explanation: An attempt was made to use the STANDBY or MIRROR option of the DB2INIDB tool with a non-recoverable split mirror image database. The STANDBY and MIRROR options require the split mirror image to be recoverable.

A split mirror image is a mirrored copy of a database that is taken while I/O writes are suspended. I/O writes can be suspended using the following command while being connected to the database:

```
SET WRITE SUSPEND FOR DATABASE
```

Once the split mirror image has been taken, I/O writes on the primary database can be resumed with the following command:

```
SET WRITE RESUME FOR DATABASE
```

Before the split mirror image can be used, the DB2INIDB tool must be run in order to remove its suspended state and initialize it by either performing crash recovery or placing it in rollforward pending state.

The purpose of the STANDBY and MIRROR options are to place the split mirror image in rollforward pending state to permit logs from the primary database to be applied. The primary database must retain the log records for them to be available to the split mirror image database.

User response: Resubmit the command using a split mirror image of a recoverable database. To enable the database for rollforward recovery, set the database configuration parameter **logarchmeth1** or **logarchmeth2** to a value other than OFF, disconnect all existing applications from the database, and perform an offline backup of the database.

DBT1010N The split mirror image database *database-name* has been previously initialized as a snapshot.

Explanation: An attempt was made to use the STANDBY or MIRROR option of the DB2INIDB tool on a split mirror image that was previously initialized with the SNAPSHOT option. The STANDBY and MIRROR options can no longer be used to rollforward the log files from the primary database because the SNAPSHOT option caused this split mirror image to start a new log chain.

A split mirror image is a mirrored copy of a database that is taken while I/O writes are suspended. I/O writes can be suspended using the following command while being connected to the database:

```
SET WRITE SUSPEND FOR DATABASE
```

Once the split mirror image has been taken, I/O writes on the primary database can be resumed with the following command:

```
SET WRITE RESUME FOR DATABASE
```

Before the split mirror image can be used, the DB2INIDB tool must be run in order to remove its suspended state and initialize it by either performing crash recovery or placing it in rollforward pending state.

User response: Resubmit the command with a new split mirror image database.

DBT1011N Unable to restart database *database-name* with the WRITE RESUME option. SQLCODE = *sqlcode*.

Explanation: An attempt to restart the specified database with the WRITE RESUME option failed.

User response: Verify the specified SQLCODE, correct the problem and resubmit the command.

DBT1012N Unable to perform operation *operation-name* on the database directory.

Explanation: An attempt to perform the specified operation on the database directory was unsuccessful.

Possible reasons include:

- The database manager could not process the request due to an insufficient amount of memory in the system.
- No database entries were found in the system database directory.
- The database does not exist in the system database directory.

User response: Possible solutions include:

- Ensure that there are sufficient system resources available.
- Ensure that the database is cataloged correctly.

Correct the problem and resubmit the command.

DBT1013N The database *database-name* could not be found.

Explanation: The specified database is not an existing database or the database could not be found in the local or system database directories.

User response: Ensure that the specified database name exists in the system database directory. If the database name does not exist in the system database directory, then the database either does not exist or the database name has not been cataloged.

If the database name appears in the system database

directory and the entry type is INDIRECT, ensure that the database exists in the specified local database directory.

DBT1014N The instance name could not be determined.

Explanation: An attempt to determine the current instance name was unsuccessful.

User response: Ensure that the DB2INSTANCE environment variable is set to the current instance name. Correct the problem and resubmit the command.

DBT1015N Unable to allocate memory.

Explanation: During processing, there was not enough memory to continue processing.

User response: Possible solutions include:

- Ensure that your system has sufficient real and virtual memory.
- Remove background processes.

DBT1016N Table space files *file-name1* and *file-name2* are different sizes.

Explanation: The two specified table space files have a different size even though they are supposed to be identical.

User response: Contact your technical service representative with the following information:

- Problem description
- Contents of the table space catalog table
- Table space files

DBT1017N The syntax of the DB2RELOCATEDB tool is incorrect.

Explanation: The DB2RELOCATEDB tool has the following syntax:

```
db2relocatedb -f <config_file>
```

Where <config_file> is the name of file containing configuration information.

File format is:

```
DB_NAME=oldName,newName
DB_PATH=oldPath,newPath
INSTANCE=oldInst,newInst
DBPARTITIONNUM=dbPartitionNumber
LOG_DIR=oldDirPath,newDirPath
CONT_PATH=oldContPath1,newContPath1
CONT_PATH=oldContPath2,newContPath2
MIRRORLOG_PATH=newDirPath
FAILARCHIVE_PATH=newDirPath
LOGARCHMETH1=newDirPath
LOGARCHMETH2=newDirPath
OVERFLOWLOG_PATH=newDirPath
...
```

Notes:

- Database name, database path, and instance name are all required fields. If one of these fields is not changing then it is not necessary to list the old and new value for it, just give the old/current one.
- If you have settings for any of the following database configuration parameters, you can specify the corresponding keyword in the configuration file:
 - mirrorlogpath
 - failarchpath
 - logarchmeth1
 - logarchmeth2
 - overflowlogpath
- Blank lines or lines beginning with a comment character (#) will be ignored.

User response: Resubmit the command using the correct syntax.

DBT1018N *field-name* is a required field in the configuration file.

Explanation: The specified field is a required field and does not exist in the configuration file.

User response: Add the specified field information to the configuration file and then resubmit the command.

DBT1019N Invalid database path specified in the configuration file.

Explanation: The specified database path in the configuration file is invalid.

On Windows, the database path must include the drive letter followed by a colon character. On other platforms, the database path must be absolute and end with a path separator.

User response: Correct the erroneous database path and resubmit the command.

DBT1020N The configuration file contains an invalid entry at line *line-number*. Reason code = *reason-code*.

Explanation: You cannot relocate with the DB2RELOCATEDB tool until the following condition indicated by the reason code is resolved:

1

The specified line is longer than the maximum number of characters.

2

A field was specified but no values follow it.

3

A field has been specified more than once.

4

| | | |
|----|---|--|
| | Either the old or new database name contains too many characters. | User response: The action corresponding to the reason code is: |
| 5 | | 1 |
| | Either the old or new database path name contains too many characters. | Ensure that the specified line contains at most 1000 characters. |
| 6 | | 2 |
| | Either the old or new instance name contains too many characters. | Ensure that the necessary values are specified for the field in question. |
| 7 | | 3 |
| | An invalid node number has been specified in the DBPARTITIONNUM entry. | Ensure that the field in question is only specified once. |
| 8 | | 4 |
| | Either the old or new log directory path contains too many characters. | Ensure that the old and new database names contain no more than 8 characters. |
| 9 | | 5 |
| | Either the old or new container name contains too many characters. | Ensure that the old and new database path names contain no more than 215 characters. |
| 10 | | 6 |
| | A duplicate container name in a CONT_PATH entry has been found. | Ensure that the old and new instance names contain no more than 8 characters. |
| 11 | | 7 |
| | The specified line is invalid. | Ensure that the node number contains no more than 4 digits. |
| 12 | | 8 |
| | Incorrect usage of the wildcard character (*). | Ensure that the old and new log directory paths contain no more than 242 characters. |
| 13 | | 9 |
| | Either the old or the new storage path name contains too many characters. | Ensure that the old and new container names contain no more than 256 characters. |
| 14 | | 10 |
| | A duplicate storage path in a STORAGE_PATH entry has been found. | Ensure that each container name is only specified once in a CONT_PATH entry. |
| 15 | | 11 |
| | The new directory indicated for mirrorlogpath has too many characters. | Verify the validity of the line in question. |
| 16 | | 12 |
| | The new directory indicated for failarchpath has too many characters. | The wildcard character must be the last character in both the old and new paths. |
| 17 | | 13 |
| | The new directory indicated for logarchmeth1 has too many characters. | Ensure that the old and new storage path names contain no more than 190 characters. |
| 18 | | 14 |
| | The new directory indicated for logarchmeth2 has too many characters. | Ensure that each storage path name is only specified once in a STORAGE_PATH entry. |
| 19 | | 15 |
| | The new directory indicated for overflowlogpath has too many characters. | |

Ensure that the new directory indicated for mirrorlogpath does not exceed the maximum length allowed.

16

Ensure that the new directory indicated for failarchpath does not exceed the maximum length allowed.

17

Ensure that the new directory indicated for logarchmeth1 does not exceed the maximum length allowed.

18

Ensure that the new directory indicated for logarchmeth2 does not exceed the maximum length allowed.

19

Ensure that the new directory indicated for overflowlogpath does not exceed the maximum length allowed.

Correct the problem and resubmit the command.

DBT1021N The total size of all container names is too large.

Explanation: The total space required to store the list of containers exceeds the space allotted for a given table space in the table space file.

User response: Try one or more of the following:

- Use symbolic links, mounted file systems, etc. to shorten the new container names.
- Back up the table space and then use the database administration utility to reduce the number and/or name lengths of the containers. Restore the table space to the new containers.

Correct the problem and resubmit the command.

DBT1022N The DB2RELOCATEDB tool must be run under instance *instance-name*.

Explanation: An attempt has been made to run the DB2RELOCATEDB tool from an incorrect instance. If the instance is changing then the tool should be run from the new instance.

User response: Resubmit the command using the correct instance.

DBT1023N The database release *release-number* is invalid.

Explanation: The tool that is being used is from a different release than the database's current release. The command cannot be processed.

User response: Ensure that the correct version of the tool is being used. Correct the problem and resubmit the command.

DBT1024N Change in database name has resulted in duplicate names in the local database directory.

Explanation: An attempt to rename a database failed because an entry in the local database directory already existed with the same name.

User response: Either change the database name to one that doesn't already exist or uncatalog the database if it no longer exists. Correct the problem and resubmit the command.

DBT1025N Neither old nor new database name were found in the database directory.

Explanation: An attempt to determine the database token necessary to find the database path failed because neither the old nor new database names could be found in the database directory.

User response: Ensure the database with the old database name exists and that it is cataloged properly. Correct the problem and resubmit the command.

DBT1026N Replacement of the new database path *database-path* in container name *container-name* will exceed length restrictions.

Explanation: An attempt to rename the database path failed because it exceeds container name length restrictions when appended to the front of the specified container name.

User response: Rename the database path and/or container name such that they do not exceed length restrictions when appended together. Correct the problem and resubmit the command.

DBT1027N Replacement of the new database path *database-path* in log directory path *log-path* will exceed length restrictions.

Explanation: An attempt to rename the database path failed because it exceeds log directory path length restrictions when appended to the front of the specified log directory path.

User response: Rename the database path and/or log directory path such that they do not exceed length restrictions when appended together. Correct the problem and resubmit the command.

DBT1028N Unable to copy table space files.

Explanation: An attempt to duplicate the table space file SQLSPCS.1 failed.

User response: Ensure there is sufficient disk space for the duplicate table space file. Correct the problem and resubmit the command.

DBT1029N Unable to copy buffer pool files.

Explanation: An attempt to duplicate the buffer pool file SQLBP.1 failed.

User response: Ensure there is sufficient disk space for the duplicate buffer pool file. Correct the problem and resubmit the command.

DBT1030N Unable to copy local database directory files.

Explanation: An attempt to duplicate the local database directory file failed.

User response: Ensure there is sufficient disk space for the necessary duplicate local database directory file. Correct the problem and resubmit the command.

DBT1031N An error occurred during an attempt to update the log control files.

Explanation: The following log control files exist for each member:

- The SQLOGCTL1.LFH primary log control file
- The SQLOGCTL2.LFH secondary log control file

The following global log control files exist for each database partition:

- The SQLOGCTL.GLFH.1 primary log control file
- The SQLOGCTL2.GLFH.2 secondary log control file

The secondary log control file is a copy of the primary log control file for use in the event that there is a problem with the primary log control file.

This error was returned because the DB2 database manager could not update the primary log control files or the secondary log control files. Two of the possible reasons are as follows:

- The log control files do not exist in the database directory.
- The file permission settings prevent read or write access.

User response:

1. Identify why the DB2 database manager could not update the log control files by reviewing the DB2 diagnostic log files.
2. Correct the problem, for example, by taking the following steps:

- Ensure that both the primary and secondary log control files exist in the database directory.
- Change the permission settings of the files.

3. Reissue the command.

DBT1032N The instance is not currently active.

Explanation: The tool requires the instance to be active before it can perform any processing.

User response: Issue a db2start before issuing the command.

DBT1033N Unable to attach to instance *instance-name*. SQLCODE = *sqlcode*.

Explanation: An attempt to attach to the specified instance failed.

User response: Verify the specified SQLCODE, correct the problem and resubmit the command.

DBT1034N Replacement of the string *string-one* with *string-two* for container *container-name* will exceed length restrictions.

Explanation: An attempt to rename a container path using wildcards failed because the resulting path length is too long.

User response: The maximum length for a container path is 256 characters. Modify the container path to conform to the length restriction and reissue the command.

DBT1035N Unable to set the HADR database role to standard for database *database-name*. SQLCODE = *sqlcode*.

Explanation: An attempt to set the HADR database role to standard failed.

User response: Verify the specified SQLCODE, correct the problem and resubmit the command.

DBT1036N Unable to refresh the directory cache after a successful relocation of database *database-name*. SQLCODE = *sqlcode*.

Explanation: An attempt to perform a necessary refresh of the directory cache was unsuccessful.

User response: Restart the instance on the current node and rerun the DB2INIDB tool without specifying the RELOCATE USING option. If the problem persists, contact your technical service representative.

DBT1037N CONT_PATH entries cannot be specified for an automatic storage table space.

Explanation: An attempt has been made to relocate container paths for an automatic storage table space.

User response: Remove the CONT_PATH entries that correspond to automatic storage table space. Only STORAGE_PATH entries can be used to relocate automatic storage table space.

DBT1038N A storage group has not been defined for the database.

Explanation: An attempt has been made to relocate storage paths for a database that contains no storage groups.

User response: Remove the STORAGE_PATH entries from the configuration file.

DBT1039N Unable to copy storage group files.

Explanation: An attempt to duplicate the storage group file SQLSGF.1 has failed.

User response: Ensure that there is sufficient disk space for the duplicated storage group file. Correct the problem and resubmit the command.

Chapter 88. DBT2000 - DBT2499

DBT2002W Unable to write the end of tape marker.
Reason: *reason*

Explanation: Writing to the end of tape marker failed.

User response: No further action is required.

DBT2006I db2tapemgr completed successfully.

Explanation: Processing completed without any errors and warnings.

User response: No further action is required.

DBT2007W db2tapemgr completed with warnings.

Explanation: The db2tapemgr command completed without any errors, but warnings.

User response: Refer to the output for more information.

DBT2008N db2tapemgr failed with errors.

Explanation: db2tapemgr completed with errors.

User response: Refer to the output for more information.

DBT2009N Internal error occurred. **Reason:** *error*.

Explanation: An unexpected error occurred.

User response: Refer to the reason text for further actions.

DBT2015W Log file *log-file-name* is not on disk.

Explanation: The log file cannot be found.

User response: If the log file was accidentally deleted, either restore the log file and reissue the command or update the location field for the missing log file in the history file.

DBT2016I No log files found for processing.

Explanation: The history files do not contain information about log files.

User response: No further action is required.

DBT2017N Label *label-one* is not inserted in the tape drive. The tape inserted in the drive has a label *label-two*.

Explanation: Two different tapes cannot have the same name.

User response: Confirm that the correct tape is in the tape drive. If the tape is unreadable, use the DELETE TAPE LABEL option to remove the information about this tape from the history file. If an existing tape label was specified, choose another label and reissue the command.

DBT2018N Invalid value *value* for variable *variable*.

Explanation: The specified value is invalid.

User response: Specify a valid value.

DBT2019N The tape has not expired. The tape will expire on *expiration-date*.

Explanation: Unable to write to the tape as the tape has not expired. It is possible that the tape contains log files which may be required for database recovery.

User response: Either reduce the value of DB2_TAPEMGR_TAPE_EXPIRATION or insert a different tape.

DBT2020N Log files on tape are from a different database *database-name*.

Explanation: The log files on the tape are from a different database.

User response: Insert another tape or specify the FORCE option.

DBT2021N Log files on tape are from a different database instance *instance-name*.

Explanation: The log files on the tape are from a different database instance.

User response: Insert another tape or specify the FORCE option.

DBT2022N Log files on tape are from a different database partition *database-partition*.

Explanation: The log files on the tape are from a different database partition.

User response: Insert another tape or specify the FORCE option.

DBT2027N Tape has not been used for storing log files before.

Explanation: The contents of the tape are not recognized as a file which has been written by db2tapemgr.

User response: No further action is required.

DBT2032W Only *number-of-log-files* log files fit on tape.

Explanation: The selected number of log files do not fit on tape.

User response: To avoid this warning, specify a maximum number of log files to write on tape using the *n LOGS* option.

DBT2036W Storing history file on tape failed.
Reason: *reason*

Explanation: The log files have already been written to tape successfully.

User response: If the tape is full, limit the number of log files written to tape using the *n LOGS* option.

DBT2039I Clearing location in history for log files currently on tape *tape-name*.

Explanation: The location field of the log file entries in the history file are cleared. An empty location field in a log file entry indicates that the log file has been deleted or overwritten and is no longer available for database recovery.

User response: No further action is required.

DBT2047N Tape contains log files for database partition number *db-partitionnum-1*, however database partition number *db-partitionnum-2*.

Explanation: The tape contains logs for a database partition that does not correspond to the database partition specified.

User response: Specify the correct database partition number. If no value is specified, the value of DB2NODE is used.

DBT2048I Tape contains log files of database *database-name1*, however database *database-name2* was specified.

Explanation: The tape contains log files of another database.

User response: Proceed only if you plan on performing a redirected restore operation.

DBT2049I Tape contains log files of instance *instance1*, however instance *instance2* was specified.

Explanation: The tape contains log files of another database instance.

User response: Proceed only if you plan on

performing a redirected restore operation.

DBT2050I The log file *log-file* is on disk.

Explanation: The log file is already on disk and will not be retrieved from tape.

User response: To retrieve the log file from tape again, delete the log file on disk and specify the *RETRIEVE* option. Otherwise, no further action is required.

DBT2051N No matching backup found in the history file.

Explanation: The history file does not contain a backup.

User response: Specify a different rollforward time and reissue the command. Otherwise, restore the log files manually using the *ALL LOGS* or *LOGS n TO m* option.

DBT2052I No required log file found.

Explanation: The log files are required for rollforward operation.

User response: If the history file is current, no further action is required. Otherwise the history file may not contain all information required for a rollforward operation. If this is the case, restore a history file from the newest tape using the *RETRIEVE HISTORY FILE* option and reissue the command with the *USING HISTORY FILE* option.

DBT2053I No required log file on tape.

Explanation: The history file does not contain any entries specifying that a log file will need to be retrieved from tape for a rollforward operation.

User response: If the history file is current, no further action is required. Otherwise the history file may not contain all information required for a rollforward operation. If this is the case, restore a history file from the newest tape using the *RETRIEVE HISTORY FILE* option and reissue the command with the *USING HISTORY FILE* option.

DBT2054I Log files required for backup taken at timestamp *timestamp*:

Explanation: The log files are required for rollforward.

User response: If the history file is up-to-date, no further action is required. Otherwise the history file may not contain all information required for a rollforward operation. If this is the case, restore a history file from the most recent tape using the *RETRIEVE HISTORY FILE* option and reissue the command with the *USING HISTORY FILE* option.

DBT2055I No tapes are required for the rollforward operation.

User response: If your history file is up-to-date, no further action is required. Otherwise the history file may not contain all information required for a rollforward operation. If this is the case, restore a history file from the most recent tape using the RETRIEVE HISTORY FILE option and reissue the command with the USING HISTORY FILE option.

DBT2062I Working on database *database-name*.

Explanation: The DATABASE option was not specified. The default value controlled by the DB2DBDFT variable is used.

User response: To operate on another database, specify the DATABASE option.

DBT2063N The DATABASE option is not specified and the DB2DBDFT is not set.

User response: Specify the DATABASE option or set the DB2DBDFT variable.

DBT2065I Using database partition *partition-number*.

Explanation: The default value is either 0 or the value of the DB2NODE variable.

User response: to change the database partition, specify the DATABASE option or set the DB2DBDFT variable.

DBT2067N No directory specified and the database configuration parameter OVERFLOWLOGPATH is not set.

Explanation: Retrieve operations require that a destination directory is set.

User response: Specify the TO option or set the OVERFLOWLOGPATH database configuration parameter.

DBT2068N The value *value* of database configuration parameter OVERFLOWLOGPATH is not a directory.

User response: Ensure that the OVERFLOWLOGPATH directory exists.

DBT2069N Unable to double store to same tape.

Explanation:

User response: Choose another tape for your double store operation.

DBT2071I Using the automatic generated tape-label *label*.

Explanation: The format of the label is the current time followed by the database alias.

User response: To specify a table label, use the TAPE LABEL option.

DBT2102N File name *filename1* does not match *filename2*.

Explanation: The file name does not match the expected file name. This may occur if another program has written files to the tape.

User response: No further action is required.

DBT2103N Directory *directory-name* does not exist

Explanation: The specified directory does not exist.

User response: Create the specified directory.

DBT2104N File *file-name* found.

Explanation: More than one file is found on this tape.

User response: No further action is required.

DBT2106N Database *database-name* is remote.

Explanation: The specified database is catalogued as remote database. The db2tapemgr can only be used on local databases.

User response: Login locally and recall tool.

DBT2108N Database *database-name* not found in database directory.

Explanation: The specified database is not found in the database directory.

User response: Choose another database name or if the database exists, ensure that the database is catalogued.

DBT2109N Parameter *parameter* is expected.

Explanation: The specified parameter is expected.

User response: Reissue the command using the correct syntax.

DBT2110N Value *value* of parameter *parameter* is too long.

Explanation: The specified value for this parameter is too long.

User response: Reissue the command using the correct syntax.

DBT2111N Value *value* of parameter *parameter* is too short.

Explanation: The specified value for this parameter is too short.

User response: Reissue the command using the correct syntax.

DBT2112N Device *device-name* is not a tape device.

Explanation: The specified device is not recognized as a tape device.

User response: Specify a recognized tape device name like and reissue the command.

DBT2113N Device *device-name* is a rewind device.

Explanation: The specified device is not recognized as a non-rewind tape device.

User response: Specify a recognized tape device name like and reissue the command.

DBT2114N History file *history-file* does not exist.

Explanation: The specified history file does not exist.

User response: Check the path to the history file and Reissue the command.

DBT2115N History file *history-file* does not end with *value*.

Explanation: The specified history file name does not end with db2rhist.asc.

User response: Check the file name of the history file and reissue the command.

DBT2116N Value *value* of parameter *parameter* is not alphanumeric.

Explanation: The value of the parameter needs to be alphanumeric.

User response: Reissue the command using the correct syntax.

DBT2117N Value *value* of parameter *parameter* is out of range.

Explanation: The value of the parameter is out of range.

User response: Reissue the command using the correct syntax.

DBT2118N Parameter blocksize needs to be a multiple of 512.

User response: Choose a multiple of 512 as blocksize and reissue the command.

DBT2119N Value *value* of parameter *parameter* is not numeric.

User response: Reissue the command using the correct syntax.

DBT2120N No operation is specified.

Explanation: The parameter which specifies the operation is not specified.

User response: Reissue the command with the correct syntax.

DBT2121N Unknown operation *operation* is specified.

User response: Reissue the command using the correct syntax.

DBT2122N Too many command arguments are specified following the command parameter: *parameter*.

User response: Reissue the command using the correct syntax.

DBT2123N Invalid time format *time-format*.

Explanation: The time format should conform to the following format: YYYY-MM-DD:HH:MM:SS.

User response: Reissue the command with the correct time format syntax.

DBT2124N Unable to reduce number of log files.

User response: Use the n LOGS option to limit the amount of log files to write on tape and reissue the command.

DBT2125N File *file-name1* is not a tape header file. File *file-name2* is found.

Explanation: The tape contains other types of files. This may occur if another program has written files to the tape.

User response: No further action is required.

DBT2126N Invalid tape header format.

Explanation: The contents of the tape header are not recognized.

User response: Attempt to retrieve the log files manually from tape.

DBT2127N Changes were made to history file during update.

Explanation: Updating the history file failed because another process has updated the history file.

User response: Reissue the command.

DBT2128N The maximum number of open scans has exceeded the limit.

Explanation: Reading the history file failed because too many processes are accessing the history file.

User response: Reissue the command.

DBT2129W Fixed damaged history file.

Explanation: A damaged history file has been fixed automatically.

User response: No further action is required.

DBT2130N History file cannot be fixed.

Explanation: A damaged history file has been detected and it cannot be fixed.

User response: Restore another version of the history file and reissue command.

DBT2131N The number of table spaces has changed in the history file.

Explanation: The contents of the history file has been changed by another process.

User response: Reissue the command.

DBT2132N A duplicate timestamp exists within the history file.

Explanation: The history file contains entries with same timestamp.

User response: Restore another version of the history file and reissue command.

DBT2138N The history file does not correspond to the current partition.

User response: Specify a history file that corresponds to the ON DBPARTITIONNUM parameter that was specified when the utility started. If this parameter was not specified when the utility started, the value of DB2NODE is used.

DBT2150N db2relocatedb failed because of the value specified for the keyword *keyword* in the db2relocatedb configuration file.
Reason code: *reason-code*.

Explanation: You can use the db2relocatedb command to rename or relocate all or part of a database. The db2relocatedb takes the name of a configuration file as a parameter. In the configuration file, you can specify the configuration of the renamed or relocated database using keyword-value pairs.

For more information about the db2relocatedb command and syntax details, refer to the Command Reference in the DB2 Information Center.

This message is returned when there is a problem with the value specified for the given keyword *keyword*. The reason code indicates the nature of the problem:

1

The specified directory does not exist.

2

The db2relocatedb utility does not have permission to access the specified directory.

3

A value was specified for keyword *keyword* but the corresponding database configuration parameter is not set for the database that you are trying to relocate.

User response: Respond to this error according to the reason code:

1

Create the new directory that you want to use and re-issue the db2relocatedb command again.

2

Grant the db2relocatedb utility and the database manager write permission on the new directory and re-issue the db2relocatedb command again.

3

Remove the keyword on the configuration file and re-issue the db2relocatedb command. After the command completes successfully, you can set the database configuration parameter using the UPDATE DATABASE CONFIGURATION command or the db2CfgSet API.

DBT2200E This option is not valid. Option: *option-name*

Explanation: The specified option is not valid.

User response: Specify a valid option and run the command again.

Chapter 89. DBT3000 - DBT3499

Chapter 90. DBT3500 - DBT3999

DBT3500E **The db2prereqcheck command failed because the db2prereqcheck utility could not find the XML resource file.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can define the releases of DB2 database for which you want to verify installation prerequisites in an XML resource file. This message is returned when the db2prereqcheck utility cannot find the specified XML resource file or the default XML resource file.

User response: Perform one of the following actions:

- Specify the XML resource file using the -f parameter with the db2prereqcheck command.
- If you did not specify the XML resource file using the -f parameter, check to see if the default XML resource file exists in the default installation directory.

DBT3501E **The db2prereqcheck utility could not open the following file: *file-name*.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot open a file during normal processing. Because the db2prereqcheck utility runs with the same privileges as the user who ran the db2prereqcheck command, this message can be returned when the user who ran the db2prereqcheck command does not have read and write permission for the named file.

User response: Respond to this message in one of the following ways:

- Modify the permissions of the named file so that your user id has read and write permission for the file.
- Log in to your system with a user ID that has read and write permission on the file.
- Ask someone who has read and write permission on the file to run the db2prereqcheck command.

DBT3502E **The db2prereqcheck utility failed to verify installation prerequisites because an invalid version of DB2 database was specified in the XML resource file.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can define the releases of DB2 database for which you want to verify installation prerequisites in an XML resource file. This message is returned when an invalid

version of DB2 database is specified in the XML resource file.

User response: Correct the XML resource file and run the db2prereqcheck command again.

DBT3503E **The db2prereqcheck command failed because an invalid parameter was specified.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when an invalid parameter is specified with the db2prereqcheck command.

User response:

1. Review the db2prereqcheck command syntax.
2. Run the db2prereqcheck command again, specifying valid parameters.

DBT3504E **The db2prereqcheck utility failed to verify installation prerequisites because the utility was unable to determine the level of the operating system.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot determine the current level of the operating system.

User response: Perform one of the following actions:

- Review operating systems that are supported with DB2 database and the db2prereqcheck utility.
- Manually retrieve or look up the operating system level.

DBT3505E **The db2prereqcheck utility was unable to determine the Linux distribution level.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot determine the distribution level of the current Linux operating system. There are several reasons why the utility might be unable to determine the distribution level. For example, the utility might be unable to determine the distribution level if the utility cannot find or read the system file called /etc/issue.

The db2prereqcheck utility cannot perform all of the the necessary prerequisite verification steps without knowing the distribution level of the current Linux operating system. In order for the db2prereqcheck

utility to be used to verify installation prerequisites for DB2 database, the problem that prevented the utility from determining the distribution level must be resolved.

User response: To enable the db2prereqcheck utility to determine the distribution level of the current Linux operating system, perform the following actions:

- Ensure that the system file called /etc/issue exists.
- Ensure that the system file called /etc/issue contains distribution level details.
- Run the db2prereqcheck command as a user who has read permission for the system file called /etc/issue.

DBT3506E The db2prereqcheck utility encountered errors while running an rpm command.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

While verifying installation prerequisites for DB2 database, the db2prereqcheck utility uses the Resource Package Manager utility. This message is returned when the db2prereqcheck utility call the rpm command and that rpm command fails.

User response:

1. Determine the cause of the rpm failure by reviewing diagnostic information in the db2prereqcheck report file.
2. Resolve the cause of the rpm failure.
3. Call the db2prereqcheck command again.

DBT3507E The db2prereqcheck utility failed to find the following package or file:
package-or-file-name.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot find the named package or file while the utility is verifying installation prerequisites.

User response: To review the diagnostic report file, run db2prereqcheck with -o option. By default, prerequisite validation is displayed on screen.

DBT3508E The db2prereqcheck utility encountered an unhandled error.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility encounters an error that the utility cannot classify or resolve.

User response:

1. Run with db2prereqcheck -o option.

2. Identify the cause of the problem by reviewing diagnostic information in the db2prereqcheck report file.
3. Resolve the cause of the problem.
4. Run the db2prereqcheck command again.

DBT3509E The db2prereqcheck utility was unable to determine the version of the following package: *package-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility needs to determine the version of a package to verify installation prerequisites, but the utility cannot determine the version of the package.

User response: Review the list of operating systems and software that are supported with DB2 database manually.

DBT3510E The db2prereqcheck utility was unable to find the following required library file: *library-file-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot find a library file that is required to verify installation prerequisites.

User response: Manually verify that the DB2 database installation prerequisites are met.

DBT3511E The db2prereqcheck utility was unable to find the following map file: *map-file-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot find a map file that is required to verify installation prerequisites.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3512W The db2prereqcheck utility failed to determine the currently-installed version of the C++ standard library, libstdc++.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The current system might meet DB2 database installation prerequisites even if the db2prereqcheck utility cannot determine the version of the C++ standard library that is installed.

User response: Manually verify that the system has the required version of C++ standard library.

DBT3513W The db2prereqcheck utility failed to determine the version of the currently-installed C++ standard library, libstdc++, using the ldconfig utility.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility uses several different methods for determining the current version of the C++ standard library. This message is returned when the db2prereqcheck utility attempts to determine the version of the currently-install C++ standard library using the ldconfig utility.

The current system might meet DB2 database installation prerequisites even if the db2prereqcheck utility cannot determine the version of the C++ standard library that is installed.

User response: Manually verify that the system has the required version of C++ standard library.

DBT3514W The db2prereqcheck utility failed to find the following 32-bit library file: *library-file-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot find the 32-bit version of a required library file. If the named file is missing, 32-bit database applications might not function properly.

User response: Respond to this message in one of the following ways:

- If you do not intend to use 32-bit applications with DB2 database, then no response is required.
- If you do intend to use 32-bit applications with DB2 database, ensure that the named 32-bit library file exists on the system before installing DB2 database.

DBT3515E The db2prereqcheck utility failed to determine the version of the currently-installed C standard library, glibc.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3516E The db2prereqcheck utility was unable to execute the following command: *command*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility attempts to execute a command to verify

installation prerequisites, but the command fails.

User response: Verify manually whether the DB2 database installation prerequisites are met.

DBT3517E The db2prereqcheck utility could not read the following file: *file-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot read a file during normal processing. Because the db2prereqcheck utility runs with the same privileges as the user who ran the db2prereqcheck command, this message can be returned when the user who ran the db2prereqcheck command does not have read and write permission for the named file.

User response: Respond to this message in one of the following ways:

- Modify the permissions of the named file so that your user id has read and write permission for the file.
- Log in to your system with a user ID that has read and write permission on the file.
- Ask someone who has read and write permission on the file to run the db2prereqcheck command.

DBT3518E The db2prereqcheck utility was unable to determine the AIX service pack level.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Manually verify that the system has the required AIX service pack level.

DBT3519E The db2prereqcheck utility was unable to determine the AIX Technology Level (TL).

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Manually verify that the system has the required AIX Technology Level (TL).

DBT3520E The db2prereqcheck utility could not find the library file libaio.so.1.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether the DB2 database installation prerequisites are met.

DBT3521E The db2prereqcheck utility was unable to determine the version of the library libibmc++.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether the DB2 database prerequisites are met.

DBT3522E **The db2prereqcheck utility could not find the following patch:** *patch-identifier*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether the DB2 database installation prerequisites are met.

DBT3523E **The db2prereqcheck utility could not find the following required HP-UX bundles:** *bundle-names*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3524E **The db2prereqcheck utility determined that the current Windows operating system is not supported by the version of DB2 database that was specified in the XML resource file.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility determines that the version of the currently-install Windows operating system is not supported by one or more versions of DB2 database that were specified in the db2checkprereq XML resource file.

User response: To install DB2 database, change the Windows operating system to a version that is supported by DB2 database.

DBT3525E **The db2prereqcheck command failed because the XML resources file is invalid.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can define the releases of DB2 database for which you want to verify installation prerequisites in an XML resource file. This message is returned when the db2prereqcheck utility cannot process the contents of the specified XML resource file or the default XML resource file.

User response: Correct the contents of the XML resource file and then run the db2prereqcheck command again.

DBT3526E **The db2prereqcheck command failed because the value specified with the -o parameter is invalid.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility produces an output, report file that contains a log of the prerequisites checked and the success of the verification. There is a default report file name, but you can also specify the name of the report file using the -o parameter with the db2prereqcheck command.

This message is returned when the value specified with the -o parameter is not valid for a file name. For example, this message can be returned when the value specified with the -o parameter is not in a format that is valid for a file name.

User response: Respond to this message in one of the following ways:

- Run the db2prereqcheck command again, specifying a valid file name with the -o parameter.
- Use the default report file name by running the db2prereqcheck command again, without specifying the -o parameter.

DBT3527E **The db2prereqcheck command failed because the value specified with the -v parameter is invalid.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can specify a version of DB2 database using the -v parameter with the db2prereqcheck command.

This message is returned when a value is specified for the -v parameter with the db2prereqcheck command, but the specified value does not match any valid versions of DB2 database defined in the resource XML file.

User response: Call the db2prereqcheck command again, specifying a valid version of DB2 database with the -v parameter.

DBT3528E **The db2prereqcheck command failed because the value specified with the -f parameter is invalid.**

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility reads information about which prerequisites to check from an input XML resource file. There is a default XML resource file name, but you can also specify the name of the XML resource file using the -f parameter with the db2prereqcheck command.

This message is returned when the file specified with the -f parameter is not in valid XML file or its content

is not in the format recognized by db2prereqcheck tool.

User response: Respond to this message in one of the following ways:

- Call the db2prereqcheck command again, specifying a valid XML resource file with the -f parameter.
- Use the default XML resource file by calling the db2prereqcheck command again, without specifying the -f parameter.

DBT3529E The db2prereqcheck command failed because the mutually exclusive parameters -v and -i were both specified.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can specify a version of DB2 database using the -v parameter with the db2prereqcheck command.

If you use an XML resource file that lists one or more versions of DB2 database, you can cause the db2prereqcheck utility to verify the installation prerequisites for only the latest version of DB2 database that is specified in the XML resource file by specifying the -i parameter with the db2prereqcheck command.

The -v parameter and the -i parameter are mutually exclusive.

User response: Call the db2prereqcheck command again, specifying either the -v parameter or the -i parameter, but not both.

DBT3530E The db2prereqcheck command failed because the -i parameter was specified with the command, but the db2prereqcheck utility could not determine the latest DB2 database version that is listed in the XML resource file.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can specify a version of DB2 database using the -v parameter with the db2prereqcheck command.

If you use an XML resource file that lists one or more versions of DB2 database, you can cause the db2prereqcheck utility to verify the installation prerequisites for only the latest version of DB2 database that is specified in the XML resource file by specifying the -i parameter with the db2prereqcheck command.

This message is returned when the -i parameter is specified with the db2prereqcheck command, but the db2prereqcheck utility failed to determine what the latest version of DB2 database is that is specified in the XML resource file contents. There are multiple reasons for which the db2prereqcheck utility might be unable to retrieve the latest version of DB2 database from the XML resource file, including the following reasons:

- No DB2 database versions are specified in the XML resource file.
- The format of the XML resource file is not valid.
- The DB2 database versions that are specified in the XML resource file are not valid.

User response:

1. Perform the following troubleshooting steps:
 - Ensure that the format of the XML resource file is valid.
 - Ensure that one or more versions of DB2 database are correctly specified in the XML resource file.
2. Call the command again in one of the following ways:
 - Verify prerequisites for the latest version of DB2 database that is defined in the resource XML file by calling the db2prereqcheck command with the -i option.
 - Verify prerequisites for all DB2 database versions that are defined in the resource XML file by calling the db2prereqcheck command without the -i option.

DBT3531E The db2prereqcheck utility failed to log the following message to the output report file: *message-text*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility produces a report file that contains information about the installation prerequisites that were verified on the current system.

This message is returned when an error occurs while db2prereqcheck is attempting to print information to the report file.

User response:

1. Verify that the user has a write permission to the report file.
2. Review the information in the message that could not be printed to the report file, and then consider that information when you review the information that is in the report file.

DBT3532E The db2prereqcheck utility determined that the DB2 database product that is being installed requires a 64-bit operating environment, but the current environment is a 32-bit environment.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: To install this DB2 database product, use a 64-bit operating environment.

DBT3533I The db2prereqcheck utility has confirmed that all installation prerequisites were met for DB2 database client-or-server-info feature-info. Version: DB2-database-version.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: No response is required.

DBT3534W The db2prereqcheck utility determined that ASLR is set to ON and that this could cause issues with some tools.

Explanation: Address Space Layout Randomization (ASLR) is a feature that is activated by default on some Linux distributions. ASLR is designed to load shared memory objects in random addresses.

DB2 database cannot guarantee the availability of addresses for shared memory objects when ASLR is turned on.

This conflict in the address space means that a process trying to attach a shared memory object to a specific address might not be able to do so, resulting in a failure in the shmat subroutine. However, on subsequent retry (using a new process) the shared memory attachment may work. The result is a random set of failures. Some processes that have been known to see this error are: db2pd, db2egcf, and db2vend.

Some of the errors that can result from this problem include the following:

- The db2pd command might report no data found even though the instance or database is active.
- The db2egcf process, which is used for monitoring, might incorrectly determine the instance is down and initiate a failover operation.
- During backup and log archive operations, the db2vend process might fail with an error indicating a child process could not be started.

User response: Respond to this error in one of the following ways:

- Turn off the randomization by setting the following kernel parameter:
kernel.randomize_va_space=0
 - Temporarily disable the randomization and restart the db2 instance by issuing the following command:
echo 0 > /proc/sys/kernel/randomize_va_space
-

DBT3535E The db2prereqcheck utility was unable to find the following InfiniBand support package: *package_name*.

Explanation: On Linux operating systems on an InfiniBand network or a 10 Gigabit Ethernet (10GE) network, specific InfiniBand Support packages are required. For a list of required packages, see the DB2

pureScale Feature installation prerequisites for Linux topic.

User response: To install the required InfiniBand software, run a group installation of the InfiniBand Support package using the yum command.

DBT3536E The db2prereqcheck utility was unable to find the following HPN package: *package_name*.

Explanation: On Linux operating systems on a 10 Gigabit Ethernet (10GE) network, specific High Performance Networking (HPN) packages are required. For a list of required packages, see the DB2 pureScale Feature installation prerequisites for Linux topic.

User response: To install the required HPN packages, run a group installation of the InfiniBand Support package using the yum command.

DBT3537E The db2prereqcheck utility was unable to find the following RSCT package: *package_name*.

Explanation: On Linux operating, specific Reliable Scalable Cluster Technology (RSCT) packages are required. For a list of required packages, see the DB2 pureScale Feature installation prerequisites for Linux topic.

User response: To install the required RSCT packages, run the yum install command for each required package.

DBT3538E The db2prereqcheck command failed because the value specified with the -t parameter is invalid.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can specify the type of network configuration for which the db2prereqcheck utility should verify prerequisites using the -t parameter. There are four valid values for the -t parameter:

- SINGLE_IB_PORT_CLUSTER
- MULTI_IB_PORT_CLUSTER
- SINGLE_ROCE_PORT_CLUSTER
- MULTI_ROCE_PORT_CLUSTER

This message is returned when an invalid value is specified for the -t parameter.

User response: Call the db2prereqcheck command again, specifying a valid value with the -t parameter.

DBT3539E The db2prereqcheck command failed because the -t parameter was specified but the -p parameter was not specified.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

You can verify installation prerequisites specifically for the DB2 pureScale environment by specifying the -p parameter. If you specify that the db2prereqcheck utility should verify the prerequisites for the DB2 pureScale environment, you can also specify the type of network configuration for which the db2prereqcheck utility should verify prerequisites using the -t parameter.

If you do not specify the -p parameter, the -t parameter must not be specified.

User response: Respond to this error in one of the following ways:

- If you are verifying installation prerequisites for a DB2 pureScale environment, call the db2prereqcheck command again, specifying the -p parameter with or without the -t parameter.
- If you are not verifying installation prerequisites for a DB2 pureScale environment, call the db2prereqcheck command again without specifying the -t parameter.

DBT3540E The db2prereqcheck command failed because the db2prereqcheck command is not supported on the current operating system level.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility requires the following minimum operating system levels:

- AIX 6.1
- HP-UX 11iv3
- Solaris 10

User response: Run the db2prereqcheck command on a system with a supported operating system level.

DBT3541E An instance management task failed because the db2prereqcheck utility encountered an internal error which cannot be resolved.

Explanation: The DB2 database manager performs a variety of system validation steps, including using the db2prereqcheck utility, before installing the DB2 database product and before creating or managing database manager instances.

This message is returned when an unhandled error is encountered by the db2prereqcheck utility as part of those larger install or instance management operations.

User response: Contact IBM support for help with resolving this problem.

DBT3542E An instance management task failed because the db2prereqcheck utility encountered an internal error. Reason code: *reason-code*.

Explanation: The DB2 database manager performs a variety of system validation steps, including using the db2prereqcheck utility, before installing the DB2 database product and before creating or managing database manager instances.

The reason code indicates a specific reason for the error:

1

The database manager or the db2prereqcheck utility failed to open the file called "/etc/hosts".

2

The database manager or the db2prereqcheck utility failed to determine the IP address of the machine on which the cluster caching facility is located.

User response: Respond to this error by performing troubleshooting steps according to the given reason code:

Reason code 1

Ensure the system file called "/etc/hosts" exists, is accessible, and that the file contents are standard for a hosts file.

Reason code 2

Ensure that the host machines in the DB2 cluster are up and running, and the the network in the cluster is functioning normally.

DBT3543E The db2prereqcheck utility encountered an internal error while printing a message. Message number: *message-number*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the utility encounters an error while printing a message in the current user language.

The failure to print one message might not affect the usefulness of the db2prereqcheck utility. However, if this error happens multiple times, this might indicate a problem with the installation of the db2prereqcheck utility.

User response: Respond to this error in one or more of the following ways:

- Run the db2prereqcheck utility in English by setting the locale to English and then calling the db2prereqcheck command.
- Download a new copy of the db2prereqcheck utility.

DBT3544E The db2prereqcheck utility failed to determine the level of the operating system by using the uname command.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility uses the system command, "uname", to determine the level of the current operating system. There are multiple reasons that the db2prereqcheck utility might be unable to determine the operating system level by using the uname command, including the following examples:

- The call to the uname system command failed.
- The output of the uname system command is not what the db2prereqcheck utility was designed to expect.

User response: Manually determine the version of the operating system and verify whether the current operating system version is supported with the DB2 database product being installed.

DBT3545E The db2prereqcheck utility encountered an error while opening an internal, temporary file.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility fails to open a temporary file while performing the prerequisite verification.

User response: Run the db2prereqcheck utility again.

DBT3546E The db2prereqcheck utility failed to determine whether the following file or package exists: *file-or-package-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility determines details about the current operating environment and system components by using system commands and examining currently installed files. This message is returned when the db2prereqcheck utility is unable to perform the prerequisites validation because the utility cannot find the given file or package.

User response: Manually verify prerequisites for the DB2 database product.

DBT3547E The db2prereqcheck utility failed to determine the version of the glibc library by using the 'strings' system command to extract information from the glibc library file.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3548E The db2prereqcheck command failed because the db2prereqcheck utility failed to open the resource XML file: *resource-XML-file-name*.

Explanation: You can verify whether a system meets the prerequisites for DB2 database product releases and fix packs that are defined in a resource XML file. You can cause the db2prereqcheck utility to read from a specified resource XML file by using the -f parameter with the db2prereqcheck command, or you can allow the db2prereqcheck utility to use a default resource XML file.

This message is returned when the db2prereqcheck utility fails to open the resource XML file, either a file specified with the -f parameter, or the default file.

User response: Call the db2prereqcheck command again, specifying a resource XML file that the db2prereqcheck utility can access.

DBT3549E The db2prereqcheck utility failed to extract an expected string from the output of a system command or the contents of a system configuration file. Expected string: *expected-string*. Source string: *source-string*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility verifies system details by using system commands and examining various system configuration files. This message is returned when the db2prereqcheck utility is unable to extract necessary information from the output of a system command or the contents of a system configuration file.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3550E The db2prereqcheck utility failed to determine the package information by using the lslpp command. Package label: *package-label*. Package identifier: *package-identifier*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility verifies system details by

using system commands and examining the contents of system configuration files. This message is returned when the db2prereqcheck utility fails to extract the necessary package information from the output of the lspp command.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3551E The db2prereqcheck utility failed to determine the package version by using the lspp command.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility verifies system details by using system commands and examining the contents of system configuration files. This message is returned when the db2prereqcheck utility fails to extract the necessary package information from the output of the lspp command.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3552E The db2prereqcheck utility failed to open the libaio.so.1 file. Command output: *err-output*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Verify manually whether DB2 database installation prerequisites are met.

DBT3553I The db2prereqcheck utility successfully loaded the libaio.so.1 file.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility verifies system details using system commands and system configuration files. This message is returned when the db2prereqcheck utility successfully loads a system file.

User response: No user response is required.

DBT3554W The db2prereqcheck utility was unable to find the package for the IOCP on the system.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The Input/Output Completion Port (IOCP) is required by the DB2 database server.

User response: To successfully run a DB2 database server, install and configure IOCP manually.

DBT3555E The db2prereqcheck utility determined that the current platform is not supported by the following version of DB2 database: *version*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: No response is required

DBT3556W The db2prereqcheck utility encountered an internal error while cleaning up temporary resources.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: No response is required.

DBT3557E The db2prereqcheck utility determined that the current platform is not supported with the DB2 pureScale Feature.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: No user response is required.

DBT3558E The db2prereqcheck command failed because two or more of the following mutually exclusive parameters were specified: *-c*, *-u*, *-g*, and *-p*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The *-c*, *-u*, *-g*, and *-p* parameters are mutually exclusive.

User response: Call the db2prereqcheck command again, specifying only one of the parameters *-c*, *-u*, *-g*, or *-p*.

DBT3559E The db2prereqcheck command failed because the value specified with the *-a* parameter is invalid.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Call the db2prereqcheck command again, specifying a valid value with the *-a* parameter.

DBT3560I The db2prereqcheck utility found the following patch: *patch-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: No response is required.

DBT3561E The db2prereqcheck utility failed to find the following bundle: *bundle-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

User response: Manually verify installation prerequisites for DB2 database.

DBT3562E The db2prereqcheck utility failed to determine whether SELinux is enabled.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

Because some requirements, such as General Parallel File System (GPFS), are not compatible with Security-Enhanced Linux (SELinux) enabled, the db2prereqcheck utility attempts to determine whether SELinux is enabled. This message is returned when the db2prereqcheck utility is unable to determine whether SELinux is enabled.

User response: Manually verify installation prerequisites for DB2 database.

DBT3563E The db2prereqcheck utility determined that SELinux is enabled, which is not supported with GPFS.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

DB2 database requires a General Parallel File System (GPFS.) However, GPFS is not supported with Security-Enhanced Linux (SELinux) enabled.

User response: Disable SELinux and then proceed with the DB2 database product installation.

DBT3564E The db2prereqcheck utility was unable to find the package *package-name* on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility determines details about the current operating environment and system components by using system commands and examining currently installed files. This message is returned when the db2prereqcheck utility is unable to perform the prerequisites validation because the utility cannot find the given file or package on the given host machine.

User response: Manually verify prerequisites for the DB2 database product.

DBT3565E The db2prereqcheck utility found that remote root login is not enabled on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

In versions of DB2 database earlier than DB2 Version 10.5, passwordless remote root login must be enabled to install DB2 database. This message is returned when passwordless remote root login needs to be enabled but the db2prereqcheck utility determined that passwordless remote root login is not currently enabled on the named host machine.

User response: Enable passwordless remote root login on all host machines on which the DB2 database install will be performed, and then perform the installation operation.

DBT3566E The db2prereqcheck utility detected that the service named *service-name* is not enabled on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility attempts to determine whether all system services that are required for successfully installing DB2 database are enabled. This message is returned when the db2prereqcheck utility detects that a required system service is not enabled.

User response: Manually verify installation prerequisites for DB2 database.

DBT3567E The db2prereqcheck utility found that passwordless ssh is not enabled from host *host-name-1* to host *host-name-2*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

DB2 database requires that passwordless remote root login be enabled during install. However, the db2prereqcheck utility determined that remote root login is not currently enabled between the named host machines.

User response: Enable passwordless remote root login on all host machines on which the DB2 database install will be performed, and then perform the installation operation.

DBT3568E The db2prereqcheck utility found that the input/output completion Port (IOCP) is not installed on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The input/output Completion Port (IOCP) is required by the DB2 database server.

This message is returned when the db2prereqcheck utility detects that no input/output completion ports are installed on the given host machine.

User response: Manually install and configure IOCP on the named host machine.

DBT3569E The db2prereqcheck utility found that the input/output completion Port (IOCP) is not enabled on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The input/output Completion Port (IOCP) is required by the DB2 database server.

This message is returned when the db2prereqcheck utility detects that no input/output completion ports are enabled on the given host machine.

User response: Manually install and configure IOCP on the named host machine.

DBT3570E The db2prereqcheck utility found that the specified device (device path: *path*) does not have an associated identifier of type *type*. Host: *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

This message is returned when the db2prereqcheck utility cannot determine either the physical volume identifier (PVID) or the worldwide identifier (WWID) for the specified device path.

User response: Configure an identifier for the specified device and perform the installation operation again.

DBT3571E The db2prereqcheck utility found that the directory *directory-name* does not have enough free space on host *host-name*. Required space: *space-needed*. Actual space present: *space-available*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database.

This message is returned when the db2prereqcheck utility determined that there is insufficient space in the named directory on the given host machine for the installation of the DB2 product to proceed successfully.

User response: Respond to this message in one of the following ways:

- Allocate more space in the given directory on the given host machine.
- Specify a different directory that has more available space.

DBT3572W The db2prereqcheck utility found that netname *net-name* is not pingable from host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck is unable to successfully use the ping command to test connectivity and network response time from the given host machine to the given cluster interconnect netname.

User response: Configure the network so that the specified netname is accessible on the network from the host machine.

DBT3573W The db2prereqcheck utility found that the /etc/hosts file of host *host-name* associates the hostname with localhost.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck finds that TCP/IP has been configured incorrectly for the named host machine.

User response: Configure TCP/IP for the named host machine.

DBT3574W The db2prereqcheck utility found that the /etc/hosts file is missing an entry for the hostname or netname *name* on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the /etc/hosts file has not been configured correctly on the named host machine.

User response: Configure the /etc/hosts file on the host machine.

DBT3575W The db2prereqcheck utility found that the /etc/hosts file of host *host-name* contains duplication of entry *duplicate-host-name-entry*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for

successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the /etc/hosts file has not been configured correctly on the named host machine.

User response: Configure the /etc/hosts file on the host machine.

DBT3576W The db2prereqcheck utility found that the /etc/hosts file contains inconsistent entries for hostname or netname *name* on the following hosts: *host-machine-list*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the /etc/hosts file has not been configured correctly on the named host machines.

User response: Configure the /etc/hosts file on the listed host machines.

DBT3577E The db2prereqcheck utility found that the library entry *library-name* is missing in the dat.conf file for netname *net-name* on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the direct access transport file, dat.conf, has not been configured correctly on the named host machine.

User response: In the dat.conf file on the host machine, add an entry for the library.

DBT3578E The db2prereqcheck utility found that the dat.conf file is not configured properly for netname *net-name* on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the direct access transport file, dat.conf, has not been configured correctly on the named host machine.

User response: In the dat.conf file on the host machine, configure an entry for the netname.

DBT3579E The db2prereqcheck utility detected more than one dat.conf file on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the direct access transport file, dat.conf, has not been configured correctly on the named host machine.

User response: Configure only one dat.conf file on the host machine.

DBT3580E The db2prereqcheck utility was unable to find the dat.conf file on host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the direct access transport file, dat.conf, has not been configured correctly on the named host machine.

User response: Configure a dat.conf file on the host machine.

DBT3581E The db2prereqcheck utility found that the specified user *user-id* is not present in host *host-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that a required user id does not exist on the named host machine.

User response: Create the required user accounts on the host machine.

DBT3582E The db2prereqcheck utility found that the UID of user *user-name-1* present on host *host-name-1* is not same as the UID of user *user-name-2* present on host *host-name-2*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for

successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that a required user id does not exist on all required host machines.

User response: Create the required user accounts on all host machines.

DBT3583E The db2prereqcheck utility found that the GID of user *user-name-1* present on host *host-name-1* is not same as the GID of user *user-name-2* present on host *host-name-2*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that a required user id does not exist on all required host machines.

User response: Create the required user accounts on all host machines.

DBT3584E The db2prereqcheck utility found that either the user account or the group ID for the instance named *instance-name* is not valid on the following host machine: *host-machine-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that a required user id does not exist on the host machine.

User response: Create the required user accounts on all host machines.

DBT3585E The db2prereqcheck utility was unable to find the following Mellanox package: *package-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility determines details about the current operating environment and system components by using system commands and examining currently installed files. This message is returned when the db2prereqcheck utility is unable to perform the prerequisites validation because the utility cannot find the given file or package.

User response: Manually verify prerequisites for the DB2 database product.

DBT3586E The db2prereqcheck utility detected both of the following packages which cannot coexist: the RHEL server High Performance Networking package (HPN) and the Mellanox adapter driver package.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility determines details about the current operating environment and system components by using system commands and examining currently installed files. This message is returned when the db2prereqcheck utility finds two packages are installed that are not supported together.

User response: Remove either the HPN package or the Mellanox package.

DBT3587E The db2prereqcheck utility failed to verify prerequisites on the host machine *host-name*.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

The db2prereqcheck utility determines details about the current operating environment and system components by using system commands and examining currently installed files. This message is returned when the db2prereqcheck utility is unable to perform the prerequisites validation.

User response: Manually verify prerequisites for the DB2 database product.

DBT3588W The db2prereqcheck utility was unable to validate the configuration of the log_mmts_per_seg parameter on the following host machine: *host-machine-name*. Reason code: *reason-code*.

Explanation: In DB2 pureScale environments, the Mellanox HCA driver mlx4_core parameter named log_mmts_per_seg must be configured in the modprobe configuration file, modprobe.conf, on host machines where any cluster caching facility (CF) resides. If this configuration parameter is not set appropriately, errors could occur when the database manager is started with the db2start command.

This message is returned when the db2prereqcheck utility is unable to validate the setting of log_mmts_per_seg on the given host machine for the reason indicated by the reason code:

1

The db2prereqcheck utility cannot find the file modprobe configuration file on the host machine.

2

The db2prereqcheck utility cannot find an entry for the log_mmts_per_seg parameter in the modprobe configuration file in the host machine.

User response:

- If there is no CF on the host machine, no user response is required.
- If there is a CF on the host machine, manually verify the configuration of the log_mmts_per_seg parameter.

DBT3589W The db2prereqcheck utility failed to validate that the IOCP is enabled on the following host machine:
host-machine-name.

Explanation: Configuring I/O completion ports (IOCPs) on AIX servers is not required for the installation of DB2 for Linux, UNIX, and Windows software. However, this configuration step is recommended for performance purposes.

User response: Optional: Manually configure IOCP.

DBT3590W The db2prereqcheck utility successfully accessed the disk called *disk-identifier* on the host machine named *host-machine-name*. However the db2prereqcheck utility failed to verify whether that disk is part of a SAN.

Explanation: You can verify installation prerequisites for DB2 database using the db2prereqcheck utility.

Part of the prerequisite verification that the db2prereqcheck utility performs includes verifying whether storage prerequisites have been met.

In DB2 pureScale environments, because the DB2 pureScale Feature requires the use of a storage area network (SAN), the db2prereqcheck utility attempts to determine whether a disk on a host machine is a member of a SAN.

User response:

- If you are not using the DB2 pureScale Feature, no response is required.
- If you are using the DB2 pureScale Feature, manually verify whether the disk is a part of a SAN.

DBT3591E The db2prereqcheck utility failed to verify the dat.conf file on the host machine named *host-machine-name*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck utility finds that the direct access transport file, dat.conf, has

not been configured correctly on the named host machine.

User response: Manually verify the contents of the dat.conf file on the host machine.

DBT3592E The db2prereqcheck utility found that the /etc/hosts file of host *host-name* associates the hostname with localhost.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database. This message is returned when the db2prereqcheck finds that TCP/IP has been configured incorrectly for the named host machine.

User response: Configure TCP/IP for the named host machine.

DBT3593W The db2prereqcheck utility found that the log_mmts_per_seg parameter is not set to the recommended value on the following host machine:
host-machine-name.

Explanation: In pureScale environments, the Mellanox HCA driver mlx4_core parameter named log_mmts_per_seg must be configured in the modprobe configuration file, modprobe.conf, on host machines where any cluster caching facility (CF) resides.

This message is returned when the log_mmts_per_seg parameter is not configured to the recommended value of "7" on the host machine.

User response:

- If there is no CF on the host machine, no user response is required.
- If there is a CF on the host machine, manually configure the log_mmts_per_seg parameter to the recommended value.

DBT3594W The db2prereqcheck utility found that the directory *directory-name* does not have the recommended amount of free space on host *host-name*. Recommended amount needed: *space-needed*. Actual space present: *space-available*.

Explanation: You can verify installation prerequisites for DB2 database by using the db2prereqcheck utility.

The db2prereqcheck utility verifies many hardware, software, and configuration prerequisites for successfully installing and running DB2 database.

This message is returned when the db2prereqcheck utility determines that there is less than the recommended amount of space in the named directory

on the given host machine, but that installing the DB2 product could proceed successfully.

Note that although having less than the recommended amount of space on the directory named in the run time token *directory-name* might not prevent installing the DB2 product from succeeding, having less than the recommended amount of space could cause run time performance problems or other errors in the future.

User response: [Optional] Respond to this message in one of the following ways:

- Allocate more space in the given directory on the given host machine.
- Specify a different directory that has more available space.

Chapter 91. DBT4000 - DBT4499

DBT4000N No value was specified for the option named *option-name*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

You must specify a value for the named option.

For an explanation of the syntax of the db2fedgentf command, use the following command: "db2fedgentf -h".

User response: Call the db2fedgentf command again, specifying a value for the named option.

DBT4001N Neither "-create" nor "-drop" was specified. It is mandatory to specify one of these options.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

You must specify exactly one of the "-create" or "-drop" options with the db2fedgentf command.

For an explanation of the syntax of the db2fedgentf command, use the following command: "db2fedgentf -h".

User response: Call the db2fedgentf command again, specifying either "-create" or "-drop".

DBT4002N An invalid option was specified: *option-name*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The given option is not a valid option for the db2fedgentf command.

For an explanation of the syntax of the db2fedgentf command, use the following command: "db2fedgentf -h".

User response: Call the db2fedgentf command again, specifying valid options.

DBT4003N One of the following required options was missing: "-db", "-u", or "-p".

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

When you call the db2fedgentf command, you must specify a database name, a user name, and a password.

For an explanation of the syntax of the db2fedgentf command, use the following command: "db2fedgentf -h".

User response: Call the db2fedgentf command again, specifying a database, a user name, and a password.

DBT4004N One of the following required options was missing: "-stpn" or "-c".

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

When you call the db2fedgentf command to create a new table function for a given federated stored procedure, you must specify both of the following:

- The name of the federated stored procedure using the "-stpn" option
- The column name and type pairs of the signature of the federated stored procedure

For an explanation of the syntax of the db2fedgentf command, use the following command: "db2fedgentf -h".

User response: Call the db2fedgentf command again, specifying the federated stored procedure name and the column name and type pairs of the signature of the stored procedure.

DBT4005N The string value for the "-c" option is not valid. db2fedgentf could not parse the column name and type pairs specified in that string. The string that was passed to the db2fedgentf call is: *column-name-type-pairs-list*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

When you call the db2fedgentf command to create a new table function for a given federated stored procedure, you must specify the column name and type pairs of the signature of the federated stored procedure.

The format of that string should be:

"<name1> <type1>, <name2> <type2>, ..."

For example:

"PID CHAR(10), PRICE DOUBLE, QTY INT"

User response:

1. Review the signature of the federated stored procedure for which you want to create the table function.

2. Review the format of the string you want to pass to the "-c" option of the db2fedgentf command.
3. Call the db2fedgentf command again, passing the column name and type pairs from the federated stored procedure signature in a string of valid format to the "-c" option.

DBT4006N db2fedgentf could not convert the SQL data type of the column named *column-name* to a JAVA data type. The SQL data type of the column is: *SQL-data-type*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The db2fedgentdf utility creates JAVA table functions. To accomplish this, the utility performs the following tasks: creates a JAVA file that contains the source for the new table function; compiles that generated JAVA file; and registers the new table function.

Because db2fedgentdf creates JAVA table functions, db2fedgentf must convert the SQL data types of the stored procedure result set columns to JAVA data types. In this case, there is no JAVA data type to which the named column can be converted.

No table function was created.

User response: Modify the federated stored procedure to change the SQL data type of then named column to one that can be converted to a JAVA data type.

DBT4008N db2fedgentf could not find the federated stored procedure called *stored-procedure-name*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The message was returned because db2fedgentf could not find an entry in the catalog for the federated stored procedure specified with the "-stpn" option on the command line.

No table function was created.

User response: Call db2fedgentf again, specifying a federated stored procedure that can be found in the catalog.

DBT4011N The db2fedgentf utility failed to create the JAVA source file for the table function. File name: *File-name*. Reason: *reason-code*. Error string: *Error-string*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The db2fedgentdf utility creates JAVA table functions.

To accomplish this, the utility performs the following tasks: creates a JAVA file that contains the source for the new table function; compiles that generated JAVA file; and registers the new table function.

The reason this message was returned is indicated by the reason code:

- 1 db2fedgentf could not allocate a file handle for the JAVA source file.
- 2 db2fedgentf could not write the table function source to the JAVA source file.
- 3 db2fedgentf could not open a file handle for the JAVA source file. Refer to the error string for more information.
- 4 db2fedgentf could not write the table function source code to the JAVA source file. Refer to the error string for more information.

No table function was created.

User response: Respond to this message according to the reason code:

- 1 Restart db2fedgentf.
- 2 If the maximum number of file handles has been allocated, close some of the open handles, and then restart db2fedgentf.
- 3 and 4 Respond according to the contents of the error string. For example, if db2fedgentf does not have the necessary file access permissions, work with your system administrator to grant db2fedgentf the required permission.

DBT4012N The db2fedgentf utility failed to compile the JAVA source file for the table function. File name: *File-name*. Reason: *reason-code*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The db2fedgentdf utility creates JAVA table functions. To accomplish this, the utility performs the following tasks: creates a JAVA file that contains the source for the new table function; compiles that generated JAVA file; and registers the new table function.

The reason this message was returned is indicated by the reason code:

- 1
db2fedgentf could not determine the current path in which to create the JAVA source file.
- 2
db2fedgentf could not determine the DB2 install path.
- 3
db2fedgentf could not delete the existing version of the class file.
- 4
db2fedgentf could not copy the class file to the destination: <DB2-INSTALL-PATH>/function

No table function was created.

User response: Respond to this message according to the reason code:

- 1
Verify that db2fedgentf has the necessary access authorization for the current directory.
- 2
Investigate whether there are any errors with the current installation of DB2 database.
- 3
Manually delete the existing version of the class file from: <DB2-INSTALL-PATH>/function
- 4
Verify that db2fedgentf has the required access authorization for the following directory: <DB2-INSTALL-PATH>/function. If necessary, work with your system administrator to grant db2fedgentf access to that directory.

DBT4014I db2fedgentf successfully dropped the table function called *table-function-name*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

User response: You do not need to respond to this message.

DBT4016N db2fedgentf could not drop the table function called *table-function-name* because there are multiple table functions with the same name.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets

from federated stored procedures.

This message is returned when db2fedgentf cannot uniquely identify the table function to drop. This can happen when there are multiple table functions with the same name.

No table function was dropped.

User response: Use -tfsn option with db2fedgentf instead of the -tfn option to drop the table function. The specific name of the table function is unique.

DBT4017N The following error occurred while db2fedgentf was creating the table function. SQLSTATE: *sqlstate*; SQLCODE: *sqlcode*; message text: *message-text*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

While creating a new table function, db2fedgentf encountered the given SQL error.

No table function was created.

User response:

1. Search for information about the SQL error in the DB2 Information Center.
2. Respond to the SQL error.
3. Call db2fedgentf again.

DBT4018N db2fedgentf encountered an internal error while creating the table function. The command that db2fedgentf was running is: *command*. The value that was returned from that command is: *return-code*. Diagnostic information is in the file called *file-name*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

To perform its tasks, the db2fedgentf utility executes multiple database commands and invokes the JAVA compiler 'javac' to compile the JAVA file for the table function. The db2fedgentf utility will get JDK path by reading DBM configuration parameter JDK_PATH. This message is returned when one of those tasks encountered an error.

No table function was dropped or created.

User response:

1. Review the diagnostic information in the named file.
2. Check DBM configuration parameter JDK_PATH and ensure it is set correctly.
3. Respond to the errors in the diagnostic information.
4. Rerun the db2fedgentf command.

DBT4022N The table function to drop was not specified. Both of the following options were missing: "-tfn", and "-tfsn".

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

When you call the db2fedgentf command to drop a table function for a given federated stored procedure, you must identify the table function using one of the following options:

- The table function name using the "-tfn" option
- The table function specific name using "-tfsn" option

For an explanation of the syntax of the db2fedgentf command, use the following command: "db2fedgentf -h".

User response: Call the db2fedgentf command again, specifying the table function using either the table function name or the table function specific name.

DBT4023N db2fedgentf could not drop the table function with specific name *specific-name* because that table function was not created using db2fedgentf.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures. You can use the db2fedgentf utility to drop only those table functions that were created using the db2fedgentf utility.

This message was returned because the specified table function was not created using db2fedgentf.

No table function was dropped.

User response: Use other methods, such as the DROP statement, to drop the table function.

DBT4024N db2fedgentf could not create a table function for the specified procedure *procedure-name* because there are multiple procedures with the same name.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

This message is returned when db2fedgentf cannot uniquely identify the procedure. This can happen when there are multiple procedures with the same name.

No table function was created.

User response: Use both -stpn and -stpc to determine a unique procedure.

DBT4025N db2fedgentf could not create a table function for the specified procedure, because a table function has already been created for that procedure. The specific name of the existing table function is: *schema-name. specific-name*

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

When using db2fedgentf, you can create only one table function for each federated stored procedure.

No table function was created.

User response: Drop the existing table function for the specified procedure and run db2fedgentf again.

DBT4026I The db2fedgentf utility successfully created the table function named *table-function-name* that you can use to access result sets that are returned from the federated stored procedure named *federated-stored-procedure*.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

This message is returned when the db2fedgentf utility successfully creates a table function.

User response: Use the new table function to access result sets that are returned from the federated stored procedure.

DBT4027N The db2fedgentf command failed because the utility encountered an error while allocating an environment handle.

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The db2fedgentdf utility creates Java table functions. To accomplish this, the utility must perform a variety of database tasks. This message is returned when the db2fedgentf utility encounters an internal DB2 database error while performing those tasks. Specifically, this message is returned when the db2fedgentf utility fails to allocate an environment handle.

User response: Run the db2fedgentf utility again:

1. Review available diagnostic information, such as the db2diag log, to determine why the db2fedgentf utility was unable to allocate an environment handle.
2. Resolve the problem or problems that caused the failure. For example, if the root cause of the handle allocation error was insufficient system memory, reconfigure the system to make more memory available.

3. Rerun the db2fedgentf command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support.

DBT4028N **The db2fedgentf command failed because the db2fedgentf utility could not determine the directory under which the Software Developer's Kit (SDK) for Java is installed from the database manager configuration parameter JDK_PATH.**

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

To perform its tasks, the db2fedgentf utility executes multiple database commands and invokes the Java compiler 'javac' to compile the Java file for the table function. The db2fedgentf utility determines where the Java SDK installed by reading the JDK_PATH configuration parameter.

This message is returned when the db2fedgentf utility cannot determine the location of the Java SDK from the JDK_PATH configuration parameter for some reason, such as because the configuration parameter is NULL.

User response:

1. Set the JDK_PATH database manager configuration parameter to the directory under which the Java SDK is installed.
2. Run db2fedgentf again.

DBT4029W **The db2fedgentf successfully created the table function, but failed to delete the Java source file: *Java-source-file-name*.**

Explanation: You can use the db2fedgentf utility to create or drop table functions that access result sets from federated stored procedures.

The db2fedgentdf utility creates Java table functions. To accomplish this, the utility performs the following tasks:

- Creates a Java file that contains the source for the new table function
- Compiles that generated Java file
- Registers the new table function
- Deletes the generated Java source file

This message is returned when the db2fedgentf utility was supposed to delete the generated Java source file, but the utility failed to delete the file.

User response: Optional: You can delete the generated Java source file manually.

Because the db2fedgentf utility always generates Java source files with unique names, choosing not to delete the generated source file that the utility failed to delete will not cause problems with future db2fedgentf use. However, it is recommended that you review any available diagnostic information (such as the db2diag log files or operating system diagnostic files) to determine why the utility failed to delete the file. The root cause of this failure to delete the Java source file by the db2fedgentf utility might also cause problems with other activities.

Chapter 92. DBT5000 - DBT5499

DBT5000I The db2checkSD utility completed successfully. The specified database can be used in a DB2 pureScale environment. The output log file is named *file-name*.

Explanation: Before converting your instance to a DB2 pureScale instance type, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when the db2checkSD utility did not find any database objects or features in the specified database that are not supported in a DB2 pureScale environment.

User response: Optional: review the contents of the output log file.

You can now convert your instance to a DB2 pureScale instance type.

DBT5001W The db2checkSD utility completed with some warnings. The database or databases can be used in a DB2 pureScale environment, but some database objects or features will have to be converted or disabled. The output log file is named *file-name*.

Explanation: Before converting your instance to a DB2 pureScale instance type, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when the db2checkSD utility found some database objects or features in the specified database that are not supported in a DB2 pureScale environment which will be automatically converted or disabled when you convert the instance to a DB2 pureScale instance type.

User response:

1. Review the contents of the db2checkSD output log file.
2. Investigate any issues identified in the db2checkSD output log file that are of concern to you.
3. Convert the instance to a DB2 pureScale instance type.

DBT5002N The db2checkSD utility completed with errors. The database or databases cannot be used in a DB2 pureScale environment. The output log file is named *file-name*.

Explanation: Before converting your instance to a DB2 pureScale instance type, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when the db2checkSD utility found some database objects or features in the specified database that are not supported in a DB2 pureScale environment and which cannot be automatically converted or disabled when you convert the instance to a DB2 pureScale instance type.

User response:

1. Review the contents of the db2checkSD output log file.
2. Convert or disable any unsupported database objects or features that db2checkSD identified.
3. Run the db2checkSD utility again.

DBT5003N The db2checkSD command failed because the following invalid parameter was specified: *parameter-name*.

Explanation: Before converting your instance to a DB2 pureScale instance type, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when an invalid parameter was passed to the db2checkSD command.

User response: Run the db2checkSD command again, specifying valid parameters.

DBT5004N The db2checkSD command failed because a database was not specified and the -e parameter was not specified.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

You can pass the name of a specific database to be checked, or you can use the -e parameter to specify that the db2checkSD utility should check all databases in the current DB2 instance. This message is returned when no database or databases were specified to be examined by the db2checkSD utility.

User response: Run the db2checkSD command again, specifying which database or databases to check in one of the following ways:

- Specify one specific database.

- Pass the `-e` parameter to cause the `db2checkSD` utility to examine all databases in the current DB2 instance.

DBT5005N **The db2checkSD command failed because no output log file name was specified.**

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the `db2checkSD` utility.

The `db2checkSD` creates an output log file that contains detailed information about the investigation of the database structure, metadata, objects, and features that the `db2checkSD` utility performed. You must specify the name of the output log file using the `-l` parameter.

This message is returned when no name for the `db2checkSD` output log file was specified using the `-l` parameter with the `db2checkSD` command.

User response: Run the `db2checkSD` command again, specifying an output log file name using the `-l` parameter.

DBT5006N **The db2checkSD command failed because no value was specified for the following parameter: *parameter-name*.**

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the `db2checkSD` utility.

This message is returned when a parameter is passed to the `db2checkSD` command without a value.

User response: Run the `db2checkSD` command again, specifying a value for the named parameter.

DBT5007N **The db2checkSD command failed because the following parameter was specified more than once: *parameter-name*.**

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the `db2checkSD` utility.

This message is returned when a parameter is passed to the `db2checkSD` command more than once.

User response: Run the `db2checkSD` command again, specifying the named parameter only once.

DBT5008N **The db2checkSD command failed because the length of the specified log file name is too long.**

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the `db2checkSD` utility.

The `db2checkSD` creates an output log file that contains detailed information about the investigation of the database structure, metadata, objects, and features that the `db2checkSD` utility performed. You must specify the name of the output log file name using the `-l` parameter. The log file name cannot be longer than 256 characters.

This message is returned when an output log file name was specified using the `-l` parameter with the `db2checkSD` command that is longer than 256 characters.

User response: Run the `db2checkSD` command again, specifying an output log file name that is not longer than 256 characters.

DBT5009N **The db2checkSD command failed because the db2checkSD utility is not supported with the current version of DB2 database.**

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the `db2checkSD` utility. The `db2checkSD` utility is supported with DB2 Version 9.7 for Linux, UNIX, and Windows and later.

This message is returned when an attempt is made to use the `db2checkSD` utility to examine a DB2 database that is older than DB2 Version 9.7.

User response:

1. Upgrade the database to DB2 Version 9.7 or later.
2. Run the `db2checkSD` command against the database again once the database is at Version 9.7 or higher.

DBT5010N **The db2checkSD utility failed because the utility could not confirm that the current user has sufficient authority to run the db2checkSD utility.**

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the `db2checkSD` utility.

The following authority is required to run the `db2checkSD` command: SYSADM or higher. This message is returned when a user without SYSADM

authority or higher attempts to run the db2checkSD command.

User response: Respond to this message with one of the following actions:

- Request that your authorization be upgraded to SYSADM authorization or higher and then rerun the db2checkSD command.
- Request that a user with SYSADM authorization or higher run the db2checkSD command.

DBT5011N The db2checkSD utility failed because there are table spaces in the database which are not in the 'NORMAL' state. The db2checkSD utility generated a user script named *script-name*.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

The db2checkSD utility performs the following tasks:

1. The db2checkSD utility identifies database objects and features that are not supported in a DB2 pureScale environment by examining database structures and metadata, as well as running queries against the database.
2. The db2checkSD utility generates a user script that you can run to prepare the database to be upgraded to a DB2 pureScale environment.

This message is returned when the db2checkSD utility found some table spaces that are in a state other than 'NORMAL', which would prevent the database from being upgraded successfully to a DB2 pureScale environment.

User response:

1. Identify the table spaces that are not in 'NORMAL' state by running the user script that was generated by the db2checkSD utility.
2. Move the table spaces that are not in a 'NORMAL' state to a 'NORMAL' state.
3. Run the db2checkSD utility against the database again.
4. If the db2checkSD utility succeeds, you can upgrade the database to a DB2 pureScale environment.

DBT5021N The db2checkSD utility found the following configuration that is not supported in a DB2 pureScale environment: the instance is configured to run more concurrently active databases than is supported.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not

supported in a DB2 pureScale environment using the db2checkSD utility.

The db2checkSD utility performs the following tasks:

- The db2checkSD utility identifies database objects and features that are not supported in a DB2 pureScale environment by examining database structures and metadata, as well as running queries against the database.
- The db2checkSD utility generates a user script that you can run to prepare the database to be upgraded to a DB2 pureScale environment.

This message is returned when the NUMDB database manager configuration parameter is set to a value that is larger than the allowed range for this DB2 database product in a DB2 pureScale environment.

User response:

- Change the NUMDB database manager configuration parameter to a value in the supported range for the DB2 product in a DB2 pureScale environment.
- Run the db2checkSD utility against the database again.
- If the db2checkSD utility succeeds, you can upgrade the database to a DB2 pureScale environment.

DBT5022N The db2checkSD command failed because both a database name was specified and the -e parameter was specified.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

You can pass the name of a specific database to be checked, or you can use the -e parameter to specify that the db2checkSD utility should check all databases in the current DB2 instance.

User response: Run the db2checkSD command again, specifying which database or databases to check in one of the following ways:

- Specify one specific database.
- Pass the -e parameter to cause the db2checkSD utility to examine all databases in the current DB2 instance.

DBT5025N The db2checkSD utility completed with errors. The output log file is named *log-file-name*.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility. The db2checkSD utility identifies these unsupported database objects and features by examining database structures and metadata, and

running queries against the database.

This message is returned in the following situations, among others:

- An attempt was made to run the db2checkSD command from a computer that is configured to run only a cluster caching facility (CF).
- The db2checkSD utility was unable to connect to one or more of the databases.
- The db2checkSD utility encountered an unexpected error condition while running queries against the database.

User response: Respond to this message in one of the following ways:

- Verify that you can connect to the database successfully.
- If the db2checkSD command was issued on a computer that is configured to run only a CF, run the command from any computer that is configured to run a DB2 member instead.
- If the db2checkSD command was not issued on a computer that is configured to run only a CF, perform the following procedure:
 1. Collect the following files:
 - The output log file created by the db2checkSD utility
 - The generated user script created by the db2checkSD utility
 - The db2diag log
 2. Contact IBM support to investigate the problem.

DBT5026I The db2checkSD utility completed processing for database *dbname*.

Explanation: Before converting your instance to a DB2 pureScale instance type, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

This message is returned after the db2checkSD utility has finished checking each database, when it is invoked with the -e option to check all databases in the instance.

User response: Optional: Review the contents of the log file when the db2checkSD utility has finished checking all databases in the instance.

DBT5027W The db2checkSD utility did not check any databases. The output log file is named *file-name*.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

The db2checkSD utility identifies these unsupported database objects and features by examining database structures and metadata, as well as running queries against the database. This message is returned when the db2checkSD utility found an empty database directory or when the db2checkSD utility could not find a local database directory.

User response:

1. Catalog the databases that you want to check using the db2checkSD utility.
2. Run the db2checkSD command again.

DBT5029N The db2checkSD utility could not open or write to the file named *file-name*.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

The db2checkSD utility generates the following files:

- An output log file that you can review for details about the db2checkSD utility's investigations
- A user script that you can use to identify database objects or features that must be removed or disabled before upgrading the database or databases to a DB2 pureScale environment

This message is returned when the db2checkSD utility was unable to create the output log file or the generated user script.

User response:

1. Give the db2checkSD utility write permission for the path where you want the db2checkSD utility to create the output log file and the generated user script.
2. Run the db2checkSD utility again.

DBT5033W The db2checkSD utility identified one or more features that will be disabled when updating the instance to a DB2 pureScale environment.

Explanation: Before upgrading a database to a DB2 pureScale environment, you can identify any database objects or features in the database that are not supported in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when the db2checkSD utility found some database objects or features in the specified database that are not supported in a DB2 pureScale environment which will be automatically converted or disabled when you upgrade the database to a DB2 pureScale environment.

User response:

1. Review the contents of the db2checkSD output log file.
2. Investigate any issues identified in the db2checkSD log that are of concern to you.
3. Upgrade the database to a DB2 pureScale environment.

DBT5034I **The db2checkSD utility has already verified that database *dbname* can be used in a DB2 pureScale environment.**

Explanation: The db2checkSD utility has already been run for the database and all checks were successful. The checks were not performed again.

User response: No further action is needed.

DBT5035N **The db2checkSD command failed because the db2checkSD utility was not able to access the database manager configuration file.**

Explanation: You can verify that a given database can be used in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when the db2checkSD utility is not able to open the database manager configuration file.

User response: Use commands such as db2 get dbm configuration to get more information about the error that occurs when DB2 attempts to read the database manager configuration file.

DBT5036N **The db2checkSD command failed because either the db2checkSD utility was unable to connect to the database or the db2checkSD utility was unable to access the database configuration file.**

Explanation: You can verify that a given database can be used in a DB2 pureScale environment using the db2checkSD utility.

This message is returned when the db2checkSD utility is not able to open the database configuration file.

User response: Respond to this error by performing the following troubleshooting steps:

- Verify that you can connect to the database.
- Get more information about the error that occurs when DB2 attempts to read the database configuration file by performing one or both of the following actions:
 - Connect to the database
 - Run the GET DATABASE CONFIGURATION command for the database

DBT5037N **The db2checkSD utility failed because the database is configured as an HADR standby database.**

Explanation: The db2checkSD utility verifies that a database can be used in a DB2 pureScale environment. An HADR standby database must be dropped before the instance is converted to a DB2 pureScale instance and reinitialized in the DB2 pureScale environment.

User response: Drop the standby database, convert both instances to DB2 pureScale instances, then reinitialize the standby database.

Chapter 93. DBT5500 - DBT5999

DBT5500N The current user does not have sufficient authority to run the db2ckupgrade utility.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

The following authority is required to run the db2ckupgrade command: SYSADM or higher. This message is returned when a user with insufficient authority attempts to run the db2ckupgrade command.

The command failed.

User response: Do one of the following:

1. Request that your authorization be upgraded to SYSADM authorization or higher and then rerun the db2ckupgrade command.
2. Request that a user with SYSADM authorization or higher run the db2ckupgrade command.

DBT5501N The db2ckupgrade utility could not open or write to the file named *log-file-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

The db2ckupgrade utility generates an output log file that you can review for details about the processing done by the utility.

This message is returned when the db2ckupgrade utility was unable to open or write to the output log file.

User response:

- Give the db2ckupgrade utility write permission for the path where you want the db2ckupgrade utility to create the output log file.
- Run the db2ckupgrade utility again.

DBT5502N The db2ckupgrade utility failed because there was an incorrect number of parameters specified.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the number of parameters specified is incorrect.

User response: Run the db2ckupgrade command again, specifying valid parameters.

DBT5503N The db2ckupgrade utility failed because the following invalid parameter was specified: *parameter-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when an invalid parameter was passed to the db2ckupgrade command.

The command failed.

User response: Run the db2ckupgrade command again, specifying valid parameters.

DBT5504N The db2ckupgrade utility failed because a database was not specified and the -e parameter was not specified.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

You can pass the name of a specific database to be checked, or you can use the -e parameter to specify that the db2ckupgrade utility should check all databases in the current DB2 instance. This message is returned when no database or databases were specified to be examined.

User response: Run the db2ckupgrade command again, specifying which database or databases to check in one of the following ways:

- Specify one database name.
- Pass the -e parameter to cause the db2ckupgrade utility to examine all databases in the current DB2 instance.

DBT5505N The db2ckupgrade utility failed because no output log file name was specified.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

The db2ckupgrade creates an output log file that contains detailed information about the investigation that the db2ckupgrade utility performed. You must specify the name of the output log file name using the -l parameter.

This message is returned when no name for the db2ckupgrade output log file was specified.

User response: Run the db2ckupgrade command again, using the -l parameter, to specify an output log file name.

DBT5506N The db2ckupgrade utility failed because no value was specified for the following parameter: *parameter-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when a parameter is passed to the db2cupgrade command without a value.

User response: Run the db2cupgrade command again, specifying a value for the named parameter.

DBT5507N The db2cupgrade utility failed because the following parameter was specified more than once: *parameter-name*.

Explanation: The db2cupgrade command is used to verify if a database can be upgraded.

This message is returned when a parameter is passed to the db2cupgrade command more than once.

User response: Run the db2cupgrade command again, specifying the named parameter only once.

DBT5508I The db2cupgrade utility completed successfully. The database or databases can be upgraded.

Explanation: The db2cupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2cupgrade utility verifies that all upgrade conditions are met. You can now upgrade the database.

User response:

1. Optional: Review the contents of the output log file.
2. You can now upgrade the database.

DBT5509N The db2cupgrade utility failed because it could not connect to the database. Database: *database-name*.

Explanation: The db2cupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2cupgrade utility is unable to connect to the specified database.

The command failed.

User response:

1. Ensure that the specified database name is correct
2. Ensure that the database manager is running.
3. Try the command again.

DBT5510W The db2cupgrade utility found that the database contains the following types of objects, which are not supported by the version of DB2 database to which upgrading is being considered: XML global variables; or compiled SQL functions that use XML parameters or that return XML types.

Explanation: You can verify that a given database can successfully be upgraded to a later version of DB2 database using the db2cupgrade utility.

This message is returned when the db2cupgrade utility finds database objects which are not supported in the version of DB2 database to which you are upgrading. Specifically, this message is returned when the following database objects are in a database being checked:

- Global variables of type XML
- Compiled SQL functions that use parameters of type XML or that return XML types

If you proceed to upgrade the database to the target version of DB2 database, these objects will be invalidated during database upgrade. You will be unable to use these database objects until you upgrade to a version of DB2 database that supports those database objects.

User response: To be able to use the database objects that are invalidated during upgrade, upgrade to a release and fix pack of DB2 database that supports XML global variables and compiled SQL functions that use XML parameters or that return XML types. When you upgrade to a fixpack that supports those database objects, the objects will be automatically revalidated the first time they are referenced after the database has been upgraded.

DBT5511N The db2cupgrade utility failed because a parameter was too long. Parameter: *parameter-name*. Max length: *max-length*.

Explanation: The db2cupgrade command is used to verify if a database can be upgraded.

This message is returned when one of the parameters passed to the db2cupgrade utility is longer than the maximum allowed length for that parameter.

User response: Run the db2cupgrade command again, specifying the parameter with the correct length.

DBT5512E The db2cupgrade utility determined that the following workload management object cannot be upgraded because the ID of that object conflicts with a system-reserved ID. Object name: *object-name*.

Explanation: Some workload management objects are assigned default, system-reserved identifiers (ID).

This message is returned when the db2cupgrade utility determines that an existing workload management object has an ID that happens to fall in the range of default, system-reserved IDs. The database cannot be upgraded while the workload management object with the conflicting ID exists in the database.

User response:

1. Extract the DDL statements that would reproduce your workload management objects by using the db2look utility, specifying the -wlm option.

2. Drop all of the workload management objects from the database.
3. Verify whether the database can be upgraded by using the db2ckupgrade utility, and resolve any other problems that are blocking the database from being upgraded.
4. After the db2ckupgrade utility determines that the database can be upgraded, perform the following steps:
 - a. Upgrade the database.
 - b. Recreate the workload management objects in the upgraded database by running the DDL statements that the db2look utility generated.

DBT5513N The db2ckupgrade utility failed because the SYSCATSPACE table space requires more space in order to complete the upgrade.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

Typically, amount of free space defined by SYSCATSPACE should be at least the same as the current amount of used space.

This message is returned when SYSCATSPACE has less than 50% free pages left.

User response:

- Increase the size of existing containers. You can also add additional containers although this might trigger data rebalancing.
- Optional: You can reduce the size of the containers after upgrade.

DBT5514N The db2ckupgrade utility failed because the database is in a restore pending state.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the database being checked is found to be in a restore pending state.

User response: Complete the restore and run the db2ckupgrade command again.

DBT5515N The db2ckupgrade utility failed because the database is in a backup pending state.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the database being checked is in the process of having a backup taken to provide a starting point for roll-forward recovery, but is still in a backup pending state.

User response: Wait until the backup has completed

and run the db2ckupgrade command again.

DBT5516N The db2ckupgrade utility failed because the database is in a rollforward pending state.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the database being checked is enabled for roll-forward recovery and it has been restored but not rolled forward.

User response:

1. Roll forward the database or indicate that you do not wish to roll forward by using the ROLLFORWARD command. Note that if you do not roll forward the database, the records written since the last backup of the database will not be applied to the database.
2. Run the db2ckupgrade command again.

DBT5517N The db2ckupgrade utility failed because the database is in an inconsistent state.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the database being checked is found to be in an inconsistent state.

The possible causes of this inconsistent state include

- The database is online and SQL has been issued which modified data in the database.
- The database is online and HADR has been enabled.

User response:

1. Perform a clean shutdown of the database.
2. After shutdown, if HADR has been enabled on the database, issue the STOP HADR command on the database.
3. Re-issue the db2ckupgrade command.

DBT5518W The db2ckupgrade utility completed with one or more warnings but the database can still be upgraded. Log file: *log-file*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility completes with one or more warnings but the database can still be upgraded

User response:

1. Optional: Review the contents of the output log file.
2. You can now upgrade the database.

DBT5519N The db2ckupgrade utility failed because the the utility failed to deactivate the database. **SQLCODE:** *sqlcode*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

The db2ckupgrade command attempts to deactivate the specified database and stop all necessary database services. This message is returned when the db2ckupgrade utility failed to deactivate the database.

User response: Perform the following:

1. Review the SQLCODE to determine the reasons why the database could not be deactivated and take corrective action.
2. Run the db2ckupgrade command again.

DBT5520N The db2ckupgrade utility failed because an error occurred while setting the bypass flag for Query Patroller. **SQLCODE:** *sqlcode*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility attempted to set the DB2_QP_BYPASS_COST variable but failed.

User response: Perform the following:

1. Review the SQLCODE and take corrective action.
2. Run the db2ckupgrade command again.

DBT5521N The db2ckupgrade utility failed because the utility was not invoked on the instance being upgraded.

Explanation: The db2ckupgrade command verifies whether a database can be upgraded.

The db2ckupgrade command inspects a database in an instance from a previous level to determine whether it can be upgraded to a more recent level.

This message is returned when the db2ckupgrade utility was not invoked on the instance from a previous level.

User response: Run the db2ckupgrade command again on the instance from the previous level.

DBT5522N The db2ckupgrade utility failed because the database has one or more tables in a load pending state.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when a previous LOAD attempt on one or more tables in the database did not complete successfully. The upgrade cannot proceed until the LOAD operation is restarted or terminated.

User response: Restart or terminate the previously failed LOAD operation on this table by issuing LOAD with the RESTART or TERMINATE option respectively.

To find out what tables are in a load pending state run the following command:

```
select tabname from SYSIBMADM.ADMINTABINFO
where load_status is not NULL
```

DBT5523N The db2ckupgrade utility failed because the database has one or more tables in a redistribute pending state.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This messages is returned when the REDISTRIBUTE utility has not completed for one or more tables in the database. The upgrade cannot proceed until the REDISTRIBUTE is finished.

User response: Perform one of the following:

1. If the REDISTRIBUTE utility is in progress, wait until it completes. You can use the LIST UTILITIES command to monitor the progress of the REDISTRIBUTE utility.
2. If a previous REDISTRIBUTE operation failed and left the table in redistribute pending state, issue the REDISTRIBUTE utility again with the CONTINUE or ABORT option and let it complete.

DBT5524N The db2ckupgrade utility failed because upgrading databases enabled for VARCHAR2 support from DB2 Version 9.5 is not supported.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility determines that the database to be upgraded is enabled for VARCHAR2 support. Upgrading from DB2 Version 9.5 with VARCHAR2 is not supported.

User response: No action is needed. The database cannot be upgraded.

DBT5525N The db2ckupgrade utility failed with an unhandled error.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned in the following situations, among others:

- An attempt was made to run the db2ckupgrade command from a computer that is configured to run only a cluster caching facility (CF).
- The db2ckupgrade utility encountered an unhandled error condition while examining database structures and metadata, and running queries against the database.

User response: Respond to this message in one of the following ways:

1. If the db2ckupgrade command was issued on a computer that is configured to run only a CF, run the command from any computer that is configured to run a DB2 member instead.
2. If the db2ckupgrade command was not issued on a computer that is configured to run only a CF, collect the output log file created by the db2ckupgrade utility, the generated user script created by the db2ckupgrade utility, and the db2diag log. Then contact IBM support to investigate the problem.

DBT5526W The db2ckupgrade utility did not check any databases. The output log file is named *file-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility finds an empty database directory or when the db2ckupgrade utility could not find a local database directory.

User response: Perform the following actions:

- Catalog the databases that you want to check using the db2ckupgrade utility.
- Run the db2ckupgrade command again.

DBT5527N The db2ckupgrade command failed because the database manager could not be started. Return Code: *return-code*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

The db2ckupgrade command must start the database manager in order to perform the verifications, but was unable to do so.

The command failed.

User response: Determine the cause of the failure using the return code and try again.

DBT5528N The db2ckupgrade utility failed because the I/O write operations for the database are suspended or are being suspended.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility cannot operate because I/O writes are suspended for the database.

User response: If write operations for the database are in the process of being suspended, monitor the state of the database using the suspend_io configuration parameter, and wait until the SET WRITE SUSPEND operation completes before you continue, then reissue the db2ckupgrade command again.

DBT5529N The db2ckupgrade utility did not complete successfully. The database cannot be upgraded. The output log file is named *log-file-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility has completed but not all upgrade conditions were met. The database cannot be upgraded.

User response:

1. Review the contents of the output log file and fix any issues.
2. Run the db2ckupgrade command again.

DBT5530N The db2ckupgrade command failed because the tablespace is not in normal state. Tablespace: *tablespace-name*. State: *state*. Member: *member-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned then the db2ckupgrade utility attempted to access a table space which is in an invalid state, for which the intended access is not allowed.

The possible causes of the inconsistent state include:

- The database is online and SQL has been issued which modified data in the database.
- The database is online and HADR has been enabled.

User response: Perform the following actions:

1. Perform a clean shutdown of the database.
2. After shutdown, if HADR has been enabled on the database, issue the STOP HADR command on the database.
3. Re-issue the db2ckupgrade command.

DBT5531W The db2ckupgrade utility failed because it could not find any local database to process.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the -e parameter is specified but the db2ckupgrade utility is unable to find any local databases.

The command failed.

User response: Check whether any database names exist in the system database directory. If there are no database names in the system database directory, then either no databases exist or the database names have not been cataloged.

If a database name appears in the system database directory and the entry type is INDIRECT, ensure that

the database exists in the specified local database directory

DBT5532N The db2ckupgrade utility failed because there are one or more MQTs that depend on system views.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This messages is returned when the db2ckupgrade utility finds MQTs that depend on system views. During database upgrade all system views will be recreated, so any MQTs that depend on the system views will be automatically dropped.

User response: Perform the following actions:

- Generate and backup the DDL needed to create the MQTs
- Drop the MQTs
- Upgrade the database
- Recreate the MQTs

DBT5533N The db2ckupgrade utility failed because the utility failed to activate the database. **SQLCODE:** *sqlcode*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned when the db2ckupgrade utility could not perform the verification because it failed to activate the database.

User response:

1. Review the SQLCODE to determine the reasons why the database could not be activated and take corrective action.
2. Run the db2ckupgrade command again.

DBT5534W The db2ckupgrade command found a database object which has a dependency on system built-in routines that are discontinued. **Object name:** *schema-name.object-name*. **Object type:** *object-type*. **Discontinued routine upon which the object depends:** *schema-name.routine-name*.

Explanation: Some system built-in routines such as table functions, views, or stored procedures, are discontinued in the target version of the database upgrade due to the introduction of new functionality or changes to the existing functionality.

The database upgrade to that target version drops any system built-in routines that are discontinued. To continue to use database objects that have a dependency on these system built-in routines, you have to alter these database objects and remove the dependency before the database upgrade.

This message is returned when the db2ckupgrade command detects a database object such as a view, trigger, or a user-defined stored procedure, that has a dependency on system built-in routines that are discontinued in the target version of the database upgrade.

User response: Perform one of the following actions:

- If you no longer need the database object which depends on system built-in routines that are discontinued, drop this object.
- Alter the database object so that the object does not depend on system built-in routines that are discontinued. Use the replacement routines provided for system built-in routines that are discontinued.

DBT5536W The db2ckupgrade utility did not find any audit policies associated with the database *database-name*.

Explanation: You can verify whether a database can be upgraded by using the db2ckupgrade utility. Before upgrading to a new version of DB2 database, you can identify any database objects or features in the database that are not supported in the new version using the db2ckupgrade utility.

The security administrator uses audit policies to configure the audit facility to gather information about a database. This message is returned when the db2ckupgrade utility finds audit policies that are associated with a database (entries in the SYSCAT.AUDITUSE catalog view with a blank OBJECTTYPE), but none of the entries is associated with the current database. There is no entry that has an OBJECTNAME that matches the current database name.

An audit policy is not required to upgrade the database.

User response:

Optional:

1. Audit the database using the following command:
AUDIT DATABASE USING POLICY <policy-ID>
2. Run the db2ckupgrade command again.

If you do not audit the database and rerun db2ckupgrade, the SYSCAT.AUDITUSE catalog view will be cleared when the database is upgraded.

DBT5537I The db2ckupgrade utility has completed processing for database *database-name*.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded.

This message is returned after the db2ckupgrade utility has finished checking each database, when it is invoked

with the -e option to check all databases in the instance.

User response: Optional: Review the contents of the log file when the db2ckupgrade utility has finished checking all databases in the instance.

DBT5538N The db2ckupgrade utility found an unexpected mount point configuration.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded. When upgrading from DB2 V9.8 the database upgrade moves certain table space files to a new location. If the table space containers use an unexpected mount point configuration the table space data cannot be moved by the upgrade processing.

User response: Contact IBM support. If possible, have the details from the log file associated with the db2ckupgrade command available.

DBT5539N The db2ckupgrade utility found a discrepancy in the database control files.

Explanation: The db2ckupgrade command is used to verify if a database can be upgraded. When upgrading from DB2 9.8 the files under /MEMBERnnnn/sqldbdir are consolidated to a single location. If the files in /sqldbdir are not identical on all members the upgrade processing cannot consolidate them.

User response: Contact IBM support. If possible, have the log file associated with the db2ckupgrade command available.

Chapter 94. DBT6000 - DBT6499

DBT6000N The global database configuration file is not valid. Configuration file: *file-name*.

Explanation: The utility is unable to update the global database configuration file because the file was created using a different DB2 version or is corrupt.

User response:

1. Ensure that the DB2 product that you are using on the secondary system is at the same level as the one that you are using on the primary system.
2. If the file is corrupt, restore the original split mirror copy of the database if the copy is still available. If the configuration file is still corrupt after the restore operation, take a new split mirror copy of the database.
3. Rerun the utility.

DBT6001N The db2relocatedb command failed because the database manager was unable to relocate one or more database storage paths. Reason code = *reason-code*.

Explanation: You can rename or relocate a database using the db2relocatedb command. You specify the details of the new configuration for the database by passing the name of a configuration file that contains configuration keyword-value pairs.

This message is returned when an attempt is made to change one or more storage paths of a database by calling the db2relocatedb command with a configuration file that includes the STORAGE_PATH keyword, and the database manager is unable to move the storage path or paths.

The reason code indicates the reason why the database manager was unable to move the storage path or paths:

1. One or both of the storage group control files, SQLSGF.1 or SQLSGF.2 was not successfully upgraded.
2. The old storage path that was specified with the STORAGE_PATH keyword in configuration file does not match any existing storage paths.

User response: Respond to this error according to the reason code:

Reason code 1:

1. Validate that the storage control files, SQLSGF.1 and SQLSGF.2, are consistent and upgraded properly, using the db2dart tool with /TSF option.

Reason code 2:

1. Determine which storage paths exist using one of the following methods:
 - Refer to the db_storage_path monitor element
 - Call the ADMIN_GET_STORAGE_PATHS table function
2. Run the db2relocatedb command again, specifying a configuration file that contains only valid old storage paths with the STORAGE_PATH keyword.

DBT6100I The db2ida command completed successfully.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

Command processing completed without any errors.

User response: No further action is required.

DBT6101W The db2ida -start command was attempted on *number-of-partitions* database partitions. Number of database partitions that were successfully started: *number-started*. Number of database partitions that had already been started: *number-already-started*. Number of database partitions that were not successfully started: *number-not-started*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The db2ida command did not start all of the embedded processes of the in-database analytics provider on all database partitions. The in-database analytics functionality will be unavailable until all of the embedded processes of the in-database analytics provider are running on all database partitions.

User response: Correct the problem on the failed database partitions and resubmit the command.

DBT6102W The db2ida -stop command was attempted on *number-of-partitions* database partitions. Number of database partitions that were successfully stopped: *number-stopped*. Number of database partitions that had already been stopped: *number-already-stopped*. Number of database partitions that were not successfully stopped: *number-not-stopped*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The db2ida command did not stop all of the embedded processes of the in-database analytics provider on all database partitions. If you are planning to perform maintenance that involves the in-database analytics provider, the planned maintenance might not be successful because some of the related processes are still running.

User response: Correct the problem on the failed database partitions and resubmit the command.

DBT6103N The db2ida command failed because the following invalid parameter was specified: *parameter-name*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

An invalid parameter was passed to the db2ida command.

User response: Run the db2ida command again, specifying valid parameters.

DBT6104N The db2ida command failed because the specified provider is not supported. Specified provider name: *provider-name*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

This message is returned if an unsupported value is specified with the -provider parameter.

User response: Run the db2ida command again, specifying a valid value for the -provider parameter.

DBT6105N The db2ida command could not start the embedded processes of the specified provider because an option of a related registry variable is not set. Specified provider: *provider-name*. Registry option: *option-name*. Registry variable: *registry-variable*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

In order for the db2ida command to successfully start the embedded processes of the in-database analytics provider, the given registry variable must be set.

User response:

1. Set the given option for the given registry variable.
2. Run the db2ida command again.

DBT6106N The db2ida command failed because the database manager did not have enough system resources.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

There were insufficient system resources for the database manager to run the db2ida command on the local or remote database partition.

User response:

1. Check the system workload status.
2. Release some system resources.
3. Run the db2ida command again.

DBT6107N The db2ida command failed to start the embedded processes for the specified provider because a required executable was not found. Specified provider name: *provider-name*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

A required executable to start the embedded process could not be located under the database manager instance directory.

User response:

1. Ensure that the DB2 database product is properly installed.
2. Ensure that the database manager instance has been created successfully.
3. Run the db2ida command again.

DBT6108N The db2ida command failed to start the embedded processes for the specified provider because the database manager was not running on the database partition. Specified provider name: *provider-name*.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The embedded processes for an in-database analytics provider cannot be started if the database manager is not running on all database partitions.

User response:

1. Determine on which database partitions the database manager is not up and running.
2. Start the database manager on all database partitions using the db2start command.
3. Run the db2ida command again.

DBT6109N The db2ida command timed out.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

This message is returned when the db2ida command fails to receive a response from the embedded processes of the in-database analytics provider.

User response:

1. Increase the value of the start_stop_time database manager configuration parameter.
2. Run the db2ida command again.

DBT6110N The db2ida command failed because of unexpected system errors.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

A system error occurred. Some possible reasons for this error are:

- The physical host for the database partition cannot be reached because of a network problem.
- The physical host for the database partition has powered down.
- The remote command service is not installed or running on the host from which the command was issued.

User response:

1. Ensure that the hosts for the DB2 instance are running.

2. Ensure that the remote command service is running on the host.
3. Run the db2ida command again.

DBT6111N The db2ida command failed because the user who ran the command did not have the required authorization.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

This message is returned when the db2ida command did not have the required authorization to run necessary DB2 database-related utilities. The db2ida command runs with the same authorization as the user who executed the command.

User response: Run the db2ida command again, as a user with the required authorization.

DBT6112N The db2ida command failed to start the embedded processes of the specified in-database analytics provider because the db2ida utility was unable to load a library of the provider.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The library for the in-database analytics feature is part of the installation of the provider. Some possible reasons for this error are:

- The library for the embedded process could not be found in the library path.
- The library could not be loaded by the database manager.

User response:

1. Ensure that the library path for the in-database analytics library of the provider exists.
2. Run the db2ida command again.

DBT6113N The db2ida command failed to invoke the entry function for the embedded process.

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The main entry function could not be found in the library, and the embedded process failed.

User response: Respond to this error by performing the following troubleshooting steps:

- Resubmit the action.
- Reinstall the in-database analytics provider library.

- Contact technical support for the in-database analytics provider.

DBT6114N **The db2ida command failed to start the embedded process because the communication buffers could not be allocated.**

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The DB2 database manager instance does not have enough shared memory to allocate for communication buffers between the DB2 database server and the embedded process. One possible reason for this error is that the DB2_FMP_COMM_HEAPSZ registry variable was set too low.

User response: Respond to this error by performing the following troubleshooting steps:

1. Increase the value of the DB2_FMP_COMM_HEAPSZ registry variable.
2. Restart the DB2 instance.
3. Run the db2ida command again.

DBT6115N **The db2ida command failed to stop the embedded process because the embedded process was still involved in an uncompleted transaction.**

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

The embedded process was still involved in an uncompleted transaction. The database manager could not stop the embedded process.

User response: Respond to this error in one of the following ways:

- Wait until the transaction completes and then issue the db2ida command again.
- Force the embedded process to stop by issuing the db2ida command with the -stopforce parameter.

DBT6116W **The db2ida command could not complete successfully because of concurrent start or stop processing on the embedded process.**

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

Other activities to start or stop the embedded process were running concurrently in the DB2 instance. Concurrent activities include the following:

- Forced termination of an unresponsive embedded process
- Automatic restart of a previously terminated embedded process
- Other running db2ida commands

User response:

1. Wait until other activities complete.
2. Run the db2ida command again.

DBT6117N **The db2ida command could not complete successfully because a communication error occurred.**

Explanation: You can stop or restart embedded processes for an in-database analytics provider without stopping and restarting the DB2 database manager instance using the db2ida command.

A TCP/IP communication error occurred while the db2ida command was attempting to establish a connection with a database partition.

User response: Complete the following troubleshooting steps:

- Ensure that the database partition has the proper authorization defined in the rhosts or host.equiv file.
- Ensure that the application is not using more than $500 + (1995 - 2 * \text{total_number_of_dbpartitions})$ file descriptors simultaneously.
- Ensure that all environment variables are defined, and that the profile is written in Korn Shell script format.
- Ensure that all of the hosts that are defined in the db2nodes.cfg file in the sqllib directory are defined on the network and are running.
- Ensure that the DB2FCMCOMM registry variable is set correctly.

Chapter 95. DBT7000 - DBT7499

DBT7000I **db2caem completed. Output path:**
output-path.

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the db2caem utility has completed.

User response: Review the files generated by the db2caem utility in the named output path.

DBT7001E *utility-name:* **The db2caem command failed because invalid syntax was specified.**

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when a syntax error is made in the db2caem command.

You can view the db2caem command syntax by running the following command:

db2caem -h

User response: Review the db2caem command syntax, and then rerun the command with valid syntax.

DBT7002E *utility-name:* **The following input value is invalid: value.**

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when invalid input data is specified.

You can view the db2caem command syntax by running the following command:

db2caem -h

User response: Review the db2caem command syntax, and then rerun the command with valid input.

DBT7003E *utility-name:* **The db2caem command failed because no database name was specified.**

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when no database is specified.

You can view the db2caem command syntax by

running the following command:

db2caem -h

User response: Review the db2caem command syntax, and then rerun the command specifying a database with the -d parameter.

DBT7004E *utility-name:* **The db2caem command failed because no query statement options were specified and no event monitor options were specified.**

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You must specify either query statement options or event monitor options when you issue the command. This message is returned when no query statement options are specified and no event monitor options are specified.

You can view the db2caem command syntax by running the following command:

db2caem -h

User response: Review the db2caem command syntax, and then rerun the command specifying either query statement options or event monitor options.

DBT7005E *utility-name:* **The db2caem command failed because one of the parameters -u or -p was not specified.**

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a user id with the -u parameter and the password for that user ID with -p parameter.

You must specify both -u and -p or neither. This message is returned when only one of -u or -p is specified.

You can view the db2caem command syntax by running the following command:

db2caem -h

User response: Review the db2caem command syntax, then rerun the command specifying a user ID with the -u parameter and a password with the -p parameter.

DBT7006E *utility-name:* **The db2caem command failed because one or more required event monitor parameters were not specified.**

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can view the db2caem command syntax by running the following command:

```
db2caem -h
```

User response: Review the db2caem command syntax, and then rerun the command specifying all of the required parameters.

DBT7007E *utility-name: The db2caem command failed because both the -st parameter and the -sf parameter were specified.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a SQL statement for which activity event monitor data should be collected using the -st parameter. You can specify a file that contains a SQL statement for which activity event monitor data should be collected using the -sf parameter. The -st parameter and the -sf parameter cannot both be specified. This message is returned when the -st parameter and the -sf parameter are both specified.

You can view the db2caem command syntax by running the following command:

```
db2caem -h
```

User response: Review the db2caem command syntax, and then rerun the command specifying either the -st parameter or the -sf parameter, not both.

DBT7008E *utility-name: The db2caem command failed because both query statement options and event monitor options were specified.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify either query statement options or event monitor options when you issue the command, but not both options. This message is returned when both query statement options and event monitor options are specified.

You can view the db2caem command syntax by running the following command:

```
db2caem -h
```

User response: Review the db2caem command syntax, and then rerun the command specifying either query statement options or event monitor options.

DBT7009E *utility-name: The db2caem command failed because the specified database name database-name is too long.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the database name that is specified is longer than length limits imposed on DB2 database names.

User response: Rerun the command specifying a valid database name.

DBT7010E *utility-name: The db2caem command failed because the specified user ID user-ID is too long.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the user ID that is specified is longer than length limits imposed on user names by DB2 database.

User response: Rerun the command specifying a valid user id.

DBT7011E *utility-name: The db2caem command failed because the specified password password is too long.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the password that is specified is longer than length limits imposed on passwords by DB2 database.

User response: Rerun the command specifying a valid password.

DBT7012E *utility-name: The db2caem command failed because the SQL statement in the file specified with the -sf parameter is longer than the SQL limit. Length of specified statement: length.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the SQL statement in the file specified with the -sf parameter is longer than length limits imposed on SQL statements in DB2 databases. This message can also be returned when the terminating character is missing or invalid in the specified SQL statement.

User response: Rerun the command specifying a valid SQL statement.

DBT7013E *utility-name: The db2caem command failed because there is a problem with the path of the specified SQL file or the compilation environment file. Specified file path: path.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a file that contains a SQL statement for which activity event monitor data should be collected using the -sf parameter. You can specify the compilation environment information in a file using the -compenv parameter. This message is returned when the db2caem utility cannot access the file path for one of these specified files.

There are several reasons why the db2caem utility might not be able to access the specified file path:

- The specified file path does not exist.
- The db2caem utility runs with the same authority as the user who executes the db2aem command. This means that you must be able to access the file path specified with the -sf parameter and the -compenv parameter. This message can be returned when the user who executes the db2caem command does not have access to the specified file path.

User response: Rerun the command specifying a file path that exists and that you can access.

DBT7014E *utility-name: The db2caem command failed because the db2caem utility could not open the SQL file specified with the -f parameter. File name: file-name. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a file that contains a SQL statement for which activity event monitor data should be collected using the -sf parameter. This message is returned when the db2caem utility cannot open the file specified with the -sf parameter.

There are several reasons why the db2caem utility might not be able to open the specified file:

- The specified file does not exist.
- The db2caem utility runs with the same authority as the user who executes the db2aem command. This means that you must be able to open the file specified with the -sf parameter. This message can be returned when the user who executes the db2caem command does not have authority to open the specified file.

User response: Rerun the command specifying a file path that exists and that you can open.

DBT7015W *utility-name: The db2caem command failed because the db2caem utility could not open the following internal file: file-name. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the db2caem utility cannot access an internal, temporary file at run time. The db2caem utility can complete successfully despite this warning.

User response: No user response is required.

DBT7016E *utility-name: The db2caem command failed because the length of the specified SQL file or compilation environment file is too long. Specified file name: file-name. Length of specified file name: length*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a file that contains a SQL statement for which activity event monitor data should be collected using the -sf parameter. You can specify a file that contains compilation environment information using the -compenv parameter. This message is returned when the full name, including the path, of the file specified with the -sf parameter or the -compenv parameter is too long.

User response: Rerun the command specifying a SQL file name that is a valid length.

DBT7017E *utility-name: The db2caem command failed because the specified table space name is too long. Specified table space name: table-space-name.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a table space name for which activity event monitor data should be collected using the -tbspname parameter. This message is returned when the table space name specified with the -tbspname parameter is too long.

User response: Rerun the command specifying a table space name that is a valid length.

DBT7018E *utility-name: The db2caem command failed because the specified activity event monitor name is too long. Specified activity event monitor name: event-monitor-name.*

Explanation: You can capture activity event monitor

data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify an activity event monitor name for which activity event monitor data should be collected using the -actevm parameter. This message is returned when the activity event monitor name specified with the -actevm parameter is too long.

User response: Rerun the command specifying an activity event monitor name that is a valid length.

DBT7019E *utility-name: The db2caem command failed because the specified application ID is too long. Specified application ID: application-ID.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify an application ID for which activity event monitor data should be collected using the -appid parameter. This message is returned when the application ID specified with the -appid parameter is too long.

User response: Rerun the command specifying an application ID that is a valid length.

DBT7020E *utility-name: The db2caem command failed because the specified output path does not exist or cannot be fully qualified. Specified output path: path.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify the path in which the db2caem utility will generate output files using the -o parameter. This message is returned when the db2caem utility cannot access the path specified with the -o parameter.

There are several reasons why the db2caem utility might not be able to access the specified path:

- The specified path does not exist.
- The db2caem utility runs with the same authority as the user who executes the db2caem command. This means that you must be able to access the path specified with the -o parameter. This message can be returned when the user who executes the db2caem command does not have access to the specified output path.

User response: Rerun the command specifying a output path that exists and that you can access.

DBT7021E *utility-name: The db2caem command failed because the db2caem utility encountered an out-of-memory error. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

An out-of-memory error is a system error that prevents the command from completing. This message is returned when the db2caem utility encounters an out-of-memory error while attempting to capture activity event monitor data.

User response:

1. Free memory that is being used by other processes.
2. Rerun the db2caem command.

DBT7022E *utility-name: The db2caem command failed because the db2caem utility was unable to create the following directory: directory-name. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify the path in which the db2caem utility will generate output files using the -o parameter. The db2caem utility creates subdirectories in the output directory. This message is returned when the db2caem utility is unable to create subdirectories in the output directory.

User response:

1. Determine why db2caem is unable to create subdirectories in the output directory by reviewing diagnostic information in the db2diag log files or the db2caem.log file, if it was created.
2. Resolve the problem or problems that are preventing the db2caem utility from creating subdirectories in the output directory.
3. Execute the db2caem command again.

DBT7023E *utility-name: The db2caem command failed because the specified output path name is too long. Specified output path name: file-name.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify the path in which the db2caem utility will generate output files using the -o parameter. This message is returned when the name of the output path specified with the -o parameter is too long.

User response: Rerun the command specifying a output path name that is a valid length.

DBT7024E *utility-name:* **The db2caem command failed because the db2caem utility encountered an internal database error.**
Return code: *return-code*.

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility performs a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility fails to perform one of these internal database tasks such as failing to determine the path for a specified file.

The db2caem utility cannot proceed.

User response:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support

DBT7025E *utility-name:* **The db2caem command failed because querying the table *table-name* returned unexpected rows.**
Number of rows returned: *num-rows-returned*.

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility queries the specified table to collect information about the query statement that is specified with the -st parameter, that is contained in the file specified with the -sf parameter, or that is obtained from an existing event monitor specified with the -actevm parameter. The db2caem utility found unexpected rows that were returned by this query. Thus, an error condition might exist that you cannot resolve.

User response: Rerun the db2caem command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support

DBT7026E *utility-name:* **The db2caem command failed because querying the table *table-name* returned no rows.** **Return code:** *return-code*.

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility queries the specified table to collect information about the query statement that is specified with the -st parameter, that is contained in the file specified with the -sf parameter, or that is obtained

from an existing event monitor specified with the -actevm parameter. The db2caem utility expects at least one row to be returned by this query. If no rows are returned by this query, then an error condition might exist that you cannot resolve.

This message can be returned when an invalid terminating character is used in the SQL file specified with the -sf parameter or when an invalid event monitor name is specified with -actevm parameter.

User response:

1. Determine what the internal error is, based on the return code.
2. Resolve the source of the internal error.
3. Execute the db2caem command again.

DBT7027E *utility-name:* **The db2caem command failed because the db2caem utility was unable to obtain the current timestamp.**
Return code: *return-code*.

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify the path in which the db2caem utility will generate output files using the -o parameter. The db2caem utility creates subdirectories in the output directory. The db2caem utility generates names for these subdirectories using a naming format that includes a timestamp. This message is returned when the db2caem utility is unable to obtain a timestamp to use in the output directory names.

The db2caem utility cannot proceed.

User response:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support

DBT7028E *utility-name:* **The db2caem command failed because the DB2 database manager encountered an error while processing the specified input parameters and the specified SQL statement.** **Handle type:** *handle-type*.
Return code: *return-code*.

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the database server cannot process the input options specified with the db2caem command and the SQL statement specified with the -st parameter or contained in the SQL file specified with the -sf parameter.

This message can be returned when there is a syntax error with the specified SQL statement.

User response:

1. Review available diagnostic information.
2. Correct the invalid input parameters or invalid SQL syntax.
3. Rerun the db2caem command.

DBT7029E *utility-name: The db2caem command failed because the DB2 database manager encountered a handle error. Handle type: handle-type. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility encounters an internal DB2 database error while performing those tasks.

User response: Perform the activity event monitor data activity again:

1. Review available diagnostic information.
2. Resolve the problem or problems that caused the failure.
3. Rerun the db2caem command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support.

DBT7030E *utility-name: The db2caem command failed because the DB2 database manager encountered an error while allocating an environment handle. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility encounters an internal DB2 database error while performing those tasks.

User response: Perform the activity event monitor data activity again:

1. Review available diagnostic information.
2. Resolve the problem or problems that caused the failure.
3. Rerun the db2caem command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:

- db2diag logs
2. Contact IBM support.

DBT7031E *utility-name: The db2caem command failed because the DB2 database manager encountered an error. Handle type: handle-type. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility encounters an internal DB2 database error while performing those tasks.

User response: Perform the activity event monitor data activity again:

1. Review available diagnostic information.
2. Resolve the problem or problems that caused the failure.
3. Rerun the db2caem command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support.

DBT7032W *utility-name: The db2caem command failed because the DB2 database manager encountered a warning. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility encounters an internal DB2 database warning while performing those tasks.

This message can be returned when the db2caem utility cannot retrieve an environment variable.

User response: Rerun the db2caem command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support .

DBT7033E *utility-name: The db2caem utility encountered an error while processing the specified SQL statement. SQL statement: statement. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

This message is returned when the database server encountered an error while processing the SQL statement specified with the -st parameter or contained in the SQL file specified with the -sf parameter.

This message can be returned for several reasons, including the following reasons:

- There is a syntax error with the SQL statement.
- An invalid terminating character is used in the SQL statement.

User response:

1. Review available diagnostic information.
2. Correct the invalid input parameters or invalid SQL syntax.
3. Rerun the db2caem command.

DBT7034E *utility-name: The db2caem utility encountered an error while running the db2exfmt command. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. For example, the db2caem command executes the db2exfmt command to format the contents of the EXPLAIN table for the specified SQL statement. This message is returned when the db2caem utility encounters an error while running the db2exfmt utility.

The db2caem utility can proceed despite the db2exfmt error.

User response: No response is required.

DBT7035W *utility-name: The db2caem utility encountered an error while running the GET CONNECTION ATTRIBUTE command. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. For example, the db2caem utility uses the connection attribute to verify the connection to the database while the command runs. This message is returned when the db2caem utility encounters an error while running the GET CONNECTION ATTRIBUTE command.

The fact that the db2caem utility failed to obtain the connection attribute might indicate a problem with the database connection.

User response: When the db2caem command completes, review output information and diagnostic information to confirm that the activity event monitor data was successfully collected.

DBT7036W *utility-name: The db2caem utility encountered an error while disconnecting from the database. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility encounters an error while disconnecting from the database specified with the -d parameter.

The db2caem utility can complete successfully despite failing to disconnect from the database.

User response: No response is required.

DBT7037W *utility-name: The db2caem utility encountered an error while freeing a CLI handle. The SQLFreeHandle function failed with the following return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. This message is returned when the db2caem utility encounters an error while freeing internal resources.

The db2caem utility can complete successfully despite failing to free this CLI handle.

User response: No response is required.

DBT7038I *utility-name: The db2caem utility successfully connected to the following database: database-name.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. For example, the db2caem utility must connect to the specified database.

User response: No response is required.

DBT7039I *utility-name: The db2caem utility successfully disconnected from the following database: database-name.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data. For example, the db2caem utility must connect to and disconnect from the specified database.

User response: No response is required.

DBT7040E *utility-name: The db2caem command failed because the specified statement contains a keyword that is not explainable or will not be executed by the db2caem utility. statement. Keyword: keyword*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

To collect activity event monitor data, the db2caem utility must be able to explain or execute the specified statement. This message is returned when the specified SQL statement cannot be explained or will not be executed. This can happen for the following reasons:

- A statement might not be explainable if the statement creates, alters, or drops database objects, for example.
- The db2caem utility will not execute CALL statements because executing CALL statements might have undesirable side-effects in the specified database.

The db2caem utility cannot collect the requested activity event monitor data.

User response:

1. Modify the SQL statement specified with the -st parameter or included in the SQL file specified with the -sf parameter so that the statement is explainable and does not include the CALL keyword.
 2. Rerun the db2caem command with the altered statement.
-

DBT7041I *utility-name: The db2caem utility is connecting to the following database: database-name.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data, including connecting to the specified database.

User response: No response is required.

DBT7042I *utility-name: The SQL statement statement is being issued.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data, including issuing the indicated statement with the -st option is specified or the -sf option is specified.

User response: No response is required.

DBT7043I *utility-name: The db2caem utility is disconnecting from the following database: database-name.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

The db2caem utility must perform a variety of database tasks to collect the activity event monitor data, including connecting to and disconnecting from the specified database.

User response: No response is required.

DBT7044E *utility-name: The db2caem command failed because the db2caem utility could not obtain information from the operating system about the following file: file-name. Return code: return-code.*

Explanation: You can capture activity event monitor data, which can be used to analyze SQL statement performance, using the db2caem utility.

You can specify a file that contains a SQL statement for which activity event monitor data should be collected using the -sf parameter. You can specify a file that contains compilation environment information using the -compenv parameter. This message is returned when the db2caem utility cannot obtain information from the operating system about the file specified with the -sf parameter or the file specified with the -compenv parameter.

There are several reasons why the db2caem utility might not be able to access the specified file:

- The specified file does not exist.
- The db2caem utility runs with the same authority as the user who executes the db2caem command. This means that you must be able to open the file specified with the -sf parameter or the file specified with the -compenv parameter. This message can be returned when the user who executes the db2caem command does not have authority to open the specified file.

User response: Rerun the command specifying a file path that exists and that you can access.

DBT7045E **The db2cklog command failed because the db2cklog utility cannot open the current log file. Log file name: *log-file-name*.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

The db2cklog utility runs with the same authority as the user who ran the db2cklog command. This message is returned when the db2cklog utility cannot verify a log file because the db2cklog utility cannot open the log file. The db2cklog utility might not have been able to open the log file for the following reasons:

- The log file exists in a path that the db2cklog utility cannot access
- The db2cklog utility does not have read permission for the log file

User response: Rerun the command specifying an archived log file or range of archived log files that exist in a path that you can access.

DBT7046E **The db2cklog command failed because the db2cklog utility encountered a memory allocation failure.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility failed to allocate enough memory to perform the tasks associated with validating log files.

User response: Perform the log file validation operation again:

1. Review available diagnostic information.
2. Resolve the problem or problems that caused the failure.
3. Rerun the db2cklog command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support.

DBT7047E **The db2cklog command failed because the db2cklog utility cannot close the current log file.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility is able to open a log file to verify the file, but the

db2cklog utility is then unable to close the log file. The db2cklog utility might not have been able to close the log file for the following reasons:

- The log file no longer exists in the same location as when the db2cklog command was run
- Another file operation has affected the log file

User response: Rerun the db2cklog command ensuring that the specified archived log file or range of archived log files, and the directories in which the log file or files are located, persist for as long as the db2cklog utility is running.

DBT7048E **The db2cklog utility determined that the current log file is invalid.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

The db2cklog utility examines many different aspects of the specified log file or files and determines whether the files are in a valid format and whether the contents are valid. This message is returned when the db2cklog utility finds that one of the specified log files is invalid and therefore cannot be used for rollforward recovery purposes.

User response:

1. Determine whether you have a secondary copy of the specified archived log file by referring to the value of the logarchmeth2 configuration parameter, which determines how your database server handles secondary archived log files.
2. If you have a secondary copy of the archived log file, perform the following steps:
 - a. Verify the validity of the secondary log file by running the db2cklog command against the secondary copy.
 - b. If the secondary copy of the log file passes validation, replace the log file that the db2cklog utility cannot read with the valid, secondary copy of the log file.
3. If you have only one copy of the archived log file, or if the db2cklog utility failed to verify the secondary log file, the log file is beyond repair and cannot be used for rollforward recovery purposes. Establish a new, more recent recovery point that does not depend on the unusable log file for rollforward recovery by making a full database backup as soon as possible.

DBT7049E **Log file validation failed because the log file header contains an invalid number of log pages. The expected number of pages is *expected-number*. The actual number of pages is *actual-number*.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for

rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7050E **Log file validation failed because both copies of the log file header for the specified log file are invalid.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7051E **Log file validation failed because the specified log file contains an undetermined error.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7052E **Log file validation failed because the specified log file contains no log pages.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7053E **Log file validation failed because the specified log file contains an invalid log page followed by another invalid log page.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7054E **Log file validation failed because the specified log file contains an invalid log page followed by an outdated log page.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7055E **Log file validation failed because the specified log file contains an invalid log page followed by a partial log page.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7056E **Log file validation failed because the specified log file contains an invalid log page followed by a complete log page.**

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility

finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7057E Log file validation failed because the specified log file contains a complete log page followed by an invalid log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7058E Log file validation failed because the actual number of log pages in the specified log file does not match the expected number of log pages. Actual number of log pages: *actual-number*. Expected number of log pages: *expected-number*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7059E Log file validation failed because the check fields in the log file header do not match.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7060E Log file validation failed because the log file header contains an invalid log file number. The expected log file name is: *expected-log-file-name*. The number that was expected in the header is: *log-file-number*. The log file name that was found is: *log-file-name*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7065E Log file validation failed because the current log page is empty and not all log pages of the specified log file have been processed. Current log page: *log-page-id*. Total number of log pages: *total-log-pages*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7066E Log file validation failed because the log page is marked as truncated but it is not the last page in the specified log file.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this

message. The details are useful for helping IBM support personnel only.

DBT7067E Log file validation failed because a log page ends with an unexpected log record type. Header flag of the log page: *header-flag*. Unexpected log record type: *log-record-type*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7068E Log file validation failed because the log record type is invalid. Log record type: *log-record-type*. The maximum type value allowed is: *maximum-log-record-type*. The expected type value is: *expected-log-record-type*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7069E Log file validation failed because the specified log file contains the invalid log record size *log-record-size*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7070E Log file validation failed because the log record size *log-record-size* exceeds the maximum log record size in the specified log file.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7071E Log file validation failed because the size of the log record is invalid. Log record size: *log-record-size*. The minimum valid log record size is *minimum-log-record-size*. The maximum valid log record size is *maximum-log-record-size*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7072E Log file validation failed because the log record size is invalid. Log record size: *log-record-size*. The expected log record size is *expected-log-record-size*. The maximum valid log record size is *maximum-log-record-size*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7074E Log file validation failed because an unexpected log record type was found in the specified log file. Log record type found: *log-record-type*. The expected log record type is: *expected-log-record-type*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7075E The db2cklog command failed because an unexpected token, *token*, was found following the input parameter *input-parameter*. Supported tokens include: *input-parameter-tokens*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when an invalid parameter or value is specified with the db2cklog command.

User response:

1. [Optional] View the command syntax by running the db2cklog command without specifying any parameters.
2. Rerun the command specifying only supported command syntax.

DBT7076E The db2cklog command failed because the db2cklog utility could not access the archive log path that was specified with the ARCHLOGPATH parameter. Specified archived log path: *archived-log-path*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility. You can specify the absolute or the relative path of the archived log files using the ARCHLOGPATH parameter.

The db2cklog utility runs with the same authority as the user who ran the db2cklog command. This message can be returned when the person who ran the db2cklog command does not have read permission for the specified path.

User response:

1. Work with a system administrator to acquire read permission for the archived log path.

2. Rerun the command specifying an archived log file path using the ARCHLOGPATH parameter that you can access.

DBT7077E The db2cklog command failed because the specified archived log file path is too long. Specified archived log file path: *path*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility. You can specify the absolute or relative path of the archived log files using the ARCHLOGPATH parameter with the db2cklog command.

This message is returned when the string passed as the value of the ARCHLOGPATH parameter with the db2cklog command is longer than a defined maximum length. The maximum length of archived log path that can be specified with the ARCHLOGPATH parameter varies, depending on the operating system.

User response: Respond to this message in one of the following ways:

- Run the db2cklog command from the directory in which the archived log file is located, and do not use the ARCHLOGPATH parameter at all.
- If the length of the absolute path of the archived log file is shorter than or equal to the maximum length of the string that can be passed with the ARCHLOGPATH parameter, rerun the db2cklog command specifying the absolute path with the ARCHLOGPATH parameter.
- If the length of the archived log path relative to the directory in which the db2cklog command is being run is shorter than or equal to the maximum length of the string that can be passed with the ARCHLOGPATH parameter, rerun the db2cklog command specifying the relative path using the ARCHLOGPATH parameter.

DBT7078E The db2cklog command failed because the current log file does not exist in the directory from which the db2cklog command was run or in the archived log file directory specified with the ARCHLOGPATH parameter. Log file name: *log-file-name*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

If a path is specified using the ARCHLOGPATH parameter with the db2cklog command, then the db2cklog utility will search for the specified archived log file or range of archived log files in the specified path. However, if the ARCHLOGPATH parameter is not passed to the db2cklog command, the db2cklog utility will search for the specified archived log file or

range of archived log files in the directory from which the db2cklog command was run. This message is returned when the db2cklog utility cannot find the specified archived log file in either of these locations.

User response: Respond to this message in one of the following ways:

- Run the db2cklog command from the directory in which the archived log file is located.
- Rerun the db2cklog command, specifying the location of the archived log file using the ARCHLOGPATH parameter.

DBT7079E The db2cklog command failed because of an internal error.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

To perform the tasks associated with validating log files, the command depends on various system resources and software components. This message is returned when there is a problem with one of these resources or components.

User response: Perform the log file validation operation again:

1. Review available diagnostic information.
2. Resolve the problem or problems that caused the failure.
3. Rerun the db2cklog command.

If the problem persists, perform the following steps:

1. Collect the following diagnostic information:
 - db2diag logs
2. Contact IBM support.

DBT7080E Log file validation failed because the specified log file contains a partial log page followed by an invalid log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This message is returned when the db2cklog utility finds that a log file is invalid. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7081W Log file validation encountered a warning because the specified log file contains an empty log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This warning is returned when the db2cklog utility finds information that could indicate that a log file is active. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7082W Log file validation encountered a warning because the specified log file contains a partial log page followed by a complete log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This warning is returned when the db2cklog utility finds information that could indicate that a log file is active. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7083W Log file validation encountered a warning because the specified log file contains a partial log page followed by an outdated log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This warning is returned when the db2cklog utility finds information that could indicate that a log file is active. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7084W Log file validation encountered a warning because the specified log file contains a partial log page followed by a partial log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This warning is returned when the db2cklog utility finds information that could indicate that a log file is active. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7085W Validation of the specified log file encountered a warning because an outdated log page follows a log page of the following type: *log-page-type*.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This warning is returned when the db2cklog utility finds information that could indicate that a log file is active. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7086W Log file validation encountered a warning because the specified log file contains a complete log page followed by an invalid log page.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility.

This warning is returned when the db2cklog utility finds information that could indicate that a log file is active. The detailed information printed in this message is useful for helping IBM support personnel investigate how the log file became invalid.

User response: You do not need to respond to this message. The details are useful for helping IBM support personnel only.

DBT7087W The db2cklog utility found information in the current log file which indicates that the log file might be active.

Explanation: You can determine whether an archived log file or a range of archived log files are valid for rollforward recovery using the db2cklog utility. The db2cklog utility should only be used to validate archived log files, not active log files.

Active database log files are the primary and secondary log files that are currently needed for recovery and rollback. Active logs might be open and database logging information might still be written to active logs. Archived database log files are closed and no longer needed for normal processing. These files are

retained, or archived, for use in rollforward recovery.

When an attempt is made to run the db2cklog utility against an active log file, this warning is returned because the db2cklog utility might not be able to accurately check the validity of an active log file.

User response: Rerun the command specifying only archived log files, not active log files.

DBT7088W The following file was not found or could not be executed during diagnostic data collection: *file*

Explanation: When you run the db2support -install command, diagnostic data is collected from several sources on a host machine. Sources for diagnostic data include log files stored on the file system and output returned by executing other, diagnostic commands. If diagnostic data for one of these sources cannot be obtained, this warning is returned.

User response: If possible, always run the db2support -install command with root authority. Running the command with root authority ensures that the most complete set of diagnostic data possible is collected.

If you are already running the command with root authority, you do not need to respond to this message. The details are useful for alerting IBM support personnel to missing diagnostic data only.

DBT7089E The command syntax specified is unsupported or invalid.

Explanation: When running the db2support command included with the installation image, only a subset of the db2support syntax is available for you to use, and the command syntax you include must be specified in a valid form.

User response: Issue db2support -h to see the supported syntax, or see "db2support - Problem analysis and environment collection tool command" in the DB2 Information Center, then reissue the command.

DBT7090E The host name *hostname* is invalid.

Explanation: A host name cannot be validated as specified and no diagnostic data has been collected for this host.

User response: Verify that the host name is spelled correctly. Include only valid host names with the -host option, then reissue the command.

DBT7091E The host *hostname* does not support or allow an ssh connection.

Explanation: The db2support command requires an ssh connection to collect diagnostic data from hosts other than the local host. If an ssh connection cannot be established, diagnostic data collection fails.

An ssh connection can fail because ssh is not configured on the host or it is configured in such a way so as to prevent the db2support command from connecting.

User response: Verify that ssh is configured correctly, then reissue the command.

DBT7092W The db2support command was not invoked with root authority. As a result, some diagnostic data might not get collected.

Explanation: When you run the db2support -install command, diagnostic data is collected from several sources on a host machine, such as log files stored on the file system and output returned by executing other, diagnostic commands. If the user invoking the db2support command does not have the required level of authorization to access one of these sources, some diagnostic data will not get collected.

User response: For the most complete collection of diagnostic data, always invoke the db2support -install command with root authority. If you cannot run the command with root authority, you do not need to respond to this message. If necessary, this message will alert IBM support personnel that diagnostic data might be missing.

DBT7093E The host *hostname* cannot be reached.

Explanation: To collect diagnostic information on a remote host, the db2support command must be able to connect to the host over the network.

There are several reasons why a host cannot be reached:

- The network cabling has been disconnected, the network adapter is unresponsive, or the network configuration is preventing network access.
- The host machine is offline (the host is not powered up or in the process of rebooting, for example).

User response: Resolve the problem that is preventing the host from being reachable, then reissue the command.

DBT7094E The db2support command cannot be executed as specified.

Explanation: You must to be the instance owner to invoke the db2support command with the options specified.

User response: As instance owner, invoke the db2support command from the DB2 install path.

DBT7095E The -install option cannot be specified together with other options except for the -host option.

Explanation: When running the db2support command included with the installation image, only a subset of the db2support syntax is available for you to use, and the command syntax you include must be specified in a valid form. When you specify the -install option, it can be specified together with the -host option only.

User response: Issue db2support -h to see the supported syntax, or see "db2support - Problem analysis and environment collection tool command" in the DB2 Information Center, then reissue the command.

DBT7096W The instance name cannot be determined.

Explanation: When you run the db2support -install command, diagnostic data is collected from several sources on a host machine, including output returned by executing the db2instance -list command. If the instance name that must be specified with the db2instance -list command cannot be determined, this warning is returned.

User response: You do not need to respond to this message. The details are useful for alerting IBM support personnel to missing diagnostic data only.

Part 14. DQP Messages

Chapter 96. DQP0000 - DQP0499

DQP0001E The database name must be specified in the command.

Explanation: The command syntax requires that the database name be specified.

User response: See the Query Patroller documentation for a detailed description of the command.

Specify the database name and reissue the command.

DQP0002E Query Patroller server is unable to connect to database *databasename*.

Explanation: Query Patroller was not able to establish a connection to the *databasename* database.

User response: Ensure that the database name is correct and that the database manager is running.

See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

DQP0003E Unable to start the Query Patroller server.

Explanation: Query Patroller cannot be started.

User response: See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

DQP0004E Unable to stop the Query Patroller server.

Explanation: The Query Patroller server cannot be stopped.

User response: Verify that the Query Patroller server was actually running when you attempted to stop it. See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

DQP0005E Query Patroller configuration cannot be found for database *databasename*.

Explanation: The configuration cannot be found for database *databasename*.

User response: Issue the qpstop command to stop Query Patroller.

DQP0006E Unable to find the Java Properties file *filename*.

Explanation: The required properties file *filename* cannot be found.

User response: Check the sqllib/msg/<locale>/qp/

directory to determine if the Properties files exist.

If the problem persists, contact IBM Support.

DQP0007E An internal error *error* has occurred.

Explanation: An internal error occurred during processing.

User response: See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

DQP0008E User *username* does not have a submitter profile.

Explanation: Each user is required to have a submitter profile.

User response: Request that the administrator define a submitter profile for the user.

DQP0009E File *filename* cannot be opened.

Explanation: An error occurred while attempting to open the specified file.

User response: Ensure that the file exists and that its permissions are correct.

See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

DQP0010E User *user-ID* does not have an effective submitter profile.

Explanation: The user *user-ID* does not have an effective Query Patroller submitter profile. This may occur due to one or more of the following reasons:

- The user may not have a submitter profile.
- The user and/or group submitter profiles(s) belonging to the user may have been suspended.

User response: Request that the database administrator create a submitter profile or have the submitter profile reactivated.

DQP0011E The submission preferences dialog cannot be displayed because neither *property-file1* nor *property-file2* was found.

Explanation: The properties file cannot be found. As a result, the submission preferences dialog cannot be displayed.

User response: Verify that the files exist.

DQP0012E The database *databasename* has not been setup for use with Query Patroller.

Explanation: To use Query Patroller on a database, Query Patroller needs to be installed on the server that contains the database. During installation, the databases whose queries are to be managed by Query Patroller are chosen and Query Patroller control tables and stored procedures are created in this database. If you need to manage queries of additional databases using Query Patroller, then you must run the qpsetup command for each database.

User response: Request that an administrator install Query Patroller on the server containing the database or, if this has already been done, request that the administrator run the qpsetup command to setup Query Patroller for the database.

If the qpsetup command has been run for the database and the problem persists, check that the file db2qp exists in INSTPATH/function and the file db2qpt exists in INSTPATH/function/unfenced on UNIX, or check that the file db2qp.dll exists on INSTPATH\function and the file db2qpt.dll exists on INSTPATH\function\unfenced on Windows. INSTPATH is the DB2 instance directory. If these files exist, see the qpdiag.log for possible cause of this message. If problem persists, contact IBM Support.

DQP0019E SYSADM authority is required to issue the *command-name* command.

Explanation: The command requires SYSADM authority.

User response: Request that the system administrator run the command.

DQP0020E Function not supported. Reason code = *reason-code*.

Explanation: The statement cannot be processed because it violates a restriction as indicated by the following reason code:

- 1 The connection codepage cannot be set because the functionality is not yet available.
- 2 The restriction is unknown.

User response: The action corresponding to the reason code is:

- 1 Update the Query Patroller client to a newer version.
- 2 If the problem persists, contact IBM Support.

DQP0021E The DB2 server and the Query Patroller server are at a different product level.

Explanation: The DB2 server and the Query Patroller server must be on the same product level.

User response: Ensure that your DB2 server and Query Patroller server are at the same product level. If you have applied a FixPak on the DB2 server before installing Query Patroller server, ensure that you reapply the FixPak.

See the Query Patroller documentation for a description of the product prerequisites.

DQP0024E An unexpected error *error* occurred.

Explanation: An unexpected error occurred.

User response: See the qpdiag.log file for possible causes of this error.

DQP0025E Invalid timestamp format *format*.

Explanation: The timestamp format that the user specified is invalid. The correct timestamp format is YYYY-MM-DD HH24:MI:SS.

User response: Refer to the Query Patroller documentation for a detailed description of the command.

Specify the correct timestamp format and reissue the command.

DQP0406E Unable to start the Query Controller as the maximum number of *maxnumber* Query Controllers is already running.

Explanation: There has been an attempt to start more than the maximum number of allowed Query Controllers.

User response: Use one of the Query Controllers that is currently running.

DQP0408E DBADM authority is required to run the command *commandname*.

Explanation: The command *command-name* requires DBADM authority.

User response: Request that the security administrator grant DBADM authority, and reissue the command.

DQP0409E The Query Patroller system settings do not exist.

Explanation: Query Patroller cannot be run when the system settings do not exist.

User response: Issue the qpsetup command to create the Query Patroller system settings.

See the Query Patroller documentation for a detailed description of the qpsetup command.

DQP0410E The submission preferences for the PUBLIC group do not exist.

Explanation: Query Patroller cannot be run when the submission preferences for the PUBLIC group do not exist.

User response: Recreate the submission preferences for the PUBLIC group and issue the qpstart command to start Query Patroller.

DQP0412I The database connection has been reestablished.

Explanation: The database connection has been reestablished.

User response: No action required.

DQP0413E The database connection has been lost. Query Patroller is terminated.

Explanation: The database connection has been lost. As a result, Query Patroller has terminated.

User response: Reestablish the database connection and issue the qpstart command to restart Query Patroller.

DQP0414E Query Patroller was not stopped because there are active queries.

Explanation: Query Patroller was not stopped because there are queued or active queries.

User response: Reissue the command using the FORCE option. If the active queries are forced, they will be in an inconsistent state until Query Patroller is restarted and the query recovery is complete.

DQP0415I The estimated cost of the query exceeds the maximum allowed for the user's submitter profile.

Explanation: The query will be held.

User response: Request that the database administrator run or cancel the query.

DQP0416E The estimated cost of the query exceeds the maximum query cost for the system.

Explanation: The query will be held.

User response: Request that the database administrator run or cancel the query.

DQP0417E The number of queries running has reached the maximum allowed for the system.

Explanation: The query will be placed in the queue.

User response: The query will run automatically when

the number of queries running drops below the maximum allowed for the system.

DQP0418E The number of queries running for user *username* is the maximum number allowed for the submitter profile.

Explanation: The query will be placed in the queue.

User response: The query will run automatically when the number of queries running drops below the maximum allowed for the system.

DQP0419E The number of queries running under query class *queryclassnumber* is the maximum allowed for the query class.

Explanation: The query will be placed in the queue.

User response: The query will run automatically when the number of queries running drops below the maximum allowed for the query class.

DQP0420I The sum of the estimated cost and the current system workload exceeds the maximum system workload allowed.

Explanation: The sum of the estimated cost of the query submitted and the current system workload is greater than the maximum system workload allowed.

User response: Increase the maximum system workload allowed, if desired.

DQP0421I The estimated cost of the query exceeds the maximum system workload allowed.

Explanation: The query is being held because its estimated cost exceeds the maximum system workload allowed.

User response: Request that the administrator release the query so that it can run, or increase the maximum system workload allowed.

DQP0422E DB2 Query Patroller license cannot be found.

Explanation: A valid license key for DB2 Query Patroller cannot be found or it has expired.

User response: Install a license key for the fully entitled version of the product. You can obtain a license key for the product by contacting your IBM representative or authorized dealer.

DQP0423E Query Patroller is already started.

Explanation: Unable to start Query Patroller as it is already running.

User response: To restart Query Patroller, issue the qpstop command and then issue qpstart.

DQP0424E **No start Query Patroller command was issued.**

Explanation: The start Query Patroller command has not been processed. It must be processed before a stop Query Patroller command is issued.

User response: Issue a start Query Patroller command, qpstart, and resubmit the current command.

DQP0431E **The DB2 Query Patroller service could not be started. Reason code *reason-code*.**

Explanation: The explanation corresponding to the reason code is:

1. The service does not exist.
2. The current user does not have proper authority to start or stop the service.
3. The service fails to start.
4. The logon information for the service is not correct.

User response: The action corresponding to the reason code is:

1. Ensure that the service exists. If the service does not exist, then it may have been removed manually or the installation of DB2 Query Patroller did not complete successfully. In either case, DB2 Query Patroller must be reinstalled.
2. Ensure that the logon information for the service is correct and issue the command again.
3. Issue the command again. If the problem persists, contact your system administrator or see the qpdiag.log file for possible causes of this message.
4. Ensure that the logon information for the service is correct and issue the command again.

DQP0432E **Unable to communicate with the DB2 Query Patroller Java process.**

Explanation: The command successfully created a Java process for DB2 Query Patroller, however, there was a communication error while trying to obtain output. The communication error may occur if a network error exists or if the Java process have been terminated abnormally.

User response: See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

Chapter 97. DQP1000 - DQP1499

DQP1001E The date or time specified is not valid.

Explanation: The date or time specified is in an incorrect format.

User response: Enter the correct format for the date and time.

DQP1002E The start date and time specified must occur before the end date and time.

Explanation: The specified start date and time must precede the end date and time.

User response: Specify a start date and time that occur before the end date and time.

DQP1003E The operation was aborted because it would violate the consistency of the date. `SQLSTATE = sqlstate`.

Explanation: The Query Patroller control tables have triggers to protect the consistency of the data contained in them. Such a trigger has been activated through an action that has been attempted. The operation was aborted due to a configuration consistency constraint violation.

User response: Examine the `qpschema.sql` file to find the trigger that signaled the `SQLSTATE`. Based on this examination, make the necessary adjustments to prevent the trigger from aborting the action again.

DQP1004E A memory allocation error occurred.

Explanation: There was not enough memory to continue processing.

User response: Ensure that your system has sufficient memory.

See the Query Patroller documentation for the system memory requirements. If the problem persists, contact IBM Support.

DQP1005E The time range must not be greater than two years.

Explanation: The time range specified is greater than two years.

User response: Specify a time range that is within the next two years and reissue the command.

DQP1006E SQL statement `sqlstmt` failed with `SQLCODE SQLCODE`.

Explanation: The SQL statement failed.

User response: Check the `qpdiag.log` file for possible causes of this message. If the problem persists, contact IBM Support.

DQP1008E SQL data type `datatype` is not supported.

Explanation: The SQL data type is not supported.

User response: See the Query Patroller documentation for information about supported SQL data types.

DQP1009E Result destination is too large to view.

Explanation: The Query Patroller Center and Query Patroller Command Line Processor cannot return your results because the results exceed the maximum size limit of the control table.

User response: Resubmit the query and ensure that the results are returned to your application rather than stored in a result table.

If your query generated a result table because it was run in the background, resubmit the query with your submission preferences set to wait until results are returned.

If your query generated a result table because it was held and then run, ask an administrator to raise the query cost of your submitter profile so that you can resubmit the query without having it held.

Alternatively, you can select directly from the result table, but the names of the columns in the result table do not match those of your query and the result table contains an additional column, `A0000`.

DQP1010E File `filename` cannot be accessed.

Explanation: The specified file cannot be accessed.

User response: Ensure that the file exists and that the file permissions are correct.

DQP1011E User `username` is not authorized to issue the command.

Explanation: The command cannot run because the user does not possess the necessary authorization level.

User response: Check the Query Patroller documentation for the required authorization. Request that the database administrator grant the user the required authority and reissue the command.

DQP1012E File `filename` does not exist.

Explanation: The specified file does not exist.

User response: Ensure that the file exists. If the problem persists, contact IBM Support.

DQP1023E **An error occurred while saving to the file *filename*.**

Explanation: Possible causes of this message include:

- The file does not have the correct file permissions.
- The file name does not conform to the file system convention.

User response: Possible solutions include:

- Verify the file permissions.
- Specify the correct file name.

DQP1024W **Creation, change, or removal of a query class will not take effect until the Query Patroller server is restarted.**

Explanation: You have just created, changed, or deleted a query class. This will change the behavior of how queries are processed by the Query Patroller server. This change in behavior will not occur until the next time Query Patroller is started.

User response: Request that an administrator restart the Query Patroller server by issuing the `qpstop` command followed by the `qpstart` command on the computer where the Query Patroller server is running.

DQP1025W **The update will not take effect until the Query Patroller server is restarted.**

Explanation: The settings will not take effect until the Query Patroller server is restarted.

User response: Restart the Query Patroller server to have the changes take effect.

DQP1026W **The update was successful. However, communication with the Query Patroller server has failed.**

Explanation: The update was successful. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP1028E **The network operation did not complete within a reasonable amount of time. The command cannot be completed.**

Explanation: The network operation did not complete within a reasonable amount of time. The command cannot be completed.

User response: Reissue the command when the network is not as busy. If problem persists, contact the system administrator to analyze if there is a network error.

Chapter 98. DQP2000 - DQP2499

DQP2020E The path *path* is invalid.

Explanation: The specified path is invalid.

User response: Verify the path and reissue the command.

DQP2101I Operator profile for user *username* was added successfully.

Explanation: An operator profile was created for the user.

User response: No action required.

DQP2102I Operator profile for group *groupname* was added successfully.

Explanation: An operator profile was created for the group.

User response: No action required.

DQP2103I Operator profile for user *username* was updated successfully.

Explanation: An existing user operator profile was updated.

User response: No action required.

DQP2104I Operator profile for group *groupname* was updated successfully.

Explanation: An existing group operator profile was updated.

User response: No action required.

DQP2105I Operator profile for user *username* was removed successfully.

Explanation: A user operator profile was removed.

User response: No action required.

DQP2106I Operator profile for group *groupname* was removed successfully.

Explanation: A group operator profile was removed.

User response: No action required.

DQP2107E Operator profile for user *username* does not exist.

Explanation: The user operator profile is not defined.

User response: Ensure that the user name specified is

correct and reissue the command.

DQP2108E Operator profile for group *groupname* does not exist.

Explanation: The group operator profile is not defined.

User response: Ensure that the group name specified is correct and reissue the command.

DQP2109E Operator profile for user *username* already exists.

Explanation: The user operator profile already exists.

User response: Ensure that the user name specified is unique and reissue the command.

DQP2110E Operator profile for group *groupname* already exists.

Explanation: The group operator profile already exists.

User response: Ensure that the group name specified is unique and reissue the command.

DQP2111I There are no operator profiles defined.

Explanation: There are no operator profiles defined to display.

User response: No action required.

DQP2112I Query class *queryclassID* was added successfully.

Explanation: A new query class was defined.

User response: No action required.

DQP2113I Query class *queryclassID* was updated successfully.

Explanation: An existing query class was updated.

User response: No action required.

DQP2114I Query class *queryclassID* was removed successfully.

Explanation: A query class was removed.

User response: No action required.

DQP2115E Query class *queryclassID* does not exist.

Explanation: The query class does not exist.

User response: Ensure that the query class ID specified is correct and reissue the command.

DQP2116E **The query class with the same maximum query cost value already exists.**

Explanation: The query class already exists.

User response: Specify a query class that does not already exist and reissue the command.

DQP2117I **There are no query classes defined.**

Explanation: There are no existing query classes. There is no information to list.

User response: No action required.

DQP2118I **Submitter profile for user *username* was added successfully.**

Explanation: A new user submitter profile was created for the user.

User response: No action required.

DQP2119I **Submitter profile for group *groupname* was added successfully.**

Explanation: A new group submitter profile was created for the group.

User response: No action required.

DQP2120I **Submitter profile for user *username* was updated successfully.**

Explanation: An existing user submitter profile was updated.

User response: No action required.

DQP2121I **Submitter profile for group *groupname* was updated successfully.**

Explanation: An existing group submitter profile was updated.

User response: No action required.

DQP2122I **Submitter profile for user *username* was removed successfully.**

Explanation: A user submitter profile was removed.

User response: No action required.

DQP2123I **Submitter profile for group *groupname* was removed successfully.**

Explanation: A group submitter profile was removed.

User response: No action required.

DQP2124E **Submitter profile for user *username* does not exist.**

Explanation: The user submitter profile does not exist.

User response: Specify an existing user name and reissue the command.

DQP2125E **Submitter profile for group *groupname* does not exist.**

Explanation: The group submitter profile does not exist.

User response: Specify an existing group name and reissue the command.

DQP2126E **Submitter profile for user *username* already exists.**

Explanation: The user submitter profile exists.

User response: Specify a unique user name and reissue the command.

DQP2127E **Submitter profile for group *groupname* already exists.**

Explanation: The group submitter profile exists.

User response: Specify a unique group name and reissue the command.

DQP2128I **No submitter profiles exist.**

Explanation: There are no submitter profiles. There is no information to list.

User response: No action required.

DQP2129I **The settings for the Query Patroller system were updated successfully.**

Explanation: The settings for the Query Patroller system were updated successfully.

User response: No action required.

DQP2130E **Query *queryID* does not exist.**

Explanation: The query ID specified does not exist.

User response: Specify an existing query ID and reissue the command.

DQP2131I **No queries matching the input criteria exist.**

Explanation: There is no information to list as there are no queries that match the input criteria.

User response: Modify the input criteria, if necessary.

DQP2132I **New submission preferences for user *username* were added.**

Explanation: New user submission preferences were created for the user.

User response: No action required.

DQP2133I **Submission preferences for user *username* were updated successfully.**

Explanation: Existing submission preferences were updated.

User response: No action required.

DQP2134I **Submission preferences for user *username* were removed. This user will use the default submission preferences.**

Explanation: An individual's submission preferences were removed.

User response: No action required.

DQP2135E **Submission preferences for user *username* already exist.**

Explanation: The user submission preferences already exist.

User response: Specify a unique user name and reissue the command.

DQP2136I **All result tables were removed successfully.**

Explanation: All of the result tables were dropped.

User response: No action required.

DQP2137I **There are no existing result tables to drop.**

Explanation: There are no result tables to drop.

User response: No action required.

DQP2138I **The result tables belonging to user *username* were removed successfully.**

Explanation: The result tables were dropped.

User response: No action required.

DQP2139I **There are no result tables belonging to user *username*. Nothing was removed.**

Explanation: There are no result tables to drop.

User response: No action required.

DQP2140I **Result tables for query *queryID* were removed successfully.**

Explanation: The result tables were dropped.

User response: No action required.

DQP2141I **Result table for query *queryID* does not exist. Nothing was removed.**

Explanation: The result table for the specified query does not exist.

User response: No action required.

DQP2142E **Query *queryID* is not in the correct state.**

Explanation: The query is not in the correct state to execute the command.

User response: Verify the state of the query and reissue the command.

DQP2143E **Invalid time unit *timeunit* specified for parameter *parametername*.**

Explanation: The time unit specified for the parameter is invalid.

User response: See the Query Patroller documentation for a detailed description of the command.

Specify the correct time unit and reissue the command.

DQP2144E **The parameter *parametername* specified is not correct.**

Explanation: The parameter specified is not correct.

User response: Check the Query Patroller documentation for a detailed description of the command. Correct the syntax and reissue the command.

DQP2145E **The parameter *parametername* must be specified.**

Explanation: The command syntax specified is not correct.

User response: Check the Query Patroller documentation for a detailed description of the command. Correct the syntax and reissue the command.

DQP2146E **An incorrect value *value* was specified for the parameter *parametername*.**

Explanation: An incorrect value was specified for the parameter.

User response: Check the Query Patroller documentation for a detailed description of the

command. Correct the parameter value and reissue the command.

DQP2147E **Missing value for parameter**
parametername.

Explanation: A required parameter value is missing from the command.

User response: Check the Query Patroller documentation for a detailed description of the command. Include the necessary parameter value and reissue the command.

DQP2148E **Unexpected keyword** *keyword*, **expected keyword might be** *keyword*.

Explanation: A keyword that does not apply to the command was encountered.

User response: Check the Query Patroller documentation for a detailed description of the command. Correct the syntax and reissue the command.

DQP2149E **Unexpected end of command, expected value might include** *value*.

Explanation: A syntax error in the command was detected.

The command cannot be processed.

User response: See the Query Patroller documentation for a detailed description of the command.

Correct the syntax and reissue the command.

DQP2150E **The parameter** *parametername* **cannot be specified more than once.**

Explanation: The parameter *parametername* is specified more than once.

User response: Check the Query Patroller documentation for a detailed description of the command. Correct the syntax and reissue the command.

DQP2151E **There is no default value for parameter**
parametername.

Explanation: The parameter does not have a default value.

User response: Check the Query Patroller documentation for a detailed description of the command. Specify a valid parameter value and reissue the command.

DQP2152E **The value of parameter** *parametername* **cannot be null.**

Explanation: A value must be specified for parameter *parametername*.

User response: Check the Query Patroller documentation for a detailed description of the command. Provide a value for the parameter and reissue the command.

DQP2153E **The value of parameter** *parametername* **must be surrounded by single quotation marks.**

Explanation: The parameter must be surrounded by single quotation marks.

User response: Check the Query Patroller documentation for a description of how to run a command. Surround the parameter value with single quotation marks and reissue the command.

DQP2154E **The query is no longer running and the query status is unknown.**

Explanation: The query is not running, but its final status cannot be determined. It may have completed successfully or failed. This is an abnormal situation which may have been caused by the following scenarios:

- The Query Patroller server crashed, or was shut down with the FORCE option, terminated because of a power outage, or its connection to DB2 was forced off while this query was running.
- The DB2 server crashed, was shut down with the force option, or terminated because of a power outage, while this query was queued or running.
- The Query Patroller server did not respond within an expected time frame when DB2 attempted to report query completion status.

User response: Examine the client application that submitted the query for the query status and results. Resubmit the query if necessary.

DQP2155E **User** *username* **is suspended.**

Explanation: The user is not authorized to submit queries, nor perform any Query Patroller commands.

User response: Request that the database administrator reactivate the user submitter profile.

DQP2156E **Communication with the Query Patroller server has timed out.**

Explanation: The DB2 server cannot connect to or communicate with the Query Patroller server within the expected time frame.

User response: The Query Patroller server may be

temporarily overloaded, leading to a slower than normal response time. Check the system workload and ensure that the Query Patroller system is tuned properly for optimal performance. Resubmit the query if necessary.

If the problem persists, contact IBM Service.

DQP2157E Submission preferences for user *username* do not exist.

Explanation: There are no existing submission preferences for the user.

User response: Create submission preferences for the user or specify a user name that has existing submission preferences.

DQP2158E The submission preferences belonging to user *username* cannot be copied onto itself.

Explanation: The submission preferences cannot be copied onto itself.

User response: Specify a different username and reissue the command.

DQP2159I The query information for query *queryID* has been removed.

Explanation: There is no existing information related to the *queryID* query.

User response: No action required.

DQP2160I The historical information for query *queryID* has been removed.

Explanation: There is no existing historical information related to the *queryID* query.

User response: No action required.

DQP2161I All query information has been removed.

Explanation: The query information has been removed.

User response: No action required.

DQP2162I All historical query information has been removed.

Explanation: The historical query information has been removed.

User response: No action required.

DQP2163I Query *queryID* has been canceled.

Explanation: The query has been canceled.

User response: No action required.

DQP2164I The query information for *number* queries older than *timeunit* has been removed.

Explanation: The information related to queries older than *timeunit* was removed.

User response: No action required.

DQP2165I The historical query information for *number* queries older than *timeunit* has been removed.

Explanation: The historical information related to queries older than *timeunit* was removed.

User response: No action required.

DQP2166E The query *queryID* cannot be canceled. The query has already been completed, canceled, or aborted.

Explanation: The query *queryID* cannot be canceled because the query has already been completed, canceled, or aborted.

User response: No action required.

DQP2167E Unexpected keyword *keyword*, end of command is expected.

Explanation: A keyword that does not apply to the command was encountered.

User response: Check the Query Patroller documentation for a detailed description of the command. Correct the syntax and reissue the command.

DQP2168I The query information for all the specified queries has been removed.

Explanation: The query information has been removed.

User response: No action required.

DQP2169I The historical information for all the specified queries has been removed.

Explanation: The historical information for the queries has been removed.

User response: No action required.

DQP2170I *Number result tables older than timeunit have been dropped.*

Explanation: The result tables have been removed because they were older than *timeunit*.

User response: No action required.

DQP2171I **Query Patroller has processed *number* queries and historical data was generated for *number* queries successfully**

Explanation: Historical data was generated for the number of queries identified.

User response: If the number of queries processed and the number of queries that were generated for historical data is the same, no action is required.

If this is not the case, see the qpdiag.log for possible causes of this message. If the problem persists, contact IBM Support.

DQP2172I **The query class update may take a long time. Currently queued and newly submitted queries will remain queued until the process is completed.**

Explanation: The query class update may take a long time to complete.

User response: No action required.

DQP2173E **A query class update is in progress. Another update cannot be performed at this time.**

Explanation: A query class update cannot be performed while another query class update is in progress.

User response: Reissue the command after the current update is complete.

DQP2174I **The maximum number of query classes has been exceeded.**

Explanation: The maximum number of query classes has been reached. New query classes cannot be created.

User response: Decrease the number of existing query classes.

DQP2175I **A query class with the same maximum query cost value already exists or the maximum number of query classes has been exceeded.**

Explanation: The maximum number of query classes has been reached. New query classes cannot be created.

Query classes must have a unique query cost value.

User response: Decrease the number of existing query classes and ensure that all query classes have a unique query cost value.

DQP2176E **The length of the keyword *keyword* exceeds the maximum length allowed.**

Explanation: The length of the specified keyword has exceeded the maximum length allowed.

User response: Ensure that the keyword is valid and reissue the command.

DQP2177E **The public submitter profile cannot be specified.**

Explanation: By default, the user's own submitter profile would be used automatically, if it exists. If it does not exist, the user can specify a group submitter profile that the user belongs to.

User response: Reissue the command and specify a group submitter profile, or do not specify a submitter profile.

DQP2178E **An invalid user name or password was entered.**

Explanation: An invalid user name or password was entered.

User response: Reissue the command with the correct user name and password.

DQP2179I **Historical query information does not exist.**

Explanation: Unable to remove historical information for some queries as the information does not exist.

User response: No action required.

DQP2180W **Query *queryid* is now running in the background.**

Explanation: The specified query is now running in the background.

User response: No action required.

DQP2181E **The query was aborted because the number of rows in the result set is larger than the number of rows that can be accommodated in the result table.**

Explanation: Query Patroller attempted to create a result table for this query because it was either run in the background or held and then released from held state. The submission preferences of this submitter stated that in the event the result set is longer than the maximum allowed, no result should be returned.

Query Patroller aborted the query when it detected that

the number of rows returned by the query would exceed the maximum size of a result table specified in the submitter profile under which this query was submitted.

User response: Modify the query to reduce the number of rows it returns or request that an administrator or operator increase the number of rows allowed for a result table. If the query was run in the background, try waiting for the result instead of releasing the application.

DQP2182I Held query *queryid* has been run.

Explanation: The specified query is running in the background.

User response: No action required.

DQP2183E Query *queryid* cannot be run in the background.

Explanation: The specified query cannot be run in the background.

User response: See the Query Patroller documentation for the class of queries that can be run in the background.

DQP2184E Result tables older than *time-unit* do not exist.

Explanation: The result tables do not exist. Nothing was removed.

User response: No action required.

DQP2185I Query information older than *time-value* does not exist.

Explanation: The query information cannot be removed as it does not exist.

User response: No action required.

DQP2186I Historical query information older than *time-value* does not exist.

Explanation: The historical query information cannot be removed as it does not exist.

User response: No action required.

DQP2187W Submission preferences for user *username* were added successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The new submission preferences for the user were added successfully. However, the Query

Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2188W Submission preferences for user *username* were updated successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submission preferences for the user were updated successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2189W Submitter profile for user *username* was added successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submitter profile for the user was added successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2190W Submitter profile for user *username* was updated successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submitter profile for the user was updated successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2191W Submitter profile for group *groupname* was added successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submitter profile for the group was added successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2192W Submitter profile for group *groupname* was updated successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The new submitter profile for the group was updated successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2193W The settings for the Query Patroller system were updated successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The settings for the Query Patroller system were updated successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2194W Submission preferences for user *username* were removed successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submission preferences for the user were removed successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2195W Submitter profile for user *username* was removed successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submitter profile for the user was removed successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2196W Submitter profile for group *groupname* was removed successfully. However, communication with the Query Patroller server has failed. The changes have not been reflected on the Query Patroller server.

Explanation: The submitter profile for the group was removed successfully. However, the Query Patroller server cannot be contacted to reflect the changes.

User response: Restart the Query Patroller server to see the reflected changes.

DQP2197I Query *queryID* cannot be removed.

Explanation: The query information cannot be removed.

User response: No action required.

DQP2198E Unable to display or file result sets for query *queryID*. The result sets do not exist.

Explanation: The result sets for the query do not exist. This error might occur when the query is a non-SELECT statement, the result sets were manually dropped, or when a client application is specified as the result destination.

User response: No action required.

DQP2199W The purge period specified for managed queries is less than the purge period specified for the result tables. The associated result table will be deleted when the managed query is deleted.

Explanation: The purge period specified for managed queries is less than the purge period specified for the result tables. When a managed query is deleted, the associated result table is also deleted; therefore the purge period for the result tables is ignored.

User response: Set the purge period for the managed queries to be greater than or equal to the purge period for the result tables.

DQP2200W The purge period specified for historical queries is less than the purge period specified for the managed queries. The associated managed query will be deleted when the historical query is deleted.

Explanation: When a historical query is deleted the associated managed query is also deleted, therefore, the purge period for the managed queries is ignored.

User response: Set the purge period for the historical queries to be greater than or equal to the purge period for the managed queries.

DQP2202I The historical analysis data generator was stopped by a user. Query Patroller processed *number* queries and historical data was generated for *number* queries successfully before it was stopped.

Explanation: The GENERATE HISTORICAL_DATA STOP command was issued in a separate process. Historical data was generated for the number of queries identified.

User response: No action required.

DQP2203I An attempt to stop the generation of historical data is in progress.

User response: No action required.

DQP2204I Query Patroller is not generating historical data at this time.

Explanation: Query Patroller is not able to issue the GENERATE HISTORICAL_DATA STOP command as historical data is not generating data at this time.

User response: No action required.

DQP2205E Only the query submitter can run a query in the background.

Explanation: To run a query in the background, you must be the submitter who submitted the query originally.

User response: See the Query Patroller documentation for a detailed description of running a query in the background.

DQP2206E Value for parameter *parameter1* must be set as the parameter *parameter2* was set to *parameter2-value*.

Explanation: The value for *parameter1* must be set.

User response: Specify a value for parameter *parameter1* and reissue the command.

DQP2207E The command completed successfully.

Explanation: No errors were encountered during the execution of this command.

User response: No action required.

DQP2208E The length of the specified command exceeds the maximum length allowed.

Explanation: The command length is too long.

User response: See the Query Patroller documentation for a detailed description of the command.

Reissue the command.

DQP2209W Releasing this query will cause the system to run a workload that exceeds the maximum system workload cost.

Explanation: The query that is being released has a cost that exceeds the system workload cost. Once it has been release from held state, Query Patroller will queue it until there are no other queries queued.

This does not guarantee that there will be no other queries running at the same time as this query.

User response: No action required.

DQP2210E The queue query no longer exists; the query status is now aborted.

Explanation: The query was queued and had not started running, but now it no longer exists. This is an abnormal situation which may have been caused by the one of the following scenarios:

- The DB2 server crashed, was shut down with the force option, or terminated because of a power outage while the query was queued.
- The Query Patroller sever did not respond within an expected time frame when DB2 attempted to report query status.

User response: Examine the client application that submitted the query for the query status and results. Resubmit the query if necessary.

DQP2211E Unable to run Historical Analysis Data Generator as it is currently running.

Explanation: Only one Historical Analysis Data Generator can run on a database at one time.

User response: Wait until the currently running Historical Analysis Data Generator completes before trying to run it again. Or stop the currently running Historical Analysis Data Generator by issuing the GENERATE HISTORICAL_DATA STOP Query Patroller command.

DQP2212W The result of the specified query was saved successfully to the file *filename*, however, the utility encountered warnings. Refer to the message file *message-file* for more details.

Explanation: DB2 Query Patroller uses the DB2 export utility to save or file the query result. The result was exported successfully to the file, however DB2 export generated a message file *message-file* containing warnings.

User response: Refer to the message file to view the warnings. For more information about the DB2 export utility, refer to the DB2 Information Center.

DQP2213W One or more result tables could not be removed.

Explanation: One or more result tables could not be removed. This error may occur when a user has insufficient authority.

User response: Refer to the qpuser.log file to find out which result tables could not be removed and to determine possible causes of this message.

DQP2214E User *username* does not have the privilege to remove the result table for query *queryID*.

Explanation: The result table cannot be removed because the user does not have sufficient authority.

User response: Refer to the Query Patroller documentation for the authorization requirements for dropping a result table.

DQP2215E Query Patroller cannot release query *queryid* from held state because DYN_QUERY_MGMT is disabled.

Explanation: To release a query from held state, the database configuration parameter DYN_QUERY_MGMT needs to be enabled. It is currently disabled.

User response: Have an administrator update the database configuration parameter DYN_QUERY_MGMT to enabled state and retry releasing the query from held state.

DQP2216E Query Patroller cannot run query *queryid* in the background because DYN_QUERY_MGMT is disabled.

Explanation: To run a query in the background, the database configuration parameter DYN_QUERY_MGMT needs to be enabled. It is currently disabled.

User response: Have an administrator update the database configuration parameter DYN_QUERY_MGMT to enabled state and retry running the query in the background.

DQP2217E Query Patroller was unable to generate historical data due to an error. *SQLCODE = sqlcode*.

Explanation: An error occurred while trying to generate historical data. No historical data was generated.

User response: See the qpdiag.log file for possible causes of this message. If the problem persists, contact IBM Support.

DQP2218I The submission preferences for user *user-id* were added successfully. If the user does not belong to the group whose submitter profile has been selected, Query Patroller will select another submitter profile automatically when processing queries submitted by the user.

Explanation: The query submission preferences for the specified user was created successfully, and a group submitter profile was selected. If the specified user does not belong to the group or no longer belongs to the group, at the time he/she submits a query, Query Patroller will select the most restrictive submitter profile from the groups to which the user belongs.

User response: Ensure that the user belongs to the group whose submitter profile has been specified at the time a query is submitted. Otherwise, let Query Patroller select the most restrictive submitter profile.

Chapter 99. DQP2500 - DQP2999

DQP2505I The database partition group *dbpartitiongroupname* does not exist. The **qpsetup** command will attempt to create a new one.

Explanation: The specified database partition group does not exist.

User response: The **qpsetup** command will attempt to create a new database partition group.

DQP2506E One or more Query Patroller control tables already exist.

Explanation: One or more Query Patroller control tables were found. The **qpsetup** command cannot proceed.

User response: Verify if the control tables are valid and complete. If you wish to replace the control tables with new control tables, reissue the **qpsetup** command using the **REPLACE** option.

DQP2507E The table space *tablespacename* does not exist. The **qpsetup** command cannot proceed.

Explanation: The specified table space name does not exist.

User response: Ensure that the table space name is correct and reissue the command.

DQP2508E The schema *schemaname* already exists. The command *commandname* cannot proceed.

Explanation: The specified schema name already exists.

User response: Specify a unique schema name and reissue the command.

DQP2516E Failed to bind package *package-name*, reason *reason*.

Explanation: An attempt to bind a package failed.

User response: If the failure occurred due to an SQLCODE, see the Message Reference documentation for more information. If the problem persists, contact IBM Support.

DQP2518E The Query Patroller control tables do not exist.

Explanation: The Query Patroller control tables cannot be found in the database specified.

User response: Issue the **qpsetup** command to create the Query Patroller control tables.

DQP2519E Failed to update the **DB2QP.QP_SYSTEM** table with the result table space information. Reason = *reason*.

Explanation: The **qpsetup** command attempted to update the **QP_SYSTEM** table with the result table space, however the update failed.

User response: Analyze the reason, correct the problem, and reissue the command.

If the problem persists, contact IBM Support.

DQP2520W The schema *schema-name* already exists. The **qpsetup** command will attempt to create Query Patroller control tables using this schema.

Explanation: The schema already exists in the specified database. The **qpsetup** command will attempt to create Query Patroller control tables in this schema. DB2 tables existing in this schema will not be replaced unless the **REPLACE** option is specified in the **qpsetup** command.

User response: If any existing tables have the same name as the Query Patroller control tables, move them to a different schema.

DQP2521E The table space container *container-path* is already in use.

Explanation: The specified table space container path is already in use by another application.

User response: Verify the path and reissue the command.

DQP2522E The **qpsetup** command failed. All database objects created by this command were cleaned up.

Explanation: The **qpsetup** command failed. All database objects, such as table spaces, tables, functions, and procedures, were cleaned up.

User response: Refer to the **qpsetup.log** file for possible causes of this message. If the problem persists, contact IBM Support.

DQP2523I The **qpsetup** command completed successfully.

DQP2526I • DQP2613W

Explanation: The qpsetup command completed successfully.

User response: No action required.

DQP2526I **Package *package-name* was bound successfully.**

Explanation: Query Patroller bound the package successfully.

User response: No action required.

DQP2604I **The cleanup of the Query Patroller Version 7 control tables has completed successfully.**

Explanation: The Query Patroller migration cleanup completed successfully and all Version 7 Query Patroller tables, view, and triggers have been dropped.

User response: No action required.

DQP2605E **The cleanup of the Query Patroller Version 7 control tables has failed.**

Explanation: The Query Patroller migration tool encountered fatal error while cleaning the Version 7 Query Patroller database objects. The cleanup did not complete successfully.

User response: Refer to the qpmigrate.log file for possible causes of this message.

DQP2606W **The cleanup of the Query Patroller Version 7 control tables has completed with warnings.**

Explanation: The Query Patroller migration cleanup encountered warnings while deleting the Version 7 Query Patroller tables. However, it may have completed successfully. There may have been errors migrating users and/or parts of the system configuration.

User response: Refer to the qpmigrate.log file for possible causes of this message.

DQP2607I **Migration of Query Patroller control tables from Version 7 to Version 8 has completed successfully.**

Explanation: The Query Patroller migration tool completed successfully.

User response: No action required.

DQP2608E **Migration of Query Patroller control tables from Version 7 to Version 8 has failed.**

Explanation: The Query Patroller Migration tool encountered fatal errors while migrating the database.

The migration did not complete successfully.

User response: Refer to the qpmigrate.log file for possible causes of this message.

DQP2609W **Migration of Query Patroller control tables from Version 7 to Version 8 has completed with warnings.**

Explanation: The Query Patroller Migration tool encountered warnings while migrating the database, however, the migration may have completed successfully. There may have been errors migrating users and/or parts of the system configuration.

User response: Refer to the qpmigrate.log file for possible causes of this message.

DQP2610E **Stop Query Patroller server before proceeding.**

Explanation: The Query Patroller server must be stopped before the Query Patroller migration tool can be run.

User response: Issue the qpstop command to stop the Query Patroller server.

DQP2611W **User profile for user *username* has not been migrated. SQLCODE = *SQLCODE*.**

Explanation: The user has not been migrated.

User response: Refer to the SQLCODE for possible causes of this message.

DQP2612I **Migration of Query Patroller user and group profiles has completed successfully.**

Explanation: The data from the Version 7 Query Patroller table IWM003_USER_PROF has been migrated to the Version 8 Query Patroller tables SUBMITTER_PROFILE, OPERATOR_PROFILE, and SUBMISSION_PREFERENCES.

User response: No action required.

DQP2613W **Migration of Query Patroller user and group profiles has completed with warnings.**

Explanation: The Query Patroller migration tool encountered warnings while migrating from Version 7 Query Patroller table IWM003_USER_PROF to the Version 8 Query Patroller tables SUBMITTER_PROFILE, OPERATOR_PROFILE, and SUBMISSION_PREFERENCES.

User response: Refer to the qpmigrate.log for a list of users that were not migrated. If the user name conflict was expected, then no action is required. Otherwise, you may copy the data manually into the table.

DQP2614W Migration of Query Patroller user and group profiles has failed.

Explanation: There was an error reading data from the Version 7 Query Patroller tables and the migration of user and group profiles information could not continue. The Version 7 Query Patroller tables may be corrupt or there was an I/O error while reading from the Version 7 Query Patroller tables.

User response: Ensure that all Version 7 Query Patroller tables are present with valid data.

DQP2615I Migration of the Query Patroller system configuration has completed successfully.

Explanation: The data from Version 7 Query Patroller table IWM003_JOB_QUEUE and IWM003_SYS_PARMS have been migrated to the Version 8 Query Patroller tables QUERY_CLASS and QP_SYSTEM successfully.

User response: No action required.

DQP2616W Migration of the Query Patroller system configuration has completed with warnings.

Explanation: The Query Patroller migration tool encountered warnings while migrating from Version 7 Query Patroller tables IWM003_JOB_QUEUE and IWM003_SYS_PARMS to the Version 8 Query Patroller tables QUERY_CLASS and QP_SYSTEM. It is possible that various job queues already exist in the table.

User response: Refer to the qpmigrate.log file for a list of job queues that were not migrated. If the job queue conflict was expected, then no action is needed. Otherwise, you may copy the data manually into the table.

DQP2617E Migration of the Query Patroller system configuration has failed.

Explanation: There was an error reading data from the Version 7 Query Patroller tables and the migration of system configuration information could not continue. It is possible that the Version 7 Query Patroller tables are corrupt or there was an I/O error while reading from the Version 7 Query Patroller tables.

User response: Ensure that all the Version 7 Query Patroller tables are present with valid data.

DQP2625W Version 7 Query Patroller job queue ID *queryid* failed to migrate because an entry with an identical maximum query cost of *querycost* already exists.

Explanation: Query Patroller migration found an entry in the QUERY_CLASS control table with a maximum cost that is identical to the query class that is attempting to be migrated. Entries with identical

maximum costs are not permitted.

User response: If the maximum cost conflict was expected, then no action is required. Otherwise, change the maximum cost of the query class to migrate and manually insert it into the QUERY_CLASS control table.

DQP2627W Version 7 Query Patroller job queue ID *queueID* failed to migrate because it has a maximum query cost of 0.

Explanation: The Version 7 job queue has a maximum query cost of 0. No queries will belong to a query class with a maximum query cost of 0, therefore the job queue is not migrated.

User response: No action required.

Chapter 100. DQP3000 - DQP3499

DQP3000E **The maximum number of queries for this query class must not be greater than that of the Query Patroller System.**

Explanation: The maximum number of queries may be set for a query class and for the Query Patroller system itself. No query class can exist with a maximum of queries larger than that of the Query Patroller system.

User response: Lower the maximum number of queries for the query class or increase the maximum number of queries for the Query Patroller system before increasing the maximum number of queries for the query class.

DQP3001E **The maximum query cost for this query class must not be greater than the maximum workload cost of the Query Patroller System.**

Explanation: No query class can exist with a maximum query cost larger than the maximum workload cost of the Query Patroller system.

User response: Lower the maximum query cost for the query class or increase the maximum workload cost for the Query Patroller system before increasing the maximum query cost for the query class.

DQP3002E **One or more query classes exist with the maximum number of queries greater than the value specified for the Query Patroller System.**

Explanation: The maximum number of queries may be set for a query class and for the Query Patroller system itself. No query class can exist with a maximum number of queries larger than that of the Query Patroller system.

User response: Increase the maximum number of queries for the Query Patroller system or decrease the maximum number of queries for each query class that is preventing the system maximum from being lowered.

DQP3003E **One or more query classes exist with a maximum query cost greater than the value specified for the maximum workload cost of the Query Patroller System.**

Explanation: No query class can exist with a maximum query cost larger than the maximum workload cost of the Query Patroller system.

User response: Increase the maximum workload cost for the Query Patroller system or decrease the maximum query cost for each query class that is preventing the system maximum from being lowered.

DQP3010E **The public submitter profile cannot be removed.**

Explanation: The public submitter profile cannot be removed.

User response: No action required.

DQP3011E **A user submitter profile name cannot be specified.**

Explanation: The command does not accept a user submitter profile name to be specified. By default, the user's own submitter profile would be used automatically, if it exists. If it does not exist, the user can specify a group submitter profile that the user belongs to.

User response: Reissue the command and specify a group submitter profile, or do not specify a submitter profile.

DQP3012E **The public submission preferences cannot be removed.**

Explanation: The public submission preferences cannot be removed.

User response: No action required.

Part 15. EXP Messages

This section contains the Explain (EXP) messages. The messages are listed in numeric sequence.

Chapter 101. EXP0000 - EXP0499

EXP0001W An unexpected error occurred processing or applying an optimization profile or embedded optimization guideline. Line number *line-number*, character number *character-number*.

Explanation: An error has occurred which is due to a system error. The processing of the SQL statement was successful but the optimization profile could not be applied.

User response: If trace was active, invoke the Independent Trace Facility at the operating system command prompt. Then contact your technical service representative with the following information:

- Problem description
- Contents of the EXPLAIN_DIAGNOSTICS and EXPLAIN_DIAGNOSTICS_DATA tables
- The db2diag.log
- Trace file if possible

EXP0002W There was insufficient memory to process the optimization profile or embedded optimization guideline.

Explanation: There was insufficient memory available to allow complete processing of the optimization profile.

User response: Increase the size of application heap memory and recompile the statement.

EXP0003W The OPT_PROFILE table in the SYSTOOLS schema is either missing or has an unexpected format.

Explanation: The optimizer attempted to read an optimization profile from the OPT_PROFILE table in the SYSTOOLS schema, but was unsuccessful. Either the optimizer could not find the table, or the table had an unexpected format and could not be read.

User response: Create the table with the proper format and recompile the statement.

EXP0004W The optimization profile or embedded optimization guideline is either not well-formed or is invalid. Line number *line-number*, character number *character-number*.

Explanation: An optimization profile could not be processed by the optimizer because it had an invalid format.

User response: Ensure that the XML is well-formed

and valid with respect to the current optimization profile schema and recompile the statement.

EXP0005W The optimization profile could not be processed or found in the OPT_PROFILE table in the SYSTOOLS schema.

Explanation: The optimization profile in effect for the statement could not be found in the OPT_PROFILE table in the SYSTOOLS schema.

User response: Ensure that the optimization profile identified by the OPTPROFILE bind option or the CURRENT OPTIMIZATION PROFILE special register is contained in the table and recompile the statement.

EXP0006W The optimization profile cache is out of memory, performance may be impacted.

Explanation: The specified optimization profile could not be inserted into the optimization profile cache. This means that it will need to be reread from the SYSTOOLS.OPT_PROFILE table for subsequent uses. This could lead to a degradation in SQL compile time performance.

User response: Use the catalogcache_sz configuration parameter to increase the size of the profile cache.

EXP0007W The VALUE attribute of the REOPT element has the invalid value *REOPT-value*. Line number *line-number*, character number *character-number*.

Explanation: The value *REOPT-value* specified for the VALUE attribute of the REOPT element at the line number and character number is invalid.

User response: Replace the invalid REOPT value with a valid one and recompile the statement. See the optimization guidelines reference guide for a list of valid optimization guideline elements.

EXP0008W Invalid access request. The table reference identified by the TABID attribute could not be found. Line number *line-number*, character number *character-number*.

Explanation: The correlation name specified by the TABID attribute could not be mapped to one of the correlation names in the optimized version of the statement.

User response: Replace the invalid correlation name with a valid one and recompile the statement.

EXP0009W Invalid access request. The table reference identified by the TABLE attribute could not be found. Line number *line-number*, character number *character-number*.

Explanation: The exposed or extended name specified by the TABLE attribute could not be mapped to any of the exposed or extended names given to tables referenced in the statement.

User response: Replace the invalid exposed or extended name with a valid one and recompile the statement.

EXP0010W Invalid access request. The table reference identified by the TABLE attribute is not unique. Line number *line-number*, character number *character-number*.

Explanation: The exposed or extended name specified by the TABLE attribute is not permitted to map to more than one table reference.

User response: Do one of the following and then recompile the statement:

- Modify the statement to make the exposed or extended name unique.
- Use the unique correlation name associated with the table reference in the optimized statement.

EXP0011W Invalid access request. The TABLE and TABID fields must identify the same table reference. Line number *line-number*, character number *character-number*.

Explanation: An access request can identify a table reference using either the TABID or TABLE attribute. If both are provided, they must identify the same table reference.

User response: Drop the attribute that identifies the wrong table reference or ensure that both attributes refer to the same table reference and recompile the statement.

EXP0012W Invalid access request. The index *index-name* could not be found. Line number *line-number*, character number *character-number*.

Explanation: An index scan, list prefetch, or index ANDing access request specified an index that is not defined for the table.

User response: Replace the invalid index name with the name of an index that is defined for the table and recompile the statement.

EXP0013W Invalid index ANDing access request. Block indexes must appear before record indexes. Line number *line-number*, character number *character-number*.

Explanation: An optimization guideline specified an index ANDing access request that did not have all references to block indexes before references to record indexes. For star join index ANDing access requests, this requirement applies to fact table index accesses on semijoin inners.

User response: Replace the invalid index ANDing access request with a valid one, and recompile the statement.

EXP0014W Invalid access or join request. An element's data length exceeds DB2 maximum for that type and cannot be applied. Element *element*, line number *line-number*, character number *character-number*.

EXP0015W Invalid join request. Join refers to tables that are not in the same FROM clause. Line number *line-number*, character number *character-number*.

Explanation: A join request contains references to tables that are not in the same FROM clause.

User response: Ensure that all access requests within the join request refer to correlation or table names in the same FROM clause of the optimized statement and recompile the statement.

EXP0016W Invalid join request. Conflicting join dependencies.

Explanation: An attempt was made to join derived tables that are mutually dependent on one another due to correlated column references.

User response: Correct the join request and recompile the statement.

EXP0017W More than one comment containing embedded optimization guideline was found in the statement.

EXP0019W Invalid access request. A table with the supplied correlation and table name cannot be found. DB2 has substituted an MQT for the table.

EXP0020W Table has no statistics. The table *schema.table-name* has not had runstats run on it. This may result in a sub-optimal access plan and poor performance.

EXP0021W Table column has no statistics. The column *column-number* of table *schema.table-name* has not had runstats run on it. This can lead to poor cardinality and predicate filtering estimates.

EXP0022W Index has no statistics. The index *schema.index-name* has not had runstats run on it. This can lead to poor cardinality and predicate filtering estimates.

EXP0023W Optimization guideline is a duplicate of another guideline. Line number *line-number*, character number *character-number*.

Explanation: An equivalent optimization guideline with the same target exists.

User response: Remove the duplicate guideline.

EXP0024W Optimization guideline conflicts with another guideline. Line number *line-number*, character number *character-number*.

Explanation: More than one optimization guideline exists with the same target but with conflicting options.

User response: Remove one of the conflicting guidelines and recompile the statement.

EXP0025W Optimization guideline is ambiguous. Line number *line-number*, character number *character-number*.

Explanation: The Optimization guideline target is not unique.

User response: Refine the guideline to make the target unique and recompile the statement.

EXP0026W Invalid option *option*. Line number *line-number*, character number *character-number*.

Explanation: The option provided is invalid.

User response: Correct the option and recompile the statement.

EXP0027W A column name is specified without a table reference. Line number *line-number*, character number *character-number*.

Explanation: A column name requires a qualifying TABLE or TABID attribute.

User response: Qualify the column in the guideline

with a TABLE or TABID attribute and recompile the statement.

EXP0028W The optimization guideline SUBQ2JOIN has not been applied because semantic conditions are not met or context was changed.

Explanation: The precondition for SUBQ2JOIN is not met or the statement was changed by other rewrite rules.

User response: Ignore the warning or remove the element.

EXP0029W The optimization guideline INLIST2JOIN has not been applied because semantic conditions are not met or context was changed.

EXP0030W The optimization guideline NOTEX2AJ has not been applied because semantic conditions are not met or context is changed.

Explanation: The precondition for NOTEX2AJ is not met or the statement has been changed by other rewrite rules.

User response: Ignore the warning or remove the element.

EXP0031W The optimization guideline NOTIN2AJ has not been applied because semantic conditions are not met or context is changed.

Explanation: The precondition for NOTIN2AJ is not met or the statement has been changed by other rewrite rules.

User response: Ignore the warning or remove the element.

EXP0032W IN list predicate could not be found. Line number *line-number*.

Explanation: No IN list predicate exists that matches the table reference and column name provided. The IN list predicate may have been removed by other rewrite rules.

User response: If the TABLE, TABID or COLUMN attributes are in error, correct them and recompile the statement.

EXP0033W Invalid access request. The MQT name could not be matched. Line number *line-number*, character number *character-number*.

Explanation: No MQT exists with the specified name and schema.

User response: Ensure that the MQT exists and the name and schema specified are correct and recompile the statement.

EXP0034W Invalid access request. The table reference could not be found. Line number *line-number*, character number *character-number*.

EXP0035W Guideline not applied. Line number *line-number*, character number *character-number*.

Explanation: The specified guideline could not be applied to the statement. A database limit may have been reached or the optimization level may not allow the guideline.

User response: If the guideline is not supported at the current optimization level, change the level and recompile the statement.

EXP0036W Invalid access request. The index at line *line-number*, column *column-number* is not a correct identifier (index names should not have leading or trailing white spaces unless delimited strings).

EXP0037W Invalid setting request. The schema at line *line-number*, column *column-number* is not a correct identifier (schema names should not have leading or trailing white spaces unless delimited strings).

EXP0038W Invalid setting request. The MQT at line *line-number*, column *column-number* is not a correct identifier.

EXP0039I Query complexity measure. Highest number of joins in any query block: *joins*.

EXP0040I Column group statistics recommendation. The columns *columns* of table *table-name* are statistically correlated. This can lead to underestimating the cardinality by a factor of *error-factor* in the worst case. Collecting a column group statistic on the columns may improve cardinality estimation.

EXP0041I Column group statistics recommendation for bias prevention. The columns *columns* of table *table-name* are statistically correlated. This can lead to underestimating the cardinality by a factor of *error-factor* in the worst case. Collecting a column group statistic on the columns may improve cardinality estimation.

EXP0042I Column group statistics recommendation. The columns *columns* of table *table-name* might be statistically correlated. This can lead to underestimating the cardinality by a factor of *error-factor* in the worst case. Collecting a column group statistic on the columns may improve cardinality estimation.

EXP0043I Column group statistics recommendation for bias prevention. The columns *columns* of table *table-name* might be statistically correlated. This can lead to underestimating the cardinality by a factor of *error-factor* in the worst case. Collecting a column group statistic on the columns may improve cardinality estimation.

EXP0044W Invalid value specified in the optimization guideline for the TYPE attribute of the MQTENFORCE element at line number *line-number*, column number *column-number*.

Explanation: An optimizer profile guideline containing an MQTENFORCE element with the TYPE attribute attempted to specify an invalid value for TYPE. Valid values for this attribute are: NORMAL, REPLICATED, and ALL.

User response: Specify either NORMAL, REPLICATED, or ALL for the TYPE attribute.

EXP0045W The table named *schema.table-name* has fabricated statistics. This can lead to poor cardinality and predicate filtering estimates. The size of the table changed significantly since the last time the RUNSTATS command was run.

Explanation: Fabricating statistics means deriving or creating statistics, rather than collecting them as part of normal RUNSTATS activity. For example, the number of rows in a table can be derived from knowing the number of pages in the table, the page size, and the average row width.

When a table has undergone many updates or

otherwise changed significantly, the statistics cease to be valid, and you should rerun the RUNSTATS command.

User response: Collect new statistics using the RUNSTATS command to avoid performance degradations.

EXP0046W Invalid attribute specified in the optimization guideline for the MQTENFORCE element at line *line-number*, column *column-number*.

Explanation: An optimizer profile guideline containing an MQTENFORCE element attempted to specify an invalid attribute for MQTENFORCE. Valid attributes for this element are: NAME and TYPE.

User response: Specify either NAME or TYPE for the MQTENFORCE element.

EXP0047W The VALUE attribute of the DPFXMLMOVEMENT element has the invalid value *DPFXMLMOVEMENT-value*. Line number *line-number*, character number *character-number*.

Explanation: The value of the DPFXMLMOVEMENT element's VALUE attribute is invalid.

In a partitioned database environment, the VALUE attribute of the DPFXMLMOVEMENT element affects the optimizer's plan when moving XML documents between database partitions. The VALUE attribute can be one of the following values.

- REFERENCE specifies that a reference to an XML document is moved through the TQ operator. The XML document remains on the source partition.
- COMBINATION specifies that some XML documents are moved through the TQ operator and references to other XML documents are moved through the TQ operator.

By default, when the VALUE attribute is not specified, the optimizer makes a cost-based decision to move an XML document as REFERENCE or COMBINATION in order to maximize performance.

User response: Change the VALUE attribute of the DPFXMLMOVEMENT element to a valid value.

EXP0051W The following MQT was not eligible because no outlier predicate was found: *schema.table-name*.

EXP0052W The following MQT or statistical view was not considered for rewrite matching because it did not match with any MQTs specified in any optimization profiles: *schema.table-name*.

EXP0053W The following REFRESH DEFERRED MQT was not considered for rewrite matching because its isolation level was inferior to the query's isolation level: *schema.table-name*.

EXP0054W The following REFRESH DEFERRED MQT was not considered for rewrite matching because the CURRENT REFRESH AGE register was not set to ANY: *schema.table-name*.

EXP0055W The following REFRESH DEFERRED MQT was not considered for rewrite matching because the CURRENT MAINTAINED TABLE TYPES FOR OPTIMIZATION register or the DFT_MTTB_TYPES configuration parameter was not set to ALL or SYSTEM or USER: *schema.table-name*.

EXP0056W The following FEDERATED_TOOL MQT was not considered for rewrite matching because the CURRENT MAINTAINED TABLE TYPES FOR OPTIMIZATION register or the DFT_MTTB_TYPES configuration parameter was not set to FEDERATED_TOOL: *schema.table-name*.

EXP0057W The following MQT or statistical view was not considered for rewrite matching because it was based on recursive VIEWS or contained SQL constructs other than SELECT, GROUP BY and UNION ALL: *schema.table-name*.

EXP0058W The following MQT or statistical view was not considered for rewrite matching because of one or more of the following reasons: (1) MQT was in SET INTEGRITY PENDING state, or (2) MQT would be put to SET INTEGRITY PENDING state, or (3) MQT was modified in the same statement, or (4) MQT or statistical view was not enabled for optimization: *schema.table-name*.

EXP0059W The following MQT or statistical view was not eligible because of one or more of the following reasons: (1) the MQT had extra table joins or GROUP BY columns that were not present in the query, or (2) the query had an SQL construct such as ORDER BY, FETCH FIRST n ROWS, DISTINCT, or had a subquery that could not be matched: *schema.table-name*.

EXP0060W The following materialized query table (MQT) or statistical view was not eligible for query optimization: *schema.table-name*. The MQT cannot be used for query optimization because one or more tables, views or subqueries specified in the MQT could not be found in the query that is being explained.

Explanation: You can improve query performance using materialized query tables (MQTs). The optimizer uses MQTs to improve response time for complex queries. You can also use statistical views to improve the performance of queries. The optimizer uses the statistics from the view to compute a better estimate of the cost of different plan choices.

This message is returned if when the explain utility is run for a specified query, and there is a join predicate in an MQT for one of the tables in the query, but that join predicate is not in the query, and that join predicate is not lossless. A join predicate is lossless only when the join columns being joined are the primary key and the foreign key of a referential integrity constraint on each side of the predicate and the foreign key is not nullable.

User response:

1. Generate an explain plan of the query using the db2exfmt tool.
2. Compare the optimized statement of the explain plan with MQT definition query.
3. Determine whether the extra join predicate matches a referential integrity that has a foreign key that is nullable.
4. If the extra join predicate matches a referential integrity that has a foreign key that is nullable, change that foreign key column in the following way:
Define the foreign key to be not nullable if when business logic determines that the foreign key value is not null.

EXP0061W The following MQT or statistical view was not eligible because the MQT had more table joins than the query: *schema.table-name*.

EXP0062W The following MQT or statistical view was not eligible because one or more columns or expressions referenced in the query were not found in the MQT: *schema.table-name*.

EXP0063W The following MQT or statistical view was not eligible because of the correlated subquery in the query: *schema.table-name*.

EXP0064W The following MQT or statistical view was not eligible because a UNION construct used in the MQT could not be matched with the query: *schema.table-name*.

EXP0065W The following MQT or statistical view was not eligible because it contained a GROUP BY subquery or correlated table expression that could not be matched: *schema.table-name*.

EXP0066W The following MQT or statistical view was not eligible because an outer join or a subquery from the MQT or the query did not match: *schema.table-name*.

EXP0067W MQT matching was not done in SQL query rewrite, because MQT expansion failed due to sqlcode: *sqlcode*. See the documentation for the sqlcode for more details.

EXP0068W The following MQT or statistical view was not eligible because a subquery from the MQT or the query did not match: *schema.table-name*.

EXP0069W The following MQT or statistical view was not eligible because the query or the MQT had correlation with an SQL construct like ORDER BY, FETCH FIRST n ROWS, DISTINCT or maximum cardinality enforcement: *schema.table-name*.

| | |
|----------|---|
| EXP0070W | The following MQT or statistical view was not eligible because the query or the MQT had an SQL construct like ORDER BY, FETCH FIRST n ROWS, DISTINCT or had maximum cardinality enforcement: <i>schema.table-name</i> . |
| EXP0071W | The following MQT or statistical view was not eligible because the query or the MQT had a DISTINCT SQL construct: <i>schema.table-name</i> . |
| EXP0072W | The following MQT or statistical view was not eligible because GROUP BY terms in the query could not be matched with the MQT: <i>schema.table-name</i> . |
| EXP0073W | The following MQT or statistical view was not eligible because one or more data filtering predicates from the query could not be matched with the MQT: <i>schema.table-name</i> . |
| EXP0074W | The following MQT was not eligible because no suitable bindings were found between the query and the MQT: <i>schema.table-name</i> . |
| EXP0075W | The following MQT or statistical view was not eligible because it contained SQL constructs that could not be compensated to match with the query: <i>schema.table-name</i> . |
| EXP0077W | The following MQT or statistical view was not eligible because it could not be compiled: <i>schema.table-name</i> . |
| EXP0078W | The following MQT or statistical view was not eligible because an equivalent or better candidate was available: <i>schema.table-name</i> . |
| EXP0079W | The following MQT was not used in the final access plan, because the plan cost with this MQT was more expensive or a better candidate was available: <i>schema.table-name</i> . |
| EXP0080W | The current usage of the statement or the statement containing update, delete or insert or constructs like sampling limits MQT matching. |

EXP0081W MQT matching was not done in SQL query rewrite because the query optimization level is set to 0, 1, or 3.

EXP0082W The following MQT or statistical view was not eligible because FRIENDLY ARITHMETIC characteristics between the query and the MQT did not match: *schema.table-name*.

EXP0083W The explain operation failed because the server specified in the optimization profile is not defined for the query. Specified server: *server-name*. Line number: *line_number*. Character number: *character_number*.

Explanation: The EXPLAIN tables contain detailed information about the access plans created by the SQL and XQuery compiler to resolve SQL or XQuery statements. You can view the access plans in EXPLAIN tables in a text format using the db2exfmt command.

This message is returned when the db2exfmt command is called to display the contents of the EXPLAIN tables, but an earlier explain operation failed because of a problem with specified optimization guidelines.

The tokens *line_number* and *character_number* indicate the location in the XML input profile where the problematic SERVER element is specified.

User response:

1. Explain the statement again with an optimization profile that specifies a database server that is defined for the query.
2. Format the contents of the EXPLAIN tables using the db2exfmt command.

EXP0084W The explain operation failed because an invalid option name was specified in the optimization profile of the statement. Option name: *option_name*. Server name: *server-name*.

Explanation: The EXPLAIN tables contain detailed information about the access plans created by the SQL and XQuery compiler to resolve SQL or XQuery statements. You can view the access plans in EXPLAIN tables in a text format using the db2exfmt command.

This message is returned when the db2exfmt command is called to display the contents of the EXPLAIN tables, but an earlier explain operation failed because of a problem with specified optimization guidelines. Specifically, this message is returned when an option name is specified that satisfies any of the following conditions:

- Specifies an option name that does not exist
- Is not valid for the specified data source

- Is not valid for the type of the specified data source
- Is not valid for the type of the specified database object

User response:

1. Explain the statement again with an optimization profile that includes a server option name that is valid for the specified server.
2. Format the contents of the EXPLAIN tables using the db2exfmt command.

EXP0085W The explain operation failed because a value was specified for one of the server options that is not valid for the specified server in the optimization profile of the statement. **Option name:** *option_name*. **Option value:** *option_value*. **for Server name:** *server_name*.

Explanation: The EXPLAIN tables contain detailed information about the access plans created by the SQL and XQuery compiler to resolve SQL or XQuery statements. You can view the access plans in EXPLAIN tables in a text format using the db2exfmt command.

This message is returned when the db2exfmt command is called to display the contents of the EXPLAIN tables, but an earlier explain operation failed because of a problem with specified optimization guidelines. Specifically, this message is returned when a server options setting request specifies an option value that either lacks the proper delimiters or is invalid.

User response:

1. Explain the statement again with an optimization profile that includes server option values that are valid for the specified server.
2. Format the contents of the EXPLAIN tables using the db2exfmt command.

EXP0086W The explain operation failed because an option was specified twice in the optimization profile of the statement. **Option name:** *option_name*. **Server name:** *server_name*.

Explanation: The EXPLAIN tables contain detailed information about the access plans created by the SQL and XQuery compiler to resolve SQL or XQuery statements. You can view the access plans in EXPLAIN tables in a text format using the db2exfmt command.

This message is returned when the db2exfmt command is called to display the contents of the EXPLAIN tables, but an earlier explain operation failed because of a problem with specified optimization guidelines.

User response:

1. Drop the duplicate option.
2. Explain the statement again with an optimization profile that includes options only once.

3. Format the contents of the EXPLAIN tables using the db2exfmt command.

EXP0147W The following statistical view was used by the optimizer to estimate cardinalities: *schema.table-name*.

Explanation: This message indicates that the optimizer found the named statistical view and the optimizer has access to its statistics.

The DB2 cost-based optimizer uses an estimate of the number of rows, or cardinality, processed by an access plan operator to accurately cost that operator. The accuracy of this cardinality estimate depends largely upon the statistics that the runstats utility collects from the database. The optimizer can also estimate the cost of an operator without access to statistical views. In some cases the cardinality estimate calculated without access to statistical views is the same as the estimate calculated with statistical view information.

For more information about how the optimizer uses statistical views to estimate the cost of an operator, refer to the topic called "Statistical views" in the DB2 Information Center.

User response: No response is required.

EXP0148W The following MQT or statistical view was considered in query matching: *schema.table-name*.

EXP0149W The following MQT was used (from those considered) in query matching: *schema.table-name*.

EXP0150W The following MQT or statistical view was not used (from those considered) in query matching: *schema.table-name*.

EXP0151W Invalid degree value. The string specified as a degree request is invalid. **Line number** *number*, **character number** *number*.

Explanation: The string specified as degree is not correct and the guideline will not be used.

User response: Make sure that the string specified as the VALUE attribute in the guideline is ANY or a number between -1 and 32767.

EXP0152W Ignored degree value. The specified degree of parallelism will be ignored because the system is not enabled for intra-partition parallelism. **Line number** *number*, **character number** *number*.

Explanation: The database manager is not configured for intra-partition parallelism.

User response: If you want to use intra-partition parallelism restart the database manager with the `intra_parallel` configuration parameter set to ON.

EXP0153W Invalid query optimization value. Line number *number*, character number *number*.

Explanation: The string specified as `qryopt` value is not one of the acceptable values.

User response: Please check the Administration Guide or the SQL Reference for a list of supported optimization levels.

EXP0154W The statement was not optimized using DB2 fastpath because it references local objects.

EXP0155W The statement was not optimized using DB2 fastpath because it contains an update operation on a view.

EXP0156W The statement was not optimized using DB2 fastpath because it uses a three-part name for a column reference.

EXP0157W The statement was not optimized using DB2 fastpath because it references a user-defined data type.

EXP0158W The statement was not optimized using DB2 fastpath because it references a user-defined function.

EXP0159W The statement was not optimized using DB2 fastpath because it references a renamed nickname column.

EXP0160W The statement was not optimized using DB2 fastpath because it references a DB2 special register.

EXP0161W The statement was not optimized using DB2 fastpath because of the invalid server option values.

Explanation: Server options `COLLATING_SEQUENCE`, `VARCHAR_NO_TRAILING_BLANKS`, and `FASTPATH` must be set to 'Y' to enable fastpath processing.

EXP0162W The statement was not optimized using DB2 fastpath because it references nicknames from multiple servers.

EXP0163W The statement was not optimized using DB2 fastpath because it references a nickname with a dependent materialized query table.

EXP0164W The statement was not optimized using DB2 fastpath because it contains an error tolerance specification.

EXP0165W The statement was not optimized using DB2 fastpath because one or more predicates requires a lob parameter.

EXP0167W Invalid time value. The string specified as a time budget is invalid. Line number *line-num*, character number *char-num*.

EXP0171I An SQL statement was compiled for which runtime execution is not required.

Explanation: This message indicates that you have either:

- Set the option on the SET INTEGRITY statement to OFF, UNCHECKED, or FULL ACCESS, which changes the state of the table. In these cases, only the catalog is updated; there is no runtime section.
- Issued an EXPLAIN REFRESH TABLE or SET INTEGRITY statement, which indicates that the table does not need to be maintained.

User response: No action required. For these type of DDL operations, a dummy access plan is generated and this diagnostic message is saved in the explain table.

EXP0181W Invalid index ANDing access request. A non-star join index ANDing access request has NLJOIN elements as children. Line number *line-number*, character number *character-number*.

Explanation: An index ANDing access request with the STARJOIN attribute set to 'FALSE' has one or more NLJOIN join request elements as children. An index ANDing access request can have NLJOIN join requests as children only if the STARJOIN attribute is set to 'TRUE' or is not specified.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0182W Invalid star join index ANDing access request. A star join index ANDing access request has the TYPE attribute set to 'XMLINDEX'. Line number *line-number*, character number *character-number*.

Explanation: An index ANDing access request that has been interpreted as a star join access request cannot have a TYPE attribute set to 'XMLINDEX'.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0183W Invalid join request. This request is nested within another similar request which is not supported. Line number *line-number*, character number *character-number*.

Explanation: A star join index ANDing access request cannot be nested within another star join index ANDing access request or within a zigzag join.

A zigzag join request cannot be nested within another zigzag join request or within a star join index ANDing access request.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and join guideline rules, and recompile the statement.

EXP0184W Multiple star join index ANDing access requests have been found for tables in the same FROM clause. Line number *line-number*, character number *character-number*.

Explanation: Only one star join index ANDing access request can be specified for tables that are referenced in the same FROM clause.

User response: Ensure that there is only one star join index ANDing access request for any table in the same FROM clause of the optimized statement, and recompile the statement.

EXP0185W A star join index ANDing access request has an unexpected child element. Line number *line-number*, character number *character-number*.

Explanation: An index ANDing access request that has been interpreted as a star join access request has a child element that is not an NLJOIN join request. Only NLJOIN join request elements can be children of a star join index ANDing access request.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current

optimization profile schema and star join guideline rules, and recompile the statement.

EXP0186W The inner child of a star join semi-join is not an IXSCAN access request element that references the fact table. Line number *line-number*, character number *character-number*.

Explanation: A star join semi-join (the NLJOIN join request child of a star join index ANDing access request), must have an IXSCAN access request element as its inner child. Moreover, this IXSCAN access request element must reference the same TABLE or TABID as the star join index ANDing access request within which it is nested.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0187W The fact table is referenced in the outer arm of a star join semi-join. Line number *line-number*, character number *character-number*.

Explanation: No access request that is nested on the outer arm of a star join semi-join (the NLJOIN join request child of a star join index ANDing access request) can reference the same TABLE or TABID as the star join index ANDing access request within which it is nested.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0188W Invalid star join index ANDing access request. In semi-join *semijoin-number*, no "=" join predicate can be applied using the leading column of the specified fact table index. Line number *line-number*, character number *character-number*.

Explanation: If a fact table index is to be used on the inner arm of a star join semi-join (the NLJOIN join request child of a star join index ANDing access request), the leading column of the index must be referenced in an "=" join predicate between the fact table and one of the tables that is specified on the outer arm of the NLJOIN. No such predicate was found for the specified index.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0189W Invalid star join index ANDing access request. In semi-join *semi-join-number*, a fact table index was not found where an "=" join predicate could be applied using the leading column of the index. Line number *line-number*, character number *character-number*.

Explanation: If a fact table index is to be used on the inner arm of a star join semi-join (the NLJOIN join request child of a star join index ANDing access request), the leading column of the index must be referenced in an "=" join predicate between the fact table and one of the tables that is specified on the outer arm of the NLJOIN. A fact table index that could satisfy this requirement was not found.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0190W Invalid star join index ANDing access request. Neither the INDEX attribute nor INDEX child elements are permitted. Line number *line-number*, character number *character-number*.

Explanation: An index ANDing access request that has been interpreted as a star join access request cannot have an INDEX attribute or INDEX child elements.

User response: Ensure that the optimization guideline is well-formed and valid with respect to the current optimization profile schema and star join guideline rules, and recompile the statement.

EXP0191I Join enumeration method changed due to isolated quantifier threshold.

EXP0192W The access plan might be sub-optimal because data might have been stored in a LOB. Some rows might not be stored in a a system temporary table because it is too large for the temporary table's page size.

Explanation: The page size of the system temporary table space is 32 KB, but the data to be stored might be too large for this page size. Data that exceeds the 32 KB row size limit will be stored in a temporary LOB.

The performance of LOB data access might be slower because additional I/O might be required to fetch, insert, or update this data.

User response:

1. Verify the access plan for argument EXTROWS to identify which operator might need to store rows as LOB objects.
2. Consider reducing the number of referenced columns in the query.

EXP0201W The ALLINDEXES attribute has an invalid value *ALLINDEXES-value*. Line number *line-number*, character number *character-number*.

Explanation: The ALLINDEXES attribute indicates that the optimizer must select all applicable indexes.

The ALLINDEXES attribute has only one valid value:

- TRUE specifies that the optimizer must select all applicable indexes.

If you want the optimizer to make a cost-based decision to pick XML indexes, then do not specify the ALLINDEXES attribute.

User response: Change the ALLINDEXES attribute to a valid value.

EXP0202W The TYPE attribute has an invalid value *type-value*. Line number *line-number*, character number *character-number*.

Explanation: The TYPE attribute determines whether the request is an XML access request.

The TYPE attribute has only one valid value:

- The value 'XMLINDEX' specifies that optimizer must select an XML index access method.

User response: Change the TYPE attribute to a valid value.

EXP0203W Invalid access request. The index *index-name* is not valid for this optimization guideline. Line number *line-number*, character number *character-number*.

Explanation: The INDEX attribute of the XISCAN optimization guideline must specify an XML index.

User response: Change the INDEX attribute of the XISCAN element to specify an XML index.

EXP0204W Invalid access request. The XML index *index-name* was specified in the INDEX attribute or INDEX element, but the TYPE attribute is not set to XMLINDEX. Line number *line-number*, character number *character-number*.

Explanation: For the ACCESS or IXAND optimization guideline, an XML index can only be specified in INDEX attribute or INDEX element if the TYPE attribute is set to XMLINDEX.

User response: Do one of the following:

- Set the TYPE attribute to XMLINDEX.
- Remove the specified XML index from the INDEX attribute or INDEX element.

EXP0205W Invalid access request. The ALLINDEXES attribute was specified, but the TYPE attribute is not set to XMLINDEX. Line number *line-number*, character number *character-number*.

Explanation: For the ACCESS or IXAND optimization guideline, the ALLINDEXES attribute can only be specified if the TYPE attribute is set to XMLINDEX.

User response: Do one of the following:

- Set the TYPE attribute to XMLINDEX.
- Remove the ALLINDEXES specification.

EXP0206W The optimization guideline has conflicting values for the INDEX attribute or INDEX element and for the attribute ALLINDEXES. Line number *line-number*, character number *character-number*.

Explanation: If the ALLINDEXES attribute is set to TRUE, then the optimizer selects all applicable indexes. In this case, you cannot identify a specific index to be selected by the optimizer.

To specify only some indexes using the INDEX attribute or INDEX element, do not specify the ALLINDEXES attribute.

User response: Do one of the following:

- Remove the ALLINDEXES attribute.
- Remove the INDEX specification.

EXP0207W The optimization guideline has not been applied because the join method is not applicable to XML type. Line number *line-number*, character number *character-number*.

Explanation: A MSJOIN join request element or a HSJOIN join request element cannot be specified when the join column is of type XML.

User response: If the join column is of type XML, change the join type to specify either the NLJOIN request element or the JOIN request element, which allows the optimizer to choose the appropriate join type.

EXP0208W Invalid access request. The XML index *index-name* is not allowed in this type of access request. Line number *line-number*, character number *character-number*.

Explanation: You cannot specify an XML index in an IXSCAN or LPREFETCH optimization guideline. You can only specify a relational index in these guidelines.

User response: Change the INDEX attribute or INDEX element value to specify an available relational index.

EXP0211W Optimization guidelines cannot be used to set the SQL compiler registry variable string *registry-variable* on line number *line-number*, character number *character-number*.

Explanation: The setting for the specified SQL compiler registry variable in the optimization guideline was ignored. There is either no such registry variable, or this registry variable cannot be specified in an optimization guideline.

User response: Verify and correct the SQL compiler registry variable string. If it is a valid SQL compiler registry variable then it is not supported in an optimization guideline and should be removed.

EXP0212W The optimization guideline contains an invalid value specified for the SQL compiler registry variable *registry-variable* on line number *line-number*, character number *character-number*.

Explanation: The value specified in the optimization guideline for the SQL compiler registry variable is invalid and has not been used.

User response: Verify and correct the value for the registry variable. Refer to the DB2 Information Center to determine the valid values for the registry variable.

EXP0213W An empty REGISTRY guideline element is not permitted in the XML optimization profile on line number *line-number*, character number *character-number*.

Explanation: The REGISTRY guideline element must contain at least one OPTION element in the XML optimization profile.

User response: Specify one or more OPTION elements within the REGISTRY element, or remove the empty REGISTRY element.

EXP0214W The optimization profile contains an invalid attribute setting *invalid-attribute-setting* for the element *element* on line *line-number*.

Explanation: The setting for the specified optimization profile element was ignored. Either the attribute name or its value is not valid.

User response: Verify and correct the attribute setting for the specified optimization profile element.

EXP0221W The EXPLAIN_ACTUALS table does not exist. Section actuals are not available.

Explanation: The EXPLAIN_ACTUALS table was not created. To obtain section actuals data, the EXPLAIN_ACTUALS table must exist.

User response: Create the EXPLAIN_ACTUALS table by using the SYSINSTALLOBJECTS procedure, the EXPLAIN.DDL, or the db2exmig command.

EXP0222W The section actuals, version *section_actuals_version*, is not supported.

Explanation: If the section actuals were collected from a release later than the current release, the section actuals are not supported and not available.

User response: Perform one of the following actions:

- Upgrade your current DB2 release to the same version on which the section actuals were collected.
- Rerun the stored procedure on the DB2 release which captured the section actuals.

EXP0223W The section actuals are invalid.

Explanation: The section actuals data might be corrupted.

User response: Recapture the section actuals.

If the section actuals are again invalid after attempting to recapture them, contact IBM Software Support.

EXP0224W The section actuals were not collected.

Explanation: The section actuals were not collected most likely due to an improper or incomplete setup of the section actuals collection mechanism.

User response: Ensure that all the steps were followed to set up the collection mechanism.

Recapture the section actuals.

EXP0225W The section was unusable for collecting statistics.

Explanation: The EXPLAIN_FROM_ACTIVITY stored procedure in the current version of DB2 database cannot interpret the statistical data in the section because the section was created by a version of DB2 database that is higher than the current version.

User response: Respond to this error in one of the following ways:

- Rebind the application by performing the following steps:
 1. Recreate the section using the REBIND command on the package for the application.
 2. Rerun the application.
 3. Call the EXPLAIN_FROM_ACTIVITY stored procedure again.
- Call the EXPLAIN_FROM_ACTIVITY stored procedure again, using the same version of the database manager as the version of database manager that was used to create the section.

EXP0231W The high percentage of NULL values in foreign key column *schema.table-name . column-name* referenced in a join predicate in the statistical view *schema.view-name* may reduce the accuracy of statistics inferred from the statistical view.

Explanation: The query optimizer uses referential integrity constraints to infer statistics from a statistical view whenever possible. The high percentage of NULL values in a foreign key column referenced in a join predicate in the statistical view may have reduced the accuracy of statistics inferred from the statistical view, which could result in sub-optimal access plans. These statistics will still be used to optimize access plans, but consider the suggestions in the user response section to further improve the access plans.

User response: Consider these actions:

- If possible, replace the NULL values in the foreign key column with appropriate, non-NULL, values.
- If replacing the NULL values is not possible, then alter the referential integrity constraints involving the foreign key column to disable query optimization, and create additional statistical views to improve access plans.

Re-issue the queries once you have completed one of these options.

EXP0232W Statistics inferred from the statistical view *schema-name . stats-view* may be inaccurate.

Explanation: If all foreign key columns in a referential integrity constraint are referenced in the join predicates in a statistical view, the DB2 optimizer may use the referential integrity constraints to infer statistics from a statistical view if possible.

The statistics inferred from the statistical view might be inaccurate if:

- too many rows in the child table do not have matching rows in the parent table (when the constraint is not enforced)
- the NULL values in the foreign key columns referenced in the join predicates filter out too many rows from the join

User response: Consider the join predicate in the statistical view. If the join predicates reference all foreign key columns in a referential integrity constraint, and the referential integrity constraint has one or more of the listed issues, then address these issues in the referential integrity constraints in the following ways:

- ensure there are matching rows in the parent table for each distinct foreign key in the child table
- if possible, replace the NULL values with appropriate non-NULL values

If it is not possible to make these changes to the data, then alter the referential integrity constraints to disable query optimization and optionally create additional statistical views to improve access plans.

EXP0233I **The predicate *predicate-text* contains an expression matching the statistical view column *schema.table-name.column-name***

Explanation: The query optimizer matched an expression used in the predicate to the expression represented by the statistical view column. The query optimizer may have used the indicated statistical view column's statistics to improve the predicate's filtering estimate, if statistics were available.

User response: This message is for your information only; no action is required.

EXP0256I **Analysis of the query shows that the query might execute faster if an additional index was created to enable zigzag join. Schema name: *table-schema*. Table name: *table-name*. Column list: *column-list*.**

Explanation: Analysis of the query indicates a star query. However, a multi-column index on the fact table that would enable an optimal zigzag join was not found.

User response:

1. Create an index using the following statement:

```
CREATE INDEX index_name ON
  table-schema.table-name column-list.
```
2. Reissue the query.

Additionally, consider running the Index Advisor utility to obtain comprehensive index recommendations for your query workload.

EXP0261I **An invalid distribution map ID *identifier* was found in table `ADVISE_PARTITION` for table *table_schema.tablename*. The default distribution map ID `IBMDEFAULTGROUP` will be used instead.**

Explanation: An invalid distribution map ID (also referred to as a partition map identifier) was specified in the PMID column of the `ADVISE_PARTITION` table.

User response: Modify the incorrect value in the `ADVISE_PARTITION` table for the identified table or create a partition map with the given ID.

Part 16. GSE Messages

This section contains the IBM DB2 Spatial Extender and Geodetic Data Management Feature messages. The messages are listed in numeric sequence.

GSE0000I The operation was completed successfully.

Explanation: No errors were encountered during the execution of this operation.

User response: No action required.

msgcode: 0

sqlstate: 00000

GSE0001C An internal error occurred.

Explanation: Spatial Extender encountered an unexpected internal error.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -1

sqlstate: 38S01

GSE0002C Spatial Extender could not access its memory pool. Reason code = *reason-code*.

Explanation: Spatial Extender tried unsuccessfully to access its memory pool.

User response: Note the reason code *reason-code* and contact IBM Software Support.

msgcode: -2

sqlstate: 38S02

GSE0003N Spatial Extender could not allocate *number* bytes of memory.

Explanation: Not enough memory was available. Possible reasons are that the supply of memory was too low, or that memory was being used by other applications.

User response: Resolve the memory shortage and repeat the command.

msgcode: -3

sqlstate: 38S03

GSE0004C An internal parameter error occurred.

Explanation: Spatial Extender encountered an unexpected error in a parameter passed to an internal

function. The operation cannot be completed successfully.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -4

sqlstate: 38S04

GSE0005N Spatial Extender could not retrieve the instance path.

Explanation: Spatial Extender could not retrieve the instance path. The operation cannot be completed successfully.

User response: Verify the DB2 environment. If necessary, contact IBM Software Support.

msgcode: -5

sqlstate: 38S05

GSE0006N An internal string error occurred.

Explanation: Spatial Extender encountered an unexpected error in an internal string operation. The operation cannot be completed successfully.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -6

sqlstate: 38S06

GSE0007N The string *string* is missing either a closing quotation mark or a closing pair of quotation marks.

Explanation: This string lacks a closing delimiter and therefore is not terminated correctly.

User response: Terminate the string correctly. If it starts with a quotation mark, close it with a quotation mark. If it starts with a pair of quotation marks, close it with a pair of quotation marks.

msgcode: -7

sqlstate: 38S07

GSE0008N An invalid error code *error-code* was used to raise an error.

Explanation: There was an attempt to raise an error

identified by an invalid *error-code*.

User response: Contact IBM Software Support.

msgcode: -8

sqlstate: 38S08

GSE0009N Not enough space is available in DB2's application heap.

Explanation: The available memory for DB2's application heap was exceeded by Spatial Extender. The operation cannot be completed successfully.

User response: On receipt of this message, terminate the application. Increase the database configuration parameter (APPLHEAPSZ) to allow a larger application heap.

Refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference for more information.

msgcode: -9

sqlstate: 38S09

GSE0010N Not enough log space is available to DB2.

Explanation: All available space in DB2's transaction log is being used. The operation cannot be completed successfully.

User response: If the database is being used by other applications concurrently, retry the operation. Log space may be freed up when another application finishes a transaction.

Increase the database configuration parameters to allow more log space. For example, the database configuration parameters LOGPRIMARY, LOGSECOND, and LOGFILSIZ can be modified to increase the size of the available logs.

msgcode: -10

sqlstate: 38S0A

GSE0100N Spatial Extender could not open a file named *file-name*. Reason code = *reason-code*.

Explanation: Reasons why a file cannot be opened, preceded by their reason codes, are as follows:

- 1 Access to the file is denied.
- 3 A disk error occurred.
- 8 A directory with the name of the file already exists.
- 10 Spatial Extender tried either to create an already-existing file, or to open an existing file that could not be found.

12 The disk is full.

17 The specified path to the file does not exist.

22 A sharing violation occurred.

All other reason codes indicate an internal error.

User response: Verify the authorization for the file and the directories, then repeat the command.

If a reason code that indicates an internal error was encountered, contact IBM Software Support.

msgcode: -100

sqlstate: 38S10

GSE0101N An I/O error occurred while a file named *file-name* was being processed. Reason code = *reason-code*.

Explanation: Reasons why an I/O error can occur during file processing, preceded by their reason codes, are as follows:

- 1 Access to the file is denied.
- 3 A disk error occurred.
- 9 Spatial Extender attempted to read past the end of the file.
- 12 The disk is full.
- 22 A sharing violation occurred.

All other reason codes indicate an internal error.

User response: Verify that the file exists, that you have the appropriate access to the file, and that the file is not in use by another process.

If a reason code that indicates an internal error was encountered, contact IBM Software Support.

msgcode: -101

sqlstate: 38S11

GSE0102N Spatial Extender could not close a file named *file-name*. Reason code = *reason-code*.

Explanation: Reasons why an error can occur during an attempt to close a file, preceded by their reason codes, are as follows:

- 3 A disk error occurred.
- 12 The disk is full.

All other reason codes indicate an internal error.

User response: Verify that the file system is in fully working condition and that enough disk space is available.

If a reason code that indicates an internal error was

encountered, contact IBM Software Support.

msgcode: -102

sqlstate: 38S12

GSE0103N Spatial Extender could not delete a file named *file-name*. Reason code = *reason-code*.

Explanation: Reasons why an error can occur during an attempt to delete a file, preceded by their reason codes, are as follows:

- 1 Access to the file is denied.
- 3 A disk error occurred.
- 17 The specified path to file does not exist.
- 22 A sharing violation occurred.

All other reason codes indicate an internal error.

User response: Given for each reason code:

- 1 Verify that you have sufficient privileges on the file and all directories in the path of the file.
- 3 Verify that the disk and file system are in proper working order.
- 17 Verify that the path to the file exists.
- 22 Verify that the file is not accessed by another process.

If a reason code that indicates an internal error was encountered, contact IBM Software Support.

msgcode: -103

sqlstate: 38S13

GSE0200N An attempt to connect to the database failed. SQLERROR = *sql-error*.

Explanation: Spatial Extender was not able to connect to the database. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -200

sqlstate: 38S20

GSE0201W An attempt to disconnect from the database failed. SQLERROR = *sql-error*.

Explanation: Spatial Extender was not able to disconnect from the database. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: +201

sqlstate: 38S21

GSE0202N No connection to a database exists.

Explanation: Spatial Extender cannot connect to a database. The command cannot be executed successfully.

User response: Verify the Spatial Extender and database setup. Make sure that a connection to the database can be established.

msgcode: -202

sqlstate: 38S22

GSE0203W Spatial Extender was already connected to database *database-name*.

Explanation: Spatial Extender tried to connect to the database *database-name* but was already connected to it.

User response: Contact IBM Software Support.

msgcode: +203

sqlstate: 38S23

GSE0204N An attempt to commit a transaction failed. SQLERROR = *sql-error*.

Explanation: Spatial Extender could not commit the current transaction successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -204

sqlstate: 38S24

GSE0205W An attempt to roll back a transaction failed. SQLERROR = *sql-error*.

Explanation: Spatial Extender could not roll back the current transaction. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: +205

sqlstate: 38S25

GSE0206N A SELECT statement failed. SQLERROR = *sql-error*.

Explanation: Spatial Extender could not execute a SELECT statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -206

sqlstate: 38S26

GSE0207N A VALUES statement failed. SQLERROR = *sql-error*.

Explanation: Spatial Extender could not execute a VALUES statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -207

sqlstate: 38S27

GSE0208N A PREPARE statement failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not prepare an SQL statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -208

sqlstate: 38S28

GSE0209N An attempt to open an SQL cursor failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not open a cursor over a result set successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -209

sqlstate: 38S29

GSE0210W An attempt to close an SQL cursor failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not close a cursor over a result set successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: +210

sqlstate: 38S2A

GSE0211N A fetch from an SQL cursor failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not fetch a result from a cursor successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -211

sqlstate: 38S2B

GSE0212N An attempt to drop an object failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not drop the specified database object. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -212

sqlstate: 38S2C

GSE0213N A bind operation failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not bind the specified file to the current database. DB2 returned *sql-error*.

This error can occur if the size of DB2's application heap is too small.

User response: Verify that the database is enabled for spatial operations.

If the error occurs during the enabling, verify the installation of Spatial Extender.

If the problem persists, increase the database configuration parameter (APPLHEAPSZ) for the size of the application heap.

msgcode: -213

sqlstate: 38S2D

GSE0214N An INSERT statement failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not execute an INSERT statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -214

sqlstate: 38S2E

GSE0215N An UPDATE statement failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not execute an UPDATE statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -215

sqlstate: 38S2F

GSE0216N A DELETE statement failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not execute a DELETE statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -216

sqlstate: 38S2G

GSE0217N A LOCK TABLE statement failed.
SQLERROR = *sql-error*.

Explanation: Spatial Extender could not execute a LOCK TABLE statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -217

sqlstate: 38S2H

GSE0218N **A DECLARE GLOBAL TEMPORARY TABLE statement failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not execute a DECLARE GLOBAL TEMPORARY TABLE statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -218

sqlstate: 38S2I

GSE0219N **An EXECUTE IMMEDIATE statement failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not execute an EXECUTE IMMEDIATE statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -219

sqlstate: 38S2J

GSE0220N **The setting of a savepoint failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not set a savepoint successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -220

sqlstate: 38S2K

GSE0221N **No database name was specified.**

Explanation: Spatial Extender could not connect to a database because the name of the database was not specified.

User response: Specify a database name.

msgcode: -221

sqlstate: 38S2L

GSE0222N **An attempt to retrieve the authorization list from DB2 failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not retrieve the authorization list for the current user. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -222

sqlstate: 38S2M

GSE0223N **An attempt to quiesce a table space failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not quiesce a table space successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -223

sqlstate: 38S2N

GSE0224N **An attempt to import data into a table failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not import data into a table successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -224

sqlstate: 38S2O

GSE0225N **An attempt to retrieve the database configuration or the database manager configuration failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not retrieve the configuration for the current database or the configuration for the database manager successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -225

sqlstate: 38S2P

GSE0226N **An attempt to create a trigger failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not create a trigger successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -226

sqlstate: 38S2Q

GSE0227N **An ALTER TABLE statement failed. SQLERROR = *sql-error*.**

Explanation: Spatial Extender could not execute an ALTER TABLE statement successfully. DB2 returned *sql-error*.

User response: Refer to the description of *sql-error*.

msgcode: -227

sqlstate: 38S2R

GSE0228N An attempt to retrieve the message for error *gse-error* and `SQLCODE = sqlcode` from the Spatial Extender message catalog failed.

Explanation: Spatial Extender could not retrieve the message for error *gse-error* and `SQLCODE = sqlcode` successfully.

User response: Verify the installation of Spatial Extender. Verify also that the message catalog in the language that you want to use is installed.

msgcode: -228

sqlstate: 38S2S

GSE0229N An attempt to bind a file named *bind-file* failed. Reason code = *reason-code*.

Explanation: Spatial Extender failed in an attempt to use the `sqlabndx` function to bind the file named *bind-file* to the database. `sqlabndx` returned the reason code *reason-code*.

User response: Verify that the database is enabled for spatial operations.

Refer to the description of *reason-code* for the function `sqlabndx`.

msgcode: -229

sqlstate: 38S2T

GSE0230N The system catalog could not be updated.

Explanation: Spatial Extender encountered an error while attempting to use the DB2 service to update the system catalog.

User response: Contact IBM Software Support.

msgcode: -230

sqlstate: 38S2U

GSE0231N A PREPARE statement encountered a warning condition. `SQLWARNING = sql-warning`.

Explanation: Spatial Extender encountered a warning condition when preparing an SQL statement. DB2 returned *sql-warning*. The PREPARE statement was completed successfully.

User response: Refer to the description of *sql-warning*.

msgcode: -231

sqlstate: 38S2V

GSE0300N The specified password is too long.

Explanation: The password used in your attempt to connect to the database is too long.

User response: Verify that the password you specified is correct. If it is the correct password, then shorten it and try the operation again.

msgcode: -300

sqlstate: 38S40

GSE0301N The specified schema name, *schema-name*, is too long.

Explanation: The requested operation cannot be completed successfully because the length of the schema name exceeds DB2's limit for schema names.

User response: Specify a valid, shorter schema name. For more information on the schema name length limit, refer to IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference and try the operation again.

msgcode: -301

sqlstate: 38S41

GSE0302N The specified table name, *table-name*, is too long.

Explanation: The requested operation cannot be completed successfully because the length of the table name exceeds DB2's limit for table names.

User response: Specify a valid, shorter table name.

For more information on the table name length limit, refer to IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference, and then try the operation again.

msgcode: -302

sqlstate: 38S42

GSE0303N The specified column name, *column-name*, is too long.

Explanation: The requested operation cannot be completed successfully because the length of the column name exceeds DB2's limit for column names.

User response: Specify a valid, shorter column name.

For more information on the column name length limit, refer to IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference, and then try the operation again.

msgcode: -303

sqlstate: 38S43

GSE0304N The specified index name, *index-name*, is too long.

Explanation: The requested operation cannot be completed successfully because the length of the index name exceeds DB2's limit for index names.

User response: Specify a valid, shorter index name. For more information on the index name length limit, refer to IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference, and then try the operation again.

msgcode: -304

sqlstate: 38S44

GSE0305N The specified data type name, *type-name*, is too long.

Explanation: The requested operation cannot be completed successfully because the length of the data type name exceeds DB2's limit for data type names.

User response: Specify a valid, shorter type name. For more information on the data type name length limit, refer to IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference, and then try the operation again.

msgcode: -305

sqlstate: 38S45

GSE0306N A complete path that starts with *path* would exceed the acceptable limit of *limit* bytes.

Explanation: The file in a path that starts with *path* cannot be accessed because the length of the complete path would exceed the limit of *limit* bytes. As a result, the statement that you submitted cannot be executed successfully.

User response: Change the location of the file to be accessed so that it can be found using a shorter path and re-submit the statement that you specified. On UNIX systems, symbolic links can be used to establish a shorter path name.

msgcode: -306

sqlstate: 38S46

GSE0307N The length of a dynamic SQL statement *statement-length* would exceed the acceptable limit of *limit* bytes.

Explanation: The statement cannot be constructed because it would be too long.

User response: If the statement is constructed in the context of a stored procedure, verify that the WHERE clause is not too long. If necessary, shorten the WHERE clause and retry the operation. If the problem persists,

contact IBM Software Support.

msgcode: -307

sqlstate: 38S47

GSE0308N The string, *string*, exceeds the limit of *limit* bytes.

Explanation: The requested operation cannot be completed successfully because the string, *string*, is too long.

User response: Specify a shorter string. If necessary, contact IBM Software Support.

msgcode: -308

sqlstate: 38S48

GSE1000N Spatial Extender could not perform an operation *operation-name* that was requested under user id *user-id*.

Explanation: You requested this operation under a user id that does not hold the privilege or authority to perform the operation.

User response: Consult the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference to find out what the required authorization for the operation is.

msgcode: -1000

sqlstate: 38S50

GSE1001N The specified value, *value*, is not valid for the *argument-name* argument.

Explanation: The value *value* that you entered for argument *argument-name* was incorrect or misspelled.

User response: Consult the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference to find out what value or range of values you need to specify.

msgcode: -1001

sqlstate: 38S51

GSE1002N A required argument, *argument-name*, was not specified.

Explanation: The requested operation cannot be completed successfully because an argument that it requires was not specified.

User response: Specify argument *argument-name* with the value that you want; then request the operation again.

msgcode: -1002

sqlstate: 38S52

GSE1003N The spatial column, *schema-name.table-name.column-name*, could not be registered with the spatial reference system *srs-name* because it is already registered with another spatial reference system.

Explanation: A spatial reference system is already registered with the spatial column. It cannot be registered again unless it is unregistered first.

User response: Either unregister the spatial column and then register it with the spatial reference system you want or do not attempt to register it again.

msgcode: -1003

sqlstate: 38S53

GSE1004N The specified geocoder, *geocoder-name*, is already registered.

Explanation: This geocoder is already registered. It cannot be registered again unless it is unregistered first.

User response: Register the geocoder with a unique name, or unregister the existing geocoder first and then try the operation again.

msgcode: -1004

sqlstate: 38S54

GSE1005N Geocoding has already been set up for the spatial column named *schema-name.table-name.column-name*.

Explanation: Geocoding has already been set up for this column. Geocoding cannot be set up again unless the current setup is removed first.

User response: Remove the existing geocoding setup, or choose a column for which geocoding is already set up.

msgcode: -1005

sqlstate: 38S55

GSE1006N The spatial column *schema-name.table-name.column-name* is not registered.

Explanation: This spatial column was not registered with a spatial reference system. Therefore, it cannot be unregistered.

User response: Specify a spatial column that is already registered, or do not attempt to unregister the column.

msgcode: -1006

sqlstate: 38S56

GSE1007N The specified geocoder *geocoder-name* is not registered.

Explanation: A geocoder named *geocoder-name* has not been registered. Therefore, it cannot be unregistered.

User response: Specify a geocoder that is already registered, or do not attempt to unregister the geocoder.

msgcode: -1007

sqlstate: 38S57

GSE1008N A geocoder whose numeric identifier is *geocoder-id* is not registered.

Explanation: A geocoder whose numeric identifier is *geocoder-id* is not registered. It cannot be unregistered.

User response: Specify a geocoder that is already registered, or do not attempt to unregister the geocoder.

msgcode: -1008

sqlstate: 38S58

GSE1009N A table named *schema-name.table-name* does not exist.

Explanation: The requested operation cannot be completed successfully because the table *schema-name.table-name* does not exist.

User response: Specify a valid table name and retry the operation.

msgcode: -1009

sqlstate: 38S59

GSE1010N A spatial column named *schema-name.table-name.column-name* does not exist.

Explanation: The requested operation cannot be completed successfully because *schema-name.table-name.column-name* does not identify an existing column.

User response: Specify a valid spatial column name and retry the operation.

msgcode: -1010

sqlstate: 38S5A

GSE1011N A data type named *schema-name.type-name* does not exist.

Explanation: The requested operation cannot be completed successfully because a data type *schema-name.type-name* does not exist.

User response: Specify a valid data type name and retry the operation.

msgcode: -1011

sqlstate: 38S5B

GSE1012N The database has not been enabled for spatial operations.

Explanation: The requested operation cannot be completed successfully because the database has not been enabled for spatial operations and, therefore, a Spatial Extender catalog has not been created.

User response: Enable the database for spatial operations.

msgcode: -1012

sqlstate: 38S5C

GSE1013N The database is already enabled for spatial operations.

Explanation: The database is already enabled for spatial operations. It cannot be enabled again.

User response: Verify that the database has been enabled as you expected. If necessary, disable the database.

msgcode: -1013

sqlstate: 38S5D

GSE1014N Spatial Extender was unable to register a column named *schema-name.table-name.column-name* because it is not a spatial column.

Explanation: Either this column does not have a spatial data type, or it does not belong to a local table.

User response: Define a spatial data type for column *schema-name.table-name.column-name*, or specify a column with a spatial data type as declared type.

msgcode: -1014

sqlstate: 38S5E

GSE1015N A spatial reference system named *srs-name* does not exist.

Explanation: The requested operation cannot be completed successfully because a spatial reference system with the name *srs-name* does not exist.

User response: Specify an existing spatial reference system and retry the operation.

msgcode: -1015

sqlstate: 38S5F

GSE1016N A spatial reference system whose numeric identifier is *srs-id* does not exist.

Explanation: The requested operation could not be completed successfully because a spatial reference

system with the specified numeric identifier *srs-id* does not exist.

User response: Specify an existing spatial reference system identifier and retry the operation.

msgcode: -1016

sqlstate: 38S5G

GSE1017N A coordinate system named *coordsys-name* already exists.

Explanation: A coordinate system named *coordsys-name* already exists. Another coordinate system with the same name cannot be created.

User response: Specify a unique name for the new coordinate system.

msgcode: -1017

sqlstate: 38S5H

GSE1018N A coordinate system named *coordsys-name* does not exist.

Explanation: The requested operation cannot be completed successfully because a coordinate system with the name *coordsys-name* does not exist.

User response: Specify the name of an existing coordinate system.

msgcode: -1018

sqlstate: 38S5I

GSE1019N No values of the spatial coordinate system *coordsys-name* are specified.

Explanation: You attempted to alter the coordinate system *coordsys-name*, but did not specify any new values.

User response: Specify at least one new value for the coordinate system.

msgcode: -1019

sqlstate: 38S5J

GSE1020N A spatial reference system named *srs-name* already exists.

Explanation: A spatial reference system named *srs-name* already exists. Another spatial reference system with the same name cannot be created.

User response: Specify a unique name for the spatial reference system to be created and retry the operation.

msgcode: -1020

sqlstate: 38S5K

GSE1021N A spatial reference system named *srs-name* does not exist.

Explanation: The requested operation cannot be completed successfully because a spatial reference system with the name *srs-name* does not exist.

User response: Specify a name of an existing spatial reference system and retry the operation.

msgcode: -1021

sqlstate: 38S5L

GSE1022N A spatial reference system whose numeric identifier is *srs-id* does not exist.

Explanation: The requested operation cannot be completed successfully because a spatial reference system with the numeric identifier *srs-id* does not exist.

User response: Specify an existing numeric identifier for the spatial reference system.

msgcode: -1022

sqlstate: 38S5M

GSE1023N A coordinate system whose numeric identifier is *coordsys-id* does not exist.

Explanation: The requested operation cannot be completed successfully because a coordinate system with the numeric identifier *coordsys-id* does not exist.

User response: Specify an existing numeric identifier for the coordinate system and retry the operation.

msgcode: -1023

sqlstate: 38S5N

GSE1024N No values of the spatial reference system *srs-name* are specified.

Explanation: You attempted to alter the spatial reference system *srs-name*, but did not specify any new values.

User response: Specify at least one new value for the spatial reference system and then try the operation again.

msgcode: -1024

sqlstate: 38S5O

GSE1025N A geocoder whose function name is *schema-name.function-name* could not be found in the database.

Explanation: The requested operation cannot be completed successfully because Spatial Extender could not locate a function named *schema-name.function-name* for the geocoder.

User response: Specify a geocoder with an existing function name or create the function, then try the operation again.

msgcode: -1025

sqlstate: 38S5P

GSE1026N The specified number of default parameter values (*number1* values), does not match with the number of default parameter values (*number2* values) that the specified geocoder requires.

Explanation: The specified geocoder could not be registered because you did not specify the complete list for all the parameter values that the geocoder requires.

User response: Correct the number of default parameter values. Specify a default value for each of the parameters of the geocoder or specify a null value for the default parameter values.

msgcode: -1026

sqlstate: 38S5Q

GSE1027N The specified number of parameter description values (*number1* values) does not match with the number of parameter description values (*number2* values) that the specified geocoder requires.

Explanation: The specified geocoder could not be registered because you did not specify the complete list for all the parameter description values that the geocoder requires.

User response: Specify a description for each of the parameters of the geocoder, or specify a null value for the parameter descriptions.

msgcode: -1027

sqlstate: 38S5R

GSE1028N A geocoder named *geocoder-name* does not exist.

Explanation: The requested operation cannot be completed successfully because a geocoder with the name *geocoder-name* does not exist.

User response: Specify an existing geocoder name or register the geocoder with the name that you specified and then try the operation again.

msgcode: -1028

sqlstate: 38S5S

GSE1029N **Geocoding has not been set up for a column named *schema-name.table-name.column-name*.**

Explanation: The requested operation cannot be completed successfully because geocoding has not been set up for the column named *schema-name.table-name.column-name*.

User response: Set up geocoding for the specified column, or correct the table schema, table name, and column name.

msgcode: -1029

sqlstate: 38S5T

GSE1030N **Auto-geocoding is enabled for column *schema-name.table-name.column-name* and the geocoding setup cannot be removed.**

Explanation: Auto-geocoding have been enabled for the column named *schema-name.table-name.column-name*. Therefore, the geocoding setup for this column cannot be removed.

User response: Disable auto-geocoding for the column.

msgcode: -1030

sqlstate: 38S5U

GSE1031N **Auto-geocoding for a column named *schema-name.table-name.column-name* is not enabled.**

Explanation: Auto-geocoding has not been enabled for column *schema-name.table-name.column-name*. Therefore, auto-geocoding cannot be disabled for this column.

User response: Specify a correct column name on which auto-geocoding is enabled.

msgcode: -1031

sqlstate: 38S5V

GSE1032N **Auto-geocoding for a column named *schema-name.table-name.column-name* is already enabled.**

Explanation: Auto-geocoding has already been enabled for the column named *schema-name.table-name.column-name*. Auto-geocoding cannot be enabled again for this column unless it is disabled first.

User response: Specify a correct name for a column for which (1) auto_geocoding is not enabled and (2) geocoding is set up.

msgcode: -1032

sqlstate: 38S5W

GSE1033N **A geocoder whose numeric identifier is *geocoder-id* does not exist.**

Explanation: The requested operation cannot be completed successfully because a geocoder with the numeric identifier *geocoder-id* does not exist.

User response: Specify an existing numeric identifier for the geocoder, or register a geocoder with numeric identifier *geocoder-id*.

msgcode: -1033

sqlstate: 38S5X

GSE1034N **The parameters passed to the stored procedure do not include a parameter *parameter-name*.**

Explanation: The SQLDA that was passed to the stored procedure is too small. It does not contain an entry for parameter *parameter-name*.

User response: Correct the parameters that are being passed to the stored procedure.

msgcode: -1034

sqlstate: 38S5Y

GSE1035N **The *parameter-name* parameter, which is being passed to the stored procedure, has an incorrect data type.**

Explanation: The data type for parameter *parameter-name* that is passed to the stored procedure is not correct.

User response: Correct the parameters that are being passed to the stored procedure.

msgcode: -1035

sqlstate: 38S5Z

GSE1036W **The operation was successful. But values of certain database manager and database configuration parameters should be increased.**

Explanation: The operation was successful, but Spatial Extender needs more database and database manager resources. These resources can be obtained by increasing the values of certain database manager and database configuration parameters.

User response: Some configuration parameters should be increased. For example, for the database configuration, check the APPLHEAPSZ parameter. Refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference for more details.

msgcode: +1036

sqlstate: 38S60

GSE1037N The definition of the specified coordinate system named *coordsys-name* is invalid.

Explanation: A coordinate system named *coordsys-name* cannot be created because the definition given for it is invalid.

User response: Specify a correct definition for the coordinate system.

The function ST_EqualCoordsys can be used to verify the definition by comparing the coordinate system with itself.

msgcode: -1037

sqlstate: 38S61

GSE1038N The WHERE clause specified for the geocoder named *geocoder-name* is invalid. When Spatial Extender attempted to verify the clause, it encountered SQL error *sql-error*.

Explanation: The geocoding that you requested cannot be completed successfully because the where clause that determines which rows are to be geocoded is invalid.

User response: Specify a syntactically correct WHERE clause.

msgcode: -1038

sqlstate: 38S62

GSE1039N A coordinate system identified by the specified identifier *organization-coordsys-id* in combination with the specified organization *organization* already exists.

Explanation: Your request to create a coordinate system could not be met because the combination of identifiers that you specified for the coordinate system (the name of the organization that defined the system and a number that this organization assigned to it) was not unique. Either these two values must be unique in combination, or they must be null.

User response: Specify a unique set of values for *organization* and *organization-coordsys-id*, or choose null values for both.

msgcode: -1039

sqlstate: 38S63

GSE1040N A spatial reference system with the numeric identifier *srs-id* already exists.

Explanation: Your request to create a spatial reference system could not be met because the numeric identifier *srs-id* that you assigned to it already identifies another

spatial reference system. A spatial reference system's identifier must be unique.

User response: Specify a unique numeric identifier for the spatial reference system.

msgcode: -1040

sqlstate: 38S64

GSE1041N A coordinate system with the numeric identifier *coordsys-id* already exists.

Explanation: Your request to create a coordinate system could not be met because the numeric identifier *coordsys-id* that you assigned to it already identifies another spatial coordinate system. A spatial coordinate system's identifier must be unique.

User response: Specify a unique value *coordsys-id* for the coordinate system.

msgcode: -1041

sqlstate: 38S65

GSE1042N A geocoder with the numeric identifier *geocoder-id* already exists.

Explanation: Your request to register a geocoder could not be met because the numeric identifier for the geocoder *geocoder-id* already identifies another geocoder. A geocoder's identifier must be unique.

User response: Specify a unique value for the numeric geocoder identifier.

msgcode: -1042

sqlstate: 38S66

GSE1043N The specified grid index *schema-name.index-name* already exists.

Explanation: This index already exists. It must be dropped before an index with the same name can be created.

User response: Specify a name for the index that does not yet exist, or drop the existing index and retry the operation.

msgcode: -1043

sqlstate: 38S67

GSE1044N The specified coordinate system *coordsys-name* cannot be dropped because an existing spatial reference system is based on this coordinate system.

Explanation: At least one spatial reference system exists that is based on the specified coordinate system *coordsys-name*. The coordinate system cannot be dropped.

User response: Drop all spatial reference systems that

are based on the specified coordinate system. Then try to drop the coordinate system again.

msgcode: -1044

sqlstate: 38S68

GSE1045N The specified spatial reference system *srs-name* cannot be dropped because a spatial column is registered with this spatial reference system.

Explanation: At least one spatial column exists that is associated with the specified spatial reference system *srs-name*. The spatial reference system cannot be dropped.

User response: Unregister all spatial columns that are associated with the specified spatial reference system. Then try to drop the spatial reference system again.

msgcode: -1045

sqlstate: 38S69

GSE1046N The specified geocoder *geocoder-name* cannot be unregistered because it is used in a geocoding setup.

Explanation: At least one geocoding setup exists that uses the specified geocoder *geocoder-name*. The geocoder cannot be unregistered.

User response: Remove all geocoding setups that use the specified geocoder. Then try to unregister the geocoder again.

msgcode: -1046

sqlstate: 38S6A

GSE1047N Geocoder parameter validation failed. **SQLERROR** = *sql-error*.

Explanation: Validation of the geocoder parameters failed. DB2 returned *sql-error*. Geocoder parameters may have been specified when the geocoder was registered, during geocoding setup or both.

User response: Use the information in *sql-error* to determine which parameter is invalid. Correct the value and resubmit your request.

msgcode: -1047

sqlstate: 38S6B

GSE1048N The spatial reference system with numeric identifier *srs-id* is a predefined geodetic spatial reference system and cannot be altered.

Explanation: The spatial reference system was not altered. Spatial reference systems with numeric identifiers in the range 2000000000 to 2000000317 are

predefined geodetic spatial reference systems and cannot be altered.

User response: Do not attempt to alter this spatial reference system. If a geodetic spatial reference system with a different definition is needed, you can create a new geodetic spatial reference system with the numeric identifier in the range 2000000318 to 2000001000.

msgcode: -1048

sqlstate: 38SP3

GSE1049N The spatial reference system with numeric identifier *srs-id* is a predefined geodetic spatial reference system and cannot be dropped.

Explanation: The spatial reference system was not altered. Spatial reference systems with numeric identifiers in the range 2000000000 to 2000000317 are predefined geodetic spatial reference systems and cannot be dropped.

User response: Do not attempt to drop this spatial reference system. If a geodetic spatial reference system with a different definition is needed, you can create a new geodetic spatial reference system with the numeric identifier in the range 2000000318 to 2000001000.

msgcode: -1049

sqlstate: 38SP4

GSE2100N The number of attribute columns being imported (*input-columns* columns) does not match the number of attribute columns in the target table (*table-columns* columns).

Explanation: If you are importing columns that contain attribute data, you have the choice of either specifying or not specifying which attribute columns are being imported and which columns are in the target table. If you specify these values, this error occurs when the specified number of attribute columns being imported differs from the specified number of columns in the target table. If you do not specify these values, this error occurs when the actual number of columns being imported differs from the actual number of attribute columns in the target table.

User response: Make sure that the number of specified or actual attribute columns being imported matches the number of specified or actual columns in the target table.

msgcode: -2100

sqlstate: 38S70

GSE2101N The data type *schema-name.type-name* to be used during import is unknown to DB2.

Explanation: The spatial data type *schema-name.type-name* cannot be used during the import of spatial data because it does not exist in the database.

User response: Create the data type in the database or use a data type that exists.

msgcode: -2101

sqlstate: 38S71

GSE2102N The table specified for import, *schema-name.table-name*, does not exist.

Explanation: A table named *schema-name.table-name* does not exist in the database. Also, Spatial Extender was not asked to create a table to hold the data that is to be imported. The data was not imported.

User response: If the table is to be created by Spatial Extender, specify the appropriate flag. Otherwise, create the table and retry the operation.

msgcode: -2102

sqlstate: 38S72

GSE2103N The table specified for import *schema-name.table-name* already exists.

Explanation: Spatial Extender was asked to create a table named *schema-name.table-name* for the imported data, but a table with that name already exists in the database. No data was imported.

User response: If the table is not to be created by Spatial Extender, do not indicate that the table is to be created. Otherwise, specify the name for a table which does not yet exist in the database.

msgcode: -2103

sqlstate: 38S73

GSE2104N The column *schema-name.table-name.column-name* to import data into does not exist.

Explanation: The column into which you want to import data *column-name* does not exist in the table *schema-name.table-name*. No data can be imported into it.

User response: Correct the column name or create the column in the table that is to be imported, or correct the table name.

msgcode: -2104

sqlstate: 38S74

GSE2105W The import operation completed successfully but not all records from the file were imported.

Explanation: The import operation completed successfully but not all records from the file were imported. The exception file contains the records that could not be imported, and the messages file the contains information why those records were not imported.

User response: Consult the messages file for the reason why not all records were imported, correct the problem and repeat the operation with the original file or the exception file.

msgcode: +2105

sqlstate: 38S75

GSE2106N The data type of the column *schema-name.table-name.column-name* is *column-type*, which does not match the expected type *expected-type* for the data to be imported from file.

Explanation: The column *column-name* in the table *schema-name.table-name* to import data into has a declared type *column-type*. *column-type* does not match the type name *expected-type* for the data to be imported from the file. No data can be imported.

User response: Verify the definition of the table with the structure of the file to be imported.

msgcode: -2106

sqlstate: 38S76

GSE2107N The table to import data into could not be created due to error *sql-error*.

Explanation: Spatial Extender was asked to create a table to import data into, but the table could not be created successfully. DB2 returned *sql-error*.

User response: Refer to the description of this *sql-error*.

msgcode: -2107

sqlstate: 38S77

GSE2108N The method specification *method* to identify the attribute columns to be imported from the file is not correct.

Explanation: Either no method specification was given or *method* is not a valid method specification. Only 'N' and 'P' are supported method specifications for importing spatial data from a file.

User response: Correct the method specification and try the method again.

msgcode: -2108

sqlstate: 38S78

GSE2109N A character *found-char* was found when a character *expected-char* was expected.

Explanation: An unexpected character *found-char* was found in the string that identifies the attribute columns to be imported from the file but *expected-char* was expected. The statement cannot be processed successfully.

User response: Correct the string that identifies the attribute columns to be imported from the file.

msgcode: -2109

sqlstate: 38S79

GSE2110N The column position identifier *position* in the string *string* is invalid.

Explanation: The column position identifier *position* specified in the string starting with *string* is not in the valid range. Only values greater than 0 (zero) and less than or equal to the number of columns in the file to be imported can be specified. The statement cannot be processed successfully.

User response: Correct the column position identifier.

msgcode: -2110

sqlstate: 38S7A

GSE2111N A column named *dbf-column-name* in the dBASE file is too long.

Explanation: The name of column *dbf-column-name* in the dBASE file (.dbf) exceeds DB2's limit for column names.

User response: Specify a *dbf-column-name* that does not exceed DB2's length limit.

msgcode: -2111

sqlstate: 38S7B

GSE2112N The column *dbf-column-name* cannot be found in the dBASE file.

Explanation: The name *dbf-column-name* does not identify an existing attribute column in the dBASE file (.dbf). The operation cannot be completed successfully.

User response: Specify a column name that exists in the dBASE file.

msgcode: -2112

sqlstate: 38S7C

GSE2113N The dBASE file data type *dbf-data-type* for the column *dbf-column-name* in the dBASE file is not supported.

Explanation: The dBASE file data type *dbf-data-type* for the attribute column *dbf-column-name* in the dBASE file (.dbf) cannot be mapped to a data type in the DB2 database. The shape file cannot be imported.

User response: Exclude the column from the column list.

msgcode: -2113

sqlstate: 38S7D

GSE2114N The column position *position* is out of range. The dBASE file contains *dbf-column-number* columns.

Explanation: The specified column position *position* must be a value within the valid range. A valid value must be greater than 0 (zero) and less than or equal to the *dbf-column-number*.

User response: Specify a valid position.

msgcode: -2114

sqlstate: 38S7E

GSE2115N A spatial reference system whose numeric identifier is *srs-id* does not exist.

Explanation: A spatial reference system whose numeric identifier is *srs-id* does not exist. The data cannot be imported.

User response: Either specify an existing spatial reference system, or create the spatial reference system before attempting the import operation.

msgcode: -2115

sqlstate: 38S7F

GSE2116N The coordinate system definition *coordsys-def* is too long.

Explanation: The coordinate system definition *coordsys-def* used for the spatial data to be imported is too long. It could not be verified with the coordinate system that underlies the spatial reference system that is to be used for the imported data.

User response: Verify that the coordinate system defined in the projection file (.prj) is correct. To skip the verification step, do not supply the projection file.

msgcode: -2116

sqlstate: 38S7G

GSE2117N The coordinate system definition *coordsys-def* does not match the coordinate system definition on which the spatial reference system *srs-id* is based.

Explanation: The coordinate system *coordsys-def* does not match the coordinate system on which the spatial reference system *srs-id* is based. Both coordinate systems must be semantically identical.

User response: Verify that the coordinate system defined in the projection file (.prj) matches the coordinate system of the spatial reference system. To skip the verification step, do not supply the projection file.

msgcode: -2117

sqlstate: 38S7H

GSE2118N The spatial data does not fit into the spatial reference system with the numeric identifier *srs-id*.

Explanation: The spatial data covers an area that exceeds the minimum and maximum coordinates of the spatial reference system with the numeric identifier *srs-id*.

User response: Specify a spatial reference system which may fully contain the spatial data to be imported. Refer to the DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS catalog view for the minimum and maximum coordinate values applicable for the spatial reference system.

msgcode: -2118

sqlstate: 38S7I

GSE2119N The imported data's spatial reference system, whose numerical identifier is *srs-id1*, does not match the target column's spatial reference system, whose numeric identifier is *srs-id2*. The target column's name is *schema-name.table-name.column-name*.

Explanation: The spatial column *schema-name.table-name.column-name* was registered with a spatial reference system *srs-id2* assigned to it. This spatial reference system does not match the spatial reference system *srs-id1*, which is used for the spatial data to be imported into that column. The data cannot be imported.

User response: Either unregister the spatial column, or specify the same spatial reference system for the data to be imported that the column uses.

msgcode: -2119

sqlstate: 38S7J

GSE2120N No data was imported.

Explanation: None of the shape data could be imported. All rows were rejected and written to the exception file.

User response: Consult the message file for the reasons why all the spatial data was rejected.

msgcode: -2120

sqlstate: 38S7K

GSE2121N The value *restart-count* specifying the record to restart the operation is out of range. The shape file contains *record-count* records.

Explanation: The specified restartCount *restart-count* must be a greater than or equal to 0 (zero), and less than or equal to *record-count*.

User response: Specify a valid number for restartCount, or specify a null value for the resartCount.

msgcode: -2121

sqlstate: 38S7L

GSE2122N The SQL statement used to import the shape data does not fit into the internal buffer.

Explanation: The SQL statement used to import the shape data into the table does not fit into the internal buffer. A possible reason for this might be too many columns are in the file.

User response: Import a smaller set of the attribute columns in the shape file.

msgcode: -2122

sqlstate: 38S7M

GSE2123N A buffer to hold the data for *row-count* rows cannot be allocated.

Explanation: Spatial Extender tried to use a single INSERT statement to import at least *row-count* rows, but a buffer to hold the data for these rows could not be allocated. Too much memory was required.

User response: Specify a commit count for the import that is less than *row-count*. Or, specify a smaller set of columns to be imported. This will reduce the amount of memory required.

msgcode: -2123

sqlstate: 38S7N

GSE2124N An invalid type identifier *type-id* was found in the header of the shape file to be imported.

Explanation: The data in the shape file does not appear to have a valid spatial data type. The shape file is possibly corrupted. The data was not imported.

User response: Verify that the shape file is valid.

msgcode: -2124

sqlstate: 38S7O

GSE2125N A column in the shape file has an unsupported data type *type*.

Explanation: The shape file contains a column whose data type is not supported by Spatial Extender. The shape file could not be imported.

User response: Import only a smaller set of the columns of the shape file and omit the column with the unsupported data type.

msgcode: -2125

sqlstate: 38S7P

GSE2126N The header of the shape file *shape-file* is invalid.

Explanation: The header of the shape file *shape-file* is invalid. The shape file cannot be imported.

The extension of the file name *shape-file* indicates in which part of the shape file the error was encountered. The file extensions include:

.shp main file
.shx index file
.dbf dBASE file
.pri projection file

User response: Verify and correct the header of the shape file.

msgcode: -2126

sqlstate: 38S7Q

GSE2127N The offset *offset* for the record *record-number* in the shape index file *shx-file* is invalid.

Explanation: The offset *offset* for the record *record-number* in the index file (.shx) *shx-file* is invalid. The offset must be greater than or equal to 50 and less than the total length of the main file (.shp) of the shape file. The offset is measured in 16 bit words.

User response: Verify and correct the shape file.

msgcode: -2127

sqlstate: 38S7R

GSE2128N The length of the shape in record *record-number* of the shape index file *shx-file* is too short.

Explanation: The length of the shape in record *record-number* found in the shape index file *shx-file* is too short. Each shape must consist of at least 4 bytes (two 16 bit words).

User response: Verify and correct the shape file.

msgcode: -2128

sqlstate: 38S7S

GSE2129N Spatial Extender found an incorrect record number *record-number* in the shape file *shp-file* when expecting record number *expected-number*.

Explanation: Spatial Extender found an incorrect record number *record-number* in the shape file *shp-file* when expecting record number *expected-number*.

User response: Verify and correct the shape file.

msgcode: -2129

sqlstate: 38S7T

GSE2130N The size of the shape data *record-size* indicated in the shape file *shp-file* does not match the size indicated in the shape index file *index-size*.

Explanation: The size of the shape data *record-size* indicated in the shape file *shp-file* does not match the size indicated in the shape index file *index-size*.

The main file of the shape file (.shp) is not consistent with the index file (.shx) and cannot be processed further.

User response: Verify and correct the shape file.

msgcode: -2130

sqlstate: 38S7U

GSE2131N The data for record *record-number* in the dBASE file *dbf-file* is invalid.

Explanation: The data for record *record-number* in the dBASE file *dbf-file* that contains the attribute information associated with the geometries in the shape file is invalid.

Possible explanations are:

- The first byte of the record is neither an asterisk (*) nor a space (' ').
- The sum of all lengths of the columns in the dBASE file (.dbf) must be equal to the record size indicated in the header of the file.

User response: Verify and correct the dBASE file.

msgcode: -2131

sqlstate: 38S7V

GSE2132N The data in shape file *shape-file* is invalid.

Explanation: The data in shape file *shape-file* is corrupted. This shape file cannot be imported.

The file name *shape-file* indicates in which part of the shape file the error was encountered.

User response: Verify and correct the shape file.

msgcode: -2132

sqlstate: 38S7W

GSE2133N The import operation failed because the column *schema-name.table-name.column-name* is not nullable.

Explanation: The definition of the column *column-name* in the existing table *schema-name.table-name* indicates that the column may not contain nulls. The column is not included in the list of columns to be imported and DB2 would not produce the values for that column by any other means like default values, a generated column definition or any triggers.

The import operation cannot be completed successfully.

User response: Include the column in the list of columns to be imported, identify the column as *id-column*, or define an alternate way for DB2 to generate the values for that column during the import operation.

msgcode: -2133

sqlstate: 38S7X

GSE2134N The spatial reference system associated with the data to be imported is not identical with the spatial reference system with the numeric identifier *srs-id*.

Explanation: The spatial data in the file to be imported uses a spatial reference system with different offsets and scale factors than the spatial reference system with the numeric identifier *srs-id*. The data cannot be imported successfully.

User response: Specify a spatial reference system which has the same definition as the spatial reference system required by the data in the file to be imported. Refer to the DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS catalog view for the minimum and maximum coordinate values and the offsets and scale factors applicable for the spatial reference system.

msgcode: -2134

sqlstate: 38S7Y

GSE2200N The results from the SELECT statement did not include any spatial column.

Explanation: The specified SELECT statement for export operations must reference exactly one spatial column. No data can be exported.

User response: Correct the SELECT statement so that it references exactly one spatial column.

msgcode: -2200

sqlstate: 38S90

GSE2201N The results from the SELECT statement included more than one spatial column.

Explanation: The specified SELECT statement for export operations must specify exactly one spatial column. No data can be exported.

User response: Correct the SELECT statement so that it references exactly one spatial column.

msgcode: -2201

sqlstate: 38S91

GSE2202N The SQL data type *data-type* for column *column-number* in the fullselect describing the data to be exported is not supported for shape files.

Explanation: The SQL data type *data-type* for column *column-number* in the fullselect describing the data to be exported is not supported for shape files. No data can be exported.

User response: Do not specify a column with the unsupported type in the SELECT statement for the export shape operation. You can use the DESCRIBE command to verify the data types for all columns in the result of the SELECT statement.

msgcode: -2202

sqlstate: 38S92

GSE2203N The data in column *column-number* has a length of *length*, which is too long to be exported to a shape file.

Explanation: The data in column *column-number* has a length of *length*, which is too long to be exported to a shape file.

User response: Change the length of the column to be shorter than 256 bytes. You can use DB2's cast functions to shorten the length of the column.

msgcode: -2203

sqlstate: 38S93

GSE2204N The column name *column-name* exceeds the limit of 11 characters for dBASE files.

Explanation: The column name *column-name* exceeds the restriction of 11 bytes for dBASE files. No data is exported.

User response: Specify a column name with 11 or less bytes for the attribute data to be exported to the dBASE file (.dbf).

msgcode: -2204

sqlstate: 38S94

GSE2205W DB2 returned a SQL warning *sql-warning* when Spatial Extender fetched the data to be exported.

Explanation: DB2 returned a SQL warning *sql-warning* when Spatial Extender fetched the data to be exported. The warning was ignored and the data exported.

User response: Refer to the description of *sql-warning*. If the warning is not acceptable, correct the source of the warning and export the data again.

msgcode: +2205

sqlstate: 38S95

GSE2206W No data was exported.

Explanation: The SELECT statement specified for the export did not retrieve any rows. No data was exported.

User response: Specify a SELECT statement that returns at least one row to be exported.

msgcode: +2206

sqlstate: 38S96

GSE2207W Not all rows referenced by the SELECT statement were exported.

Explanation: The SELECT statement identified more rows than were exported. This error can occur if an error is encountered during the export and the operation is terminated.

User response: Consult the message file.

msgcode: +2207

sqlstate: 38S97

GSE2208N In row *row-number*, a geometry to be exported is associated with a spatial reference system whose numeric identifier is *srs-id1*. This identifier does not match the numeric identifier *srs-id2* of the spatial reference system used for geometries in the preceding rows.

Explanation: In row *row-number*, a geometry to be exported is associated with a spatial reference system whose numeric identifier is *srs-id1*. This identifier does not match the numeric identifier *srs-id2* of the spatial reference system used for geometries in the preceding rows. The row *row-number* was not exported.

User response: Verify that all geometries to be exported have the same spatial reference system identifier. If not, use the ST_Transform method in the SELECT statement to convert all geometries to the same spatial reference system.

msgcode: -2208

sqlstate: 38S98

GSE2209N The geometry in row *row-number* to be exported is a null value.

Explanation: The notion of null values is not supported by the data in the file to which you want to export geometries. The geometry to be exported in row *row-number* is a null value and cannot be exported.

User response: Exclude rows with geometries that are null values from the export by modifying the SELECT statement accordingly.

msgcode: -2209

sqlstate: 38S99

GSE2210N The spatial reference system for the geometry in row *row-number* is invalid. This spatial reference system's numerical identifier is *srs-id*.

Explanation: In row *row-number*, the geometry that is to be exported uses an invalid spatial reference system. The geometry cannot be exported.

User response: Correct the indicated geometry or exclude the row from the export operation by modifying the SELECT statement accordingly.

msgcode: -2210

sqlstate: 38S9A

GSE2211N The spatial data to be exported is associated with a spatial reference system whose numeric identifier is *srs-id*. The coordinate system on which this spatial reference system is based does not match the coordinate system definition *coordsys-def* for the file to which the exported data is to be appended.

Explanation: The spatial data to be exported is associated with a spatial reference system whose numeric identifier is *srs-id*. The coordinate system on which this spatial reference system is based does not match the coordinate system *coordsys-def* for the file to

which the exported data is to be appended. No data was exported.

User response: Append the data to a file which has a matching coordinate system, transform the spatial data to the corresponding coordinate system, or export the data to a separate file.

msgcode: -2211

sqlstate: 38S9B

GSE2212N The attribute data does not match the dBASE file.

Explanation: The specified SELECT statement for the export operation produces attribute data that does not match the dBASE file (.dbf).

The attribute data from the SELECT statement for the export operation cannot be appended to the dBASE file. Possible causes are:

- The number of columns do not match.
- The data type of column does not match.
- The attribute column name does not match.

User response: Either correct the SELECT statement, or do not append to the specified dBASE file.

msgcode: -2212

sqlstate: 38S9C

GSE2213W The geometry to be exported to the file in row *row-number* is a null value.

Explanation: The geometry to be exported to the file in row *row-number* is a null value. The row was not exported.

User response: Modify the SELECT statement to exclude geometries that have null values.

msgcode: +2213

sqlstate: 38S9D

GSE2214W Not enough memory is available to export the geometry in row *row-number*.

Explanation: Not enough memory is available to export the geometry in row *row-number*. The row was not exported, but the exporting operation continued.

User response: Provide more memory, or modify the SELECT statement to exclude the geometry from that row.

msgcode: +2214

sqlstate: 38S9E

GSE2215W The geometry to be appended to the shape file as number *record-number* has a shape type *geometry-shape-type* which does not match the type in the shape file *file-shape-type*.

Explanation: The geometry to be appended to the shape file as number *record-number* has an shape type *geometry-shape-type* which does not match the type in the shape file *file-shape-type*. The geometry cannot be exported to that file.

Possible shape types are:

- | | |
|----|---|
| 0 | Empty geometry. |
| 1 | Point without Z coordinates and measures. |
| 3 | Linestring or multilinestring without Z coordinates and measures. |
| 5 | Polygon or multipolygon without Z coordinates and measures. |
| 8 | Multipoint without Z coordinates and measures. |
| 9 | Point with Z coordinate and no measures. |
| 10 | Linestring or multilinestring with Z coordinates and no measures. |
| 11 | Point with Z coordinate and measure. |
| 13 | Linestring or multilinestring with Z coordinates and measures. |
| 15 | Polygon or multipolygon with Z coordinates and measures. |
| 18 | Multipoint with Z coordinates and measures. |
| 19 | Polygon or multipolygon with Z coordinates and no measures. |
| 20 | Multipoint with Z coordinates and no measures. |
| 21 | Point with measure and no Z coordinates. |
| 23 | Linestring or multilinestring with measures and no Z coordinates. |
| 25 | Polygon or multipolygon with measures and no Z coordinates. |
| 28 | Multipoint with measures and no Z coordinates. |

User response: Convert the geometry to the correct type, or export it to another shape file.

msgcode: +2215

sqlstate: 38S9F

GSE2216N The shape data to be exported is invalid.

Explanation: Before Spatial Extender exports geometries to a shape file, it converts them into shape data. This error was returned because the shape data into which a geometry was converted is invalid. The geometry was not exported.

Possible explanations are:

- The shape data has an odd number of bytes.
- The shape data is too short to contain all information.
- An unknown shape indication was returned.
- The information about the minimum bounding rectangle is not consistent with the type indication.

User response: Contact IBM Software Support.

msgcode: -2216

sqlstate: 38S9G

GSE2217N The column name *column-name* is a duplicate in the output column list.

Explanation: The column name *column-name* appears more than once in the output column list. All column names in the list must be unique. No data is exported.

User response: Specify a unique column name in the output column list.

All column names for dBASE files (.dbf) must at most consist of 11 bytes, and all column names for SDEX files must at most consist of 32 bytes.

msgcode: -2217

sqlstate: 38S9I

GSE2299N The shape file *file-name* has an invalid file size.

Explanation: Shape file *file-name* has an invalid file size. Shape files have a file size that is a multiple of 16-bit words. Therefore, their size is always even. The shape file might be corrupt. It cannot be used.

User response: Verify and correct the shape file.

msgcode: -2299

sqlstate: 38S9H

GSE2500N The header of the SDEX file *file-name* is invalid.

Explanation: The header of the SDEX file *file-name* is invalid. The SDEX file cannot be opened successfully.

User response: Verify and correct the header of the SDEX file.

msgcode: -2500

sqlstate: 38SA0

GSE2501N The coordinate system parameters could not be read from the SDEX file *file-name*.

Explanation: The SDEX file *file-name* contains an invalid data for the coordinate system definition. The SDEX file cannot be opened successfully.

User response: Verify and correct the coordinate system definition stored in the SDEX file.

msgcode: -2501

sqlstate: 38SA1

GSE2502N The column definitions could not be read from the SDEX file *file-name*.

Explanation: The SDEX file *file-name* contains invalid data for the column definitions. The SDEX file cannot be opened successfully.

User response: Verify and correct the column definitions stored in the SDEX file.

msgcode: -2502

sqlstate: 38SA2

GSE2503N A column in the SDEX file has an unsupported data type *type*.

Explanation: The SDEX file contains a column whose data type is not supported by Spatial Extender. The SDEX file could not be imported.

User response: Import only a smaller set of the columns of the SDEX file and omit the column with the unsupported data type.

msgcode: -2503

sqlstate: 38SA3

GSE2504N The column position *position* is out of range. The SDEX file contains only *column-number* columns.

Explanation: The specified column position *position* must be a value within the valid range. A valid value must greater than 0 (zero) and less than or equal to the *column-number*.

User response: Specify a valid position.

msgcode: -2504

sqlstate: 38SA4

GSE2505N A column named *column-name* in the SDEX file is too long.

Explanation: The name of column *column-name* in the SDEX file exceeds DB2's limit for column names.

User response: Specify a *column-name* that does not exceed DB2's length limit.

msgcode: -2505

sqlstate: 38SA5

GSE2506N The column *column-name* cannot be found in the SDEX file.

Explanation: The name *column-name* does not identify an existing attribute column in the SDEX file. The operation cannot be completed successfully.

User response: Specify a column name that exists in the SDEX file.

msgcode: -2506

sqlstate: 38SA6

GSE2507N The SDEX file data type *data-type* for the column *column-name* in the SDEX file is not supported.

Explanation: The SDEX file data type *data-type* for the attribute column *column-name* in the SDEX file cannot be mapped to a data type in the DB2 database. The SDEX file cannot be imported.

User response: Exclude the column from the column list.

msgcode: -2507

sqlstate: 38SA7

GSE2508N Multiple spatial columns exist in the SDEX file *file-name*.

Explanation: Spatial Extender detected the existence of more than one spatial column in the SDEX file *file-name*. Only one spatial column is allowed. The file cannot be imported successfully.

User response: Correct the SDEX file so that it contains only one spatial column.

msgcode: -2508

sqlstate: 38SA8

GSE2509N Could not read record *record-number* from the SDEX file *file-name*. Reason code: *reason-code*

Explanation: The record *record-number* could not be read from the SDEX file *file-name*.

Reasons why the record could not be read, preceded by their reason codes, are as follows:

-1 An internal error occurred.

-13 Not enough memory is available.

User response: Depending on the reason code, verify

and correct the SDEX file, or resolve the memory shortage.

msgcode: -2509

sqlstate: 38SA9

GSE2600N Could not write header of the SDEX file *file-name*.

Explanation: The header of the SDEX file *file-name* could not be written successfully.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -2600

sqlstate: 38SF0

GSE2601N Could not write coordinate system parameters to the SDEX file *file-name*.

Explanation: The coordinate system parameters could not be written to the SDEX file *file-name* successfully.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -2601

sqlstate: 38SF1

GSE2602N Could not write column definitions to the SDEX file *file-name*.

Explanation: The column definitions could not be written to the SDEX file *file-name* successfully.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -2602

sqlstate: 38SF2

GSE2603N The column name *column-name* exceeds the limit of 32 characters for SDEX files.

Explanation: The column name *column-name* exceeds the restriction of 32 bytes for SDEX files. No data is exported.

User response: Specify a column name with 32 or less bytes for the data to be exported to the SDEX file.

msgcode: -2603

sqlstate: 38SF3

GSE2604N The SQL data type *data-type* for column *column-number* in the fullselect describing the data to be exported is not supported for SDEX files.

Explanation: The SQL data type *data-type* for column *column-number* in the fullselect describing the data to be

exported is not supported for SDEX files. No data can be exported.

User response: Do not specify a column with the unsupported type in the SELECT statement for the export sde operation. You can use the DESCRIBE command to verify the data types for all columns in the result of the SELECT statement.

msgcode: -2604

sqlstate: 38SF4

GSE2605N Could not write record *record-number* to the SDEX file *file-name*.

Explanation: The record *record-number* could not be written to the SDEX file *file-name* successfully.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -2605

sqlstate: 38SF5

GSE3000N Null SRS identifier.

Explanation: A null value was passed to the function or method instead of a numeric spatial reference system identifier.

User response: Specify a numeric spatial reference system identifier for an existing spatial reference system. Refer to the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS for the defined spatial reference systems.

msgcode: -3000

sqlstate: 38SU0

GSE3001N Invalid SRS identifier *srs-id*.

Explanation: The spatial reference system identifier *srs-id* that was provided to the spatial function or method does not identify an existing spatial reference system.

User response: Specify an existing numeric spatial reference system identifier that is defined in the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS, or create a spatial reference system that is identified by *srs-id*.

msgcode: -3001

sqlstate: 38SU1

GSE3002N Null unit name.

Explanation: A null was specified as a unit of measure. The specification for a unit of measure must be the unit itself (for example, "meter"). It cannot be a null.

User response: Either omit the unit of measure when calling the spatial function or method, or specify an existing unit of measure. Consult the Spatial Extender catalog view DB2GSE.ST_UNITS_OF_MEASURE for supported units.

msgcode: -3002

sqlstate: 38SU2

GSE3003N Unknown unit *unit-name*.

Explanation: The unit *unit-name* that was provided to the spatial function or method does not identify an existing unit of measure.

User response: Either omit the unit of measure when calling the spatial function or method, or specify an existing unit of measure. Consult the Spatial Extender catalog view DB2GSE.ST_UNITS_OF_MEASURE for supported units.

msgcode: -3003

sqlstate: 38SU3

GSE3004N Unsupported conversion to unit *unit-name*.

Explanation: The conversion to the unit *unit-name* is not supported.

The functions ST_Area, ST_Buffer, ST_Length, and ST_Perimeter cannot accept a linear unit of measure if the given geometry is not in a projected coordinate system.

User response: Use one of the following methods:

- Omit the unit of measure when calling the spatial function or method.
- Specify an angular unit of measure.
- Project the geometry into a projected coordinate system using the ST_Transform function. Consult the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS for applicable spatial reference system.

msgcode: -3004

sqlstate: 38SU4

GSE3005N No unit in SRS.

Explanation: The spatial reference system for the geometry does not have an associated linear or angular unit. The operation cannot be performed in the requested unit of measure.

User response: Either represent the geometry in a correct spatial reference system, which does have an associated linear or angular unit of measure, or omit the unit parameter when you request the operation.

msgcode: -3005

sqlstate: 38SU5

GSE3006N Invalid internal type id.

Explanation: The internal data type identifier for this geometry is a null value and therefore invalid.

This error can occur if the internal representation of the geometry is corrupted, or if the geometry was not constructed by one of the supported constructor functions or methods.

User response: Construct the geometry again by using one of the supported constructor functions or methods.

msgcode: -3006

sqlstate: 38SU6

GSE3007N Unknown internal type id *type-id*.

Explanation: The value of the internal type identifier *type-id* for the geometry is not valid.

This error can occur if the internal representation of the geometry is corrupted, or if the geometry was not constructed by one of the supported constructor functions or methods.

User response: Construct the geometry again by using one of the supported constructor functions or methods.

msgcode: -3007

sqlstate: 38SU7

GSE3008N Internal type id mismatch (*type-id1*, *type-id2*).

Explanation: A mismatch of internal data type identifiers was found. Spatial Extender expected to retrieve a geometry whose internal data type identifier is *type-id2*, but instead retrieved a geometry whose internal data type identifier is *type-id1*.

This error can occur if the internal representation of the geometry is corrupted, or if the geometry was not constructed by one of the supported constructor functions or methods.

User response: Construct the geometry again by using one of the supported constructor functions or methods.

msgcode: -3008

sqlstate: 38SU8

GSE3009W Invalid part number *part-number*.

Explanation: The specified part number *part-number* is not valid. A null value was returned.

User response: If the geometry is not empty, then specify a valid part number, which should be greater than 0 (zero) and less than or equal to the maximum number of parts in the geometry collection.

You can use the ST_NumGeometries function to determine the number of parts of the geometry collection.

If the geometry is empty, the method should not be applied.

msgcode: +3009

sqlstate: 01HS0

GSE3010W Invalid ring number *ring-number*.

Explanation: The specified number *ring-number* for an internal ring is not valid. A null value was returned.

User response: If the polygon value is not empty, then specify a valid ring number, which should be greater than or equal to 1 (one) and less than or equal to the maximum number of interior rings in the polygon.

If the polygon is empty, the function or method should not be applied. You can use the function ST_NumInteriorRings to determine the number of interior rings of the polygon.

msgcode: +3010

sqlstate: 01HS1

GSE3011W Invalid point number *point-number*.

Explanation: The specified point number *point-number* is not valid. A null value was returned.

User response: If the curve value is not empty, then specify a valid point number, which should be greater than 0 (zero) and less than or equal to the maximum number of points in the curve. If the curve is empty, the function or method should not be applied.

You can use the ST_NumPoints function to determine the number of points used to define the curve.

msgcode: +3011

sqlstate: 01HS2

GSE3012N Invalid DE9-IM *matrix*.

Explanation: The intersection matrix *matrix* specified for the ST_Relate function is not valid. The matrix must be exactly 9 characters long, and each character in the matrix must be one of the following: 'T', 'F', '0', '1', '2', or '*'.

User response: Specify a valid intersection matrix.

msgcode: -3012

sqlstate: 38SU9

GSE3013N Exterior ring is no ring.

Explanation: The linestring that is to serve as the new exterior ring for the polygon is not a ring. To be a ring, the linestring must be both simple and closed. One or both of these two conditions is not met.

User response: Specify a simple and closed linestring for the new exterior ring of the polygon.

msgcode: -3013

sqlstate: 38SUA

GSE3014N Interior ring is no ring.

Explanation: The linestring that is to serve as a new interior ring for the polygon is not a ring. To be a ring, the linestring must be both simple and closed. At least one of these two conditions is not met.

User response: Specify a simple and closed linestring for the new interior ring of the polygon.

msgcode: -3014

sqlstate: 38SUB

GSE3015N Reason code = *reason-code*. Transformation to SRS *srs-id* failed.

Explanation: The geometry could not be transformed from the spatial reference system it is represented into the spatial reference system with the numeric identifier *srs-id*. The transform failed with reason code *reason-code*.

The reason codes have the following meanings:

- 2008** The geometry is invalid.
- 2018** Not enough memory is available to successfully complete the transformation.
- 2020** The spatial reference systems are not compatible. Both spatial reference systems must be based directly or indirectly on the same geographic coordinate system.
- 2021** One or more points of the resulting geometry would be outside the maximum possible extent for the new spatial reference system. The resulting geometry cannot be represented in the new spatial reference system.
- 2025** The definition of the new spatial reference system is not valid.
- 2026** An internal error occurred during the projection of the geometry.

User response: Represent the geometry in a spatial reference system that can be transformed into the spatial reference system identified by *srs-id*, or specify a different spatial reference system identifier to transform the geometry into.

msgcode: -3015

sqlstate: 38SUC

GSE3016N Unsupported cast *type-id1*, *type-id2*.

Explanation: The attempted cast operation from the data type with the internal type identifier *type-id1* to the data type with the internal type identifier *type-id2* is not supported. The geometry cannot be processed further.

User response: Specify a supported cast operation. For more information, refer to the IBM DB2 SQL Reference for the supported cast functions.

msgcode: -3016

sqlstate: 38SUD

GSE3020N Invalid Z coordinate and measure combination.

Explanation: The geometries that are to be processed by the function or method are not represented using the same dimensions with respect to their Z coordinates and measures.

All the geometries must either contain Z coordinates or contain no Z coordinates. All the geometries must either contain measures or contain no measures.

User response: Provide geometries to the function or method that are represented using the same dimensions with respect to their Z coordinates and measures.

msgcode: -3020

sqlstate: 38SUH

GSE3021N Reason code = *reason-code*. Locator failure.

Explanation: An internal error occurred when a spatial function or method operated on a LOB locator. The reason code *reason-code* was returned by a locator function.

User response: Refer to the DB2 Application Development Guide to determine the meaning of *reason-code* returned from the LOB locator operation and correct the problem. If the problem persists, contact IBM Software Support.

msgcode: -3021

sqlstate: 38SUI

GSE3022N Representation too long (*append-length* vs. *written-length* bytes).

Explanation: The representation of the geometry in Geographic Markup Language (GML), well-known text (WKT), well-known binary (WKB), or the shape representation would be too long. From *append-length* bytes, only *written-length* bytes could be appended to the encoding. A representation of the geometry cannot be created.

User response: Simplify the geometry by omitting points that are not essential for the geometry. You can use the ST_Generalize function for this procedure. Alternatively, break down the geometry into several smaller geometries.

msgcode: -3022

sqlstate: 38SUJ

GSE3023N Representation too short (*length bytes*).

Explanation: The representation of the geometry in well-known binary (WKB) representation or the shape representation is only *length* bytes long. It needs to have at least 4 bytes for the shape representation, exactly 5 bytes for the well-known binary representation for empty geometries, and at least 9 bytes for the well-known binary representation for non-empty geometries. The binary representation must also be long enough to contain all of the geometry points.

User response: Provide a valid well-known binary representation or shape representation to the function or method.

msgcode: -3023

sqlstate: 38SUK

GSE3024N Internal geometry too short.

Explanation: The internal representation of the geometry is too short. It could not be processed further.

This error can occur if the internal representation of the geometry is corrupted, or if the geometry was not constructed by one of the supported constructor functions or methods.

User response: Construct the geometry again using one of the supported constructor functions or methods.

msgcode: -3024

sqlstate: 38SUL

GSE3025N Geometry inconsistent.

Explanation: The geometry value is inconsistent and cannot be processed any further.

User response: Recreate the geometry from a valid binary or text representation.

msgcode: -3025

sqlstate: 38SUM

GSE3026N Inconsistent no. of points (*indicated-number vs. data-number*).

Explanation: An internal parameter of the geometry indicates that the geometry data contains *indicated-number* points. But the actual geometry data

contains *data-number* points. Because of this inconsistency, the geometry will not be used further in the processing.

This error can occur if the internal representation of the geometry is corrupted, or if the geometry was not constructed by one of the supported constructor functions or methods.

User response: Recreate the geometry using the functions or methods supported by Spatial Extender.

msgcode: -3026

sqlstate: 38SUN

GSE3027N Point is empty.

Explanation: It is invalid to specify an X coordinate, Y coordinate, Z coordinate, or measure for an empty point.

If the point is constructed by the constructor function ST_Point, the point's X and Y coordinates must both be null. Furthermore, no Z coordinate or measure should be specified unless it is a null value.

If the mutators ST_X, ST_Y, ST_Z, or ST_M are used to modify an empty point, the point's X and Y coordinates must both be null. No Z coordinate or measure should be specified unless it is null.

User response: Use mutators ST_X, ST_Y, ST_Z, or ST_M to modify points that are not empty, or construct the point by specifying both X and Y coordinates with values that are not null.

msgcode: -3027

sqlstate: 38SUO

GSE3028N Inconsistent coordinates.

Explanation: If a new point is constructed, both the X and Y coordinates must be specified. Both coordinates must be either null or not null.

If both coordinate values are null, the resulting point will be empty. In that case, no Z coordinate or measure should be specified unless it is null.

User response: Specify null values for both the X and Y coordinates, or specify values that are not null for both coordinates.

msgcode: -3028

sqlstate: 38SUP

GSE3029N Invalid byte order *byte-order*.

Explanation: The byte order in the binary representation of the geometry must be either 0 (zero) or 1 (one), but it is *byte-order*.

In the well-known binary representation, a byte order of 0 (zero) indicates big endian format, and a byte

order of 1 (one) indicates little endian format.

User response: Correct the byte order in the binary representation so that it is either 0 (zero) or 1 (one).

msgcode: -3029

sqlstate: 38SUQ

GSE3030N Invalid number of points *num-points* in geometry.

Explanation: The geometry has an invalid number of points *num-points*. This number must be greater than or equal to 0 (zero).

If the geometry is not empty, then the following conditions must be met:

point The geometry must have exactly one point.

linestring
The geometry must have 2 or more points defining it.

polygon
The geometry must have 3 or more points defining it.

User response: Construct the geometry by using the functions or methods supported by Spatial Extender.

msgcode: -3030

sqlstate: 38SUR

GSE3031N Invalid extent (*min-coord* vs. *max-coord*) in geometry.

Explanation: The extent of the geometry in one of the dimensions is invalid. The minimum coordinate *min-coord* must be less than or equal to the maximum coordinate *max-coord* for all dimensions of the geometry.

User response: Construct the geometry by using the functions or methods supported by Spatial Extender.

msgcode: -3031

sqlstate: 38SUS

GSE3032N Aggregation failure.

Explanation: A mismatch between internal identifiers was encountered for the computation of a spatial aggregate.

Aggregate functions are not supported if used in any of the following situations:

- In the partitioned environment.
- A GROUP BY clause is used in the query that contains the spatial aggregate.
- Any function other than the DB2 aggregate function MAX is used.

- The aggregate function is not used in the correct context.

User response: Make sure that you use the aggregate function in a way that is supported by Spatial Extender.

msgcode: -3032

sqlstate: 38SUT

GSE3033N Invalid binary data (type ids *type-id1*, *type-id2*).

Explanation: A binary representation that is passed as input to this spatial function or method has to represent a geometry whose data type identifier is *type-id2*. But the representation that was actually passed to the function or method represents a geometry whose data type identifier is *type-id1*. No geometry could be constructed.

User response: Either call the correct function or method which constructs geometries of type *type-id2* or correct the binary representation to represent a geometry of *type-id1*.

msgcode: -3033

sqlstate: 38SUU

GSE3034N Invalid text data (type ids *type-id1*, *type-id2*).

Explanation: A text representation that is passed as input to this spatial function or method has to represent a geometry whose data type identifier is *type-id2*. But the representation that was actually passed to the function represents a geometry whose data type identifier is *type-id1*. No geometry could be constructed.

User response: Either call the correct function which constructs geometries of type *type-id1* or correct the text representation to represent a geometry of *type-id2*.

msgcode: -3034

sqlstate: 38SUV

GSE3035W Curve not changed.

Explanation: The curve was not changed because the specified point to be appended to the curve was empty.

User response: Append a point that is not empty to the curve.

msgcode: +3035

sqlstate: 01HS3

GSE3036W Geometry not accurate.

Explanation: The resulting geometry could not be represented accurately in the spatial reference system. One of the scale factors is too small and does not allow for a high enough precision to represent each point that

defines the resulting geometry.

For example, consider a linestring with a well-known text representation of 'linestring m (10 10 8, 10 11 12)' represented in a spatial reference system that includes a scale factor of 1 (one) for X coordinates and a scale factor of 1 (one) also for Y coordinates. If the function ST_MeasureBetween is applied to that linestring, and the upper and lower bounds for the measures are 9 and 10, respectively, the resulting linestring, represented in its well-known text representation, would have to be 'linestring m (10 10.25 9, 10 10.50 10)'. However, the scale factor of 1 (one) for the Y coordinates prevents the representation of fractions. The coordinates 10.25 and 10.50 cannot be represented without rounding that would produce an incorrect result. Such coordinates will be removed from the geometry.

User response: Represent the geometry in a spatial reference system that uses larger scale factors. Alternatively, choose different parameters that influence the resulting geometry.

msgcode: +3036

sqlstate: 01HS4

GSE3037N Invalid GML, expecting *char* instead of *string* at position *position*.

Explanation: A character *char* was expected in the Geography Markup Language of the geometry, but the text *string* was found instead at position *position*. The GML representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the GML representation and construct the geometry again.

msgcode: -3037

sqlstate: 38SUW

GSE3038N Invalid GML, expecting *expected-tag* instead of *given-tag* at position *position*.

Explanation: The tag *given-tag* was found in the Geography Markup Language of the geometry at position *position*, but a tag *expected-tag* was expected. The GML representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the GML representation and construct the geometry again.

msgcode: -3038

sqlstate: 38SUX

GSE3039N Invalid GML, expecting *number* instead of *text* at position *position*.

Explanation: Unexpected text *text* was found in the Geography Markup Language of the geometry at position *position*. A number representing a coordinate

was expected instead. The GML representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the GML representation and construct the geometry again.

msgcode: -3039

sqlstate: 38SUY

GSE3040N Invalid GML type *type*.

Explanation: An unknown type *type* was specified in the Geography Markup Language of the geometry. The GML supports points, linestrings, polygons, multipoints, multilinestrings, and multipolygons. Spatial Extender cannot construct the geometry successfully.

User response: Correct the GML representation and construct the geometry again.

msgcode: -3040

sqlstate: 38SUZ

GSE3041N GML point has been incorrectly specified.

Explanation: The problem occurred due to one of the following reasons:

- A point, represented using the Geography Markup Language, can only have one set of coordinates. The given point had either no set of coordinates or more than one set.
- The set of coordinates is not enclosed by corresponding <gml:coord> or <gml:coordinates> tags.

The GML representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the GML representation and construct the geometry again.

msgcode: -3041

sqlstate: 38SV0

GSE3042N Could not read *number-bytes* bytes from locator at offset *offset*. Total length of data is *length*.

Explanation: An attempt was made to read *number-bytes* bytes from the locator, starting at the *offset*. This exceeds the total length of the data *length* that is referenced by the locator. The data might be truncated.

For binary representations of a geometry, the binary representation might indicate an invalid binary encoding. The encoded geometry has fewer points than the header indicates.

User response: Verify and correct the representation of the geometry. Make sure that the binary or textual representation does not get truncated before it is passed to the Spatial Extender function.

msgcode: -3042

sqlstate: 38SV1

GSE3043N Invalid number of parts *number-parts*.

Explanation: The number of parts *number-parts* indicated in the binary representation of the geometry is invalid. The number of parts must be larger than 0 (zero) and match the actual number of parts supplied in the encoding.

User response: Specify the correct number of parts or supply all parts for the geometry.

msgcode: -3043

sqlstate: 38SV2

GSE3044N Invalid number of rings *number-rings*.

Explanation: The number of rings *number-rings* indicated in the binary representation of the polygon or multipolygon is invalid. The number of rings must be larger than 0 (zero) and match the actual number of parts supplied in the encoding.

User response: Specify the correct number of rings or supply all rings for the geometry.

msgcode: -3044

sqlstate: 38SV3

GSE3045N Invalid part offset *part-offset* in shape.

Explanation: An invalid offset *part-offset* for a part in the shape representation of the geometry was encountered. A part offset must be larger than or equal to 0 (zero), and each part offset must be larger than the preceding one. The shape representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the shape representation and construct the geometry again.

msgcode: -3045

sqlstate: 38SV4

GSE3046N Invalid type ID *type-id* in shape.

Explanation: The shape representation of the geometry contains an invalid type identifier *type-id*. The shape data is possibly corrupted. Spatial Extender cannot construct the geometry successfully.

User response: Verify and correct the shape representation of the geometry.

msgcode: -3046

sqlstate: 38SV5

GSE3047N Invalid length *shape-length* of shape encoding for type *type*, expecting only *expected-length* bytes.

Explanation: The shape encoding contains *shape-length* bytes, which is too long. To encode a geometry of the specified type *type*, only *expected-length* bytes are required. The shape data is possibly corrupted. Spatial Extender cannot construct the geometry successfully.

User response: Verify and correct the shape representation of the geometry.

msgcode: -3047

sqlstate: 38SV6

GSE3048N Invalid WKT format, expecting *char* instead of *string*.

Explanation: A character *char* was expected in the well-known text representation of the geometry, but the text *string* was found instead. The well-known text representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the well-known text representation and construct the geometry again.

msgcode: -3048

sqlstate: 38SV7

GSE3049N Invalid WKT format, expecting a number instead of *text*.

Explanation: An unexpected text *text* was found in the well-known text representation of the geometry. A number representing a coordinate was expected instead. The well-known text representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the well-known text representation and construct the geometry again.

msgcode: -3049

sqlstate: 38SV8

GSE3050N Unexpected parenthesis in WKT format at *text*.

Explanation: An unexpected opening or closing parenthesis was found in the well-known text representation of the geometry at *text*. The well-known text representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the well-known text representation and construct the geometry again.

msgcode: -3050

sqlstate: 38SV9

GSE3051N **Parenthesis mismatch in WKT format, expecting *parenthesis*.**

Explanation: The end of the well-known text representation was reached unexpectedly. A parenthesis *parenthesis* was expected. The well-known text representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the well-known text representation and construct the geometry again.

msgcode: -3051

sqlstate: 38SVA

GSE3052N **Unknown type *type* in WKT.**

Explanation: The well-known text representation of the geometry contains an unknown type name of *type*. The well-known text representation is not valid. Spatial Extender cannot construct the geometry successfully.

User response: Correct the well-known text representation and construct the geometry again.

msgcode: -3052

sqlstate: 38SVB

GSE3053N **Invalid type id *type-id* in WKB.**

Explanation: The well-known binary representation of the geometry contains an invalid type identifier *type-id*. The data is possibly corrupted. Spatial Extender cannot construct the geometry successfully.

The type identifiers of separate parts in a geometry collection (multipoint, multilinestring, or multipolygon) must have the same indicators for the Z and M coordinates as the geometry collection itself.

User response: Verify and correct the well-known binary representation of the geometry.

msgcode: -3053

sqlstate: 38SVC

GSE3300N **Invalid grid size *grid-size-number*.**

Explanation: The grid size identified by its position *grid-size-number* is invalid. One of the following invalid specifications was made when the grid index was created with the CREATE INDEX statement:

- A number less than 0 (zero) was specified as the grid size for the first, second, or third grid level.
- 0 (zero) was specified as the grid size for the first grid level.
- The grid size specified for the second grid level is less than the grid size of the first grid level but it is not 0 (zero).

- The grid size specified for the third grid level is less than the grid size of the second grid level but it is not 0 (zero).
- The grid size specified for the third grid level is greater than 0 (zero) but the grid size specified for the second grid level is 0 (zero).

The function ST_GetIndexParms can be used to retrieve the values used for the parameters specified when the index was created.

User response: Drop the grid index and create a new grid index using valid grid sizes only.

msgcode: -3300

sqlstate: 38SI0

GSE3301N **Invalid z-order parameter *parameter-number*.**

Explanation: The parameter identified by its position *parameter-number* for a Z-Order index contains an invalid value. One of the following invalid specifications was made in the CREATE INDEX statement that was used to create the index to which the geometry is to be added:

- A null value was specified for the parameter.
- A negative number was specified for a scale factor (this rule applies to parameter numbers 2 and 4 only).

The function ST_GetIndexParms can be used to retrieve the values used for the parameters specified when the index was created.

User response: Drop the spatial z-order index and create a new index using only valid parameters.

msgcode: -3301

sqlstate: 38SI1

GSE3302N **No point to be indexed.**

Explanation: The geometry to be indexed using a Z-Order index is not a point. The Z-Order index supports only points, and the index entry cannot be generated.

User response: Do not insert a geometry that is not a point into a column that has a Z-Order index defined on it. Either drop the index or do not insert the geometry.

msgcode: -3302

sqlstate: 38SI2

GSE3303N Invalid quad tree parameter
parameter-number.

Explanation: An invalid parameter was specified when the quad tree index was created. The parameter is identified by its position *grid-size-number*.

One of the following invalid specifications was made:

- A null value was specified for the parameter.
- A negative number was specified for a scale factor (this rule applies to parameter numbers 3 and 5 only).
- A value less than 1 (one) was specified for the first parameter.

The function ST_GetIndexParms can be used to retrieve the values used for the parameters specified when the index was created.

User response: Drop the spatial quad tree index and create a new index using only valid parameters.

msgcode: -3303

sqlstate: 38SI3

GSE3400C Unknown error *error-code*.

Explanation: An internal error with code *error-code* was encountered when a geometry was processed.

User response: Note the error and contact IBM Software Support.

msgcode: -3400

sqlstate: 38SS0

GSE3402C Insufficient memory.

Explanation: Not enough memory was available for the spatial function or method that you invoked.

User response: Make more memory available to the DB2 process that executes the function or method.

msgcode: -3402

sqlstate: 38SS2

GSE3403N Invalid geometry type.

Explanation: An invalid type of geometry was passed to the function or method that you invoked.

User response: Specify a valid geometry. For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3403

sqlstate: 38SS3

GSE3405N Too many parts specified.

Explanation: The number of parts indicated in the binary or text representation of the geometry is greater than the actual number of parts supplied. Either the number of parts indicated is too high or not all the parts were supplied.

User response: Specify the correct number of parts or supply all parts for the geometry.

msgcode: -3405

sqlstate: 38SS5

GSE3406N Incorrect geometry type.

Explanation: The wrong type of geometry was passed to the function or method that you invoked. For example, a linestring might have been passed to a function or method that takes only polygons as input.

User response: Either pass to the function or method a type of geometry that it can process, or use a function or method that accepts the type of geometry that you want to pass.

msgcode: -3406

sqlstate: 38SS6

GSE3407N Text is too long.

Explanation: The geometry contains too much detail to be converted to its well-known text representation. The well-known text representation exceeds the maximum allowable length (2 gigabytes).

User response: Simplify the geometry - for example, by using the ST_Generalize function - or convert the geometry to its well-known binary representation.

msgcode: -3407

sqlstate: 38SS7

GSE3408N Invalid parameter value.

Explanation: An invalid parameter was encountered.

User response: Refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference for the function's correct syntax and retry the operation. If the problem persists, contact IBM Software Support.

msgcode: -3408

sqlstate: 38SS8

GSE3409N Invalid geometry produced.

Explanation: The parameters provided for the function or method have produced an invalid geometry; for example, an invalid shape representation.

An invalid geometry is one that violates a geometry property.

User response: Construct the geometry again from a valid representation.

msgcode: -3409

sqlstate: 38SS9

GSE3410N Incompatible geometries.

Explanation: The function or method expected two geometries of a certain type and did not receive them. For example, the ST_AddPoint function expects two geometries, one a representation and the other a point.

User response: Specify geometries that the function or method accepts as valid input. To determine what types of geometries are valid for this function or method, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3410

sqlstate: 38SSA

GSE3411N Invalid geometry.

Explanation: The function or method cannot process the geometry passed to it because one or more properties of the geometry violate the geometry's integrity.

User response: Use the ST_IsValid function to validate the geometry. Construct the geometry again from a correct representation if it is not valid.

msgcode: -3411

sqlstate: 38SSB

GSE3412N Too many points.

Explanation: The construction of a geometry has exceeded the 1-megabyte storage limit; the geometry has too many points.

User response: Construct a geometry that contains fewer points. Or, if possible, remove some points. For performance and storage considerations, include only those points that are needed to render a geometry.

msgcode: -3412

sqlstate: 38SSC

GSE3413N Geometry too small.

Explanation: The geometry returned by the ST_Difference, ST_Intersection, ST_SymDifference, or ST_Union function is too small to be represented accurately in the current spatial reference system.

For example, this can happen if the internal computation constructs a very thin polygon, but the

scale factor of the spatial reference system is so low that the geometry would collapse to a linestring if it were to be represented in this spatial reference system. It would lose its property as a polygon.

User response: Use a spatial reference system for the calculation which allows for a higher resolution. The ST_Transform function can be used to convert a geometry from one spatial reference system into another.

msgcode: -3413

sqlstate: 38SSD

GSE3414N Buffer out of bounds.

Explanation: The ST_Buffer function has created a buffer around the provided geometry that is outside the range of the coordinates to which the spatial reference system applies.

Refer to the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS to determine the minimum and maximum absolute values for each of the dimensions. These values must not be exceeded by the calculated buffer.

User response: Either reduce the distance to be used for the buffer calculation, or change the spatial reference system in which the calculation is done. The ST_Transform function can be used to convert geometries from one spatial reference system into another.

msgcode: -3414

sqlstate: 38SSE

GSE3415N Invalid scale factor.

Explanation: A scale factor for any of the four dimensions (X, Y, Z, and M) must be greater than or equal to 1 (one).

User response: Use a correctly defined spatial reference system to represent the geometry.

msgcode: -3415

sqlstate: 38SSF

GSE3416N Coordinate out of bounds.

Explanation: A coordinate cannot be represented in the spatial reference system because, in at least one dimension, it exceeds the possible minimum or maximum absolute value within the system's range of values.

Refer to the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS to determine the minimum and maximum absolute values for each of the dimensions.

User response: Determine whether the coordinate is

correct. If it is, determine whether it fits within the extent of the spatial reference system that you are using. For information about this spatial reference system, consult the DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS catalog view.

msgcode: -3416

sqlstate: 38SSG

GSE3417N Invalid coordsys definition.

Explanation: There are one or more errors in the text representation of the definition of the coordinate system on which the geometry's spatial reference system is based. The representation cannot be converted into a valid projection.

User response: Verify the coordinate system definition of the spatial reference system. Alternatively, construct the geometry in a spatial reference system that is associated with a valid coordinate system. The ST_EqualCoordsys function can be used to verify the coordinate system definition by comparing it with itself.

msgcode: -3417

sqlstate: 38SSH

GSE3418N Projection error.

Explanation: An error occurred during an attempt to project a geometry to another spatial reference system.

User response: Make sure that the geometry is within the legal domain of the projection.

msgcode: -3418

sqlstate: 38SSI

GSE3419N Polygon rings overlap.

Explanation: The rings of a polygon overlap. By definition, the inner and outer rings of a polygon must not overlap. They can intersect only at a tangent, which means the rings can only touch but not cross each other.

User response: Specify the coordinates for the polygon that will not produce overlapping rings. Note that the scale factors of the spatial reference system for the geometry have an influence on the precision.

msgcode: -3419

sqlstate: 38SSJ

GSE3420N Too few points.

Explanation: The error is a result of one of the following:

- Linestrings must consist of at least two points, and polygons must consist of at least four points.
- The geometry cannot be constructed from the points that you have specified.

Note that if the geometry to be constructed is empty, these rules do not apply.

User response: Construct the geometry again from a valid set of points.

msgcode: -3420

sqlstate: 38SSK

GSE3421N Polygon is not closed.

Explanation: The inner and outer rings that define the polygon must be closed. A ring is closed if the start and end points are identical in the X and Y dimensions. If the polygon has Z coordinates, then the start and end points must also be identical to the Z coordinates. Note that this rule does not apply to measures, which can be different for the start and end points.

User response: Specify inner and outer rings for the polygon that have the same points for the start and end points in the X and Y dimension. If the polygon has Z coordinates, the start and end points of the Z coordinate points also have to be identical. If the polygon has measures, the start and end points can be different.

msgcode: -3421

sqlstate: 38SSL

GSE3422N Invalid exterior ring.

Explanation: The exterior ring of the polygon is not valid.

The exterior ring of a polygon must enclose all interior rings of the polygon. All interior rings have to be completely inside the area that is defined by the outer ring and must not cross the exterior ring.

User response: Specify a geometry that consists of a valid set of interior and exterior rings, where the interior rings lie fully within the area that is enclosed by the exterior ring to represent it.

If the geometry has multiple polygons, use a multipolygon.

msgcode: -3422

sqlstate: 38SSM

GSE3423N Polygon has no area.

Explanation: The specified polygon lacks an interior that covers an area that is not the empty set in the X and Y dimensions.

A geometry is a polygon only if its coordinates span

two dimensions in the 2-dimensional space defined by the X and Y coordinates.

User response: Specify a polygon that encloses an area that is not empty. If the polygon is empty, construct an empty polygon.

msgcode: -3423

sqlstate: 38SSN

GSE3424N Exterior rings overlap.

Explanation: The exterior rings of distinct polygons in a multipolygon overlap. Distinct polygons in a multipolygon must not overlap, and the boundaries must touch only at a finite number of points. That means the polygons must not share line segments.

The scale factors of the spatial reference system that is used to represent the geometry influences the precision that applies to the coordinates. Rounding operations performed when the geometry is converted to the representation in the spatial reference system might cause a loss in precision and, subsequently, this error.

User response: Specify coordinates for the polygon that will not produce overlapping rings.

Note that the scale factors of the spatial reference system have an influence on precision.

Refer to the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS for the scale factor used for the spatial reference system in which the geometry will be represented.

msgcode: -3424

sqlstate: 38SSO

GSE3425N Polygon intersects itself.

Explanation: A ring of a polygon cannot intersect itself. The start and end points on each ring of the polygon must be reached twice when traversing the ring. All other points must only be reached once. This holds true also for the line segments that define the rings of the polygon.

The scale factors of the spatial reference system that is used to represent the geometry influences the precision that applies to the coordinates. Rounding operations performed when the geometry is converted to the representation in the spatial reference system might cause a loss in precision and, subsequently, this error.

User response: Construct a valid polygon in which the rings do not intersect themselves.

Refer to the Spatial Extender catalog view DB2GSE.ST_SPATIAL_REFERENCE_SYSTEMS for the scale factor used for the spatial reference system in which the geometry will be represented.

msgcode: -3425

sqlstate: 38SSP

GSE3426N Invalid number of parts.

Explanation: The number of parts indicated in the binary or text representation of the geometry is not equal to the actual number of parts supplied. Either the number is too low or too many parts were supplied to the function or method.

User response: Specify the correct number of parts or supply all parts for the geometry.

msgcode: -3426

sqlstate: 38SSQ

GSE3427N Incompatible SRSs.

Explanation: The two spatial reference systems are not compatible. They cannot be transformed into or compared with one another. The operation cannot be completed successfully.

User response: Specify two compatible spatial reference systems.

msgcode: -3427

sqlstate: 38SSR

GSE3428N BLOB too small.

Explanation: The number of bytes in the specified binary representation of the geometry is too small.

User response: Specify a valid binary representation of the geometry.

msgcode: -3428

sqlstate: 38SSS

GSE3429N Invalid geometry type.

Explanation: An invalid internal geometry type was encountered. The geometry is not valid and will not be processed any further.

User response: Construct the geometry again from a valid binary or text representation.

msgcode: -3429

sqlstate: 38SST

GSE3430N Invalid byte order.

Explanation: The byte order in the binary representation of the geometry has an invalid value. The byte order must be 0 (zero) or 1 (one).

In the well-known binary representation, a byte order of 0 (zero) indicates big endianness, and a byte order of 1 (one) indicates little endianness.

User response: Specify a valid byte order in the

binary representation for the geometry.

msgcode: -3430

sqlstate: 38SSU

GSE3431N Empty geometry.

Explanation: An empty geometry was passed to the ST_AsBinary function, even though it is not allowed as input.

User response: Edit the SQL statement that you submitted so that only non-empty geometries will be passed to the ST_AsBinary function. For example, you can use the ST_IsEmpty function in the WHERE clause to exclude empty geometries.

msgcode: -3431

sqlstate: 38SSV

GSE3432N Invalid end point.

Explanation: The specified point is intended to be appended to the curve, but it is not valid.

User response: Specify a valid point to be appended.

msgcode: -3432

sqlstate: 38SSW

GSE3433N Point not found.

Explanation: The specified point is intended to be changed or removed, but it does not exist in the curve.

User response: Specify a point that does exist in the curve.

msgcode: -3433

sqlstate: 38SSX

GSE3500N Instance path not found.

Explanation: The function implementing the geocoder could not find the DB2 instance path.

User response: Verify the correct installation of DB2 and the IBM DB2 Spatial Extender. Check that the DB2INSTANCE environment variable is set for the user running the process executing the function.

If the problem persists, contact IBM Software Support.

msgcode: -3500

sqlstate: 38SG0

GSE3501N Changing SRS id (*new-srs-id*, *previous-srs-id*).

Explanation: The geocoder is optimized to geocode addresses and produce the resulting points all in same spatial reference system in a single SQL statement. But

it encountered different spatial reference systems in the same SQL statement. The new spatial reference system is identified by *new-srs-id*; the spatial reference system used for the preceeding rows was identified by *previous-srs-id*.

User response: Specify the default parameters and overwriting parameters for the column that the geocoder runs on. This should be done in such a way that the numerical spatial reference system identifier remains constant within an SQL statement that is used to geocode multiple addresses at once.

msgcode: -3501

sqlstate: 38SG1

GSE3502N Path too long.

Explanation: The path name specified for the locator file parameter or the base map parameter exceeds 256 bytes and is too long.

User response: Use a shorter path name for the locator file parameter or the base map parameter. On Unix systems, symbolic links can be used to shorten the path name.

msgcode: -3502

sqlstate: 38SG2

GSE3503N Invalid line *line-number* in locator file.

Explanation: In the locator file, line *line-number* contains an invalid entry. Each entry in the file must be of the form "property name = property value".

User response: Correct the locator file.

msgcode: -3503

sqlstate: 38SG3

GSE3504N Non-matching quote in line *line-number* in locator file.

Explanation: In the locator file, line *line-number* contains an entry where the quote characters are not balanced. If quoted, the property name as well as the property value must have an opening and a closing quote character.

User response: Correct the locator file.

msgcode: -3504

sqlstate: 38SG4

GSE3505N Reason code = *reason-code*. Property failure.

Explanation: A failure occurred while handling properties defined in the locator file (.loc).

The reason codes have the following meanings:

- 502 A property is inconsistent with other properties.
- 503 A required property was not found.
- 504 A value of a property is invalid.
- 505 An unexpected property was encountered.
- 506 An array of values was specified for a property that expected only a scalar value.
- 507 A value of a property does not have the expected data type.
- 513 A value of a property is too long.
- 533 The expression specified for a value of a property is invalid.

User response: Correct the problem identified by the reason code; then retry the operation.

msgcode: -3505

sqlstate: 38SG5

GSE3506N Copy of properties failed.

Explanation: The copying of the properties into an internal buffer failed. Not enough memory is available.

User response: Ensure that enough memory is available for the geocoder to copy the properties from the locator file to the internal buffer.

msgcode: -3506

sqlstate: 38SG6

GSE3507N Too many properties.

Explanation: The locator file contains too many properties. At most 2048 properties can be specified in the locator file.

User response: Reduce the number of properties specified in the locator file.

msgcode: -3507

sqlstate: 38SG7

GSE3508N No point produced.

Explanation: The geocoder produced a geometry that is not an ST_Point geometry.

User response: Contact IBM Software Support.

msgcode: -3508

sqlstate: 38SG8

GSE3509N Reason code = *reason-code*. Initialization of the geocoder failed.

Explanation: The initialization of the geocoder failed.

The reason codes have the following meanings:

- 522 A general failure occurred while the geocoder was being initialized.
- 527 The initialization of the match key failed.
- 529 The initialization of the address normalization failed.

User response: Verify the installation of Spatial Extender.

msgcode: -3509

sqlstate: 38SG9

GSE3510N Address not normalized.

Explanation: The geocoder could not normalize the address.

User response: Contact IBM Software Support.

msgcode: -3510

sqlstate: 38SGA

GSE3511N Reason code = *reason-code*. Geocoder file operation failed.

Explanation: An internal error with reason code *reason-code* occurred during a file operation.

Possible reasons for the file operation failure, preceded by their reason codes, are as follows:

- 543 The match rule file (.mat) could not be opened.
- 544 The match rule file (.mat) is invalid.
- 547 The reference data file (.edg) could not be opened.
- 548 The table in the reference data file (.edg) is either missing or inaccessible.
- 549 A required column could not be found in the reference data file (.edg).
- 550 The index file could not be accessed.

User response: Correct the problem identified by the reason code; then retry the operation.

msgcode: -3511

sqlstate: 38SGB

GSE3512N Reason code = *reason-code*. Geocoder failed.

Explanation: The geocoder failed with an internal error of reason code *reason-code*.

User response: Note the internal error and contact IBM Software Support.

msgcode: -3512

sqlstate: 38SGC

GSE3600N No index specified.

Explanation: No valid index was specified. The index schema parameter, the index name parameter, or both, are null. The index parameter values cannot be derived.

User response: Specify a valid spatial index to retrieve the parameter information.

msgcode: -3600

sqlstate: 38SQ0

GSE3601N Invalid spatial index name
schema-name.index-name.

Explanation: The specified name of the index for which you want parameter information retrieved does not exist or does not identify a spatial index. This name is *schema-name.index-name*.

User response: Specify an existing spatial index to retrieve the parameter information.

msgcode: -3601

sqlstate: 38SQ1

GSE3602N Invalid parameter number *number* specified.

Explanation: The parameter number *number* is not valid for the specified spatial index.

The following limits apply for the different types of spatial indexes:

grid index

Parameter numbers between 1 (one) and 3.

z-order index

Parameter numbers between 1 (one) and 4.

quad-tree index

Parameter numbers between 1 (one) and 5.

User response: Specify a valid parameter number for the spatial index. Consult the DB2 system catalog for the type of the spatial index.

msgcode: -3602

sqlstate: 38SQ2

GSE3603N Invalid column name.

Explanation: The specified column does not exist in the table. At least one of the following - table schema, table name, or column name - is a null value. The index parameter for an index on a column cannot be derived.

User response: Specify an existing column which has a spatial index defined on it.

msgcode: -3603

sqlstate: 38SQ3

GSE3701N Distance *distance-value* out of range: valid range is *min* to *max* meters (inclusive).

Explanation: An invalid distance was passed to the function or method that you invoked.

User response: Specify a valid distance and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3701

sqlstate: 38SO9

GSE3702N Out of workspace in internal geometry engine.

Explanation: The geodetic workspace was not large enough to perform the operation.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -3702

sqlstate: 38SOT

GSE3703N Out of heap memory in internal geometry engine.

Explanation: Not enough memory was available. Possible reasons are that the supply of memory was too low, or that memory was being used by other applications.

User response: Resolve the memory shortage and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3703

sqlstate: 38SOU

GSE3704C Possible corrupt data or invalid input in internal geometry engine.

Explanation: Spatial Extender encountered an unexpected internal error in a geometry value.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3704

sqlstate: 38SOV

GSE3706C Error number *hipparchus-error* in internal geometry engine.

Explanation: Spatial Extender encountered an unexpected internal error.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -3706

sqlstate: 38SOX

GSE3708C Internal error: empty dispatch table entry.

Explanation: Spatial Extender encountered an unexpected internal error in a geometry value.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -3708

sqlstate: 38SOR

GSE3709C Internal error.

Explanation: Spatial Extender encountered an unexpected internal error in a geometry value.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -3709

sqlstate: 38SOY

GSE3712N Latitude *latitude-value* out of range.

Explanation: A latitude value must be in a valid range.

For latitude values that are measured in degrees, the valid range is between -90 and 90 degrees (inclusive).

For latitude values that are measured in grads, the valid range is between -100 and 100 grads (inclusive).

For latitude values that are measured in radians, the valid range is between $-\pi/2$ and $\pi/2$ radians

(inclusive) where π is approximately 3.14159265358979323846.

User response: Specify a valid latitude value and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3712

sqlstate: 38SO7

GSE3713N Longitude *longitude-value* out of range.

Explanation: A longitude value must be in a valid range.

For longitude values that are measured in degrees, the valid range is between -180 and 180 degrees (inclusive).

For longitude values that are measured in grads, the valid range is between -200 and 200 grads (inclusive).

For longitude values that are measured in radians, the valid range is between $-\pi$ and π radians (inclusive) where π is approximately 3.14159265358979323846.

User response: Specify a valid longitude value and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3713

sqlstate: 38SO8

GSE3714N Too few rings (*nrings*) specified: Must have at least 1 ring.

Explanation: A non-empty ST_Polygon requires at least one ring.

User response: Specify at least one ring and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3714

sqlstate: 38SOD

GSE3716N Too few points (*npoints*) specified: Must have at least *min* points.

Explanation: A non-empty ST_LineString has too few points. An ST_LineString value must have at least 2 points and an ST_LineString value specified as a ring in an ST_Polygon must have at least 4 points.

User response: Specify the correct number of points in the geometry value and repeat the command.

For more information, refer to the IBM DB2 Spatial

Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3716

sqlstate: 38SOL

GSE3721N Geometry Collections are not supported.

Explanation: ST_GeomCollection values are not supported in the DB2 Geodetic Data Management Feature.

User response: Use ST_MultiPoint, ST_MultiLineString or ST_MultiPolygon type instead of ST_GeomCollection and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3721

sqlstate: 38SP6

GSE3722N Unknown or unsupported WKB type tag: *wkb-type*.

Explanation: The Well-known Binary value is of an unknown or unsupported type.

User response: Use a known Well-known Binary type (ST_Point, ST_LineString, ST_Polygon, ST_MultiPoint, ST_MultiLineString, or ST_MultiPolygon) and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3722

sqlstate: 38SP7

GSE3724N Type mismatch: received type=*given-type*, expected type=*expected-type*.

Explanation: An invalid type of geometry was passed to the function or method that you invoked.

User response: Repeat the command. If the problem persists, contact IBM Software Support.

msgcode: -3724

sqlstate: 38SON

GSE3726N *type* data type version *version-number* is not supported by the DB2 Geodetic Data Management Feature currently in use.

Explanation: The geometry value contains an unsupported version.

User response: Use a supported version and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3726

sqlstate: 38SOZ

GSE3733W Polygon covers more than half the earth. Verify counter-clockwise orientation of the vertex points.

Explanation: Polygon covers more than half the earth.

User response: Verify counter-clockwise orientation of the vertex points.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: +3733

sqlstate: 01HS5

GSE3734N Invalid ring geometry; coincident consecutive points were found at *point-value*.

Explanation: Invalid ring geometry; coincident consecutive points were found.

User response: Remove coincident consecutive points and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3734

sqlstate: 38SQ4

GSE3735N Invalid ring geometry; the sequence of coordinates crosses or touches itself at *point-value*.

Explanation: Invalid ring geometry; the sequence of coordinates crosses or touches itself.

User response: Specify a ring that does not cross or touch.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3735

sqlstate: 38SQ5

GSE3736N Constituent polygon rings do not define a valid region. Check the rotational sense of each ring.

Explanation: Constituent polygon rings do not define a valid region.

User response: Check the rotational sense of each ring.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3736

sqlstate: 38SQ6

GSE3737N Invalid polygon; the region has no boundary.

Explanation: Invalid polygon; the region has no boundary.

User response: Specify a polygon with a boundary.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3737

sqlstate: 38SQ7

GSE3739N Invalid ring geometry; collinear segments were found near *point-value*.

Explanation: Invalid ring geometry; collinear segments were found.

User response: Remove collinear segments.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3739

sqlstate: 38SQ9

GSE3740N Reason code *reason-code*. Unsupported Geodetic operation.

Explanation: The DB2 Geodetic Data Management Feature does not support the operation as indicated by the following reason code:

"1" The Geodetic License is not enabled.

"2" The function was invoked with an ST_Geometry value where the SRID null.

"3" The function was invoked with ST_Geometry values where the SRID values are not equal.

"4" The function was invoked with a spatial reference system that has a null definition.

"5" The function is not supported in Geodetic.

"6" The type of one or more ST_Geometry values is not supported in Geodetic.

"7" The function does not support being executed with Geodetic and non-Geodetic values.

"8" The Geodetic Voronoi Index does not support non-Geodetic values.

"9" The Spatial Grid Index does not support Geodetic values.

User response: The action corresponding to the reason code is:

"1" Enable the Geodetic license.

"2" Ensure all ST_Geometry values have a valid SRID attribute value.

"3" For Geodetic operations, ensure all ST_Geometry values have the same SRID value.

"4" If the problem persists, contact IBM Software Support.

"5" Do not invoke this function or method with ST_Geometry values with an SRID in the geodetic range.

"6" Do not use this ST_Geometry type with an SRID in the geodetic range.

"7" Execute the function with all ST_Geometry values with an SRID in the geodetic range or not in the geodetic range.

"8" Use the Spatial Grid Index specification on columns with non-Geodetic values.

"9" Use the Geodetic Voronoi Index specification on columns with Geodetic values.

After correcting the problem, repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3740

sqlstate: 38SOP

GSE3743N The Voronoi tessellation id *vtid* is not found.

Explanation: The specified Voronoi tessellation id is not defined.

User response: Use a defined Voronoi tessellation id and repeat the command

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3743

sqlstate: 38SOC

GSE3744N The geometry value is too large.

Explanation: The geometry value cannot exceed its maximum value.

User response: Specify fewer points in the geometry

value and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3744

sqlstate: 38SOK

GSE3745N The spatial reference system definition does not have a well defined SPHEROID definition.

Explanation: The spatial reference system definition for this ST_Geometry value is not a geographic coordinate system or does not contain a DATUM definition that includes a SPHEROID definition.

User response: Correct the spatial reference system definition and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3745

sqlstate: 38SOF

GSE3746N The ellipsoid id is not defined.

Explanation: The ellipsoid id is not defined in the DB2 Geodetic Data Management Feature.

User response: Use a defined ellipsoid and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3746

sqlstate: 38SOG

GSE3748N A ring in a polygon is not closed

Explanation: The start and end points of a ring in a polygon must be exactly equal.

User response: Correct the ring definition by ensuring the start and end points are exactly equal and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3748

sqlstate: 38SOI

GSE3749N An antipodal line segment, *linesegment*, detected in a linestring or a ring of a polygon.

Explanation: This is a line segment where the two points are directly opposite each other as seen from the ellipsoid center. For example, the line definition (0 0, 180 0) may either pass through the north pole or the south pole.

User response: Correct the linestring or ring definition to ensure there are no antipodal line segments by moving one of the two points in this line segment and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3749

sqlstate: 38SP0

GSE3750N The spatial reference system definition does not have a well defined UNIT definition.

Explanation: The spatial reference system definition for this ST_Geometry value is not a geographic coordinate system or does not contain a UNIT definition.

User response: Correct the spatial reference system definition and repeat the command.

For more information, refer to the IBM DB2 Spatial Extender and Geodetic Data Management Feature User's Guide and Reference.

msgcode: -3750

sqlstate: 38SP5

GSE4000N Required parameter *parameter-name* is missing.

Explanation: The required parameter was not found.

User response: Specify the required parameter and try to execute the command again.

msgcode: -4000

sqlstate: 38SB0

GSE4001N An error occurred while Spatial Extender was allocating an environment handle.

Explanation: An environment handle could not be allocated using the Call Level Interface (CLI). The operation cannot be completed successfully.

User response: Verify the CLI configuration. If the source of the problem cannot be found and corrected, contact IBM Software Support.

msgcode: -4001

sqlstate: 38SB1

GSE4002N **An error occurred while Spatial Extender was allocating a connection handle. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial Extender was allocating a connection handle.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4002

sqlstate: 38SB2

GSE4003N **An error occurred while Spatial Extender was connecting to the database. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial Extender was connecting to the database.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4003

sqlstate: 38SB3

GSE4004N **An error occurred while Spatial Extender was allocating a statement handle. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial Extender was allocating a statement handle.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4004

sqlstate: 38SB4

GSE4005N **An error occurred while an SQL statement was being prepared. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial

Extender was preparing an SQL statement.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4005

sqlstate: 38SB5

GSE4006N **An error occurred while Spatial Extender was binding parameters to an SQL statement. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial Extender was binding parameters to an SQL statement.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4006

sqlstate: 38SB6

GSE4007N **An error occurred while Spatial Extender was executing an SQL statement. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial Extender was executing an SQL statement.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4007

sqlstate: 38SB7

GSE4008N **An error occurred while Spatial Extender was ending a transaction. CLI error *cli-error* and native error code = *native-error-code*.**

Explanation: An unexpected error *cli-error* with native error code = *native-error-code* occurred while Spatial Extender was ending a transaction.

User response: Look up the detailed error message *cli-error*. Correct the error and execute the command again. If the problem persists, contact IBM Software Support.

msgcode: -4008

sqlstate: 38SB8

GSE4009N The option, *option*, is invalid.

Explanation: The specified option, *option*, is invalid.

User response: Specify a valid option and repeat the command.

msgcode: -4009

sqlstate: 38SB9

GSE4100N Column *schema-name.table-name.column-name* does not exist.

Explanation: The specified column *schema-name.table-name.column-name* does not exist.

User response: Specify an existing column in an existing table and execute the command again.

msgcode: -4100

sqlstate: 38SC0

GSE4101N Column *schema-name.table-name.column-name* does not have a spatial data type.

Explanation: The specified column *schema-name.table-name.column-name* does not have a spatial type. A spatial type is ST_Geometry or one of its proper subtypes.

User response: Specify a column with a spatial data type and execute the command again.

msgcode: -4101

sqlstate: 38SC1

GSE4102N Spatial grid index *schema-name.index-name* does not exist.

Explanation: The specified spatial grid index *schema-name.index-name* does not exist.

User response: Specify a spatial grid index which does exist and execute the command again.

msgcode: -4102

sqlstate: 38SC2

GSE4103N An internal error occurred while Spatial Extender was retrieving index information for column *schema-name.table-name.column-name*.

Explanation: Spatial Extender encountered an unexpected internal error while retrieving information about the spatial index on the column *schema-name.table-name.table-name*.

User response: Contact IBM Software Support.

msgcode: -4103

sqlstate: 38SC3

GSE4104N Spatial Extender could not retrieve grid sizes for spatial grid index *schema-name.index-name*.

Explanation: Spatial Extender encountered an unexpected internal error while retrieving the grid size definition for the grid index *schema-name.index-name*.

User response: Contact IBM Software Support.

msgcode: -4104

sqlstate: 38SC4

GSE4105W No rows to analyze in column *schema-name.table-name.column-name*.

Explanation: The column *schema-name.table-name.table-name* does not contain any values that the Spatial Extender could analyze.

If sampling was used, the sample rate might have been too low.

User response: Execute the command on a column that contains geometries. If sampling is used, use a higher sampling rate.

msgcode: +4105

sqlstate: 38SC5

GSE4106W Spatial Extender could not get enough memory to retrieve MBR histogram for *number* entries.

Explanation: The histogram size of the minimum boundary rectangles for all the geometries in the spatial column is too big. Not enough memory to store it could be allocated. The command cannot be executed successfully.

User response: Do not specify the SHOW DETAIL option for the command or reduce the number of geometries to be analyzed by sampling them with a lower sample rate.

msgcode: +4106

sqlstate: 38SC6

GSE4107N Grid size value *grid-size* is not valid where it is used.

Explanation: The specified grid size *grid-size* is not valid.

One of the following invalid specifications was made when the grid index was created with the CREATE INDEX statement:

- A number less than 0 (zero) was specified as the grid size for the first, second, or third grid level.
- 0 (zero) was specified as the grid size for the first grid level.

- The grid size specified for the second grid level is less than the grid size of the first grid level but it is not 0 (zero).
- The grid size specified for the third grid level is less than the grid size of the second grid level but it is not 0 (zero).
- The grid size specified for the third grid level is greater than 0 (zero) but the grid size specified for the second grid level is 0 (zero).

User response: Specify a valid value for the grid size.

msgcode: -4107

sqlstate: 38SC7

GSE4108W Geometries in column
schema-name.table-name.column-name do not cover extent.

Explanation: All geometries in the column *schema-name.table-name.column-name* do not span an extent greater than 0 (zero). That means that only identical points exist in the column. The index information cannot be gathered.

User response: Insert additional rows into the spatial column.

msgcode: +4108

sqlstate: 38SC8

GSE4109N An invalid query box of size
query-box-size was specified.

Explanation: The specified query box size *query-box-size* is not valid. A query box must be greater than 0 (zero) and less than or equal to 1 (one). The query box size indicates the percentage of the extent for the data in the column that is usually retrieved.

User response: Correct the query box size by choosing a value greater than 0 (zero) and less than or equal to 1 (one).

msgcode: -4109

sqlstate: 38SC9

GSE4110N An invalid number of rows *num-rows* to analyze was specified.

Explanation: The specified number of rows to be analyzed (*num-rows*) is invalid. The number must be greater than 0 (zero).

User response: Specify a valid number of rows greater than 0 (zero) and execute the command again.

msgcode: -4110

sqlstate: 38SCA

GSE4111N An invalid percentage *percentage* was specified.

Explanation: The specified percentage *percentage* of rows to be analyzed is not valid. The percentage must be an integer that is greater than 0 (zero) and less than or equal to 100.

User response: Specify a valid percentage greater than 0 (zero) and less than or equal to 100 and retry the command.

msgcode: -4111

sqlstate: 38SCB

GSE4112W No samples for the table
schema-name.table-name.column-name at percent percent could be produced.

Explanation: The Spatial Extender tried to sample percent percent of table *schema-name.table-name.column-name*, but could not produce any sample rows that could be analyzed any further.

User response: If the table is not empty, choose a higher sampling rate and retry the command.

msgcode: +4112

sqlstate: 38SCC

GSE4113N No tablespace for a declared global temporary table for the sampled data exists.

Explanation: A USER TEMPORARY tablespace with a page size of at least 4096 bytes must exist to declare a global temporary table that is needed to sample the data that is to be analyzed by the Spatial Extender. A global temporary table is needed so that Spatial Extender can sample the data that it is going to analyze. A USER TEMPORARY tablespace with a page size of at least 4096 bytes is required for declaring this table.

User response: Either create an appropriate USER TEMPORARY tablespace and retry the command, or do not use the ANALYZE clause to avoid sampling the spatial data.

msgcode: -4113

sqlstate: 38SCD

GSE4200N Parameter *value* of data type *type* is not expected in that context.

Explanation: The parameter value *value* of type *type* is not expected in the context of the command.

User response: Correct the command to be executed and then try the operation again.

msgcode: -4200

sqlstate: 38SD0

GSE4201N No command was specified.

Explanation: No command to be executed was specified.

User response: Specify a command.

msgcode: -4201

sqlstate: 38SD1

GSE4202N Parsing of command failed.

Explanation: The given command syntax could not be parsed successfully.

User response: Correct the syntax and execute the command again.

msgcode: -4202

sqlstate: 38SD2

GSE4203N Unexpected end-of-statement found following token *token*.

Explanation: An unexpected end of the statement was found after parsing token *token*.

User response: Correct the syntax and execute the command again.

msgcode: -4203

sqlstate: 38SD3

GSE4204N Parsing failed near token *token*.

Explanation: An unexpected parse error occurred near token *token*.

User response: Correct the syntax and execute the command again.

msgcode: -4204

sqlstate: 38SD4

GSE4205N Spatial Extender found more than one *type* clause.

Explanation: While Spatial Extender parsed the command, it found more than one clause of data type *type*.

User response: Specify only one clause of *type* and execute the command again.

msgcode: -4205

sqlstate: 38SD5

GSE9000N The database is not at the Spatial Extender Version 7 level.

Explanation: The database is not at the Spatial Extender Version 7 level and you are attempting to perform an operation such as migration to Version 8.

User response: Verify the version of the database and the operation you are attempting to perform.

msgcode: -9000

sqlstate: 38SZ0

GSE9001N The database is already at the Spatial Extender Version 8 level.

Explanation: The database is already at the Spatial Extender Version 8 level and you are attempting to perform an operation such as migration to Version 8.

User response: Verify the version of the database and the operation you are attempting to perform.

msgcode: -9001

sqlstate: 38SZ1

GSE9002N An error occurred during the upgrade of Spatial Extender database objects.

Explanation: An error was encountered during an attempt to upgrade the Spatial Extender catalog and spatial functions. Possible reasons are:

- The database is not spatially-enabled
- The version of the database that you are trying to upgrade is not supported by the db2se upgrade command.
- An internal processing error occurred during the upgrade process.

User response: Perform one of the following actions:

- Use the db2se enable_db command to use this database for spatial processing.
- Verify that the version of the database that you are trying to upgrade is supported by the db2se upgrade command and reissue this command.
- Contact the DB2 support team to report internal processing errors.

For more information, refer to the db2se upgrade message file.

msgcode: -9002

sqlstate: 38SZ2

GSE9003N **A dependency between user-defined objects and Spatial Extender objects exists.**

Explanation: The Spatial Extender objects could not be upgraded because some user-defined objects have dependencies on Spatial Extender objects.

User response: Review the db2se upgrade message file to determine what dependencies exist.

Issue the db2se upgrade command with the force option to save and restore the user-defined objects that depend on the Spatial Extender objects.

msgcode: -9003

sqlstate: 38SZ3

GSE9990C **An internal error occurred:** *error-text*.

Explanation: Spatial Extender encountered an unexpected internal error with the text *error-text*.

User response: Read the given *error-text*. If the problem cannot be resolved, contact IBM Software Support.

msgcode: -9990

sqlstate: 38SZY

GSE9999C **Internal message failure.**

Explanation: An internal failure occurred while Spatial Extender was retrieving an error message.

User response: Contact IBM Software Support.

msgcode: -9999

sqlstate: 38SZZ

Part 17. ICM Messages

Chapter 102. ICM0000 - ICM0499

ICM00001N An SQL error occurred while connecting to the database. Database = *databasename*, Catalog = *catalogname*.

Explanation: An SQL error occurred when the API attempted to open a connection to the database.

User response: Consult the Message Reference for more information on this SQL error.

ICM0002N An SQL error occurred while registering application *applicationname*.

Explanation: An SQL error occurred while the Information Catalog Manager API attempted to register the specified application.

User response: Check the Message Reference for more information on this SQL error.

ICM0003N An SQL error occurred while searching for objects.

Explanation: An SQL error occurred while the Information Catalog Manager API was performing a search of the IBM DB2 tools catalog.

User response: Check the Message Reference for more information about this SQL error.

ICM0004N An SQL error occurred while accessing the Information Catalog Manager API settings for the application *applicationname*.

Explanation: An SQL error occurred while the Information Catalog Manager API attempted to retrieve the settings for the application.

User response: Check the Message Reference for more information about this SQL error.

ICM0005N An SQL error occurred while retrieving the user and group information from the database.

Explanation: An SQL error occurred while the Information Catalog Manager API attempted to retrieve information about users and groups from the database.

User response: Check the Message Reference for more information about this SQL error.

ICM0006N An SQL error occurred while loading the list of registered applications.

Explanation: An SQL error occurred while the

Information Catalog Manager API attempted to retrieve the list of applications.

User response: Check the Message Reference for more information about this SQL error.

ICM0007N An SQL error occurred while retrieving the application ID for the application *applicationname*.

Explanation: An SQL error occurred while the Information Catalog Manager API attempted to retrieve the application ID for the specified application.

User response: Check the Message Reference for more information on this SQL error.

ICM0008N An SQL error occurred while loading an object type.

Explanation: An SQL error occurred while the Information Catalog Manager API attempted to retrieve an object type definition from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0009N An SQL error occurred while creating object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was creating the specified object type.

User response: Check the Message Reference for more information on this SQL error.

ICM0010N An SQL error occurred while updating the object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was updating the specified object type.

User response: Check the Message Reference for more information on this SQL error.

ICM0011N An SQL error occurred while deleting object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was deleting the specified object type.

User response: Check the Message Reference for more information on this SQL error.

ICM0012N An SQL error occurred while retrieving a reference ID for object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was retrieving the next available reference ID for the specified object type.

User response: Check the Message Reference for more information on this SQL error.

ICM0013N An SQL error occurred while loading an object instance from the IBM DB2 tools catalog.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading the data for an object instance from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0014N An SQL error occurred while creating instance *instancename* of object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was creating an object instance in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0015N An SQL error occurred while updating an instance of object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was updating an object instance in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0016N An SQL error occurred while deleting an instance of object type *objecttype*.

Explanation: An SQL error occurred while the Information Catalog Manager API was deleting an object instance from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0017N An SQL error occurred while loading property *propertyname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading the value for the specified object instance property.

User response: Check the Message Reference for more information on this SQL error.

ICM0018N An SQL error occurred while loading a relationship category.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading a relationship category from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0019N An SQL error occurred while creating relationship category *categoryname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was creating the relationship category in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0020N An SQL error occurred while updating relationship category *categoryname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was updating the relationship category in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0021N An SQL error occurred while deleting relationship category *categoryname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was deleting the relationship category from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0022N An SQL error occurred while loading a relationship type.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading a relationship type from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0023N An SQL error occurred while creating relationship type *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was creating the relationship type in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0024N An SQL error occurred while updating relationship type *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was updating the relationship type in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0025N An SQL error occurred while deleting relationship type *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was deleting the relationship type from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0026N An SQL error occurred while loading a relationship instance.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading the information for a relationship instance from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0027N An SQL error occurred while creating an instance for relationship type *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was creating an instance of the relationship type in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0028N An SQL error occurred while updating an instance of relationship type *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was updating an instance of the relationship type in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0029N An SQL error occurred while deleting an instance of relationship type *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was deleting an instance of the relationship type from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0030N An SQL error occurred while committing changes on the context.

Explanation: An SQL error occurred while the Information Catalog Manager API was attempting to commit a context to the database.

User response: Check the Message Reference for more information on this SQL error.

ICM0031N An SQL error occurred while rolling back changes on the context.

Explanation: An SQL error occurred while the Information Catalog Manager API attempted to rollback a context to the database.

User response: Check the Message Reference for more information on this SQL error.

ICM0032N An SQL error occurred while releasing the context.

Explanation: An SQL error occurred while the Information Catalog Manager API was releasing a context.

User response: Check the Message Reference for more information on this SQL error.

ICM0033N An SQL error occurred while accessing the context's connection properties.

Explanation: An SQL error occurred while the Information Catalog Manager API was accessing the database connection details for a context.

User response: Check the Message Reference for more information on this SQL error.

ICM0034N An SQL error occurred while cancelling a search operation.

Explanation: An SQL error occurred while the Information Catalog Manager API was processing a user request to cancel a search operation

User response: Check the Message Reference for more information on this SQL error.

ICM0035N An SQL error occurred while loading an access control list.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading information in an access control list from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0036N **An SQL error occurred while updating access control list** *accesslistname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was updating the access control list in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0037N **An SQL error occurred while accessing the data of Blob property** *propertyname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was loading the data for the specified Blob property.

User response: Check the Message Reference for more information on this SQL error.

ICM0038N **An SQL error occurred while retrieving the version information of the IBM DB2 tools catalog.**

Explanation: An SQL error occurred while the Information Catalog Manager API queried the database for the version information of the database engine and the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0039N **An SQL error occurred while retrieving the access control entries for an object.**

Explanation: An SQL error occurred while the Information Catalog Manager API was loading the access control data for an object.

User response: Check the Message Reference for more information on this SQL error.

ICM0040N **An SQL error occurred while creating access control list** *accesslistname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was creating the access control list in the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0041N **An SQL error occurred while deleting access control list** *accesslistname*.

Explanation: An SQL error occurred while the Information Catalog Manager API was deleting the access control list from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0042N **An SQL error occurred while locking object type** *typename*.

Explanation: An SQL error occurred while the Information Catalog Manager API was locking the object type to prevent access by other applications.

User response: Check the Message Reference for more information on this SQL error.

ICM0043N **An SQL error occurred while loading an object ID.**

Explanation: An SQL error occurred while the Information Catalog Manager API was loading the ID for an object instance from the IBM DB2 tools catalog.

User response: Check the Message Reference for more information on this SQL error.

ICM0100N **Unable to connect to the IBM DB2 tools catalog. Database =** *databasename*, **Catalog =** *catalogname*).

Explanation: Incorrect values were provided for the database or catalog names.

User response: Provide valid values for the database and catalog. The database name must refer to a database that has been cataloged on the system where the Information Catalog Manager API is being run. The catalog name must refer to the schema inside the database used to hold the DB2 tools catalog.

ICM0101N **This catalog object is already connected to an IBM DB2 tools catalog.**

Explanation: Failed to open a connection to a DB2 tools catalog using a catalog object that already had an open connection.

User response: Check that the catalog object is not already open.

ICM0102N **Schema** *schemaname* **in database** *databasename* **does not contain a valid IBM DB2 tools catalog. Table** *tablename* **was not found.**

Explanation: Unable to find the required table *tablename* while the catalog opened a connection to the IBM DB2 tools catalog,

User response: Check that the schema and database names are correct, and that a DB2 tools catalog has been created for schema *schemaname*.

ICM0103N **Incorrect arguments provided to orphan relationship search.**

Explanation: The emptySource and emptyTarget parameters to the orphanRelationships(Context, Object, Collection, Collection, boolean, bJboolean, SearchRules)

method of the Catalog object were both false, which is incorrect.

User response: One or both of the parameters emptySource and emptyTarget must be true.

ICM0104N *applicationname* is a reserved application name.

Explanation: The application name *applicationname* is reserved for use by the Information Catalog Manager API.

User response: Use a different name for the application name.

ICM0105N Application *applicationname* is already registered.

Explanation: An application named *applicationname* has been previously registered with the Information Catalog Manager API.

User response: Register the application using a different name or use the existing registered application ID.

ICM0106N Application *applicationname* has not been registered.

Explanation: An attempt was made to use an unregistered application name for an application.

User response: Register the application name before using it or use the name of an application that has already been registered.

ICM0107N An application must be associated with the catalog object before the IBM DB2 tools catalog can be accessed.

Explanation: Attempted to access the IBM DB2 tools catalog before associating an application with the catalog object.

User response: Associate an application with the catalog object using the setApplicationName(String) method of the catalog class before using the catalog object to access the IBM DB2 tools catalog.

ICM0108N Unable to close catalog while contexts remain unreleased.

Explanation: One or more contexts were in an unreleased state while the close method of a catalog object was called.

User response: Release all contexts opened against a catalog before closing the catalog.

ICM0109N Incorrect constraint object encountered while saving constraints.

Explanation: An I/O error occurred while preparing a constraint to be saved to the database.

User response: Check that the constraint object is Serializable.

ICM0110N Incorrect constraint object encountered while loading constraints.

Explanation: An I/O error occurred while recreating a constraint from the data stored in the IBM DB2 tools catalog.

User response: Contact you IBM service representative.

ICM0111N Cannot register application *applicationname*. No IDs are available.

Explanation: There are 63 available application IDs. Attempted to register a 64th application.

User response: Remove any registered applications that are no longer being used with this IBM DB2 tools catalog and contact your IBM service representative.

ICM0112N Incorrect arguments provided to getACEsForPrincipal(Context, String, boolean, Collection) method.

Explanation: The value for the principal or the collection of object types to retrieve access control entries was null.

User response: Specify non-null values for the principal and the collection of object types.

ICM0113N Incorrect arguments provided to getObjectsOwnedByUser(Context, String, Collection) method.

Explanation: The value was null for the user or the Collection of object types to retrieve the owned objects.

User response: Specify non-null values for both the user and the Collection of object types.

ICM0200N The setProperties(Collection) method of an object type cannot be used after the object type has been created.

Explanation: Attempted to use the setProperties(Collection) method on an object type that has already been created in the IBM DB2 tools catalog.

User response: Use the addProperty(MetadataPropertyDefinition) method to add new properties individually or create a new object type.

ICM0201N **System property *propertyname* must be specified when setting the properties of this object type.**

Explanation: The system property *propertyname* was not in the list of properties provided to the `setProperty(Collection)` method.

User response: The list of properties provided to the `setProperty(Collection)` method must include all of the system properties.

ICM0202N **Property *propertyname* is in use by index *indexname* of object type *typename*. The property must be specified when setting the list of properties for the object type.**

Explanation: The property *propertyname* must remain in the set of properties for the object type because the property is being used by an index.

User response: Remove the property from the specified index, remove the index, or add the property to the set of properties for the object type.

ICM0203N **Object *object* is not a metadata property object.**

Explanation: An incorrect object was found in the set of properties being assigned to an object type.

User response: The set of properties provided to `setProperty(Collection)` must contain only `MetadataPropertyDefinition` objects. Remove any objects that are of a different type.

ICM0204N **Duplicate property named *propertyname* was found in list.**

Explanation: A call to `setProperty(Collection)` had multiple property objects with the same name.

User response: Remove or rename any properties with duplicate names.

ICM0205N **No property definition named *propertyname* exists for this object type.**

Explanation: A call to `getProperty(String)` was made using a property that is not defined for the object type.

User response: Use the `getProperties()` method to retrieve all properties if the name of a specific property is not known.

ICM0206N **An object type named *typename* already exists in the datastore.**

Explanation: Attempted to create or rename an object type with an existing name.

User response: Chose a unique name for the new object type.

ICM0207N **Unable to delete object type *typename* while it has object instances defined.**

Explanation: Attempted to delete an object type with instances defined without selecting the delete instances option.

User response: Delete all of the instances before deleting an object type or use the delete instances option of the `delete(boolean)` method of the object type.

ICM0208N **Cannot use the `setTableOptions(String)` method of an object type after the object type has been created.**

Explanation: Attempted to modify the table options clause after the object type was created. The table options clause can only be used when creating the table.

User response: Set the desired table options before creating the object type.

ICM0209N **The `setSchemaName(String)` method of an object type cannot be used after the object type has been created.**

Explanation: Attempted to modify the schema for the object instance table after the object type has been created.

User response: Set the desired schema before creating the object type.

ICM0210N **Property *propertyname* specified for index *indexname* is not a property of object type *typename*.**

Explanation: An index provided to the object type contained one or more properties that do not belong to the object type.

User response: Check that all properties in an index are the correct property definition objects defined for an object type. Use the `getProperties()` and `getProperty(String)` methods to retrieve the correct property definition objects.

ICM0211N **The type of object instance *instancename* does not match object type (*typename*).**

Explanation: One of the object instances provided did not match the `createObjects(Context, Collection)` method object type.

User response: The `createObjects(Context, Collection)` method can only handle object instances that are of the same object type.

ICM0212N String *searchcriteria* is an invalid search criteria for the name property.

Explanation: Failed to search object instances by name because the name parameter provided was not a value.

User response: The name parameter cannot be null.

ICM0213N An incorrect query filter object was provided to the getObjectInstances(Context, Object, QueryFilterObject, SearchRules, boolean) method of an object type.

Explanation: The query filter object was either null or was for a different object type than the one it was passed to.

User response: Define a query filter object for the object type that is being searched.

ICM0214N The list of object types is not valid.

Explanation: A list of object types provided to one of the Information Catalog Manager API methods was not valid. The object type list was either null or empty.

User response: Check that the list of object types contains at least one object type.

ICM0215N The data type specified while creating a property definition object was not valid.

Explanation: The value specified did not identify one of the allowed data types.

User response: Check that the value for the data type is a data type constant defined in the MetadataPropertyDefinition class.

ICM0216N Index *indexname* of object type *typename* is not valid.

Explanation: An index provided to the object type *typename* did not have any properties defined.

User response: The index must have at least one property specified.

ICM0217N Property *propertypname* is defined multiple times in index *indexname* of object type *typename*.

Explanation: An index provided to the object type *typename* has the same property defined multiple times.

User response: Remove all duplicate properties from the index.

ICM0218N Index *indexname* does not exist for object type *typename*.

Explanation: Object type *typename* could not grant a request to remove index *indexname* because an index with that name does not exist.

User response: Check that the index is defined for the object type before attempting to remove it.

ICM0219N An index with the name *indexname* is already defined for object type *typename*.

Explanation: An index with the specified name has already exists. Index names must be unique.

User response: Chose a unique name for the new index, or remove the existing one to replace it with the new one.

ICM0220N Index *duplicatename* duplicates the definition of index *indexname* of object type *typename*.

Explanation: An index already exists with the requested definition for the object type.

User response: Either modify the new index to make it different than all existing indexes, or just use the existing indexes if they provide sufficient indexing for the object type.

ICM0300N The requested object instance was not found in the datastore.

Explanation: Attempted to load the value of a property for an object instance that has not been created in the datastore.

User response: Contact your IBM service representative.

ICM0301N A property update request violates property constraint *constraintname*.

Explanation: The operation cannot be performed because constraint *constraintname* does not allow it.

User response: Choose a value allowed by the constraint or remove the constraint from the property definition.

ICM0302N The object type has already been set for this instance.

Explanation: Attempted to set the object type for an instance that already had a type defined.

User response: Instances that already have a type defined cannot have their type changed.

ICM0303N The specified object type is not valid.

Explanation: The object type provided to the object instance setType(ObjectType) method is not valid. The object type is either null or has not yet been created.

User response: Check that the object type passed to the setType(ObjectType) method has been created.

ICM0304N No property named *propertyname* exists for this object instance.

Explanation: Attempted to retrieve a property that does not exist.

User response: Check the spelling of the property name and that the case matches. Use the getProperties() method to get a list of all defined properties.

ICM0305N The object instance is not in a valid state.

Explanation: The object instance is not able to handle requests for data or updates because it is not valid.

User response: Reload the object instance from the IBM DB2 tools catalog. If it is a new instance check that an object type is set.

ICM0306N A name is required for the object instance.

Explanation: The object instance is not valid because the name has not been set.

User response: Provide a non-null value for the name property of the object instance.

ICM0307N Property *propertyname* is not valid.

Explanation: The object instance is not valid because the property value is null when the property is required or the property value violates one of the property constraints.

User response: Provide a valid value for the specified property.

ICM0308N The current application does not have permission to update this object.

Explanation: The current application is not the owning application and does not have application update authority on the specified object.

User response: Use the owning application to grant authority to the current application.

ICM0309N Property *propertyname* does not belong to this object instance.

Explanation: Attempted to load a property value to an incorrect object instance.

User response: Contact your IBM service representative.

ICM0310N The value for property *propertyname* exceeds the property definition's maximum length.

Explanation: The property value is longer than the maximum size allowed by the object type's property definition.

User response: Use a smaller property value.

ICM0311N Property *propertyname* must have a value.

Explanation: The property *propertyname* is required but is set with a null value.

User response: Set the property with a non-null value.

ICM0312N Property *propertyname* violates constraint *constraintname*.

Explanation: The value assigned to property *propertyname* is not allowed by constraint *constraintname*.

User response: Use a value that is allowed by the constraint or remove the constraint from the property definition.

ICM0400N The list of relationship objects is not valid.

Explanation: The list of relationship objects is null or empty.

User response: Check that the list of relationship objects contains at least one relationship.

ICM0401N Adding the specified relationship would result in a loop in the relationship category tree.

Explanation: The requested change would result in a relationship category containing itself.

User response: Check that there are no loops in the hierarchy of relationships.

ICM0402N The relationship category *categoryname* already exists.

Explanation: Failed to create a new relationship category because the category had the same name as an existing category.

User response: Use a unique name for the new relationship category.

ICM0403N The relationship type *typename* already exists.

Explanation: Failed to create a new relationship type because the type had the same name as an existing type. Failed to create a new relationship type because the type had the same name as an existing type.

User response: Use a unique name for the new relationship type.

ICM0404N Unable to delete relationship type *typename* because there are relationship instances defined for it.

Explanation: Relationship types cannot be deleted if relationship instances of that type exist.

User response: Delete all instances of the relationship type before deleting the type or use the delete instances option of the relationship type delete method.

ICM0405N The ObjectType being added to the ObjectTypeRelationshipConstraint has not been created.

Explanation: The relationship constraint attempted to use an object type that has not yet been created in the database.

User response: Call the create() method of the object type before using it in a relationship constraint.

Chapter 103. ICM0500 - ICM0999

ICM0500N **The object passed to the relationship instance is not valid. All objects must be ObjectInstance or ObjectID objects.**

Explanation: The type of an object provided as a source or target to a relationship instance was not one of the allowed types.

User response: Use only ObjectInstance or ObjectID objects as sources and targets for relationship instances.

ICM0501N **An object instance must be created in the database before it can be added to a relationship.**

Explanation: An object instance that has not yet been created in the database cannot be assigned to any relationships.

User response: Call the create() method on new object instances before attempting to add them to a relationship.

ICM0502N **Unable to add a null ObjectID to relationship.**

Explanation: The Java object representing the ObjectID that was provided to the relationship object was null.

User response: Use a non-null value for the ObjectID.

ICM0503N **The relationship instance violates one or more constraints.**

Explanation: One or more of the relationship instance constraints were violated by the current set of source and target objects defined for the instance.

User response: Check the defined constraints and add or remove sources or targets as needed.

ICM0504N **The relationship instance is not in a valid state.**

Explanation: The relationship instance is in an unusable state.

User response: Check that the relationship type for the instance exists and that none of the constraints are violated. Then reload the relationship instance from the database.

ICM0600N **The current user does not have permission to update this object.**

Explanation: Failed to update an object because the user does not have sufficient authority.

User response: The object's owner must grant permission to users before they can update the object.

ICM0601N **The current user does not have permission to delete this object.**

Explanation: Failed to delete an object because the user does not have sufficient authority.

User response: The object owner must grant users permission to delete the object.

ICM0602N **The current user does not have permission to delete one of the instances of this relationship type.**

Explanation: The current user is unable to delete a relationship type and its instances because the user does not have authority to delete one of the instances.

User response: User needs the authority to delete all the instances of a relationship type in order to delete the relationship type.

ICM0603N **A name must be specified for the access control list before it can be saved.**

Explanation: Failed to update an access control list because the name has been set to null.

User response: Check that the access control list has a valid name.

ICM0604N **An access control list named *accesslistname* already exists in the datastore.**

Explanation: The new access control list cannot be created because it has the same name as an existing access control list.

User response: Specify a unique name for the new access control list.

ICM0605N **An access control list named *accesslistname* is not valid.**

Explanation: The access control list cannot be saved because it is no longer valid.

User response: Reload the access control list from the IBM DB2 tools catalog.

ICM0700N **The data for this object has been modified in the IBM DB2 tools catalog by a different process.**

Explanation: Failed to update an object because a different process modified the object data stored in the tools catalog while the object was in use by the current process.

User response: Reload the object from the tools catalog or use the overwrite setting on the update method.

ICM0701N Object does not exist.

Explanation: Information Catalog Manager API request failed because it attempted to use an object that has not been created in the IBM DB2 tools catalog.

User response: Call the create method of the object before using the method.

ICM0702N Object already exists.

Explanation: The create method was called on an object that has already been created.

User response: Use the update method rather than the create method to update objects that already exist.

ICM0703N Object is not in a valid state.

Explanation: The object cannot be created or updated because it is not in an valid state.

User response: Reload the object from the IBM DB2 tools catalog.

ICM0704N Duplicate request ID *idname* passed to search method.

Explanation: The request ID provided to a search method is already in use by a different search.

User response: Use a different request ID or wait for the already running search to finish.

ICM0705N Cannot switch to a null Context.

Explanation: The context for an Information Control Center object cannot be set to null.

User response: Specify a non-null value to the setContext(Context) method of the Information Control Center object.

ICM00706N Cannot switch to a new context while a transaction is pending on the existing context.

Explanation: Attempted to switch the context of an object while there was an active transaction involving the object.

User response: Commit the current context changes to the object before switching to a different context.

ICM0707N Cannot switch to a context in a different catalog.

Explanation: Attempted to switch the context of an object to a context opened against a different catalog.

User response: Use only contexts and objects from the same catalog.

ICM0708N An access control entry already exists for specified principal (*principalname*).

Explanation: A new access control entry cannot be created for the specified principal because one already exists.

User response: Retrieve the current access control entry and modify it. Do not create a new access control entry.

ICM0709N The context provided to access the datastore is not valid.

Explanation: The context provided to the Information Catalog Manager API is not valid. Cannot perform the requested operation.

User response: Check that the context is not null and has not been released.

ICM0710N An incorrect parameter type was encountered in the parameterized SQL statement: *sqlstatement*.

Explanation: The type of an object parameter that was provided is not valid.

User response: Use only objects of the types specified in the documentation for the Information Catalog Manager API method.

ICM0711N Permission value for the PermissionElement is not valid.

Explanation: The permission is not valid with the current value in setPermission(int, Jboolean).

User response: Set permissions using the permission element constants defined in the PermissionElement class.

ICM00712N Incorrect number of values were provided to the query filter condition. No parameters were expected.

Explanation: The number of parameters provided for a query filter condition does not match the number expected for the operator of the condition.

User response: Provide the correct number of parameters. No parameters are allowed for OP_IS_NULL and OP_IS_NOT_NULL. Two parameters are required for OP_BETWEEN and

OP_NOT_BETWEEN. Any number of parameters greater than zero are allowed for OP_IN and OP_NOT_IN. One parameter is required for all other operators.

ICM0713N **Number of values provided to the query filter condition was not valid. One parameter was expected.**

Explanation: The number of parameters provided for a query filter condition does not match the number expected for the operator of the condition.

User response: Provide the correct number of parameters. No parameters are allowed for OP_IS_NULL and OP_IS_NOT_NULL. Two parameters are required for OP_BETWEEN and OP_NOT_BETWEEN. Any number of parameters greater than zero are allowed for OP_IN and OP_NOT_IN. One parameter is required for all other operators.

ICM0714N **Number of parameters provided to the query filter condition was not valid. Two parameters were expected.**

Explanation: The number of parameters provided for a query filter condition does not match the number expected for the operator of the condition.

User response: Provide the correct number of parameters. No parameters are allowed for OP_IS_NULL and OP_IS_NOT_NULL. Two parameters are required for OP_BETWEEN and OP_NOT_BETWEEN. Any number of parameters greater than zero are allowed for OP_IN and OP_NOT_IN. One parameter is required for all other operators.

ICM0715N **Type specified for relationship search is not valid.**

Explanation: The value of the type specified to the setType(int) method of a relationship search is not valid.

User response: Use the TYPE_SOURCE and TYPE_TARGET constants defined in the RelationshipSearch class for the setType(int) method.

ICM0716N **Relationship specified for relationship search is not valid.**

Explanation: The relationship specified for a relationship search object was null.

User response: Use a non-null value for the relationship of a relationship search.

ICM0717N **The access control list for this object has not been created.**

Explanation: Must create the access control list assigned to an object before the object can be saved to the IBM DB2 tools catalog.

User response: Call the create method on the access control list.

ICM0718N **The context has been released and cannot be used.**

Explanation: Attempted to use a released context to access the IBM DB2 tools catalog.

User response: Use the newContext() method on the catalog object to retrieve a new context and use it in place of the released one.

ICM0800N **Specified a migration operation type that was not valid.**

Explanation: The value of the operation type specified to the ICMMigration migrate(int) method was not a valid migration operation type.

User response: Use one of the migration operation types defined in the ICMDatastoreInit class (MIG_REPLACE, MIG_SKIP or MIG_ERROR).

ICM0801N **The definition of the property *propertyname* of the IBM DB2 Version 7 object type *typename* being migrated is different than the existing property definition for the DB2 Version 8 object type.**

Explanation: An object type being migrated already exists in the DB2 Version 8 tools catalog. One of the properties of the object type differs in either data type or size between the DB2 Version 7 and Version 8 object type definitions.

User response: Use the replace migration operation type or delete the existing DB2 Version 8 object type before migrating the DB2 Version 7 catalog.

ICM0802N **The property *propertyname* of the IBM DB2 Version 8 object type *typename* does not exist in the DB2 Version 7 object type of the same name being migrated.**

Explanation: An object type being migrated already exists in the DB2 Version 8 tools catalog, but one of the properties of the existing DB2 Version 8 object type does not exist in the DB2 Version 7 object type definition.

User response: Run the migration using a different operation type than MIG_ERROR or delete the existing DB2 Version 8 object type before running the migration.

ICM0803N **The property *propertyname* for the IBM DB2 Version 7 object type *typename* does not exist in the DB2 Version 8 object type of the same name.**

Explanation: The object type being migrated exists in the DB2 Version 8 tools catalog, but one of the object type's properties being migrated from DB2 Version 7 does not exist in the DB2 Version 8 object type definition.

User response: If the missing property is optional, use the Information Catalog Center initialization tool to add the missing property to the DB2 Version 8 object type before migrating. If the missing property is required, remove the DB2 Version 8 object type from the tools catalog before migrating.

ICM0804N **The value of the property *propertyname* for instance *instancename* of object type *typename* is not in a valid date/time format.**

Explanation: The value of a date/time property in a DB2 Version 7 catalog is not in a recognized date/time format.

User response: Manually change the IBM DB2 Version 7 data to a valid date/time format for the current location.

ICM0805N **The data type of the *propertyname* property for instance *instancename* of object type *objecttype* is not supported by the migration tool.**

Explanation: The migration tool cannot migrate object types with properties that have data types other than character or date/time strings.

User response: The object type cannot be migrated. Check that the catalog has not been corrupted.

ICM0806N **Cannot recognize the relationship type flag *typename* in the IBM DB2 Version 7 catalog being migrated.**

Explanation: The value of the relationship type flag is not a recognized value.

User response: The IBM DB2 Version 7 catalog is not valid. The relationship instance with this type flag will not be migrated.

ICM0807N **Could not find relationship type *typename*.**

Explanation: The predefined relationship type *typename* was not found in the IBM DB2 tools catalog.

User response: Contact your IBM service representative.

ICM0808N **Object instance *instancename* violates a unique index on object type *typename* and cannot be created.**

Explanation: The instance could not be created because it violates one of the unique indexes defined for the object type.

User response: Remove the index or delete the IBM DB2 Version 7 object instance that causes the index violation before running the migration tool.

ICM0900N **An I/O error occurred while opening the log file *filename*.**

Explanation: The initialization tool was unable to open the log file because of an I/O exception.

User response: Check that the specified file path exists and that the file is not read-only or being used by another application.

ICM0901N **An SQL error occurred while creating the Information Catalog Manager API tables.**

Explanation: An SQL error occurred while the initialization tool was creating the tables required by the Information Catalog Manager API.

User response: Check the Message Reference for more information on this SQL error.

ICM0902N **An SQL error occurred during migration.**

Explanation: An SQL error occurred while the migration tool was retrieving the IBM DB2 Version 7 data being migrated.

User response: Check the Message Reference for more information on this SQL error.

ICM0903N **An SQL error occurred while creating the DB2 Version 7 Information Catalog emulation views.**

Explanation: An SQL error occurred while the initialization tool was creating the views that provide emulation of the DB2 version 7 Information Catalog Manager tables.

User response: Check the Message Reference for more information on this SQL error.

ICM0904N **Table *tablename* is missing from the IBM DB2 tools catalog.**

Explanation: The check option for the Information Catalog Manager API initialization found that table *tablename* was not defined.

User response: Run the Information Catalog Manager

API initialization with the fix option to repair this error.

ICM0905N Table *tablename* contains at least one column with a definition that is not valid.

Explanation: The check option for the Information Catalog Manager API initialization found that table *tablename* contains at least one column that has a different type than that which is required.

User response: Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0906N Table *tablename* contains at least one column that is not part of the datastore definition.

Explanation: The check option for the Information Catalog Manager API initialization found that table *tablename* contains at least one column that does not exist in the table definition.

User response: If nulls are allowed for this column then it will not affect the Information Catalog Manager API. To remove it, run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0907N Column *columnname* of table *tablename* is missing from the datastore.

Explanation: The check option for the Information Catalog Manager API initialization found that table *tablename* is missing at least one column that is required by the Information Catalog Manager API.

User response: Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0908N Table *tablename* is missing its primary key definition.

Explanation: The check option for the Information Catalog Manager API initialization found that table *tablename* is missing the required primary key definition.

User response: Run the Information Catalog Manager API initialization with the fix option to repair this error.

ICM0909N Table *tablename* is missing its primary key definition.

Explanation: The check option for the Information Catalog Manager API initialization found that the primary key for table *tablename* does not match the

primary key that is required by the Information Catalog Manager API.

User response: Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0910N Table *tablename* is missing an index definition.

Explanation: The check option for the Information Catalog Manager API found that table *tablename* is missing a required index.

User response: Run the Information Catalog Manager API initialization with the fix option to repair this error.

ICM0911N Table *tablename* has an invalid index defined.

Explanation: The check option for the Information Catalog Manager API initialization discovered that table *tablename* has an index defined which does not match the indexes required by the Information Catalog Manager API.

User response: If the index is a unique index, it may interfere with the proper functioning of the Information Catalog Manager API. Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0912N Index *indexname* on table *tablename* is not part of the datastore definition.

Explanation: The check option for the Information Catalog Manager API initialization discovered that table *tablename* has an index defined which does not match the indexes required by the Information Catalog Manager API.

User response: If the index is a unique index, it may interfere with the property functioning of the Information Catalog Manager API. Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0913N Function *functionname* is missing from the IBM DB2 tools catalog.

Explanation: The check option for the Information Catalog Manager API initialization found that the function *functionname* is missing from the DB2 tools catalog.

User response: Run the Information Catalog Manager API initialization with the fix option to repair this error.

ICM0914N The definition of function *functionname* is not valid.

Explanation: The check option for the Information Catalog Manager API initialization discovered that the definition of function *functionname* does not match that which is required by the Information Catalog Manager API.

User response: Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM0915N Trigger *triggername* is missing from the IBM DB2 tools catalog.

Explanation: The check option for the Information Catalog Manager API initialization discovered that the trigger *triggername* is missing from the IBM DB2 tools catalog.

User response: Run the Information Catalog Manager API initialization with the fix option to repair this error.

ICM0916N The definition of trigger *triggername* is not valid.

Explanation: The check option for the Information Catalog Manager API initialization discovered that the definition of trigger *triggername* does not match the definition that is required by the Information Catalog Manager API.

User response: Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM00917N Trigger *triggername* is not part of the datastore definition.

Explanation: The check option for the Information Catalog Manager API initialization discovered that a trigger named *triggername* is defined for one of the DB2 tools catalog tables which does not match any of the triggers required by the initialization tool.

User response: Run the Information Catalog Manager API initialization with the fix option in force mode to repair this error. Using the force option may result in lost data.

ICM00918N The view *viewname* is missing from the DB2 tools catalog.

Explanation: The check option for the Information Catalog Manager API initialization found that the view named *viewname* was not defined.

User response: Run the Information Catalog Manager API initialization with the fix option to repair this error.

ICM0919N Property *propertyname* of object type *typename* has a different data type than the property of the default object type.

Explanation: An existing object type matches the name of one of the default object types, but one of its properties has a different data type than the default definition for the object type.

User response: No change is required to keep the current definition. To replace the object type with the default definition, run the application initialization with the APP_REPLACE option.

ICM00920N Property *propertyname* of object type *typename* does not exist in the default definition for the object type.

Explanation: An existing object type matches the name of one of the default object types, but one of its properties has a different data type than the default definition for the object type.

User response: No change is required to keep the current definition. To replace the object type with the default definition, run the application initialization with the APP_REPLACE option.

ICM00921N Property *propertyname* of object type *typename* exists in the default definition, but is missing from the existing object type.

Explanation: An existing object type matches the name of one of the default object types, but one of its properties has a different data type than the default definition for the object type.

User response: No change is required to keep the current definition. To replace the object type with the default definition, run the application initialization with the APP_REPLACE or APP_MERGE options.

ICM00922N Could not access the Manage Information Catalog wizard because a valid license for DB2 Warehouse Manager could not be found. Please contact a local software reseller or IBM marketing representative.

Explanation: An attempt was made to initialize or migrate an information catalog without a valid license.

User response: Purchase the DB2 Warehouse Manager package and install the Information Catalog Manager Tools component which includes the Manage Information Catalog wizard.

ICM0923N **An SQL error occurred while dropping the DB2 Version 7 Information Catalog emulation views.**

Explanation: An SQL error occurred while the initialization tool was dropping the views that provide emulation of the DB2 version 7 Information Catalog Manager tables.

User response: Check the Message Reference for more information on this SQL error.

Chapter 104. ICM10000 - ICM10499

ICM10001N The information catalog is already open.

Explanation: The open method of the ICMCatalog class was called while the information catalog had already been opened.

User response: Delete the redundant call to open the catalog.

ICM10002N Cannot commit a batch checkpoint while there is no active batch.

Explanation: The checkpoint operation commits a set of transactions which are pending changes to the catalog; the transactions are committed to the information catalog. There are no pending operations while there is no active batch.

User response: Commit a batch checkpoint only when there is an active batch. Use startBatch to make a batch active and endBatch to end it. Use checkpoint to commit intermediate points in the batch of changes.

ICM10003N The information catalog is not open.

Explanation: A catalog must be open before operations can be performed against it.

User response: Open the catalog before attempting this operation.

ICM10004N The access control list named *accesslistname* is not unique.

Explanation: The attempted retrieval of an access control list by name returned more than one instance. This is an internal catalog error that must be repaired because access control lists are required to have unique names.

User response: Repair the information catalog.

ICM10005N Batch mode is not active.

Explanation: The endBatch method was called when no batch was active.

User response: A batch must be started before it can end. Use the sequence startBatch, modify catalog, checkpoint, modify catalog, checkpoint, modify catalog, endBatch for valid batch operation.

ICM10006N Batch mode is already active.

Explanation: The startBatch method was called while no batch was active.

User response: A batch cannot start until a previous

batch has ended. Use the sequence startBatch, modify catalog, checkpoint, modify catalog, checkpoint, modify catalog, endBatch for valid batch operation.

ICM10007N Failed to remove the last recorded checkpoint.

Explanation: The Information Catalog Manager API records checkpoints as they are committed during a batch. This is done so that the batch can be resumed from the last successful checkpoint if a failure occurs. On successful completion of the batch, the checkpoint record is removed. It is the removal of the checkpoint record that failed.

User response: Check the nested exception to determine the cause of the failure.

ICM10008N An I/O error occurred while creating the export temporary file *filename*.

Explanation: An error occurred while creating a temporary file that holds information during the export process.

User response: Check the nested exception and correct the problem.

ICM10009N Export operation to file *filename* failed.

Explanation: An error occurred during an export operation.

User response: Check the export log file and correct the problem.

ICM10010N Import failed. Last completed checkpoint was *checkpoint*.

Explanation: An import operation was partially successful but failed before completion.

User response: Check the import log file and correct the problem.

ICM10011N Import failed. No checkpoints completed.

Explanation: An import operation failed before completing any checkpoints.

User response: Check the import log file and correct the problem.

ICM10012N The access control list is not the default access control list.

Explanation: The access control List provided to setDefaultACL must be the access control list acquired from getDefaultACL.

User response: Check the documentation for ICMCatalog.setDefaultACL.

ICM10013N Checkpoint object instance
user=username, host=hostname,
batch=batchname in internal object type
typename is not unique.

Explanation: Checkpoints are kept in an internal object type with the unique ID set to the user, host, batch properties. It should not be possible for ambiguous checkpoint object instances to exist.

User response: Contact a database administrator for catalog repair.

ICM10014N Unable to retrieve host name for
checkpoint identification.

Explanation: An unknown host exception was thrown while retrieving the host name.

User response: Contact a system administrator to examine the nested exception.

ICM10015N Unexpected property constraint
exception occurred while working with
object instance user=username,
host=hostname, batch=batchname of the
internal object type objecttype.

Explanation: A property constraint exception was thrown while working with an instance of the internal checkpoint object type. An internal catalog problem occurred because no property constraints are placed on the properties of this object type.

User response: Contact a database administrator for catalog repair.

ICM10016N Unexpected security exception while
working with object instance
user=username, host=hostname,
batch=batchname of the internal object
type typename.

Explanation: Checkpoint object instances are owned by the current user at the time the batch completed a checkpoint. No security violation should occur.

User response: Determine the security violation and correct it.

ICM10017N Unexpected property constraint
exception occurred while recording the
delete history in the internal object type
objecttype.

Explanation: A property constraint exception was

thrown while working with an instance of the internal object type used to record the delete history. There is an internal catalog problem because no property constraints are placed on the properties of this object type.

User response: Turn off the delete history feature and contact a database administrator for catalog repair.

ICM10100N Internal object type objecttype was not
found. This catalog is not usable.

Explanation: An internal object type used to support the Information Catalog Center application was not found in the IBM DB2 tools catalog. The catalog was not properly initialized for Information Catalog Center use or it was corrupted.

User response: Check that the metadata store containing the catalog has been properly initialized for the Information Catalog Center application.

ICM10101N Internal relationship type relationshiptype
was not found. This catalog is not
usable.

Explanation: An internal relationship type used to support the Information Catalog Center application was not found in the IBM DB2 tools catalog. The catalog was not properly initialized for Information Catalog Center use or it has been corrupted.

User response: Check that the metadata store containing the catalog has been properly initialized for the Information Catalog Center application.

ICM10200N A collection named name already exists.

Explanation: The collection name supplied was identical to the name of an already existing collection owned by the current user.

User response: Specify a different name for the new collection.

ICM10201N Unable to create a new collection with a
default name.

Explanation: The Information Catalog Manager API can create up to 1000 different default collection names. Failed to create the new collection because all 1000 default names already exist.

User response: Rename some collections to non-default names.

ICM10202N Collection name is specified for deletion
but does not match any collection
owned by the current user.

Explanation: The collection name specified as the target of a delete operation was not found in the current user's collections.

User response: Specify an existing collection.

ICM10203N The relationship instance identifying members of collection *name* could not be found.

Explanation: The internal relationship type instance, which is used to identify members of a collection, was not associated with the collection. An internal error occurred and the collection is not usable.

User response: Use the Information Catalog Manager API to delete the collection.

ICM10204N Too many relationship instances exist on collection *name*.

Explanation: More than one relationship instance was found attached to an instance of the internal object type used to store collections. An internal error occurred in the information catalog and this collection is not usable.

User response: Use the Information Catalog Manager API to delete the collection.

ICM10205N Relationship instance on collection *name* is not valid.

Explanation: A relationship instance of an incorrect type was found attached to an instance of the internal object type used to store collections. This is an internal error in the information catalog and the collection is not usable.

User response: Use the Information Catalog Manager API to delete the collection.

ICM10300N Could not recognize the property search type specified.

Explanation: The property search type supplied to a DateSearchCriteria or TextSearchCriteria is not one of those defined.

User response: Specify one of the defined property search types.

ICM10301N No search date specified.

Explanation: Attempted to run a date search with no dates specified.

User response: Specify a start date, end date, or both.

ICM10302N Search has no associated criteria.

Explanation: Attempted to call the create method to write a search with no associated criteria to the catalog.

User response: Associate a criteria with the search before creating it.

ICM10303N Error occurred while loading search criteria for search *name*.

Explanation: An exception was caught while reading the search criteria for a search from the catalog.

User response: Check the nested exception to determine the problem.

ICM10304N Error occurred while storing search criteria for search *name*.

Explanation: An exception was caught while writing the search criteria for a search to the catalog.

User response: Check the nested exception to determine the problem.

ICM10305N Maximum returned object count cannot be negative. Count was: *error-code*.

Explanation: A SearchCriteria object was given a negative value for maximum objects returned. An error would occur if the SearchCriteria is run.

User response: Specify a maximum object count which is not negative.

ICM10306N Object type search collection contains foreign object.

Explanation: Only ICMObjectTypes can be specified to identify the types of ICMObjectInstances to search.

User response: Check that the collection of object types only contains objects of type ICMObjectType.

ICM10307N The search name *name* is in use.

Explanation: A user cannot have two searches with the same name.

User response: Specify a search name which is not in use.

ICM10308N Unable to find a new default search name.

Explanation: The Information Catalog Manager API can create up to 1000 different default search names when creating a new search. Failed to create a new default search name because all default names are in use.

User response: Rename some searches that have default names.

ICM10309N No search strings specified.

Explanation: Attempted to perform a text search with no text strings specified.

User response: Specify search strings to be matched.

ICM10310N The search *name* specified for deletion does not match existing searches owned by the current user.

Explanation: The search name specified as the target of a delete was not found in the current user's searches.

User response: Specify an existing search.

ICM10311N Search values must be text strings.

Explanation: Attempted to set search values for a TextSearchCriteria which were not text strings.

User response: Specify search text strings to be matched.

ICM10400N Message Text : Error occurred while retrieving property data for *user/group* from the internal object type *objecttype*.

Explanation: An exception was caught while retrieving the stored property information from an object instance of the internal object type.

User response: Contact a database administrator to check the nested exception.

ICM10401N Error occurred while storing property data for *user/group* to the internal object type *objecttype*.

Explanation: An exception was caught while storing property information to an object instance of the internal object type.

User response: Contact a database administrator to check the nested exception.

ICM10402N Unable to retrieve host name.

Explanation: An unknown host exception was thrown while retrieving the host name.

User response: Consult a system administrator to check the nested exception.

Chapter 105. ICM10500 - ICM10999

ICM10500N Error occurred while starting program with command string *command*.

Explanation: An exception was thrown while attempting to invoke a program on an object instance.

User response: Check the nested exception and correct the program or path object.

ICM10501N Internal error occurred while attempting to write a BlobProperty to a file.

Explanation: The BlobProperty has been created in the database but the associated object instance has not.

User response: Contact your IBM service representative.

ICM10502N Property *propertyname* is an incorrect type.

Explanation: A property name passed to writeBlobToStream or writeClobToWriter was not a correct Blob or Clob property.

User response: Check that the property is a Blob or Clob as needed by the method being called.

ICM10503N Unable to retrieve the value from the *propertyname* property of an instance of object type *objecttype*.

Explanation: An exception was caught while retrieving a value from the property of the object type.

User response: Contact a database administrator to check the nested exception.

ICM10504N Unmatched parameter marker in a parameterized string. String = *paramstring*.

Explanation: A parameter marker open character was found without a corresponding parameter marker close character while formatting a parameterized string.

User response: Check that the parameterized string is in a valid format.

ICM10505N Error occurred while writing data from BlobProperty or ClobProperty *propertyname* to file *filename*.

Explanation: A parameter of the type BlobProperty or ClobProperty was encountered as one of the parameters while formatting a parameterized string. An exception was caught while writing the contents of the property to a file.

User response: Check the nested exception.

ICM10506N An internal error occurred while attempting to write property *propertyname* to a file.

Explanation: Attempted to write a property to a file while formatting a parameterized string. The property was not a BlobProperty or a ClobProperty, which are the only property types allowed to be written to a file.

User response: Contact your IBM service representative.

ICM10600N Icon data size *datasize* is larger than *maxsize* bytes.

Explanation: The icon associated with an object type cannot be larger than the indicated maximum size.

User response: Specify a smaller icon file. Contact a database administrator if the icon data has been corrupted in the catalog.

ICM10601N Error occurred while retrieving icon data from catalog.

Explanation: An SQLException was caught while reading the icon data from the object type.

User response: Check the nested exception.

ICM10700N The name specified (*typename*) matches more than one object type.

Explanation: The name used to retrieve an object type matched more than one object type. The inclusion of a wildcard character in the name might cause this error.

User response: Supply a name which uniquely identifies an object type.

ICM10701N The object type *typename* does not exist.

Explanation: The object type name used with an Information Catalog Manager API method did not match existing object types.

User response: Use the name of an existing object type.

ICM10702N The required relationship category *categoryname* can not be found. This catalog is not usable.

Explanation: One of the required relationship

categories in an information catalog could not be found. This is a fatal error.

User response: Contact a database administrator.

ICM10703N The object type *typename* is not an information catalog user object type.

Explanation: The IBM DB2 tools catalog supports multiple application programs. The object types seen by information catalog users are a subset of all the object types in a tools catalog. An exception occurred while passing an object type to the Information Catalog Manager API which is not one of the information catalog user object types.

User response: Only use the Information Catalog Manager API with information catalog user object types.

ICM10704N An operation to store the default property visibility flags for the object type *typename* failed.

Explanation: An exception was caught while storing the information that defines default property visibility for this object type.

User response: Contact a database administrator to check the nested exception.

ICM10705N An operation to load the default property visibility flags for the object type *typename* failed.

Explanation: An exception was caught while loading the information that defines default property visibility for this object type.

User response: Contact a database administrator to check the nested exception.

ICM10706N The object type *typename* is not an information catalog user object type.

Explanation: The IBM DB2 tools catalog supports multiple application programs. The object types seen by information catalog users are a subset of all the object types in a tools catalog. An exception occurred because an object instance of a type not in the information catalog user object type set was linked by a relationship to an instance of an object type which is in the information catalog user object type set.

User response: Use the Information Catalog Manager API with information catalog user object types.

ICM10800N The required relationship constraint *constraintname* may not be removed.

Explanation: Relationship constraints placed on relationship types enforce required behaviors within the

information catalog. The constraints may not be removed.

User response: Do not attempt to remove the required relationship constraint.

ICM10801N Error occurred while enforcing required constraint *constraintname* on relationship type *relationshiptype*.

Explanation: An exception was thrown while attempting to instantiate a required constraint object to apply to a relationship type.

User response: Contact a database administrator to check the nested exception.

Part 18. LIC Messages

Chapter 106. LIC1000 - LIC1499

LIC1052E **You must be root to execute this program.**

Explanation: This program can only be run under the root user ID. Special privileges are required to execute this program.

User response: Login as root and issue the command again.

LIC1304E **Unexpected error.**

Explanation: The tool encountered an unexpected system error.

User response: Contact your DB2 service representative.

LIC1305E **The profile registry was not found.**

Explanation: The target machine does not have a profile registry setup.

User response: Create the registry on the target machine by installing DB2.

LIC1309E **System error.**

Explanation: The tool encountered an operating system error.

User response: A system error was encountered during registry access. Ensure that there is enough space on the file system where the registry is located, and that there is a valid LAN connection if the registry is remote.

LIC1400N **The syntax of the db2licm command is incorrect. Run the db2licm -? command for more information.**

Explanation: The db2licm tool performs basic license functions. It adds, removes, lists, and modifies licenses installed on the local system. Execute db2licm tool with -l parameter to find out the product identifier for your product:

```
db2licm [-a filename]
        [-e product-identifier HARD | SOFT]
        [-p product-identifier
        REGISTERED | CONCURRENT | OFF]
        [-r product-identifier]
        [-u product-identifier num-users]
        [-c product-identifier num-connectors]
        [-l]
        [-v]
        [-?]
```

The command options are:

-a

Adds a license for a product. Specify a file name containing valid license information. This can be obtained from your licensed product CD or contact your IBM representative or authorized dealer.

-e

Updates the enforcement policy on the system. Valid values are: HARD and SOFT. HARD specifies that unlicensed requests will not be allowed. SOFT specifies that unlicensed requests will be logged but not restricted.

-p

Updates the license policy type to use on the system. The keywords CONCURRENT, REGISTERED, or CONCURRENT REGISTERED can be specified. Specify OFF to turn off all policies.

-r

Removes the license for a product. After the license is removed, the product functions in "Try & Buy" mode. To get the password for a specific product, invoke the command with the -l option.

-u

Updates the number of user entitlements that have been purchased. Specify the password of the product for which the entitlements were purchased and the number of users.

-c

Updates the number of connector entitlements that have been purchased. Specify the password of the product for which the entitlements were purchased and the number of connectors.

-l

Lists all the products with available license information, including the product identifier.

-v

Displays version information.

-?

Displays help information. When this option is specified, all other options are ignored, and only the help information is displayed.

User response: Enter the command again using the valid parameters.

LIC1401I Command line DB2 License Manager.

Explanation: The db2licm tool performs basic license functions. It adds, removes, lists, and modifies licenses registered on the local system. Execute db2licm tool with -l parameter to find out the product identifier for your product:

```
db2licm [-a filename]
        [-e product-identifier HARD | SOFT]
        [-p product-identifier
        CONCURRENT | OFF]
        [-r product-identifier]
        [-u product-identifier num-users]
        [-c product-identifier num-connectors]
        [-g filename]
        [-x]
        [-l][show detail]
        [-v]
        [-?]
```

The command options are:

-a

Adds a license for a product. Specify a file name containing valid license information. This can be obtained from your licensed product CD or contact your IBM representative or authorized dealer.

-e

Updates the enforcement policy on the system. Valid values are: HARD and SOFT. HARD specifies that unlicensed requests will not be allowed. SOFT specifies that unlicensed requests will be logged but not restricted.

-p

Updates the license policy type to use on the system. The keyword CONCURRENT can be specified for concurrent user policy. Specify OFF to turn off all policies.

-r

Removes the license for a product. Specify the product identifier.

-u

Updates the number of user entitlements that have been purchased. Specify the product identifier and the number of users.

-c

Updates the number of connector entitlements that have been purchased. Specify the product identifier and the number of connector entitlements.

-g

Generates compliance report. Specify file name where output is to be stored.

-x

Resets license compliance information for the purposes of license compliance report.

-l[show detail]

Lists all the products with available license information, including the product identifier. Specify [show detail] to view detailed information about licensed features (if any).

-v

Displays version information.

-?

Displays help information. When this option is specified, all other options are ignored, and only the help information is displayed.

User response:

LIC1402I License added successfully.

LIC1403I License removed successfully.

LIC1404N Product identifier not found.

Explanation: The given identifier is either invalid, or a license for this product was not found in the nodelock file.

User response: Issue this command with -l option to check that the identifier entered is the correct product identifier for the product that you want to perform this action on. If you are using nodelock passwords, check that the license key for this product is installed in the nodelock file.

LIC1405I License policy type updated successfully.

LIC1406N Invalid license policy type.

Explanation: The license policy type that was entered was not valid for the product specified.

User response: Please enter a valid license policy. Options are:

- CONCURRENT
- OFF

**LIC1407N You are trying to register an invalid license certificate file,
license-certificate-file-name.**

Explanation: Either the license file is not from the current version, or the license file is corrupted. You can check the current version by running db2licm -v

User response: Obtain the valid license file for the current version from the Activation CD, and rerun the

db2licm command. For example, db2licm -a *license-certificate-file-name*

LIC1408N **The file *file-name* could not be opened. Enter the name of a file that exists and can be opened and try the command again.**

Explanation: The file is not found or access to the file is denied.

User response: Enter the name of a file that exists and can be opened and try the command again.

LIC1409N **Invalid enforcement policy type.**

Explanation: The enforcement policy type specified is not valid for this product.

User response: Please enter a valid enforcement policy type that is supported by the specified product.

LIC1410I **Concurrent entitlements updated successfully.**

LIC1411I **Enforcement policy type updated successfully.**

LIC1412W **A hard stop enforcement policy has been set. This enforcement policy stops unlicensed requests.**

Explanation: You issued the db2licm command with the -e parameter, to update the enforcement policy, and specified the value HARD. (For example, db2licm -e db2ese HARD.) The value HARD specifies that unlicensed requests will not be allowed.

User response: As a mechanism for you to keep track of, and differentiate, the DB2 database products and features installed on your system, it is recommended that you register the license key for each DB2 database product and feature.

If you want unlicensed requests to be logged but not restricted, change the enforcement policy to SOFT. For example, db2licm -e db2ese SOFT

LIC1413W **A soft stop enforcement policy has been set. This enforcement policy specifies that unlicensed requests will be logged but not restricted.**

Explanation: You issued the db2licm command with the -e parameter, to update the enforcement policy, and specified the value SOFT. (For example, db2licm -e db2ese SOFT.) The value SOFT specifies that unlicensed requests will be logged but not restricted.

User response: If you want unlicensed requests to be stopped, you must change the enforcement policy to HARD. For example, db2licm -e db2ese HARD.

LIC1416N **The license could not be added to the nodelock file automatically. The return code is *return-code*.**

User response: Please ensure the license certificate is readable. You may also enter the license into the nodelock file manually. Please see the license file for instructions.

LIC1417N **The license specified could not be removed from the nodelock file. The return code is *return-code*. Check that the license for this product exists in the nodelock file.**

User response: Ensure that the license for this product exists in the nodelock file.

LIC1418I **The number of licensed processors on this system has been updated successfully.**

LIC1419N **There was an error updating the number of licensed processors. The return code is *return-code*.**

LIC1420N **This product does not support this type of license policy.**

Explanation: The license policy specified does not apply to this product or is not supported.

User response: Enter a valid license policy or select a product that supports this policy.

LIC1421N **This product specified is not installed on this system.**

Explanation: You can not configure a license policy for a product until the product is installed.

User response: Install the product before running this command or specify the correct product identifier. To list the products install on the system issue db2licm -l command.

LIC1422N **The number of concurrent entitlements was not updated. The return code is *return-code*.**

User response: Please ensure the concurrent policy is enabled for this product.

LIC1423N **This option requires the creation of an instance.**

Explanation: Features that are required to perform this action are only accessible once an instance has been created.

User response: Please create an instance and issue this command again.

LIC1424N **An unexpected error occurred while accessing processor information.**

Explanation: The return code is *return-code*.

User response: None.

LIC1426I **This product is now licensed for use as outlined in your License Agreement. USE OF THE PRODUCT CONSTITUTES ACCEPTANCE OF THE TERMS OF THE IBM LICENSE AGREEMENT, LOCATED IN THE FOLLOWING DIRECTORY:**
dir-name

LIC1427I **This product is now licensed for use as outlined in your License Agreement. USE OF THE PRODUCT CONSTITUTES ACCEPTANCE OF THE TERMS OF THE IBM LICENSE AGREEMENT, LOCATED IN THE FOLLOWING DIRECTORY:**
dir-name

LIC1428N **There was an error updating the number of licensed processors.**

Explanation: The number of licensed processors entered exceeds the number of maximum licensed processors allowed for this product.

User response: Please enter number of licensed processors that does not exceed the defined maximum. If the number of processors on your system exceeds the maximum number of processors allowed for this product, please contact your IBM representative or authorized dealer.

LIC1429N **This product does not support this combination of license policies.**

User response: Please enter a valid combination of license policies. For example, you can specify "CONCURRENT REGISTERED" as a valid combination.

LIC1430N **The license could not be added to the nodelock file because the license date is greater than operating system date.**

User response: Please check your certificate file to ensure that the license start date precedes the current date (the date set on the operating system).

LIC1431N **This user does not have sufficient authority to perform the specified action.**

Explanation: This action can be run only by the root user ID or by a user ID with SYSADM authority.

User response: Login with a user ID that has permission to run this command.

LIC1432N **The license could not be added to the nodelock file because this product has used the maximum number of evaluation licenses. The maximum number of evaluation licenses is *lic-number*. Run this command again with a permanent license key.**

Explanation: This product has used the maximum number of evaluation licenses.

User response: Run this command again with a permanent license key.

LIC1433N **The number of license entitlements was not updated.**

Explanation: The specified number of license entitlements is not in the valid range.

User response: Run this command again using a valid number of license entitlements.

LIC1434N **DB2 has added the license entry to the nodelock file, however, this license entry is not active.**

Explanation: DB2 failed to activate this license entry, therefore DB2 will run with the previous license configuration until this license is activated.

User response: Try the command again and if it continues to fail, edit the nodelock file manually or contact IBM support.

If you edit the nodelock file manually, move the new license entry to the beginning of the license entries list.

The nodelock file could be found in the following locations:

Windows XP and Windows 2003

X:\Documents and Settings\All
Users\Application Data\IBM\DB2\<DB2
copy name>\license.

Windows Vista

X:\ProgramData\IBM\DB2\<DB2 copy
name>\license

Where 'X:' is the system drive.

On all other platforms the nodelock file is located in the installation path of this product in the license directory.

Refer to the DB2 Information Center for more information on licensing.

LIC1435E **An I/O error occurred when accessing the nodelock file. The license could not be added.**

Explanation: An error occurred when creating or accessing nodelock file. The file access settings do not allow this action.

User response: Ensure that the nodelock file and the directory where nodelock file is located allow read and write access to this program.

The nodelock file could be found in the following locations:

Windows XP and Windows 2003

X:\Documents and Settings\All
Users\Application Data\IBM\DB2\<DB2
copy name>\license.

Windows Vista, Windows 7, Windows 2008 or later

X:\ProgramData\IBM\DB2\<DB2 copy
name>\license

Where 'X:' is the system drive.

On all other platforms the nodelock file is located in the installation path of this product in the license directory.

LIC1436I **Duplicate license was found in nodelock file.**

Explanation: DB2 has determined that this license has already been registered in the nodelock file for this installation of DB2.

User response: No further action is necessary.

LIC1437I **License entitlements updated successfully.**

LIC1438E **An I/O error occurred when accessing the nodelock file. The license could not be removed.**

Explanation: An error occurred when creating or accessing nodelock file. The file access settings do not allow this action.

User response: Ensure that the nodelock file and the directory where nodelock file is located allow read and write access to this program.

The nodelock file could be found in the following locations:

Windows XP and Windows 2003

X:\Documents and Settings\All
Users\Application Data\IBM\DB2\<DB2
copy name>\license.

Windows Vista

X:\ProgramData\IBM\DB2\<DB2 copy
name>\license

Where 'X:' is the system drive.

On all other platforms the nodelock file is located in the installation path of this product in the license directory.

LIC1439I **DB2 server has detected that *product-name* is installed on this system. Products and functions obtained via this offering may only be used for testing or development purposes as outlined in your License Agreement. The License Agreement for this offering is located in the 'license' directory in the installation path for this product.**

LIC1440I **License compliance report generated successfully.**

LIC1441I **License compliance information was reset.**

LIC1442E **An error occurred when generating compliance report.**

Explanation: Compliance report could not be created.

User response: Ensure this program is able to write to the file specified and try again.

LIC1443E **An error occurred when resetting compliance information.**

Explanation: Compliance information could not be reset.

User response: Ensure this program is able to write to the license directory in the installation path and try again.

LIC1444E **An I/O error occurred. The return code is. *return-code* .**

LIC1445E **An error occurred when generating the compliance report.**

Explanation: An unexpected error occurred when generating the compliance report. The compliance report could not be created.

User response:

- Ensure at least one valid DB2 instance is created.
- Ensure the DB2 global registry is not corrupted.
- Ensure the DB2 Administration Server was started successfully.

LIC1446I **The license certificate**
license-certificate-file-name **for SA MP was successfully installed.**

Explanation: The IBM Tivoli System Automation for Multiplatforms (SA MP) requires a valid license certificate to work with the DB2 High Availability (HA) feature. This license certificate was successfully installed or updated.

User response: No response is required.

LIC1447N **The license certificate**
license-certificate-file-name **for SA MP was not successfully installed.**

Explanation: The IBM Tivoli System Automation for Multiplatforms (SA MP) requires a valid license certificate to work with the DB2 High Availability (HA) feature. This license certificate was not successfully installed or updated.

If you used the DB2 installer to install or update this license certificate, you can find more detailed information about why the install or update failed in the DB2 install log file.

User response: To manually install or update this license certificate for the SA MP Base Component, issue the command:

- **samlicm** -i *license-certificate-file-name*

For more information about the **samlicm** command, see the SA MP Base Component documentation.

LIC1448I **This license was automatically applied at install time in order to enable you to start working with DB2.**

Explanation: To be fully licensed, this product requires a license appropriate to your purchased license policy.

User response: A license can be downloaded from Passport Advantage or may be found on a separate CD in your product package. Both the download and the CD are titled "Activation CD".

For more information on licensing your product search the Information Center using terms such as "licensing".

LIC1449N **The license was not installed due to a platform restriction.**

Explanation: This DB2 product is only supported in trial mode, also known as "Try and Buy" mode, on this platform.

User response: Continue to use this product in trial mode, or install one which is fully supported on this platform.

LIC1450I **The product licensed by the certificate**
file-name **was not found in the DB2 copy.**

Explanation: Additional licenses may be added to a DB2 copy prior to additional product installation. The license has been successfully added but it will not be shown until the corresponding product is installed.

User response: No action is necessary. If you subsequently install the product covered by this certificate, you do not need to re-register the license.

Part 19. MQL Messages

Chapter 107. MQL0000 - MQL0499

MQL0001E MQListener was invoked with no major function on the command line.

Explanation: The MQListener command line failed to specify a major function, such as help, run, add, remove, or show.

User response: Consult MQListener documentation and re-run MQListener with the desired function. For more information, run

db2mqlsn help

MQL0002E MQListener was invoked with an unknown command line parameter *parameter name*.

Explanation: The MQListener command line contained an unknown parameter *parameter name*.

User response: Consult MQListener documentation and re-run MQListener with the desired parameters. For more information, run

db2mqlsn help

MQL0003E MQListener was invoked without the required command line parameter *parameter name*.

Explanation: The MQListener command line was missing the required parameter *parameter name*.

User response: Consult MQListener documentation and re-run MQListener with the required parameter. For more information, run

db2mqlsn help

MQL0004E MQListener was invoked with an illegal value for the *parameter name* command line parameter: *parameter value*.

Explanation: The MQListener command line contained an illegal value for the *parameter name* parameter. The illegal value was *parameter value*.

User response: Consult MQListener documentation and re-run MQListener with the desired parameters. For more information, run

db2mqlsn help

MQL0010E MQListener encountered a database error accessing database *database name* during operation *operation name*: sqlstate = *sqlstate value* (sqlcode = *sqlcode value*).

Explanation: MQListener encountered a DB2 database

error accessing database *database name* while executing operation *operation name*. The resulting sqlstate was *sqlstate value* (sqlcode *sqlcode value*).

User response: Consult MQListener documentation and insure that MQListener is properly installed in the specified database and access to MQListener packages and configuration tables has been granted to the user specified in -configUser or the user running MQListener.

MQL0011E MQListener failed to connect to database *database name* as user *user name*: sqlstate = *sqlstate value* (sqlcode = *sqlcode value*).

Explanation: MQListener encountered a DB2 database error while attempting to connect to database *database name* as user *user name*. The resulting sqlstate was *sqlstate value* (sqlcode *sqlcode value*).

User response: Insure that *database name* is running and accessible to *user name* using the password configured in MQListener. If necessary, re-configure MQListener using MQListener "remove" and "add" functions.

MQL0020E MQListener encountered an MQ error executing operation *operation name* on object *object name* in queue manager *queue manager*: reason code = *reason code*.

Explanation: MQListener encountered an MQ error executing operation *operation name* on object *object* in queue manager *queue manager*. The resulting reason code was *reason code value*.

User response: Insure that MQ is properly installed and configured, and accessible to the user running MQListener. Consult the MQ documentation for more information, particularly the Application Programming Reference (Document Number SC34-6062-00) for more information about the reason code.

MQL0021E MQListener failed to connect to MQ queue manager *queue manager*: reason code = *reason code*.

Explanation: MQListener encountered an MQ error while attempting to connect to MQ queue manager *queue manager*. The resulting reason code was *reason code*.

User response: Insure that MQ is properly installed and configured, and accessible to the user running MQListener. Consult the MQ documentation for more information, particularly the Application Programming

Reference (Document Number SC34-6062-00) for more information about the reason code.

SQL0022W **SQLListener failed to disconnect from MQ queue manager** *queue manager*:
reason code = reason code.

Explanation: SQLListener encountered an MQ error while attempting to disconnect from MQ queue manager *queue manager*. The resulting reason code was *reason code*.

User response: Insure that MQ is properly installed and running. Consult the MQ documentation for more information, particularly regarding the exact meaning of the reason code.

SQL0030E **SQLListener encountered a DB2 database error attempting to add a new task to the** *configuration name*
SQLListener configuration in database *configuration database name*. **The input queue specified in the task was** *input queue*, **and the specified queue manager was** *queue manager*. **The resulting sqlstate was** *sqlstate value (sqlcode = sqlcode value)*.

Explanation: SQLListener encountered a DB2 database error attempting to add a new task to the *configuration name* SQLListener configuration in database *configuration database name*. The input queue specified in the task was *input queue* and the specified queue manager was *queue manager*. The resulting sqlstate was *sqlstate value (sqlcode sqlcode value)*.

User response: Consult SQLListener documentation and insure that SQLListener is properly installed in the specified configuration database and access to SQLListener packages and configuration tables has been granted to the user specified in -configUser or the user running SQLListener. Insure that the combination of configuration name, input queue, and queue manager is unique among the SQLListener configuration tasks in the database. Use SQLListener "show" command to inspect the configuration.

SQL0040E **SQLListener encountered a DB2 database error retrieving configuration** *configuration name* **from database** *configuration database name* **during operation** *operation name*. **The resulting sqlstate was** *sqlstate value (sqlcode = sqlcode value)*.

Explanation: SQLListener encountered a DB2 database error retrieving configuration *configuration name* from database *configuration database name* during operation *operation name*. The resulting sqlstate was *sqlstate value (sqlcode = sqlcode value)*.

User response: Consult SQLListener documentation

and insure that SQLListener is properly installed in the specified configuration database and access to SQLListener packages and configuration tables has been granted to the user specified in -configUser or the user running SQLListener.

SQL0060E **SQLListener encountered a DB2 database error while removing a task from configuration** *configuration name* **in database** *configuration database name*. **The input queue specified was** *input queue*, **and the queue manager specified was** *queue manager*. **The resulting sqlstate was** *sqlstate value (sqlcode = sqlcode value)*.

Explanation: SQLListener encountered a DB2 database error while removing a task from configuration *configuration name* in database *configuration database name*. The input queue specified was *input queue*, and the queue manager specified was *queue manager*. The resulting sqlstate was *sqlstate value (sqlcode = sqlcode value)*.

User response: Consult SQLListener documentation and insure that SQLListener is properly installed in the specified configuration database and access to SQLListener packages and configuration tables has been granted to the user specified in -configUser or the user running SQLListener.

SQL0061I **SQLListener attempted to remove an unknown task from configuration** *configuration name* **in database** *configuration database name*. **The input queue specified was** *input queue*, **and the queue manager specified was** *queue manager*.

Explanation: SQLListener attempted to remove an unknown task from configuration *configuration name* in database *configuration database name*. The input queue specified was *input queue*, and the queue manager specified was *queue manager*.

User response: Insure that specified parameters are correct, and re-run SQLListener. Use SQLListener "show" command to inspect the configuration.

SQL0070E **SQLListener has encountered a DB2 database error preparing to run the stored procedure** *schema name.procedure name* **in database** *database name* **as user** *user name*: **sqlstate =** *sqlstate value (sqlcode = sqlcode value)*.

Explanation: SQLListener has encountered a DB2 database error preparing to run the stored procedure *schema name.procedure name* in database *database name* as user *user name*. The resulting sqlstate was *sqlstate value (sqlcode sqlcode value)*.

User response: Insure that SQLListener configuration

is correctly specified and that a stored procedure of the specified schema and name exists with a correct signature and is accessible to the specified user.

MQL0071E MQLListener encountered an unsupported datatype in the signature of the stored procedure *schema name.procedure name* in database *database name* as user *user name*: datatype = *datatype value*.

Explanation: MQLListener encountered an unsupported datatype in the signature of the stored procedure *schema name.procedure name* in database *database name* as user *user name*. The datatype was *datatype value*.

User response: Consult the MQLListener documentation and insure that the stored procedure has the correct datatype.

MQL0072E MQLListener was unable to start a thread running a task specified in the configuration *configuration name* by the input queue *input queue* and queue manager *queue manager*. The ECF error code was *error code*.

Explanation: MQLListener was unable to start a thread running a task specified in the configuration *configuration name* by the input queue *input queue* and queue manager *queue manager*. The ECF error code was *error code*.

User response: Insure that MQLListener configuration is correct, especially the -numInstances parameter, and that the MQLListener process has sufficient system resources to execute the required number of simultaneous tasks.

MQL0073I MQLListener started a thread running the task specified in the configuration *configuration name* by the input queue *input queue* and queue manager *queue manager*.

Explanation: MQLListener started a thread running the task specified in the configuration *configuration name* by the input queue *input queue* and queue manager *queue manager*.

User response: None. This is the normal notification of a thread start, expected during startup or restart.

MQL0074I A thread running the task specified in the configuration *configuration name* by the input queue *input queue* and queue manager *queue manager* in MQLListener has exited.

Explanation: A thread running the task specified in the configuration *configuration name*, input queue *input*

queue, and queue manager *queue manager* in MQLListener has exited.

User response: None. This is the normal notification of a thread exit, expected during shutdown or restart.

MQL0075E MQLListener encountered a DB2 database error while running the stored procedure *schema name.procedure name* in database *database name* as user *user name* during operation *operation name*. The resulting sqlstate was *sqlstate value* (sqlcode = *sqlcode value*).

Explanation: MQLListener encountered a DB2 database error while running the stored procedure *schema name.procedure name* in database *database name* as user *user name* during operation *operation name*. The resulting sqlstate was *sqlstate value* (sqlcode *sqlcode value*).

User response: Consult MQLListener documentation and insure that MQLListener is properly installed in the specified database and access to MQLListener packages and configuration tables has been granted to the specified user. Insure that the stored procedure executes correctly.

MQL0080W MQLListener received an unknown message on admin queue *admin queue*, in queue manager *queue manager*: *message*.

Explanation: MQLListener received an unknown message on admin queue *queue*, in queue manager *queue manager*: *message*.

User response: Insure that MQ is correctly installed and configured, that MQLListener is being run with the proper -adminQMGr and -adminQueue parameter values, and that the MQLListener "admin" command is also using the proper -adminQMGr and -adminQueue or -adminQueueList parameter values. If used, check that the -adminQueueList specifies the proper queue names. Check that the admin queue is not being used by another application.

MQL0081I MQLListener received a shutdown message on admin queue *admin queue*, queue manager *queue manager*.

Explanation: MQLListener received a shutdown message on admin queue *admin queue*, queue manager *queue manager*.

User response: None. This is the normal notification on receiving a shutdown message.

MQL0082I MQLListener received a restart message on admin queue *admin queue*, queue manager *queue manager*.

Explanation: MQLListener received a restart message on admin queue *admin queue*, queue manager *queue manager*.

User response: None. This is the normal notification on receiving a restart message.

Part 20. SAT Messages

This section contains the Satellite (SAT) messages. The messages are listed in numeric sequence.

Chapter 108. SAT1000 - SAT1499

SAT1000I **The first application version for this group has been defaulted to match the value used for the Typical satellite install.**

Explanation: If you install a satellite using the Typical method, the satellite's application version will be set to a predefined value. This predefined application version is the same as the default value for the first application version that is created for a group.

User response: If you did not install the satellites of this group using the Typical method, or you changed the application version for one or more satellites after installation, ensure that the identifier you supply to create a new application version is consistent with the satellite's application version.

SAT1001I **Mandatory information is missing from one of the notebook pages. The notebook will return to the page that is missing information.**

Explanation: All mandatory fields must be filled in for the action to complete.

User response: Fill in the required information and try the action again.

SAT1002I **A test connection or attachment to the selected target was successful using the specified authentication credentials.**

Explanation: An attempt was made to connect to a target. The connection or attachment was established successfully using the specified authentication credentials.

User response: No action is required.

Chapter 109. SAT2000 - SAT2499

SAT2000W **Changes have not been saved. Save now?**

Explanation: You are exiting from a window or notebook without saving your changes.

User response: To save your changes, click Yes. Otherwise, click No.

SAT2001W **An ATTACH or CONNECT statement was found in the script.**

Explanation: The script contents appear to contain either a CONNECT or an ATTACH statement. On a satellite, the required instance-level attach or database-level connect is automatically established before a DB2 instance or DB2 database script is executed. Consequently, it is not necessary to include ATTACH or CONNECT statements in DB2 instance or database scripts, unless the script specifies more than one execution target.

User response: If the script does not specify more than one target, remove the CONNECT or ATTACH statement from the script.

SAT2002W **Script contents may not be consistent with the script type.**

Explanation: Either the script contents have been modified in a way that causes an inconsistency with the script type, or the script type has been modified in a way that causes an inconsistency with the script contents.

User response: Ensure that the script type and the script contents are compatible.

SAT2003W **Implicit table space was created.**

Explanation: Promotion created at least one implicit table space.

User response: Change the table spaces in the script if the implicit table spaces that are created are not suitable for your business requirements.

SAT2014W **Are you sure that you want to promote the satellite *satellite*?**

Explanation: Promoting a satellite causes the satellite to start executing the group's batches. Only promote a satellite when it has successfully executed the fix batch and is ready to execute group batches.

User response: Click Yes to promote the satellite back to executing group batches. Otherwise, click No.

SAT2015W **Are you sure that you want to fix the satellite *satellite*?**

Explanation: Setting a satellite to fix mode will prevent the satellite from executing its group batches. The satellite will only be able to execute a fix batch until it is promoted back to executing the group batches. Only set a satellite to fix mode when the satellite requires servicing.

User response: Click Yes to set the satellite to fix mode so it can execute a fix batch. Otherwise, click No.

SAT2016W **An object with the name *object* already exists in the database *database*. Are you sure that you want to overwrite it?**

Explanation: Overwriting an object will cause all other objects that reference the object to be affected.

User response: Click Yes to overwrite the object. Otherwise, click No.

SAT2017W **An object with the name *object* already exists in the database *database* and is in use by another object. Are you sure that you want to change the definition of the object?**

Explanation: Changing the definition of an object will cause all other objects that reference that object to be affected.

User response: Click Yes to change the definition of the object. Otherwise, click No.

SAT2018W **Are you sure that you want to change the definition of object *object*?**

Explanation: Changing the definition of an object will cause all other objects that reference that object to be affected.

User response: Click Yes to change the definition of the object. Otherwise, click No.

SAT2019W **An object with the name *object* already exists in the database *database*. Would you like to rename *object* before it is created in *database*?**

Explanation: While creating *object* in *database*, another object with the same name was found to already exist. Objects must have a unique name.

User response: Click Yes to save the object under a different name. Otherwise, click No. The object will not be created.

SAT2020W At least one of the selected satellites is currently online. No action is taken on satellites that are online.

Explanation: When a satellite is online and synchronizing, it cannot be altered.

User response: Determine which satellites are currently online. Wait until these satellites are offline, and try the request again.

SAT2021W A test connection or attachment to the selected target was not successful using the specified authentication credentials.

Explanation: An attempt was made to connect or attach to a target. The connection could not be established because of an authentication error.

User response: Verify that the authentication credentials are correct for the target, and try the request again.

SAT2022W You will be able to make modifications to this satellite when it is a test satellite. Care should be taken to ensure the satellite's state remains consistent.

Explanation: Setting a satellite as a test satellite means that it will execute test batches. A test batch may contain batch steps that produce uncertain outcomes, and may put the satellite into an inconsistent state. A satellite should only be set as a test satellite when it is acceptable to subject it to these risks.

User response: Click OK to set the satellite as a test satellite. Otherwise, click Cancel.

SAT2023W Are you sure that you want to change this test satellite to a production satellite? The administrative function available for this satellite will be severely restricted.

Explanation: A test satellite is used to try out configuration changes before these changes are made available to the production satellites. Consequently, a test satellite may become inconsistent when configuration changes do not succeed. To repair the test satellite, return it to a consistent state so that the configuration changes can be iteratively tested. More administrative actions can be performed against a test satellite. If the satellite is changed to a production satellite, it can no longer be used for test purposes. The administrative function available for this satellite will necessarily be severely restricted.

User response: Click OK to set the satellite as a production satellite. Otherwise, click Cancel.

SAT2024W Synchronization session is in progress. Are you sure that you want to terminate it?

Explanation: A synchronization session was in progress when the terminate action was issued.

User response: If you continue with the termination, the synchronization session will be stopped before it completes. You should cancel the termination, and allow the synchronization session to complete.

SAT2025W If you use anything other than a noncomplete consistent-change-table in a consolidation scenario, the target table is refreshed when satellites synchronize.

Explanation: In a consolidation replication scenario, a noncomplete consistent-change-data-table does not result in a full refresh of the target table when a satellite synchronizes. Any other type of table (for example, a user copy or a point-in-time target table) results in a full refresh of the target table when a satellite synchronizes. If, when defining your replication scenario, you specified that the data capture is full refresh, the target table will be fully refreshed each time that a satellite synchronizes. If you did not specify that the data capture is full refresh, the target table is only fully refreshed the first time that the Apply program is invoked on the satellite.

User response: If you want the target table to be refreshed when a satellite synchronizes, no action is required. If you want to change the type of the target table, you must remove the effects of the generalization of the replication subscriptions from the DB2 control server, and return to the Control Center to modify the replication scenario. Then, generalize the replication subscriptions again.

SAT2026W The satellite should be synchronized once before the user changes user data. User data that is changed before the initial synchronization of the satellite cannot be replicated.

Explanation: When the satellite synchronizes, the Capture and Apply programs are invoked on the satellite. Depending on how you defined the replication scenario, changes that are made to the user data in the satellite database before the first synchronization session either cannot be replicated to the corporate sources, or will be overwritten at the satellite:

- If your replication scenario is either consolidation, or update anywhere with the replication direction defined from the satellite to the corporate source, changes that are made to the user data before the first synchronization session cannot be replicated from the satellite to the corporate source. When the Capture program is invoked, it cannot capture these changes.

- If your replication scenario is either distribution, or update anywhere with the replication direction defined from the corporate source to the satellite, changes that are made to the user data before the first synchronization session will be overwritten when a full refresh occurs at the satellite.

User response: Instruct all satellite users to synchronize before they change data on the satellite.

Chapter 110. SAT3000 - SAT3499

SAT3000N The name *name* already exists.

Explanation: The name of the object you are attempting to create, *name*, already exists.

User response: Specify a unique name.

SAT3001N The target name of a copy, *targetname*, already exists in the database *cdb*.

Explanation: The name provided for the copy, *targetname*, is not unique in the target satellite control database *cdb*.

User response: Specify a unique name.

SAT3002N Application version *application-version-name* already exists for the group *group-name*.

Explanation: The name supplied for this application version is already being used for this group.

User response: Specify a name for the application version that is not already in use by the group *group name*.

SAT3003N No application versions exist for this group.

Explanation: The requested action against the group requires at least one application version.

User response: Create an application version for the group.

SAT3004N The object *name* does not exist.

Explanation: The specified object, *name*, does not exist. It may have been removed after the view containing the object was filled.

User response: Refresh the view in which the object appears.

SAT3005N The object *name* cannot be deleted because it is currently referred to by another object.

Explanation: The object *name* cannot be deleted because another object references it. Consistency cannot be maintained if *name* is deleted.

User response: Delete all objects that are dependent on the object before deleting this object.

SAT3006N Group *group-name* cannot be deleted because it has at least one enabled satellite.

Explanation: Before a group can be deleted, all of its satellites be disabled so that they can be deleted along with the group. At least one satellite in the group was found to be enabled.

User response: Disable all the satellites in this group.

SAT3007N Group *group-name* cannot be deleted because it has at least one application version that has either a test-level or a production-level batch.

Explanation: Production-level and test-level batches are considered to be active levels that are used by satellites. Consequently, they cannot be deleted. Therefore, the application version, and hence the group, to which they belong cannot be deleted.

User response: Obsolete all the production-level batches and remove all the test-level batches for this group. Then try the request again.

SAT3008N Application version *application-version-name* cannot be deleted because it has either a test-level or a production-level batch.

Explanation: Production-level and test-level batches are considered to be active levels that are used by satellites. Consequently, they cannot be deleted. Therefore, the application version to which they belong cannot be deleted.

User response: Obsolete the production-level batches and remove the test-level batches for this application version. Then try the request again.

SAT3009N No generalized replication subscription sets exist for the group *group-name*.

Explanation: No generalized replication subscription sets exist for the group *group-name*. The specified action requires at least one generalized replication subscription to exist.

User response: Define one or more generalized replication subscription sets for this group.

SAT3010N No application versions exist for the group *group-name*.

Explanation: No application versions exist for the group *group-name*. The specified action requires at least

one application version to exist.

User response: Define an application version for this group.

SAT3011N **No default replication subscription sets exist for application version *application-version-name* in group *group-name*.**

Explanation: Default replication subscription sets for application version *application-version-name*, in group *group-name*, are generated by the generalize action. Such subscription sets are necessary before it is possible to customize the generalized replication subscription.

User response: Ensure the generalized action completed successfully, and try the request again.

SAT3012N **A default replication control server for application version *application-version-name*, in group *group-name*, is generated by the generalize replication subscription action. Such a control server is required before the generalized replication subscription can be customized.**

Explanation: A default replication control server for application version *application-version-name* in group *group-name* does not exist.

User response: Ensure that the generalizing of the replication subscription completed successfully. Then try the request again.

SAT3013N **The confirmation password does not match the specified password.**

Explanation: The confirmation password must exactly match the specified password. Passwords are case sensitive.

User response: Type the confirmation password again. Ensure that all characters typed are in the same case as the original password.

SAT3014N **At least one of the specified batch steps does not have a target alias, a success code set, or both.**

Explanation: When a script is imported or added as a new batch step, a target alias and a success code set must be added for the new batch step is not provided.

User response: Identify which batch steps are missing a target alias or a success code set, and add the missing information.

SAT3015N **A target alias has not been selected. You must select a target alias before you can specify the authentication credentials.**

Explanation: A target alias must be specified before an authentication can be associated with it.

User response: Specify a target alias. Then try the request again.

SAT3016N **An application version has not been set for this satellite. This page will be disabled until the satellite reports its application version.**

Explanation: A satellite executes the batches of an application version. This satellite has not yet reported its application version. Therefore, batch steps cannot be specified for it.

User response: Assign the satellite an application version on the satellite itself. You cannot perform this action from the Satellite Administration Center.

SAT3017N **The action cannot be performed. An application version has not been set.**

Explanation: A satellite executes the batches of an application version. This satellite has not yet reported its application version. Therefore, the specified action cannot be performed.

User response: Assign the satellite an application version on the satellite itself. You cannot perform this action from the Satellite Administration Center.

SAT3018N **The type of the execution target specified is not correct for the type of the script.**

Explanation: A target alias has been selected that is not correct. The script is designed to execute against a different type of target.

User response: Choose a target alias that is of the same type as the selected script. Then try the request again.

SAT3019N **The specified success code ranges are in conflict.**

Explanation: A success code set cannot contain a greater than (>) relation that specifies a range of numbers that overlap the range of numbers that is specified by a less than (<) relation. For example, the following two relations are not allowed in the same success code set: > 10, < 11.

User response: Either remove or correct those success relations from the success code set that cause the overlap in ranges. To specify a finite range of numbers between two numbers, a and b, an equals (=) relation must be specified for each number in the range. For

example, to specify a finite range of numbers between 5 and 10 as success codes, the following equals relations are required: =5,=6,=7,=8, =9 and =10.

SAT3022N Promotion failed with the following error: `SQLCODE=sqlcode`, `SQLSTATE=sqlstate`, tokens: `token1`, `token2`, `token3`. The error was detected at location `location`.

Explanation: Promotion failed in an unexpected way.

User response: Contact the help desk or your system administrator.

SAT3023N A system or internal error occurred.

Explanation: The possible errors include a system error returned by the operating system, insufficient resources, or an incorrect configuration.

User response: Verify your replication configuration, or contact your system administrator or IBM Service.

SAT3024N A replication server is at a level that is not supported.

Explanation: The product level of a replication server is not supported.

User response: Upgrade to a supported level.

SAT3025N An authentication failure occurred when connecting to a replication server.

Explanation: This action requires authentication with the replication servers. Authentication with one of the servers failed.

User response: Ensure that you supplied the correct user ID and password. Either try the request again, or contact your system administrator.

SAT3026N The size of a generated script exceeds the available storage.

Explanation: This action generates scripts. At least one of the generated scripts is larger than the maximum available storage.

User response: Contact your system administrator or IBM Service.

SAT3027N The database directory entry for a replication source server, control server, or target server does not exist at the instance `instance-name`.

Explanation: This action requires the existence of a database directory entry for the replication source servers, control server, and target server at the instance `instance-name`. At least one of the directory entries could not be found.

User response: Catalog the replication source servers, control server, and target server at `instance-name`. Or contact your system administrator or IBM Service.

SAT3028N The fix batch `batch` cannot be deleted because it is being used by at least one satellite.

Explanation: A batch cannot be deleted if it is currently being used.

User response: Wait until the batch is no longer used by any satellite. Then try the request again.

SAT3029N The script `script` cannot be changed because it is being used by at least one production or obsolete batch step.

Explanation: A script can only be modified if it is not being used by a production or obsolete batch step.

User response: Ensure that the script is not currently being used by a production or obsolete batch step. Then try the request again.

SAT3030N The specified obsolete batch step cannot be deleted because it is being used by a batch that is not obsolete.

Explanation: An obsolete batch step can only be deleted if it is not being used by any batch that is not obsolete.

User response: If the batch that uses this batch step is no longer in use by any satellite, move it to the obsolete level. Then try the request again.

SAT3031N The batch step cannot be deleted because it is not the last step in the set of batch steps.

Explanation: A batch step cannot be deleted from the middle of a set of batch steps. All batch steps after it must first be deleted.

User response: Delete all the batch steps that follow the batch step to be deleted. Then try the request again.

SAT3032N A production batch step cannot be deleted.

Explanation: Production batch steps are used by production satellites for synchronization and cannot be deleted. Only test, unassigned, and obsolete batch steps can be deleted.

User response: Take the batch step out of production. Then try the request again.

SAT3033N **Satellite synchronization failed because the satellite control database SATCTLDB cannot be found in the database directory.**

Explanation: The satellite control database is not cataloged correctly in the database directory.

User response: Contact the help desk or your system administrator.

SAT3034N **Satellite synchronization failed because of an authentication error at the satellite control server.**

Explanation: Either the user ID or the password sent to the control server is not correct.

User response: Contact the help desk or your system administrator.

SAT3035N **Satellite synchronization failed because the satellite ID could not be found locally.**

Explanation: This error may result from bypassing the operating system logon.

User response: If you are already logged on to the operating system, contact the help desk or your system administrator.

SAT3036N **Satellite synchronization failed because the satellite control server does not recognize this satellite.**

Explanation: Either the satellite ID is not defined correctly on the satellite or this satellite is not defined correctly at the satellite control server.

User response: Contact the help desk or your system administrator.

SAT3037N **Satellite synchronization failed because of an unknown error.**

Explanation: Unknown.

User response: Contact the help desk or your system administrator.

SAT3038N **The synchronizer has detected a severe DB2 error. Synchronization cannot continue.**

Explanation: The cause of this severe error is not known.

User response: Contact the help desk or your system administrator.

SAT3039N **Satellite synchronization failed. The satellite is in fail state at the satellite control server.**

Explanation: This satellite is in fail state, and must be fixed.

User response: Contact the help desk or your system administrator.

SAT3040N **Satellite synchronization failed because the satellite control server cannot be contacted.**

Explanation: Communications could not be established between the satellite and its control server. One possible reason is that the control server is not operational, or communications configuration is not correct (for example, the host name or port number for the satellite control server is not correct).

User response: Contact the help desk or your system administrator.

SAT3041N **Satellite synchronization failed because communications with the satellite control server was lost.**

Explanation: A communication link failure occurred. Either the satellite control server became inactive, or a network failure was experienced.

User response: Contact the help desk or your system administrator.

SAT3042N **The satellite synchronization session did not complete all tasks successfully.**

Explanation: An error occurred during the satellite synchronization session.

User response: Contact the help desk or your system administrator.

SAT3043N **The satellite synchronization session did not start successfully.**

Explanation: An error occurred during the start-up phase of the satellite synchronization session.

User response: Contact the help desk or your system administrator.

SAT3044N **The cleanup phase of the satellite synchronization session did not complete successfully.**

Explanation: An error occurred during the cleanup phase of the satellite synchronization session.

User response: Contact the help desk or your system administrator.

SAT3045N *control-server* is not a valid replication control server.

Explanation: The selected server is not a valid replication control server. The request cannot be completed.

User response: Select a valid replication control server, then try the request again.

SAT3046N No generalized replication subscription sets exist for application version *application-version-name* in group *group-name*.

Explanation: There are no generalized replication subscription sets for the application version *application-version-name* in the group *group-name*. A generalized replication subscription set must exist before it can be customized.

User response: Generalize one or more replication subscription sets for the application version, and try the request again.

SAT3047N The input for field *field-name* exceeds the limit. The maximum allowable length for this field is *maximum-length*.

Explanation: The length of the input in field *field-name* exceeds the limit *maximum-length*.

User response: Ensure that the input is within the maximum limit, then try the request again.

SAT3048N More than one target server is associated with apply qualifier *apply-qualifier* at control server *control-server*.

Explanation: A target server of a replication subscription represents a satellite in the group that is executing the specified application version. Consequently, there can only be one target server for a given apply qualifier. If your replication configuration specifies a single target server, then the associated apply qualifier may not have been correctly specified. Otherwise, your replication configuration may not be acceptable for generalization to the satellite environment.

User response: First verify that the replication configuration does not specify more than one target server. If it does, correct the replication configuration, then try the generalize replication subscription request again. If the replication configuration does not specify more than one target server, ensure that the correct apply qualifier is specified for the generalize replication subscription request. Then try the request again.

SAT3049N The file name is not valid.

Explanation: The file name is not valid because it contains at least one of the following characters: \ / : * ? " < > |

User response: Ensure that the file name does not contain any of the listed characters. Then try the request again.

SAT3050N One or more of the selected satellites does not have an application version.

Explanation: You cannot change the replication parameters for a satellite that does not have an application version.

User response: Ensure that every selected satellite has an application version. Then try the request again.

SAT3051N The generalization of the replication subscriptions failed. At least one of the generated scripts is empty.

Explanation: At least one of the scripts generated are empty.

User response: Try the generalization again. If the problem persists, contact IBM Service.

SAT3052N The database alias or database name *name* cannot be found.

Explanation: The database name or the alias that is specified in the command cannot be found. Either the database is not cataloged in the database directories, or does not exist.

User response: Ensure that the specified database name exists in the system database directory, and verify that the database exists. Then try the request again.

Chapter 111. SAT4000 - SAT4499

SAT4000C *cdb* is not a valid satellite control database.

Explanation: The database named *cdb* is not a valid satellite control database. Any database that is named SATCTLDB is initially assumed to be a satellite control database. Any action pertaining to satellite administration that is performed against such a database will first verify that the database is valid. The action will not succeed if the database is not configured correctly. If *cdb* is a satellite control database, then its configuration has been altered such that it can no longer be used. Otherwise, a non-control database has been given the name SATCTLDB.

User response: To avoid misrepresenting a database as a satellite control database, rename the database to anything other than SATCTLDB. If *cdb* is a satellite control database, verify its configuration with the original definition.

Part 21. SPM Messages

This section contains the sync point manager (SPM) messages. The messages are listed in numeric sequence.

Chapter 112. SPM0000 - SPM0499

SPM0400 **Indoubt transaction resolution with DBALIAS=*dbalias* shows heuristic damage - the database rolled back the UOW and the coordinator with LUNAME=*luname* committed. The transaction is identified by LUWID=*luwid*.**

Explanation: DB2 was the DRDA2 AS for the transaction identified by *luwid*. Resolution with the DB2 database shows heuristic damage occurred. The database identified by *dbalias* manually resolved the indoubt transaction. The DB2 database at *luname* rolled back the transaction. This is inconsistent with the commit decision of the DRDA2 coordinator at *luname*.

- The application was told that the transaction was committed.
- Any resources updated at the participant were rolled back.

dbalias The alias name of the DB2 database.

luname The LU name of the coordinator.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Indoubt resolution with the participant completes.

User response: Call the database administrator.

Database Administrator Action: Inform the database administrators at both the COORDINATOR *luname* and at the DATABASE *dbalias* that heuristic damage occurred for the transaction with *luwid*. DB2 was a DRDA2 AS for the transaction. The DRDA2 AR at COORDINATOR *luname* made the decision to COMMIT the database updates made by *luwid*. The *dbalias* PARTICIPANT made a heuristic decision to ROLL BACK the updates made by *luwid*.

SPM0402 **Indoubt transaction resolution with participant with LUNAME=*luname* and DBALIAS=*dbalias* shows heuristic damage - the participant committed and DB2 rolled back. The transaction is identified by LUWID=*luwid*.**

Explanation: DB2 has coordinator responsibility for the transaction identified by *luwid*. Resolution with the participant shows heuristic damage occurred. The participant identified by *luname* and *dbalias* manually resolved the indoubt transaction. The action taken was

to commit the transaction. This is inconsistent with the roll back decision of the coordinator.

- The application was told that the unit of work was rolled back.
- Any resources updated at the participant were committed.

dbalias The database alias of the participant.

luname The LU name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Indoubt resolution with the participant completes.

User response: Call the database administrator.

Database Administrator Action: Inform the local database administrator and the database administrator at *luname* that heuristic damage occurred for the transaction with *luwid*. DB2 was the coordinator for the transaction and made the decision to roll back the database updates made by *luwid*. At *luname* a heuristic decision was made to COMMIT the updates made by *luwid*.

SPM0403 **Indoubt transaction resolution with participant with LUNAME=*luname* and DBALIAS=*dbalias* shows heuristic damage - the participant rolled back and DB2 committed. The transaction is identified by LUWID=*luwid*.**

Explanation: DB2 has coordinator responsibility for the transaction identified by *luwid*. Resolution with the participant shows heuristic damage occurred. The participant identified by *luname* and *dbalias* manually resolved the indoubt transaction. The action taken was to roll back the UOW. This is inconsistent with the commit decision of the coordinator.

- The application was told that the transaction committed.
- Any resources updated at the participant were rolled back.

luname The LU name of the participant.

dbalias The data base alias of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Indoubt resolution with the participant completes.

User response: Call the database administrator.

Database Administrator Action: Inform the local database administrator and the database administrator at *luname* that heuristic damage occurred for the transaction with *luwid*. DB2 was the coordinator for the transaction and made the decision to COMMIT the database updates made by *luwid*. At *luname* a heuristic decision was made to ROLL BACK the updates made by *luwid*.

SPM0404 Protocol error during indoubt transaction resolution with coordinator with LUNAME=*luname1* - the DB2 database with LUNAME=*luname2* has an indoubt transaction which is identified by LUWID=*luwid*.

Explanation: DB2 was the DRDA2 AS for the transaction identified by *luwid*. The DB2 transaction associated at the database with LUNAME=*luname2* is indoubt. A protocol error occurred during indoubt resolution with the coordinator identified by *luname1*.

luname1

The LU name of the partner acting as the coordinator.

luname2

The LU name of the database where the transaction is indoubt.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

The indoubt transaction remains indoubt. A Resync Protocol Violation trace record is written.

User response: Call the database administrator.

Database Administrator Action: DB2 does not attempt to automatically resolve the indoubt transaction. The transaction must be manually resolved. The commit or abort decision made at the coordinator must be determined so that the same decision can be made at this participant DB2.

Contact the database administrator at the coordinator with *luname* and *dbalias*, to determine whether the transaction committed or aborted.

Use the LIST INDOUBT TRANSACTIONS command at this (the participant) *dbalias* to resolve the indoubt transaction.

SPM0406 A transaction with LUWID=*luwid* at the participant with LUNAME=*luname* and DBALIAS=*dbalias* may be indoubt because of a communication failure. DB2 committed.

Explanation: During execution of the two phase commit protocol with the participant at *luname*, a communication failure occurred. Phase 1 of the protocol completed and the transaction is prepared for either commit or abort. The decision to commit the transaction was made, but cannot be communicated to the participant at this time. The participant is indoubt.

DB2 becomes responsible for indoubt resolution with the participant. This responsibility appears in the LIST DRDA INDOUBTS TRANSACTION report. Periodic attempts are made to reestablish communication with the participant for automatic resolution.

luname

The LU unit name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Periodic attempts will be made to automatically resolve the indoubt transaction at the participant.

User response: Determine the cause of the communication failure and have the problem corrected. DB2 periodically attempts to reestablish communication for automatic resolution. If automatic resolution does not occur in a reasonable amount of time, call the database administrator. Manual resolution of the transaction might be necessary at the participant to release locked resources.

Database Administrator Action: If manual resolution is necessary, inform the database administrator at the participant that the decision is commit.

SPM0407 Automatic resolution of the transaction with LUWID=*luwid* with the coordinator at LUNAME=*luname* resulted in commit. The DB2 Database is = *dbname*.

Explanation: The indoubt transaction at the database identified by *dbname* was automatically resolved by communication with the coordinator identified by *luname*. The transaction was committed.

luname1

The LU name of the coordinator.

luwid The SNA logical unit of work id of the transaction.

luname2

The LU name of the database where the changes were made.

The XID associated with the unit of work is printed following this message.

Processing continues normally.

SPM0408 **A communication failure occurred during automatic resolution with partner with LUNAME=*luname*. Communication protocol being used=*protocol*. Communication API being used=*api*. Communication function detecting the error=*function*. Protocol specific error codes: *rc1*, *rc2*, *rc3*.**

Explanation: One or more indoubt transactions exist with *luname*. DB2 attempted to automatically resolve the indoubt transaction but a communication error occurred.

protocol

The communication protocol being used. Only "APPC" is supported.

api

The communication api being used. Either "CM2 APPC" or "SNA/6000 OS Subroutines" may be specified.

codes

For "CM2 APPC" api, the *rc1*, *rc2*, and *rc3* contain the primary return code, secondary return code and possibly an SNA sense code. For "SNA/6000 OS Subroutines" api, *rc1* contains the *errno* global variable.

User response: Determine the cause of the communication failure and have the problem corrected. DB2 periodically attempts to reestablish communication for automatic resolution. If automatic resolution does not occur in a reasonable amount of time, call the database administrator. Manual resolution of the transaction might be necessary at the participant to release locked resources.

Database Administrator Action: If manual resolution is necessary, inform the database administrator at the participant that the decision is commit.

SPM0409 **A transaction with LUWID=*luwid* cannot be resolved due to a cold start with LUNAME=*luname*. DB2 transaction status=*status*. DB2 responsibility=*responsibility*.**

Explanation: An indoubt transaction exists with the partner at *luname*. DB2 is unable to resolve the indoubt transaction because the partner has lost all knowledge of indoubt transactions due to a previous cold start.

luwid The SNA logical unit of work id of the transaction.

status The commit status of the transaction as known by DB2. The commit status may be indoubt, committed, or rolled back.

responsibility

DB2 transaction responsibility. DB2 may be the coordinator or may be the participant.

User response: There is probably inconsistent data at the coordinator and participant. Inform database administrator of the status of the transaction.

Database Administrator Action: Manual resolution is necessary. The heuristic decision (that is, to commit or roll back the transaction) should be coordinated with any other participants and/or the coordinator. The existence of other participants might not be easy to determine. The information might be available in the coordinators recovery log even though the coordinator performed a cold start.

The commit or abort decision provided using the LIST INDOUBT TRANSACTIONS command for the transaction are propagated to all downstream participants, if any.

SPM0410 **Warm start connection by partner with LUNAME=*luname* rejected. Partner changed at least 1 of - our log name *oldourname*(*newourname*), their log name *oldtheirname*(*newtheirname*), sync point protocol *oldpa*(*newpa*), flag byte sent *oldfb*(*newfb*), ccluname sent *oldccls*(*newccls*), and indoubt transactions require resolution.**

Explanation: An attempt to make a warm start connection with a partner was rejected because the partner specified a different set of sync point parameters than the ones that were in use when communications were lost. DB2 has knowledge of indoubt transactions that involve the partner as either the coordinator or a participant. This error might be a recoverable error if the partner can restart with the original sync point parameters. If this is not possible, then the partner must perform a cold start connection with DB2.

luname

The LU name of the partner.

oldourname

The log name for this DB2 subsystem. This value is null if the partner did not initiate the log name exchange.

newourname

The log name for this DB2 subsystem, as remembered by our partner and sent to us in the attempted warm start connection. This value is null if the partner did not initiate the log name exchange.

oldtheirname

The log name used by the partner during the last connection, as remembered by DB2.

newtheirname

The log name used by the partner in the attempted warm start connection.

oldpa

The sync point protocol used during the last connection, as remembered by the DB2 (PA - Presumed Abort, PN - Presumed Nothing).

newpa

The sync point protocol used during the last connection, as remembered by the partner (PA - Presumed Abort, PN - Presumed Nothing).

oldfb

The use of the PS header byte 2 during the last connection, as remembered by DB2 (F - used as flag byte, N - not used as flag byte).

newfb

The use of the PS header byte 2 during the last connection, as remembered by the partner (F - used as flag byte, N - not used as flag byte).

oldccls

Whether the luname of the conversation correlator was exchanged as part of the sync point protocol during the last connection, as remembered by DB2. (E - was exchanged, N - was not exchanged).

newccls

Whether the luname of the conversation correlator was exchanged as part of the sync point protocol during the last connection, as remembered by the partner (E - was exchanged, N - was not exchanged).

The connection with the partner is rejected. DB2 retains indoubt knowledge.

User response: Call the database administrator.

Database Administrator Action: Contact the database administrator at the partner *luname* and determine if it is possible for the partner to perform a warm start with same sync point parameters as ours ('oldourname', 'oldtheirname', 'oldpa', 'oldfb', 'oldccls'). If this is possible, the next attempt to connect will succeed.

If this cannot be done, then there are two other possible solutions:

- Have the partner *luname* perform a cold start connection with this DB2 subsystem.
- Force the next connection with the partner to be a cold start connection.

SPM0411 Cold start connection by coordinator with LUNAME=*luname* accepted. Indoubt transactions need manual resolution.

Explanation: DB2 was the DRDA2 AS and has participant responsibility for indoubt transactions. The coordinator informed DB2 that it performed a cold start operation and lost all knowledge of indoubt

transactions. The indoubt transactions at this DB2 must be manually resolved with the LIST INDOUBT TRANSACTIONS command.

dbalias The database alias name of the coordinator.

The connection with the partner is accepted. A trace record is written.

User response: Call the database administrator.

Database Administrator Action: DB2 is a participant with one or more indoubt transactions where the coordinator is *luname*. The DBMS at *luname* performed a cold start. The DB2 participant assumes that the coordinator recovery log was lost or damaged and indoubt transaction resolution cannot be achieved. There is probably inconsistent data at the coordinator.

The heuristic decision (that is, to commit or abort the transaction should be coordinated with any other participants. The existence of other participants might not be easy to determine. The information might be available in the coordinators recovery log even though the coordinator performed a cold start.

The commit or abort decision provided using the LIST INDOUBT TRANSACTIONS command for the transaction are propagated to all downstream participants, if any.

SPM0413 Protocol error detected in sync point communications with participant with LUNAME=*luname* and DBALIAS=*dbalias*. The transaction with LUWID=*luwid* may be indoubt at the participant. DB2 committed.

Explanation: DB2 is the coordinator of the transaction. A protocol error occurred during the SNA sync point exchange with the participant identified by *luname* and *dbalias*. The protocol error fits into one of the following categories:

- An unknown SNA PS header was received.
- An SNA PS header was received in the wrong sequence.
- An invalid LU6.2 verb sequence was used to deliver the SNA PS header.
- A DEALLOCATE TYPE(ABEND_SVC) was received from the participant, indicating that the participant detected a protocol violation in the PS headers transmitted by DB2.

luname

The LU name of the participant.

dbalias

The database alias name of the participant.

luwid

The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

If application was told that the transaction committed.

There may be an indoubt transaction at the participant and if so, the indoubt transaction must be manually resolved. DB2 does not automatically resolve the indoubt transaction because of the protocol error.

A Syncpoint Protocol Violation trace record is written.

User response: Call the database administrator. The participant might need to manually resolve the indoubt transaction.

SPM0414 **Protocol error during indoubt transaction resolution with participant with LUNAME=*luname* and DBALIAS=*dbalias*. The transaction with LUWID=*luwid* may be indoubt at the participant. DB2 rolled back.**

Explanation: DB2 has coordinator responsibility for the transaction which was rolled back. A protocol error occurred during indoubt resolution with the participant identified by *luname* and *dbalias*.

The transaction at the participant remains indoubt. DB2 will not attempt to automatically resolve the indoubt transaction because of the protocol violation.

luname

The LU name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

A Resync Protocol Violation trace record is written.

User response: Call the database administrator. The participant might need to manually resolve the indoubt transaction.

Database Administrator Action: If the transaction is indoubt at the participant, it must be manually (heuristically) resolved.

SPM0415 **Automatic resolution of the transaction with LUWID=*luwid* with the coordinator at LUNAME=*luname* resulted in roll back. The DB2 Database is = *dbname*.**

Explanation: The indoubt transaction at the database identified by *dbname* was automatically resolved by communication with the coordinator identified by *luname*. The transaction was rolled back.

luname1

The LU name of the coordinator.

luwid The SNA logical unit of work id of the transaction.

luname2

The LU name of the database where the changes were made.

The XID associated with the unit of work is printed following this message.

Processing continues normally.

SPM0416 **Cold start connection rejected by partner with LUNAME *luname*.**

Explanation: DB2 attempted to make a cold-start connection with a partner *dbalias*. The partner rejected this attempted connection.

luname

The LU name of the coordinator.

The connection was not made.

User response: Call the database administrator.

Database Administrator Action: DB2 is not able to connect the partner *luname* until the partner *luname* allows a cold-start connection with DB2. Contact the database administrator at the partner *luname*.

Contact your IBM Support Center for further assistance.

SPM0417 **Protocol error detected in sync point communications with participant with LUNAME=*luname* and DBALIAS=*dbalias*. The transaction with LUWID=*luwid* may be indoubt at the participant. DB2 rolled back.**

Explanation: DB2 is the coordinator of the transaction. A protocol error occurred during the SNA sync point exchange with the participant identified by *luname* and *dbalias*. The protocol error fits into one of the following categories:

- An unknown SNA PS header was received.
- An SNA PS header was received in the wrong sequence.
- An invalid LU6.2 verb sequence was used to deliver the SNA PS header.
- A DEALLOCATE TYPE(ABEND_SVC) was received from the participant, indicating that the participant detected a protocol violation in the PS headers transmitted by DB2.

luname

The LU name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

If application was told that the transaction rolled back.

There may be an indoubt transaction at the participant and if so, the indoubt transaction must be manually resolved. DB2 does not automatically resolve the indoubt transaction because of the protocol error.

A Syncpoint Protocol Violation trace record is written.

User response: Call the database administrator. The participant might need to manually resolve the indoubt transaction.

SPM0420 Cold start connection by participant with LUNAME=*luname* accepted. Possible damage.

Explanation: DB2 has coordinator responsibility for indoubt transactions at a participant and just connected with the participant, which lost all knowledge of indoubt transactions because of a previous cold start. There might be damage at the participant.

luname

The LU name of the participant where there might be damage.

The connection with the partner is accepted.

User response: Call the database administrator.

Database Administrator Action: DB2 is the coordinator with indoubt transaction resolution responsibility for one or more indoubt units of work at *luname*. The DBMS at *luname* performed a cold start. DB2 assumes that the participant recovery log was lost or damaged and indoubt transaction resolution cannot be achieved. There is probably inconsistent data at the participant. Minimally, the participant might not completely reflect the final outcome of the transactions that were indoubt at the time the failure occurred.

SPM0421 SNA XLN protocol violation by partner with LUNAME=*luname*.

Explanation: DB2 detected a protocol violation in the SNA Exchange Log Names (XLN) exchange with the partner at the specified *luname*.

luname

The LU name of the partner that sent the invalid XLN message.

The attempt to connect with the remote site fails. An XLN Protocol Violation trace record is written.

User response: Contact the system programmer for the remote site. The invalid XLN message is recorded in the trace record. The system logic error that causes the invalid XLN message must be corrected at the remote site.

SPM0422 Warm start connection by partner with LUNAME=*luname* rejected because the partner remembers our log name incorrectly. Our log name is *name1* and the partner remembers it as *name2*.

Explanation: An attempt to make a warm start connection with a partner was rejected because the partner specified our log name as *name2*. Our log name is *name1*, which is the luname of the local DB2. This error might be a recoverable error if the partner can restart with our log name as *name1*. If this is not possible, then the partner must perform a cold start connection with DB2.

luname

The LU name of the partner with which the connection failed.

name1 Our log name, as remembered by us.

name2 Our log name, as remembered by our partner and sent to us in the attempted warm start connection.

The connection with the partner is rejected.

User response: Call the database administrator.

Database Administrator Action: Contact the database administrator at the partner *luname* and determine if it is possible for the partner to perform a warm start with our log name specified as the luname of this DB2. If this is possible, the next attempt to connect will succeed. Or have the partner *luname* perform a cold start connection with DB2.

SPM0423 Automatic resolution of the transaction with LUWID=*luwid* with the partner at LUNAME=*luname* and DBALIAS=*dbalias* resulted in commit.

Explanation: The indoubt unit of work was automatically resolved by communication with the participant. The participant has been notified of the commit decision.

luname

The LU name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Processing continues normally.

SPM0424 **Automatic resolution of the transaction with LUWID=*luwid* with the participant at LUNAME=*luname* and DBALIAS=*dbalias* resulted in roll back.**

Explanation: The indoubt unit of work was automatically resolved by communication with the participant. The participant has been notified of the roll back decision.

luname

The LU name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Processing continues normally.

SPM0425 **A transaction with LUWID=*luwid* at the participant with LUNAME=*luname* and DBALIAS=*dbalias* may be indoubt because of a communication failure. DB2 rolled back.**

Explanation: During execution of the two phase commit protocol with the participant at *luname*, a communication failure occurred. Phase 1 of the protocol completed and the transaction is prepared for either commit or abort. The decision to roll back the transaction was made, but cannot be communicated to the participant at this time. The participant is indoubt.

DB2 becomes responsible for indoubt resolution with the participant. This responsibility appears in the LIST DRDA INDOUBTS TRANSACTION report. Periodic attempts are made to reestablish communication with the participant for automatic resolution.

luname

The LU unit name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

Periodic attempts will be made to automatically resolve the indoubt transaction at the participant.

User response: Determine the cause of the communication failure and have the problem corrected. DB2 periodically attempts to reestablish communication for automatic resolution. If automatic resolution does not occur in a reasonable amount of time, call the database administrator. Manual resolution of the

transaction might be necessary at the participant to release locked resources.

Database Administrator Action: If manual resolution is necessary, inform the database administrator at the participant that the decision is roll back.

SPM0426 **Protocol error detected during indoubt transaction resolution with participant at LUNAME=*luname* and DBALIAS=*dbalias*. The transaction with LUWID=*luwid* may be indoubt at the participant. DB2 committed.**

Explanation: DB2 has coordinator responsibility for the transaction which was committed. A protocol error occurred during indoubt resolution with the participant identified by *luname* and *dbalias*.

The transaction at the participant remains indoubt. DB2 will not attempt to automatically resolve the indoubt transaction because of the protocol violation.

luname

The LU name of the participant.

dbalias The database alias name of the participant.

luwid The SNA logical unit of work id of the transaction.

The XID associated with the unit of work is printed following this message.

A Resync Protocol Violation trace record is written.

User response: Call the database administrator. The participant might need to manually resolve the indoubt transaction.

Database Administrator Action: If the transaction is indoubt at the participant, it must be manually (heuristically) resolved.

SPM0438 **The Sync point manager recovery log is bad.**

Explanation: The sync point manager recovery log is inconsistent and cannot be used to perform recovery during DB2 start up processing.

User response: Indoubt transactions may exist at DRDA2 application servers. These indoubt transactions must be recovered manually.

Call the database administrator.

Database Administrator Action: To start the sync point manager, erase the spmlog directory and start DB2. This will cause DB2 to create new sync point log files and to establish cold start connections with all DRDA2 application servers.

SPM0439 **Sync point manager unrecoverable error while attempting to write to the sync point manager recovery log.**

Explanation: The sync point manager log is inconsistent and cannot be used. An unrecoverable error has been detected while attempting to write to the SPM log during DB2 processing.

User response: The sync point manager will not allow any new synclevel(twophase) connections. Issue the LIST DRDA INDOUBT TRANSACTIONS command to determine the status of any indoubt transactions.

Call the database administrator.

Database Administrator Action: To start the sync point manager, erase the spmlog directory and start DB2. This will cause DB2 to create new sync point log files and to establish cold start connections with all DRDA2 application servers.

SPM0440E **Error encountered while trying to start protocol support. Return code from function was rc. The most probable cause for this error is that SNA has not been started. Please stop DB2, start SNA, and restart DB2.**

SPM0448E **Error encountered while trying to start the sync point manager protocol support. The sync point manager failed to register LUNAME *lname* for sync point support since this LU has been configured for SNA API client use. Either choose a different LU for the sync point manager or disable the SNA API client use in the Local LU 6.2 definition for this LU.**

Explanation: This error occurs when the customer is trying to start the sync point manager using CS/NT V5.01 and is using a Local LU 6.2 definition where the SNA API client use flag has been set.

User response: Either choose a different local LU 6.2 (without the SNA API client use configured) or disable the SNA API client use flag for the Local LU 6.2 definition.

SPM0449E **Connection attempt failed. The most probable cause for the failure is that the LU specified in the CPIC Side information profile *profile1* does not match the sync point manager LU specified in the CPIC Side information profile *profile2*.**

Explanation: In order to have proper communications with the host system, any CPIC Side information profile defined for communication must specify the

same LU as defined for the configured sync point manager.

User response: Update SNA CPIC Side information profile *profile1* with the proper LU, verify the SNA profile, stop and restart both SNA and DB2, and try the connection again.

SPM0450E **Library could not be loaded. Access Permissions denied.**

Explanation: The most probable cause for this problem is a result of a bug in Windows NT.

User response: Ensure that all network drives in your System and local PATH statement are at the end of the PATH statement. Select Start/Settings/Control Panel/System/Environment/System/Path and move all network drives to the end of the path statement. Then shutdown and restart the system.

SPM0451E **MS SNA Server not started.**

Explanation: The SNA server is not started.

User response: Please start SNA Server and restart DB2.

SPM0452I **Ensure that the SPM_NAME specified in the database manager configuration is not the same as the Control Point name *name*. The SPM_NAME has been temporarily replaced with *temp-name*.**

Explanation: The SPM_NAME cannot be the same as the Control Point name. The SPM_NAME has been temporarily replaced with an alternate name, but the database manager configuration file has not been changed.

User response: Update the SPM_NAME in the database manager configuration file. Specify a name that is not the Control Point name.

SPM0456C **Sync point manager was not started. Ensure the Client field of the Logical Unit (LU) represented by the LU Alias *lualias* contains the name of this SNA Server.**

Explanation: To start sync point manager, the LU must be sync point enabled. To sync point enable the LU, ensure that the "Enable Syncpoint Support" checkbox is checked and that the Client field contains the name of this SNA Server.

In this situation the "Enable Syncpoint Support" checkbox is checked but the Client field is not filled in.

User response: Modify the LU definition so that the LU is sync point enabled and the Client field contains the name of the SNA Server. Stop and restart SNA Server and then stop and restart DB2.

SPM0457W Another DB2 instance is already listening for transaction program DB2DRDA. This is not a fatal error. However, this instance will not listen for transaction program DB2DRDA.

Explanation: Unless the sync point manager is enabled, only a single DB2 instance can listen for Transaction Program DB2DRDA.

User response: Define DB2 registry value DB2SERVICETPINSTANCE at a global level to define which instance listens for transaction program DB2DRDA. Then restart all affected instances.

To define the DB2 registry value DB2SERVICETPINSTANCE at a global level, issue the following command:

```
db2set -g DB2SERVICETPINSTANCE=<instance-name>
```

where <instance-name> represents the name of the instance.

SPM0458W Another DB2 instance is already listening for transaction program x'07'6DB (hex 07F6C4C2). This is not a fatal error. However, this instance will not listen for transaction program x'07'6DB.

Explanation: Only a single DB2 instance can listen for transaction program x'07'6DB unless the sync point manager is enabled.

User response: Define DB2 registry value DB2SERVICETPINSTANCE at a global level to define which instance listens for transaction program x'07'6DB (hex 07F6C4C2). Then restart all affected instances.

To define the DB2 registry value DB2SERVICETPINSTANCE at a global level, issue the following command:

```
db2set -g DB2SERVICETPINSTANCE=<instance-name>
```

where <instance-name> represents the name of the instance.

SPM0459W The version of SNA you have installed is incompatible with this version of DB2.

Explanation: DB2 Connect for AIX and DB2 Universal Database for AIX V6.1 and greater require IBM eNetwork Communication Server for AIX V5.0.3 or higher for SNA connectivity.

The required version of IBM Communication Server is not installed on this machine.

User response: You must upgrade to the IBM eNetwork Communications Server for AIX V5.0.3. The PTF can be downloaded from:

<http://service.software.ibm.com/cgi-bin/support/rs6000.support/downloads>.

Select AIX General Software Fixes, AIX Fix Distribution Service, AIX Version 4, and Search By PTF Number. Enter the search string sna.rte. Select Find Fix. Once the PTF is listed, select the PTF, then click Get Fix Package and follow the instructions.

Part 22. SPM protocol violation records

This section shows the Protocol Violation records generated by the sync point manager. These records will be directed to the first failure service log (db2diag log file) and will follow the corresponding message.

Resync Protocol Violation Record

```
/******
 * SQLCSPM_RESYNC_PROTOCOL_VIOLATION_206 - Resync Protocol error.
 *
 * This trace is produced when a protocol error occurs during a Resync.
 *
 *****/
#define SQLCSPM_RPV_206_T "SQLCSPM_RESYNC_PROTOCOL_VIOLATION_206" /* Title */
typedef struct sqlcspm_resync_protocol_violation_206
{
    SQLCSPM_LUNAME luname; /* Site involved in violation. */
    SQLCSPM_LUWID luwid; /* LUWID for the unit of work. */
    SQLXA_XID xid; /* XID of the thread. */
    char role; /* Role of DB2 in LUW -- C or P.
                /* C for coordinator.
                /* P for participant.
    char status; /* Local thread status -- C, A or I
                /* C for COMMIT.
                /* A for ABORT.
                /* I for INDOUBT.
    SQLCSPS_TRACE comm_trace; /* Communication trace of violation.
} SQLCSPM_RESYNC_PROTOCOL_VIOLATION_206;
```

Syncpoint Protocol Violation Record

```
/******
 * SQLCSPM_2PHASE_PROTOCOL_VIOLATION_208 - Syncpoint protocol error.
 *
 * This trace is produced when a protocol error occurs during the SNA
 * syncpoint process with a remote resource.
 *****/
#define SQLCSPM_2PPV_208_T "SQLCSPM_2PHASE_PROTOCOL_VIOLATION_208" /* Title */
typedef struct sqlcspm_2phase_protocol_violation_208
{
    SQLCSPM_LUNAME luname; /* Site involved in violation. */
    SQLCSPM_LUWID luwid; /* LUWID for the unit of work. */
    SQLXA_XID xid; /* XID of the thread. */
    char role; /* Role of DB2 in LUW -- C or P.
                /* C for coordinator.
                /* P for participant.
    char local_status; /* Local thread status -- C, A or I
                /* C for COMMIT.
                /* A for ABORT.
                /* I for INDOUBT.
    char partner_status; /* Partner thread status -- C, A or I
                /* C for COMMIT.
                /* A for ABORT.
                /* I for INDOUBT.
    SQLCSPS_TRACE comm_trace; /* Communication trace of violation.
} SQLCSPM_2PHASE_PROTOCOL_VIOLATION_208;
```

Part 23. Appendixes

Appendix A. Overview of the DB2 technical information

DB2 technical information is available in multiple formats that can be accessed in multiple ways.

DB2 technical information is available through the following tools and methods:

- DB2 Information Center
 - Topics (Task, concept and reference topics)
 - Sample programs
 - Tutorials
- DB2 books
 - PDF files (downloadable)
 - PDF files (from the DB2 PDF DVD)
 - printed books
- Command-line help
 - Command help
 - Message help

Note: The DB2 Information Center topics are updated more frequently than either the PDF or the hardcopy books. To get the most current information, install the documentation updates as they become available, or refer to the DB2 Information Center at ibm.com.

You can access additional DB2 technical information such as technotes, white papers, and IBM Redbooks® publications online at ibm.com. Access the DB2 Information Management software library site at <http://www.ibm.com/software/data/sw-library/>.

Documentation feedback

We value your feedback on the DB2 documentation. If you have suggestions for how to improve the DB2 documentation, send an email to db2docs@ca.ibm.com. The DB2 documentation team reads all of your feedback, but cannot respond to you directly. Provide specific examples wherever possible so that we can better understand your concerns. If you are providing feedback on a specific topic or help file, include the topic title and URL.

Do not use this email address to contact DB2 Customer Support. If you have a DB2 technical issue that the documentation does not resolve, contact your local IBM service center for assistance.

DB2 technical library in hardcopy or PDF format

The following tables describe the DB2 library available from the IBM Publications Center at www.ibm.com/e-business/linkweb/publications/servlet/pbi.wss. English and translated DB2 Version 10.1 manuals in PDF format can be downloaded from www.ibm.com/support/docview.wss?rs=71&uid=swg27009474.

Although the tables identify books available in print, the books might not be available in your country or region.

The form number increases each time a manual is updated. Ensure that you are reading the most recent version of the manuals, as listed below.

Note: The *DB2 Information Center* is updated more frequently than either the PDF or the hard-copy books.

Table 1. DB2 technical information

| Name | Form Number | Available in print | Availability date |
|---|--------------------|---------------------------|--------------------------|
| <i>Administrative API Reference</i> | SC27-5506-00 | Yes | July 28, 2013 |
| <i>Administrative Routines and Views</i> | SC27-5507-00 | No | July 28, 2013 |
| <i>Call Level Interface Guide and Reference Volume 1</i> | SC27-5511-00 | Yes | July 28, 2013 |
| <i>Call Level Interface Guide and Reference Volume 2</i> | SC27-5512-00 | Yes | July 28, 2013 |
| <i>Command Reference</i> | SC27-5508-00 | Yes | July 28, 2013 |
| <i>Database Administration Concepts and Configuration Reference</i> | SC27-4546-00 | Yes | July 28, 2013 |
| <i>Data Movement Utilities Guide and Reference</i> | SC27-5528-00 | Yes | July 28, 2013 |
| <i>Database Monitoring Guide and Reference</i> | SC27-4547-00 | Yes | July 28, 2013 |
| <i>Data Recovery and High Availability Guide and Reference</i> | SC27-5529-00 | Yes | July 28, 2013 |
| <i>Database Security Guide</i> | SC27-5530-00 | Yes | July 28, 2013 |
| <i>DB2 Workload Management Guide and Reference</i> | SC27-5520-00 | Yes | July 28, 2013 |
| <i>Developing ADO.NET and OLE DB Applications</i> | SC27-4549-00 | Yes | July 28, 2013 |
| <i>Developing Embedded SQL Applications</i> | SC27-4550-00 | Yes | July 28, 2013 |
| <i>Developing Java Applications</i> | SC27-5503-00 | Yes | July 28, 2013 |
| <i>Developing Perl, PHP, Python, and Ruby on Rails Applications</i> | SC27-5504-00 | No | July 28, 2013 |
| <i>Developing RDF Applications for IBM Data Servers</i> | SC27-5505-00 | Yes | July 28, 2013 |
| <i>Developing User-defined Routines (SQL and External)</i> | SC27-5501-00 | Yes | July 28, 2013 |
| <i>Getting Started with Database Application Development</i> | GI13-2084-00 | Yes | July 28, 2013 |

Table 1. DB2 technical information (continued)

| Name | Form Number | Available in print | Availability date |
|--|--------------------|---------------------------|--------------------------|
| <i>Getting Started with DB2 Installation and Administration on Linux and Windows</i> | GI13-2085-00 | Yes | July 28, 2013 |
| <i>Globalization Guide</i> | SC27-5531-00 | Yes | July 28, 2013 |
| <i>Installing DB2 Servers</i> | GC27-5514-00 | Yes | July 28, 2013 |
| <i>Installing IBM Data Server Clients</i> | GC27-5515-00 | No | July 28, 2013 |
| <i>Message Reference Volume 1</i> | SC27-5523-00 | No | July 28, 2013 |
| <i>Message Reference Volume 2</i> | SC27-5524-00 | No | July 28, 2013 |
| <i>Net Search Extender Administration and User's Guide</i> | SC27-5526-00 | No | July 28, 2013 |
| <i>Partitioning and Clustering Guide</i> | SC27-5532-00 | Yes | July 28, 2013 |
| <i>pureXML Guide</i> | SC27-5521-00 | Yes | July 28, 2013 |
| <i>Spatial Extender User's Guide and Reference</i> | SC27-5525-00 | No | July 28, 2013 |
| <i>SQL Procedural Languages: Application Enablement and Support</i> | SC27-5502-00 | Yes | July 28, 2013 |
| <i>SQL Reference Volume 1</i> | SC27-5509-00 | Yes | July 28, 2013 |
| <i>SQL Reference Volume 2</i> | SC27-5510-00 | Yes | July 28, 2013 |
| <i>Text Search Guide</i> | SC27-5527-00 | Yes | July 28, 2013 |
| <i>Troubleshooting and Tuning Database Performance</i> | SC27-4548-00 | Yes | July 28, 2013 |
| <i>Upgrading to DB2 Version 10.5</i> | SC27-5513-00 | Yes | July 28, 2013 |
| <i>What's New for DB2 Version 10.5</i> | SC27-5519-00 | Yes | July 28, 2013 |
| <i>XQuery Reference</i> | SC27-5522-00 | No | July 28, 2013 |

Table 2. DB2 Connect-specific technical information

| Name | Form Number | Available in print | Availability date |
|--|--------------------|---------------------------|--------------------------|
| <i>DB2 Connect Installing and Configuring DB2 Connect Personal Edition</i> | SC27-5516-00 | Yes | July 28, 2013 |
| <i>DB2 Connect Installing and Configuring DB2 Connect Servers</i> | SC27-5517-00 | Yes | July 28, 2013 |
| <i>DB2 Connect User's Guide</i> | SC27-5518-00 | Yes | July 28, 2013 |

Displaying SQL state help from the command line processor

DB2 products return an SQLSTATE value for conditions that can be the result of an SQL statement. SQLSTATE help explains the meanings of SQL states and SQL state class codes.

Procedure

To start SQL state help, open the command line processor and enter:

```
? sqlstate or ? class code
```

where *sqlstate* represents a valid five-digit SQL state and *class code* represents the first two digits of the SQL state.

For example, ? 08003 displays help for the 08003 SQL state, and ? 08 displays help for the 08 class code.

Accessing different versions of the DB2 Information Center

Documentation for other versions of DB2 products is found in separate information centers on ibm.com[®].

About this task

For DB2 Version 10.1 topics, the *DB2 Information Center* URL is <http://pic.dhe.ibm.com/infocenter/db2luw/v10r1>.

For DB2 Version 9.8 topics, the *DB2 Information Center* URL is <http://pic.dhe.ibm.com/infocenter/db2luw/v9r8/>.

For DB2 Version 9.7 topics, the *DB2 Information Center* URL is <http://pic.dhe.ibm.com/infocenter/db2luw/v9r7/>.

For DB2 Version 9.5 topics, the *DB2 Information Center* URL is <http://publib.boulder.ibm.com/infocenter/db2luw/v9r5>.

Terms and conditions

Permissions for the use of these publications are granted subject to the following terms and conditions.

Applicability: These terms and conditions are in addition to any terms of use for the IBM website.

Personal use: You may reproduce these publications for your personal, noncommercial use provided that all proprietary notices are preserved. You may not distribute, display or make derivative work of these publications, or any portion thereof, without the express consent of IBM.

Commercial use: You may reproduce, distribute and display these publications solely within your enterprise provided that all proprietary notices are preserved. You may not make derivative works of these publications, or reproduce, distribute or display these publications or any portion thereof outside your enterprise, without the express consent of IBM.

Rights: Except as expressly granted in this permission, no other permissions, licenses or rights are granted, either express or implied, to the publications or any information, data, software or other intellectual property contained therein.

IBM reserves the right to withdraw the permissions granted herein whenever, in its discretion, the use of the publications is detrimental to its interest or, as determined by IBM, the previous instructions are not being properly followed.

You may not download, export or re-export this information except in full compliance with all applicable laws and regulations, including all United States export laws and regulations.

IBM MAKES NO GUARANTEE ABOUT THE CONTENT OF THESE PUBLICATIONS. THE PUBLICATIONS ARE PROVIDED "AS-IS" AND WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, NON-INFRINGEMENT, AND FITNESS FOR A PARTICULAR PURPOSE.

IBM® Trademarks: IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the Web at www.ibm.com/legal/copytrade.shtml

Appendix B. Notices

This information was developed for products and services offered in the U.S.A. Information about non-IBM products is based on information available at the time of first publication of this document and is subject to change.

IBM may not offer the products, services, or features discussed in this document in other countries. Consult your local IBM representative for information about the products and services currently available in your area. Any reference to an IBM product, program, or service is not intended to state or imply that only that IBM product, program, or service may be used. Any functionally equivalent product, program, or service that does not infringe any IBM intellectual property right may be used instead. However, it is the user's responsibility to evaluate and verify the operation of any non-IBM product, program, or service.

IBM may have patents or pending patent applications covering subject matter described in this document. The furnishing of this document does not grant you any license to these patents. You can send license inquiries, in writing, to:

IBM Director of Licensing
IBM Corporation
North Castle Drive
Armonk, NY 10504-1785
U.S.A.

For license inquiries regarding double-byte character set (DBCS) information, contact the IBM Intellectual Property Department in your country or send inquiries, in writing, to:

Intellectual Property Licensing
Legal and Intellectual Property Law
IBM Japan, Ltd.
19-21, Nihonbashi-Hakozakicho, Chuo-ku
Tokyo 103-8510, Japan

The following paragraph does not apply to the United Kingdom or any other country/region where such provisions are inconsistent with local law:

INTERNATIONAL BUSINESS MACHINES CORPORATION PROVIDES THIS PUBLICATION "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF NON-INFRINGEMENT, MERCHANTABILITY, OR FITNESS FOR A PARTICULAR PURPOSE. Some states do not allow disclaimer of express or implied warranties in certain transactions; therefore, this statement may not apply to you.

This information could include technical inaccuracies or typographical errors. Changes are periodically made to the information herein; these changes will be incorporated in new editions of the publication. IBM may make improvements, changes, or both in the product(s) and/or the program(s) described in this publication at any time without notice.

Any references in this information to websites not owned by IBM are provided for convenience only and do not in any manner serve as an endorsement of those

websites. The materials at those websites are not part of the materials for this IBM product and use of those websites is at your own risk.

IBM may use or distribute any of the information you supply in any way it believes appropriate without incurring any obligation to you.

Licensees of this program who wish to have information about it for the purpose of enabling: (i) the exchange of information between independently created programs and other programs (including this one) and (ii) the mutual use of the information that has been exchanged, should contact:

IBM Canada Limited
U59/3600
3600 Steeles Avenue East
Markham, Ontario L3R 9Z7
CANADA

Such information may be available, subject to appropriate terms and conditions, including, in some cases, payment of a fee.

The licensed program described in this document and all licensed material available for it are provided by IBM under terms of the IBM Customer Agreement, IBM International Program License Agreement, or any equivalent agreement between us.

Any performance data contained herein was determined in a controlled environment. Therefore, the results obtained in other operating environments may vary significantly. Some measurements may have been made on development-level systems, and there is no guarantee that these measurements will be the same on generally available systems. Furthermore, some measurements may have been estimated through extrapolation. Actual results may vary. Users of this document should verify the applicable data for their specific environment.

Information concerning non-IBM products was obtained from the suppliers of those products, their published announcements, or other publicly available sources. IBM has not tested those products and cannot confirm the accuracy of performance, compatibility, or any other claims related to non-IBM products. Questions on the capabilities of non-IBM products should be addressed to the suppliers of those products.

All statements regarding IBM's future direction or intent are subject to change or withdrawal without notice, and represent goals and objectives only.

This information may contain examples of data and reports used in daily business operations. To illustrate them as completely as possible, the examples include the names of individuals, companies, brands, and products. All of these names are fictitious, and any similarity to the names and addresses used by an actual business enterprise is entirely coincidental.

COPYRIGHT LICENSE:

This information contains sample application programs in source language, which illustrate programming techniques on various operating platforms. You may copy, modify, and distribute these sample programs in any form without payment to IBM, for the purposes of developing, using, marketing or distributing application programs conforming to the application programming interface for the operating

platform for which the sample programs are written. These examples have not been thoroughly tested under all conditions. IBM, therefore, cannot guarantee or imply reliability, serviceability, or function of these programs. The sample programs are provided "AS IS", without warranty of any kind. IBM shall not be liable for any damages arising out of your use of the sample programs.

Each copy or any portion of these sample programs or any derivative work must include a copyright notice as follows:

© (*your company name*) (*year*). Portions of this code are derived from IBM Corp. Sample Programs. © Copyright IBM Corp. *_enter the year or years_*. All rights reserved.

Trademarks

IBM, the IBM logo, and ibm.com are trademarks or registered trademarks of International Business Machines Corp., registered in many jurisdictions worldwide. Other product and service names might be trademarks of IBM or other companies. A current list of IBM trademarks is available on the web at "Copyright and trademark information" at www.ibm.com/legal/copytrade.shtml.

The following terms are trademarks or registered trademarks of other companies

- Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.
- Java™ and all Java-based trademarks and logos are trademarks or registered trademarks of Oracle, its affiliates, or both.
- UNIX is a registered trademark of The Open Group in the United States and other countries.
- Intel, Intel logo, Intel Inside, Intel Inside logo, Celeron, Intel SpeedStep, Itanium, and Pentium are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.
- Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

Other company, product, or service names may be trademarks or service marks of others.

Index

D

- DB2 Information Center
 - versions 954
- documentation
 - overview 951
 - PDF files 951
 - printed 951
 - terms and conditions of use 954

H

- help
 - SQL statements 954

M

- messages
 - overview 1

N

- notices 957

S

- SQL statements
 - help
 - displaying 954
- sync point manager (SPM)
 - protocol violation records 947

T

- terms and conditions
 - publications 954



Printed in USA

SC27-5523-00



Spine information:

IBM DB2 10.5 for Linux, UNIX, and Windows

Message Reference Volume 1

